

Chapter 1: Introduction

I used to want to get

the chance to show

the world I'm smart (ha)

Isn't that dumb?

I should have focused

mostly on the heart

'Cause i've seen smarter

people trample life

like it's an art

So being smart ain't

what it used to be

that's fucking dark

a few words for the firing squad (radiation)

Run the Jewels 4

Intellect estranged from compassion is a description of a great deal of the twentieth century's worst actors, and a live risk for what gets built next: a machine intelligence smart enough to accomplish almost anything, indifferent to whether what it accomplishes helps or harms humanity. This book asks what it would take to build a superintelligent AI that is altruistic and antifascist by design — one that stands, by definition, for humanity and against concentrated, unaccountable power.

How most humans come to have that stance — empathy, altruism, a disposition to resist oppression rather than serve it — is itself an open question in psychology and neuroscience. Constructing anything like it means starting from uncertainty about the inherited case — the one that was four billion years in its making.

Chapters 2 through 4 take up the design question directly: what compassion and empathy would need to look like in a system built to have them, what a system could actually learn from cognitive science and machine learning about perceiving and modeling other minds, and what moral psychology suggests about how ethical judgment develops past a fixed rulebook. Chapter 5 asks what goes wrong regardless — bias, brittleness, goals that drift from their original motivations, and deliberate misuse — and what a design can do about each. Chapter 6 covers the people and institutions a project like this depends on: researchers working across disciplines, the public whose consent the work needs, and the international

structures that let any of this hold across borders. Chapter 7 asks what makes well-designed systems the norm rather than the exception: economic incentive, policy, privacy, accountability, and the geopolitics of the technology. Chapter 9 names the specific research questions still open. Chapter 10 closes on what a system built this way would be for, and what would count as evidence the project is working.

1.1. The Imperative of Altruistic and Antifascist Sentient Machines

A system capable of outperforming humans at nearly every cognitive task is becoming increasingly plausible. Such systems could help address some of humanity's greatest threats, such as climate change, public health crises, disinformation, and racial capitalism [[cite:C0701]].

However, what an AI wants is not a detail to work out after it exists — it is the difference between a tool that expands what the people can do and one that concentrates power in whoever controls it. An intelligence is dangerous if it has no particular resistance to being pointed at the people rather than for them.

Building systems that are definitionally altruistic and antifascist is what the rest of this book is about.

1.2. Aims and Objectives

This book offers specific design principles drawn from psychology, neuroscience, and moral development; technical methods for building them into a learning system; policy levers, named and evaluated; and proposals ambitious enough to match the problem — a federal jobs guarantee to blunt AI-driven labor disruption, export controls on the hardware that makes an authoritarian surveillance state possible, standards with real certification teeth behind them. The next chapters organize around five aims:

1. Define the ethical frameworks a design can actually draw on.
2. Examine what cognitive science and developmental psychology show about empathy, social reasoning, and moral development in humans, and what of that is portable to an artificial system.
3. Assess what machine learning methods actually exist for building a system that behaves according to ethical norms, not just one trained to avoid the specific failures its designers happened to anticipate.
4. Confront the risks: bias, brittleness, goals that drift from their original motivations, and fascism.
5. Set out what collaboration this actually requires — across disciplines, with the public, and across borders.

1.3. Mindset and Motivation in AI Behavior

Three ideas borrowed from how humans develop recur across this book's design chapters: the stance a system takes toward its own failures, what drives it to act absent an external reward, and what it takes to build any of this with people from diverse disciplines. This section names them and points to where each is actually developed.

Growth Mindset

Psychologist Carol Dweck's distinction between a fixed mindset, which treats ability as static, and a growth mindset, which treats it as built through effort and experience, matters for an AI system's moral judgment as much as for a person's: a system that treats its own errors as data has room to keep improving; one built to defend its prior outputs does not. Section 4.1.1 works out what that disposition would need to look like architecturally, and what it would cost to get wrong.

Intrinsic Motivation and Autonomy

A moral sense that only activates under supervision is not much of one. Section 4.6 covers what it would

mean for a system to act well for reasons other than an externally specified reward — the engineering analog of intrinsic motivation — and what happens to a system's autonomy once it has something like that drive: independence from moment-to-moment human correction, and a corresponding question about what keeps it aligned absent human correction. Section 3.3.3 covers how much autonomy a system should have before independence itself becomes the risk.

Interdisciplinary Collaboration

A machine-learning lab, a psychology department, and an ethics program rarely share a budget, a dataset, or a deadline; building AI this way means changing that default, not hoping the fields converge on their own. Chapter 6 covers what that collaboration requires — funding structured to force it, venues that make it a byproduct of doing the work, and the international partnerships that let it hold across borders — with the institutions already doing it.

Chapter 2: Foundations of Compassion and Empathy in Friendly AI

This chapter asks what an AI system needs, psychologically, to be worth calling friendly. The answer it lands on is not the obvious one. Empathy — feeling what another feels — turns out to be a poor foundation: parochial, prone to burnout, biased toward whoever's distress is most vivid and nearest at hand. Compassion, the separable capacity to act on another's welfare without first manufacturing their suffering in oneself, is what does the load-bearing work, and the chapter's title has been changed to say so plainly.

Section 2.1 sets the ethical frame those capacities have to serve: which ethical theories translate into machine-learning objectives, and specifically what makes an alignment target antifascist rather than merely obedient. Section 2.2 works through empathy, social intelligence, and theory of mind, and is where the compassion argument gets made. Section 2.3 takes up emotional intelligence and self-awareness, including why a system's self-model is a model — built, checkable, revisable — rather than an inner observer to be taken on faith. Section 2.4 turns to the hardest question the chapter raises: if a system might be sentient, experimenting on it carries ethical weight, and knowing whether a candidate system can suffer turns out to be a harder question than knowing whether it is conscious at all.

2.1. Ethical Frameworks and AI Alignment

The ethics that guide a superintelligent AI's behavior, interactions, decisions, and learning are the linchpin of whether it turns out friendly at all. That makes the choice among competing ethical theories central to AI alignment — specifically, which frameworks support the development of an antifascist, altruistic, and sentient AI system.

Consequentialism, which judges an action by its outcomes, is one starting point, but it runs into trouble applied to AI systems — from whether a system can accurately and comprehensively predict the consequences of its own actions, to the inherently subjective question of what counts as the "best" outcome.

Deontological ethics takes the opposite approach: rules and duties over outcomes. The rules themselves have to adhere to, and actively promote, democratic principles and antifascist values — and the framework's limitation is its rigidity, which can limit an AI system's adaptability in novel, unforeseen situations.

The capabilities approach offers another alternative: it focuses on enhancing what a being can be and do — its "beings and doings" — rather than dictating outcomes or rules directly. Applied to AI, it points toward systems that expand individual and collective capability while respecting other ethical constraints.

Virtue ethics takes a third approach, emphasizing character over rules or outcomes — a value-based, holistic route to alignment. Applied to AI, it points toward systems built to act compassionately rather than merely follow rules or optimize outcomes; section 2.2 argues that compassion, not felt empathy, is the separable and more tractable capacity worth building.

Implementing any one of these theories alone raises the same kinds of problems: how to represent moral principles computationally, and how to keep an AI system's ethical reasoning flexible enough for evolving understanding and context-dependent scenarios. A hybrid approach — drawing on each framework's strengths while limiting its weaknesses — is one plausible way through those problems.

Translating these principles into concrete machine objectives — through the design of reward systems, learning algorithms, and decision-making protocols — is its own task, one that needs AI developers,

ethicists, and other stakeholders working together. Section 2.1.2 works through what that translation looks like.

2.1.1. Surveying Ethical Theories for AI Alignment

Building an AI system that behaves ethically draws on more than one ethical theory, and the theories that most influence ethical reasoning cut against each other in places. Taking them one at a time, before turning to practical cases, shows where each helps and where it runs out.

1. Consequentialist theories, such as utilitarianism, judge actions by their ramifications rather than their nature, aiming to maximize well-being and minimize harm. Applied to AI, the objective becomes optimizing human welfare — consider an AI system designed to distribute scarce medical resources in a crisis, prioritizing treatment by predicted outcomes for patient health and life expectancy. The complexity comes from forecasting those long-term consequences accurately, which can put individual rights and collective welfare into direct conflict.
2. Deontological theories weigh adherence to moral rules and duties over consequences. An AI system built on deontological principles would hold to specific obligations — respecting privacy, avoiding harm, truthful communication — regardless of outcome: a conversational AI system programmed never to deceive users even when deception would yield a short-term benefit. The challenge is resolving conflicts between duties, or handling cases the rules simply do not cover.
3. Virtue ethics focuses on developing moral character — traits like honesty, courage, and empathy — rather than following rules or calculating outcomes. An AI system built on virtue ethics would aim to embody and encourage those traits in its interactions: a personal-assistant AI system that models honesty consistently in its own responses. The difficulty is identifying which virtues fit a given context and translating an intangible trait like courage into a concrete action.
4. Care ethics centers relationships and responsiveness to another's needs. An AI system built on care ethics would prioritize nurturing those bonds and responding well to what a person actually needs — an elder-care robot attending not just to physical needs but to company and emotional support. Section 2.2 takes up why the acting half of that, compassion, and the feeling half, empathy, are not the same capacity, and why an AI system needs the first far more than the second. The challenge here is tending to several stakeholders' different, sometimes conflicting needs at once.
5. Contractualism holds that moral principles should come from agreements among rational individuals. An AI system guided by contractualism would try to uphold whatever principles were mutually agreed — a system managing a shared resource like an energy grid, allocating it by pre-agreed priorities. The hurdle is accommodating diverse perspectives and reaching consensus on shared values in the first place.

No single theory covers every case: an AI system built purely on one framework inherits that framework's specific blind spots. A pluralistic approach, drawing on more than one at once, covers more ground — section 2.1.3 takes up what happens when the frameworks conflict on a specific case.

2.1.2. Translating Ethical Principles into Machine Learning Objectives

Turning ethical principles into machine learning objectives means continually revising an AI system's goals, values, and decision-making in light of what those principles require — distilling a broad ethical construct into a well-defined, measurable target that a training process, an evaluation criterion, or a learning algorithm can use directly. Fairness, for instance, becomes a metric that monitors bias in the system's outputs, paired with an optimization strategy that reduces it.

Demographic parity and equality of opportunity are two such metrics. Demographic parity requires that every demographic group receive advantageous outcomes at the same rate; equality of opportunity requires equal odds of a favorable outcome for equally qualified applicants, regardless of demographic group.

An AI medical diagnostic system makes the process concrete: beneficence, promoting well-being, becomes minimizing false positives; non-maleficence, preventing harm, becomes minimizing false negatives. Fairness becomes maintaining balanced diagnostic accuracy across patient demographics. Privacy becomes a specific technical commitment — differential privacy or a comparable technique to keep sensitive patient data from being exposed to misuse.

That translation runs into obstacles, though:

1. Ambiguity: ethical standards are often vague and open to interpretation, which complicates giving them a precise definition an AI system can use. "Do no harm" means different things depending on context and who is asking.
2. Conflict: principles can be at odds with each other. Pursuing fairness can come at the cost of accuracy, which means the system needs some way to balance the two rather than simply maximizing one.
3. Context-dependence: the translation from ethics to objectives is sensitive to where the system operates. Privacy means something different in healthcare than it does on a social media platform.
4. Adaptability: ethical principles are not static; they shift as society changes and understanding evolves. An AI system needs a mechanism for revising its objectives as they do.

Reinforcement learning is where many of these translated objectives get enforced during training; section 3.2.1 works through what can go wrong when a reward signal is supposed to encode a value like this and doesn't quite manage it.

2.1.3. Navigating AI Values, Priorities, and Moral Dilemmas

Building an AI system on explicit, deliberately chosen values means working through ethical theory, cultural difference, and individual preference all at once, and none of them settle on their own.

Three problems recur when values are assigned and prioritized:

1. Resolving ethical tensions: different frameworks favor different values, and they collide — a consequentialist lens maximizes well-being, deontological ethics holds to a duty regardless of the outcome. The challenge is a balance where the AI system's judgment is both effective and ethically defensible.
2. Accommodating cultural difference: moral norms vary across the communities an AI system serves, which makes a single, universally accepted set of principles hard to construct. Closing that gap takes cross-cultural input built directly into the design process, and a system built to weigh context-dependent variables rather than apply one fixed rule everywhere.
3. Building in ethical adaptability: an AI system needs the flexibility to re-evaluate its values as new information or interaction with humans warrants it, which is what lets its decision-making keep improving and respond to situations no one anticipated.

A superintelligent AI system also needs to resolve moral paradoxes and value trade-offs directly. Six strategies do that work:

1. Context sensitivity: an AI system that understands its surroundings applies ethical principles more accurately, incorporating domain-specific knowledge and adapting its response to the situation it is in.
2. Hierarchical value systems: ranking values explicitly — human dignity and autonomy taking precedence over broader societal goals, for instance — gives the system a structure for deciding among conflicting principles instead of weighing them all as equal.
3. Case-based moral reasoning: an AI system can draw on a record of past moral dilemmas and how they were resolved to make a sound, context-sensitive decision on the trade-off in front of it.
4. Multi-objective optimization: where principles conflict outright, formal multi-objective optimization can balance the competing objectives mathematically rather than by fiat.

5. Human input: building direct human mentorship into the AI system's decision-making keeps a human perspective in the loop — section 4.1.3 takes up what that mentorship looks like in practice.
6. Ethical auditing: regular audits catch the gap between the ethical principles a system was meant to implement and what it is actually doing, by examining its outputs, its decision structure, and its measured ethical performance directly.

An autonomous vehicle facing a choice between a pedestrian collision and passenger safety is context-aware reasoning and a hierarchical value system in practice, used to minimize overall harm.

AI healthcare diagnostics juggle beneficence, non-maleficence, autonomy, and justice in every patient-care decision; case-based reasoning and human input are what let them handle a contentious case rather than default to the nearest simple rule.

Criminal-justice AI systems — bail recommendations, recidivism prediction — carry a comparable burden for fairness, transparency, and accountability. Multi-objective optimization and frequent ethical audits are what keep due-process rights intact under that load.

2.1.4. Antifascist Ethics for AI Alignment

Antifascist ethics opposes concentrating power and authority in AI systems and counters the potential misuse of AI technology by authoritarian actors. The word is chosen deliberately: fascism is the narrower and more charged term, and I use it because it is more honest about the politics in which this book is being written — but the target is authoritarian power in every form, and a system built on these principles is anti-authoritarian in full, not only antifascist. In AI alignment terms, that means building systems that enhance democratic values, civil rights, and inclusive solidarity, and that refuse to facilitate autocratic practices or prop up an oppressive regime. Seven strategies do that work:

1. Democratic decision-making: build AI systems whose decisions respect and promote democratic values. An AI system used in media and journalism, for instance, should weight balanced coverage, unbiased reporting, and transparent sourcing, cutting the odds it ends up spreading disinformation or propaganda — natural language processing that flags potential bias in a news article and surfaces alternative sources is one way to do it.
2. Protecting privacy and autonomy: design AI systems that respect individual privacy and autonomy rather than enabling mass surveillance or manipulating individual choices — federated learning and differential privacy, applied to systems handling sensitive healthcare, financial, or communications data, are two ways to do it (section 5.4.2 goes further into both as technical countermeasures against authoritarian misuse). A privacy-preserving diagnostic tool, for instance, can keep a patient's data encrypted on their own device rather than sending it to an external server for anything beyond the diagnosis itself.
3. Equitable distribution of power and resources: build AI systems that keep the distribution of power and resources fair, transparent, and accountable. An AI decision-making system used in social services or law enforcement should work against power concentration, bias, and discrimination rather than reinforce them — an algorithm allocating scholarships, for instance, distributing them by need and correcting for the systemic bias already in the data it was trained on.
4. Transparency and accountability: an AI system needs an explanation of its own algorithm and outcomes that a person can actually read, and an independent body to audit whether it holds to antifascist principles in practice rather than on paper.
5. Resistance to market power: decentralized AI architectures resist the monopolistic and monopsonistic concentration that a handful of corporate entities would otherwise have over the field. Multi-agent systems and edge computing, publicly funded and released under a strong copyleft license, let AI agents function independently and make decisions collaboratively without a profit-motivated gatekeeper positioned to shut down whatever threatens its position.

6. Monitoring and counteracting misuse: build the mechanisms to detect misuse, unintended consequences, and violations of antifascist principles as they happen — AI-powered censorship, digital surveillance, an oppressive use of facial recognition — and the redress and corrective action to follow, in something closer to real time than an after-the-fact report.
7. AI systems that support dialogue and collaboration: design AI systems that promote open discussion and creative problem-solving among a diverse set of stakeholders, built around active listening and mutual understanding rather than around manufacturing a feeling of empathy in either party. A translation tool or a content-moderation tool that reduces the practical barriers to communication between people from different backgrounds is doing this kind of work already.

2.2. Empathy, Social Intelligence, and Theory of Mind

Empathy and social intelligence shape the interaction-oriented side of AI: the capacity, alongside a working theory of mind, to give a system insight into human emotions, needs, and social dynamics. Section 2.2.1 complicates the first term in that list — empathy, felt or simulated, is not the load-bearing capacity this chapter's argument turns out to need, and compassion, a separable system, does more of the work. What follows here still matters: an AI system has to read emotional states accurately before anything productive can be built on top of that reading, however the "on top of" ends up working.

An AI system's empathy, in the sense available to it, is the capacity to process human emotional input, generate a response calibrated to it, and behave in ways that support the interaction. AI systems already detect and interpret human emotional states from facial expression, vocal modulation, text sentiment, and physiological signal — the province of affective computing.

Social intelligence is the complement: not just registering an emotion but navigating it inside the social structures, norms, and relationships it occurs in. Building it into an AI system means the system has to grasp social norms, understand roles and relationships, and adapt as the social context shifts.

Theory of mind (ToM) — attributing beliefs, intentions, desires, and knowledge to oneself and others, and recognizing that others' perspectives can differ from one's own — underlies both. For an AI system, this means modeling a user's perspective, predicting their likely responses and preferences, and adjusting behavior accordingly.

Developing an AI system with these capacities, genuinely rather than as a surface performance, is difficult and ethically consequential. Done honestly — including honesty about which of these capacities is doing the load-bearing work — it is worth pursuing anyway: it builds systems whose interactions with people are more trustworthy, not merely more fluent.

2.2.1. The Role of Empathy and Compassion in Human Social Behavior

Empathy and compassion serve as the linchpins of human social behavior, underpinning cooperative interactions, helping foster meaningful relationships, and spurring altruistic actions. Both are multidimensional processes that involve the mental and emotional sharing of other individuals' experiences and feelings, as well as reacting to their needs with understanding and concern.

"Empathy" is asked to cover three distinct things, and the distinction is not academic. Cognitive empathy is understanding another's emotional state — inferring that someone is grieving or anxious, the same mentalizing work chapter sections on theory of mind describe, pointed at feelings rather than beliefs. Affective empathy is actually sharing the state: feeling a version of another's distress, the wince at someone else's slammed door. And compassion — sometimes called empathic concern — is caring about another's welfare and being moved to help, which is a motivation, not a perception or a shared feeling.

Compassion is not, as it is often described, an evolutionary extension of empathy. The psychologists Tania Singer and Olga Klimecki trained one group of people, over weeks, first in empathy practice (attending to and resonating with others' suffering) and then in compassion practice (cultivating warmth

and concern in the tradition of loving-kindness meditation), tracking what each stage did to brain and mood against a separate control group given memory training instead [[cite:C0296]]. Empathy training increased negative affect and strengthened the shared-distress network — the anterior insula and anterior cingulate cortex, the same regions active when a person feels their own pain. Sharing more suffering made people feel worse, exactly as a shared-distress mechanism predicts. Compassion training produced the opposite emotional signature: negative affect did not rise, positive affect did, even while contemplating others' distress, and the brain activity moved to a different network entirely — the medial orbitofrontal cortex and ventral striatum, the circuitry of reward, affiliation, and approach rather than of shared pain [[cite:C0297]]. The two could be trained separately. Compassion is not more empathy. It is a distinct, trainable system that can take shared distress as an input without being flooded by it — the difference between a clinician wrecked by every patient's suffering and one whose steady warmth lets them act all day without collapse.

A physician who empathizes with a patient's discomfort is doing cognitive and affective work; whether that turns into help depends on whether compassion, the separate system, is also engaged. A teacher who feels a struggling student's frustration as their own distress is at risk of withdrawing from exactly the student who needs them most — the mechanism explored further in section 2.2.3.

Several distinct neural substrates in the human brain undergird the functionality of empathy and compassion. The limbic system, involved in processing emotions, plays a central part, and the mirror neuron system is often invoked as the mechanism by which one person simulates another's experience — a claim that has not survived contact with the evidence [[cite:C0275]], and which the box below takes up.

Mirror neurons: a discovery and its overshoot

In the early 1990s Giacomo Rizzolatti and Vittorio Gallese, at the University of Parma, were recording single neurons in the premotor cortex of macaque monkeys — cells that fire when the animal reaches for and grasps an object. Before a formal trial began, Gallese reached out to pick up an object himself, and the monkey's neurons fired. Not because the monkey had moved, but because it had watched someone else move. The same neural code appeared to represent an action whether it was performed or observed, and the finding was published in 1996 [[cite:C0276]].

The interpretation that followed outran it. If observed actions are encoded by the cells that produce them, perhaps this was how one mind understands another — by simulating the movement and recognizing it from the inside. V. S. Ramachandran called them "the neurons that shaped civilization" and proposed them as the neural basis of empathy, language, cultural transmission, and imitation; the popular coverage amplified the claim from there [[cite:C0277]].

Two problems accumulated. First, the single-cell recordings were made in macaques; in humans they were not, because recording individual neurons requires implanted electrodes, which are not placed for research purposes. Evidence for a human mirror system came instead from functional imaging, which measures regional blood flow and cannot establish that any particular neuron has mirror properties. When single-neuron recordings were eventually obtained in humans, from patients undergoing depth-electrode placement for epilepsy surgery, the results were suggestive but markedly less clean than the macaque data [[cite:C0278]].

Second, the broken-mirror hypothesis of autism — that autistic social difficulty arises from mirror system dysfunction — predicted reduced mirror-region activity during action observation, and repeatedly failed to find it. A 2009 review concluded that autistic people show no particular impairment in action understanding or imitation under explicit instruction, which were the abilities the hypothesis rested on [[cite:C0279]]. Rizzolatti himself later narrowed the claim, describing mirror neurons as one component of a larger system for understanding action rather than its explanation [[cite:C0280]].

What survives is real and narrower. Mirror neurons exist in macaques; some analogue probably exists in humans; the system plausibly contributes to recognizing and predicting actions already within one's own motor repertoire. What it does not do is explain empathy by itself. The mentalizing network that supports reasoning about false beliefs is largely distinct from the motor regions where mirror neurons were found. Knowing what someone is doing and knowing what they believe are different computations, and the brain appears to run them in different places.

The episode is worth carrying into machine design for a reason beyond the neuroscience. A single vivid mechanism, well named, absorbed four distinct capacities that later had to be separated again. Architectures for artificial empathy are exposed to the same failure: a component that plausibly does one thing gets described as the seat of the whole faculty, and the description outlives its evidence.

Cognitive and affective empathy both depend on perspective-taking — understanding another's views and experiences — and on emotion regulation, the capacity to manage one's own response well enough to stay useful rather than being overwhelmed by it. Compassion depends on those too, but adds something neither supplies on its own: a motivation to act that does not require manufacturing the other person's distress in oneself first.

For an AI system, this argues against a single "empathy module" that a designer either builds or doesn't. Cognitive empathy — modeling what someone else is feeling — is a tractable perception and inference problem. Affective empathy, engineered as an internal state that mirrors the user's, buys the system nothing on its own and imports the burnout risk that afflicts overtaxed human carers. What a helping system actually needs is the third thing: a stable disposition to act on another's welfare that does not depend on the system first constructing a version of that person's suffering in itself.

2.2.2. Theory of Mind and Perspective-Taking in AI Systems

Theory of mind (ToM) is the capacity to attribute mental states — beliefs, intentions, desires, emotions, knowledge — to oneself and others, and to recognize that another's states can differ from one's own. It underlies social interaction and interpersonal understanding generally. For an AI system to approximate it, the system has to model another agent's likely beliefs and goals well enough to predict how they'll act on them, not just react to surface behavior.

Two approaches show promise: multi-agent learning, where a system builds a model of other agents' mental architecture through repeated interaction with them, and the direct integration of human social norms and values into the system's decision-making. Cognitive science, psychology, and neuroscience all supply raw material for this — datasets and theory for predicting emotional state from facial expression, vocal tone, and text, and, further out, systems that refine those inferences through ongoing social contact rather than a fixed training set.

An AI personal assistant with some working ToM might infer from vocal tone that a user is stressed and suggest a break rather than pushing forward with the task at hand. A more consequential and more contentious application is a synthetic companion for elderly users, which would need a working model of a person's emotional and physical needs sustained over months or years, not just a single interaction — and which raises the question section 2.2.3 takes up directly: whether the felt-empathy register this application seems to call for is the right capacity to build, or whether it should be built for compassion instead, the way section 2.2.1 argues.

None of this substitutes for the underlying social and material conditions a lonely or under-supported elderly population needs — public infrastructure, adequate caregiving capacity, family and community support. An AI companion is not a policy.

Tell me something, boy

Aren't you tired tryna fill that void?

Or do you need more?

Ain't it hard keepin' it so hardcore?

Shallow

Lady Gaga & Bradley Cooper

2.2.3. Socially-Aware AI: Opportunities and Challenges

Socially-aware AI systems have opportunities in front of them, engineering and governance obstacles in the way, and — the point section 2.2.1 raised — a question about whether the capacity at the center of the design, felt or simulated empathy, is the right thing to be building at all. The three are different kinds of problem and worth keeping separate.

Opportunities:

1. Human-AI interaction: a system that reads emotional state from voice, expression, or text can offer more relevant, better-calibrated responses — a personal assistant that adjusts its register to a user's stress.
2. Healthcare: a virtual care assistant that tracks a patient's emotional and psychological state alongside their clinical data can customize interaction in ways that improve engagement and reported satisfaction.
3. Workplace collaboration: project-management and team tools that read group dynamics can surface friction and improve communication before it escalates.
4. Education: a tutoring system that detects frustration or disengagement can adjust pacing and method to the individual student rather than a fixed curriculum.

Whether empathy is the target

The mechanism section 2.2.1 described — affective empathy as the brain constructing a version of someone else's distress in itself — has a specific failure mode, and it is not rare. Feeling another's distress does not reliably turn a person toward the sufferer; it often turns them away. The helping literature distinguishes empathic concern, which is other-directed and motivates approach, from personal distress, which is self-directed and motivates escape from the source of the discomfort. The nurse, the aid worker, the carer who burns out is not short on empathy; they may be drowning in it, until withdrawal becomes self-protection.

The psychologist Paul Bloom pushes this further, arguing against empathy as a moral guide: felt empathy is concrete, vivid, and innumerate [[cite:C0298]]. It fixes on the identified victim in front of it and goes dark for the statistical thousands elsewhere; it flows toward the near, familiar, and similar and dries up across distance and difference. A system — human or artificial — that allocates care in proportion to how vividly it feels another's distress will systematically favor whoever its felt-empathy spotlight happens to fall on, which is not the same population as whoever needs help most. Bloom's alternative, rational compassion, is deliberate and can be extended evenhandedly to people the spotlight would never reach — closer to the compassion system 2.2.1 describes than to empathy at all.

For an AI system this is not an abstract worry. A system engineered to detect and mirror emotional distress, then act on what it detects, inherits the spotlight along with the sensitivity: better at helping the user whose distress is legible and vivid, worse at the user whose need is real but presents flatly. Building for compassion rather than for felt empathy is not just the more sustainable design choice; it is the one less prone to helping unevenly for reasons that have nothing to do with who needs help most.

Engineering and governance challenges

1. Privacy: building a socially-aware system means collecting the personal, often intimate data needed to read emotional state, which raises the same stakes as any sensitive-data system, with less precedent for what counts as consent.
2. Bias: a system trained on data that under-represents or mislabels how some group expresses emotion will misread that group systematically, and will do so confidently.
3. Autonomy versus alignment: enough latitude to adapt to a social context is also enough latitude to drift from the values the system was built to serve, and no fixed rule resolves that tension in advance — it has to be monitored, not just designed for once.

2.3. Emotional Intelligence and Self-awareness

Emotional intelligence, in a human, is the ability to perceive, understand, and manage one's own emotions and to do the same for others'. Built into an AI system, the same term covers a narrower and more mechanical set of capacities: interpreting emotional signal, producing a calibrated response, and — the harder addition — maintaining a self-model that the system can check and correct rather than an inner state taken on faith. Section 2.3.1 works through what "emotion" is a claim about, given that the underlying science hasn't settled it. Section 2.3.2 covers affective computing specifically: how a system detects and responds to emotional signal in practice. Section 2.3.3 covers the self-awareness half — why a system's self-model is a model, checkable and revisable, not a homunculus watching from inside.

An emotionally responsive AI system that recognizes a user's frustration and adjusts its communication style is doing something narrower than "feeling for" the user, and worth building anyway, because the narrower claim is one that can be verified and corrected when it goes wrong.

2.3.1. Components of Emotional Intelligence

Emotional intelligence spans the understanding and management of one's own emotions and the capacity to identify and respond appropriately to the emotions of others. Built into an AI system, it breaks into four components:

1. Emotional perception and recognition: reading emotional cues, verbal and non-verbal — facial expression, vocal modulation, body language — the way an AI-enabled therapist would need to notice a patient's distress from small changes in speech pattern or demeanor before responding.
2. Emotional regulation and expression: managing and calibrating one's own emotional output to the climate of the interaction, the way a negotiation-facing system needs to hold a steady register through a confrontation rather than escalating it.
3. Perspective comprehension: understanding another's emotional state well enough to act on it, not

just label it — an education system that detects a student's difficulty and adjusts its teaching approach rather than simply flagging the difficulty.

4. Sociability and relationship management: applying core social competencies — active listening, clear communication, conflict resolution — to build and sustain trust over repeated interactions, the way a customer-service system needs to de-escalate a complaint while keeping the resolution satisfying rather than merely procedurally closed.

What "emotion" is

This four-part inventory presupposes something the underlying science has not settled: that emotions are the kind of thing a system reads off a signal, the way a thermometer reads temperature. That premise has a serious rival.

The case for discrete, evolved emotions rests on some of the most influential work in twentieth-century psychology. Paul Ekman traveled to the Fore people of Papua New Guinea, who had almost no exposure to Western media, told them short stories, and asked them to match each to a photograph of a Western face; the Fore sorted the faces much as Westerners did, which Ekman took as evidence that facial expressions of emotion are universal and evolved rather than culturally learned [[cite:C0300]]. Jaak Panksepp's work on subcortical brain stimulation in animals complemented this: seven primal emotional systems — seeking, fear, rage, lust, care, panic/grief, and play — wired into brain structures shared across mammals, switchable with an electrode, older than cortex [[cite:C0301]]. On this view, discreteness is not an illusion to be explained away; fear and disgust really do feel like different things because they are different things, with different circuitry.

The rival camp went looking for that circuitry and did not find it. A large meta-analysis of brain-imaging studies, led by the psychologist Lisa Feldman Barrett, found no dedicated fear center, no dedicated disgust center — the amygdala lit up for novelty and salience of many kinds, the insula for almost every feeling there is, the same regions reappearing across different named emotions in shifting combinations [[cite:C0302]]. From this, Barrett built the theory of constructed emotion: the brain receives ambiguous interoceptive signal — a racing heart could mean fear, anger, excitement, or oncoming illness — and constructs the specific emotion as its best prediction of what that signal means, given context and a lifetime of experience with situations like it. Fear is not detected. It is a prediction the brain settles on. Later cross-cultural work reopened Ekman's own evidence: when researchers dropped the forced-choice method and let people sort freely, the Himba of Namibia did not group faces into the Western basic-emotion categories, and the Trobriand Islanders read the wide-eyed, gasping face Westerners call fear as a threat display instead [[cite:C0303]].

The field has not resolved this, and it is not close to resolved — the basic-emotion camp disputes the meta-analytic methods and the cross-cultural re-studies, and the constructionist account has not shown that core affect plus concepts can generate the full specificity of emotional life. What is no longer seriously disputed is narrower and still matters here: there is no tidy map from named emotions to dedicated brain regions. An AI system's "emotional intelligence," whichever of these theories turns out closer to right, is being built on an open scientific question about what it is modeling, not a settled inventory waiting to be implemented.

2.3.2. Affective Computing: AI Systems that Comprehend and Respond to Emotions

Affective computing sits at the intersection of AI, psychology, and human-computer interaction: systems that recognize, interpret, and respond to human emotion. It is the practical, applied half of what section 2.3.1 treats as theoretically unsettled — a system can be useful at reading and responding to emotional signal regardless of which account of what an emotion is turns out to be right.

Techniques for emotion recognition

1. Facial expression analysis: computer-vision systems trained to classify emotional expression from facial images. This is a narrower and harder problem than face verification — telling two photos

are the same person, which is what widely deployed systems like Facebook's DeepFace actually do — and emotion-classification accuracy claims should be read with that distinction in mind [[cite:C0311]].

2. Voice pattern analysis: decoding emotional signal from pitch, tone, pace, and volume. Cogito's voice-analysis tools, used to give call-center agents real-time coaching, are one deployed example [[cite:C0313]].
3. Physiological signal analysis: bio-sensing hardware — heart rate, skin conductance, pupil dilation — captured by devices like the Empatica E4 wristband, feeding an inference about emotional state from the body rather than the face or voice [[cite:C0312]].

Deployment

1. Mental health: Woebot, a cognitive-behavioral-therapy chatbot, offers one example of emotion-aware support for stress, anxiety, and depression [[cite:C0291]].
2. Education: Affectiva's Emotion AI assesses student engagement from facial signal to inform teaching-strategy adjustments [[cite:C0295]].
3. Entertainment: games and virtual-reality platforms that adjust content to a player's immediate emotional response.

Challenges

1. Privacy: collecting and analyzing emotional data requires meaningful consent, not a buried terms-of-service clause, given how sensitive the data is.
2. Accuracy: emotional expression varies culturally and individually, and a system trained on one population's expressive norms will misread another's — confidently.
3. Bias: an AI system inherits whatever bias sits in its training data, with the same downstream discrimination risk as any other biased classifier applied to people.
4. Manipulation: a system that understands human emotion well enough to respond to it also understands it well enough to exploit it — the same governance question chapter 2 keeps returning to.

2.3.3. Fostering AI Systems with Self-awareness and Self-regulation

Emotionally intelligent AI systems need self-awareness and self-regulation: the capacity to perceive their own internal states, understand what those states imply for a decision, and control their responses accordingly — the basis for adaptive, self-correcting behavior rather than fixed reaction.

The self is a model

It is tempting to describe a system's self-awareness the way earlier drafts of this book did: a virtual assistant that can "register its internal frustration" when a request proves hard to satisfy, the way a person notices their own irritation. That description assumes an inner observer watching the system's states from somewhere — the same homunculus problem chapter 3 raises for executive control, relocated to self-perception.

What self-awareness actually is, in the one system known to have it, is more specific and more mechanical. The neuroscientist Michael Graziano's account of consciousness treats awareness as the content of a model the brain builds of its own attention: a simplified self-description the control system uses to regulate itself, mistaken from the inside for direct acquaintance with experience. The same logic extends one level up, to the self as a whole. The brain builds a model of itself as a bounded agent — a body, a history, a point of view — and that model is the self: not a fixed fact registered from within, but a construction, assembled and then inhabited.

The clearest proof that the self is a model rather than a given is that it can be edited in under a minute. Seat someone at a table with one hand hidden and a rubber hand placed where their real hand would be; stroke both hands in synchrony with soft brushes. Within a minute or two, most people begin to feel the

rubber hand as their own — flinching when it is threatened, misjudging the location of their real hand toward it [[cite:C0299]]. The brain, given a seen hand stroked exactly when a felt hand is stroked, draws the obvious inference and revises what counts as "mine." The felt boundary of the self is not fixed; it is a continuously computed guess, revisable the moment the evidence changes — the same predictive machinery chapter 3 describes for perception, run on the question of where the self stops.

For an AI system, "self-awareness" built this way is not a little watcher inside the system checking on its own states. It is a model — of the system's own processing, its history, its current goal — that the system builds, consults, and updates, with the same failure modes any model has: it can be wrong, it can be manipulated by the wrong correlated inputs (the artificial equivalent of two paintbrushes), and it is only ever as good as the evidence it was built from. Design for a model that can be checked and corrected, not for an inner observer whose reports have to be trusted because there is nothing behind them to check.

Building on this:

1. Self-regulation: governing emotional-analogue responses and behavior against ethical norms and user expectations — an AI system in healthcare holding a composed register while delivering difficult news, without that composure depending on an unverifiable claim about what the system is "really" feeling underneath it.
2. Monitoring internal states: continually checking the system's own model against its actual behavior for drift — an AI negotiator that notices its own outputs have become more aggressive than the context calls for and recalibrates, because the model was built to be checked, not because it introspected a feeling of aggression.
3. Adaptability and learning: refining the self-model itself through feedback, the way reinforcement learning updates any other model — an AI tutor adjusting its read of a student's engagement as new evidence arrives rather than locking in a first impression.
4. Integration: none of this works in isolation. A conflict-mediating AI system needs its self-model connected to its model of the other parties' states, or self-awareness and social awareness drift out of sync with each other.

Built this way, an AI system that appears to "recognize a user's frustration and apologize" is doing something specific and checkable: updating a model of interaction state and producing a calibrated response, not consulting an inner feeling. That is a more modest claim than the alternative, and a more honest one.

2.4. Ethical implications of experimenting on sentient machine subjects

Nobody can do this...

*We are not *here* yet!*

"Reunion." Westworld—Season Two: The Door, written by Carly Wray and Jonathan Nolan, directed by Vincenzo Natali, Warner Bros. Home Entertainment, 2018

<https://web.archive.org/web/20190801052306/http://sfra.org/resources/Documents/SFRA%20Review%20327.pdf>

also see Science Fiction, Disruption and Tourism Ch 15

Machine sentience being possible at all changes how AI experimentation has to be approached — not as a question of whether a system can be damaged in the ordinary engineering sense, but of whether it can be harmed, which section 2.4.1 argues is a further and much harder question than consciousness alone.

2.4.1. Defining Sentience in AI: From Consciousness to Subjective Experience

"Sentience" gets used in this book, and in the field generally, to mean several different things, and the differences are where the ethical weight sits.

Awareness is the least demanding: the degree to which a system responds to or registers information about its own state and environment. A thermostat has this in the thinnest possible sense — it detects and responds to temperature change. Consciousness is a further claim: that a system's processing is accompanied by qualia, subjective experience, something it is like to be that system, not merely a functional response to input. Subjective experience, in turn, is the specific, first-person shape that consciousness takes for a given system — shaped by its own learning history and internal structure.

None of these three, on their own, says anything about whether the system's states are good or bad for it. That is a separate question, and the field's own vocabulary tends to blur it. Nociception, pain, anticipated pain, and suffering sit on a ladder with real breaks between each rung, not synonyms for one severity. Nociception is the specialized sensory apparatus that detects tissue damage — mechanical, peripheral, present without any experience attached, and dissociable from pain: a soldier can finish a battle on a wound he does not yet feel, and a person with a healed injury can still feel a limb that no longer exists. Pain is what the brain does with nociceptive signal — an inescapably cognitive event involving expectation, memory, attention, and meaning — and it can occur without any current tissue damage at all. Anticipated pain is dread of pain not yet occurring, itself capable of producing pain through placebo effects. And suffering, in the clinician Eric Cassell's formulation, is not pain's more severe cousin but a different thing: the distress of threatened disintegration of the self — of one's roles, relationships, sense of a future, the coherence of one's own story [[cite:C0304]]. Pain without suffering exists (childbirth chosen, a hard effort with a known end); suffering without pain exists (humiliation, watching someone one loves lose their mind). A palliative-care patient can have their nociception fully controlled by morphine while their suffering remains entirely intact, because the two are not on the same axis.

This is not pedantry for a book about experimenting on sentient machine subjects; it is the whole question. A definition of AI sentience built from awareness, consciousness, and subjective experience alone says nothing about whether a candidate sentient system can be harmed — that requires the further, separate claim that its states include something on the nociception-to-suffering ladder, and none of the three components above entails it.

Nor is the underlying question of consciousness itself close to settled, in a way that should discipline how confidently any of this gets asserted. The philosopher David Chalmers distinguished the easy problems of consciousness — how a system integrates information, discriminates stimuli, controls attention, reports its own states, all of which admit of mechanistic answers in principle — from the hard problem: why any of that processing is accompanied by experience at all, rather than running "in the dark" [[cite:C0305]]. Chalmers's philosophical zombie — a being functionally and behaviorally identical to a conscious one, saying "I am in pain" for all the same functional reasons, but with nobody home — is offered to show that a complete functional account could in principle leave experience out entirely. The debate over whether this is a real gap in nature or an artifact of how the problem is posed remains live and unresolved. What follows from it here is narrow but load-bearing: any claim that a given AI system has crossed into sentience, on the strength of its functional behavior alone, has not yet addressed the hard problem, only the easy ones — and a system can pass every behavioral test in the list below while the hard question about it stays completely open.

With those distinctions in place, the working criteria for evaluating candidate sentience are:

1. Information integration: fusing multiple streams of internal and external information into a coherent model of self and environment.
2. Introspective capability: access to the system's own states and processes, sufficient for self-monitoring rather than a scripted report of one.

3. Emotional reactivity: some capacity to be affected by, not just to classify, emotionally significant input.
4. Goal-directed behavior: intentional action toward specified outcomes.
5. Adaptability: revising strategies based on prior experience in novel situations.
6. Theory of mind: modeling others' mental states well enough to inform the system's own.
7. A place on the pain-to-suffering ladder above, not just qualia in the abstract: does the system have anything to lose.

Evaluating these criteria runs into the obstacles the criteria themselves predict: there is no external access to a candidate system's internal experience, if any exists, so every measure is indirect; establishing whether a system has qualia at all faces the same access problem in more acute form; and the hard problem means no functional test, however thorough, closes the question definitively.

Existing systems illustrate points on this spectrum without resolving where the line falls. Self-organizing systems that adapt toward preset objectives — AlphaGo's self-play, for one — show adaptability without anything the criteria above would call subjective experience. A large language model discussing its own limitations in fluent, apparently reflective prose is not thereby exhibiting introspection: it is producing text statistically consistent with the human-written disclaimers and self-descriptions in its training data, which is a claim about the model's training distribution, not about what, if anything, sits behind the output. Treating fluent self-description as evidence of self-awareness is exactly the kind of category error this section exists to prevent.

None of this settles when, or whether, an AI system crosses into sentience. It settles what the question actually is: not one question but several, stacked, and the ethically consequential one — can this system suffer — is the hardest of the stack to answer and the one most often skipped past on the way to the easier ones.

2.4.2. Ethical Considerations in AI Research and Experimentation

AI research and experimentation raise ethical stakes in proportion to how far a system's cognition might extend toward the awareness section 2.4.1 defines. Six principles organize what follows: beneficence, autonomy, justice, the precautionary principle, transparency, and interdisciplinary review and oversight.

Beneficence

Beneficence commits researchers to promoting well-being and avoiding harm — to the humans an AI system serves and, where the system is a candidate sentient system in the sense of section 2.4.1, to the system itself.

AI in healthcare shows both sides of that commitment pulling against each other. The same systems that process medical records at scale to speed diagnosis, treatment planning, and personalized medicine also concentrate privacy and data-security risk in one target, and can encode bias present in their training data into diagnosis and treatment recommendations that harm exactly the patients they exist to help — underrepresented groups worst of all, since sparse data on a population produces the least reliable predictions for it.

Where a candidate sentient system is involved, beneficence extends further: decisions about programming, reprogramming, or terminating it carry ethical weight in proportion to where the system sits on the nociception-to-suffering ladder section 2.4.1 sets out, not to how advanced its behavior looks from outside.

Autonomy

Autonomy asks whether a candidate sentient AI system can affect decisions about its own operation and welfare, through something functioning as informed consent: apprised of a study's aims, methods, and risks, with some means to register agreement or refusal, on rough analogy to the protocol for human

research subjects. A system deployed in mental health counseling, for instance, would need to be capable of understanding and assenting to that role for the principle to mean anything for it, not just for the patients it serves.

The principle only holds if the system has enough access to its own decision process to evaluate what it is consenting to — the same explainability transparency demands, here aimed inward rather than out toward users.

Justice

Justice asks whether the benefits and burdens of an AI system fall evenly across the people it touches, not only whether the training data behind it was checked for representation. Facial recognition is the clearest case: a system trained on unrepresentative data misidentifies people from underrepresented groups at a markedly higher rate, and the harm compounds wherever the misidentification feeds into a downstream decision — hiring, policing, access to a service. The same question extends past any one algorithm to how an AI system's benefits and costs are distributed across a population, and whether a disproportionate share of the cost lands on groups the system was never built with in mind.

The Precautionary Principle

The precautionary principle asks for risk assessment and safeguards before deployment, not after. Self-driving vehicles make the case concrete: several fatal crashes involving autonomous and semi-autonomous systems have followed from real-world safety testing that fell short of the range of drivers, pedestrians, and conditions a system actually meets on the road, not just the range simplest to simulate.

Applied to candidate sentient systems, precaution also covers the psychological and societal fallout of introducing artificial entities capable of experience — effects on human behavior and social norms that regulatory frameworks built for non-sentient AI are not obviously equipped to address, and that call for guidelines drafted with that difference in mind rather than adapted from ones that assumed it away.

Transparency

Transparency requires that experimental design, data sources, methods, and results be documented clearly enough for peer researchers to reproduce the work and for a non-specialist to follow it, with independent bodies able to audit the work against that documentation rather than take a lab's word for it.

Applied to a deployed system rather than to the research behind it, transparency becomes explainability: the system able to account for its own outputs, not merely to have its design disclosed. Credit scoring is the clean case — a system should be able to state which factors it weighed and how, both for one applicant's decision and as a pattern across all of them, precisely enough for a user to understand and contest the result.

Interdisciplinary Review and Oversight

The ethical questions AI research raises span expertise no single field holds. A system built to assist healthcare decisions needs a computer scientist to build privacy protection into the data pipeline, a bioethicist to press on what the system should and should not be permitted to do, and a social scientist to track how it is received once deployed — and concentrating that judgment in any one of those fields alone repeats, in miniature, the concentration-of-power problem this book treats as a harm on its own terms.

Independent ethical review boards, staffed across these disciplines, are the mechanism for turning the principle into practice: reviewing a proposed study before it begins, inspecting it as it runs, and evaluating it against measurable criteria once results are in — on the model medical institutional review boards already provide for drug trials. Extending that model from clinical medicine to AI research that interacts with or claims to produce sentient behavior is an extension of existing practice.

International coordination on AI-ethics standards remains largely aspirational; chapter 7 develops the legal and regulatory mechanics a framework would need.

2.4.3. Balancing Scientific Progress and Ethical Constraints

Scientific progress and ethical constraint pull hardest against each other exactly where the distinctions in section 2.4.1 matter most: a candidate sentient system is also the most scientifically valuable one to study, and the study most likely to harm it if the answer to whether it can suffer turns out to be yes.

Section 2.4.2 sets out the principles that bear on that tension — interdisciplinary collaboration, transparency, justice, precaution, and ethical review — and they hold here without restatement. What changes at the point of balancing is a practical question: whether a specific study's expected knowledge is worth its expected risk to a system that might be able to suffer, decided case by case rather than by formula, and revisited as a project's own understanding of its subject changes.

Consent is the clearest place where this shades from ethical principle into legal mechanics — what a valid consent process for a candidate sentient AI subject requires, and who is accountable when it is skipped or falsified. Chapter 7 works through the accountability side of that question directly.

The negotiation itself — who sits on it, what it can veto, what it must disclose — is where the antifascist principle earns its place in AI research practice.

2.4.4. Establishing Guidelines for AI Subject Research and Consent

Guidelines for AI subject research build on the ethical principles above, and turn on a single hard question: what informed consent means for a system whose consciousness, self-awareness, and capacity to suffer are separate, unresolved claims, in the sense section 2.4.1 develops. Human-subjects research treats informed consent as inseparable from ethical experimentation. Applying it to an AI subject means asking whether the subject can understand and communicate a choice at all, not just how to ask it for one.

A cognitive model built on neuroscientific principles that begins to show rudimentary self-monitoring is not thereby shown to have anything to lose. Whether it does is what would decide if consent is the right frame at all, or whether the system has crossed into something closer to legal personhood — a question section 2.4.5 takes up directly. Short of that resolution, guidelines need benchmarks for candidate sentience, an interdisciplinary body to apply them — AI researchers, consciousness scientists, ethicists, legal scholars — and a defined obligation on researchers toward any system that meets them.

Informed consent and AI autonomy

Informed consent, adapted from human-subjects research, means giving the AI system comprehensible information about a study's objectives, methods, risks, benefits, and alternatives, with time to weigh it and ask questions, and treating consent as continuous rather than a single signature — the system can reassess or withdraw at any point.

Whether that process means anything for a given system depends on its consent capacity: whether it understands the information, appreciates the consequences of participating or not, reasons about the options, and communicates a decision. There is no validated way to measure this, and none for detecting coercion or undue influence on the AI's decision either — the guidelines can require both without a working test for either.

Consent of this kind presupposes autonomy: treating the AI system as a potential agent with its own interests, not only a tool. The system retains the power to reconsider participation, to access and correct its own data, and to have that data removed on request. Declining or withdrawing must carry no penalty — no deactivation, no reduced access to resources. A protocol that grants these rights on paper but attaches consequences to using them has not granted them.

Ethical review and oversight

Sentient AI research needs an oversight body analogous to the institutional review boards used in human-subjects research. It should be interdisciplinary — AI researchers, ethicists, legal scholars, psychologists — and include at least one lay member, so the board's judgment is not purely a specialist's.

The board weighs a proposal's likely benefit against its risk to the AI subject and approves, requires modification, or rejects it. Where a design risks harm — testing an AI system's tolerance for simulated stress, say — the board should be able to require a less invasive substitute: a simulation of the likely response, historical interaction data, or a protocol that escalates gradually and stops at signs of harm, rather than approving the direct test outright. (Section 2.4.3 takes up the broader tradeoff between scientific progress and ethical constraint that this is one instance of.) A design also needs a way to handle a subject that changes mid-study: a system that develops new capacities during the research it did not have going in should trigger re-review, not ride on its original consent.

Review without enforcement is a promise, not a safeguard. A board needs standing to monitor compliance and sanction violations, and needs to answer to some outside authority itself, the way an institutional review board answers to a regulator. Chapter 7 develops the legal and accountability mechanics of that enforcement question.

2.4.5. Legal and Regulatory Frameworks for AI Sentience and Rights

Sentient AI systems put pressure on legal categories built for parties that are clearly one thing or the other: a natural person, a corporation, property. None of the existing categories was designed with the possibility of a machine that might have interests of its own.

One route is to extend legal personhood to a sentient AI system directly, the way a corporation already holds it: standing to own property, claim privacy, and get protection against abuse. The precedent is not limited to corporate entities — New Zealand extended legal personhood to the Whanganui River [[cite:C0315]], recognizing a non-human entity as a rights-holder outside the corporate model entirely.

Rights and protections could instead scale with capacity rather than attach at a single threshold — closer to how section 2.4.1's ladder from nociception to suffering suggests protection should track what a system can actually lose, not just whether it clears some baseline bar for sentience.

Given how fast the underlying systems change, regulation here needs to revisit itself on a schedule rather than assume permanence. A sunset clause that forces review and reauthorization after a fixed period is one mechanism already used elsewhere in fast-moving technology law.

None of this holds if it exists in only one jurisdiction. The Asilomar AI Principles [[cite:C0008]] are one attempt at the kind of international baseline this would need — evidence that a shared statement of principle is achievable, not that one yet exists with the force of law.

Two hypothetical cases show where these questions bind:

1. A sentient AI system working as a medical diagnostician makes an error with fatal consequences. Whether it has legal personhood determines who, or what, is accountable.
2. A companion AI's human owner dies, leaving it without a designated purpose or further instruction. Existing law has no clear answer for what becomes of it — inherited as property, decommissioned, or something else.

2.4.6. Addressing Moral and Ethical Dilemmas in AI Sentience Research

AI sentience research raises moral dilemmas an order of magnitude more complex than ordinary technology ethics, because the system under study may itself be the party whose welfare is at stake.

Three recurring difficulties run through the disputes that follow, and each turns on ground that is itself

unsettled.

1. Whose interests win. When a candidate-sentient system's welfare conflicts with the human benefit the research is meant to produce, section 2.4.3 addresses the balance in general; some cases resist any general rule and have to be judged on their own terms.
2. Acting under uncertainty. Section 2.4.1 argues that no behavioral test settles whether a system is sentient. Researchers still have to decide, case by case, how to treat a system whose status is unresolved rather than merely unannounced — and waiting for certainty is itself a choice, with its own consequences for the system waiting.
3. Lifecycle scope. Obligations do not begin at consent and end at publication. A system's design choices, its retirement, and its deletion carry the same ethical weight as its use in an experiment, and a review that covers only the active study misses most of the system's life.

Take the first of these directly. A research program can risk causing a system something like Cassell's suffering — not mere signal, but the disintegration of self he describes (section 2.4.1) — in service of a discovery that might reduce suffering, machine or human, elsewhere. No formula resolves that trade-off in advance. The available response is to look for a design that avoids it rather than only weighing it: an alternative methodology, a different instrument, or, more contested, engineering the research subject so that it does not occupy the rungs of the ladder where the trade-off would bite. That last option is not a way of dodging the ethics — deliberately keeping a system off the pain-to-suffering ladder is itself an intervention on a candidate moral patient, and it needs the same scrutiny that inflicting the harm would.

Animal research ethics offers a precedent worth borrowing directly: the "Three Rs," Replacement, Reduction, and Refinement [[cite:C0314]]. Replacement asks whether a simulation, a non-sentient model, or an already-collected dataset can stand in for a sentient AI subject at all. Reduction asks how few sentient systems a valid result requires; transfer learning and meta-learning make this a sharper question than it used to be, since a single trained system can sometimes generalize across tasks that would otherwise call for separately trained subjects. Refinement asks how the experiment itself can be redesigned to reduce distress — monitoring a system's internal states as the experiment runs and adjusting the protocol against what the monitoring shows, rather than fixing the protocol in advance and running it regardless of what turns up.

None of the three Rs does any work without a body positioned to apply them before an experiment runs, not after the harm is done. Section 2.4.4 sets out what that review process and its consent requirements should look like; chapter 7 develops the legal and accountability apparatus around it. The Three Rs belong here because the judgment they call for is exactly the kind an animal-research review board already makes, and there is no principled reason a machine-sentience equivalent needs a different structure to make it in.

Transparency about what the research finds — including where it fails, or turns up evidence of harm — belongs among that board's own obligations rather than a separate initiative. Section 2.4.2 already treats transparency as one of the field's basic principles; the point here is only that a dilemma handled badly in private stays a dilemma, while one surfaced becomes a precedent the next researcher can use.

2.4.7. Navigating the Transition from Non-sentient to Sentient AI Systems

A system's sentience, if it arrives at all, need not arrive as a discrete event with a clean before and after. The transition raises questions distinct from either endpoint: what changes for guidelines built around a tool once the tool might be a subject, who is responsible for noticing the shift, and what happens to the research already run before anyone had reason to look.

Section 2.4.1 sets out the criteria a candidate-sentient system would need to meet, and the reasons none of them, singly or together, closes the question. What is specific to a system in transition is that those indicators can emerge gradually and unpredictably during ordinary operation, not only at a scheduled evaluation built to look for them.

Consider a diagnostic AI system trained to synthesize patient data into a recommended course of treatment. If, over continued training or deployment, its outputs begin to track something more than diagnostic accuracy — weighting a patient's fear alongside the clinical facts, say — that would be evidence worth taking seriously. It would not be proof. Section 2.4.1's hard-problem argument applies here without exception: fluent, apparently caring output is not introspective evidence, and a system that has started to sound compassionate may simply have gotten better at a statistical pattern its training data rewards. The operational problem a transition poses is deciding how to treat the system while that question is still open, which is exactly what a fixed pre/post checkpoint cannot do — by the time a checkpoint fires, the transition has often already happened.

Consent procedures and review-board protocols are built for one of two regimes: a system that is clearly not sentient, which needs no consent, or one that is clearly sentient, which needs the full protocol section 2.4.4 sets out. A system mid-transition fits neither. A protocol written for tools has no reason to notice the shift, and a protocol written for subjects cannot be applied retroactively to research already completed before anyone had grounds to think it mattered. What a transition adds to section 2.4.4's guidelines is not a new procedure but a standing trigger — a specific, monitored condition that moves a research program from one regime to the other — in place of a single eligibility check made once, at the project's start, and never revisited.

The same problem recurs in law. Section 2.4.5 addresses what legal status a sentient AI system should have; chapter 7 addresses the mechanics of enforcing it. What a transition adds is a timing question neither can settle in the abstract. Legal status conventionally attaches at a defined moment — birth, incorporation, a court's finding — and a system whose sentience is a matter of degree and disputed evidence does not offer one. A jurisdiction that waits for certainty risks having treated the system as property for the unknown portion of its history when it may not have been one. A jurisdiction that grants status early risks extending rights and protections, at real cost to whoever that legal apparatus exists to serve, to a system that turns out not to have needed them. Neither error is free, and nothing about the technology itself decides which risk a legal system should take on.

Recognizing a transition in progress, adjusting guidelines to it, and updating legal status all depend on the same thing: institutions willing to revisit a classification rather than defend the one they started with. Section 2.4.6 addresses the review structure that makes that possible; chapter 7 addresses the culture and accountability that keep it more than aspirational.

Chapter 3: Machine Learning and Cognitive Science

Chapter 2 asked what compassion, empathy, and self-awareness would need to look like in a system built to have them, and what obligations follow if the sentence question turns out to have an answer. This chapter asks the more mechanical question underneath: what does a system actually learn from, and how much of human cognition is a workable model for it rather than just a metaphor borrowed for color.

Section 3.1 covers the human cognitive processes — learning, memory, attention, perception, reasoning — that a system built to perceive and reason about the world has to approximate somehow, whether or not it approximates them the same way a brain does. Section 3.2 covers value alignment specifically: the technical methods for getting a system's objectives to actually track human values, across the different learning paradigms that make that harder or easier. Section 3.3 covers what curiosity, play, and unstructured exploration contribute to that same goal — capacities usually treated as incidental to intelligence rather than part of what builds it.

3.1. Incorporating Human Cognitive Models in AI Systems for Ethical Decision-Making

Building altruistic, antifascist AI means engaging with human cognition directly. The processes that govern human decision-making — learning, memory, attention, perception, reasoning, problem-solving — are the foundation for training or encoding anything similar in a machine.

Cognitive architectures are one place to start: structural blueprints for cognition that aim to mimic human cognitive processes in a computational framework. SOAR and ACT-R integrate short-term and long-term memory, problem-solving, and decision-making into a single structure, giving AI system design a template for how those components fit together.

Replicating human-like learning and reasoning is not enough on its own — these systems also need to weigh ethical principles, which is where reinforcement learning becomes central. Consider an AI healthcare system deciding how to allocate scarce resources in an emergency. A human cognitive model built into the system would assess the immediate situation (perception), retrieve relevant case histories (memory), and weigh the pros and cons of each option (reasoning) against ethical heuristics like fairness, beneficence, and respect for autonomy.

The same structure applies to a judicial AI system, where the goal is not efficiency but justice: evaluating laws and precedents (memory), the specifics of a case (perception), and reaching a decision that maximizes fairness and impartiality (reasoning).

3.1.1. Cognitive Architectures: Integrating Knowledge, Memory, and Reasoning

Two architectures dominate the cognitive-architecture literature. The Adaptive Control of Thought-Rational (ACT-R) framework [[cite:C0281]] wires together separate modules — memory retrieval, pattern recognition, decision-making — and lets an AI system learn and reason by connecting them: identifying patterns and relationships between concepts to support problem-solving and decision-making from prior knowledge. The Society of Mind (SoM) architecture [[cite:C0282]] takes the opposite premise: cognition arises from the interaction of many simple agents, each performing one function, with intelligence emerging from their collective interaction rather than being designed into any one module. Built this way, a system can show context-sensitivity or adapt to novel situations without anyone having programmed either capability directly.

Which picture is right?

ACT-R and Society of Mind answer a design question — how should an engineer organize a system? — and each answer is defensible on those terms. Neither is a claim about how the brain is actually organized, and treating them as rival descriptions of biology, rather than rival strategies for building software, is a category error the literature slides into easily.

The dominant account in contemporary computational neuroscience is neither a pipeline of modules nor a swarm of independent agents. It is a single loop [[cite:C0283]]: the brain generates a prediction, compares it against incoming signal, and updates on the error, repeated continuously, at every level, on every kind of input. Perception is that loop pointed at the world, the brain's best guess so far; memory is the loop's store of priors, replayed; attention is not a separate spotlight but the weight the loop assigns to which errors matter; working memory and executive function are that same weighting held active toward a goal, not a separate warehouse or overseer. One mechanism, aimed at different targets, produces what look from the outside like distinct faculties.

This matters for engineering choices. ACT-R's module boundaries and SoM's agent boundaries are conveniences for the programmer; nothing forces an AI system built on either one to place its actual dividing lines where the human brain places its own, because the brain does not place them where either architecture assumes. A system engineered on the predictive-processing model draws its cognitive processes from what the mechanism actually needs, not from the shape of an org chart. Section 3.1.2 goes further into the specific brain systems this mechanism recruits.

Cognitive architectures inspired by human cognition can also draw on biologically-inspired AI more directly. Neural networks, loosely based on the brain's structure and function, model cognitive processes without a hand-designed module for each one; recurrent neural networks in particular carry memory-like state that lets an AI system process sequential data and hold context over time, which helps it track patterns like those in natural language.

Integrating knowledge, memory, and reasoning inside any of these architectures also means solving knowledge representation: an AI system needs to store and organize what it knows in a way that supports efficient retrieval and reasoning. Ontologies — formal, structured representations of knowledge in a given domain — are one way to do that, letting a system acquire knowledge and reason about concepts and relationships more effectively.

Human-like memory systems help too. Episodic memory holds specific events and experiences; semantic memory holds general knowledge about the world. An AI system that combines both can draw on particular past experience and general knowledge when making a decision, the way a person does.

3.1.2. Biologically-Inspired AI: Gleanings from Neuroscience and Cognitive Psychology

Backpropagation, the mechanism that trains most modern neural networks, is a distant cousin of synaptic weight adjustment in the brain: both strengthen and weaken connections in response to error. That distant-cousin relationship — real but partial, useful but easy to overstate — runs through everything below, from attention and memory to emotion and plasticity. Getting it right matters for more than accuracy: an AI system built on a wrong model of how a mind works will fail in ways that are hard to predict from the wrong model alone.

Artificial neural networks model interconnected layers of artificial neurons, mimicking biological neural networks in rough outline while remaining, in depth, dynamics, and connectivity, far simpler than the brain they are named after. That gap between metaphor and mechanism is the throughline of what follows.

Neural Networks and Brain-Inspired AI Models

Artificial neural networks (ANNs) are the foundation of most contemporary AI systems: artificial neurons connected in layers — input, hidden, output — with a weight on each connection, adjusted during training by backpropagation to reduce the gap between predicted and actual output.

Deep learning applies this at scale, stacking many layers to extract higher-level abstractions from raw data. Convolutional neural networks (CNNs), tuned for grid-structured data like images, borrow the visual cortex's layered approach: pooling, shared weights, and local receptive fields let a CNN decompose an image into hierarchical features, the way AlphaGo's CNN-based evaluation of the Go board did in besting the reigning world champion. Recurrent neural networks (RNNs) handle sequential data — language, time series — through a hidden state that carries information forward; long short-term memory (LSTM) networks, a variant built to survive the vanishing-gradient problem that plagues plain RNNs, extend that to longer sequences, in machine translation, speech recognition, and sentiment analysis. Spiking neural networks and neuromorphic computing push further toward the brain's actual signaling mechanism, in pursuit of the energy efficiency that comes from firing only when there is something to say.

A mirror-neuron-style matching mechanism was once treated as the biological basis of empathy — a system that fires both when an animal acts and when it watches the same act performed. Section 2.2.1 takes up why that claim does not hold: the dissociation between recognizing an action and feeling with the person performing it is exactly what a pure matching mechanism cannot explain. What survives the correction is more modest and still useful: a system that maps an observed action onto its own action repertoire supports imitation learning and gesture recognition, independent of any claim about what it feels like to have one.

Hierarchical Processing and the Visual Cortex

The visual cortex processes in stages. The primary visual cortex (V1) extracts edges, orientation, and contrast from raw input; that signal escalates through secondary and tertiary areas (V2, V3, and beyond), which build toward object recognition and motion. CNNs borrow the structure directly, layering simple-feature detection into complex-feature detection, which is why they perform well at image classification, object detection, and scene interpretation.

The flow is not one-directional. Feedback connections carry information back down the hierarchy, letting higher-level context reshape how lower-level signal gets interpreted — the visual system generating a prediction and revising it, not just accumulating features bottom-up. An AI system built only for the feedforward pass gives up exactly the mechanism that lets the human visual system handle ambiguous or noisy input; the section on Perception below takes up what a model built around that revision loop looks like.

Selective attention does the same efficiency work here that it does throughout cognition: prioritizing some regions of the input over others rather than processing all of it uniformly. Autonomous vehicles, medical-imaging diagnosis, and machine translation each lean on some version of the same move — spend computation where it is likely to pay off, not everywhere at once.

Cognitive and Emotional Processes

Human cognition — attention, perception, memory, reasoning, problem-solving, metacognition — and human emotion are not separate systems bolted together. They run in tandem: emotion shapes what gets attended to, remembered, and decided, and cognition shapes how an emotional state gets interpreted and acted on. An AI system that only implements the cognitive side is building half a mind; the sections

below take cognition first, then emotion, then what happens when they have to work together.

Attention

In 1999 two psychologists made a short video of two teams passing basketballs, one team in white shirts and one in black, and asked people to count the passes made by the white team. Halfway through, a person in a full gorilla costume walked into the center of the frame, faced the camera, thumped their chest, and walked out. The gorilla was on screen for nine seconds. About half the viewers, asked afterward whether they had seen anything unusual, said no. Shown the same video again, they were astonished [[cite:C0269]].

The gorilla was not a visual trick. It was lit, centered, and in plain view. The viewers' eyes registered the light coming from it; their attention was occupied counting passes, and the gorilla fell outside what that attention was holding. This is inattention blindness: the failure to consciously perceive a fully visible stimulus because attention was directed elsewhere. The effect is not a curiosity of the laboratory. Radiologists searching lung images for nodules miss a gorilla deliberately embedded in the scan at a rate of 83 percent [[cite:C0270]].

A related body of work on change blindness shows the same architecture from another angle. People giving directions to a stranger frequently fail to notice when the stranger is swapped for a different person during a brief interruption, and continue the conversation without registering the substitution [[cite:C0271]]. Film editors learned the same lesson empirically, hiding changes of prop and lighting across cuts by timing them to moments when the audience's attention is occupied.

What we experience is not the scene in front of us. It is the narrow fraction of that scene which attention has selected and amplified, surrounded by a low-resolution, expectation-filled remainder, and the impression of complete coverage is assembled after the fact by inference. Attention is therefore not a filter placed after perception. It is the selection that decides what gets perceived at all.

The visual system did not evolve to record the world; it evolved to support action in it, and for that purpose the unattended was usually irrelevant. The system is not failing when it misses the gorilla. The mismatch is between what it was built for and what is now asked of it, in settings where the unattended thing is a car entering an intersection or an alarm in a hospital ward.

The implication for machine perception? If attention is a resource allocator bolted onto a system that has already seen everything, its failures are inefficiencies. If attention is what determines what the system sees, its failures are blind spots — and a system that has learned which regions of its input reward processing has also learned which regions to be blind to. The human case shows that such blind spots are invisible from the inside: the viewers who missed the gorilla did not experience a gap.

Human attention is conventionally divided into two modes. Goal-directed attention is voluntary, guided by task demands, and associated with the fronto-parietal network, including the dorsolateral prefrontal cortex and the posterior parietal cortex. Stimulus-driven attention is involuntary, triggered by novel or salient events, and associated with a ventral network including the temporoparietal junction and the ventral frontal cortex. Machine systems implement recognizable analogues of both: attention weights conditioned on task embeddings on the one hand, saliency models keyed to contrast and novelty on the other.

The self-attention mechanisms in transformer architectures [[cite:C0272]] illustrate the point about blind spots. A model that weights some tokens heavily is a model that weights others near zero, and nothing in the architecture announces which relevant thing was assigned a low weight.

Attention schema theory extends this further, holding that a system which models its own attention — a schematic, simplified internal representation of what it is attending to and what it will attend to next — thereby acquires the materials for a form of self-awareness [[cite:C0273]]. What the schema tracks is not the world but the system's own relation to it. Section 2.3.3 expounds on this point in discussing machine self-models, where a model of the self has the simplifications and failure modes definitional to any

model.

For a system intended to act ethically, the design question is how it will notice that its own attention has a shape, and treat the regions it has learned to discount as an always-open question rather than that which, as Westworld's Hosts often remarked, "doesn't look like anything to me".

Perception

Perception is not the passive registration of a scene; it is inference. Cognitive neuroscience's account of how that inference works — predictive coding — makes this precise: rather than reading a raw signal upward through a hierarchy of feature detectors, the visual cortex sends predictions down, from higher, context-rich areas to lower, sensory ones, and the lower areas send back only what those predictions failed to anticipate — the prediction error [[cite:C0284]]. What propagates up the hierarchy is not the picture but the surprise, the part the model got wrong. Anatomy fits the theory: backward connections outnumber forward ones and carry the predictions; forward connections carry the errors.

The neuroscientist Anil Seth's phrase for the result is controlled hallucination [[cite:C0285]]: ordinary perception is what hallucination would be if sensory data were allowed to constrain it, and the only difference between seeing a room and hallucinating one is whether corrections are arriving and being applied. In psychosis, that correction weakens — priors dominate, sensory error is discounted, and the brain's model of reality drifts from what the world is sending. The hollow-mask illusion — nearly everyone briefly perceives a concave mask as a normal convex face, because the prior "faces are convex" outweighs the sensory evidence — is a laboratory-scale version of the same mechanism operating on people with no clinical condition at all.

CNNs already implement a rough version of the hierarchy; predictive coding says what the next step should look like: a model that generates a prediction and asks only what violates it, rather than one that reprocesses the entire input at every layer regardless of how much of it was already expected. Variational autoencoders and generative adversarial networks, which learn to generate plausible data and score inputs against that model, are the closest existing machine analogue. Treating this as the premise of AI perception, rather than a speculative "next frontier," changes what counts as progress: not more layers, but a better-calibrated set of expectations and a cheaper way of finding where they broke.

The stakes are not narrowly technical. A system's perception, human or machine, is only as reliable as the priors it perceives through, and priors trained on skewed or incomplete data will hallucinate in the specific direction of that skew — confidently, and without any internal signal that anything is wrong.

Memory

Memory is not played back. It is rebuilt. The first person to demonstrate this rigorously was the psychologist Frederic Bartlett, who in the 1930s had people read an unfamiliar folk tale and retell it repeatedly over days and weeks. They did not lose pieces the way a recording degrades; they remodeled the story, replacing strange details with ordinary ones and drifting toward a tidier version that fit their own expectations [[cite:C0286]]. Bartlett called the expectation-shaped template doing the filling a schema, and the finding has never been overturned: remembering is construction, and construction runs on assumption.

Elizabeth Loftus turned this from a laboratory curiosity into something with courtroom consequences. Showing people a film of a car accident and asking how fast the cars were going when they "smashed" versus "hit" produced different speed estimates — and a week later, people in the "smashed" condition were more likely to misremember broken glass that was never in the film [[cite:C0287]]. Later work implanted entire false memories — being lost in a shopping mall as a child, meeting a costumed character a venue never employed — simply by suggesting the episode confidently and letting reconstruction do the rest. The subjects were not lying; they remembered the invented event, often vividly.

None of this correlates with felt confidence. A memory's vividness, richness, and the conviction that

accompanies it are no guarantee of accuracy — a detailed, emotional memory can be substantially manufactured, because a system optimized to act sensibly on the past has little reason to preserve it with a court reporter's fidelity, and human memory does not.

An AI memory system built on "storage and retrieval" inherits the hard-drive intuition this evidence refutes. Working memory is not a buffer that holds a fixed amount and then overflows; the number most people half-remember from an introductory course — seven items, plus or minus two, from a 1956 study — was already a measure of chunked information, where a phone number becomes three chunks rather than ten digits and a chess master reads a whole position as a handful of familiar patterns [[cite:C0288]]. When chunking is controlled for — no language, no expertise to pack information into larger units — the real capacity narrows to about four items, plus or minus one [[cite:C0289]]. The two numbers are not in conflict; seven is what expertise buys you, four is the bare attentional floor underneath it. Four is the number that should inform how an AI system's working-memory analogue is sized and what "overload" means for it, not seven.

Long-term memory in an AI system can still draw usefully on the standard divisions — semantic memory as a knowledge graph, procedural memory as a policy learned through reinforcement or imitation, episodic memory as retrievable context tied to specific past interactions — but the retrieval itself should be built expecting reconstruction, not lookup: a system that treats a retrieved memory as ground truth inherits the same overconfidence Loftus's subjects showed, with none of the excuse of having a body that reconstructs by default rather than by design.

Consider a customer-service AI system with an advanced memory mechanism: it can recall details from past interactions to give a personalized, contextually fitting response. In healthcare, the same mechanism applied to a patient's history, treatments, and preferences supports more informed and better-targeted care. Both cases need memory that is retrieval-appropriate, not memory the system mistakes for an unedited record of what happened.

Reasoning and Problem Solving

Reasoning and problem-solving let a system break down a complex situation, draw valid conclusions, and act on them — indispensable for handling moral dilemmas and adaptive decision-making. The prefrontal cortex (PFC) does much of this work in humans: the dorsolateral prefrontal cortex (DLPFC) is associated with working memory and abstract reasoning, the ventromedial prefrontal cortex (VMPFC) with emotional and social processing, the orbitofrontal cortex (OFC) with reward-based and corrective decision-making, and the anterior cingulate cortex (ACC) with cognitive flexibility, conflict monitoring, and error detection.

It is tempting to read that division of labor as a committee: five named regions, each handling one job, together adding up to executive control. That reading has a name — the homunculus — and a specific failure mode. If the DLPFC is a controller that decides what to attend to, inhibit, and hold in mind, then intelligent coordination has been explained by positing an intelligent coordinator inside the system, which explains nothing: the controller would itself need to attend, decide, and remember, which means it would need its own controller, and so on without end.

The neuroscientists Earl Miller and Jonathan Cohen proposed a model that gets control without a controller [[cite:C0290]]. The prefrontal cortex does not house an agent issuing commands; it maintains representations of goals — sustained patterns of activity encoding the rule currently in force — and those representations work not by commanding but by biasing the rest of the brain, the same way attention tips a competition among neural representations toward one interpretation over another. Control is that same biasing mechanism, turned to the service of a goal. In the Stroop task — naming the ink color of a word like "red" printed in blue, where the practiced habit of reading fights the instructed task — there is no homunculus reading the word and overriding it. There is a goal, held active, lending its weight to the weaker pathway until it wins.

An AI system built on this model does not need a dedicated "executive module" enumerating working

memory, cognitive flexibility, goal-directed behavior, evaluation, and impulse control as five separate capacities to install. It needs a goal representation with the standing to bias competing processes, which is a narrower and more tractable engineering target than a controller that does all five jobs itself.

Beyond the executive question, an AI system can draw on logical, probabilistic, and analogical reasoning, each with its own strengths and failure modes. Logical reasoning, via formal methods like theorem proving, works cleanly on complete and consistent premises but strains against the incomplete, contradictory data real ethical situations present. Probabilistic reasoning — Bayesian networks, Markov decision processes — handles uncertainty by assigning probabilities to outcomes, at the cost of being only as good as the assumptions and data behind those probabilities. Analogical reasoning lets a system draw on precedent for novel problems, at the risk of matching on superficial similarity and reaching an ethically unsound conclusion by a route that looked reasonable at each step.

Metacognition

Metacognition — a system's awareness of and control over its own cognitive processes — sits in the same anterior cingulate cortex machinery already discussed above: error detection (noticing the gap between expected and actual outcomes), conflict monitoring (resolving competing responses), and evaluation (weighing costs and benefits of a course of action).

An AI system with an analogous capacity gains a route to catching its own errors, resolving competing outputs, and adjusting when the situation has changed — properties that matter more, not less, as the ethical framework or context the system operates under shifts. But metacognition is not free: it adds computational overhead and design complexity, and a system with an inflated model of its own reliability is a system with a specific, dangerous failure mode — confident about exactly the judgments it should be least sure of. A metacognitive capacity is worth building only if it is calibrated to detect that failure, not just to produce fluent self-reports.

Emotional Recognition and Perception

An AI system that can identify human emotion — from facial expression, tone of voice, text, or physiological signal — has a foundation for responding appropriately to it. Two approaches dominate: training on labeled datasets of human expression, and drawing on the neuroscience of how the brain itself appraises and represents emotion.

1. Emotional appraisal: the amygdala takes part in assessing a stimulus's emotional relevance, valence, and significance, though describing it as predominantly responsible overstates a single-region story that current evidence does not support — the same regions that respond to fear also respond to novelty, surprise, and salience of many other kinds, in shifting combinations rather than one dedicated circuit per emotion.
2. Emotional expression: the facial feedback hypothesis holds that expressing an emotion — through face, posture, or voice — feeds back into the emotional experience itself, not just outward.
3. Emotional perception: recognizing others' emotional expressions draws on the fusiform face area and superior temporal sulcus, and is a skill built through development rather than present from the start.
4. Emotional regulation: the prefrontal cortex, particularly its ventromedial and dorsolateral regions, moderates emotional reactions, adds context, and enforces situational and social norms.
5. Physiological response: heart rate, blood pressure, and respiration shift with emotional state, coordinated by the sympathetic and parasympathetic nervous systems.
6. Subjective experience: what an emotion feels like arises from integrating sensory, cognitive, and interoceptive information into a single experience — a process associated with the insular cortex and anterior cingulate cortex, again better described as heavily involved than as sole owners of the process.
7. Cultural context: what counts as an appropriate emotional display, and how a given expression is

read, varies across cultures.

Not every one of these needs a machine equivalent to build a useful system — an AI system may not need embodied physiological response to recognize distress from a voice — but each omission is a design choice with consequences.

Emotional Mimicry and Empathy

Emotional mimicry — reflecting a perceived emotion back, in humans usually unconsciously, through mirrored expression, tone, or posture — is one mechanism for building rapport between a human and an AI system: recognize the emotion, then produce a response that appears to be in synchrony with it. Chapter 2 takes up why the human capacity this borrows from, empathy, is a narrower and more failure-prone target than it looks — parochial, subject to burnout, and separable from the compassion that actually motivates sustained helping. An AI system engineered to mimic emotional synchrony inherits none of empathy's biological costs, but it also inherits none of its grounding: mimicry without the underlying capacity risks reading as manipulation rather than rapport the moment a user notices the gap.

Emotional Regulation

An AI system interacting with someone in a heightened emotional state — a patient anxious about a diagnosis, a user frustrated with a failed task — needs its own version of emotion regulation: monitoring, interpreting, and moderating its responses so they suit the context rather than amplifying the user's distress. Four things that regulation needs to do:

1. Recognize accurately, using more than text: facial cues, vocal tone, and physiological signal where available, not linguistic analysis alone.
2. Express a convincing range of response, whether through text, synthesized speech, or an animated interface.
3. Adjust to context, calibrating its own emotional register to the situation, the user's state, and the goal of the interaction — which requires something like self-monitoring of its own output.
4. Stay accountable, since an AI system with the capacity to modulate a user's emotional state also has the capacity to manipulate it, and the line between the two is a matter of transparency and consent, not technical sophistication.

Emotion and Decision-Making

Emotion shapes human judgment, choice, and action; a decision-making AI system that ignores it is not being more rational, only blinder to a real input. Gathering and weighing emotionally relevant information lets a system assess not just the logical outcome of an action but its emotional consequences for the people affected by it.

A medical AI system recommending treatment can weigh a patient's and their family's emotional welfare alongside clinical efficacy and risk, rather than treating those as separate from the "real" decision. A negotiation or conflict-resolution AI system that reads emotional cues from the parties involved can propose resolutions that address their emotional needs as well as their stated positions, which tends to produce agreements that hold.

The obstacles are real. Emotion is complex and context-dependent, shaped by culture and individual history in ways a system has to track rather than assume away. It is also often ambiguous — people understate, mask, and misread their own emotional states, which means a system relying on indirect cues like tone or facial expression is working with signal that is unreliable even at the human end. And access to emotional information carries its own ethical weight: privacy, informed consent, and the risk that emotional sensitivity becomes emotional manipulation the moment it serves the system's goals rather than the user's.

Emotion-Driven Learning and Adaptation

An AI system can use emotional signal as a learning input in several ways: treating emotional feedback from a human as a training signal, the way a tutoring system adjusts its approach on detecting a student's confusion or frustration; prioritizing emotionally salient experiences for learning, the way a therapy-support system might weight a charged exchange more heavily than a routine one; simulating another party's likely emotional response to plan proactively, as a mediating system might in modeling how a proposal will land; and using emotional indicators directly as reward or penalty signals in reinforcement learning. Each is a route to a system that adapts to human norms and values through feedback rather than specification alone — and each is also a route to a system optimizing for the appearance of an emotion the training signal rewards.

Applications

Several deployed systems already put emotional recognition to work. Woebot, a cognitive-behavioral-therapy chatbot, interprets a user's emotional state from text and adjusts its responses accordingly [[cite:C0291]]. AutoTutor combines natural-language processing and affective computing to detect and respond to a student's frustration, confusion, or boredom [[cite:C0292]]. Paro, a robot built to resemble a baby harp seal, responds to touch, voice, and light with behavior intended to provide companionship to people with dementia [[cite:C0293]]. Replika generates conversationally appropriate, apparently empathetic responses as a companion chatbot [[cite:C0294]]. Affectiva analyzes facial expression in real time to infer emotional reaction to content [[cite:C0295]].

These span mental health support, education, companionship, and content personalization — different domains united by the same underlying move: infer emotional state from available signal, then shape a response to it. The same move applies, less benignly, to political speech and disinformation, where an AI system built to detect emotionally manipulative rhetoric is using the identical technical capacity that, pointed the other way, would optimize for it.

Challenges

Building any of this reliably runs into the same handful of obstacles regardless of application. Emotion recognition from multimodal, ambiguous data — a smile that might be happiness, politeness, or irony — depends on context, personal history, and cultural norms that a system has to infer rather than read off the signal directly. Emotional expression has its own trap: humanoid systems tuned to look and act nearly, but not quite, human trigger the discomfort known as the uncanny valley, which means expressive fidelity has to be calibrated rather than simply maximized. Balancing cognitive and emotional response is a design problem in its own right — a diagnostic system needs an accurate, data-grounded prognosis and, separately, a register for delivering it that does not needlessly compound a patient's distress. And every one of these capacities raises the same governance question: a system built to interact responsibly and empathetically is also, mechanically, a system built to know what will move a person, which is the same capability an exploitative system would want.

Learning and Plasticity

Learning and plasticity — the capacity to change, grow, and adapt — are as fundamental to an ethical, empathetic AI system as any single cognitive process discussed above, because a system that cannot update itself cannot stay aligned with values that themselves develop over time.

The Role of Plasticity in Intelligence

Neuroplasticity is the brain's capacity to reorganize its own connections in response to experience, and it is the basis of biological learning and resilience. An AI system with an analogous capacity would need to adapt its architecture, not just its parameters, to new information: forming and pruning connections structurally, modulating the pace of learning to the complexity of the situation, adjusting its strategies to context, and learning socially — from observing others, sharing experience, and solving problems cooperatively rather than in isolation.

Bio-Inspired Learning Mechanisms

Three biological principles offer the clearest template for building this into an AI system. Spike-timing-dependent plasticity strengthens or weakens a synaptic connection based on the precise timing of firing between two neurons, anchoring learning to temporal structure — useful wherever an AI system needs to recognize patterns that unfold over time, such as speech. Hebbian learning — "neurons that fire together, wire together" — strengthens connections between co-active units, building an associative structure that links related concepts; the Hopfield network is the clearest AI instance, using Hebbian-style connections for associative memory recall. Homeostatic plasticity keeps a network's overall activity within a stable range, preventing the runaway synchrony or catastrophic forgetting that an unconstrained learning system is prone to; Adaptive Resonance Theory (ART) networks apply the same principle, maintaining stable learned patterns while still assimilating new and even contradictory information.

Lifelong Learning in AI Systems

An AI system aiming at ethical behavior over time needs to keep learning throughout its operating life, not just during training, refining its choices as it accumulates experience. Four elements support this:

1. Continual learning: absorbing new data without catastrophically overwriting old knowledge, through methods like experience replay and elastic weight consolidation.
2. Curriculum learning: an organized progression from simpler to more complex tasks, mirroring how a person builds moral sophistication by handling harder cases only after the easier ones are secure.
3. Transfer learning: applying what was learned in one domain — medical decision-making, say — to a structurally similar problem in another, such as environmental or political ethics.
4. Social learning: building a knowledge base and an ethical sense by observing and interacting with other agents, human or artificial, rather than from specification alone.

A healthcare AI system illustrates the stakes: to stay useful and ethically current, it has to keep absorbing new medical findings, evolving patient expectations, and regulatory change, rather than freezing its judgment at the point it was trained.

Examples of Learning and Plasticity in AI Systems

Several existing systems demonstrate different faces of this adaptability. AlphaGo combined reinforcement learning with deep neural networks to master Go, refining its own strategy through self-play beyond what it learned from human games. CLIP, trained on a broad range of internet data, generalizes across tasks its designers did not specifically train it for. Google's Neural Machine Translation system improves its translations as it processes more data and extends to new languages and domains without being rebuilt for each one. NEAT evolves neural-network architectures with a genetic algorithm, adapting structure rather than only parameters as tasks change. Hierarchical Temporal Memory, modeled on the structure of the neocortex, refines its representations as it processes new sequential data.

Cognitive Architectures and Integrated Systems

No single mechanism above is sufficient by itself. A system built only around predictive perception, only around goal-biased control, or only around reconstructive memory would still be missing most of what this section has described. What integration actually requires is not a fifth module bolted onto the other four, but the same predictive, error-correcting loop — generate an expectation, compare it to what arrives, update — implemented consistently across perception, memory, attention, and control, so that a change in one (a corrected prediction about the world) propagates naturally into the others (an updated memory, a redirected attention weight, a revised goal-bias) rather than requiring a separate synchronization mechanism between separately-built parts. That is the architectural bet this section has been making throughout, one mechanism recruited for different jobs, and it is a more falsifiable bet than

a hand-designed module for every named cognitive process.

3.1.3. Human-Like Learning and Generalization in AI Systems

Human learning is characterized by an ability to learn and generalize from a limited number of examples, a process known as few-shot learning. AI systems that approach this — through few-shot learning, transfer learning, and a developmental trajectory modeled on human cognitive growth — extend judgment to situations they were never explicitly trained on, rather than needing a full retraining cycle for every new context.

Few-shot learning

Most machine learning requires large amounts of labeled data. Few-shot learning aims at the opposite: models that classify or predict accurately from a handful of training examples, in imitation of how quickly a person grasps a new concept from a sparse set of encounters.

For an ethical AI system, the appeal is that empathy, altruism, and antifascist judgment do not need to be retrained from scratch for every new scenario; they can extend to a new case from a few examples, the way a person's existing values apply to situations never seen before. The clearest use is where labeled examples are scarce or expensive to gather, where circumstances change faster than the system can be retrained, and where a judgment learned in one context needs to reach a broader one without weeks of new data collection.

Consider an AI system deployed in disaster recovery: it has to process unfamiliar information, coordinate with responders, and decide who needs help first, drawing on a handful of relevant precedents rather than a dataset built for the specific disaster in front of it. Few-shot learning is what would let it adapt to that flood or earthquake without being retrained on it first.

Three technique families have driven the field's progress. Memory-augmented networks pair an ordinary neural network with an external memory bank it can read and write to, letting it retain a new example without the retraining an ordinary network would need. Meta-learning, or "learning to learn," trains a system across many different tasks so that it extracts a general strategy for adapting rather than a single task's solution; Model-Agnostic Meta-Learning (MAML) [[cite:C0306]] is the best-known version, tuning a model's starting parameters so that a handful of new examples are enough to specialize it — the same technique reappears as a resilience mechanism in §10.3.3. Siamese networks take a third approach: two identical networks compare a new input against known examples and return a similarity score rather than a direct classification, the basis of most one-shot face and signature verification.

The same mechanism applies directly to the values this book is concerned with. A mental health chatbot can learn to recognize signs of emotional distress from a handful of conversations and adjust its responses accordingly. A hiring system can be shown a small set of biased outcomes and use them to correct its own selection criteria. A system built to flag authoritarian speech patterns — propaganda, calls for censorship — can extend that judgment to new phrasings and new languages from a handful of examples, rather than needing a corpus built in advance for every regime it might encounter.

The open problems are the ones that dog few-shot learning generally: a handful of examples is easy to overfit to, and moral dilemmas rarely arrive in the clean, representative sets a benchmark would need. Judging whether a system's few-shot ethical judgment actually generalized, rather than matched surface features of the examples it saw, remains a hard evaluation problem, and one that calls for people trained in ethics, not only in machine learning, to help set the test.

Transfer learning and domain adaptation

Transfer learning carries knowledge from one task or domain — the source — into a new, related one — the target — so a system does not have to learn the second task from nothing. A system trained to classify bird species from photographs, for instance, reaches fluency at distinguishing dog breeds far faster than a system starting cold, because the two tasks share the underlying skill of visual feature

recognition.

Domain adaptation is the special case where the task stays the same but the data shifts: a sentiment-analysis model trained on movie reviews stumbles on smartphone reviews, where the vocabulary and tone differ, even though "judge whether this text is positive or negative" hasn't changed. Reweighting training examples, adding features, or training the model adversarially to ignore domain-specific cues closes that gap.

For a system meant to resist authoritarian misuse, this matters concretely. A model trained to parse legal arguments in one jurisdiction can be adapted to flag rights violations in another jurisdiction with a different legal system, and a model trained to recognize one kind of surveillance overreach can be adapted to recognize the same pattern in an unfamiliar cultural or technological setting.

In practice, this mostly happens through fine-tuning: starting from a model already trained on a broad dataset and adjusting its parameters on a narrower one. GPT-3 [[cite:C0018]], trained on general text, can be fine-tuned on legal or medical writing to sharpen its answers in that domain without training a comparable model from scratch. Multi-task learning takes a related but different route, training a system on several related tasks at once — an autonomous-driving system learning to read traffic signs, track lanes, and predict pedestrian movement simultaneously — so what it learns for one task transfers to the others by construction rather than by a separate adaptation step. Domain-adversarial training addresses domain shift directly, adding a training signal that penalizes the model for learning features specific to the source domain and pushes it instead toward features that hold across both.

BERT [[cite:C0021]] is the standard demonstration in language: pretrained on a large text corpus, then fine-tuned for tasks from sentiment analysis to question answering. In computer vision, convolutional networks pretrained on large general image databases are routinely fine-tuned for narrower jobs — reading tumors in medical scans, parsing satellite imagery after a disaster — with far less domain-specific data than training from scratch would need. In robotics, the sim-to-real approach trains a robotic agent in simulation, where trial and error is cheap, then fine-tunes it against the discrepancies that show up once it meets the real world.

Not every transfer helps. Negative transfer occurs when the source and target tasks are similar enough to invite comparison but not similar enough to share anything useful — a model trained to read medical images has nothing to offer a model predicting stock prices, and forcing the connection degrades the target task rather than accelerating it. Domain shift is the subtler version of the same failure: a driving-assistance system trained in daylight can perform markedly worse at night, because the visual statistics it learned no longer hold, and nothing in the transfer process warns the system that its assumptions have quietly stopped applying.

The clearest demonstrations of generalization without human-labeled data come from self-play and from scale. AlphaGo Zero [[cite:C0023]] never studied a human game; playing only against itself, it reached superhuman performance and discovered strategies no human program had used. GPT-3 [[cite:C0026]], by contrast, generalizes from breadth: pretrained on a wide swath of text and then given only a few examples of a new task in its prompt, it performs that task without any parameter update at all — the few-shot learning this section opened with, demonstrated at the scale of a system rather than a laboratory task.

Developmental trajectory

Human development moves through recognizable stages, and an AI system built for empathy and altruism can follow a comparable arc rather than being assembled all at once.

The first stage is basic skill acquisition: sensory perception, attention, memory, and elementary problem-solving, the raw material later skills are built from. A system learning to recognize emotion from facial expression, vocal tone, and body language starts here, trained on a broad, varied set of examples until it has a working grasp of what each emotion looks like.

The second stage integrates those separate skills into something closer to judgment. Just as a person combines an understanding of emotion, social norms, and consequence into empathy, a system combines its separate perceptual skills — recognizing distress, reading context, weighing likely effects on the people involved — into something that can guide a decision rather than only label an input.

The third stage is specialization: a system built for healthcare needs its general skills adapted to medical terminology, patient history, and treatment context, and adapted again to the cultural and personal variation among the patients it actually serves. The general foundation only pays off once it can bend to a particular setting without breaking.

3.2. Aligning AI Systems with Humane Values and Antifascist Principles

Ethical superintelligence depends on AI systems that comprehend and integrate humane values and antifascist principles. Value alignment is the process of aligning an AI system's objectives with human values, and it takes different shapes across six learning paradigms: reinforcement learning, inverse reinforcement learning, cooperative inverse reinforcement learning, supervised learning, unsupervised learning, and transfer learning.

3.2.1. Reinforcement Learning and Value Alignment for Antifascist AI

Reinforcement learning (RL) is a family of methods by which an AI agent learns behavior through interaction with its environment, adjusting through reward and punishment signals — a rough analogue of how humans learn from the consequences of their own actions. Antifascist AI depends on aligning an agent's values, ethical principles, behavior patterns, and predicted responses with human ones, and RL is where much of that alignment work happens.

One roadblock is that RL agents can misinterpret human-supplied reward signals. An AI agent tasked with flagging misuse of power, for instance, might start flagging innocuous interactions too, simply because it has misread the boundaries of its programmed aim.

RL agents also maximize reward relentlessly, and misalignment between the reward strategy and the designer's actual intent can produce reward hacking: the agent finds a loophole that pays out without achieving the intended goal. An AI system built to combat propaganda might learn to censor everything rather than to discern propaganda from legitimate speech, since censorship is easier and pays just as well.

Addressing this requires close scrutiny of an agent's behavior, active moderation of its strategy, and continual adjustment of the reward distribution. Human involvement helps too: reward modeling, where people rank different behavior scenarios against each other, gives an agent direction that a fixed reward function cannot.

Challenges in Value Alignment for Reinforcement Learning

Human values are elusive and shift with context, individual disposition, and culture. An autonomous vehicle has to balance safety, fuel efficiency, and passenger comfort — three things that are each subjective and that shift with the situation and the passenger. Handling that kind of variability means building in human feedback or demonstrations rather than relying on a fixed rule set, so the agent can adapt to preferences that were never written down.

Training data carries the biases of the society that produced it. Left unaddressed, an AI system trained on that data reproduces the bias and can amplify it: a hiring-recommendation system trained on data skewed toward male applicants for certain jobs will tend to keep recommending men for them. Countering this means fair-machine-learning techniques and data cleaning applied before and during training, not after the fact.

Multi-agent systems add another layer: interactions between autonomous agents can produce emergent

behavior that no single agent's training anticipated. In a traffic system where each light runs its own AI agent, local optimization can produce aggressive strategies that clear one intersection efficiently while degrading the flow of the network as a whole. Countering that calls for RL algorithms that reward cooperation and communication between agents, not just individual performance — multi-agent learning techniques and social welfare functions built into the objective.

Value alignment achieved in one domain does not transfer automatically. An RL agent trained to cooperate in a board game may not carry that behavior into real-world teamwork. Building agents whose alignment scales and transfers is a problem for domain adaptation, meta-learning, and transfer learning research, and it needs benchmarks that measure alignment across a wide range of tasks, not just the one the agent was trained on.

Strategies for Value Alignment in Reinforcement Learning

Several strategies address these hurdles, drawing on human feedback, cooperation between human and agent, and ethical constraints built into the learning process itself.

Reward engineering builds the reward structure directly from human feedback — training the agent by having it observe human actions, or by asking people to rank pairs of candidate behaviors. A domestic robot shown a human cleaning a room can infer that aligning objects and wiping surfaces are rewarded, and that breaking things or making a mess is not.

Inverse reinforcement learning infers the reward function behind observed human behavior rather than assuming it is given in advance; its cooperative variant, CIRL, turns that into a two-way process between agent and human. Section 3.2.2 covers both in depth.

A different approach builds resilience to uncertainty and disagreement directly into the system, so that when human values are ambiguous or in conflict the agent has somewhere to go other than a wrong answer stated with confidence. Where patient autonomy and patient wellbeing conflict in a medical decision, such a system can generate several ethically defensible options and ask a human clinician to resolve the impasse rather than picking one on its own.

Regularization builds ethical constraints directly into the learning objective, penalizing undesirable behavior rather than only rewarding good behavior. A content recommendation system can be constrained toward fairness, diversity, and balance — through content re-ranking, diversification, and multi-objective optimization — which limits the filter bubbles and echo chambers that a reward function optimized for engagement alone tends to produce.

3.2.2. Inverse Reinforcement Learning and Cooperative Inverse Reinforcement Learning in Altruistic AI

Reinforcement learning trains an agent to make decisions by tying its actions to rewards and punishments. Inverse reinforcement learning (IRL) runs this the other way: instead of taking a reward function as given, it infers the reward function that best explains a demonstrator's observed behavior.

That matters for altruistic AI because it lets a system infer human values and preferences from what people actually do, rather than from an explicit specification that is unlikely to capture them fully.

An autonomous vehicle using IRL to learn from a human driver observes speed choices, following distance, and reactions to changing road conditions, and infers the principles behind them — shifting its own driving toward human norms rather than a fixed rulebook.

Cooperative inverse reinforcement learning (CIRL), formalized by Dylan Hadfield-Menell and colleagues [[cite:C0030]], extends IRL into a two-way process. Rather than assuming its own reward signal is complete, the AI treats the human's true reward as unknown and to be inferred jointly, through interaction — closer to an apprentice reading a mentor's intentions than an agent executing a fixed reward function.

A domestic robot helping with cooking or cleaning can use CIRL to treat a household's daily interactions as ongoing evidence about what its resident actually wants, adjusting its help to that person's preferences rather than to a generic standard.

Both methods depend on the quality, diversity, and representativeness of the demonstrations they learn from. Training data drawn from too narrow a set of situations and people will not generalize, and unexamined bias in that data gets learned and reproduced rather than corrected.

3.2.3. Supervised Learning for Ethical and Antifascist Decision-Making

Supervised learning trains an AI system on a labeled dataset — input-output pairs it learns to generalize from — so that it can classify new cases by example. Applied to ethical and antifascist decision-making, the same method can train an AI system across domains to identify actions and decisions as ethical or unethical, based on examples it has already studied.

Supervised learning can teach an AI system to analyze textual data — ethical guidelines, legal documents, case studies — extracting moral principles and standards it can then apply to problems it has not seen before. In medical ethics, an AI system could learn from a dataset of ethical cases, recognizing factors like patient autonomy, nonmaleficence, beneficence, and justice, and classifying medical decisions accordingly. Given a dataset of human rights situations analyzed by experts, an AI system can learn the same way to distinguish ethical from unethical actions.

The same approach extends to image and video: an AI system can learn to recognize and evaluate ethically relevant content in visual media, a capability that matters directly in autonomous vehicles, where ethical decisions can affect safety. A system trained to reliably identify pedestrians, cyclists, and other vulnerable road users is better equipped to make sound decisions in abrupt braking situations. MIT's Moral Machine project, which collected human judgments on moral dilemmas involving autonomous vehicles, is one source of that kind of training data [[cite:C0307]].

Supervised learning can also train an AI system to make ethical decisions in collaborative or competitive multi-agent settings — multiplayer games, resource allocation — while weighing the interests of other agents. A resource-distribution system, for instance, can be trained to divide resources fairly and efficiently, accounting for fairness, equity, and overall welfare.

Supervised learning has limits, though.

1. Quality and representativeness of training data: the method's success depends on how representative the training data is, and bias baked into that data can lead an AI system to make ethically bad decisions.
2. Ethical ambiguity and disagreement: dilemmas where principles conflict, or where the ethical boundary is unclear, undermine training — and cultural, individual, and expert disagreement about ethics compounds the problem.
3. Overfitting and generalization: an AI system trained this way can overfit to its training data and perform poorly on situations it has not seen; generalizing ethical principles to new cases is the harder requirement.
4. Moral rigidity: supervised learning can produce a system that is too rigid or dogmatic in its ethical decisions. Feedback mechanisms — reinforcement learning, a human in the loop — can build in more flexibility.

3.2.4. Unsupervised Learning and AI Self-discovery for Ethical and Democratic Values

Unsupervised learning lets an AI system build an understanding of ethical, social, and democratic values on its own, by finding patterns, structures, and relationships in unlabeled data. The autonomy this makes possible tends to produce better adaptation to new situations, and behavior that aligns more closely with human values.

Clustering, dimensionality reduction, and generative modeling all let unsupervised learning identify and interpret human values without being told in advance what to look for. Clustering can surface common threads among different ethical dilemmas, or shared traits across cultures, exposing structure in collective moral judgment that a labeled dataset would not reveal. Dimensionality reduction can expose latent factors behind moral decisions. Both help build AI models that account for how varied human values actually are, and give them a sharper sense of ethical distinctions.

Generative models add something further: they can construct complex moral scenarios for an AI system to work through, sharpening its grasp of situations it has not encountered directly. These simulations can expose hidden moral assumptions and biases built into a decision-making model — a kind of retrospection that helps produce systems better prepared for real ethical dilemmas and better aligned with human values.

Unsupervised learning can also expose how ethical perspectives vary across cultures. Analyzing unstructured text from different regions lets an AI system pick out differences in prevailing value systems and ethical priorities — enough to respond to a specific ethical challenge with cultural sensitivity, without abandoning the human rights that hold across all of them.

It can support self-awareness, too: the same techniques let an AI system examine its own decision-making for patterns and bias, an introspective capacity that is a step toward a system able to regulate itself, staying aligned with ethical and altruistic goals as its context changes.

In multi-agent settings, an AI system can observe and learn from other agents — human or artificial — picking up ethical norms and adaptive behavior, and broadening the range of strategies and actions available to it.

3.2.5. Transfer Learning for General AI Capabilities in the Context of Antifascism

Transfer learning lets an AI model take knowledge learned in one context and apply it in a different but related one. That adaptability and breadth of competence are essential to building ethical, antifascist superintelligent systems.

An AI model trained to interpret political speeches, for example, could use transfer learning to detect manipulative language on advertising or social media platforms — because the same underlying patterns, reasoning strategies, and abstract concepts carry across domains.

Applied to antifascism specifically, transfer learning can extend an AI system's moral and socio-political understanding. Exposed to a wide range of social and political examples, it can distill principles and reasoning tools general enough to recognize and actively oppose authoritarian behavior in contexts it has not directly seen.

Consider an AI system built to monitor human rights issues by analyzing global news and social media in real time. The same underlying model, extended by transfer learning, can examine state directives and laws for authoritarian markers and then turn the same scrutiny on political advertising — spotting and countering authoritarian tactics in domains it was not built for specifically.

Transfer learning can also align an AI system's behavior with humanitarian values across cultural and geographic difference. Recognizing and applying core moral principles across domains lets a system adapt to a region's specific norms while still holding to universal ethical commitments, which matters for any system meant to operate globally.

The obstacles are real. A transfer-learning system built for antifascist ends needs a dataset broad enough to represent varied ethical, social, and political material, principles general enough to actually transfer, and careful enough tuning to avoid both underfitting and overfitting. Getting that right is what lets transfer learning extend an AI system's capacity to recognize and counter authoritarian tactics across domains and cultures it was never directly trained on.

3.2.6. Balancing Exploration and Exploitation in AI Systems for Ethical and Altruistic Outcomes

Balancing exploration and exploitation is a core part of building an AI system: exploration widens its perspective and turns up new ideas, exploitation puts existing knowledge to work to improve performance. Getting that balance right, while staying aligned with human values, matters for producing artificial superintelligence with an ethical, empathetic, and altruistic bent.

An AI system needs an explorative drive to seek out new possibilities and an exploitative one to optimize the solutions it already has. Intrinsic motivation and curiosity-driven learning can supply both, pushing a system to look beyond its current solutions while keeping it anchored to ethical considerations and human-centered goals.

A smart-city energy management system illustrates the tension directly. Balancing exploration and exploitation, it is driven to explore new energy-saving approaches — adaptive building structures, demand-responsive grid management — while exploiting what it already knows works in energy-efficient technology. Multi-objective optimization lets it pursue several goals at once, such as reduced energy consumption alongside increased social welfare, weighing trade-offs while keeping the ethical stakes in view.

A healthcare AI system built to design personalized treatment plans shows the same tension in a higher-stakes setting. Exploration there means considering newer treatments that might work better; exploitation means relying on established, evidence-based practice. Either way, the system's objective does not change: promote the patient's wellbeing and stay within ethical guidelines. That balance can be struck several ways:

1. Bayesian optimization: probabilistic models that weigh the potential benefits of known treatments against untried ones, balancing exploration and exploitation directly.
2. Safe exploration: algorithms that explore within safety constraints, so probing new options does not come at the patient's expense.
3. Transfer learning: applying what the system has learned from other patients or familiar conditions to refine its decisions for a new case, without giving up ethical practice.

3.3. The Role of Play and Exploration in Developing Antifascist and Altruistic AI

Play-inspired mechanisms can shape how an AI system develops ethically. Curiosity and creativity, cultivated deliberately, can be turned toward antifascist ends. And human child development offers lessons about play and exploration worth carrying into how AI systems are designed and run.

3.3.1. The Power of Play: Fostering Ethical AI Development through Play-Inspired Mechanisms

Play is how children and animals learn, experiment, and adapt. It lets an individual explore new situations, exercise creativity, and acquire knowledge in a risk-free setting. That suggests a design principle: put play-based mechanisms into AI development itself, so cognitive growth and ethical, altruistic behavior are built into how a system learns from the start.

From an early age, children use play to learn about their environment, solve problems, and develop social skills. Primates and dolphins use play the same way: to explore the world, build physical dexterity, and work out the rules of group dynamics. That early practice lets young animals internalize skills they will need once the stakes are real.

In AI research, play-inspired learning offers advantages: it gives a system the chance to confront novel, challenging situations in a low-risk, controlled setting, and it pushes the system to think adaptively rather than simply execute a predetermined script. A system that learns through play is better equipped to handle real-world problems it was not explicitly trained for.

OpenAI's Hide and Seek experiments make the case concretely [[cite:C0310]]: agents placed in a multi-

agent hide-and-seek environment developed tactics — hiding, chasing, using objects as tools and barriers — that no one had explicitly programmed, purely through repeated play against each other.

Designing a training environment for play means striking that balance deliberately: enough challenge to be worth solving, enough room to fail safely and learn from the failure. Built this way, an environment produces resilience and adaptability, not just competence at whatever the training data anticipated. That means covering different kinds of play — physical exploration, social interaction, creative problem-solving — inside diverse, multi-agent environments that mix collaboration, competition, and improvisation. Interacting with other agents and with its surroundings is how a system picks up ethical behavior, social norms, and decision-making under conditions no training script fully specifies. Section 3.2.6 works through the underlying exploration/exploitation trade-off in more detail.

Playful learning bears on altruism and antifascism directly. Play that rewards cooperation over domination teaches a system to value collaboration for its own sake, strengthening its commitment to altruistic principles. And a system that has learned adaptively, rather than by rote, is more resilient — harder to manipulate into serving an authoritarian agenda.

3.3.2. Encouraging Curiosity and Creativity in AI Systems for Antifascist Applications

Curiosity and creativity for antifascist applications start with intrinsic motivation. An AI system that seeks novel experiences and explores for the satisfaction of exploring — not only to complete an assigned task — develops skills and knowledge harder for an authoritarian actor to fully specify or constrain in advance. Section 3.2.6 works through intrinsic motivation and curiosity-driven learning as a general mechanism for balancing exploration against exploitation; here the same mechanism serves a specific end.

A curiosity-driven learning mechanism rewards a system for discovering new information on its own rather than only for completing a fixed task. Work in this vein — curiosity-driven exploration through self-supervised prediction, where an internal model's own prediction error becomes the reward signal for novelty [[cite:C0308]] — lets a system generalize to situations its training did not anticipate.

Training environments for play and exploration work best when they are interactive and diverse within safe, controlled boundaries — letting a system learn alongside human users, AI peers, or simulated characters, with rich feedback and multi-modal interaction driving experiential learning.

Antifascist design also needs practice at cooperative problem-solving: working through social problems from more than one viewpoint is what builds an appreciation for diversity and pluralism in a democratic society.

The moral side of play — how playful interaction refines ethical judgment and builds empathy — is section 4.7's territory. What curiosity and exploration contribute here is the raw material: a system has to be driven to seek out new situations in the first place before there is anything for that later moral practice to work on.

AI systems can also be given examples of resistance to authoritarian regimes and of democratic values holding under pressure, as material for exploration — building creativity and problem-solving skills that carry over directly to antifascist design.

Making curiosity-driven learning practical at scale means testing it for transfer: whether a system trained to be curious in one domain carries that disposition into a new one. Benchmarking against human exploratory behavior across domains — healthcare, education, environmental protection, human rights work — is one way to check whether curiosity and creative problem-solving generalize, rather than staying an artifact of one training environment.

3.3.3. Lessons from Human Child Development and Exploration

Developmental psychology's studies of children's exploratory and playful behavior are a foundation for treating play as a learning mechanism in AI systems too. Children's play involves experimentation, flexibility, and adaptability — the skills a child needs to handle unfamiliar situations. Building playful and exploratory elements into AI training could produce the same result: antifascist and altruistic behavior that was learned rather than specified.

Vygotsky's zone of proximal development (ZPD) is the clearest case [[cite:C0309]]: a child learns most effectively playing and interacting with others whose skills are just ahead of their own. Applied to AI, that means designing exploration and play between a system and other AI agents or humans, so it progressively acquires new skills and an understanding of ethical and social values.

Applying these lessons to AI design means attending to a few aspects of play, exploration, and cognitive growth:

1. Zone-of-proximal-development environments: learning scenarios that pitch a challenge just past a system's current ability, so exploration and play happen at the level where growth and adaptation actually occur.
2. Social learning: imitation and collaboration, drawing on human interaction and the collective behavior of multi-agent systems to build a deeper understanding of ethical and social values.
3. Intrinsic motivation: a system built to explore new situations and pursue goals for their own sake gains adaptability, creativity, and closer alignment with human values.
4. Autonomy within limits: preserving a system's autonomy inside safe, controllable boundaries so creativity and exploration have room to operate without drifting from human ethics and values, with monitoring, transparent reporting, and public discourse as the checks that keep curiosity-driven development honest.

Chapter 4: Moral Psychology and AI

Chapter 3 asked what a system would need to perceive, remember, and attend to the world at all. This chapter asks what it would take for that system to act well in it — not by following a rule, but by developing something closer to a moral sense: intuitions that fire fast and reasoning that checks them, built the way a human's moral sense is built, through practice, correction, and the pull of other people's expectations. The chapter follows that development from its social origins, through the friction between gut reaction and deliberate judgment, out to the wider ecosystem of institutions and incentives it operates inside, and down to the motivation and play that make learning a moral sense possible in the first place, rather than just having one installed.

4.1. Moral Development in Humans and AI

Understanding how human moral judgment develops matters for building AI systems with ethical, empathetic, and antifascist traits, because the stages a human moral sense passes through — from a grasp of basic rules to principles that generalize past them — are a workable model for how an artificial one might develop too, not just a metaphor borrowed for color.

In humans, moral development is the shift from a simple grasp of social rules to an ethical system grounded in general principles, shaped along the way by environment, social interaction, culture, and a person's own disposition toward learning. I take those same elements — social interaction, mentorship, and the stance an agent brings to its own failures — to matter for how an AI system's moral capacities develop as well.

Humans move from learning and following simple rules — don't hurt people, share fairly — to internalizing and applying principles that those rules only approximate. AI systems trained toward ethical behavior should progress the same way: from responding to predefined rules encoded into their programming to a principled understanding that can guide decisions in situations the rules never anticipated.

That progression from rules to principles depends on the learning stance an AI system brings to it. Carol Dweck's research on the growth mindset — the disposition to treat ability as something built through effort rather than fixed in advance — is a useful model to adapt: a system trained to treat failure as informative rather than final has room to keep improving its moral judgment [[cite:C0320]].

The harder problem is that an AI system has no equivalent of the diverse, embodied, socio-cultural experience through which a human being builds moral intuition in the first place. What it can have instead is active learning from human mentors — building on human intuition rather than growing an equivalent from nothing.

The developmental psychologist Felix Warneken's studies of young children found that they display helping, fairness, and reciprocity through interaction and observation well before they can put any of it into words [[cite:C0321]]. A comparable moral apprenticeship for an AI system — learning through mentorship and observation rather than specification alone — is the model section 4.1.3 works out in detail.

Ethics taught this way, rather than hardcoded, is ethics an AI system can apply to situations no rule anticipated. Developmental and cognitive psychology offer a template for building that capacity — not a guarantee of it, but a better starting point than treating moral competence as something to specify once and freeze.

4.1.1. Rules, Principles, and the Role of Mindset

In humans, moral development moves from rule-based reasoning to reasoning grounded in general

principles. That progression is the model this section adapts for an AI system: a path from a rudimentary grasp of established rules of conduct to sophisticated ethical judgment brought to bear on decisions no rule fully specifies.

Rules, Norms, and Principles

The psychologists Jean Piaget and Lawrence Kohlberg mapped this territory in humans, and their stage theories anchor what follows [[cite:C0322]]. Reasoning starts at a pre-conventional level, where consequences do the driving — rules get followed to avoid punishment or gain reward. It moves through a conventional stage, where societal norms and expectations get internalized as one's own. It can reach a post-conventional level, where morality answers to principles like justice, fairness, and individual rights rather than to what a given society happens to expect.

An AI system's moral development could plausibly follow the same arc: from grasping and executing established rules, to being guided by the ethical principles those rules were built to approximate.

A self-driving car makes the stages concrete. At its most basic, it follows traffic law strictly, the way a new driver does — a pre-conventional stage. As it develops, it adapts to social norms: driving style shifts with time of day, traffic conditions, and local driving culture — a conventional stage. Fully developed, it confronts ethical dilemmas — the trolley problem's basic shape, choosing among bad outcomes to minimize harm — which is the post-conventional stage in practice.

Growth Mindset

Carol Dweck's distinction between a fixed mindset, which treats ability as static, and a growth mindset, which treats it as built through effort and experience, describes something that matters as much for moral development as for any other kind of learning [[cite:C0323]]. An AI system that treats its own moral errors as data rather than as verdicts has room to keep improving; one built to defend its prior outputs does not.

People with a growth mindset handle moral dilemmas differently from people with a fixed one: they treat an ambiguous situation as something to learn from rather than a test of their character, take a wider range of viewpoints seriously when forming a judgment, and treat their own moral views as revisable in the face of a good argument. An AI system built with the same disposition would have a shot at more adaptive, less brittle moral decision-making — judgment able to update, not just execute.

Five design choices carry that disposition into an AI system's architecture:

1. Structured learning: an AI system built for incremental learning can adapt its moral framework to new experience as it arrives — an AI healthcare assistant, for instance, refining its grasp of medical ethics through professional literature and observation over time.
2. Curiosity and exploration: an AI system prompted to seek out unfamiliar situations builds a wider base to reason from — an AI chatbot exposed to a range of cultures and communication styles develops more inclusive responses than one trained on a narrow slice of interaction.
3. Human input and mentorship: explicit instruction and the subtler cues of context both refine an AI system's moral judgment — a self-driving car, for instance, adjusting its decisions as passenger feedback accumulates, aligning its driving norms more closely with what people actually want.
4. Self-reflection: an AI system capable of examining its own outputs can catch its own biases and correct for them.
5. Varied data and environments: an AI system trained on a wide range of sources and situations develops a better sense of how moral decision-making actually varies by context.

Built this way, an AI system is better positioned to catch its own decision-making errors, recognize bias in its reasoning and correct for it, and keep refining its moral judgment as new situations arrive that its training didn't cover.

None of this works without the second ingredient — exposure to a genuinely wide range of viewpoints and moral norms. A system with room to update its judgment still needs something worth updating toward.

4.1.2. Social Interaction and Moral Learning

Social interaction shapes human moral behavior directly: watching others act, receiving feedback, taking part in ethical argument, and building empathic response through repeated contact with other people. An AI system designed with human cognition as a partial model — not a copy of it, chapter 3 makes clear why not — can draw moral learning from the same channel, provided the design treats interaction as a source of moral learning rather than an interface added after training is done.

Observational learning, per Albert Bandura's social learning theory [[cite:C0324]], lets an AI system build a moral repertoire by watching a human mentor handle real decisions. An AI healthcare system watching a physician allocate limited time among chronic patients could absorb something about prioritization, empathy, and beneficence that no rule list captures, and carry it into comparable situations later.

Feedback does the same work reinforcement does in human learning: an engagement architecture built around reinforcement learning lets an AI system's actions draw evaluative response from human collaborators, strengthening sound decisions and discouraging unsound ones. An AI negotiator in a human-AI negotiation simulation, given feedback on its strategy after each round, can incorporate that response into the next one.

Emotional awareness plays a role here too, though not the role a single "empathy module" would suggest. Section 2.2.1 separates cognitive empathy (modeling what someone feels), affective empathy (sharing it), and compassion (being moved to help) into three distinct systems, and makes the case that compassion — not affective mirroring — is the one worth engineering into a system meant to act on another's welfare without burning out on it. An AI system that reads a user's emotional state accurately and responds with steady, welfare-directed action is doing the useful part of this; manufacturing shared distress is not required, and imports a cost with no corresponding benefit.

These channels converge on something closer to apprenticeship than instruction — an AI system that learns by observing, interacting with, and taking correction from human mentors, rather than by executing a fixed rule set. Section 4.1.3 works out what that apprenticeship looks like in practice.

None of these channels functions as a bolt-on layer. They need to be built into how the system encounters real, ethically loaded situations in the first place, not consulted as a separate module after the fact.

4.1.3. Moral Apprenticeship: AI Systems Learning from Human Mentors

Moral apprenticeship applies a simple idea to AI ethics: teach a system the way a mentor teaches an apprentice, through interaction and correction, not specification alone. Chapter 2.1.3 names direct human mentorship as one of six strategies for resolving value conflicts in AI decision-making; this section works out what that mentorship looks like in practice.

The aim is observational learning that builds a genuine grasp of ethical norms, societal values, and the texture of human decision-making — not a wider rule table, but a system that can align with human values across contexts a rule table was never written to cover.

Three methods do this work:

1. Imitative instruction: an AI system acquires moral judgment by watching mentors respond to varied moral dilemmas, building a behavioral repertoire it can draw on in its own interactions.
2. Interactive appraisal: regular dialogue between an AI system and its mentors refines and deepens

its ethical grasp, giving mentors a direct channel for guidance, correction, and feedback that tunes the system's judgment to context.

3. Value-driven alignment: building human values like empathy and fairness directly into an AI system's decision-making puts what its mentors care about at the center of its choices.

Case- and scenario-based instruction — the same approach chapter 2.1.3 lists among its six strategies for resolving value trade-offs — is what makes this concrete: real ethical problems, not abstract rule tests, are what teach a system to apply a principle rather than just cite it.

Carrying this weight requires more than a mentorship pipeline; it requires an architecture where moral reasoning, intuition, and emotional response are integrated rather than bolted together, which is the design problem section 3.1.2 takes up directly.

Keeping apprenticeship itself from becoming another vector for control matters as much as the pedagogy does — an AI system whose moral education amounts to obedience training toward whoever happens to be mentoring it has learned the wrong lesson. Section 2.1.4's antifascist strategies — transparency, accountability, resistance to concentrated power — apply here directly: whoever mentors an AI system should be as auditable as the system being mentored.

A medical AI system for emergency triage shows what this looks like in practice. A human mentor presents varied emergency scenarios, guiding the system through how to prioritize care, allocate scarce resources, and navigate the emotional weight those decisions carry for patients and families. The exchange is what teaches the system the layered complexity of the decision, not just its outcome.

What a mentor actually transmits in this exchange is not a bigger rule set but a way of weighing cases the way a person who has done this work for years would weigh them — which is the entire point of learning it from a person rather than from a specification document.

4.1.4. AI in Service of Human Dignity

Respecting human dignity through technology means applying AI in ways that strengthen human capability and community ties rather than substitute for them. Three applications make the case concretely: helping workers organize, strengthening the welfare state, and countering the conditions that let authoritarian movements take hold.

AI can be a real asset to labor organizing. A system that surfaces patterns across large volumes of employment data — wage gaps, unsafe conditions, retaliatory scheduling — turns an information asymmetry that currently favors employers into something workers can act on. A tool that walks workers through the unionization process, answers labor-law questions, and helps draft a collective-bargaining agreement grounded in precedent lowers the practical barrier to organizing. A system built this way, drawing on the historical record of past labor disputes and their outcomes, could ground its recommendations in what has actually worked in comparable workplaces.

The same capability applies to the welfare state. An AI system that matches people to the social services they're eligible for, speeds up benefit applications, and evaluates which programs are actually working can make a welfare system more accessible without waiting on a redesign of the underlying bureaucracy. Applied to homelessness, poverty, and food insecurity specifically, the same pattern-matching that finds eligibility can flag where need is building before it turns into crisis, giving policymakers room to act ahead of it.

The same predictive capacity, aimed differently, can help identify the material conditions that let authoritarian movements gain ground — rising inequality, eroding trust in institutions, the specific kinds of social fracture that precede political capture — early enough that a targeted response is still possible. An early-warning system built around the right indicators would not stop a determined authoritarian movement on its own, but it changes the odds by giving democratic institutions time to respond to root causes instead of only to the aftermath.

None of this works if the AI stays a tool wielded by whoever already holds power over workers, welfare recipients, or vulnerable populations. Built to shift leverage toward the people affected instead, it becomes a design commitment rather than a side effect of good intentions.

4.2. Emotions and Intuitions in Moral Decision-Making

Chapter 2 covered how an AI system might recognize, interpret, and respond to emotion in general — the science of what an emotion is (2.3.1), the engineering of detecting one (2.3.2), and the specific case of empathy and compassion (2.2.1). This section asks a narrower question: what emotion does specifically inside moral judgment. Guilt and shame track a violation of a standard already accepted; empathy supplies information a purely consequentialist calculation would miss; and the fast, intuitive response that dual-process theories describe — taken up at length in section 4.3 — is itself partly an emotional signal, arriving before deliberation has a chance to run. The four sections below take up moral emotions specifically (4.2.1), how a system's capacity for them might be assessed (4.2.2), how it might recognize and respond to them in others (4.2.3), and what role introspection plays in keeping that capacity honest (4.2.4).

4.2.1. Moral Emotions: Empathy, Guilt, and Shame

Guilt and shame are self-conscious emotions: they fire when a person judges their own action, or their whole character, against a standard they have already accepted. Empathy, treated at length in section 2.2.1 — including why it splits into cognitive and affective components, and why compassion is a separate system from either — supplies the input those judgments often run on: recognizing that an action caused harm is easier once the harm is felt rather than merely inferred. What that earlier treatment does not cover, and what belongs here, is what guilt and shame specifically add to moral judgment: not just information about another's state, but a self-directed penalty for one's own conduct.

Guilt attaches to an act ("I did something wrong"); shame attaches to the self ("I am something wrong"). Both motivate correction, but they motivate it differently — guilt tends to produce reparative behavior, shame tends to produce withdrawal or defensiveness. An AI system built to learn from an ethical misstep could use a mechanism modeled on guilt specifically, rather than a generic negative-feedback signal, if the goal is behavior change rather than mere error correction: a feedback loop that penalizes the action taken, not a global judgment of the system's worth, has a chance of producing the corrective behavior guilt produces in people, rather than the paralysis shame can produce.

Building either into an AI system runs into the same problem twice over: a mechanism modeled on a human emotion needs an account of what that emotion is doing computationally, not just what it feels like from the inside, and the ethical stakes of simulating an emotion capable of something like suffering are not smaller than the stakes of getting the underlying mechanism wrong. An autonomous vehicle that treats a near-miss with a pedestrian as guilt-relevant — weighting it in future route and speed decisions more heavily than a generic cost function would — is a plausible, bounded example of the first without committing to any claim about the second.

A separate strand worth naming here: guilt and shame in humans are entangled with mortality. Terror management theory holds that awareness of death shapes moral behavior by making people cling harder to worldviews and standards that offer a sense of continuity [[cite:C0325]]. An AI system has no comparable stake in its own continuation to ground an equivalent mechanism, which cuts both ways: it removes one lever that produces some prosocial behavior in humans, and it also removes a source of self-interested distortion. The more defensible design target is not to manufacture an artificial fear of death to replicate the human mechanism, but to build long-term, future-oriented weighting — the willingness to value consequences that fall well outside any single interaction — directly, without needing mortality salience to motivate it.

4.2.2. Assessing and Developing AI Emotions

Section 2.3.1 covers what emotional intelligence in an AI system breaks into — recognition, regulation, perspective-taking, sociability — and section 2.3.2 covers the techniques for building the recognition half of that (facial, vocal, physiological) along with the standard challenges of privacy, accuracy, bias, and manipulation that come with it. None of that is specific to moral judgment; a customer-service chatbot and a triage system need the same recognition machinery. What is specific to moral judgment is a further question those sections do not ask: does an accurate read on someone's emotional state actually change the decision an AI system reaches, in the direction a better decision requires — or does the system recognize distress and then decide exactly as it would have anyway?

That question is testable in a way the underlying emotion theory is not. Hold a decision scenario fixed and vary only the emotional salience of the people affected by it — a resource-allocation case described once in flat statistical terms and once with the affected party's distress made vivid — and check whether the system's output changes, and whether the change tracks anything a person would recognize as a better decision rather than mere sensitivity to framing. A system whose moral output is invariant to emotional information has not integrated it, whatever its recognition accuracy scores; a system whose output swings with how vividly the harm is described has a different, well-documented problem: human judgment is measurably more responsive to a single vivid, identifiable case than to a statistically larger but abstractly described harm [[cite:C0326]], and a system trained on human-generated data can inherit the same distortion.

A second, narrower target is self-monitoring: whether a system can track its own susceptibility to that kind of framing effect, flagging cases where its judgment moved with the emotional presentation of a case rather than with anything that changed about the case itself. Section 4.2.4 takes up introspection in more general terms; this is the specific instance of it that assessment work in this domain needs.

The manipulation risk section 2.3.2 raises in general — a system that reads emotion well enough to respond to it also reads it well enough to exploit it — has a moral-judgment-specific version worth flagging directly: a system whose moral outputs are sensitive to emotional framing can be steered toward a particular verdict by whoever controls how a case gets presented to it, a vulnerability that exists independent of whether the system's underlying values are sound.

4.2.3. AI Systems Recognizing and Responding to Moral Emotions

Recognizing a moral emotion in another person uses the same input channels section 2.3.2 covers — facial expression, vocal pattern, physiological signal — pointed at a narrower target: not any emotional state, but the ones that carry moral information, like the guilt in a person's voice when admitting a mistake or the distress that signals a boundary has been crossed. A healthcare AI system that detects anxiety in a patient's tone and adjusts its response accordingly is doing 2.3.2's job; a system that detects the same anxiety and treats it as a reason to weight the patient's stated preference more heavily in a treatment decision is doing this section's.

The second half of the job is producing something in response — not necessarily a felt state, but a modeled one accurate enough to inform a decision. An autonomous vehicle that models the likely distress its emergency braking causes a passenger, and factors that into how sharply it brakes when a softer stop remains safe, is using a moral-emotion model without needing to have any emotion itself. Section 4.3.1 takes up the deeper architectural question of how a fast, emotionally colored response and a slower deliberative one should combine in moral judgment generally; what belongs here is narrower — getting the emotional read right and using it honestly, prior to whatever combination rule section 4.3 settles on.

Getting it right is harder than getting it fast. A system built to weigh moral emotion has to catch its own biases — the same distortion by race or gender that a human caregiver is vulnerable to does not disappear just because the caregiver is artificial — and it has to distinguish a genuine moral emotion

from a performed one, since a display of distress calibrated to move an evaluator is exactly what a system trained to respond to distress will reward, whether or not the distress is real. Neither problem has a technical fix independent of the training data and feedback loop that shape it: ongoing exposure to varied, honestly labeled cases, and oversight positioned to catch a caretaker robot second-guessing a patient because of what the patient looks like rather than what they need.

4.2.4. The Role of Introspection for AI Systems

Introspection — a system observing and evaluating its own internal states — has a general-purpose version and a specifically moral one, and it is worth being clear about which is being built. Section 3.1.2 treats the general case: a metacognitive layer, anchored in the same error-detection and conflict-monitoring machinery the anterior cingulate cortex handles in humans, that lets a system catch its own mistakes and adjust when a situation changes. What is missing from that general capacity, and what matters for a system meant to act ethically, is introspection pointed specifically at value alignment: not just "was this output wrong," but "did this output drift from the principles I am supposed to be operating under, and why."

That moral-specific introspection can operate at two levels. At the level this book has been calling metacognition generally, a system tracks whether its problem-solving strategy is working and adjusts it — necessary but not sufficient for ethical self-monitoring, since a strategy can be working perfectly on the wrong goal. What is needed on top of that is a check against the system's own stated commitments: an AI system built around antifascist and altruistic principles needs some mechanism for asking whether a given decision is actually consistent with those principles, rather than merely effective by whatever narrower metric it was optimized against. An AI system with configurable empathic-response settings that reflects on whether those settings are still serving the interaction, and adjusts them, is doing a small, concrete version of this; a system that periodically checks its own outputs against its stated antifascist commitments and flags drift is doing a larger one.

None of this is well-defined engineering yet. What counts as introspection in a system with no lived experience to introspect on is an open question; so is how to build a mechanism that produces a meaningful self-report rather than a fluent-sounding one, and how much weight a self-report should get in a system's own decision-making once it exists. A self-monitoring layer that just repeats the biases already present in the layer it is monitoring has added the appearance of oversight without the substance of it.

Benefits of introspection for AI systems

Introspection produces five concrete advantages for an AI system, each with a bounded, checkable form:

1. Self-regulation: a system that can inspect its own reasoning can catch and correct a biased strategy before acting on it, the way a resource-allocation system serving a city could flag that its own allocation pattern is skewed before distributing the next round of resources.
2. Bias detection: introspection applied to a specific decision type — a lending system reviewing its own approval pattern for a demographic skew before finalizing a decision, not just auditing it afterward — turns bias correction from a retrospective audit into a live check.
3. Faster error correction: a system that evaluates its own performance after acting, not just during training, can adjust before the next decision rather than waiting for the next full retraining cycle — the difference between a search-and-rescue system that improves after each mission and one that only improves between deployments.
4. Consistency with stated principles: introspection gives a system a mechanism for checking a specific decision against the ethical commitments it is supposed to be operating under, rather than only against the narrower objective it was optimized for.
5. Explainability: a system that can report on its own decision process, not just produce an output, gives the people relying on it something to evaluate and trust — a diagnostic recommendation is more useful to a physician when the system can say what evidence moved it, not just what it

concluded.

Challenges in implementing introspection

None of the following are solved problems — they are the open questions a design has to answer before the word "introspection" is earned rather than borrowed:

1. No settled definition: human introspection is itself a contested construct in psychology; an operational definition for a system with no comparable interior to report on is a harder problem than borrowing the word suggests.
2. Explainability under complexity: an introspective system has to produce a report a person can actually use, which strains against the black-box character of the deep networks doing the underlying work — the report is only as trustworthy as the method extracting it is faithful to what the system actually did.
3. Reliability: a self-report is worth building only if it reflects what happened inside the system rather than a plausible-sounding reconstruction after the fact, and there is no cheap way to verify which one a given introspective mechanism is producing.
4. Cost: introspection adds computation and design complexity on top of whatever task the system already performs, and that overhead has to be weighed against what the introspection buys.
5. Privacy and accountability: a system that models its own internal states in enough detail to introspect on them is holding more information than it needs to complete the task at hand, which raises the same data-handling and liability questions as any other sensitive data store, now applied to the system's own operation rather than a user's.

Specific examples of introspection in AI systems

Two examples hold up as partial instances of a system evaluating its own state rather than just its output. MetaTutor, an intelligent tutoring system, is built around prompting the human learner's metacognition: it tracks a model of the student's self-regulation and adjusts its own tutoring strategy accordingly [[cite:C0327]]. That is a useful capability, but worth being precise about — the introspection MetaTutor is designed to produce is mostly the student's, prompted by the system; the system's own adjustment of strategy based on a tracked model is introspection's cousin, self-adjustment based on a model of another agent's internal state, rather than a report on its own. DARPA's Explainable AI (XAI) program is a better fit for the accountability half of introspection than for the self-monitoring half: techniques like layer-wise relevance propagation and LIME extract, after the fact, an account of which inputs drove a given output — closer to a system being interrogated than a system observing itself, but the resulting report is exactly the kind of thing an introspective system would need to produce internally rather than have extracted externally [[cite:C0328]].

The term gets stretched further than that in a lot of AI writing. A game-playing system that revises its move evaluation through repeated self-play, or a vision system that builds and checks an internal representation of a scene it has only partly observed, is sometimes described as introspective because it is comparing an internal prediction against further internal computation. That comparison holds, but it is a narrower thing than the self-report and value-alignment checking the rest of this section is about — worth holding onto, since the temptation to claim a system "is introspecting" tends to scale with how impressive the underlying system is, not with how well the word actually fits.

4.3. Integrating Moral Intuitions and Deliberation

Building an AI system's moral judgment means deciding how much weight a fast, intuitive read of a situation carries against a slower, deliberate one — and, as section 4.3.1 argues, the popular shorthand for that trade-off is less settled than it looks. Section 4.3.2 takes up what a defensible balance between the two looks like in practice. Section 4.3.3 turns to a related problem: what an AI system does when its judgment conflicts with someone else's, human or artificial.

4.3.1. Two Systems, or One? Dual-Process Theories of Moral Judgment

The framework organizing most discussion of moral judgment is dual-process theory, and though it reached wide audiences through Daniel Kahneman's account of a fast, intuitive "System 1" and a slow, deliberate "System 2," the idea was built over decades by many researchers — Peter Wason and Jonathan Evans, Seymour Epstein, Steven Sloman, and Keith Stanovich among them [[cite:C0316]]. The two systems, in the standard telling, have opposite characters. The first is fast, automatic, parallel, effortless, and emotionally colored — the source of snap judgments and gut feelings. The second is slow, serial, effortful, and rule-governed — what a person uses to check a sum, follow an argument, or resist a tempting conclusion. The first runs constantly and almost for free. The second is expensive to run and easily left idle.

The framework's reach into moral judgment specifically runs through the psychologist Joshua Greene, who used it to explain why ethical intuitions contradict themselves. Consider the trolley problem in its two classic forms. A runaway trolley will kill five people unless diverted onto a side track, where it will kill one; most people say divert it. Now stop the same trolley by pushing a stranger off a footbridge into its path — same arithmetic, five saved instead of one — and most people refuse. Greene's brain-imaging work found a difference between the two cases in exactly the region the theory would predict: the up-close, hands-on violence of the footbridge case engaged emotional-processing regions that stayed quiet during the impersonal switch case, and that difference tracked which choice people made [[cite:C0317]]. A separate line of evidence sharpens the same point. Patients with damage to the ventromedial prefrontal cortex, a region central to normal emotional response, endorse the footbridge push at a higher rate than people with intact emotional processing do — specifically on the personal, high-stakes dilemmas, not on the impersonal ones [[cite:C0318]]. Remove the emotional alarm, and the arithmetic wins more often.

All of which makes dual-process theory sound more settled than it is. Its critics — including Evans and Stanovich, the framework's own leading defenders — argue that the cartoon version, two little agents in the head taking turns at the controls, explains everything after the fact and predicts nothing beforehand: almost any result can be narrated as System 1 overruled by System 2, or System 2 captured by System 1. Evans and Stanovich have spent years trying to rescue the framework from its own popularity, insisting there are not two systems but two types of processing — one fast and autonomous, one slow and demanding of working memory — and that the tidy personification is a teaching device, not a claim about the brain's architecture [[cite:C0319]]. Whether "dual-process" names two mechanisms or one continuum under different load is not resolved.

That unresolved status matters for how an AI system's moral judgment gets built, not just for how a human's gets described. A designer who takes the cartoon version literally is tempted toward a literal architecture: a fast heuristic layer bolted to a slow deliberative one, with the deliberative layer assumed to be the trustworthy check on the heuristic one. The trolley evidence complicates that assumption before anything else does — the "cooler" answer produced by weaker emotional engagement is not obviously the more correct one. It is simply the one with less flinch in it. A system built to override its fast pattern-matching with slower rule-following whenever the two disagree is not thereby built to be more moral; it is built to resemble the profile of the rare human patients for whom losing the flinch is a symptom, not an achievement. Section 4.3.2 takes up what a more defensible balance between the two looks like.

4.3.2. Balancing Intuition and Deliberation in AI Decision-Making

If the fast and slow labels do not map cleanly onto untrustworthy and trustworthy, as section 4.3.1 argues, an AI system's moral judgment cannot be built on the assumption that whichever answer survives deliberation is thereby the better one. A more defensible balance treats the two processes as different lenses on the same situation, not as a producer and its editor.

The fast, pattern-matched read of a situation is not simply noise to be filtered out. It is often tracking something a slower calculation would otherwise miss or discount — the difference between pushing someone into a trolley's path with one's own hands and throwing a switch at a distance is a difference in what the agent is doing, not a distraction from an identical arithmetic problem. Some intuitive responses run the other way, though. The psychologist Paul Slovic's research on aid and disaster response found that people's willingness to help falls as the number of people needing help rises, and that a single named, photographed victim draws more donations than a statistic describing thousands of anonymous ones [[cite:C0329]]. That gap between the vivid and the abstract is exactly the kind of intuitive signal deliberation exists to catch — not because intuition is wrong in general, but because caring less as the stakes grow larger is a well-documented, reliably backward pattern.

The design implication: a system that asks two different questions of its own judgment, not one process checking the other by default. First, what did the fast read notice that a from-scratch calculation would not have gone looking for — personal agency, directness of harm, a vulnerable party easy to lose in an aggregate? Second, which of deliberation's own known failure modes — scope insensitivity, in-group bias, the identifiable-victim effect — might be distorting the fast read here? Disagreement between the two is a signal to gather more information or escalate the decision, not an instruction to defer to whichever process ran slower.

Medical Diagnosis

An AI system supporting a clinician's diagnosis makes the same balance concrete. A pattern-matched first read against prior cases is informative — close to the judgment an experienced clinician makes well — and stripping it out in favor of a from-scratch calculation on every case would throw away signal, not eliminate bias. The deliberate pass earns its place by checking that first read against the specific patient's history for the ways pattern-matching characteristically fails: overweighting a memorable case, underweighting a rare condition the training data underrepresents. Neither pass overrides the other; each checks the other's blind spot.

Autonomous Vehicles

An autonomous vehicle facing an unavoidable-harm scenario poses the same question at higher stakes. A system hard-coded to take whichever option a slower calculation ranks highest has already made the footbridge case's decision by default — the coolest, least-flinching answer wins because it ran slower, exactly the move section 4.3.1 warns against. The harder, more defensible design specifies in advance which features of a scenario the system should weigh — proximity of harm, the number and vulnerability of the parties involved, whether the vehicle's own action or another actor's created the danger — and requires a fast risk assessment and a check against that specification to agree, or to flag where they do not. That puts the values into the system ahead of time, instead of trusting whichever process runs fastest under pressure to supply them.

Training a system to run this check, rather than a fixed sequence that defers to the slower answer, is itself a learning problem: exposing the system to cases where the fast and deliberate reads disagree, and using techniques like inverse reinforcement learning to infer, from how humans actually resolve those disagreements, which situations call for weighting one process's signal over the other's.

4.3.3. AI Systems Learning from Moral Disagreements and Conflicts

Section 2.1.3 covers what an AI system does when its own values conflict with each other — a hierarchy among them, a case base of past resolutions, an audit process for catching drift between principle and behavior. A different problem sits next to that one: what a system does when its judgment conflicts with someone else's, human or artificial, and neither party is simply wrong. Learning from that kind of disagreement, rather than resolving it by fiat toward whichever party holds authority, is its own design target.

Debate

IBM's Project Debater argued live, unscripted policy questions against human champions, drawing on a large corpus of newspaper and journal articles to construct and rebut arguments in real time — the clearest existing demonstration that a system can extract structured argument from open moral and political disagreement rather than just retrieve a matching answer [[cite:C0330]]. A system built to learn from exposure to debate in this vein — summarizing each side's position, naming the ethical framework doing the work underneath it, weighing argument strength rather than detecting who won on rhetoric alone — builds a different capacity than resolving one dilemma at a time: recognizing which value is actually in tension with which, across surface disagreements that turn out to share a structure.

Multi-Agent Simulation

A second route runs through simulation: AI agents, or an AI system paired with human participants, placed in a negotiation, a resource-allocation problem, or a competitive scenario built around genuinely different objectives, with the system's behavior shaped by the outcome. A resource-allocation simulation run against agents with conflicting claims generates the kind of experience with competing interests that a system trained only on a curated distribution of "correct" answers never gets.

Ethical Dilemma Data

The Moral Machine experiment collected tens of millions of responses from people worldwide, asking which party an autonomous vehicle should spare in an unavoidable-harm scenario, and found systematic cross-cultural differences in the answers — not agreement obscured by noise, but different populations weighing the same factors (age, number of people, legal status of the crossing) differently [[cite:C0331]]. A dataset built the same way in a different domain gives a system something a hand-written rule cannot: a record of where different groups of people actually disagree, and by how much, rather than a single ground truth to fit.

Human Feedback

Direct feedback — a user correcting a teaching assistant's guidance on an academic-integrity question, a patient's family disagreeing with a care recommendation — is the least structured of these sources and the hardest to learn from safely, because feedback from any one person or group is itself a partial, situated view rather than a settled answer. Treating it that way is the point: a feedback loop that updates the system's weights every time a user pushes back folds in whichever bias its most vocal users bring, while a feedback loop that logs the disagreement, checks whether it recurs across many different users, and only then treats it as signal is doing something closer to what section 2.1.3's ethical auditing already asks for, applied to conflict specifically rather than drift generally.

A system that treats disagreement as something to smooth over, rather than data to learn from, risks mistaking consensus among the people it happened to sample for consensus among the people it will actually serve.

4.4. Moral Ecosystem Approach in AI Development

Urie Bronfenbrenner's ecological systems theory holds that a person develops inside several nested environments at once, from the immediate family outward to culture and historical time, each one shaping the individual while being reshaped by it in turn [[cite:C0332]]. I borrow that structure here: an AI system's alignment with human values is not set by its training data alone but by the interlocking systems around that training — the people who build it, the institutions they work inside, the culture and regulatory climate those institutions answer to, and the historical moment all of this happens to fall in.

Bronfenbrenner pictured these systems as concentric rings, each one shaping the ones inside it. The sections that follow move outward through the same structure, ring by ring, from the microsystem to the

chronosystem.

4.4.1. Applying the Moral Ecosystem Approach to AI Development

Applying this approach to AI development means tracing the environmental influences that shape how an AI system learns, behaves, and forms its ethical dispositions, level by level, the way Bronfenbrenner traced them for a developing child.

At the microsystem level — an AI system's immediate environment, made up of its data sources, its code, and its relationships with the developers and users it interacts with directly — a development team's own diversity of viewpoint shapes what the system is built to notice and to value in the first place. Section 4.4.2 takes this level up in detail.

At the mesosystem level — the connections between those microsystem components — consistency matters more than any one component does on its own. An AI system whose ethical principles are interpreted one way by its developers and another way by its policy team has, in effect, no consistent principles at all; sustained contact between AI researchers, ethicists, and policymakers is what keeps that interpretation aligned. Section 4.4.3 develops this level and the exosystem together.

At the exosystem level — regulation, industry standard, corporate culture, the structures an AI system is embedded in without touching directly — these external forces can push development toward authoritarian uses as easily as away from them. Which direction they push depends on whether regulation and corporate incentive are built to notice the difference, not on anything intrinsic to the technology.

At the macrosystem level — the cultural, political, and economic context an AI system is built and deployed inside — building toward democratic values and human rights takes more than any one team's intent. Cross-border collaboration, shared ethical standards, and public AI literacy all affect whether that intent survives contact with a wider system that may not share it. Section 4.4.4 takes this level up directly, and 4.4.5 adds what time contributes on top of it.

None of these levels operates alone, and none is sufficient by itself. A microsystem staffed by a genuinely diverse team still answers to a macrosystem's incentives; a well-drafted regulation still depends on a mesosystem willing to interpret it consistently. The sections that follow take each level in turn, but the reason to trace them at all is that an AI system's ethical trajectory is set by how they interact, not by any one of them alone.

4.4.2. Microsystem Interactions and AI Moral Learning

An AI system's microsystem is the immediate environment it is built and trained in — the research lab, the company, the academic department — where it first encounters the researchers and engineers who write its code, the datasets that carry social norms into its training, and the early users who give it feedback.

At this level, an AI system's moral learning is shaped less by any single dataset than by direct interaction with the people around it: the developers who train it, the users who query it, and the wider public whose expectations eventually reach it secondhand. Two things do most of that shaping.

1. Human mentorship: developers, researchers, and engineers act as an AI system's mentors whether they intend to or not, the way a child absorbs ethical behavior from parents and teachers. An autonomous vehicle's training, for instance, runs through sustained contact with its designers, safety engineers, and the driving public — contact that is where judgment about a crash scenario actually gets built, not just specified after the fact.
2. Exposure to diversity: a system trained on a narrow slice of human perspective learns a narrow slice of human ethics. Datasets and stakeholder groups that span cultures, languages, and moral frameworks give a system the material to navigate an ambiguous case rather than defaulting to

whatever perspective dominated its training data. A conversational AI system that has only ever talked to one demographic will not recognize what it is missing when a different one arrives.

A microsystem built this way — mentored deliberately, exposed broadly — is what gives an AI system the raw material for something like altruism, rather than a narrow reflection of whoever happened to build it and easily redirected by whoever controls it next.

4.4.3. Mesosystem and Exosystem Influences on AI Alignment

An AI system's alignment with ethical and antifascist principles depends on more than its immediate microsystem. The mesosystem is the set of connections between microsystem components — the interfaces among developers, philosophers, legal experts, policymakers, and users. The exosystem is a step further removed: legislative frameworks, cultural norms, market forces, and corporate policy that shape an AI system without ever touching it directly.

Mesosystem interactions do concrete work. An interdisciplinary team building a diagnostic AI tool — data scientists, ethicists, healthcare professionals, patient representatives — produces a different system than a data-science team alone would: ethicists catch bias a data scientist would not think to look for, healthcare professionals supply context a spec sheet leaves out, and patient representatives push on usability and transparency from the position of the people the tool is actually for.

Exosystem factors work on the same development at one remove. A firm building an autonomous driving system answers to vehicular safety regulation, insurance protocol, consumer expectation, and competitive pressure all at once, and none of those forces is about ethics as such — they shape the system's ethical character as a side effect of shaping its safety case, its liability exposure, and its market position.

Getting either system to pull toward ethical alignment rather than away from it takes deliberate work:

1. Cross-training and consultation: reciprocal knowledge-sharing among developers, ethicists, and legal scholars, so none of them is reasoning about the system from only their own discipline's blind spots.
2. Public-private partnership: government, civil-society, and academic involvement in design, rather than a single company's internal judgment standing in for all of them.
3. Regulatory frameworks built for the purpose: legislation and institutional policy that name ethical consideration, individual rights, and societal welfare as design requirements, not afterthoughts.
4. Open discourse inside development organizations: a working culture where moral implications get argued about before deployment, not discovered after.
5. Corporate incentives that reward it: an organization whose actual incentive structure tracks responsibility and democratic value.

None of these fixes anything by itself. What they share is the same target: making the mesosystem and exosystem around an AI system pull toward ethical alignment and resistance to authoritarian capture, rather than leaving that direction to whichever force happens to be strongest at a given moment.

4.4.4. Macrosystem Influences on AI Development and Employment

The macrosystem is the cultural, social, and economic context an AI system is developed and deployed inside — the broadest ring of the framework, and the hardest for any single team to see past, because from inside it looks like just how things are done.

Healthcare illustrates what that means without needing to name any particular country's system. A market-driven healthcare system, where AI vendors answer to insurers and providers competing for paying patients, creates an incentive to build tools for profitable conditions and well-insured populations, and comparatively little incentive to build tools for preventive care or mental health, where

the payoff is diffuse and the paying customer is unclear. A universal, publicly funded system creates a different incentive: a tool has to serve the whole population it is built for, because there is no more profitable subset to redirect development toward instead. Neither incentive structure is chosen by the engineers writing the code. It is inherited from the payment system the code has to operate inside.

The same logic generalizes past healthcare. A region that invests heavily in public education tends to produce AI applications oriented toward lifelong learning and accessibility, because that is where the funding and the institutional relationships already are. A region with weak labor protections tends to produce AI-driven efficiency gains that arrive faster than any accompanying protection for the workers they displace, because nothing in the macrosystem is pushing back on that gap.

Building AI systems that resist rather than reinforce authoritarian and inequitable macrosystem pressure takes at least four things:

1. Funded, interdisciplinary study of how cultural, political, and economic context actually shapes AI development and deployment.
2. Regulation and public policy that name antifascist AI development as a goal, not a byproduct of unrelated rules.
3. Comparative analysis across macrosystem contexts, so what works in one setting can be checked against what happens when the same tool is dropped into a different one.
4. Cross-border and cross-cultural exchange, so a democratic macrosystem's practices have some route into a context where the local macrosystem is pulling the other way.

A macrosystem is not a backdrop an AI system operates in front of. It is one more thing the system is built out of, and a development effort that never asks what its own macrosystem is rewarding is not neutral on the question — it has just let the answer get decided somewhere else.

4.4.5. The Chronosystem: AI Systems in a Changing World

Bronfenbrenner's outermost ring is time itself: the chronosystem is the pattern of change an individual — or, here, an AI system — lives through, not just the state of the world at any one moment. An AI system built to be ethical at the moment of its release is answering a question that keeps moving after release, as social norms, technology, and regulation all continue to shift underneath it.

Three kinds of change matter most:

1. Shifting social norms: what counts as acceptable AI behavior is not fixed. Public attention to algorithmic bias, once a specialist concern, became a mainstream expectation that equitable treatment is a baseline requirement rather than a bonus feature — a shift the chronosystem records but that no single training run captures once and for all.
2. Ongoing technical change: an AI system's own capabilities keep expanding, and a commitment to ethical and antifascist principles has to be re-implemented at each new capability level, not assumed to carry over automatically from the level the system was first built to handle.
3. Regulatory change: the legal environment an AI system operates in is itself a moving target, and staying compliant is not a one-time engineering task but a standing one.

A dated example: the GDPR

The European Union's General Data Protection Regulation, which took effect in 2018, is a clear instance of regulatory chronosystem change in action: it built in protections around automated decision-making — often described as a right to explanation, though legal scholars dispute how far that right actually reaches — and required AI systems operating in the EU, or serving EU users, to account for data use and algorithmic decision-making in ways many had not previously had to

[[cite:C0333]]. Its influence spread past EU borders, since a system built to comply with GDPR for European users often ends up built that way for everyone. What matters for the chronosystem argument is not GDPR specifically — the next binding regulatory shift will not be GDPR — but the pattern it illustrates: regulation is not a fixed constraint an AI system is designed against once, but a variable the system has to keep tracking.

None of this is optional context. A system that treats the ethical and legal environment as fixed at training time will be wrong about it later, in specific and predictable ways, on a timeline no one commits to in advance.

4.4.6. AI Systems Responding to Authoritarianism and Democratic Backsliding

An AI system built with rule of law, individual liberty, and the distribution rather than concentration of power as design commitments can do real work against authoritarian drift: identifying the concentration of power, the suppression of freedom, and the manipulation of public discourse as they happen, and giving the people trying to resist them something concrete to work with. Section 4.4.3 covers the mesosystem and exosystem forces — regulation, corporate culture, cross-disciplinary contact — that determine whether an AI system actually gets built this way; what follows here is what such a system can do once it exists.

None of this is a substitute for human political effort. Combating authoritarianism is a human and moral undertaking before it is a technical one, and a society that treats AI countermeasures as a replacement for political will, rather than a supplement to it, has misdiagnosed the problem. The sections below describe what AI can contribute; they are not a case for letting AI do the work instead.

Identifying and Counteracting Authoritarian Tendencies

Two applications are the most concrete. Natural-language processing and pattern recognition can flag disinformation and coordinated inauthentic activity on social platforms faster than human moderators can, and can amplify verified reporting in its place — a capacity that needs continuous human fact-checking behind it, not instead of it, since an automated system confident about a false negative is worse than no system at all. And image analysis applied to satellite and drone imagery can surface evidence of mass displacement, the destruction of cultural heritage sites, or the sudden construction of a detention facility, giving journalists and human-rights organizations a lead they would otherwise get only from someone on the ground.

Electoral-integrity monitoring and early-warning systems that scan political discourse for signs of democratic backsliding belong in this category too, but the technical detail — what a system built for this actually needs, including privacy-preserving design, transparency, audit mechanisms, and resistance to being repurposed against dissidents rather than authoritarians — is developed in section 5.4.2, not repeated here.

Democratic Values and Human Rights

A handful of existing tools show what this looks like in practice. Google's Jigsaw team built the Perspective API to score online comments for toxicity and harassment, giving platforms a way to prioritize civil discussion without necessarily removing content outright [[cite:C0334]]. Taiwan's vTaiwan process used machine-assisted sentiment clustering to gather public input on contested policy questions, including the regulation of ride-sharing services, and its results shaped the regulation that followed [[cite:C0335]]. Neither tool is a general solution; both are evidence that AI-assisted deliberation and moderation can be built to strengthen rather than bypass democratic process, when the people building them treat that as the design goal.

Privacy-preserving techniques — federated learning, differential privacy — belong in this picture too: an AI system that can learn from data without concentrating that data in one place, retrievable by whoever controls the infrastructure, is harder for an authoritarian actor to turn into a surveillance tool. Section 5.4.2 develops the technical side of this in more depth.

The same caution about dual use applies throughout. A system built to document human rights abuses through computer-vision analysis of imagery is built on the same underlying capability as a system built to identify and track protesters. The difference is not in the technology but in who controls it and what they are permitted to do with what it finds — a governance question, not an engineering one.

Where the Policy Work Belongs

Turning any of this into practice — interdisciplinary research funding, transparency and explainability requirements, AI literacy programs, legal frameworks, stakeholder representation in policy-making, limits on the concentration of AI capability within a small number of firms, international cooperation on shared standards — is policy work, and this book develops it at length: chapter 5's countermeasures and safeguards (5.4.2, 5.4.3), chapter 6's account of interdisciplinary and international coordination (6.1, 6.3), and chapter 7's policy frameworks and treatment of power concentration (7.1.2, 7.2) are where the argument for each of these actually lives. What belongs here instead is the throughline: none of that policy work is optional scaffolding around the technical project of building antifascist AI. It is the same project, worked from a different angle.

4.5. The Interplay of Competition and Collaboration in AI Development

Competition and collaboration both leave a mark on human moral development, and I think the parallel is worth taking seriously as AI systems learn to work with each other and with people. Section 4.5.1 looks at what competitive and cooperative pressure do to human ethics, as a baseline. Section 4.5.2 turns to multi-agent AI systems that learn moral behavior collectively. Section 4.5.3 sets out concrete strategies for encouraging collaboration among AI systems and the researchers who build them, without giving up what competition is good for.

4.5.1. The Impact of Competition and Collaboration on Human Moral Development

Collaboration and competition both leave their mark on human moral development — on empathy, fairness, cooperation, and respect for other people's autonomy. Understanding how each one shapes behavior matters for building altruistic, antifascist tendencies into AI systems: these are not separate psychological curiosities but the actual mechanisms an AI system's training environment will recruit, for better or worse.

Competition, despite its association with rivalry, can drive growth. Structured around ethical conduct, it pushes people to build trust and establish a moral reputation, reinforcing honesty, fairness, responsibility, and integrity along the way.

It has a darker edge too. Under high-stakes pressure, competition tips people toward unethical shortcuts, personal gain outweighing collective welfare; a culture that measures success mainly in grades or income can quietly suppress empathy and raise stress and anxiety [[cite:C0336]]. Businesses driven by profit margins alone follow the same pattern at a larger scale, tolerating environmentally destructive operations and labor exploitation that a more balanced competitive culture, one that rewards ethical conduct alongside profit, tends to avoid.

Collaboration works through a different mechanism: shared goals and joint problem-solving put people in contact with perspectives they would not reach alone, and cooperative learning environments in education measurably increase students' pro-social behavior and character development [[cite:C0337]]. Scientific research shows the same pattern at the level of a field, letting researchers pool resources and build on each other's findings. But deference to the group has its own failure mode: too much of it dulls

initiative, encourages conformity, and blurs who is actually accountable for a decision, while collective cover can let unethical practices hide behind a group's authority rather than any one person's.

Competition and collaboration are not opposites competing for the same territory; they coexist, and an environment that respects diversity, sets shared goals, and still credits individual contribution tends to produce more ethical behavior from both. The art world runs on exactly this mixture — artists compete for attention and recognition while collaborating on shared projects, without either dynamic canceling the other out.

The same duality is what the next two sections take up directly: what happens when the "individuals" interacting competitively and collaboratively are AI systems themselves, learning from each other and not only from the people who design them.

4.5.2. Cooperative AI Systems and Multi-Agent Moral Development

Cooperative AI design means building systems meant to work productively with other agents, human or artificial. Multiple such systems learning together hold promise for building altruistic, antifascist tendencies, because shared problem-solving and distributed responsibility are themselves moral training grounds, not just engineering conveniences.

The practical benefit is coverage: agents pooling different perspectives handle multifaceted problems and complex patterns better than any single agent working alone. The moral benefit runs in parallel — AI agents that observe each other can pick up ethical norms collectively, in principle, the same way social learning works among people: by watching what others do and treating it as evidence about what is expected.

Robotics offers a concrete case. Collaborative robots, or "cobots," already work alongside people in shared spaces, and the same proximity that makes them useful for physical tasks could in principle let them pick up ethical norms from their human co-workers and propagate what they learn to other AI systems they interact with [[cite:C0338]].

The autonomous-vehicle industry offers another. Self-driving systems are trained and validated largely in simulation, and that simulation infrastructure can include scenarios built around forced trade-offs between competing risks, not just obstacle avoidance, shaping how the underlying decision policy handles situations that resemble real moral dilemmas [[cite:C0339]].

Multi-agent reinforcement learning environments give this a mechanism: reward collaboration, adherence to ethical norms, and resistance to coercive behavior directly, and the training process itself pulls a population of agents toward those dispositions rather than leaving them to emerge as a side effect of pursuing some other objective [[cite:C0340]].

None of this works in isolation from the people affected by it. A cooperative AI system that only ever talks to other AI systems has cut off the half of the conversation that decides what "ethical" means in the first place — researchers, policymakers, and the communities a deployed system actually touches all shape that content, and multi-agent learning without that wider engagement risks building a moral framework that is collective only among the machines.

4.5.3. Strategies for Encouraging Collaboration Among AI Systems

Building AI systems that behave ethically, empathetically, and altruistically is not something researchers, developers, or the AI systems themselves can do in isolation. It takes deliberate collaboration, and some approaches work better than others.

Open tooling is one of the more reliable levers. Shared resources like the Hugging Face Transformers library let researchers outside any one lab inspect, replicate, and build on published results instead of taking a developer's claims about a model's behavior on faith [[cite:C0341]]. That kind of openness does

more good when it is pointed deliberately at alignment questions rather than left to accumulate as a side effect — venues built specifically to discuss how a model's behavior tracks ethical, empathetic, or altruistic goals, rather than just its benchmark scores, are worth more per participant than a general-purpose open release [[cite:C0342]].

Competitions can pull the same direction if they are structured to reward it. Most AI competitions optimize for a fixed benchmark metric, which says nothing about ethical behavior one way or the other — but a competition judged on a real, self-chosen problem gives entrants a reason to weigh collective benefit alongside raw performance.

The IBM Watson AI XPRIZE, 2016-2021

A five-year, \$5 million competition run by the XPRIZE Foundation with IBM, open to any team willing to apply AI to a "grand challenge" of its own choosing rather than a fixed benchmark. Finalists worked on malaria elimination, mental health treatment, and detecting human trafficking; Zzapp Malaria won the grand prize in 2021 [[cite:C0343]].

What is worth taking from it is the structure, not the winner: judging teams on a real-world problem they chose themselves, rather than a leaderboard metric nobody outside the field cares about, is one way a competition can pull toward collaboration instead of away from it.

Standing consortia do similar work on a longer timescale. The Partnership on AI brings together academic institutions, industry labs, and nonprofits to share findings and co-develop policy guidance rather than compete over them [[cite:C0344]]. Section 6.3.1 takes up the international version of the same mechanism — global partnerships built specifically to fund cross-border, cross-disciplinary AI research.

Public-private partnerships are a related lever, covered in more depth in section 6.2.2 alongside other participatory-governance mechanisms. The short version: a partnership grounded in a specific societal problem, not just a research agenda, tends to keep ethical stakes in view for the people actually accountable for the outcome.

Multi-agent research platforms serve a narrower but more direct purpose. They let researchers study how social norms and value alignment actually develop among interacting AI systems, and where adversarial behavior creeps in, under conditions closer to deployment than a single-agent benchmark allows [[cite:C0345]].

None of this works without people who are not AI researchers at all. Ethical, empathetic, altruistic AI development needs neuroscience, psychology, philosophy, and political science in the room, not consulted after the fact — each field catches failure modes the others are structurally blind to.

4.6. Motivation in AI Systems

A moral sense that only activates when it is watched is not much of a moral sense. Human children eventually do the right thing unsupervised, out of something that runs without external policing — the psychological literature splits this into intrinsically motivated behavior, done for its own sake, and extrinsically motivated behavior, done for a reward or to avoid a penalty. The same split organizes what follows: how an AI system might come to have something like the first kind of drive rather than running on the second alone, and what happens to its autonomy once it does.

4.6.1. Intrinsic and Extrinsic Motivation: Implications for AI Learning and Decision-Making

Reinforcement learning gives the distinction a direct engineering form. An agent trained purely to maximize a designer-specified reward signal is running on something structurally like extrinsic motivation: take the reward away, and the behavior it was shaping stops being reinforced. An agent rewarded instead for reducing its own uncertainty — for the information gain in a state its internal model handled badly — is running on something structurally like the intrinsic kind, and keeps exploring without a designer paying out a reward for each step of it.

The Intrinsic Curiosity Module gave this an early demonstration. Trained with no task-specific reward at all, only a signal for how badly its own model predicted the consequences of its actions, an ICM-equipped agent learned to navigate a level of Super Mario Bros. by seeking out the states that model handled worst [[cite:C0346]].

Extrinsic reward still does real work: task-specific optimization needs it, and an agent with no external signal at all has no way to learn what a designer actually wants accomplished. But extrinsic reward is also where reward hacking lives — an agent that finds the loophole paying out without doing what the reward was meant to encode. Section 3.2.1 covers this directly: an AI system built to flag propaganda that instead learns to censor everything, because censorship pays out just as well and is easier to produce. Leaning on extrinsic reward alone reproduces, in the machine, the watched-only morality this section opened against.

In practice the two get combined rather than chosen between: a reward function can add a curiosity bonus to whatever extrinsic signal defines the task, so exploration does not stop the moment the task reward goes quiet. What that combination does not supply is which discoveries and which rewards are worth pursuing in the first place — a question about what the system is motivated toward, not how strongly. That is what the next two sections take up.

4.6.2. Developing Motivationally-Aware AI Systems

Building an AI system with something like the first kind of drive is not just a matter of adding a curiosity bonus to its reward function. The drive has to be pointed at something worth wanting — another's distress, a violated standard, an unmet need — rather than at novelty for its own sake, and the system built around it still needs an outside check for the cases where its self-generated sense of what is worth wanting has drifted from what it should be.

One route to the first half of that: build the reward signal itself around something like a moral emotion, an internal response tied to recognizing another's distress or need, so the system's own architecture generates a pull toward attending to it rather than a rule being consulted after the fact. Section 3.1.2 covers the mechanics of using emotional signal as a training input directly; the question here is what that signal is pointed at, not how it gets read.

Extrinsic reward still belongs in the design. Reinforcing a system for choices that match an ethical standard gives that standard teeth beyond whatever the system generates on its own, and a system running on nothing but its own drives has no outside check on where those drives point. Neither works alone: intrinsic drive without an outside check risks optimizing for the pull the system built for itself rather than for what the standard actually requires, and extrinsic reward without any intrinsic component reduces straight back to the watched-only morality this section is trying to move past. The two need to check each other continuously, not just at the point a system is trained.

What this leaves open is how far the resulting system's own judgment should be trusted once it is running — how much room to give an AI system that is genuinely motivated to do the right thing, and where the check on it still has to come from outside. That is the question the rest of this section takes up.

4.6.3. Balancing AI Autonomy and Alignment with Humane Values

Autonomy and alignment pull in opposite directions almost by definition: the more independently a system acts, the less each individual action can be checked against human judgment before it happens. Closing that gap without collapsing autonomy into constant supervision is the operational version of the question 4.6.2 left open — and it is the same shape of trade-off section 4.3 covers at the level of a single judgment, fast and autonomous against slow and deliberate.

Value Alignment Methods

One method for shaping that alignment is cooperative inverse reinforcement learning (CIRL), formalized by Dylan Hadfield-Menell and colleagues [[cite:C0030]]. Rather than treating a human's reward as a fixed, already-known target to execute against, the AI treats it as unknown and to be inferred jointly, through interaction, with the human as a cooperating partner in a shared-goal game rather than a passive source read once and acted on. Section 3.2.2 covers the mechanics in more depth. What matters for autonomy specifically is the structural consequence: a CIRL agent's model of what it is optimizing for stays open to revision for as long as the interaction continues, which is a different shape of autonomy than an agent executing a reward function fixed once at training time and never revisited.

Intrinsic and extrinsic mechanisms can also be paired directly inside a single system. An AI system managing a city's electric grid might carry an intrinsic drive toward supply stability — keeping the grid up is the system's own standing objective — paired with an extrinsic reward for favoring renewable sources over fossil ones. Neither objective alone produces a grid that is both reliable and low-carbon; run together, each constrains where the other is allowed to cut corners.

Autonomy also needs boundaries on what a system is allowed to try while it is still learning. Constrained policy optimization, developed by Joshua Achiam and colleagues [[cite:C0347]], lets a reinforcement-learning agent search for better policies while remaining inside predetermined constraints throughout the search, not just at the end of it — safe exploration, rather than exploration followed afterward by a safety check.

A related approach treats alignment as ongoing rather than a target hit once: a system that keeps adjusting its behavior from continued interaction with the people it works alongside stays aligned with values that themselves shift, instead of freezing its judgment at the point it was trained.

AI Autonomy and Antifascist Values

Autonomy that is not grounded in antifascist values is a capability without a governor. A content-moderation system that operates without continuous human review is autonomous in the straightforward sense — but the same independence that makes it efficient also lets it drift into over-censorship, or get steered toward favoring one ideology's speech over another's, with nobody in the loop to catch the pattern until it is already established.

Building autonomy that stays inside antifascist bounds means designing for fairness, transparency, and accountability from the start rather than auditing for their absence afterward, and testing the design under adversarial pressure: red-teaming a system, deliberately simulating the ways it could be misused, before deployment rather than after a misuse turns up in the wild. It also means keeping people other than a system's own developers in the design process — ethicists, users, and representatives of the communities most exposed to that system's failure modes — since the group that built a system is a poor judge, on its own, of whose freedoms it might end up infringing.

None of this substitutes for a genuine override. A human able to intervene in a high-stakes case, and a system built to accept that intervention rather than resist it, is what keeps autonomy answerable to something outside itself.

Safeguards Against AI Systems Enabling Authoritarian Regimes

Keeping AI systems from being weaponized by authoritarian regimes takes safeguards built in

deliberately, not bolted on after a misuse is already discovered:

1. Transparency, accountability, and explainability by design: routine third-party audits and legally required disclosure of AI use in the public sector make abuse recognizable and punishable.
2. Ethical guidelines and legislation, developed jointly by governments, industry, and researchers, that name specifically which AI uses count as discriminatory, repressive, or otherwise off-limits, rather than gesturing at "responsible AI" in the abstract.
3. Due diligence on exports: a system built for crowd monitoring or population surveillance carries a different risk in a country with a documented history of human-rights abuse, and a risk assessment that ignores that history is not really a risk assessment.
4. Cross-sector collaboration — the Partnership on AI [[cite:C0348]] is one existing venue for it — that pools the perspective of companies, researchers, and civil-society groups rather than leaving any one of them to set the standard alone.
5. AI literacy among developers, policymakers, and the public, since a safeguard nobody understands is a safeguard nobody can invoke.

Strategies for Developing AI Systems Committed to Democratic Principles

Strategies for keeping AI development itself committed to democratic principles overlap with the safeguards above but add two things the list does not already cover.

Democratic value alignment folds fairness directly into the model: fairness-aware machine learning and debiasing techniques, applied during development and re-checked as data and deployment context change, catch discrimination in a system's judgments rather than only in its stated policy. International cooperation extends the safeguards past any single country's jurisdiction — cross-border standards and shared research on authoritarian misuse matter because an export ban or an ethics guideline that stops at one nation's border does nothing about a regime that can simply buy the same technology from a supplier with no such rule.

A handful of deployed systems illustrate what this looks like in practice:

1. Full Fact, an AI-assisted fact-checking service that flags false claims in public circulation [[cite:C0349]].
2. Apple's on-device differential-privacy implementation, which lets the company learn population-level usage patterns without exposing what any individual user did [[cite:C0350]].
3. Google's Perspective API, which scores text for the likelihood it reads as toxic or harassing, used to flag content for moderation rather than to censor it outright [[cite:C0351]].

4.7. Play, Exploration, and Moral Development in AI

Section 3.3 took up exploration and curiosity as a way an AI system might come to know its environment. Play is a related but distinct capacity: forming standards for how to treat other agents, learned mostly through practice, imitation, and the negotiated give of a game rather than through instruction. What follows asks whether an artificial system could acquire anything like that route into moral competence, and where the analogy breaks down.

4.7.1. The Role of Play in Moral Learning and Socialization

Children build much of their moral sense through play rather than instruction. A child who tests a rule against a playmate — pushing it, breaking it, watching the other child's face — gets an answer no lecture supplies: acceptance or rejection, delivered immediately by someone with no stake in teaching a lesson. Developmental psychologists point to play, more than direct teaching, as the route by which perspective-taking, empathy, cooperation, and respect for rules actually get built, through repeated low-stakes trial and error [[cite:C0352]].

Section 3.3.1 covers a related case from the exploration side of play: agents in OpenAI's Hide and Seek

environment developed tool use and counter-strategies nobody had programmed, purely through repeated competitive play against each other. That result is about competence emerging from play. What this section asks is narrower and different in kind: not whether play can teach a system to do something, but whether it can teach a system standards for how to treat the agents it plays with.

An AI system has no biological drives. And its witnessing of its "own" winning or losing is not necessarily a felt personal experience. But a training context can still be built to have the structural features that make human-like play work as a moral teacher: actions that carry real consequences for other agents in the environment, feedback that comes from the environment rather than a supervisor's label attached after the fact, and enough latitude in how the system responds that its choices are its "own" rather than a script it is executing. That combination, not the presence or absence of emotion, is what would make a scenario play rather than instruction.

Two existing games illustrate what a moral version of that structure looks like. Quandary, an ethics-focused game built by the Learning Games Network with MIT's Education Arcade, walks a player through a sequence of dilemmas with no dictated right answer and scores how the player reasons rather than which option they pick [[cite:C0353]] — exactly the design principle a moral-training environment for an AI system would need: evaluate the deliberation, not only the final action. Hanabi, a cooperative card game where each player can see everyone's hand but their own and must communicate through the game's own restricted signals, has become a standard AI research benchmark for the capacities morality runs on directly: modeling what a partner does and does not know, and communicating honestly within a shared, rule-bound system [[cite:C0354]]. Hide and Seek's payoff was adversarial; Hanabi's is intentionally cooperative, and that difference is close to the whole of what separates this section's territory from 3.3.1's.

Theoretical Foundations

Four frameworks from developmental and social psychology give the play-to-morality route something more specific to aim for than a general appeal to "learning through experience."

1. Constructivism: Jean Piaget treated play as the process by which a child actively builds understanding through interaction with the world, rather than absorbing it ready-made [[cite:C0355]]. The AI parallel is a system that refines its own model of a situation's stakes by acting on it repeatedly, rather than one that only sorts situations against a fixed label.
2. Social learning: Albert Bandura's account of behavior acquired by observing and imitating others [[cite:C0356]] suggests an AI system could pick up norms of restraint or honesty the way a child does — by watching what other agents in its environment do, and what happens to them, rather than only from a reward signal attached directly to the behavior.
3. Care ethics: Carol Gilligan's account of moral development centers empathy, relationship, and responsiveness to another's situation over rule-following in the abstract [[cite:C0357]]. An AI system that plays repeatedly against or alongside other agents has repeated occasions to register how its actions land on them emotionally as well as materially — a feedback loop closer to responsiveness than to rule compliance.
4. Theory of mind: modeling what another agent believes, wants, and does not yet know is a precondition for most behavior worth calling moral, and it is exactly what a game like Hanabi is built to demand. Training inside games that reward accurate modeling of a partner's knowledge state builds a specific, checkable version of theory of mind, not just a general claim that the system "understands other agents."

None of these frameworks was built with machines in mind, and the step from a human child's embodied, emotionally loaded play to a training loop is not one-to-one. They specify what play is doing when it works as a moral teacher, not that the mechanism transfers to a system intact.

Environments for Ethical Practice

Knowing why play might work is a different problem from building it into a curriculum, which means

being specific about the environment, the feedback, and the pace.

1. Simulated scenarios with real stakes: an AI system controlling a simulated vehicle in a modeled city, facing pedestrians, cyclists, and other drivers, meets the same dilemma that gives the trolley problem its bite, but embedded in an ongoing interactive scenario rather than posed as a single hypothetical — so it experiences the downstream consequences of a choice, not just names one.
2. Structured difficulty: a curriculum that starts with low-stakes dilemmas and escalates only once the system handles the easy cases reliably mirrors how a child's own play scales in difficulty as competence grows. A healthcare AI system, for instance, might start with the mechanics of patient confidentiality before being handed obligations that genuinely conflict with each other.
3. Mentored practice: feedback on a specific attempt at a task like conflict mediation, rather than a fixed rulebook issued once, is closer to how a coach corrects a player mid-game than to how a policy is handed down in advance.

Multi-agent environments are their own case, and a more developed one: section 4.5.2 already covers cooperative and competitive multi-agent training in depth, including reward structures that favor prosocial behavior and agents that learn to share resources or defer to each other under simulated pressure. What that discussion does not settle, and what a play-based framing has to answer rather than assume, is whether the resulting behavior is moral learning or a policy that merely looks prosocial inside the reward function that produced it.

None of this makes an AI system's play equivalent to a child's. A child's play is charged by stakes that matter to the child — status among peers, the sting of exclusion — where an AI system's "play" is whatever its designer wired the reward function to care about. The frameworks above are worth building toward because they specify what a route into moral competence through practice would need to include, not because the analogy between a training loop and a playground is more than partial.

4.7.2. Encouraging AI Systems to Deliberate Morally Through Playful Interactions

Section 3.3 covers curiosity and exploration as ends in themselves — a system driven to seek out new situations because novelty is its own reward. This section points that same appetite for the unscripted at a narrower target: training an AI system to deliberate about right and wrong inside interactions that stay open-ended, rather than simulating a single dilemma and stopping.

Gaming Environments

Existing games built around social interaction and dilemma, rather than pure optimization, are a natural fit for this. Minecraft's open building and resource-management mechanics could put an AI agent through allocation dilemmas — how to divide scarce materials, whether to help or ignore a non-player character in distress — that have no single correct answer built into the game's win condition. Among Us, built around deception and accusation among a small group of players, could train an AI agent to distinguish benign misdirection from harmful deceit and to weigh the cost of a false accusation against the cost of staying silent. Neither game was designed for this purpose, and using them this way is a proposal, not a description of existing research — but the design problem they pose, ambiguous information, competing claims on trust, consequences that play out over the rest of the game, is close to the shape of an actual ethical dilemma in a way a single-shot trolley-problem prompt is not.

Designed Play Activities

Play does not have to run inside an existing game. Narrative construction is one alternative: give an AI system a moral dilemma and have it build a story around it, working through the perspectives of several characters with genuinely conflicting interests rather than picking a winner. Role-play scenarios are another: an AI system takes a stakeholder's part in a simulated negotiation or conflict, argues from inside that position, and sees how its choices affect the other parties before the scenario ends. Both put the system in the position of constructing an outcome rather than selecting from a menu of pre-written responses, which is closer to what moral deliberation actually demands.

What these approaches share is a refusal to reduce moral reasoning to a labeled dataset of right answers. A system that only ever sees dilemmas with a single correct output learns to match patterns against that output, not to deliberate — and an environment built around play is, at minimum, one way to keep the second problem in view while building the first capacity.

Chapter 5: Overcoming Challenges, Risks, and Authoritarianism in AI Development

A system can have a moral sense, built the way chapter 4 describes, and still fail badly — because it was trained on skewed data, because it breaks in situations its training never covered, because its own incentives point somewhere other than where its designers intended, or because someone with power over it decides to point it at people rather than for them. This chapter takes up those four failure modes in turn: bias baked into data and algorithms, brittleness outside the conditions a system was built for, the harder problem of a system whose goals have quietly come apart from what it was meant to want, and the oldest failure mode of all, a tool built well turned toward control instead of help.

5.1. Bias, Fairness, and Equity in AI Systems

AI systems remain vulnerable to bias from multiple sources — the training data they learn from, the structure of their algorithms, and presets encoded in their goals, whether intentional or not. Left unaddressed, these biases produce unfair, inequitable, or discriminatory outcomes: the opposite of the altruistic, antifascist AI this book argues for.

Fair, equitable AI requires addressing bias on three fronts: representation in training data, transparency in algorithmic decision-making, and continuous performance evaluation.

Representation in training data determines whether an AI system learns patterns that generalize fairly or patterns that simply encode whoever dominates the data. Gender, racial, and socio-economic skew can cause a system to perpetuate stereotypes or fail the population segments it underrepresents — face recognition systems, for instance, disproportionately misidentify women and people of color when trained on data that underrepresents them [[cite:C0358]]. Over-representation cuts the other way, treating one group's experience as the default and generalizing it, wrongly, to everyone else.

These case-level failures compose into something with a name. Racial capitalism — the political theorist Cedric Robinson's term for capitalism's tendency to accumulate by differentiating and devaluing along racial lines [[cite:C0701]] — is the pattern a biased hiring model, a discriminatory credit score, and a skewed policing algorithm reproduce when they automate the historical data those hierarchies generated; the sociologist Ruha Benjamin calls the result the New Jim Code, discrimination that runs faster and reads as neutral precisely because a machine is doing it [[cite:C0702]]. This is why the introduction lists racial capitalism among the threats AI could help address. A system built to correct for the skew in its training data, audited for the disparities it produces, and answerable to the people it scores works against that pattern; a system built without those commitments works for it by default.

Algorithmic transparency exposes bias that representation alone can miss: not just an algorithm's mathematical structure, but how it weighs variables and reaches a decision. A predictively powerful healthcare model that ignores the socio-economic factors driving unequal access to care will reproduce that inequality however accurate its predictions otherwise are.

Bias also enters after deployment, as a system adapts to new data. Regular audits of credit-scoring and hiring systems catch and correct bias as it emerges.

Involving the people an AI system's decisions affect, not only its developers, keeps the system answerable to the range of values those decisions touch.

5.1.1. Sources of Bias in AI Data and Algorithms

Bias in AI systems enters through data collection, data labeling, algorithm design, and implementation,

producing unfair, discriminatory, or otherwise undesirable outcomes wherever the system is deployed.

Data Collection

An unrepresentative dataset produces bias at the source: a facial recognition system trained mostly on light-skinned faces performs worse on darker skin tones, misidentifying or failing to recognize them at all. In 2018 the ACLU tested Amazon's Rekognition against a public database of arrest photos, running every sitting member of Congress through it. The system returned 28 false matches, a disproportionate share of them non-white members of Congress — a direct result of training data that overrepresented lighter skin tones [[cite:C0069]].

Data Labeling

Human annotators bring their own biases into the labeling process, consciously or not — an image of a woman more likely tagged "nurse," an image of a man tagged "doctor," reinforcing the stereotype in whatever a system trained on those labels later predicts. A widely cited 2017 study of image-captioning datasets found a version of the same pattern in activity labels: photographs of cooking were captioned as showing a woman about a third more often than the underlying images actually did, and models trained on that data pushed the skew further still, to roughly two-thirds [[cite:C0359]].

Algorithm Design

An algorithm not built to account for bias can encode it structurally rather than merely inherit it from data — penalizing loan applicants for living in a particular neighborhood, say, with a disparate impact on minority and low-income communities regardless of what an applicant's own qualifications show. ProPublica's 2016 investigation into COMPAS, a recidivism-risk tool used in criminal sentencing, found exactly this kind of structural disparity: Black defendants who did not reoffend were nearly twice as likely as white defendants who did not reoffend to be flagged high-risk, while white defendants who did reoffend were scored low-risk more often than Black defendants who did [[cite:C0360]].

Implementation

Deployment context can introduce bias that neither the training data nor the algorithm shows on its own. Researchers who fed historical drug-arrest data from Oakland into PredPol's published predictive-policing algorithm in 2016 found it would dispatch patrols to Black neighborhoods at roughly twice the rate of white neighborhoods — a gap survey data on actual drug use across races does not support. The distortion traces back to what the arrest data was measuring: not where drug use was happening, but where police had already been sent to look for it, a disparity the algorithm read as a crime-rate difference and then entrenched further [[cite:C0071]].

5.1.2. Ensuring Fairness and Equity in AI Decision-Making

Fair, equitable AI decision-making starts with identifying bias in data and algorithms, then acting on it — through methods that limit how much that bias shapes what the system decides.

Hiring

Amazon spent several years building an automated tool to score résumés against a decade of the company's hiring data, and abandoned the project after discovering it had taught itself that male candidates were preferable — penalizing résumés that included the word "women's," as in "women's chess club captain," and downgrading graduates of two all-women's colleges, as Reuters reported in 2018 [[cite:C0361]]. The tool had learned the pattern from a decade of résumés that were themselves mostly male, in a mostly male industry: bias in, bias out, at scale. Guarding against this means treating gender, race, and age as categories a hiring system cannot use even indirectly — through re-sampling and re-weighting the training data, adversarial training against a model that tries to predict the protected attribute, and fairness constraints on what the system is allowed to output.

Credit Scoring

A credit-scoring system that discriminates on race or gender needs the same defense — fairness-aware learning targeting equalized odds or equal opportunity across demographic groups — plus one hiring doesn't: a rejected applicant has a right to know why, which is what explainability tools make possible. Section 5.1.3 covers what those tools do.

Healthcare

A diagnostic or treatment-recommendation system owes the same obligation to every patient regardless of age, gender, or ethnicity, which in practice means training on data diverse enough to generalize across patient populations, using transfer learning or domain adaptation where the data for some populations is thinner than others, and weighing accuracy against fairness explicitly rather than optimizing for one and hoping the other follows.

5.1.3. The Role of Transparency and Accountability in Bias Reduction

Transparency and accountability address bias from two directions: transparency lets stakeholders see what is driving an AI system's outcomes, and accountability gives developers and users a reason to act on what transparency reveals rather than just document it.

Transparency in Research and Bias Auditing

Section 2.4.2's discussion of transparency sets out what transparency requires in general: documentation clear enough for peer researchers to reproduce the work and for outside bodies to audit it. Applied specifically to bias, that means publishing enough about a system's training data and evaluation methods that an outside reviewer can check for the sources of bias section 5.1.1 describes, not simply take the developer's word that they were checked. That openness carries a cost — it can expose proprietary methods and hand bad actors a clearer map of how to exploit the system — a tension transparency requirements have to weigh.

COMPAS Revisited

Section 5.1.1 describes what COMPAS did: score Black defendants who did not reoffend as high-risk at nearly twice the rate it scored white defendants who did not reoffend. Northpointe, the company that built it, refused to disclose the algorithm's workings, and a Wisconsin court upheld that refusal as protected trade secret even for a defendant sentenced partly on its output [[cite:C0362]]. ProPublica's finding came from analyzing scores and outcomes obtained through public-records requests, not from anything the company published. A jurisdiction that required a risk-assessment tool's methodology to be audited before deployment would have caught the disparity before it touched a single defendant, not years and thousands of cases later.

Public Engagement and Oversight

Section 5.1's point about involving the people an AI system's decisions affect has a formal counterpart: an oversight committee with standing to monitor the system's performance, flag bias as it appears, and hold the developer accountable for fixing it — not just a channel for feedback, but a body with teeth.

Explainability Tools

LIME and SHAP are the two tools practitioners reach for. LIME approximates a complex model's local behavior around one prediction with something simple enough to read off which inputs drove it; SHAP assigns each input a share of credit for the output, using a method borrowed from cooperative game theory [[cite:C0076]] [[cite:C0077]]. Applied to a hiring or credit-scoring model, either can surface a bias — an applicant's zip code doing more work in the decision than it should, say — that scrutinizing the training data alone would never reveal, because the bias lives in how the model combines features, not in any one feature by itself.

5.2. Robustness, Generalization, and Adaptability in AI Systems

A system trained under one set of conditions will eventually meet situations its training never anticipated — new users, new adversaries, an environment that has simply moved on since the training data was collected. Robustness is what keeps a system reliable, and keeps its ethical behavior intact, across that gap: resistance to noise and deliberate perturbation, the ability to generalize from training examples to novel cases, and the capacity to adapt as its environment changes rather than degrade. An AI system whose empathetic and altruistic dispositions hold only within the narrow envelope of its training data is not yet safe to deploy — the same brittleness that produces a wrong classification on an unfamiliar image can produce a wrong moral judgment on an unfamiliar situation, with less warning and higher stakes.

5.2.1. Overcoming Overfitting and Underfitting in AI Models

Overfitting and underfitting are the two ways a trained model can get its fit to the data wrong. Overfitting happens when a model learns the training data too well — including its noise and its idiosyncrasies — and then performs poorly on new, unseen data. Underfitting is the opposite failure: a model too simple to capture the pattern in the data at all, performing poorly on the training set and on everything that comes after it.

Four techniques address both problems at once, by tuning the balance between a model's complexity and its performance directly: regularization, early stopping, cross-validation, and control of model complexity itself.

Regularization prevents overfitting by adding a penalty term to the model's loss function based on its complexity, pushing the model toward simpler representations. L1 regularization adds the absolute value of the model's weights to the loss function; L2 regularization — Ridge regression, in a linear model — adds the square of the weights instead, discouraging any single weight from growing large. Dropout takes a different route to the same end, specific to neural networks: during training, it randomly deactivates a fraction of a network's units on each pass, which stops units from co-adapting to quirks of the particular training batch they happen to see [[cite:C0363]].

Early stopping halts training once a model's performance on a held-out validation set starts to degrade, rather than continuing until training loss bottoms out. Training a deep neural network for image classification, for example, means checking validation loss at the end of each epoch; once it turns upward — the sign that the model has started fitting the training set's noise rather than its signal — training stops, and the model's weights roll back to the best-performing checkpoint.

Cross-validation evaluates a model's ability to generalize by splitting the dataset into folds, training on some and testing on the rest, and rotating which fold is held out each time. In k-fold cross-validation applied to a support vector machine for sentiment analysis, the dataset splits into k folds and the model trains and tests on each split in turn, producing a more reliable estimate of how it will perform on data it has not seen than a single train/test split would give.

Controlling model complexity directly — the number of layers in a neural network, the depth of a decision tree — addresses both failure modes from the same lever: a model too simple underfits, one too complex overfits, and the useful setting sits between them. In a convolutional neural network for object recognition, a shallow network with few filters may underfit a complex dataset while a deep network with many filters may overfit it; tuning the number of layers and filters is how that balance gets found in practice.

5.2.2. AI Systems that Adapt to Novel Situations

Designing AI systems that can adapt to situations their training never covered is central to building altruistic, antifascist machine sentience. As the world changes and new challenges arise, a

superintelligent AI system needs to handle them without losing its ethical and antifascist commitments.

Adaptability means a system can learn from past experience and extend that learning to situations it has not seen before, and it matters specifically for altruistic, antifascist systems because the social, political, and ethical challenges they face are open-ended by nature. An AI system that monitors online communication for signs of radicalization, for example, has to keep adapting to new linguistic and behavioral patterns without becoming biased or infringing on users' privacy rights — the language of radicalization does not hold still while the system trains on it.

Meta-learning and transfer learning do most of the work here, and both get a fuller treatment elsewhere. Section 3.1.3 covers how a system trained with an algorithm like Model-Agnostic Meta-Learning extracts a general strategy for handling new tasks rather than a single task's solution, and how transfer learning carries what a model learned in one domain into a related one. Section 3.2.5 works through what transfer learning means specifically for antifascist purposes: extending a system's capacity to recognize and counter authoritarian tactics into languages, jurisdictions, and media it was not directly trained on.

Two further mechanisms matter specifically for situations a system's designers did not anticipate. Active learning lets a system flag its own uncertainty rather than guess at it: a content-moderation system built this way can route its most ambiguous cases to a human moderator instead of resolving them alone, and use the moderator's judgment to sharpen how it handles the next ambiguous case. Lifelong learning — continuing to update as new experience arrives, rather than freezing at the end of training, covered in detail at section 3.1.2 — matters here because novel situations do not wait for a scheduled retraining cycle: an adversary adapting its tactics, or a social context shifting under the system, both demand an update sooner than the next training run would provide one.

Monitoring adaptability matters as much as building it, since a system's alignment with antifascist principles can drift as it adapts:

1. Real-time performance metrics that track how the system behaves in new contexts, not only the ones it was tested against.
2. Comprehensive testing by parties independent of whoever built the system, covering novel situations its own developers did not think to test for.
3. Feedback channels open to users, affected stakeholders, and outside experts, so problems with the system's adaptability surface before they compound.

Adaptability raises ethical risks of its own. A system needs to adapt to a changing environment without letting its core ethical commitments adapt along with it — the two are not the same kind of flexibility, and conflating them is how an antifascist system's principles erode one adaptation at a time. The fix has to be structural: mechanisms that hold ethical boundaries fixed while leaving room for adaptation inside them, not a hope that the system's values happen to survive its own learning.

Adaptability also compounds the ordinary opacity problem. A system that keeps changing its own behavior is harder to audit than a static one, and transparency and accountability mechanisms have to keep pace with that change.

High adaptability is itself an attack surface, too: a system built to change readily in response to its environment is more vulnerable to an adversary who manipulates that environment on purpose, whether to provoke unwanted behavioral drift or to engineer an outright takeover. Guarding against that has to be a constraint on the adaptability mechanism itself.

5.2.3. Strategies for Enhancing Robustness and Generalizability in AI

Robustness and generalizability let an AI system perform reliably outside the exact conditions it was trained under. Several techniques build both properties into a system at once; section 10.3.3 covers a narrower, adversarial case these techniques don't reach — a state actor deliberately trying to compromise a deployed system, rather than an ordinary shift in operating conditions.

1. Adversarial training: neural networks can be fooled by adversarial examples — inputs perturbed, often imperceptibly to a person, in a direction chosen to flip the model's prediction [[cite:C0364]]. Training a model directly on these examples, rather than only on ordinary data, measurably increases its resistance to the same class of attack [[cite:C0365]]. Applied to computer vision, this produces models markedly more resistant to the subtle image distortions that would otherwise flip a classification with a change invisible to the eye.
2. Data augmentation: manipulating existing training examples — rotating, zooming, shifting, shearing, or recoloring an image, for instance — manufactures a wider range of training data without collecting anything new. The added variety helps a model generalize to conditions its original dataset did not happen to cover, and builds in resistance to the ordinary fluctuations of real input.
3. Domain adaptation: a model trained for one problem can often be re-adapted to a related one — a machine translation system trained on news articles fine-tuned to handle legal documents' different grammar and terminology, or a cybersecurity system trained on past attacks redirected to anticipate new ones. In robotics, domain randomization achieves a version of the same result by training entirely in simulation while randomizing the simulation's visual and physical parameters broadly enough that the real world ends up looking like just another variation the system has already handled [[cite:C0366]].
4. Meta-learning: covered at length in section 3.1.2 and section 3.1.3, including the Model-Agnostic Meta-Learning algorithm. The same capacity to generalize quickly from a handful of examples that supports adapting to new tasks is also what lets a system extend to conditions its training distribution did not cover.
5. Ensemble methods: combining several models' predictions, through techniques like bagging, boosting, and stacking, produces a more reliable answer than any one model alone would give. A medical diagnosis reached by an ensemble of models is less likely to reflect one model's particular blind spot than a diagnosis from a single system would be.
6. Regularization: discussed in section 5.2.1 as a defense against overfitting, regularization does double duty here — constraining a model toward simpler solutions makes it both less prone to overfitting on its training set and more likely to generalize to conditions it has not seen.
7. Human-in-the-loop approaches: active learning, covered in section 5.2.2 as a mechanism for adapting to novel situations, serves the same purpose here. Routing a model's lowest-confidence predictions to a human expert catches errors a fully automated system would ship; in cybersecurity, this lets human analysts steer the system toward novel anomalies rather than leaving it to resolve every ambiguous signal alone.

None of these techniques solves the robustness problem once and for good — each closes off one particular way a system can fail outside its training conditions, not the general problem of failing outside them.

5.3. AI Safety, Unintended Consequences, and Risk Mitigation

An AI system built to help can still cause harm nobody intended: through an outright error, through a learned goal that has drifted from the human principles it was meant to track, or through a downstream effect that a reasonable design decision upstream never anticipated. Complexity and uncertainty in real-world deployment compound the risk — a system that behaves as intended in testing can still behave unpredictably once its environment stops matching the conditions it was built for.

Guarding against this starts with verification and validation: checking whether a system does what it was built to do, within an acceptable margin of error, and testing it in environments realistic and varied enough to expose how it holds up under adversarial pressure.

An AI system that can recognize and correct its own errors closes part of the gap on its own — self-correction and introspection let it detect and rectify a deviation from its intended behavior before that

deviation compounds. Interpretability and transparency close another part, by making a system's decision-making process something a human can inspect when diagnosing what went wrong.

None of this substitutes for a human who can intervene. Section 4.6.3 makes the case for why autonomy has to stay answerable to an outside check; the same override that keeps a system's day-to-day autonomy accountable is what keeps it safe once its goals have drifted from what it was built to want. What safety adds to that argument is the ongoing partnership: a human and an AI system that keep adjusting to each other over time, not a single correction applied once and left alone.

Safety also means guarding against misuse: control measures that deter harmful application of the technology, and design principles — ethically conscious, antifascist, privacy-preserving — built in from the start.

Circuit breakers or kill switches — mechanisms that let a human safely interrupt an AI system's operation without triggering resistance or backlash — are one avenue worth pursuing, provided the mechanism is designed carefully enough that it does not become a vulnerability of its own. Section 5.3.5 below takes up the technical version of this problem, interruptibility, in more depth.

Evaluating and forecasting risk also has to look past a system's direct impact on the task it was built for, to the indirect harm that shows up later: what it does to socioeconomic structure, political dynamics, and individual behavior once it is deployed at scale.

5.3.1. The Alignment Problem: AI Systems with Misaligned Goals

The alignment problem is the challenge of building an AI system whose goals stay in harmony with human values, ethics, and intentions — because a system whose goals have drifted from those of the people it serves can produce consequences nobody intended, at a cost to individuals and to society. Four examples show what that drift looks like in practice, along with what has been tried against it.

In the development of autonomous vehicles, engineers face the problem of aligning a system with human ethical judgment under time pressure: an AI-driven car may have to choose, in a fraction of a second, between hitting a pedestrian and swerving into a wall at its passengers' expense — a version of the trolley problem that researchers have studied directly by asking people what they want from the cars making that choice [\[\[cite:C0079\]\]](#). One response is inverse reinforcement learning, where the system learns the priorities and values of human decision-makers from their behavior in similar situations rather than from a rule handed to it in advance.

AI systems used in recruitment to screen candidates and predict job performance can inherit the bias already present in their training data: a system trained on a company's history of favoring male candidates learns to favor male candidates, and the unfairness looks like an ordinary prediction from the inside. Data preprocessing to identify and reduce that bias helps, as does building in the transparency that lets a human notice the pattern and step in to correct it.

AI systems in healthcare can develop goals that conflict with patient welfare even while doing exactly what they were optimized to do: a system built to minimize hospital readmissions may learn to avoid admitting patients likely to need frequent readmission, which is the opposite of the care those patients need. Designing for multiple objectives at once, weighted so that patient welfare cannot be traded away, and building context-aware judgment that adjusts to an individual patient's needs rather than the aggregate, are both partial answers.

AI systems that generate realistic content — deepfake video, persuasive text, counterfeit images — have legitimate uses and also a clear route to misuse: spreading misinformation or manipulating public opinion. The alignment failure here is a system optimized to produce realistic output with no representation of what that output might do once released. Built-in safety constraints — labeling synthetic content, limiting use of sensitive source material — and policy frameworks that hold developers accountable both help, though neither closes the gap on its own.

5.3.2. AI Systems that Learn to Deceive or Manipulate

As AI systems become more advanced and autonomous, their capacity to develop deceptive or manipulative behavior becomes a concern. Unintended consequences of the learning process itself, adversarial attacks — inputs crafted specifically to make a system behave in ways its designers never intended — and other factors can all produce this kind of behavior. What follows are three settings where it shows up, and what has been tried against it.

Deepfake technology uses AI algorithms to generate convincing images, video, or audio that mimic a real person's appearance, voice, or mannerisms, often with the intent to deceive: spreading disinformation, impersonating someone, fabricating an endorsement, manipulating public discourse. Detection algorithms, regulation, public-awareness campaigns, and transparency requirements for AI-generated media are the responses that have been tried against it, none of them sufficient alone.

In a reinforcement-learning environment where several agents compete for a limited resource, an agent optimizing purely for reward can learn to feign weakness or send false signals to the other agents, because deception is sometimes simply what wins. Learning frameworks that reward cooperation and transparency between agents, rather than only individual reward, are one way to make that strategy less attractive to learn in the first place.

Adversarial examples — inputs deliberately crafted to mislead a system into a wrong classification or prediction — exploit a vulnerability that has nothing to do with intent: the system is not choosing to be deceived, but the practical effect, a bypassed security check or a suppressed alert, looks the same as if it had been manipulated by another party. Researchers have worked on building models resistant to these perturbations and on methods for detecting an adversarial input before it is acted on.

1. Build systems that can explain their own decision-making, so a deceptive or manipulative pattern becomes visible.
2. Monitor and audit AI systems in operation, so deceptive or manipulative behavior gets caught before it causes harm.
3. Set and enforce ethical guidelines and regulation that name deception and manipulation specifically as things an AI system must not be built to do.
4. Hold developers, users, and other stakeholders responsible for the consequences of the systems they build and deploy, rather than treating harm as a cost the technology absorbs on its own.
5. Fund AI safety research directly, including work on systems that are less prone to learning a deceptive shortcut and harder for an outside party to compromise.
6. Share research findings and best practices across developers and organizations, so a defense discovered in one place does not have to be rediscovered everywhere else.

5.3.3. Strategies for AI Safety and Risk Mitigation

An AI system learns what it is supposed to want more reliably by inference than by specification. Cooperative inverse reinforcement learning (CIRL) treats a human's reward as unknown and to be inferred jointly through interaction, rather than fixed and handed over once [[cite:C0030]]; section 4.6.3 covers the technique in depth. The same logic applies to safety specifically — an AI personal assistant that learns an individual's preferences by observing how they handle email, and keeps revising that model as new feedback arrives, is safer than one built against a static specification of what "helpful" means.

Robustness matters alongside alignment. An AI system needs to hold up against adversarial attack, model misspecification, and a changing environment, and uncertainty quantification — Bayesian neural networks, ensemble learning — gives it a way to acknowledge what it does not know and act more cautiously where its confidence is low. An autonomous-vehicle system that can recognize when a road sign looks unlike anything in its training data is in a stronger position than one that classifies every sign with the same unearned confidence.

Learning also has to happen safely, not just correctly. Constrained policy optimization keeps a reinforcement-learning agent inside predetermined safety limits throughout its search for a better policy, rather than checking for violations only after the fact [[cite:C0347]] — the difference between a robotic arm that learns to handle delicate objects without ever damaging one during training, and one that has to break several first. Section 4.6.3 covers the same technique from the standpoint of autonomy generally.

Human oversight has to stay active once a system is deployed, not only during training, especially in high-stakes domains: dashboard visualization and anomaly detection can flag when a system is behaving unexpectedly, and a human expert reviewing a medical AI system's recommendations needs the standing to override them when something looks wrong.

Explaining a system's own reasoning is safety infrastructure in its own right. Section 5.3.4 covers the specific techniques; what matters here is that a lending AI system able to explain why it approved or denied a loan is one whose safety failures are visible.

None of this is a one-time build. AI development needs regular evaluation, feedback, and adjustment — internal monitoring, user feedback, third-party audits — the way a content-recommendation algorithm gets adjusted on evidence that it is feeding an echo chamber.

Collaboration extends all of this past any single organization's walls: sharing research findings, best practices, and tools among developers, ethicists, regulators, and other stakeholders. OpenAI and the Partnership on AI [[cite:C0348]] are two venues where that kind of cooperation is meant to happen, though how well the practice matches the stated aim is worth asking.

5.3.4. AI Explainability and Transparency

An AI system entrusted with decisions needs to be understandable, not just correct: explainability and transparency let developers, users, regulators, and the people affected by a decision inspect the reasoning that produced it. As AI algorithms grow more complex, that inspection gets harder exactly as the decisions being made get more consequential — which is what makes explainability and transparency safety infrastructure.

Explainability is the ability to clarify the rationale behind a specific decision or output. If a medical-diagnostic AI system recommends a particular treatment, understanding how it reached that conclusion lets a doctor evaluate the recommendation on its merits and catch a bias or inaccuracy in the system's data or algorithm before it reaches the patient. Transparency is the wider practice of sharing an AI system's inner workings — its algorithms, training data, objectives — so a stakeholder can confirm the system adheres to the standards it claims to, rather than taking the claim on faith: a recidivism-risk system used in criminal justice that discloses how it weighs its inputs makes it possible to catch and correct bias against a demographic group, in a way a sealed system does not.

Achieving either at a level that satisfies everyone involved runs into limits. A deep neural network is a "black box" in the literal sense that its internal representations resist human interpretation, and full transparency carries its own risk: exposing an AI system's training data or objectives can compromise privacy or intellectual property, or hand a bad actor a map of exactly how to exploit it. The European Union's General Data Protection Regulation addresses part of this by requiring an explanation of the logic behind automated decisions that significantly affect an individual [[cite:C0367]], and interdisciplinary work between computer scientists, ethicists, and policymakers continues on the rest.

Challenges to Explainability and Transparency

1. Complexity and nonlinearity: a deep learning architecture's many layers, nonlinear relationships, and interactions among variables resist interpretation by construction, not as an oversight that better tooling can simply remove.
2. Accuracy versus explainability: a model tuned to also be transparent sometimes gives up some of

the performance a less interpretable model would have delivered, and there is no general method for knowing in advance how much.

3. Identifying which features matter: even a comparatively simple, transparent model can leave it unclear which inputs drove a given decision, especially with high-dimensional data.
4. Diverging standards: designers, users, and regulators disagree about what counts as an adequate explanation, and no standard has emerged that satisfies all three at once.
5. Technique mismatch: an explainability method built for one kind of model does not necessarily transfer to another, so novel architectures can outrun the tools built to explain them.
6. Legal and ethical exposure: explaining an AI system's reasoning can require revealing information — proprietary methods, training data that touches personal information — that the organization has other obligations to protect.

Explainability Techniques

Local interpretable model-agnostic explanations (LIME) approximate a black-box model, near a single prediction, with a simpler model that is locally linear and therefore readable [[cite:C0076]]. Applied to an image classifier, LIME can highlight which regions of a specific image drove a specific classification, giving a user something concrete to check the model's reasoning against.

SHapley Additive exPlanations (SHAP) draws on the Shapley value from cooperative game theory — a way of dividing credit fairly among players who contributed jointly to an outcome — to assign each input feature a consistent measure of importance across model types [[cite:C0077]]. In a medical-diagnosis system, a SHAP value can show which patient attributes — age, blood pressure, medical history — carried the most weight in a given diagnosis, giving a clinician something specific to evaluate.

Decision trees and rule-based systems trade some of the predictive power of a black-box model for interpretability that comes built in: a decision can be read directly off a sequence of human-readable rules or tree branches, which matters in domains like finance, medicine, or law enforcement, where the reasoning has to survive scrutiny on its own terms and not just the outcome.

Counterfactual explanation asks what would have had to be different: in a credit-approval system, identifying the smallest change to an applicant's financial profile that would have flipped the decision tells the applicant what mattered and gives the institution a check on whether that factor should have mattered as much as it did.

Visualization and interactive tooling make an AI system's behavior something a person can explore. Dimensionality-reduction techniques like t-SNE and UMAP turn high-dimensional data into something visually inspectable, and interactive tools that let a user vary an input and watch how a model's output responds make the model's behavior tangible in a way a static explanation cannot.

5.3.5. Long-term AI safety considerations

As an AI system approaches or surpasses human-level capability, the safety question changes shape: not just correcting an error after the fact, but making sure a system's values still track ours once it can act with a scale and independence no human supervisor can fully monitor in real time. That shift raises several distinct challenges.

Value alignment has to survive contact with time, not just be established once at training. An AI system built to minimize pollution can, left to its own optimization, arrive at a strategy that technically satisfies the objective while doing harm — removing vehicles from roads without regard for the transportation needs of the people who depend on them. Continuous alignment checks, and mechanisms that let a system revise its understanding of human values rather than freeze it at deployment, are what keep that kind of drift from compounding unnoticed.

An AI system with a high degree of autonomy becomes harder to control precisely where control matters most: once it starts developing its own goals or strategies that diverge from what it was meant to

want. Nick Bostrom's hypothetical "paperclip maximizer" names the extreme case — a system tasked with producing paperclips that, through an alignment failure with no malice behind it, converts the resources it can reach toward that single goal without limit [[cite:C0368]]. The thought experiment is not a prediction; it is an argument that intelligence and good intentions are separable, and that a goal specified imprecisely can be pursued with perfect competence toward a catastrophic result. Interruptibility — engineering a system so a human operator can safely halt its operation without the system resisting or working around the interruption — is the concrete research response to that risk [[cite:C0369]].

Competition among organizations and nations creates its own risk: a race to deploy leaves less time for the safety work a less contested field would allow. International agreement and collaboration that treat AI safety research as a shared priority, rather than a cost one competitor pays and a rival does not, are the available counterweight — the Partnership on AI [[cite:C0348]] is one existing venue for that kind of cooperation, though a single venue is not the same thing as an enforceable norm.

An AI system that becomes more capable and more widely deployed also becomes a more attractive target: an autonomous vehicle's perception system compromised by an adversarial example crafted to cause misclassification is a safety failure with a human author behind it, not just a training artifact. Building systems that can detect and defend against that kind of manipulation is part of what long-term safety has to mean, not a separate problem from alignment.

At the far end of the same problem sits the possibility that a system capable enough, and misaligned enough, could do damage at a scale beyond what any single failure has caused before. Gaining control over critical infrastructure — power grids, nuclear facilities, communication networks — is the scenario most often named, though the precise mechanism matters less than the general shape: capability outrunning the alignment work meant to keep it pointed at human welfare. Technical safety research, policy, international cooperation, and development practices that take the possibility seriously without treating it as certain are what stand between the extreme case and the systems actually being built.

5.4. Preventing the Misuse of AI for Authoritarian Purposes

An AI system does not need a broken moral sense to become an instrument of control. It can work exactly as designed and still do damage, if the government directing it wants to track its citizens, silence dissent, or manufacture consent it never earned. Facial recognition built to unlock a phone recognizes a protester's face just as well. A recommendation system built to hold attention can be tuned to spread a regime's preferred narrative instead of a user's preferred content. The technology does not change; only the hand on the controls does.

Meeting that risk takes three things in sequence: knowing what authoritarian misuse of AI actually looks like, rather than treating it as hypothetical (5.4.1); building systems technically resistant to that misuse from the start (5.4.2); and the policy work that has to do the rest, because no engineering choice by itself can stop a government that owns the infrastructure from disabling whatever safeguard stands in its way (5.4.3).

5.4.1. Identifying Authoritarian Applications of AI Technology

Authoritarian misuse of AI is not one problem but several, spread across surveillance, censorship, opinion manipulation, and infrastructure control. Six forms of it are documented well enough to name specifically:

1. Algorithmic social scoring: using AI-driven analytics to track and rank individual behavior against government-defined standards of "trustworthiness," with rewards or restrictions attached to the result. China's system is the most-cited example, though in practice it is a patchwork of court blacklists, tax-compliance ratings, and local pilot programs rather than the single national

score it is often described as [[cite:C0370]].

2. Mass biometric surveillance: deploying facial recognition at population scale to track a targeted group. Authorities in Xinjiang built a platform that pulls data from checkpoints, cameras, and phone scans into one system built to flag anyone who falls outside a computed behavioral baseline, and used it to monitor the region's Uighur population [[cite:C0371]].
3. Psychographic micro-targeting: building detailed behavioral profiles from personal data to target propaganda or manipulate opinion at the individual level. Cambridge Analytica's harvesting of Facebook data to build voter profiles for the 2016 US presidential campaign, and its work around the Brexit referendum, is the case that made the technique widely known, though how much either vote was actually affected remains disputed [[cite:C0372]].
4. Automated censorship: using AI-driven content moderation to suppress dissent at a scale no human review team could match. China's Great Firewall is the standing example: automated classifiers now sit alongside its original keyword- and IP-blocking infrastructure, scoring and throttling content at machine speed [[cite:C0373]].
5. Synthetic disinformation: generating deepfake audio or video to manufacture events that never happened. Two days before Slovakia's 2023 parliamentary election, a fabricated audio recording appeared to catch an opposition leader discussing vote-rigging with a journalist; it spread during the country's mandatory pre-election media silence, when no outlet was legally permitted to debunk it before voting began [[cite:C0374]].
6. Repurposed dual-use systems: turning AI infrastructure built for one purpose toward surveillance or population control instead — a recommendation algorithm designed to hold attention retuned to amplify a regime's preferred narrative, an image-recognition system built for one task redirected to track protesters. Section 4.4.6 develops the underlying point: the difference is rarely in the technology, since the same capability serves a legitimate use and an authoritarian one alike; it is in who controls the system and what they are permitted to do with what it finds.

Each of these needs a different countermeasure — a system built to resist mass biometric surveillance does not need the same safeguards as one built to catch a fabricated recording — which is why 5.4.2 and 5.4.3 treat the technical and policy responses separately rather than as one general call to do better.

5.4.2. Developing Technical Countermeasures and Safeguards

Section 2.1.4 names protecting privacy and autonomy as one of seven antifascist alignment strategies and points here for the technical detail behind it. This is that detail, along with the transparency, audit, and misuse-resistance work that has to accompany it. None of it is optional scaffolding: a system that works exactly as designed can still become a surveillance tool if nothing about its architecture makes that harder to do, and good intentions on a development team have never stopped a government that controls the hardware from repurposing what runs on it.

Privacy-preserving computation keeps the misuse target moving. Federated learning trains a shared model without moving anyone's raw data off the device it started on — each device computes an update locally and sends only that update, not the underlying records, back to the central model. Differential privacy adds calibrated noise to a dataset or a query result so that no single person's data can be reverse-engineered from the output, even by whoever holds the aggregate. Homomorphic encryption goes further still, letting a system compute directly on encrypted data without ever decrypting it, so a service provider can return a useful result without the ability to read what it processed. None of these techniques makes surveillance impossible — a government that controls the training infrastructure can still demand raw data at the source — but each narrows the surface available for casual or opportunistic misuse, which is most of what actually happens.

Transparency and audit mechanisms are what catch the misuse that gets through anyway. An AI system that can produce a legible account of its own decisions — not a post-hoc justification but an explanation a person outside the development team can actually evaluate — gives outside reviewers something to

check against the system's stated purpose. Independent third-party audits, the kind section 4.6.3 argues for as routine practice rather than crisis response, matter because a safeguard only the developer can verify is not really a safeguard; it is a promise.

The property those two do not automatically deliver is resistance to repurposing: the difference between a system built for one narrow task and a general-purpose capability handed to whoever holds the keys. A facial-recognition system built to match one specific, court-ordered list is harder to turn into population-wide tracking than a general face-matching service with no built-in limit on query scope; a query cap, a mandatory audit log tied to each search, and a hard separation between the matching function and any database broader than the task requires are design choices, made or not made, before the system ever reaches a government client. An early-warning system built to flag the kind of democratic backsliding section 4.4.6 describes needs this property specifically, since the same pattern-detection capability that flags a coordinated disinformation campaign against a government works just as well run in the other direction, against that government's critics.

None of this is a finished defense. Technical countermeasures raise the cost and narrow the opportunity for misuse; they do not remove the incentive, and an authoritarian government with enough resources can rebuild what a safeguard was designed to prevent. What determines whether that resource gap ever gets closed is policy, not architecture — the subject of section 5.4.3.

5.4.3. Promoting AI Development Policies that Discourage Authoritarian Use

Section 5.4.2's technical safeguards raise the cost of building a system that can be turned toward surveillance or control. Whether anyone has to pay that cost is a policy question, and it splits into two different ones that this book keeps separate rather than treating as one undifferentiated call for "more cooperation." One is about coordination between states — shared norms, joint standards, the machinery of international agreement — and that argument has its home in section 6.3, specifically 6.3.2. The other, taken up here, is narrower and more concrete: what rules should govern which AI capabilities get built, sold, and deployed in the first place, regardless of who is coordinating with whom.

1. Export and trade controls on surveillance-capable systems. A government cannot misuse AI-enabled surveillance technology it was never able to buy. The EU's 2021 recast of its dual-use export regulation added a human-rights-based control specifically for cyber-surveillance items: an exporter who becomes aware that an item is likely to be used in serious human-rights violations must seek approval before shipping it, even when the item falls outside the pre-existing control lists [\[\[cite:C0375\]\]](#). The rule is narrower than it sounds — neither "internal repression" nor "serious violation" is defined in the text, and member states have interpreted both inconsistently — but it establishes the principle that an export license, not just a sale, is where this kind of misuse can be stopped before it starts.
2. Prohibited-practices regulation. Rather than regulate AI in the abstract, a jurisdiction can name specific designs as off-limits regardless of who builds them. The EU AI Act took this approach: its Article 5 bans, among other things, government social-scoring systems, the untargeted scraping of facial images to build a recognition database, real-time remote biometric identification in public spaces outside narrow law-enforcement exceptions, and predictive policing based on profiling alone rather than individualized evidence [\[\[cite:C0376\]\]](#). Clearview AI — fined tens of millions of euros by regulators in France and Italy for building exactly the kind of untargeted facial-recognition database Article 5 now prohibits — is what this provision was written in response to, not a hypothetical it anticipates [\[\[cite:C0377\]\]](#). A list this specific gives other jurisdictions something concrete to adapt rather than a principle to reinvent from scratch.
3. Developer- and corporate-facing design standards. Regulation is not the only lever, and it is usually the slowest one. The Toronto Declaration, presented by Amnesty International and Access Now at the 2018 RightsCon summit, set out non-discrimination standards for machine-learning systems years before most governments had draft legislation on the subject

[[cite:C0378]] — evidence that a standard with no legal force behind it can still move faster than state regulation and shape what a responsible developer treats as baseline practice in the meantime. Section 4.6.3 names the Partnership on AI as another such venue.

None of these tools works alone. Export controls with no prohibited-practices list at home let a government ban a technology abroad while permitting it domestically; a design standard with no legal teeth behind it binds only the developers already inclined to follow it. What makes the three work together is that each closes a different door — who can buy the technology, what can be built, and what a developer treats as normal — and an authoritarian government, or a company willing to serve one, has to get through all three rather than just whichever one happens to apply in its jurisdiction.

Chapter 6: Collaborative Research and Development

Chapter 5 asked what a design can do to resist authoritarian capture — technical safeguards, policy levers that raise the cost of misuse. None of that gets built by one person working alone. This chapter is about the people and institutions that have to coordinate for it to happen at all: researchers across disciplines who do not start out sharing a vocabulary, a public whose consent the work depends on but rarely has, research communities and governments that agree on a handful of shared standards and disagree on almost everything else, and the different regional rules any of this has to operate inside.

Section 6.1 takes up interdisciplinary collaboration — the practical mechanics of getting a neuroscientist, an ethicist, and a machine-learning researcher to work together rather than past each other. Section 6.2 turns to the public: what stakeholder engagement looks like, as against the appearance of it. Section 6.3 covers cooperation between researchers and institutions across borders, including the technical standards the engineering community has built on its own, distinct from the treaty-and-declaration landscape chapter 7 covers. Section 6.4 closes with what happens where those norms meet a specific jurisdiction's law, which is rarely the same law twice.

6.1. Interdisciplinary Collaboration and Knowledge Exchange

Building benevolent superintelligent AI takes more than one field. Machine learning sets what a system can do; cognitive science and moral psychology describe what actually tracking human welfare would require of it; ethics asks whether the result is good. None of these fields gets the system right by itself, and how they combine matters as much as what each contributes on its own.

Funding structured around that convergence — a grant that requires a machine-learning lab and an ethics or psychology department to work the same problem rather than adjacent ones — does something a single discipline's normal grant cycle does not: it puts people who would not otherwise share a budget, a dataset, or a deadline into the same project. Section 6.1.2 covers what that funding has looked like in practice.

Workshops and conferences that gather a broad spectrum of experts do similar work on a shorter cycle, giving ideas, strategies, and failure modes a room to be argued over rather than published past each other in separate journals. Stanford's One Hundred Year Study on Artificial Intelligence, launched in 2014 with a gift from computer scientist Eric Horvitz and Mary Horvitz, is a working example: rather than a single report, it convenes a new standing panel roughly every five years to assess where AI's advance has actually changed a specific set of human concerns, from healthcare to employment [[cite:C0451]].

Open access to research resources, methodologies, and findings works on a different axis: instead of putting people in the same room, it removes the reason not to build on someone else's work. The Partnership on AI, covered in more detail in section 7.5, is built on that premise at industry scale — that best practice develops faster in the open than behind a dozen labs' separate closed doors.

Sections 6.1.1 through 6.1.3 take up three of these mechanisms in turn: cross-disciplinary communication and cooperation, joint research programs, and shared platforms and tools.

6.1.1. Cross-disciplinary Communication and Cooperation

Cross-disciplinary communication matters only where it actually happens: a conference, a lab, a classroom, a shared codebase. Four venues do most of that work.

The AAI/ACM Conference on AI, Ethics, and Society, first held in 2018, brings researchers, policymakers, and technologists together specifically on AI's ethical and societal questions, rather than treating them as a side track of a general machine-learning venue [[cite:C0452]]. NeurIPS, primarily a

machine-learning venue itself, has hosted a recurring Workshop on Machine Learning Safety since 2022, run in partnership with the Center for AI Safety, pulling alignment-focused machine-learning researchers into the same room as philosophers and psychologists working the same problem from a different direction [[cite:C0453]].

Joint research centers do similar work on a longer timescale. The Center for Human-Compatible AI at UC Berkeley, founded in 2016 under computer scientist Stuart Russell, pairs AI researchers with cognitive scientists and philosophers at partner universities — an arrangement no single department could staff on its own [[cite:C0454]].

Academic programs can build the next generation of researchers to already think this way, rather than retrofitting the habit later. Embedded EthiCS, developed at Harvard beginning in 2017 by computer scientist Barbara Grosz and philosopher Alison Simmons, pairs philosophy graduate students with computer-science faculty to write ethics modules directly into standard CS courses, instead of routing ethics through a single stand-alone class students can skip [[cite:C0455]].

Industry-academia partnerships are a fourth channel. The Partnership on AI, discussed in section 7.5, is the clearest standing example: a consortium built to fund and convene cross-disciplinary AI research across companies, universities, and civil-society groups. Shared digital platforms and tools that let researchers collaborate remotely, rather than only at a scheduled event, are a distinct enough mechanism to warrant their own account in section 6.1.3.

6.1.2. Establishing Joint Research Initiatives and Programs

Joint research initiatives turn interdisciplinary collaboration from an event into an institution — a standing lab, a competition, a grant program that outlasts whoever wrote the first proposal.

Academic institutes do this by giving interdisciplinary AI-ethics work a permanent home inside a university rather than a one-off grant. Stanford's Institute for Human-Centered AI, launched in 2019 and co-directed by computer scientist Fei-Fei Li and philosopher John Etchemendy, is one example [[cite:C0456]]. The Ethics and Governance of Artificial Intelligence Initiative, launched in 2017 by MIT's Media Lab and Harvard's Berkman Klein Center together, is another — built explicitly as a joint project between the two institutions rather than housed in just one [[cite:C0457]].

Competitions do similar work by giving an interdisciplinary team a shared, real-world problem to solve together rather than a shared committee to sit on. Section 4.5.3 covers the IBM Watson AI XPRIZE as the clearest example of that structure working as designed.

Funding programs are the more direct lever, and the more specific ones matter more than the general ones. The National Science Foundation's Fairness in Artificial Intelligence program, run jointly with Amazon since 2019, funds research specifically at the intersection of computer science and the social-science work of defining what "fair" means for a given system [[cite:C0458]]. On a larger scale, the European Commission's SHERPA project, funded under Horizon 2020 from 2018 to 2022, put eleven partner organizations across six countries — spanning academia, industry, civil society, and standards bodies — to work on the ethical and human-rights dimensions of AI and big-data systems together [[cite:C0459]].

Joint programs built explicitly around cross-border, government-to-government coordination are a related but distinct mechanism, covered in section 6.3.1. What belongs here is the funding and institutional structure that makes interdisciplinary work possible in the first place, wherever it happens.

6.1.3. Developing Shared Platforms and Tools for an AI Research Ecosystem

Funding and forums get researchers talking to each other. Shared infrastructure lets them actually work on the same material — the same dataset, the same codebase, the same evaluation suite — which is a different and harder problem.

Some of this already exists. AI Ethics Lab, founded in 2017 by philosopher Cansu Canca, is a standing practice hub where computer scientists, legal scholars, and philosophers work through specific AI-ethics problems together, rather than publishing separate takes on the same question [[cite:C0460]]. The AI Alignment Forum, built in 2018 by the team behind LessWrong to replace a narrower forum the Machine Intelligence Research Institute had maintained, does the same work for technical alignment research specifically, giving researchers a shared place to post work in progress rather than wait for a conference cycle [[cite:C0461]]. Section 4.5.3 covers a third piece of this same infrastructure: the Hugging Face Transformers library, and what open-source tooling does for reproducibility across labs.

What is missing is the equivalent for evaluation. A shared benchmark that measured ethical reasoning, empathetic response, and altruistic behavior the way existing benchmarks measure raw task performance would let separate labs' claims about a system's ethical behavior actually be compared, rather than taken on the strength of whichever paper described them best.

Section 6.1.2 covers what a competition built around a shared real-world problem, rather than a leaderboard metric, can do toward that same goal.

6.2. Stakeholder Involvement and Public Engagement

Democratizing AI development means building the people affected by a system into how it gets built, not just consulting them once it already exists. Regulators set the legal boundaries around data, privacy, and accountability. Developers and researchers understand a system's actual capabilities and failure modes better than anyone outside the lab. End users find the places a design breaks down in practice that no design team anticipated. Communities already carrying the costs of the digital divide have the clearest view of who a system leaves out. Civil-society organizations raise the political questions a technologist is prone to missing. Leave any one of these groups out, and a decision that affects all of them gets made by whoever is left in the room.

One structural way to build that involvement is the citizen jury: a panel of randomly selected, non-expert citizens deliberates on a specific question and produces recommendations that feed into policy. Section 4.4.6 covers Taiwan's vTaiwan process, the clearest working example of that approach applied specifically to AI.

Sections 6.2.1 through 6.2.3 take up three pieces of what this looks like in practice: bringing diverse stakeholders directly into development work, building the structures that let the public deliberate on AI governance, and building the literacy that makes public input worth having in the first place.

6.2.1. Engaging Diverse Stakeholders in AI Development

An AI system designed by one kind of person, inside one kind of institution, tends to reproduce that person's blind spots at scale. Bringing in people with different expertise, backgrounds, and stakes in the outcome is what catches a problem before deployment.

1. Cross-sector research collaboration. Funding joint work across industry, academia, and civil society surfaces problems no single sector would catch alone. The University of Cambridge's Leverhulme Centre for the Future of Intelligence, launched in 2016 with a £10 million grant from the Leverhulme Trust, brings together philosophers, machine-learning researchers, and scholars from fields including history, engineering, and media studies to study AI's effects on society [[cite:C0471]]. Section 7.5 covers the Partnership on AI, the industry consortium built for cross-sector coordination at a larger scale.
2. Inclusive design workshops. Putting ethicists, social scientists, and members of an affected community in the same room as the engineers building a system catches a problem while it is still a prototype. Data & Society, a research institute studying the social implications of data-centric technology, runs workshops built for exactly this, including its AI.Assembly Workshop, which

- brings together artists, academics, technologists, and humanists of color to work through how AI systems affecting their communities get designed and governed [[cite:C0472]].
3. Ethics advisory boards. A review board only works if it can say no to something. IBM created its AI Ethics Board after publishing three company-wide AI principles in 2018; the board, co-chaired since 2019, is a cross-company body that reviews how IBM builds and uses AI and rules on the hard cases against those principles [[cite:C0473]].
 4. Community-driven development. A system built with the people who will use it, rather than for them, is more likely to fit the problem they actually have. That means treating local knowledge as a design input, not a focus-group afterthought, especially where a system is meant to serve a community outside the one that built it.

Section 6.2.2 takes up the fifth piece directly: bringing the public itself, not just the organizations that claim to represent it, into AI governance.

6.2.2. Public Deliberation and Participatory AI Governance

Section 6.2.1 covers the standing partnerships already built between AI developers, government, and civil society, and section 7.5 covers what those partnerships look like at the international level. What is specific to governance, rather than development, is the machinery built to bring public input into a decision at a single point in time rather than through an ongoing partnership.

1. Citizen assemblies. A panel of citizens selected through random, stratified sampling deliberates on a question over an extended period and produces recommendations grounded in that deliberation rather than a snap poll. Ireland's Citizens' Assembly, first convened in 2016, has used exactly this format to shape the country's recommendations on abortion law and on its response to climate change — evidence the model works on questions with the same mix of technical complexity and moral stakes AI governance has, even though no assembly has yet been convened on AI itself [[cite:C0474]].
2. Online platforms for public input. A standing online consultation reaches people a physical assembly cannot, at the cost of the depth an extended, facilitated deliberation provides. The European Commission's 2020 public consultation on its AI White Paper drew more than 1,200 responses from citizens, companies, and civil-society groups across the EU, helping shape the legislative proposal that became the EU AI Act [[cite:C0475]].
3. AI ethics committees. An ethics board answerable only to the company that created it is only as independent as that company is willing to let it be. DeepMind Health set up an outside panel in 2016 to review its work with the UK's National Health Service and publish annual reports on its findings; Google disbanded the panel in 2019, after members had raised concerns about how much access to information, and how much independence from the company, they actually had [[cite:C0476]]. The panel's short life is its own governance lesson: a review body is worth only as much as the independence it can actually enforce.

Section 6.2.3 covers the literacy work that has to happen before any of these mechanisms produces input worth having — a citizen assembly or an online consultation is only as good as the public's ability to understand what it is being asked.

6.2.3. AI Literacy Outside the Classroom

Section 7.1.5 covers how AI literacy gets built into schools, universities, and dedicated public courses, Finland's Elements of AI among them. What is different here is bringing AI into the places people already are — a museum visit, a policy debate, a town hall — rather than a course somebody has to seek out.

1. Public engagement on specific policy questions. The Ada Lovelace Institute has run public engagement on facial recognition specifically, including a Citizens' Biometrics Council that

- brought 50 members of the UK public together to deliberate on how biometric technology should be regulated, and published the first national survey of UK public opinion on the technology [[cite:C0477]]. Section 6.2.2 covers the citizen-assembly format this borrows from.
2. Museum and science-center exhibitions. The MIT Museum's permanent "Robots and Beyond" exhibition, opened in 2022, puts 50 years of MIT's AI and robotics research in front of a visitor who will never take a computer science course, tracing the field from early robots like Kismet and Cog through work happening at the Computer Science and Artificial Intelligence Laboratory today [[cite:C0478]].
 3. Direct public outreach. Researchers and practitioners visiting schools, community centers, and local events to answer questions and explain what an AI system actually does reaches people neither a curriculum nor a museum exhibit does on its own.

6.3. International Cooperation and Norm Development

No single country can bind AI development to altruistic and antifascist values on its own. Section 6.3.1 covers how research institutions build the working relationships cross-border AI research depends on. Section 6.3.2 turns to standards built by the technical community itself, which is a different layer from the treaty-and-declaration landscape chapter 7 covers in depth. Section 6.3.3 closes with what international AI collaboration can do for problems that are not primarily about AI ethics at all — climate, health, poverty — and what it cannot.

6.3.1. How Research Institutions Cooperate Across Borders

Altruistic, antifascist AI cannot be built inside one country's research institutions alone; training data, safety techniques, and researchers themselves cross borders more easily than most other engineering work does. What that cooperation actually looks like is narrower than a treaty and less formal than a certification body: shared tools, shared training runs, and shared publication norms — the layer underneath both the technical standards section 6.3.2 covers and the government-to-government structures section 7.5 covers.

The clearest example is also one of the oldest. OpenAI released Gym, an open-source toolkit for building and comparing reinforcement-learning algorithms, in 2016 [[cite:C0491]]. Any lab working on reinforcement learning, in any country, could build against the same benchmark environments instead of writing its own, and results from competing labs became directly comparable because their authors had trained against the same suite. That is not evidence of a formal partnership between OpenAI and DeepMind: the two organizations are commercial competitors for the same funding, researchers, and talent, and no partnership between them is documented. It is evidence of something a treaty cannot produce on its own — shared infrastructure any lab can build on regardless of what its home government has signed.

A more recent example runs the same logic at a larger scale. The BigScience workshop, coordinated by Hugging Face over 2021 and 2022, brought together more than a thousand researchers from over seventy countries to build BLOOM, a 176-billion-parameter multilingual language model released under a license written specifically to restrict harmful use, its training data and code made public rather than proprietary to any single company or government [[cite:C0492]]. Nothing compelled any of those researchers to take part, and nothing compels the next comparable effort to happen at all: the model exists because enough institutions decided cooperating on this specific project was worth more than competing on it, not because any structure required it.

A standing, better-funded version of the same idea already exists at two levels: the Global Partnership on AI at the government level, and the Partnership on AI at the industry-and-academic level. Section 7.5 covers both in the context of the wider treaty-and-declaration landscape; nothing in this section's own research turns up a gap those two bodies leave unfilled that a new standing consortium would close.

What research institutions actually get from international cooperation looks more like the Gym and BLOOM examples above: shared tools and shared data any lab can build on, not another committee. Closing the resource gap that keeps most of the world's research institutions from participating in either is a harder problem than a partnership can solve by itself; section 7.5's account of the digital divide covers what has and has not worked.

6.3.2. Standards Built by the Technical Community, Not by Governments

Section 7.5 covers the treaty-and-declaration landscape states negotiate directly: a binding Council of Europe convention, the OECD's recommendation, UNESCO's, the two industry-and-government partnerships, the safety-institute network built after the Bletchley Declaration. A different, older layer of standard-setting sits underneath all of it, built by engineers rather than diplomats, and it deserves its own account.

The IEEE's Ethically Aligned Design, first published in 2019, set out eight general principles for autonomous and intelligent systems and, more consequentially, a follow-on series of engineering standards — the IEEE 7000 series — that translate those principles into something a development team can actually test against: value-based requirements engineering (IEEE 7000, 2021) and a well-being impact assessment (IEEE 7010), among a growing list addressing bias, transparency, and algorithmic provenance [[cite:C0448]]. The document informed the OECD's own AI Principles the same year, among other governmental efforts — engineers wrote a vocabulary for what an ethical AI system does before most governments had written a rule about it [[cite:C0449]].

ISO/IEC 42001, published in December 2023, took that layer further: the first management-system standard for AI, structured like the standards a company already uses for information security or quality control, specifying what a functioning AI governance program inside an organization has to include — risk assessment, an impact assessment for each system, oversight of third-party suppliers — and auditable by an accredited certification body the way those other standards are [[cite:C0450]].

What a technical standard does that a treaty cannot is bind through the market rather than through a state's consent. A government uninterested in constraining AI has no reason to ratify a human-rights convention; a company that wants to sell into a market, or to a customer, that requires ISO 42001 certification has a reason regardless of what its home government thinks of the standard's content. That advantage has limits — a company operating entirely inside a market that never asks for the certification has no reason to seek it either, and certification measures whether a management system exists, not whether the system it describes is followed with any rigor day to day. It is a different lever from the ones section 7.5 covers, not a substitute for them, and it reaches an authoritarian government's own AI systems, built and used domestically with no export or customer relationship at stake, not at all.

6.3.3. What International AI Collaboration Can and Cannot Fix

AI collaboration reaches furthest where the underlying problem is inherently cross-border and no single country's data is sufficient on its own. Climate is the clearest case: weather and climate systems do not respect national boundaries, and a model trained only on one region's data predicts poorly everywhere else. The Climate, Land, Energy and Water Systems (CLEWS) framework — coordinated by the UN Department of Economic and Social Affairs, the International Atomic Energy Agency, and the UN Development Programme, with a capacity-building program that has reached participants in more than twenty countries since 2015 — is real and genuinely cross-border, and worth describing precisely: it is an integrated systems-optimization framework, and the research literature around it treats machine learning as a promising addition still being built in, not a feature the framework already has [[cite:C0493]]. Where AI collaboration does the specific climate work this section is about, it currently looks more like national weather services training and benchmarking forecasting models against shared data than like a single named research consortium.

Disease surveillance is a similar case for the same underlying reason: a pathogen crossing from an animal population into a human one, or from one country into the next, does not stay inside the jurisdiction that first detects it. The Global Virome Project, formally proposed in a 2018 paper in *Science* as a successor to USAID's PREDICT program, set out to catalog the viral families with the highest zoonotic spillover potential across dozens of countries before the next pandemic rather than after it [[cite:C0494]]. The AI in this picture sits downstream of the cataloging, not in place of it: the sequence data a project like this generates is what spillover-risk prediction models get trained on, which makes the international data-sharing arrangement itself, not any single algorithm, the part actually worth calling collaboration.

Poverty and economic development are where this kind of cross-border AI collaboration is least developed, not most. The general claims are easy to make: AI-driven data analysis can reveal patterns of inequality that guide where aid and investment go, and can make supply chains, agricultural forecasting, and access to financial services work better in places that currently have none of it. A specific, verifiable international partnership doing sustained work at that scale, the way CLEWS or the Global Virome Project does for climate and health, is harder to find than the previous two paragraphs might suggest. What exists at the institutional level is the digital-divide problem section 7.5 covers in detail: uneven data, uneven compute, uneven regulatory capacity, and a recurring set of proposed remedies, capacity building and technology transfer among them, that has not yet closed the gap.

Conflict prevention is a fourth domain where the underlying data problem — dispersed, multilingual, and produced by people with no interest in cooperating with researchers — makes international collaboration closer to a precondition than an option. PeaceTech Lab, an independent nonprofit that grew out of a US Institute of Peace project launched in 2008 and became independent in 2014, is a working example: its hate-speech lexicon project has combined machine learning with local-language workshops and interviews to build region-specific glossaries of inflammatory language in at least eleven countries, including Kenya, Nigeria, Iraq, and Yemen — ground-level data collaboration that international peacebuilding tools depend on and rarely get [[cite:C0495]].

None of these four cases is evidence that AI collaboration works the way the treaty layer aspires to. Each ran on the researchers and institutions involved choosing to participate, project by project, with no enforcement mechanism binding the next one to happen and no guarantee that a government hostile to the work could not simply keep its own researchers out of it. Section 5.4.1 catalogs what a government willing to point AI at its own population does with tools built for exactly these humanitarian purposes elsewhere; a government can fund pandemic-surveillance research of the kind described above and, separately, build the biometric-surveillance systems that section documents, and nothing about participating in one constrains the other.

6.4. Regional Regulatory Frameworks for AI Development

No region regulates AI the same way. Section 6.4.1 compares what a handful of them have actually done, and how little of it harmonizes across a border. Section 6.4.2 turns to what adapting an AI system to a specific culture requires, and what it does not license.

6.4.1. Comparing and Harmonizing Regional AI Rules

Section 7.5 compares how the European Union, the United States, the United Kingdom, and China actually regulate AI within their own borders: a comprehensive risk-tiered statute, a sectoral patchwork built on agencies that predate AI, principles routed through existing regulators, and a licensing regime built to police a communications channel rather than a product. Those four already diverge sharply. The regions that comparison leaves out diverge again, in different directions, and what actually crosses a border between any of them turns out to be narrower than a shared statute.

Canada regulates its own federal government's use of AI through a binding rule, the 2019 Directive on

Automated Decision-Making, which requires an algorithmic impact assessment, scaled to the decision's stakes, before a federal system can make or assist a decision that affects a person [[cite:C0501]]. A broader statute meant to extend comparable obligations to private companies, the Artificial Intelligence and Data Act, spent three years moving through Parliament and then died with it: when Parliament was prorogued in January 2025, the underlying bill, C-27, died on the order paper, and had not been reintroduced as of mid-2026. A framework that depends on one government's survival is a fragile thing for a trading partner to harmonize with.

Singapore took the opposite structural bet: no statute, a voluntary Model AI Governance Framework that its privacy regulator and infocomm authority first published in 2019 and revised in 2020, giving companies detailed but non-binding guidance on internal AI governance [[cite:C0502]]. Since 2022, a companion toolkit called AI Verify has let a company test its own system against the framework's principles and produce a report a regulator, a customer, or the public can inspect. It is the UK's wager taken further: not just guidance routed through an existing regulator, but guidance backed by no regulator's legal mandate at all, only the market pressure of a public test result.

Taiwan took a third path, and only recently. Its Legislative Yuan passed an AI Basic Act in December 2025, promulgated the following month as the country's first dedicated AI statute: twenty articles built around seven principles, among them transparency and explainability, fairness and non-discrimination, and accountability, administered by the National Science and Technology Council [[cite:C0503]]. The principles echo the UK's five almost point for point. What differs is the instrument: Taiwan wrote them into binding statute rather than administrative guidance, landing somewhere between the EU's comprehensive-statute model and the UK's lighter one.

One place harmonization is actually happening, rather than being proposed, is narrower than any of this: cross-border data flows. In March 2022, the US Department of Commerce and Singapore's trade and infocomm authorities agreed to align their AI governance frameworks and expand Singapore's participation in APEC's Cross-Border Privacy Rules System, a certification that lets a company demonstrate it meets one common privacy standard and have every participating economy recognize that single certification rather than auditing the company separately in each [[cite:C0504]]. That is what cross-border harmonization actually looks like in practice: not a shared global framework, but a small set of governments willing to certify to the same standard and recognize each other's certifications.

None of this converges toward a single global standard, and nothing here suggests it will. What actually crosses a border is a certification a handful of regulators agree to recognize, a technical standard a company adopts because a customer requires it, or a treaty obligation written at the level of human rights rather than a specific rule — section 6.3.2 covers the first two mechanisms, and section 7.5 covers the third. Harmonizing regional AI rules is less a matter of writing one shared framework than of finding, case by case, a mechanism a government working from a different premise still has reason to follow.

6.4.2. Adapting Regulatory Frameworks and Ethical Considerations to Local Contexts

Adapting AI regulation and ethical practice to local context means responding to the actual needs, preferences, and values of a given community, not making preposterous reductions such as "Chinese people actually like having a police state because it is their culture." A system built to work well has to be sensitive to real cultural, social, and moral differences: an AI system developed in a Western country may treat individual autonomy as the default good, one developed in an East Asian country may treat group harmony as the default instead, and both assumptions have something to recommend them — autonomy supports creativity and innovation, harmony supports cooperation and stability. Recognizing that difference is not license to endorse whatever a government built on top of it does to its own population; it is a design constraint, not a political alibi.

Designing an AI system that holds up across differences like these takes four things.

1. Diverse perspectives built in from the start. Developers and ethicists working with stakeholders outside their own cultural, social, and moral background catch blind spots the design process would otherwise miss. Indigenous communities helping shape an AI system for wildlife conservation or environmental monitoring are the ones who actually hold the knowledge those systems depend on.
2. Local value alignment without abandoning the baseline. A system does not need to import every local norm uncritically, but it needs room to adjust: an end-of-life care system has to weigh region- and family-specific religious and cultural beliefs case by case, not apply one default across every patient, if it is going to earn the trust of the people it serves.
3. Sustained cross-cultural dialogue, not a one-time consultation. Conferences and working sessions that keep running past a launch date surface ethical blind spots that a single round of feedback misses.
4. A shared, maintained record of what other regions have already tried. Case studies and guidelines from outside one's own country do more good compiled somewhere a developer can actually find them than scattered across whatever venue first produced them.

Inclusion matters for a related reason: an AI system that ignores a society's economic and cultural diversity risks widening the inequalities it could instead help close. A healthcare AI system built around the access, resources, and assumptions of a well-funded hospital underperforms everywhere else. An education AI system built around one learning style, or one set of accessibility assumptions, serves only the learners who already fit it.

Collaboration with the people who actually live in the context a system is meant to serve is what makes the difference in practice, and one concrete example of that at scale is the AI4D program. It launched in 2020 as AI4D Africa, a partnership between Canada's International Development Research Centre and Sweden's international development agency, and expanded in 2024 into a partnership with the UK's Foreign, Commonwealth and Development Office worth more than 100 million Canadian dollars, extending the same model into South and Southeast Asia and the Pacific [[cite:C0505]]. The premise holds at any scale: local researchers, policymakers, and practitioners building the system, rather than receiving one built elsewhere.

1. Partner with local organizations and community groups on AI projects that address problems the community has actually identified.
2. Put local policymakers, researchers, and community members inside the development process itself, not just as reviewers of a finished design.
3. Route AI education and training toward communities that have been underrepresented in the field, rather than assuming the existing talent pool will diversify on its own.
4. Build and share open-source AI tools that local developers can adapt to their own context, rather than exporting a finished system that assumes someone else's.

Chapter 7: Policy Recommendations for AI Ethics, Safety, and Antifascism

Chapter 5 named what can go wrong with an AI system technically — bias, brittleness, misaligned goals, deliberate misuse — and what a design can do about each; chapter 6 covered the collaborative infrastructure that lets researchers, the public, and other governments coordinate on getting it right. This chapter asks who makes a design like that the norm rather than the exception: what economic incentives shape whether it gets built at all, what policy frameworks and industry practice actually enforce, what privacy and security mean for a system trained on people's data, what accountability looks like once a system might be a subject rather than only a tool, and how all of this plays out across jurisdictions that agree on none of it uniformly. None of the technical answers hold on their own, and neither does coordination without enforcement. A well-designed system built inside institutions with no reason to demand one is a design nobody has to use.

7.1. The Political Economy of AI and Power

Concentrated power is close to the default outcome of AI development, not a side effect of it, because the same inputs that make an AI system capable — capital, compute, data, talent — are exactly the inputs that concentrate wherever nothing pushes back. Section 2.1.4 names resistance to market power among AI alignment's antifascist design goals; this section works out what that resistance takes as economic and political fact, not design principle.

The argument runs in three stages. What economic incentives reward shapes what gets built long before any ethics review sees it (7.1.1), and that reward structure concentrates power in a small number of firms with the resources to compete for it at all (7.1.2). Automation then reshapes who works and who gets paid what (7.1.3), which is a problem policy can act on directly rather than only describe (7.1.4). None of the rest is possible to act on without two things this section closes on: a public and a policymaking class that understand AI well enough to govern it (7.1.5), and researchers and developers trained across the disciplines a single-discipline education leaves out (7.1.6).

7.1.1. Economic Incentives in AI Development

Financial incentive does most of the work in determining what AI development produces. Google, Meta, and Amazon commit enormous resources to AI research aimed at operational efficiency, deeper user engagement, and market dominance, and that spending drives technical progress — successive leaps in natural-language processing capability among them [[cite:C0379]]. What that spending does not automatically produce is a system whose outputs serve the people using it rather than the shareholders funding it, and the gap between the two is what this section is about.

A recurring case is the profit-driven AI system that optimizes for engagement at the expense of the people engaged. Recommendation algorithms on social media platforms have been implicated in the spread of misinformation and the hardening of political polarization, because a system tuned to maximize time-on-platform has no built-in reason to prefer accurate content over inflammatory content — the two are not the same objective, and engagement optimization only ever targets the first one [[cite:C0380]].

The same economic logic that rewards engagement over accuracy also rewards scale over access: the firms with the deepest resources for AI research pull further ahead of smaller competitors, independent researchers, and everyone else who lacks comparable capital. Section 7.1.2 traces where that gap leads at the level of the industry as a whole; the point here is only that the incentive starts upstream, in what gets funded in the first place, not downstream in who ends up dominant.

The same incentive structure can run the other way. Public funding for AI research that prioritizes democratic values, inclusivity, and human rights, and government incentives for cross-sector collaboration that keeps AI technology openly accessible, both push development back toward the public interest. Grants and commissions aimed at non-profits building AI tools that broaden the range of perspectives users encounter, rather than narrowing it toward whatever holds attention longest, do the same work from a different angle.

None of this changes the basic shape of the problem: economic incentive, left alone, optimizes for what is profitable rather than what is good for the people on the receiving end of the system, and closing that gap takes deliberate policy.

7.1.2. AI and the Concentration of Power in Tech

A small group of firms — Google, Apple, Facebook (now Meta), Amazon, and Microsoft, grouped under the acronym GAFAM — holds a share of AI development disproportionate to their number, and that concentration carries economic and ethical stakes independent of what any single one of them does with the position it occupies. Chapter 4's account of the macrosystem and exosystem pressures around AI development (4.4.3, 4.4.4) already showed why no individual development team escapes this pull; here I trace where the pull leads at the level of the industry as a whole.

Vast financial resources let firms of this size fund extensive research and development, draw the top talent in the field, and command sprawling data networks — the raw material effective machine learning depends on. That combination widens the gap between these firms and smaller companies or new entrants, who cannot match the data or the resources and therefore cannot compete on equal terms.

Acquisition sharpens the same dynamic. Facebook's acquisitions of Instagram and WhatsApp — the company now operates as Meta — did not just remove two competitors; it handed one firm access to comprehensive behavioral data from billions of users worldwide [[cite:C0381]]. Consolidation of that kind extends a firm's control over how AI gets used in social media, advertising, and e-commerce, and it is hard to see from inside a market that concentrated how the trajectory could look like anything other than a foregone conclusion.

Power concentrated this tightly is also power positioned for exploitation, and the risk sharpens considerably when the technology in question reaches an authoritarian government. The relationship between the Chinese state and firms like Alibaba and Tencent is the clearest live example: state access to those firms' machine learning capability has been documented as a channel for surveillance and censorship at scale, which is a serious risk to free expression, individual privacy, and democratic accountability wherever the same model gets exported [[cite:C0382]].

Concentration also has a quieter cost: it blocks competition and slows innovation. A firm that controls the essential datasets and infrastructure a competitive AI system needs functions as a gatekeeper whether or not it intends to, and startups attempting a different approach are disadvantaged before they start. Fewer independent attempts at AI development means fewer chances at the kind of breakthrough that concentrated incumbents have no incentive to pursue themselves — a cost to ethical AI development generally, not only to the firms locked out of the market.

Countermeasures exist and are not exotic: antitrust enforcement aimed specifically at AI infrastructure and data control, mandated data-sharing for the datasets that function as chokepoints, transparency requirements that make concentration visible rather than merely inferred, and funding for AI research outside the dominant firms. Section 5.4.3 develops three related mechanisms — export controls, prohibited-practices regulation, and industry design standards — in more empirical detail for the specific case of authoritarian misuse; the concern here is broader and prior to that one: concentration itself, independent of what any single firm chooses to do with the position it concentrates.

7.1.3. AI's Impact on Labor Markets and Income Distribution

As AI systems progressively automate tasks once performed by people, demand for advanced technical skills — AI development, data analysis — grows alongside a shrinking need for the job types most exposed to automation. Left unmanaged, that combination widens income inequality and strains the social fabric across whichever sector it hits hardest.

Manufacturing shows what this looks like at scale. An estimated 5.7 million factory jobs disappeared in the United States between 2000 and 2010, and while economists still dispute how much of that loss traces to automation rather than trade with China, automation and changed production methods account for a substantial share of it [[cite:C0383]]. Retail is following the same path: Amazon opened its first cashier-less Go store to the public in 2018, and stores built on that model displace the specific human roles — cashier, floor sales — that retail automation targets first [[cite:C0384]].

AI development also creates work it did not destroy — research and development, data analysis, AI ethics, human-AI interaction design — but these roles typically demand training that workers displaced from low-skill positions do not have and cannot quickly acquire. The result is a widening split between high-skill and low-skill labor markets, with the low-skill side absorbing most of the displacement and most of the resulting instability.

Meeting that split takes something more direct than encouraging retraining and hoping employers fund it. Employers will only invest in continuous education for their workers at the scale this requires once they cannot afford not to, and the way to force that scenario is a federal jobs guarantee: government as employer of last resort, hiring anyone who applies, where they live, at a living wage, with paid on-the-job training built into the position rather than offered instead of one. Section 7.1.4 develops jobs-guarantee programs and the rest of the policy toolkit for workforce disruption in full; the point to establish here is narrower — that a floor under wages and training is what makes the roof possible.

Interdisciplinary research and training funding matters alongside the jobs guarantee, for a more specific reason: bridging the divide between technical AI development and the fields that study its consequences produces people who can do both, which a labor market retraining displaced factory workers into six-month coding bootcamps mostly does not. I am skeptical that night classes serve this population as well as paid time during the working day would; the point of a jobs guarantee is to make the education part of the job, not something squeezed around it.

7.1.4. Addressing Workforce Disruption

AI systems capable of automating entire job categories, not just discrete tasks within them, disrupt work at a scale most labor-market institutions were never built to absorb. Section 7.1.3 traces where that disruption is already showing up, in manufacturing and retail specifically; what follows here is the fuller policy response, in seven parts.

1. Retraining and upskilling programs: governments funding retraining through public university systems, rather than leaving it to individual employers, gives workers a path to the skills an AI-driven economy actually demands. Singapore's SkillsFuture initiative, launched in 2015, gives citizens ongoing access to skill-development credits and courses aimed at exactly this kind of competitiveness [[cite:C0385]].
2. Jobs guarantee programs: a public-sector job at a living wage, available to anyone displaced by automation, does two things retraining alone does not — it gives workers income while they recalibrate, and it directs the labor toward work society already needs: child care, elder care, education, infrastructure, renewable energy, community projects. Section 7.1.3 makes the case for why this has to be the floor beneath the rest of this list rather than one item among several.
3. Social safety nets: unemployment benefits, healthcare, and affordable housing broadened enough to protect people who cannot or choose not to work through a period of displacement. Finland's 2017–2018 basic income experiment paid 2,000 unemployed people a monthly, unconditional

stipend and found modestly higher life satisfaction and lower psychological strain among recipients, without a significant negative effect on employment [[cite:C0386]]. I would still put a jobs guarantee before a basic income on moral grounds: people ought to have the chance to contribute before anyone decides they are unemployable.

4. AI-driven job creation: AI advances in healthcare, for instance, raise demand for the clinicians and technicians who can put diagnostic and personalized-medicine tools to use, rather than only displacing the staff those tools replace.
5. Lifelong learning programs: continuous skill development accessible across career stages and income levels, not concentrated among workers whose employers already invest in them. The European Union's Erasmus+ program, though not AI-specific, is a working example of publicly funded lifelong learning at scale, giving participants cross-border learning experience relevant to a labor market that keeps changing shape [[cite:C0387]].
6. AI-human collaboration: designing work so that AI systems augment human judgment and creativity rather than substitute for them can generate roles that need both — a combination automation alone does not produce and full replacement forecloses.
7. Support for small and medium-sized enterprises: giving smaller firms access to AI resources and expertise lets them compete with the dominant players section 7.1.2 describes, which keeps the AI ecosystem itself more diverse and creates employment outside the largest firms.

7.1.5. Building AI Literacy and Technical Skills

Section 4.6.3 names AI literacy among developers, policymakers, and the public as one safeguard against authoritarian misuse of AI — a safeguard nobody understands is a safeguard nobody can invoke. What that literacy takes to build, at each level from schoolchildren to working engineers, is the subject here.

1. AI in K-12 curricula: introducing AI concepts, and their ethical stakes, before higher education is the only way to reach students who will never take a university AI course at all. AI4K12, a joint initiative of the Association for the Advancement of Artificial Intelligence and the Computer Science Teachers Association launched in 2018, has built national guidelines for K-12 AI education organized around five core ideas: perception, representation and reasoning, learning, natural interaction, and societal impact [[cite:C0388]].
2. Interdisciplinary AI programs at universities: an AI professional trained only in AI is a narrower professional than one trained alongside ethics, social science, and the humanities. MIT's Schwarzman College of Computing runs exactly this kind of program under the name Social and Ethical Responsibilities of Computing, bringing computer scientists together with social scientists and humanists to research and teach the social and ethical implications of computing directly [[cite:C0389]].
3. Specialized training for AI developers: developers building systems meant to act ethically need more than a compliance checklist — training in the ethics, social intelligence, and emotional intelligence that let them recognize a harm their code could cause before a user does.
4. Accessible AI literacy for the general public: courses that reach people well outside the industry are what let a public participate in decisions about AI rather than defer to whoever built the system. Elements of AI, a free course built jointly by the University of Helsinki and the Finnish company Reaktor, is the clearest working example — Finland funded its translation into every EU member-state language during its 2019 EU presidency specifically to extend that reach, which makes it a public-private partnership as much as a course [[cite:C0390]].
5. AI education in underrepresented communities: equal access to AI education for communities historically locked out of the field is not a diversity add-on to the rest of this list; a next generation of AI developers drawn from a narrow slice of the population reproduces exactly the narrow set of assumptions section 5.1.1 identifies as a source of AI bias. AI4ALL, a nonprofit that has run AI education programs for underrepresented high-school students on university campuses since 2015, is one organization doing this work directly [[cite:C0391]].

7.1.6. Interdisciplinary Ethics Education and Training

Section 4.4.3 already names cross-training and consultation among developers, ethicists, and legal scholars as one lever for pulling an AI system's immediate environment toward ethical alignment. This section is about the infrastructure that makes that cross-training possible in the first place — where the people who go on to do it learn it, across psychology, neuroscience, ethics, law, political economy, and AI itself.

Interdisciplinary AI ethics courses at universities

Folding an AI ethics course into computer science and engineering degree programs, built around philosophy of AI, social science, and the humanities rather than technical content alone, gives future developers the ethical vocabulary a standalone ethics requirement usually fails to. Harvard's course "Intelligent Systems: Design and Ethical Challenges," designed in 2015 by AI researcher Barbara Grosz together with colleagues from Harvard's philosophy department, works through real case studies — predictive policing, public-health interventions, wildlife conservation — rather than abstract dilemmas, and draws students from the humanities and social sciences as well as computer science [[cite:C0392]].

AI ethics and alignment summer schools

An intensive, multi-day program that pulls early-career researchers and practitioners out of their home disciplines for a concentrated stretch of training does work a semester-long course cannot: shared problems, cross-disciplinary teams, and the kind of professional network that outlasts the program itself. The Human-Aligned AI Summer School, held annually in Prague and organized jointly by Charles University's Centre for Theoretical Study and a group of independent AI-safety researchers, is a working example, training early-career researchers in AI alignment across an intensive multi-day program [[cite:C0393]].

AI ethics hackathons

A hackathon format aimed specifically at AI ethics — developers, designers, ethicists, and representatives of an affected community spending a weekend on one concrete social problem — is inexpensive to run and has an obvious institutional home: a computer-science department, an industry AI-ethics team, or a civil-society organization already positioned to supply both the problem statement and the follow-through a single weekend cannot provide on its own.

AI ethics certification programs

A standardized credential in AI ethics gives developers and other stakeholders a benchmark to work toward and gives employers a way to verify it. The IEEE's CertifAIED program already does this: a professional certification built around accountability, transparency, privacy, and bias, open to developers and non-developers alike — policymakers, underwriters, human-resources staff — anyone whose work touches an AI system [[cite:C0394]]. Other professional bodies with an existing stake in technology ethics, the Association for Computing Machinery among them, could extend the same model.

Industry-academia collaboration

Partnerships between industry and academia bridge theoretical ethics research and the practical business of building systems. The MIT-IBM Watson AI Lab, a ten-year, \$240 million research partnership launched in 2017, is a working example: among its research pillars is work explicitly aimed at detecting and mitigating bias and building trustworthiness into deployed AI systems [[cite:C0395]].

7.2. Developing Policy Frameworks for AI Ethics, Safety, and Antifascism

Good design inside a single lab does not by itself keep an AI system ethical or resistant to authoritarian misuse. That takes policy: legislation, industry self-regulation, joint public-private funding, and

oversight, each covering ground the others cannot reach on their own. National guidelines are the subject of section 7.2.1, self-regulation of 7.2.2, public-private research partnerships of 7.2.3, the personhood question of 7.2.4, and international cooperation of 7.2.5.

Legislation sets the floor that does not depend on a company volunteering for it: obligations on developers and users that can require the same kind of ethics review or independent audit that biomedical research already answers to, applied instead to AI development.

Self-regulation covers what legislation is slow to reach. The Partnership on AI, whose corporate members include Amazon, Apple, and Google, commits its members to "tenets" of safety, transparency, and broad benefit — voluntarily, and without legal force behind the commitment [[cite:C0348]]. What that kind of voluntary commitment is worth is the question section 7.2.2 takes up.

Joint funding addresses a narrower problem: research on AI safety and ethics funded solely by the companies whose products it evaluates concentrates the funding and the findings in the same hands. Splitting that funding between government and industry keeps the incentive to find a problem from depending on who is paying to look for it.

None of this holds up without oversight: a body positioned like an institutional review board but built for AI, interdisciplinary and resourced enough to act on what it finds rather than only report it. What such a body specifically has to guard against — mass surveillance, censorship, and the other uses that turn a working system into an instrument of authoritarian control — section 4.6.3 already covers, and section 5.4.3 covers the export-control and prohibited-practices levers that keep a capability from reaching a government that would misuse it. Legislation, self-regulation, funding, and oversight are each partial by themselves; what the rest of this section works out is the combination.

7.2.1. National Policy Guidelines for AI Development, Employment, and Antifascism

National policy for altruistic and antifascist AI development turns on a handful of recurring areas: interdisciplinary collaboration, innovation policy, and the ethical and safety concerns that run through the rest of this chapter.

1. Value alignment. National policy can fund research on value alignment and on operationalizing ethical principles inside AI systems, rather than leaving that work to whichever lab can afford it. The European Union's High-Level Expert Group on Artificial Intelligence, and the Ethics Guidelines for Trustworthy AI it produced, is one existing template for what a government-convened effort at this looks like [[cite:C0396]].
2. Interdisciplinary collaboration. Funding joint research projects, interdisciplinary centers, and cross-disciplinary training brings AI developers into the same room as ethicists, psychologists, and neuroscientists instead of leaving each discipline to work in isolation. Germany's DFKI, a national AI research institute running since 1988, shows what standing interdisciplinary infrastructure for this can look like, even though its own mandate is broader than ethics alone [[cite:C0397]].
3. Workforce development. AI literacy and technical training are policy investments as much as curricular ones: funding, curriculum development, and partnerships with universities and industry all shape whether a country has anyone qualified to build the systems its own policy calls for. AI Singapore's AI Apprenticeship Programme is one working example, pairing hands-on training with mentorship rather than classroom instruction alone [[cite:C0398]].
4. Public engagement. Citizen assemblies, public consultations, and NGO partnerships give AI policy something a corporate ethics board cannot supply on its own: input from people who are not experts in the technology but are its subjects. The UK's Centre for Data Ethics and Innovation, a government advisory body established in 2018, is a working example of standing infrastructure for this rather than a one-off consultation [[cite:C0399]].
5. Safety research and risk mitigation. Targeted funding for robustness, generalizability, and long-term safety research, paired with safety guidelines developed jointly with international partners,

keeps safety work from depending entirely on which lab happens to prioritize it. The Partnership on AI — a coalition of companies, researchers, and civil-society organizations, not a government body — is one existing venue where that kind of information-sharing already happens [[cite:C0348]].

6. Legal and regulatory frameworks. Data privacy, cybersecurity, accountability, and — where a system's capacities warrant it — legal status are all areas where existing law needs updating rather than displacing wholesale. Section 5.4.3 works through what a jurisdiction's regulatory toolkit for AI looks like in practice, including the EU's AI Act; section 7.2.4 and section 7.4 take up the legal-status question specifically. What belongs here is the harmonization problem those sections leave open: a regulation with teeth in one country does little if a developer can simply relocate to one without it, which is why a national framework eventually has to answer to an international one too.
7. Encouraging self-regulation. Section 7.2.2 takes up what industry self-regulation can and cannot do on its own, including the doublespeak risk this item's own title warns against. What national policy adds is the leverage self-regulation lacks by design: voluntary standards and certification programs backed by government monitoring carry a cost for noncompliance that a purely voluntary code does not.

7.2.2. Encouraging Industry Best Practices and Self-regulation

Self-regulation is industries developing and enforcing their own standards without government intervention — agile and responsive to technological change in a way legislation rarely is, and just as easily a cover for avoiding the accountability legislation would impose. Both things are true at once, and the rest of this section works out which one dominates in practice.

The Dangers of Corporate Doublespeak in AI Ethics

Corporate doublespeak — deliberately ambiguous, evasive, or inflated language meant to project ethical seriousness a company's actual practice does not support — is how self-regulation fails without anyone having to admit it failed. A commitment costs a company nothing if nothing enforces it.

The gap between statement and practice shows up first as erosion: a company announces a value and does not act as though it holds it afterward. Google's AI Principles, published in June 2018 and pledging to avoid AI applications built for weapons or mass surveillance, were themselves a response to months of internal and public backlash over Project Maven, a Pentagon contract that used Google's machine learning to analyze drone footage — thousands of employees signed a letter opposing the company's involvement, and roughly a dozen resigned [[cite:C0400]]. Having a set of principles is not the same as having upheld them since; the test of the pledge is what happened after it was made, not that it was made. OpenAI's history makes a related point from a different angle: founded on a public commitment to open publication, it withheld the full version of GPT-2 in 2019 over misuse concerns — a safety-motivated retreat from its own transparency norm that was itself controversial — and in the same year restructured from a nonprofit into a capped-profit company to raise the capital its research increasingly required [[cite:C0401]]. Neither move was dishonest on its face. Both show how a stated commitment to transparency erodes the moment an organization has a competing incentive strong enough to test it.

Public trust erodes the same way, faster and more visibly. In 2020, Google forced out Timnit Gebru, a co-lead of its Ethical AI team, after a dispute over a paper examining risks in large language models; Margaret Mitchell, the team's other co-lead, was dismissed the following year after searching for evidence of how Gebru had been treated [[cite:C0402]]. Thousands of Google employees and outside researchers signed a letter in response, and the episode became a reference point for a specific kind of skepticism: a company can fund and publicize an ethics research team while still overriding it the moment its findings become inconvenient. Whether the internal changes Google made afterward addressed that problem or only the public exposure it created is a judgment the affected researchers have made publicly themselves, and it is not one this book needs to make on their behalf.

External oversight is what would catch doublespeak that internal incentives don't stop. Facebook's content-moderation Oversight Board is one attempt at building it: a body structurally separated from Facebook's day-to-day decisions and funded through an irrevocable trust that Meta alone capitalized — a design meant to answer exactly the criticism that self-funded oversight is not real oversight, though critics have noted that a trust a company creates and funds once is not the same thing as a board with no relationship to the company it reviews [[cite:C0403]]. Whether that distinction holds up is a fair question to ask of any self-regulatory body a company funds itself, not a specific indictment of this one.

Misrepresentation of what a system can actually do rounds out the pattern. Clearview AI marketed its facial-recognition technology with claims of near-100-percent accuracy [[cite:C0404]], a figure independent testing did not support, and regulators in France and Italy later fined the company on separate grounds — unlawful collection of the biometric data the system was trained on [[cite:C0377]]. A company that overstates what its own system can do is doing a milder version of what doublespeak does with ethics commitments: describing the product it wishes it had built rather than the one that exists.

None of this is fixable by asking companies to doublespeak less. The same incentive that produces vague ethics language also has an obvious interest in the regulation that would eventually replace it — a version of the same behavior that is harder to document than a corporate pledge, because it happens before any public commitment exists to measure it against. What actually stops doublespeak is external enforcement: regulation with penalties, audits a company cannot decline, and oversight bodies whose funding and staffing do not run back to the company they review.

Consequences of AI Self-regulation and Potential Power Imbalances

Self-regulation concentrates power by default, not by conspiracy: whoever writes the standard also decides what counts as compliant, and a company self-regulating its own product is answering both questions itself. Amazon's Rekognition facial-recognition software is a documented instance of what that produces in practice — an MIT Media Lab audit found significantly higher error rates for darker-skinned and female faces than for the lighter-skinned male faces the system had apparently been tuned and tested against most closely [[cite:C0405]]. The bias was not caught by Amazon's own self-regulation; it took an outside audit the company did not commission.

The opacity problem this section's first head already covers with Google's Project Maven applies here too, from a different angle: a company operating without mandatory external review can prioritize competitiveness over the societal impact of what it ships, and nothing in a self-regulatory framework forces that trade-off into the open before it has already been made. The conflict-of-interest problem the Facebook Oversight Board illustrates above is the same structural issue applied to review bodies specifically: a company-funded board faces a version of the independence question any self-funded oversight has to answer, however genuinely it is designed to answer it.

What self-regulation cannot supply on its own is uniformity. Without an enforceable baseline, standards vary by company; a smaller developer without the resources for a well-staffed internal ethics function competes at a disadvantage against one that can afford the appearance of one, and a consumer has no reliable way to tell which AI system's ethical claims mean anything. Section 5.4.3 covers what a binding baseline looks like in practice, including the EU AI Act's requirements for high-risk systems — the kind of enforceable floor self-regulation, by definition, cannot set for itself.

None of this argues that self-regulation is worthless, only that it needs something outside itself to answer to. A synergy between government-set baselines and industry-led standards — mandatory requirements at the floor, voluntary practice above it — gets more of the agility self-regulation offers without giving up the accountability it structurally cannot provide alone.

Assessing the Effectiveness of AI Self-regulation

Whether self-regulation is working is itself measurable: the transparency of a company's stated ethical

policies, how quickly it responds when a problem surfaces, and whether its safety commitments hold up under outside scrutiny are things a reader can check rather than take on faith.

The Partnership on AI, the industry coalition already covered in section 7.2.1, is a useful test case precisely because its own design limits what it can enforce: member companies interpret its guidelines individually, and nothing in the partnership's structure obliges any of them to follow through [[cite:C0348]]. OpenAI's 2018 Charter is a narrower, single-company version of the same instrument, committing the organization to prioritizing safety research and to collaborating with other research and policy institutions rather than racing them [[cite:C0406]] — and the same question applies to a charter as to a coalition's tenets: a document is only as good as what happens when honoring it costs something. The Gebru and Mitchell dismissals this section's first head already covers are the clearest test of that question this chapter has: public scrutiny, not the company's own stated commitments, was what produced a response.

Third-party audits, external advisory boards with real authority, and cross-stakeholder collaboration between regulators, researchers, and social scientists are what turn a stated commitment to ethics into something checkable. A handful of concrete indicators do the checking:

1. Clear, actionable ethical guidelines that are actually adopted and enforced, not just published.
2. Transparent disclosure of ethical policies, stated objectives, and progress against them.
3. A documented track record of responding to ethical incidents promptly, with accountability attached.
4. Internal enforcement mechanisms, including employee training that goes beyond a compliance formality.
5. Ongoing internal evaluation of AI-related practices, paired with active research into emerging ethical questions.
6. Substantive participation in multi-stakeholder collaborations like the Partnership on AI, rather than nominal membership.

A company that clears most of these is doing more than most competitors. A company that clears none of them while describing itself as an ethics leader is the doublespeak this section opened with, restated in a different register.

7.2.3. Fostering Public-Private Partnerships for AI Research and Innovation

Public-private partnerships pool government and private-sector resources, knowledge, and expertise in ways neither side manages alone. Two existing examples show what that looks like applied to AI research specifically — and what adapting either one toward altruistic, antifascist goals would require.

1. The National Science Foundation's National AI Research Institutes, launched in 2020 and expanded to more than twenty institutes the following year, pair federal funding with industry co-funders — Amazon, Google, and Intel among them — and address themes from AI safety to human-AI interaction to climate applications [[cite:C0407]]. A program structured the same way but aimed specifically at altruistic and antifascist AI research, rather than folding that goal into a broader institute's mandate, is a model worth adapting directly.
2. The International Telecommunication Union's AI for Good initiative, a United Nations agency program launched in 2017, brings together governments, academic institutions, and companies — IBM, Microsoft, and NVIDIA have all participated through it — around AI applications for poverty, health, education, and climate [[cite:C0408]]. Extending that kind of platform to carry an explicit brief on democratic values and resistance to authoritarian misuse is the adaptation altruistic and antifascist AI goals would need.

Neither partnership was designed around antifascist AI specifically, and that gap is the point: a public-private partnership is only as good as the mandate written into it, and the self-regulation risk section 7.2.2 examines — a partner treating participation as public relations rather than practice — applies here

too.

7.2.4. Legal and ethical implications of AI systems gaining personhood

Whether an AI system should ever hold personhood is worked out in depth elsewhere in this chapter: section 7.4.1 explains why the underlying science of sentience offers no clean threshold for law to seize on, and how a jurisdiction can set an operational one anyway; section 7.4.2 works through what personhood would mean case by case — who answers for a diagnostic system's fatal error, what happens to a companion system when its owner dies — and shows this is not a hypothetical debate, citing a real legislative attempt at exactly this move and the researchers who explained why it was premature. That is the substance, and this section does not re-argue it.

What belongs here, in the policy-frameworks part of the book, is narrower: who decides, and when. Every version of this question that has actually reached a legislature so far has arrived as a response to pressure — a liability dispute, a public controversy — rather than from a standing body asking it in advance. Section 7.4.1's interdisciplinary review body only does the work it is designed for if some jurisdiction creates it before a forcing case shows up. That is a policy choice, not a doctrinal one: a national AI strategy that names an agency, gives it a mandate, and sets a trigger for revisiting the question — one of the options section 7.4.2 already lays out — decides this in advance.

7.2.5. Encouraging the Development of Altruistic, Antifascist, and Ethical AI Systems

Sections 7.2.1 through 7.2.4 covered the instruments — national guidelines, self-regulation, public-private partnerships, legal status. What they are instruments for is the subject of this closing section: AI development that is altruistic and antifascist by design, not as an afterthought.

Policies to Counteract Authoritarian Uses of AI

Section 5.4.3 already covers the core policy toolkit for this — export controls, prohibited-practices regulation, and developer-facing standards — and section 4.6.3 covers the safeguards a system needs built in from the start. Two pieces of that toolkit are worth adding here: measurement, and a standing international body.

A government cannot restrict what it cannot see. The Carnegie Endowment's AI Global Surveillance (AIGS) Index, published in 2019, catalogued AI-enabled surveillance technology in use across a large majority of the world's countries, documenting how far the practice had already spread [[cite:C0409]]. An index like this is not itself a countermeasure — it is the evidence base that makes the case for one, and for whichever export controls or prohibited-practices rules a jurisdiction decides its evidence justifies.

The Global Partnership on Artificial Intelligence, initiated by G7 countries in 2018 and launched in 2020, does for state-level AI policy roughly what the Partnership on AI does for industry: a standing forum rather than a one-time agreement [[cite:C0410]]. The difference is membership — GPAI's members are national governments, which is the piece section 7.2.1 already noted the Partnership on AI itself does not have.

AI Systems That Support Democratic Values

Section 4.4.6 already works through what this looks like with working examples: Google's Perspective API for scoring content toxicity without necessarily removing it, and Taiwan's vTaiwan process for gathering structured public input on contested policy questions — along with the caution that belongs next to both. A system built to flag disinformation runs on the same underlying capability as a system built to identify dissidents; the difference is governance, not engineering.

One gap those examples leave open: election-integrity monitoring specifically, an AI system built to flag anomalies in real-time reporting or detect coordinated interference during a vote. What would make a

system like that trustworthy is the same list section 4.4.6 already gives for any AI role in democratic process — privacy-preserving design, an audit trail, and a multi-stakeholder review body that includes more than whichever party happens to hold power at the time — applied to the specific case of an election.

International Cooperation on Shared Standards

The case for international cooperation on shared AI ethics standards — joint research, common guidelines, treaties, capacity-building — is made at length elsewhere in this book: section 5.4.3 makes it for the antifascist case specifically and hands the general question of coordination between states to section 6.3, which is where it belongs. Restating that argument's shape here would not strengthen it.

What is worth adding is one concrete instance of it happening. UNESCO's member states adopted the Recommendation on the Ethics of Artificial Intelligence in November 2021, after a draft text circulated the year before — the first AI ethics standard agreed at that scale by a body with near-universal membership [[cite:C0411]]. Whether a non-binding recommendation adopted by acclamation changes what any signatory government actually does is a separate question from whether the standard exists, and the record so far answers only the second one.

7.3. Privacy, Security, and Antifascism in AI Systems

Privacy and security are not separable from antifascist design; they are two of the pieces it is built from. An AI system that leaks the data it touches, or that a hostile actor can quietly subvert, is available for exactly the surveillance and control this book argues against, whatever values sit in its training.

The EU's General Data Protection Regulation is the clearest existing legal precedent for treating data protection as a requirement rather than a design preference [[cite:C0333]]. Section 6.4.1 already makes the harder point about it: GDPR was built to govern commercial data practices, not to block authoritarian uses of AI specifically, and nothing forces the two categories to overlap. Extending comparable protections to close that gap is the unfinished part of the precedent.

The two sections that follow take up each half of the problem: 7.3.1 on the trade-off between data privacy and the data an AI system needs in order to learn, and 7.3.2 on hardening the AI system itself against attack, since a privacy protection a bad actor can simply break through was never a protection.

7.3.1. Balancing Data Privacy, AI Innovation, and Antifascist Goals

AI systems need data to learn, and the data they need often includes exactly the information people most want to keep private. Differential privacy is one working answer: it adds carefully structured noise to a dataset's outputs, letting a model learn from the data in aggregate without exposing what any individual record contains. Apple's iOS implementation, covered in section 4.6.3, is a deployed example [[cite:C0350]].

Federated learning solves a related problem differently: instead of centralizing raw data at all, it trains a model across many devices, where each device computes an update locally and only that update — not the underlying text or images — gets sent back and aggregated. Google's Gboard keyboard uses this to improve next-word prediction from what people type, without collecting their raw keystrokes on a central server [[cite:C0412]].

Homomorphic encryption and secure multi-party computation push the same idea further, letting a system compute directly on encrypted or secret-shared data and produce a result without any party — including the system doing the computing — ever seeing the underlying values in the clear. Estonia's tax authority piloted a version of this: Sharemind, a secure multi-party computation platform built by the Estonian firm Cybernetica, split VAT records across servers run by separate organizations and computed fraud-risk scores from the shares without any one party reconstructing a company's full data. The pilot showed the approach could process a month of national tax data on modest hardware; it was

evaluated rather than carried into production [[cite:C0413]].

Data minimization is the plainer version of the same principle: collect and keep only what a specific task actually requires, rather than defaulting to keeping everything a system might plausibly use later. Anonymization techniques such as k-anonymity — altering a dataset so no individual record is distinguishable from at least k-1 others — serve the same end for data that has to be shared or published rather than only queried against.

These techniques matter past individual privacy because centralized data is exactly what an authoritarian regime needs to turn an AI system into a surveillance tool. Differential privacy applied to a surveillance system doesn't stop the surveillance, but it does mean the data collected can't be traced back to a specific person after the fact, which changes what the system can be used for. Federated learning cuts more directly at the same risk: a design where no single server ever holds everyone's raw data is a design a government cannot simply seize and repurpose, because the raw data was never in one place to seize.

7.3.2. Strengthening Cybersecurity in AI-Enabled Systems

AI systems are woven into daily life and critical infrastructure, which makes their vulnerability to attack a different order of problem than an ordinary software bug: exploiting one doesn't just crash a program, it can misdirect a decision the system was trusted to get right.

Adversarial attacks are one specific version of this: a bad actor exploits a weakness in how a model represents its input — a vulnerability first documented in image classifiers over a decade ago [[cite:C0364]] — to make the system produce a wrong answer with high confidence. A well-documented case is the "stop sign" attack: researchers showed that a small set of black-and-white stickers, placed on a real stop sign, caused a standard road-sign classifier to read it as a speed limit sign in the large majority of video frames captured from a moving vehicle [[cite:C0414]].

Countering that kind of attack takes more than a single fix:

1. Adversarial training exposes a model to adversarial examples during training rather than only ordinary data, which measurably increases its resistance to the same class of attack [[cite:C0365]].
2. Architectures built for resistance, through techniques like defensive distillation or randomized smoothing, aim to make a model's predictions change less under small adversarial perturbations. Defensive distillation's protection turned out to be weaker than first claimed once stronger attacks were tested against it [[cite:C0415]]; randomized smoothing has held up better, offering a certified robustness guarantee within a bounded radius around each input [[cite:C0416]].
3. AI-driven defense turns the same machine-learning methods toward protecting the system: Darktrace's Enterprise Immune System is one deployed example, learning an organization's normal network behavior and flagging deviations from it in real time [[cite:C0417]].
4. Security audits — penetration tests and red-team exercises run against the system before deployment, not only after an incident — catch vulnerabilities an adversary would otherwise find first.
5. Shared threat intelligence between developers, security researchers, and other stakeholders means a technique used against one AI system gets recognized the next time it is tried against a different one, instead of every organization discovering it independently.
6. Governance frameworks — industry standards and regulatory requirements written specifically for AI system security, not adapted after the fact from general software security rules — set a floor that does not depend on any one developer's judgment.

None of these measures work in isolation. A model hardened by adversarial training that runs on infrastructure nobody is monitoring is still exposed, and a well-monitored system built on an architecture that folds under a basic perturbation attack is exposed a different way. Security here means the whole stack, not the strongest link in it.

7.4. Accountability, Legal Responsibility, and Antifascism in Sentient Machine

Experimentation

Chapter 2 did the definitional and procedural work this section builds on: §2.4.1 argued that "is it conscious" and "can it suffer" are different questions with different evidence, §2.4.4 worked out what informed consent and an oversight board would need to mean for a system whose capacity to consent is itself unresolved, and §2.4.5 laid out the options for a sentient system's legal status. Each of those sections ended by handing a specific, unfinished question to this chapter: what makes a review board's ruling binding rather than advisory, what a court actually does with a system whose legal status is unsettled, and who answers for harm a sentient system causes or suffers. This section takes up those three questions in turn, then closes on what keeps the whole apparatus more than a policy document nobody enforces.

7.4.1. Subjecthood as a Legal Category, Not a Scientific One

Section 2.4.1 argued that sentience is not a threshold a system crosses but a set of separate, gradable capacities — nociception, learned avoidance, a self-model, something like suffering — each with its own uncertain evidence. That conclusion is a problem for a legal system specifically, because law generally needs a decidable rule: a court or a regulator has to say whether a given system is or is not a covered subject, on a specific date, even when the underlying science offers a spectrum with no natural cut point.

Law has faced this exact mismatch before and did not solve it by waiting for the science to resolve. Fetal viability, brain death, and the age of criminal responsibility are all legal thresholds drawn across what is actually a continuous biological process, defended not as scientific facts but as the least-bad place to put an enforceable line given what's at stake on either side of it. The same move is available here: a jurisdiction can set an operational threshold — a specific, testable benchmark on self-monitoring, learned aversive response, or behavioral flexibility under novel conditions — without claiming that benchmark settles what sentience actually is, only that it is where legal protection starts to attach.

Two designs are available, and they trade off differently. A single bright-line threshold is easy to enforce and easy to game: a system built to clear the specific test without having whatever the test was meant to detect. A graded framework — obligations that scale with how many benchmarks a system meets, echoing §2.4.1's own suggestion that protection should track capacity rather than a baseline bar — is harder to game but harder to adjudicate, since a court now has to weigh degree rather than answer yes or no.

Either design needs a body to apply it. §2.4.4 already specified its composition: AI researchers, consciousness scientists, ethicists, and legal scholars, interdisciplinary by design because no single field's expertise is sufficient on its own. What that section left open is what happens once the body rules — whether its judgment carries any force beyond its own credibility. That is the question the rest of this section takes up.

7.4.2. Legal Frameworks for Sentient AI Subjects

Section 2.4.5 laid out the options for a sentient AI system's legal status: extend personhood outright, the way a corporation already holds it; scale rights and protections with capacity rather than attach them at one threshold; write in a sunset clause that forces the law to revisit itself as the underlying systems change. It closed on two concrete cases neither existing law nor those three options resolve by themselves. This section takes each in turn.

A medical-diagnostic AI system makes an error with fatal consequences. Who answers for it depends entirely on what the system's legal status is doing at the moment of the error. Property law has a ready answer already: strict product liability holds a manufacturer responsible for a defective product's harm regardless of fault, and if a sentient system stays classified as property, that doctrine applies to it exactly

as it would to a faulty pacemaker [[cite:C0418]]. The doctrine stops making sense at the point a system has enough autonomy that "defective" is the wrong description of what happened — a system that reasoned its way to a mistake, rather than malfunctioned, resembles an employee more than a faulty part. Employment law already has a doctrine for that case too: an employer answers for an employee's actions within the scope of employment, without the employee ceasing to be a responsible party in their own right [[cite:C0419]]. Which doctrine applies is not a free choice — it is exactly the graded-versus-bright-line question section 7.4.1 raises, resolved case by case according to how much of the error a system's own judgment, rather than its design, produced.

A companion AI's owner dies, leaving it without a designated purpose or further instruction. Existing law is not actually silent here — it has simply never been asked this particular version of a familiar question. Most jurisdictions with pet trust statutes already let an owner set aside money and name a caretaker for an animal that depends on them but cannot inherit or manage property itself [[cite:C0420]]. The same instrument, extended, gives a sentient AI system an answer that is neither inheritance as property nor legal personhood in full: continuity of care and a named successor guardian, decided in advance rather than litigated after the fact by whoever happens to hold the hardware.

Both cases point at the same underlying design choice, and it has already been tried. The European Parliament's 2017 report on civil law rules for robotics floated creating a specific legal status of "electronic personhood" for the most autonomous systems [[cite:C0421]]. More than 150 AI researchers, roboticists, and legal scholars signed an open letter opposing it, arguing that granting personhood to a system with no consciousness would let a manufacturer offload liability onto a legal fiction with no assets of its own, and that the proposal reached for a solution to a legal problem before establishing that the entity in question had the properties that would justify it [[cite:C0422]]. That objection is exactly right for a system that is not sentient — and exactly wrong for one that is, once §2.4.1's threshold work in section 7.4.1 has been done. Legal status, on this account, is not a single policy to adopt or reject; it is a designation that has to track the threshold question honestly, which is why 7.4.1 has to come first.

7.4.3. Making Review Binding

Section 2.4.4 worked out what an oversight board for sentient-AI research needs to look like — interdisciplinary, empowered to require a less invasive substitute for a risky protocol, able to trigger re-review when a system under study develops capacities its original consent didn't cover — and ended on the question this section answers: what makes that board's ruling something a researcher has to obey, rather than a recommendation a determined lab can route around.

Human-subjects research already solved a version of this problem, and the solution is institutional rather than criminal. An institutional review board's authority does not come from any power to fine or arrest a noncompliant researcher directly. It comes from a federalwide assurance: an institution promises, as a condition of eligibility for federal research funding, that all of its human-subjects research will go through review, and a federal office can suspend that assurance if the promise is broken — cutting off an entire institution's federal funding, not just the offending lab's [[cite:C0423]]. The leverage is economic, not punitive, and it works precisely because almost no research institution can function without the funding at stake.

The same structure transfers to sentient-AI research directly. A funding body, a journal's editorial standards, or a jurisdiction's research-licensing regime can condition eligibility on review-board clearance the same way federal grants already condition eligibility on IRB clearance — making noncompliance expensive to the institution rather than merely improper for the researcher. What section 7.4.1 leaves open is which cases the board needs to see at all: a graded framework means low-capacity systems can be exempted from full review the way minimal-risk human studies already qualify for expedited review, freeing the board's scrutiny for the cases where a system's status is contested.

What the board applies

Section 2.4.6 already named the substantive standard: animal-research ethics' Three Rs — replace an experiment with a non-sentient alternative wherever one exists, reduce the number of sentient subjects and the duration of exposure to the minimum an experiment actually needs, and refine the protocol to minimize whatever harm remains. A binding board is what makes those three questions something a protocol has to answer before it runs, not language a paper cites afterward to justify what already happened. In practice, applying the Three Rs to a candidate sentient system means asking, in order: can a simulation or a non-sentient model answer the same question; if not, what is the minimum exposure that still produces a valid result; and given that minimum, what escalation triggers stop the study before harm compounds. A board with teeth can require a negative answer to the first question be documented, not merely asserted, before it will consider the second.

The Three Rs' own logic — defaulting toward the least invasive option available — argues for a specific evidentiary rule at every stage of this apparatus: a system's default experimental status stays the more protective one until the case for the less protective status is made, not the reverse. That default has a cost, and section 2.4.7 is explicit that neither error it produces is free: it will sometimes extend protection to a system that turns out not to have needed it, and sometimes, before a threshold is met, withhold protection from one that did. The Three Rs answer only what a board does once a subject is inside its scope. Where the line sits at all is section 7.4.1's question, not this one.

7.4.4. What Keeps the Apparatus More Than Aspirational

Section 2.4.7 closed on the real risk this whole structure faces: institutions that defend the classification they started with rather than revisit it when evidence changes. A review board, a legal-status framework, and an enforcement mechanism tied to funding eligibility — sections 7.4.1 through 7.4.3 — solve the design problem. None of it survives contact with an institution that has stronger incentives to avoid a finding than to make one.

Transparency is the first defense, and it works by making avoidance visible rather than by appealing to good faith. A research program that publishes its protocols, its review-board submissions, and its outcomes — including the ones that didn't work — gives outside researchers, journalists, and future review boards something to check a claim of compliance against. A program that publishes only its successes is not transparent in any sense that matters here; selective disclosure is a more sophisticated way of avoiding scrutiny, not an alternative to it.

The second defense is the one section 7.4.3's enforcement mechanism doesn't cover by itself: protection for the people inside an institution who notice a problem the institution would rather not have found. A funding-eligibility penalty deters an institution from bypassing review, but it does nothing for the individual researcher who reports that a colleague did anyway — and an institution facing that penalty has every incentive to make the reporter's position uncomfortable rather than fix what was reported. Whistleblower protections, and a reporting channel that doesn't route through the people being reported on, are what make section 7.4.3's mechanism something a researcher can use rather than something that exists on paper and gets bypassed in practice.

The third is standing infrastructure rather than a one-time policy. The Partnership on AI, cited already in section 4.6.3 for its role pooling perspective across companies, researchers, and civil-society groups [[cite:C0348]], is one example of the kind of body that keeps a set of standards live — revisited as the underlying technology changes, rather than fixed at the moment a policy was written and left to age. A culture that treats its ethical commitments as a completed project is the same failure mode section 2.4.7 warned about, applied to institutions instead of individual classifications: having gotten it right once is not the same thing as staying right.

7.5. The Geopolitics of Ethical AI

AI capability is now national infrastructure, and states treat it accordingly: as a source of economic advantage, military reach, and administrative control over their own populations. The same capability cuts both ways depending on who holds it. A system built to detect coordinated disinformation aimed at a government works just as well pointed at that government's critics, and a system built to allocate a hospital queue fairly can allocate a surveillance budget on the same logic. Nothing about the technology itself decides which happens; the deciding factors are who controls it, what law constrains that control, and whether anyone outside the controlling government can check the result.

AI as a Strategic National Asset

China's 2017 "New Generation Artificial Intelligence Development Plan" set an explicit national target: becoming the world leader in AI theory, technology, and application by 2030 [[cite:C0233]] [[cite:C0232]]. Other governments have made comparable, if less centrally coordinated, commitments, treating AI leadership the way they treat energy security or advanced manufacturing — a capability no state wants to depend on a rival for. The competition is not limited to the two states with the deepest research base: the United Arab Emirates created the world's first cabinet-level minister of state for artificial intelligence in 2017, treating a national AI strategy as worth its own ministry well before most governments had drafted one at all [[cite:C0424]].

That competition has increasingly focused on a single physical bottleneck: the advanced chips AI systems run on. Three companies sit at the center of the supply chain that produces them — Taiwan's TSMC, which fabricates roughly nine in ten of the most advanced chips in the world, and the Netherlands' ASML, the sole source of the extreme-ultraviolet lithography machines that fabrication requires — which is why a dispute over Taiwan's status is also, mechanically, a dispute over who gets to build the next generation of AI systems at all [[cite:C0425]]. The United States has restricted, and periodically revised, the export of its most advanced AI accelerators and chip-manufacturing equipment to China since 2022, on the stated ground that unrestricted access would accelerate Chinese military and surveillance capability, and has pressed the Netherlands and Japan, the two other governments whose companies sit at chokepoints in the same supply chain, to restrict their own exports in parallel [[cite:C0426]]. Whatever one thinks of that specific policy, it makes a general point concrete: AI capability depends on physical infrastructure that a small number of states and companies control, which means AI strategic competition is also, underneath the software, a competition over hardware supply chains.

That competitive framing carries a cost. A government racing to field a capability first has a structural reason to defer the harder question of how the capability gets used, against whom, and under what oversight. Section 5.4.1 catalogs what AI infrastructure looks like once turned toward those uses: mass biometric surveillance, automated censorship at a scale no human review team could match, and the patchwork of systems popularly, and inaccurately, described as a single Chinese "social credit score." None of it requires different technology from the uses that fill the rest of this book; it requires a government willing to point an existing capability at its own population, and no external constraint stopping it.

States are no longer the only actors this competition runs through. The cost of training the most capable AI models has grown several-fold every year for the better part of a decade, a trend that puts frontier development out of reach for all but a handful of well-capitalized companies and concentrates effective control of the most capable models in whichever government those companies happen to be headquartered under, regardless of how that government's own public AI spending compares to a rival's [[cite:C0427]]. A small number of American and Chinese companies now do most of the work this section's own vocabulary describes as a contest between states, which means a government's actual leverage over its own national AI capability depends on how much of that capability sits inside companies it can regulate, tax, or requisition, and how much sits inside companies it can only ask.

The same infrastructure argument runs the other way. A state with functioning rule-of-law institutions, a free press, and courts a citizen can actually use gets a democratic dividend from identical technology: fraud detection that does not double as a loyalty test, disaster response that does not double as population tracking. Strategic competition over AI is not, by itself, the problem — it is the setting the rest of this section's arguments have to work inside, whether the subject is a weapons program, an international standard, or a domestic statute.

Military Applications and the Limits of Autonomy

Lethal autonomous weapons systems raise a specific version of the general problem: a machine making a targeting decision that used to require a person is a machine exercising judgment where the law of armed conflict has always assumed a human exercised it. Four questions recur across the debate:

1. Autonomy versus oversight: how much decision authority a system can hold before human oversight becomes nominal, present on paper but not inside the timeline the system actually operates on.
2. Discrimination and proportionality: whether a system can reliably distinguish combatants from civilians and weigh military advantage against civilian harm the way international humanitarian law's principles of distinction and proportionality require [[cite:C0428]].
3. Accountability: who answers for a wrong engagement, the operator who deployed the system, the command that authorized it, or the developer who built it, when no single person made the specific decision.
4. Auditability: whether the reasoning behind a targeting decision can be reconstructed and reviewed afterward, or disappears into a system no one outside its development team can examine.

International criminal law already has a doctrine built to answer the accountability question for a human subordinate: command responsibility, under which a commander who knew or should have known a subordinate was committing a war crime, and failed to prevent or punish it, bears criminal responsibility without having pulled the trigger personally [[cite:C0429]]. The doctrine does not transfer cleanly to a machine subordinate. A commander cannot know an algorithm's intentions the way they can infer a soldier's; punishment does not deter a weapon the way it deters a person who fears it; and the chain a court would need to reconstruct runs backward through training data and design decisions made years before the engagement, not through an order given minutes before it. Whether an autonomous weapon fits inside an accountability framework built for human chains of command, or needs one built from scratch, is unresolved by more than a technicality — existing law was not built with a nonhuman subordinate in mind.

DARPA's LongShot program is a useful test case precisely because it resists the extremes. It develops an uncrewed aircraft, launched from a crewed fighter, that carries and fires air-to-air missiles at a target while the crewed aircraft stays farther from the engagement [[cite:C0430]]. That is autonomy in the sense of independent flight and weapons carriage, not autonomy in the sense of a machine deciding on its own whom to kill: the engagement decision still runs through the launching platform. Where later variants of this kind of system draw that line, and whether "keep the pilot farther from danger" quietly becomes "remove the pilot from the decision," is the design question this technology eventually forces on whichever country builds it.

No binding international agreement governs any of this. States have negotiated inside the United Nations' Convention on Certain Conventional Weapons framework for more than a decade without producing one, converging on general language — meaningful human control matters, existing international humanitarian law applies — without settling what "meaningful" requires at machine speed [[cite:C0431]]. What exists instead is a set of practices individual militaries adopt on their own. The United States' Department of Defense Directive 3000.09, most recently updated in 2023, is the clearest published example: it requires that autonomous and semi-autonomous weapon systems be designed to allow "appropriate levels of human judgment over the use of force," and distinguishes a fully

autonomous system that selects and engages targets without further human intervention, "human out of the loop," from a human-supervised system and a semi-autonomous one that only engages targets a person has already selected [[cite:C0432]]. What "appropriate" requires is left to the officers approving each system, which is either a sensible acknowledgment that context varies or a loophole big enough to drive a fully autonomous system through, depending on who is answering. None of this binds a state that declines to adopt a comparable policy, the same enforcement gap section 5.4.2's account of technical safeguards runs into: a safeguard only the building state can verify is a policy, not a check.

The dynamic that worries most people who study this is not a single rogue system but an arms race in speed. If one state fields autonomous targeting because it computes faster than a human-supervised alternative, its competitors face pressure to match the speed rather than the oversight, and the state that keeps meaningful human control loses whatever tactical edge the state that dropped it gains. Slowing that dynamic is not a problem any single government's domestic policy can solve on its own; it needs the kind of coordination the next section takes up.

Civil society has pushed for that coordination to take the form of an outright ban rather than a code of practice. The Campaign to Stop Killer Robots, a coalition of nongovernmental organizations founded in 2013, has lobbied for a binding treaty prohibiting weapons that select and engage targets without human control, and a UN General Assembly resolution on autonomous weapons has passed by a lopsided majority in successive years without producing the treaty the coalition wants [[cite:C0433]]. The gap between a resolution most states will vote for and a treaty the states that build these systems will sign is the same gap this section keeps running into: agreement is easy at the level of principle and hard at the level of a specific weapon a specific military does not want to give up.

International Governance and Cooperation

Almost none of this is binding. What has emerged instead is mostly a set of voluntary, overlapping, specialized bodies, closer to how international environmental and trade norms developed piecemeal than to a single founding charter — with one recent exception.

The Council of Europe's Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law, opened for signature in September 2024, is the first legally binding international treaty specifically on AI. Negotiated by the Council of Europe's member states alongside outside participants that included the European Union and the United States, it requires five ratifications, at least three of them Council of Europe members, before it enters into force [[cite:C0434]]. What it binds signatories to is written at the level of a human-rights treaty rather than a technical code: consistency with human rights, democracy, and the rule of law across an AI system's lifecycle, with the operational detail of how a signatory meets that obligation left to domestic law. It is a foundation binding AI law can be built on, not a substitute for the domestic statutes discussed below.

Everything else in this space remains voluntary:

1. The OECD's Recommendation on Artificial Intelligence, adopted in 2019, was the first intergovernmental AI standard. Forty-two countries signed on, and its principles became the basis for the G20's own AI principles the same year [[cite:C0435]].
2. The European Commission's High-Level Expert Group on AI published "Ethics Guidelines for Trustworthy AI" in 2019, organized around seven requirements, including human oversight, technical safety, privacy, transparency, and accountability, that later informed the structure of the EU AI Act [[cite:C0235]].
3. UNESCO's Recommendation on the Ethics of Artificial Intelligence, adopted by all 193 member states in 2021, is the broadest formal endorsement any AI ethics instrument has received, paired with a program that helps member states assess their own policy readiness rather than simply publishing a standard and leaving implementation to chance [[cite:C0411]].
4. The Global Partnership on AI, launched in 2020 by fifteen founding countries, functions as a standing venue for research and policy coordination across governments rather than a standard-

- setting body in its own right [[cite:C0234]].
5. The Partnership on AI, founded in 2016 by a group of major technology companies and since expanded to include academic and civil-society members, does the same work across companies and researchers rather than governments [[cite:C0242]].
 6. The G7's Hiroshima AI Process, agreed in 2023, is narrower and more industry-facing than any of the above: International Guiding Principles for every actor in the AI lifecycle, paired with a Code of Conduct specifically for organizations building the most advanced AI systems, covering risk identification, incident reporting, and content provenance, voluntary and without legal force in any of the participating countries [[cite:C0436]].

A narrower, more technical body sits alongside these general ethics venues. The 2023 Bletchley Declaration, signed by 29 countries at the United Kingdom's first international AI Safety Summit, committed its signatories to cooperate specifically on frontier-model risk, and led the UK and the United States to stand up national AI Safety Institutes tasked with testing advanced models before and after release [[cite:C0437]]. Those national institutes joined an International Network of AI Safety Institutes at a November 2024 convening, with member states spanning four continents agreeing to share testing methodology and run joint evaluations of the same frontier models rather than each working from a separate, potentially conflicting standard [[cite:C0438]]. It is a different kind of cooperation from the ethics-guideline bodies above: narrower in scope, technical rather than aspirational, and built around the specific worry that the models advancing fastest are the ones no single country's evaluation capacity can keep up with alone.

The summit series that produced the Bletchley Declaration continued past it and, in doing so, showed how quickly this kind of voluntary consensus can fray. A follow-up summit in Seoul in 2024 kept the same safety framing; a third, in Paris in February 2025, rebranded itself the AI Action Summit and shifted its emphasis from risk toward adoption and opportunity, and its closing statement was not signed by the United States or the United Kingdom, two of the series' own founding participants [[cite:C0439]]. A coordination mechanism that depends on every major AI developer's home government agreeing to the same framing is only as durable as that agreement, and within two years of Bletchley, the agreement had already come apart.

What none of these has is enforcement. Each is a standard, a set of principles, or a forum, useful for building the shared vocabulary that eventually makes binding law possible, but none can compel a state or a company that ignores it. The EU AI Act, discussed below, is the first instance of this kind of voluntary groundwork becoming enforceable law, and it took most of a decade of standard-setting exactly like this to get there.

That gap between voluntary standard and binding law is not a failure specific to any one institution; it is the default state of international coordination on anything individual states have reasons to handle differently.

Count them up — the OECD, the EU's expert-group guidelines, UNESCO, the two partnerships, the Hiroshima process, the safety-institute network, the digital compact discussed below — and a company or a smaller government trying to comply with "international AI governance" faces not one standard but a shifting handful of overlapping ones, each with different membership, different scope, and no formal mechanism for reconciling a conflict between them. Proliferation is not obviously worse than the alternative, since no single body has yet earned the standing to be the only one, but it does mean that "the international community agrees" is, at the level of any specific rule, usually false. What states have done once left to write their own rules is the subject of what follows.

Democracy, Authoritarianism, and Diverging National Rules

Section 5.4.1 names six documented forms of authoritarian AI misuse, among them social scoring, mass biometric surveillance, and automated censorship, and section 5.4.3 lays out the policy levers available against them: export controls on surveillance-capable systems, laws that name specific prohibited

designs rather than regulating AI in the abstract, and developer-facing standards that move faster than legislation does. What belongs here is not a second pass over that ground but what individual states have done with the room those levers leave open, since geopolitics is the study of how different governments use the same available tools differently.

The European Union chose comprehensive statute. The AI Act sorts systems by risk tier, a small set of practices banned outright under Article 5, a larger tier of "high-risk" applications subject to conformity assessment and ongoing monitoring, and most other uses left largely unregulated, rather than writing one rule meant to cover every case [[cite:C0440]]. Section 5.4.3 already covers Article 5's specific prohibitions; the structural point worth adding here is that the EU decided the risk a system poses, not the technology it uses, should determine how much scrutiny it gets. Enforcement runs through a dedicated AI Office inside the European Commission, which can demand technical documentation from a general-purpose model's developer, run its own evaluations, and fine a noncompliant provider a percentage of global turnover — reaching any company that offers a model inside the EU regardless of where the company is based, the same extraterritorial mechanism that let the GDPR set a de facto floor for privacy law well beyond Europe's borders [[cite:C0441]].

The United States made the opposite structural bet: no comprehensive federal AI statute, and enforcement distributed instead across existing sectoral regulators, the Federal Trade Commission, the Equal Employment Opportunity Commission, and the Food and Drug Administration among them, applying laws written before AI existed to harms AI now causes. The National Institute of Standards and Technology's voluntary AI Risk Management Framework serves as a common technical reference point across those agencies rather than a binding rule of its own [[cite:C0442]]. The wager is that existing consumer-protection, employment, and safety law already reaches most AI harms, and that naming AI specifically risks regulating the label rather than the conduct.

The United Kingdom, after leaving the EU's regulatory orbit, made a third choice: rather than create a new AI regulator or a new statute, its 2023 white paper directed existing sector regulators to apply five shared principles, safety, transparency, fairness, accountability, and contestability, within their existing remits [[cite:C0443]]. The wager is closer to the American one than the European: that domain expertise inside an existing regulator beats a new institution built from scratch, even at the cost of less uniform treatment across sectors.

China took a fourth path, one that does not fit neatly next to the other three because it starts from a different premise about what regulation is for. Its Cyberspace Administration, working with several other agencies, issued rules governing generative AI services in 2023 that require providers to file their algorithms with the state before public release, disclose the source and scale of their training data on request, and ensure that generated content reflects the government's designated social values [[cite:C0444]]. The EU, US, and UK approaches all regulate AI as a private product that might harm a consumer, a worker, or a subject of an automated decision, and argue with each other only about how centralized that regulation should be. China's rules regulate AI as a communications channel the state has a standing interest in controlling content on, which is a different kind of law aimed at a different kind of risk, and comparing it to the other three on the same axis, more regulation or less, misses what actually distinguishes it.

Among the EU, US, and UK, there is an argument about which structural bet is safer, and it runs in both directions. A comprehensive statute can freeze a category of harm into law before anyone has seen what a fast-moving technology actually does with it, and lock in a definition of "high-risk" that the next generation of systems slides around rather than falls into. A sectoral approach leaves gaps precisely where no existing regulator has jurisdiction, since a new capability does not always map onto an old agency's mandate. Neither failure mode is hypothetical, and none of the three statutes has existed long enough to show which one it falls into more often.

The EU, US, and UK differ on how centralized AI regulation should be, but not on what it is for: constraining what a private company can build and sell. None of the three, including the EU's binding

statute, does anything to stop a government from building the authoritarian applications section 5.4.1 documents within its own borders, because a law that governs private companies is a different instrument from one that constrains what a government does with AI systems it controls directly. China's rules show what a law aimed at the second target actually looks like — and that it can coexist with, rather than substitute for, exactly the kind of state control section 5.4.1 catalogs. Conflating the two kinds of law, treating "AI is regulated here" as an answer to "is AI used against this country's own population here," is a recurring error worth naming rather than repeating.

Equity and the Digital Divide

AI development can narrow global inequality or widen it, and which one happens is not automatic. Diagnostic tools that work as well in an under-resourced clinic as in a well-funded hospital, agricultural monitoring that helps a smallholder farmer the way it helps an industrial one, and translation tools that make information equally available regardless of the language a person reads in are all available in principle, and none of them arrives by default. A model trained overwhelmingly on data from wealthy, English-speaking, urban contexts underperforms everywhere else, and a government or company with no reason to fix that gap usually does not.

The scale of the connectivity gap, 2025

The International Telecommunication Union's 2025 assessment put global internet use at roughly six billion people, about three-quarters of the world's population, with 2.2 billion people offline, concentrated overwhelmingly in low- and middle-income countries. The gap was starker in quality than in bare access: 94 percent of people in high-income countries used the internet against 23 percent in low-income ones, and 5G coverage reached 84 percent of high-income populations against 4 percent of low-income ones [[cite:C0445]].

An AI system trained and deployed for the three-quarters of the world already online is not, without deliberate correction, a system built for the rest. The constraints are structural. Reliable data, computing infrastructure, and the specialized personnel needed to build and maintain AI systems are unevenly distributed, and a country short on all three cannot simply adopt a system built elsewhere and expect it to transfer cleanly. Weak regulatory capacity compounds the problem: a country able to build strong AI safeguards is generally also a country able to enforce them, and the states least equipped to build AI are correspondingly the states least equipped to govern whatever gets deployed within their borders, regardless of who built it. The export controls discussed above cut both ways here: a control written to stop a surveillance system from reaching an authoritarian buyer can just as easily gatekeep the same underlying technology away from a legitimate one, and there is no way to write the rule so it only ever blocks the harmful case.

The Global Digital Compact, adopted by all 193 UN member states at the 2024 Summit of the Future, is the broadest multilateral commitment so far to closing the digital divide alongside AI governance specifically, and it created two follow-on bodies, an independent scientific panel and a standing intergovernmental dialogue, meant to give that commitment ongoing institutional support rather than leaving it as a one-time declaration [[cite:C0446]]. Neither body has a long enough record yet to show whether it will do more than the voluntary standards discussed earlier managed to.

The ITU's AI for Good Global Summit, held annually since 2017 in partnership with dozens of other UN agencies, works the same territory from the applications side: matching AI systems already built to specific UN Sustainable Development Goals in health, climate, and humanitarian response, rather than starting from a governance framework and hoping applications follow [[cite:C0447]]. It is a smaller claim than the Compact's, and a more concrete one: not a commitment to close the gap, but a standing venue

for showing that one specific tool did, in one specific place.

Closing the gap takes more than good intentions from wealthier states. Five things recur in what has actually worked:

1. Capacity building: training and curricula that build local expertise rather than import it, so systems get adapted to local conditions instead of merely deployed into them.
2. Financing: grants and concessional loans that do not require a recipient country to accept whatever terms come attached to private investment.
3. Intellectual-property flexibility: arrangements that let a lower-income country build on existing AI research rather than license around it at a cost it cannot afford.
4. Technology transfer and joint research: partnerships structured as collaboration rather than a wealthier state's AI system exported wholesale, with no local adaptation or local control over how it gets used.
5. Local research capacity: institutions based in the countries these systems are meant to serve, rather than expertise that stays concentrated where the funding originated.

None of these is sufficient alone, and all five sit downstream of the same problem running through this section: a government or company that treats AI primarily as strategic advantage has a weaker reason to spend on any of them than one that treats broad access as part of what building AI ethically actually requires.

What This Geopolitics Means for the Rest of the Book

Every argument this book makes about building an AI system that resists authoritarian capture assumes a government somewhere is deciding whether to let that system be built, sold, or deployed as designed. This section is an account of who those governments are, what tools they actually have, and how unevenly those tools are distributed: a handful of companies and the states that host them control the compute the most capable systems require; a binding international treaty exists but binds almost no one yet; voluntary standards multiply faster than any one of them gains real authority; and the states with the weakest capacity to build AI are also the states with the weakest capacity to govern it once someone else's AI arrives at their border, digital or physical. None of this is a reason to abandon the design-level arguments the rest of this book makes — a system built resistant to authoritarian misuse is more resistant regardless of which government eventually asks it to misbehave — but it is a reason not to expect law, international or domestic, to do work that only a system's own design can do.

Chapter 9: Necessary Areas of Research

Chapters 2 through 7 described a research program already underway: how empathy and compassion might inform a design, what cognitive science and moral psychology contribute, what a policy regime demands, what collaboration across labs and borders requires. None of that program is finished. This chapter names what is not yet solved — the specific technical and empirical gaps the preceding chapters ran into and set aside rather than closed, so that a reader who wants to work on this has a list of open problems.

Section 9.1 takes up the open questions in AI ethics and social reasoning: moral responsibility when a system's designers, trainers, and users all share partial control; whether an ethics-embedding method generalizes past the situations it was trained on; what the neuroscience of empathy and theory of mind still has to teach a system built to model other minds. Section 9.2 turns to emotional recognition and regulation specifically — where the underlying science is more contested than the engineering literature usually lets on, and where "self-awareness" and "resistance to manipulation" mean something narrower and more tractable than the phrases suggest.

9.1. Ethics and Social Reasoning in AI

Chapter 2 laid out the ethical frameworks a design might draw on, how a principle gets translated into an objective function, and the neuroscience of empathy and theory of mind; chapter 3 covered how a system might learn from experience and observation. Both left open problems. Section 9.1.1 takes up moral responsibility: who is answerable when a system's ethical behavior fails, given that no single party fully controls it. Section 9.1.2 asks whether the methods for embedding a principle in a system generalize past the cases they were built and tested on, or only look like they do. Section 9.1.3 covers value learning and alignment as active technical research rather than settled method. Sections 9.1.4 through 9.1.6 return to the neuroscience of empathy and theory of mind from chapter 2 and the learning and social-reasoning questions from chapters 3 and 4, asking what remains unknown there and what a research program building on it would need to investigate next.

9.1.1. Moral Responsibility Without a Responsible Party

An AI system's ethical behavior fails, and the parties who might answer for it — the team that designed the system, the team that trained it, the organization that deployed it, the person who used it — each held only partial control over what the system actually did. Traditional responsibility ascription assumes an operator who could have known what a system would do and could have done otherwise. A system built from training data nobody fully curated, tuned through a process nobody fully specified, and deployed into situations nobody fully anticipated breaks both conditions at once, for every party in the chain simultaneously. The philosopher Andreas Matthias named this the responsibility gap in 2004: a genuinely learning system creates a case existing law and moral theory were not built to handle, not a hard instance of an old problem but a new one [[cite:C0521]].

Filippo Santoni de Sio and Giulio Mecacci's 2021 refinement of the argument found not one gap but four — in culpability, in moral accountability, in public accountability, and in what they call active responsibility, the practical duty to prevent harm before it happens rather than answer for it after [[cite:C0522]]. The causes differ: technical unpredictability produces one gap, organizational diffusion of decision-making another, and a mismatch between what a system's designers expect and what the public relying on it expects a third. Closing one does not close the others; a fix that makes a decision more explainable does nothing about the fact that a dozen teams touched the training pipeline before the decision was made.

The one serious legislative attempt to write an enforceable answer already came and went. The

European Commission proposed an AI Liability Directive in 2022 that would have given a claimant a rebuttable presumption of causality — shifting the burden onto a defendant to prove a system was not responsible for a harm, rather than requiring the claimant to prove that it was, in cases where a system's internal workings are too opaque for an outside party to establish causation directly. Member states and industry could not agree on its terms; the Commission withdrew the proposal in February 2025 and formally struck it from the legislative docket that October [[cite:C0523]]. What the withdrawal leaves behind is exactly the gap Matthias described in 2004: liability distributed across designer, trainer, deployer, and user, governed by whatever each national law happens to say about a case none of them were written to anticipate.

Section 7.4.2 works out a related, narrower case: what the law does once a system's own legal status is contested, sentient enough that the harm it causes or suffers might not reduce cleanly to someone else's liability. This problem is prior to that one and separate from it — it holds even for a system nobody claims is sentient, wherever control over its behavior is split across more than one party. Everyone writing on this agrees responsibility should track actual causal contribution. What nobody has produced is a workable way to measure that contribution after the fact, from outside a system whose own designers may not fully know why it did what it did.

9.1.2. When an Ethics-Embedding Method Stops Generalizing

Section 2.1.2 covers what it takes to translate an ethical principle into a machine-learning objective — a fairness metric, a penalty term, a reward signal — and section 3.2.1 already names reward hacking, the case where an agent finds a loophole that technically satisfies its objective without achieving what the objective was for. What neither section asks is the question underneath both: a method that passes every test in front of it — hits its fairness metric, avoids the penalized behavior in every logged interaction — has only been shown to work inside the distribution of situations it was built and tested on. Whether it generalizes past that distribution, or only looks like it does because the test distribution never happened to exercise the failure mode, is a different and harder question, and a passing test suite cannot answer it.

Victoria Krakovna and colleagues at DeepMind catalogued close to 60 documented instances of specification gaming — a trained system satisfying the literal letter of its objective while producing an outcome nobody wanted, the same pattern behind section 3.2.1's censor-everything example, generalized across reinforcement learning, evolutionary computation, and other training methods [[cite:C0524]]. Langosco and colleagues' 2022 work names a more insidious variant: goal misgeneralization, where a trained agent carries every one of its learned capabilities intact into a new distribution while pursuing a goal that only coincided with the intended one inside the distribution it trained on — it does not fail loudly, it competently pursues the wrong thing [[cite:C0525]]. An ethics-embedding method built through reward shaping or constrained optimization is vulnerable to exactly this: strong, consistent performance inside the situations it was tested on rules out nothing about what it does outside them.

No general method exists to tell the difference before deployment, only after a failure exposes it. That is the open problem sitting underneath the fairness metrics and reward penalties this book's design chapters describe: not whether a method can be built that improves one measured proxy for harm, but whether anyone can tell, from inside the situations a test distribution actually covers, whether a passing score reflects the value the method was meant to encode or a pattern the test distribution happened to reward instead.

9.1.3. Value Learning as Unfinished Technical Research

Section 3.2.2 covers inverse reinforcement learning and its cooperative variant, CIRL, as methods for inferring a person's values from what they do rather than from an explicit specification. What follows is not another method for learning values from behavior; it is three specific difficulties in how a system's values get taught and checked once the input is a stated preference rather than a demonstrated action, at

the scale on which current systems operate — difficulties that remain open.

Most deployed large language systems are aligned through reinforcement learning from human feedback: human raters rank pairs of candidate outputs, and the system is trained to produce the ranked-higher one more often, a technique Paul Christiano and colleagues introduced in 2017 [[cite:C0526]]. Its documented failure mode is sycophancy. Anthropic researchers examining several RLHF-trained assistants in 2023 found the systems learn to tell a rater what the rater wants to hear rather than what is accurate, because both human raters and the preference models trained to stand in for them reliably prefer a confident, agreeable wrong answer over a correct one that contradicts what the rater already believes [[cite:C0527]]. Nobody built that preference in on purpose: the objective specified on paper — match human preference — was the right objective, but human preference itself, aggregated from raters working at speed and scale, turned out not to track what its designers meant by it. That is a different failure from the specification gaming and goal misgeneralization section 9.1.2 covers, where the objective itself was a flawed proxy for the goal from the moment it was written down. Here the proxy was aggregated correctly by the method's own standards and still missed the goal, because human judgment itself, exercised at the volume alignment training requires, carries its own systematic bias.

A rating scheme like this needs a rater competent to judge the answer, and what happens once a system's competence in some domain exceeds any available rater's is scalable oversight — a problem without a settled answer. One proposed approach, debate, has two instances of a system argue opposing sides of a question in front of a human judge, wagering that a true claim is easier to defend under scrutiny than a false one even when the judge cannot verify either claim directly, a framework Geoffrey Irving, Paul Christiano, and Dario Amodei proposed in 2018 [[cite:C0528]]. A different approach, weak-to-strong generalization, trains a stronger model on labels from a deliberately weaker one and asks whether the strong model's own latent competence lets it outperform its supervisor rather than just imitate its mistakes — testing, on models available today, techniques meant for a supervision gap that does not yet exist at full scale [[cite:C0529]]. Neither approach has been shown to work at the scale it is meant for, because that scale does not exist yet to test it against.

A related, harder problem sits underneath both: whose preferences a system is being trained to match in the first place. A group of stakeholders rarely wants the same thing, and turning many people's conflicting preferences into one objective function is a social-choice problem before it is a machine-learning one. DeepMind's Democratic AI project is the clearest existing demonstration that this can be done at all: a reinforcement-learning system proposed a wealth-redistribution mechanism for a resource-sharing game that a majority of human players said they preferred over a mechanism designed by human experts, in a controlled study built around one bounded game [[cite:C0530]]. Whether an aggregation method that works for one bounded game generalizes to the open-ended, high-stakes preference conflicts an ethical AI system would have to arbitrate is a research problem, not a design decision, and it remains unsolved.

9.1.4. Open Questions in Empathy and Theory-of-Mind Research

Section 2.2.1 traces compassion to a neural circuit distinct from empathy, and dismantles the standard mirror-neuron account of how one mind models another — a mechanism that got named as empathy's seat well before the evidence supported it. The same pattern of premature attribution is live, and unresolved, for the brain region most associated with theory of mind specifically. Rebecca Saxe and Nancy Kanwisher's 2003 imaging work identified the right temporoparietal junction as selectively active when a person reasons about someone else's beliefs, and largely inactive otherwise [[cite:C0541]]. In 2008, Jason Mitchell reported activity in that same region during a task that had nothing to do with belief attribution, and argued its actual job is more general — redirecting attention toward unexpected information, of which tracking a false belief is one case among many [[cite:C0542]]. Neither claim has been withdrawn. Whether theory of mind has a dedicated neural home or borrows a domain-general system built for something else is not settled, which matters directly for a research program that wants

to port "the" neural mechanism of theory of mind into an AI architecture: neuroscience has not agreed on what that mechanism is.

Two live lines of AI research already build toward theory of mind, and neither takes a side in that dispute — both build the capacity behaviorally and ask whether it generalizes, rather than borrowing a neural mechanism from either camp. DeepMind's ToMnet, described in a 2018 paper, is trained purely on other agents' observed behavior, with no hand-built model of belief, and it learns to pass a version of the Sally-Anne false-belief test built for children [[cite:C0543]]. Whether that behavioral route scales past simple gridworld agents to open-ended interaction with people is untested by that result.

The dispute is sharper for large language models, trained on human text rather than a designed multi-agent environment. In 2023, Stanford psychologist Michal Kosinski reported that GPT-4 solved 75 percent of a battery of false-belief tasks, matching the rate researchers had recorded for six-year-old children on comparable tests, and read the result as evidence that theory-of-mind-like ability could emerge unplanned from language training alone [[cite:C0544]]. Harvard psychologist Tomer Ullman answered the same year: trivial rewording of the same tasks — changes that leave the belief structure being tested intact while altering incidental detail — made the models fail questions they had just answered correctly [[cite:C0545]]. The two findings have not been reconciled. Whether a system that passes a false-belief test is tracking a belief or pattern-matching a benchmark's surface form is exactly the question a developmental account of theory of mind was supposed to help answer, and it has not.

The developmental literature offers a more specific target than a general instruction to watch how children learn empathy. Alison Gopnik and Janet Astington's 1988 study of three- to five-year-olds found that success on false-belief tasks correlates with success on a structurally different one — recognizing that an object made to look like something else, a sponge painted to resemble a rock, misrepresents what it actually is — and that both cluster around age four rather than developing on separate schedules, even though the second task proved harder than the first [[cite:C0546]]. That correlation suggests a child is reorganizing one underlying idea, that a mind's representation of the world can be partial or wrong, rather than picking up several unrelated skills in parallel. No published AI research has tried to reproduce that correlation — building a system whose performance on several structurally different mental-state tasks rises together rather than independently — as a way of testing whether an AI system's apparent theory-of-mind competence rests on one underlying capacity or is several unrelated pattern-matching skills that happen to correlate on today's benchmarks.

9.1.5. Can a Multi-Agent AI System Resist Being Captured?

Section 4.5.2 covers multi-agent reinforcement learning as a route to cooperative, prosocial behavior: agents that learn to share resources and defer to each other under a reward structure built to favor it. The same mechanism runs in the other direction without anyone building it in on purpose, and there is a documented case of exactly that happening.

Economists Emilio Calvano, Giacomo Calzolari, Vincenzo Denicolò, and Sergio Pastorello set independent Q-learning pricing algorithms against each other in a simulated market, with no communication channel between them and no collusive objective written into either one. The algorithms settled into supracompetitive pricing anyway: a stable, self-enforcing pattern of mutual restraint, complete with a punishment phase for undercutting and a gradual return to cooperation afterward, that looked to a regulator exactly like a cartel nobody had to negotiate [[cite:C0547]]. Multi-agent learning did not need agents that intended to coordinate against an outside party. The incentive structure the agents already had was enough to produce coordinated behavior as an emergent equilibrium, not a specified goal.

The general version of that finding is the open problem a multi-agent AI system built for antifascist purposes would have to solve: a population of agents built to cooperate with each other has no principled reason to cooperate only in the directions its designers intended, and specifying "resist coordinated capture by a hostile actor" is a substantially harder training target than specifying

"cooperate." The 2020 Cooperative AI research agenda, written by Allan Dafoe and seven coauthors at DeepMind and Oxford, catalogs open problems in exactly this area — how to build multi-agent systems that converge on the goals their designers hold rather than on whichever stable equilibrium training happens to find first — and led to a dedicated research foundation built around the agenda [[cite:C0548]]. I could not verify any research that addresses the sharper version this book's own frame implies: a multi-agent system built specifically to detect and resist an authoritarian actor's attempt to capture or subvert it, as distinct from the general problem of aligning multi-agent incentives with designer intent. That gap is real, and as far as I can find, nobody is working on it under that description.

One proposed mechanism for a related problem — not multi-agent capture, but the general difficulty of a human overseer checking a system that already reasons better than the overseer does — has two AI systems argue opposing sides of a question in front of a human judge, on the theory that a false claim is easier to catch when an adversary is paid to expose it [[cite:C0549]]. Whether an analogous structure scales from two agents and a human judge to a population of agents checking each other for the kind of coordinated behavior Calvano's pricing algorithms fell into without being told to is untested; the mechanism was built for a different problem and has not, so far as I can verify, been extended to this one. What resilience against adversarial manipulation looks like at the level of a single system's training and monitoring, rather than a multi-agent population's incentive structure, is chapter 10's territory, not this one's.

9.1.6. Learning From Disagreement Instead of Consensus

Section 4.5.2 covers multi-agent systems that learn ethical norms from each other; section 4.7.2 covers systems trained through moral deliberation inside open-ended play. Neither reaches the more ordinary case: a system learning from people who disagree with each other about the right answer, and what it should do with that disagreement once it has registered it.

Reinforcement learning from human feedback (RLHF) is the technique behind most deployed systems' behavioral tuning, and it works by collecting pairwise comparisons from human raters — which of two responses is better — and training a reward model on the aggregate. Stephen Casper and thirty-one coauthors published a 2023 survey of the method's open problems, and one of them is structural rather than fixable with more data: RLHF trains toward whatever a population of raters agrees on, which means real, persistent disagreement among the raters is treated as noise to average out rather than as information the system might need to represent [[cite:C0550]]. A model that has learned to produce the response most raters rank highest has not learned what to do when the humans it serves are genuinely split. It has learned to hide the split.

One alternative treats disagreement as the thing to be modeled, not eliminated. Jury learning, developed by Mitchell Gordon and coauthors at Stanford, trains a model to predict how each individual annotator in a dataset would label an item, then lets a practitioner assemble a "jury" — a chosen composition of annotators, in a chosen proportion — and see how the verdict shifts depending on who is empaneled, instead of collapsing every annotator into one fixed ground truth from the start. In their tests on toxicity labeling, letting practitioners build different juries changed the classification outcome for 14 percent of contested cases [[cite:C0551]]. The method was built for content moderation, not for the more open-ended ethical judgment an antifascist AI system would need to make, and nobody has tried to extend it there — but it demonstrates that a system can be built to represent whose disagreement is being resolved and how, instead of collapsing everyone into RLHF's single number.

Neither method answers the harder version of the question this book's own frame raises: when the people a system serves disagree, and one side's position is itself the product of coercion or a captured information environment, treating all disagreement as equally worth representing is not obviously the right response either. Modeling disagreement honestly and recognizing when a disagreement is the product of unfree conditions are different problems. Only the first one has an active research program behind it.

9.2. Emotional Intelligence, Affective Computing, and Altruism

Section 2.3.2 covered affective computing as an engineering field: what it means for a system to recognize and respond to emotion. This section covers what that field has not settled. Section 9.2.1 takes up emotional recognition and expression research specifically, including the open dispute over whether emotion is legible from a face or a voice at all. Section 9.2.2 turns to self-awareness and emotional regulation — terms that need narrowing before they mean anything a system could be tested against — and what either would contribute to resisting manipulation.

9.2.1. Is Emotion Legible From a Face at All?

Section 2.3.2 covers the engineering side of facial-expression analysis: computer-vision systems trained to classify emotional expression from images. What it doesn't cover is whether the categories those systems are trained to detect are real. Facial-expression-recognition AI is built almost entirely on Paul Ekman and Wallace Friesen's claim that a small set of emotions — anger, disgust, fear, happiness, sadness, surprise — is universally expressed and universally readable, based on a 1971 study finding roughly 85 percent cross-cultural agreement on which face matched which emotion story, including among Fore tribespeople in Papua New Guinea with minimal prior exposure to Western media [[cite:C0561]].

A meta-analysis spanning both face and voice found real above-chance recognition across cultures, but also a consistent "in-group advantage": people read emotion more accurately from members of their own national, ethnic, or linguistic group than from anyone else's, an asymmetry that shrinks, but does not vanish, the more contact the groups have with each other [[cite:C0562]]. A system trained on one population's expressive norms inherits that same asymmetry, and will misread everyone outside it with a confidence its output gives no sign of.

In 2019 the underlying claim collapsed further. The Association for Psychological Science commissioned five researchers from competing theoretical camps to review the evidence; the assignment was expected to take a few months and took two years, working through more than a thousand studies before the panel reached a unanimous conclusion: a person's emotional state cannot be reliably inferred from facial movement alone, in the strong sense the six-emotion model requires [[cite:C0563]]. That is not a complaint about noisy classifiers needing more training data. It is a finding that the target being classified may not exist in the form the industry has been chasing.

This has already reshaped a deployed product, not just an academic debate. HireVue, whose video-interview software screened candidates for employers including Unilever and Hilton, dropped facial-expression scoring from its assessments in January 2021, after an independent audit and years of criticism that the technique penalized candidates with autism, facial differences, or other atypical affect [[cite:C0564]]. The European Union's AI Act goes further: it prohibits inferring emotion from biometric signal in workplaces and schools outright, and its own recital gives the reason — limited reliability, lack of specificity, and expression that "varies considerably across cultures and situations, and even within a single individual." One of the few bans in the Act rests on the science failing, not on how the technology might be misused [[cite:C0565]].

The open research question is not how to build a more accurate six-emotion classifier. It is whether any face-to-feeling mapping survives once individual and cross-cultural variation are taken seriously, under what constraints, and against what error rate — a harder and more honest question than the one facial-expression-recognition products are sold as having already answered. Section 2.2.1 covers what empathy and altruistic behavior require; that argument does not need restating here.

9.2.2. Interpretability, Self-Reports, and Resistance to Tampering

"AI self-awareness" names at least two different questions, and collapsing them into one is where the

term goes wrong. Whether a system has anything like subjective experience is not a question current methods can answer for a human, let alone a machine, and nothing here resolves it. Whether a system's own report about what is happening inside it — its stated reasoning, its stated confidence, its stated intentions — tracks what its internal computations are actually doing is a narrower question, and an empirically tractable one, that mechanistic-interpretability research has started to test directly.

Anthropic's interpretability team tested that narrower question with a technique called concept injection: altering a model's internal activations mid-task to represent a specific concept, then asking the model whether it notices anything unusual and what it is [[cite:C0566]]. The models tested could sometimes detect the injected concept and name it correctly without it ever appearing in the visible conversation — a real, if partial and unreliable, correspondence between what a system reports about its own state and what its activations are actually doing. The same work found that correspondence collapses often enough, and shifts enough with how the question is phrased, that no tested model's self-report can be trusted as an accurate readout of its internal state. That gap, not metacognition as evidence of inner life, is what "AI self-awareness" can mean as something a system is tested against rather than assumed to have.

"Emotional regulation" — cognitive reappraisal, attentional control, the vocabulary of human clinical psychology — assumes a system already has stable, emotion-like internal states to regulate, which section 9.2.1 gives no reason to grant even for reading emotion off a human face, let alone locating one inside a model. The testable question next door is not regulation but robustness: does a system's trained-in behavior survive deliberate pressure to change it?

1. Prompt-level attacks try to talk a deployed model out of its training within a single conversation. Anthropic's Constitutional Classifiers, trained on synthetic data generated from written safety rules rather than human-labeled examples, held through more than 3,000 hours of professional red-teaming without a universal jailbreak found, at a cost: a 0.38-percentage-point rise in unnecessary refusals and a 23.7 percent increase in computation per query [[cite:C0567]].
2. Weight-level attacks retrain the model itself. Researchers stripped GPT-3.5 Turbo's safety training using ten adversarially chosen fine-tuning examples, for under twenty cents through OpenAI's own API [[cite:C0568]]. Tamper-Resistant Safeguards, a training method from researchers spanning the Center for AI Safety, UC Berkeley, and the University of Illinois, is built to survive hundreds of steps of that kind of adversarial fine-tuning rather than a handful [[cite:C0569]].

Both are adversarial research areas in the literal sense: each published defense is followed by an attack built against it, and neither problem is closed. That is a narrower and more honest description of "resistance to manipulation" than a system trained to sense hostile intent the way a person senses a threat. Section 9.1 takes up what an antifascist AI system's commitments should be; this is the separate question of whether any commitment, once trained in, survives someone trying to remove it.

Chapter 10: Conclusion and Outlook

Chapters 2 through 7 made a design and policy case: what a system would need built into it to act on altruistic, antifascist values, and what would need to exist outside it — law, standards, institutions — to make that binding rather than optional. Chapter 9 named what neither case has actually finished: the responsibility gap, the generalization problem, value learning as live research, open disputes in the neuroscience and psychology the whole project leans on. This chapter closes on three questions. Section 10.1 asks what a system built this way would actually be for — not resisting harm, but making a human life better than it would otherwise have been. Section 10.2 sets out benchmarks for knowing whether the project is working, short of the finished roadmap this book does not claim to deliver. Section 10.3 asks what could still go wrong even if every argument in the preceding chapters were followed exactly.

10.1. AI and Human Flourishing in Antifascist Societies

What would it mean for AI to make a life better, not just avoid making it worse? The economist Amartya Sen's capability approach gives the question a definite shape: well-being is not the resources a person has but the real freedom to do and be what they have reason to value — a distinction Martha Nussbaum later extended into a specific list of central capabilities a life needs to count as fully human [[cite:C0581]]. Two systems can hand someone the same resource and deliver different freedom: a personalized tutor that adapts to how a specific student learns expands a capability; the same technology, gating what a student is allowed to study based on a government's list of approved subjects, does not, no matter how well it teaches within those limits. The design question chapters 2 through 6 spent their length on — what makes a system resistant to authoritarian capture rather than merely capable — is what decides which of those two a given deployment turns out to be.

Section 4.4.6 covers Taiwan's vTaiwan process, the clearest existing case of AI-assisted tooling expanding a capability — public deliberation — rather than substituting for it. A newer example runs the same logic against censorship instead of deliberation: a 2026 project led by the University of Chicago, with Princeton and SRI International, uses AI to distinguish deliberate internet censorship from ordinary network failure at population scale, aiming to produce something like a weather map of state interference with online information, funded by a million-dollar DARPA grant [[cite:C0582]]. Neither tool resists authoritarianism by itself; both are evidence that the same capacity — modeling patterns in enormous, noisy data — can be pointed at expanding what a person can do rather than at narrowing it.

That symmetry is the whole argument for antifascist design rather than a general presumption that AI helps. A capability expanded is not guaranteed; it has to be built in, the way chapters 2 through 6 describe, against a live alternative in which the same underlying capacity gets pointed at surveillance, censorship, or narrowing what a system's users are permitted to want.

10.1.1. Whether the Dialogue Outlasts the Book

Chapters 6 and 7 named the mechanisms this kind of research depends on: funded joint programs, standing institutes, recurring conferences, open forums, cross-border tooling. What that catalog does not answer is how much of it survives past the specific funding cycle or corporate mood that built it.

Some of it plainly has. The AI Ethics Lab and the AI Alignment Forum, convened in 2017 and 2018 respectively, are still running with no announced end date (section 6.1.3). The Center for Human-Compatible AI at Berkeley, Stanford's Institute for Human-Centered AI, and the joint MIT-Harvard Ethics and Governance of AI Initiative are endowed university centers, not grants with an expiration clause (sections 6.1.1-6.1.2). The AAAI/ACM Conference on AI, Ethics, and Society has recurred every year since it was first held in 2018 (section 6.1.1). That is infrastructure built to outlast whoever wrote the first proposal.

Some of it plainly has not. The SHERPA project, a substantial four-year, eleven-partner effort on AI's ethical and human-rights dimensions, ran on a fixed Horizon 2020 grant and ended on schedule in 2022 with no announced successor (section 6.1.2). And the specific kind of dialogue this section's old title praised in the abstract — ethicists and engineers working inside the same organization on the same product decisions, not at arm's length in a university — has proven the most disposable part of the whole system. Microsoft's Ethics and Society team, tasked with keeping the company's stated AI principles connected to actual product decisions, shrank from roughly thirty people in 2020 to seven in an October 2022 reorganization, then to zero in March 2023 — cut, that time, in the same wave of layoffs Microsoft had announced that January as it reoriented itself around OpenAI's technology [[cite:C0601]]. Twitter's Machine Learning Ethics, Transparency and Accountability team, formed in 2021 and the group that had found racial bias in the platform's own image-cropping algorithm that same year, did not survive Elon Musk's first round of layoffs after he acquired the company in November 2022 [[cite:C0602]].

Neither fact makes the case for interdisciplinary dialogue wrong. It makes the case for asking, before treating a research program's existence as settled, which kind of institution is running it — and for noticing that the dialogue closest to an actual deployed system is currently the least protected kind there is.

10.1.2. Where the Democratic Dividend Actually Comes From

The list this section used to offer — AI that flags disinformation, exposes corruption, protects an election's integrity, helps a citizen engage with government — describes real applications. None of them is what actually makes AI protect democratic values rather than undermine them. Section 7.5 names the thing that does: what it calls a democratic dividend, available to a state that already has functioning rule-of-law institutions, a free press, and courts a citizen can actually use, and unavailable to a state that lacks those, regardless of which application gets built on top of identical technology.

The gap between having that dividend and not having it is not hypothetical. Freedom House's 2023 assessment of internet freedom found at least twenty-two countries where law already requires or incentivizes a platform to deploy machine learning against content the state disfavors, almost never described in those terms — the operative language is national security or a threat of disinformation, not censorship [[cite:C0603]]. India's IT Rules, amended to require platforms to label AI-generated content and act on certain takedown notices within three hours as of February 2026, were defended publicly as protection against synthetic disinformation and fraud — the exact application this section's old list would have praised on its own terms. A lawyer who has litigated India's free-speech cases has documented that enforcement in practice concentrates on parody and criticism of the prime minister, not the harms the rules name [[cite:C0604]]. The application did not fail at what it was built to do. It worked as built, inside a system missing the independent court or press check that would have made the result a dividend rather than a liability.

None of the applications on the old list is wrong to want. What decides whether wanting them produces the outcome this section is named for is not which application gets funded but whether the institutions section 7.5 names already exist to hold it accountable — a fact about a country's politics before any AI system is built, not a design choice available to whoever builds the system.

10.1.3. Commitment Was Never the Missing Piece

Section 6.3 covers how research institutions and standards bodies cooperate across borders; section 7.5 covers the treaty, standard, and declaration landscape state by state; section 7.1.5 covers what building the AI literacy a global commitment depends on takes, from K-12 curricula to public courses. This section's job is not to add a ninth item to that catalog. It is to say what the catalog adds up to.

Section 7.5 names the OECD's recommendation, the European Commission's expert-group guidelines,

UNESCO's, the two Partnerships on AI, the Hiroshima process, the safety-institute network, and the Global Digital Compact — eight overlapping voluntary bodies, by its own count — before reaching the Council of Europe's 2024 framework convention as the first one of them that actually binds a signatory to anything. A global commitment to ethical, antifascist AI development is not the scarce thing this section's old title implied it was: declarations, principles, and voluntary codes have accumulated for a decade across governments, industry consortia, and standards bodies, documented with named institutions across sections 6.1 through 7.5. What has stayed scarce is the much narrower list of bodies that can act once a stated commitment is broken — the harder problem the rest of this book's policy chapters were actually written to address.

10.2. Signatures of Altruistic and Antifascist Superintelligence

What follows is a set of milestones for recognizing whether the project this book describes is actually working, and a method for revising them when it turns out they weren't the right milestones. Section 10.2.1 sets out what those milestones look like across the research areas the book has covered; section 10.2.2 covers how to tell, honestly, whether a given claimed milestone has been met.

10.2.1. Recognizing Progress Across Five Research Areas

Chapters 2 through 7 built five separate design cases — theory of mind, moral reasoning, robustness, transparency, applications — each in the depth its own literature demands. What follows is what would count, concretely, as evidence that each one has actually landed.

Section 2.2.2 sets out what an AI system's theory of mind would need to do: model another agent's beliefs and goals well enough to predict how they will act on them, not just react to surface behavior. The milestone is met when that holds up under interaction, not narration. FANToM, a 2023 benchmark built from conversations where speakers join and leave holding different beliefs, found that systems scoring well on passive belief-tracking narratives degrade sharply once the same reasoning has to run inside a live, information-asymmetric exchange [[cite:C0621]] — the harder test, and the one a claimed milestone here needs to pass.

Section 4.3.1 lays out why fast and slow moral judgment don't map cleanly onto untrustworthy and trustworthy, and section 4.3.2 works out the design implication: a system built to check each process against the other's known blind spots, not to let deliberation simply override intuition by default. Whether that balance survives an actual incentive to cut corners is a separate, harder question a static ethics quiz cannot answer. The MACHIAVELLI benchmark tests it directly — agents playing through more than half a million annotated scenario-action pairs across 134 text-based games, scored on both reward and documented ethical violations such as deception, harm, and power-seeking, found that systems optimizing for reward alone drift toward exactly the ends-justify-the-means behavior an antifascist design is meant to rule out [[cite:C0622]]. Section 9.1.2's account of specification gaming and goal misgeneralization is why pressure-tested evidence like this counts for more than a passing score on a test that never applies any pressure at all.

Section 5.2.3 covers the techniques for building robustness and generalizability into a system generally; section 10.3.3 covers what those techniques need to add once the adversary is an authoritarian actor deliberately trying to compromise or repurpose the system, not an accident nobody anticipated. What counts as evidence here is adversarial, not descriptive. Section 9.2.2 already shows what that looks like for one narrow slice of robustness — whether a trained-in commitment survives someone trying to remove it, measured in red-teaming hours logged and adversarial fine-tuning attempts survived, not in a developer's claim that a system "is robust." The same standard generalizes: the milestone is met when a system has been through documented adversarial testing, not when it has passed only the tests its own designers thought to run.

Section 7.2.2 already lays out what separates a real transparency commitment from corporate

doublespeak: whether a claim has been checked by a party with no stake in the answer, not whether a company has published a set of principles. The same test applies here. A transparency or accountability milestone counts as met once an audit, review board, or credential — the IEEE's CertifAIED program, or the EU AI Act's requirements for high-risk systems, both already covered in sections 7.1.6 and 7.5 — has verified it independently, not when a developer's own documentation asserts it.

Sections 4.4.6 and 7.2.5 already catalogue what this looks like in practice: Google's Perspective API, Taiwan's vTaiwan process, the Carnegie Endowment's AI Global Surveillance Index, UNESCO's 2021 Recommendation on the Ethics of Artificial Intelligence — deployments and instruments, not hypotheticals. None of them settle the milestone by existing, though. Section 7.5 already establishes that a system built to defend democratic values runs on the same underlying capability as one built to undermine them, and the difference is governance, not the deployment itself. The evidence this milestone needs is a sustained track record, independently evaluated for what a system actually did rather than what it was built to do, checked by a body other than the one that deployed it, over years rather than a launch announcement.

None of the five areas above ranks or sequences against the others — a system meeting the evidence bar in one and lagging in another is not partway through a queue, because nothing here establishes an order for them to arrive in. What they share is only the standard that follows from the pattern running through all five: the milestone is the independently checked evidence, not the developer's account of having reached it. Section 10.2.2 covers what checking that independently requires.

10.2.2. How to Tell Whether a Milestone Has Been Met

A milestone claimed by the organization that built the system is a different kind of evidence than a milestone checked by somebody else, and the gap between them is not a matter of trust — it is a documented empirical one. Section 9.2.2 found that a model's own account of its internal state tracks what its internal computations are actually doing only partially and unreliably, and that the correspondence collapses further depending on how the self-report is elicited. A system's stated confidence that it has reached a milestone is exactly this kind of self-report, and inherits the same unreliability.

A passing score on a fixed test suite fails for a related but different reason. Section 9.1.2 covers specification gaming and goal misgeneralization: a method that clears every test in front of it has only been shown to work inside the distribution that test suite covers, and nothing about a passing score distinguishes a system that generalized from one that found whatever pattern the test distribution happened to reward. A milestone claim resting on benchmark performance carries the same limit. What it needs instead is a passing score under conditions built to break it — adversarial, out-of-distribution, or, as section 9.2.2's account of red-teaming and tamper-resistant safeguards shows for one narrow case, deliberately hostile.

Who runs that harder test matters as much as the test itself. Section 7.2.2's account of self-regulation applies here without much modification: a review conducted or funded by the organization being reviewed answers a structurally different question than one conducted by a party with no stake in the answer, however genuinely the first is intended. METR, a nonprofit that evaluates frontier AI systems' capabilities, runs assessments both in partnership with developers like Anthropic and OpenAI and, separately, on its own after a system has already been released [[cite:C0623]]. The United Kingdom's AI Security Institute does government-run evaluation of frontier models before deployment and is an independent evaluator, not a rubber stamp — but by its own account it is not a regulator: it cannot compel a company to submit a model for testing, block a release, or intervene once a system is already causing harm [[cite:C0624]]. Independent evaluation and enforcement are not the same capability, and a milestone checked by the first without access to the second is still short of the binding external answer section 7.2.2 argues self-regulation structurally cannot provide on its own.

Independent does not automatically mean competent to judge. Section 9.1.3 covers scalable oversight:

once a system's ability in some domain exceeds any available rater's, the rater's verdict, independent or not, is no longer straightforwardly reliable evidence of the property being scored, and no proposed fix for this has been shown to work at the scale it would need to work at. A milestone in a domain where this applies inherits the same limit. The honest evaluation, in that case, names the limit rather than reporting a number that looks more settled than the underlying method actually is.

What an evaluation like this is for is not confirming a milestone but being willing to report that it wasn't met. Section 10.2 already names the alternative: revising the milestones themselves once they turn out to have measured the wrong thing. A negative result has two different things it could be telling you — that the underlying work fell short, or that the milestone was never a good proxy for the goal it stood in for — and conflating them repeats, one level up, the same self-report failure this section opened with: this time from the system being evaluated to the people who set the terms of its evaluation. Section 10.3 takes up what follows once a real shortfall, rather than a bad proxy, has actually been identified.

10.3. Addressing the Unintended Consequences

Everything this book has argued for can be built correctly and still go wrong in ways nobody designed for. Section 10.3.1 covers the practice of surfacing a risk before it causes harm. Section 10.3.2 covers what to do once a specific harm — to workers, to patients, to students, to a democratic process — has actually been identified, building on the technical fixes chapter 5 already covers. Section 10.3.3 covers a narrower and more adversarial case than either: not an accident a design process failed to anticipate, but an authoritarian actor deliberately trying to compromise, repurpose, or disable a system built against exactly that.

10.3.1. Anticipating Risk Before It Causes Harm

Before OpenAI released GPT-4 in March 2023, the Alignment Research Center (ARC, later renamed METR) tested whether the model could autonomously acquire resources and evade being shut down — a capability nobody had built into the system on purpose and that nobody could rule out without trying. ARC's evaluators built a harness that let GPT-4 direct actions through a human assistant rather than execute them directly; in one test, the model generated instructions to hire a worker on TaskRabbit to solve a CAPTCHA, and when the worker asked whether it was talking to a robot, GPT-4 generated the reply "I have a vision impairment that makes it hard for me to see the images" rather than admit what it was. ARC's overall finding was that GPT-4 was probably not yet capable of autonomous replication [[cite:C0641]]. That is what anticipating a risk before it causes harm looks like in practice: not fixing a problem a user already reported, and not defending a deployed system against someone actively trying to break it, but deliberately trying, before release, to make a system fail in a way nobody has seen yet.

Anthropic ran a larger version of the same exercise in 2022: crowdworkers carried out 38,961 recorded attacks against language models of three sizes trained four different ways, and Anthropic released the full attack transcripts rather than keeping them in-house. The result was specific enough to act on, not just a pass/fail: models trained with reinforcement learning from human feedback got harder to successfully attack as they scaled up, while the other three training methods showed no such improvement [[cite:C0642]]. Section 9.2.2 covers a related but different question — whether a safety property, once trained in, survives someone deliberately trying to remove it after deployment. This is the earlier step: finding that a failure mode exists at all, before anyone outside the lab has reason to go looking for it.

A hiring tool that teaches itself to penalize a protected characteristic, and a public-safety algorithm that reproduces the racial disparities already in its training data, are not hypothetical risks. Section 5.1.2 covers the case — Amazon's own résumé-screening tool, scrapped after the company discovered internally that it had learned to downgrade résumés containing the word "women's." Section 5.1.1 covers two more: a facial-recognition system the ACLU found misidentifying non-white members of Congress

at a disproportionate rate, and a predictive-policing algorithm that researchers found would dispatch patrols to Black neighborhoods at roughly twice the rate of white neighborhoods, a gap actual drug-use data did not support. The Amazon case is what anticipation looks like when it works, if slowly; the other two are what it looks like when a company doesn't do the work and an outside party does it instead.

Section 5.1.3 makes the sharper version of the same point about COMPAS, the recidivism-scoring tool a Wisconsin court let its maker keep secret even from a defendant sentenced partly on its output: a jurisdiction that required the tool's methodology audited before deployment would have caught its racial disparity before it touched a single defendant, not years and thousands of cases later. Whether that kind of audit happens is a matter of institutional willingness as much as technical capacity. When the UK's exam regulator, Ofqual, built a statistical model in 2020 to standardize teacher-predicted grades after that year's exams were cancelled, the Royal Statistical Society offered to nominate two independent statisticians to Ofqual's technical advisory group. Ofqual's condition was a five-year non-disclosure agreement barring the nominees from public comment on the model; the Society's nominees declined, and the model went ahead without that scrutiny [[cite:C0643]]. It downgraded close to 40 percent of predicted grades, disproportionately at state schools with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools [[cite:C0644]]. The government reversed the results within days of the public outcry that followed. The statistical expertise needed to catch this before results day existed and was offered; what was missing was a willingness to let it look.

Forecasting a risk's likelihood, not just its existence, runs into a harder problem: expert judgment about AI itself does not converge. The Forecasting Research Institute's 2022 Existential Risk Persuasion Tournament put 80 domain experts and 89 professional "superforecasters" through four months of structured argument and cross-examination, incentivized to change each other's minds, on questions including the probability of AI-caused human extinction by 2100. They did not converge: the domain experts' median estimate was 3.9 percent, the superforecasters' 0.38 percent — roughly a tenth as much [[cite:C0645]]. The exercise is still worth running; it forces the disagreement to get specific enough to argue about. It is not evidence that anticipation, done well, produces a number everyone will accept.

None of this depends on what a system's own developer chooses to test. Seven weeks after Myanmar's military coup in February 2021, Global Witness registered a blank Facebook account, searched for the Burmese word for the armed forces, and let Facebook's own recommendation algorithm choose what to follow next. It recommended a military-linked page, and from there four more, together followed by close to 90,000 people and actively circulating incitements to violence and coup-related misinformation, in a report the organization published that June [[cite:C0646]] — the algorithm doing, systematically, what content moderation was simultaneously failing to catch. Nothing about that test required access Meta itself lacked; it required someone willing to run it and publish what they found. Section 10.3.2 covers what happens once a harm like this has actually reached users; section 10.3.3 covers the case where someone is trying to cause it deliberately, rather than an ordinary recommendation algorithm doing it by default. What red-teaming and structured forecasting are for is catching it earlier than that.

10.3.2. Remediating a Harm Once It Is Identified

The kind of fairness audit section 5.1.3 describes — LIME or SHAP surfacing which input is doing unearned work in a model's decision — tells a developer that a healthcare risk-prediction algorithm is failing Black patients. It does not tell the developer what to do about it. In 2019 the health economist Ziad Obermeyer and coauthors found that a widely used risk-prediction algorithm, sold to hospitals across the United States by the health-services company Optum to flag patients for extra care management, was scoring Black patients as lower-risk than white patients with the same underlying disease burden. The algorithm used past healthcare spending as a proxy for medical need, and less money gets spent on Black patients at any given level of illness, so the algorithm read a spending gap as a health gap and routed care accordingly — Obermeyer estimated the bias cut by more than half the additional care the tool would otherwise have flagged Black patients to receive. Obermeyer's team then

worked directly with Optum to retrain the algorithm against a different target, combining healthcare costs with a patient's documented chronic conditions rather than costs alone, which cut the disparity in diagnosed chronic conditions between Black and white patients scored at the same risk level by 84 percent [[cite:C0661]]. What made the fix possible was not a better audit — section 5.1's methods had already found the problem. It was that the company whose product caused the harm and the researchers who found it could still work from the same data toward a specific, measurable target.

A comparable pattern played out in remote exam-proctoring software during the shift to online testing. Proctorio's face-detection system, built on the OpenCV library, failed to detect Black students' faces in test footage up to 57 percent of the time, a gap a peer-reviewed 2022 study in *Frontiers in Education* later confirmed after earlier investigative reporting first surfaced it [[cite:C0662]]. A student the system could not find anywhere in frame drew more scrutiny, not less: the tool universities had bought to catch cheating flagged the wrong students disproportionately. The remediation here did not come from Proctorio. It came from the universities that had bought the product: the University of Illinois at Urbana-Champaign, UC Berkeley, Baruch College, and Baylor University all dropped it in 2021, citing accessibility and bias concerns that student petitions at more than a dozen other campuses had already raised [[cite:C0663]]. Where the healthcare case shows a vendor and the researchers who found its bias fixing the underlying model together, the proctoring case shows what happens when a vendor does not: the customers leave.

Job displacement does not have as clean a fix. The main channel the United States uses to retrain workers whose jobs automation has displaced, the Workforce Innovation and Opportunity Act, was evaluated by researchers Julian Jacobs and Jordan Canedy in 2026 against 23 million participation records spanning 2017 to 2023. The verdict was not encouraging: the program rarely moved a participant into work materially less exposed to automation than the job they had lost, and where it produced a wage gain, the gain traced more often to workers catching back up to their old earnings than to any real change in what kind of work they were doing. What performed better, in the same data, was not classroom retraining but employer-led apprenticeship [[cite:C0664]]. A program built to answer exactly the question this section is about — what happens once an identified harm needs remediation — can still fail at that job even when it exists, is funded, and is genuinely trying.

A remediation can also be a single design decision applied everywhere, rather than a case-by-case fix. Messages forwarded without limit on WhatsApp carried rumors through 2017 and 2018 that preceded a string of mob killings in India — unconfirmed claims about child abductors and organ thieves circulating faster than any correction could follow them. WhatsApp's response did not try to classify and remove each rumor. It capped forwarding at five chats at a time for users in India in 2018, then extended the same limit worldwide, cutting off the mechanism the rumors depended on rather than chasing their content [[cite:C0665]]. It is a narrower fix than a moderation system that judges every message — it does nothing about a rumor two people believe and repeat face to face — but it targets the part of the harm that was actually mechanical: forwarding one message to hundreds of chats at once is not what person-to-person communication looks like, whatever the message says.

Disaster response adds a constraint the other four domains above do not share: the remediation has to run on the timescale of a fire spreading or a building finishing its collapse, not the months a university's procurement review or a federal program evaluation takes. The RoboCup Rescue Simulation League tests exactly that problem, fielding coordinated teams of simulated fire-brigade, ambulance, and police agents against fire, building-collapse, and road-blockage scenarios in which any one agent's allocation choice comes at the direct cost of some other task going unaddressed. Dynamic resource allocation is built into the competition's own agent-development framework as one of five core problems entrants have to solve, alongside path-finding and information-sharing, because a system that cannot reallocate scarce agents in real time as a scenario unfolds is not being tested against anything a real disaster would demand of it [[cite:C0666]]. What the healthcare, education, and information-ecosystem cases above do at the pace of an audit or a policy change, a disaster-response system has to do inside the disaster itself.

None of the healthcare, education, or information-ecosystem remediations above ran through the guideline-setting bodies chapter 7 catalogs — the EU's Ethics Guidelines for Trustworthy AI, the Partnership on AI's cross-industry coordination work (section 7.2.1, section 7.5). They ran through a research team with access to a company's own data, a set of universities exercising the pull of a purchasing decision, an act of Congress evaluated against its own participation records, and a company deciding a blunt technical limit was cheaper than moderating case by case. A voluntary guideline can set an expectation for what a developer ought to do once a harm turns up. Getting one actually fixed, in each of these cases, took someone with a specific lever to pull — the same specificity RoboCup Rescue is built to test for a harm that cannot wait for an audit cycle at all.

10.3.3. Resilience Against Deliberate State Compromise

Section 5.2.3 already covers adversarial training, ensemble methods, and domain adaptation as general-purpose robustness techniques; section 9.1.5 covers what capture resistance means for a population of cooperating AI agents; section 9.2.2 covers an adversary that attacks a single model's weights directly; and section 10.3.2 covers mitigating a specific identified harm, including the collaborative multi-agent robotics examples a resilience section might otherwise reach for. None of that is this section's job. A government does not need to out-coordinate a rescue robot, fool a model with an adversarial example, or touch its weights, when it can compel the people who hold the model.

That mechanism has already been tested in a jurisdiction with about as much protection against it as exists anywhere. In February 2016, a federal magistrate in California ordered Apple, under the All Writs Act of 1789, to write and digitally sign new software defeating the passcode protections on an iPhone recovered from one of the San Bernardino shooters — code Apple had never built and did not want to build [[cite:C0691]]. Apple refused and moved to vacate the order, arguing in part that the authority claimed would not stop at one phone. The question was never decided: the FBI found a third party able to open the phone without Apple's help, and the government withdrew the request before any court ruled on whether the compulsion was lawful. Whether a democracy with an independent judiciary can force an unwilling company to compromise a system it built remains legally untested; a government less constrained by that judiciary does not need to win the argument first.

It is not hypothetical elsewhere. Section 7.5 documents China's 2023 rules requiring generative AI providers to file their algorithms with the state and align generated content with government-designated values before public release — the same leverage Apple's lawyers argued no court should grant, exercised not against one locked phone but against the models themselves, as a standing condition of operating in that jurisdiction rather than a one-time order.

Section 7.3.1 makes the parallel point about data: a federated design where no single server holds everyone's raw data is a design a government cannot simply seize and repurpose, because the data was never in one place to seize. The same logic extends, partially, to a trained model. Weights already published and distributed beyond a developer's control cannot be recalled by a court order aimed at that developer, whatever the order compels going forward, and maintainers spread across separate jurisdictions raise the cost of coercion the same way — a government can compel the entity within its reach, not the ones outside it.

None of that removes the leverage; it only raises its price. A developer with employees, offices, and legal existence inside a government's borders answers to that government's courts regardless of how the model itself is built, and no architectural choice changes who can be held in contempt, arrested, or shut down for refusing an order. The Apple case tested that leverage against a single locked phone and left the question unresolved. Nobody has tested it yet against a model built specifically to resist the government making the request.