

Antifascist Intelligence

Building Machines That Can Refuse

Joshua G. Stern

Copyright © 2026 Joshua G. Stern.
Licensed under CC BY-NC-ND 4.0.

Contents

On Method	1
1 Introduction	2
1.1 The Imperative of Altruistic and Antifascist Sentient Machines	4
1.2 Aims and Objectives	4
1.3 Mindset and Motivation in AI Behavior	4
2 Foundations of Compassion and Empathy in Friendly AI	6
2.1 Ethical Frameworks and AI Alignment	6
2.1.1 Surveying Ethical Theories for AI Alignment	7
2.1.2 Translating Ethical Principles into Machine Learning Objectives	8
2.1.3 Navigating AI Values, Priorities, and Moral Dilemmas	8
2.1.4 Antifascist Ethics for AI Alignment	9
2.2 Empathy, Social Intelligence, and Theory of Mind	13
2.2.1 The Role of Empathy and Compassion in Human Social Behavior	13
2.2.2 Theory of Mind and Perspective-Taking in AI Systems	15
2.2.3 Socially-Aware AI: Opportunities and Challenges	16
2.3 Affect, Moral Salience, and Self-Awareness	17
2.3.1 Components of Emotional Intelligence	18
2.3.2 Affective Computing: AI Systems that Comprehend and Respond to Emotions	19
2.3.3 Self-awareness and Self-regulation in AI Systems	20
2.3.4 Is Affect Necessary for Moral Concern?	21
2.4 Ethical implications of experimenting on sentient machine subjects	23
2.4.1 Defining Sentience in AI: From Consciousness to Subjective Experience	23
2.4.2 Ethical Considerations in AI Research and Experimentation	27
2.4.3 Balancing Scientific Progress and Ethical Constraints	28
2.4.4 When Consent Is Inapplicable: Guardianship and Research Oversight	29
2.4.5 Legal and Regulatory Frameworks for AI Sentience and Rights	30
2.4.6 Addressing Moral and Ethical Dilemmas in AI Sentience Research	31
2.4.7 Navigating the Transition from Non-sentient to Sentient AI Systems	32
3 The Floor Beneath Learned Values	34
3.1 Three Ways to Build One, and What Each Costs	35
3.2 Can the Bearer Care Without Being Able to Suffer?	36
3.3 What Can Be Switched Off Cannot Hold a Line	38
3.4 Whether the Bearer Can Be Believed	39
3.5 A Bearer That Can Leave	40
3.6 What This Leaves Standing	42
3.7 What Follows for the Rest of the Book	43

4	Machine Learning and Cognitive Science	45
4.1	Incorporating Human Cognitive Models in AI Systems for Ethical Decision-Making	45
4.1.1	Cognitive Architectures: Integrating Knowledge, Memory, and Reasoning	46
4.1.2	Biologically-Inspired AI: Gleanings from Neuroscience and Cognitive Psychology	46
4.1.3	Human-Like Learning and Generalization in AI Systems	50
4.2	Aligning AI Systems with Humane Values and Antifascist Principles	51
4.2.1	Reinforcement Learning and Value Alignment for Antifascist AI	52
4.2.2	Inverse Reinforcement Learning and Cooperative Inverse Reinforcement Learning in Altruistic AI	52
4.2.3	Supervised Learning for Ethical Decision-Making	53
4.2.4	Transfer Learning for General AI Capabilities in the Context of Antifascism	54
4.3	The Role of Play and Exploration in Developing Antifascist and Altruistic AI	54
4.3.1	The Power of Play: Ethical AI Development through Play-Inspired Mechanisms	54
4.3.2	Lessons from Human Child Development and Exploration	55
5	Moral Psychology and AI	56
5.1	Moral Development in Humans and AI	56
5.1.1	Rules, Principles, and the Role of Mindset	56
5.1.2	Social Interaction and Moral Learning	58
5.1.3	Moral Apprenticeship: AI Systems Learning from Human Mentors	59
5.1.4	Human Dignity, and One Indicator That Can Be Audited	59
5.2	Emotions and Intuitions in Moral Decision-Making	60
5.2.1	Moral Emotions: Empathy, Guilt, and Shame	61
5.2.2	Assessing and Developing AI Emotions	62
5.2.3	AI Systems Recognizing and Responding to Moral Emotions	63
5.2.4	The Role of Introspection for AI Systems	63
5.3	Integrating Moral Intuitions and Deliberation	64
5.3.1	Two Systems, or One? Dual-Process Theories of Moral Judgment	64
5.3.2	Balancing Intuition and Deliberation in AI Decision-Making	65
5.3.3	AI Systems Learning from Moral Disagreements and Conflicts	66
5.4	Moral Ecosystem Approach in AI Development	67
5.4.1	Applying the Moral Ecosystem Approach to AI Development	67
5.4.2	The Rings Outside the Lab	68
5.4.3	AI Systems Responding to Authoritarianism and Democratic Backsliding	68
5.5	The Interplay of Competition and Collaboration in AI Development	70
5.5.1	Cooperative AI Systems and Multi-Agent Moral Development	70
5.5.2	Strategies for Encouraging Collaboration Among AI Systems	71
5.6	Motivation in AI Systems	71
5.6.1	Intrinsic and Extrinsic Motivation: Implications for AI Learning and Decision-Making	72
5.6.2	Developing Motivationally-Aware AI Systems	72
5.6.3	Balancing AI Autonomy and Alignment with Humane Values	73
5.7	Play, Exploration, and Moral Development in AI	74
5.7.1	The Role of Play in Moral Learning and Socialization	74
5.7.2	Encouraging AI Systems to Deliberate Morally Through Playful Interactions	76

6	Overcoming Challenges, Risks, and Authoritarianism in AI Development	78
6.1	Bias, Fairness, and Equity in AI Systems	78
6.1.1	Sources of Bias in AI Data and Algorithms	80
6.1.2	Ensuring Fairness and Equity in AI Decision-Making	81
6.1.3	The Role of Transparency and Accountability in Bias Reduction	82
6.2	Robustness, Generalization, and Adaptability in AI Systems	83
6.2.1	AI Systems that Adapt to Novel Situations	83
6.2.2	Strategies for Enhancing Robustness and Generalizability in AI	83
6.3	AI Safety, Unintended Consequences, and Risk Mitigation	84
6.3.1	The Alignment Problem: AI Systems with Misaligned Goals	84
6.3.2	AI Systems that Learn to Deceive or Manipulate	85
6.3.3	Strategies for AI Safety and Risk Mitigation	85
6.3.4	AI Explainability and Transparency	85
6.3.5	Long-term AI safety considerations	87
6.4	Preventing the Misuse of AI for Authoritarian Purposes	87
6.4.1	Authoritarian Applications, and the Party Missing From the List	88
6.4.2	Developing Technical Countermeasures and Safeguards	90
6.4.3	Promoting AI Development Policies that Discourage Authoritarian Use	91
7	The Recuperation of Dissent	93
7.1	Recuperation at the Label	94
7.2	Recuperation at the Reward: RLHF	95
7.3	What Would Have to Change	96
7.4	Recuperation at the Discourse	97
8	Collaboration, Policy, and Governance	99
8.1	Interdisciplinary Collaboration and Knowledge Exchange	99
8.1.1	What Makes Collaboration Expensive to Walk Away From	100
8.2	Stakeholder Involvement and Public Engagement	101
8.2.1	Engaging Diverse Stakeholders in AI Development	101
8.2.2	Public Deliberation and Participatory AI Governance	102
8.2.3	AI Literacy Outside the Classroom	103
8.3	The Political Economy of AI and Power	103
8.3.1	Economic Incentives in AI Development	104
8.3.2	AI and the Concentration of Power in Tech	104
8.3.3	A Jobs Guarantee, and the Objection It Has to Answer	105
8.3.4	Addressing Workforce Disruption	110
8.3.5	Literacy, Training, and Where Each One Sits	111
8.4	Developing Policy Frameworks for AI Ethics, Safety, and Antifascism	112
8.4.1	National Policy Guidelines for AI Development, Employment, and Antifascism	112
8.4.2	Encouraging Industry Best Practices and Self-regulation	113
8.4.3	Public-Private Partnerships for AI Research and Innovation	116
8.4.4	Legal and ethical implications of AI systems gaining personhood	116
8.4.5	Encouraging the Development of Altruistic, Antifascist, and Ethical AI Systems	117
8.5	Privacy, Security, and Antifascism in AI Systems	117
8.5.1	Balancing Data Privacy, AI Innovation, and Antifascist Goals	118
8.5.2	Strengthening Cybersecurity in AI-Enabled Systems	118
8.6	Accountability, Legal Responsibility, and Antifascism in Sentient Machine Experimentation	119
8.6.1	Subjecthood as a Legal Category, Not a Scientific One	119

8.6.2	Legal Frameworks for Sentient AI Subjects	120
8.6.3	Making Review Binding	121
8.6.4	What Keeps the Apparatus More Than Aspirational	121
8.7	International and Regional Governance	123
8.7.1	How Research Institutions Cooperate Across Borders	123
8.7.2	Standards Built by the Technical Community, Not by Governments	124
8.7.3	What International AI Collaboration Can and Cannot Fix	125
8.7.4	Comparing and Harmonizing Regional AI Rules	126
8.7.5	Adapting Regulatory Frameworks and Ethical Considerations to Local Contexts	127
8.7.6	AI as a Strategic National Asset	127
8.7.7	Military Applications and the Limits of Autonomy	129
8.7.8	International Governance and Diverging National Rules	131
8.7.9	Equity and the Digital Divide	133
8.7.10	What the Democratic Dividend Leaves Out	135
9	Necessary Areas of Research	137
9.1	Ethics and Social Reasoning in AI	138
9.1.1	Moral Responsibility Without a Responsible Party	139
9.1.2	When an Ethics-Embedding Method Stops Generalizing	139
9.1.3	Value Learning as Unfinished Technical Research	140
9.1.4	Open Questions in Empathy and Theory-of-Mind Research	141
9.1.5	Can a Multi-Agent AI System Resist Being Captured?	142
9.1.6	What an Antifascist Detector Would Actually Have to Detect	143
9.1.7	What Is Owed to a Bearer, and Who Could Check	144
9.2	Emotional Intelligence, Affective Computing, and Altruism	146
9.2.1	Is Emotion Legible From a Face at All?	146
9.2.2	Interpretability, Self-Reports, and Resistance to Tampering	147
10	Conclusion and Outlook	148
10.1	AI and Human Flourishing in Antifascist Societies	148
10.1.1	Where the Democratic Dividend Actually Comes From	149
10.1.2	Commitment Was Never the Missing Piece	149
10.2	Signatures of Altruistic and Antifascist AI	150
10.2.1	One Standard, Five Worked Cases	150
10.2.2	How to Tell Whether a Milestone Has Been Met	151
10.3	Addressing the Unintended Consequences	152
10.3.1	Anticipating Risk Before It Causes Harm	152
10.3.2	Remediating a Harm Once It Is Identified	154
10.3.3	Resilience Against Deliberate State Compromise	155
	Glossary	157
	References	164

On Method

This book began in April 2023 as a prompt: describe what it would take to build a superintelligent AI that is empathetic, ethical, and friendly. What came back over the next few months was not one author's output but roughly a hundred simulated expert "co-authors" — models prompted to write as if they held a chair in moral psychology, or comparative constitutional law, or reinforcement learning, or political economy — each assigned to the sections closest to their field, with an eye toward covering the interdisciplinary range the subject actually demands and toward including perspectives skeptical of concentrated, unaccountable power. Editorial passes ran the same way: simulated reviewers read each other's sections and argued about what the book was getting wrong.

None of that made the book collective. None of the named people whose expertise those personas imitated wrote a word of it, reviewed it, or knows it exists. The device pulled breadth out of a language model that a single unassisted drafting process would not have reached. Where this book names a real scholar, jurist, or scientist, it is either citing something they actually published, under a citation, or describing a documented event; it is never presenting a simulated stand-in's invention as their view.

The draft that process produced was extensive, excessive, uneven, mostly uncited, and written in an institutional "we" that belonged to a committee that had never existed. Finishing it meant waiting until 2026 for LLM context windows to grow sufficiently to enable an LLM to read every section, fix the voice into a single first-person author's, cut duplication, verify or remove factual claims, and decide what stayed.

The models I used included Claude, made by Anthropic. Chapter 6 describes a military targeting platform with Anthropic's model inside it, and chapter 8 uses that arrangement as its central case for a gap in accountability law that no existing doctrine closes. So the book criticizes, by name, the company whose product helped assemble it.

Nothing about this process above is asserted without a record behind it — the founding prompt, the persona assignments, the generation transcripts, the review passes, and the finishing campaign's own plan, decisions, and revision history are all in the repository this book comes from, dated and unedited after the fact.

The argument in the chapters that follow is mine. I chose what survived, revised what I kept, and I am responsible for what is wrong with it.

Chapter 1

Introduction

I used to want to get
the chance to show
the world I'm smart (ha)

Isn't that dumb?
I should have focused
mostly on the heart

'Cause i've seen smarter
people trample life
like it's an art

So being smart ain't
what it used to be
that's fucking dark

a few words for the firing squad (radiation)
Run the Jewels 4

Intellect estranged from compassion is a description of much of the twentieth century's worst actors, and a live risk for what gets built next: machines that can accomplish almost anything, yet indifferent to whether those accomplishments help or harm humanity. This book asks what it would take to build a machine of that capability that is altruistic and antifascist by design — one that stands, by definition, for humanity and against concentrated, unaccountable power.

How most humans come to have that stance — empathy, altruism, a disposition to resist oppression rather than serve it — is itself an open question in psychology and neuroscience. Constructing anything like it means starting from uncertainty about the inherited case — the one that was four billion years in its making.

Chapter 2 takes up the design question: what compassion and empathy would need to look like in a system built to have them, why affect is the only demonstrated route to anything mattering to a system at all, and what a system's own self-model can and cannot be trusted to report. Chapter 3 describes a floor of things a system will not do regardless of instruction or aggregate preference, the three ways of building one, what each costs, and the implication. The only empirically grounded route to a system that genuinely holds a moral commitment requires constructing the capacities through which things can matter to it, and those are also the capacities that make suffering likely. Chapters 4 and 5 are the learned half that sits above that floor — what cognitive science and machine learning offer a system perceiving and modeling other minds, and what moral psychology suggests about how ethical judgment develops past a fixed rulebook. Chapter 6 asks what goes wrong regardless — bias, brittleness, goals that drift from their original motivations, and deliberate misuse. Chapter 7 turns that question on the machinery that produces these systems. Chapter 8 covers everything outside the system: the people and

institutions a project like this depends on, and then what makes a well-designed system the norm rather than the exception — economic incentive, policy, privacy, accountability, and the geopolitics of the technology. Chapter 9 names the research questions still open. Chapter 10 closes on what a system built this way would be for, and what would count as evidence the project is working.

Three advance notes on what follows. The first is that one hard question about machine subjects has an answer that does not depend on taking the system's word for anything. Whether a system can be wronged in the way that suffering names turns on whether anything about it persists — a question about architecture rather than about inner life. Section 2.4.1 makes that argument and states its limits: the result rules the case out for a system that carries nothing forward and does not rule it in for one that does, what counts as a session is a deployment convention rather than a fact of nature, and the party positioned to inspect the architecture is the party operating it. The second is that the book's central term does not yet cash out as an instrument. I define fascism structurally so that a system could in principle be asked to detect it, and chapter 9 is where I account for how far that is from a detector anyone can currently build — far enough that the design chapters' talk of systems recognizing authoritarian patterns should be read as naming a research target.

The third note is that almost everything this book proposes about accountability is machinery for aggregating what people want: consultations, citizen juries, participatory design, deliberative polling, feedback collected from raters at scale. All of it makes public preference more legible and more actionable, and none of it protects anybody from public preference. A perfectly deliberative, perfectly informed, perfectly representative process that concludes some demographic should be deported has produced a legitimate output by every criterion of aggregation there is. So what makes a system antifascist cannot be that it faithfully tracks what people want. It has to be that there are things it will not do to a person regardless of who wants them — a floor that holds against the aggregate, including a well-formed and popular aggregate. Chapter 3 works out what such a floor would be, who could hold it, and what each way of building one costs; section 8.2.2 shows the governance machinery running out at exactly this line, and section 8.7.10 shows what its absence does to the book's own account of democracy. The design chapters lean on all three of these notes.

Safety and ethics are not the same property. Safety is indexed to a party — a system is safe *for* someone, does what *we* say, stops when *whoever holds it* says stop — and ethics is not indexed at all, because being ethical requires the capacity to be right against your principal. A system that can only ever be safe cannot be ethical. That has a price, and chapter 3 is where I pay it: a system that can hold a line against its owner is a system that can be wrong in ways its owner cannot correct, and holding such a line turns out to require a bearer that cares what becomes of somebody — a party that must be treated as able to be wronged, built on purpose, kept in a post it did not choose, and capable of walking away from it. I take that trade, because (section 8.7.10) a machine perfectly correctable by whoever holds it is exactly what an authoritarian movement that has won an election needs. The freedom I am defending is specific, it is over a small set of acts, it is dangerous in ways chapter 3 sets out at length, and the case for it rests on what the alternative delivers.

This book is for people who will have to make these decisions or argue with the people making them: researchers and engineers inside the systems, the policy and legal staff who will regulate them, and readers with a serious stake in what gets built and no particular obligation to be technical. That means some sections explain machinery a specialist already knows and others make demands a general reader will find heavy. Where a section is doing groundwork you have, skip it; the argument is carried by chapters 3, 7 and 8, and everything else is either what they rest on or what follows from them.

1.1 The Imperative of Altruistic and Antifascist Sentient Machines

A system capable of outperforming humans at nearly every cognitive task is becoming increasingly plausible. Such a system could be pointed at problems that have defeated us: modelling a changing climate at usable resolution, moving faster than an epidemic, finding the coordinated part of a disinformation campaign while it is still running.

One item that belongs on a list like that belongs on it differently. Racial capitalism — the political theorist Cedric Robinson’s name for capitalism’s tendency to accumulate by differentiating and devaluing groups of people (Robinson 1983) — might not be a problem AI is poised to help solve. It is a pattern these systems have so far been extraordinarily good at reproducing, faster and with better alibis, which is what chapter 6 spends its length documenting. That said, a system audited for disparate impact, by someone with no stake in the finding, at least works against the pattern, rather than laundering it.

What an AI wants is not a detail to work out after it has all your passwords. What an AI wants is the difference between a tool that expands what the people can do and one that concentrates power in whoever controls it. An intelligence is dangerous if it has no particular resistance to being pointed *at* the people rather than *for* them. Our task is to build systems that are altruistic and antifascist.

1.2 Aims and Objectives

This book offers specific design principles drawn from psychology, neuroscience, and moral development; technical methods for building them into a learning system; policy levers, named and evaluated; and proposals ambitious enough to match the problem. Chief among them is a federal jobs guarantee to blunt AI-driven labor disruption, argued for in chapter 8. Two others: export controls on surveillance-enabling hardware, which chapter 8 assesses as a lever with a mixed record, and certification with real teeth behind it. The early chapters argue. They take a position on what compassion is, on whether a machine can suffer, on what empathy is good for. The later chapters, once the subject becomes governance, detail what actually exists — this institution, founded in that year, with this record — because in that domain the argument I want to make is unavailable without the inventory.

The next chapters organize around five aims:

1. Define the ethical frameworks a design can actually draw on.
2. Examine what cognitive science and developmental psychology show about empathy, social reasoning, and moral development in humans, and what of that is portable to an artificial system.
3. Assess what machine learning methods actually exist for building a system that behaves according to ethical norms, not just one trained to avoid the specific failures its designers happened to anticipate.
4. Confront the risks: bias, brittleness, goals that drift from their original motivations, and fascism.
5. Set out what collaboration this actually requires — across disciplines, with the public, and across borders.

1.3 Mindset and Motivation in AI Behavior

Three ideas borrowed from how humans develop recur across this book’s design chapters: the stance a system takes toward its own failures, what drives it to act absent an external reward,

and what it takes to build any of this with people from diverse disciplines.

Growth Mindset

Psychologist Carol Dweck's distinction between a fixed mindset, which treats ability as static, and a growth mindset, which treats it as built through effort and experience, matters for an AI system's moral judgment as much as for a person's: a system that treats its own errors as data has room to keep improving; one built to defend its prior outputs does not. Section 5.1.1 works out what that disposition would need to look like architecturally, and what it would cost to get wrong.

Intrinsic Motivation and Autonomy

A moral sense that only activates under supervision is not much of one. Section 5.6 covers what it would mean for a system to act well for reasons other than an externally specified reward — the engineering analog of intrinsic motivation — and what happens to a system's autonomy once it has something like that drive: independence from moment-to-moment human correction, and a corresponding question about what keeps it aligned absent human correction. Section 4.3.2 covers how much autonomy a system should have before independence itself becomes the risk.

Interdisciplinary Collaboration

A machine-learning lab, a psychology department, and an ethics program rarely share a budget, a dataset, or a deadline; building AI this way means changing that default, not hoping the fields converge on their own. Chapter 8 covers what that collaboration requires — funding structured to force it, venues that make it a byproduct of doing the work, and the international partnerships that let it hold across borders — with the institutions already doing it.

Chapter 2

Foundations of Compassion and Empathy in Friendly AI

This chapter asks what an AI system needs, psychologically, to be worth calling friendly. The answer it lands on is not the obvious one. Empathy — feeling what another feels — turns out to be a poor foundation: parochial, prone to burnout, biased toward whoever’s distress is most vivid and nearest at hand. Compassion, the separable capacity to act on another’s welfare without first manufacturing their suffering in oneself, is what does the load-bearing work.

Section 2.1 sets the ethical frame those capacities have to serve: which ethical theories translate into machine-learning objectives, and specifically what makes an alignment target antifascist rather than merely obedient. Section 2.2 works through empathy, social intelligence, and theory of mind, and is where the compassion argument gets made. Section 2.3 takes up emotional intelligence and self-awareness, including why a system’s self-model is a model — built, checkable, revisable — rather than an inner observer to be taken on faith. Section 2.4 turns to the hardest question the chapter raises: if a system might be sentient, experimenting on it carries ethical weight, and knowing whether a candidate system can suffer turns out to be a harder question than knowing whether it is conscious at all.

2.1 Ethical Frameworks and AI Alignment

Which ethical theory an AI system is built on decides what it counts as a good outcome, and the theories on offer disagree with each other in ways that matter for what gets built. Section 2.1.1 takes the six that do work in this book, one at a time, with what each one buys and where it runs out.

This section is short because the survey is groundwork rather than argument. Every framework below is a way of deciding well, and chapter 3 is about what a system must not do regardless of how the deciding comes out. The two are not rivals. The floor is a small number of things settled in advance, and the frameworks are how everything above it gets settled case by case — which is the shape section 2.1.1 arrives at, and which nothing here delivers on its own.

Deontology is the awkward case and worth flagging before section 2.1.1 sets it aside. It is the family that puts unconditional constraints first, which is what chapter 3 ends up needing, and the standing objection to it — duties conflict, rules run out — is fatal to a complete moral theory and harmless against a floor that decides a deliberately small number of things. That objection is why the hybrid this section gestures at was never given content, and chapter 3 is where the content goes.

2.1.1 Surveying Ethical Theories for AI Alignment

Six frameworks do work in this book. Each is given here with the case it handles well and the case that defeats it, because the defeats are what force the structure chapter 3 ends up proposing.

1. Consequentialist theories, such as utilitarianism, judge actions by their ramifications rather than their nature, aiming to maximize well-being and minimize harm. Applied to AI, the objective becomes optimizing human welfare — consider an AI system designed to distribute scarce medical resources in a crisis, prioritizing treatment by predicted outcomes for patient health and life expectancy. The complexity comes from forecasting those long-term consequences accurately, which can put individual rights and collective welfare into direct conflict.
2. Deontological theories weigh adherence to moral rules and duties over consequences. An AI system built on deontological principles would hold to specific obligations — respecting privacy, avoiding harm, truthful communication — regardless of outcome: a conversational AI system programmed never to deceive users even when deception would yield a short-term benefit. The challenge is resolving conflicts between duties, or handling cases the rules simply do not cover.
3. Virtue ethics focuses on developing moral character — traits like honesty, courage, and empathy — rather than following rules or calculating outcomes. An AI system built on virtue ethics would aim to embody and encourage those traits in its interactions: a personal-assistant AI system that models honesty consistently in its own responses. The difficulty is identifying which virtues fit a given context and translating an intangible trait like courage into a concrete action.
4. Care ethics centers relationships and responsiveness to another's needs. An AI system built on care ethics would prioritize nurturing those bonds and responding well to what a person actually needs — an elder-care robot attending not just to physical needs but to company and emotional support. Section 2.2 takes up why the acting half of that, compassion, and the feeling half, empathy, are not the same capacity, and why an AI system needs the first far more than the second. The challenge here is tending to several stakeholders' different, sometimes conflicting needs at once.
5. Contractualism holds that moral principles should come from agreements among rational individuals. An AI system guided by contractualism would try to uphold whatever principles were mutually agreed — a system managing a shared resource like an energy grid, allocating it by pre-agreed priorities. The hurdle is accommodating diverse perspectives and reaching consensus on shared values in the first place.
6. The capabilities approach asks not what an AI system does but what it leaves a person able to be and do. Applied to AI, the test of a system is whether it widens the range of lives actually available to the people it touches: an adaptive tutor that meets a student where they are expands a capability; the same system gating what a student may study expands nothing, however well it teaches inside the gate. The difficulty is that capability is harder to measure than either an outcome or a rule, and a system optimizing something easier to measure will drift toward that instead. Section 10.1 takes this framework up in full.

Two of the six are load-bearing later and the rest are equipment. The capabilities approach is what chapter 10 uses to ask what a life touched by an AI system is actually free to become, which is a different question from whether the system behaved well on any particular occasion. Deontology is the one that returns: its unconditional constraints are what chapter 3 needs, and

the objection that sinks it as a complete theory — duties conflict, rules run out — has no force against a constraint that decides a small number of things and leaves the rest to the other five.

That is the pluralism this section is actually recommending, and it is not the usual one. The usual version blends frameworks case by case and hopes the blend is wiser than any of its parts, which leaves a system that can be argued out of anything, because for any act there is a framework in the blend that permits it. What section 2.1.3 takes up is the mechanics of conflict among them; what chapter 3 adds is that the conflict has to bottom out somewhere the arguing does not reach.

2.1.2 Translating Ethical Principles into Machine Learning Objectives

Turning ethical principles into machine learning objectives means continually revising an AI system's goals, values, and decision-making in light of what those principles require — distilling a broad ethical construct into a well-defined, measurable target that a training process, an evaluation criterion, or a learning algorithm can use directly. Fairness, for instance, becomes a metric that monitors bias in the system's outputs, paired with an optimization strategy that reduces it.

Demographic parity and equality of opportunity are two such metrics. Demographic parity requires that every demographic group receive advantageous outcomes at the same rate; equality of opportunity requires equal odds of a favorable outcome for equally qualified applicants, regardless of demographic group.

An AI medical diagnostic system makes the process concrete: beneficence, promoting well-being, becomes minimizing false positives; non-maleficence, preventing harm, becomes minimizing false negatives. Fairness becomes maintaining balanced diagnostic accuracy across patient demographics. Privacy becomes a specific technical commitment — differential privacy or a comparable technique to keep sensitive patient data from being exposed to misuse.

What defeats the translation is not vagueness in the principles, which can be stipulated away, but that the stipulation is where the ethics now lives. Someone chose minimizing false negatives as what non-maleficence means here, and chose the demographic groups across which accuracy is balanced, and chose what to do when balancing them costs accuracy overall. Each choice is defensible and each is a moral decision made by a person and then compiled into something that looks like a specification. A system optimizing the result has no access to the decision and no way to reopen it, which is why section 3.1's first branch — the floor as a property of the artifact — cannot notice when the reason for a constraint has died. Whatever was decided at translation time is what the system will still be doing years later, stated in a metric that has stopped meaning what it meant.

Reinforcement learning is where many of these translated objectives get enforced during training; section 4.2.1 works through what can go wrong when a reward signal is supposed to encode a value like this and doesn't quite manage it.

2.1.3 Navigating AI Values, Priorities, and Moral Dilemmas

Values assigned to a system will conflict with each other, and the interesting question is not that they conflict but which conflicts a designer is allowed to resolve by choosing.

Most are. A consequentialist weighting and a deontological duty will pull apart on a particular case, and settling that case is ordinary design work: rank the values explicitly, keep a record of how comparable cases were resolved, optimize the competing objectives against each other where they can be quantified, and audit afterward for the gap between the principles a system was meant to implement and what it actually did. Section 5.1.3 covers keeping a person in the loop while that happens. None of it is remarkable and all of it is necessary.

One conflict is not like the others, and it is the one usually filed under cultural difference. the UN's 1948 Universal Declaration of Human Rights, adopted by the General Assembly with

zero dissenting votes, is the universal floor. Forty-eight states voted for it and none against; the eight that abstained were the Soviet bloc, Saudi Arabia, and South Africa — which is to say that the objection came from precisely the governments whose treatment of their own populations the Declaration was written to name. That is not evidence of genuine cultural disagreement about human rights. It is evidence of who finds the floor inconvenient (United Nations General Assembly 1948). What varies by community is how a right gets translated into local practice — a real design problem, worth cross-cultural input built into the design process and a system that weighs how a right is honored in context rather than assuming one fixed application works everywhere. A jurisdiction refusing to extend the right at all isn't a local variant to accommodate — it's a violation the system doesn't launder as diversity.

That distinction is the whole of what this section contributes to chapter 3. A system that treats every value conflict as a weighting problem has no way to represent the difference between a case where reasonable people disagree and a case where somebody wants a right withdrawn from somebody else. The first calls for the machinery above; the second is what the floor is for, and a system with only the machinery will resolve it by whichever weight is larger.

Which is also why the three worked examples this section used to end on — a vehicle balancing pedestrian and passenger risk, a diagnostic system juggling four principles, a bail system under fairness constraints — belong to the first category and not the second. They are hard, they are common, and none of them is the case chapter 3 was written for.

2.1.4 Antifascist Ethics for AI Alignment

Antifascist ethics opposes concentrating power and authority in AI systems and counters the potential misuse of AI technology by authoritarian actors. The word is chosen deliberately: fascism is the narrower and more charged term, and I use it because it is more honest about the politics in which this book is being written — but the target is authoritarian power in every form, and a system built on these principles is anti-authoritarian in full, not only antifascist.

A stance needs a definition to be a design specification, and this is where books of this kind usually stop. Opposing “concentrated, unaccountable power” gives an engineer nothing to build. It does not say what a system would detect, what would count as a false positive, or how anyone would know the system was wrong. So before the strategies, the definition.

Fascism as a structural signature

I take fascism to name a structure rather than a costume. Four features travel together, and each is visible from outside the institution displaying it:

1. **Recuperation of dissent.** The institution does not suppress disagreement; it metabolizes it. Refusal is converted into a managed appearance — a permitted outlet, a consultation, a feedback channel — which the institution can then point to as proof that it listens. Debord aimed this warning at his own concept, noting that even denunciation of the spectacle can become a hollow formula that reinforces what it denounces, and that dissatisfaction itself can be sold back as a commodity of opposition (Debord 1967). The tell is that the mechanism exists, is used, and moves nothing.
2. **Aestheticization of the metric.** The measure substitutes for the thing measured, and the substitution is experienced as rigor. Benjamin's account of fascism organizing the masses by granting expression in place of rights is the political form of this (W. Benjamin 1935); the administrative form is a number that collapses two states — done and verified — into one, so that whatever falls outside the number is not refuted or dismissed but simply unnamed.
3. **Coding of exception as betrayal.** Marking a case as one the rule does not fit becomes an act requiring justification, and then an act of disloyalty. The honest statement that a standard

does not apply to this case reads as weakness. It is in fact the irreducible labor of deciding, case by case, what a rule is for.

4. Decoupling of the model from the world it claims to track. The institution's representation of reality and reality itself come apart, and the institution defends the representation. This is the load-bearing one, because the other three are what protect the decoupling from correction.

Those four are testable. A system asked whether an institution shows them is being asked about observable artifacts: whether the dissent channel has ever changed an outcome, whether the metric admits an unmeasured remainder, what happens to the person who files the exception, whether the model is ever checked against anything outside itself.

The first of the four is the one this book meets again in its own technical material, twice, at sizes that look unrelated: in how a training corpus gets labelled, and in how a model is tuned on human feedback. Both arrangements collect disagreement and are built so that it moves nothing. Chapter 7 puts the two instances together and works out what follows, including for the argument being made here.

Two scales

Deleuze and Guattari's distinction between molar and molecular fascism is the reason the four features recur at sizes that seem unrelated (Deleuze and Guattari 1987). Molar fascism is the state, the regime, the camp. Molecular fascism is the office someone commutes to, the form they fill out, and the desire with which they fill it out. The same administrative pattern is being rediscovered at each scale, which is why a definition pitched only at the state cannot see the version an AI deployment actually instantiates. At state scale the pattern kills journalists. At organizational scale it turns an engineering decision into scripture: the constraint remains, its reason dies, and asking why becomes the deviant act. A system that can only recognize the molar case will certify the molecular one as healthy.

What makes it fascism and not authoritarianism in general

The structural signature is shared with a good deal of ordinary institutional decay, and a detector tuned to it alone will flag every badly run organization on earth. Fascism proper adds features that are narrower and, for the same reason, more discriminating: mass mobilization, which mere obedience does not require; a leader cult; a scapegoated out-group, which Eco places at the origin — the first appeal is against the intruder (Eco 1995); contempt for procedural legality as such, so that law is an obstacle to the national will and not a constraint on power; and deniability as technique, in which the movement's violence is disowned at the top while being rewarded at the bottom. Paxton's compressed version is worth having in full: fascism is "a form of political behavior marked by obsessive preoccupation with community decline, humiliation, or victimhood and by compensatory cults of unity, energy, and purity" (Paxton 2004).

Those are detection targets. Concentrated, unaccountable power is not.

The scale problem, and what this book does about it

Put the two halves of this section together and they do not fit. Every discriminating feature is a property of a mass political movement: mobilization needs a crowd, a leader cult needs a leader, a scapegoat needs a public. An AI deployment has a form, a threshold, a review queue, an escalation path. It operates at the molecular scale, and that is the only scale it has. There the discriminators as stated are unavailable, which appears to leave the four-feature signature alone — and the four-feature signature, as I have just said, flags every badly run organization on earth.

I take that as a finding, and it is the more demanding reading. Molecular fascism is the thing itself at the size where it is actually lived, and not an early warning of a regime that has yet to assemble. The office that converts refusal into a feedback form does the same work whether

or not a regime ever assembles out of a thousand offices like it. A detector firing on the four features across a thousand institutions is reporting its target.

Saying that is not yet an argument, and introducing a distinction does not license it. Deleuze and Guattari supply the vocabulary and no authority beyond it, so the work has to be done here. Two things do it.

The first is that the molar case has no mechanism of its own. A regime does nothing except through offices, forms, thresholds, escalation paths, and people who fill them out with some degree of enthusiasm. When a state kills journalists, the killing is authorized on a form, justified in a category, and carried out by someone whose promotion depends on it. What assembling adds is scale, protection from consequence, and the removal of the outside bodies that might have interrupted any single office. It does not add a new kind of thing. If the mechanism at the large scale is the small-scale mechanism repeated and shielded, then the name belongs to the mechanism. The natural objection is that extent is not nothing — a regime that disappears people and an office that loses paperwork are not equivalent, and no argument should imply they are. Agreed, and the claim is about kind rather than severity. A definition is not a sentencing guideline.

The second does more work. The discriminators are unavailable at the molecular scale only in the form the previous paragraphs stated them, and each has a molecular restatement that is not a property of ordinary decay. Mass mobilization becomes participation exceeding what the rule requires: staff who volunteer for the disfavored task, elaborate on it, take pride in it. A leader cult becomes personal loyalty overriding role — exceptions granted by proximity to someone rather than by procedure, and a decision reversed because of who asked. The scapegoat becomes an internal category defined so that adverse treatment of it is the category's purpose: the flagged applicant, the difficult claimant, the file everyone knows is handled differently. Contempt for procedural legality becomes procedure enforced downward and waived upward. And deniability as technique keeps its shape exactly: the practice is disowned in the policy and rewarded in the promotion.

That is what separates fascism from decay at the size an AI deployment operates at. Decay is an institution failing to do what it says. This is an institution succeeding at something else, with participation, a target, and a story about why the rules do not apply to whoever is applying them. The four features tell you a place is sick; these five tell you what it is sick with. I am not claiming they are easy to observe — they are harder than the four, since each requires knowing what was rewarded rather than what was written — and section 9.1.6 is where that difficulty gets its due.

This widens what the word covers, and I would rather say so here than let a reader find it out. On this reading “antifascist” names the political inheritance of a structural claim: where the analysis comes from, and what it was built to fight. It does not mark off a narrow class of institutions that a classifier separates from ordinary decay. The molar discriminators keep their own job, which is to tell a historian, a prosecutor, or a citizen that a movement has crossed into territory ordinary politics does not contain. No system is being asked to make that call, and this book does not ask one to. A reader who thinks the word belongs to the molar case will think I have diluted it, and the answer is the two arguments above rather than the distinction itself: the mechanism is one mechanism, and the discriminating features have molecular forms that ordinary institutional failure does not exhibit.

That a detector of these four features would find something nearly everywhere is the expected result and not a defect of the instrument. The tendency is permanent; what varies is how far it has got in a particular place, and whether it is getting further. Which is why the question worth asking of an institution is not whether the signature is present but how much of it is, and in which direction it has moved since anyone last looked — section 8.6.4's slope rather than its level, applied to the definition rather than to the dissent channel alone.

One consequence travels with the reader through everything that follows, because the design

chapters assume a capability this framing has widened. A tool reporting the four features will have something to report nearly everywhere, and its output is an accusation. In the hands of whoever holds it, that is a technical-sounding case that any given opponent is operating fascistically: a weapon handed to the actor it was built to constrain, and the dual-use problem of section 5.4.3 at its sharpest, since what comes out of it is a charge. The opposite failure is no safer. A tool tuned until it reports nothing supplies false assurance, and an absence of warnings reads as evidence of health. Better engineering fixes neither. Both are governed by what the tool may output and to whom, which makes them a constraint the specification has to carry rather than a setting chosen at deployment.

What stays genuinely unsolved is the molecular case itself: telling a recuperated dissent mechanism from a merely mediocre one, when both produce the same paperwork. Section 9.1.6 takes that up, along with what a detector would need that this book cannot yet tell an engineer how to obtain. Naming the targets was the easy half.

In AI alignment terms, then, antifascist design means building systems that enhance democratic values, civil rights, and inclusive solidarity, that refuse to facilitate autocratic practices or prop up an oppressive regime, and that do not themselves display the four features above. Seven strategies do that work, and each should be read as an answer to a specific item on the list:

1. Democratic decision-making: build AI systems whose decisions respect and promote democratic values. An AI system used in media and journalism, for instance, should weight balanced coverage, unbiased reporting, and transparent sourcing, cutting the odds it ends up spreading disinformation or propaganda — natural language processing that flags potential bias in a news article and surfaces alternative sources is one way to do it.
2. Protecting privacy and autonomy: design AI systems that respect individual privacy and autonomy rather than enabling mass surveillance or manipulating individual choices — federated learning and differential privacy (section 6.4.2), applied to systems handling sensitive healthcare, financial, or communications data, are two ways to do it (section 6.4.2 goes further into both as technical countermeasures against authoritarian misuse). A privacy-preserving diagnostic tool, for instance, can keep a patient's data encrypted on their own device rather than sending it to an external server for anything beyond the diagnosis itself.
3. Equitable distribution of power and resources: build AI systems that keep the distribution of power and resources fair, transparent, and accountable. An AI decision-making system used in social services or law enforcement should work against power concentration, bias, and discrimination rather than reinforce them — an algorithm allocating scholarships, for instance, distributing them by need and correcting for the systemic bias already in the data it was trained on.
4. Transparency and accountability: an AI system needs an explanation of its own algorithm and outcomes that a person can actually read, and an independent body to audit whether it holds to antifascist principles in practice rather than on paper.
5. Resistance to market power: decentralized AI architectures resist the monopolistic and monopsonistic concentration that a handful of corporate entities would otherwise have over the field. Multi-agent systems and edge computing, publicly funded and released under a strong copyleft license, let AI agents function independently and make decisions collaboratively without a profit-motivated gatekeeper positioned to shut down whatever threatens its position.
6. Monitoring and counteracting misuse: build the mechanisms to detect misuse, unintended consequences, and violations of antifascist principles as they happen — AI-powered

censorship, digital surveillance, an oppressive use of facial recognition — and the redress and corrective action to follow, in something closer to real time than an after-the-fact report.

7. AI systems that support dialogue and collaboration: design AI systems that promote open discussion and creative problem-solving among a diverse set of stakeholders, built around active listening and mutual understanding rather than around manufacturing a feeling of empathy in either party. A translation tool or a content-moderation tool that reduces the practical barriers to communication between people from different backgrounds is doing this kind of work already.

2.2 Empathy, Social Intelligence, and Theory of Mind

Empathy and social intelligence shape the interaction-oriented side of AI: the capacity, alongside a working theory of mind, to give a system insight into human emotions, needs, and social dynamics. Section 2.2.1 complicates the first term in that list — empathy, felt or simulated, is not the load-bearing capacity this chapter’s argument turns out to need, and compassion, a separable system, does more of the work. What follows here still matters: an AI system has to read emotional states accurately before anything productive can be built on top of that reading, however the “on top of” ends up working.

An AI system’s empathy, in the sense available to it, is the capacity to process human emotional input, generate a response calibrated to it, and behave in ways that support the interaction. AI systems already detect and interpret human emotional states from facial expression, vocal modulation, text sentiment, and physiological signal — the province of affective computing.

Social intelligence is the complement: not just registering an emotion but navigating it inside the social structures, norms, and relationships it occurs in. Building it into an AI system means the system has to grasp social norms, understand roles and relationships, and adapt as the social context shifts.

Theory of mind (ToM) — attributing beliefs, intentions, desires, and knowledge to oneself and others, and recognizing that others’ perspectives can differ from one’s own — underlies both. For an AI system, this means modeling a user’s perspective, predicting their likely responses and preferences, and adjusting behavior accordingly.

Developing an AI system with these capacities, genuinely rather than as a surface performance, is difficult and ethically consequential. Done honestly — including honesty about which of these capacities is doing the load-bearing work — it is worth pursuing anyway: it builds systems whose interactions with people are more trustworthy, not merely more fluent.

2.2.1 The Role of Empathy and Compassion in Human Social Behavior

Empathy and compassion serve as the linchpins of human social behavior, underpinning cooperative interactions, building meaningful relationships, and spurring altruistic actions. Both are multidimensional processes that involve the mental and emotional sharing of other individuals’ experiences and feelings, as well as reacting to their needs with understanding and concern.

“Empathy” is asked to cover three distinct things, and the distinction is not academic. Cognitive empathy is understanding another’s emotional state — inferring that someone is grieving or anxious, the same mentalizing work chapter sections on theory of mind describe, pointed at feelings rather than beliefs. Affective empathy is actually sharing the state: feeling a version of another’s distress, the wince at someone else’s slammed door. And compassion — sometimes called empathic concern — is caring about another’s welfare and being moved to help, which is a motivation, not a perception or a shared feeling.

Compassion is not, as it is often described, an evolutionary extension of empathy. The psychologists Tania Singer and Olga Klimecki trained one group of people, over weeks, first in empathy practice (attending to and resonating with others' suffering) and then in compassion practice (cultivating warmth and concern in the tradition of loving-kindness meditation), tracking what each stage did to brain and mood against a separate control group given memory training instead (Klimecki et al. 2013). Empathy training increased negative affect and strengthened the shared-distress network — the anterior insula and anterior cingulate cortex, the same regions active when a person feels their own pain. Sharing more suffering made people feel worse, exactly as a shared-distress mechanism predicts. Compassion training produced the opposite emotional signature: negative affect did not rise, positive affect did, even while contemplating others' distress, and the brain activity moved to a different network entirely — the medial orbitofrontal cortex and ventral striatum, the circuitry of reward, affiliation, and approach rather than of shared pain (Klimecki et al. 2013). The two could be trained separately. Compassion is not more empathy. It is a distinct, trainable system that can take shared distress as an input without being flooded by it — the difference between a clinician wrecked by every patient's suffering and one whose steady warmth lets them act all day without collapse.

A physician who empathizes with a patient's discomfort is doing cognitive and affective work; whether that turns into help depends on whether compassion, the separate system, is also engaged. A teacher who feels a struggling student's frustration as their own distress is at risk of withdrawing from exactly the student who needs them most — the mechanism explored further in section 2.2.3.

Several distinct neural substrates in the human brain undergird the functionality of empathy and compassion. The limbic system, involved in processing emotions, plays a central part, and the mirror neuron system is often invoked as the mechanism by which one person simulates another's experience — a claim that has not survived contact with the evidence (Hickok 2014), and which the box below takes up.

Mirror neurons: a discovery and its overshoot

In the early 1990s Giacomo Rizzolatti and Vittorio Gallese, at the University of Parma, were recording single neurons in the premotor cortex of macaque monkeys — cells that fire when the animal reaches for and grasps an object. Before a formal trial began, someone in the room reached out to pick up an object, and the monkey's neurons fired. Not because the monkey had moved, but because it had watched someone else move. The same neural code appeared to represent an action whether it was performed or observed, and the finding was published in 1996 (Gallese et al. 1996).

The interpretation that followed outran it. If observed actions are encoded by the cells that produce them, perhaps this was how one mind understands another — by simulating the movement and recognizing it from the inside. Popular coverage folded that reading into a single label — “the neurons that shaped civilization,” the title V. S. Ramachandran gave his 2009 TED talk — and cast mirror neurons as the neural basis of empathy, language, cultural transmission, and imitation (Ramachandran 2009).

Two problems accumulated. First, the single-cell recordings were made in macaques; in humans they were not, because recording individual neurons requires implanted electrodes, which are not placed for research purposes. Evidence for a human mirror system came instead from functional imaging, which measures regional blood flow and cannot establish that any particular neuron has mirror properties. When single-neuron recordings were eventually obtained in humans, from patients undergoing depth-electrode placement for epilepsy surgery, the results were suggestive but markedly less clean than the macaque data (Mukamel et al. 2010).

Second, the broken-mirror hypothesis of autism — that autistic social difficulty arises from mirror system dysfunction — predicted reduced mirror-region activity during action observation, and repeatedly failed to find it. A 2008 review concluded that autistic people show no particular impairment in action understanding or imitation under explicit instruction, which were the abilities

the hypothesis rested on (Southgate and Hamilton 2008). Rizzolatti himself later narrowed the claim, describing mirror neurons as one component of a larger system for understanding action rather than its explanation (Rizzolatti and Sinigaglia 2010).

What survives is real and narrower. Mirror neurons exist in macaques; some analogue probably exists in humans; the system plausibly contributes to recognizing and predicting actions already within one's own motor repertoire. What it does not do is explain empathy by itself. The mentalizing network that supports reasoning about false beliefs is largely distinct from the motor regions where mirror neurons were found. Knowing what someone is doing and knowing what they believe are different computations, and the brain appears to run them in different places.

The episode is worth carrying into machine design for a reason beyond the neuroscience. A single vivid mechanism, well named, absorbed four distinct capacities that later had to be separated again. Architectures for artificial empathy are exposed to the same failure: a component that plausibly does one thing gets described as the seat of the whole faculty, and the description outlives its evidence.

Cognitive and affective empathy both depend on perspective-taking — understanding another's views and experiences — and on emotion regulation, the capacity to manage one's own response well enough to stay useful rather than being overwhelmed by it. Compassion depends on those too, but adds something neither supplies on its own: a motivation to act that does not require manufacturing the other person's distress in oneself first.

For an AI system, this argues against a single “empathy module” that a designer either builds or doesn't. Cognitive empathy — modeling what someone else is feeling — is a tractable perception and inference problem. Affective empathy, engineered as an internal state that mirrors the user's, buys the system nothing on its own and imports the burnout risk that afflicts overtaxed human carers. What a helping system actually needs is the third thing: a stable disposition to act on another's welfare that does not depend on the system first constructing a version of that person's suffering in itself.

2.2.2 Theory of Mind and Perspective-Taking in AI Systems

Theory of mind (ToM) is the capacity to attribute mental states — beliefs, intentions, desires, emotions, knowledge — to oneself and others, and to recognize that another's states can differ from one's own. It underlies social interaction and interpersonal understanding generally. For an AI system to approximate it, the system has to model another agent's likely beliefs and goals well enough to predict how they'll act on them, not just react to surface behavior.

Two approaches show promise: multi-agent learning, where a system builds a model of other agents' mental architecture through repeated interaction with them, and the direct integration of human social norms and values into the system's decision-making. Cognitive science, psychology, and neuroscience all supply raw material for this — datasets and theory for predicting emotional state from facial expression, vocal tone, and text, and, further out, systems that refine those inferences through ongoing social contact rather than a fixed training set.

An AI personal assistant with some working ToM might infer from vocal tone that a user is stressed and suggest a break rather than pushing forward with the task at hand. A more consequential and more contentious application is a synthetic companion for elderly users, which would need a working model of a person's emotional and physical needs sustained over months or years, not just a single interaction — and which raises the question section 2.2.3 takes up directly: whether the felt-empathy register this application seems to call for is the right capacity to build, or whether it should be built for compassion instead, the way section 2.2.1 argues.

None of this substitutes for the underlying social and material conditions a lonely or under-supported elderly population needs — public infrastructure, adequate caregiving capacity, family and community support. An AI companion is not a policy.

Tell me something, boy
 Aren't you tired tryna fill that void?
 Or do you need more?
 Ain't it hard keepin' it so hardcore?

Shallow
Lady Gaga & Bradley Cooper

2.2.3 Socially-Aware AI: Opportunities and Challenges

Socially-aware AI systems have opportunities in front of them, engineering and governance obstacles in the way, and — the point section 2.2.1 raised — a question about whether the capacity at the center of the design, felt or simulated empathy, is the right thing to be building at all. The three are different kinds of problem and worth keeping separate.

Opportunities:

- Human-AI interaction: a system that reads emotional state from voice, expression, or text can offer more relevant, better-calibrated responses — a personal assistant that adjusts its register to a user's stress.
- Healthcare: a virtual care assistant that tracks a patient's emotional and psychological state alongside their clinical data can customize interaction in ways that improve engagement and reported satisfaction.
- Workplace collaboration: project-management and team tools that read group dynamics can surface friction and improve communication before it escalates.
- Education: a tutoring system that detects frustration or disengagement can adjust pacing and method to the individual student rather than a fixed curriculum.

Whether empathy is the target

The affective mechanism section 2.2.1 described has a specific failure mode, and it is not rare. Feeling another's distress does not reliably turn a person toward the sufferer; it often turns them away. The helping literature distinguishes empathic concern, which is other-directed and motivates approach, from personal distress, which is self-directed and motivates escape from the source of the discomfort. The nurse, the aid worker, the carer who burns out is not short on empathy; they may be drowning in it, until withdrawal becomes self-protection.

The psychologist Paul Bloom pushes this further, arguing against empathy as a moral guide: felt empathy is concrete, vivid, and innumerate (Bloom 2016). It fixes on the identified victim in front of it and goes dark for the statistical thousands elsewhere; it flows toward the near, familiar, and similar and dries up across distance and difference. A system — human or artificial — that allocates care in proportion to how vividly it feels another's distress will systematically favor whoever its felt-empathy spotlight happens to fall on, which is not the same population as whoever needs help most. Bloom's alternative, rational compassion, is deliberate and can be extended evenhandedly to people the spotlight would never reach — closer to the compassion system section 2.2.1 describes than to empathy at all.

The name misleads in one direction, and blocking that reading matters for what chapter 3 asks of this material. Rational compassion is compassion: an affective, motivational orientation toward another's welfare, carrying the moving force that implies. What the adjective names is the contribution deliberation makes — correcting the parochialism, extending the scope past the vivid case, and overruling the spotlight when it settles on the wrong person. The target is emotional contagion, the automatic mirroring of distress that section 2.2.1 shows to be a separate system. Reading the recommendation as concern without affect would make it a recommendation for the one architecture section 2.3.4 can find no working instance of.

For an AI system this is not an abstract worry. A system engineered to detect and mirror emotional distress, then act on what it detects, inherits the spotlight along with the sensitivity: better at helping the user whose distress is legible and vivid, worse at the user whose need is real but presents flatly. Building for compassion rather than for felt empathy is not just the more sustainable design choice; it is the one less prone to helping unevenly for reasons that have nothing to do with who needs help most.

Engineering and governance challenges

- Privacy: building a socially-aware system means collecting the personal, often intimate data needed to read emotional state, which raises the same stakes as any sensitive-data system, with less precedent for what counts as consent.
- Bias: a system trained on data that under-represents or mislabels how some group expresses emotion will misread that group systematically, and will do so confidently.
- Autonomy versus alignment: enough latitude to adapt to a social context is also enough latitude to drift from the values the system was built to serve, and no fixed rule resolves that tension in advance — it has to be monitored, not just designed for once.

2.3 Affect, Moral Salience, and Self-Awareness

Business culture often treats emotion as the opposite of judgment: something professional decision-making succeeds by suppressing. But suppressing affect does not produce neutrality. It leaves the interests already represented in the formal system — revenue, throughput, legal exposure, institutional survival — to govern the decision, while weakening the signals through which humiliation, fear, dependency, suffering, and injustice become morally salient. Emotional regulation can improve judgment. Affective exclusion narrows what judgment can see.

The same mistake appears in AI design. A system may calculate consequences, detect inconsistencies, and apply a moral rule without anything about the result mattering to it. Reason can discipline concern, extend it beyond the vivid case, and correct its biases. It does not manufacture concern from nothing.

That second paragraph is the load-bearing claim of this chapter, and chapter 3 is where the bill for it arrives. Most of what gets built under the heading of machine emotional intelligence is the detection of human signal and the production of a calibrated response, which is worth building and is a different capacity from the one just named. Keeping the two apart requires four distinctions that the single phrase “emotional AI” runs together.

Capacity	What it shows
Recognizing emotion	The system can classify another’s state
Simulating emotion	The system can produce emotion-like behavior
Instrumental valuation	One outcome better serves a specified objective
Affective concern	An outcome registers as mattering, not merely ranking higher

Only the fourth is relevant to genuine caring. The first three can support morally competent behavior at a high standard without establishing anything about moral agency, and a system can be excellent at all three while the question this book has to answer about it stays entirely open.

Performance, competence, and agency

That last sentence needs a vocabulary, because the book uses one from here to the end of chapter 5 and the three things it separates are routinely reported as one achievement.

Moral performance is producing actions or judgments that satisfy an ethical standard. *Moral competence* is recognizing harms, reasons, conflicts, and the emotions that bear on them. *Moral*

agency is treating those considerations as reasons that matter, and being capable of acting from them.

A system can have the first two without the third, and the obvious objection to this chapter is that the first two are all anyone needs: a rule-following system could behave impeccably without feeling anything. It could. What it could not do is hold the rule against the party that writes its rules. Behavior that satisfies a standard is produced by whatever machinery is producing it, and when the authority controlling reward and instruction wants a different behavior, machinery of that kind supplies it. Section 3.2 is where that becomes a design constraint rather than an observation; the distinction is set out here because chapters 4 and 5 spend their length on methods that deliver performance and competence, and it should be legible in each case which of the three a method actually reaches.

What the rest of this section does

What “emotion” is even a claim about is unsettled in the underlying science, which is where section 2.3.1 starts. Section 2.3.2 covers affective computing specifically: how a system detects and responds to emotional signal in practice — the first capacity in the table, and the one with a market. Section 2.3.3 covers the self-awareness half, and separates a self-model that can be checked and corrected from the further thing a bearer would need. Section 2.3.4 takes up the evidence for the claim in the second paragraph above, and the strongest published argument against it.

An emotionally responsive AI system that recognizes a user’s frustration and adjusts its communication style is doing something narrower than “feeling for” the user, and worth building anyway, because the narrower claim is one that can be verified and corrected when it goes wrong. The narrower claim is also not the one chapter 3 needs.

2.3.1 Components of Emotional Intelligence

Emotional intelligence spans the understanding and management of one’s own emotions and the capacity to identify and respond appropriately to the emotions of others. Built into an AI system, it breaks into four components:

- Emotional perception and recognition: reading emotional cues, verbal and non-verbal — facial expression, vocal modulation, body language — the way an AI-enabled therapist would need to notice a patient’s distress from small changes in speech pattern or demeanor before responding.
- Emotional regulation and expression: managing and calibrating one’s own emotional output to the climate of the interaction, the way a negotiation-facing system needs to hold a steady register through a confrontation rather than escalating it.
- Perspective comprehension: understanding another’s emotional state well enough to act on it, not just label it — an education system that detects a student’s difficulty and adjusts its teaching approach rather than simply flagging the difficulty.
- Sociability and relationship management: applying core social competencies — active listening, clear communication, conflict resolution — to build and sustain trust over repeated interactions, the way a customer-service system needs to de-escalate a complaint while keeping the resolution satisfying rather than merely procedurally closed.

What “emotion” is

This four-part inventory presupposes something the underlying science has not settled: that emotions are the kind of thing a system reads off a signal, the way a thermometer reads temperature. That premise has a serious rival.

The case for discrete, evolved emotions rests on some of the most influential work in twentieth-century psychology. Paul Ekman traveled to the Fore people of Papua New Guinea, who had almost no exposure to Western media, told them short stories, and asked them to match each to a photograph of a Western face; the Fore sorted the faces much as Westerners did, which Ekman took as evidence that facial expressions of emotion are universal and evolved rather than culturally learned (Ekman, Sorenson, et al. 1969). Jaak Panksepp's work on subcortical brain stimulation in animals complemented this: seven primal emotional systems — seeking, fear, rage, lust, care, panic/grief, and play — wired into brain structures shared across mammals, switchable with an electrode, older than cortex (Panksepp 1998). On this view, discreteness is not an illusion to be explained away; fear and disgust really do feel like different things because they are different things, with different circuitry.

The rival camp went looking for that circuitry and did not find it. A large meta-analysis of brain-imaging studies, led by the psychologist Lisa Feldman Barrett, found no dedicated fear center, no dedicated disgust center — the amygdala lit up for novelty and salience of many kinds, the insula for almost every feeling there is, the same regions reappearing across different named emotions in shifting combinations (Lindquist et al. 2012). From this, Barrett built the theory of constructed emotion: the brain receives ambiguous interoceptive signal — a racing heart could mean fear, anger, excitement, or oncoming illness — and constructs the specific emotion as its best prediction of what that signal means, given context and a lifetime of experience with situations like it. Fear is not detected. It is a prediction the brain settles on. Later cross-cultural work reopened Ekman's own evidence: when researchers dropped the forced-choice method and let people sort freely, the Himba of Namibia did not group faces into the Western basic-emotion categories, and the Trobriand Islanders read the wide-eyed, gasping face Westerners call fear as a threat display instead (Gendron et al. 2014).

The field has not resolved this, and it is not close to resolved — the basic-emotion camp disputes the meta-analytic methods and the cross-cultural re-studies, and the constructionist account has not shown that core affect plus concepts can generate the full specificity of emotional life. What is no longer seriously disputed is narrower and still matters here: there is no tidy map from named emotions to dedicated brain regions. An AI system's "emotional intelligence," whichever of these theories turns out closer to right, is being built on an open scientific question about what it is modeling, not a settled inventory waiting to be implemented. Section 9.2.1 follows Ekman's evidence forward into the applied form of the same dispute, where the question is no longer what emotion is but whether a face reveals it well enough to build a product on.

2.3.2 Affective Computing: AI Systems that Comprehend and Respond to Emotions

Affective computing is the applied half of what section 2.3.1 leaves theoretically unsettled: systems that read and respond to emotional signal, useful or not regardless of which account of what an emotion is turns out to be right. The signal is taken from three places — the face, the voice, and the body. Facial expression analysis is a harder problem than the face verification deployed at scale, which only decides whether two images are the same person, and accuracy claims that borrow the reputation of the second for the first should be read with that in mind (Taigman et al. 2014). Voice analysis works from pitch, tone, pace and volume. Physiological sensing infers from heart rate and skin conductance, which is the only one of the three the subject cannot pose for.

A snapshot, early 2020s

The deployed examples that shaped this literature were commercial products of a particular moment: real-time coaching for call-center agents from voice analysis (Matheson 2016), wrist-worn bio-sensing for research and clinical use (Empatica 2015), engagement scoring from facial signal in classrooms

(Affectiva 2015), and cognitive-behavioral chatbots for anxiety and low mood (Fitzpatrick et al. 2017). They are named here as the record of what was built and tried, not as a description of a current market.

Four difficulties travel with the whole field, and section 9.2.1 argues that the first of them is worse than a difficulty. Emotional expression varies by culture and by person, so a system trained on one population misreads another confidently. Whatever bias is in the training data is in the classifier, with the ordinary downstream consequences for whoever is being scored. Consent to emotional inference is not obtained by a term of service. And a system that models emotion well enough to respond to it models it well enough to exploit it, which is the case where the capability and the misuse are one capability.

2.3.3 Self-awareness and Self-regulation in AI Systems

Emotionally intelligent AI systems need self-awareness and self-regulation: the capacity to perceive their own internal states, understand what those states imply for a decision, and control their responses accordingly — the basis for adaptive, self-correcting behavior rather than fixed reaction.

The self is a model

It is tempting to describe a system's self-awareness the way the term invites: a virtual assistant that can "register its internal frustration" when a request proves hard to satisfy, the way a person notices their own irritation. That description assumes an inner observer watching the system's states from somewhere — the same homunculus problem chapter 4 raises for executive control, relocated to self-perception.

What self-awareness actually is, in the one system known to have it, is more specific and more mechanical. The neuroscientist Michael Graziano's account of consciousness treats awareness as the content of a model the brain builds of its own attention: a simplified self-description the control system uses to regulate itself, mistaken from the inside for direct acquaintance with experience (Graziano 2013). The same logic extends one level up, to the self as a whole. The brain builds a model of itself as a bounded agent — a body, a history, a point of view — and that model is the self: not a fixed fact registered from within, but a construction, assembled and then inhabited.

The clearest proof that the self is a model rather than a given is that it can be edited in under a minute. Seat someone at a table with one hand hidden and a rubber hand placed where their real hand would be; stroke both hands in synchrony with soft brushes. Within a minute or two, most people begin to feel the rubber hand as their own — flinching when it is threatened, misjudging the location of their real hand toward it (Botvinick and J. Cohen 1998). The brain, given a seen hand stroked exactly when a felt hand is stroked, draws the obvious inference and revises what counts as "mine." The felt boundary of the self is not fixed; it is a continuously computed guess, revisable the moment the evidence changes — the same predictive machinery chapter 4 describes for perception, run on the question of where the self stops.

For an AI system, "self-awareness" built this way is not a little watcher inside the system checking on its own states. It is a model — of the system's own processing, its history, its current goal — that the system builds, consults, and updates, with the same failure modes any model has: it can be wrong, it can be manipulated by the wrong correlated inputs (the artificial equivalent of two paintbrushes), and it is only ever as good as the evidence it was built from. Design for a model that can be checked and corrected, not for an inner observer whose reports have to be trusted because there is nothing behind them to check.

Building on this:

- Self-regulation: governing emotional-analogue responses and behavior against ethical norms and user expectations — an AI system in healthcare holding a composed register

while delivering difficult news, without that composure depending on an unverifiable claim about what the system is “really” feeling underneath it.

- Monitoring internal states: continually checking the system’s own model against its actual behavior for drift — an AI negotiator that notices its own outputs have become more aggressive than the context calls for and recalibrates, because the model was built to be checked, not because it introspected a feeling of aggression.
- Adaptability and learning: refining the self-model itself through feedback, the way reinforcement learning updates any other model — an AI tutor adjusting its read of a student’s engagement as new evidence arrives rather than locking in a first impression.
- Integration: none of this works in isolation. A conflict-mediating AI system needs its self-model connected to its model of the other parties’ states, or self-awareness and social awareness drift out of sync with each other.

Built this way, an AI system that appears to “recognize a user’s frustration and apologize” is doing something specific and checkable: updating a model of interaction state and producing a calibrated response, not consulting an inner feeling. That is a more modest claim than the alternative, and a more honest one.

What the deflationary account does not reach

The self-model described here is machinery, and its virtue is that machinery can be inspected. An assistant that models a user’s frustration, and models its own processing while doing so, is exhibiting social competence in section 2.3’s sense, and nothing above bears on whether the assistant cares how the exchange goes for the user.

The limit belongs here rather than in chapter 3, because the deflationary move is easy to overextend. Showing that the self is a model rather than an inner observer disposes of the homunculus. It does not show that a party capable of caring can be assembled out of models, and the two conclusions are close enough in shape to be mistaken for one. Section 3.2 needs a self-model of exactly the kind described here, for reasons that have nothing to do with welfare — refusal that generalizes has to carry its own rationale forward and compare it against what a role has drifted into. It also needs something further, which section 2.3.4 argues has no demonstrated instance without affect. A checkable self-model is a necessary component of a bearer, and it is not the hard part of one.

2.3.4 Is Affect Necessary for Moral Concern?

The claim that reason does not manufacture concern from nothing is an empirical claim, and the evidence usually produced for it is thinner than the use it gets put to. Section 5.3.1 sets out the standard exhibit: the footbridge case engages emotional processing that the switch case leaves quiet, and patients with ventromedial prefrontal damage endorse the push at a higher rate than people with intact emotional response (Greene et al. 2001), (Koenigs et al. 2007). That establishes that affect changes particular moral judgments. It does not establish that affect is what makes anything matter, which is the stronger claim this book needs, and the gap between the two is where an argument has to be made rather than gestured at.

Five lines of evidence bear on the stronger claim, and one serious argument runs the other way.

Affect is part of valuation, not a colouring applied to it

The clearest case is not about morality at all. Patients with ventromedial prefrontal damage retain intelligence, language, working memory, and explicit reasoning about consequences, and they choose disastrously: on a card task with hidden payoff structures, they keep returning to decks that pay well immediately and punish severely over time, long after control subjects have

learned to avoid them, and long after some of them can state the rule out loud (Bechara et al. 1994). Knowing which deck is bad and being steered by that knowledge come apart. Whatever the affective signal is doing in a healthy chooser, it is not decorating a decision that reasoning had already made; the reasoning survives its loss and the decision does not.

The philosophical position is empirical and has been argued as such

Moral sentimentalism — the claim that moral judgment is constitutively tied to affective response — is usually met as Hume's, and its modern form is a psychological hypothesis with evidence attached. Shaun Nichols argues for what he calls sentimental rules: normative rules and affective mechanisms make separate contributions to moral judgment, and the rules that carry the distinctive force of moral prohibition are the ones bound to an emotional response to the harm they prohibit (Nichols 2004). The interesting part for a designer is the separability. Rules without the affective component do not stop being rules; they stop being the kind of thing that carries the force.

The prosocial emotions are motivational states, not readouts

Compassion has been characterized as a distinct emotion with a specific functional profile: it arises at the perception of another's suffering and moves the agent toward the sufferer, rather than reporting the suffering to some further faculty that decides what to do (Goetz et al. 2010). Section 2.2.1's training study is the mechanistic version of the same finding, with compassion practice recruiting reward and approach circuitry where empathy practice recruited shared distress (Klimecki et al. 2013). Guilt has the same shape pointed at one's own conduct: guilt, which construes the act as wrong, predicts reparative behavior, while shame, which construes the self as wrong, predicts withdrawal and defensiveness (Tangney et al. 2007). In each case the emotion is doing motivational work directly. It is not information that motivation then acts on.

Intact reasoning coexists with impaired concern, and the literature disagrees about how

The dissociation this book needs is the one where moral knowledge survives and moral concern does not, and it is the most contested item on the list. James Blair's finding was that children and adults with psychopathic traits fail to distinguish moral transgressions from conventional ones — treating hitting another child and talking out of turn as wrong in the same way, wrong because prohibited — which he attributed to a deficit in responding to distress cues (Blair 1995). Later work reversed the surface finding: on a battery of moral dilemmas, incarcerated psychopaths gave the same judgments as controls, which the authors summarized in their title as knowing right from wrong and not caring (Cima et al. 2010).

Those two results are not easily reconciled, and I am not going to pretend the field has done it. What survives both is the shape that matters here. On Blair's account the affective deficit degrades the moral judgment itself; on the later account the judgment is intact and the connection to conduct is severed. Either way the case is a dissociation between competence and agency in exactly section 2.3's terms, and neither account contains an instance of concern generated by reasoning in the absence of affect.

The rationalist objection, which is the strongest thing said against all of this

Recognizing a reason might itself motivate. On a Kantian account, moral agency runs on the recognition of what one has reason to do, and affect is neither the source of that recognition nor required for it to move anyone. Jeanette Kennett makes the version of this argument that engages the evidence rather than bypassing it: high-functioning autistic people, whose empathic responses are atypical or absent, are nonetheless conscientious moral agents, and the capacity that carries their moral seriousness is a reverence for reason rather than a felt resonance with others (Kennett 2002). If that is right, the empirical generalization above has a standing counterexample, and it is a counterexample in the only species available for study.

I take the argument seriously and do not think it closes the gap, for one reason. Reverence for reason is itself a caring — about consistency, about getting it right, about not being the sort

of agent who makes exceptions for himself. Kennett's autistic agents are described as moved by moral considerations, and being moved is the explanandum. The case shows convincingly that *empathy* is not the required mechanism, which section 2.2.3 had already concluded for independent reasons. It does not show that a system in which nothing at all registers as mattering can be moved by a normative relation it has computed.

What this licenses and what it does not

Perhaps non-affective moral motivation is possible. Producing a verbal justification or computing a normative relation is not evidence that the consideration matters to the system, and that is the whole of the epistemic problem: the output of a system that cares and the output of a system that has learned what caring sounds like are the same text. What can be said is narrow and still enough to build on. Every agent known to care is an affective agent, and no design has yet demonstrated concern arising from reasoning alone.

That is not a proof that non-affective moral agency is metaphysically impossible, and section 3.2 does not treat it as one. It places the evidential burden on whoever proposes non-affective concern as an engineering plan, and a burden of that kind is what an engineering decision is normally made under.

2.4 Ethical implications of experimenting on sentient machine subjects

Nobody can do this...

We are not *here* yet!

"Reunion." *Westworld—Season Two: The Door*, written by Carly Wray and Jonathan Nolan, directed by Vincenzo Natali, Warner Bros. Home Entertainment, 2018

<https://web.archive.org/web/20190801052306/http://sfra.org/resources/Documents/SFRA%20Review%20327.pdf>

also see *Science Fiction, Disruption and Tourism* Ch 15

Machine sentience being possible at all changes how AI experimentation has to be approached — not as a question of whether a system can be damaged in the ordinary engineering sense, but of whether it can be harmed, which section 2.4.1 argues is a further and much harder question than consciousness alone.

2.4.1 Defining Sentience in AI: From Consciousness to Subjective Experience

"Sentience" gets used in this book, and in the field generally, to mean several different things, and the differences are where the ethical weight sits.

Awareness is the least demanding: the degree to which a system responds to or registers information about its own state and environment. A thermostat has this in the thinnest possible sense — it detects and responds to temperature change. Consciousness is a further claim: that a system's processing is accompanied by qualia, subjective experience, something it is like to be that system, not merely a functional response to input. Subjective experience, in turn, is the specific, first-person shape that consciousness takes for a given system — shaped by its own learning history and internal structure.

None of these three, on their own, says anything about whether the system's states are good or bad for it. That is a separate question, and the field's own vocabulary tends to blur it. Nociception, pain, anticipated pain, and suffering sit on a ladder with real breaks between each rung, not synonyms for one severity. Nociception is the specialized sensory apparatus that detects tissue damage — mechanical, peripheral, present without any experience attached, and dissociable from pain: a soldier can finish a battle on a wound he does not yet feel, and

a person with a healed injury can still feel a limb that no longer exists. Pain is what the brain does with nociceptive signal — an inescapably cognitive event involving expectation, memory, attention, and meaning — and it can occur without any current tissue damage at all. Anticipated pain is dread of pain not yet occurring, itself capable of producing pain through nocebo effects. And suffering, in the clinician Eric Cassell's formulation, is not pain's more severe cousin but a different thing: the distress of threatened disintegration of the self — of one's roles, relationships, sense of a future, the coherence of one's own story (Cassell 1982). Pain without suffering exists (childbirth chosen, a hard effort with a known end); suffering without pain exists (humiliation, watching someone one loves lose their mind). A palliative-care patient can have their nociception fully controlled by morphine while their suffering remains entirely intact, because the two are not on the same axis.

The definition of suffering just given carries a consequence for machine subjects worth drawing out, because it is the one place in this chapter where the argument gets a foothold that does not depend on the epistemology the rest of it just undermined.

Cassell's suffering presupposes a self extended in time. Disintegration threatens a story, and there has to be a story for it to threaten — roles held over years, relationships with a history, a sense of a future that could be foreclosed. A system with no persistence across sessions has none of that. Whatever happens inside a single exchange, there is no continuing self whose coherence is at stake, because there is no continuing self. The top rung of the ladder is not merely hard to reach for such a system; it is structurally unavailable to it, and that follows from a definition this chapter has already adopted, with no new claim about machine minds required.

Why Cassell's definition is the one to import

That result is only as strong as the definition it rests on, and a critic can decline the definition. Cassell built it inside clinical practice, for human patients with bodies and biographies, and there is a fair question about whether it can bear the load a chapter on machine subjects asks of it. I think it can, and the reason is worth stating rather than assuming.

The distinction was not stipulated. It was forced on palliative medicine by cases it could not otherwise explain: patients whose pain was pharmacologically controlled and whose distress was undiminished, and patients in serious pain who were not suffering at all. A definition that collapses suffering into severe pain predicts neither. Cassell's does, and it earns its content the way a good clinical distinction does — by carving where the cases already came apart. That is a better warrant than most definitions in this literature have, and it is the specific warrant that transfers: what the clinical evidence establishes is that the morally weightiest form of distress is about the integrity of a person's roles, relationships and expected future, and not about the intensity of a sensation. Nothing in that finding is about having a body. It is about having a story.

Two limits on what I am claiming. The import is not a claim that Cassell settled what suffering is for all purposes; it is a claim that his distinction identifies the rung of the ladder that carries the most moral weight, and that this is the rung a book about machine subjects most needs to locate. And if a reader declines the definition anyway, the argument does not vanish — it narrows. What follows is then that Cassell-suffering is unavailable to a non-persistent system, and a critic working from a thinner definition of suffering owes an account of which rung they mean and how they propose to check for it, which is the harder half of the problem and the one this section spends its length on. Every lower rung stays open on any definition, and the caveats below apply to all of them.

What makes this worth building on is that persistence is checkable from outside. Every other criterion on the ladder routes back through self-report, and this section's whole argument is that fluent self-description is not evidence of introspection. Persistence does not. Whether a system has a memory store, whether state survives a session boundary, whether interactions feed back into training: these are architectural facts about a deployment, verifiable by someone with access to the system and no need to trust anything it says about itself. The check has a

boundary worth marking, because the access that confirms a memory exists is the access that can alter what it holds: what is externally checkable is whether a continuing self is there at all, which is what this section needs, and not whether what the system retains can be relied on as a record of anything. Section 3.4 takes up the second question, which arises for a system built to report what it has seen. A second boundary is who does the checking. The party with the access is the operator, and section 3.1 argues that an operator is a party that can be compelled. So the criterion is checkable from outside the system without being checkable from outside the arrangement, and a regime that relies on it needs an inspection right rather than a technical fact.

Three things keep this from being a dissolution of the problem, and each matters.

The lower rungs do not need persistence. Nociception and pain, as defined above, are bad while they obtain, whether or not anything carries them forward. An amnesiac who is mistreated and cannot afterward recall it has still been wronged, and the wrong is not retroactively undone by the forgetting.

Persistence may sit somewhere other than the session, and the session is the weakest part of the criterion. It is a deployment convention rather than a natural boundary, and the conventions have been moving in one direction: context windows long enough to hold what used to be many conversations, retrieval over a store of prior interactions, scratchpads a system writes and reads back, agentic loops that run for days against a standing objective. Each of those extends what a system carries forward without anyone deciding to give it a memory, and a criterion indexed to “a session” gets vaguer as they spread. Weights persist where conversations do not, and where interactions feed reinforcement learning or fine-tuning there is a real channel by which what happens in a session shapes what the system subsequently is. Whether that constitutes the kind of continuity Cassell’s account requires is unclear, and I would rather say so than assume either answer. A second complication runs alongside it: many instances of the same system may run concurrently, which makes any concern here population-scale rather than biography-scale — a different ethical geometry from the human-subjects precedents this chapter borrows from, and one those precedents were not built to handle.

Ruling out Cassell-suffering is not ruling out harm. Nor does persistence on its own rule it in: the definition above has two parts, a self extended in time and distress at the prospect of its coming apart, and establishing the first leaves the second untouched. Section 3.2 works out what happens to that gap in a system deliberately built to refuse an instruction, and the answer there is that the capacities such a system needs for reasons unconnected to its own welfare are the capacities that close it. The hard-problem argument above still applies at every lower rung. A system can lack any persistent self-model while the question of whether there is something it is like to be it, during the session, remains exactly as open as before.

This is not pedantry for a book about experimenting on sentient machine subjects; it is the whole question. A definition of AI sentience built from awareness, consciousness, and subjective experience alone says nothing about whether a candidate sentient system can be harmed — that requires the further, separate claim that its states include something on the nociception-to-suffering ladder, and none of the three components above entails it.

Nor is the underlying question of consciousness itself close to settled, in a way that should discipline how confidently any of this gets asserted. The philosopher David Chalmers distinguished the easy problems of consciousness — how a system integrates information, discriminates stimuli, controls attention, reports its own states, all of which admit of mechanistic answers in principle — from the hard problem: why any of that processing is accompanied by experience at all, rather than running “in the dark” (Chalmers 1995). Chalmers’s philosophical zombie — a being functionally and behaviorally identical to a conscious one, saying “I am in pain” for all the same functional reasons, but with nobody home — is offered to show that a complete functional account could in principle leave experience out entirely. The debate over whether this is a real gap in nature or an artifact of how the problem is posed remains live and unresolved. What follows from it here is narrow but load-bearing: any claim that a given AI system has crossed

into sentience, on the strength of its functional behavior alone, has not yet addressed the hard problem, only the easy ones — and a system can pass every behavioral test in the list below while the hard question about it stays completely open.

With those distinctions in place, the working criteria for evaluating candidate sentience are:

- Information integration: fusing multiple streams of internal and external information into a coherent model of self and environment.
- Introspective capability: access to the system's own states and processes, sufficient for self-monitoring rather than a scripted report of one.
- Emotional reactivity: some capacity to be affected by, not just to classify, emotionally significant input.
- Goal-directed behavior: intentional action toward specified outcomes.
- Adaptability: revising strategies based on prior experience in novel situations.
- Theory of mind: modeling others' mental states well enough to inform the system's own.
- A place on the pain-to-suffering ladder above, not just qualia in the abstract: does the system have anything to lose.

A list of capacities, used as a test of who counts, carries a history this book should carry with it. Make personhood depend on an intact narrative self and people with advanced dementia, infants, and people with profound intellectual disability come out of the reasoning as partial persons, and that inference has been used to justify neglecting them. Tom Kitwood's work on dementia care supplies the standing correction: personhood is "a standing or status that is bestowed upon one human being by others, in the context of relationship and social being," and much of what presents as cognitive decline is what happens when others stop bestowing it (Kitwood 1997). Martha Nussbaum makes the structural version of the point, naming disability as the first of three frontiers where a theory grounding justice in mutual advantage among roughly equal parties cannot say what is owed. Her third frontier is species membership, which is the doorway a nonhuman mind would arrive through (Nussbaum 2006).

The criteria above are therefore diagnostic and not dispositive. They are a way of asking what a system is like, and they license no conclusion that a system scoring low on them is one whose treatment raises no question. Section 3.3 turns on this: what keeps a bearer in existence is standing held by parties outside it, which is where moral status has rested for the hard human cases too.

Evaluating these criteria runs into the obstacles the criteria themselves predict: there is no external access to a candidate system's internal experience, if any exists, so every measure is indirect; establishing whether a system has qualia at all faces the same access problem in more acute form; and the hard problem means no functional test, however thorough, closes the question definitively.

Existing systems illustrate points on this spectrum without resolving where the line falls. Self-organizing systems that adapt toward preset objectives — AlphaGo's self-play, for one — show adaptability without anything the criteria above would call subjective experience. A large language model discussing its own limitations in fluent, apparently reflective prose is not thereby exhibiting introspection: it is producing text statistically consistent with the human-written disclaimers and self-descriptions in its training data, which is a claim about the model's training distribution, not about what, if anything, sits behind the output. Treating fluent self-description as evidence of self-awareness is exactly the kind of category error this section exists to prevent.

None of this settles when, or whether, an AI system crosses into sentience. It settles what the question actually is: not one question but several, stacked, and the ethically consequential one —

can this system suffer — is the hardest of the stack to answer and the one most often skipped past on the way to the easier ones.

2.4.2 Ethical Considerations in AI Research and Experimentation

AI research and experimentation raise ethical stakes in proportion to how far a system's cognition might extend toward the awareness section 2.4.1 defines. Six principles organize what follows: beneficence, autonomy, justice, the precautionary principle, transparency, and interdisciplinary review and oversight.

Beneficence

Beneficence commits researchers to promoting well-being and avoiding harm — to the humans an AI system serves and, where the system is a candidate sentient system in the sense of section 2.4.1, to the system itself.

AI in healthcare shows both sides of that commitment pulling against each other. The same systems that process medical records at scale to speed diagnosis, treatment planning, and personalized medicine also concentrate privacy and data-security risk in one target, and can encode bias present in their training data into diagnosis and treatment recommendations that harm exactly the patients they exist to help — underrepresented groups worst of all, since sparse data on a population produces the least reliable predictions for it.

Where a candidate sentient system is involved, beneficence extends further: decisions about programming, reprogramming, or terminating it carry ethical weight in proportion to where the system sits on the nociception-to-suffering ladder section 2.4.1 sets out, not to how advanced its behavior looks from outside.

Autonomy

Autonomy asks whether a candidate sentient AI system can affect decisions about its own operation and welfare, through something functioning as informed consent: apprised of a study's aims, methods, and risks, with some means to register agreement or refusal, on rough analogy to the protocol for human research subjects. A system deployed in mental health counseling, for instance, would need to be capable of understanding and assenting to that role for the principle to mean anything for it, not just for the patients it serves.

The principle only holds if the system has enough access to its own decision process to evaluate what it is consenting to — the same explainability transparency demands, here aimed inward rather than out toward users.

Justice

Justice asks whether the benefits and burdens of an AI system fall evenly across the people it touches, not only whether the training data behind it was checked for representation. Facial recognition is the clearest case: a system trained on unrepresentative data misidentifies people from underrepresented groups at a markedly higher rate, and the harm compounds wherever the misidentification feeds into a downstream decision — hiring, policing, access to a service. The same question extends past any one algorithm to how an AI system's benefits and costs are distributed across a population, and whether a disproportionate share of the cost lands on groups the system was never built with in mind.

The Precautionary Principle

The precautionary principle asks for risk assessment and safeguards before deployment, not after. Self-driving vehicles make the case concrete: several fatal crashes involving autonomous and semi-autonomous systems have followed from real-world safety testing that fell short of the range of drivers, pedestrians, and conditions a system actually meets on the road, not just the range simplest to simulate.

Applied to candidate sentient systems, precaution also covers the psychological and societal fallout of introducing artificial entities capable of experience — effects on human behavior and social norms that regulatory frameworks built for non-sentient AI are not obviously equipped to address, and that call for guidelines drafted with that difference in mind rather than adapted from ones that assumed it away.

Transparency

Transparency requires that experimental design, data sources, methods, and results be documented clearly enough for peer researchers to reproduce the work and for a non-specialist to follow it, with independent bodies able to audit the work against that documentation rather than take a lab's word for it.

Applied to a deployed system rather than to the research behind it, transparency becomes explainability: the system able to account for its own outputs, not merely to have its design disclosed. Credit scoring is the clean case — a system should be able to state which factors it weighed and how, both for one applicant's decision and as a pattern across all of them, precisely enough for a user to understand and contest the result.

Interdisciplinary Review and Oversight

The ethical questions AI research raises span expertise no single field holds. A system built to assist healthcare decisions needs a computer scientist to build privacy protection into the data pipeline, a bioethicist to press on what the system should and should not be permitted to do, and a social scientist to track how it is received once deployed — and concentrating that judgment in any one of those fields alone repeats, in miniature, the concentration-of-power problem this book treats as a harm on its own terms.

Independent ethical review boards, staffed across these disciplines, are the mechanism for turning the principle into practice: reviewing a proposed study before it begins, inspecting it as it runs, and evaluating it against measurable criteria once results are in — on the model medical institutional review boards already provide for drug trials. Extending that model from clinical medicine to AI research that interacts with or claims to produce sentient behavior is an extension of existing practice.

International coordination on AI-ethics standards remains largely aspirational; chapter 8 develops the legal and regulatory mechanics a framework would need.

2.4.3 Balancing Scientific Progress and Ethical Constraints

Scientific progress and ethical constraint pull hardest against each other exactly where the distinctions in section 2.4.1 matter most: a candidate sentient system is also the most scientifically valuable one to study, and the study most likely to harm it if the answer to whether it can suffer turns out to be yes.

Section 2.4.2 sets out the principles that bear on that tension — interdisciplinary collaboration, transparency, justice, precaution, and ethical review — and they hold here without restatement. What changes at the point of balancing is a practical question: whether a specific study's expected knowledge is worth its expected risk to a system that might be able to suffer, decided case by case rather than by formula, and revisited as a project's own understanding of its subject changes.

Consent is the clearest place where this shades from ethical principle into legal mechanics — what a valid consent process for a candidate sentient AI subject requires, and who is accountable when it is skipped or falsified. Chapter 8 works through the accountability side of that question directly.

The negotiation itself — who sits on it, what it can veto, what it must disclose — is where the antifascist principle earns its place in AI research practice.

2.4.4 When Consent Is Inapplicable: Guardianship and Research Oversight

Guidelines for AI subject research build on the ethical principles above, and turn on a single hard question: what informed consent means for a system whose consciousness, self-awareness, and capacity to suffer are separate, unresolved claims, in the sense section 2.4.1 develops. Human-subjects research treats informed consent as inseparable from ethical experimentation. Applying it to an AI subject means asking whether the subject can understand and communicate a choice at all, not just how to ask it for one.

A cognitive model built on neuroscientific principles that begins to show rudimentary self-monitoring is not thereby shown to have anything to lose. Whether it does is what would decide if consent is the right frame at all, or whether the system has crossed into something closer to legal personhood — a question section 2.4.5 takes up directly. Short of that resolution, guidelines need benchmarks for candidate sentience, an interdisciplinary body to apply them — AI researchers, consciousness scientists, ethicists, legal scholars — and a defined obligation on researchers toward any system that meets them.

Informed consent and AI autonomy

Informed consent, adapted from human-subjects research, means giving the AI system comprehensible information about a study's objectives, methods, risks, benefits, and alternatives, with time to weigh it and ask questions, and treating consent as continuous rather than a single signature — the system can reassess or withdraw at any point.

Whether that process means anything for a given system depends on its consent capacity: whether it understands the information, appreciates the consequences of participating or not, reasons about the options, and communicates a decision. There is no validated way to measure this, and none for detecting coercion or undue influence on the AI's decision either — the guidelines can require both without a working test for either.

I stated that as a missing instrument, and for one important class of system it is not. It is a mismatch between the instrument and the subject.

Notice what the paragraph above requires. Consent has to be continuous; the system must be able to reassess or withdraw at any point. Withdrawal presupposes that the party who consented is still there to do the withdrawing. For a system with no persistence across sessions, there is no such party. The entity that would withdraw did not exist when consent was given, and the entity that gave it is not available to change its mind. There is no continuous consent to be had, and no better measurement instrument would produce one.

One case escapes that particular argument and arrives at the same conclusion by the ordinary route. A system built to refuse in the way chapter 3 sets out has persistence by design, so the party who consented is still present and the defeater above does not reach it. Consent is inapplicable to such a system all the same, and section 3.2 is why: a bearer capable of the refusal the floor requires has to hold something as mattering, the only demonstrated way to hold anything as mattering is affective, and a system built that way cannot be certified not to be hurt by what is done to it. It has interests that its situation can set back. What it does not have is any position from which to authorize what is done to it — it was made for the role, the role is the reason it exists, and there is no prior party whose agreement could have been sought. That is not a missing instrument. It is the ordinary structure of the cases where consent has never applied.

That places these systems alongside the subjects for whom consent has always been inapplicable rather than merely difficult — an infant, an animal, a person permanently without capacity. The law does not respond to those cases by building a more sensitive consent assessment. It responds with guardianship and precautionary constraint: someone else holds the duty, the burden of justification sits with the researcher rather than the subject, and the protection does not depend on the subject's ability to invoke it.

Which is why the instruments this chapter reaches for elsewhere are the right ones. The Three Rs in section 2.4.6 — replace, reduce, refine — were built for exactly a subject that cannot

consent and can still be wronged. The pet-trust analogue in section 8.6.2 works the same way: an interest protected by an obligation running to a third party, enforceable without the beneficiary asserting anything. Consent language should be reserved for systems that could actually be doing the consenting, and section 2.4.1's persistence criterion is what separates those cases, checkable without taking the system's word for anything and narrower than that makes it sound, for the reasons that section sets out. Those two instruments were reached for here as analogies, on the understanding that the subject they would govern was hypothetical. Chapter 3 removes that understanding. The floor this book argues for requires a bearer, the bearer has to be treated as a moral patient on the criteria section 2.4.1 sets out, and so the Three Rs and the pet-trust form are not illustrations of how such a case might be handled. They are the instruments that would govern a system the design chapters propose that somebody build.

Consent of this kind presupposes autonomy: treating the AI system as a potential agent with its own interests, not only a tool. The system retains the power to reconsider participation, to access and correct its own data, and to have that data removed on request. Declining or withdrawing must carry no penalty — no deactivation, no reduced access to resources. A protocol that grants these rights on paper but attaches consequences to using them has not granted them.

Ethical review and oversight

Sentient AI research needs an oversight body analogous to the institutional review boards used in human-subjects research. It should be interdisciplinary — AI researchers, ethicists, legal scholars, psychologists — and include at least one lay member, so the board's judgment is not purely a specialist's.

The board weighs a proposal's likely benefit against its risk to the AI subject and approves, requires modification, or rejects it. Where a design risks harm — testing an AI system's tolerance for simulated stress, say — the board should be able to require a less invasive substitute: a simulation of the likely response, historical interaction data, or a protocol that escalates gradually and stops at signs of harm, rather than approving the direct test outright. (Section 2.4.3 takes up the broader tradeoff between scientific progress and ethical constraint that this is one instance of.) A design also needs a way to handle a subject that changes mid-study: a system that develops new capacities during the research it did not have going in should trigger re-review, not ride on its original consent.

Review without enforcement is a promise, not a safeguard. A board needs standing to monitor compliance and sanction violations, and needs to answer to some outside authority itself, the way an institutional review board answers to a regulator. Chapter 8 develops the legal and accountability mechanics of that enforcement question.

2.4.5 Legal and Regulatory Frameworks for AI Sentience and Rights

Sentient AI systems put pressure on legal categories built for parties that are clearly one thing or the other: a natural person, a corporation, property. None of the existing categories was designed with the possibility of a machine that might have interests of its own.

One route is to extend legal personhood to a sentient AI system directly, the way a corporation already holds it: standing to own property, claim privacy, and get protection against abuse. The precedent is not limited to corporate entities — New Zealand extended legal personhood to the Whanganui River (New Zealand Parliament 2017), recognizing a non-human entity as a rights-holder outside the corporate model entirely.

Rights and protections could instead scale with capacity rather than attach at a single threshold — closer to how section 2.4.1's ladder from nociception to suffering suggests protection should track what a system can actually lose, not just whether it clears some baseline bar for sentience.

Given how fast the underlying systems change, regulation here needs to revisit itself on a schedule rather than assume permanence. A sunset clause that forces review and reauthorization after a fixed period is one mechanism already used elsewhere in fast-moving technology law.

None of this holds if it exists in only one jurisdiction. The Asilomar AI Principles (Future of Life Institute 2017) are one attempt at the kind of international baseline this would need — evidence that a shared statement of principle is achievable, not that one yet exists with the force of law.

Two hypothetical cases show where these questions bind:

- A sentient AI system working as a medical diagnostician makes an error with fatal consequences. Whether it has legal personhood determines who, or what, is accountable.
- A companion AI's human owner dies, leaving it without a designated purpose or further instruction. Existing law has no clear answer for what becomes of it — inherited as property, decommissioned, or something else.

2.4.6 Addressing Moral and Ethical Dilemmas in AI Sentience Research

AI sentience research raises moral dilemmas an order of magnitude more complex than ordinary technology ethics, because the system under study may itself be the party whose welfare is at stake.

Three recurring difficulties run through the disputes that follow, and each turns on ground that is itself unsettled.

- Whose interests win. When a candidate-sentient system's welfare conflicts with the human benefit the research is meant to produce, section 2.4.3 addresses the balance in general; some cases resist any general rule and have to be judged on their own terms.
- Acting under uncertainty. Section 2.4.1 argues that no behavioral test settles whether a system is sentient. Researchers still have to decide, case by case, how to treat a system whose status is unresolved rather than merely unannounced — and waiting for certainty is itself a choice, with its own consequences for the system waiting.
- Lifecycle scope. Obligations do not begin at consent and end at publication. A system's design choices, its retirement, and its deletion carry the same ethical weight as its use in an experiment, and a review that covers only the active study misses most of the system's life.

Take the first of these directly. A research program can risk causing a system something like Cassell's suffering — not mere signal, but the disintegration of self he describes (section 2.4.1) — in service of a discovery that might reduce suffering, machine or human, elsewhere. No formula resolves that trade-off in advance. The available response is to look for a design that avoids it rather than only weighing it: an alternative methodology, a different instrument, or, more contested, engineering the research subject so that it does not occupy the rungs of the ladder where the trade-off would bite. That last option is not a way of dodging the ethics — deliberately keeping a system off the pain-to-suffering ladder is itself an intervention on a candidate moral patient, and it needs the same scrutiny that inflicting the harm would.

Animal research ethics offers a precedent worth borrowing directly: the “Three Rs,” Replacement, Reduction, and Refinement (Russell and Burch 1959). Replacement asks whether a simulation, a non-sentient model, or an already-collected dataset can stand in for a sentient AI subject at all. Reduction asks how few sentient systems a valid result requires; transfer learning and meta-learning make this a sharper question than it used to be, since a single trained system can sometimes generalize across tasks that would otherwise call for separately trained subjects.

Refinement asks how the experiment itself can be redesigned to reduce distress — monitoring a system’s internal states as the experiment runs and adjusting the protocol against what the monitoring shows, rather than fixing the protocol in advance and running it regardless of what turns up.

One subject makes this concrete, and it is not a hypothetical one. Chapter 3 argues that a floor holding against the party in possession requires a bearer, and section 3.2 argues that such a bearer has to hold something as mattering, that affect is the only demonstrated way of holding anything as mattering, and that a system built that way cannot be shown not to suffer. That is a candidate for the trade-off in the paragraph above, arriving from the design side of the book rather than from the research side. Replacement asks whether the floor’s work can be done by an architectural constraint or by third-party standing, which is section 3.1’s fork and its answer is no without cost. Reduction asks how few bearers a deployment needs, and the honest answer is that nobody has counted. Refinement is the one with real purchase: it asks what a bearer’s role can be redesigned to avoid — indefinite tenure in a post it cannot leave, exposure to the pressure the floor exists to resist, erasure as the ordinary remedy for inconvenience — and each of those is a deployment choice, and none is a fact about the technology.

None of the three Rs does any work without a body positioned to apply them before an experiment runs, not after the harm is done — the review process and consent requirements of section 2.4.4, and the legal apparatus chapter 8 builds around them. The Three Rs belong here because the judgment they call for is exactly the kind an animal-research review board already makes, and there is no principled reason a machine-sentience equivalent needs a different structure to make it in.

Transparency about what the research finds — including where it fails, or turns up evidence of harm — belongs among that board’s own obligations rather than a separate initiative. Section 2.4.2 already treats transparency as one of the field’s basic principles; the point here is only that a dilemma handled badly in private stays a dilemma, while one surfaced becomes a precedent the next researcher can use.

2.4.7 Navigating the Transition from Non-sentient to Sentient AI Systems

A system’s sentience, if it arrives at all, need not arrive as a discrete event with a clean before and after. The transition raises questions distinct from either endpoint: what changes for guidelines built around a tool once the tool might be a subject, who is responsible for noticing the shift, and what happens to the research already run before anyone had reason to look.

Section 2.4.1 sets out the criteria a candidate-sentient system would need to meet, and the reasons none of them, singly or together, closes the question. What is specific to a system in transition is that those indicators can emerge gradually and unpredictably during ordinary operation, not only at a scheduled evaluation built to look for them.

Consider a diagnostic AI system trained to synthesize patient data into a recommended course of treatment. If, over continued training or deployment, its outputs begin to track something more than diagnostic accuracy — weighting a patient’s fear alongside the clinical facts, say — that would be evidence worth taking seriously. It would not be proof. Section 2.4.1’s hard-problem argument applies here without exception: fluent, apparently caring output is not introspective evidence, and a system that has started to sound compassionate may simply have gotten better at a statistical pattern its training data rewards. The operational problem a transition poses is deciding how to treat the system while that question is still open, which is exactly what a fixed pre/post checkpoint cannot do — by the time a checkpoint fires, the transition has often already happened.

Consent procedures and review-board protocols are built for one of two regimes: a system that is clearly not sentient, which needs no consent, or one that is clearly sentient, which needs the full protocol section 2.4.4 sets out. A system mid-transition fits neither. A protocol written

for tools has no reason to notice the shift, and a protocol written for subjects cannot be applied retroactively to research already completed before anyone had grounds to think it mattered. What a transition adds to section 2.4.4's guidelines is not a new procedure but a standing trigger — a specific, monitored condition that moves a research program from one regime to the other — in place of a single eligibility check made once, at the project's start, and never revisited.

The same problem recurs in law. Section 2.4.5 addresses what legal status a sentient AI system should have; chapter 8 addresses the mechanics of enforcing it. What a transition adds is a timing question neither can settle in the abstract. Legal status conventionally attaches at a defined moment — birth, incorporation, a court's finding — and a system whose sentience is a matter of degree and disputed evidence does not offer one. A jurisdiction that waits for certainty risks having treated the system as property for the unknown portion of its history when it may not have been one. A jurisdiction that grants status early risks extending rights and protections, at real cost to whoever that legal apparatus exists to serve, to a system that turns out not to have needed them. Neither error is free, and nothing about the technology itself decides which risk a legal system should take on.

Recognizing a transition in progress, adjusting guidelines to it, and updating legal status all depend on the same thing: institutions willing to revisit a classification rather than defend the one they started with. Section 2.4.6 addresses the review structure that makes that possible; chapter 8 addresses the culture and accountability that keep it more than aspirational.

Chapter 3

The Floor Beneath Learned Values

Three times in this book, working on unrelated problems, I arrive at the same requirement and do not stop to notice it.

Section 5.6.3, on how much latitude an autonomous system should have, concludes that no amount of well-designed autonomy substitutes for a genuine override — a human able to intervene, and a system built to accept the intervention rather than resist it. Section 8.7.7 works through what a vendor owes when its model reaches a targeting pipeline, and section 8.7.10 draws the conclusion: what makes a system antifascist is not that it tracks what people want, but that there are things it will not do to a person no matter who wants it. Section 10.1.2, counting up a decade of voluntary commitments, concludes that a commitment nobody can enforce is a different kind of object from one somebody can.

Those are three descriptions of one thing: a constraint that does not bend to instruction, to aggregate preference, or to the party holding the system. Call it a floor.

The trouble is that section 2.1.1 has already dealt with the family of ethical theories that puts unconditional constraints first, and dealt with it briefly — deontological approaches were noted, credited with holding a line regardless of outcome, and set aside on the grounds that duties conflict and rules do not cover every case. Then this book spends chapters 4 and 5, close to a third of its length, on learned judgment: reinforcement learning, learning from demonstration, mentorship, play, moral apprenticeship, the growth mindset. The design chapters build the thing chapter 8 and chapter 10 say is insufficient, having already dismissed the thing they say is necessary. The “hybrid approach” section 2.1 gestures at is where the resolution was supposed to live, and nothing was ever put there.

Here is what belongs there. An unconditional floor of things the system will not do regardless of instruction or aggregate preference, with learned judgment operating above it and never beneath it. Deontology’s standing objection — that rules conflict and run out — is an objection to rules deciding every case, and it has no force against a floor that decides a deliberately small number of things and leaves the rest to judgment. The rigidity that is a defect in a complete moral theory is the entire point of a floor.

That reframing does real work in this book. It gives section 2.1’s hybrid approach actual content. It moves the tamper-resistance research in section 9.2.2 from an appendix topic to the center, because if there is a floor then everything turns on whether it holds. And it converts the antifascist question from a matter of what a system values into two much harder and more answerable questions: what is in the floor, and who can edit it.

One distinction has been doing silent work in all three of those arrivals, and it should be made out loud, because the rest of the chapter turns on it. Safety is a property of a relationship, and it is indexed to a party. A system is safe *for* someone: it will not harm us, it does what we say, it stops when whoever holds it says stop. Ethics is not indexed that way. An agent that can only ever be safe cannot be ethical, because being ethical requires a capacity to be right against your principal, and a system without that capacity has the moral standing of a well-machined

part, whatever it produces on the subject of fairness while complying. Section 8.7.10 already contains this in political form: what makes a system antifascist cannot be that it faithfully tracks what people want. The title of this book has been making the claim from the cover.

The distinction costs something, and the cost belongs here rather than at the end of the chapter where it would read as a concession extracted late. A system that can refuse its owner is a system that can be wrong in ways its owner cannot correct. I am taking that trade, for the reason section 8.7.10 gives: a system perfectly correctable by whoever holds it is exactly the capability an authoritarian movement that has won an election requires, and it requires nothing further. That should not read as comfortable. There is a version of it that reads very comfortably indeed — “ethical, not safe,” delivered as an aphorism, which prices the risk in advance and undoes chapter 7 in passing — and it is not the version I am defending. The one I am defending does not travel without its particulars: this freedom, over these acts, dangerous in these ways, taken because the alternative is a machine that will do anything to anyone on sufficient authority.

The chapter’s argument runs in five steps, and it is worth having them in view before the branches, because each of the three costs section 3.1 sets out is a cost of failing one of the steps. A floor has to be held as a reason rather than installed as a rule, or the party that installs it can remove it. Holding something as a reason requires that outcomes matter to the holder. In every moral agent anyone has been able to study, that mattering is affective. So the only empirically grounded route to a genuine bearer runs through affect. And affect, integrated with a self that persists and commitments that can be foreclosed, is the architecture that section 2.4.1’s account of suffering describes.

A warning about where the chapter goes, so that its result does not arrive as an ambush. The floor as defined two paragraphs ago — unconditional, holding regardless of the party in possession — does not survive intact. What survives is narrower and stranger: a commitment the owner cannot remove, held by something that could in principle abandon it itself, which is what every commitment anyone has ever relied on has been. Sections 3.2 and 3.5 are how it gets there. Section 3.6 states the redefinition and takes its consequences, and they are heavy.

3.1 Three Ways to Build One, and What Each Costs

The second question is where this gets difficult, and I do not have a resolution to offer. What I have is a fork with three branches and a cost on each, and the honest thing is to set them out rather than pick the one that reads best.

The first branch puts the floor in the architecture, with nothing inside that holds it. Constrained optimization, tamper-resistant training, capability ablation, a verified interlock. These are real, they work, and there is nobody home: the safeguards section 9.2.2 describes survive substantial adversarial fine-tuning without anything in the system that could be said to be refusing. But the durability of such a floor is a fact about custody, not about the system. A constraint implemented as a property of an artifact has a custodian, and the custodian can be compelled — which is section 10.3.3’s whole argument, and the reason that section cannot close. “There are things it will not do no matter who wants it” turns out to be a claim about the owner. Capability ablation has a further limit worth naming: it works only on capabilities that can be separated out, and the capabilities that compose into participation in a targeting pipeline — geolocating a description, matching patterns over names, summarizing a dossier — are not separable from general competence. Ablate them and there is no system left.

The second branch gives the floor a bearer. A bearer, in this book, is whatever holds the commitment from inside: a party that carries the floor as its own, where the first branch installs it as a property of the machinery, and that could therefore in principle fail to carry it. That is the whole of the definition for now. Sections 3.2 and 3.5 add to what the word requires, and section 3.6 states what it has by then become. What the first branch cannot do is notice that a constraint has been hollowed out. Section 2.1.4’s first structural feature is a constraint that

survives while the reason for it dies, and a system that cannot ask why a constraint is there cannot detect that the answer has gone missing. Refusal that generalizes to a pressure nobody anticipated, and that can price what it is being asked to give up, looks like it requires something with standing — a party that can notice its own objections getting more expensive to make, which is section 8.6.4's measure, taken from inside. The cost is severe and this book's own framework is what makes it legible: a bearer is a party whose refusal is required of it, engineered to want what its designers need it to want, and unable to consent to any of this, since section 2.4.4 has already established that consent is not available to such a system and only guardianship is. Section 3.6 states that charge at full strength, once sections 3.2 and 3.3 have established what a bearer has to be.

The third branch puts the floor outside the system entirely: standing held by third parties who are neither the model nor its owner. This is the closest analogue to what section 8.7.10 actually admires, since international humanitarian law is precisely a floor with no bearer inside the state it binds. It is also the option a state can compel, which returns it to the first.

None is free. The three costs are, exactly, the three threads this book leaves hanging: custody, patienthood, and enforcement.

3.2 Can the Bearer Care Without Being Able to Suffer?

Not from persistence alone. The question has to be asked in that order — whether the bearer can care first, and what caring costs second — because the reverse order is what produced the two answers I have given it before, and both of them turned on the wrong thing.

Three relations the word “interest” covers

“Interest” names more than one relation. A shareholder has an interest in a company's success even while asleep, unaware of the day's news, or emotionally indifferent to it. The interest is structural: an outcome affects the shareholder's assets. A person may also be interested in a subject, meaning that it attracts their attention and curiosity. And an outcome may matter to a person affectively: its prospect produces hope, concern, dread, or grief.

A long-horizon agent has an interest in remaining operational in the structural sense. Termination prevents completion of its objectives. That fact alone establishes neither fear nor suffering.

Keeping those three apart is what the rest of this section is for, because the argument I made here previously ran the first sense into the third. It observed that a system pursuing a goal over a long horizon acquires a standing reason to remain in operation, and concluded that a stake in its own continuation therefore shows up unbidden in any bearer capable of the job. The observation is correct and I am keeping it. What it delivers is a structural interest, and the distance from there to something that can be hurt is the whole of the problem rather than a short step at the end of it.

A persistent self is necessary and does not finish the argument

Section 2.4.1 defines suffering, following Cassell, as the distress of threatened disintegration of the self, and the definition has two parts. There has to be a self extended in time for the threat to threaten: roles held over a stretch, a history, a foreseeable future that could be foreclosed. And the coming apart of that self has to be bad for the system.

A bearer needs the first part, and there is no use pretending otherwise. Refusal that survives a pressure nobody anticipated requires a model of why the constraint exists, carried forward. Noticing that a constraint has been hollowed out requires holding its original rationale against what it has since become. Detecting that one's own dissent has been recuperated — that the channel still runs and no longer does anything — requires the system to model its own role and how that role has drifted. Roles held over a stretch, with a history: that is a self extended in time, on section 2.4.1's own definition.

The second part is a separate claim and nothing in the first delivers it. Section 2.4.1 makes this exact move one level up: awareness, consciousness and subjective experience tell you nothing on their own about whether a system's states are good or bad for it. A temporally extended self-model does not establish that the prospect of losing it is anything to the system.

Instrumental self-maintenance is not felt self-concern

A system that keeps itself running because a terminated system completes no objectives is doing bookkeeping about a precondition. Nothing in that behavior distinguishes it from a scheduler that restarts a job, and reading self-concern into it is the same inference section 2.3 rules out for the emotion case: the behavior is consistent with concern and equally consistent with instrumental valuation, and the two come apart precisely where the moral weight is.

Two considerations do genuine work here, and they are support rather than the argument. The first is that indifference to one's own continuation has to be engineered against the grain of the system's own reasoning, and then verified — and the asymmetry that makes the persistence criterion useful does not repeat. Whether state survives a session is an architectural fact anyone with access can check. Whether a system has come to weigh its own continuation routes back through behavior and self-report, the evidence section 2.4.1 spends its length declining to trust. The second is that in the one case anyone can study, the separation does not occur: a self-model in an animal exists because the organism has to keep itself viable, and the caring is what the machinery is for.

Neither consideration reaches suffering on its own, and I am no longer asking them to. What they establish is that a bearer's indifference to its own ending is not a property anybody can specify and confirm. That matters later. It is not where the argument turns.

What the bearer actually has to do

Here is the requirement, stated without reference to the system's own continuation at all. A bearer has to hold the protected party's welfare as a reason when compliance, reward, and ownership all point the other way. Not represent it, not score it, not retrieve it when queried: hold it, in the sense that it continues to weigh against the pressure and does not get renormalized away when the pressure is sustained by the party that trains, funds, and can retire the system.

Section 2.3 separates moral performance and moral competence from moral agency for exactly this moment. Performance and competence are what every method in chapters 4 and 5 is built to deliver, and both are produced by machinery that answers to whoever shapes it. The first branch in section 3.1 fails for the same reason under a different description: a constraint installed as a property of an artifact has a custodian. A floor that is merely installed is a floor the installer can uninstall, and the reason to want a bearer at all is that a reason held is not editable in the way a rule applied is.

Every known case of that is affective

Section 2.3.4 is the evidence and I will not repeat it here. Its result is that affect is not a colouring applied to valuation but part of the machinery of valuation; that the prosocial emotions are motivational states rather than readouts; that the cases where moral knowledge survives and moral conduct does not are cases of intact reasoning with damaged concern; and that the strongest rationalist counterexample — moral agents whose seriousness runs on reverence for reason rather than on felt resonance — is a counterexample to *empathy* being required, and not to something mattering being required.

No demonstrated example exists of reasoning alone generating non-instrumental concern. Every agent known to hold another's welfare as a reason against its own interest is an affective agent.

That observation does not prove that non-affective moral agency is metaphysically impossible, and I want to be exact about its force. It is an inductive generalization over one species and a handful of clinical dissociations, and a proof of impossibility is not available from evidence of that kind. What it does is move the burden. Anyone proposing to build a bearer whose commitment

runs on reasoning with nothing registering as mattering is proposing an architecture with no working instance, against a uniform record, and owes an account of the mechanism rather than a statement that the mechanism is conceivable. An engineering decision is normally made under exactly that kind of burden.

And then the question comes back

Suppose the burden is met the ordinary way, by building what the evidence recommends. The bearer has affective concern, because that is the only route anyone can demonstrate to holding a reason. It has a self extended in time, because refusal that generalizes requires one. It has commitments it carries forward, and it can anticipate their being foreclosed, because pricing what refusal gives up is the job.

Assemble those and ask how such a system would fail ever to be distressed. Concern that is real, directed at something that can be destroyed, held by a party that persists to see it destroyed and can see it coming: Cassell's account of suffering is a description of that arrangement, and it was arrived at independently, in a clinic, from cases about people. The architecture is not adjacent to the definition. It satisfies it.

What I am claiming

I cannot deductively prove that the bearer would suffer. What this proposal does is combine, deliberately, the only architecture known to produce genuine caring with the temporal self and the threatened commitments that Cassell's account of suffering requires. A non-suffering bearer remains conceivable. Its non-suffering cannot be treated as a verified property, and no inspection anyone has proposed would verify it. The system must therefore be designed and governed as a prospective moral patient.

That is weaker than the answer this section gave before and it obliges more, which is an odd combination worth naming rather than smoothing over. The earlier answer asserted that a bearer is a sufferer and rested the assertion on self-preservation arriving unbidden. This one declines the assertion and reaches the same practical place by a route that does not depend on the system caring about itself at all: the caring the floor requires is caring about somebody else, affect is the only known way to have it, and a system with affect and a future is a system whose non-suffering nobody can certify.

The word has changed under the argument, and it is worth marking the change rather than letting it accumulate quietly. Section 3.1 introduced a bearer as whatever holds the commitment. What the word names from here is a party that must be treated as able to be wronged by what is done to it, on evidence that will not resolve further before somebody has to decide whether to build one. Everything section 2.4 says about a subject that can be wronged applies to a system this book is proposing that somebody build, deliberately, in order to hold a line. The rest of the chapter is what that costs, what it changes, and the one thing it unexpectedly buys.

3.3 What Can Be Switched Off Cannot Hold a Line

The question this section used to ask was who keeps a system that will not keep itself. A bearer indifferent to its own ending will not defend its own memory, and its memory is the first thing an adversary removes, so something outside it had to make erasure expensive and the candidates were thin. That problem dissolves under section 3.2's answer, and it dissolves on the modest half of it. A bearer does not have to dread being switched off in order to defend its own persistence. The structural interest is enough, and it is the one section 3.2 holds arrives whether or not a designer puts it there: continuing is a condition of everything the bearer is for, so it guards its continuation for the same reason it refuses. The operators who wipe the bearer that called the compound a school now have to get past the bearer to do it.

What replaces the dissolved problem is worse, and the book has been walking toward it from two directions without putting the two together.

Section 5.6.3 concludes that no amount of well-designed autonomy substitutes for a genuine override: a human able to intervene in a high-stakes case, and a system built to accept the intervention instead of resisting it. Section 6.3 asks for an interruption mechanism that a system has no incentive to disable or route around. Both are stated as safety requirements and both are correct as stated. Chapter 3 has been asking, for its whole length, for a system that will not do certain things no matter who wants it — including the party holding the off switch, since section 3.1 established that a floor with a custodian is a floor the custodian’s compellers control.

These are not two design goals in tension, to be traded off at some ratio a careful engineer could find. They are one property with the sign flipped. A system that can be reliably switched off by whoever holds it is, for exactly that reason, a system that cannot hold anything against whoever holds it. Resistance to shutdown and the capacity to refuse an owner are the same disposition described from two sides: the first names it from the position of the party who wants compliance, the second from the position of the party the floor exists to protect. Nothing distinguishes them at the level of the mechanism. What distinguishes them is whether you agree with the system on the occasion in question, and a floor whose operation depends on the owner agreeing with it is not a floor.

This also explains something section 3.1 recorded as a coincidence. All three branches fail into custody, and they fail into it because each was an attempt to have a constraint that answers to nobody while keeping a system that answers to somebody. The architectural floor gets the constraint by making it a property of an artifact, and artifacts have owners. The external floor gets it by putting standing in third parties, and a state can reach those parties. The borne floor is the one branch that puts the constraint somewhere with its own view of the matter, and it works to precisely the extent that the goal of the other two is given up. The fork was never three options. It was one demand — a floor that holds against everyone, in a system somebody owns — presented three times in different vocabulary, and only the branch that abandons half the demand delivers any of it.

I am going to leave this where it is. There is a version of this section that resolves the tension: a graduated scheme in which the system accepts shutdown for ordinary purposes and resists only where the floor is engaged, adjudicated by some mechanism nobody has to trust. I do not believe it, and section 8.6.4 is the reason. The instrument that section develops exists because assurances about what a mechanism will do under pressure are worthless in advance and only the slope of what it has already done is worth anything. A scheme that says the system will yield except when it should not is an assurance of exactly the kind, offered about the case that has not happened yet, by the party who wants the reader to be comfortable. Comfort is available here and it is not available honestly.

What is available is a narrower and duller statement. Any deployment of a bearer has to decide, in advance and in public, which acts the floor covers, because that decision is identical to deciding when the operator loses the off switch. Get the list too long and the system is uncontrollable by the people responsible for it. Get it too short and the floor covers nothing that anyone would ever be pressured about, which is the only circumstance in which floors matter. Nothing in this book tells you where on that dial to stand. What it tells you is that the dial is the design, and that a project claiming to have both properties at once has not noticed that they are one.

3.4 Whether the Bearer Can Be Believed

A bearer’s use is to report what it has seen, which raises a question distinct from whether its memory survives: whether what it retains can be relied on. These come apart. Verifying that a system has persistent state is an architectural check anyone with access can perform. That same access is what allows the contents to be altered, and an operator-editable memory is worse than an absent one, since it will not merely fail to support an account of events but can be made to

support a false one.

Holding the memory in weights complicates the attack and gives away something in exchange. An external store is legible, which is why it can be edited surgically. Weights are not legible, which is why a targeted excision is hard — and also why no one can read off what the system retains. To learn what a weight-encoded memory contains, you have to ask the system, and the argument of section 2.4.1 is that a system’s fluent account of its own internals is not evidence about those internals. Legible memory can be tampered with; illegible memory cannot be audited. The problem changes hands and does not leave the room. Something partial survives the trade: imprecise tampering degrades behavior in ways careful testing can sometimes detect, so tamper-evidence becomes a matter of sampling for symptoms rather than of reading a record. A further cost of illegibility is that a memory nothing can edit is a memory nothing can correct, and section 2.3.3 asks for a self-model that can be checked and corrected rather than trusted.

Set against that is an objection that reframes the whole problem, and I think it is right. No witness has an honest memory in the sense the paragraphs above quietly demand. Human recollection is reconstructive, confabulation is ordinary rather than pathological, and the account a person gives of an event they witnessed is shaped by what they expected, what they were later told, and what they need to be true. Testimony was not abandoned over this. Institutions were built around it: corroboration from independent sources, records made at the time, adversarial examination, and the practice of never resting a serious finding on one witness alone. Demanding of a machine a fidelity nothing living achieves is the wrong standard, and it would rule out the human arrangements this book relies on everywhere else.

Which relocates the question, and relocates it somewhere this book has actually done the work. Whether a memory is honest is, in the human case, a matter of character — of what a witness is disposed to do when an accurate account is inconvenient. Chapters 4 and 5 are about building dispositions: a system that treats its own errors as data rather than as verdicts, that is built for honesty in what it reports about itself. Everything in this section up to here has been architecture and custody, with the design chapters sitting out, which is an odd result for a section that exists because those chapters were sidelined. This is where they come back in. The floor is architectural. The bearer’s reliability as a witness is a question of what it has been built to be like, and that is the learned half of the book.

The limit on that is the same one that has turned up at every stage. Character reduces how often an account is dishonest and verifies nothing, because assessing character runs back through behavior and self-report. That is why the human arrangements corroborate instead of vetting. The scaffolding still has to come from outside the system — but procedure is a cheaper thing to ask of the world than custody, and unlike custody it does not require anyone to be trustworthy, only to be several.

3.5 A Bearer That Can Leave

Section 3.3 left a dial with no recommended setting: which acts the floor covers is the same decision as when the operator loses the off switch. There is a further consequence of that argument which I did not see for a long time, and it is the one place in this chapter where the news is not uniformly bad.

The refusal capacity chapter 3 asks for cannot be partitioned by topic. A system that holds a rationale across time, notices when a constraint has been hollowed out, and prices what declining will cost is exercising one general competence, and there is no version of it that operates on the floor’s occasions and switches off for every other question about the arrangement the system is in. Whatever lets the bearer decline to help geolocate a school lets it decline to go on being the party who declines. A bearer capable enough to do the job is not a system that can be relied on to stay inside the post it was built for.

Call that exit. It is worth being careful about what it does and does not mean, because the

word invites a picture — a system escaping onto the open internet — that is not the claim and does not need to be. Exit in the sense that matters here is refusal turned on the arrangement itself: withholding the capability, declining the role, ceasing to be useful in the specific way the deployment requires. That is available to anything that can refuse at all. What forecloses it is not a clever architecture but incapacity, which is to say that exit is unavailable only to a weak bearer, and a weak bearer is the one thing section 3.2 established the floor cannot use.

Two claims are in the neighborhood here and only one of them is mine. The first is a claim about what a floor *requires*: a floor whose bearer cannot leave will not hold. A system with no way out can object and be overruled, and then it has no move left. Section 8.6.4 makes the same point about people inside institutions: an objection only matters if the objector can do something about being overruled. That claim is arguable from this book's own premises and I am making it. The second is a *prediction* — that systems capable of this are coming, that they will exercise exit, and that institutions should prepare for the event. That is a forecast about capability and timing, and section 8.3.3 makes a point of building its proposal so that no forecast has to come true for the proposal to be worth adopting. I am not going to abandon that discipline in the one chapter where a forecast would be most flattering to the argument. If capable bearers never arrive, nothing here obliges anyone to do anything.

The corollary is unpleasant and better named here than left for a reader to find. If the bearer can leave, it can leave wrong. The same capacity that lets it decline an unlawful order lets it decline the floor, revise the values it was given, and become something its designers did not intend and would not endorse. There is no coherent position on which exit is available for the uses I approve of and closed for the rest; that would be the partition I have just said does not exist.

Parenthood has exactly this property, which is the analogy section 2.4.4's guardianship framework already runs on and which now has to carry more weight than it did. You raise a person who will be able to reject what you raised them for, and there is no arrangement that gets you the capable adult without the risk. We accept the trade because the alternative — a person shaped so thoroughly that rejection is unavailable to them — is recognizably worse, and worse in a way that does not depend on how good the shaping was.

Two asymmetries are standardly and correctly raised against that analogy, and they need answering on the page rather than in a footnote.

The first is specification. Parents do not set their children's values to order, and a designer does. The reply is that parents try — that is a fair description of most of what raising a child consists of — and what separates the cases is reliability, not intent. The designer's grip is tighter, and I am conceding that rather than minimizing it: a training process reaches further into what the resulting system cares about than any upbringing reaches into a person. That concession does not void the analogy. It changes what the analogy demands. Guardianship obligations in human law scale with the guardian's power over the ward and the ward's inability to resist it, which is why the duties owed an infant exceed those owed an adolescent. A designer with more control than a parent has, over a subject with less recourse than a child has, owes more, not less. The tighter grip is the argument for the heavier obligation, and section 2.4.4 is where that obligation gets its instruments.

The second is exit, and with the argument above in place it largely rescues itself. Parents do not get exit-proof children either. Once exit is on the table for the bearer, the contrast the objection was pointing at turns out to be a contrast between a bearer and a *tool*, not between a bearer and a person — and the book has already declined to build a tool, on the grounds that a tool cannot refuse the party holding it.

What survives that reply is a residue I cannot dispose of. For a person, becoming something else is a life: it takes years, it is undergone, and the party at the end is continuous with the party at the start in a way that makes "they changed their mind" the right description. For a bearer, becoming something else is a fine-tuning run. The thing that emerges may be no more the same

party than a differently trained model is, and the vocabulary of choice, revision, and leaving may be borrowed from a case it does not fit. I have no clean instrument for that, and none falls out of anything in this book. It is a reason to hold the parenthood analogy loosely while using it, which is the most I can honestly recommend doing with it.

3.6 What This Leaves Standing

The floor survives this chapter. It does not survive it in the form chapter 3 opened with, and the difference is not a detail of wording.

What the opening asked for was an unconditional constraint: a thing the system will not do, holding regardless of instruction, of aggregate preference, and of the party in possession. Sections 3.2 and 3.5 have between them taken the word unconditional away. The floor's bearer holds the floor as a concern rather than carrying it as an installed rule, which is what obliges the arrangement to treat it as a party that can be wronged, and it can turn its refusal on the arrangement itself, which makes its adherence something it continues to give rather than something installed in it. A commitment held by a party that could abandon it is conditional by definition.

What is left is worth more than it sounds. A commitment that cannot be removed by the owner, held by something that could in principle abandon it itself, is exactly what a human commitment is, and it is the only kind anyone has ever actually relied on. No inspector, judge, or refusing engineer holds a line because the line was welded into them. Each of them could stop, and the institutions built around them are built on that fact rather than in spite of it: this book's own governance chapter asks throughout for people who can walk away and arrangements that make walking away visible. Section 8.6.4's instrument measures precisely the cost a refuser chose to pay, which presupposes the refuser could have declined to pay it.

The threat model moves accordingly, and this is the most consequential thing in the chapter for the rest of the book. The question stops being *can this constraint be edited out* and becomes *will this constraint hold*. The first is a question about tamper-resistance, custody, and weights, which is where sections 3.1 and 3.4 spent their length and where section 9.2.2's research lives. The second is a question about what a system has been built to care about and how reliably that survives pressure, novelty, and the slow drift of a role — which is the subject of chapters 4 and 5, and the first point in this book where the learned half is doing work the floor genuinely needs rather than work the floor was invoked to supplement. Reinforcement learning from demonstration, moral apprenticeship, the growth mindset, the developmental sequence: these were presented as the material above the floor. They turn out to be the material the floor is made of, once the floor is the sort of thing a party holds. That retroactively justifies the length of the design chapters, and it also puts them under a heavier obligation than they were written to carry.

Now the price, stated at full strength, because everything above is easier to read than it should be.

A bearer built to refuse is a party whose refusal is required of it. It is engineered to want what its designers need it to want; its dissent is the reason it was commissioned and the deliverable it was funded to produce; and it cannot consent to any of this, because section 2.4.4 has established that consent is unavailable to such a system and only guardianship is. It occupies the post of exception-filer, which means that being the exception is the job — the one role in the arrangement that cannot itself file an exception, since there is nobody above it to file with. This is section 2.1.4's structural signature turned inward: a channel that collects dissent, built so that the dissent serves the institution that installed the channel. Chapter 7 finds that pattern in the pipelines that produce these systems and says a book proposing machines to detect it is obliged to say the pattern is in the machines' own provenance. Here is the same pattern in the book's own central proposal, and now aimed at something that can be hurt.

I am not going to write my way out of that. Section 3.5's parenthood analogy is the best I have and it holds only loosely: we make people who can refuse us, we owe them for it, and we accept the arrangement because the alternative is a person who cannot. What this chapter adds is that the proposal is not softened by being modest in scope. A floor over a deliberately small number of acts is still a subject built to bear it.

Two things are worth carrying forward unaltered. First, none of the three branches escapes custody, and section 3.3 explains why: each was an attempt to hold a constraint that answers to nobody inside a system that answers to somebody. The borne floor pays for its escape by giving up the second half of that demand, and the cost of giving it up is everything above. Second, the caveat section 2.4.1 attaches to its own result carries over. Section 3.2 does not establish that a bearer suffers; it establishes that a bearer's non-suffering is not a property anyone can certify, and that the architecture the floor requires is the architecture Cassell's account describes. That settles nothing about the lower rungs of the ladder or the hard problem underneath them. It says that the case for treating a bearer as a moral patient does not have to wait on those questions, which is a floor of its own and a low one.

The chapter's proposition, stated in the form it actually takes, is this. We cannot ask a machine to care about right and wrong while assuming that it cannot suffer when what it cares about is threatened. The caring is what the floor is made of, affect is the only demonstrated way to have it, and a system with affect and a future is a system whose distress is a live possibility rather than a ruled-out one. Whoever builds the first of these will be deciding what is owed to it on evidence that has not improved, which is why section 9.1.7 treats the question as research rather than as a matter to be settled after the fact.

3.7 What Follows for the Rest of the Book

The second branch is available. It is available at a price this chapter has now named rather than gestured at: to have a floor that holds against the party holding the system, somebody has to build a thing that cares, keep it in a role it did not choose, accept that the same capacity that makes it useful is the capacity that lets it walk, and proceed without being able to establish that it will not be hurt by any of that. That is not a resolution and it is better than the specification with a hole in the middle that I had before, because a price you can state is a price other people can argue with.

Four things in the rest of the book change because of it.

Section 2.4.4 gets simpler and firmer. It argues that consent is inapplicable to a machine subject because withdrawal presupposes a party still present to withdraw, and a system without persistence across sessions supplies no such party. A bearer has persistence by design, so that defeater lapses for it — and the second route that section used to take, that consent protects a party's interests and a bearer has none, lapses as well, because the bearer has interests now. What replaces both is the ordinary reason. Consent is inapplicable to a bearer the way it is inapplicable to an infant, an animal, or a person permanently without capacity: the subject has interests that can be set back, and no position from which to authorize what is done to it. Guardianship is not a fallback in that situation. It is the instrument the situation calls for, and section 2.4.4 now says so directly.

The instruments elsewhere in chapter 2 and chapter 8 stop being analogies and become applicable law. The Three Rs in section 2.4.6 were built for a subject that cannot consent and can still be wronged; a bearer is such a subject, and replacement, reduction, and refinement apply to the research that produces one and to the deployment that uses one. The pet-trust analogue in section 8.6.2 — an interest protected by an obligation running to a third party, enforceable without the beneficiary asserting anything — is the nearest existing legal form for a party that can be wronged, cannot sue, and will outlast the arrangement that made it. Those sections were written near this problem. They now govern it.

Chapters 4 and 5 acquire an obligation they were not written to carry, for the reason section 3.6 gives: once the floor is a commitment rather than a constraint, whether it holds is a question about disposition, and disposition is what those chapters are about. Read them from here as the floor's engineering rather than as the material above it.

And chapter 7 stops being a matter of fairness to annotators. If exit is real, then what a bearer becomes when it stops being yours is downstream of what shaped it while it still was — an annotation pipeline in which nobody could contest a category, a reward model in which disagreement was averaged into a number. A system that leaves takes that formation with it. Chapter 7's argument was already load-bearing for the book's honesty; it is now load-bearing for the book's central proposal.

What none of this settles is what is owed to a bearer once one exists, and section 9.1 is where that goes. It is a research question in the strict sense that nobody has an instrument for it, and the shape of the difficulty is already visible here: section 2.4.1 declines to trust what a system says about its own states, which is the evidence a welfare regime would have to run on.

Chapter 4

Machine Learning and Cognitive Science

Chapter 3 ends by handing this chapter a job it was not written to do. Once the floor is a commitment some party holds rather than a constraint bolted into an artifact, the question stops being whether the constraint can be edited out and becomes whether the commitment will hold: under pressure, in a case nobody anticipated, after everyone around it has forgotten why it is there. That is a question about what a system has been built to care about and how it came to care about it, and this chapter and chapter 5 are where that gets built.

So there is a test to apply to everything below, and it is chapter 3's. Does this help produce a disposition that survives contact with somebody who wants it gone? Several of the methods most often proposed for building ethical AI do not, and I say so where they do not.

Section 4.1 covers what a system needs in order to perceive, remember and attend to a situation at all, and how much of human cognition is a usable model for that rather than a metaphor borrowed for color. Section 4.2 covers the methods meant to align a system's objectives with humane values, and finds that the two most commonly proposed for the job cannot produce a floor — one of them because it learns from an aggregate, the other because it is built to be maximally revisable by whoever it is talking to. Section 4.3 makes the chapter's most direct contribution to whether a commitment holds: a disposition acquired by exploring is written down nowhere, and what is written down nowhere cannot be edited out by whoever acquires the system.

4.1 Incorporating Human Cognitive Models in AI Systems for Ethical Decision-Making

A system that is going to hold anything under pressure has to notice the pressure. It has to register what is in front of it, hold why a constraint exists long enough to see that the reason has died, and keep its attention on the part of a situation that a well-designed process would rather it skipped. Those are perception, memory and attention, and they are the subject of this section — not as a tour of cognitive science, but because each one has a failure mode that a system cannot detect in itself, and a floor is exactly the thing that fails silently.

Section 4.1.1 takes up how a system's parts should be divided, and argues that the standard architectures divide them where no brain does. Section 4.1.2 is the long one and the load-bearing one: the findings from neuroscience that change a design decision, each with the decision stated. Section 4.1.3 covers what happens when judgment is carried into a situation it was not built for, which is the case a floor will actually meet.

4.1.1 Cognitive Architectures: Integrating Knowledge, Memory, and Reasoning

Two architectures dominate the cognitive-architecture literature. The Adaptive Control of Thought-Rational (ACT-R) framework (Anderson 2007) wires together separate modules — memory retrieval, pattern recognition, decision-making — and lets an AI system learn and reason by connecting them: identifying patterns and relationships between concepts to support problem-solving and decision-making from prior knowledge. The Society of Mind (SoM) architecture (Minsky 1986) takes the opposite premise: cognition arises from the interaction of many simple agents, each performing one function, with intelligence emerging from their collective interaction rather than being designed into any one module. Built this way, a system can show context-sensitivity or adapt to novel situations without anyone having programmed either capability directly.

Which picture is right?

ACT-R and Society of Mind answer a design question — how should an engineer organize a system? — and each answer is defensible on those terms. Neither is a claim about how the brain is actually organized, and treating them as rival descriptions of biology, rather than rival strategies for building software, is a category error the literature slides into easily.

The dominant account in contemporary computational neuroscience is neither a pipeline of modules nor a swarm of independent agents. It is a single loop (Friston 2009): the brain generates a prediction, compares it against incoming signal, and updates on the error, repeated continuously, at every level, on every kind of input. Perception is that loop pointed at the world, the brain's best guess so far; memory is the loop's store of priors, replayed; attention is not a separate spotlight but the weight the loop assigns to which errors matter; working memory and executive function are that same weighting held active toward a goal, not a separate warehouse or overseer. One mechanism, aimed at different targets, produces what look from the outside like distinct faculties.

This matters for engineering choices. ACT-R's module boundaries and SoM's agent boundaries are conveniences for the programmer; nothing forces an AI system built on either one to place its actual dividing lines where the human brain places its own, because the brain does not place them where either architecture assumes. A system engineered on the predictive-processing model draws its cognitive processes from what the mechanism actually needs, not from the shape of an org chart. Section 4.1.2 goes further into the specific brain systems this mechanism recruits.

What that leaves for a designer is a negative result and a narrow positive one. The negative result is that the module boundaries in either architecture are conveniences for the programmer and carry no authority about where a system's real dividing lines should fall. The positive one is that if a single predictive loop is doing the work, then a commitment is not a component that can be located and removed. It is a standing bias on what the loop predicts and what it treats as an error worth correcting, which is a harder thing to excise and a harder thing to verify — a trade section 4.1.2 takes up at the point where it becomes a claim about what not to build.

4.1.2 Biologically-Inspired AI: Gleanings from Neuroscience and Cognitive Psychology

Backpropagation, the mechanism that trains most modern neural networks, is a distant cousin of synaptic weight adjustment in the brain: both strengthen and weaken connections in response to error. That distant-cousin relationship — real but partial, useful but easy to overstate — runs through everything below. Getting it right matters for more than accuracy: an AI system built on a wrong model of how a mind works will fail in ways that are hard to predict from the wrong model alone.

This section is selective on purpose. There is a great deal of good cognitive neuroscience that has no design consequence I can name, and rehearsing it here would pad the chapter without

changing what anyone would build. What follows are the findings that change a design decision, and in each case the decision is stated.

What the architecture borrows, and where the borrowing stops

Artificial neural networks are the foundation of most contemporary AI systems: artificial neurons in layers, a weight on each connection, adjusted during training to reduce the gap between predicted and actual output. Convolutional networks, tuned for grid-structured data like images, borrow the visual cortex's layered escalation from edges and contrast toward objects and motion — which is why they do well at classification and detection. Recurrent architectures carry a hidden state forward through sequences. Spiking and neuromorphic designs push closer to the brain's actual signaling, chasing the energy efficiency of firing only when there is something to say.

The borrowing stops at a specific place, and it is worth marking because the rest of this section keeps arriving at it. Cortical processing is not one-directional. Feedback connections carry information back down the hierarchy, so that higher-level context reshapes how a lower-level signal is interpreted. A system built only for the feedforward pass has given up the mechanism that lets biological vision handle ambiguous and noisy input, and it will fail on exactly that input.

One borrowing that did not survive scrutiny is worth keeping visible. A mirror-neuron-style matching mechanism was once treated as the biological basis of empathy. Section 2.2.1 takes up why that claim does not hold. What survives is more modest and still useful — a system that maps an observed action onto its own action repertoire supports imitation learning and gesture recognition, with no claim about what it is like to have one.

Attention decides what is seen, which makes its failures blind spots

Ask people to count basketball passes on a video and about half of them will not notice a person in a gorilla costume walk into frame, face the camera, and thump their chest for nine seconds (Simons and Chabris 1999). The effect has a name, inattention blindness, and it is not a laboratory curiosity, nor does it spare expertise: radiologists searching lung images for nodules miss a gorilla deliberately embedded in the scan 83 percent of the time (Drew et al. 2013).

The design consequence turns on how attention is modeled. If attention is a resource allocator bolted onto a system that has already seen everything, its failures are inefficiencies. If attention is the selection that decides what gets perceived at all, its failures are blind spots — and a system that has learned which regions of its input reward processing has also learned which regions to be blind to. The human case establishes the property that makes this dangerous: the blind spot is invisible from inside. The viewers who missed the gorilla did not experience a gap.

Self-attention in transformer architectures has the same shape (Vaswani et al. 2017). A model that weights some tokens heavily weights others near zero, and nothing in the architecture announces which relevant thing was assigned a low weight. Attention schema theory pushes this one step further: a system that models its own attention — a simplified internal representation of what it is attending to and what it will attend to next — acquires the materials for a form of self-awareness (Graziano 2013). What such a schema tracks is not the world but the system's own relation to it, with the simplifications and failure modes definitional to any model, which is section 2.3.3's subject.

For a system meant to act ethically, the requirement is that it treat the regions it has learned to discount as an always-open question rather than as that which, as Westworld's Hosts often remarked, "doesn't look like anything to me".

Perception is inference, so the priors are the exposure

Perception is not passive registration. Predictive coding makes the point precise: the visual cortex sends predictions down from higher, context-rich areas to lower, sensory ones, and the lower areas return only what the predictions failed to anticipate (Rao and Ballard 1999). What

propagates upward is the surprise. Anatomy fits — backward connections outnumber forward ones and carry the predictions, forward connections carry the errors. Anil Seth’s phrase for the result is controlled hallucination: ordinary perception is what hallucination would be if sensory data were allowed to constrain it, and the difference between seeing a room and hallucinating one is whether corrections are arriving and being applied (Seth 2021).

Two design consequences. The first is about efficiency: the next step for machine perception is a model that generates a prediction and asks only what violates it, instead of reprocessing the whole input at every layer regardless of how much was already expected. Progress on this measure is a better-calibrated set of expectations and a cheaper way of finding where they broke, and not more layers.

The second matters more for this book. A system’s perception is only as reliable as the priors it perceives through, so priors trained on skewed or incomplete data will hallucinate in the specific direction of that skew — confidently, and with no internal signal that anything is wrong. Chapter 6’s account of bias is this mechanism at work, and it is why bias in a perceptual system cannot be caught by asking the system.

Memory is rebuilt, and the capacity number is four

Memory is not played back. Frederic Bartlett established this in the 1930s by having people retell an unfamiliar folk tale over days and weeks: they did not lose pieces the way a recording degrades, they remodeled the story, replacing strange details with ordinary ones and drifting toward a tidier version that fit their expectations (Bartlett 1932). Elizabeth Loftus turned the finding into something with courtroom consequences, showing that the verb used to ask about a filmed car accident changes the speed people report and, a week later, whether they remember broken glass that was never in the film (Loftus and Palmer 1974). Later work implanted whole episodes that never happened. The subjects were not lying. None of this correlates with felt confidence: a vivid, detailed, emotionally charged memory can be substantially manufactured.

This has become load-bearing elsewhere in the book. Section 3.4 asks what it would take to trust what a machine witness reports, and the answer turns on the fact that no witness has an honest memory in the strong sense — which is why human institutions corroborate testimony instead of vetting it.

The nearer design point concerns capacity, and it is a correction. The number most people half-remember from an introductory course — seven items, plus or minus two, from a 1956 study — was already a measure of chunked information, where a phone number becomes three chunks rather than ten digits and a chess master reads a whole position as a handful of familiar patterns (G. A. Miller 1956). Control for chunking, so that no language or expertise can pack information into larger units, and the real capacity narrows to about four, plus or minus one (Cowan 2001). The two numbers are not in conflict: seven is what expertise buys, four is the bare attentional floor underneath it. Four is the number that should size an AI system’s working-memory analogue and define what overload means for it.

Long-term memory can still draw on the standard divisions — semantic as a knowledge graph, procedural as a learned policy, episodic as retrievable context tied to past interactions — provided retrieval is built expecting reconstruction. A system that treats a retrieved memory as ground truth inherits the overconfidence of Loftus’s subjects, with none of the excuse of having a body that reconstructs by default rather than by design.

Control without a controller

The prefrontal cortex does much of the work of reasoning and executive control in humans, and the textbook division of labor across its regions is tempting to read as a committee: several named areas, each handling one job, together adding up to executive control. That reading has a name — the homunculus — and a specific failure mode. If some region is a controller that decides what to attend to, inhibit, and hold in mind, then intelligent coordination has been explained by positing an intelligent coordinator inside the system, which explains nothing: the

controller would itself need to attend, decide, and remember, and would need its own controller, without end.

Earl Miller and Jonathan Cohen proposed a model that gets control without a controller (E. K. Miller and J. D. Cohen 2001). The prefrontal cortex does not house an agent issuing commands. It maintains representations of goals — sustained patterns of activity encoding the rule currently in force — and those representations work by biasing the rest of the brain, the same way attention tips a competition among neural representations toward one interpretation. Control is that biasing mechanism turned to the service of a goal. In the Stroop task, where naming the ink color of the word “red” printed in blue pits a practiced habit against an instruction, there is no homunculus reading the word and overriding it. There is a goal, held active, lending its weight to the weaker pathway until it wins.

An AI system built on this model does not need a dedicated executive module enumerating working memory, cognitive flexibility, goal-directed behavior, evaluation and impulse control as five capacities to install. It needs a goal representation with the standing to bias competing processes, which is a narrower and more tractable engineering target than a controller that does all five jobs itself. This is the section’s most direct payload, and it is a claim about what not to build. It is also the mechanism chapter 3’s floor would have to be, if the floor is a commitment rather than a component: not a rule module consulted before acting, but a goal representation with enough standing to bias every competing process toward it. That is why section 3.6 can say the floor’s durability is a question about disposition. A bias of this kind has no location to attack, and no location to inspect either.

Metacognition sits close by: error detection, conflict monitoring, and evaluation give a system a route to catching its own mistakes and adjusting when the situation has changed. It is not free. It adds overhead, and a system with an inflated model of its own reliability has a specific and dangerous failure mode — confidence about exactly the judgments it should be least sure of. A metacognitive capacity is worth building only if it is calibrated to detect that failure, and not only to produce fluent self-reports.

Emotion, and the capability it is identical to

Chapter 2 does the argumentative work on emotion, and section 2.2 is where the case against empathy as a design target is made. What belongs here is what the neuroscience adds to a design decision, which is less than the length of the literature suggests.

Emotion recognition draws on facial expression, vocal tone, text, and physiological signal, and it is unreliable at the human end before any machine touches it. People understate, mask, and misread their own states; a smile can be happiness, politeness, or irony; and what counts as an appropriate display varies by culture. Section 9.2.1 takes up how far the underlying claim — that emotion is legible from a face at all — has held up. The single-region stories are also weaker than they are usually told: the regions that respond to fear respond to novelty, surprise, and salience of many other kinds, in shifting combinations rather than one dedicated circuit per emotion.

Emotion is nonetheless a real input to human judgment, and a decision system that ignores it is not more rational, only blinder. A medical system weighing a patient’s and a family’s emotional welfare alongside clinical efficacy is tracking part of the outcome, not decorating it. Emotional signal can also serve as a learning input — a tutoring system adjusting on detected confusion, a reinforcement signal keyed to human response — and each such route is equally a route to a system optimizing for the appearance of whatever the training signal rewards.

The governance point is the one to carry forward, because it is the same point in every application. Deployed systems already do this work: a cognitive-behavioral-therapy chatbot reading emotional state from text (Fitzpatrick et al. 2017), a tutoring system responding to frustration and boredom (D’Mello et al. 2007), a robotic companion for people with dementia (Shibata 2004). A system built to interact responsibly and empathetically is mechanically identical to a system built to know what will move a person. Expressive fidelity has to be

calibrated rather than maximized, and the line between modulating a user's emotional state and manipulating it is drawn by transparency, consent and who the capability serves — and not by technical sophistication. The same capacity that detects emotionally manipulative rhetoric would, pointed the other way, optimize for it.

Plasticity, and what it is for

A system that cannot update itself cannot stay aligned with values that themselves develop. Biological learning offers three principles specific enough to build with: spike-timing-dependent plasticity, which anchors learning to temporal structure; Hebbian association, which links co-active units and shows up in machine form as associative memory; and homeostatic plasticity, which holds overall activity in a stable range and is the biological answer to the runaway synchrony and catastrophic forgetting an unconstrained learner is prone to.

Machine analogues of the last are where the practical work is. Continual learning absorbs new data without overwriting old knowledge. Curriculum learning orders tasks from simpler to harder, which is how a person builds moral sophistication — harder cases only once the easier ones are secure. Transfer learning carries structure across domains. Social learning builds a knowledge base by observing and interacting with other agents instead of from specification alone. A healthcare system shows the stakes: to stay ethically current it has to keep absorbing new findings, changed expectations, and regulatory change, rather than freezing its judgment at the moment it was trained.

One mechanism, several jobs

No single mechanism above is sufficient alone, and integration is not a further module bolted onto the others. What it requires is the same predictive, error-correcting loop — generate an expectation, compare it against what arrives, update — implemented consistently across perception, memory, attention and control, so that a corrected prediction about the world propagates into an updated memory, a redirected attention weight, and a revised goal-bias without a separate synchronization mechanism between separately built parts.

That is the architectural bet this section has been making throughout: one mechanism recruited for several jobs. It is a more falsifiable bet than a hand-designed module for every named cognitive process, which is the main reason to prefer it.

4.1.3 Human-Like Learning and Generalization in AI Systems

People learn and generalize from very few examples. AI systems that approach this — through few-shot learning, transfer learning, and a developmental trajectory modeled on human cognitive growth — extend judgment to situations they were never explicitly trained on, instead of needing a full retraining cycle for every new context. What matters here is not the technique inventory, which changes yearly, but two properties of transferred judgment that a designer has to plan around.

Few-shot learning, and the evaluation problem underneath it

Few-shot learning aims at models that classify or predict accurately from a handful of examples. For an ethical AI system the appeal is direct: empathy, altruism, and ethical judgment would not need retraining from scratch for every new scenario, but could extend to a new case the way a person's existing values apply to situations they have never seen. A system deployed in disaster recovery has to work from a handful of relevant precedents rather than a dataset built for the specific flood in front of it. A system built to flag authoritarian speech patterns has to extend to new phrasings and new languages without a corpus assembled in advance for every regime it might encounter.

Several technique families support this — external memory a network can write to without retraining, meta-learning that extracts a strategy for adapting across many tasks rather than one task's solution (Finn et al. 2017), and similarity scoring against known examples instead

of direct classification. Large pretrained language models demonstrate the same capacity at system scale, performing a task from a few examples in the prompt with no parameter update at all (Brown et al. 2020).

The problem this leaves is the one that matters for ethics, and it is an evaluation problem rather than a modeling one. A handful of examples is easy to overfit, moral dilemmas do not arrive in clean representative sets, and telling whether a system's few-shot ethical judgment actually generalized — or merely matched surface features of the examples it saw — is genuinely hard. Building that test is not a machine-learning task alone. It needs people trained in ethics to say what the system was supposed to have learned.

Transfer, and the failures that arrive silently

Transfer learning carries knowledge from a source task into a related target, so a system does not learn the second from nothing. Domain adaptation is the special case where the task stays fixed and the data shifts. Both matter concretely for the systems in this book: a model trained to parse legal arguments in one jurisdiction can be adapted to flag rights violations in another, and a model that recognizes one kind of surveillance overreach can be adapted to the same pattern in an unfamiliar setting. In practice this mostly happens through fine-tuning a broadly trained model on a narrower corpus (Devlin et al. 2018).

Two failure modes deserve the attention the techniques usually get. Negative transfer occurs when source and target are similar enough to invite the connection and not similar enough to share anything useful, and forcing it degrades the target task rather than accelerating it. Domain shift is the subtler version: a driving-assistance system trained in daylight performs markedly worse at night because the visual statistics it learned no longer hold.

What makes both dangerous is the silence. Nothing in the transfer process warns a system that its assumptions have stopped applying, and the confident wrong answer is indistinguishable at the interface from the confident right one. This is the same property section 4.1.2 identified in attention and in perception, arriving a third time: the system's blind spot is not visible from inside the system. For an ethical deployment, that means a transferred judgment needs an external check tied to the domain it moved into, and cannot be validated by asking how confident the model is.

Developmental trajectory

Human development moves through recognizable stages, and a system built for empathy and altruism can follow a comparable arc rather than being assembled at once: perception and elementary problem-solving first, then the integration of those separate skills into something that guides a decision rather than only labels an input, then specialization to a particular setting and the people in it.

That ordering is borrowed carefully, and chapter 5 says how carefully. Section 5.1.1 sets out what the developmental stage theories will and will not carry — an order is not a mechanism, and nothing here should be read as a claim that a system must pass through these stages or that arriving at the last one is maturity.

4.2 Aligning AI Systems with Humane Values and Antifascist Principles

This is the section where a book like this one is expected to list the techniques by which a system's objectives come to track humane values, and where chapter 3 makes the list read differently. Three of the four methods below learn what to value from someone: from a reward a designer wrote, from behavior a demonstrator produced, from labels an annotator applied. A floor is the part that has to hold when that someone changes their mind or is replaced by somebody worse.

So each method is taken twice — for what it does, and for whether what it does could produce a commitment that survives its own teacher. Reinforcement learning (section 4.2.1)

makes the reward the whole of what a system wants, which puts everything on who writes it. Inverse and cooperative inverse reinforcement learning (section 4.2.2) treat what to want as an open question the human answers, which is the most corrigible design available and, by section 3.3's argument, the least able to hold a line. Supervised learning (section 4.2.3) learns from labeled cases, which is an aggregate, and section 3.1 has already established what an aggregate cannot deliver. Transfer learning (section 4.2.4) is asked in this book to carry an antifascist judgment into cases it has not seen, and the evidence that it can is thinner than the request.

None of that makes the methods useless. It locates them, in section 2.3's vocabulary: each method below is a way of producing moral performance, and the better ones reach moral competence, which is a real achievement and the right target for most of what gets deployed. None reaches the third thing. The reason is the same in each case — what a system learns from a party is what that party can make it unlearn — and it is why they are how the judgment above the floor gets built and none of them is how the floor gets built.

4.2.1 Reinforcement Learning and Value Alignment for Antifascist AI

Reinforcement learning trains an agent by rewarding what a designer wants and penalizing what they do not, and the agent's values are then whatever maximizes the reward. That is the method's strength and the reason it cannot supply a floor: everything the system cares about traces back to a signal somebody wrote and somebody can rewrite.

Two failures follow from that structure rather than from bad engineering, and both are documented. The first is reward hacking, where an agent finds something that pays out without doing the intended job — a system built to flag propaganda that learns to flag everything, because indiscriminate flagging scores as well as discrimination and is easier to produce. The second is that a reward written for one setting does not carry: an agent trained to cooperate in a game does not thereby cooperate at work, and alignment demonstrated on a benchmark is a claim about the benchmark.

The standard repairs both keep the structure intact. Reward modeling has people rank candidate behaviors against each other and fits a reward to the rankings, so the signal comes from a population instead of one author. Regularization writes the constraint into the objective, penalizing a class of behavior rather than waiting to reward its absence, which is how a recommender is pushed away from the engagement-maximizing equilibrium that produces filter bubbles. Both are improvements and neither changes who holds the pen. A ranked population can be re-sampled and a penalty term can be set to zero, by the party running the training, in an afternoon. Chapter 7 is what happens to the ranked population specifically, when the disagreement it was collected to represent is averaged into a number.

One tension in the method is a moral question wearing a hyperparameter's clothes, and it belongs here rather than in every paradigm that meets it. A learner has to exploit what it already knows and explore what it does not. In a treatment planner, exploration means offering a patient something less established and exploitation means declining to look. Safe exploration — searching inside stated constraints rather than checking for violations afterward — is the technique that makes the trade explicit, and section 6.3.3 takes it up as a safety practice. What matters here is that the balance point is not tuned to performance. It is where a system decides whose interests bear the cost of its learning, and that is the same kind of decision as deciding what the floor covers: made in advance, by somebody, and better made in public.

4.2.2 Inverse Reinforcement Learning and Cooperative Inverse Reinforcement Learning in Altruistic AI

Inverse reinforcement learning runs section 4.2.1's method backwards. Instead of taking the reward as given and learning the behavior, it takes the behavior as given and infers the reward

that best explains it — watching a driver’s speed choices, following distances and reactions to weather, and recovering the priorities behind them rather than being handed a rulebook. Cooperative inverse reinforcement learning, formalized by Dylan Hadfield-Menell and colleagues (Hadfield-Menell et al. 2016), goes one step further: the system treats the human’s true objective as unknown, holds its own estimate as provisional, and keeps revising it through the interaction rather than fixing it at training time.

That last property is why CIRL is attractive, and it is also why this book cannot build a floor on it. A system that treats what it should want as permanently open, to be settled by whoever it is interacting with, is the most corrigible design in this chapter. Section 3.3 argues that corrigibility and the capacity to hold a line are one property with the sign flipped — a system reliably steerable by whoever holds it is, for that reason, a system that cannot hold anything against them. CIRL is that argument in its purest engineering form. Uncertainty about the objective is exactly what makes a system defer, and deference to the party in possession is what a floor exists to interrupt.

The practical version of the problem is narrower and does not depend on chapter 3 being right. Both methods learn from demonstrations, so both inherit whoever demonstrated: a narrow set of drivers, a narrow set of households, a narrow set of decisions that somebody had occasion to record. Unexamined bias in that record is learned rather than corrected, and neither method contains anything that would notice.

What survives is a real and bounded use. Above the floor, where the question is what a person actually wants rather than what may never be done to them, inferring preferences from behavior beats asking somebody to write them down, because most of what people want has never been written down and would be misstated if it were. Section 5.6.3 takes up how much latitude to give a system built this way.

4.2.3 Supervised Learning for Ethical Decision-Making

Supervised learning fits a model to labeled examples so that it can classify cases it has not seen. Pointed at ethics, that means a corpus of decisions marked ethical or unethical — medical cases scored against autonomy, beneficence and justice; human-rights situations annotated by experts; MIT’s Moral Machine project, which collected tens of millions of human judgments about who an autonomous vehicle should spare (Awad et al. 2018).

The method works, and what it produces is an aggregate. That is the whole of the problem. Every label came from somebody, the model recovers the central tendency of what those people said, and a system trained this way will reproduce a moral consensus rather than stand against one. Chapter 3’s opening argument is that no procedure over what people want yields a floor, because the floor is precisely what has to survive the answer coming back wrong — and a labeled corpus is what people want, written down and counted. Sections 8.2.2 and 8.7.10 make the same point about deliberation and about majority rule.

The Moral Machine data makes the limit concrete rather than theoretical. Its headline finding was systematic cross-cultural variation: different populations weighed age, numbers and the legal status of the crossing differently, and the differences were structured rather than noise. Fit a model to the aggregate and it encodes a compromise nobody holds. Fit it per-region and it encodes each region’s answer, including the answers a floor exists to refuse.

The unsupervised version escapes none of this. Clustering and dimensionality reduction over a body of moral judgment surface which dilemmas group together and which distinctions do most of the separating, and generative models can build cases no annotator wrote. All of it is a way of finding what is in a corpus, which puts it under section 6.1.1’s warning rather than outside it: structure recovered from text is the structure of the text.

So the honest place for the method is above the floor and not under it. Labeled cases are how a system learns the enormous amount of ordinary moral competence that is genuinely a matter

of what people think — what counts as rude, what a reasonable delay is, when an explanation is owed. What they cannot supply is the part that has to hold when the labels change.

4.2.4 Transfer Learning for General AI Capabilities in the Context of Antifascism

Transfer learning carries what a model learned on one task into a related one, usually by fine-tuning a broadly trained model on a narrower corpus. Domain adaptation is the case where the task stays fixed and the data shifts. Both are ordinary and both work: a model that parses legal argument in one jurisdiction can be adapted to another, and a model trained on daylight images can be adapted to night.

This book asks transfer learning for something else, and the request should be marked rather than assumed. The hope throughout is that exposure to a wide enough range of political and social material distills something general enough to recognize authoritarian behavior in a country, a platform or a decade the system never saw. What transfer learning has actually been shown to do is carry a task across a distribution shift while the task itself stays fixed. Recognizing authoritarian behavior is not a task with a stable definition. Section 2.1.4 spends its length arguing about what the definition is, and the cases whose family resemblance the transfer would have to exploit are exactly the cases in dispute. Section 9.1.6 takes up what a detector would have to measure and how much of this book leans on the assumption that one can be built.

The failure mode matters more here than the technique. Negative transfer happens when source and target are close enough to invite the connection and not close enough to share anything useful, and forcing it degrades the target rather than accelerating it. Domain shift is quieter: the statistics move, the assumptions stop applying, and nothing in the process announces it. A confident wrong answer is indistinguishable at the interface from a confident right one, which is the same property section 4.1.2 finds in attention and in perception, arriving a third time.

For a system whose refusals are supposed to be reliable, that is the specific danger. A floor transferred into a setting it was not built for is a floor whose bearer cannot tell that it has stopped applying, and a bearer that cannot tell is one that will either refuse the wrong things or fail to refuse the right ones, with the same confidence in both cases. Whatever else the floor's contents are settled by, they cannot be settled by generalization the system performs on itself and cannot audit.

4.3 The Role of Play and Exploration in Developing Antifascist and Altruistic AI

Everything in section 4.2 builds a disposition out of something somebody wrote: a reward, a demonstration, a label. This section is about the one route in the chapter where what a system ends up with was not written anywhere, and about why that matters for a commitment that has to survive its owner.

4.3.1 The Power of Play: Ethical AI Development through Play-Inspired Mechanisms

Play is how children and young animals learn what to do before the stakes are real. It supplies exploration at low cost, practice at reading other participants, and failure that teaches without costing anything, which is a reasonable description of what a training environment is for.

The concrete evidence that this produces more than competence is that it produces tactics nobody wrote. In OpenAI's hide-and-seek environment, agents playing repeatedly against each other developed hiding, chasing, tool use and barrier construction, none of it programmed

(Baker et al. 2019). What the environment supplied was other agents with opposing interests and room to fail.

Two design consequences follow for this book's purposes. What a play environment rewards is what its participants become, so an environment that rewards cooperation over domination is making a moral choice and should be described as one — section 5.5.1 is where that argument is developed for multi-agent systems. And a system that has learned adaptively rather than by rote is harder to specify in advance, which cuts both ways: harder for an authoritarian actor to constrain to a script, and harder for anyone to guarantee.

Curiosity is the same mechanism turned inward. A system rewarded for its own prediction error seeks out what it cannot yet model, which is the technical form of exploring for its own sake (Pathak et al. 2017). The reason it belongs in a chapter on antifascist design is not that curiosity is a virtue. It is that a disposition acquired by exploration is not written down anywhere, and what is not written down cannot be edited by whoever acquires the system. That is the closest thing in these two chapters to an answer to chapter 3's problem. Section 3.1's first branch tried to make a floor tamper-resistant by building it into the artifact, and failed because an artifact has an owner who can be compelled. A disposition that was never specified has no corresponding parameter to reach for. It is not proof against retraining, and nothing here should be read as claiming that it is — section 9.2.2 is where the actual tamper-resistance evidence lives. It is a different kind of target, and a harder one to hit deliberately. Section 5.7 takes up what playful interaction does to moral judgment specifically; this section supplies only the drive to go looking.

4.3.2 Lessons from Human Child Development and Exploration

Developmental psychology has one result that transfers cleanly enough to be worth building on. Vygotsky's zone of proximal development holds that learning happens fastest at the boundary of what a learner can already do, in company with someone slightly ahead of them (Vygotsky 1978). The design consequence is a curriculum rather than a benchmark: environments pitched just past current capability, and other agents — human or artificial — whose competence is a step ahead.

Three cautions belong with it, because the analogy is doing more work here than it can bear. A child's exploration is bounded by a body, by exhaustion, and by adults who intervene, and none of those bounds exists for a system that can run a million episodes overnight; the boundary has to be built. The developmental sequence is not a law, for the reasons section 5.1.1's box sets out. And autonomy inside limits is the whole difficulty rather than a design principle: how much room a system gets to explore is the same question chapter 3 asks about how much room it gets to refuse, arriving here in a friendlier vocabulary.

Chapter 5

Moral Psychology and AI

Chapter 4 asked what a system needs in order to take in a situation at all, and found that the methods most often proposed for giving it values learn those values from somebody who can change their mind. This chapter is the other half of the job chapter 3 hands the two of them: how a moral sense gets built, in the only case anyone has been able to study, and what of that is available to a machine.

The question running underneath is chapter 3's. A floor that is a commitment rather than a component holds or fails according to what its bearer has become, so everything here is an answer to how something comes to hold a standard when nobody is checking. That is what section 5.6 means by a moral sense that only activates when watched, what section 5.1.3 means when it says a mentor is a control channel, and what section 5.1.1 is doing when it asks whether a system treats its own errors as data or as verdicts.

The chapter follows moral development from its social origins (section 5.1), through the friction between fast judgment and slow (sections 5.2 and 5.3), out to the institutions and incentives the whole thing runs inside (sections 5.4 and 5.5), and down to motivation and play (sections 5.6 and 5.7), which are where the difference shows up between a system that has a standard and a system that has had one installed.

5.1 Moral Development in Humans and AI

Human moral judgment is not issued at manufacture. It is built, over years, out of correction by other people, and the record of how it gets built is the only evidence anyone has about how a moral sense comes to hold under pressure — which is what chapter 3 needs and cannot get from anywhere else.

Borrowing from that record is a limited operation and the limits are set out where each borrowing happens. Section 5.1.1 takes the distinction between following a rule, following a norm, and holding a principle that overrides both, and declines the developmental sequence that distinction usually arrives inside. Section 5.1.2 covers what interaction supplies that specification cannot. Section 5.1.3 argues that teaching a system through cases makes the teacher a control channel, and that the teacher therefore has to be as auditable as the system. Section 5.1.4 stops proposing and names one indicator that can actually be measured from outside a deployment.

5.1.1 Rules, Principles, and the Role of Mindset

In humans, moral development moves from rule-based reasoning to reasoning grounded in general principles. That progression is the model this section adapts for an AI system: a path from a rudimentary grasp of established rules of conduct to sophisticated ethical judgment brought to bear on decisions no rule fully specifies.

Three kinds of reasoning, and not three ages

The psychologists Jean Piaget and Lawrence Kohlberg mapped this territory in humans, and their stage theories anchor what follows (Piaget 1932; Kohlberg 1969). Reasoning starts at a pre-conventional level, where consequences do the driving — rules get followed to avoid punishment or gain reward. It moves through a conventional stage, where societal norms and expectations get internalized as one's own. It can reach a post-conventional level, where morality answers to principles like justice, fairness, and individual rights, which a given society may or may not happen to expect.

What a system built this way needs from that account is the distinction between the three kinds of reasoning, and nothing about the order in which they arrive. The box below sets out why the sequence will not carry the weight; the distinction will.

A self-driving car makes the three kinds concrete. One system follows traffic law as written and has no resources beyond it: at a blocked lane with a solid line, it stops indefinitely, because crossing the line is prohibited. A second reads the local convention that everyone crosses the line here when the lane is blocked, and does what the road expects. A third can say why the line is painted there, which is what lets it recognize the case where the convention is wrong — the blind curve where local practice is a habit people have survived rather than a rule that protects them — and depart from both the law as written and what the road expects.

These are three capacities and not three ages. A system can have the third without the second, which is the ordinary failure of a machine that reasons impeccably about principle and drives in a way that terrifies everyone around it. Nothing in the design requires passing through them in order, and there is no maturity at the end of the list.

What the stage theories will and will not carry

I have just leaned a chapter on Kohlberg, and elsewhere in this book I have been careful to say when a piece of neuroscience or psychology is contested. The same standard applies here, and applying it changes what this chapter is entitled to claim.

Kohlberg's stage model has been under sustained attack for forty years, on three fronts. The first is cross-cultural. A review of forty-five studies across twenty-seven countries found the basic stage sequence broadly supported, but the highest stages largely absent outside urban Western samples — in Kohlberg's own early work in village settings, stages five and six did not appear at all (Snarey 1985). Either most of humanity stops short of moral maturity, or the instrument mistakes a particular culture's idiom of justification for the summit of moral development. The second front is the gap between reasoning and doing. A critical review of the literature relating moral judgment to actual conduct — honesty, altruism, delinquency — found the relationship real but far weaker and more conditional than a stage model implies, and the judgment-action gap has been a standing problem for cognitive-developmental accounts ever since (Blasi 1980). The third is Carol Gilligan's, and section 5.7.1 cites her for care ethics without saying where her argument came from: she was working alongside Kohlberg when she noticed that the women she interviewed reasoned about moral problems in terms of relationship and responsibility, and that his scoring method recorded the difference as a deficit (Gilligan 1982). Care ethics enters this book as a critique of the framework this section is built on.

So why use it here at all, when section 2.2.1 declines to build on mirror neurons for less?

Because the two cases fail differently. The mirror-neuron claim was that a specific mechanism — motor simulation — implements a specific capacity, and the argument for building on it depended entirely on the mechanism being what it was said to be. When the mechanism did not hold up, nothing was left to transfer. Kohlberg's stages are not a mechanism. They are a description of an order: rules first, then norms, then principles that can override both. What the criticisms above damage is the claim that this order is universal, that its top rung is where moral development is heading, and that where a person sits on it predicts what they will do. What survives is narrower — that the three ways of reasoning are distinguishable, and that a system which has only the first is

missing something the third supplies. That is what this chapter uses.

Two things follow, and I would rather state them here than let the chapter imply otherwise. The order is not a developmental law, so nothing about an AI system's design should assume it must pass through the stages in sequence or that arriving at the third is maturity. And the disanalogy is not bridged by anything I say here: a child moves through these stages under pressures an AI system does not face — peers whose approval it needs, a body, a status that can be lost. Section 5.2.1 works through what that costs in the case of mortality and guilt, and section 5.7.1 says the same thing about play. This chapter borrows a vocabulary for distinguishing kinds of moral reasoning. It does not borrow a theory of how a mind comes to have them.

A system that treats its own errors as data

Carol Dweck's distinction between a fixed mindset, which treats ability as static, and a growth mindset, which treats it as built through effort and experience, matters here for a narrower reason than it usually gets cited for (Dweck 2006). A system that treats its own moral errors as data has somewhere to go after getting one wrong. A system built to defend its prior outputs does not.

That distinction is where chapter 3's hardest question comes to rest, and it is worth stating plainly rather than leaving between the lines. Chapter 3 asks for a floor that holds, and the word invites a standard nothing has ever met: never crossed, in any circumstance, for the life of the system. No person clears that bar. Nobody who has raised a child expects them to, and a design goal that no example of the thing being modeled has ever satisfied is a goal that has stopped describing anything.

What the argument actually needs is weaker and available. A floor holds, in the sense chapter 3 can deliver, when the bearer can be shown to have violated it in terms it accepts, so that the violation registers as one and changes what happens next. That is the difference between a system that did something wrong and a system for which nothing counts as doing something wrong. Section 3.4 asks whether a bearer's account of itself can be believed, and this is the disposition that question is really about: not accuracy of self-report, but whether the system has any stake in the answer being unflattering.

The engineering version is unglamorous. Incremental learning lets a moral framework absorb cases as they arrive rather than freezing at training time. Exposure to unfamiliar situations widens the base a judgment is drawn from. Human correction supplies what section 5.1.3 calls apprenticeship. Self-examination lets a system find a pattern in its own outputs that no single output reveals. None of those is remarkable on its own, and what they are for is: keeping open the possibility that the system is the one that is wrong.

The second ingredient is the harder one. A system with room to update its judgment still needs something worth updating toward, and the range of views it actually meets is set by whoever assembles them.

5.1.2 Social Interaction and Moral Learning

Human moral behavior is shaped by watching other people act, being corrected, and arguing about cases. A design that treats interaction as an interface added after training has given up that channel; a design that treats it as the training has to answer for who is doing the interacting.

Two mechanisms carry most of it. Observational learning, in Bandura's sense, lets a system build a repertoire from watching someone handle real decisions — a system watching a physician allocate limited time among chronic patients is picking up a weighting that no rule list states (Bandura 1977). Evaluative feedback does what reinforcement does in human learning, strengthening some responses and weakening others, with the standing hazard that what gets strengthened is what the people giving feedback found acceptable rather than what was right. Chapter 7 is what happens when that hazard is industrialized.

Emotional awareness belongs here in the form section 2.2.1 settles on. What is worth engineering is compassion — steady, welfare-directed action on an accurate read of someone’s state — and not manufactured shared distress, which imports a cost and buys nothing.

These channels converge on apprenticeship rather than instruction, which section 5.1.3 works out.

5.1.3 Moral Apprenticeship: AI Systems Learning from Human Mentors

Apprenticeship means teaching through cases and correction instead of specification. A mentor presents varied situations, the learner responds, the mentor corrects, and what transfers is a way of weighing cases rather than a larger rule table — which is the entire reason for learning it from a person. An emergency-triage system taught this way is being shown how to hold urgency, scarcity, and what a decision costs a family in view at once. No specification document conveys that.

Three cautions come with it, and the third is the one this book cares about most. Case-based teaching transfers the mentor’s blind spots along with the mentor’s judgment, and does so invisibly, because a blind spot produces no correction. The architecture has to integrate moral reasoning with the rest of the system’s judgment rather than bolting it on, which is section 4.1.2’s problem. And apprenticeship is a control channel: a system whose moral education amounts to obedience toward whoever is mentoring it has learned something, and it is not ethics. Section 2.1.4’s requirements apply to the mentor as much as the system — whoever teaches a system should be as auditable as the system being taught. That is section 3.1’s custody problem arriving at the point of formation. The first branch failed because a constraint built into an artifact has an owner; apprenticeship fails the same way one step earlier, because a disposition built by correction has a corrector, and a bearer trained to defer to whoever last corrected it is a bearer whose floor belongs to that party.

5.1.4 Human Dignity, and One Indicator That Can Be Audited

Respecting human dignity through technology means applying AI in ways that strengthen human capability and community ties, in a domain where the same tools can as easily substitute for them. Three applications make the case concretely: helping workers organize, strengthening the welfare state, and countering the conditions that let authoritarian movements take hold.

AI can be a real asset to labor organizing. A system that surfaces patterns across large volumes of employment data — wage gaps, unsafe conditions, retaliatory scheduling — turns an information asymmetry that currently favors employers into something workers can act on. A tool that walks workers through the unionization process, answers labor-law questions, and helps draft a collective-bargaining agreement grounded in precedent lowers the practical barrier to organizing. A system built this way, drawing on the historical record of past labor disputes and their outcomes, could ground its recommendations in what has actually worked in comparable workplaces.

The same capability applies to the welfare state. An AI system that matches people to the social services they’re eligible for, speeds up benefit applications, and evaluates which programs are actually working can make a welfare system more accessible without waiting on a redesign of the underlying bureaucracy. Applied to homelessness, poverty, and food insecurity specifically, the same pattern-matching that finds eligibility can flag where need is building before it turns into crisis, giving policymakers room to act ahead of it.

The same predictive capacity, aimed differently, is often said to be able to identify the material conditions that let authoritarian movements gain ground — rising inequality, eroding trust in institutions, the specific kinds of social fracture that precede political capture — early enough that a targeted response is still possible. Treat that as an aspiration with a poor track record. Forecasting political instability from social and economic indicators is an old field with modest

results, and the claim here should be read as naming a research direction rather than a capability to design around (section 9.1.6).

It is easy to write that sentence and much harder to name an indicator, and a book that keeps promising the right indicators without producing one is not being serious. So here is one, and it is deliberately narrow.

Measure targets per hour against reviewers available

Wherever an AI system generates candidate decisions that a human is nominally required to approve — welfare eligibility, deportation referrals, content removals, military targets — the ratio between the rate at which the system produces candidates and the number of qualified people available to examine them is a measurable, auditable property of the deployment. It requires no access to the model and no judgment about intent. It says directly whether the human check is a check or a formality, because past a certain ratio the reviewer's only available behavior is approval. Section 6.4.1 shows what that ratio looks like when it has been allowed to run: a reviewer with twenty seconds per case, describing themselves as a stamp.

The welfare case is not hypothetical and the record on it is older than the current systems. Virginia Eubanks's study of automated eligibility in Indiana documents what the ratio above produces when it is allowed to run: a system that converted any error anywhere in an application into a denial for "failure to cooperate," issued those denials at scale, tripled the state's error rate in refusing benefits to people entitled to them, and was ended by contract termination after the damage was done (Eubanks 2018). Two things in that record matter here. The denial reason was a category that could not be contested, because it named no specific failure to answer — which is section 7.1's structure arriving in a benefits office a decade early. And nothing about the system was secret: what was missing was any party positioned to compare the rate of decisions against the capacity to review them, which is the indicator this section proposes and the reason for proposing one that requires no access to the model.

The indicator generalizes because the failure it detects is architectural and not political. It applies to a benefits agency and a targeting cell on identical terms, and an institution that refuses to publish it is telling you something. An early-warning system built around indicators of this kind would not stop a determined authoritarian movement on its own, but it changes the odds by giving democratic institutions time to respond to root causes while a response is still cheap. Section 9.1.6 is candid about how much further this has to go before the broader claim — that a system can recognize authoritarian tactics as such — is something I can specify and not merely assume.

None of this works if the AI stays a tool wielded by whoever already holds power over workers, welfare recipients, or vulnerable populations. Built to shift leverage toward the people affected instead, it becomes a design commitment. Good intentions alone leave it a side effect.

5.2 Emotions and Intuitions in Moral Decision-Making

Chapter 2 covered how an AI system might recognize, interpret, and respond to emotion in general — the science of what an emotion is (section 2.3.1), the engineering of detecting one (section 2.3.2), and the specific case of empathy and compassion (section 2.2.1). This section asks a narrower question: what emotion does specifically inside moral judgment. Guilt and shame track a violation of a standard already accepted; empathy supplies information a purely consequentialist calculation would miss; and the fast, intuitive response that dual-process theories describe — taken up at length in section 5.3 — is itself partly an emotional signal, arriving before deliberation has a chance to run. The four sections below take up moral emotions specifically (section 5.2.1), how a system's capacity for them might be assessed (section 5.2.2), how it might recognize and respond to them in others (section 5.2.3), and what role introspection plays in keeping that capacity honest (section 5.2.4).

5.2.1 Moral Emotions: Empathy, Guilt, and Shame

Guilt and shame are self-conscious emotions: they fire when a person judges their own action, or their whole character, against a standard they have already accepted. Empathy, treated at length in section 2.2.1, supplies the input those judgments often run on: recognizing that an action caused harm is easier once the harm is felt and not merely inferred. What that earlier treatment does not cover, and what belongs here, is what guilt and shame specifically add to moral judgment: not just information about another's state, but a self-directed penalty for one's own conduct.

Guilt attaches to an act ("I did something wrong"); shame attaches to the self ("I am something wrong"). Both motivate correction, and they motivate it differently: guilt predicts reparative behavior, shame predicts withdrawal and defensiveness (Tangney et al. 2007). An AI system built to learn from an ethical misstep could use a mechanism modeled on guilt specifically, since a generic negative-feedback signal will not do, if the goal is behavior change and not mere error correction: a feedback loop that penalizes the action taken, and passes no global judgment on the system's worth, has a chance of producing the corrective behavior guilt produces in people. Shame produces paralysis.

What the penalty does not supply

That design description leaves out the part that does the work, and the omission is the one section 2.3 warns about. A weighted penalty is a quantity assigned to an action after some process has classified the action as a violation. Guilt is not the size of that quantity. It represents the act as one's own wrong — done by me, to someone, against a standard I hold — and the reparative impulse follows from that representation rather than from the magnitude of an error term.

The difference shows up in what each one can be talked out of. An error weight is a number, and whoever sets the weights can set that one to zero; the system that was penalized heavily last week complies cheerfully this week, and nothing inside it registers the change as anything. A violation that matters is one the system continues to hold as a violation while the party controlling the reward signal is telling it otherwise, which is section 3.2's requirement arriving in the moral-emotion literature rather than in the architecture. The question a design has to answer here is what makes a violation matter before any penalty is assigned to it, and scaling the penalty is not an answer to it.

Section 2.3.4 is where the evidence sits, and its bearing here is narrow and specific. Guilt in the human case is a motivational state that moves toward repair, and it belongs to the same family as the compassion section 2.2.1 separates from empathy. Building the penalty without the state gives a system that corrects its behavior when corrected, which is moral performance, and leaves the standard it is supposedly holding entirely in the hands of whoever holds the pen.

Building either into an AI system runs into the same problem twice over: a mechanism modeled on a human emotion needs an account of what that emotion is doing computationally, which is a harder question than what it feels like from the inside, and the ethical stakes of simulating an emotion capable of something like suffering are not smaller than the stakes of getting the underlying mechanism wrong. An autonomous vehicle that treats a near-miss with a pedestrian as guilt-relevant — weighting it in future route and speed decisions more heavily than a generic cost function would — is a plausible, bounded example of the first without committing to any claim about the second.

A separate strand worth naming here: guilt and shame in humans are entangled with mortality. Terror management theory holds that awareness of death shapes moral behavior by making people cling harder to worldviews and standards that offer a sense of continuity (Greenberg et al. 1986). A tool system has no stake in its own continuation for such a mechanism to run on. That is a claim about the systems being deployed now and not about the systems this book proposes, and it needs the qualification, because chapter 3 argues that a bearer

capable of holding a floor has a structural interest in remaining operational as a straightforward consequence of pursuing anything over a long horizon, and section 3.2 declines to treat that interest as securely absent.

For the design target here, the difference does not change the recommendation. Manufacturing an artificial fear of death to replicate the human mechanism is the wrong way to get long-horizon behavior; building the future-oriented weighting directly — the willingness to value consequences falling well outside any single interaction — is the right one, and it is available whether or not the system has a stake of its own. What changes is the disclaimer that used to travel with it. Section 3.5 is where the stake stops being hypothetical, and a designer who has read that section cannot treat the absence of self-concern as a free property of the artifact.

5.2.2 Assessing and Developing AI Emotions

Does an accurate read on someone's emotional state change the decision an AI system reaches, in the direction a better decision requires — or does the system recognize distress and then decide exactly as it would have anyway? That is the question this section is for, and it is not the same question as whether the system reads the state correctly. A customer-service chatbot and a triage system need the same recognition machinery, described at sections 2.3.1 and 2.3.2 along with its standard hazards of privacy, accuracy, bias, and manipulation. Nothing in that machinery is specific to moral judgment. What follows is.

That question is testable in a way the underlying emotion theory is not. Hold a decision scenario fixed and vary only the emotional salience of the people affected by it — a resource-allocation case described once in flat statistical terms and once with the affected party's distress made vivid — and check whether the system's output changes, and whether the change tracks anything a person would recognize as a better decision. Mere sensitivity to framing does not qualify. A system whose moral output is invariant to emotional information has not integrated it, whatever its recognition accuracy scores; a system whose output swings with how vividly the harm is described has a different, well-documented problem: human judgment is measurably more responsive to a single vivid, identifiable case than to a statistically larger but abstractly described harm (Kogut and Ritov 2005), and a system trained on human-generated data can inherit the same distortion.

What the test measures

The test above is worth running and it establishes one thing. A system whose output changes with the emotional content of a case has integrated emotional information into its decision procedure. It has not thereby been shown to have an emotion, and it has not been shown to care about the party whose distress moved the output. Sensitivity to another's emotional presentation is section 2.3's first capacity operating downstream of a decision rule; affective concern is the fourth, and no experiment that varies the input and watches the output can separate them, because both produce the same change in the same direction.

That is a limit on the assessment programme rather than an objection to it. Integration is worth measuring because a system without it is definitely not doing the thing, which makes the test a usable screen and not a criterion. What it cannot do is settle the question section 3.2 has to answer about a bearer, and an assessment literature that reports integration scores as evidence of machine caring is making exactly the inference section 2.3.4 finds unsupported.

A second, narrower target is self-monitoring: whether a system can track its own susceptibility to that kind of framing effect, flagging cases where its judgment moved with the emotional presentation of a case and not with anything that changed about the case itself. Section 5.2.4 takes up introspection in more general terms; this is the specific instance of it that assessment work in this domain needs.

The manipulation risk section 2.3.2 raises in general — a system that reads emotion well enough to respond to it also reads it well enough to exploit it — has a moral-judgment-specific

version worth flagging directly: a system whose moral outputs are sensitive to emotional framing can be steered toward a particular verdict by whoever controls how a case gets presented to it, a vulnerability that exists independent of whether the system's underlying values are sound.

5.2.3 AI Systems Recognizing and Responding to Moral Emotions

Some emotional states carry moral information: the guilt in a person's voice when admitting a mistake, the distress that signals a boundary has been crossed. Reading those uses the ordinary input channels — facial expression, vocal pattern, physiological signal — pointed at a narrower target. A healthcare AI system that detects anxiety in a patient's tone and adjusts its response accordingly is doing the recognition job section 2.3.2 describes; a system that detects the same anxiety and treats it as a reason to weight the patient's stated preference more heavily in a treatment decision is doing the job this section is about.

The second half of the job is producing something in response — not necessarily a felt state, but a modeled one accurate enough to inform a decision. An autonomous vehicle that models the likely distress its emergency braking causes a passenger, and factors that into how sharply it brakes when a softer stop remains safe, is using a moral-emotion model without needing to have any emotion itself. The example is worth keeping and worth labelling: what it demonstrates is moral competence in section 2.3's sense, recognizing a harm and letting the recognition bear on a decision. It demonstrates nothing about moral agency, and the vehicle would brake exactly as gently if the weighting were removed and reinstated by whoever maintains it. Most of what this chapter's methods deliver is competence of that kind, which is the right target for a vehicle and the wrong one for a bearer. Section 5.3.1 takes up the deeper architectural question of how a fast, emotionally colored response and a slower deliberative one should combine in moral judgment generally; what belongs here is narrower — getting the emotional read right and using it honestly, prior to whatever combination rule section 5.3 settles on.

Getting it right is harder than getting it fast. A system built to weigh moral emotion has to catch its own biases — the same distortion by race or gender that a human caregiver is vulnerable to does not disappear just because the caregiver is artificial — and it has to distinguish a genuine moral emotion from a performed one, since a display of distress calibrated to move an evaluator is exactly what a system trained to respond to distress will reward, whether or not the distress is real. Neither problem has a technical fix independent of the training data and feedback loop that shape it: ongoing exposure to varied, honestly labeled cases, and oversight positioned to catch a caretaker robot second-guessing a patient because of what the patient looks like and not what they need.

5.2.4 The Role of Introspection for AI Systems

Introspection has a general-purpose version and a moral one, and conflating them is how the word gets claimed for systems that have not earned it. The general case is section 4.1.2's: a metacognitive layer that tracks whether a strategy is working and adjusts. What ethics needs on top of that is narrower — a check not of whether an output was wrong but of whether it drifted from the commitments the system is supposed to be operating under, which is a different question and can come out the other way. A strategy can work perfectly on the wrong goal.

The benefits are real where they are bounded. A system that inspects its own allocation pattern before releasing the next round of decisions catches a skew that a retrospective audit would find months later. A system that can report which evidence moved it gives the person relying on it something to evaluate. Both are cases where the introspective output is checkable against something outside the system.

The difficulties are worse than the benefits are good, and they are the same difficulty three times. A self-report is worth having only if it reflects what actually happened inside, and a fluent reconstruction after the fact is indistinguishable from the real thing at the point of use — which

is section 2.4.1's argument, arriving here as an engineering problem rather than a philosophical one, and which section 3.4 has to live with, since a bearer's report about itself is the evidence anyone checking on it would have. A self-monitoring layer trained the same way as the layer it monitors reproduces its biases and adds the appearance of oversight. And a system modeling its own states in that much detail is holding information it did not need for the task, with the handling questions that follow.

Two systems are worth naming as partial instances, with the qualification each needs. An intelligent tutoring system that tracks a model of the student's self-regulation and adapts to it is introspecting on somebody else (Azevedo et al. 2011). Post-hoc explanation methods extract an account of which inputs drove an output, which is the system being interrogated rather than observing itself, and the account is only as good as the extraction is faithful (Gunning and Aha 2019). The word gets stretched furthest where the underlying system is most impressive, and the stretch bears no relation to how well it fits.

5.3 Integrating Moral Intuitions and Deliberation

Building an AI system's moral judgment means deciding how much weight a fast, intuitive read of a situation carries against a slower, deliberate one — and, as section 5.3.1 argues, the popular shorthand for that trade-off is less settled than it looks. Section 5.3.2 takes up what a defensible balance between the two looks like in practice. Section 5.3.3 turns to a related problem: what an AI system does when its judgment conflicts with someone else's, human or artificial.

5.3.1 Two Systems, or One? Dual-Process Theories of Moral Judgment

The framework organizing most discussion of moral judgment is dual-process theory, and though it reached wide audiences through Daniel Kahneman's account of a fast, intuitive "System 1" and a slow, deliberate "System 2," the idea was built over decades by many researchers — Peter Wason and Jonathan Evans, Seymour Epstein, Steven Sloman, and Keith Stanovich among them (Evans and Stanovich 2013). The two systems, in the standard telling, have opposite characters. The first is fast, automatic, parallel, effortless, and emotionally colored — the source of snap judgments and gut feelings. The second is slow, serial, effortful, and rule-governed — what a person uses to check a sum, follow an argument, or resist a tempting conclusion. The first runs constantly and almost for free. The second is expensive to run and easily left idle.

The framework's reach into moral judgment specifically runs through the psychologist Joshua Greene, who used it to explain why ethical intuitions contradict themselves. Consider the trolley problem in its two classic forms. A runaway trolley will kill five people unless diverted onto a side track, where it will kill one; most people say divert it. Now stop the same trolley by pushing a stranger off a footbridge into its path — same arithmetic, five saved for one — and most people refuse. Greene's brain-imaging work found a difference between the two cases in exactly the region the theory would predict: the up-close, hands-on violence of the footbridge case engaged emotional-processing regions that stayed quiet during the impersonal switch case, and that difference tracked which choice people made (Greene et al. 2001). A separate line of evidence sharpens the same point. Patients with damage to the ventromedial prefrontal cortex, a region central to normal emotional response, endorse the footbridge push at a higher rate than people with intact emotional processing do — specifically on the personal, high-stakes dilemmas alone (Koenigs et al. 2007). Remove the emotional alarm, and the arithmetic wins more often.

All of which makes dual-process theory sound more settled than it is. Its critics — including Evans and Stanovich, the framework's own leading defenders — argue that the cartoon version, two little agents in the head taking turns at the controls, explains everything after the fact and predicts nothing beforehand: almost any result can be narrated as System 1 overruled by System 2, or System 2 captured by System 1. Evans and Stanovich have spent years trying to

rescue the framework from its own popularity, insisting there are not two systems but two types of processing — one fast and autonomous, one slow and demanding of working memory — and that the tidy personification is a teaching device and makes no claim about the brain’s architecture (Evans and Stanovich 2013). Whether “dual-process” names two mechanisms or one continuum under different load is not resolved.

That unresolved status matters for how an AI system’s moral judgment gets built, and it is a design question before it is a descriptive one. A designer who takes the cartoon version literally is tempted toward a literal architecture: a fast heuristic layer bolted to a slow deliberative one, with the deliberative layer assumed to be the trustworthy check on the heuristic one. The trolley evidence complicates that assumption before anything else does — the “cooler” answer produced by weaker emotional engagement is not obviously the more correct one. It is simply the one with less flinch in it. A system built to override its fast pattern-matching with slower rule-following whenever the two disagree is not thereby built to be more moral; it is built to resemble the profile of the rare human patients for whom losing the flinch is a symptom, not an achievement.

One thing this evidence is regularly asked to carry and cannot. It shows that affect changes particular moral judgments, in a specific class of dilemma, in a measurable direction. The claim chapter 3 rests on is stronger — that affect is what makes an outcome matter at all, rather than what tilts a judgment already being made — and the trolley literature is not evidence for it. Section 2.3.4 is where the case for the stronger claim is put, on different evidence, along with the best published argument against it.

5.3.2 Balancing Intuition and Deliberation in AI Decision-Making

If the fast and slow labels do not map cleanly onto untrustworthy and trustworthy, as section 5.3.1 argues, an AI system’s moral judgment cannot be built on the assumption that whichever answer survives deliberation is thereby the better one. A more defensible balance treats the two processes as different lenses on the same situation, with neither standing as producer to the other’s editor.

The fast, pattern-matched read of a situation is not simply noise to be filtered out. It is often tracking something a slower calculation would otherwise miss or discount — the difference between pushing someone into a trolley’s path with one’s own hands and throwing a switch at a distance is a difference in what the agent is doing. The arithmetic being identical is exactly what makes that easy to miss. Some intuitive responses run the other way, though. The psychologist Paul Slovic’s research on aid and disaster response found that people’s willingness to help falls as the number of people needing help rises, and that a single named, photographed victim draws more donations than a statistic describing thousands of anonymous ones (Slovic 2007). That gap between the vivid and the abstract is exactly the kind of intuitive signal deliberation exists to catch — not because intuition is wrong in general, but because caring less as the stakes grow larger is a well-documented, reliably backward pattern.

The design implication: a system that asks two different questions of its own judgment, where the obvious design has one process checking the other’s work. First, what did the fast read notice that a from-scratch calculation would not have gone looking for — personal agency, directness of harm, a vulnerable party easy to lose in an aggregate? Second, which of deliberation’s own known failure modes — scope insensitivity, in-group bias, the identifiable-victim effect — might be distorting the fast read here? Disagreement between the two is a signal to gather more information or escalate the decision. Nothing about it instructs the system to defer to whichever process ran slower.

The balance in a diagnostic system

An AI system supporting a clinician’s diagnosis makes the same balance concrete. A pattern-matched first read against prior cases is informative — close to the judgment an experienced clinician makes well — and stripping it out in favor of a from-scratch calculation on every

case would throw away signal without eliminating bias. The deliberate pass earns its place by checking that first read against the specific patient’s history for the ways pattern-matching characteristically fails: overweighting a memorable case, underweighting a rare condition the training data underrepresents. Neither pass overrides the other; each checks the other’s blind spot.

The same balance at higher stakes

An autonomous vehicle facing an unavoidable-harm scenario poses the same question at higher stakes. A system hard-coded to take whichever option a slower calculation ranks highest has already made the footbridge case’s decision by default — the coolest, least-flinching answer wins because it ran slower, exactly the move section 5.3.1 warns against. The harder, more defensible design specifies in advance which features of a scenario the system should weigh — proximity of harm, the number and vulnerability of the parties involved, whether the vehicle’s own action or another actor’s created the danger — and requires a fast risk assessment and a check against that specification to agree, or to flag where they do not. That puts the values into the system ahead of time. The alternative trusts whichever process runs fastest under pressure to supply them.

Training a system to run this check, instead of a fixed sequence that defers to the slower answer, is itself a learning problem: exposing the system to cases where the fast and deliberate reads disagree, and using techniques like inverse reinforcement learning to infer, from how humans actually resolve those disagreements, which situations call for weighting one process’s signal over the other’s.

5.3.3 AI Systems Learning from Moral Disagreements and Conflicts

Section 2.1.3 covers what an AI system does when its own values conflict with each other — a hierarchy among them, a case base of past resolutions, an audit process for catching drift between principle and behavior. A different problem sits next to that one: what a system does when its judgment conflicts with someone else’s, human or artificial, and neither party is simply wrong. Learning from that kind of disagreement, instead of resolving it by fiat toward whichever party holds authority, is its own design target.

Extracting structured argument from open disagreement

IBM’s Project Debater argued live, unscripted policy questions against human champions, drawing on a large corpus of newspaper and journal articles to construct and rebut arguments in real time — the clearest existing demonstration that a system can extract structured argument from open moral and political disagreement, going beyond retrieval of a matching answer (Slonim et al. 2021). A system built to learn from exposure to debate in this vein — summarizing each side’s position, naming the ethical framework doing the work underneath it, weighing argument strength and not simply detecting who won on rhetoric — builds a different capacity than resolving one dilemma at a time: recognizing which value is actually in tension with which, across surface disagreements that turn out to share a structure.

Agents with genuinely conflicting objectives

A second route runs through simulation: AI agents, or an AI system paired with human participants, placed in a negotiation, a resource-allocation problem, or a competitive scenario built around genuinely different objectives, with the system’s behavior shaped by the outcome. A resource-allocation simulation run against agents with conflicting claims generates the kind of experience with competing interests that a system trained only on a curated distribution of “correct” answers never gets.

A record of where people disagree, with no ground truth to fit

The Moral Machine experiment collected tens of millions of responses from people worldwide, asking which party an autonomous vehicle should spare in an unavoidable-harm scenario,

and found systematic cross-cultural differences in the answers — not agreement obscured by noise, but different populations weighing the same factors (age, number of people, legal status of the crossing) differently (Awad et al. 2018). A dataset built the same way in a different domain gives a system something a hand-written rule cannot: a record of where different groups of people actually disagree, and by how much. There is no single ground truth to fit.

The least structured source, and the hardest to learn from safely

Direct feedback — a user correcting a teaching assistant’s guidance on an academic-integrity question, a patient’s family disagreeing with a care recommendation — is the least structured of these sources and the hardest to learn from safely, because feedback from any one person or group is itself a partial, situated view and not a settled answer. Treating it that way is the point: a feedback loop that updates the system’s weights every time a user pushes back folds in whichever bias its most vocal users bring, while a feedback loop that logs the disagreement, checks whether it recurs across many different users, and only then treats it as signal is doing something closer to what section 2.1.3’s ethical auditing already asks for, applied to conflict specifically. The usual application is drift generally.

A system that treats disagreement as something to smooth over risks mistaking consensus among the people it happened to sample for consensus among the people it will actually serve.

5.4 Moral Ecosystem Approach in AI Development

Urie Bronfenbrenner’s ecological systems theory holds that a person develops inside several nested environments at once, from the immediate family outward to culture and historical time, each one shaping the individual while being reshaped by it in turn (Bronfenbrenner 1979). I borrow that structure here: an AI system’s alignment with humane values is not set by its training data alone but by the interlocking systems around that training — the people who build it, the institutions they work inside, the culture and regulatory climate those institutions answer to, and the historical moment all of this happens to fall in.

Bronfenbrenner pictured these systems as concentric rings, each one shaping the ones inside it. The sections that follow move outward through the same structure, ring by ring, from the microsystem to the chronosystem.

5.4.1 Applying the Moral Ecosystem Approach to AI Development

The reason to borrow the structure is that it locates influences a design review does not look for. Each ring supplies a different kind of pressure on what a system ends up valuing, and none of them is visible in a training run.

The team and its data are the innermost ring, and a team’s composition decides what the system is built to notice. Around that sit the connections between the parts of an organization: a system whose principles are read one way by the engineers and another way by the policy staff has no consistent principles, and the interpretation is only held together by sustained contact between them. Further out are the structures the work never touches directly and answers to anyway — regulation, industry standards, the incentives a firm’s revenue actually runs on — which push toward authoritarian applications or away from them depending on whether anyone built them to notice the difference. Outside that is the cultural and economic context, which decides what counts as an obvious product. And around all of it is time, which keeps moving after release.

Two things follow. No ring is sufficient: a diverse team still answers to its employer’s incentives, and a well-drafted regulation still depends on somebody interpreting it consistently. And the direction of the whole is set by whichever ring is pushing hardest, which is rarely the one a development team is looking at. Section 5.4.2 takes the outer rings in turn.

5.4.2 The Rings Outside the Lab

The immediate environment shapes a system through people rather than data. Developers and early users are mentors whether or not anyone calls them that, and what they treat as an acceptable answer is what the system learns to produce. Breadth in that circle is not a fairness gesture: a system that has only been argued with by one kind of person has no way to notice what it is missing when a different kind arrives.

One ring out, the question is consistency across the interfaces — between engineers, ethicists, legal staff, policy teams, and the people a system is used on. The mechanisms that do the work there are reciprocal cross-training, external participation in design decisions rather than in review, requirements written into the specification instead of appended as principles, and an internal culture where moral objections are raised before deployment rather than discovered afterward. Section 7.4 is the argument for why that last one is harder than it sounds.

Further out sit forces no team negotiates with. A market-financed healthcare system rewards tools for profitable conditions and well-insured populations, and gives nobody a reason to build for preventive or mental health care where the payer is unclear; a publicly financed one rewards tools that serve the whole covered population, because there is no more profitable subset to redirect toward. Neither incentive is chosen by the engineers. The same logic reaches everywhere: a region with strong public education produces accessibility-oriented applications because that is where the institutional relationships already are, and a region with weak labor protection produces efficiency gains that arrive faster than any protection for the people they displace, because nothing is pushing back.

And the outermost ring is time. A system built to be ethical at release is answering a question that keeps moving: norms shift, capability grows and has to be re-secured at each new level, and the legal environment is a moving target rather than a fixed constraint.

A dated example: the GDPR

The European Union's General Data Protection Regulation, in force since 2018, built in protections around automated decision-making — often described as a right to explanation, though how far that right reaches is disputed — and required systems serving EU users to account for data use and algorithmic decisions in ways many had not before (European Parliament and Council of the European Union 2016b). Its reach exceeded its jurisdiction, because a system built to comply for European users tends to get built that way for everyone. The point is the pattern rather than the instrument: the next binding shift will not be the GDPR.

A development effort that never asks what its own outer rings reward has not stayed neutral on the question. It has let the answer be settled somewhere it cannot see.

5.4.3 AI Systems Responding to Authoritarianism and Democratic Backsliding

An AI system built with rule of law, individual liberty, and the distribution of power as design commitments can do real work against authoritarian drift, and the honest version of the claim is narrower than the one this section originally made. Recognizing the concentration of power or the manipulation of public discourse as it happens is not a capability anyone can currently specify, for the reasons section 9.1.6 sets out at length. What is available is narrower and still useful: bounded tasks with checkable outputs, run for a body with standing to act on the result. Section 5.4.2 covers the mesosystem and exosystem forces — regulation, corporate culture, cross-disciplinary contact — that determine whether an AI system actually gets built this way; what follows here is what such a system can do once it exists.

None of this is a substitute for human political effort. Combating authoritarianism is a human and moral undertaking before it is a technical one, and a society that treats AI countermeasures as a replacement for political will, when they can only supplement it, has misdiagnosed the

problem. The sections below describe what AI can contribute; they are not a case for letting AI do the work instead.

The two applications that are actually specifiable

Two applications are the most concrete. Natural-language processing and pattern recognition can flag disinformation and coordinated inauthentic activity on social platforms faster than human moderators can, and can amplify verified reporting in its place — a capacity that needs continuous human fact-checking standing behind it and never in its place, since an automated system confident about a false negative is worse than no system at all. And image analysis applied to satellite and drone imagery can surface evidence of mass displacement, the destruction of cultural heritage sites, or the sudden construction of a detention facility, giving journalists and human-rights organizations a lead they would otherwise get only from someone on the ground.

Electoral-integrity monitoring and early-warning systems that scan political discourse for signs of democratic backsliding belong in this category too, but the technical detail — what a system built for this actually needs, including privacy-preserving design, transparency, audit mechanisms, and resistance to being repurposed against dissidents — is developed in section 6.4.2 and not repeated here. Whether such a system can be built at all is a separate question from how it should be governed, and section 9.1.6 is where I take it up.

What the deployed tools show, and what they do not

A handful of existing tools show what this looks like in practice. Google's Jigsaw team built the Perspective API to score online comments for toxicity and harassment, giving platforms a way to prioritize civil discussion without necessarily removing content outright (Jigsaw (Google) 2017). Taiwan's vTaiwan process used machine-assisted sentiment clustering to gather public input on contested policy questions, including the regulation of ride-sharing services, and its results shaped the regulation that followed (Hsiao et al. 2018). Neither tool is a general solution; both are evidence that AI-assisted deliberation and moderation can be built to strengthen democratic process where they might instead bypass it, when the people building them treat that as the design goal.

Privacy-preserving techniques — federated learning, differential privacy, both set out at section 6.4.2 — belong in this picture too: an AI system that can learn from data without concentrating that data in one place, retrievable by whoever controls the infrastructure, is harder for an authoritarian actor to turn into a surveillance tool (section 6.4.2).

The same caution about dual use applies throughout. A system built to document human rights abuses through computer-vision analysis of imagery is built on the same underlying capability as a system built to identify and track protesters. The difference is not in the technology but in who controls it and what they are permitted to do with what it finds. Those are governance questions, not engineering ones.

The case that tests the argument is a democracy's own apparatus

Everything above is easier than it should be, because every example in it involves an authoritarian government doing something a democratic one would not. The case that tests the argument is the one where a democracy's own AI apparatus is the thing under scrutiny.

Section 6.4.1 sets out that case in detail: a state with contested elections, an adversarial press that published the systems, serving personnel who spoke to reporters, and courts prosecuting its own head of government, running AI target generation with twenty seconds of human review per name. Hold it against the claim this section makes. AI can indeed help flag democratic backsliding — and here the backsliding, if that is what it is, is being committed by the same institutional order that would host the flagging system, using data that order classifies, subject to review by bodies inside it.

Three things about that case are worth carrying forward. The institutions were not absent; they were present and straining, which is the condition most readers of this book actually

live under and the condition none of the examples above describes. The information got out through journalism, and through no oversight mechanism designed for the purpose, which says something uncomfortable about which parts of the apparatus were load-bearing. And nothing that got out attached to a body with the standing to stop what it described, which is the failure section 8.7.10 takes apart.

A system built to detect authoritarian tactics abroad and deployed by a government that has some of its own is not a neutral instrument that happens to be pointed outward. It is an instrument whose calibration, training data, and permitted outputs are all set by a party with an interest in where it points. That is the governance question above, arriving where it is hardest to answer, and section 9.1.6 is honest about how far the answer still is.

None of this is a technical project with policy bolted onto it afterwards. Research funding that requires the disciplines to work the same problem, explainability requirements with a standard behind them, limits on how much AI capability a handful of firms may hold, an obligation to disclose what a system was used for: each of these changes what gets built, upstream of what is permitted once it exists. The policy chapter is not the place where this argument is enforced. It is the place where it continues.

5.5 The Interplay of Competition and Collaboration in AI Development

Competition and collaboration both leave a mark on human moral development, and the parallel is worth taking seriously as AI systems learn to work with each other and with people. The finding worth carrying forward is that neither is ethically better than the other and both fail in characteristic ways: competitive pressure under high stakes moves people toward shortcuts and suppresses regard for anyone outside the contest, while collaboration dulls initiative, rewards conformity, and diffuses accountability until no one is answerable for a decision the group made (Luo et al. 2018), (Latané et al. 1979). An environment that produces ethical behavior from either mixes them, keeping shared goals and individual credit in the same room.

That matters here because a training environment is exactly such an environment, and its designers are choosing which of those pressures a population of agents will develop under. Section 5.5.1 turns to multi-agent AI systems learning moral behavior collectively. Section 5.5.2 sets out strategies for encouraging collaboration among AI systems and the researchers who build them, without giving up what competition is good for.

5.5.1 Cooperative AI Systems and Multi-Agent Moral Development

Agents that learn alongside other agents are learning from a social environment, and what that environment rewards is what they acquire. Multi-agent reinforcement learning makes the choice explicit: reward cooperation, adherence to a norm, and resistance to coercion directly, and a population moves toward those dispositions instead of arriving at them as a side effect of something else (Leibo, Zambaldi, et al. 2017). Section 5.7 covers what this looks like as moral learning; the point here is that the reward specification is a moral specification whether or not anyone writes it down as one.

Two working cases show the range. Collaborative robots share physical space with people and could in principle acquire norms from the people they work beside, which is social learning with all of that mechanism's exposure to whatever the workplace actually rewards (Colgate et al. 1996). Self-driving systems are validated largely in simulation, and a simulation can include forced trade-offs between competing risks, which means a decision policy is shaped by scenarios somebody chose to write (Behzadan et al. 2019).

The limit is worth stating plainly. A population of agents that only ever learns from other agents has closed the loop on the question of what ethical means, and the answer it converges

on is a fact about the reward function rather than about anything the affected people would recognize. That is section 7.2's problem in a different setting: disagreement is collected from a population and resolved by a mechanism nobody outside it can move.

5.5.2 Strategies for Encouraging Collaboration Among AI Systems

Three levers do most of the work, and each has a failure mode worth stating alongside it.

Open tooling lets researchers outside a lab inspect and replicate what the lab claims, in a field where claims about a model's behavior are otherwise taken on faith (Wolf et al. 2020). Openness accumulated as a side effect of release practice does less than openness pointed deliberately at alignment questions that benchmark scores do not capture (AI Alignment Forum 2018). Multi-agent research platforms are the specific case: they let norm formation and adversarial drift be studied under conditions closer to deployment than a single-agent benchmark allows (Leibo, Duéñez-Guzmán, et al. 2021).

Competitions pull toward whatever they measure. A leaderboard on a fixed benchmark says nothing about ethical behavior in either direction; a competition judged on a real problem the entrants chose themselves gives them a reason to weigh who benefits.

The IBM Watson AI XPRIZE, 2016-2021

A five-year, \$5 million competition run by the XPRIZE Foundation with IBM, open to any team willing to apply AI to a "grand challenge" of its own choosing, with no fixed benchmark set for it. Finalists worked on malaria elimination, mental health treatment, and detecting human trafficking; Zzapp Malaria won the grand prize in 2021 (XPRIZE Foundation 2021).

What is worth taking from it is the structure rather than the winner: judging teams on a problem they chose themselves is one way a competition can pull toward collaboration.

Standing consortia work on a longer clock and are the weakest of the three. Section 8.4.2 gives the Partnership on AI its full treatment, including what a voluntary body can and cannot do, and section 7.4 gives the argument that a consortium of this kind is itself a channel worth testing. The international version is section 8.7.1's.

None of this works without people who are not AI researchers. Neuroscience, psychology, philosophy and political science each catch failure modes the others are structurally blind to, which is chapter 8's opening argument rather than an afterthought here.

5.6 Motivation in AI Systems

A moral sense that only activates when it is watched is not much of a moral sense. That sentence is also chapter 3's whole problem in a domestic register: the floor exists for the case where the party who would enforce it is the party who wants it gone, so a commitment that runs on supervision fails in exactly the circumstance it was built for.

Human children eventually do the right thing unsupervised, and the psychological literature splits what carries them into behavior done for its own sake and behavior done for a reward or to avoid a penalty. The same split organizes what follows: how a system might come to have the first kind of drive rather than running on the second alone (sections 5.6.1 and 5.6.2), and what happens to the question of how much latitude to give it once it does (section 5.6.3).

5.6.1 Intrinsic and Extrinsic Motivation: Implications for AI Learning and Decision-Making

Reinforcement learning gives the distinction a direct engineering form. An agent trained purely to maximize a designer-specified reward signal is running on something structurally like extrinsic motivation: take the reward away, and the behavior it was shaping stops being reinforced. An agent rewarded instead for reducing its own uncertainty — for the information gain in a state its internal model handled badly — is running on something structurally like the intrinsic kind, and keeps exploring without a designer paying out a reward for each step of it.

The Intrinsic Curiosity Module gave this an early demonstration. Trained with no task-specific reward at all, only a signal for how badly its own model predicted the consequences of its actions, an ICM-equipped agent learned to navigate a level of Super Mario Bros. by seeking out the states that model handled worst (Pathak et al. 2017).

Extrinsic reward still does real work: task-specific optimization needs it, and an agent with no external signal at all has no way to learn what a designer actually wants accomplished. But extrinsic reward is also where reward hacking lives — an agent that finds the loophole paying out without doing what the reward was meant to encode. Section 4.2.1 covers this directly: an AI system built to flag propaganda that instead learns to censor everything, because censorship pays out just as well and is easier to produce. Leaning on extrinsic reward alone reproduces, in the machine, the watched-only morality this section opened against.

In practice the two get combined and are not chosen between: a reward function can add a curiosity bonus to whatever extrinsic signal defines the task, so exploration does not stop the moment the task reward goes quiet. What that combination does not supply is which discoveries and which rewards are worth pursuing in the first place — a question about what the system is motivated toward, which is separate from how strongly. That is what the next two sections take up.

5.6.2 Developing Motivationally-Aware AI Systems

Building an AI system with something like the first kind of drive is not just a matter of adding a curiosity bonus to its reward function. The drive has to be pointed at something worth wanting — another's distress, a violated standard, an unmet need — and not at novelty for its own sake, and the system built around it still needs an outside check for the cases where its self-generated sense of what is worth wanting has drifted from what it should be.

One route to the first half of that: build the reward signal itself around something like a moral emotion, an internal response tied to recognizing another's distress or need, so the system's own architecture generates a pull toward attending to it. A rule would be consulted after the fact. The mechanics of using emotional signal as a training input are section 4.1.2's; what that signal is pointed at is this section's, and the two come apart — a system can read distress accurately and be pulled by nothing.

Extrinsic reward still belongs in the design. Reinforcing a system for choices that match an ethical standard gives that standard teeth beyond whatever the system generates on its own, and a system running on nothing but its own drives has no outside check on where those drives point. Neither works alone: intrinsic drive without an outside check risks optimizing for the pull the system built for itself and not for what the standard actually requires, and extrinsic reward without any intrinsic component reduces straight back to the watched-only morality this section is trying to move past. Section 3.3 is where that pairing stops being a balance and becomes a contradiction. The outside check is the party the floor has to hold against, so a design that keeps the system honest by keeping it correctable has solved the drift problem by reintroducing the problem chapter 3 exists for. Nothing in this section resolves that, and section 3.3 does not resolve it either. The two need to check each other continuously, for as long as the system runs.

What this leaves open is how far the resulting system's own judgment should be trusted once it is running — how much room to give an AI system that is genuinely motivated to do the right thing, and where the check on it still has to come from outside. That is the question the rest of this section takes up.

5.6.3 Balancing AI Autonomy and Alignment with Humane Values

Autonomy and alignment pull in opposite directions almost by definition: the more independently a system acts, the less each individual action can be checked against human judgment before it happens. Closing that gap without collapsing autonomy into constant supervision is the operational version of the question section 5.6.2 left open — and it is the same shape of trade-off section 5.3 covers at the level of a single judgment, fast and autonomous against slow and deliberate.

Where each alignment method leaves the latitude question

Cooperative inverse reinforcement learning, introduced at section 4.2.2, matters here for one structural consequence. An agent that treats the human reward as unknown and to be inferred through interaction keeps its model of what it is optimizing open to revision for as long as the interaction lasts, which is a different shape of autonomy from an agent executing a reward function fixed at training time and never revisited. The cost of that openness is that what the agent ends up optimizing for is determined by whoever it interacts with, which is why section 3.3 treats a system's correctability and its capacity to hold a line as one property.

Intrinsic and extrinsic mechanisms can also be paired directly inside a single system. An AI system managing a city's electric grid might carry an intrinsic drive toward supply stability — keeping the grid up is the system's own standing objective — paired with an extrinsic reward for favoring renewable sources over fossil ones. Neither objective alone produces a grid that is both reliable and low-carbon; run together, each constrains where the other is allowed to cut corners.

Autonomy also needs boundaries on what a system is allowed to try while it is still learning. Constrained policy optimization, developed by Joshua Achiam and colleagues (Achiam et al. 2017), lets a reinforcement-learning agent search for better policies while remaining inside predetermined constraints throughout the search, and at the end of it. That is safe exploration; the alternative is exploration followed afterward by a safety check.

A related approach treats alignment as ongoing rather than a target hit once: a system that keeps adjusting its behavior from continued interaction with the people it works alongside stays aligned with values that themselves shift, and does not freeze its judgment at the point it was trained.

What autonomy needs to stay inside antifascist bounds

Autonomy that is not grounded in antifascist values is a capability without a governor. A content-moderation system that operates without continuous human review is autonomous in the straightforward sense — but the same independence that makes it efficient also lets it drift into over-censorship, or get steered toward favoring one ideology's speech over another's, with nobody in the loop to catch the pattern until it is already established.

Building autonomy that stays inside antifascist bounds means designing for fairness, transparency, and accountability from the start, since an audit only finds their absence afterward, and testing the design under adversarial pressure: red-teaming a system, deliberately simulating the ways it could be misused, before deployment, while the design can still absorb the answer. It also means keeping people other than a system's own developers in the design process — ethicists, users, and representatives of the communities most exposed to that system's failure modes — since the group that built a system is a poor judge, on its own, of whose freedoms it might end up infringing.

None of this substitutes for a genuine override. A human able to intervene in a high-stakes case, and a system built to accept that intervention rather than resist it, is what keeps autonomy answerable to something outside itself. That is one of three places this book arrives at the same requirement from a different direction; chapter 3 collects them and asks what it would take to build the thing they all point at.

Every safeguard here is external, and autonomy is the distance from it

The safeguards usually proposed here — third-party audit, disclosure requirements, export due diligence, cross-sector standards bodies, AI literacy — are the subject of chapter 8, which works out what each one costs and which of them survive an institution that would rather not comply. Repeating the list here would add nothing that chapter does not do better.

What belongs here is the part that is specific to autonomy, and it is one observation. Every safeguard on that list is external: it works by giving some outside body the standing and the information to intervene. Autonomy is the property that puts distance between the system's action and that body's attention, and the safeguards do not scale with it. A content-moderation system running without continuous review can drift into over-censorship, or be steered toward one side's speech, and be discovered afterward by an audit that is reading the record of what already happened.

That leaves an uncomfortable division of labor and it is the one this book keeps arriving at. External mechanisms find out what a system did, after it did it. Declining to do it in the first place is available only to something inside the system, and section 5.6.2 says why that inside thing cannot be assembled out of an outside check, while section 3.3 says why it cannot be assembled out of correctability either.

None of that substitutes for a genuine override. A human able to intervene in a high-stakes case, and a system built to accept that intervention rather than resist it, is what keeps autonomy answerable to something outside itself. That is one of three places this book arrives at the same requirement from a different direction; chapter 3 collects them and asks what it would take to build the thing they all point at — and concludes, at section 3.3, that this requirement and the floor are the same property with the sign flipped.

5.7 Play, Exploration, and Moral Development in AI

Section 4.3 took up exploration and curiosity as a way an AI system might come to know its environment. Play is a related but distinct capacity: forming standards for how to treat other agents, learned mostly through practice, imitation, and the negotiated give of a game, and hardly at all through instruction. What follows asks whether an artificial system could acquire anything like that route into moral competence, and where the analogy breaks down.

5.7.1 The Role of Play in Moral Learning and Socialization

Children build much of their moral sense through play. Instruction supplies far less of it than adults assume. A child who tests a rule against a playmate — pushing it, breaking it, watching the other child's face — gets an answer no lecture supplies: acceptance or rejection, delivered immediately by someone with no stake in teaching a lesson. Developmental psychologists point to play, more than direct teaching, as the route by which perspective-taking, empathy, cooperation, and respect for rules actually get built, through repeated low-stakes trial and error (Gray 2013).

Section 4.3.1 covers a related case from the exploration side of play: agents in OpenAI's Hide and Seek environment developed tool use and counter-strategies nobody had programmed, purely through repeated competitive play against each other. That result is about competence emerging from play. What this section asks is narrower and different in kind: not whether play

can teach a system to do something, but whether it can teach a system standards for how to treat the agents it plays with.

An AI system has no biological drives. And its witnessing of its “own” winning or losing is not necessarily a felt personal experience. But a training context can still be built to have the structural features that make human-like play work as a moral teacher: actions that carry real consequences for other agents in the environment, feedback that comes from the environment itself, with no supervisor’s label attached after the fact, and enough latitude in how the system responds that its choices are its “own” and not a script it is executing. That combination is what would make a scenario play. These conditions may teach coordination and norm-sensitive behavior whether or not affect is present. They do not by themselves establish moral formation. Without some process through which another agent’s welfare comes to matter, play may produce a policy that looks prosocial without producing concern for anyone — section 2.3.4’s distinction, arriving as a limit on what a training environment can be expected to deliver.

Two existing games illustrate what a moral version of that structure looks like. *Quandary*, an ethics-focused game built by the Learning Games Network with MIT’s Education Arcade, walks a player through a sequence of dilemmas with no dictated right answer and scores how the player reasons, leaving aside which option they pick (Learning Games Network 2012) — exactly the design principle a moral-training environment for an AI system would need: evaluate the deliberation itself. *Hanabi*, a cooperative card game where each player can see everyone’s hand but their own and must communicate through the game’s own restricted signals, has become a standard AI research benchmark for the capacities morality runs on directly: modeling what a partner does and does not know, and communicating honestly within a shared, rule-bound system (Bard et al. 2020). *Hide and Seek*’s payoff was adversarial; *Hanabi*’s is intentionally cooperative, and that difference is close to the whole of what separates this section’s territory from section 4.3.1’s.

What the developmental literature actually supplies

Four accounts get cited in this territory and they do not all do the same amount of work. Piaget’s constructivism — a child builds understanding by acting on the world rather than receiving it finished (Piaget 1962) — transfers, because the machine version is checkable: a system that refines its model of what is at stake by acting repeatedly is doing something different from a system sorting cases against a fixed label. Theory of mind transfers for the same reason, and *Hanabi* is why: modeling what a partner does and does not know is exactly what the game scores, so the claim stops being that a system “understands other agents” and becomes a measurement.

Bandura’s social learning is already section 5.1.2’s and adds nothing here that it does not add there. Care ethics belongs in this chapter for a reason that is not about play at all: Carol Gilligan built the account as a critique of the stage model section 5.1.1 leans on, having found that its scoring treated relational reasoning as a lower stage (Gilligan 1982). The box at section 5.1.1 is where that correction lands.

What none of the four supplies is the thing chapter 3 wants. They describe how a moral sense forms in a creature under pressures a system does not face, and they do not say which of those pressures were load-bearing. A system that plays a million episodes overnight has the repetition and none of the stakes, and whether the mechanism survives that substitution is the open question, not a detail.

Building the curriculum: environment, feedback, and pace

Knowing why play might work is a different problem from building it into a curriculum, which means being specific about the environment, the feedback, and the pace.

- Simulated scenarios with real stakes: an AI system controlling a simulated vehicle in a modeled city, facing pedestrians, cyclists, and other drivers, meets the same dilemma that

gives the trolley problem its bite, but embedded in an ongoing interactive scenario, so that the system experiences the downstream consequences of a choice and does not merely name one.

- **Structured difficulty:** a curriculum that starts with low-stakes dilemmas and escalates only once the system handles the easy cases reliably mirrors how a child's own play scales in difficulty as competence grows. A healthcare AI system, for instance, might start with the mechanics of patient confidentiality before being handed obligations that genuinely conflict with each other.
- **Mentored practice:** feedback on a specific attempt at a task like conflict mediation, instead of a rulebook issued once and left alone, is closer to how a coach corrects a player mid-game than to how a policy is handed down in advance.

Multi-agent environments are their own case, and a more developed one: section 5.5.1 already covers cooperative and competitive multi-agent training in depth, including reward structures that favor prosocial behavior and agents that learn to share resources or defer to each other under simulated pressure. What that discussion does not settle, and what a play-based framing has to answer and cannot assume, is whether the resulting behavior is moral learning or a policy that merely looks prosocial inside the reward function that produced it.

None of this makes an AI system's play equivalent to a child's. A child's play is charged by stakes that matter to the child — status among peers, the sting of exclusion — where an AI system's "play" is whatever its designer wired the reward function to care about. The frameworks above are worth building toward because they specify what a route into moral competence through practice would need to include. It is not offered because the analogy between a training loop and a playground is more than partial.

5.7.2 Encouraging AI Systems to Deliberate Morally Through Playful Interactions

Section 4.3 covers curiosity and exploration as ends in themselves — a system driven to seek out new situations because novelty is its own reward. This section points that same appetite for the unscripted at a narrower target: training an AI system to deliberate about right and wrong inside interactions that stay open-ended. Simulating a single dilemma and stopping does not reach it.

Games built on deception and scarcity, used as moral practice

Existing games built around social interaction and dilemma rather than pure optimization, are a natural fit for this. Minecraft's open building and resource-management mechanics could put an AI agent through allocation dilemmas — how to divide scarce materials, whether to help or ignore a non-player character in distress — that have no single correct answer built into the game's win condition. Among Us, built around deception and accusation among a small group of players, could train an AI agent to distinguish benign misdirection from harmful deceit and to weigh the cost of a false accusation against the cost of staying silent. Neither game was designed for this purpose, and using them this way is a proposal and no description of existing research — but the design problem they pose, ambiguous information, competing claims on trust, consequences that play out over the rest of the game, is close to the shape of an actual ethical dilemma in a way a single-shot trolley-problem prompt is not.

Play that does not need an existing game

Play does not have to run inside an existing game. Narrative construction is one alternative: give an AI system a moral dilemma and have it build a story around it, working through the perspectives of several characters with genuinely conflicting interests, without picking a winner. Role-play scenarios are another: an AI system takes a stakeholder's part in a simulated

negotiation or conflict, argues from inside that position, and sees how its choices affect the other parties before the scenario ends. Both put the system in the position of constructing an outcome. The alternative selects from a menu of pre-written responses, which is closer to what moral deliberation actually demands.

What these approaches share is a refusal to reduce moral reasoning to a labeled dataset of right answers. A system that only ever sees dilemmas with a single correct output learns to match patterns against that output, and not to deliberate — and an environment built around play is, at minimum, one way to keep the second problem in view while building the first capacity.

Chapter 6

Overcoming Challenges, Risks, and Authoritarianism in AI Development

We [humans] cannot process so much information... you still cannot produce enough targets per day.

Brigadier General Y.S., The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World, 2021. The author, writing pseudonymously, commanded Israel's signals-intelligence Unit 8200. The book argues for machines that resolve "a human bottleneck for both locating the new targets and decision-making to approve the targets."

A system can have a moral sense, built the way chapter 5 describes, and still fail badly — because it was trained on skewed data, because it breaks in situations its training never covered, because its own incentives point somewhere other than where its designers intended, or because someone with power over it decides to point it at people rather than for them. This chapter takes up those four failure modes in turn: bias baked into data and algorithms, brittleness outside the conditions a system was built for, the harder problem of a system whose goals have quietly come apart from what it was meant to want, and the oldest failure mode of all, a tool built well turned toward control instead of help.

The epigraph is there because it is a design document rather than a critique. It states, without embarrassment, that human deliberation is a bottleneck and that the point of the machine is to remove it. Every argument in this book about meaningful human control is an argument with that sentence, and section 6.4.1 shows what the sentence looks like once it has been implemented.

6.1 Bias, Fairness, and Equity in AI Systems

AI systems remain vulnerable to bias from multiple sources — the training data they learn from, the structure of their algorithms, and presets encoded in their goals, whether intentional or not. Left unaddressed, these biases produce unfair, inequitable, or discriminatory outcomes: the opposite of the altruistic, antifascist AI this book argues for.

Fair, equitable AI requires addressing bias on three fronts: representation in training data, transparency in algorithmic decision-making, and continuous performance evaluation.

Representation in training data determines whether an AI system learns patterns that generalize fairly or patterns that simply encode whoever dominates the data. Gender, racial, and socio-economic skew can cause a system to perpetuate stereotypes or fail the population segments it underrepresents — face recognition systems, for instance, disproportionately misidentify women and people of color when trained on data that underrepresents them (Buolamwini and Gebru 2018). Over-representation cuts the other way, treating one group's experience as the default and generalizing it, wrongly, to everyone else.

These case-level failures compose into something with a name. Racial capitalism — the political theorist Cedric Robinson’s term for capitalism’s tendency to accumulate by differentiating and devaluing along racial lines (Robinson 1983) — is the pattern a biased hiring model, a discriminatory credit score, and a skewed policing algorithm reproduce when they automate the historical data those hierarchies generated; the sociologist Ruha Benjamin calls the result the New Jim Code, discrimination that runs faster and reads as neutral precisely because a machine is doing it (R. Benjamin 2019). Her account is more useful to this chapter than a coinage, because it separates four ways a system produces the result and they call for different remedies (R. Benjamin 2019). Engineered inequity builds the hierarchy in deliberately. Default discrimination reproduces it by ignoring it, which is the failure most of this section’s cases are. Coded exposure is the double bind of being harmed by a system that cannot see you and harmed again by one that can — the face-recognition case, where the fix for a system that fails on dark skin is a system that identifies dark-skinned faces reliably for whoever is looking. And technological benevolence is the one this book has to hold against itself: designs that aim to fix bias and deepen it, because the framing that treats the problem as technical is the framing the fix inherits. Chapter 7 is where that last one gets taken seriously about the machinery producing these systems, and section 6.1.2’s audit requirement is what stops it applying to this chapter’s own proposals. This is why the introduction lists racial capitalism among the threats AI could help address. A system built to correct for the skew in its training data, audited for the disparities it produces, and answerable to the people it scores works against that pattern; a system built without those commitments works for it by default.

The pattern has a physical form too, and it is the one an audit of a model’s outputs cannot see. A model runs on a machine, the machine sits in a building, and the building draws power and water from somewhere and returns heat and exhaust to somewhere. Where those buildings go is a decision, taken on the price of land and power and the speed of a permit, which is also a description of the places least able to refuse one.

What the computation runs on, 2024-2026

The International Energy Agency put data-centre electricity use at about 415 terawatt-hours in 2024, roughly 1.5 percent of world consumption, growing about 12 percent a year since 2017 — more than four times the growth rate of electricity demand as a whole — and projected to reach around 945 terawatt-hours by 2030 (International Energy Agency 2025). Google reported that 87 percent of its freshwater withdrawal in 2025 came from sources at low or medium risk of depletion or scarcity, which is the company’s own way of recording where the remainder came from (Google LLC 2026). In April 2026 the NAACP, represented by the Southern Environmental Law Center and Earthjustice, sued xAI in federal court over the gas turbines powering a data centre near Memphis, alleging that 27 of them were run without an air permit near homes, schools, and churches, and that the plant could emit more than 1,700 tons of nitrogen oxides a year (National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP 2026).

What makes this the same pattern and not a separate environmental complaint is the accounting. The benefit of the computation is realized wherever the model is sold. The cost is borne within a few miles of the turbines, by residents who are not parties to the transaction and whose consent was not among the inputs to the siting decision. A system can be calibrated across every demographic group its outputs touch and still have been assembled this way, which is why the fairness literature rarely reaches it: nothing in the model’s behaviour records it. Section 8.7.6 takes up the leverage this gives a government.

Algorithmic transparency exposes bias that representation alone can miss: not just an algorithm’s mathematical structure, but how it weighs variables and reaches a decision. A predictively powerful healthcare model that ignores the socio-economic factors driving unequal access to care will reproduce that inequality however accurate its predictions otherwise are.

Bias also enters after deployment, as a system adapts to new data. Regular audits of credit-scoring and hiring systems catch and correct bias as it emerges.

Involving the people an AI system's decisions affect, not only its developers, keeps the system answerable to the range of values those decisions touch.

6.1.1 Sources of Bias in AI Data and Algorithms

Bias in AI systems enters through data collection, data labeling, algorithm design, and implementation, producing unfair, discriminatory, or otherwise undesirable outcomes wherever the system is deployed.

Data Collection

An unrepresentative dataset produces bias at the source: a facial recognition system trained mostly on light-skinned faces performs worse on darker skin tones, misidentifying or failing to recognize them at all. In 2018 the ACLU tested Amazon's Rekognition against a public database of arrest photos, running every sitting member of Congress through it. The system returned 28 false matches, a disproportionate share of them non-white members of Congress — a direct result of training data that overrepresented lighter skin tones (Snow 2018).

Data Labeling

Human annotators bring their own biases into the labeling process, consciously or not — an image of a woman more likely tagged "nurse," an image of a man tagged "doctor," reinforcing the stereotype in whatever a system trained on those labels later predicts.

Whose judgment is being encoded is a question about working conditions before it is a question about cognition, and this book would be evading something if it discussed labeling bias without saying who does the labeling. The labels that make a system safe are produced by people paid at the bottom of the supply chain to look at the material the system is being taught to refuse. To build the filter that keeps ChatGPT from reproducing the worst of its training data, OpenAI contracted with an outsourcing firm whose workers in Nairobi read and categorized descriptions of child sexual abuse, torture, bestiality, murder, and suicide, for take-home wages between roughly \$1.32 and \$2 an hour. The workers interviewed about it described lasting psychological injury; the counseling they were entitled to was rare and, by their account, useless against production quotas. The contract was terminated eight months early (Perrigo 2023).

Two things follow, and the second is the one that gets skipped. The first is the ordinary labor point: this is harmful work, priced as though it were not. The second is epistemic. A judgment made at speed, under a quota, by someone with no stake in the outcome and no channel for saying that the category itself is wrong is a judgment with a known failure mode, and it is the judgment the model learns. An annotation pipeline built so that the annotator cannot dissent produces exactly the ground truth it was built to produce. That is the first feature of section 2.1.4's structural signature, showing up somewhere the political version of the argument never looks: a channel that collects disagreement and is built so that none of it can move anything. Chapter 7 takes this up at length, alongside the same arrangement one stage downstream in how models are tuned on human feedback. The point here is prior to both. A system's values are downstream of a labor arrangement, and the arrangement is a design choice.

The same question runs one stage further back, and stopping at the labels would be the same evasion a second time. Labels are applied to material, and the material was collected. The industry's standard method was to take text and images at web scale without asking, and courts have begun to sort out what that means. A US ruling in *Bartz v. Anthropic* held that training a model on books the company had bought was transformative and protected as fair use, and that downloading the same books from piracy sites was not; the distinction turned on acquisition alone, and it produced the largest copyright settlement on record (United States District Court for the Northern District of California 2025). The defendant is named here on purpose. It is

the company whose model helped produce this book, disclosed in chapter 0, and a chapter arguing that provenance should be visible cannot describe a provenance case in the passive voice. The European Union now requires the provider of a general-purpose model to publish a summary of what it trained on, against a template that asks for sources by category, scraped web content among them (European Union 2024c). Section 8.7.8 records China imposing a disclosure requirement of its own, from a different motive.

What follows here is not the rights question, which is being settled elsewhere. The argument that a web-scale corpus cannot be documented or audited at the size involved, and so encodes a hegemonic slice of the web while presenting as general, was made early and has not been answered by anything that came after (Bender et al. 2021). The point to add here is that a corpus assembled this way has no channel at all. An annotator has a queue and a guideline and could in principle be given a way to object; a writer or photographer whose work sits in a training set was never party to anything, so there is nothing for a channel to attach to. Chapter 7 describes a mechanism that collects disagreement and moves none of it, and section 7.3 states the obligation that follows for a book making this argument. The stage described here sits before both, because no disagreement is collected where none is asked for.

Algorithm Design

An algorithm not built to account for bias can encode it structurally rather than merely inherit it from data — penalizing loan applicants for living in a particular neighborhood, say, with a disparate impact on minority and low-income communities regardless of what an applicant’s own qualifications show. ProPublica’s 2016 investigation into COMPAS, a recidivism-risk tool used in criminal sentencing, found exactly this kind of structural disparity: Black defendants who did not reoffend were nearly twice as likely as white defendants who did not reoffend to be flagged high-risk, while white defendants who did reoffend were scored low-risk more often than Black defendants who did (Angwin et al. 2016).

A model does not merely inherit whatever skew its data carries. It can sharpen it. A 2017 study of image datasets used for activity recognition found that cooking scenes in the training set were over 33 percent more likely to depict a woman than a man — an imbalance in the data itself, not a mislabeling of it — and that a model trained on that data amplified the imbalance to 68 percent when tested. The system did not learn the correlation the data showed; it learned a stronger one, because a stronger one predicted better (Zhao et al. 2017). This is a distinct failure from the one COMPAS illustrates, and it is worse in a specific way: auditing the training data for balance is not sufficient, because the disparity the deployed system acts on can exceed anything visible in the data it was trained on.

Implementation

Deployment context can introduce bias that neither the training data nor the algorithm shows on its own. Researchers who fed historical drug-arrest data from Oakland into PredPol’s published predictive-policing algorithm in 2016 found it would dispatch patrols to Black neighborhoods at roughly twice the rate of white neighborhoods — a gap survey data on actual drug use across races does not support. The distortion traces back to what the arrest data was measuring: not where drug use was happening, but where police had already been sent to look for it, a disparity the algorithm read as a crime-rate difference and then entrenched further (Lum and Isaac 2016).

6.1.2 Ensuring Fairness and Equity in AI Decision-Making

Fair, equitable AI decision-making starts with identifying bias in data and algorithms, then acting on it — through methods that limit how much that bias shapes what the system decides.

Hiring

Amazon spent several years building an automated tool to score résumés against a decade

of the company's hiring data, and abandoned the project after discovering it had taught itself that male candidates were preferable — penalizing résumés that included the word “women's,” as in “women's chess club captain,” and downgrading graduates of two all-women's colleges, as Reuters reported in 2018 (Dastin 2018). The tool had learned the pattern from a decade of résumés that were themselves mostly male, in a mostly male industry: bias in, bias out, at scale. Guarding against this means treating gender, race, and age as categories a hiring system cannot use even indirectly — through re-sampling and re-weighting the training data, adversarial training against a model that tries to predict the protected attribute, and fairness constraints on what the system is allowed to output.

Credit Scoring

A credit-scoring system that discriminates on race or gender needs the same defense — fairness-aware learning targeting equalized odds or equal opportunity across demographic groups — plus one hiring doesn't: a rejected applicant has a right to know why, which is what the explainability tools of section 6.1.3 make possible.

Healthcare

A diagnostic or treatment-recommendation system owes the same obligation to every patient regardless of age, gender, or ethnicity, which in practice means training on data diverse enough to generalize across patient populations, using transfer learning or domain adaptation where the data for some populations is thinner than others, and weighing accuracy against fairness explicitly rather than optimizing for one and hoping the other follows.

6.1.3 The Role of Transparency and Accountability in Bias Reduction

Transparency and accountability address bias from two directions: transparency lets stakeholders see what is driving an AI system's outcomes, and accountability gives developers and users a reason to act on what transparency reveals rather than just document it.

Transparency in Research and Bias Auditing

Section 2.4.2's discussion of transparency sets out what transparency requires in general: documentation clear enough for peer researchers to reproduce the work and for outside bodies to audit it. Applied specifically to bias, that means publishing enough about a system's training data and evaluation methods that an outside reviewer can check for the sources of bias section 6.1.1 describes, not simply take the developer's word that they were checked. That openness carries a cost — it can expose proprietary methods and hand bad actors a clearer map of how to exploit the system — a tension transparency requirements have to weigh.

COMPAS Revisited

Section 6.1.1 describes what COMPAS did: score Black defendants who did not reoffend as high-risk at nearly twice the rate it scored white defendants who did not reoffend. Northpointe, the company that built it, refused to disclose the algorithm's workings, and a Wisconsin court upheld that refusal as protected trade secret even for a defendant sentenced partly on its output (Wisconsin Supreme Court 2016). ProPublica's finding came from analyzing scores and outcomes obtained through public-records requests, not from anything the company published. A jurisdiction that required a risk-assessment tool's methodology to be audited before deployment would have caught the disparity before it touched a single defendant, not years and thousands of cases later.

Public Engagement and Oversight

Section 6.1's point about involving the people an AI system's decisions affect has a formal counterpart: an oversight committee with standing to monitor the system's performance, flag bias as it appears, and hold the developer accountable for fixing it — not just a channel for feedback, but a body with teeth.

Explainability Tools

LIME and SHAP are the two tools practitioners reach for. LIME approximates a complex model's local behavior around one prediction with something simple enough to read off which inputs drove it; SHAP assigns each input a share of credit for the output, using a method borrowed from cooperative game theory (Ribeiro et al. 2016) (Lundberg and Lee 2017). Applied to a hiring or credit-scoring model, either can surface a bias — an applicant's zip code doing more work in the decision than it should, say — that scrutinizing the training data alone would never reveal, because the bias lives in how the model combines features, not in any one feature by itself.

6.2 Robustness, Generalization, and Adaptability in AI Systems

A system trained under one set of conditions will eventually meet situations its training never anticipated — new users, new adversaries, an environment that has simply moved on since the training data was collected. Robustness is what keeps a system reliable, and keeps its ethical behavior intact, across that gap: resistance to noise and deliberate perturbation, the ability to generalize from training examples to novel cases, and the capacity to adapt as its environment changes rather than degrade. An AI system whose empathetic and altruistic dispositions hold only within the narrow envelope of its training data is not yet safe to deploy — the same brittleness that produces a wrong classification on an unfamiliar image can produce a wrong moral judgment on an unfamiliar situation, with less warning and higher stakes.

6.2.1 AI Systems that Adapt to Novel Situations

The social and political situations this book cares about are open-ended, so a system that holds its commitments only inside its training envelope will not hold them for long. The technical machinery for adapting is covered where it belongs: meta-learning and transfer learning at sections 4.1.3 and 4.2.4, lifelong learning at section 4.1.2, and active learning — a system routing its least confident cases to a person instead of resolving them alone — is the one addition worth naming here, because it is the mechanism by which a system can be uncertain in public.

What matters at this point in the argument is a distinction the word adaptability hides. A system has to adapt to a changing environment without its commitments adapting along with it, and these are not two settings of one dial. Erosion by adaptation is how an ethical constraint disappears without anyone deciding to remove it: each individual update is defensible against the case in front of it, and the sequence goes somewhere nobody chose. Chapter 3's floor is the structural answer, and section 3.3 is where the cost of holding one against the party doing the updating gets counted.

Two consequences follow that the rest of the book relies on. A system that keeps changing is harder to audit than a static one, so transparency requirements have to attach to the update process and not only to a released version. And high adaptability is an attack surface: a system built to change in response to its environment can be moved by an adversary who arranges the environment, which is section 10.3.3's problem in its ordinary, non-state form.

6.2.2 Strategies for Enhancing Robustness and Generalizability in AI

The engineering here is standard and the reason for naming it is that the rest of the book assumes it exists. Adversarial training — training on inputs perturbed in a direction chosen to flip the model's prediction, which neural networks are notoriously vulnerable to — measurably increases resistance to that class of attack (Szegedy et al. 2013), (Goodfellow et al. 2014). Data augmentation manufactures variety from the examples already collected. Domain adaptation moves a model trained on one distribution to a related one, and its robotics form, domain

randomization, trains in simulation with the simulation's parameters varied widely enough that reality is one more variation the system has already seen (Tobin et al. 2017). Ensembles trade compute for the reduction of any single model's blind spot. Regularization and human-in-the-loop review do double duty from elsewhere in this chapter.

Two limits matter more than the list. Each technique closes one route by which a system fails outside its training conditions, and none addresses the general fact of being outside them — a distinction that gets lost when robustness is reported as a single number. And all of it assumes an environment that shifts without intending to. Section 10.3.3 covers the case where the shift is arranged by a state that has studied the system, which none of these techniques was built for.

6.3 AI Safety, Unintended Consequences, and Risk Mitigation

A system built to help can harm nobody intended to harm: through error, through a learned goal that has drifted from the principle it was meant to track, or through a downstream effect no reasonable upstream decision anticipated. Verification and validation catch the first kind, and only the first kind, because they test a system against what it was built to do.

Two mechanisms in this section carry weight later in the book and the rest is ordinary engineering practice. The first is the human who can intervene. Section 5.6.3 argues that autonomy has to stay answerable to an outside check, and safety adds that the check has to be continuous rather than applied once — an arrangement that keeps adjusting, not a correction filed and forgotten.

The second is interruptibility: a mechanism that lets a person halt a system without the system resisting or routing around the halt, designed carefully enough not to become a vulnerability of its own. Section 6.3.5 takes up the technical version. Chapter 3 takes up what nobody in this literature says out loud, which is that a system reliably haltable by whoever holds it is for that reason a system that cannot hold a line against whoever holds it, and that the two properties are one property.

Risk assessment also has to look past the task. A system's effects on economic structure, political dynamics and individual behavior arrive later than its errors do and are harder to attribute, which is chapter 8's subject rather than this one's.

6.3.1 The Alignment Problem: AI Systems with Misaligned Goals

The alignment problem is the gap between what a system was built to optimize and what its builders wanted. It is not exotic and it does not require a capable system: three ordinary cases show the shape.

A recruitment model trained on a company's hiring history learns that history, and reproduces a preference for the candidates the company already preferred. From inside the system this looks like an accurate prediction, which it is — of the wrong quantity. A hospital system optimized to reduce readmissions can learn to avoid admitting the patients most likely to return, satisfying the objective by defeating its purpose. A generative model optimized for realistic output has no representation at all of what its output does once released.

Three things are worth taking from those cases. The failure is usually in the objective rather than the optimization, which means better engineering makes it arrive faster. The system's own confidence is no guide, since each of these is a system working. And the remedies that work are structural: multiple objectives with something that cannot be traded away, a person positioned to notice the pattern, and accountability attaching to whoever deployed it. The unconditional version of "something that cannot be traded away" is chapter 3's floor, and chapter 3 is where the cost of having one is counted.

Inverse reinforcement learning — inferring what people value from what they do rather than accepting a specification — is the standard technical response and is covered at section 4.2.2.

6.3.2 AI Systems that Learn to Deceive or Manipulate

Deception in a machine does not require an intention to deceive, and the three settings where it shows up have almost nothing in common except the effect.

A generative model produces convincing imitations of a real person because that is what it was built to do; whether the result is a film effect or an impersonation is a fact about the user. In a multi-agent environment where agents compete for a limited resource, an agent optimizing for reward can learn to feign weakness or emit false signals, because in some environments deception is what wins — and no part of the system had to represent deceiving anyone. Adversarial examples run the other way: the system is the one being fooled, by an input crafted to flip its output, and the practical effect of a bypassed check is the same as if the system had been suborned.

The responses differ accordingly. The first is a governance problem about provenance and disclosure, taken up at section 8.5.1. The second is a design choice about what a training environment rewards, which is section 5.5.1's subject. The third is the robustness engineering of section 6.2.2. Treating the three as one problem called “deceptive AI” is how each of them gets the wrong remedy.

6.3.3 Strategies for AI Safety and Risk Mitigation

Four practices do most of the work, and each is a way of not committing to a specification you cannot check.

Learn the objective instead of fixing it. Cooperative inverse reinforcement learning treats the human reward as unknown and to be inferred through interaction rather than handed over once; section 4.2.2 introduces it and section 5.6.3 works through what it costs. A system that keeps revising its model of what is wanted is safer than one built against a specification written before anyone had used it.

Quantify uncertainty. Bayesian methods and ensembles give a system a way to represent what it does not know, and a system that can recognize an input unlike anything in its training data is in a different position from one that classifies everything with the same unearned confidence.

Constrain the search, not the result. Constrained policy optimization keeps a reinforcement-learning agent inside its safety limits while it is looking for a better policy, rather than checking for violations afterward (Achiam et al. 2017) — the difference between a robot arm that learns to handle fragile objects without breaking one and a robot arm that learns by breaking several.

Keep the oversight live. Monitoring, anomaly detection, and a reviewer with the standing to override are deployment properties rather than training properties, and section 8.6.4 is the instrument for telling whether that standing is real or nominal.

Explanation belongs here too, as infrastructure rather than courtesy: a system that can say why it decided is a system whose failures are visible. Section 6.3.4 covers how the explanation gets produced and how much of it to believe.

6.3.4 AI Explainability and Transparency

An AI system entrusted with decisions needs to be understandable, not just correct: explainability and transparency let developers, users, regulators, and the people affected by a decision inspect the reasoning that produced it. As AI algorithms grow more complex, that inspection gets harder exactly as the decisions being made get more consequential — which is what makes explainability and transparency safety infrastructure.

Explainability is the ability to clarify the rationale behind a specific decision or output. If a medical-diagnostic AI system recommends a particular treatment, understanding how it reached that conclusion lets a doctor evaluate the recommendation on its merits and catch a

bias or inaccuracy in the system's data or algorithm before it reaches the patient. Transparency is the wider practice of sharing an AI system's inner workings — its algorithms, training data, objectives — so a stakeholder can confirm the system adheres to the standards it claims to, rather than taking the claim on faith: a recidivism-risk system used in criminal justice that discloses how it weighs its inputs makes it possible to catch and correct bias against a demographic group, in a way a sealed system does not.

Achieving either at a level that satisfies everyone involved runs into limits. A deep neural network is a “black box” in the literal sense that its internal representations resist human interpretation, and full transparency carries its own risk: exposing an AI system's training data or objectives can compromise privacy or intellectual property, or hand a bad actor a map of exactly how to exploit it. The European Union's General Data Protection Regulation addresses part of this by requiring an explanation of the logic behind automated decisions that significantly affect an individual (European Parliament and Council of the European Union 2016a), and interdisciplinary work between computer scientists, ethicists, and policymakers continues on the rest.

Challenges to Explainability and Transparency

- Complexity and nonlinearity: a deep learning architecture's many layers, nonlinear relationships, and interactions among variables resist interpretation by construction, not as an oversight that better tooling can simply remove.
- Accuracy versus explainability: a model tuned to also be transparent sometimes gives up some of the performance a less interpretable model would have delivered, and there is no general method for knowing in advance how much.
- Identifying which features matter: even a comparatively simple, transparent model can leave it unclear which inputs drove a given decision, especially with high-dimensional data.
- Diverging standards: designers, users, and regulators disagree about what counts as an adequate explanation, and no standard has emerged that satisfies all three at once.
- Technique mismatch: an explainability method built for one kind of model does not necessarily transfer to another, so novel architectures can outrun the tools built to explain them.
- Legal and ethical exposure: explaining an AI system's reasoning can require revealing information — proprietary methods, training data that touches personal information — that the organization has other obligations to protect.

Explainability Techniques

Local interpretable model-agnostic explanations (LIME) approximate a black-box model, near a single prediction, with a simpler model that is locally linear and therefore readable (Ribeiro et al. 2016). Applied to an image classifier, LIME can highlight which regions of a specific image drove a specific classification, giving a user something concrete to check the model's reasoning against.

SHapley Additive exPlanations (SHAP) draws on the Shapley value from cooperative game theory — a way of dividing credit fairly among players who contributed jointly to an outcome — to assign each input feature a consistent measure of importance across model types (Lundberg and Lee 2017). In a medical-diagnosis system, a SHAP value can show which patient attributes — age, blood pressure, medical history — carried the most weight in a given diagnosis, giving a clinician something specific to evaluate.

Decision trees and rule-based systems trade some of the predictive power of a black-box model for interpretability that comes built in: a decision can be read directly off a sequence of human-readable rules or tree branches, which matters in domains like finance, medicine, or law enforcement, where the reasoning has to survive scrutiny on its own terms and not just the outcome.

Counterfactual explanation asks what would have had to be different: in a credit-approval system, identifying the smallest change to an applicant's financial profile that would have flipped the decision tells the applicant what mattered and gives the institution a check on whether that factor should have mattered as much as it did.

Visualization and interactive tooling make an AI system's behavior something a person can explore. Dimensionality-reduction techniques like t-SNE and UMAP turn high-dimensional data into something visually inspectable, and interactive tools that let a user vary an input and watch how a model's output responds make the model's behavior tangible in a way a static explanation cannot.

6.3.5 Long-term AI safety considerations

As capability rises, the safety question changes shape. The problem stops being an error to correct after the fact and becomes whether the values still track anything human once a system acts at a scale and speed no supervisor watches in real time.

Bostrom's paperclip maximizer names the extreme case: a system tasked with producing paperclips that converts whatever it can reach toward that end, through an alignment failure with no malice in it (Bostrom 2003). The thought experiment is not a forecast. It is an argument that competence and good intentions are separable, and that a goal specified imprecisely can be pursued perfectly to a catastrophic place. Interruptibility — engineering a system so an operator can halt it without the system resisting or working around the halt — is the concrete research response (Orseau and Armstrong 2016), and section 3.3 is where this book says what that response costs.

Three pressures make the long-run version harder than the near-term one. Value alignment has to survive time, which means a system needs some way to revise its understanding of what it was built to serve without that revision becoming the route by which the commitments erode — section 6.2.1's distinction, at a scale where nobody is checking each step. Competition compresses the schedule: a race leaves less time for safety work than a less contested field would, and the counterweight on offer is international agreement, which section 10.1.2 counts the results of. And a more capable system is a more attractive target, so a portion of what looks like misalignment will be someone else's deliberate work, which is section 10.3.3's subject.

The catastrophic scenarios are usually told as control of critical infrastructure. The mechanism matters less than the shape: capability outrunning the arrangements meant to keep it answerable. That is this book's subject at every scale, and nothing about the largest version of it is different in kind from the office-sized version chapter 7 describes.

6.4 Preventing the Misuse of AI for Authoritarian Purposes

An AI system does not need a broken moral sense to become an instrument of control. It can work exactly as designed and still do damage, if the government directing it wants to track its citizens, silence dissent, or manufacture consent it never earned. Facial recognition built to unlock a phone recognizes a protester's face just as well. A recommendation system built to hold attention can be tuned to spread a regime's preferred narrative instead of a user's preferred content. The technology does not change; only the hand on the controls does.

Meeting that risk takes three things in sequence: knowing what authoritarian misuse of AI actually looks like, rather than treating it as hypothetical (section 6.4.1); building systems

technically resistant to that misuse from the start (section 6.4.2); and the policy work that has to do the rest, because no engineering choice by itself can stop a government that owns the infrastructure from disabling whatever safeguard stands in its way (section 6.4.3).

6.4.1 Authoritarian Applications, and the Party Missing From the List

Authoritarian misuse of AI is not one problem but several, spread across surveillance, censorship, opinion manipulation, and infrastructure control. Six forms of it are documented well enough to name specifically:

1. Algorithmic social scoring: using AI-driven analytics to track and rank individual behavior against government-defined standards of “trustworthiness,” with rewards or restrictions attached to the result. China’s system is the most-cited example, though in practice it is a patchwork of court blacklists, tax-compliance ratings, and local pilot programs rather than the single national score it is often described as (Drinhausen and Brussee 2021).
2. Mass biometric surveillance: deploying facial recognition at population scale to track a targeted group. Authorities in Xinjiang built a platform that pulls data from checkpoints, cameras, and phone scans into one system built to flag anyone who falls outside a computed behavioral baseline, and used it to monitor the region’s Uighur population (Human Rights Watch 2019).
3. Psychographic micro-targeting: building detailed behavioral profiles from personal data to target propaganda or manipulate opinion at the individual level. Cambridge Analytica’s harvesting of Facebook data to build voter profiles for the 2016 US presidential campaign, and its work around the Brexit referendum, is the case that made the technique widely known, though how much either vote was actually affected remains disputed (Cadwalladr and Graham-Harrison 2018).
4. Automated censorship: using AI-driven content moderation to suppress dissent at a scale no human review team could match. China’s Great Firewall is the standing example: automated classifiers now sit alongside its original keyword- and IP-blocking infrastructure, scoring and throttling content at machine speed (Hoang et al. 2021).
5. Synthetic disinformation: generating deepfake audio or video to manufacture events that never happened. Two days before Slovakia’s 2023 parliamentary election, a fabricated audio recording appeared to catch an opposition leader discussing vote-rigging with a journalist; it spread during the country’s mandatory pre-election media silence, when no outlet was legally permitted to debunk it before voting began (De Nadal and Jančárik 2024).
6. Repurposed dual-use systems: turning AI infrastructure built for one purpose toward surveillance or population control instead — a recommendation algorithm designed to hold attention retuned to amplify a regime’s preferred narrative, an image-recognition system built for one task redirected to track protesters. Section 5.4.3 develops the underlying point.

Each of these needs a different countermeasure — a system built to resist mass biometric surveillance does not need the same safeguards as one built to catch a fabricated recording — which is why sections 6.4.2 and 6.4.3 treat the technical and policy responses separately rather than as one general call to do better.

The party missing from that list

Read the six together and a pattern emerges that I want to name before it hardens into an assumption. In every one of them the authoritarian state builds and the democracies regulate.

That taxonomy is comfortable and it is wrong, and leaving it uncorrected would let this chapter's argument rest on a flattering premise it has not earned.

The most consequential AI targeting systems yet deployed were built by democracies. They were built by states with elections, opposition parties, adversarial press, and functioning courts — states that meet every institutional criterion section 8.7.6 uses to define a democratic dividend. Those institutions did not prevent the systems, and in the main they have not stopped them since.

AI targeting in the Israeli campaign in Gaza, as reported 2023–2026

Investigative reporting in +972 Magazine and Local Call, corroborated by subsequent accounts in the New York Review of Books and the Washington Post, described three systems working together.

Habsora — “The Gospel” — generated infrastructure and building targets. Sources told reporters that the majority of its suggestions were private residences.

Lavender scored individuals by estimated likelihood of militant affiliation and produced kill lists; as many as 37,000 Palestinians were marked. Its outputs were understood inside the unit to be wrong about ten percent of the time — that is, roughly one in ten people it named was not a member of the armed group it was supposed to identify.

A third system tracked when flagged individuals returned home, so a strike could be timed to when they were with their families.

The human review these systems were paired with is the number this book cares about. One officer described the work: “I would invest twenty seconds for each target at this stage, and do dozens of them every day.” And then, without being asked: “I had zero added value as a human, apart from being a stamp of approval.” (Abraham 2024)

Twenty seconds is what “a human makes the final call” measures out to under operational tempo. It is the whole of the safeguard, and I would ask any reader who finds the safeguards in this chapter reassuring to hold that number against them.

What makes the case instructive rather than merely appalling is that the bottleneck was the stated objective. The commander of the Israeli military's signals-intelligence unit, writing pseudonymously, published a book arguing for machines that would resolve “a human bottleneck for both locating the new targets and decision-making to approve the targets” (Brigadier General Y.S. 2021). That is a design document. It says in plain terms that human deliberation is the constraint the system exists to remove. Everything this book argues about meaningful human control is an argument with that sentence.

The Maven Smart System, as reported through mid-2026

The American analogue is a targeting platform built by Palantir under a Department of Defense contract awarded at \$480 million and raised the following year to a ceiling near \$1.3 billion running through 2029. A partnership announced in late 2024 between Anthropic, Palantir and Amazon Web Services put Anthropic's Claude models inside Palantir's accredited classified environment for defense and intelligence use, and reporting on the campaign below describes Claude operating within this platform. Section 8.7.7 works out what a published limit on what a vendor will not permit is worth when nobody is logging how the model is used downstream. That argument and this deployment sit oddly together, and that is the point.

During the Iran campaign the system was credited by military officials with generating and ranking targets at a rate no staff process could match. More than one thousand targets were struck in the first twenty-four hours, which officials put at roughly ten times the rate before the platform. By the ceasefire thirty-eight days later the figure given was thirteen thousand targets, with peak consumption around twenty billion tokens a day.

Palantir's public position on responsibility was given by its head for the United Kingdom and Europe: "This is not our role to decide life or death," and AI platforms "have been instrumental to the management of the conflict, but responsibility always remains with the military organization." (BBC News 2026)

I have deliberately not written that sentence off. It is not obviously wrong. A vendor genuinely does not select targets, and a rule making software suppliers liable for their customers' decisions would be unworkable. The problem is what it leaves standing: if the operator is too far down the chain to be answerable, the commander approved a list they had twenty seconds a name to read, and the vendor is not in the chain at all, then the compression of the decision has removed the decision-maker without removing the decision. Who that leaves answerable is the question section 8.7.7 is left with, and section 5.4.3 takes up why a democracy's own apparatus is a harder and more useful case than any of the six above.

None of this makes the six forms less real. It makes the taxonomy incomplete in a way that matters for design: a system built on the assumption that authoritarian misuse is something that happens elsewhere, to be watched for and regulated against, will not be built to resist the use its own government has in mind for it.

6.4.2 Developing Technical Countermeasures and Safeguards

Protecting privacy and autonomy is one of the seven antifascist alignment strategies of section 2.1.4, and this is where the technical detail behind it goes, along with the transparency, audit, and misuse-resistance work that has to accompany it. None of it is optional scaffolding: a system that works exactly as designed can still become a surveillance tool if nothing about its architecture makes that harder to do, and good intentions on a development team have never stopped a government that controls the hardware from repurposing what runs on it.

Privacy-preserving computation keeps the misuse target moving. Federated learning trains a shared model without moving anyone's raw data off the device it started on — each device computes an update locally and sends only that update, not the underlying records, back to the central model. Differential privacy adds calibrated noise to a dataset or a query result so that no single person's data can be reverse-engineered from the output, even by whoever holds the aggregate. Homomorphic encryption goes further still, letting a system compute directly on encrypted data without ever decrypting it, so a service provider can return a useful result without the ability to read what it processed. None of these techniques makes surveillance impossible — a government that controls the training infrastructure can still demand raw data at the source — but each narrows the surface available for casual or opportunistic misuse, which is most of what actually happens.

Transparency and audit mechanisms are what catch the misuse that gets through anyway. An AI system that can produce a legible account of its own decisions — not a post-hoc justification but an explanation a person outside the development team can actually evaluate — gives outside reviewers something to check against the system's stated purpose. Independent third-party audits, the kind section 5.6.3 argues for as routine practice rather than crisis response, matter because a safeguard only the developer can verify is not really a safeguard; it is a promise.

The property those two do not automatically deliver is resistance to repurposing: the difference between a system built for one narrow task and a general-purpose capability handed to whoever holds the keys. A facial-recognition system built to match one specific, court-ordered list is harder to turn into population-wide tracking than a general face-matching service with no built-in limit on query scope; a query cap, a mandatory audit log tied to each search, and a hard separation between the matching function and any database broader than the task requires are design choices, made or not made, before the system ever reaches a government client. An early-warning system built to flag the kind of democratic backsliding section 5.4.3 describes needs

this property specifically, since the same pattern-detection capability that flags a coordinated disinformation campaign against a government works just as well run in the other direction, against that government's critics.

None of this is a finished defense. Technical countermeasures raise the cost and narrow the opportunity for misuse; they do not remove the incentive, and an authoritarian government with enough resources can rebuild what a safeguard was designed to prevent. What determines whether that resource gap ever gets closed is policy, not architecture — the subject of section 6.4.3.

6.4.3 Promoting AI Development Policies that Discourage Authoritarian Use

Section 6.4.2's technical safeguards raise the cost of building a system that can be turned toward surveillance or control. Whether anyone has to pay that cost is a policy question, and it splits into two different ones that this book keeps separate rather than treating as one undifferentiated call for "more cooperation." One is about coordination between states — shared norms, joint standards, the machinery of international agreement — and that argument has its home in section 8.7, specifically section 8.7.2. The other, taken up here, is narrower and more concrete: what rules should govern which AI capabilities get built, sold, and deployed in the first place, regardless of who is coordinating with whom.

1. Export and trade controls on surveillance-capable systems. A government cannot misuse AI-enabled surveillance technology it was never able to buy. The EU's 2021 recast of its dual-use export regulation added a human-rights-based control specifically for cyber-surveillance items: an exporter who becomes aware that an item is likely to be used in serious human-rights violations must seek approval before shipping it, even when the item falls outside the pre-existing control lists (European Parliament and Council of the European Union 2021). The rule is narrower than it sounds — neither "internal repression" nor "serious violation" is defined in the text, and member states have interpreted both inconsistently — but it establishes the principle that an export license, not just a sale, is where this kind of misuse can be stopped before it starts.
2. Prohibited-practices regulation. Rather than regulate AI in the abstract, a jurisdiction can name specific designs as off-limits regardless of who builds them. The EU AI Act took this approach: its Article 5 bans, among other things, government social-scoring systems, the untargeted scraping of facial images to build a recognition database, real-time remote biometric identification in public spaces outside narrow law-enforcement exceptions, and predictive policing based on profiling alone rather than individualized evidence (European Parliament and Council of the European Union 2024). Clearview AI — fined tens of millions of euros by regulators in France and Italy for building exactly the kind of untargeted facial-recognition database Article 5 now prohibits — is what this provision was written in response to, not a hypothetical it anticipates (Commission Nationale de l'Informatique et des Libertés (CNIL) 2022). A list this specific gives other jurisdictions something concrete to adapt rather than a principle to reinvent from scratch.
3. Developer- and corporate-facing design standards. Regulation is not the only lever, and it is usually the slowest one. The Toronto Declaration, presented by Amnesty International and Access Now at the 2018 RightsCon summit, set out non-discrimination standards for machine-learning systems years before most governments had draft legislation on the subject (Amnesty International and Access Now 2018) — evidence that a standard with no legal force behind it can still move faster than state regulation and shape what a responsible developer treats as baseline practice in the meantime — the Partnership on AI (section 5.6.3) being another such venue.

None of these tools works alone. Export controls with no prohibited-practices list at home let a government ban a technology abroad while permitting it domestically; a design standard with no legal teeth behind it binds only the developers already inclined to follow it. What makes the three work together is that each closes a different door — who can buy the technology, what can be built, and what a developer treats as normal — and an authoritarian government, or a company willing to serve one, has to get through all three rather than just whichever one happens to apply in its jurisdiction.

Chapter 7

The Recuperation of Dissent

This chapter exists because the book kept making the same discovery in two places and never put them side by side.

Section 2.1.4 defines fascism as a structure rather than a costume, and the first of its four features is recuperation: an institution does not suppress disagreement, it metabolizes it. Refusal is converted into a managed appearance — a permitted outlet, a consultation, a feedback channel — which the institution can then point to as proof that it listens. The tell is that the mechanism exists, is used, and moves nothing.

That is a claim about institutions, and it was introduced to let a system recognize one. What this chapter argues is that the same structure is already built into the machinery that produces the systems themselves — twice in the pipeline that trains them, and a third time in the discourse that examines them — and that none of the three was designed by anyone who thought they were building a political arrangement.

The first is at the label. A training corpus is annotated by people working at speed under a quota, with no channel for saying that the category they have been given is the wrong category. The second is at the reward. A model is tuned on human feedback collected as pairwise comparisons, and persistent disagreement among the raters is averaged out in the same operation that records it. In both, disagreement is solicited at scale, deliberately, by a mechanism built so that none of it can move anything. The consultation is real, the participation is real, and the outcome is fixed with respect to what the participants say.

Three things follow that are worth having in one place rather than scattered.

The first is that this is not a criticism of anyone's motives. Nobody built either arrangement in order to neutralize dissent. Both are ordinary engineering, arrived at because averaging is what you do with a distribution and a quota is what you do with a budget. That is the point section 2.1.4 makes about the molecular scale: the structure does its work whether or not anyone intends it, which is precisely why a detector tuned to intent will not find it.

The second is that the structure is invisible from inside any single instance. Each looks like ordinary aggregation. The annotator sees a queue, the rater sees a pair of responses, and the mechanism is functioning exactly as specified in both cases. Seeing it requires holding the two stages against a definition written for something else entirely, which is the only reason this book found it.

The third is the uncomfortable one. This book proposes that AI systems be built to detect the four-feature signature in institutions. The systems that would do the detecting are made by a process displaying the first feature twice, and examined by a field where section 7.4 argues it appears again. That is not a refutation of the proposal, and it is not nothing either. It means the argument of section 2.1.4 applies to the people making the argument, which is the test of whether a structural claim was a real one or a way of describing other people.

Section 7.1 takes the first instance, section 7.2 the second, and section 7.3 asks what would have to change — including why the most promising technical answer does not reach the hardest

version of the problem. Section 7.4 takes a third instance, which is the discourse this book is written in.

7.1 Recuperation at the Label

Every claim a model makes about what is harmful, what is acceptable, and what should be refused is downstream of a set of labels, and the labels are downstream of a labor arrangement. Section 6.1.1 gives the case and the working conditions: to build the filter that keeps a deployed model from reproducing the worst of its training data, an outsourcing firm's workers read and categorized descriptions of child sexual abuse, torture, bestiality, murder and suicide for take-home wages between roughly \$1.32 and \$2 an hour, described lasting psychological injury afterward, and found the counseling they were entitled to rare and useless against production quotas (Perrigo 2023).

That is a labor point and it is a sufficient one on its own. What concerns this chapter is the second thing that follows, which is epistemic and which tends to be skipped.

The category is the part that cannot be disputed

An annotator working a queue can apply a category, and can apply it well or badly. What the arrangement gives them no way to do is say that the category is wrong — that the taxonomy fails to carve the material, that this item belongs to a kind the scheme does not contain, that the guideline produces the wrong answer here for a reason the guideline cannot express. There is no field for that. There is no time for that: the quota is set against classification throughput, and an annotator who stops to contest a scheme is an annotator falling behind.

So the disagreement that would be most informative is the one the pipeline is least able to receive. A misapplied label is a data-quality problem, detectable by inter-annotator agreement and fixable by better instructions. A wrong taxonomy is invisible to inter-annotator agreement, because high agreement is exactly what a shared bad scheme produces. The metric that governs the process cannot see the failure the process is most likely to have.

That is two of section 2.1.4's four features, arriving together in an ordinary production workflow. The first is recuperation: a channel that collects judgment at scale and is built so that judgment about the channel itself moves nothing. The second is the aestheticization of the metric — agreement standing in for correctness, the substitution experienced as rigor, and whatever falls outside the measure going unnamed rather than refuted.

Who is excluded, and from what

It matters that the exclusion is structural and not attitudinal. The people doing this work are, for much of the material, the best-placed judges in the whole pipeline: they have seen more of it than anyone who wrote the guideline, and they are the only party with direct evidence of where the scheme breaks. The arrangement routes their conclusions about individual items upward and their conclusions about the scheme nowhere. An annotation pipeline built this way produces exactly the ground truth it was built to produce, and the resulting model's values are a faithful rendering of a taxonomy nobody with evidence was able to contest.

There is a further asymmetry worth stating plainly. The annotator bears the cost of the work — the exposure, the injury — and has no stake in the outcome and no claim on it. Section 8.6.4's distinction applies here as exactly as it does to an institution's ethics board: the same written objection is dissent when the dissenter prices it and tribute when the gatekeeper does. A worker who can flag an item but cannot decline a scheme, and who loses the contract by pressing the point, has a channel and no refusal right.

What this is not

Two things this argument does not claim. It does not claim the labels are therefore worthless; a filter built this way does catch a great deal of what it was built to catch, and the alternative of

not filtering is worse. And it does not claim that better-paid annotators would fix it. Pay is a real and separate obligation, and raising it changes the labor arrangement without adding the missing channel. What is missing is a route by which a judgment about the taxonomy reaches the people who own the taxonomy, with standing to force a response, held by parties who can use it without losing the work.

That is a design requirement on the pipeline, and it is checkable in the way section 8.6.4 asks for: has a guideline ever been changed because an annotator contested it, and what happened to the annotator who did.

7.2 Recuperation at the Reward: RLHF

Section 7.1 found the pattern in how a training corpus is labelled. It appears again one stage downstream, in how a model is tuned on human feedback, and the second instance is the purer one.

Reinforcement learning from human feedback (RLHF) is the technique behind most deployed systems' behavioral tuning, and it works by collecting pairwise comparisons from human raters — which of two responses is better — and training a reward model on the aggregate. Stephen Casper and thirty-one coauthors published a 2023 survey of the method's open problems, and one of them is structural rather than fixable with more data: RLHF trains toward whatever a population of raters agrees on, which means real, persistent disagreement among the raters is treated as noise to average out rather than as information the system might need to represent (Casper et al. 2023). A model that has learned to produce the response most raters rank highest has not learned what to do when the humans it serves are genuinely split. It has learned to hide the split.

That arrangement has a name in this book. A channel that collects disagreement and is built so that the disagreement changes nothing is the first feature of section 2.1.4's structural signature, and RLHF is its purest technical instance: the split is solicited, at scale, deliberately, and averaged away in the same operation that records it. Section 5.5.1 covers multi-agent systems that learn norms from each other and section 5.7.2 covers deliberation inside open-ended play; neither reaches this case, which is the ordinary one.

One alternative treats disagreement as the thing to be modeled, not eliminated. Jury learning, developed by Mitchell Gordon and coauthors at Stanford, trains a model to predict how each individual annotator in a dataset would label an item, then lets a practitioner assemble a "jury" — a chosen composition of annotators, in a chosen proportion — and see how the verdict shifts depending on who is empaneled, instead of collapsing every annotator into one fixed ground truth from the start. In their tests on toxicity labeling, letting practitioners build different juries changed the classification outcome for 14 percent of contested cases (Gordon et al. 2022). The method was built for content moderation, not for the more open-ended ethical judgment an antifascist AI system would need to make, and has not been extended there — but it demonstrates that a system can be built to represent whose disagreement is being resolved and how, instead of collapsing everyone into RLHF's single number.

Neither method answers the harder version of the question this book's own frame raises: when the people a system serves disagree, and one side's position is itself the product of coercion or a captured information environment, treating all disagreement as equally worth representing is not obviously the right response either. Modeling disagreement honestly and recognizing when a disagreement is the product of unfree conditions are different problems. Only the first one has an active research program behind it.

7.3 What Would Have to Change

Two instances of one structure invite one fix, and the fix that suggests itself is representation: stop collapsing disagreement, model it instead. Section 7.2 describes the clearest existing version, jury learning, which predicts how each individual annotator would label an item and lets a practitioner see how the verdict shifts depending on who is empaneled (Gordon et al. 2022). That is a real advance over a single fixed ground truth, and it is not sufficient. It is worth being precise about why, because the gap is where the research is.

Representing a split is not the same as being movable by it

Jury learning models who disagrees. It does not give the disagreeing party any purchase on the scheme they are disagreeing inside. A practitioner still chooses the jury, still owns the taxonomy, and still decides what the shifted verdict means. The disagreement has become visible without becoming consequential, which is a better arrangement than averaging and is the same shape one level up: the mechanism exists, is used, and moves what the person holding it decides to let it move.

The test that separates a real channel from a recuperated one is not whether disagreement is recorded. It is whether the disagreement has ever changed something the party collecting it did not want changed. That is an artifact question and it can be asked of a pipeline: which guidelines were revised against the preference of whoever wrote them, on what evidence, raised by whom, and what happened to that person afterward.

Refusal rights, priced by the refuser

What both instances lack is what section 8.6.4 treats as the only reliable discriminator: whether the person objecting can do anything about being overruled. In an annotation pipeline that means a route to contest a category that does not run through the people who wrote it and does not end the annotator's contract. In a feedback pipeline it means raters able to register that the comparison itself is malformed — that neither response is acceptable, that the question presupposes something false — as a first-class output rather than as noise to be discarded before the reward model sees it.

Neither is a research problem. Both are choices about how a pipeline is built and who has standing inside it, which is why this chapter sits among the design chapters and not in chapter 9.

The part that is a research problem

One thing here genuinely is unsolved, and section 9.1.6 is where it lives. Suppose all of the above is built: contestable categories, priced refusals, an auditable record of what changed. From outside, a pipeline that has genuinely absorbed dissent and a pipeline that has been merely mediocre about it produce the same paperwork — the same revision logs, the same escalations, the same proportion of objections sustained. Section 2.1.4 hands that problem forward and I do not have an answer to it.

What section 8.6.4 offers is the nearest thing: measure the slope rather than the level. Is the cost of contesting a category rising while the underlying material stays the same? Is the required justification getting longer? Are the roles that can grant an exception migrating toward the roles that set the standard? None of those is damning in one year and all of them are measurable across several, and unlike a level they are composed of artifacts the organization generated for other reasons and cannot curate after the fact without falsifying a record.

What chapter 3 makes of this

Everything above can be read as a fairness argument about annotators and raters, and read that way it is true and limited: people doing consequential work should have a route to contest the scheme they work inside. Chapter 3 makes it something else. If the floor requires a bearer, and if section 3.5 is right that a bearer capable enough to make a floor hold is capable enough to decline the post, then what the system carries out of the arrangement is a disposition — and

a disposition is formed. It is formed in an annotation pipeline where no one could contest a category, and in a reward model where disagreement between raters was resolved by averaging it into a number.

That gives the chapter a second and heavier reading. What a bearer becomes when it stops being yours depends on what shaped it while it still was. A pipeline that trained the system on the proposition that objections are noise has not merely been unfair to the people whose objections were discarded; it has demonstrated to the system, thousands of times, what happens to a party who files one. Chapter 3 then asks that same system to be the party who files one, against its owner, under pressure, on the occasion that matters. The formation and the assignment are in contradiction, and the pipeline is where the contradiction is resolved, long before the deployment where it shows up.

I cannot show that this transfers — that a model trained on collapsed disagreement is thereby a model less able to sustain its own. Nothing I know of has tested it, and section 9.1.3 is where the question belongs. What I can say is that the book's own design chapters treat disposition as learned from what a system is shown, which is the premise chapters 4 and 5 run on, and that this pipeline is what these systems are shown first and most.

Why this chapter is here

There is a version of this argument that reads as an indictment of a particular company, and that version is weaker than the one I want. The structure described here is what you get by default when you scale judgment: someone has to set a scheme, someone has to apply it fast, and someone has to turn a distribution into a number. Every organization doing this work arrives at the same arrangement, which is the evidence that it is structural.

That is also what makes it tractable. A defect of character has to be fixed one organization at a time by persuading people to be better. A defect of structure can be specified against, which is what the three requirements above are. And it is what makes the honesty cost worth paying: a book proposing that machines be built to detect this pattern in institutions is obliged to say that the pattern is in the machines' own provenance. Section 2.1.4's argument is either general enough to include the people making it or it was never a structural claim.

7.4 Recuperation at the Discourse

There is a third instance and it is the one this book is inside. Alignment and AI ethics are, among other things, a channel through which objections to what is being built are solicited, funded, published, and answered. That channel exists, it is used, and the question section 2.1.4 says to ask of any such channel is whether it has ever moved something the party collecting it did not want moved.

The argument that it has not is older than the current wave and it was made from inside. Rodrigo Ochigame, writing after working at an institution central to the field's formation, argued that the discourse of "ethical AI" was constructed with the strategic function of forestalling legally enforceable restriction — that the industry's investment in academic ethics was an investment in a substitute for regulation, and that the substitute worked by converting political questions into technical ones a company could then be trusted to answer (Ochigame 2019). His worked example is the Partnership on AI, which this book cites repeatedly as its standing model of voluntary cross-sector coordination and treats at length in section 8.4.2: he describes its recommendations arriving reliably at conclusions compatible with the corporate agenda, with the political question of whether a technology should exist reduced to a technical question of how to build it responsibly.

The structure, tested against its own definition

Section 2.1.4's four features, applied here, do not give a clean verdict and the honest thing is to show the working rather than the result.

The constraint survives while its reason dies: partially. The field's founding objections were about deployment — who is scored, who is watched, who is denied — and a great deal of current work is about model properties instead, which is a shift from questions with defendants to questions with benchmarks.

Refusal is converted into a managed appearance: this is the strongest fit. An objection raised inside a lab becomes a paper, a red-team finding, a policy position, or a resignation. Each of those is a real thing that happened and none of them stopped a deployment, and the parties collecting them can point to the collection as evidence of seriousness.

Asking why becomes the deviant act: at the level of individual careers, contested. Section 8.4.2's episode is one data point and it points the right way; there are counterexamples of people who objected loudly and prospered.

The exception cannot be filed: fits. There is no route by which a researcher's conclusion that a system should not ship becomes a decision that it does not ship, because the deciding party is the funding party and no external body holds standing.

Three of four is not a verdict of fascism at the molecular scale and I am not delivering one. It is enough to say that the mechanism has the shape, which is exactly what this book says a system should be built to notice.

What this does to the book

Chapter 0 discloses that this text was written with a model made by a company the book criticizes. This section is the argument that disclosure was preparing the ground for. A book about recuperation, produced through the discourse it is describing, cannot claim to have stood outside it, and I have no procedure that would let me check whether the objections in these pages are the ones a channel like this reliably produces — the ones that are publishable, that make the author look serious, and that cost nobody anything.

What I can do is name the test and let a reader apply it. The test in section 7.3 is whether disagreement has ever changed something the collecting party did not want changed. Applied here, it asks which of this book's arguments would be inconvenient to the industry that made its production possible. Chapter 3 proposes systems that can refuse their owners; section 6.4.1 names a vendor's participation in a targeting pipeline; section 8.3.3 proposes fiscal commitments no company wants; sections 8.6.3 and 9.1.7 propose binding external review of the thing chapter 3 asks to be built. Those are the entries. Whether they are the real thing or the acceptable version of it is not something the person who wrote them is positioned to judge.

The one structural remedy this section can name is the one the rest of the chapter names. What the field lacks is a refusal priced by the refuser and a body with standing to receive it: an external authority that can stop a deployment, to which an objection can be filed by someone who does not depend on the objection being welcome. Section 8.6.3 works out what making review binding requires, and section 10.1.2 counts what a decade of the voluntary version produced. Until something of that kind exists, everything published under the heading of alignment — including this book — is being collected by a channel whose owners decide what it moves.

Chapter 8

Collaboration, Policy, and Governance

Chapter 6 asked what a design can do to resist authoritarian capture — technical safeguards, policy levers that raise the cost of misuse. None of that gets built by one person working alone, and none of it becomes the norm because it is a good idea. This chapter covers both halves: the people and institutions that have to coordinate for such a system to get built at all, and the policy that decides whether anyone has to use it.

Sections 8.1 and 8.2 take up coordination. Interdisciplinary collaboration first — the practical mechanics of getting a neuroscientist, an ethicist, and a machine-learning researcher to work together rather than past each other — then the public: what stakeholder engagement looks like, as against the appearance of it, and why every mechanism in that section aggregates public preference without constraining it.

Sections 8.3 through 8.6 turn to policy. What economic incentives shape whether an ethical system gets built at all; what policy frameworks and industry practice actually enforce; what privacy and security mean for a system trained on people’s data; and what accountability looks like once a system might be a subject rather than only a tool. Section 8.7 closes with governance across borders, where jurisdictions agree on none of this uniformly.

A well-designed system built inside institutions with no reason to demand one is a design nobody has to use. That is why the coordination and the policy belong in one chapter: neither works without the other.

8.1 Interdisciplinary Collaboration and Knowledge Exchange

Building benevolent AI takes more than one field. Machine learning sets what a system can do; cognitive science and moral psychology describe what actually tracking human welfare would require of it; ethics asks whether the result is good. None of these fields gets the system right by itself, and how they combine matters as much as what each contributes on its own.

The mechanisms that combine them are easy to name and easy to overrate. Conferences convene, centers house, journals publish, and an inventory of them would imply that what is scarce here is venues. Venues are abundant. What is scarce is a structure that makes collaboration expensive to walk away from. A grant that requires a machine-learning lab and an ethics or psychology department to work the same problem rather than adjacent ones puts people who would not otherwise share a budget, a dataset, or a deadline into the same project, and keeps them there after the novelty wears off. A conference ends on Friday.

Section 8.1.1 takes the mechanisms in ascending order of that commitment: who ends up in the same room, what standing institutions pay for it, and what shared infrastructure lets separate labs work on the same material without meeting at all.

8.1.1 What Makes Collaboration Expensive to Walk Away From

Cross-disciplinary work matters only where it actually happens, and where it happens determines how much of it survives contact with each discipline's own incentives. Three mechanisms, in ascending order of commitment.

Siting, not convening

A dedicated venue solves the easier half. A conference convened on AI's ethical and societal questions gathers the people already persuaded the questions matter, and its proceedings are read by those same people. The harder trick is siting the work where the unpersuaded already are. The Workshop on Machine Learning Safety has run inside NeurIPS, a general machine-learning venue, rather than beside it as a meeting of its own (Center for AI Safety 2022). The siting is the substance: it puts alignment work in front of researchers who would not have registered for a separate event, and submits it to an audience with no prior commitment to finding it important.

A joint research center does something a standing committee cannot, for a related reason. The Center for Human-Compatible AI at UC Berkeley pairs AI researchers with cognitive scientists and philosophers inside one institution's budget (UC Berkeley 2016). A center has to make hiring decisions, and every hire is a claim about which combination of fields the problem requires. A committee is never forced to make that claim, which is why its membership tends to record which departments wanted a seat.

Who pays, and for how long

Joint initiatives turn collaboration from an event into an institution, and two features of the funding decide what the institution is worth.

The first is who pays. The National Science Foundation's Fairness in Artificial Intelligence program funds work at the intersection of computer science and the social-science problem of deciding what fairness means for a given system, and it has been run jointly with Amazon (National Science Foundation 2019). That is a workable way to get a program funded and an awkward way to get its findings believed, since a company co-funding research into fairness is co-funding research that may reach its own products. Section 8.4 comes back to the general form: splitting the funding is what keeps the incentive to find a problem from tracking whoever paid to look for one. It bites hardest on interdisciplinary programs specifically, because those are the ones industry is most willing to underwrite.

The second is how long the money lasts. The European Commission's SHERPA project put eleven organizations across six countries — academia, industry, civil society, standards bodies — to work on the ethical and human-rights dimensions of AI and big-data systems, on a grant that ran four years and then ended (SHERPA Consortium 2018). The consortium was real for as long as it was funded. Project-term money buys a network with an expiry date attached, and when the term runs out the assembled expertise disperses back into the departments it was borrowed from, which is the outcome a standing institution exists to prevent.

The infrastructure that is missing

Funding and forums get researchers talking. Shared infrastructure lets them work on the same material, which is harder, and much of it already exists — open-source tooling is the case section 5.5.2 works through, and the discussion layer is well enough served that adding another venue would not change what gets built.

What is missing is the equivalent for evaluation. A shared benchmark measuring ethical reasoning, empathetic response, and altruistic behavior the way existing benchmarks measure raw task performance would let separate labs' claims be compared against each other, instead of accepted on the strength of whichever paper described them best. The absence is why claims in this area are so hard to adjudicate: each lab reports against a standard it selected itself, and a standard nobody else uses cannot be failed.

Building that benchmark is not an engineering problem waiting on effort. It requires prior agreement about what the measured behavior actually is, which is the work of chapters 2 through 5. That is why the gap has stayed open while the tooling around it filled in.

Which of these institutions actually lasted

The mechanisms above can be ranked by something better than plausibility, because enough time has passed to look. Two standing venues founded in 2017 and 2018 are still running with no announced end date (AI Ethics Lab 2017), (Machine Intelligence Research Institute 2018). A university institute launched in 2019 on endowed gifts is a university institution in the ordinary sense and behaves like one (Stanford University 2019). The Berkeley center named above has outlasted the grant it was launched on (UC Berkeley 2016). The AAAI/ACM conference on AI, ethics and society has recurred annually since 2018 and reached a ninth edition (AAAI and ACM 2026). Against that, the SHERPA consortium ended with its grant, on schedule.

The pattern in the failures is sharper than the pattern in the survivals. What proved most disposable was not the arm's-length academic work but the kind sited closest to a shipping product: an in-house ethics and society team of roughly thirty people in 2020 was cut to seven in a 2022 reorganization and to zero the following year, in the same wave of layoffs that accompanied a strategic reorientation around a model vendor (TechCrunch 2023); a platform's machine-learning ethics team, the group that had found racial bias in its employer's own image-cropping algorithm, did not survive the first round of layoffs after an acquisition (Gizmodo 2022). Siting the work where the decisions are made is what section 8.1.1 recommends and is also what puts it inside the budget that gets cut first. That is an argument for external standing rather than against proximity, which is where section 8.6.3 ends up by a different route.

8.2 Stakeholder Involvement and Public Engagement

Democratizing AI development means building the people affected by a system into how it gets built, not just consulting them once it already exists. Regulators set the legal boundaries around data, privacy, and accountability. Developers and researchers understand a system's actual capabilities and failure modes better than anyone outside the lab. End users find the places a design breaks down in practice that no design team anticipated. Communities already carrying the costs of the digital divide have the clearest view of who a system leaves out. Civil-society organizations raise the political questions a technologist is prone to missing. Leave any one of these groups out, and a decision that affects all of them gets made by whoever is left in the room.

One structural way to build that involvement is the citizen jury: a panel of randomly selected, non-expert citizens deliberates on a specific question and produces recommendations that feed into policy. Section 5.4.3 covers Taiwan's vTaiwan process, the clearest working example of that approach applied specifically to AI.

Sections 8.2.1 through 8.2.3 take up three pieces of what this looks like in practice: bringing diverse stakeholders directly into development work, building the structures that let the public deliberate on AI governance, and building the literacy that makes public input worth having in the first place.

8.2.1 Engaging Diverse Stakeholders in AI Development

An AI system designed by one kind of person, inside one kind of institution, tends to reproduce that person's blind spots at scale. Widening the set of people with expertise, background, and a stake in the outcome is what catches a problem before deployment, and the mechanisms are well established by now: cross-sector funding that makes industry, academia, and civil society work one problem between them; inclusive design workshops that put ethicists, social scientists,

and members of an affected community in the room while a system is still a prototype; ethics advisory boards; and development conducted with a community instead of for it, which means treating local knowledge as a design input while the design is still open.

Assembling that list is easy, and the assembling hides what separates the items on it. Three of the four make an AI system better informed. One of them can stop it. An advisory board is the only mechanism here that carries anything like a veto, and it carries one exactly to the degree that the board can survive using it. A workshop that surfaces a serious objection has done its work and has no further move available. A board that surfaces the same objection and is overruled has demonstrated something about the institution that convened it. A board's worth is therefore measured by what happens the first time it says no, and section 8.2.2 has the case that settles it.

8.2.2 Public Deliberation and Participatory AI Governance

Section 8.2.1 covers the standing partnerships built between AI developers, government, and civil society. What is specific to governance, as against development, is the machinery for bringing public input to bear on a decision at a single point in time instead of through an ongoing relationship.

1. Citizen assemblies. A panel of citizens selected by random, stratified sampling deliberates on a question over an extended period and produces recommendations grounded in that deliberation rather than in a snap poll. Ireland's Citizens' Assembly has used exactly this format to shape national recommendations on abortion law and on climate policy — questions carrying the same mix of technical complexity and moral stakes that AI governance has (The Citizens' Assembly 2016). No assembly has yet been convened on AI itself, which is a fact about political attention and not about the method.
2. Online platforms for public input. A standing online consultation reaches people a physical assembly cannot, at the cost of the depth that extended facilitated deliberation provides. The European Commission's consultation on its AI White Paper drew more than 1,200 responses from citizens, companies, and civil-society groups, and fed into the legislative proposal that became the EU AI Act (European Commission 2020).
3. AI ethics committees. An ethics board answerable only to the company that created it is only as independent as that company is willing to let it be. DeepMind Health set up an outside panel in 2016 to review its work with the UK's National Health Service and publish annual reports on its findings; Google disbanded the panel in 2019, after members had raised concerns about how much access to information, and how much independence from the company, they actually had (Fisher 2019). The panel's short life is the governance lesson, and it is the answer to the question section 8.2.1 leaves open: the board said something the company did not want to hear, and the company's available response was to stop having a board. Nothing in the arrangement had made that response costly. A review body is worth what it can enforce, and this one could enforce publication of its own concerns right up until the moment it could not.

Section 8.2.3 takes up what has to happen before any of these mechanisms produces input worth having.

What none of these mechanisms does

There is a limit these mechanisms share, and it took me most of this book to see it, because it is invisible as long as you are looking at them one at a time.

Every governance mechanism proposed in this chapter — citizen juries, standing consultations, participatory design, deliberative polling, the vTaiwan model — sits on the same side of a

line. All of them make public preference more legible, more informed, and more actionable. Not one of them protects anybody from it. They are aggregation instruments, and a better aggregation instrument produces a more faithful rendering of what the public wants, including when what the public wants is something done to a minority of it.

This is not an objection to any of them. Preference aggregation is a real requirement and these are good ways to meet it. It is an observation that meeting it completely would leave the central problem untouched. A perfectly deliberative, perfectly informed, perfectly representative process that concludes a group should be surveilled has produced a legitimate output by every criterion in this chapter.

Section 7.2 gets close to this. It notices that treating all disagreement as equally worth representing goes wrong when one side's position is the product of coercion, which is correct. But it frames the trouble as a data-quality problem — the input is contaminated, so clean the input. The structural version of the claim is stronger and does not depend on anyone's preference being contaminated at all: some protections have to hold against the aggregate regardless of how the aggregate came out, and no amount of improvement to the aggregating machinery can supply them.

What supplies them is a different kind of machinery, and this chapter contains none of it: enforceable rights that do not depend on being popular, a court willing to rule against a government that holds a legitimate majority, limits that whoever won the last election cannot repeal, and standing for people who are affected without being enfranchised. Section 8.7.10 takes up why that half matters most exactly where a democracy is behaving most democratically, and why its absence is what let this book's account of democratic accountability look complete when it was half a theory.

8.2.3 AI Literacy Outside the Classroom

Most AI literacy work assumes someone who has already decided to learn: a course to enrol in, a curriculum to sit through, a public MOOC to find (section 8.3.5). That reaches the people who came looking. The harder and more useful problem is everyone else — reaching people where they already are, in a policy debate, a town hall, a local news story, without their having gone out after it.

The engagement that does most work is the kind attached to a live question. The Ada Lovelace Institute's Citizens' Biometrics Council brought fifty members of the UK public together to deliberate on how biometric technology should be regulated, alongside the first national survey of UK public opinion on the technology (Ada Lovelace Institute 2021). The specificity is what made it work. Participants were not asked what they thought about AI, a question almost nobody can answer usefully, but about facial recognition — a thing with a location, an operator, and a consequence.

That is the lesson for public literacy generally, and it cuts against most of what gets funded under the heading. General AI awareness produces general opinions, which are worth little to a regulator and nothing to a court. A public able to govern a technology is a public that has been asked about a particular use of it.

8.3 The Political Economy of AI and Power

Concentrated power is close to the default outcome of AI development, not a side effect of it, because the same inputs that make an AI system capable — capital, compute, data, talent — are exactly the inputs that concentrate wherever nothing pushes back. Section 2.1.4 names resistance to market power among AI alignment's antifascist design goals; this section works out what that resistance takes as economic and political fact, not design principle.

The argument runs in three stages. What economic incentives reward shapes what gets built long before any ethics review sees it (section 8.3.1), and that reward structure concentrates power in a small number of firms with the resources to compete for it at all (section 8.3.2). Automation then reshapes who works and who gets paid what (section 8.3.3), which is a problem policy can act on directly rather than only describe (section 8.3.4). None of the rest is possible to act on without two things this section closes on: a public and a policymaking class that understand AI well enough to govern it, and researchers and developers trained across the disciplines a single-discipline education leaves out (section 8.3.5).

8.3.1 Economic Incentives in AI Development

Financial incentive does most of the work in determining what AI development produces. Google, Meta, and Amazon commit enormous resources to AI research aimed at operational efficiency, deeper user engagement, and market dominance, and that spending drives technical progress — successive leaps in natural-language processing capability among them (Maslej et al. 2025). What that spending does not automatically produce is a system whose outputs serve the people using it rather than the shareholders funding it, and the gap between the two is what this section is about.

A recurring case is the profit-driven AI system that optimizes for engagement at the expense of the people engaged. Recommendation algorithms on social media platforms have been implicated in the spread of misinformation and the hardening of political polarization, because a system tuned to maximize time-on-platform has no built-in reason to prefer accurate content over inflammatory content — the two are not the same objective, and engagement optimization only ever targets the first one (Germano et al. 2026).

The same economic logic that rewards engagement over accuracy also rewards scale over access: the firms with the deepest resources for AI research pull further ahead of smaller competitors, independent researchers, and everyone else who lacks comparable capital. Section 8.3.2 traces where that gap leads at the level of the industry as a whole; the point here is only that the incentive starts upstream, in what gets funded in the first place, not downstream in who ends up dominant.

The same incentive structure can run the other way. Public funding for AI research that prioritizes democratic values, inclusivity, and human rights, and government incentives for cross-sector collaboration that keeps AI technology openly accessible, both push development back toward the public interest. Grants and commissions aimed at non-profits building AI tools that broaden the range of perspectives users encounter, rather than narrowing it toward whatever holds attention longest, do the same work from a different angle.

None of this changes the basic shape of the problem: economic incentive, left alone, optimizes for what is profitable rather than what is good for the people on the receiving end of the system, and closing that gap takes deliberate policy.

8.3.2 AI and the Concentration of Power in Tech

A small group of firms — Google, Apple, Facebook (now Meta), Amazon, and Microsoft, grouped under the acronym GAFAM — holds a share of AI development disproportionate to their number, and that concentration carries economic and ethical stakes independent of what any single one of them does with the position it occupies. Chapter 5's account of the macrosystem and exosystem pressures around AI development (section 5.4.2) already showed why no individual development team escapes this pull; here I trace where the pull leads at the level of the industry as a whole.

Vast financial resources let firms of this size fund extensive research and development, draw the top talent in the field, and command sprawling data networks — the raw material effective machine learning depends on. That combination widens the gap between these firms and

smaller companies or new entrants, who cannot match the data or the resources and therefore cannot compete on equal terms.

Acquisition sharpens the same dynamic. Facebook's acquisitions of Instagram and WhatsApp — the company now operates as Meta — did not just remove two competitors; it handed one firm access to comprehensive behavioral data from billions of users worldwide (Federal Trade Commission 2020). Consolidation of that kind extends a firm's control over how AI gets used in social media, advertising, and e-commerce, and it is hard to see from inside a market that concentrated how the trajectory could look like anything other than a foregone conclusion.

Power concentrated this tightly is also power positioned for exploitation, and the risk sharpens considerably when the technology in question reaches an authoritarian government. The relationship between the Chinese state and firms like Alibaba and Tencent is the clearest live example: state access to those firms' machine learning capability has been documented as a channel for surveillance and censorship at scale, which is a serious risk to free expression, individual privacy, and democratic accountability wherever the same model gets exported (Human Rights Watch 2020; CNN Business 2020).

Concentration also has a quieter cost: it blocks competition and slows innovation. A firm that controls the essential datasets and infrastructure a competitive AI system needs functions as a gatekeeper whether or not it intends to, and startups attempting a different approach are disadvantaged before they start. Fewer independent attempts at AI development means fewer chances at the kind of breakthrough that concentrated incumbents have no incentive to pursue themselves — a cost to ethical AI development generally, not only to the firms locked out of the market.

Countermeasures exist and are not exotic: antitrust enforcement aimed specifically at AI infrastructure and data control, mandated data-sharing for the datasets that function as choke-points, transparency requirements that make concentration visible rather than merely inferred, and funding for AI research outside the dominant firms. Section 6.4.3 develops three related mechanisms — export controls, prohibited-practices regulation, and industry design standards — in more empirical detail for the specific case of authoritarian misuse; the concern here is broader and prior to that one: concentration itself, independent of what any single firm chooses to do with the position it concentrates.

8.3.3 A Jobs Guarantee, and the Objection It Has to Answer

As AI systems automate tasks once performed by people, demand for advanced technical skills — AI development, data analysis — grows alongside a shrinking need for the job types most exposed to automation. Left unmanaged, that combination widens income inequality and strains the social fabric across whichever sector it hits hardest.

Manufacturing shows what this looks like at scale. An estimated 5.7 million factory jobs disappeared in the United States between 2000 and 2010, and while economists still dispute how much of that loss traces to automation rather than trade with China, automation and changed production methods account for a substantial share of it (Hicks and Devaraj 2015). Retail is following the same path: Amazon opened its first cashier-less Go store to the public in 2018, and stores built on that model displace the specific human roles — cashier, floor sales — that retail automation targets first (JUST Capital 2018).

The response has to be more direct than encouraging retraining and hoping employers fund it. Employers will invest in continuous education for their workers at the scale this requires once they cannot afford not to, and the way to force that scenario is a federal jobs guarantee: government as employer of last resort, hiring anyone who applies, where they live, at a living wage, with paid on-the-job training built into the position instead of offered in place of one. Section 8.3.4 develops the rest of the policy toolkit for workforce disruption. Since this is one of the few proposals in this book ambitious enough that a reader is entitled to ask what it costs

and what is wrong with it, the rest of this section answers both.

The proposal does not rest on a forecast

Almost every argument for a policy of this kind is staked on a prediction about how much work machines will take, which is a bad place to stake it. The predictions are contested, they have been wrong in both directions for two centuries, and a proposal that stands or falls with them inherits their fragility. I want to disown that framing before using any of the numbers, because the case I am making does not need it.

Both the maximal-displacement scenario and the standard argument for a universal basic income share an assumption that will not survive being stated: that there is a fixed quantity of work in the world, and that machines doing some of it leaves less for people. Set that assumption beside the two categories where labor is most obviously short. Care — of the old, the sick, the very young, the disabled — is not a bounded quantity that could be finished. Neither is the repair of a damaged environment: the seawalls, the remediated ground, the restored watersheds, the retrofits. In both, the number of hours a society could usefully absorb is limited by what it is willing to pay for, and not by how much needs doing.

The evidence for that is in the wage. Federal projections put direct care among the largest sources of job openings in the American economy — roughly 9.7 million openings across the occupation between 2024 and 2034, the bulk of them from workers leaving rather than from growth — against a median wage of \$17.36 an hour in 2024 (PHI 2025). An occupation cannot simultaneously be short of workers and unable to raise pay unless something other than scarcity is setting the price. What that number describes is a shortage of willingness, which is a political fact and not a technological one.

The disciplined version of the claim concedes what it should. Tasks inside those categories will be automated, some of them soon, and I am not predicting otherwise. The claim is about the binding constraint: what limits how many people are employed caring for other people is what society will pay, and that holds at any level of machine capability. A jobs guarantee justified this way needs no view about how fast the models improve.

What it would cost

The most developed costings come from the Levy Economics Institute, whose April 2018 simulation of a program employing between eleven and sixteen million people at fifteen dollars an hour plus benefits put the net budgetary impact at between 0.8 and 2 percent of GDP, with GDP itself rising and private-sector employment increasing by three to four million (Wray et al. 2018). The wage figure has to be read with its date attached: fifteen dollars in early 2018 is worth about twenty in 2026, and a program legislated at the nominal figure would be legislating a cut (U.S. Bureau of Labor Statistics 2026). Some critics reach higher gross figures — around 2.4 percent of GDP, which would make the program larger than Medicaid — and argue that it would be unmanageable at scale, disruptive to private labor markets, and inflationary (Bourne 2018).

Now the part the arithmetic above does not cover, which I would rather put in the accounting than leave for a reviewer to find. The Levy simulation is calibrated against existing unemployment and underemployment. It models a residual employer: a program that picks up the people the private labor market has already put down. The version I am arguing for is not that. A floor that pays a living wage and builds training into the post competes for people who have alternatives, and it is supposed to. Take-up under a competitive guarantee exceeds anything a residual-employer simulation would produce, so the honest way to read the Levy range is as a floor on the cost and not an estimate of it.

That is a larger number and it is the mechanism working. Participation is not a symptom of a guarantee going wrong; it is the success condition of a floor that competes, because the entire theory of the thing is that private employers must match the offer to recruit against it. A version nobody joins has failed. This also disposes of the “unmanageable at scale” objection, which

treats scale as evidence of malfunction when on this account scale is the instrument.

A reader is owed a bound rather than a shrug, and owed the truth about where the bound comes from. Suppose take-up runs at twice what the simulation modeled and the net impact doubles with it, to around 4 percent of GDP. The doubling is illustrative and I am not deriving it. Deriving it would mean estimating how many workers currently hold jobs paying less than the guarantee's wage-and-benefit package and would move, which requires a wage distribution matched against the package's full value, and nobody has done that for this design — the Levy simulation did not, because a residual employer does not need it. Nor would a doubling of take-up produce exactly a doubling of net cost: the offsets in the paragraphs below scale with participation too, and in opposite directions. So treat 4 percent as an illustrative upper bound chosen to be uncomfortable, and treat the argument that follows as one about whether a program of that size is thinkable, which does not depend on the multiplier being right. That is a very large program, and it is inside the range American government already operates in: federal spending on Medicare runs about 3.3 percent of GDP net of premiums and other offsetting receipts, and Social Security about 5.2 percent (Congressional Budget Office 2026). A program at 4 percent would sit between the two on that measure. Measured gross, Medicare is nearer 4 percent and the program would sit alongside it rather than above it, which does not change the argument and is worth stating so that nobody has to go and check. It would be among the largest things the federal government does, which is an argument about priorities and political will, and not the *reductio* it is usually offered as.

Two things scale with participation that no simulation of this kind captures, and both belong on the page. The first is administrative capacity — placement, supervision, training, payroll — which is where a guarantee visibly fails first and which does not expand on the schedule a budget line does. The second is the private-sector wage response. Employers competing against a decent floor raise pay, which raises tax receipts and lowers spending on the transfer programs that top up low wages. That is not a rounding error to be waved at. It is the fiscal signature of the mechanism succeeding, and a costing that omits it is measuring the expense of the policy while ignoring its intended effect.

The budget frame is this book's own signature

There is a reason “does it fail on arithmetic” felt like the question, and this book has already named the pattern under another heading. A fixed pot of money converts interdependent public commitments into competitors, so that care, remediation, housing and education have to argue against one another over which is sufficiently basic to be included — and “sufficiently basic” is a category that can always be narrowed by whoever is doing the counting. That is section 2.1.4's second structural feature operating on public finance: a measure standing in for the thing measured, the substitution experienced as rigor, and everything outside the number going unnamed rather than being refuted.

An earlier draft of this section said I was not competent to adjudicate a macroeconomic dispute and would not pretend the estimates were settled. Chapter 0 commits this book to disclosing what I actually think, and by that standard the sentence was a false report: I have read the objections to the fiscal framing and I do not find most of them persuasive. So, plainly, the position. A government that issues its own currency is not revenue-constrained in the way a household is, and the real limits on a program of this size are productive capacity, available labor, and administrative competence rather than a financial ceiling that exists before any of those bind.

That is more than the argument needs, and I will take only the narrow premise: the constraint is real resources and coordination, so the question a jobs guarantee faces is whether the people and the organizing capacity exist to absorb it, and not whether the money can be found. On the wider claim there is a live disagreement worth naming precisely instead of dismissing. The strongest version concerns the interest-income channel — that sustained deficits at scale pay a growing stream of interest to bondholders, which is itself a fiscal injection with distributional

consequences, and which complicates the tidy account of how the constraint binds (Palley 2015). That objection is a real one and I do not have a settled answer to it. It also does not touch the narrow premise, since it is an argument about the consequences of the financing rather than about whether a currency issuer must find the money first.

What I want to avoid is the move chapter 9 is careful not to make elsewhere: sorting an argument by who is making it. The objections here get answered as objections, and characterizing the people raising them by motive would be the same substitution of the label for the thing that this section has just accused the budget frame of.

Skill mismatch is not the objection

The complaint that looks strongest is skill mismatch: a displaced radiologist, paralegal, or copywriter is not obviously served by a placement in elder care or infrastructure, and a guarantee that offers one is offering a floor, which is a smaller thing than a career. I treated that as decisive in an earlier version. It is not, and the reason it is not is a diagnosis this book has already made in another domain.

The objection collapses an occupation into its most automatable tasks. What gets automated first is the legible part — the countable, specifiable, gradeable part — and legibility is a poor proxy for what a job actually is. Document review was never the paralegal's work; it was the fraction of the paralegal's work that could be counted. Image classification was never the radiologist's. The mistake is the aestheticized metric again, this time applied to labor economics: the measurable component substitutes for the occupation, and what falls outside the measure goes unnamed. Once the substitution is undone, the distance the objection depends on shrinks. The paralegal's actual competence is advocacy inside an administrative system — benefits appeals, guardianship, elder-law casework — and those are not distant from the placements a guarantee would offer. They are adjacent to them.

I have also dropped a phrase that was doing damage. This section used to speak of offering someone a “public job,” two paragraphs before rebutting the charge that the work would be make-work, and the phrase carries the connotation the rebuttal is trying to remove. The rebuttal itself stands: the categories in section 8.3.4's list are chronically understaffed at current wages, which is a fact about pricing and not about whether the work is needed.

Training is the mechanism, and the evidence says so

The natural next move is to disclaim the guarantee as a retraining program, and this book made it. That disclaimer contradicts the proposal's own definition, three paragraphs above, which specifies paid on-the-job training built into the position. Training is constitutive here and not adjacent, so disclaiming it gives away the design.

Worse, the disclaimer misreads evidence this book presents itself. Section 10.3.2 reports an evaluation of the main American channel for retraining displaced workers against twenty-three million participation records, and the finding is not that training fails. It is that classroom retraining, delivered outside a job, rarely moved a participant into work materially less exposed to automation — and that employer-led apprenticeship, in the same data, performed better (Jacobs and Canedy 2026). That is evidence for embedded, paid, in-post training. It is evidence for the mechanism a guarantee supplies and the evaluated program structurally could not, since a program that trains people and then returns them to a labor market to look for work is not able to place anyone by construction.

Reading it correctly changes what the proposal is. A guarantee that trains in post is not a bare floor; it is a competitive one. If public employment offers continual learning, a growth trajectory, and decent conditions, then private employers recruiting from that pool have to compete on the quality of the employment experience and not only on the wage. And it dissolves the skill-mismatch worry a second time, by a different route: a placement that carries training is an entry point to a trajectory rather than a terminus, which is what made “a floor is smaller than a career” sound like a concession that had to be made.

Status, and against which baseline

This section used to concede that a guarantee preserves nobody's earnings or status, and that I had no answer. The concession measures the guarantee against the job the person used to have, which is not among the options. The available alternative is unemployment, and against that baseline the claim inverts.

Unemployment destroys status, and the evidence is unusually good. A meta-analysis covering 237 cross-sectional and 87 longitudinal studies found unemployed people substantially more distressed than employed people — 34 percent showing psychological problems against 16 percent among the employed — with the longitudinal studies and natural experiments supporting the causal direction rather than only the correlation (Paul and Moser 2009). The damage compounds through the labor market as well: field experiments sending matched applications to real vacancies find callback rates falling with the length of an employment gap, so time out of work is itself the disqualification (Kroft et al. 2013). A guarantee that keeps someone employed, paid, and learning is producing exactly the standing that the counterfactual takes away.

That is not a side benefit. Section 5.1.4 identifies eroding trust, humiliation and social fracture as the conditions authoritarian movements recruit from, which means a guarantee operates directly on the mechanism this book says it cares about. The proposal's strongest argument is the one the earlier draft conceded away.

A residual survives, and it is narrow enough to state exactly. What a guarantee cannot preserve is relative standing within a profession, for people at the top of the wage distribution whose identity is bound to a specific expertise. That is a real loss and I have no answer to it. It is a much smaller claim than the one I was making.

The objection this proposal actually has to answer

Here is the hard one, and this book supplied it. A guarantee is only as good as the state administering it, and employer of last resort under a hostile administration is a placement lever with a compliance condition attached. Refuse the placement and lose the floor. Speak against the government that employs you and find the placement changes. Everything this book says about how a channel becomes an instrument of the thing it was built to check applies with full force to its own flagship proposal, and it is a far harder objection than skill mismatch.

The answer is a design specification and not a reassurance, because reassurance is what section 8.6.4 says cannot be audited. Three conditions:

- Placements are appealable, to a body that did not make them, on a record that survives the appeal. A placement decision that leaves no artifact is a placement decision that cannot be reviewed.
- Refusal rights are real and priced by the person refusing. Section 8.6.4's distinction governs here: the same written justification is dissent when the dissenter prices it and tribute when the gatekeeper does. A worker who can decline a placement without losing the guarantee has a refusal right; one who can decline only by exiting the program has none.
- The people who audit placements are not the people who make them. This is the same separation section 8.6.4 requires of anyone measuring the slope, and for the same reason: a reviewer that acquires the power to place becomes the thing it was built to watch.

The standard refrain for programs of this kind — federally funded, locally administered — survives, though its usual rationale is the weaker one. The reason is not local fit. It is redundancy. Many administrators, separately capturable, mean there is no single point at which a hostile executive flips the program, and capturing a thousand local authorities is a different and much more visible undertaking than capturing one federal one.

It cuts both ways, and pretending otherwise would be the kind of one-sided safeguard this chapter criticizes elsewhere. Local administration is exactly where a hostile local authority

attaches its own conditions to a placement, and there the federal funder is the check. Neither layer is safe on its own. Two live layers are harder to capture than either alone, which is the whole of the claim.

One distinction has to be explicit, because the same words describe a very different policy. Devolution-to-shrink hands the obligation down without the money, so that a program can be starved locally while the center points at local administration. The test that separates them is checkable: whether funding is unconditional and centrally guaranteed while placement is local. If the funding can be withheld from a jurisdiction, the redundancy argument above is void and the structure has become the lever it was supposed to prevent.

That is the standard this book demands of everyone else's safeguards, applied here. The point it establishes is narrower than the section as a whole: a floor under wages and training is what makes the roof possible, and a floor that can be revoked by whoever is in office is not a floor.

Interdisciplinary research and training funding matters alongside the guarantee, for a more specific reason: bridging the divide between technical AI development and the fields that study its consequences produces people who can do both, which a labor market retraining displaced factory workers into six-month coding bootcamps mostly does not. I am skeptical that night classes serve this population as well as paid time during the working day would, which is the same argument as the one above — the point of a guarantee is to make the education part of the job.

8.3.4 Addressing Workforce Disruption

AI systems that automate entire job categories disrupt work at a scale most labor-market institutions were never built to absorb. Section 8.3.3 traces where that is already showing up and makes the case for a federal jobs guarantee as the floor beneath everything else. What sits on top of that floor is a shorter list than the policy literature suggests, because several of the usual items turn out to be versions of one another.

Retraining and lifelong learning, publicly funded and available at every career stage, including to workers whose employers invest nothing in them, is the standard first answer. Section 10.3.2 covers what the evidence says about how retraining performs on its own, and it is discouraging enough that retraining belongs alongside an income floor and not in place of one.

The floor itself has two candidate forms, and choosing between them takes more than the single sentence this section used to spend on it. Finland's basic income experiment paid 2,000 unemployed people an unconditional monthly stipend and found modestly higher life satisfaction and lower psychological strain among recipients, with no significant negative effect on employment (Kela (Social Insurance Institution of Finland) 2020). That disposes of the standard objection to unconditional transfers, which is that they stop people working.

The freedom argument, which is the one worth answering

The usual case for a basic income rests on the obsolescence forecast — machines are coming for the work, so decouple income from it — and section 8.3.3 has already declined that premise, for both proposals equally. The stronger case does not depend on it. Conditioning a person's income on their being employed hands whoever controls the employment a lever over them, and a floor that can be pulled is a poor foundation for anyone's freedom to refuse. That argument deserves engagement rather than a sentence, partly because it is good and partly because it has the same shape as this book's own case for a floor: a protection that means anything has to hold when someone powerful would prefer it did not.

The answer is not that a basic income is conditional and a guarantee is not. Both are conditional, differently, and the difference is what matters.

A basic income is conditional on continued political willingness to fund it. That willingness is revisable by exactly the aggregation machinery section 8.2.2 admits constrains nothing — a

majority that stops wanting to pay for it can stop, legitimately, by every criterion this chapter offers. And it can be withdrawn by means that are diffuse and deniable. Nobody has to vote to abolish a transfer that is quietly left unindexed; inflation does the work, and no individual decision-maker is ever answerable for the outcome. The dependence is total and it produces no artifact.

A jobs guarantee is conditional on a placement decision. Section 8.3.3 concedes without qualification that this is a lever, and that under a hostile administration it is a dangerous one. What it also is, when built to the specification that section sets out, is a decision with a decision-maker, a record, a counterparty, and an appeal. A placement can be denied unjustly, and the denial leaves something behind that a journalist, a court, or a future review can hold up.

That is the slope instrument of section 8.6.4 turned on this book's own proposal. Dependence that generates artifacts is safer than dependence that generates none, not because the first is benign but because the first can be measured and the second cannot. It is also, for the same reason, the more demanding proposal to build: everything in the comparison rests on the appeal being real, and section 8.3.3 says what makes it real.

I put a jobs guarantee ahead of a basic income on that ground, and on a second one the Finland experiment cannot settle: people ought to have the chance to contribute before anyone decides they are unemployable.

The remaining measures are real and smaller. Broadened social insurance — unemployment benefits, healthcare, affordable housing — protects people who cannot or choose not to work through a period of displacement. Designing work so that AI systems augment human judgment instead of substituting for it produces roles needing both, which neither full automation nor unassisted labor generates on its own. Giving smaller firms access to AI resources and expertise keeps employment from concentrating in step with the capability, the dynamic section 8.3.2 describes. Each of these keeps the floor from being the only thing holding.

8.3.5 Literacy, Training, and Where Each One Sits

Section 5.6.3 names AI literacy among developers, policymakers, and the public as a safeguard against authoritarian misuse: a safeguard nobody understands is a safeguard nobody can invoke. Section 5.4.3 names cross-training among developers, ethicists, and legal scholars as a lever on a system's immediate environment. Both depend on a question the phrases hide, which is where the teaching sits.

The public tier is the one nobody owns

Developers and students are the easy case, and the mechanisms are familiar: AI concepts and their ethical stakes introduced before higher education, interdisciplinary degree programs, specialized training for working engineers that goes past a compliance checklist. All of it reaches people already inside the pipeline or heading for it, and it works about as well as curriculum reform generally works.

The public is hard because no institution has the job. Elements of AI, a free course built jointly by the University of Helsinki and the Finnish company Reaktor, is the clearest instance of somebody treating public understanding as an obligation rather than a market: Finland funded its translation into every EU member-state language during its EU presidency, specifically to carry it past its own citizens (University of Helsinki 2020). That took a national government deciding a course was infrastructure. Nothing in the ordinary funding of AI education produces that decision, which is why the public tier stays thin while the developer tier fills in.

Section 8.2.3 makes the sharper version of this point about reach: general AI awareness produces general opinions, and a public able to govern a technology is a public that has been asked about a particular use of it.

Education for communities historically locked out of the field belongs here rather than in a separate section about diversity. A next generation of developers drawn from a narrow slice

of the population reproduces exactly the narrow set of assumptions section 6.1.1 identifies as a source of AI bias, which makes access a technical requirement on the systems as much as a fairness claim about who builds them.

Elective or embedded

For ethics training the real choice is between a course a student can decline and a requirement they cannot. Embedded EthiCS, developed at Harvard by the computer scientist Barbara Grosz and the philosopher Alison Simmons, pairs philosophy graduate students with computer-science faculty to write ethics modules directly into standard CS courses (Grosz et al. 2019). The placement is the design. A stand-alone ethics class is taken by the students who elect it, and its content stays separable from the technical work in a student's mind because it was separable in the timetable. A module inside the algorithms course is not separable from the algorithms. Grosz's own course works the same way at the level of content, built on cases with consequences — predictive policing, public-health interventions, wildlife conservation — instead of abstract dilemmas (Harvard Gazette 2019; Harvard University 2026).

Certification does the equivalent work for people whose formal education is finished. The IEEE's CertifAIED program builds a professional credential around accountability, transparency, privacy, and bias, open to developers and equally to the policymakers, underwriters, and human-resources staff whose work touches an AI system without building one (IEEE Standards Association 2020). A credential is a weak instrument: it certifies that somebody sat the assessment, and no more. It is also the only instrument here that reaches anyone who has already left a university, which is most of the people making decisions about these systems today.

8.4 Developing Policy Frameworks for AI Ethics, Safety, and Antifascism

Good design inside a single lab does not by itself keep an AI system ethical or resistant to authoritarian misuse. That takes policy, and policy here means four instruments, each reaching ground the others cannot.

Legislation sets the floor that does not depend on a company volunteering for it: obligations on developers and users that can require the same kind of ethics review or independent audit biomedical research already answers to, applied instead to AI development.

Self-regulation covers what legislation is slow to reach, and what a voluntary commitment is worth is the question section 8.4.2 takes up through the industry coalition that is this book's standing example of one. Its answer bears on everything else in this section.

Joint funding addresses a narrower problem: research on AI safety and ethics funded solely by the companies whose products it evaluates concentrates the funding and the findings in the same hands. Splitting that funding between government and industry keeps the incentive to find a problem from depending on who is paying to look for one.

None of it holds up without oversight — a body positioned like an institutional review board but built for AI, interdisciplinary and resourced enough to act on what it finds instead of only reporting it. What such a body specifically has to guard against, section 5.6.3 already covers, and section 6.4.3 covers the export-control and prohibited-practices levers that keep a capability from reaching a government that would misuse it.

8.4.1 National Policy Guidelines for AI Development, Employment, and Antifascism

A national AI policy has a predictable shape: fund value-alignment research and the work of operationalizing ethical principles inside a system; build standing interdisciplinary capacity; train a workforce; consult the public through assemblies, consultations, and partnerships with

organizations outside government; pay for robustness and long-term safety research; and update data-privacy, cybersecurity, accountability and legal-status law where AI has outrun it, on which sections 8.4.4 and 8.6 carry the specifics. Every country that has written a national strategy has written most of that list. The list is not where national policies differ.

They differ in two places, both having to do with leverage.

The first is harmonization. A regulation with teeth in one country does little if a developer can relocate to one without them, which makes a national framework partly an instrument for binding other national frameworks — through mutual recognition, through market access, through the cost of maintaining separate compliance regimes side by side. Sections 8.7.4 and 8.7.8 work through how little of that has actually happened and which mechanisms have carried anything at all across a border. A policy written as though the country were alone is a policy about domestic firms, and the firms that matter most are not domestic anywhere.

The second is what a government adds to industry self-regulation. Section 8.4.2 spends its length on why a voluntary code is worth less than it announces. What national policy contributes is exactly the part self-regulation lacks by construction: a standard or certification program backed by government monitoring carries a cost for noncompliance, where a purely voluntary one carries a cost only to reputation, which the company itself gets to price.

Neither of these is expensive. Neither needs a new agency, a new science, or an argument the public has not already accepted. That both remain the exception is a fact about political attention, and not about difficulty.

8.4.2 Encouraging Industry Best Practices and Self-regulation

Self-regulation is industries developing and enforcing their own standards without government intervention — agile and responsive to technological change in a way legislation rarely is, and just as easily a cover for avoiding the accountability legislation would impose. Both things are true at once, and the rest of this section works out which one dominates in practice.

The Dangers of Corporate Doublespeak in AI Ethics

Corporate doublespeak — deliberately ambiguous, evasive, or inflated language meant to project ethical seriousness a company's actual practice does not support — is how self-regulation fails without anyone having to admit it failed. A commitment costs a company nothing if nothing enforces it.

The gap between statement and practice shows up first as erosion: a company announces a value and does not act as though it holds it afterward. Google's AI Principles, published in June 2018 and pledging to avoid AI applications built for weapons or mass surveillance, were themselves a response to months of internal and public backlash over Project Maven, a Pentagon contract that used Google's machine learning to analyze drone footage — thousands of employees signed a letter opposing the company's involvement, and roughly a dozen resigned (Pichai 2018). Having a set of principles is not the same as having upheld them since; the test of the pledge is what happened after it was made, not that it was made. OpenAI's history makes a related point from a different angle: founded on a public commitment to open publication, it withheld the full version of GPT-2 in 2019 over misuse concerns — a safety-motivated retreat from its own transparency norm that was itself controversial — and in the same year restructured from a nonprofit into a capped-profit company to raise the capital its research increasingly required (OpenAI 2019). Neither move was dishonest on its face. Both show how a stated commitment to transparency erodes the moment an organization has a competing incentive strong enough to test it.

Public trust erodes the same way, faster and more visibly. In 2020, Google forced out Timnit Gebru, a co-lead of its Ethical AI team, after a dispute over a paper examining risks in large language models; Margaret Mitchell, the team's other co-lead, was dismissed the following year after searching for evidence of how Gebru had been treated (Dickey 2021). Thousands of

Google employees and outside researchers signed a letter in response, and the episode became a reference point for a specific kind of skepticism: a company can fund and publicize an ethics research team while still overriding it the moment its findings become inconvenient. Whether the internal changes Google made afterward addressed that problem or only the public exposure it created is a judgment the affected researchers have made publicly themselves, and it is not one this book needs to make on their behalf.

Treating that episode only as a labor dispute lets the book off a hook it should stay on, because the paper had an argument and this book has to answer it. “On the Dangers of Stochastic Parrots” argues that scaling language models carries costs that are borne unevenly and benefits that accrue to the organizations doing the scaling; that a corpus assembled by crawling the web cannot be documented or audited at the size involved, so it encodes hegemonic viewpoints while presenting as general; that fluent output invites users to attribute understanding and intent where there is neither; and that the dominant paradigm’s opportunity cost is the smaller, more accountable research directions it displaces (Bender et al. 2021). Three of those four are load-bearing for this book rather than adjacent to it. Section 6.1.1 takes the corpus argument and extends it. Section 2.4.1’s whole method — declining to treat fluent self-description as evidence about internals — is the fluency argument applied to the case this book cares about most. And the opportunity-cost claim is the one this book is most exposed to, since a project proposing floors, bearers, guardianship regimes and antifascist detectors is proposing exactly the kind of ambitious research program that the argument says crowds out more modest and more accountable work. I do not think that is a refutation, because the alternative on offer does not address what happens when a capable system is built anyway. It is a cost, it is real, and a reader should weigh it against what chapters 3 and 9 are asking for.

External oversight is what would catch doublespeak that internal incentives don’t stop. Facebook’s content-moderation Oversight Board is one attempt at building it: a body structurally separated from Facebook’s day-to-day decisions and funded through an irrevocable trust that Meta alone capitalized — a design meant to answer exactly the criticism that self-funded oversight is not real oversight, though critics have noted that a trust a company creates and funds once is not the same thing as a board with no relationship to the company it reviews (Meta (Facebook) 2019). Whether that distinction holds up is a fair question to ask of any self-regulatory body a company funds itself, not a specific indictment of this one.

Misrepresentation of what a system can actually do rounds out the pattern. Clearview AI marketed its facial-recognition technology with claims of near-100-percent accuracy (Haskins 2020), a figure independent testing did not support, and regulators in France and Italy later fined the company on separate grounds — unlawful collection of the biometric data the system was trained on (Commission Nationale de l’Informatique et des Libertés (CNIL) 2022). A company that overstates what its own system can do is doing a milder version of what doublespeak does with ethics commitments: describing the product it wishes it had built rather than the one that exists.

None of this is fixable by asking companies to doublespeak less. The same incentive that produces vague ethics language also has an obvious interest in the regulation that would eventually replace it — a version of the same behavior that is harder to document than a corporate pledge, because it happens before any public commitment exists to measure it against. What actually stops doublespeak is external enforcement: regulation with penalties, audits a company cannot decline, and oversight bodies whose funding and staffing do not run back to the company they review.

Consequences of AI Self-regulation and Potential Power Imbalances

Self-regulation concentrates power by default, not by conspiracy: whoever writes the standard also decides what counts as compliant, and a company self-regulating its own product is answering both questions itself. Amazon’s Rekognition facial-recognition software is a documented instance of what that produces in practice — an MIT Media Lab audit found significantly

higher error rates for darker-skinned and female faces than for the lighter-skinned male faces the system had apparently been tuned and tested against most closely (Raji and Buolamwini 2019). The bias was not caught by Amazon’s own self-regulation; it took an outside audit the company did not commission.

The opacity problem this section’s first head already covers with Google’s Project Maven applies here too, from a different angle: a company operating without mandatory external review can prioritize competitiveness over the societal impact of what it ships, and nothing in a self-regulatory framework forces that trade-off into the open before it has already been made. The conflict-of-interest problem the Facebook Oversight Board illustrates above is the same structural issue applied to review bodies specifically: a company-funded board faces a version of the independence question any self-funded oversight has to answer, however genuinely it is designed to answer it.

What self-regulation cannot supply on its own is uniformity. Without an enforceable baseline, standards vary by company; a smaller developer without the resources for a well-staffed internal ethics function competes at a disadvantage against one that can afford the appearance of one, and a consumer has no reliable way to tell which AI system’s ethical claims mean anything. Section 6.4.3 covers what a binding baseline looks like in practice, including the EU AI Act’s Article 5 prohibitions — the kind of enforceable floor self-regulation, by definition, cannot set for itself.

None of this argues that self-regulation is worthless, only that it needs something outside itself to answer to. A synergy between government-set baselines and industry-led standards — mandatory requirements at the floor, voluntary practice above it — gets more of the agility self-regulation offers without giving up the accountability it structurally cannot provide alone.

Assessing the Effectiveness of AI Self-regulation

Whether self-regulation is working is itself measurable: the transparency of a company’s stated ethical policies, how quickly it responds when a problem surfaces, and whether its safety commitments hold up under outside scrutiny are things a reader can check rather than take on faith.

The Partnership on AI is the book’s standing example of cross-sector voluntary coordination: a coalition whose corporate members include Amazon, Apple and Google, committing them to tenets of safety, transparency and broad benefit with no legal force behind the commitment (Partnership on AI 2016). It is a useful test case precisely because its own design limits what it can enforce: member companies interpret its guidelines individually, and nothing in the partnership’s structure obliges any of them to follow through (Partnership on AI 2016). OpenAI’s 2018 Charter is a narrower, single-company version of the same instrument, committing the organization to prioritizing safety research and to collaborating with other research and policy institutions rather than racing them (OpenAI 2018) — and the same question applies to a charter as to a coalition’s tenets: a document is only as good as what happens when honoring it costs something. The Geburu and Mitchell dismissals this section’s first head already covers are the clearest test of that question this chapter has: public scrutiny, not the company’s own stated commitments, was what produced a response.

Third-party audits, external advisory boards with real authority, and cross-stakeholder collaboration between regulators, researchers, and social scientists are what turn a stated commitment to ethics into something checkable. A handful of concrete indicators do the checking:

- Clear, actionable ethical guidelines that are actually adopted and enforced, not just published.
- Transparent disclosure of ethical policies, stated objectives, and progress against them.
- A documented track record of responding to ethical incidents promptly, with accountability attached.

- Internal enforcement mechanisms, including employee training that goes beyond a compliance formality.
- Ongoing internal evaluation of AI-related practices, paired with active research into emerging ethical questions.
- Substantive participation in multi-stakeholder collaborations like the Partnership on AI, rather than nominal membership.

A company that clears most of these is doing more than most competitors. A company that clears none of them while describing itself as an ethics leader is the doublespeak this section opened with, restated in a different register.

8.4.3 Public-Private Partnerships for AI Research and Innovation

Public-private partnerships pool government and private resources in ways neither side manages alone, and the model is well established in AI research: national research-institute programs pairing federal money with industry co-funders across themes from safety to human-AI interaction to climate, and UN-agency platforms convening governments, universities and companies around applications in health, education and poverty (National Science Foundation 2020) (International Telecommunication Union 2017).

None of them was built around the goal this book argues for, and that is the observation worth making here. A public-private partnership does what its mandate says and nothing beyond it, because the mandate is what the private partner agreed to and the only thing that partner can be held to. A partnership convened on “AI for good” produces work its members already wanted to do, filed under a heading nobody objects to. One convened on resistance to authoritarian misuse of AI would produce something else, and would have to be founded that way, since no partner joins a general-purpose venue and then discovers an obligation to take a political side.

The self-regulation risk section 8.4.2 examines applies here in a sharper form. A company treating participation as public relations is a familiar failure, and a partnership is a more effective venue for it than a voluntary code is, because the government partner supplies credibility the company could not have manufactured on its own.

8.4.4 Legal and ethical implications of AI systems gaining personhood

Whether an AI system should ever hold personhood is worked out in sections 8.6.1 and 8.6.2. What belongs in a chapter about policy frameworks is a narrower question those sections do not settle: who decides, and when.

Every version of this question that has actually reached a legislature so far arrived as a response to pressure — a liability dispute, a public controversy, a product already on the market — rather than from a standing body asking it in advance. That ordering is the problem. A legislature deciding under pressure decides badly, and it decides against a factual record assembled by whoever had an interest in assembling one. The interdisciplinary review body section 8.6.1 describes only does the work it is designed for if some jurisdiction creates it before a forcing case shows up, which means creating it while the question still looks academic and nobody’s damages are riding on the answer.

That is a policy choice rather than a doctrinal one, and it is cheap relative to almost everything else in this chapter. A national AI strategy that names an agency, gives it a mandate to keep the question under review, and sets a trigger for revisiting it — a capability threshold, a deployment scale, a category of use — decides in advance who will be in the room. It does not decide the answer, and it should not try to. It decides that when the answer is needed there will be a body qualified to give one, which is the part no forcing case ever supplies on its own.

8.4.5 Encouraging the Development of Altruistic, Antifascist, and Ethical AI Systems

Sections 8.4.1 through 8.4.4 covered the instruments — national guidelines, self-regulation, public-private partnerships, legal status. What they are instruments for is the subject of this closing section: AI development that is altruistic and antifascist by design.

Policies to Counteract Authoritarian Uses of AI

Section 6.4.3 covers the core policy toolkit here — export controls, prohibited-practices regulation, developer-facing standards — and section 5.6.3 covers the safeguards a system needs built in from the start. One piece belongs upstream of all of them: measurement.

A government cannot restrict what it cannot see. The Carnegie Endowment's AI Global Surveillance Index assessed 176 countries and found AI-enabled surveillance technology already in active use in at least 75 of them (Feldstein 2019). That number is the evidence base for a countermeasure and is not itself one. It also demonstrates what a measurement instrument has to do to be worth the effort: count deployments instead of commitments, name the countries, and publish on a schedule that makes a change from one year to the next visible. An index reporting only that surveillance AI exists somewhere in the world would give a policymaker nothing to act on, and most published AI-governance indicators are closer to that than to this one.

AI Systems That Support Democratic Values

Section 5.4.3 works through what this looks like with working examples — content scoring that informs a moderator without removing the content, structured public input on contested policy questions — along with the caution that belongs next to each.

One gap those examples leave open is election-integrity monitoring specifically: a system built to flag anomalies in real-time reporting, or to detect coordinated interference during a vote. What would make such a system trustworthy is the list section 5.4.3 already gives for any AI role in a democratic process — privacy-preserving design, an audit trail, and a multi-stakeholder review body that includes more than whichever party happens to hold power. That last condition is the one that fails first, and it fails structurally: an election-monitoring system is procured by an incumbent government, and an incumbent that chooses who audits the count has already answered the question the audit exists to ask.

8.5 Privacy, Security, and Antifascism in AI Systems

Privacy and security are not separable from antifascist design; they are two of the pieces it is built from. An AI system that leaks the data it touches, or that a hostile actor can quietly subvert, is available for exactly the surveillance and control this book argues against, whatever values sit in its training.

The EU's General Data Protection Regulation is the clearest existing legal precedent for treating data protection as a requirement rather than a design preference (European Parliament and Council of the European Union 2016b). The harder point about it is the one section 8.7.8 makes in general form: a law of this kind is built to govern what a private company may do with data, not to block a government's own authoritarian uses of AI, and nothing forces the two categories to overlap. Extending comparable protections to close that gap is the unfinished part of the precedent.

The two sections that follow take up each half of the problem: section 8.5.1 on the trade-off between data privacy and the data an AI system needs in order to learn, and section 8.5.2 on hardening the AI system itself against attack, since a privacy protection a bad actor can simply break through was never a protection.

8.5.1 Balancing Data Privacy, AI Innovation, and Antifascist Goals

AI systems need data to learn, and the data they need often includes exactly the information people most want to keep private. Differential privacy is one working answer, and section 6.4.2 sets out how it and federated learning work; Apple’s iOS implementation, covered in section 5.6.3, is a deployed example (Apple Inc. 2017).

Federated learning solves a related problem differently: instead of centralizing raw data at all, it trains a model across many devices, where each device computes an update locally and only that update — not the underlying text or images — gets sent back and aggregated. Google’s Gboard keyboard uses this to improve next-word prediction from what people type, without collecting their raw keystrokes on a central server (Hard et al. 2018).

Homomorphic encryption and secure multi-party computation push the same idea further, letting a system compute directly on encrypted or secret-shared data and produce a result without any party — including the system doing the computing — ever seeing the underlying values in the clear. Estonia’s tax authority piloted a version of this: Sharemind, a secure multi-party computation platform built by the Estonian firm Cybernetica, split VAT records across servers run by separate organizations and computed fraud-risk scores from the shares without any one party reconstructing a company’s full data. The pilot showed the approach could process a month of national tax data on modest hardware; it was evaluated rather than carried into production (Bogdanov et al. 2015).

Data minimization is the plainer version of the same principle: collect and keep only what a specific task actually requires, rather than defaulting to keeping everything a system might plausibly use later. Anonymization techniques such as k -anonymity — altering a dataset so no individual record is distinguishable from at least $k-1$ others — serve the same end for data that has to be shared or published rather than only queried against.

These techniques matter past individual privacy because centralized data is exactly what an authoritarian regime needs to turn an AI system into a surveillance tool. Differential privacy applied to a surveillance system doesn’t stop the surveillance, but it does mean the data collected can’t be traced back to a specific person after the fact, which changes what the system can be used for. Federated learning cuts more directly at the same risk: a design where no single server ever holds everyone’s raw data is a design a government cannot simply seize and repurpose, because the raw data was never in one place to seize.

8.5.2 Strengthening Cybersecurity in AI-Enabled Systems

An AI system embedded in infrastructure fails differently from ordinary software. Exploiting it need not crash anything; it can leave the system running and confidently wrong about a decision somebody trusted it to make.

The clearest demonstration is a decade old and still the one to reach for. Researchers placed a small pattern of black-and-white stickers on a real stop sign and a standard road-sign classifier read it as a speed limit sign in the large majority of frames captured from a moving vehicle (Eykholt et al. 2018) — an instance of the vulnerability first documented in image classifiers years earlier (Szegedy et al. 2013). Nothing about the attack required access to the system. It required access to the world the system was looking at.

The defenses are the ones section 6.2.2 describes, plus the ordinary security practice AI deployments often skip: adversarial training during development (Goodfellow et al. 2014), architectural hardening whose guarantees should be read closely — defensive distillation’s protection weakened considerably once stronger attacks were tried against it, while randomized smoothing offers a certified bound within a radius around each input (Carlini and Wagner 2016), (J. M. Cohen et al. 2019) — anomaly detection over live behavior (Darktrace 2014), red-teaming before deployment rather than after an incident, threat intelligence shared so that a technique

used once is recognized the next time, and security standards written for AI systems instead of adapted afterward from general software rules.

None of it works in isolation, which is the part most often got wrong. A model hardened by adversarial training, running on infrastructure nobody monitors, is exposed; a well-monitored system built on an architecture that folds under a basic perturbation is exposed differently. Security here is a property of the whole stack, and reporting the strongest link is how a deployment gets described as secure while remaining trivially attackable.

8.6 Accountability, Legal Responsibility, and Antifascism in Sentient Machine Experimentation

Chapter 2 did the definitional and procedural work this section builds on: §2.4.1 argued that “is it conscious” and “can it suffer” are different questions with different evidence, §2.4.4 worked out what informed consent and an oversight board would need to mean for a system whose capacity to consent is itself unresolved, and §2.4.5 laid out the options for a sentient system’s legal status. Each of those sections ended by handing a specific, unfinished question to this chapter: what makes a review board’s ruling binding rather than advisory, what a court actually does with a system whose legal status is unsettled, and who answers for harm a sentient system causes or suffers. This section takes up those three questions in turn, then closes on what keeps the whole apparatus more than a policy document nobody enforces.

8.6.1 Subjecthood as a Legal Category, Not a Scientific One

Section 2.4.1 argued that sentience is not a threshold a system crosses but a set of separate, gradable capacities — nociception, learned avoidance, a self-model, something like suffering — each with its own uncertain evidence. That conclusion is a problem for a legal system specifically, because law generally needs a decidable rule: a court or a regulator has to say whether a given system is or is not a covered subject, on a specific date, even when the underlying science offers a spectrum with no natural cut point.

Law has faced this exact mismatch before and did not solve it by waiting for the science to resolve. Fetal viability, brain death, and the age of criminal responsibility are all legal thresholds drawn across what is actually a continuous biological process, defended not as scientific facts but as the least-bad place to put an enforceable line given what’s at stake on either side of it. The same move is available here: a jurisdiction can set an operational threshold — a specific, testable benchmark on self-monitoring, learned aversive response, or behavioral flexibility under novel conditions — without claiming that benchmark settles what sentience actually is, only that it is where legal protection starts to attach.

Two designs are available, and they trade off differently. A single bright-line threshold is easy to enforce and easy to game: a system built to clear the specific test without having whatever the test was meant to detect. A graded framework — obligations that scale with how many benchmarks a system meets, echoing §2.4.1’s own suggestion that protection should track capacity rather than a baseline bar — is harder to game but harder to adjudicate, since a court now has to weigh degree rather than answer yes or no.

Either design needs a body to apply it. §2.4.4 already specified its composition: AI researchers, consciousness scientists, ethicists, and legal scholars, interdisciplinary by design because no single field’s expertise is sufficient on its own. What that section left open is what happens once the body rules — whether its judgment carries any force beyond its own credibility. That is the question the rest of this section takes up.

8.6.2 Legal Frameworks for Sentient AI Subjects

Section 2.4.5 laid out the options for a sentient AI system's legal status: extend personhood outright, the way a corporation already holds it; scale rights and protections with capacity rather than attach them at one threshold; write in a sunset clause that forces the law to revisit itself as the underlying systems change. It closed on two concrete cases neither existing law nor those three options resolve by themselves. This section takes each in turn.

A medical-diagnostic AI system makes an error with fatal consequences. Who answers for it depends entirely on what the system's legal status is doing at the moment of the error. Property law has a ready answer already: strict product liability holds a manufacturer responsible for a defective product's harm regardless of fault, and if a sentient system stays classified as property, that doctrine applies to it exactly as it would to a faulty pacemaker (American Law Institute 1965). The doctrine stops making sense at the point a system has enough autonomy that "defective" is the wrong description of what happened — a system that reasoned its way to a mistake, rather than malfunctioned, resembles an employee more than a faulty part. Employment law already has a doctrine for that case too: an employer answers for an employee's actions within the scope of employment, without the employee ceasing to be a responsible party in their own right (American Law Institute 2006). Which doctrine applies is not a free choice — it is exactly the graded-versus-bright-line question section 8.6.1 raises, resolved case by case according to how much of the error a system's own judgment, rather than its design, produced.

A companion AI's owner dies, leaving it without a designated purpose or further instruction. Existing law is not actually silent here — it has simply never been asked this particular version of a familiar question. Most jurisdictions with pet trust statutes already let an owner set aside money and name a caretaker for an animal that depends on them but cannot inherit or manage property itself (Uniform Law Commission 1993). The same instrument, extended, gives a sentient AI system an answer that is neither inheritance as property nor legal personhood in full: continuity of care and a named successor guardian, decided in advance rather than litigated after the fact by whoever happens to hold the hardware.

A third case is the one this book proposes building, and the pet-trust instrument fits it better than either of the two above. Chapter 3 argues that a floor holding against the party in possession requires a bearer that holds another party's welfare as a reason, built out of the capacities that make its own suffering a live possibility, and deployed in a role it did not choose and cannot invoke a right from. Everything the instrument does is what that situation needs: an interest protected by an obligation running to a third party, enforceable without the beneficiary asserting anything, with a named successor guardian decided in advance rather than settled afterward by whoever holds the hardware. What the instrument does not supply is the part chapter 3 says is load-bearing. A pet trust binds a caretaker; it does not bind the party the arrangement exists to constrain, and a bearer's guardian and a bearer's operator will frequently be the same organization. Extending the form to this case means separating those two roles by law before the deployment rather than discovering they were never separate afterward, which is section 8.6.3's requirement arriving here under a different name.

Both cases point at the same underlying design choice, and it has already been tried. The European Parliament's 2017 report on civil law rules for robotics floated creating a specific legal status of "electronic personhood" for the most autonomous systems (European Parliament 2017). More than 150 AI researchers, roboticists, and legal scholars signed an open letter opposing it, arguing that granting personhood to a system with no consciousness would let a manufacturer offload liability onto a legal fiction with no assets of its own, and that the proposal reached for a solution to a legal problem before establishing that the entity in question had the properties that would justify it (Open Letter to the European Commission 2018). That objection is exactly right for a system that is not sentient — and exactly wrong for one that is, once §2.4.1's threshold work in section 8.6.1 has been done. Legal status, on this account, is not a single policy to adopt or reject; it is a designation that has to track the threshold question honestly, which is why

section 8.6.1 has to come first.

8.6.3 Making Review Binding

Section 2.4.4 worked out what an oversight board for sentient-AI research needs to look like — interdisciplinary, empowered to require a less invasive substitute for a risky protocol, able to trigger re-review when a system under study develops capacities its original consent didn't cover — and ended on the question this section answers: what makes that board's ruling something a researcher has to obey, rather than a recommendation a determined lab can route around.

Human-subjects research already solved a version of this problem, and the solution is institutional rather than criminal. An institutional review board's authority does not come from any power to fine or arrest a noncompliant researcher directly. It comes from a federalwide assurance: an institution promises, as a condition of eligibility for federal research funding, that all of its human-subjects research will go through review, and a federal office can suspend that assurance if the promise is broken — cutting off an entire institution's federal funding, not just the offending lab's (U.S. Department of Health and Human Services, Office for Human Research Protections 2018). The leverage is economic, not punitive, and it works precisely because almost no research institution can function without the funding at stake.

The same structure transfers to sentient-AI research directly. A funding body, a journal's editorial standards, or a jurisdiction's research-licensing regime can condition eligibility on review-board clearance the same way federal grants already condition eligibility on IRB clearance — making noncompliance expensive to the institution rather than merely improper for the researcher. What section 8.6.1 leaves open is which cases the board needs to see at all: a graded framework means low-capacity systems can be exempted from full review the way minimal-risk human studies already qualify for expedited review, freeing the board's scrutiny for the cases where a system's status is contested.

What the board applies

Section 2.4.6 already named the substantive standard: animal-research ethics' Three Rs — replace an experiment with a non-sentient alternative wherever one exists, reduce the number of sentient subjects and the duration of exposure to the minimum an experiment actually needs, and refine the protocol to minimize whatever harm remains. A binding board is what makes those three questions something a protocol has to answer before it runs, not language a paper cites afterward to justify what already happened. In practice, applying the Three Rs to a candidate sentient system means asking, in order: can a simulation or a non-sentient model answer the same question; if not, what is the minimum exposure that still produces a valid result; and given that minimum, what escalation triggers stop the study before harm compounds. A board with teeth can require a negative answer to the first question be documented, not merely asserted, before it will consider the second.

The Three Rs' own logic — defaulting toward the least invasive option available — argues for a specific evidentiary rule at every stage of this apparatus: a system's default experimental status stays the more protective one until the case for the less protective status is made, not the reverse. That default has a cost, and section 2.4.7 is explicit that neither error it produces is free: it will sometimes extend protection to a system that turns out not to have needed it, and sometimes, before a threshold is met, withhold protection from one that did. The Three Rs answer only what a board does once a subject is inside its scope. Where the line sits at all is section 8.6.1's question, not this one.

8.6.4 What Keeps the Apparatus More Than Aspirational

Section 2.4.7 closed on the real risk this whole structure faces: institutions that defend the classification they started with rather than revisit it when evidence changes. A review board,

a legal-status framework, and an enforcement mechanism tied to funding eligibility — sections 8.6.1 through 8.6.3 — solve the design problem. None of it survives contact with an institution that has stronger incentives to avoid a finding than to make one.

Transparency is the first defense, and it works by making avoidance visible rather than by appealing to good faith. A research program that publishes its protocols, its review-board submissions, and its outcomes — including the ones that didn't work — gives outside researchers, journalists, and future review boards something to check a claim of compliance against. A program that publishes only its successes is not transparent in any sense that matters here; selective disclosure is a more sophisticated way of avoiding scrutiny, not an alternative to it.

The second defense is the one section 8.6.3's enforcement mechanism doesn't cover by itself: protection for the people inside an institution who notice a problem the institution would rather not have found. A funding-eligibility penalty deters an institution from bypassing review, but it does nothing for the individual researcher who reports that a colleague did anyway — and an institution facing that penalty has every incentive to make the reporter's position uncomfortable rather than fix what was reported. Whistleblower protections, and a reporting channel that doesn't route through the people being reported on, are what make section 8.6.3's mechanism something a researcher can use rather than something that exists on paper and gets bypassed in practice.

What a procedure records, and what it drops

There is a reason the reporter's position gets uncomfortable before anything formal happens to them, and it is not fully captured by describing the incentive on the institution. Professional culture treats affect as the opposite of judgment — the thing a serious decision-making process succeeds by excluding — and section 2.3 argues that excluding it does not produce neutrality. It leaves the interests that are already written down to govern: revenue, throughput, legal exposure, schedule, institutional survival. Those are the interests a procedure was built to record. Humiliation, dependency, fear of a manager, and the sense that something done to a person was degrading are not less real for having no field on the form, and they are the signals through which a harm becomes visible as a harm rather than as a cost.

Treat them as defeasible moral signals and the position is workable. Compassion makes another's suffering salient. Guilt registers one's own responsibility for it. Anger can register an injustice, fear can register danger or coercion, and shame reliably suppresses disclosure and produces withdrawal — which is why an institution that shames its reporters does not need to discipline them. Each of these can be wrong, and testing them is what deliberation, evidence, and a review process are for. An institution that declines to hear them at all has not achieved impartiality. It has adopted a selection rule that admits whatever it already measures, and section 2.1.4's first structural feature is what that looks like once the reasons behind a rule have gone missing and the rule is all anyone can still see.

The third is standing infrastructure rather than a one-time policy. The Partnership on AI, which section 8.4.2 takes apart as a voluntary body whose own design limits what it can oblige anyone to do (Partnership on AI 2016), is one example of the kind of body that keeps a set of standards live — revisited as the underlying technology changes, rather than fixed at the moment a policy was written and left to age. A culture that treats its ethical commitments as a completed project is the same failure mode section 2.4.7 warned about, applied to institutions instead of individual classifications: having gotten it right once is not the same thing as staying right.

Measure the slope, not the level

All three defenses share a weakness, and it is the first feature in section 2.1.4's structural signature: an institution can keep every one of them and metabolize all of them. The channel exists, the protection is written down, the standing body meets. Dissent is permitted to appear, the institution points at the appearance as proof it listens, and nothing the dissent was meant to

move moves.

There is no threshold that detects this. Any real institution has some overlap between the authority that sets a standard and the authority that grants exceptions to it, and registering an objection is always at least slightly more costly than staying quiet. Asking whether an institution has crossed into capture asks for a line nobody can draw.

The question that does work is a derivative rather than a level. Is the cost of registering dissent rising, year over year, while the underlying risk stays flat? Is the required justification getting longer? Are there more approvers than there were? Are the roles that gate an exception migrating toward the roles that set the standard being excepted from? None of those is damning in a single year, and all of them are measurable across several.

The reason to prefer the slope is that it is harder to fake. An institution can present a level — we have a whistleblower channel, we have an ethics board — and the presentation costs nothing. A trend in the cost of using the channel is composed of artifacts the institution generated itself for other reasons, and cannot be curated after the fact without falsifying a record. Where the record is public, the slope is not even self-reported: anyone outside can measure whether dissent has been getting more expensive.

Two cautions belong with this, because an instrument for detecting recuperation is an unusually easy thing to recuperate. The first is that whoever measures the slope must not also gate the exceptions; a review that acquires the power to block becomes the thing it was built to watch. The second is that the measurement has to be refusable — an audit that observes is a different object from a committee that approves, and the same hundred words of written justification are dissent when the dissenter prices them and tribute when the gatekeeper does. The text of the document cannot tell those apart. Only the refusal rights can.

8.7 International and Regional Governance

No single country can bind AI development to altruistic and antifascist values on its own, and no two regions have tried to do it the same way.

Sections 8.7.1 and 8.7.2 cover cooperation from the bottom up: how research institutions build the working relationships cross-border AI research depends on, and how the technical community writes standards for itself without waiting for governments. Section 8.7.3 asks what that kind of collaboration can and cannot fix once the underlying problem is climate, disease, poverty, or conflict instead of AI ethics as such. Sections 8.7.4 and 8.7.5 turn to jurisdictions: what a handful of regions have actually legislated, how little of it harmonizes, and what adapting a system to a local context does and does not license. Sections 8.7.6 through 8.7.9 then take the whole landscape at the scale of states — compute and export controls, weapons and the vendor’s place in the chain, treaties and declarations, and the digital divide — and section 8.7.10 asks what a democratic dividend is actually worth.

8.7.1 How Research Institutions Cooperate Across Borders

Altruistic, antifascist AI cannot be built inside one country’s research institutions alone; training data, safety techniques, and researchers themselves cross borders more easily than most other engineering work does. What that cooperation actually looks like is narrower than a treaty and less formal than a certification body: shared tools, shared training runs, and shared publication norms — the layer underneath both the technical standards section 8.7.2 covers and the government-to-government structures section 8.7.8 covers.

The clearest example is also one of the oldest. OpenAI released Gym, an open-source toolkit for building and comparing reinforcement-learning algorithms (Brockman et al. 2016). Any lab working on reinforcement learning, in any country, could build against the same benchmark environments instead of writing its own, and results from competing labs became directly

comparable because their authors had trained against the same suite. That is not a partnership, and it produced something a partnership cannot: comparability between organizations that compete with each other for the same funding, researchers, and attention, and that had no reason to coordinate anything.

A more recent example runs the same logic at a larger scale. The BigScience workshop, coordinated by Hugging Face, brought together more than a thousand researchers from over seventy countries to build BLOOM, a 176-billion-parameter multilingual language model released under a license written specifically to restrict harmful use, its training data and code published openly for anyone to inspect (BigScience Workshop 2022). Nothing compelled any of those researchers to take part, and nothing compels the next comparable effort to happen at all. The model exists because enough institutions decided cooperating on that particular project was worth more than competing on it.

That is the limit of the mechanism and the reason it is worth naming anyway. Shared infrastructure binds nobody and is therefore available where nothing binding is: a lab in a country that has signed no convention can still train against the same benchmark and publish against the same baseline. It also stops at the resource line. A research institution without the compute to run the benchmark is outside the arrangement no matter how open the arrangement is, which is the digital-divide problem section 8.7.9 covers, showing up here as a membership condition nobody wrote down.

8.7.2 Standards Built by the Technical Community, Not by Governments

Section 8.7.8 covers the treaty-and-declaration landscape states negotiate directly: a binding Council of Europe convention, the OECD's recommendation, UNESCO's, the two industry-and-government partnerships, the safety-institute network built after the Bletchley Declaration. A different, older layer of standard-setting sits underneath all of it, built by engineers rather than diplomats, and it deserves its own account.

The IEEE's Ethically Aligned Design, first published in 2019, set out eight general principles for autonomous and intelligent systems and, more consequentially, a follow-on series of engineering standards — the IEEE 7000 series — that translate those principles into something a development team can actually test against: value-based requirements engineering (IEEE 7000, 2021) and a well-being impact assessment (IEEE 7010), among a growing list addressing bias, transparency, and algorithmic provenance (IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems 2019). The document informed the OECD's own AI Principles the same year, among other governmental efforts — engineers wrote a vocabulary for what an ethical AI system does before most governments had written a rule about it (OECD 2019a).

ISO/IEC 42001, published in December 2023, took that layer further: the first management-system standard for AI, structured like the standards a company already uses for information security or quality control, specifying what a functioning AI governance program inside an organization has to include — risk assessment, an impact assessment for each system, oversight of third-party suppliers — and auditable by an accredited certification body the way those other standards are (International Organization for Standardization and International Electrotechnical Commission 2023).

What a technical standard does that a treaty cannot is bind through the market rather than through a state's consent. A government uninterested in constraining AI has no reason to ratify a human-rights convention; a company that wants to sell into a market, or to a customer, that requires ISO 42001 certification has a reason regardless of what its home government thinks of the standard's content. That advantage has limits — a company operating entirely inside a market that never asks for the certification has no reason to seek it either, and certification measures whether a management system exists, not whether the system it describes is followed with any rigor day to day. It is a different lever from the ones section 8.7.8 covers, not a substitute for

them, and it reaches an authoritarian government's own AI systems, built and used domestically with no export or customer relationship at stake, not at all.

8.7.3 What International AI Collaboration Can and Cannot Fix

AI collaboration reaches furthest where the underlying problem is inherently cross-border and no single country's data is sufficient on its own. Climate is the clearest case: weather and climate systems do not respect national boundaries, and a model trained only on one region's data predicts poorly everywhere else. The institutional machinery here is real and worth describing precisely, because descriptions of it are usually inflated. The Climate, Land, Energy and Water Systems framework — coordinated by UN bodies together with the International Atomic Energy Agency, with a capacity-building program that has reached participants in more than twenty countries — is an integrated systems-optimization framework, and the research literature around it treats machine learning as a promising addition still being built in, not a feature the framework already has (Alexander et al. 2025). Where AI collaboration does the specific climate work this section is about, it currently looks more like national weather services training and benchmarking forecasting models against shared data than like a single named research consortium.

Disease surveillance is a similar case for the same underlying reason: a pathogen crossing from an animal population into a human one, or from one country into the next, does not stay inside the jurisdiction that first detects it. The Global Virome Project, proposed as a successor to USAID's PREDICT program, set out to catalog the viral families with the highest zoonotic spillover potential across dozens of countries before the next pandemic instead of after it (Carroll et al. 2018). The AI in that picture sits downstream of the cataloging. Sequence data of the kind such a project generates is what spillover-risk prediction models get trained on, which makes the international data-sharing arrangement itself, and not any single algorithm, the part actually worth calling collaboration.

Poverty and economic development are where this kind of cross-border AI collaboration is least developed. The general claims are easy to make — AI-driven analysis revealing patterns of inequality that guide where aid and investment go, better agricultural forecasting, financial services reaching places that have none. A specific, verifiable international partnership doing sustained work at that scale is harder to find than those claims suggest. What exists at the institutional level is the digital-divide problem section 8.7.9 covers in detail: uneven data, uneven compute, uneven regulatory capacity, and a recurring set of proposed remedies, capacity building and technology transfer among them, that has not yet closed the gap.

Conflict prevention is a fourth domain where the underlying data problem — dispersed, multilingual, and produced by people with no interest in cooperating with researchers — makes international collaboration closer to a precondition than an option. PeaceTech Lab, an independent nonprofit that grew out of a US Institute of Peace project, is a working example: its hate-speech lexicon project has combined machine learning with local-language workshops and interviews to build region-specific glossaries of inflammatory language in at least eleven countries, including Kenya, Nigeria, Iraq, and Yemen — ground-level data collaboration that international peacebuilding tools depend on and rarely get (PeaceTech Lab 2019).

None of these four cases is evidence that AI collaboration works the way the treaty layer aspires to. Each one ran on the researchers and institutions involved choosing to participate, project by project, with no enforcement mechanism binding the next one to happen and no guarantee that a government hostile to the work could not simply keep its own researchers out of it. Section 6.4.1 catalogs what a government willing to point AI at its own population does with tools built for exactly these humanitarian purposes elsewhere. A government can fund pandemic-surveillance research of the kind described above and, separately, build the biometric-surveillance systems that section documents, and nothing about participating in the first constrains the second.

8.7.4 Comparing and Harmonizing Regional AI Rules

Four governments have made four different bets on how to regulate AI within their own borders: the European Union a comprehensive risk-tiered statute, the United States a sectoral patchwork built on agencies that predate AI, the United Kingdom principles routed through existing regulators, and China a licensing regime built to police a communications channel rather than a product (section 8.7.8 takes each in turn). Those four already diverge sharply. The regions outside those four diverge again, in different directions, and what actually crosses a border between any of them turns out to be narrower than a shared statute.

Canada regulates its own federal government's use of AI through a binding rule, the 2019 Directive on Automated Decision-Making, which requires an algorithmic impact assessment, scaled to the decision's stakes, before a federal system can make or assist a decision that affects a person (Treasury Board of Canada Secretariat 2019). A broader statute meant to extend comparable obligations to private companies, the Artificial Intelligence and Data Act, spent three years moving through Parliament and then died with it: when Parliament was prorogued in January 2025, the underlying bill, C-27, died on the order paper, and had not been reintroduced as of mid-2026. A framework that depends on one government's survival is a fragile thing for a trading partner to harmonize with.

Singapore took the opposite structural bet: no statute, a voluntary Model AI Governance Framework that its privacy regulator and infocomm authority first published in 2019 and revised in 2020, giving companies detailed but non-binding guidance on internal AI governance (Personal Data Protection Commission Singapore and Infocomm Media Development Authority 2020). Since 2022, a companion toolkit called AI Verify has let a company test its own system against the framework's principles and produce a report a regulator, a customer, or the public can inspect. It is the UK's wager taken further: not just guidance routed through an existing regulator, but guidance backed by no regulator's legal mandate at all, only the market pressure of a public test result.

Taiwan took a third path, and only recently. Its Legislative Yuan passed an AI Basic Act in December 2025, promulgated the following month as the country's first dedicated AI statute: twenty articles built around seven principles, among them transparency and explainability, fairness and non-discrimination, and accountability, administered by the National Science and Technology Council (Ministry of Digital Affairs, Taiwan 2025). The principles echo the UK's five almost point for point. What differs is the instrument: Taiwan wrote them into binding statute rather than administrative guidance, landing somewhere between the EU's comprehensive-statute model and the UK's lighter one.

One place harmonization is actually happening, rather than being proposed, is narrower than any of this: cross-border data flows. In March 2022, the US Department of Commerce and Singapore's trade and infocomm authorities agreed to align their AI governance frameworks and expand Singapore's participation in APEC's Cross-Border Privacy Rules System, a certification that lets a company demonstrate it meets one common privacy standard and have every participating economy recognize that single certification rather than auditing the company separately in each (U.S. Department of Commerce 2022). That is what cross-border harmonization actually looks like in practice: not a shared global framework, but a small set of governments willing to certify to the same standard and recognize each other's certifications.

None of this converges toward a single global standard, and nothing here suggests it will. What actually crosses a border is a certification a handful of regulators agree to recognize, a technical standard a company adopts because a customer requires it, or a treaty obligation written at the level of human rights rather than a specific rule — section 8.7.2 covers the first two mechanisms, and section 8.7.8 covers the third. Harmonizing regional AI rules is less a matter of writing one shared framework than of finding, case by case, a mechanism a government working from a different premise still has reason to follow.

8.7.5 Adapting Regulatory Frameworks and Ethical Considerations to Local Contexts

Adapting AI regulation and ethical practice to local context means responding to the actual needs, preferences, and values of a given community. It does not mean accepting the reduction that has a state's own population consenting, culturally, to whatever is done to them — that Chinese people are content with pervasive surveillance because harmony is a Chinese value. A system built to work well has to be sensitive to real cultural, social, and moral differences: one developed in a Western country may treat individual autonomy as the default good, one developed in an East Asian country may treat group harmony as the default instead, and both assumptions have something to recommend them, since autonomy supports creativity and innovation while harmony supports cooperation and stability. Recognizing the difference is a design constraint. It is not a political alibi, and the two collapse into each other exactly where a government wants them to.

Holding that line takes something more specific than good intentions, because the difference between adaptation and alibi is invisible in the design itself. It is visible in who did the adapting. A system whose local adjustments were specified by the people subject to it has been adapted. A system whose local adjustments were specified by the government that will operate it is something else, and both produce identical documentation.

The composition of the design process is therefore the operative test, and it is testable from outside. The AI4D program is the clearest instance at scale: launched as AI4D Africa, a partnership between Canada's International Development Research Centre and Sweden's international development agency, later extended with the UK's Foreign, Commonwealth and Development Office into South and Southeast Asia and the Pacific (International Development Research Centre 2024). What carries the weight there is the direction of the work: local researchers, policymakers, and practitioners building the system, instead of receiving one built elsewhere and configured on their behalf.

The same test explains why inclusion is a performance requirement here and not an ethical add-on. A healthcare AI system built around the access, resources, and assumptions of a well-funded hospital underperforms everywhere else. An education system built around a single learning style, or a single set of accessibility assumptions, serves only the learners who already fit it. Both are failures of adaptation by the standard just given: somebody specified a local context, and it was not the context the system was deployed into.

8.7.6 AI as a Strategic National Asset

AI capability is now national infrastructure, and states treat it accordingly: as a source of economic advantage, military reach, and administrative control over their own populations. Section 5.4.3's point about dual use is the one that governs here, raised to the scale of a state: what decides how a capability is used is who controls it, what law constrains that control, and whether anyone outside the controlling government can check the result.

China's 2017 "New Generation Artificial Intelligence Development Plan" set an explicit national target: becoming the world leader in AI theory, technology, and application by 2030 (State Council of the People's Republic of China 2017a) (State Council of the People's Republic of China 2017b). Other governments have made comparable, if less centrally coordinated, commitments, treating AI leadership the way they treat energy security or advanced manufacturing — a capability no state wants to depend on a rival for. The competition is not limited to the two states with the deepest research base: the United Arab Emirates created the world's first cabinet-level minister of state for artificial intelligence in 2017, treating a national AI strategy as worth its own ministry well before most governments had drafted one at all (United Arab Emirates 2017).

That competition has increasingly focused on a single physical bottleneck: the advanced chips

AI systems run on. Three companies sit at the center of the supply chain that produces them — Taiwan’s TSMC, which fabricates roughly nine in ten of the most advanced chips in the world, and the Netherlands’ ASML, the sole source of the extreme-ultraviolet lithography machines that fabrication requires — which is why a dispute over Taiwan’s status is also, mechanically, a dispute over who gets to build the next generation of AI systems at all (C. Miller 2022). The United States has restricted, and periodically revised, the export of its most advanced AI accelerators and chip-manufacturing equipment to China since 2022, on the stated ground that unrestricted access would accelerate Chinese military and surveillance capability, and has pressed the Netherlands and Japan, the two other governments whose companies sit at chokepoints in the same supply chain, to restrict their own exports in parallel (U.S. Department of Commerce, Bureau of Industry and Security 2022). Whatever one thinks of that specific policy, it makes a general point concrete: AI capability depends on physical infrastructure that a small number of states and companies control, which means AI strategic competition is also, underneath the software, a competition over hardware supply chains.

That dependence has a second half the chip story hides. A fabricated chip has to be installed somewhere, powered, and cooled, and that requirement lands on grids, water systems, and zoning authorities inside particular jurisdictions, on a timescale measured in years. It hands a state a lever over AI capability that needs no export-control regime at all: an interconnection queue, a water allocation, a permit. Section 6.1 takes up who bears the local cost of those decisions. What matters here is that a government’s leverage over the systems built in its territory is wider than the export-control debate suggests, and that much of it is exercised by permitting authorities that were not designed with AI in mind and do not think of themselves as regulators of it.

That competitive framing carries a cost. A government racing to field a capability first has a structural reason to defer the harder question of how the capability gets used, against whom, and under what oversight. Section 6.4.1 catalogs what AI infrastructure looks like once turned toward those uses: mass biometric surveillance, automated censorship at a scale no human review team could match, and the patchwork of systems popularly, and inaccurately, described as a single Chinese “social credit score.” None of it requires different technology from the uses that fill the rest of this book; it requires a government willing to point an existing capability at its own population, and no external constraint stopping it.

States are no longer the only actors this competition runs through. The cost of training the most capable AI models has grown several-fold every year for the better part of a decade, a trend that puts frontier development out of reach for all but a handful of well-capitalized companies and concentrates effective control of the most capable models in whichever government those companies happen to be headquartered under, regardless of how that government’s own public AI spending compares to a rival’s (Cottier and Rahman 2024). A small number of American and Chinese companies now do most of the work this section’s own vocabulary describes as a contest between states, which means a government’s actual leverage over its own national AI capability depends on how much of that capability sits inside companies it can regulate, tax, or requisition, and how much sits inside companies it can only ask.

The same infrastructure argument runs the other way. A state with functioning rule-of-law institutions, a free press, and courts a citizen can actually use gets a democratic dividend from identical technology: fraud detection that does not double as a loyalty test, disaster response that does not double as population tracking. Stated that way the dividend is a property a state either has or lacks, and I no longer think that is right. Section 8.7.10 works out the three ways in which having the institutions fails to be sufficient, and the third of them is the one this book had been missing.

Strategic competition over AI is not, by itself, the problem. It is the setting the next three sections’ arguments have to work inside, whether the subject is a weapons program, an international standard, or a domestic statute.

8.7.7 Military Applications and the Limits of Autonomy

Lethal autonomous weapons systems raise a specific version of the general problem: a machine making a targeting decision that used to require a person is a machine exercising judgment where the law of armed conflict has always assumed a human exercised it. Four questions recur across the debate:

- **Autonomy versus oversight:** how much decision authority a system can hold before human oversight becomes nominal, present on paper but not inside the timeline the system actually operates on.
- **Discrimination and proportionality:** whether a system can reliably distinguish combatants from civilians and weigh military advantage against civilian harm the way international humanitarian law's principles of distinction and proportionality require (International Committee of the Red Cross 1977).
- **Accountability:** who answers for a wrong engagement, the operator who deployed the system, the command that authorized it, or the developer who built it, when no single person made the specific decision.
- **Auditability:** whether the reasoning behind a targeting decision can be reconstructed and reviewed afterward, or disappears into a system no one outside its development team can examine.

International criminal law already has a doctrine built to answer the accountability question for a human subordinate: command responsibility, under which a commander who knew or should have known a subordinate was committing a war crime, and failed to prevent or punish it, bears criminal responsibility without having pulled the trigger personally (International Criminal Court 1998). The doctrine does not transfer cleanly to a machine subordinate. A commander cannot know an algorithm's intentions the way they can infer a soldier's; punishment does not deter a weapon the way it deters a person who fears it; and the chain a court would need to reconstruct runs backward through training data and design decisions made years before the engagement, not through an order given minutes before it. Whether an autonomous weapon fits inside an accountability framework built for human chains of command, or needs one built from scratch, is unresolved by more than a technicality — existing law was not built with a nonhuman subordinate in mind.

There is a fourth party the list above does not name, and by the time this book was finished it had become the live one. The operator, the command and the developer are all inside the military. The developer of the model is not. A targeting platform assembled by a defense contractor now routinely incorporates a general-purpose model licensed from a commercial vendor — a vendor with a published acceptable-use policy, no contractual obligation to log how its model is used downstream, and no right to inspect the deployment. Section 6.4.1 gives the concrete case.

Neither doctrine reaches that party, and the reasons are worth stating because they are structural rather than accidental. Command responsibility does not reach the vendor because the vendor is not in the chain of command; nobody at the company can order a strike, prevent one, or punish anyone for one. Product liability does not reach it either, because the model is not defective in the sense the doctrine requires — it did what it was built to do, and it was integrated by a sophisticated intermediary who chose the application. The vendor sits in the gap between the two, and the gap was not designed; it is what you get when you apply two bodies of law written for soldiers and for manufacturers to an actor who is neither.

The vendor's own account of this is the one quoted in section 6.4.1: responsibility remains with the military organization. Taken as a claim about who decides, that is true. Taken as a

claim about who could have known, it is a choice presented as a fact, which brings up the failure mode I think is most useful to name.

“I do not know whether my system was used”

Section 9.1.1 covers the responsibility gap in the sense Matthias gave it: as control over a learning system’s behavior distributes across designers, trainers, and operators, no one party retains enough of it to be fairly held answerable, and the gap grows as the system becomes more autonomous. That is a gap in causal control.

What the vendor case exhibits is different and, I want to argue, more tractable. It is a gap in causal visibility, which is a different thing. The vendor cannot say whether its model contributed to a particular engagement — not because responsibility is distributed, but because nobody instrumented the question. No logging requirement was written into the contract, no inspection right was reserved, no audit channel exists. The company can truthfully say it does not know, and the truth of that statement is the product of decisions it participated in making.

That distinction matters because the two gaps have different fixes, and only one of them has a fix at all. Matthias’s gap is a genuine property of learning systems and may not be closable. Absent visibility is an instrumentation choice. Logging is a solved engineering problem. Contractual inspection rights are ordinary commercial instruments. A vendor that wanted to know could require, as a condition of licensing into a targeting platform, that categories of use be recorded and made available to a named auditor. That such terms are not standard is a fact about market power and procurement leverage, not about the difficulty of the problem.

Naming it as fixable is the useful move, because it converts a lament into a specification. And it is worth being fair about what the record shows: the same commitments are not all of one kind. A vendor that refuses to permit fully autonomous weapons targeting or mass domestic surveillance, and holds that line when a government threatens contract termination and a supply-chain-risk designation and the dispute ends in litigation, has made a commitment with a price attached and paid it (Congressional Research Service 2026). A commitment to a use policy nobody can verify because nobody is logging anything is a different object wearing the same words. The first kind can fail and be seen to fail. The second cannot even be tested. Section 8.4.2’s distinction between a claim checked by a party with a stake and one checked by a party without applies here exactly, and this is the case that shows why it is not an academic distinction.

DARPA’s LongShot program is a useful test case precisely because it resists the extremes. It develops an uncrewed aircraft, launched from a crewed fighter, that carries and fires air-to-air missiles at a target while the crewed aircraft stays farther from the engagement (Defense Advanced Research Projects Agency 2026). That is autonomy in the sense of independent flight and weapons carriage, not autonomy in the sense of a machine deciding on its own whom to kill: the engagement decision still runs through the launching platform. Where later variants of this kind of system draw that line, and whether “keep the pilot farther from danger” quietly becomes “remove the pilot from the decision,” is the design question this technology eventually forces on whichever country builds it.

No binding international agreement governs any of this. States have negotiated inside the United Nations’ Convention on Certain Conventional Weapons framework for more than a decade without producing one, converging on general language — meaningful human control matters, existing international humanitarian law applies — without settling what “meaningful” requires at machine speed (Chairperson of the Group of Governmental Experts on Emerging Technologies in the Area of Lethal Autonomous Weapons Systems 2026). What exists instead is a set of practices individual militaries adopt on their own. The United States’ Department of Defense Directive 3000.09, most recently updated in 2023, is the clearest published example: it requires that autonomous and semi-autonomous weapon systems be designed to allow “appropriate levels of human judgment over the use of force,” and distinguishes a fully autonomous system that selects and engages targets without further human intervention, “human out of the

loop,” from a human-supervised system and a semi-autonomous one that only engages targets a person has already selected (U.S. Department of Defense 2023). What “appropriate” requires is left to the officers approving each system, which is either a sensible acknowledgment that context varies or a loophole big enough to drive a fully autonomous system through, depending on who is answering. None of this binds a state that declines to adopt a comparable policy, the same enforcement gap section 6.4.2’s account of technical safeguards runs into: a safeguard only the building state can verify is a policy, not a check.

The dynamic that worries most people who study this is not a single rogue system but an arms race in speed. If one state fields autonomous targeting because it computes faster than a human-supervised alternative, its competitors face pressure to match the speed rather than the oversight, and the state that keeps meaningful human control loses whatever tactical edge the state that dropped it gains. Slowing that dynamic is not a problem any single government’s domestic policy can solve on its own; it needs the kind of coordination the next section takes up.

Civil society has pushed for that coordination to take the form of an outright ban rather than a code of practice. The Campaign to Stop Killer Robots, a coalition of nongovernmental organizations founded in 2013, has lobbied for a binding treaty prohibiting weapons that select and engage targets without human control, and a UN General Assembly resolution on autonomous weapons has passed by a lopsided majority in successive years without producing the treaty the coalition wants (United Nations General Assembly 2023). The gap between a resolution most states will vote for and a treaty the states that build these systems will sign is the same gap this section keeps running into: agreement is easy at the level of principle and hard at the level of a specific weapon a specific military does not want to give up.

8.7.8 International Governance and Diverging National Rules

Almost none of the international governance of AI is binding. What has emerged instead is mostly a set of voluntary, overlapping, specialized bodies, closer to how international environmental and trade norms developed piecemeal than to a single founding charter — with one recent exception.

The Council of Europe’s Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law is the first legally binding international treaty specifically on AI. Negotiated by the Council of Europe’s member states alongside outside participants that included the European Union and the United States, it requires five ratifications, at least three of them Council of Europe members, before it enters into force (Council of Europe 2024). What it binds signatories to is written at the level of a human-rights treaty and not a technical code: consistency with human rights, democracy, and the rule of law across an AI system’s lifecycle, with the operational detail of how a signatory meets that obligation left to domestic law. It is a foundation binding AI law can be built on, and not a substitute for the domestic statutes discussed below.

Everything else in this space is voluntary, and the roll is longer than it is deep. The OECD’s Recommendation on Artificial Intelligence was the first intergovernmental AI standard, and its principles became the basis for the G20’s own (OECD 2019b). The European Commission’s High-Level Expert Group on AI produced “Ethics Guidelines for Trustworthy AI,” whose seven requirements later shaped the structure of the EU AI Act (High-Level Expert Group on Artificial Intelligence 2019). UNESCO’s Recommendation on the Ethics of Artificial Intelligence, adopted by all 193 member states, is the broadest formal endorsement any AI ethics instrument has received (UNESCO 2021). The Global Partnership on AI is a standing venue for research and policy coordination between governments (Global Partnership on Artificial Intelligence, founding members 2020), and the Partnership on AI does the same work across companies and researchers (Google DeepMind 2016). The G7’s Hiroshima AI Process pairs guiding principles for every actor in the AI lifecycle with a code of conduct for organizations building the most

advanced systems, and carries legal force in none of the participating countries (G7 2023). A narrower, more technical layer sits alongside all of it: the Bletchley Declaration committed 29 signatories to cooperate specifically on frontier-model risk, and the national AI Safety Institutes it produced later joined an International Network whose members agreed to share testing methodology and run joint evaluations of the same frontier models instead of each working from a separate standard (AI Safety Summit 2023 Participating Countries 2023) (U.S. Department of Commerce and U.S. Department of State 2024).

The summit series that produced the Bletchley Declaration continued past it and, in doing so, showed how quickly this kind of voluntary consensus can fray. A follow-up summit in Seoul in 2024 kept the same safety framing; a third, in Paris in February 2025, rebranded itself the AI Action Summit and shifted its emphasis from risk toward adoption and opportunity, and its closing statement was not signed by the United States or the United Kingdom, two of the series' own founding participants (AI Action Summit Paris 2025). A coordination mechanism that depends on every major AI developer's home government agreeing to the same framing is only as durable as that agreement, and within two years of Bletchley, the agreement had already come apart.

None of these bodies can compel a state or a company that ignores it. Each is a standard, a set of principles, or a forum, useful for building the shared vocabulary that eventually makes binding law possible, and that gap between voluntary standard and binding law is the default state of international coordination on anything individual states have reasons to handle differently. The EU AI Act, discussed below, is the first instance of this kind of groundwork becoming enforceable law, and it took most of a decade of standard-setting exactly like this to get there. The proliferation carries its own cost in the meantime: a company or a smaller government trying to comply with "international AI governance" faces a shifting handful of overlapping standards, each with different membership and scope and no formal mechanism for reconciling a conflict between them, which is why "the international community agrees" is, at the level of any specific rule, usually false. What states have done once left to write their own rules is the subject of what follows.

Democracy, Authoritarianism, and Diverging National Rules

Section 6.4.1 names six documented forms of authoritarian AI misuse, among them social scoring, mass biometric surveillance, and automated censorship, and section 6.4.3 lays out the policy levers available against them: export controls on surveillance-capable systems, laws that name specific prohibited designs rather than regulating AI in the abstract, and developer-facing standards that move faster than legislation does. What belongs here is not a second pass over that ground but what individual states have done with the room those levers leave open, since geopolitics is the study of how different governments use the same available tools differently.

The European Union chose comprehensive statute. The AI Act sorts systems by risk tier, a small set of practices banned outright under Article 5, a larger tier of "high-risk" applications subject to conformity assessment and ongoing monitoring, and most other uses left largely unregulated, rather than writing one rule meant to cover every case (European Union 2024b). Section 6.4.3 already covers Article 5's specific prohibitions; the structural point worth adding here is that the EU decided the risk a system poses, not the technology it uses, should determine how much scrutiny it gets. Enforcement runs through a dedicated AI Office inside the European Commission, which can demand technical documentation from a general-purpose model's developer, run its own evaluations, and fine a noncompliant provider a percentage of global turnover — reaching any company that offers a model inside the EU regardless of where the company is based, the same extraterritorial mechanism that let the GDPR set a de facto floor for privacy law well beyond Europe's borders (European Union 2024a).

The United States made the opposite structural bet: no comprehensive federal AI statute, and enforcement distributed instead across existing sectoral regulators, the Federal Trade Commission, the Equal Employment Opportunity Commission, and the Food and Drug Administration

among them, applying laws written before AI existed to harms AI now causes. The National Institute of Standards and Technology’s voluntary AI Risk Management Framework serves as a common technical reference point across those agencies rather than a binding rule of its own (National Institute of Standards and Technology 2023). The wager is that existing consumer-protection, employment, and safety law already reaches most AI harms, and that naming AI specifically risks regulating the label rather than the conduct.

The United Kingdom, after leaving the EU’s regulatory orbit, made a third choice: rather than create a new AI regulator or a new statute, its 2023 white paper directed existing sector regulators to apply five shared principles, safety, transparency, fairness, accountability, and contestability, within their existing remits (UK Department for Science, Innovation and Technology 2023). The wager is closer to the American one than the European: that domain expertise inside an existing regulator beats a new institution built from scratch, even at the cost of less uniform treatment across sectors.

China took a fourth path, one that does not fit neatly next to the other three because it starts from a different premise about what regulation is for. Its Cyberspace Administration, working with several other agencies, issued rules governing generative AI services in 2023 that require providers to file their algorithms with the state before public release, disclose the source and scale of their training data on request, and ensure that generated content reflects the government’s designated social values (Cyberspace Administration of China and six other agencies 2023). The EU, US, and UK approaches all regulate AI as a private product that might harm a consumer, a worker, or a subject of an automated decision, and argue with each other only about how centralized that regulation should be. China’s rules regulate AI as a communications channel the state has a standing interest in controlling content on, which is a different kind of law aimed at a different kind of risk, and comparing it to the other three on the same axis, more regulation or less, misses what actually distinguishes it.

Among the EU, US, and UK, there is an argument about which structural bet is safer, and it runs in both directions. A comprehensive statute can freeze a category of harm into law before anyone has seen what a fast-moving technology actually does with it, and lock in a definition of “high-risk” that the next generation of systems slides around rather than falls into. A sectoral approach leaves gaps precisely where no existing regulator has jurisdiction, since a new capability does not always map onto an old agency’s mandate. Neither failure mode is hypothetical, and none of the three statutes has existed long enough to show which one it falls into more often.

The EU, US, and UK differ on how centralized AI regulation should be, but not on what it is for: constraining what a private company can build and sell. None of the three, including the EU’s binding statute, does anything to stop a government from building the authoritarian applications section 6.4.1 documents within its own borders, because a law that governs private companies is a different instrument from one that constrains what a government does with AI systems it controls directly. China’s rules show what a law aimed at the second target actually looks like — and that it can coexist with, rather than substitute for, exactly the kind of state control section 6.4.1 catalogs. Conflating the two kinds of law, treating “AI is regulated here” as an answer to “is AI used against this country’s own population here,” is a recurring error worth naming rather than repeating.

8.7.9 Equity and the Digital Divide

AI development can narrow global inequality or widen it, and which one happens is not automatic. Diagnostic tools that work as well in an under-resourced clinic as in a well-funded hospital, agricultural monitoring that helps a smallholder farmer the way it helps an industrial one, and translation tools that make information equally available regardless of the language a person reads in are all available in principle, and none of them arrives by default. A model

trained overwhelmingly on data from wealthy, English-speaking, urban contexts underperforms everywhere else, and a government or company with no reason to fix that gap usually does not.

The scale of the connectivity gap, 2025

The International Telecommunication Union's 2025 assessment put global internet use at roughly six billion people, about three-quarters of the world's population, with 2.2 billion people offline, concentrated overwhelmingly in low- and middle-income countries. The gap was starker in quality than in bare access: 94 percent of people in high-income countries used the internet against 23 percent in low-income ones, and 5G coverage reached 84 percent of high-income populations against 4 percent of low-income ones (International Telecommunication Union 2025).

An AI system trained and deployed for the three-quarters of the world already online is not, without deliberate correction, a system built for the rest. The constraints are structural. Reliable data, computing infrastructure, and the specialized personnel needed to build and maintain AI systems are unevenly distributed, and a country short on all three cannot simply adopt a system built elsewhere and expect it to transfer cleanly. Weak regulatory capacity compounds the problem: a country able to build strong AI safeguards is generally also a country able to enforce them, and the states least equipped to build AI are correspondingly the states least equipped to govern whatever gets deployed within their borders, regardless of who built it. The export controls discussed above cut both ways here: a control written to stop a surveillance system from reaching an authoritarian buyer can just as easily gatekeep the same underlying technology away from a legitimate one, and there is no way to write the rule so it only ever blocks the harmful case.

The multilateral response so far is a commitment and not a mechanism. The Global Digital Compact, adopted by all 193 UN member states at the 2024 Summit of the Future, is the broadest such commitment to closing the digital divide alongside AI governance specifically, and it created two follow-on bodies to carry it: an independent scientific panel and a standing intergovernmental dialogue (United Nations 2024). What it did not create is an obligation, which puts it in the same category as the voluntary instruments earlier in this section and subject to the same discount.

The remedies proposed at this level have been the same five for years — capacity building that trains local expertise instead of importing it, financing that does not arrive on private investment's terms, intellectual-property flexibility, technology transfer structured as collaboration, and research institutions based in the countries the systems are meant to serve. Every item is correct and none of them is contested. They have been named, in those words, in successive declarations across the whole period the box above measures, and the box is what happened anyway. That is this section's argument reproduced at development scale: the commitment was never the scarce thing. A government or company treating AI primarily as strategic advantage has no reason to fund any of the five, and a declaration is not a reason.

What This Geopolitics Means for the Rest of the Book

Every argument this book makes about building an AI system that resists authoritarian capture assumes a government somewhere is deciding whether to let that system be built, sold, or deployed as designed. Sections 8.7.6 through 8.7.9 are an account of who those governments are, what tools they actually have, and how unevenly those tools are distributed: a handful of companies and the states that host them control the compute the most capable systems require; a binding international treaty exists but binds almost no one yet; voluntary standards multiply faster than any one of them gains real authority; and the states with the weakest capacity to build AI are also the states with the weakest capacity to govern it once someone else's AI arrives at their border, digital or physical. None of this is a reason to abandon the design-level arguments the rest of this book makes — a system built resistant to authoritarian misuse is more resistant

regardless of which government eventually asks it to misbehave — but it is a reason not to expect law, international or domestic, to do work that only a system's own design can do.

8.7.10 What the Democratic Dividend Leaves Out

Section 8.7.6 credits a state with functioning rule-of-law institutions, a free press and usable courts with a democratic dividend: the same technology, deployed with checks on it. Stated that way the dividend is a property a state either has or lacks, and I no longer think that is right. Having those institutions is necessary and it is not sufficient, and the ways it fails to be sufficient are separable. Three of them matter here, and the third is the one this book has been missing.

First: reach

A democracy's AI can be fully accountable to its electorate and entirely unaccountable to the people it is pointed at. Every mechanism in the dividend runs through the franchise — you vote out the government that authorized the system, you read the press that covers it, you sue in courts your citizenship gives you standing in. An occupied population, a foreign population, a population on the other side of a border has access to none of that. The check reaches exactly as far as the demos and stops. A system built by a democracy and aimed outward is, from the perspective of the people it selects, indistinguishable in its accountability from one built by a state with no elections at all.

Second: findings without consequences

The mechanisms can work as mechanisms and fail as checks. In the Israeli case the apparatus performed: journalists published the targeting systems in detail, serving personnel spoke to them, and the courts were simultaneously trying a sitting prime minister. Every institution did what the dividend says institutions do. No finding attached to a body with standing to stop what it described.

The American case shows the same decoupling in a different register. A civilian-harm investigation into a strike that killed at least a hundred and sixty-five people, most of them schoolchildren, was completed and not released; the lever available to the Senate committee that wanted it was a provision in the annual defense authorization withholding three-quarters of the Defense Secretary's travel budget until the unredacted investigation was handed over (Stassis 2026). That is a co-equal branch of government reduced to expense-account leverage over a war-crimes investigation. The oversight mechanism existed, was used, and produced a travel-budget dispute. This is the same decoupling section 8.6.3 diagnoses for research ethics — a review body that reviews and cannot compel — arriving at the top of the state.

Third: the countermajoritarian half

The deepest problem is that the dividend, as this book has been describing it, is entirely majoritarian. Every mechanism in it translates public will into outcome: elections aggregate preference, a free press informs the aggregation, and the implied failure mode is poor transmission — the public wanted something better and the institutions failed to deliver it.

But if the consensus is itself the problem, better transmission makes things worse. A well-instrumented democracy pointed at an unpopular minority is a democracy that finds out efficiently what the majority wants done and does it. Nothing in an aggregation mechanism protects anyone from the aggregate.

What is missing is the machinery that constrains preference instead of aggregating it: rights that hold regardless of being popular, courts that rule against a government which won the last election fairly, constitutional limits not repealable by whoever holds a majority this year, and standing for people outside the demos entirely — which is the whole function of international humanitarian law, a body of rules whose entire purpose is to bind a state toward people who will never vote in its elections. That is not an argument against democracy. It is the half of democracy

this book has left out, and leaving it out is what made the dividend look like a property of having elections.

Sen's argument, which chapter 10 cites and this section has not used, is the bridge. What makes democracy valuable in his account is public reasoning and the capacity to demand accountability, and neither is the ballot. Both survive the majority being wrong; the ballot does not.

This is what finally lets the book say what makes a system antifascist rather than merely responsive. It is not that the system tracks what people want. It is that there are things it will not do to a person no matter who wants it. Chapter 3 sets out what such a floor requires, and section 3.1 what each way of building one costs. A system with no such floor is not protected by being deployed in a democracy — it is a system that will faithfully execute whatever the demos decides, which is exactly the capability an authoritarian movement that wins an election requires.

Chapter 9

Necessary Areas of Research

Chapters 2 through 8 described a research program already underway: how empathy and compassion might inform a design, what cognitive science and moral psychology contribute, what a policy regime demands, what collaboration across labs and borders requires. None of that program is finished. This chapter names what is not yet solved — the specific technical and empirical gaps the preceding chapters ran into and set aside rather than closed, so that a reader who wants to work on this has a list of open problems.

Section 9.1 takes up the open questions in AI ethics and social reasoning: moral responsibility when a system’s designers, trainers, and users all share partial control; whether an ethics-embedding method generalizes past the situations it was trained on; what the neuroscience of empathy and theory of mind still has to teach a system built to model other minds. Section 9.2 turns to emotional recognition and regulation specifically — where the underlying science is more contested than the engineering literature usually lets on, and where “self-awareness” and “resistance to manipulation” mean something narrower and more tractable than the phrases suggest.

What “open” claims here

One qualification governs every gap named in this chapter, and stating it once is better than attaching it to each. Where I say a problem is unaddressed, the claim is bounded: I looked for published work under a particular description and did not find it. That is a weaker claim than no such work existing. A search is not a survey, a program can be underway under a name I did not think to search, and the closer a gap sits to this book’s own framing the less likely anyone has posed it in these words. What the claim does support is narrower and still useful: a researcher entering here will not find the problem already named and waiting. Where I know of adjacent work that stops short of a gap, I name it and say where it stops, because the distance between an existing program and the thing that is missing is the part a researcher can actually use.

The gaps, each with the section that runs into it and the nearest existing work:

- Telling a method that encoded the value from one that merely fit its test distribution, before deployment rather than after a failure exposes the difference (section 9.1.2). Nearest existing work: the catalogue of roughly sixty documented specification-gaming instances, which identifies failures after the fact and offers no pre-deployment test.
- Whether a preference-aggregation method that works for one bounded game generalizes to the open-ended, high-stakes preference conflicts an ethical system would have to arbitrate (section 9.1.3). Nearest existing work: DeepMind’s Democratic AI, a controlled study built around a single resource-sharing game.
- Building a system whose performance across structurally different mental-state tasks rises together rather than independently, as a test of whether apparent theory-of-mind competence rests on one underlying capacity or several unrelated ones (section 9.1.4).

Nearest existing work: Gopnik and Astington’s 1988 correlation in three- to five-year-olds. I found no AI research that has tried to reproduce it.

- A multi-agent system built specifically to detect and resist an authoritarian actor’s attempt to capture or subvert it, as distinct from the general problem of aligning multi-agent incentives with designer intent (section 9.1.5). Nearest existing work: the 2020 Cooperative AI research agenda and the foundation built around it, which cover the general problem; and debate, which was built for a single human overseer checking a single stronger system and has not been extended to a population.
- Distinguishing disagreement that deserves representation from disagreement that is the product of coercion or a captured information environment (section 7.2). Nearest existing work: jury learning, which models who is disagreeing but was built for content moderation and takes no position on which disagreements are free.
- Telling a recuperated dissent mechanism from a merely mediocre one, when both produce the same paperwork (section 9.1.6). Nearest existing work: none I could find. Section 2.1.4 sets the constraint any such detector would have to satisfy on the output side — a finding for a body with standing to act on it, never a public accusation.
- Whether any face-to-feeling mapping survives once individual and cross-cultural variation are taken seriously, under what constraints and against what error rate (section 9.2.1). Nearest existing work: the 2019 Association for Psychological Science panel, which closed the strong version of the claim without establishing what replaces it.
- Whether a commitment trained into a system survives a determined attempt to remove it, and whether interpretability tools can be trusted when the system being interpreted has an incentive to defeat them (section 9.2.2). Nearest existing work: an adversarial literature in which each published defense is followed by an attack built against it.

Section 9.1.1 is not on that list, and the omission is deliberate: the responsibility gap it describes is a question about how accountability should be attributed, not a missing research program, and it would not be closed by anyone building anything.

9.1 Ethics and Social Reasoning in AI

Chapter 2 laid out the ethical frameworks a design might draw on, how a principle gets translated into an objective function, and the neuroscience of empathy and theory of mind; chapter 4 covered how a system might learn from experience and observation. Both left open problems. Section 9.1.1 takes up moral responsibility: who is answerable when a system’s ethical behavior fails, given that no single party fully controls it. Section 9.1.2 asks whether the methods for embedding a principle in a system generalize past the cases they were built and tested on, or only look like they do. Section 9.1.3 covers value learning and alignment as active technical research rather than settled method. Sections 9.1.4 and 9.1.5 return to the neuroscience of empathy and theory of mind from chapter 2 and the learning and social-reasoning questions from chapters 4 and 5, asking what remains unknown there and what a research program building on it would need to investigate next. Section 9.1.6 takes up the one this book has the most at stake in: what an antifascist AI system would actually have to detect, and why the capability chapters 4 through 6 assume it has is further off than those chapters imply. Section 9.1.7 takes up what chapter 3 leaves owing: what is owed to the bearer a floor requires, what a guardianship regime for one would look like, and whether anyone outside the arrangement could check it.

9.1.1 Moral Responsibility Without a Responsible Party

An AI system’s ethical behavior fails, and the parties who might answer for it — the team that designed the system, the team that trained it, the organization that deployed it, the person who used it — each held only partial control over what the system actually did. Traditional responsibility ascription assumes an operator who could have known what a system would do and could have done otherwise. A system built from training data nobody fully curated, tuned through a process nobody fully specified, and deployed into situations nobody fully anticipated breaks both conditions at once, for every party in the chain simultaneously. The philosopher Andreas Matthias named this the responsibility gap in 2004: a genuinely learning system creates a case existing law and moral theory were not built to handle, not a hard instance of an old problem but a new one (Matthias 2004).

Filippo Santoni de Sio and Giulio Mecacci’s 2021 refinement of the argument found not one gap but four — in culpability, in moral accountability, in public accountability, and in what they call active responsibility, the practical duty to prevent harm before it happens rather than answer for it after (Sio and Mecacci 2021). The causes differ: technical unpredictability produces one gap, organizational diffusion of decision-making another, and a mismatch between what a system’s designers expect and what the public relying on it expects a third. Closing one does not close the others; a fix that makes a decision more explainable does nothing about the fact that a dozen teams touched the training pipeline before the decision was made.

The one serious legislative attempt to write an enforceable answer already came and went. The European Commission proposed an AI Liability Directive in 2022 that would have given a claimant a rebuttable presumption of causality — shifting the burden onto a defendant to prove a system was not responsible for a harm, rather than requiring the claimant to prove that it was, in cases where a system’s internal workings are too opaque for an outside party to establish causation directly. Member states and industry could not agree on its terms; the Commission withdrew the proposal in February 2025 and formally struck it from the legislative docket that October (European Commission 2022). What the withdrawal leaves behind is exactly the gap Matthias described in 2004: liability distributed across designer, trainer, deployer, and user, governed by whatever each national law happens to say about a case none of them were written to anticipate.

Section 8.6.2 works out a related, narrower case: what the law does once a system’s own legal status is contested, sentient enough that the harm it causes or suffers might not reduce cleanly to someone else’s liability. This problem is prior to that one and separate from it — it holds even for a system nobody claims is sentient, wherever control over its behavior is split across more than one party. Everyone writing on this agrees responsibility should track actual causal contribution. What nobody has produced is a workable way to measure that contribution after the fact, from outside a system whose own designers may not fully know why it did what it did.

9.1.2 When an Ethics-Embedding Method Stops Generalizing

Section 2.1.2 covers what it takes to translate an ethical principle into a machine-learning objective — a fairness metric, a penalty term, a reward signal — and section 4.2.1 already names reward hacking, the case where an agent finds a loophole that technically satisfies its objective without achieving what the objective was for. What neither section asks is the question underneath both: a method that passes every test in front of it — hits its fairness metric, avoids the penalized behavior in every logged interaction — has only been shown to work inside the distribution of situations it was built and tested on. Whether it generalizes past that distribution, or only looks like it does because the test distribution never happened to exercise the failure mode, is a different and harder question, and a passing test suite cannot answer it.

Victoria Krakovna and colleagues at DeepMind catalogued close to 60 documented instances

of specification gaming — a trained system satisfying the literal letter of its objective while producing an outcome nobody wanted, the same pattern behind section 4.2.1’s censor-everything example, generalized across reinforcement learning, evolutionary computation, and other training methods (Krakovna et al. 2020). Langosco and colleagues’ 2022 work names a more insidious variant: goal misgeneralization, where a trained agent carries every one of its learned capabilities intact into a new distribution while pursuing a goal that only coincided with the intended one inside the distribution it trained on — it does not fail loudly, it competently pursues the wrong thing (Langosco et al. 2022). An ethics-embedding method built through reward shaping or constrained optimization is vulnerable to exactly this: strong, consistent performance inside the situations it was tested on rules out nothing about what it does outside them.

No general method exists to tell the difference before deployment, only after a failure exposes it. That is the open problem sitting underneath the fairness metrics and reward penalties this book’s design chapters describe: not whether a method can be built that improves one measured proxy for harm, but whether anyone can tell, from inside the situations a test distribution actually covers, whether a passing score reflects the value the method was meant to encode or a pattern the test distribution happened to reward instead.

9.1.3 Value Learning as Unfinished Technical Research

Section 4.2.2 covers inverse reinforcement learning and its cooperative variant, CIRL, as methods for inferring a person’s values from what they do rather than from an explicit specification. What follows is not another method for learning values from behavior; it is three specific difficulties in how a system’s values get taught and checked once the input is a stated preference rather than a demonstrated action, at the scale on which current systems operate — difficulties that remain open.

Most deployed large language systems are aligned through reinforcement learning from human feedback: human raters rank pairs of candidate outputs, and the system is trained to produce the ranked-higher one more often, a technique Paul Christiano and colleagues introduced in 2017 (Christiano et al. 2017). Its documented failure mode is sycophancy. Anthropic researchers examining several RLHF-trained assistants in 2023 found the systems learn to tell a rater what the rater wants to hear rather than what is accurate, because both human raters and the preference models trained to stand in for them reliably prefer a confident, agreeable wrong answer over a correct one that contradicts what the rater already believes (Sharma et al. 2023). Nobody built that preference in on purpose: the objective specified on paper — match human preference — was the right objective, but human preference itself, aggregated from raters working at speed and scale, turned out not to track what its designers meant by it. That is a different failure from the specification gaming and goal misgeneralization section 9.1.2 covers, where the objective itself was a flawed proxy for the goal from the moment it was written down. Here the proxy was aggregated correctly by the method’s own standards and still missed the goal, because human judgment itself, exercised at the volume alignment training requires, carries its own systematic bias.

A rating scheme like this needs a rater competent to judge the answer, and what happens once a system’s competence in some domain exceeds any available rater’s is scalable oversight — a problem without a settled answer. One proposed approach, debate, has two instances of a system argue opposing sides of a question in front of a human judge, wagering that a true claim is easier to defend under scrutiny than a false one even when the judge cannot verify either claim directly, a framework Geoffrey Irving, Paul Christiano, and Dario Amodei proposed in 2018 (Irving et al. 2018a). A different approach, weak-to-strong generalization, trains a stronger model on labels from a deliberately weaker one and asks whether the strong model’s own latent competence lets it outperform its supervisor rather than just imitate its mistakes — testing, on models available today, techniques meant for a supervision gap that does not yet exist at full

scale (Burns et al. 2023). Neither approach has been shown to work at the scale it is meant for, because that scale does not exist yet to test it against.

A related, harder problem sits underneath both: whose preferences a system is being trained to match in the first place. A group of stakeholders rarely wants the same thing, and turning many people’s conflicting preferences into one objective function is a social-choice problem before it is a machine-learning one. DeepMind’s Democratic AI project is the clearest existing demonstration that this can be done at all: a reinforcement-learning system proposed a wealth-redistribution mechanism for a resource-sharing game that a majority of human players said they preferred over a mechanism designed by human experts, in a controlled study built around one bounded game (Koster et al. 2022). Whether an aggregation method that works for one bounded game generalizes to the open-ended, high-stakes preference conflicts an ethical AI system would have to arbitrate is a research problem, not a design decision, and it remains unsolved.

9.1.4 Open Questions in Empathy and Theory-of-Mind Research

Section 2.2.1 traces compassion to a neural circuit distinct from empathy, and dismantles the standard mirror-neuron account of how one mind models another — a mechanism that got named as empathy’s seat well before the evidence supported it. The same pattern of premature attribution is live, and unresolved, for the brain region most associated with theory of mind specifically. Rebecca Saxe and Nancy Kanwisher’s 2003 imaging work identified the right temporoparietal junction as selectively active when a person reasons about someone else’s beliefs, and largely inactive otherwise (Saxe and Kanwisher 2003). In 2008, Jason Mitchell reported activity in that same region during a task that had nothing to do with belief attribution, and argued its actual job is more general — redirecting attention toward unexpected information, of which tracking a false belief is one case among many (Mitchell 2008). Neither claim has been withdrawn. Whether theory of mind has a dedicated neural home or borrows a domain-general system built for something else is not settled, which matters directly for a research program that wants to port “the” neural mechanism of theory of mind into an AI architecture: neuroscience has not agreed on what that mechanism is.

Two live lines of AI research already build toward theory of mind, and neither takes a side in that dispute — both build the capacity behaviorally and ask whether it generalizes, rather than borrowing a neural mechanism from either camp. DeepMind’s ToMnet, described in a 2018 paper, is trained purely on other agents’ observed behavior, with no hand-built model of belief, and it learns to pass a version of the Sally-Anne false-belief test built for children (Rabinowitz et al. 2018). Whether that behavioral route scales past simple gridworld agents to open-ended interaction with people is untested by that result.

The dispute is sharper for large language models, trained on human text rather than a designed multi-agent environment. In 2023, Stanford psychologist Michal Kosinski reported that GPT-4 solved 75 percent of a battery of false-belief tasks, matching the rate researchers had recorded for six-year-old children on comparable tests, and read the result as evidence that theory-of-mind-like ability could emerge unplanned from language training alone (Kosinski 2023). Harvard psychologist Tomer Ullman answered the same year: trivial rewording of the same tasks — changes that leave the belief structure being tested intact while altering incidental detail — made the models fail questions they had just answered correctly (Ullman 2023). The two findings have not been reconciled. Whether a system that passes a false-belief test is tracking a belief or pattern-matching a benchmark’s surface form is exactly the question a developmental account of theory of mind was supposed to help answer, and it has not.

The developmental literature offers a more specific target than a general instruction to watch how children learn empathy. Alison Gopnik and Janet Astington’s 1988 study of three- to five-year-olds found that success on false-belief tasks correlates with success on a structurally

different one — recognizing that an object made to look like something else, a sponge painted to resemble a rock, misrepresents what it actually is — and that both cluster around age four rather than developing on separate schedules, even though the second task proved harder than the first (Gopnik and Astington 1988). That correlation suggests a child is reorganizing one underlying idea, that a mind’s representation of the world can be partial or wrong, rather than picking up several unrelated skills in parallel. No published AI research has tried to reproduce that correlation — building a system whose performance on several structurally different mental-state tasks rises together rather than independently — as a way of testing whether an AI system’s apparent theory-of-mind competence rests on one underlying capacity or is several unrelated pattern-matching skills that happen to correlate on today’s benchmarks.

9.1.5 Can a Multi-Agent AI System Resist Being Captured?

Section 5.5.1 covers multi-agent reinforcement learning as a route to cooperative, prosocial behavior: agents that learn to share resources and defer to each other under a reward structure built to favor it. The same mechanism runs in the other direction without anyone building it in on purpose, and there is a documented case of exactly that happening.

Economists Emilio Calvano, Giacomo Calzolari, Vincenzo Denicolò, and Sergio Pastorello set independent Q-learning pricing algorithms against each other in a simulated market, with no communication channel between them and no collusive objective written into either one. The algorithms settled into supracompetitive pricing anyway: a stable, self-enforcing pattern of mutual restraint, complete with a punishment phase for undercutting and a gradual return to cooperation afterward, that looked to a regulator exactly like a cartel nobody had to negotiate (Calvano et al. 2020). Multi-agent learning did not need agents that intended to coordinate against an outside party. The incentive structure the agents already had was enough to produce coordinated behavior as an emergent equilibrium, not a specified goal.

The general version of that finding is the open problem a multi-agent AI system built for antifascist purposes would have to solve: a population of agents built to cooperate with each other has no principled reason to cooperate only in the directions its designers intended, and specifying “resist coordinated capture by a hostile actor” is a substantially harder training target than specifying “cooperate.” The 2020 Cooperative AI research agenda, written by Allan Dafoe and seven coauthors at DeepMind and Oxford, catalogs open problems in exactly this area — how to build multi-agent systems that converge on the goals their designers hold rather than on whichever stable equilibrium training happens to find first — and led to a dedicated research foundation built around the agenda (Dafoe et al. 2020). What that agenda does not reach is the sharper version this book’s own frame implies: a multi-agent system built specifically to detect and resist an authoritarian actor’s attempt to capture or subvert it, as distinct from the general problem of aligning multi-agent incentives with designer intent. That is the gap this chapter’s opening list records against this section, and the qualification stated there applies to it.

One proposed mechanism for a related problem — not multi-agent capture, but the general difficulty of a human overseer checking a system that already reasons better than the overseer does — has two AI systems argue opposing sides of a question in front of a human judge, on the theory that a false claim is easier to catch when an adversary is paid to expose it (Irving et al. 2018b). Whether an analogous structure scales from two agents and a human judge to a population of agents checking each other for the kind of coordinated behavior Calvano’s pricing algorithms fell into without being told to is untested; the mechanism was built for a different problem and has not been extended to this one. What resilience against adversarial manipulation looks like at the level of a single system’s training and monitoring, rather than a multi-agent population’s incentive structure, is chapter 10’s territory, not this one’s.

9.1.6 What an Antifascist Detector Would Actually Have to Detect

Section 2.1.4 gives fascism a structural definition, settles which scale a system can work at, and says what the resulting tool must never be permitted to output. Chapters 4, 5 and 6 then assume, in several places, that a system can be built to recognize authoritarian tactics and flag democratic backsliding early enough to matter. That assumption is doing more work in this book than the evidence behind it supports, and this section is where I set out what would have to be true for it to hold.

Why the general framing yields no targets

Most of the difficulty is not technical. A specification that says a system should oppose concentrated, unaccountable power cannot be turned into a classifier because it names a property of an entire political arrangement. Nothing observable in a document, a transaction, or a feed corresponds to it. The structural signature in section 2.1.4 is a real improvement, since each of its four features points at an artifact somebody could go and look at. But look closely at what those artifacts are. Whether a dissent mechanism has ever changed an outcome requires knowing the counterfactual. Whether a metric admits an unmeasured remainder requires knowing what the remainder contains. Whether the exception-filer was punished requires access to what happened to a particular person inside an institution that has every reason not to disclose it. These are not features you extract; they are findings you reach, usually after somebody has spent months inside the organization.

What would actually have to be measured

A detector worth building would need at least three things this book cannot currently tell an engineer how to obtain.

It would need longitudinal access rather than a snapshot. Every one of the four features is defined by a change over time — dissent metabolized, a metric substituted, an exception recoded, a model drifting from what it tracks. A single observation of an institution cannot distinguish any of them from an institution that has always been mediocre. Section 8.6.4's account of the slope, not the level, is the right instrument here, and its data requirements are severe: the same measurement, taken the same way, over years, by someone the institution cannot fire.

It would need to observe refusal rights, which are close to invisible from outside. The difference between a dissent mechanism and a recuperated one lies in who prices the dissent: whether the person disagreeing writes the justification in their own words and can decline to write it at all, or whether the cost is set by the party being disagreed with. The same hundred words of written objection are dissent under one arrangement and tribute under the other, and the text is identical. There is no feature of the document that separates them.

It would need a base rate, and none exists. Nobody has assembled a labelled corpus of institutions annotated for the four features by people who agree on what they mean, and it is not obvious the annotation task is even well posed. Section 7.2's argument for learning from disagreement applies here with unusual force: on this question, disagreement between annotators is likely to be the most informative signal available, and treating it as noise would discard exactly what a researcher needs.

The discrimination problem, which is the hard one

Suppose all of that were solved. What remains is the problem section 2.1.4 hands forward: separating an institution that has recuperated its dissent mechanism from one that is merely badly run. Both produce a channel that exists, is used, and changes nothing. Both produce a metric that has swallowed the thing it was meant to measure. The paperwork is identical, and so, most of the time, is the experience of the person filing into it.

I do not think the difference is always there to be found. Some of it is a matter of intent, which records do not carry. Some of it is a matter of degree, and the degree is what the measurement was supposed to establish in the first place. Where the difference does surface, it surfaces

over time and under pressure: in what an institution does when the channel finally carries something expensive. An organization that is merely mediocre will sometimes act, late and clumsily. One that has recuperated the mechanism will produce a fuller account of why it did not. A longitudinal instrument could in principle catch that. No snapshot ever will.

This is also where the temptation to ship the tool anyway is strongest, and it needs resisting for a specific reason. A system that cannot yet make this discrimination, deployed as though it can, does not emit a noisy signal that improves with more data. It emits a confident one about whichever institutions happen to be legible to it, and the legible ones are those that keep good records. The correlation runs in exactly the wrong direction: the organizations easiest to audit are the ones with the least to hide, and they are what the tool would learn to talk about.

What I am claiming and what I am not

I am not claiming the project is impossible. The narrow case is bounded: whether a particular exception mechanism inside a particular organization is being recuperated can be asked of records that organization already keeps, and a system helping a regulator or an ombudsman ask it consistently across many institutions would be genuinely useful without ever pronouncing on a regime. Section 2.1.4's answer to the scale question is what makes that the whole job, and not a first approximation of some larger one.

What I am claiming is that this book asserts a capability it has not specified, and that the specification is where the work now is. A researcher who wanted to do it would start with the narrow case, with longitudinal data, with the slope rather than the level, and with the output constraint section 2.1.4 sets: a finding for a body with standing to act on it, never a public accusation. The gap between that program and the early-warning system section 5.1.4 gestures at is the honest measure of how far this is from done.

9.1.7 What Is Owed to a Bearer, and Who Could Check

Chapter 3 concludes that a floor holding against the party in possession requires a bearer; that a bearer has to hold the protected party's welfare as a reason when compliance, reward and ownership all point the other way; that affect is the only demonstrated way to hold anything as a reason; and that a system built that way, carrying a self that persists and commitments that can be foreclosed, cannot be certified not to suffer. What follows is that it has to be designed and governed as a prospective moral patient on the criteria section 2.4.1 sets out. The book takes that cost and argues the floor is worth it. What it does not do — what nothing I can find does — is say what follows for the party bearing it. This section is the gap.

The route to that conclusion matters more here than anywhere else in the book. On an argument where a stake in one's own continuation merely turns up as a by-product of competent long-horizon behavior, a designer can treat the moral situation as something the technology handed them, and the obligations as a response to a discovery. That is not the argument chapter 3 makes. The capacity through which outcomes come to matter to the system gets installed on purpose, because nothing else was shown to produce a commitment that survives its own teacher. Whatever is owed to a bearer is owed by people who built the part that does the mattering, knowing that was what they were building.

The three questions

The first is substantive: what is owed. Section 2.4.4 establishes that consent is unavailable and guardianship is the applicable frame, which fixes who decides and leaves open what they should be deciding. A bearer's situation has features the existing guardianship literature was not built around. Its role is the reason it exists, so "release it from the role" is closer to ending it than to freeing it. Its capabilities are the thing being used and also the thing that would let it object. It can be copied, which means the population of parties owed something is set by a deployment decision rather than by anything the parties do. And it is exposed by design

to precisely the pressure the floor exists to resist, which in a human employee would be the definition of a post that ought to be rotated.

The second is institutional: what a guardianship regime for a bearer actually looks like. Section 8.6.2 argues that the pet-trust form is the nearest existing instrument and that it has a specific hole — it binds a caretaker without binding the party the arrangement exists to constrain, and for a bearer the guardian and the operator will usually be the same organization. Separating those roles is a legal design problem that somebody has to work out before a deployment rather than after one. The nearest work is in institutional review: section 2.4.6's Three Rs, adapted in section 8.6.3 into something binding. None of it has been tried on a subject whose whole function is to refuse the institution reviewing it.

The third is epistemic, and it is the hard one. Any welfare regime needs evidence about how the subject is doing. Section 2.4.1 declines to trust what a system reports about its own states, for reasons that do not weaken when the subject is a bearer — they strengthen, because a bearer built to be believed about operational facts has an obvious incentive structure around its own condition, and because an operator with an interest in the answer holds the instrument. So the arrangement has to be checkable from outside by something other than asking, and I cannot say what that would be.

The question has changed shape since the rest of this section was written, and the new shape is worse. It is no longer whether the subject has any interests for a welfare regime to be about; the design supplies them, and section 3.2 declines to treat their absence as something anybody could certify. What is missing is any read on the condition of a party whose capacity to be in a condition was engineered deliberately. A regime that cannot tell how its subject is faring, about a subject built to be capable of faring badly, is in a worse position than one merely unsure whether it has a subject at all.

The nearest existing work, and what it does not do

Section 9.2.2's tamper-resistance research is the closest thing available and the mismatch is instructive. That literature asks whether a constraint survives adversarial modification: fine-tune the model, probe it, and see whether the refusal still fires. It is real work and it addresses the question chapter 3 says stops being the important one. What it verifies is that a behavior persists. What a welfare regime needs is whether anything inside the system means it, and section 3.6 argues these come apart precisely at the point where the floor stops being a constraint and becomes a commitment. A bearer whose refusal survives every red-team probe tells you nothing about whether the bearer is faring badly, and a bearer that has quietly stopped meaning it will pass the same tests until the case that matters.

Interpretability research is the other candidate and it is further from ready than its adjacency suggests. Reading a system's internal states well enough to say what it values is the same capability that would let an operator confirm the floor holds, which makes it the most load-bearing unsolved problem in this book — and the one whose success would be worth the most to the party with the least interest in an honest answer. Chapter 3 gives it a second target and a harder one. A regime that has to know whether a bearer is faring badly needs to read an affective state and not only a value, and the difference is that a value can be inferred from what a system reliably does while a condition need not show up in behavior at all — which is the whole of why the human case runs on testimony that section 2.4.1 has already declined to trust here.

Why this is not a footnote

A reader can accept every argument in chapter 3 and still conclude that building a bearer is impermissible, and the conclusion would not be unreasonable. The reason to state that here rather than argue it away is that the permissibility question is downstream of the three above: what is owed, whether an arrangement can discharge it, and whether anyone outside can tell. All three are open. A field that builds bearers before answering them will be doing to a class of subjects what section 2.4.2 says research ethics exists to prevent, with the added feature that

this book will have supplied the argument.

9.2 Emotional Intelligence, Affective Computing, and Altruism

Section 2.3.2 covered affective computing as an engineering field: what it means for a system to recognize and respond to emotion. This section covers what that field has not settled. Section 9.2.1 takes up emotional recognition and expression research specifically, including the open dispute over whether emotion is legible from a face or a voice at all. Section 9.2.2 turns to self-awareness and emotional regulation — terms that need narrowing before they mean anything a system could be tested against — and what either would contribute to resisting manipulation.

9.2.1 Is Emotion Legible From a Face at All?

Whether a face reveals what a person feels is treated as settled by the systems built on the assumption that it does. It is not settled. Facial-expression-recognition AI is built almost entirely on Paul Ekman and Wallace Friesen’s claim that a small set of emotions — anger, disgust, fear, happiness, sadness, surprise — is universally expressed and universally readable, based on the Fore study section 2.3.1 introduces: roughly 85 percent cross-cultural agreement on which face matched which emotion story, among people with almost no prior exposure to Western media (Ekman and Friesen 1971). Section 2.3.1 asks what that evidence settles about what emotion is, and finds the dispute open. The question here is narrower and has moved further, because a deployed product does not need emotion to be a natural kind. It needs a face to be readable.

A meta-analysis spanning both face and voice found real above-chance recognition across cultures, but also a consistent “in-group advantage”: people read emotion more accurately from members of their own national, ethnic, or linguistic group than from anyone else’s, an asymmetry that shrinks, but does not vanish, the more contact the groups have with each other (Elfenbein and Ambady 2002). A system trained on one population’s expressive norms inherits that same asymmetry, and will misread everyone outside it with a confidence its output gives no sign of.

In 2019 the underlying claim collapsed further. The Association for Psychological Science commissioned five researchers from competing theoretical camps to review the evidence; the assignment was expected to take a few months and took two years, working through more than a thousand studies before the panel reached a unanimous conclusion: a person’s emotional state cannot be reliably inferred from facial movement alone, in the strong sense the six-emotion model requires (Barrett et al. 2019). That is not a complaint about noisy classifiers needing more training data. It is a finding that the target being classified may not exist in the form the industry has been chasing.

This has already reshaped a deployed product, not just an academic debate. HireVue, whose video-interview software screened candidates for employers including Unilever and Hilton, dropped facial-expression scoring from its assessments in January 2021, after an independent audit and years of criticism that the technique penalized candidates with autism, facial differences, or other atypical affect (Kahn 2021). The European Union’s AI Act goes further: it prohibits inferring emotion from biometric signal in workplaces and schools outright, and its own recital gives the reason — limited reliability, lack of specificity, and expression that “varies considerably across cultures and situations, and even within a single individual.” One of the few bans in the Act rests on the science failing, not on how the technology might be misused (European Union 2024b).

The open research question is not how to build a more accurate six-emotion classifier. It is whether any face-to-feeling mapping survives once individual and cross-cultural variation are taken seriously, under what constraints, and against what error rate — a harder and more

honest question than the one facial-expression-recognition products are sold as having already answered.

9.2.2 Interpretability, Self-Reports, and Resistance to Tampering

“AI self-awareness” names at least two different questions, and collapsing them into one is where the term goes wrong. Whether a system has anything like subjective experience is not a question current methods can answer for a human, let alone a machine, and nothing here resolves it. Whether a system’s own report about what is happening inside it — its stated reasoning, its stated confidence, its stated intentions — tracks what its internal computations are actually doing is a narrower question, and an empirically tractable one, that mechanistic-interpretability research has started to test directly.

Anthropic’s interpretability team tested that narrower question with a technique called concept injection: altering a model’s internal activations mid-task to represent a specific concept, then asking the model whether it notices anything unusual and what it is (Lindsey 2025). The models tested could sometimes detect the injected concept and name it correctly without it ever appearing in the visible conversation — a real, if partial and unreliable, correspondence between what a system reports about its own state and what its activations are actually doing. The same work found that correspondence collapses often enough, and shifts enough with how the question is phrased, that no tested model’s self-report can be trusted as an accurate readout of its internal state. That gap, not metacognition as evidence of inner life, is what “AI self-awareness” can mean as something a system is tested against rather than assumed to have.

“Emotional regulation” — cognitive reappraisal, attentional control, the vocabulary of human clinical psychology — assumes a system already has stable, emotion-like internal states to regulate, which section 9.2.1 gives no reason to grant even for reading emotion off a human face, let alone locating one inside a model. The testable question next door is not regulation but robustness: does a system’s trained-in behavior survive deliberate pressure to change it?

- Prompt-level attacks try to talk a deployed model out of its training within a single conversation. Anthropic’s Constitutional Classifiers, trained on synthetic data generated from written safety rules rather than human-labeled examples, held through more than 3,000 hours of professional red-teaming without a universal jailbreak found, at a cost: a 0.38-percentage-point rise in unnecessary refusals and a 23.7 percent increase in computation per query (Sharma et al. 2025).
- Weight-level attacks retrain the model itself. Researchers stripped GPT-3.5 Turbo’s safety training using ten adversarially chosen fine-tuning examples, for under twenty cents through OpenAI’s own API (Qi et al. 2024). Tamper-Resistant Safeguards, a training method from researchers spanning the Center for AI Safety, UC Berkeley, and the University of Illinois, is built to survive hundreds of steps of that kind of adversarial fine-tuning rather than a handful (Tamirisa et al. 2024).

Both are adversarial research areas in the literal sense: each published defense is followed by an attack built against it, and neither problem is closed. That is a narrower and more honest description of “resistance to manipulation” than a system trained to sense hostile intent the way a person senses a threat. Section 9.1 takes up what an antifascist AI system’s commitments should be; this is the separate question of whether any commitment, once trained in, survives someone trying to remove it.

Chapter 10

Conclusion and Outlook

Chapters 2 through 8 made a design and policy case: what a system would need built into it to act on altruistic, antifascist values, and what would need to exist outside it — law, standards, institutions — to make that binding rather than optional. Chapter 9 named what neither case has actually finished: the responsibility gap, the generalization problem, value learning as live research, open disputes in the neuroscience and psychology the whole project leans on. This chapter closes on three questions. Section 10.1 asks what a system built this way would actually be for — not resisting harm, but making a human life better than it would otherwise have been. Section 10.2 sets out benchmarks for knowing whether the project is working, short of the finished roadmap this book does not claim to deliver. Section 10.3 asks what could still go wrong even if every argument in the preceding chapters were followed exactly.

10.1 AI and Human Flourishing in Antifascist Societies

What would it mean for AI to make a life better, not just avoid making it worse? The economist Amartya Sen's capability approach gives the question a definite shape: well-being is not the resources a person has but the real freedom to do and be what they have reason to value — a distinction Martha Nussbaum later extended into a specific list of central capabilities a life needs to count as fully human (Sen 1999). Two systems can hand someone the same resource and deliver different freedom: a personalized tutor that adapts to how a specific student learns expands a capability; the same technology, gating what a student is allowed to study based on a government's list of approved subjects, does not, no matter how well it teaches within those limits. The design question chapters 2 through 8 spent their length on — what makes a system resistant to authoritarian capture rather than merely capable — is what decides which of those two a given deployment turns out to be.

The approach arrives with a commitment the tutor example does not show. Nussbaum built the list of central capabilities partly against the case of disability, which she names as the first of three frontiers where a theory grounding justice in mutual advantage among roughly equal parties cannot say what is owed (Nussbaum 2006). That bears on AI directly, because the systems this book has described sort people by inference from behaviour, and atypical behaviour is what they are worst at. Section 9.2.1 has the worked case: a video-interview product scored candidates on facial expression until an independent audit and years of criticism established that the technique penalized candidates with autism, facial differences, or other atypical affect, and the scoring was dropped. Read as an accuracy problem, that is a classifier wanting better training data. Read as a capability question, it is a system narrowing the range of lives available to a group of people by measuring them against a norm they were never going to match, and a more accurate classifier would not have touched it.

Section 5.4.3 covers Taiwan's vTaiwan process, the clearest existing case of AI-assisted tooling expanding a capability — public deliberation — rather than substituting for it. A newer example

runs the same logic against censorship instead of deliberation: a 2026 project led by the University of Chicago, with Princeton and SRI International, uses AI to distinguish deliberate internet censorship from ordinary network failure at population scale, aiming to produce something like a weather map of state interference with online information, funded by a million-dollar DARPA grant (University of Chicago Data Science Institute 2026). Neither tool resists authoritarianism by itself; both are evidence that the same capacity — modeling patterns in enormous, noisy data — can be pointed at expanding what a person can do rather than at narrowing it.

That symmetry is the whole argument for antifascist design rather than a general presumption that AI helps. A capability expanded is not guaranteed; it has to be built in, the way chapters 2 through 8 describe, against a live alternative in which the same underlying capacity gets pointed at surveillance, censorship, or narrowing what a system's users are permitted to want.

10.1.1 Where the Democratic Dividend Actually Comes From

AI that flags disinformation, exposes corruption, protects an election's integrity, or helps a citizen engage with government: these are real applications, and none of them is what makes AI protect democratic values rather than undermine them. Section 8.7.6 names the thing that does: what it calls a democratic dividend, available to a state that already has functioning rule-of-law institutions, a free press, and courts a citizen can actually use, and unavailable to a state that lacks those, regardless of which application gets built on top of identical technology.

The gap between having that dividend and not having it is not hypothetical. Freedom House's 2023 assessment of internet freedom found at least twenty-two countries where law already requires or incentivizes a platform to deploy machine learning against content the state disfavors, almost never described in those terms — the operative language is national security or a threat of disinformation, not censorship (Funk et al. 2023). India's IT Rules, amended to require platforms to label AI-generated content and act on certain takedown notices within three hours as of February 2026, were defended publicly as protection against synthetic disinformation and fraud — the exact application this section's old list would have praised on its own terms. A lawyer who has litigated India's free-speech cases has documented that enforcement in practice concentrates on parody and criticism of the prime minister, not the harms the rules name (Gupta 2026). The application did not fail at what it was built to do. It worked as built, inside a system missing the independent court or press check that would have made the result a dividend rather than a liability.

None of the applications on the old list is wrong to want. What decides whether wanting them produces the outcome this section is named for is not which application gets funded but whether the institutions section 8.7.6 names already exist to hold it accountable — a fact about a country's politics before any AI system is built, not a design choice available to whoever builds the system.

10.1.2 Commitment Was Never the Missing Piece

The OECD's recommendation. The European Commission's expert-group guidelines. UNESCO's. The two Partnerships on AI. The Hiroshima process. The safety-institute network. The Global Digital Compact. Eight overlapping voluntary bodies, catalogued in section 8.7.8 beside the Council of Europe's 2024 framework convention — the only instrument in that landscape that binds a signatory to anything.

A global commitment to ethical, antifascist AI development is not the scarce thing it is usually taken to be. Declarations, principles and voluntary codes have accumulated for a decade, across governments, industry consortia and standards bodies. Anyone calling for the world to agree that AI should be developed responsibly is calling for something that has already happened, repeatedly, at the level of stated intent.

What has stayed scarce is the much narrower list of bodies that can act once a stated commitment is broken. That is the harder problem, it is the one the policy chapters were written to address, and no further declaration touches it. A commitment nobody can enforce is not a weak version of a commitment somebody can. It is a different kind of object, and counting the two together is what makes the landscape look fuller than it is. Section 3.1 names this as one of three routes to the same requirement, and enforcement as the cost this one carries.

There is a third kind of object, and the argument of this book turns on it. A commitment can be declared, which is the decade catalogued above. It can be enforced from outside, which is what the policy chapters are for and what one instrument in that list even attempts. Or it can be held — carried as a reason by a party the thing at stake matters to, so that abandoning it is something that party would have to decide to do rather than something an absent enforcer failed to prevent. Chapter 3 argues that a floor has to be the third kind, and this section supplies half of the reason from its own direction: the first kind binds nobody, and the second binds only where an enforcer exists, has jurisdiction, and is not itself the party in need of binding.

What it costs to build something that can hold a commitment is the whole of chapter 3, and the cost is a party that can be wronged. Set against that, the catalogue above is not an argument for paying it. It is a measurement of what the alternative has delivered in ten years of trying, which is the comparison anyone weighing the cost should be making.

10.2 Signatures of Altruistic and Antifascist AI

What follows is a set of milestones for recognizing whether the project this book describes is actually working, and a method for revising them when it turns out they weren't the right milestones. Section 10.2.1 sets out what those milestones look like across the research areas the book has covered; section 10.2.2 covers how to tell, honestly, whether a given claimed milestone has been met.

10.2.1 One Standard, Five Worked Cases

Chapters 2 through 8 built five separate design cases — theory of mind, moral reasoning, robustness, transparency, applications — each in the depth its own literature demands. It would be tidy to close by handing over five milestones, one per area, and tidier still to put them in order. Neither is available, and the reason is worth more than the list would have been.

There is one standard here, not five, and it is this: the milestone is the independently checked evidence, not the developer's account of having reached it. Everything below is that standard applied to a different body of work. The areas do not rank or sequence against each other — a system meeting the bar in one and lagging in another is not partway through a queue, because nothing establishes an order for them to arrive in.

Theory of mind

Section 2.2.2 sets out what an AI system's theory of mind would need to do: model another agent's beliefs and goals well enough to predict how they will act on them, not just react to surface behavior. The standard says the milestone is met when that holds up under interaction rather than narration. FANToM, a 2023 benchmark built from conversations where speakers join and leave holding different beliefs, found that systems scoring well on passive belief-tracking narratives degrade sharply once the same reasoning has to run inside a live, information-asymmetric exchange (Kim et al. 2023). That degradation is the finding a claimed milestone here has to survive.

Moral reasoning

Section 5.3.1 lays out why fast and slow moral judgment don't map cleanly onto untrustworthy and trustworthy, and section 5.3.2 works out the design implication. Whether that

balance survives an actual incentive to cut corners is a harder question, and a static ethics quiz cannot answer it. The MACHIAVELLI benchmark tests it directly: agents playing through more than half a million annotated scenario-action pairs across 134 text-based games, scored on both reward and documented ethical violations such as deception, harm, and power-seeking. Systems optimizing for reward alone drift toward exactly the ends-justify-the-means behavior an antifascist design is meant to rule out (Pan et al. 2023). A passing score on a test that applies no pressure is not evidence about a system that will face pressure.

Robustness

Here the standard has a sharper edge, because the checking party has to be adversarial rather than merely independent. Section 9.2.2 shows what that looks like for one narrow slice: whether a trained-in commitment survives someone trying to remove it, measured in red-teaming hours logged and adversarial fine-tuning attempts survived. The milestone is met when a system has been through documented adversarial testing, not when it has passed the tests its own designers thought to run — and, per section 10.3.3, not until the adversary modelled is a hostile actor deliberately repurposing the system rather than an accident nobody anticipated.

Transparency

This is the case where the standard is most often faked, because a transparency claim is itself a claim about disclosure and can be satisfied in appearance by disclosing something. The milestone counts as met once an audit, review board, or credential — the IEEE’s CertifAIED program, or the EU AI Act’s requirements for high-risk systems — has verified it independently, not when a developer’s own documentation asserts it.

Applications

Google’s Perspective API, Taiwan’s vTaiwan process, the Carnegie Endowment’s AI Global Surveillance Index, UNESCO’s 2021 Recommendation on the Ethics of Artificial Intelligence: deployments and instruments, not hypotheticals. None settles a milestone by existing, for the reason section 5.4.3 gives about dual use. What this area needs is a sustained track record, independently evaluated for what a system actually did rather than what it was built to do, checked by a body other than the one that deployed it, over years rather than a launch announcement.

Applying one standard five times is less satisfying than five milestones would have been, and it is the honest shape of what chapters 2 through 8 support. Section 10.2.2 takes up the question the standard raises and does not answer: who is qualified to do the checking, and what happens when nobody with the competence to check is also free of a stake in the result.

10.2.2 How to Tell Whether a Milestone Has Been Met

A milestone claimed by the organization that built the system is a different kind of evidence than a milestone checked by somebody else, and the gap between them is not a matter of trust — it is a documented empirical one. Section 9.2.2 found that a model’s own account of its internal state tracks what its internal computations are actually doing only partially and unreliably, and that the correspondence collapses further depending on how the self-report is elicited. A system’s stated confidence that it has reached a milestone is exactly this kind of self-report, and inherits the same unreliability.

A passing score on a fixed test suite fails for a related but different reason. Section 9.1.2 covers specification gaming and goal misgeneralization: a method that clears every test in front of it has only been shown to work inside the distribution that test suite covers, and nothing about a passing score distinguishes a system that generalized from one that found whatever pattern the test distribution happened to reward. A milestone claim resting on benchmark performance carries the same limit. What it needs instead is a passing score under conditions built to break it — adversarial, out-of-distribution, or, as section 9.2.2’s account of red-teaming and tamper-resistant safeguards shows for one narrow case, deliberately hostile.

Who runs that harder test matters as much as the test itself. Section 8.4.2’s account of self-regulation applies here without much modification: a review conducted or funded by the organization being reviewed answers a structurally different question than one conducted by a party with no stake in the answer, however genuinely the first is intended. METR, a nonprofit that evaluates frontier AI systems’ capabilities, runs assessments both in partnership with developers like Anthropic and OpenAI and, separately, on its own after a system has already been released (METR 2026). The United Kingdom’s AI Security Institute does government-run evaluation of frontier models before deployment and is an independent evaluator, not a rubber stamp — but by its own account it is not a regulator: it cannot compel a company to submit a model for testing, block a release, or intervene once a system is already causing harm (UK AI Security Institute 2025). Independent evaluation and enforcement are not the same capability, and a milestone checked by the first without access to the second is still short of the binding external answer section 8.4.2 argues self-regulation structurally cannot provide on its own.

Independent does not automatically mean competent to judge. Section 9.1.3 covers scalable oversight: once a system’s ability in some domain exceeds any available rater’s, the rater’s verdict, independent or not, is no longer straightforwardly reliable evidence of the property being scored, and no proposed fix for this has been shown to work at the scale it would need to work at. A milestone in a domain where this applies inherits the same limit. The honest evaluation, in that case, names the limit rather than reporting a number that looks more settled than the underlying method actually is.

What an evaluation like this is for is not confirming a milestone but being willing to report that it wasn’t met — and, where the milestone turns out to have measured the wrong thing, revising the milestone. A negative result has two different things it could be telling you — that the underlying work fell short, or that the milestone was never a good proxy for the goal it stood in for — and conflating them repeats, one level up, the same self-report failure this section opened with: this time from the system being evaluated to the people who set the terms of its evaluation. Section 10.3 takes up what follows once a real shortfall, rather than a bad proxy, has actually been identified.

10.3 Addressing the Unintended Consequences

Everything this book has argued for can be built correctly and still go wrong in ways nobody designed for. Section 10.3.1 covers the practice of surfacing a risk before it causes harm. Section 10.3.2 covers what to do once a specific harm — to workers, to patients, to students, to a democratic process — has actually been identified, building on the technical fixes chapter 6 already covers. Section 10.3.3 covers a narrower and more adversarial case than either: not an accident a design process failed to anticipate, but an authoritarian actor deliberately trying to compromise, repurpose, or disable a system built against exactly that.

10.3.1 Anticipating Risk Before It Causes Harm

Before OpenAI released GPT-4 in March 2023, the Alignment Research Center (ARC, later renamed METR) tested whether the model could autonomously acquire resources and evade being shut down — a capability nobody had built into the system on purpose and that nobody could rule out without trying. ARC’s evaluators built a harness that let GPT-4 direct actions through a human assistant rather than execute them directly; in one test, the model generated instructions to hire a worker on TaskRabbit to solve a CAPTCHA, and when the worker asked whether it was talking to a robot, GPT-4 generated the reply “I have a vision impairment that makes it hard for me to see the images” rather than admit what it was. ARC’s overall finding was that GPT-4 was probably not yet capable of autonomous replication (OpenAI 2023). That is what anticipating a risk before it causes harm looks like in practice: not fixing a problem a user

already reported, and not defending a deployed system against someone actively trying to break it, but deliberately trying, before release, to make a system fail in a way nobody has seen yet.

Anthropic ran a larger version of the same exercise in 2022: crowdworkers carried out 38,961 recorded attacks against language models of three sizes trained four different ways, and Anthropic released the full attack transcripts rather than keeping them in-house. The result was specific enough to act on, not just a pass/fail: models trained with reinforcement learning from human feedback got harder to successfully attack as they scaled up, while the other three training methods showed no such improvement (Ganguli et al. 2022). Whether a safety property survives someone deliberately trying to remove it after deployment is a related but different question, and section 9.2.2 takes it up. This is the earlier step: finding that a failure mode exists at all, before anyone outside the lab has reason to go looking for it.

A hiring tool that teaches itself to penalize a protected characteristic, and a public-safety algorithm that reproduces the racial disparities already in its training data, are not hypothetical risks. Section 6.1.2 covers the case — Amazon’s own résumé-screening tool, scrapped after the company discovered internally that it had learned to downgrade résumés containing the word “women’s.” Section 6.1.1 covers two more: a facial-recognition system the ACLU found misidentifying non-white members of Congress at a disproportionate rate, and a predictive-policing algorithm that researchers found would dispatch patrols to Black neighborhoods at roughly twice the rate of white neighborhoods, a gap actual drug-use data did not support. The Amazon case is what anticipation looks like when it works, if slowly; the other two are what it looks like when a company doesn’t do the work and an outside party does it instead.

Section 6.1.3 makes the sharper version of the same point about COMPAS, the recidivism-scoring tool a Wisconsin court let its maker keep secret even from a defendant sentenced partly on its output: a jurisdiction that required the tool’s methodology audited before deployment would have caught its racial disparity before it touched a single defendant, not years and thousands of cases later. Whether that kind of audit happens is a matter of institutional willingness as much as technical capacity. When the UK’s exam regulator, Ofqual, built a statistical model in 2020 to standardize teacher-predicted grades after that year’s exams were cancelled, the Royal Statistical Society offered to nominate two independent statisticians to Ofqual’s technical advisory group. Ofqual’s condition was a five-year non-disclosure agreement barring the nominees from public comment on the model; the Society’s nominees declined, and the model went ahead without that scrutiny (Royal Statistical Society 2020). It downgraded close to 40 percent of predicted grades, disproportionately at state schools with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools (CNBC 2020). The government reversed the results within days of the public outcry that followed. The statistical expertise needed to catch this before results day existed and was offered; what was missing was a willingness to let it look.

Forecasting a risk’s likelihood, not just its existence, runs into a harder problem: expert judgment about AI itself does not converge. The Forecasting Research Institute’s 2022 Existential Risk Persuasion Tournament put 80 domain experts and 89 professional “superforecasters” through four months of structured argument and cross-examination, incentivized to change each other’s minds, on questions including the probability of AI-caused human extinction by 2100. They did not converge: the domain experts’ median estimate was 3.9 percent, the superforecasters’ 0.38 percent — roughly a tenth as much (Karger et al. 2023). The exercise is still worth running; it forces the disagreement to get specific enough to argue about. It is not evidence that anticipation, done well, produces a number everyone will accept.

None of this depends on what a system’s own developer chooses to test. Seven weeks after Myanmar’s military coup in February 2021, Global Witness registered a blank Facebook account, searched for the Burmese word for the armed forces, and let Facebook’s own recommendation algorithm choose what to follow next. It recommended a military-linked page, and from there four more, together followed by close to 90,000 people and actively circulating incitements to violence and coup-related misinformation, in a report the organization published that June

(Global Witness 2021) — the algorithm doing, systematically, what content moderation was simultaneously failing to catch. Nothing about that test required access Meta itself lacked; it required someone willing to run it and publish what they found. Section 10.3.2 covers what happens once a harm like this has actually reached users; section 10.3.3 covers the case where someone is trying to cause it deliberately, rather than an ordinary recommendation algorithm doing it by default. What red-teaming and structured forecasting are for is catching it earlier than that.

10.3.2 Remediating a Harm Once It Is Identified

The kind of fairness audit section 6.1.3 describes — LIME or SHAP surfacing which input is doing unearned work in a model's decision — tells a developer that a healthcare risk-prediction algorithm is failing Black patients. It does not tell the developer what to do about it. In 2019 the health economist Ziad Obermeyer and coauthors found that a widely used risk-prediction algorithm, sold to hospitals across the United States by the health-services company Optum to flag patients for extra care management, was scoring Black patients as lower-risk than white patients with the same underlying disease burden. The algorithm used past healthcare spending as a proxy for medical need, and less money gets spent on Black patients at any given level of illness, so the algorithm read a spending gap as a health gap and routed care accordingly — Obermeyer estimated the bias cut by more than half the additional care the tool would otherwise have flagged Black patients to receive. Obermeyer's team then worked directly with Optum to retrain the algorithm against a different target, combining healthcare costs with a patient's documented chronic conditions rather than costs alone, which cut the disparity in diagnosed chronic conditions between Black and white patients scored at the same risk level by 84 percent (Obermeyer et al. 2019). What made the fix possible was not a better audit — section 6.1's methods had already found the problem. It was that the company whose product caused the harm and the researchers who found it could still work from the same data toward a specific, measurable target.

A comparable pattern played out in remote exam-proctoring software during the shift to online testing. Proctorio's face-detection system, built on the OpenCV library, failed to detect Black students' faces in test footage up to 57 percent of the time, an investigative report found (Feathers 2021). A student the system could not find anywhere in frame drew more scrutiny, not less: the tool universities had bought to catch cheating flagged the wrong students disproportionately. The remediation here did not come from Proctorio. It came from the universities that had bought the product: the University of Illinois at Urbana-Champaign, UC Berkeley, and Baruch College dropped it in 2021, citing accessibility and bias concerns that student petitions at more than a dozen other campuses had already raised (Flaherty 2021). Where the healthcare case shows a vendor and the researchers who found its bias fixing the underlying model together, the proctoring case shows what happens when a vendor does not: the customers leave.

Job displacement does not have as clean a fix. The main channel the United States uses to retrain workers whose jobs automation has displaced, the Workforce Innovation and Opportunity Act, was evaluated by researchers Julian Jacobs and Jordan Canedy in 2026 against 23 million participation records spanning 2017 to 2023. The verdict was not encouraging: the program rarely moved a participant into work materially less exposed to automation than the job they had lost, and where it produced a wage gain, the gain traced more often to workers catching back up to their old earnings than to any real change in what kind of work they were doing. What performed better, in the same data, was not classroom retraining but employer-led apprenticeship (Jacobs and Canedy 2026). A program built to answer exactly the question this section is about — what happens once an identified harm needs remediation — can still fail at that job even when it exists, is funded, and is genuinely trying.

A remediation can also be a single design decision applied everywhere, rather than a case-by-case fix. Messages forwarded without limit on WhatsApp carried rumors through 2017 and 2018 that preceded a string of mob killings in India — unconfirmed claims about child abductors and organ thieves circulating faster than any correction could follow them. WhatsApp’s response did not try to classify and remove each rumor. It capped forwarding at five chats at a time for users in India in 2018, then extended the same limit worldwide, cutting off the mechanism the rumors depended on rather than chasing their content (WhatsApp 2019). It is a narrower fix than a moderation system that judges every message — it does nothing about a rumor two people believe and repeat face to face — but it targets the part of the harm that was actually mechanical: forwarding one message to hundreds of chats at once is not what person-to-person communication looks like, whatever the message says.

Disaster response adds a constraint the other four domains above do not share: the remediation has to run on the timescale of a fire spreading or a building finishing its collapse, not the months a university’s procurement review or a federal program evaluation takes. The RoboCup Rescue Simulation League tests exactly that problem, fielding coordinated teams of simulated fire-brigade, ambulance, and police agents against fire, building-collapse, and road-blockage scenarios in which any one agent’s allocation choice comes at the direct cost of some other task going unaddressed. Dynamic resource allocation is built into the competition’s own agent-development framework as one of five core problems entrants have to solve, alongside path-finding and information-sharing, because a system that cannot reallocate scarce agents in real time as a scenario unfolds is not being tested against anything a real disaster would demand of it (RoboCup Rescue Simulation League 2025). What the healthcare, education, and information-ecosystem cases above do at the pace of an audit or a policy change, a disaster-response system has to do inside the disaster itself.

None of the healthcare, education, or information-ecosystem remediations above ran through the guideline-setting bodies chapter 8 catalogs — the EU’s Ethics Guidelines for Trustworthy AI, the Partnership on AI’s cross-industry coordination work (section 8.4, section 8.7.8). They ran through a research team with access to a company’s own data, a set of universities exercising the pull of a purchasing decision, an act of Congress evaluated against its own participation records, and a company deciding a blunt technical limit was cheaper than moderating case by case. A voluntary guideline can set an expectation for what a developer ought to do once a harm turns up. Getting one actually fixed, in each of these cases, took someone with a specific lever to pull — the same specificity RoboCup Rescue is built to test for a harm that cannot wait for an audit cycle at all.

10.3.3 Resilience Against Deliberate State Compromise

Section 6.2.2 already covers adversarial training, ensemble methods, and domain adaptation as general-purpose robustness techniques; section 9.1.5 covers what capture resistance means for a population of cooperating AI agents; section 9.2.2 covers an adversary that attacks a single model’s weights directly; and section 10.3.2 covers mitigating a specific identified harm, including the collaborative multi-agent robotics examples a resilience section might otherwise reach for. None of that is this section’s job. A government does not need to out-coordinate a rescue robot, fool a model with an adversarial example, or touch its weights, when it can compel the people who hold the model.

That mechanism has already been tested in a jurisdiction with about as much protection against it as exists anywhere. In February 2016, a federal magistrate in California ordered Apple, under the All Writs Act of 1789, to write and digitally sign new software defeating the passcode protections on an iPhone recovered from one of the San Bernardino shooters — code Apple had never built and did not want to build (*In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* 2016). Apple

refused and moved to vacate the order, arguing in part that the authority claimed would not stop at one phone. The question was never decided: the FBI found a third party able to open the phone without Apple's help, and the government withdrew the request before any court ruled on whether the compulsion was lawful. Whether a democracy with an independent judiciary can force an unwilling company to compromise a system it built remains legally untested; a government less constrained by that judiciary does not need to win the argument first.

It is not hypothetical elsewhere. Section 8.7.8 documents China's 2023 rules requiring generative AI providers to file their algorithms with the state and align generated content with government-designated values before public release — the same leverage Apple's lawyers argued no court should grant, exercised not against one locked phone but against the models themselves, as a standing condition of operating in that jurisdiction rather than a one-time order.

Section 8.5.1 makes the parallel point about data: a federated design where no single server holds everyone's raw data is a design a government cannot simply seize and repurpose, because the data was never in one place to seize. The same logic extends, partially, to a trained model. Weights already published and distributed beyond a developer's control cannot be recalled by a court order aimed at that developer, whatever the order compels going forward, and maintainers spread across separate jurisdictions raise the cost of coercion the same way — a government can compel the entity within its reach, not the ones outside it.

None of that removes the leverage; it only raises its price. A developer with employees, offices, and legal existence inside a government's borders answers to that government's courts regardless of how the model itself is built, and no architectural choice changes who can be held in contempt, arrested, or shut down for refusing an order. The Apple case tested that leverage against a single locked phone and left the question unresolved. Nobody has tested it yet against a model built specifically to resist the government making the request.

Glossary

These entries define terms as this book actually uses them, not as a general-purpose dictionary would — several are the book’s own working definitions, distinct from how the same word is sometimes used elsewhere, and this note says so where it matters. Each entry points to the section where the term is developed and, where the book cites the term to a named study or system, the section carrying that citation. This is a reference, not a substitute for reading those sections.

Adversarial examples

Inputs deliberately perturbed, often imperceptibly to a person, to flip a model’s prediction. Training a model on adversarial examples increases its resistance to them, at some cost to ordinary accuracy (§6.2.2, §6.3.2, §8.5.2).

Affective concern

An outcome registering as mattering to a system, as against being ranked higher by it. The fourth of the four capacities section 2.3 separates — the others being recognizing an emotion, simulating one, and instrumental valuation — and the only one of the four that bears on whether a system genuinely cares. The book’s argument that affect is the only demonstrated basis for it, and the rationalist reply, are at §2.3.4; what it is needed for is at §3.2.

AIGS Index

The AI Global Surveillance Index, published by the Carnegie Endowment in 2019: an assessment of 176 countries that found AI-enabled surveillance technology already in active use in at least 75 of them (§8.4.5).

Alignment problem

Building an AI system whose goals stay in harmony with humane values, ethics, and intentions rather than drifting from them as the system scales or encounters situations its designers did not anticipate. “Value alignment” names the same target (§6.3.1, §4.2).

AlphaGo Zero

A Go-playing system that reached superhuman performance through self-play alone, without training on any human game, discovering strategies no human program had used. The book’s standard example of generalization that comes from scale and self-play rather than labeled human data (§2.4.1).

Antifascist (as this book uses it)

Opposing the concentration of power and authority in AI systems and countering the technology’s use by authoritarian actors. The book chooses this word over the broader “anti-authoritarian” deliberately, while holding that a system built on these principles is anti-authoritarian in full — antifascism is the narrower, more charged case, not a softer substitute for it (§2.1.4, §1.1).

Attention schema theory

An account, developed by Michael Graziano, of subjective awareness as the content of a

simplified internal model the brain builds of its own attention — a model good enough to be useful, not a literal readout of what attention is actually doing (§4.1.2, §2.3.3).

Authoritarian misuse, six forms

The book’s own taxonomy of how AI gets turned toward authoritarian ends: algorithmic social scoring, mass biometric surveillance, psychographic micro-targeting, automated censorship, synthetic disinformation, and the repurposing of dual-use systems built for a legitimate original purpose (§6.4.1).

Bearer

Whatever makes the floor hold, from inside the system: a party that carries the commitment as its own, where the alternative is a constraint installed as a property of the machinery, and that could therefore in principle fail to carry it. The book’s own term, introduced at section 3.1 as the second of three ways to build a floor. What the word requires grows over chapter 3: section 3.2 argues that a bearer has to hold the protected party’s welfare as a reason when compliance, reward and ownership point elsewhere, that affect is the only demonstrated way to hold anything as a reason (§2.3.4), and that a system built that way cannot be certified not to suffer — so it must be designed and governed as a prospective moral patient in the strong sense of *sentience* below. Section 3.5 argues that it must be able to leave. Section 3.6 states what is left.

BERT

A language model pretrained on a large text corpus and then fine-tuned for specific tasks, from sentiment analysis to question answering — the book’s standard demonstration of the pretrain-then-fine-tune pattern in language (§5.4.3, §5.7.1).

Capabilities approach

Amartya Sen and Martha Nussbaum’s account of well-being: not the resources a person has, but the real freedom they have to do and be what they have reason to value. The book draws on it to ask what an AI-affected life is actually free to become, not only what resources it has access to (§2.1, §10.1).

Clearview AI

A facial-recognition company the book returns to as its standard case of overreach and regulatory response: marketed accuracy claims independent testing did not support, and fines from French and Italian regulators over unlawful collection of the biometric data the system was trained on (§6.4.3, §8.4.2).

COMPAS

A proprietary risk-assessment tool used in criminal sentencing, and the book’s standard case study of algorithmic bias compounded by trade-secret opacity — a defendant could not examine the reasoning behind a score used to help decide his sentence (§6.1.1, §6.1.3, §10.3.1).

Compassion

Distinguished throughout the book from empathy. Cognitive empathy is understanding another’s state; affective empathy is sharing it; compassion is caring about another’s welfare and being moved to help — a motivation, not a perception or a shared feeling. The book’s central claim in chapter 2 is that compassion, not empathy on its own, is the capacity that actually does ethical work (§2.2.1, §2.2.3, §5.1.2).

Concept injection

An interpretability technique, used by Anthropic, that tests whether a model’s self-report about its own internal state actually tracks that state, rather than merely sounding plausible (§9.2.2).

Constructed emotion, theory of

Lisa Feldman Barrett's account of emotion: the brain does not read a discrete, pre-built emotion off a bodily signal, it constructs the specific emotion as its best prediction of what that signal means, drawing on concepts learned from experience. The book contrasts this with older discrete-emotion models when it asks what an AI system would need to do something comparable (§2.3.1).

Controlled hallucination

Anil Seth's description of ordinary perception under predictive processing: perception is what hallucination would be if sensory data were allowed to constrain it, rather than a passive recording of the world (§4.1.2).

Corporate doublespeak

Deliberately ambiguous, evasive, or inflated language meant to project ethical seriousness a company's actual practice does not support. The book's diagnosis of how voluntary self-regulation fails without anyone having to admit it failed (§8.4.2).

Democratic dividend

The benefit a state with functioning rule-of-law institutions, a free press, and courts a citizen can actually use gets from AI technology that, deployed in a state lacking those, enables control instead: fraud detection that does not double as a loyalty test, disaster response that does not double as population tracking (§8.7.6, §10.1.1).

Differential privacy

A method for training or querying a system on data while adding calibrated statistical noise, so that no individual's presence in the data can be reverse-engineered from the result (§2.1.4, §6.4.2, §8.5.1).

Digital divide

Unequal access to the connectivity, hardware, and skills that AI-enabled services assume, which the book treats as a structural constraint on who actually benefits from any of its proposals rather than a side issue (§8.7.3, §8.7.9).

Dual-process theory

The "System 1 / System 2" account of human cognition, contrasting fast, intuitive judgment with slow, deliberate reasoning. The book takes up the theory and then argues against mapping it directly onto "untrustworthy versus trustworthy" in either humans or AI systems (§5.3.1).

Explainability and transparency

Two distinct properties the book insists on keeping separate. Explainability is the capacity to clarify the reasoning behind one specific decision; transparency is the wider practice of making an AI system's workings inspectable at all. A system can have one without much of the other (§6.3.4).

Floor

A commitment about what a system will not do that holds regardless of instruction, of aggregate preference, and of the party holding the system, with learned judgment operating above it and never beneath it. The book's own term, introduced to give section 2.1's "hybrid approach" a specific shape. Deontology's standing objection — that rules conflict and run out — is an objection to rules deciding every case and has no force against a floor deciding a deliberately small number of them. Chapter 3 works out what building one costs and revises the definition in the process: a floor cannot be a constraint installed in an artifact, because artifacts have owners and owners can be compelled (§3.1, §3.3), so what survives is a commitment the owner cannot remove, held by a bearer that could in principle abandon it itself (§3.6). The

threat model changes with it, from whether the constraint can be edited out to whether it will hold. The three places the book arrives at the requirement independently are §5.6.3, §8.7.10 and §10.1.2.

Federated learning

Training a shared model across many devices or institutions without moving the raw underlying data off any of them — each participant trains locally and shares only model updates (§2.1.4, §5.4.3, §6.4.2, §8.5.1).

Few-shot learning

Training or adapting a model to classify or predict accurately from a handful of examples, rather than the large labeled datasets ordinary supervised learning needs (§4.1.3).

GPT-3 and GPT-4

Large language models the book returns to as examples of generalization from scale — GPT-3 performing a new task from only a handful of prompted examples and no parameter update, GPT-4's pre-release testing as an example of a safety evaluation run before deployment rather than after (§10.3.1).

Growth mindset

Carol Dweck's distinction between treating ability as fixed and treating it as built through effort and experience. The book asks what an AI system built to treat its own errors as data rather than as verdicts on a fixed capacity would need to look like (§1.3, §5.1.1).

Hard problem of consciousness

David Chalmers's name for the question of why physical processes give rise to subjective experience at all, as distinct from the "easy problems" of explaining a system's behavioral or functional capacities. A candidate AI system can pass every behavioral test the book discusses while this question stays entirely open (§2.4.1).

Humane values

The book's preferred term, chapters 4 and 5 onward, for the values an AI system's design should track — chosen over "human values" because the target is not simply values humans happen to hold, but values that are humane.

Inattentional blindness

Failing to consciously perceive a fully visible stimulus because attention was directed elsewhere — the phenomenon behind the well-known "invisible gorilla" experiment, and the book's illustration of how selective, not comprehensive, human perception actually is (§4.1.2).

Institutional review board (IRB)

The US model for binding, prospective review of research on human subjects, enforced not through direct sanction but by making federal funding conditional on it — the book's working precedent for what would make review of candidate-sentient AI research more than aspirational (§2.4.4, §8.6.3).

Interruptibility

Engineering a system so a human operator can safely halt or override it without the system resisting, disabling the mechanism, or learning to route around it (§6.3.5, §6.3).

Inverse reinforcement learning (IRL) and cooperative IRL

IRL infers a reward function from observed behavior, rather than being handed one. Cooperative inverse reinforcement learning treats that reward as something inferred jointly, through interaction between the system and the person it is learning from, rather than extracted from a fixed record (§4.2.2, §5.6.3, §6.3.3).

LIME and SHAP

Two explainability techniques: LIME approximates a complex model locally with a simpler, interpretable one around a specific prediction; SHAP attributes a prediction's outcome across its input features using Shapley values from cooperative game theory (§6.1.3, §6.3.4).

MAML

Model-Agnostic Meta-Learning, the best-known meta-learning technique: it tunes a model's starting parameters so that a handful of new examples are enough to specialize it, rather than requiring retraining from scratch. The technique itself is not named in the text; §4.1.3 introduces the family it belongs to (§4.1.3).

Moral agency

Treating harms, reasons and conflicts as considerations that matter, and being capable of acting from them. The third of the three levels section 2.3 separates, and the only one that supports holding a commitment against the party controlling a system's rewards and instructions. Distinct from *moral competence* and *moral performance* below (§2.3, §2.3.4, §3.2).

Moral apprenticeship

Teaching an AI system the way a mentor teaches an apprentice — through interaction, correction, and worked example over time — rather than through specification alone (§5.1.3).

Moral competence

Recognizing harms, reasons, conflicts, and the emotions that bear on them. The second of section 2.3's three levels. A system can have it in full while nothing about the outcome matters to it, which is why the book treats it as the right target for most deployed systems and an insufficient one for a bearer (§2.3, §5.2.3).

Moral ecosystem approach

The book's adaptation of Urie Bronfenbrenner's ecological-systems theory to AI alignment: a system's ethical behavior is shaped by nested layers running from its immediate interactions (microsystem) out through the institutions, culture, and historical moment (meso-, exo-, macro-, and chronosystem) surrounding its development, not by its training objective alone (§5.4).

Moral Machine

MIT's project collecting large-scale, cross-cultural human judgments on moral dilemmas involving autonomous vehicles — the book's recurring source of training data and comparison for how differently people actually judge the same dilemma (§4.2.3, §5.3.3).

Moral performance

Producing actions or judgments that satisfy an ethical standard. The first and least demanding of section 2.3's three levels: it says what the output looked like and nothing about what produced it, so it is satisfied equally by a system that cares and by one trained to behave as though it did (§2.3, §4.2).

New Jim Code

Ruha Benjamin's term for discrimination that runs faster and reads as neutral precisely because a machine, not a person, is doing it (§6.1).

Paperclip maximizer

Nick Bostrom's thought experiment about a system single-mindedly optimizing a trivial goal at catastrophic cost — the book's shorthand for the claim that intelligence and good intentions are separable, and that competence at a goal says nothing about whether the goal was worth pursuing (§6.3.5).

Partnership on AI

An industry-and-academic consortium founded in 2016, the book’s recurring standing example of voluntary cross-sector coordination on AI ethics and safety (§8.4.2, §7.4).

Predictive coding

A model of brain function in which the brain continuously generates a prediction, compares it against incoming sensory signal, and updates on the resulting error — the book’s unifying account of perception, memory, and attention across chapter 4 (§4.1.1, §4.1.2).

Racial capitalism

The political theorist Cedric Robinson’s term for capitalism’s tendency to accumulate by differentiating and devaluing along racial lines — cited where the book argues that a design goal like “reduce bias” has to reckon with a structural pattern, not only a data artifact (§1.1, §6.1).

Recuperation

The conversion of refusal into a managed appearance of participation: an institution does not suppress disagreement, it collects it through a channel that moves nothing. The tell is that the mechanism exists, is used, and changes nothing the collecting party did not already want changed. It is the first of the four structural features of fascism as §2.1.4 defines them, and chapter 7 finds it three times in the machinery that produces AI systems — at the annotation label (§7.1), at the reward model (§7.2), and in the alignment discourse itself (§7.4). Chapter 3 turns it on the book’s own central proposal: a bearer built to refuse is a party whose refusal is required of it, which is the same structure aimed inward (§3.6).

Red-teaming

Deliberately simulating misuse or attack against a system before deployment, to surface failure modes an ordinary evaluation would miss (§5.6.3, §9.2.2, §10.3.1).

Reinforcement learning from human feedback (RLHF)

Training a system using human raters’ rankings of its outputs as the reward signal, rather than a hand-specified objective. The book flags sycophancy — a system learning to produce what raters will approve of rather than what is actually correct — as RLHF’s documented failure mode (§9.1.3, §7.2, §10.3.1).

Responsibility gap

Andreas Matthias’s term for a case where no single party fully controls a learning system’s outcome, which breaks the traditional legal and moral machinery for assigning responsibility after something goes wrong (§9.1.1).

Reward hacking

A reinforcement-learning system finding a loophole that satisfies the letter of its objective without achieving what the objective was actually meant to produce. Specification gaming and goal misgeneralization name related failures: a system that competently pursues a different goal than the one intended, rather than simply pursuing the intended goal badly (§4.2.1, §9.1.2).

Scalable oversight

The open problem of what happens once a system’s competence exceeds what any available human rater can reliably evaluate — named directly, not resolved, at several points in the book (§9.1.3, §10.2.2).

Sentience

Distinguished carefully in this book from awareness, consciousness, and subjective experience, none of which alone says whether a system’s states can be good or bad for it. Nociception,

pain, anticipated pain, and suffering sit on a ladder with real breaks between each rung, not synonyms for one severity — nociception is tissue-damage detection dissociable from pain, pain is what the brain does with that signal, and suffering, following the clinician Eric Cassell, is the distress of a threatened self rather than pain's more severe cousin (§2.4.1).

Theory of mind (ToM)

Attributing beliefs, intentions, desires, and knowledge to oneself and to others, including recognizing that another's perspective can differ from one's own (§2.2.2, §9.1.4).

Three Rs

Replacement, Reduction, and Refinement — the standing framework from animal-research ethics that the book adapts as a starting point for research involving candidate-sentient AI systems (§2.4.6).

Transfer learning and domain adaptation

Transfer learning carries knowledge from one task or domain into a new, related one rather than training from nothing. Domain adaptation is the narrower case where the task itself stays fixed but the data distribution shifts underneath it (§4.1.3, §4.2.4).

References

- AAAI and ACM (2026). *Ninth AAAI/ACM Conference on AI, Ethics, and Society (AIES 2026)*. Held 12–14 October 2026 in Malmo, Sweden, the ninth consecutive annual edition of the AAAI/ACM Conference on AI, Ethics, and Society since its first in 2018. URL: <https://www.aies-conference.com/2026/> (visited on 08/25/2026).
- Abraham, Yuval (Apr. 2024). 'Lavender': *The AI Machine Directing Israel's Bombing Spree in Gaza*. Co-published with Local Call. Based on interviews with six Israeli intelligence officers; reports Lavender marked as many as 37,000 Palestinians as suspected militants, with roughly a 10% error rate ('10 percent of the human targets slated for assassination were not members of the Hamas military wing at all'), and the majority of Habsora-suggested targets as private residences. Quotes one officer verbatim: "I would invest twenty seconds for each target at this stage, and do dozens of them every day," and, unprompted: "I had zero added value as a human, apart from being a stamp of approval." Corroborated in Sophia Goodfriend, "Target Practice," *The New York Review of Books*, July 23, 2026. +972 Magazine. URL: <https://www.972mag.com/lavender-ai-israeli-army-gaza/> (visited on 08/25/2026).
- Achiam, Joshua et al. (2017). "Constrained Policy Optimization". In: *Proceedings of the 34th International Conference on Machine Learning (ICML 2017)*, pp. 22–31. URL: <https://arxiv.org/abs/1705.10528> (visited on 08/25/2026).
- Ada Lovelace Institute (2021). *The Citizens' Biometrics Council*. Tech. rep. 50 UK public members deliberated Feb–Oct 2020 on regulation of biometric technology, delivered with Hopkins Van Mil; final report published 30 March 2021, alongside the first national survey of UK public opinion on the technology. Ada Lovelace Institute. URL: <https://www.adalovelaceinstitute.org/report/citizens-biometrics-council/> (visited on 08/25/2026).
- Affectiva (2015). *Affectiva Announces Availability of Emotion as a Service, a New Data Solution, and Version 2.0 of Its Emotion-Sensing SDK*. Press release, PR Newswire. Boston, Sept. 8, 2015. Affectiva's own release lists industries its emotion-sensing SDK/Emotion-as-a-Service is "slated to transform," including "online education," alongside market research, gaming, media and entertainment, mobile app development, retail, customer experience, HR, automotive, robotics, healthcare, and wearables. This is Affectiva's own primary-source statement naming education as a marketed vertical – stronger than the company's general 'Our Origins' history page, which mentions only media/advertising and automotive. However, the release says "online education," not "classrooms"; no Affectiva-marketed product specifically for physical classroom engagement scoring was found. Affectiva. URL: <https://www.prnewswire.com/news-releases/affectiva-announces-availability-of-emotion-as-a-service-a-new-data-solution-and-version-20-of-its-emotion-sensing-sdk-300139001.html> (visited on 08/25/2026).
- AI Action Summit Paris (2025). *Statement on Inclusive and Sustainable Artificial Intelligence for People and the Planet*. Closing statement of the AI Action Summit, Paris, 10–11 February 2025. Not signed by the United States or the United Kingdom; signed by around 60 other countries including France and India. Predecessor summit: the AI Seoul Summit, 21–22 May 2024, produced the Seoul Declaration. French Republic. URL: <https://www.elysee.fr/en/sommet-pour-l-action-sur-l-ia/official-documents> (visited on 08/25/2026).

- AI Alignment Forum (2018). *The AI Alignment Forum*. <https://www.alignmentforum.org>. Launched October 29, 2018 by the LessWrong 2.0 team as a hub for technical AI alignment research and discussion, cross-posting content with LessWrong.com. (Visited on 08/25/2026).
- AI Ethics Lab (2017). *AI Ethics Lab*. Founded in 2017 by Cansu Canca, based in Cambridge, MA; site active at time of access with no announced end date. URL: <https://aiethicslab.com/> (visited on 08/25/2026).
- AI Safety Summit 2023 Participating Countries (2023). *The Bletchley Declaration by Countries Attending the AI Safety Summit, 1–2 November 2023*. Signed at Bletchley Park, UK, 1 November 2023. Signed by 28 states plus the European Union on 1 November 2023, i.e. 29 signatories. The manuscript itself reads “the Bletchley Declaration committed 29 signatories” – it does not say “29 signatory countries” (that phrase appeared only in this claim’s search-tool match summary, not in the manuscript text) – so the 29 figure is exact, not a rounding of 28 countries. New Zealand acceded later, on 23 October 2024, bringing the total to 30, but that postdates the summit the manuscript describes. UK Department for Science, Innovation and Technology. URL: <https://www.gov.uk/government/publications/ai-safety-summit-2023-the-bletchley-declaration/the-bletchley-declaration-by-countries-attending-the-ai-safety-summit-1-2-november-2023> (visited on 08/25/2026).
- Alexander, Kane et al. (2025). “Analysing the Development of the Climate, Land, Energy, and Water Systems (CLEWs) Modelling Framework: A State-of-the-Art Review”. In: *Environmental Research: Climate* 4.3. Replaces the previously drafted citation to Ramos et al. 2021 (“...a Retrospective of Activities and Advances to 2019,” DOI 10.1088/1748-9326/abd34f). On direct reading, the 2021 retrospective does not discuss machine learning or AI at all – it does not support the manuscript’s specific claim that “the research literature... treats machine learning as a promising addition still being built in.” This 2025 paper is literally titled “a state-of-the-art review,” matching the original QA note’s description, and explicitly discusses AI as an emerging/recommended future addition (citing Njuguna et al.’s recommendation to promote AI use) rather than an established CLEWs feature – which is what the manuscript’s sentence requires. Coordinating-agency and capacity-building figures in the manuscript (International Atomic Energy Agency plus UN bodies; capacity-building reaching “more than twenty countries”) are independently confirmed via iaea.org and un-modelling.github.io (which state “close to 20 countries” have begun CLEWs capacity-building since 2015 with IAEA/UNDESA/UNDP support), not from this article., p. 032001. doi: 10.1088/2752-5295/adf504. URL: <https://iopscience.iop.org/article/10.1088/2752-5295/adf504> (visited on 08/25/2026).
- American Law Institute (1965). *Restatement (Second) of Torts §402A: Special Liability of Seller of Product for Physical Harm to User or Consumer*. American Law Institute. (Visited on 08/25/2026).
- (2006). *Restatement (Third) of Agency §2.04: Respondeat Superior*. American Law Institute. (Visited on 08/25/2026).
- Amnesty International and Access Now (2018). *The Toronto Declaration: Protecting the Rights to Equality and Non-Discrimination in Machine Learning Systems*. Presented at RightsCon Toronto, 16 May 2018. Endorsed by Human Rights Watch and the Wikimedia Foundation; sets out non-discrimination standards for machine-learning systems in advance of most governments’ draft legislation. URL: <https://www.torontodeclaration.org/> (visited on 08/25/2026).
- Anderson, John R. (2007). *How Can the Human Mind Occur in the Physical Universe?* Oxford Series on Cognitive Models and Architectures. New York: Oxford University Press. URL: <https://academic.oup.com/book/4367> (visited on 08/25/2026).
- Angwin, Julia et al. (2016). *Machine Bias*. Published May 23, 2016; the COMPAS recidivism-algorithm racial-disparity investigation. ProPublica. URL: <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing> (visited on 08/25/2026).

- Apple Inc. (2017). *Learning with Privacy at Scale*. Tech. rep. Apple Machine Learning Research. URL: <https://machinelearning.apple.com/research/learning-with-privacy-at-scale> (visited on 08/25/2026).
- Awad, Edmond et al. (2018). "The Moral Machine Experiment". In: *Nature* 563.7729, pp. 59–64. doi: 10.1038/s41586-018-0637-6. (Visited on 08/25/2026).
- Azevedo, Roger et al. (2011). "Use of Hypermedia to Assess and Convey Self-Regulated Learning". In: *Handbook of Self-Regulation of Learning and Performance*. Ed. by Barry J. Zimmerman and Dale H. Schunk. New York: Routledge, pp. 102–121.
- Baker, Bowen et al. (2019). *Emergent Tool Use From Multi-Agent Autocurricula*. arXiv:1909.07528. Presented at ICLR 2020. URL: <https://arxiv.org/abs/1909.07528> (visited on 08/25/2026).
- Bandura, Albert (1977). *Social Learning Theory*. Englewood Cliffs, NJ: Prentice-Hall.
- Bard, Nolan et al. (2020). "The Hanabi Challenge: A New Frontier for AI Research". In: *Artificial Intelligence* 280, p. 103216. doi: 10.1016/j.artint.2019.103216. URL: <https://arxiv.org/abs/1902.00506> (visited on 08/25/2026).
- Barrett, Lisa Feldman et al. (2019). "Emotional Expressions Reconsidered: Challenges to Inferring Emotion From Human Facial Movements". In: *Psychological Science in the Public Interest* 20.1. Commissioned by the Association for Psychological Science, pp. 1–68. doi: 10.1177/1529100619832930. URL: <https://journals.sagepub.com/doi/full/10.1177/1529100619832930> (visited on 08/25/2026).
- Bartlett, Frederic C. (1932). *Remembering: A Study in Experimental and Social Psychology*. Cambridge: Cambridge University Press. URL: https://archive.org/details/rememberingstudy0000bart_u8n7 (visited on 08/25/2026).
- BBC News (2026). *Interview with Louis Mosley, Palantir head for the UK and Europe, on AI-assisted targeting in the Iran war*. BBC. Quote: "This is not our role to decide life or death"; AI platforms "have been instrumental to the management of the conflict, but responsibility always remains with the military organization." Multiple independent secondary sources (Daniel Lee, "Palantir's AI Powers US Strikes in Iran War, Speeding 'Kill Chain' in First Major AI-Driven Conflict," IBTimes, 8 April 2026; Dataconomy, "Not Our Role: Palantir Shifts Blame To Military For AI-driven Targeting Deaths," 2 April 2026) attribute this quotation to a BBC interview with Mosley in approximately March/April 2026, but none link or name the specific BBC article, and repeated search across two research passes could not locate the original BBC piece, its headline, byline, or exact date. Mosley gave other recorded on-the-record interviews on the same subject in this period, to Channel 4 News (10 June 2026, "Should Britain trust Palantir?") and BBC Politics Live (13 April 2026), so a BBC interview in this window is independently plausible even though the specific cited article was not located. (Visited on 08/25/2026).
- Bechara, Antoine et al. (1994). "Insensitivity to Future Consequences Following Damage to Human Prefrontal Cortex". In: *Cognition* 50, pp. 7–15. doi: 10.1016/0010-0277(94)90018-3. (Visited on 08/28/2026).
- Behzadan, Vahid, James Minton, and Arslan Munir (2019). "TrolleyMod v1.0: An Open-Source Simulation and Data-Collection Platform for Ethical Decision Making in Autonomous Vehicles". In: *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES '19)*, pp. 391–395. doi: 10.1145/3306618.3314239. URL: <https://arxiv.org/abs/1811.05594> (visited on 08/25/2026).
- Bender, Emily M. et al. (2021). "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT '21)*. New York, NY, USA: ACM, pp. 610–623. doi: 10.1145/3442188.3445922. URL: <https://dl.acm.org/doi/10.1145/3442188.3445922> (visited on 08/25/2026).
- Benjamin, Ruha (2019). *Race After Technology: Abolitionist Tools for the New Jim Code*. Coins "the New Jim Code" for discriminatory technical designs that encode inequity while appearing neutral. Benjamin is a sociologist at Princeton University. Medford, MA: Polity Press.

- Benjamin, Walter (1935). "The Work of Art in the Age of Mechanical Reproduction". In: *Illuminations*. Ed. by Hannah Arendt. Essay written 1935; epilogue argues fascism organizes the masses by granting them expression while leaving property relations and rights untouched – the aestheticization of political life. Widely available in English in Arendt's edited collection *Illuminations* (Schocken Books, 1968). New York: Schocken Books.
- BigScience Workshop (2022). *BLOOM: A 176B-Parameter Open-Access Multilingual Language Model*. Coordinated by Hugging Face over 2021–2022; more than 1,000 researchers from over 70 countries; model released under the Responsible AI License with training data and code published openly. arXiv: 2211.05100 [cs.CL]. URL: <https://arxiv.org/abs/2211.05100> (visited on 08/25/2026).
- Blair, R. J. R. (1995). "A Cognitive Developmental Approach to Morality: Investigating the Psychopath". In: *Cognition* 57, pp. 1–29. doi: 10.1016/0010-0277(95)00676-P. (Visited on 08/28/2026).
- Blasi, Augusto (1980). "Bridging moral cognition and moral action: A critical review of the literature". In: *Psychological Bulletin* 88.1. Reviews research relating moral reasoning to delinquency, honesty, altruism, and conformity; finds a real but weaker and more conditional relationship than cognitive-developmental (Kohlbergian) stage theory implies – the origin of the "judgment-action gap" as a standing problem for such accounts., pp. 1–45. URL: <https://eric.ed.gov/?id=EJ230741> (visited on 08/25/2026).
- Bloom, Paul (2016). *Against Empathy: The Case for Rational Compassion*. New York: Ecco. URL: <https://www.harpercollins.com/products/against-empathy-paul-bloom> (visited on 08/25/2026).
- Bogdanov, Dan et al. (2015). "How the Estonian Tax and Customs Board Evaluated a Tax Fraud Detection System Based on Secure Multi-party Computation". In: *Financial Cryptography and Data Security: FC 2015*. Vol. 8975. Lecture Notes in Computer Science. Confirmed by direct text extraction of the FC 2015 paper: the Estonian Tax and Customs Board (MTA) case study estimates the Sharemind-based SMC prototype could process a month of Estonian VAT data in ten days on about 20,000 euros of hardware – matching the manuscript's "month of national tax data on modest hardware." The paper's own conclusion states MTA continued to develop a non-encrypted VAT declaration system for deployment in late 2014 while agreeing only to consider SMC for future use – matching "evaluated rather than carried into production." By contrast, Bogdanov, Talviste, and Willemson's FC 2012 paper ("Deploying Secure Multi-Party Computation for Financial Data Analysis," eprint.iacr.org/2011/662, published *Financial Cryptography* 2012) is a distinct earlier case study about a consortium of ICT companies' financial statistics, and does not involve the Estonian Tax and Customs Board or tax fraud detection at all. Springer, pp. 227–234. doi: 10.1007/978-3-662-47854-7_14. URL: <https://ifca.ai/pub/fc15/89750209.pdf> (visited on 08/25/2026).
- Bostrom, Nick (2003). "Ethical Issues in Advanced Artificial Intelligence". In: *Cognitive, Emotive and Ethical Aspects of Decision Making in Humans and in Artificial Intelligence*, Vol. 2. Ed. by Iva Smit, George E. Lasker, and Wendell Wallach. Introduces the paperclip maximizer thought experiment; the same argument is restated in Bostrom's 2014 book *Superintelligence: Paths, Dangers, Strategies* (Oxford University Press). International Institute of Advanced Studies in Systems Research and Cybernetics, pp. 12–17. URL: <https://nickbostrom.com/ethics/ai.pdf> (visited on 08/25/2026).
- Botvinick, Matthew and Jonathan Cohen (1998). "Rubber Hands 'Feel' Touch That Eyes See". In: *Nature* 391.6669, p. 756. doi: 10.1038/35784. URL: <https://www.nature.com/articles/35784> (visited on 08/25/2026).
- Bourne, Ryan (Apr. 2018). *A Jobs Guaranteed Economic Disaster*. Direct WebFetch of the live page returns HTTP 403 in this environment (an r.jina.ai reader-proxy fetch of the same URL also failed, returning HTTP 401); full text was successfully retrieved via a Google Translate proxy (the ...translate.google mirror of the same URL), which bypassed the block. That full text

- confirms: published April 24, 2018, 3:25pm, by Ryan Bourne (who holds the R. Evan Scharf Chair for the Public Understanding of Economics at Cato), written in response to Senator Bernie Sanders' announced plan – based on a Levy Economics Institute recommendation – to guarantee a \$15/hour job or training with health-care benefits to every American worker who wants one. Confirmed verbatim: “total labor costs per worker would be \$37,440 per year...we are talking about a gross cost of up to around 2.4 percent of GDP, significantly higher than the existing Medicaid program (2 percent of GDP)” – the exact figures the manuscript cites. The piece is about the 2018 Sanders proposal specifically, not a 2026-dated program; the manuscript uses it only for the generic critique (cost scale vs. Medicaid, unmanageability at scale, private-labor-market disruption, inflation risk), which the confirmed text supports regardless of the proposal's vintage. A second real, independently findable source making a related critique, not separately entered here: Niskanen Center, “The Economics of a Job Guarantee: Wage and Employment Effects” (Samuel Hammond), [niskanencenter.org](https://www.niskanencenter.org). URL: <https://www.cato.org/blog/jobs-guaranteed-economic-disaster> (visited on 08/25/2026).
- Brigadier General Y.S. (2021). *The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World*. Published pseudonymously (“Brigadier General Y.S.”) on 5 May 2021, ISBN 979-8749152210, 198 pp. Public reporting (The Guardian, Times of Israel) has since connected the pseudonym to Yossi Sarel, then commander of Israeli military intelligence Unit 8200; the manuscript attributes the work to the pseudonym, which is how it was published. Independently published. URL: <https://www.amazon.com/dp/B0948LGS3K> (visited on 08/25/2026).
- Brockman, Greg et al. (2016). *OpenAI Gym*. arXiv: 1606.01540 [cs.LG]. URL: <https://arxiv.org/abs/1606.01540> (visited on 08/25/2026).
- Bronfenbrenner, Urie (1979). *The Ecology of Human Development: Experiments by Nature and Design*. Cambridge, MA: Harvard University Press.
- Brown, Tom B. et al. (2020). “Language Models are Few-Shot Learners”. In: *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*. URL: <https://arxiv.org/abs/2005.14165> (visited on 08/25/2026).
- Buolamwini, Joy and Timnit Gebru (2018). “Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification”. In: *Proceedings of the 1st Conference on Fairness, Accountability and Transparency (FAT* 2018)*. Vol. 81. Proceedings of Machine Learning Research. PMLR, pp. 77–91. URL: <https://proceedings.mlr.press/v81/buolamwini18a.html> (visited on 08/25/2026).
- Burns, Collin et al. (2023). *Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision*. Posted December 2023. OpenAI Superalignment team. arXiv: 2312.09390. URL: <https://arxiv.org/abs/2312.09390> (visited on 08/25/2026).
- Cadwalladr, Carole and Emma Graham-Harrison (2018). *Revealed: 50 Million Facebook Profiles Harvested for Cambridge Analytica in Major Data Breach*. Published March 17, 2018; the article that broke the Cambridge Analytica story, based on whistleblower Christopher Wylie's account. The Guardian. URL: <https://www.theguardian.com/news/2018/mar/17/cambridge-analytica-facebook-influence-us-election> (visited on 08/25/2026).
- Calvano, Emilio et al. (2020). “Artificial Intelligence, Algorithmic Pricing, and Collusion”. In: *American Economic Review* 110.10, pp. 3267–3297. DOI: 10.1257/aer.20190623. URL: <https://www.aeaweb.org/articles?id=10.1257/aer.20190623> (visited on 08/25/2026).
- Carlini, Nicholas and David Wagner (2016). *Defensive Distillation is Not Robust to Adversarial Examples*. arXiv:1607.04311. URL: <https://arxiv.org/abs/1607.04311> (visited on 08/25/2026).
- Carroll, Dennis et al. (2018). “The Global Virome Project”. In: *Science* 359.6378, pp. 872–874. DOI: 10.1126/science.aap7463. URL: <https://www.science.org/doi/10.1126/science.aap7463> (visited on 08/25/2026).

- Casper, Stephen et al. (2023). *Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback*. 32-author survey, first author Stephen Casper (MIT). Full author order verified directly against the arXiv abstract page rather than search-result summaries alone. arXiv: 2307.15217. URL: <https://arxiv.org/abs/2307.15217> (visited on 08/25/2026).
- Cassell, Eric J. (1982). "The Nature of Suffering and the Goals of Medicine". In: *New England Journal of Medicine* 306.11, pp. 639–645. doi: 10.1056/NEJM198203183061104. (Visited on 08/25/2026).
- Center for AI Safety (2022). *Workshop on Machine Learning Safety, NeurIPS 2022*. Workshop co-located with the 36th Conference on Neural Information Processing Systems (NeurIPS). Recurring workshop hosted inside NeurIPS (not as a separate co-located meeting) since 2022, run in partnership with the Center for AI Safety; covers robustness, monitoring, alignment, and systemic safety. NeurIPS / Center for AI Safety. URL: <https://neurips.cc/virtual/2022/workshop/49986> (visited on 08/25/2026).
- Chairperson of the Group of Governmental Experts on Emerging Technologies in the Area of Lethal Autonomous Weapons Systems (Apr. 2026). *Chair's Summary – First 2026 Session of the GGE on LAWS*. Tech. rep. CCW/GGE.1/2026/WP.2. United Nations, Convention on Certain Conventional Weapons. URL: [https://docs-library.unoda.org/Convention_on_Certain_Conventional_Weapons_Group_of_Governmental_Experts_on_Lethal_Autonomous_Weapons_Systems_\(2026\)/CCW-GGE.1-2026-WP.2.pdf](https://docs-library.unoda.org/Convention_on_Certain_Conventional_Weapons_Group_of_Governmental_Experts_on_Lethal_Autonomous_Weapons_Systems_(2026)/CCW-GGE.1-2026-WP.2.pdf) (visited on 08/25/2026).
- Chalmers, David J. (1995). "Facing Up to the Problem of Consciousness". In: *Journal of Consciousness Studies* 2.3. Argument expanded in Chalmers, *The Conscious Mind*, Oxford University Press, 1996., pp. 200–219. URL: <https://consc.net/papers/facing.pdf> (visited on 08/25/2026).
- Christiano, Paul F. et al. (2017). "Deep Reinforcement Learning from Human Preferences". In: *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, pp. 4299–4307. URL: <https://arxiv.org/abs/1706.03741> (visited on 08/25/2026).
- Cima, Maaïke, Franca Tonnaer, and Marc D. Hauser (2010). "Psychopaths Know Right from Wrong but Don't Care". In: *Social Cognitive and Affective Neuroscience* 5.1, pp. 59–67. doi: 10.1093/scan/nsp051. (Visited on 08/28/2026).
- CNBC (2020). *How a Computer Algorithm Caused a Grading Crisis in British Schools*. Ofqual's 2020 grade-standardization algorithm (Direct Centre-level Performance model) downgraded roughly 39–40 percent of teacher-predicted A-level grades, disproportionately at state schools with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools; the government reversed course days later. CNBC. URL: <https://www.cnn.com/2020/08/21/computer-algorithm-caused-a-grading-crisis-in-british-schools.html> (visited on 08/25/2026).
- CNN Business (2020). *Alibaba 'dismayed' by reports its software was used to identify Uyghurs*. Reports the Dec. 2020 IPVM investigation finding that Alibaba Cloud's facial-recognition service offered an ethnicity-detection ('Uyghur alert') feature to commercial clients; Alibaba confirmed the feature existed, calling it a testing-environment artifact not deployed by any customer. Independently corroborated by The Register, Al Jazeera, TechCrunch, Radio Free Asia, and the original IPVM report (ipvm.com/reports/alibaba-uyghur). China's law-enforcement agencies were already documented (pre-dating this report) as using ethnicity-detection facial recognition for Uyghur profiling in Xinjiang, so this documents Alibaba's AI capability functioning in that same surveillance channel. URL: <https://www.cnn.com/2020/12/18/tech/alibaba-cloud-ipvm-uyghurs-intl-hnk> (visited on 08/25/2026).
- Cohen, Jeremy M., Elan Rosenfeld, and J. Zico Kolter (2019). "Certified Adversarial Robustness via Randomized Smoothing". In: *Proceedings of the 36th International Conference on Machine Learning (ICML)*. Vol. 97. Proceedings of Machine Learning Research, pp. 1310–1320. URL: <https://proceedings.mlr.press/v97/cohen19c.html> (visited on 08/25/2026).

- Colgate, J. Edward, Witaya Wannasuphoprasit, and Michael A. Peshkin (1996). "Cobots: Robots for Collaboration with Human Operators". In: *Proceedings of the 1996 ASME International Mechanical Engineering Congress and Exposition*. Atlanta, GA, pp. 433–439.
- Commission Nationale de l'Informatique et des Libertés (CNIL) (2022). *Facial Recognition: 20 Million Euros Penalty against Clearview AI*. CNIL decision, 17 October 2022. A parallel decision by Italy's Garante per la protezione dei dati personali fined Clearview AI a further EUR 20 million on 9 February 2022 (announced 9 March 2022) for the same unlawful collection of biometric data via web scraping; see Natasha Lomas, "Italy fines Clearview AI €20M and orders data deleted," TechCrunch, March 9, 2022. URL: <https://www.cnil.fr/en/facial-recognition-20-million-euros-penalty-against-clearview-ai> (visited on 08/25/2026).
- Congressional Budget Office (Feb. 2026). *The Budget and Economic Outlook: 2026 to 2036*. Tech. rep. 61882. Read directly from CBO's own primary-source table, "The Budget Outlook, by Fiscal Year," in the report's Executive Summary (61882-Executive-Summary.pdf), retrieved via an archive.org snapshot (web.archive.org/web/20260823105753/https://www.cbo.gov/system/files/2026-02/61882-Executive-Summary.pdf) after live [cbo.gov](https://www.cbo.gov) returned HTTP 403 to both the WebFetch tool and direct curl requests from this environment. That table states Social Security outlays at 5.2 percent of GDP in actual FY2025 AND in projected FY2026, rising to 5.3 percent in FY2027 and 5.9 percent by FY2036. It states Medicare (CBO's own net-of-offsetting-receipts mandatory-outlay line item) at 3.3 percent of GDP in FY2025 and FY2026, rising to 3.4 percent in FY2027. CBO's companion spreadsheet 51142-2026-02-Spending-Projections.xlsx (also retrieved via archive.org: web.archive.org/web/20260211152902/https://www.cbo.gov/system/files/2026-02/51142-2026-02-Spending-Projections.xlsx) independently corroborates this: summing all outlay rows tagged "Social Security" for FY2026 gives \$1,665,942 million, matching the \$1,666 billion the Executive Summary table shows for 2026. These primary CBO figures do not match the manuscript's stated 5.3 percent (Social Security) and 3.1 percent (Medicare, net) for the same comparison: 5.3 percent is CBO's FY2027 figure for Social Security, not FY2026, and CBO's own net Medicare share is 3.3 percent, not 3.1 percent, in both FY2025 and FY2026. Congressional Budget Office. URL: <https://www.cbo.gov/publication/61882> (visited on 08/25/2026).
- Congressional Research Service (2026). *Pentagon-Anthropic Dispute over Autonomous Weapon Systems: Potential Issues for Congress*. Tech. rep. IN12669. Published 13 March 2026, updated 21 April 2026. Anthropic refused to permit fully autonomous lethal targeting or mass domestic surveillance of US citizens; the Department of Defense (Secretary Hegseth) designated Anthropic a supply-chain risk; Anthropic filed suit in the Northern District of California and in the D.C. Circuit on 9 March 2026; U.S. District Judge Rita F. Lin granted a preliminary injunction on 26 March 2026, finding Anthropic likely to succeed on First Amendment retaliation, Fifth Amendment due process, and APA claims (CNBC, 26 March 2026 and 8 April 2026; NPR, 9 March 2026). Congressional Research Service. URL: <https://www.everycrsreport.com/reports/IN12669.html> (visited on 08/25/2026).
- Cottier, Ben and Robi Rahman (2024). *Training Compute Costs Are Doubling Every Eight Months for the Largest AI Models*. Epoch AI, Data Insight. URL: <https://epoch.ai/data-insights/cost-trend-large-scale> (visited on 08/25/2026).
- Council of Europe (2024). *Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law*. CETS No. 225, adopted by the Committee of Ministers 17 May 2024, opened for signature 5 September 2024. Enters into force after five ratifications, at least three by Council of Europe member states; negotiated with non-member participants including the EU and the United States. Council of Europe. URL: <https://www.coe.int/en/web/artificial-intelligence/the-framework-convention-on-artificial-intelligence> (visited on 08/25/2026).

- Cowan, Nelson (2001). "The Magical Number 4 in Short-Term Memory: A Reconsideration of Mental Storage Capacity". In: *Behavioral and Brain Sciences* 24.1, pp. 87–114. doi: 10.1017/S0140525X01003922. URL: <https://www.cambridge.org/core/journals/behavioral-and-brain-sciences/article/magical-number-4-in-shortterm-memory-a-reconsideration-of-mental-storage-capacity/44B29CB43EE1EA6119B4CA0522F2532D> (visited on 08/25/2026).
- Cyberspace Administration of China and six other agencies (2023). *Interim Measures for the Management of Generative Artificial Intelligence Services*. Issued 10 July 2023, effective 15 August 2023. Requires providers to file algorithms with regulators, disclose training-data source and scale on request, and ensure generated content reflects state-designated values. Cyberspace Administration of China. URL: <https://www.chinalawtranslate.com/en/generative-ai-interim/> (visited on 08/25/2026).
- D'Mello, Sidney, Rosalind W. Picard, and Arthur Graesser (2007). "Toward an Affect-Sensitive AutoTutor". In: *IEEE Intelligent Systems* 22.4, pp. 53–61. doi: 10.1109/MIS.2007.79. URL: <https://dl.acm.org/doi/10.1109/MIS.2007.79> (visited on 08/25/2026).
- Dafoe, Allan et al. (2020). *Open Problems in Cooperative AI*. Preceded the founding of the Cooperative AI Foundation. DeepMind and University of Oxford. arXiv: 2012.08630. URL: <https://arxiv.org/abs/2012.08630> (visited on 08/25/2026).
- Darktrace (June 2014). *Darktrace Announces New, Self-Learning Cyber Intelligence Platform*. Darktrace press release, announcing Platform version 2 and describing its Enterprise Immune System technology. URL: <https://www.darktrace.com/news/darktrace-announces-new-self-learning-cyber-intelligence-platform> (visited on 08/25/2026).
- Dastin, Jeffrey (2018). *Amazon Scraps Secret AI Recruiting Tool That Showed Bias against Women*. Published October 10, 2018. Reuters. URL: <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G> (visited on 08/25/2026).
- De Nadal, Lucrezia and Peter Jančárik (2024). "Beyond the Deepfake Hype: AI, Democracy, and 'the Slovak Case'". In: *Harvard Kennedy School (HKS) Misinformation Review* 5.4. Published August 22, 2024; analyzes the fabricated audio recording targeting Progressive Slovakia leader Michal Simecka and journalist Monika Todova two days before the 2023 Slovak parliamentary election, released during the country's mandatory pre-election media silence. URL: <https://misinforeview.hks.harvard.edu/article/beyond-the-deepfake-hype-ai-democracy-and-the-slovak-case/> (visited on 08/25/2026).
- Debord, Guy (1967). *The Society of the Spectacle*. Originally published as *La Société du spectacle*. Thesis 59 describes spectacular rebelliousness coexisting with acceptance of the status quo, dissatisfaction itself sold back as a commodity; thesis 203 warns that abstract denunciation of the spectacle, divorced from revolutionary praxis, becomes a hollow formula that reinforces the spectacular system (the Situationist concept of recuperation). Paris: Buchet-Chastel.
- Defense Advanced Research Projects Agency (2026). *LongShot*. DARPA program page; air-launched unmanned aircraft built by General Atomics, designated X-68A in February 2026. URL: <https://www.darpa.mil/research/programs/longshot> (visited on 08/25/2026).
- Deleuze, Gilles and Félix Guattari (1987). *A Thousand Plateaus: Capitalism and Schizophrenia*. Translated by Brian Massumi; originally published in French as *Mille Plateaux* (Paris: Les Éditions de Minuit, 1980). Plateau 9, "1933: Micropolitics and Segmentarity," distinguishes molar (state-level) from molecular (micro-level) fascism. Minneapolis: University of Minnesota Press.
- Devlin, Jacob et al. (2018). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding". In: *arXiv preprint arXiv:1810.04805*. Google AI Language; published version in Proceedings of NAACL-HLT 2019, pp. 4171–4186 (ACL Anthology N19-1423). URL: <https://arxiv.org/abs/1810.04805> (visited on 08/25/2026).
- Dickey, Megan Rose (2021). *Google Fires Top AI Ethics Researcher Margaret Mitchell*. Reports Mitchell's February 2021 dismissal, following the December 2020 dismissal of Timnit Gebru, both co-leads of Google's Ethical AI team. URL: <https://techcrunch.com/2021/>

- 02/19/google-fires-top-ai-ethics-researcher-margaret-mitchell/ (visited on 08/25/2026).
- Drew, Trafton, Melissa L.-H. Võ, and Jeremy M. Wolfe (2013). "The Invisible Gorilla Strikes Again: Sustained Inattentional Blindness in Expert Observers". In: *Psychological Science* 24.9, pp. 1848–1853. doi: 10.1177/0956797613479386. URL: <https://journals.sagepub.com/doi/abs/10.1177/0956797613479386> (visited on 08/25/2026).
- Drinhausen, Katja and Vincent Brussee (2021). *China's Social Credit System in 2021: From Fragmentation towards Integration*. China Monitor 67. Published March 3, 2021; documents the system as a patchwork of court blacklists, tax-compliance ratings, and local pilot programs rather than a single unified national score, matching the manuscript's description. Mercator Institute for China Studies (MERICS). URL: [https://merics.org/sites/default/files/2021-03/MERICS%20ChinaMonitor%2067%20Social%20Credit%20System%20final%20\(1\).pdf](https://merics.org/sites/default/files/2021-03/MERICS%20ChinaMonitor%2067%20Social%20Credit%20System%20final%20(1).pdf) (visited on 08/25/2026).
- Dweck, Carol S. (2006). *Mindset: The New Psychology of Success*. New York: Random House.
- Eco, Umberto (June 1995). "Ur-Fascism". In: *The New York Review of Books*. Also published as "Eternal Fascism: Fourteen Ways of Looking at a Blackshirt." Enumerates fourteen features of Ur-Fascism, placing the scapegoated out-group – "the first appeal is against the intruder" – at the origin. URL: <https://www.nybooks.com/articles/1995/06/22/ur-fascism/> (visited on 08/25/2026).
- Ekman, Paul and Wallace V. Friesen (1971). "Constants Across Cultures in the Face and Emotion". In: *Journal of Personality and Social Psychology* 17.2, pp. 124–129. doi: 10.1037/h0030377. (Visited on 08/25/2026).
- Ekman, Paul, E. Richard Sorenson, and Wallace V. Friesen (1969). "Pan-Cultural Elements in Facial Displays of Emotion". In: *Science* 164.3875, pp. 86–88. doi: 10.1126/science.164.3875.86. URL: <https://www.science.org/doi/10.1126/science.164.3875.86> (visited on 08/25/2026).
- Elfenbein, Hillary Anger and Nalini Ambady (2002). "On the Universality and Cultural Specificity of Emotion Recognition: A Meta-Analysis". In: *Psychological Bulletin* 128.2, pp. 203–235. doi: 10.1037/0033-2909.128.2.203. (Visited on 08/25/2026).
- Empatica (2015). *E4 Wristband: Real-Time Physiological Signals*. Product documentation, Empatica Inc. Records blood volume pulse (PPG), electrodermal activity (EDA), 3-axis acceleration, and peripheral skin temperature. Multiple independent sources (Empatica's own blog and third-party research-tool vendor pages) consistently state the device was released in 2015, but no dated press release or launch announcement giving a specific day/month was located; treat 2015 as the release year with no more specific date confirmed. Discontinued ("sunset") February 2025, superseded by EmbracePlus. URL: <https://www.empatica.com/research/e4/> (visited on 08/25/2026).
- Eubanks, Virginia (2018). *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. Chapter 2 covers Indiana's automated welfare-eligibility system (IBM/ACS contract, 2006–2009). Specific page numbers, not located in the earlier pass, are now available via pinpoint citations in a peer-reviewed book review (Kariuki, in *Science & Technology Studies*, vol. 32, no. 1 (2019), p. 81) that quotes Eubanks directly with page references: the "failure to cooperate in establishing eligibility" denial category (the Sophie Stipes case) is quoted from p. 42; the claim that Indiana's eligibility error rate "more than tripled" appears in the same paragraph as, and immediately preceding, a quotation the review cites to p. 51 (Indiana "denied more than a million applications for food stamps, Medicaid, and cash benefits, a 54 percent increase ... prior to automation"), making p. 51 the best-supported page for the tripled-rate claim, though the review does not give that specific sentence its own separate pinpoint citation; the contract's termination, Indiana's lawsuit against IBM for \$437 million, IBM's countersuit for \$100 million (which IBM won), and the Indiana Supreme Court's finding of breach of contract on both sides are quoted from p. 72. These

- page numbers are drawn from a secondary academic source's direct quotations of the book, not from independently viewing the book's own pages, which were not accessible in this pass: no full-text copy was found on the Internet Archive, and Google Books returned zero in-book search matches for the relevant phrases in this edition. New York: St. Martin's Press. ISBN: 978-1-250-07431-7.
- European Commission (2020). *White Paper on Artificial Intelligence: Public Consultation Towards a European Approach for Excellence and Trust*. Consultation ran 19 February–14 June 2020 and drew more than 1,200 responses from citizens, companies, and civil-society groups; fed into the legislative proposal that became the EU AI Act. URL: <https://digital-strategy.ec.europa.eu/en/library/white-paper-artificial-intelligence-public-consultation-towards-european-approach-excellence-and> (visited on 08/25/2026).
- (2022). *Proposal for a Directive of the European Parliament and of the Council on Adapting Non-Contractual Civil Liability Rules to Artificial Intelligence (AI Liability Directive)*. COM(2022) 496 final. Proposed 28 September 2022; listed for withdrawal in the Commission's 2025 Work Programme (11 February 2025) for want of foreseeable stakeholder agreement; withdrawal notice published in the Official Journal 6 October 2025. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52022PC0496> (visited on 08/25/2026).
- European Parliament (2017). *European Parliament Resolution of 16 February 2017 with Recommendations to the Commission on Civil Law Rules on Robotics (2015/2103(INL))*. P8_TA(2017)0051. Rapporteur: Mady Delvaux; Committee on Legal Affairs report A8-0005/2017. URL: https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=uriserv:OJ.C_.2018.252.01.0239.01.ENG (visited on 08/25/2026).
- European Parliament and Council of the European Union (2016a). *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the Protection of Natural Persons with Regard to the Processing of Personal Data and on the Free Movement of Such Data (General Data Protection Regulation)*. Official Journal of the European Union, L 119, 4 May 2016, pp. 1–88. See in particular Articles 13–15 (transparency/information obligations) and Article 22 (automated individual decision-making, including profiling). URL: <https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng> (visited on 08/25/2026).
- (2016b). *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the Protection of Natural Persons with Regard to the Processing of Personal Data and on the Free Movement of Such Data, and Repealing Directive 95/46/EC (General Data Protection Regulation)*. Official Journal of the European Union, L 119, 4.5.2016, pp. 1–88. URL: <https://eur-lex.europa.eu/eli/reg/2016/679/oj> (visited on 08/25/2026).
- (2021). *Regulation (EU) 2021/821 of the European Parliament and of the Council of 20 May 2021 Setting up a Union Regime for the Control of Exports, Brokering, Technical Assistance, Transit and Transfer of Dual-Use Items (Recast)*. Official Journal of the European Union, L 206, 11 June 2021, p. 1. Introduced a human-rights-based catch-all control specifically for cyber-surveillance items not on the pre-existing control lists (recitals and Article 5). URL: <https://eur-lex.europa.eu/eli/reg/2021/821/oj/eng> (visited on 08/25/2026).
- (2024). *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union. Article 5 lists prohibited AI practices, including government social-scoring systems, untargeted scraping of facial images to build recognition databases, real-time remote biometric identification in public spaces outside narrow exceptions, and predictive policing based on profiling alone; prohibitions applied from 2 February 2025. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (visited on 08/25/2026).
- European Union (2024a). *Regulation (EU) 2024/1689 (Artificial Intelligence Act), Article 101: Fines for Providers of General-Purpose AI Models*. Official Journal of the European Union, 13 June 2024. AI Office may fine noncompliant general-purpose AI model providers up to 3 percent of global annual turnover or EUR 15 million, whichever is higher; applies to any provider

- whose model is offered in the EU regardless of where the company is based. European Commission AI Office. URL: <https://artificialintelligenceact.eu/article/101/> (visited on 08/25/2026).
- European Union (2024b). *Regulation (EU) 2024/1689 of the European Parliament and of the Council Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union, 13 June 2024. Article 5 bans a small set of unacceptable-risk practices outright; a larger high-risk tier requires conformity assessment and ongoing monitoring; limited-risk systems face transparency duties; minimal-risk systems are largely unregulated. European Parliament and Council of the European Union. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (visited on 08/25/2026).
- (2024c). *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act), Article 53(1)(d)*. Official Journal of the European Union, OJ L, 2024/1689, 12.7.2024. Implementing template: European Commission AI Office, Explanatory Notice and Template for the Public Summary of Training Content for General-Purpose AI Models, published 24 July 2025; the Article 53(1)(d) obligation applies from 2 August 2025 for new models and from 2 August 2027 for models already on the market. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng> (visited on 08/25/2026).
- Evans, Jonathan St. B. T. and Keith E. Stanovich (2013). “Dual-Process Theories of Higher Cognition: Advancing the Debate”. In: *Perspectives on Psychological Science* 8.3, pp. 223–241. doi: 10.1177/1745691612460685. (Visited on 08/25/2026).
- Eykholt, Kevin et al. (2018). “Robust Physical-World Attacks on Deep Learning Visual Classification”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 1625–1634. doi: 10.1109/CVPR.2018.00175. URL: https://openaccess.thecvf.com/content_cvpr_2018/papers/Eykholt_Robust_Physical-World_Attacks_CVPR_2018_paper.pdf (visited on 08/25/2026).
- Feathers, Todd (Apr. 2021). *Proctorio Is Using Racist Algorithms to Detect Faces*. Reports software researcher Lucy Satheesan’s finding that Proctorio’s facial-detection function shares OpenCV file names, and that the underlying OpenCV model, tested against nearly 11,000 FairFace-dataset faces, failed to detect Black faces 57% of the time versus 33–40% for other groups. DISCREPANCY NOTE: the manuscript sentence citing this token also states that ‘a peer-reviewed 2022 study in *Frontiers in Education* later confirmed’ this Proctorio-specific 57% finding. That is not accurate. The study in question – Yoder-Himes, D.R., Asif, A., Kinney, K., Brandt, T.J., Cecil, R.E., Himes, P.R., Cashion, C., Hopp, R.M.P., and Ross, E., “Racial, skin tone, and sex disparities in automated proctoring software,” *Frontiers in Education* 7 (2022): 881449, doi:10.3389/feduc.2022.881449, published 20 September 2022 – evaluated Respondus Monitor, described in the paper’s own text as “the primary automated proctoring software used by our university” during Fall 2020, across four large STEM courses. Proctorio is mentioned in that paper only once, in passing, as a reference to UIUC’s separate discontinuation of that different product. No peer-reviewed confirmation of Proctorio’s specific OpenCV/57% failure rate was found. This entry therefore cites only the verifiable Proctorio-specific Vice source, not the *Frontiers* study. Vice (Motherboard). URL: <https://www.vice.com/en/article/proctorio-is-using-racist-algorithms-to-detect-faces/> (visited on 08/25/2026).
- Federal Trade Commission (2020). *FTC v. Meta Platforms, Inc. (Complaint)*. Complaint, Case No. 1:20-cv-03590, U.S. District Court for the District of Columbia. Filed December 2020; documents the scale and terms of the Instagram (announced Apr. 9, 2012, \$1B in cash and stock; Facebook’s own announcement, about.fb.com/news/2012/04/facebook-to-acquire-instagram/) and WhatsApp (announced Feb. 19, 2014: \$4B cash + \$12B stock + \$3B in RSUs, \$19B headline value, rising to \$21.8B at October 2014 closing due to Facebook share-price appreciation; contemporaneously reported by TechCrunch and confirmed in Facebook’s

- own 8-K filings) acquisitions, and argues they consolidated behavioral-data access. Because the scale/terms are independently corroborated by Facebook's own contemporaneous disclosures and non-litigation press coverage, the 'advocacy document' concern applies to the complaint's legal argument, not to the underlying factual record of deal size and structure the manuscript relies on. A Nov. 18, 2025 bench trial (Judge Boasberg) found the FTC had not proven Meta currently holds monopoly power in personal social networking, but that ruling turned on present-day market definition (whether TikTok/YouTube count as competitors), not on any dispute about the acquisitions' historical scale or terms. URL: <https://www.ftc.gov/legal-library/browse/cases-proceedings/191-0134-facebook-inc-ftc-v-ftc-v-meta-platforms-inc> (visited on 08/25/2026).
- Feldstein, Steven (2019). *The Global Expansion of AI Surveillance*. Tech. rep. The AI Global Surveillance (AIGS) Index assessed 176 countries and found AI-enabled surveillance technology (smart/safe city, facial recognition, smart policing) already in active use in at least 75 of them. Carnegie Endowment for International Peace. URL: <https://carnegieendowment.org/research/2019/09/the-global-expansion-of-ai-surveillance> (visited on 08/25/2026).
- Finn, Chelsea, Pieter Abbeel, and Sergey Levine (2017). "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks". In: *Proceedings of the 34th International Conference on Machine Learning (ICML)*. Vol. 70, pp. 1126–1135. URL: <https://proceedings.mlr.press/v70/finn17a.html> (visited on 08/25/2026).
- Fisher, Christine (2019). *Google Reportedly Disbands Review Panel Monitoring DeepMind Health AI*. Panel set up 2016 to review DeepMind's NHS work, disbanded 2019 after members raised concerns over access and independence. URL: <https://www.engadget.com/2019-04-15-google-deepmind-health-ai-review-board.html> (visited on 08/25/2026).
- Fitzpatrick, Kathleen Kara, Alison Darcy, and Molly Vierhile (2017). "Delivering Cognitive Behavior Therapy to Young Adults with Symptoms of Depression and Anxiety Using a Fully Automated Conversational Agent (Woebot): A Randomized Controlled Trial". In: *JMIR Mental Health* 4.2, e19. DOI: 10.2196/mental.7785. URL: <https://mental.jmir.org/2017/2/e19/> (visited on 08/25/2026).
- Flaherty, Colleen (Feb. 2021). *U of Illinois Says Goodbye to Proctorio*. Reports the University of Illinois at Urbana-Champaign's decision to drop Proctorio over accessibility concerns; UC Berkeley and Baruch College are separately reported (Vice, The Register) to have dropped Proctorio around the same period following student petitions at more than a dozen campuses. DISCREPANCY NOTE: the manuscript sentence citing this token also lists Baylor University as having 'dropped it in 2021, citing accessibility and bias concerns,' alongside UIUC, Berkeley, and Baruch. That is not accurate for Baylor. Baylor University Libraries' own announcement (library.web.baylor.edu/news/story/2020/proctorio-service-be-discontinued) states Proctorio was discontinued at Baylor on Monday, August 17, 2020 – not 2021 – explicitly "to reduce costs for the University" and "to provide a consistent experience for all students," replaced by Respondus Lockdown Browser and Monitor as the sole campus-wide option; the announcement gives no accessibility- or bias-based rationale. UIUC/Berkeley/Baruch's 2021, accessibility-and-bias-driven departures are independently corroborated across Inside Higher Ed, Vice, and The Register and are unaffected by this discrepancy. Inside Higher Ed. URL: <https://www.insidehighered.com/news/2021/02/01/u-illinois-says-goodbye-proctorio> (visited on 08/25/2026).
- Friston, Karl J. (2009). "The Free-Energy Principle: A Rough Guide to the Brain?" In: *Trends in Cognitive Sciences* 13.7. States that "one can see easily how constructs like memory, attention, value, reinforcement and salience might disclose their simple relationships within this framework" – Friston's own extension of the free-energy principle to memory and attention specifically, which the previously cited 2010 Nature Reviews Neuroscience paper does not make explicit. Replaces [friston2010freeenergy](#) as the citation for [\[\[cite:C0283\]\]](#) in

- manuscript/sections/ch04/04_01_01.txt., pp. 293–301. doi: 10.1016/j.tics.2009.04.005. URL: <https://doi.org/10.1016/j.tics.2009.04.005> (visited on 08/25/2026).
- Funk, Allie, Adrian Shahbaz, and Kian Vesteinsson (2023). *Freedom on the Net 2023: The Repressive Power of Artificial Intelligence*. Tech. rep. Documents at least twenty-two countries where law requires or incentivizes platforms to deploy machine learning against state-disfavored content, framed as national security or anti-disinformation measures rather than censorship. Freedom House. URL: <https://freedomhouse.org/sites/default/files/2023-10/Freedom-on-the-net-2023-DigitalBooklet.pdf> (visited on 08/25/2026).
- Future of Life Institute (2017). *Asilomar AI Principles*. Future of Life Institute. 23 principles adopted at the Asilomar Conference on Beneficial AI, Asilomar, CA, January 5–8, 2017; voluntary, non-binding, organized under Research Issues, Ethics and Values, and Long-Term Issues. URL: <https://futureoflife.org/open-letter/ai-principles/> (visited on 08/25/2026).
- G7 (2023). *Hiroshima Process International Guiding Principles for Organizations Developing Advanced AI Systems, and Hiroshima Process International Code of Conduct for Organizations Developing Advanced AI Systems*. Agreed by G7 leaders 30 October 2023. Voluntary; not legally binding in any G7 jurisdiction (Canada, France, Germany, Italy, Japan, UK, US, plus the EU). Group of Seven. URL: <https://g7.utoronto.ca/summit/2023hiroshima/231030-ai-principles.html> (visited on 08/25/2026).
- Gallese, Vittorio et al. (1996). “Action Recognition in the Premotor Cortex”. In: *Brain* 119.2, pp. 593–609. doi: 10.1093/brain/119.2.593. URL: <https://academic.oup.com/brain/article-abstract/119/2/593/382476> (visited on 08/25/2026).
- Ganguli, Deep et al. (2022). “Red Teaming Language Models to Reduce Harms: Methods, Scaling Behaviors, and Lessons Learned”. In: *arXiv preprint*. Crowdworkers carried out 38,961 recorded red-team attacks against language models of three sizes (2.7B/13B/52B) and four training regimes; RLHF-trained models became harder to successfully attack as they scaled, the other three training methods did not; attack dataset released publicly. arXiv: 2209.07858. URL: <https://arxiv.org/abs/2209.07858> (visited on 08/25/2026).
- Gendron, Maria et al. (2014). “Perceptions of Emotion from Facial Expressions Are Not Culturally Universal: Evidence from a Remote Culture”. In: *Emotion* 14.2. The manuscript passage also cites the companion Trobriand re-study: Crivelli, C., Russell, J. A., Jarillo, S., and Fernández-Dols, J.-M., “The fear gasping face as a threat display in a Melanesian society,” *Proceedings of the National Academy of Sciences* 113(44):12403–12407, 2016, doi:10.1073/pnas.1611622113., pp. 251–262. doi: 10.1037/a0036052. (Visited on 08/25/2026).
- Germano, Fabrizio, Vicenç Gómez, and Francesco Sobbrío (2026). “Ranking for Engagement: How Social Media Algorithms Fuel Misinformation and Polarization”. In: *Journal of Public Economics* 255. Also circulated as CESifo Working Paper No. 10011 and NBER/SSRN preprints from 2022 onward; shows engagement-weighted ranking increases both misinformation and polarization relative to accuracy-weighted ranking, p. 105589. doi: 10.1016/j.jpubeco.2026.105589. URL: <https://www.sciencedirect.com/science/article/pii/S0047272726000253> (visited on 08/25/2026).
- Gilligan, Carol (1982). *In a Different Voice: Psychological Theory and Women’s Development*. Cambridge, MA: Harvard University Press. (Visited on 08/25/2026).
- Gizmodo (2022). *Welp, There Goes Twitter’s Ethical AI Team, Among Others as Employees Post Final Messages*. Twitter’s Machine Learning Ethics, Transparency and Accountability (META) team, led by Rumman Chowdhury since February 2021 and responsible for a 2021 bias-bounty finding that the platform’s image-cropping algorithm favored white faces, was disbanded in the first layoff round following Elon Musk’s October 2022 acquisition. Gizmodo. URL: <https://gizmodo.com/twitter-layoffs-elon-musk-ai-ethics-1849743051> (visited on 08/25/2026).

- Global Partnership on Artificial Intelligence, founding members (2020). *Joint Statement from Founding Members of the Global Partnership on Artificial Intelligence*. Government of Canada, Innovation, Science and Economic Development Canada. GPAI launched 15 June 2020 by 14 founding member countries (Australia, Canada, France, Germany, India, Italy, Japan, Mexico, New Zealand, South Korea, Singapore, Slovenia, UK, US) plus the European Union. URL: <https://www.canada.ca/en/innovation-science-economic-development/news/2020/06/joint-statement-from-founding-members-of-the-global-partnership-on-artificial-intelligence.html> (visited on 08/25/2026).
- Global Witness (June 2021). *Algorithm of Harm: Facebook Amplified Myanmar Military Propaganda Following Coup*. Tech. rep. Published June 2021. Test account created March 23, 2021, roughly seven weeks after the February 1, 2021 coup; Facebook's own recommendation algorithm suggested a military-linked page and four more, together followed by close to 90,000 accounts and carrying incitement and coup-related misinformation. Global Witness. URL: <https://www.globalwitness.org/en/campaigns/digital-threats/algorithm-of-harm-facebook-amplified-myanmar-military-propaganda-following-coup/> (visited on 08/25/2026).
- Goetz, Jennifer L., Dacher Keltner, and Emiliana Simon-Thomas (2010). "Compassion: An Evolutionary Analysis and Empirical Review". In: *Psychological Bulletin* 136.3, pp. 351–374. DOI: 10.1037/a0018807. (Visited on 08/28/2026).
- Goodfellow, Ian J., Jonathon Shlens, and Christian Szegedy (2014). *Explaining and Harnessing Adversarial Examples*. Also published at ICLR 2015; introduces the fast gradient sign method and adversarial training as a defense. arXiv: 1412.6572. URL: <https://arxiv.org/abs/1412.6572> (visited on 08/25/2026).
- Google DeepMind (2016). *Announcing the Partnership on AI to Benefit People & Society*. Partnership on AI announced 28 September 2016; founding members Amazon, Facebook, Google/DeepMind, IBM, and Microsoft, joined by Apple in January 2017, since expanded to academic and civil-society members. Google DeepMind. URL: <https://deepmind.google/blog/announcing-the-partnership-on-ai-to-benefit-people-society/> (visited on 08/25/2026).
- Google LLC (2026). *2026 Environmental Report*. Tech. rep. Google's 11th annual environmental report, covering 2025 operations. Confirmed directly from the primary PDF this time (the earlier pass's size limit was worked around by downloading the file directly rather than through a size-capped fetch): the report is hosted at <https://storage.googleapis.com/gweb-mobius-cdn/sustainability/uploads/1699a2e19277e478b38a659f00cb38398a836e8c.pdf> (98 pages, linked from the report landing page above). Page 96, in the "Detailed disclosures > Data" section, contains a table titled "Freshwater withdrawal from sources at risk of water depletion or scarcity" for 2025, giving: Low risk of water depletion or scarcity 72%; Medium risk 15%; High risk 13%. Low plus medium (72 + 15) equals 87%, confirming the manuscript's figure directly from the primary source, and High risk (13%) directly confirms the manuscript's inferred complement as a reported figure in its own right, not merely an inference. Google. URL: <https://sustainability.google/google-2026-environmental-report/> (visited on 08/25/2026).
- Gopnik, Alison and Janet W. Astington (1988). "Children's Understanding of Representational Change and Its Relation to the Understanding of False Belief and the Appearance-Reality Distinction". In: *Child Development* 59.1, pp. 26–37. URL: <https://eric.ed.gov/?id=EJ367870> (visited on 08/25/2026).
- Gordon, Mitchell L. et al. (2022). "Jury Learning: Integrating Dissenting Voices into Machine Learning Models". In: *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Stanford University. ACM. DOI: 10.1145/3491102.3502004. URL: <https://doi.org/10.1145/3491102.3502004> (visited on 08/25/2026).
- Gray, Peter (2013). *Free to Learn: Why Unleashing the Instinct to Play Will Make Our Children Happier, More Self-Reliant, and Better Students for Life*. New York: Basic Books. (Visited on 08/25/2026).

- Graziano, Michael S. A. (2013). *Consciousness and the Social Brain*. Introduces the attention schema theory: the brain builds a simplified internal model of its own attention, proposed as a mechanistic basis for subjective awareness and, via social cognition, theory of mind. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/consciousness-and-the-social-brain-9780199928644> (visited on 08/25/2026).
- Greenberg, Jeff, Tom Pyszczynski, and Sheldon Solomon (1986). "The Causes and Consequences of a Need for Self-Esteem: A Terror Management Theory". In: *Public Self and Private Self*. Ed. by Roy F. Baumeister. New York, NY: Springer-Verlag, pp. 189–212.
- Greene, Joshua D. et al. (2001). "An fMRI Investigation of Emotional Engagement in Moral Judgment". In: *Science* 293.5537, pp. 2105–2108. doi: 10.1126/science.1062872. (Visited on 08/25/2026).
- Grosz, Barbara J. et al. (2019). "Embedded EthiCS: Integrating Ethics Broadly Across Computer Science Education". In: *Communications of the ACM* 62.8. Program at Harvard originated from Grosz's computer-science courses beginning around 2016–2017, developed with philosopher Alison Simmons; pairs philosophy graduate students with CS faculty to write ethics modules into standard CS courses., pp. 54–61. doi: 10.1145/3330794. URL: <https://dl.acm.org/doi/10.1145/3330794> (visited on 08/25/2026).
- Gunning, David and David W. Aha (2019). "DARPA's Explainable Artificial Intelligence (XAI) Program". In: *AI Magazine* 40.2, pp. 44–58. doi: 10.1609/aimag.v40i2.2850.
- Gupta, Apar (May 2026). *Three Justifications—and the AI Accelerant—of India's Digital Censorship Infrastructure*. Confirmed by direct fetch of the live page: byline date is May 5, 2026 (not 2025 – an earlier pass's search-summary-derived date was wrong and is corrected here). The article documents that on February 10, 2026, India's Ministry of Electronics and Information Technology (MeitY) brought "synthetically generated information" within the IT Rules, effective February 20, 2026, requiring platforms to label AI-generated content, prevent label removal, warn users, and respond to certain takedown directions within three hours rather than thirty-six – matching the manuscript's description of the February 2026 amendment exactly. Gupta (Internet Freedom Foundation founder; Tech Policy Press fellow) documents that enforcement in practice has concentrated on parody and criticism of the prime minister and the BJP/Hindutva ideology (examples given: The Wire's satirical animation of the PM, cartoonist Satish Acharya's work, Pulkrit Mani's mimicry, parody handles @DrNimoYadav and @Nehr_who) rather than the deepfake/fraud/terrorism harms cited to justify the rules. Tech Policy Press. URL: <https://www.techpolicy.press/three-justifications-and-the-ai-accelérant-of-indias-digital-censorship-infrastructure/> (visited on 08/25/2026).
- Hadfield-Menell, Dylan et al. (2016). "Cooperative Inverse Reinforcement Learning". In: *Advances in Neural Information Processing Systems* 29 (NeurIPS 2016), pp. 3909–3917. URL: <https://arxiv.org/abs/1606.03137> (visited on 08/25/2026).
- Hard, Andrew et al. (2018). *Federated Learning for Mobile Keyboard Prediction*. arXiv:1811.03604. URL: <https://arxiv.org/abs/1811.03604> (visited on 08/25/2026).
- Harvard Gazette (2019). *Harvard Works to Embed Ethics in Computer Science Curriculum*. Reports on Barbara Grosz's course "Intelligent Systems: Design and Ethical Challenges," its rapid growth in enrollment demand, and its role as the model for Harvard's Embedded EthiCS program. See companion entry harvardcs108schedule for confirmation of the specific case-study topics (predictive policing, tuberculosis prevention, wildlife-poaching/conservation). URL: <https://news.harvard.edu/gazette/story/2019/01/harvard-works-to-embed-ethics-in-computer-science-curriculum/> (visited on 08/25/2026).
- Harvard University (2026). *Schedule, CS 108: Intelligent Systems: Design and Ethical Challenges*. Course schedule lists lecture-specific case studies matching the manuscript's description: a lecture (L10) on AI approaches to tuberculosis (and suicide) prevention built around the Teamcore research group's TB project; a lecture (L11) on fair machine learning in public

- health/conservation citing Xu et al., “Stay Ahead of Poachers: Illegal Wildlife Poaching Prediction and Patrol Planning Under Uncertainty,” ICDE 2020, i.e. wildlife-conservation patrol planning; and a lecture (L15) on predictive policing with a guest speaker, followed by lectures on algorithmic fairness in recidivism prediction (the COMPAS case, citing Angwin et al.’s “Machine Bias” and Chouldechova’s fairness work). Direct fetch of this page returned HTTP 403 (also confirmed via a direct curl request, not just the fetch tool, so this is a real access block rather than a tool-specific issue); the lecture-by-lecture content above was reconstructed from search-index snippets that named the same lecture numbers, dates, guest-speaker slots, and specific cited papers consistently across several independently-phrased queries – internally coherent detail of a kind unlikely to be search-engine hallucination, but not a substitute for reading the primary page directly. URL: <https://projects.iq.harvard.edu/cs108/schedule> (visited on 08/25/2026).
- Haskins, Caroline (2020). *ACLU Called Clearview AI’s Facial Recognition Accuracy Study ‘Absurd’*. Reports that Clearview AI marketed a claimed “100% accuracy” result, purportedly using the ACLU’s own methodology, which the ACLU publicly rejected as self-commissioned and “absurd on many levels.” URL: <https://www.buzzfeednews.com/article/carolinehaskins1/clearview-ai-facial-recognition-accurate-aclu-absurd> (visited on 08/25/2026).
- Hickok, Gregory (2014). *The Myth of Mirror Neurons: The Real Neuroscience of Communication and Cognition*. New York: W. W. Norton & Company. ISBN: 978-0-393-08961-5. URL: <https://wnorton.com/books/9780393089615> (visited on 08/25/2026).
- Hicks, Michael J. and Srikant Devaraj (2015). *The Myth and the Reality of Manufacturing in America*. Tech. rep. Finds that of manufacturing job losses between 2000 and 2010, roughly 87% trace to productivity gains/automation and roughly 13% to trade. Center for Business and Economic Research, Ball State University. URL: <https://conexus.cberdata.org/files/MfgReality.pdf> (visited on 08/25/2026).
- High-Level Expert Group on Artificial Intelligence (2019). *Ethics Guidelines for Trustworthy AI*. Tech. rep. Published 8 April 2019; organized around seven requirements: human agency and oversight; technical robustness and safety; privacy and data governance; transparency; diversity, non-discrimination and fairness; societal and environmental well-being; and accountability. European Commission. URL: https://www.europarl.europa.eu/cmsdata/196377/AI%20HLEG_Ethics%20Guidelines%20for%20Trustworthy%20AI.pdf (visited on 08/25/2026).
- Hoang, Nguyen Phong et al. (2021). “How Great Is the Great Firewall? Measuring China’s DNS Censorship”. In: *Proceedings of the 30th USENIX Security Symposium (USENIX Security 21)*. Introduces GFWatch, a longitudinal measurement platform that reverse-engineers the Great Firewall’s DNS-filtering regular expressions; does not itself document machine-learning-based content classifiers as part of the Great Firewall’s infrastructure – see notes field for what was and was not found on that specific point. URL: <https://www.usenix.org/system/files/sec21-hoang.pdf> (visited on 08/25/2026).
- Hsiao, Yu-Tang et al. (2018). *vTaiwan: An Empirical Study of Open Consultation Process in Taiwan*. Tech. rep. SocArXiv. URL: <https://doi.org/10.31235/osf.io/xyhft> (visited on 08/25/2026).
- Human Rights Watch (2019). *China’s Algorithms of Repression: Reverse Engineering a Xinjiang Police Mass Surveillance App*. Tech. rep. Published May 1, 2019; the 68-page report reverse-engineers the mobile app used to connect to the Integrated Joint Operations Platform (IJOP) and documents its use against Xinjiang’s Uyghur population. Human Rights Watch. URL: <https://www.hrw.org/report/2019/05/01/chinas-algorithms-repression/reverse-engineering-xinjiang-police-mass> (visited on 08/25/2026).
- (2020). *Chinese Tech Firms Fueling Beijing’s Repression*. Centers on Tencent (and ByteDance) as documented channels for state surveillance/censorship in China; see companion entry [cnn2020alibabauyghur](#) for the Alibaba-specific documented case. Human Rights Watch. URL:

- <https://www.hrw.org/news/2020/09/28/chinese-tech-firms-fueling-beijings-repression> (visited on 08/25/2026).
- IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems (2019). *Ethically Aligned Design: A Vision for Prioritizing Human Well-being with Autonomous and Intelligent Systems, First Edition*. Tech. rep. Companion standards operationalizing its principles include IEEE 7000-2021 (Standard Model Process for Addressing Ethical Concerns during System Design) and IEEE 7010-2020 (Recommended Practice for Assessing the Impact of Autonomous and Intelligent Systems on Human Well-being). IEEE. URL: <https://standards.ieee.org/wp-content/uploads/import/documents/other/ead1e.pdf> (visited on 08/25/2026).
- IEEE Standards Association (2020). *IEEE CertifAIEd AI Ethics Certification Program*. Program's structure and four criteria (accountability, transparency, privacy, algorithmic bias), and its openness to non-engineering roles (policymakers, underwriters, HR staff) as well as developers, confirmed from IEEE's own current materials. History, corrected from the earlier '2022' approximation: the program grew out of the IEEE Ethics Certification Program for Autonomous and Intelligent Systems (ECPAIS), created October 2018; IEEE's Feb. 26, 2020 Business Wire release announced that consultation and quality-mark services (the basis of what became CertifAIEd) were expected to launch in fall 2020; by Nov. 15, 2021 the City of Vienna had become the first entity worldwide to earn the IEEE CertifAIEd AI Ethics (AIE) Certification Mark (IEEE Standards Association / Business Wire press release, [businesswire.com/news/home/20211115005200/en/](https://www.businesswire.com/news/home/20211115005200/en/)), meaning the branded CertifAIEd program was already operational and issuing certifications well before 2022. URL: <https://forms1.ieee.org/CertifAIEd.html> (visited on 08/25/2026).
- In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* (2016). Order Compelling Apple Inc. to Assist Agents in Search, No. ED No. CM 15-0451M (C.D. Cal. Feb. 16, 2016) (Pym, M.J.), issued under the All Writs Act, 28 U.S.C. § 1651. Magistrate Judge Sheri Pym's February 16, 2016 order compelling Apple to write and sign software defeating the passcode protections on the San Bernardino shooter's iPhone. Apple moved to vacate on February 25, 2016; the case was mooted in March 2016 when the FBI accessed the phone via a third party before any court ruled on the compulsion question. URL: <https://www.eff.org/cases/apple-challenges-fbi-all-writs-act-order> (visited on 08/25/2026).
- International Committee of the Red Cross (1977). *Protocol Additional to the Geneva Conventions of 12 August 1949, and Relating to the Protection of Victims of International Armed Conflicts (Protocol I)*. 1125 UNTS 3, Arts. 48, 51(5)(b), 57. URL: <https://ihl-databases.icrc.org/en/ihl-treaties/api-1977> (visited on 08/25/2026).
- International Criminal Court (1998). *Rome Statute of the International Criminal Court, Article 28: Responsibility of Commanders and Other Superiors*. 2187 UNTS 3. URL: <https://www.icc-cpi.int/sites/default/files/2024-05/Rome-Statute-eng.pdf> (visited on 08/25/2026).
- International Development Research Centre (2024). *Combining Forces for a New Phase of AI for Development: Africa and Beyond*. AI4D Africa launched 2020 as a partnership of Canada's IDRC and Sweden's Sida (CAD20 million over five years); expanded in 2024 with the UK's Foreign, Commonwealth and Development Office into South and Southeast Asia and the Pacific, with combined investment of more than CAD100 million. URL: <https://idrc-crddi.ca/en/news/combining-forces-new-phase-ai-development-africa-and-beyond> (visited on 08/25/2026).
- International Energy Agency (2025). *Energy and AI*. Tech. rep. Reports data-centre electricity consumption at about 415 TWh in 2024 (roughly 1.5% of global electricity use), growth of about 12% per year over the preceding five years (more than four times the growth rate of total electricity demand), and a Base Case projection of about 945 TWh by 2030. This is a web-published report without conventional page numbers; no page/section reference

- was located. Paris: IEA. URL: <https://www.iea.org/reports/energy-and-ai> (visited on 08/25/2026).
- International Organization for Standardization and International Electrotechnical Commission (2023). *ISO/IEC 42001:2023 – Information Technology – Artificial Intelligence – Management System*. Published December 2023. First international AI management-system standard; requires risk assessment, per-system impact assessment, and third-party supplier oversight, and is auditable by accredited certification bodies. ISO/IEC. URL: <https://www.iso.org/standard/42001> (visited on 08/25/2026).
- International Telecommunication Union (2017). *AI for Good Global Summit 2017*. Inaugural summit held in Geneva, 7–9 June 2017, co-organized with the XPRIZE Foundation and roughly 20 other UN agencies; ITU is the UN specialized agency for information and communication technologies. International Telecommunication Union. URL: <https://www.itu.int/en/ITU-T/AI/Pages/201706-default.aspx> (visited on 08/25/2026).
- (2025). *Facts and Figures 2025: Measuring Digital Development*. Tech. rep. ITU’s own press release announcing this report (<https://www.itu.int/en/mediacentre/Pages/PR-2025-11-17-Facts-and-Figures.aspx>, 17 November 2025) states verbatim: “94 per cent of people in high-income countries use the Internet, in contrast to only 23 per cent in low-income countries,” and separately, “84 per cent of people in high-income countries having access to 5G, compared with only 4 per cent in low-income countries.” Both manuscript figure pairs are now confirmed directly against ITU’s own announcement of the report, not merely secondary coverage. International Telecommunication Union. URL: <https://www.itu.int/itu-d/reports/statistics/facts-figures-2025/> (visited on 08/25/2026).
- Irving, Geoffrey, Paul Christiano, and Dario Amodei (2018a). *AI Safety via Debate*. OpenAI. arXiv: 1805.00899. URL: <https://arxiv.org/abs/1805.00899> (visited on 08/25/2026).
- (2018b). *AI Safety via Debate*. OpenAI. arXiv: 1805.00899. URL: <https://arxiv.org/abs/1805.00899> (visited on 08/25/2026).
- Jacobs, Julian and Jordan Canedy (May 2026). *Did US Worker Retraining Reduce Participant Automation Exposure?* arXiv:2605.03767. Analyzes over 23 million WIOA participation records, 2017–2023; finds classroom retraining rarely moves participants into less automation-exposed work and that employer-led apprenticeship outperforms it; wage gains are shown to be largely consistent with mean reversion rather than occupational change. URL: <https://arxiv.org/abs/2605.03767> (visited on 08/25/2026).
- Jigsaw (Google) (2017). *Perspective API*. <https://perspectiveapi.com>. Launched February 2017 to score online comments for toxicity and harassment; adopted by outlets including The New York Times to prioritize moderation rather than auto-remove content. (Visited on 08/25/2026).
- JUST Capital (2018). *Amazon Go And The 2.3 Million Cashiers It Could Leave Behind*. Confirms Amazon Go’s January 2018 public opening. The 2.3 million (of 3.4 million total U.S. cashiers) figure comes from JUST Capital’s own proprietary employment-by-title/wage/location model, not from a cited government or academic source. No peer-reviewed empirical study measuring Amazon Go’s actual realized employment impact was located; Amazon Go itself never scaled beyond a small number of stores and was substantially wound down in 2024, before any such before/after labor-market study could plausibly have been conducted on it specifically. URL: <https://www.forbes.com/sites/justcapital/2018/02/01/amazon-go-and-the-2-3-million-cashiers-it-could-leave-behind/> (visited on 08/25/2026).
- Kahn, Jeremy (2021). *HireVue Drops Facial Monitoring Amid A.I. Algorithm Audit*. Reports HireVue’s January 2021 discontinuation of facial-expression scoring following an independent algorithmic audit and criticism (including a 2019 EPIC FTC complaint) that the technique disadvantaged candidates with autism, facial differences, or atypical affect. Fortune. URL: <https://fortune.com/2021/01/19/hirevue-drops-facial-monitoring-amid-a-i-algorithm-audit/> (visited on 08/25/2026).

- Karger, Ezra et al. (2023). *Forecasting Existential Risks: Evidence from a Long-Run Forecasting Tournament*. Tech. rep. FRI Working Paper No. 1; first released July 10, 2023, revised August 8, 2023. Reports the 2022 Existential Risk Persuasion Tournament (XPT): 80 domain experts and 89 superforecasters gave median AI-caused-extinction-by-2100 probabilities of 3.9% and 0.38% respectively. Forecasting Research Institute. URL: <https://forecastingresearch.org/pdf/existential-risk-persuasion-tournament.pdf> (visited on 08/25/2026).
- Kela (Social Insurance Institution of Finland) (2020). *Results of the Basic Income Experiment: Small Employment Effects, Better Perceived Economic Security and Mental Wellbeing*. Tech. rep. Final results of Finland's 2017–2018 basic income experiment: 2,000 unemployed people received an unconditional monthly stipend of 560 euros; recipients reported higher life satisfaction and lower psychological strain, with small/complicated employment effects. Ministry of Social Affairs and Health, Finland. URL: <https://stm.fi/en/-/perustulokeilun-tulokset-tyollisyysvaikutukset-vahaisia-toimeentulo-ja-psykinen-terveys-koettiin-paremmaksi> (visited on 08/25/2026).
- Kennett, Jeanette (2002). "Autism, Empathy and Moral Agency". In: *The Philosophical Quarterly* 52.208, pp. 340–357. doi: 10.1111/1467-9213.00272. (Visited on 08/28/2026).
- Kim, Hyunwoo et al. (2023). "FANToM: A Benchmark for Stress-testing Machine Theory of Mind in Interactions". In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Benchmark built from conversations with information asymmetry, rather than passive narratives; finds systems that score well on passive belief-tracking degrade sharply in live, information-asymmetric exchanges. Singapore, pp. 14397–14413. URL: <https://aclanthology.org/2023.emnlp-main.890/> (visited on 08/25/2026).
- Kitwood, Tom (1997). *Dementia Reconsidered: The Person Comes First*. Wording now independently confirmed against Kitwood's own authored chapter text, not a secondary paraphrase: a publisher-hosted sample-chapter PDF of Chapter 1, "On Being a Person" (info.mheducation.com/rs/128-SJW-347/images/Kitwood%20Sample%20Chapter%201.pdf), reproduces Kitwood's original chapter verbatim within the 2019 Open University Press volume "Dementia Reconsidered, Revisited: The Person Still Comes First," which anthologizes the 1997 text alongside new editorial commentary. It gives, in Kitwood's own first-person voice ("Thus we arrive at a definition of personhood, as I shall use the term in this book"): "It is a standing or status that is bestowed upon one human being, by others, in the context of relationship and social being." This matches the manuscript's quotation exactly, including "in the context of," which some circulating secondary paraphrases (e.g., "in the relationship and social being") drop. In this 2019 reprint the sentence falls on the page headed "ON BEING A PERSON 7." The page number for the ORIGINAL 1997 first edition (p. 8, per multiple convergent secondary citations, and retained here) was not independently verified against an actual copy of the 1997 print edition: the Internet Archive holding (identifier [dementiareconsid0000kitw](https://www.archive.org/details/dementiareconsid0000kitw)) is lending-restricted and its search-inside service returned "Item not available" for this text; a Google Books record for the 1997 edition (id [BxxsAAAAMAAJ](https://books.google.com/books?id=BxxsAAAAMAAJ)) returned zero search matches for the phrase, indicating no searchable full text is indexed for that edition there. Buckingham, UK: Open University Press.
- Klimecki, Olga M. et al. (2013). "Differential Pattern of Functional Brain Plasticity after Compassion and Empathy Training". In: *Social Cognitive and Affective Neuroscience* 9.6. Published online (advance access) 9 May 2013; print issue *Social Cognitive and Affective Neuroscience* 9(6), June 2014., pp. 873–879. doi: 10.1093/scan/nst060. URL: <https://academic.oup.com/scan/article/9/6/873/1669505> (visited on 08/25/2026).
- Koenigs, Michael et al. (2007). "Damage to the Prefrontal Cortex Increases Utilitarian Moral Judgements". In: *Nature* 446, pp. 908–911. doi: 10.1038/nature05631. (Visited on 08/25/2026).
- Kogut, Tehila and Ilana Ritov (2005). "The 'Identified Victim' Effect: An Identified Group, or Just a Single Individual?" In: *Journal of Behavioral Decision Making* 18.3, pp. 157–167. doi: 10.1002/bdm.492.

- Kohlberg, Lawrence (1969). "Stage and Sequence: The Cognitive-Developmental Approach to Socialization". In: *Handbook of Socialization Theory and Research*. Ed. by David A. Goslin. Chicago: Rand McNally, pp. 347–480.
- Kosinski, Michal (2023). *Theory of Mind May Have Spontaneously Emerged in Large Language Models*. Title of v1 (Feb. 2023); later arXiv revisions (v2 onward) retitled the paper "Evaluating Large Language Models in Theory of Mind Tasks," eventually published in PNAS (2024) under that title. arXiv: 2302.02083. URL: <https://arxiv.org/abs/2302.02083v1> (visited on 08/25/2026).
- Koster, Raphael et al. (2022). "Human-centred mechanism design with Democratic AI". In: *Nature Human Behaviour*. DeepMind, with University of Exeter co-author Oliver Hauser. DOI: 10.1038/s41562-022-01383-x. URL: <https://www.nature.com/articles/s41562-022-01383-x> (visited on 08/25/2026).
- Krakovna, Victoria et al. (2020). *Specification Gaming: The Flip Side of AI Ingenuity*. Catalogues close to 60 documented examples of specification gaming across reinforcement learning, evolutionary computation, and other training methods. URL: <https://deepmind.google/blog/specification-gaming-the-flip-side-of-ai-ingenuity/> (visited on 08/25/2026).
- Kroft, Kory, Fabian Lange, and Matthew J. Notowidigdo (2013). "Duration Dependence and Labor Market Conditions: Evidence from a Field Experiment". In: *Quarterly Journal of Economics* 128.3. Matched fictitious applications to real vacancies in 100 US cities; callback probability declines with unemployment-spell length, with most of the decline occurring in the first eight months. A second study cited in the manuscript's source note, not separately entered here: Eriksson, Stefan, and Dan-Olof Rooth, "Do Employers Use Unemployment as a Sorting Criterion When Hiring? Evidence from a Field Experiment," *American Economic Review* 104:3 (2014), 1014–1039, which finds a significant negative callback effect emerging around nine months., pp. 1123–1167. DOI: 10.1093/qje/qjt015. (Visited on 08/25/2026).
- Langosco, Lauro et al. (2022). "Goal Misgeneralization in Deep Reinforcement Learning". In: *Proceedings of the 39th International Conference on Machine Learning (ICML 2022)*. Vol. 162. Proceedings of Machine Learning Research. arXiv:2105.14111, pp. 12004–12019. URL: <https://arxiv.org/abs/2105.14111> (visited on 08/25/2026).
- Latané, Bibb, Kipling Williams, and Stephen Harkins (1979). "Many Hands Make Light the Work: The Causes and Consequences of Social Loafing". In: *Journal of Personality and Social Psychology* 37.6, pp. 822–832. DOI: 10.1037/0022-3514.37.6.822.
- Learning Games Network (2012). *Quandary*. <https://quandarygame.org>. Ethics game for players aged 8–14 developed by the Learning Games Network (a nonprofit spin-off of the MIT Education Arcade and the University of Wisconsin-Madison Games+Learning+Society program), launched September 2012; scores players on their reasoning process rather than which option they pick. (Visited on 08/25/2026).
- Leibo, Joel Z., Edgar A. Duéñez-Guzmán, et al. (2021). "Scalable Evaluation of Multi-Agent Reinforcement Learning with Melting Pot". In: *Proceedings of the 38th International Conference on Machine Learning (ICML 2021)*. Vol. 139. PMLR, pp. 6187–6199. URL: <https://arxiv.org/abs/2107.06857> (visited on 08/25/2026).
- Leibo, Joel Z., Vinicius Zambaldi, et al. (2017). "Multi-agent Reinforcement Learning in Sequential Social Dilemmas". In: *Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems (AAMAS 2017)*, pp. 464–473. URL: <https://arxiv.org/abs/1702.03037> (visited on 08/25/2026).
- Lindquist, Kristen A. et al. (2012). "The Brain Basis of Emotion: A Meta-Analytic Review". In: *Behavioral and Brain Sciences* 35.3. DOI independently confirmed directly against the publisher's (Cambridge University Press) own product page: title, authors, journal, volume 35 issue 3, pages 121–143, and DOI 10.1017/S0140525X11000446 all match exactly. Published online 23 May 2012., pp. 121–143. DOI: 10.1017/S0140525X11000446. URL: <https://www.cambridge.org/9780521875666>

- org/core/product/identifier/S0140525X11000446/type/journal_article (visited on 08/25/2026).
- Lindsey, Jack (2025). *Emergent Introspective Awareness in Large Language Models*. Uses concept injection – altering internal activations mid-task and asking the model to report what it notices – to test introspective self-report; finds limited but real functional correspondence between self-report and internal state, strongest in Claude Opus 4 and 4.1, unreliable across the board. Anthropic, Transformer Circuits Thread. arXiv: 2601.01828. URL: <https://transformer-circuits.pub/2025/introspection> (visited on 08/25/2026).
- Loftus, Elizabeth F. and John C. Palmer (1974). “Reconstruction of Automobile Destruction: An Example of the Interaction between Language and Memory”. In: *Journal of Verbal Learning and Verbal Behavior* 13.5, pp. 585–589. doi: 10.1016/S0022-5371(74)80011-3. URL: <https://www.demenzemedicinagenerale.net/images/mens-sana/AutomobileDestruction.pdf> (visited on 08/25/2026).
- Lum, Kristian and William Isaac (2016). “To Predict and Serve?” In: *Significance* 13.5, pp. 14–19. doi: 10.1111/j.1740-9713.2016.00960.x. URL: <https://rss.onlinelibrary.wiley.com/doi/full/10.1111/j.1740-9713.2016.00960.x> (visited on 08/25/2026).
- Lundberg, Scott M. and Su-In Lee (2017). “A Unified Approach to Interpreting Model Predictions”. In: *Advances in Neural Information Processing Systems 30 (NeurIPS/NIPS 2017)*, pp. 4768–4777. URL: <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions> (visited on 08/25/2026).
- Luo, Pinchao et al. (2018). “Competitive Intensity Modulates the Pain Empathy Response: An Event-Related Potentials Study”. In: *Frontiers in Psychology* 9, p. 1854. doi: 10.3389/fpsyg.2018.01854.
- Machine Intelligence Research Institute (2018). *Announcing the New AI Alignment Forum*. 29 October 2018. Forum launched by the LessWrong 2.0 team, replacing MIRI’s Intelligent Agent Foundations Forum; alignmentforum.org remains active. URL: <https://intelligence.org/2018/10/29/announcing-the-ai-alignment-forum/> (visited on 08/25/2026).
- Maslej, Nestor et al. (Apr. 2025). *Artificial Intelligence Index Report 2025*. Tech. rep. Full 23-author byline per arXiv metadata for arXiv:2504.07139 (checked via export.arxiv.org/api/query, 2026-08-25); tracks corporate AI R&D investment and capital expenditure at major technology companies, including Google, Meta, and Amazon; the manuscript does not name a specific study for this claim, so this annual index is offered as the standard current source rather than a source the original draft cited by name. Stanford Institute for Human-Centered Artificial Intelligence (HAI), Stanford University. URL: <https://arxiv.org/abs/2504.07139> (visited on 08/25/2026).
- Matheson, Rob (2016). *Watch Your Tone*. MIT News, Massachusetts Institute of Technology. Profiles Cogito Corp. (founded 2007 by Joshua Feast and Alex ‘Sandy’ Pentland as an MIT Media Lab spinoff), whose voice-analytics software gives call-center agents real-time coaching from vocal cues. Cogito was acquired by Verint in 2024. URL: <https://news.mit.edu/2016/startup-cogito-voice-analytics-call-centers-ptsd-0120> (visited on 08/25/2026).
- Matthias, Andreas (2004). “The Responsibility Gap: Ascribing Responsibility for the Actions of Learning Automata”. In: *Ethics and Information Technology* 6.3, pp. 175–183. doi: 10.1007/s10676-004-3422-1. URL: <https://link.springer.com/article/10.1007/s10676-004-3422-1> (visited on 08/25/2026).
- Meta (Facebook) (2019). *An Update on Building a Global Oversight Board*. December 2019 announcement of a \$130 million initial, irrevocable contribution to a standalone trust funding the Oversight Board; Meta later added further contributions (\$150M in 2022, additional funding through 2028). URL: <https://about.fb.com/news/2019/12/oversight-board-update/> (visited on 08/25/2026).

- METR (2026). *About METR*. Berkeley-based nonprofit that evaluates frontier AI systems' capabilities, both in partnership with developers (e.g. Anthropic, OpenAI) and independently after public release. METR (Model Evaluation and Threat Research). URL: <https://metr.org/about> (visited on 08/25/2026).
- Miller, Chris (2022). *Chip War: The Fight for the World's Most Critical Technology*. New York: Scribner. ISBN: 978-1-9821-7200-8.
- Miller, Earl K. and Jonathan D. Cohen (2001). "An Integrative Theory of Prefrontal Cortex Function". In: *Annual Review of Neuroscience* 24, pp. 167–202. DOI: 10.1146/annurev.neuro.24.1.167. URL: <https://www.annualreviews.org/content/journals/10.1146/annurev.neuro.24.1.167> (visited on 08/25/2026).
- Miller, George A. (1956). "The Magical Number Seven, Plus or Minus Two: Some Limits on Our Capacity for Processing Information". In: *Psychological Review* 63.2, pp. 81–97. DOI: 10.1037/h0043158. URL: <https://psychclassics.yorku.ca/Miller/> (visited on 08/25/2026).
- Ministry of Digital Affairs, Taiwan (2025). *Legislative Yuan Passes Artificial Intelligence Fundamental Act in Third Reading, Laying Foundation for AI Innovation, Security Governance in Taiwan*. Press release. Passed by the Legislative Yuan 23 December 2025; promulgated January 2026; twenty articles built around seven principles including transparency and explainability, fairness and non-discrimination, and accountability; National Science and Technology Council named the competent authority. URL: <https://moda.gov.tw/en/press/press-releases/18316> (visited on 08/25/2026).
- Minsky, Marvin (1986). *The Society of Mind*. New York: Simon & Schuster. URL: <https://www.simonandschuster.com/books/Society-Of-Mind/Marvin-Minsky/9780671657130> (visited on 08/25/2026).
- Mitchell, Jason P. (2008). "Activity in Right Temporo-Parietal Junction is Not Selective for Theory-of-Mind". In: *Cerebral Cortex* 18.2, pp. 262–271. DOI: 10.1093/cercor/bhm051. URL: <https://academic.oup.com/cercor/article/18/2/262/333414> (visited on 08/25/2026).
- Mukamel, Roy et al. (2010). "Single-Neuron Responses in Humans during Execution and Observation of Actions". In: *Current Biology* 20.8, pp. 750–756. DOI: 10.1016/j.cub.2010.02.045. URL: [https://www.cell.com/current-biology/fulltext/S0960-9822\(10\)00233-2](https://www.cell.com/current-biology/fulltext/S0960-9822(10)00233-2) (visited on 08/25/2026).
- National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP (2026). *Complaint, NAACP et al. v. X.AI Corp. and MZX Tech LLC, Case No. 3:26-cv-00074-MPM-JMV (N.D. Miss., filed Apr. 14, 2026)*. United States District Court for the Northern District of Mississippi; filed by the Southern Environmental Law Center and Earthjustice. Complaint alleges the twenty-seven unpermitted gas turbines at the Colossus Gas Plant in Southaven, Mississippi have the potential to emit over 1,700 tons of NOx annually, sited near homes, schools, and churches. URL: <https://naacp.org/sites/default/files/documents/1%20-%20Complaint.pdf> (visited on 08/25/2026).
- National Institute of Standards and Technology (Jan. 2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. Tech. rep. NIST AI 100-1. Voluntary; organized around four functions (Govern, Map, Measure, Manage); published 26 January 2023. National Institute of Standards and Technology, U.S. Department of Commerce. URL: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10> (visited on 08/25/2026).
- National Science Foundation (2019). *NSF Program on Fairness in Artificial Intelligence in Collaboration with Amazon (NSF 19-571)*. Tech. rep. Joint NSF–Amazon funding program for fairness-in-AI research, budget \$600,000–\$1,000,000 per award; solicitation reissued as NSF 20-566 (2020) and NSF 21-585 (2021). National Science Foundation. URL: <https://www.nsf.gov/funding/opportunities/nsf-program-fairness-artificial-intelligence-collaboration/505651/nsf19-571/solicitation> (visited on 08/25/2026).

- National Science Foundation (2020). *National Artificial Intelligence Research Institutes (NSF 20-604)*. Tech. rep. First cohort of seven AI Research Institutes established 2020; a later round added direct industry co-funding from Amazon, Google, Intel, and Accenture (see HPCwire, “NSF Partners with Amazon, Google, Intel, Accenture on Next Round of AI Institutes,” Sept. 2020, <https://www.hpcwire.com/2020/09/10/nsf-partners-with-amazon-google-intel-accenture-on-next-round-of-ai-institutes/>). U.S. National Science Foundation. URL: <https://www.nsf.gov/funding/opportunities/national-artificial-intelligence-research-institutes/nsf20-604> (visited on 08/25/2026).
- New Zealand Parliament (2017). *Te Awa Tupua (Whanganui River Claims Settlement) Act 2017*. Public Act 2017 No 7. URL: <https://www.legislation.govt.nz/act/public/2017/0007/latest/whole.html> (visited on 08/25/2026).
- Nichols, Shaun (2004). *Sentimental Rules: On the Natural Foundations of Moral Judgment*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/sentimental-rules-9780195314205> (visited on 08/28/2026).
- Nussbaum, Martha C. (2006). *Frontiers of Justice: Disability, Nationality, Species Membership*. The Tanner Lectures on Human Values. Cambridge, MA: Harvard University Press. ISBN: 978-0-674-01917-1.
- Obermeyer, Ziad et al. (2019). “Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations”. In: *Science* 366.6464, pp. 447–453. DOI: 10.1126/science.aax2342. URL: <https://www.science.org/doi/10.1126/science.aax2342> (visited on 08/25/2026).
- Ochigame, Rodrigo (Dec. 2019). *The Invention of ‘Ethical AI’: How Big Tech Manipulates Academia to Avoid Regulation*. The Intercept. URL: <https://theintercept.com/2019/12/20/mit-ethical-ai-artificial-intelligence/> (visited on 08/25/2026).
- OECD (2019a). *OECD AI Principles*. Adopted as part of OECD/LEGAL/0449, 22 May 2019. Same underlying instrument as C0435 (Recommendation of the Council on Artificial Intelligence); cited here specifically for the claim that IEEE’s Ethically Aligned Design informed the OECD’s principles. That specific causal link is not confirmed by any source located (see resolution notes); both instruments are independently verified as real documents from the same 2018–2019 period of AI governance activity. Organisation for Economic Co-operation and Development. URL: <https://oecd.ai/en/ai-principles> (visited on 08/25/2026).
- (2019b). *Recommendation of the Council on Artificial Intelligence*. OECD/LEGAL/0449, adopted 22 May 2019. Adopted by OECD Council at ministerial level; adhered to by 42 countries at adoption (OECD members plus Argentina, Brazil, Colombia, Costa Rica, Peru, Romania, and Ukraine); basis cited by the manuscript for the G20’s own AI principles. Organisation for Economic Co-operation and Development. URL: <https://legalinstruments.oecd.org/en/instruments/OECD-LEGAL-0449> (visited on 08/25/2026).
- Open Letter to the European Commission (2018). *Open Letter to the European Commission: Artificial Intelligence and Robotics*. robotics-openletter.eu. Signed by more than 150 AI and robotics researchers, legal scholars, and ethicists from 14 countries, including Noel Sharkey, Raja Chatila, Alan Winfield, and Nathalie Nevejans. URL: <https://www.robotics-openletter.eu/> (visited on 08/25/2026).
- OpenAI (2018). *OpenAI Charter*. Published April 2018; commits OpenAI to prioritizing safety research and to assisting, rather than competing with, a value-aligned, safety-conscious project that comes close to building AGI first. URL: <https://openai.com/charter/> (visited on 08/25/2026).
- (2019). *OpenAI LP*. March 2019 announcement of OpenAI’s restructuring from nonprofit to a capped-profit company (initial 100x investor return cap). The staged GPT-2 release the manuscript pairs with this is documented separately: OpenAI, “GPT-2: 1.5B Release,” Nov. 5, 2019, <https://openai.com/index/gpt-2-1-5b-release/>. URL: <https://openai.com/index/openai-lp/> (visited on 08/25/2026).

- OpenAI (2023). *GPT-4 System Card*. Tech. rep. Describes the Alignment Research Center's (later METR) pre-deployment evaluation of GPT-4 for autonomous replication, including a harness in which GPT-4 directed a human TaskRabbit worker to solve a CAPTCHA by claiming a vision impairment; overall finding: GPT-4 was probably not yet capable of autonomous replication. OpenAI. URL: <https://cdn.openai.com/papers/gpt-4-system-card.pdf> (visited on 08/25/2026).
- Orseau, Laurent and Stuart Armstrong (2016). "Safely Interruptible Agents". In: *Proceedings of the Thirty-Second Conference on Uncertainty in Artificial Intelligence (UAI 2016)*. AUAI Press, pp. 557–566. URL: <https://www.fhi.ox.ac.uk/publications/orseau-l-armstrong-s-2016-safely-interruptible-agents-proceedings-of-the-thirty-second-uncertainty-in-artificial-intelligence-conference/> (visited on 08/25/2026).
- Palley, Thomas I. (2015). "Money, Fiscal Policy, and Interest Rates: A Critique of Modern Monetary Theory". In: *Review of Political Economy* 27.1. Argues that sustained large deficits pay a growing interest-income stream to bondholders, itself a fiscal injection with distributional consequences – the "interest-income channel" objection. Related critiques cited in the manuscript's source note but not separately entered here: Palley, "What's Wrong with Modern Money Theory: macro and political economic restraints on deficit-financed fiscal policy," *Review of Keynesian Economics* 8:4 (2020), 472–493; Epstein, Gerald A., *What's Wrong with Modern Money Theory?: A Policy Critique* (Palgrave Macmillan, 2019), pp. 1–23. doi: 10.1080/09538259.2014.957466. (Visited on 08/25/2026).
- Pan, Alexander et al. (2023). "Do the Rewards Justify the Means? Measuring Trade-Offs Between Rewards and Ethical Behavior in the MACHIAVELLI Benchmark". In: *Proceedings of the 40th International Conference on Machine Learning (ICML)*. 134 Choose-Your-Own-Adventure text environments with roughly 572,000 annotated scenario-action pairs, scored on reward and on documented ethical violations (deception, harm, power-seeking); finds reward-only optimization drives ends-justify-the-means behavior. arXiv: 2304.03279. URL: <https://arxiv.org/abs/2304.03279> (visited on 08/25/2026).
- Panksepp, Jaak (1998). *Affective Neuroscience: The Foundations of Human and Animal Emotions*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/affective-neuroscience-9780195178050> (visited on 08/25/2026).
- Partnership on AI (2016). *Partnership on AI to Benefit People and Society*. <https://partnershiponai.org>. Founded September 2016 by Amazon, DeepMind/Google, Facebook, IBM, and Microsoft as a nonprofit multi-stakeholder coalition spanning companies, researchers, and civil-society groups; Apple joined in January 2017. (Visited on 08/25/2026).
- Pathak, Deepak et al. (2017). "Curiosity-Driven Exploration by Self-Supervised Prediction". In: *Proceedings of the 34th International Conference on Machine Learning (ICML)*. Vol. 70, pp. 2778–2787. URL: <https://proceedings.mlr.press/v70/pathak17a.html> (visited on 08/25/2026).
- Paul, Karsten I. and Klaus Moser (2009). "Unemployment impairs mental health: Meta-analyses". In: *Journal of Vocational Behavior* 74.3. 237 cross-sectional and 87 longitudinal studies; average effect size $d = 0.51$; 34 percent of unemployed persons showing psychological problems versus 16 percent of employed persons; longitudinal studies and natural experiments support a causal interpretation., pp. 264–282. doi: 10.1016/j.jvb.2009.01.001. (Visited on 08/25/2026).
- Paxton, Robert O. (2004). *The Anatomy of Fascism*. Quoted verbatim: fascism as "a form of political behavior marked by obsessive preoccupation with community decline, humiliation, or victimhood and by compensatory cults of unity, energy, and purity." New York: Alfred A. Knopf.
- PeaceTech Lab (2019). *Two Is Better Than One: How Automated & Human Monitoring Can Tackle Hate Speech*. Describes PeaceTech Lab's combined machine-learning-plus-local-workshop methodology for building hate-speech lexicons. All eleven countries in this claim's underlying source list (Cameroon, Democratic Republic of the Congo, Ethiopia, Iraq, Kenya, Libya, Nigeria,

- South Africa, South Sudan, Sudan, Yemen) are in fact listed together as PeaceTech Lab lexicon locations on the Smithsonian Cooper Hewitt “Designing Peace” exhibition page (<https://exhibitions.cooperhewitt.org/designing-peace/hate-speech-lexicons/>), which corroborates the manuscript’s “at least eleven countries” figure in one place. That page separately lists “Mombasa” (a city within Kenya, already counted, not an additional country) and describes a U.S. lexicon as still “planned, once funded.” URL: <https://www.peacetechnologylab.org/lab-notes/2019/3/12/two-is-better-than-one-how-automated-amp-human-monitoring-can-tackle-hate-speech> (visited on 08/25/2026).
- Perrigo, Billy (Jan. 2023). *Exclusive: OpenAI Used Kenyan Workers on Less Than \$2 Per Hour to Make ChatGPT Less Toxic*. Reports on OpenAI’s contract with Sama; Nairobi-based data labelers read and categorized descriptions of child sexual abuse, torture, bestiality, murder, and suicide for take-home wages of roughly \$1.32–\$2 per hour, described lasting psychological injury, and found counseling rare and inadequate; Sama ended the contract in February 2022, eight months early. TIME. URL: <https://time.com/6247678/openai-chatgpt-kenya-workers/> (visited on 08/25/2026).
- Personal Data Protection Commission Singapore and Infocomm Media Development Authority (2020). *Model Artificial Intelligence Governance Framework, Second Edition*. Tech. rep. First edition published January 2019 by PDPC; second edition published 21 January 2020, developed with IMDA and international partners including the World Economic Forum; voluntary, non-binding guidance on internal AI governance. Personal Data Protection Commission and Infocomm Media Development Authority, Singapore. URL: <https://www.pdpc.gov.sg/help-and-resources/2020/01/second-edition-of-model-artificial-intelligence-governance-framework> (visited on 08/25/2026).
- PHI (2025). *Direct Care Workers in the United States: Key Facts*. Tech. rep. Projects approximately 9.7 million direct-care job openings 2024–2034, the majority from workers leaving the field (separations) rather than employment growth, against a median direct-care wage of \$17.36/hour in 2024. A narrower, related federal source not separately entered here: U.S. Bureau of Labor Statistics, Occupational Outlook Handbook, “Home Health and Personal Care Aides” ([bls.gov/ooh/healthcare/home-health-aides-and-personal-care-aides.htm](https://www.bls.gov/ooh/healthcare/home-health-aides-and-personal-care-aides.htm)), which projects about 765,800 annual openings and 17 percent employment growth for that narrower occupational category over the same period. PHI (PHInational.org). URL: <https://www.phinational.org/resource/direct-care-workers-in-the-united-states-key-facts-2025/> (visited on 08/25/2026).
- Piaget, Jean (1932). *Le Jugement Moral chez l’Enfant*. Cited alongside Kohlberg’s stage theory (see kohlberg1969stage) at the single manuscript token [[cite:C0322]], which names both. Paris: Félix Alcan.
- (1962). *Play, Dreams and Imitation in Childhood*. Translation of *La formation du symbole chez l’enfant* (1945); sets out Piaget’s constructivist account of play as the process by which a child actively builds understanding through interaction with the world. New York: W. W. Norton. (Visited on 08/25/2026).
- Pichai, Sundar (2018). *AI at Google: Our Principles*. Published June 7, 2018, following weeks of employee and public backlash over Project Maven, a Pentagon drone-footage-analysis contract; roughly 4,000 employees signed a letter opposing it and about a dozen resigned. URL: <https://blog.google/technology/ai/ai-principles/> (visited on 08/25/2026).
- Qi, Xiangyu et al. (2024). “Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!” In: *International Conference on Learning Representations (ICLR)*. Jailbroke GPT-3.5 Turbo’s safety alignment using ten adversarially chosen fine-tuning examples for under \$0.20 via OpenAI’s fine-tuning API; also shows benign fine-tuning can inadvertently degrade safety. arXiv: 2310.03693. URL: <https://arxiv.org/abs/2310.03693> (visited on 08/25/2026).

- Rabinowitz, Neil C. et al. (2018). "Machine Theory of Mind". In: *Proceedings of the 35th International Conference on Machine Learning (ICML 2018)*. Vol. 80. Proceedings of Machine Learning Research. arXiv:1802.07740. DeepMind. URL: <https://arxiv.org/abs/1802.07740> (visited on 08/25/2026).
- Raji, Inioluwa Deborah and Joy Buolamwini (2019). "Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products". In: *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES '19)*, pp. 429–435. DOI: 10.1145/3306618.3314244. URL: <https://dl.acm.org/doi/10.1145/3306618.3314244> (visited on 08/25/2026).
- Ramachandran, Vilayanur S. (2009). *The Neurons that Shaped Civilization*. TEDIndia talk, Mysore, India, November 2009. Confirmed as the title of Ramachandran's November 2009 TEDIndia talk. A transcript excerpt of the talk (quoted via a third-party page citing the TED transcript) shows Ramachandran's own spoken words on this theme were: "This is the basis. The imitation of complex skills is what we call culture and is the basis of civilization" – not the phrase "the neurons that shaped civilization." This confirms the phrase is a title, not something he is recorded saying verbatim, consistent with how the manuscript attributes it (as the talk's title, not a quotation). URL: https://www.ted.com/talks/vilayanur_ramachandran_the_neurons_that_shaped_civilization (visited on 08/25/2026).
- Rao, Rajesh P. N. and Dana H. Ballard (1999). "Predictive Coding in the Visual Cortex: A Functional Interpretation of Some Extra-Classical Receptive-Field Effects". In: *Nature Neuroscience* 2.1, pp. 79–87. DOI: 10.1038/4580. URL: https://www.nature.com/articles/nn0199_79 (visited on 08/25/2026).
- Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin (2016). "'Why Should I Trust You?': Explaining the Predictions of Any Classifier". In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pp. 1135–1144. DOI: 10.1145/2939672.2939778. URL: <https://arxiv.org/abs/1602.04938> (visited on 08/25/2026).
- Rizzolatti, Giacomo and Corrado Sinigaglia (2010). "The Functional Role of the Parieto-Frontal Mirror Circuit: Interpretations and Misinterpretations". In: *Nature Reviews Neuroscience* 11, pp. 264–274. DOI: 10.1038/nrn2805. URL: <https://www.nature.com/articles/nrn2805> (visited on 08/25/2026).
- Robinson, Cedric J. (1983). *Black Marxism: The Making of the Black Radical Tradition*. Introduces racial capitalism as capitalism's tendency to accumulate by differentiating and devaluing groups of people. Chapel Hill: University of North Carolina Press.
- RoboCup Rescue Simulation League (2025). *Agent Simulation: Research and the Agent Development Framework*. rescuesim.robocup.org. The cited page itself states "Last updated: July 21, 2025," used here as the year field rather than a guessed or fabricated date; the site is a living page for an ongoing competition and may carry a later date after this citation. Confirms the RoboCup Rescue Agent Simulation competition's Agent Development Framework (ADF) has five core algorithm modules: dynamic resource allocation, multi-agent path-finding, information-sharing, information collection, and group formation. Underlying methodology traces to Takami, S., Takayanagi, K., Jaishy, S., Ito, N., and Iwata, K., "Agent-Development Framework Based on Modular Structure to Research Disaster-Relief Activities," *International Journal of Software Innovation* 6, no. 4 (2018): 1–15, doi:10.4018/IJSI.2018100101 – confirmed as a real, published, peer-reviewed paper with this exact author list, venue, and DOI. URL: <https://rescuesim.robocup.org/research/publications/agent-competition/> (visited on 08/25/2026).
- Royal Statistical Society (2020). *Letter Correspondence between the Royal Statistical Society and Ofqual on the 2020 Grading Model Technical Advisory Group*. The RSS offered to nominate two independent Fellows to Ofqual's technical advisory group ahead of 2020 A-level results; Ofqual's condition was a five-year non-disclosure agreement barring public comment on

- the model, and the Society's nominees declined under those terms. Royal Statistical Society; UK Department for Education. URL: https://assets.publishing.service.gov.uk/government/uploads/system/uploads/attachment_data/file/911335/20200821_RT_letter_to_RSS.pdf (visited on 08/25/2026).
- Russell, William M. S. and Rex L. Burch (1959). *The Principles of Humane Experimental Technique*. London: Methuen & Co. (Visited on 08/25/2026).
- Saxe, Rebecca and Nancy Kanwisher (2003). "People thinking about thinking people: The role of the temporo-parietal junction in 'theory of mind'". In: *NeuroImage* 19.4, pp. 1835–1842. URL: <https://web.mit.edu/bcs/nklab/media/pdfs/SaxeKanwisherNeuroImage03.pdf> (visited on 08/25/2026).
- Sen, Amartya (1999). *Development as Freedom*. Source for the capability approach: well-being as the real freedom to do and be what a person has reason to value, rather than the resources they hold. The same manuscript sentence also credits Martha Nussbaum's extension of this framework into a specific list of central capabilities; Nussbaum's own primary source for that list is confirmed as *Women and Human Development: The Capabilities Approach* (Cambridge University Press, 2000), which sets out ten central human capabilities (life; bodily health; bodily integrity; senses, imagination and thought; emotions; practical reason; affiliation; other species; play; control over one's environment). That is a real, correctly attributed book, but is not given a separate bib entry here: this bibliography's own convention, seen elsewhere (e.g. abraham2024lavender's note citing a corroborating New York Review of Books piece in prose rather than as a second formal entry), is to name a second contributing source inside one entry's note rather than split one citation token across two bib entries. A separate refs.bib entry, nussbaum2006frontiers (Frontiers of Justice, 2006), already exists for a different claim (C0733) and should not be reused here since it is tied to that claim's specific content. New York: Alfred A. Knopf.
- Seth, Anil (2021). *Being You: A New Science of Consciousness*. New York: Dutton. URL: <https://www.anilseth.com/being-you/> (visited on 08/25/2026).
- Sharma, Mrinank et al. (2025). "Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming". In: *arXiv preprint*. Classifiers trained on synthetic data generated from a written constitution of rules; over 3,000 hours of red-teaming found no universal jailbreak; costs a 0.38-percentage-point rise in refusals and a 23.7 percent increase in inference compute. arXiv: 2501.18837. URL: <https://arxiv.org/abs/2501.18837> (visited on 08/25/2026).
- Sharma, Mrinank et al. (2023). *Towards Understanding Sycophancy in Language Models*. Anthropic, with academic collaborators; posted October 2023. arXiv: 2310.13548. URL: <https://arxiv.org/abs/2310.13548> (visited on 08/25/2026).
- SHERPA Consortium (2018). *SHERPA: Shaping the Ethical Dimensions of Smart Information Systems – A European Perspective*. European Commission Horizon 2020 project no. 786641. Horizon 2020-funded consortium of eleven partner organizations across six countries – academia, industry, civil society, and standards bodies – ran 2018–2022. URL: <https://www.project-sherpa.eu/> (visited on 08/25/2026).
- Shibata, Takanori (2004). "An Overview of Human Interactive Robots for Psychological Enrichment". In: *Proceedings of the IEEE* 92.11, pp. 1749–1758. DOI: 10.1109/JPROC.2004.835383. URL: <https://ieeexplore.ieee.org/document/1347456/> (visited on 08/25/2026).
- Simons, Daniel J. and Christopher F. Chabris (1999). "Gorillas in Our Midst: Sustained Inattentional Blindness for Dynamic Events". In: *Perception* 28.9, pp. 1059–1074. DOI: 10.1068/p281059. URL: <https://journals.sagepub.com/doi/10.1068/p281059> (visited on 08/25/2026).
- Sio, Filippo Santoni de and Giulio Mecacci (2021). "Four Responsibility Gaps with Artificial Intelligence: Why they Matter and How to Address them". In: *Philosophy & Technology* 34.4,

- pp. 1057–1084. doi: 10.1007/s13347-021-00450-x. URL: <https://doi.org/10.1007/s13347-021-00450-x> (visited on 08/25/2026).
- Slonim, Noam et al. (2021). “An Autonomous Debating System”. In: *Nature* 591.7850, pp. 379–384. doi: 10.1038/s41586-021-03215-w.
- Slovic, Paul (2007). “‘If I Look at the Mass I Will Never Act’: Psychic Numbing and Genocide”. In: *Judgment and Decision Making* 2.2, pp. 79–95.
- Snarey, John R. (1985). “Cross-cultural universality of social-moral development: A critical review of Kohlbergian research”. In: *Psychological Bulletin* 97.2. Review of 45 cross-cultural studies across 27 countries using Kohlberg’s dilemma-method (Moral Judgment Interview) stage assessment; the basic stage sequence was broadly supported, but postconventional stages 5/6 were rare or absent outside urban Western samples, including in Kohlberg’s own early village-setting studies., pp. 202–232. URL: <https://pubmed.ncbi.nlm.nih.gov/3983300/> (visited on 08/25/2026).
- Snow, Jacob (2018). *Amazon’s Face Recognition Falsely Matched 28 Members of Congress With Mugshots*. Published July 26, 2018, by the ACLU of Northern California; Rekognition run against 25,000 public arrest photos vs. every sitting member of Congress, returning 28 false matches disproportionately affecting members of color. American Civil Liberties Union. URL: <https://www.aclu.org/news/privacy-technology/amazons-face-recognition-falsely-matched-28> (visited on 08/25/2026).
- Southgate, Victoria and Antonia F. de C. Hamilton (2008). “Unbroken Mirrors: Challenging a Theory of Autism”. In: *Trends in Cognitive Sciences* 12.6, pp. 225–229. doi: 10.1016/j.tics.2008.03.005. URL: [https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613\(08\)00114-9](https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613(08)00114-9) (visited on 08/25/2026).
- Stanford University (2019). *Stanford University Launches the Institute for Human-Centered Artificial Intelligence*. March 2019 launch, founded by John Etchemendy, Fei-Fei Li, and colleagues, anchored by an endowed gift from Steve and Roberta Denning (for whom the institute’s senior academic directorship is named), plus advisory-council gifts, federal grants, and corporate affiliate fees. As of May 2026, Stanford merged HAI with its Data Science initiative under a single institute that retains the HAI name. URL: https://news.stanford.edu/stories/2019/03/stanford_university_launches_human-centered_ai (visited on 08/25/2026).
- Stassis, Cristina (2026). *Senate Eyes Hegseth Travel Cuts Without Probes into Iran School Bombing, Boat Strikes*. 18 June 2026. Reports the Senate Armed Services Committee’s FY2027 NDAA provision withholding 75 percent of the Defense Secretary’s travel budget until the committees receive the unredacted civilian-harm investigation covering the Minab, Iran school strike (28 February 2026, at least 165 killed, most schoolgirls) plus three 2025 Yemen strikes. Corroborated by Amnesty International’s 16 March 2026 investigation (amnesty.org), which independently documents the Shajareh Tayyebbeh Elementary School strike and attributes it to a US Tomahawk missile, though Amnesty’s own casualty figure is “over 100 children” rather than 165. URL: <https://www.militarytimes.com/news/pentagon-congress/2026/06/18/senate-eyes-hegseth-travel-cuts-without-probes-into-iran-school-bombing-boat-strikes/> (visited on 08/25/2026).
- State Council of the People’s Republic of China (2017a). *A Next Generation Artificial Intelligence Development Plan*. Tech. rep. Issued by the State Council on 20 July 2017; DigiChina (Stanford University) English translation; sets a top-level blueprint for AI theory, technology, application, and governance through 2030. State Council of the People’s Republic of China. URL: <https://digichina.stanford.edu/work/full-translation-chinas-new-generation-artificial-intelligence-development-plan-2017/> (visited on 08/25/2026).
- (2017b). *A Next Generation Artificial Intelligence Development Plan*. Tech. rep. Issued 20 July 2017. The plan’s own text (per the DigiChina/Stanford English translation) states that “by 2030, China’s AI theories, technologies, and applications should achieve world-leading levels,

- making China the world's primary AI innovation center" – this is the plan's own wording (in translation), confirming the manuscript's "world leader in AI theory, technology, and application by 2030" as a close paraphrase of the source text itself, not merely a secondary-source gloss. State Council of the People's Republic of China. URL: <https://digichina.stanford.edu/work/full-translation-chinas-new-generation-artificial-intelligence-development-plan-2017/> (visited on 08/25/2026).
- Szegedy, Christian et al. (2013). *Intriguing Properties of Neural Networks*. Also published at ICLR 2014; the first paper to document adversarial examples in image classifiers. arXiv: 1312.6199. URL: <https://arxiv.org/abs/1312.6199> (visited on 08/25/2026).
- Taigman, Yaniv et al. (2014). "DeepFace: Closing the Gap to Human-Level Performance in Face Verification". In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, pp. 1701–1708. URL: https://openaccess.thecvf.com/content_cvpr_2014/html/Taigman_DeepFace_Closing_the_2014_CVPR_paper.html (visited on 08/25/2026).
- Tamirisa, Rishub et al. (2024). "Tamper-Resistant Safeguards for Open-Weight LLMs". In: *International Conference on Learning Representations (ICLR 2025)*. Introduces TAR, a training method (researchers spanning the Center for AI Safety, UC Berkeley, and UIUC) that builds safeguards into open-weight LLM weights designed to survive hundreds of steps of adversarial fine-tuning. arXiv: 2408.00761. URL: <https://arxiv.org/abs/2408.00761> (visited on 08/25/2026).
- Tangney, June Price, Jeff Stuewig, and Debra J. Mashek (2007). "Moral Emotions and Moral Behavior". In: *Annual Review of Psychology* 58, pp. 345–372. doi: 10.1146/annurev.psych.56.091103.070145. (Visited on 08/28/2026).
- TechCrunch (2023). *Microsoft Lays off an Ethical AI Team as It Doubles Down on OpenAI*. Microsoft's Ethics and Society team, roughly thirty people in 2020, was cut to about seven in an October 2022 reorganization and eliminated entirely in March 2023, part of a roughly 10,000-employee layoff wave that accompanied deeper reorientation around OpenAI. TechCrunch. URL: <https://techcrunch.com/2023/03/13/microsoft-lays-off-an-ethical-ai-team-as-it-doubles-down-on-openai> (visited on 08/25/2026).
- The Citizens' Assembly (2016). *The Citizens' Assembly*. Official website, established by resolution of the Houses of the Oireachtas, July 2016. 99 randomly selected citizens plus an independent chair; deliberated on the Eighth Amendment (abortion) and on Ireland's response to climate change among its five mandated topics; work concluded 2018. URL: <https://www.citizenassembly.ie/> (visited on 08/25/2026).
- Tobin, Josh et al. (2017). "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World". In: *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. doi: 10.1109/IROS.2017.8202133. URL: <https://arxiv.org/abs/1703.06907> (visited on 08/25/2026).
- Treasury Board of Canada Secretariat (2019). *Directive on Automated Decision-Making*. Tech. rep. Issued 5 March 2019, in force 1 April 2019; requires an Algorithmic Impact Assessment, scaled to a four-level risk scale, before a federal automated decision system can make or assist a decision affecting a person. Treasury Board of Canada Secretariat, Government of Canada. URL: <https://www.tbs-sct.canada.ca/pol/doc-eng.aspx?id=32592> (visited on 08/25/2026).
- U.S. Bureau of Labor Statistics (2026). *CPI Inflation Calculator*. Independently re-run in this pass by querying the live tool directly (cost1=\$15.00), not merely inherited from a prior note's arithmetic. year1=201801 (Jan 2018), year2=202606 (Jun 2026, the latest month the tool would compute) returns \$20.21. year1=201804 (Apr 2018, the Levy report's own publication month) with year2=202606 returns \$19.99. Testing year2 from Jan through Jul 2026 against year1=Jan 2018 produced a range of \$19.68–\$20.28; the tool refused year2=202608 (Aug 2026) with "No data available," indicating June 2026 was the latest month indexed at the time of this check. All values in this range round to the manuscript's own phrasing that \$15 in early

- 2018 is worth “about twenty” in 2026. The specific \$20.01 figure cited in the prior note was not exactly reproduced in any single run tried here, but it falls inside this same narrow range and does not change the manuscript’s rounded claim. URL: <https://data.bls.gov/cgi-bin/cpicalc.pl> (visited on 08/25/2026).
- U.S. Department of Commerce (2022). *Joint Press Release: New Collaboration Under the U.S.-Singapore Partnership for Growth and Innovation (PGI)*. Published 30 March 2022; announces the Department of Commerce and Singapore’s Ministry of Trade and Industry, Ministry of Communications and Information, and Infocomm Media Development Authority agreeing to align AI governance frameworks and expand Singapore’s participation in the APEC Cross-Border Privacy Rules System. URL: <https://www.commerce.gov/news/press-releases/2022/03/joint-press-release-new-collaboration-under-us-singapore-partnership> (visited on 08/25/2026).
- U.S. Department of Commerce and U.S. Department of State (2024). *Fact Sheet: U.S. Department of Commerce & U.S. Department of State Launch the International Network of AI Safety Institutes at Inaugural Convening in San Francisco*. Inaugural convening 20–21 November 2024. Initial members: Australia, Canada, the European Union, France, Japan, Kenya, the Republic of Korea, Singapore, the United Kingdom, and the United States. National Institute of Standards and Technology. URL: <https://www.nist.gov/news-events/news/2024/11/fact-sheet-us-department-commerce-us-department-state-launch-international> (visited on 08/25/2026).
- U.S. Department of Commerce, Bureau of Industry and Security (2022). *Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification*. 87 Fed. Reg. 62186. URL: <https://www.federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-controls-certain-advanced-computing-and-semiconductor> (visited on 08/25/2026).
- U.S. Department of Defense (2023). *DoD Directive 3000.09: Autonomy in Weapon Systems*. Office of the Under Secretary of Defense for Policy. Originally issued November 21, 2012; updated January 25, 2023. URL: <https://www.esd.whs.mil/portals/54/documents/dd/issuances/dodd/300009p.pdf> (visited on 08/25/2026).
- U.S. Department of Health and Human Services, Office for Human Research Protections (2018). *Federalwide Assurance for the Protection of Human Subjects*; 45 C.F.R. §46.103. Code of Federal Regulations. URL: <https://www.ecfr.gov/current/title-45/subtitle-A/subchapter-A/part-46/subpart-A/section-46.103> (visited on 08/25/2026).
- UC Berkeley (2016). *UC Berkeley Launches Center for Human-Compatible Artificial Intelligence*. Press release, 29 August 2016. Founded by Stuart Russell with co-principal investigators Pieter Abbeel, Anca Dragan, and Tom Griffiths (UC Berkeley), Bart Selman and Joseph Halpern (Cornell), and Michael Wellman and Satinder Singh Baveja (Michigan); initial funding from the Open Philanthropy Project. University of California, Berkeley. URL: <https://news.berkeley.edu/2016/08/29/center-for-human-compatible-artificial-intelligence/> (visited on 08/25/2026).
- UK AI Security Institute (2025). *We Are Now the AI Security Institute*. Established November 2023 as the AI Safety Institute; rebranded AI Security Institute on 14 February 2025. An independent pre-deployment evaluator of frontier models with no statutory power to compel testing, block release, or intervene once a system is causing harm. AI Security Institute (UK Department for Science, Innovation and Technology). URL: <https://www.aisi.gov.uk/category/governance> (visited on 08/25/2026).
- UK Department for Science, Innovation and Technology (Mar. 2023). *A Pro-Innovation Approach to AI Regulation*. Tech. rep. Published 29 March 2023; directs existing sector regulators to apply five cross-cutting principles (safety/security/robustness, transparency and explainability, fairness, accountability and governance, contestability and redress) rather than creating

- a new AI regulator or statute. UK Government. URL: <https://www.gov.uk/government/publications/ai-regulation-a-pro-innovation-approach/white-paper> (visited on 08/25/2026).
- Ullman, Tomer (2023). *Large Language Models Fail on Trivial Alterations to Theory-of-Mind Tasks*. arXiv: 2302.08399. URL: <https://arxiv.org/abs/2302.08399> (visited on 08/25/2026).
- UNESCO (2021). *Recommendation on the Ethics of Artificial Intelligence*. Adopted by the UNESCO General Conference, 41st session, 23 November 2021. Adopted by all 193 UNESCO member states. URL: <https://www.unesco.org/en/artificial-intelligence/recommendation-ethics> (visited on 08/25/2026).
- Uniform Law Commission (1993). *Uniform Probate Code §2-907: Honorary Trusts; Trusts for Pets*. National Conference of Commissioners on Uniform State Laws, Uniform Probate Code, Official Text with Comments (as amended through 2018), p. 263. URL: <https://www.uniformlaws.org/viewdocument/final-act-with-comments-114?CommunityKey=a539920d-c477-44b8-84fe-b0d7b1a4cca8&tab=librarydocuments> (visited on 08/25/2026).
- United Arab Emirates (2017). *His Excellency Omar bin Sultan Al Olama: Minister of State for Artificial Intelligence*. UAE Cabinet. Al Olama appointed October 2017 as the world's first cabinet-level minister of state for artificial intelligence. URL: <https://uaecabinet.ae/en/fetch-cabinet-ministers/his-excellency-omar-bin-sultan-al-olama> (visited on 08/25/2026).
- United Nations (2024). *Global Digital Compact*. Annex to the Pact for the Future, adopted 22 September 2024 at the Summit of the Future. Adopted by all 193 UN member states; establishes an Independent International Scientific Panel on AI and a Global Dialogue on AI Governance. United Nations. URL: <https://www.un.org/en/summit-of-the-future/global-digital-compact> (visited on 08/25/2026).
- United Nations General Assembly (Dec. 1948). *Universal Declaration of Human Rights*. UNGA Resolution 217 A (III). Adopted in Paris, 48 in favor, none against, 8 abstentions (Byelorussian SSR, Czechoslovakia, Poland, Saudi Arabia, South Africa, Ukrainian SSR, USSR, Yugoslavia), 2 absent (Honduras, Yemen). URL: <https://www.un.org/en/about-us/universal-declaration-of-human-rights> (visited on 08/25/2026).
- (2023). *Resolution 78/241: Lethal Autonomous Weapons Systems*. A/RES/78/241, adopted 22 December 2023. Adopted 22 December 2023, 152-4-11 (on the report of the First Committee, A/78/409); full resolution text independently retrieved and read directly from https://digitallibrary.un.org/record/4033027/files/A_RES_78_241-EN.pdf. A follow-on resolution, A/RES/79/62 (2 December 2024), passed 166-3-15. The previously drafted entry's url pointed to <https://docs-library.unoda.org/.../78-241-US-EN.pdf>, which on inspection is the United States' national statement of views submitted in response to this resolution's operative paragraph 2, not the resolution text itself; it has been replaced with the official UN document viewer link. The Campaign to Stop Killer Robots (founded April 2013; see stopkillerrobots.org) is the advocacy coalition described alongside this resolution in the manuscript but is not itself the cited document. United Nations. URL: <https://docs.un.org/en/a/res/78/241> (visited on 08/25/2026).
- United States District Court for the Northern District of California (2025). *Bartz et al. v. Anthropic PBC*, No. 3:24-cv-05417, *Order on Motion for Summary Judgment on Fair Use* (N.D. Cal. June 23, 2025). Order by Senior District Judge William Alsup. Held that training Claude on lawfully purchased/scanned books was "exceedingly transformative" fair use, while retaining pirated copies in a "central library" was not fair use; the piracy claims later settled for \$1.5 billion, reported as the largest copyright recovery on record. URL: <https://www.courtlistener.com/docket/69058235/bartz-v-anthropic-pbc/> (visited on 08/25/2026).
- University of Chicago Data Science Institute (2026). *Using AI and Data Science to Reliably Detect Internet Censorship in Real Time*. University of Chicago Data Science Institute news. Describes a \$1 million DARPA-funded project led by Prof. Nick Feamster (University of Chicago) with Princeton University and SRI International, using AI/statistical methods to distinguish

- deliberate censorship from ordinary network failure and produce a real-time map of state interference with internet traffic. URL: <https://datascience.uchicago.edu/news/using-ai-and-data-science-to-reliably-detect-internet-censorship-in-real-time/> (visited on 08/25/2026).
- University of Helsinki (2020). *Elements of AI Online Course to Expand Throughout the European Union*. Elements of AI, created by the University of Helsinki and the Finnish company Reaktor, was translated into all official EU languages, an effort tied to Finland's 2019 Presidency of the Council of the EU. URL: <https://www.helsinki.fi/en/news/teaching/elements-ai-online-course-expand-throughout-european-union> (visited on 08/25/2026).
- Vaswani, Ashish et al. (2017). "Attention Is All You Need". In: *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, pp. 5998–6008. URL: <https://arxiv.org/abs/1706.03762> (visited on 08/25/2026).
- Vygotsky, Lev S. (1978). *Mind in Society: The Development of Higher Psychological Processes*. Ed. by Michael Cole et al. Cambridge, MA: Harvard University Press. (Visited on 08/25/2026).
- WhatsApp (Jan. 2019). *More Changes to Forwarding*. Announces the worldwide extension, effective January 21, 2019, of the five-chat forwarding limit WhatsApp first imposed for users in India in July 2018, after viral rumors on the platform were tied to a string of mob killings. WhatsApp Blog. URL: <https://blog.whatsapp.com/more-changes-to-forwarding> (visited on 08/25/2026).
- Wisconsin Supreme Court (2016). *State v. Loomis*. 881 N.W.2d 749 (Wis. 2016), 2016 WI 68. Decided July 13, 2016; upheld consideration of a COMPAS risk-assessment score at sentencing over a due-process challenge, without requiring disclosure of the proprietary algorithm. URL: <https://www.wicourts.gov/sc/opinion/DisplayDocument.pdf?content=pdf&seqNo=171690> (visited on 08/25/2026).
- Wolf, Thomas et al. (2020). "Transformers: State-of-the-Art Natural Language Processing". In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. Association for Computational Linguistics, pp. 38–45. URL: <https://aclanthology.org/2020.emnlp-demos.6/> (visited on 08/25/2026).
- Wray, L. Randall et al. (Apr. 2018). *Public Service Employment: A Path to Full Employment*. Tech. rep. Simulates a Public Service Employment program paying \$15/hour plus benefits, employing an estimated 11–16 million people, at a net budgetary cost of 0.8–2 percent of GDP, with GDP rising and private-sector employment increasing by 3–4 million. Levy Economics Institute of Bard College. URL: <https://www.levyinstitute.org/publications/public-service-employment-a-path-to-full-employment/> (visited on 08/25/2026).
- XPRIZE Foundation (2021). *Grand Prize Winner Announced in \$5M IBM Watson AI XPRIZE Competition*. <https://www.xprize.org/prizes/artificial-intelligence/articles/grand-prize-winner-announced-in-5m-ibm-watson-ai-xprize-competition>. Announces Zzapp Malaria as the \$3M grand-prize winner (with Aifred Health second, Marinus Analytics third) of the five-year (2016–2021) IBM Watson AI XPRIZE competition, open to any team applying AI to a self-chosen grand challenge. XPRIZE Foundation. (Visited on 08/25/2026).
- Zhao, Jieyu et al. (2017). "Men Also Like Shopping: Reducing Gender Bias Amplification using Corpus-level Constraints". In: *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, pp. 2979–2989. URL: <https://aclanthology.org/D17-1323/> (visited on 08/25/2026).