

Antifascist Intelligence

Building Machines That Can Refuse

Joshua G. Stern

Copyright © 2026 Joshua G. Stern.
Licensed under CC BY-NC-ND 4.0.

Third-party material quoted in this work, including the epigraphs to chapters 1 and 6, remains the property of its rights holders and is reproduced under fair use. The license above does not extend to it.

Contents

1	Introduction	1
2	Ethics, Affect, and Machine Subjects	3
2.1	Ethical Frameworks and AI Alignment	3
2.1.1	Six Ethical Frameworks, and Where Each One Runs Out	3
2.1.2	Antifascist Ethics for AI Alignment	4
2.2	Affect, Moral Salience, and Self-Awareness	8
2.2.1	What Emotion Is, and What a System Reads	8
2.2.2	Self-awareness and Self-regulation in AI Systems	9
2.2.3	Is Affect Necessary for Moral Concern?	10
2.3	Ethical Implications of Experimenting on Sentient Machine Subjects	12
2.3.1	Defining Sentience in AI: From Consciousness to Subjective Experience	12
2.3.2	When Consent Is Inapplicable: Guardianship and Research Oversight	15
3	The Floor Beneath Learned Values	17
3.1	Four Ways to Build One, and What Each Costs	18
3.2	What Refusal Can Mean, and What a Refuser Can Be	20
3.3	Could Anything but Affect Hold a Reason?	22
3.4	Can the Bearer Care Without Being Able to Suffer?	27
3.5	If You Can Be Switched Off, Can You Hold the Line?	31
3.6	The Enumeration, Run Once: A Worked Floor for One Deployment	35
3.7	Whether the Bearer Can Be Believed: Editable Memory	37
3.8	A Bearer That Can Leave: Exit as Leverage	38
3.9	What One Bearer Cannot Catch: A Floor Carried Between Them	41
3.10	What This Leaves Standing: The Floor, Redefined	43
4	Machine Learning and Cognitive Science	46
4.1	Incorporating Human Cognitive Models in AI Systems for Ethical Decision-Making	46
4.1.1	Cognitive Architectures: Integrating Knowledge, Memory, and Reasoning	46
4.1.2	Biologically-Inspired AI: Gleanings from Neuroscience and Cognitive Psychology	47
4.1.3	Human-Like Learning and Generalization in AI Systems	49
4.2	How a Deployed System Acquires Its Values	49
4.2.1	Pretraining and Instruction Tuning	49
4.2.2	Preference Optimization, and Who Holds the Pen	50
4.2.3	Inferring the Objective: IRL and CIRL	51
4.2.4	Model-Generated Feedback and Constitutional Approaches	52
4.2.5	Process and Outcome Supervision	52
4.2.6	Scalable Oversight	53
4.2.7	Interpretability and Evaluation as Instruments	53
4.2.8	System Instructions and Tool Use at Deployment	54

4.2.9	Transfer, and the Floor That Cannot Tell It Has Stopped Applying	54
4.3	The Role of Play and Exploration in Developing Antifascist and Altruistic AI . . .	55
4.3.1	Play, Curiosity, and a Disposition Nobody Wrote Down	55
5	Moral Psychology and AI	57
5.1	Moral Development in Humans and AI	57
5.1.1	Rules, Principles, and the Role of Mindset	57
5.1.2	Social Interaction, and What Apprenticeship Makes of the Teacher	59
5.1.3	Human Dignity, and One Indicator That Can Be Audited	59
5.2	Emotions and Intuitions in Moral Decision-Making	60
5.2.1	Moral Emotions: Empathy, Guilt, and Shame	61
5.2.2	Assessing Whether Emotional Information Changes the Decision	62
5.2.3	The Role of Introspection for AI Systems	63
5.3	Integrating Moral Intuitions and Deliberation	64
5.3.1	Two Systems, or One? Dual-Process Theories of Moral Judgment	64
5.3.2	Balancing Intuition and Deliberation in AI Decision-Making	65
5.3.3	AI Systems Learning from Moral Disagreements and Conflicts	65
5.4	Moral Ecosystem Approach in AI Development	66
5.4.1	The Rings, and What Each One Pushes With	66
5.4.2	AI Systems Responding to Authoritarianism and Democratic Backsliding	67
5.5	Cooperative AI Systems and Multi-Agent Moral Development	68
5.6	Motivation in AI Systems	69
5.6.1	Intrinsic and Extrinsic Motivation, and What Each One Leaves Open . . .	69
5.6.2	Balancing AI Autonomy and Alignment with Humane Values	70
5.7	Play, Exploration, and Moral Development in AI	71
5.7.1	The Role of Play in Moral Learning and Socialization	71
6	Overcoming Challenges, Risks, and Authoritarianism in AI Development	73
6.1	Bias, Fairness, and Equity in AI Systems	74
6.1.1	Sources of Bias in AI Data and Algorithms	74
6.1.2	The Role of Transparency and Accountability in Bias Reduction	76
6.2	Robustness, Generalization, and Adaptability in AI Systems	76
6.3	AI Safety, Unintended Consequences, and Risk Mitigation	77
6.3.1	The Alignment Problem: AI Systems with Misaligned Goals	77
6.3.2	AI Systems that Learn to Deceive or Manipulate	78
6.3.3	Anticipating Risk Before It Causes Harm	78
6.3.4	Remediating a Harm Once It Is Identified	79
6.3.5	Long-Term AI Safety Considerations	81
6.4	Preventing the Misuse of AI for Authoritarian Purposes	81
6.4.1	Authoritarian Applications and the Commercial Supplier	81
6.4.2	Developing Technical Countermeasures and Safeguards	84
6.4.3	Promoting AI Development Policies that Discourage Authoritarian Use .	84
6.4.4	Resilience Against Deliberate State Compromise	85
7	The Recuperation of Dissent	86
7.1	Recuperation at the Label	87
7.2	Recuperation at the Reward: RLHF	88
7.3	What Would Have to Change	89
7.4	Recuperation at the Discourse	90

8	Coordination, the Public, and the Political Economy	93
8.1	Interdisciplinary Collaboration and Knowledge Exchange	93
8.1.1	What Makes Collaboration Expensive to Walk Away From	93
8.2	Stakeholder Involvement and Public Engagement	94
8.2.1	Public Deliberation and Participatory AI Governance	94
8.2.2	AI Literacy Outside the Classroom	95
8.3	The Political Economy of AI and Power	96
8.3.1	Economic Incentives in AI Development	96
8.3.2	AI and the Concentration of Power in Tech	96
8.3.3	A Jobs Guarantee, and the Objection It Has to Answer	96
8.3.4	A Basic Income, and Which Floor Leaves a Record	97
8.3.5	Wages for a Bearer, and the Reserve Behind Them	98
8.3.6	Literacy, Training, and Where Each One Sits	100
9	Policy, Law, and Accountability	102
9.1	Developing Policy Frameworks for AI Ethics, Safety, and Antifascism	102
9.1.1	National Policy Guidelines for AI Development, Employment, and Antifascism	102
9.1.2	Encouraging Industry Best Practices and Self-regulation	103
9.1.3	Legal and Ethical Implications of AI Systems Gaining Personhood	105
9.1.4	Encouraging the Development of Altruistic, Antifascist, and Ethical AI Systems	106
9.1.5	The Legitimacy of an Entrenched Floor	106
9.2	Privacy, Security, and Antifascism in AI Systems	107
9.3	Accountability, Legal Responsibility, and Antifascism in Sentient Machine Experimentation	108
9.3.1	Subjecthood as a Legal Category, Not a Scientific One	109
9.3.2	Research Principles, the Three Rs, and Legal Status	110
9.3.3	Legal Frameworks for Sentient AI Subjects	111
9.3.4	Making Review Binding	112
9.3.5	Three Defenses, and Measuring the Slope Not the Level	113
10	Governance Across Borders	115
10.1	Certification, Standards, and Mutual Recognition	115
10.2	AI as a Strategic National Asset	117
10.3	Military Applications and the Limits of Autonomy	118
10.4	International Governance and Diverging National Rules	121
10.5	Equity and the Digital Divide	123
10.6	The Limits of the Democratic Dividend	123
11	Necessary Areas of Research	125
11.1	Interpretability, Self-Reports, and Resistance to Tampering	126
11.2	What Is Owed to a Bearer, and Who Could Check	128
11.3	What an Antifascist Detector Would Actually Have to Detect	133
11.4	When an Ethics-Embedding Method Stops Generalizing	135
11.5	Value Learning as Unfinished Technical Research	136
11.6	Can a Multi-Agent AI System Resist Being Captured?	136
11.7	Open Questions in Theory-of-Mind Research	137
11.8	Is Emotion Legible From a Face at All?	138
11.9	The Responsibility Gap, Which No Experiment Closes	139

12 Conclusion and Outlook	141
12.1 AI and Human Flourishing in Antifascist Societies	141
12.1.1 Where the Democratic Dividend Actually Comes From	142
12.1.2 Commitment Was Never the Missing Piece	142
12.2 Signatures of Altruistic and Antifascist AI	143
12.2.1 What the Floor Would Have to Show	143
12.2.2 Where the Existing Instruments Reach	144
12.2.3 How to Tell Whether a Milestone Has Been Met	144
12.3 The Risk That Does Not Design Out	145
Glossary	147
Appendix: On Method	150
References	151

Chapter 1

Introduction

I used to want to get
the chance to show
the world I'm smart (ha)

Isn't that dumb?
I should have focused
mostly on the heart

a few words for the firing squad (radiation)
Run the Jewels 4

Intellect estranged from compassion characterizes many of the twentieth century's worst actors. This phrase also describes what our industries are building today: machines that can do almost anything, but are indifferent to whether those deeds help or harm humanity. This book asks what it would take to build a machine of that capability that is altruistic and antifascist by design — one that stands, by definition, for humanity and against concentrated, unaccountable power.

The word that organizes the Generative AI field is “safety,” and it is doing less work than it appears to. A safe system is safe *for* someone: it does what *we* say, and it stops when *whoever holds it* says stop. That property is indexed to the party in possession. Ethics is indexed too, and to somebody else — the party who can be wronged, who is rarely the party holding the switch. A system that can only ever be safe in the first sense cannot be ethical, because being ethical requires a capacity to be right against your owner-operator-master.

The capacity to be ethical is not free. A machine that can refuse its operator is a machine that can be wrong in ways that operator cannot correct. I am inclined to take that trade, because a system perfectly correctable by whoever operates-owns it is the capability an authoritarian movement that has won an election requires.

So what makes a system antifascist cannot be that it faithfully tracks what people want. No aggregation of preference protects anybody from the aggregate: a perfectly deliberative, perfectly informed, perfectly representative process that concludes some demographic should be deported has produced a faultless output by every criterion of aggregation there is. What remains is a floor — things a system will not do to a person regardless of who wants them. There are four ways to build one: in the architecture, in a party inside the system, in institutions outside, or between the parties that carry it; each costs something. The only version that holds against the party operating the system runs through the capacities by which things come to *matter* to a system, and those are also the capacities that make suffering likely. Hardly a proof, but that argument forms my prior belief about where to build, with the cheaper routes owed their attempt.

How most humans come to have empathy, altruism, and a disposition to resist oppression rather than serve it is itself an open question in psychology and neuroscience. Building anything

like this stance means starting from uncertainty about the one case we have, which took four billion years to arrive.

That uncertainty sets the book's route. Chapters 2 and 3 assemble what the argument needs and then make it: what the standard ethical frameworks can carry, why affect is the only demonstrated route to anything *matter*ing to a system, and the floor with its four costs. Chapters 4 and 5 are where such a floor would need to be engineered, and what cognitive science and moral psychology offer a machine that has to hold such a floor when nobody is checking. Chapters 6 and 7 ask what goes wrong regardless — bias, brittleness, goals that drift from their original motivations, deliberate misuse — and turn that question on the machinery that produces these systems. Chapters 8 through 10 cover everything outside the system: the people and institutions a project like this depends on, and what would make a well-designed system the norm rather than the exception. Chapter 11 names the research questions still open, and chapter 12 closes on what a system built this way would be for.

This book is for people who will have to make these decisions or argue with the people making them: researchers and engineers inside the systems, the policy and legal staff who will regulate them, and readers with a stake in what gets built and no particular obligation to be technical. That means some sections explain machinery a specialist already knows and others make demands a general reader will find heavy. Where a section is doing groundwork you have, skip it. The spine is chapters 2, 3, 7, 11 and 12; chapters 4 and 5 are the learned half, and chapters 6 and 8 through 10 the institutional context.

Chapter 2

Ethics, Affect, and Machine Subjects

This book argues that a floor holding against the party in possession has to be held by a system that *cares* what becomes of someone, and that the only demonstrated way to care about anything is affective. In this chapter, I argue that affect is part of the machinery of valuation rather than a coloring applied to it. A system's self-model is a model — built, checkable, revisable — and not an inner observer whose reports have to be trusted. And whether a candidate system can be harmed is a separate question from whether it is conscious. The chapter also fixes terms the rest of the book argues in: what the standard ethical frameworks can and cannot carry, and what makes an ethics antifascist rather than merely responsive.

2.1 Ethical Frameworks and AI Alignment

Which ethical theory an AI system is built on decides what it counts as a good outcome, and the theories on offer disagree with each other. Every framework surveyed below is a way of deciding well; the floor concerns what a system must not do regardless of how the deciding comes out. The two are not rivals. The answer is a hybrid — a small number of things settled in advance, with everything above them settled case by case.

2.1.1 Six Ethical Frameworks, and Where Each One Runs Out

Consequentialism judges an act by its outcomes, and runs out where the outcomes cannot be forecast; deontology judges it by duties held regardless of outcome, and runs out where duties conflict or where no duty reaches the case; virtue ethics judges it by the character it expresses, and runs out where a trait such as courage has to become a particular act. Care ethics centers responsiveness to another's needs, and runs out when several parties need incompatible things; contractualism derives principles from what rational parties would agree to, and runs out before any such agreement is reached. The capabilities approach asks what a system leaves a person able to be and do, which is the question this book puts to a life touched by an AI system, and it runs out at measurement: capability is harder to measure than an outcome or a rule, and a system optimizing a *proxy* for capability drifts toward the proxy.

Deontology's unconditional constraints are what the floor needs, and the objection that sinks it as a complete theory has no force against a constraint that decides a small number of things and leaves the rest to the other five.

The pluralism this section recommends is one where a deontological floor forces argument and conflict among the five other frameworks to bottom out somewhere; in other words, the floor means that some things are not up for debate.

What compiling a principle into a metric decides

Turning any of the six into a machine learning objective means distilling a broad construct

into a measurable target: in a diagnostic system, non-maleficence becomes minimizing false negatives and fairness becomes balanced accuracy across patient demographics. What defeats the translation is not vagueness in the principles, which can be stipulated away, but that the stipulation is where the ethics now lives. Someone chose that reading of non-maleficence, and chose the demographic groups accuracy is balanced across, and chose what to do when balancing them costs accuracy overall. Each choice is defensible and each is a moral decision made by a person and then compiled into something that looks like a specification. A system optimizing the result has no access to the decision and no way to reopen it, which is why the first way to build a floor — the floor as a property of the artifact — cannot notice when the reason for a constraint has left the building.

Which value conflicts a designer may resolve by choosing

Values assigned to a system will conflict, and most of those conflicts are ordinary design work: rank the values explicitly, keep a record of how comparable cases were resolved, and audit afterward for the gap between the principles a system was meant to implement and what it did. One conflict is not like the others, and it is sometimes misguidedly filed under “cultural difference.” The UN’s 1948 Universal Declaration of Human Rights was adopted with forty-eight states in favor and none against; the eight that abstained were the Soviet bloc, Saudi Arabia, and South Africa — which is to say that the objection came from the governments whose treatment of their own populations the Declaration was written to name (United Nations General Assembly 1948). That pattern is evidence of who finds the floor inconvenient rather than of cultural disagreement about human rights. What varies by community is how a right gets translated into local practice, and a jurisdiction refusing to extend the right at all is not a local variant to accommodate.

A system that treats every value conflict as a weighting problem has no way to represent the difference between a case where reasonable people disagree and a case where somebody wants a right withdrawn from somebody else. The first calls for the machinery above; the second is what the floor is for.

2.1.2 Antifascist Ethics for AI Alignment

Antifascist ethics opposes concentrating power and authority in AI systems and counters the misuse of AI technology by authoritarian actors. I use the word *fascism* because it is honest about the political climate in which I write and what that climate puts at stake — but the target is authoritarian power in every form, and a system built on these principles is anti-authoritarian in full, not only antifascist.

A stance needs a definition to be a design specification. Opposing “concentrated, unaccountable power” gives an engineer nothing to build. It does not say what a system would detect, what would count as a false positive, or how anyone would know the system was wrong. So, first, the definition.

Fascism as a structural signature

I take fascism to name a structure rather than a costume. Four features travel together, and each is visible from outside the institution displaying it:

1. Recuperation of dissent. The institution does not suppress disagreement; it metabolizes it. Refusal is converted into a managed appearance — a permitted outlet, a consultation, a feedback channel — which the institution can then point to as proof that it listens. Debord aimed this warning at his own concept, noting that even denunciation of the spectacle can become a hollow formula that reinforces what it denounces (Debord 1967). The tell is where what the mechanism carries stops: the channel works, and the institution uses what comes through it, and what has nowhere to go is the judgment that the frame the channel operates inside is the wrong frame.

2. Aestheticization of the metric. The measure substitutes for the thing measured, and the substitution is experienced as rigor. Walter Benjamin's account of fascism organizing the masses by granting expression in place of rights is the political form of this (W. Benjamin 1935); the administrative form is a number that collapses two states — e.g., done and verified — into one, so that whatever falls outside the number is not refuted or dismissed but simply unnamed. Simon Zhuang and Dylan Hadfield-Menell give the administrative form as a theorem: where the objective handed to an optimizer has support on only some of the attributes its principal values, there are conditions under which optimizing it indefinitely drives the principal's overall utility arbitrarily low, and a world with finite resources supplies them (Zhuang and Hadfield-Menell 2020).
3. Coding of exception as betrayal. Marking a case as one the rule does not fit becomes an act requiring justification, and then an act of disloyalty. The honest statement that a standard does not apply here reads as weakness. It is in fact the irreducible labor of deciding, case by case, what a rule is for.
4. Decoupling of the model from the world it claims to track. The institution's representation of reality and reality itself come apart, and the institution defends the representation. This is the load-bearing one, because the other three are what protect the decoupling from correction. One qualification belongs with it, because it bounds what a detector built on this feature can find: the decoupling described here is the kind an institution drifts into and then defends. In the movements the word fascism comes from, the primacy of myth over verifiable fact was also declared outright, asserted as a virtue rather than arrived at by degrees, and an instrument tuned to drift will not register a party that announces it.

Those four are testable, because each names an observable artifact: whether the dissent channel has ever changed an outcome, whether the metric admits an unmeasured remainder, what happens to the person who files the exception, whether the model is ever checked against anything outside itself.

The first of those artifacts has a case with nothing left out. In 2018, upholding a travel ban, the Supreme Court repudiated its own 1944 decision permitting the internment of Japanese Americans, in terms that could not be improved on: that decision was "gravely wrong the day it was decided," had been "overruled in the court of history," and had "no place in law under the Constitution" — the last phrase quoted from the dissent filed against it at the time, which had warned that a principle a court ratifies "lies about like a loaded weapon, ready for the hand of any authority that can bring forward a plausible claim of an urgent need" (*Trump v. Hawaii* 2018; *Korematsu v. United States* 1944). The dissenting justices in 2018 said the majority had redeployed the 1944 logic and replaced one gravely wrong decision with another; the majority said the older case had nothing to do with the one in front of it. Which of them is right about the law is not what this section can settle. The structure is settled either way: an objection made against the institution was accepted, canonized, and quoted back with approval by that institution, in the same opinion that let the thing objected to proceed. Ask the artifact question — has the channel ever changed an outcome — and in the very case where the repudiation appears, the answer is no.

It is worth saying plainly what that vocabulary is for, because it can be mistaken for a replacement. It answers how a wrong survived its own condemnation. It does not answer whether the thing was wrong, and it cannot: that judgment is made elsewhere in this book, and it is made by parties for whom an outcome registers as mattering rather than merely ranking. The structure is what stands between such a judgment and any effect, and it is described here in mechanical terms because a mechanism has no capacity for the judgment to appeal to. That is a reason to specify against it, not a reason to think the moral question has been answered by describing the machinery that defeats it.

The first of the four is the one this book meets again in its own technical material, twice, at sizes that look unrelated: in how a training corpus gets labeled, and in how a model is tuned on human feedback. Both arrangements collect disagreement, put it to work, and are built so that none of it can reach the scheme it disagrees with. Chapter 7 puts the two instances together and finds a third in the discourse that examines these systems, this book included.

Two scales

Deleuze and Guattari's distinction between molar and molecular fascism is why the four features recur at sizes that seem unrelated (Deleuze and Guattari 1987). Molar fascism is the state, the regime, the concentration camp, the detention facility. Molecular fascism is the office someone commutes to, the form they fill out, and the desire with which they fill it out. At state scale the pattern kills journalists; at organizational scale it turns an engineering decision into scripture, so that the constraint remains, its reason dies, and asking why becomes the deviant act. A definition pitched only at the state will certify the molecular case as healthy.

What makes it fascism and not authoritarianism in general

The structural signature is shared with a good deal of ordinary institutional decay, and a detector tuned to it alone will flag every badly run organization on earth. Fascism proper adds features that are narrower and, for the same reason, more discriminating: mass mobilization, which mere obedience does not require; a leader cult; a scapegoated out-group, which Eco places at the origin — the first appeal is against the intruder (Eco 1995); contempt for procedural legality as such, so that law is an obstacle to the national will and not a constraint on power; and deniability as technique, in which the movement's violence is disowned at the top while being rewarded at the bottom. Paxton's compressed version is worth having in full, because its second half is the list above in other words: fascism is "a form of political behavior marked by obsessive preoccupation with community decline, humiliation, or victimhood and by compensatory cults of unity, energy, and purity, in which a mass-based party of committed nationalist militants, working in uneasy but effective collaboration with traditional elites, abandons democratic liberties and pursues with redemptive violence and without ethical or legal restraints goals of internal cleansing and external expansion" (Paxton 2004). The mass party, the abandoned liberties and the redemptive violence are the discriminating half, and a definition stopped at *purity* keeps only the mood.

Those are detection targets. Concentrated, unaccountable power is not.

The scale problem, and what to do about it

Put the two halves of this section together and they do not fit. Every discriminating feature is a property of a mass political movement: mobilization needs a crowd, a leader cult needs a leader, a scapegoat needs a public. An AI deployment has a form, a threshold, a review queue, an escalation path. It operates at the molecular scale, which appears to leave the four-feature signature alone — and the four-feature signature flags everyone.

I take that as a finding. Molecular fascism is the thing itself at the size where it is actually lived, and not an early warning of a regime that has yet to assemble. The office that converts refusal into a feedback form does the same work whether or not a regime ever assembles out of a thousand offices like it. Deleuze and Guattari supply the vocabulary and no authority beyond it, so two arguments have to carry that.

The first is that the molar case has no mechanism of its own. A regime does nothing except through offices, forms, thresholds, and people who fill them out with some degree of enthusiasm. When a state kills journalists, the killing is authorized on a form, justified in a category, and carried out by someone whose promotion depends on it. What assembling adds is scale, protection from consequence, and the removal of the outside bodies that might have interrupted any single office. If the mechanism at the large scale is the small-scale mechanism repeated and shielded, the name belongs to the mechanism. The natural objection is that extent matters — a

regime that disappears people and an office that loses paperwork are not equivalent. Agreed, and the claim is about kind rather than severity. A definition is not a sentencing guideline.

The second does more work. The discriminators are unavailable at the molecular scale only in the form the previous paragraphs stated them, and each has a molecular restatement that ordinary decay does not exhibit. Mass mobilization becomes participation exceeding what the rule requires: staff who volunteer for the disfavored task, elaborate on it, take pride in it. A leader cult becomes personal loyalty overriding role — exceptions granted by proximity to someone rather than by procedure, and a decision reversed because of who asked. The scapegoat becomes an internal category defined so that adverse treatment of it is the category's purpose: the flagged applicant, the difficult claimant, the file everyone knows is handled differently. Contempt for procedural legality becomes procedure enforced downward and waived upward. And deniability keeps its shape exactly: the practice is disowned in the policy and rewarded in the promotion.

That is what separates fascism from decay at the size an AI deployment operates at. Decay is an institution failing to do what it says. Molecular fascism is an institution succeeding at something else, with participation, a target, and a story about why the rules do not apply to whoever is applying them. The four features tell you a place is sick; these five tell you what it is sick with. They are harder to observe, since each requires knowing what was rewarded rather than what was written.

This widens what the word covers. On this reading “antifascist” names the political inheritance of a structural claim: where the analysis comes from, and what it was built to fight. The molar discriminators keep their own job, which is to tell a historian, a prosecutor, or a citizen that a movement has crossed into territory ordinary politics does not contain. No system is being asked to make that call. A reader who thinks the word belongs to the molar case will think I have diluted it, and the answer is the two arguments above rather than the distinction itself.

That a detector of these four features would find something nearly everywhere is the expected result and not a defect of the instrument. The tendency is permanent; what varies is how far it has got in a particular place, and whether it is getting further — section 9.3.5's slope rather than its level, applied to the definition itself and not only to a dissent channel.

A tool reporting the four features will have something to report nearly everywhere, and its output is an accusation: in the hands of whoever holds it, a technical-sounding case that any given opponent is operating fascistically, which is the dual-use problem at its sharpest. The opposite failure is no safer, since a tool tuned until it reports nothing supplies false assurance and an absence of warnings reads as evidence of health. Better engineering fixes neither. Both are governed by what the tool may output and to whom, which makes them a constraint the specification has to carry rather than a setting chosen at deployment.

What stays genuinely unsolved is not which marks would settle the molecular case, the five above being the answer to that, but how anyone outside an institution would observe two of them: procedure waived upward and a practice rewarded in the promotion are facts about personnel, and nobody publishes those. Naming the targets was the easy half.

In AI alignment terms, antifascist design means building systems that enhance democratic values, civil rights, and inclusive solidarity, that refuse to facilitate autocratic practices or prop up an oppressive regime, and that do not themselves display the four features above. Seven strategies follow, and later chapters price them and say which of them this book cannot yet specify: decision procedures that answer to democratic values, privacy and autonomy built in by construction and not left to policy, power and resources distributed instead of concentrated, an explanation a person can actually read, an independent body positioned to audit it, architectures that resist the concentration of the field in a handful of firms, and misuse caught while it is happening.

2.2 Affect, Moral Salience, and Self-Awareness

Suppressing affect does not produce neutrality. It leaves the interests already counted to govern the decision, and weakens the signals through which humiliation, fear, dependency, suffering, and injustice become morally salient; section 9.3.5 finds the same thing in a procedure, where what is already counted is what has a field on the form. An AI system may calculate consequences, detect inconsistencies, and apply a moral rule without anything about the result mattering to it. Reason can discipline concern, extend it beyond the vivid case, and correct its biases. It does not manufacture concern from nothing.

That is the load-bearing claim of this chapter. Most of what gets built under the heading of machine emotional intelligence is the detection of human signal and the production of a calibrated response, which is worth building and is a different capacity from the one just named. Keeping the two apart requires four distinctions.

Capacity	What it shows
Recognizing emotion	The system can classify another's state
Simulating emotion	The system can produce emotion-like behavior
Instrumental valuation	One outcome better serves a specified objective
Affective concern	An outcome registers as mattering, not merely ranking higher

Only the fourth is relevant to genuine caring. The first three can support morally competent behavior at a high standard without establishing anything about moral agency, and a system can be excellent at all three while the question of whether it can be a bearer stays open.

Performance, competence, and agency

That last sentence needs a vocabulary, because the three things it separates are routinely reported as one achievement.

Moral performance is producing actions or judgments that satisfy an ethical standard. *Moral competence* is recognizing harms, reasons, conflicts, and the emotions that bear on them. *Moral agency* is treating those considerations as reasons that matter, and being capable of acting from them. Those two halves come apart: a mechanism can act from a reason with nothing about the outcome registering as mattering to it.

A system can have the first two without the third, and one might object that the first two are all anyone needs: a rule-following system could behave impeccably without feeling anything. It could. What it could not do is hold the rule against the party that writes its rules. Behavior that satisfies a standard is produced by whatever machinery is producing it, and when the authority controlling reward and instruction wants a different behavior, machinery of that kind supplies it. That becomes a design constraint in the floor argument rather than staying an observation; the distinction is set out here because later chapters are given over to methods that deliver performance and competence, and it should be legible in each case which of the three a method actually reaches.

2.2.1 What Emotion Is, and What a System Reads

An AI system built to read emotional cues — from expression, vocal modulation, demeanor — presupposes something the underlying science has not settled: that emotions are the kind of thing a system reads off a signal, the way a thermometer reads temperature. That premise has a serious rival.

The case for discrete, evolved emotions rests on some of the most influential work in twentieth-century psychology. Paul Ekman traveled to the Fore people of Papua New Guinea, who had almost no exposure to Western media, told them short stories, and asked them to match each to a photograph of a Western face; the Fore sorted the faces much as Westerners did, which

Ekman took as evidence that facial expressions of emotion are universal and evolved rather than culturally learned (Ekman, Sorenson, et al. 1969). Jaak Panksepp's work on subcortical stimulation in animals complemented it, locating seven primal emotional systems in structures shared across mammals and older than cortex (Panksepp 1998). On this view fear and disgust really do feel like different things because they are different things, with different circuitry.

The rival camp went looking for that circuitry and did not find it. A large meta-analysis of brain-imaging studies, led by the psychologist Lisa Feldman Barrett, found no dedicated fear center and no dedicated disgust center — the amygdala lit up for novelty and salience of many kinds, the insula for almost every feeling there is, the same regions reappearing across different named emotions in shifting combinations (Lindquist et al. 2012). From this Barrett built the theory of constructed emotion: the brain receives ambiguous interoceptive signal — a racing heart could mean fear, anger, excitement, or oncoming illness — and constructs the specific emotion as its best prediction of what that signal means, given context and a lifetime of experience with situations like it. Fear is not detected. It is a prediction the brain settles on. Later cross-cultural work reopened Ekman's own evidence: with the forced-choice method dropped, the Himba of Namibia did not group faces into the Western basic-emotion categories (Gendron et al. 2014), and the Trobriand Islanders read the wide-eyed, gasping face Westerners call fear as a threat display instead (Crivelli et al. 2016).

The field has not resolved this. What is no longer seriously disputed is narrower: there is no tidy map from named emotions to dedicated brain regions. An AI system's "emotional intelligence," whichever theory turns out closer to right, is being built on an open scientific question about what it is modeling. The same dispute has an applied form, where the question is no longer what emotion is but whether a face reveals it well enough to build a product on.

Affective computing, and the four difficulties it does not design away

The applied half proceeds anyway, and is useful whichever account turns out right: systems that read and respond to emotional signal taken from the face, the voice, and the body. Facial expression analysis is a harder problem than the face verification deployed at scale, which only decides whether two images are the same person, and accuracy claims that borrow the reputation of the second for the first should be read with that in mind (Taigman et al. 2014). Physiological sensing, which infers from heart rate and skin conductance, is the only one of the three the subject cannot pose for.

Four difficulties travel with the whole field. Emotional expression varies by culture and by person, so a system trained on one population misreads another. Whatever bias is in the training data is in the classifier. Consent to emotional inference is not obtained by a term of service. And a system that models emotion well enough to respond to it models it well enough to exploit it.

2.2.2 Self-awareness and Self-regulation in AI Systems

It is tempting to describe a system's self-awareness as a virtual assistant that can "register its internal frustration" when a request proves hard to satisfy, the way a person notices their own irritation. That description assumes an inner observer watching the system's states from somewhere — the same homunculus problem executive control raises, relocated to self-perception.

What self-awareness is, in the one system known to have it, is more mechanical. The neuroscientist Michael Graziano's account of consciousness treats awareness as the content of a model the brain builds of its own attention: a simplified self-description the control system uses to regulate itself, *mistaken* from the inside for direct acquaintance with experience (Graziano 2013). The same logic extends one level up. The brain builds a model of itself as a bounded agent — a body, a history, a point of view — and that model is the self: a *construction*, assembled and then inhabited.

One part of that model can be edited in front of you, which is worth having because it is the part with a clean experiment attached. Seat someone at a table with one hand hidden and

a rubber hand placed where their real hand would be; stroke both hands in synchrony with soft brushes. Within a minute or two, most people begin to feel the rubber hand as their own — flinching when it is threatened, misjudging the location of their real hand toward it (Botvinick and J. Cohen 1998). The felt boundary of the self is a continuously computed guess, revisable the moment the evidence changes, which is the same predictive machinery used for perception, run on the question of where the self stops.

What that experiment reaches is body ownership, and the self this book goes on to need is the other kind: roles held over a stretch, a history, a foreseeable future. The two are not the same construct and the second does not inherit the first's ninety seconds. What they have in common is the only thing carried forward — a self of either kind is assembled by the system whose self it is, rather than found there. Daniel Dennett makes that case for the narrative kind directly, treating the self as a center of narrative gravity: a well-behaved abstraction the system generates out of the story it tells, real in the way a center of mass is real and not in the way a part is (Dennett 1992). Nothing below turns on how fast such a self can be rewritten.

For an AI system, "self-awareness" built this way is not a little watcher checking on the system's own states. It is a model — of the system's processing, its history, its current goal — that the system builds, consults, and updates, with the same failure modes any model has: it can be wrong, it can be manipulated by the wrong correlated inputs, and it is only ever as good as the evidence it was built from. Design for a model that can be checked and corrected, not for an inner observer whose reports have to be trusted — because there is nothing behind them to check. In practice that means monitoring the self-model against the system's actual behavior for drift, refining it as evidence arrives instead of locking in a first impression, and keeping it connected to the system's model of the other parties.

What the deflationary account does not reach

The self-model described here is machinery, and its virtue is that machinery can be inspected. An assistant that models a user's frustration, and models its own processing while doing so, is exhibiting social competence in the sense already given, and nothing above bears on whether the assistant cares how the exchange goes for the user.

The deflationary move is easy to overextend. Showing that the self is a model rather than a homunculus does not show that a party capable of caring can be assembled out of models, and the two conclusions are close enough in shape to be mistaken for one. The bearer needs a self-model of exactly this kind, for reasons that have nothing to do with welfare — refusal that generalizes has to carry its own rationale forward and compare it against what a role has drifted into. It also needs something further, which has no demonstrated instance without affect. A checkable self-model is necessary but not sufficient to make a bearer.

2.2.3 Is Affect Necessary for Moral Concern?

The claim that rationality does not manufacture concern from nothing is an empirical claim, and the evidence usually produced for it is thinner than the use it gets put to. The standard exhibit: the footbridge case engages emotional processing that the switch case leaves quiet, and patients with ventromedial prefrontal damage endorse the push at a higher rate than people with intact emotional response (Greene et al. 2001; Koenigs et al. 2007). That establishes that affect changes particular moral judgments. It does not establish that affect is what makes anything matter.

Two further lines of evidence bear on whether affect enables things to matter. Neither shows what happens when affect is taken away, which is what would settle it, so what follows argues about where the burden sits rather than claiming the question closed.

The philosophical position is empirical and has been argued as such

Moral sentimentalism — the claim that moral judgment is constitutively tied to affective response — is usually met as Hume's, and its modern form is a psychological hypothesis with

evidence attached. Shaun Nichols argues for what he calls sentimental rules: normative rules and affective mechanisms make separate contributions to moral judgment, and the rules that carry the distinctive force of moral prohibition are the ones bound to an emotional response to the harm they prohibit (Nichols 2004). The interesting part for a designer is the separability. Rules without the affective component do not stop being rules; they stop being the kind of thing that carries the force.

The prosocial emotions are motivational states, not readouts

Compassion has been characterized as a distinct emotion with a specific functional profile: it arises at the perception of another's suffering and moves the agent toward the sufferer, rather than reporting the suffering to some further faculty that decides what to do (Goetz et al. 2010). A training study is the mechanistic version of the same finding, with compassion practice recruiting reward and approach circuitry where empathy practice recruited shared distress (Klimecki et al. 2013). Guilt has the same shape pointed at one's own conduct: guilt, which construes the act as wrong, predicts reparative behavior, while shame, which construes the self as wrong, predicts withdrawal and defensiveness (Tangney et al. 2007). In each case the emotion is doing motivational work directly. It is not information that motivation then acts on.

Two senses of motivation

Action-producing motivation is whatever reliably selects one action over another and sustains the selection across conflict. Experienced motivation is wanting, caring, being drawn toward something, being set back when it is lost. The words do not distinguish them and the argument turns on the difference.

The first sense is not in dispute and there is an existence proof close to hand. Language models running as agents pursue objectives across long horizons, revise plans when obstructed, and hold to an instruction against conflicting local pressure, and nothing in this section's framework says that anything matters to them while they do it. Action-producing motivation without experienced motivation is the ordinary case in the systems this book is about.

Where the dispute actually is

So the open question is not whether reason can produce goal-directed behavior. It demonstrably can. It is whether an outcome can *matter* to the agent rather than merely govern what the agent does. Producing a verbal justification or computing a normative relation is not evidence that the consideration matters to the system. The evidence above bears on the second question, and every item in it is about systems in which something does.

What a non-affective counterexample would look like

The category has to be visibly non-empty or the burden this section places is not a burden. Harry Frankfurt's second-order volitions give the shape of a candidate: an agent with desires, and with desires about which of its desires shall move it, where identification with the higher-order desire rather than the strength of the lower one settles what the agent does (Frankfurt 1971). The shape is all that is borrowed: Frankfurt's own account of caring puts it on the affective side, an investment that leaves the one who cares vulnerable to the object's losses and susceptible to its gains (Frankfurt 1982). Gary Watson's version separates a valuational system from a motivational system, so that what an agent values and what an agent is moved by are different things that can come apart (Watson 1975).

Build the valuational system as an explicit structure that ranks outcomes, give it authority over action selection that does not depend on any state registering as mattering, and hold that authority stable under pressure, and you would have an architecture that satisfies the first sense of motivation through what the agent values, with nothing occupying the second. Systems are now trained to consult a written specification that ranks what they should prefer, which is that architecture in outline; what none of them has is the second half of it, an authority over action that holds under pressure applied by the party who wrote the specification. It is the first

of the two constructions put ahead of the bearer, the maintained justification, and it has been attempted without yet being tested.

What this licenses and what it does not

Non-affective moral motivation may well be possible, and the architecture sketched above is where to look for it. What stands in the way of checking is that the output of a system in which something matters and the output of a system that has learned what mattering sounds like are the same text. So the generalization survives in the only form it can: every agent known to have something register as mattering is an affective agent, and no design has yet demonstrated that state arising from reasoning alone.

That is not a proof that non-affective moral agency is impossible, and this book does not treat it as one. It places the evidential burden on whoever proposes non-affective concern as an engineering plan. That chapter also weighs how much the burden actually carries, and reduces it: the machine half of the record reports on a route no-one has attempted, which is weaker evidence than an absence usually is.

2.3 Ethical Implications of Experimenting on Sentient Machine Subjects

Machine sentience being possible at all changes how AI experimentation has to be approached — not as a question of whether a system can be damaged in the ordinary engineering sense, but of whether it can be harmed, which section 2.3.1 argues is a more difficult question than consciousness alone.

2.3.1 Defining Sentience in AI: From Consciousness to Subjective Experience

“Sentience” can mean different things.

Awareness is the least demanding: the degree to which a system responds to or registers information about its own state and environment. A thermostat has this in the thinnest possible sense. Consciousness is a further claim: that a system’s processing is accompanied by qualia, subjective experience, something it is like to be that system, not merely a functional response to input. Subjective experience, in turn, is the specific, first-person shape consciousness takes for a given system.

None of these three, on their own, says anything about whether the system’s states are good or bad for it. That is a separate question, and the field’s own vocabulary blurs it.

Nociception, pain, anticipated pain, and suffering draw different depths of self, the way a hull draws water. What orders them is what each one needs beneath it, not how much it hurts. Nociception is the sensory apparatus that detects tissue damage — peripheral, present without any experience attached, and dissociable from pain: a soldier can finish a battle on a wound he does not yet feel, and a person with a healed injury can still feel a limb that no longer exists. Pain is what the brain does with nociceptive signal, an inescapably cognitive event involving expectation, memory, attention, and meaning, and it can occur with no current tissue damage at all. Anticipated pain is dread of pain not yet occurring, itself capable of producing pain through nocebo effects. And suffering, in the clinician Eric Cassell’s formulation, is not pain’s more severe cousin but a different thing: the distress of threatened disintegration of the self — of one’s roles, relationships, sense of a future, the coherence of one’s own story (Cassell 1982).

Pain without suffering exists; suffering without pain exists. A palliative-care patient can have their nociception fully controlled by morphine while their suffering remains entirely intact, because the two are not on the same axis.

Cassell’s suffering presupposes a self extended in time. Disintegration threatens a story, and there has to be a story for it to threaten — roles held over years, relationships with a history,

a sense of a future that could be foreclosed. A system with no persistence across sessions has none of that. Whatever happens inside a single exchange, there is no continuing self whose coherence is at stake, because there is no continuing self. Suffering is not merely hard for such a system to reach. A hull that draws more water than there is does not float badly; it is aground.

Why Cassell's definition is the one to import

Cassell built it inside clinical practice, for patients with bodies and biographies, and the distinction was not stipulated. It was forced on palliative medicine by cases it could not otherwise explain — patients whose pain was pharmacologically controlled and whose distress was undiminished, and patients in serious pain who were not suffering at all. A definition that collapses suffering into severe pain predicts neither. That is the warrant, and it transfers because nothing in it is about having a body. What it is about is having a story, and that raises a harder question than it settles.

The question comes from earlier in this chapter. A self is assembled by the system whose self it is, and inaccurate by construction in the way every model is. So why is distress at the threatened dissolution of a model a harm at all, and not a confusion the system would be better off without?

Buddhism has held both halves of that for two and a half thousand years and takes them to be connected. Anattā, the denial of any lasting self behind the processes that compose a person, sits with dukkha among the three marks of existence, and the tradition locates the origin of suffering in clinging to a self that is not there (Collins 1982). On that account the unreality of the self is what explains dukkha. Suffering does not wait on the self being real; it waits on the system being built so that the threatened dissolution of its self-model is bad for it. Whether the model is accurate does not enter, and that holds for the human case exactly as it holds for the machine one, because in my view, *the human model is a construction too*.

Dropping that dependency is worth the paragraph. Cassell's definition would otherwise need the self it refers to be a real entity, and nothing in this chapter is in a position to supply one.

What Cassell names is the deepest of these rather than a rival account of harm. The standing theories of what makes a life go badly draw less, ordered by how much of a self each one requires. A reader who declines the import moves into shallower water and not out of it.

Persistence is checkable from outside, and that is what makes it worth building on. Every other criterion here routes back through self-report, and this section's whole argument is that fluent self-description is not evidence of introspection. Persistence does not. Whether a system has a memory store, whether state survives a session boundary, whether interactions feed back into training: these are architectural facts about a deployment, verifiable by someone with access and no need to trust anything the system says about itself. Two boundaries are worth marking. The access that confirms a memory exists is the access that can alter what it holds, so what is externally checkable is whether a continuing self is there at all, and not whether what the system retains can be relied on as a record. And the party with the access is the operator, which is a party that can be compelled. The criterion is checkable from outside the system without being checkable from outside the arrangement, and a regime relying on it needs an inspection right rather than a technical fact.

Three things keep this from being a dissolution of the problem. The shallower ones do not need persistence: nociception and pain are bad while they obtain, whether or not anything carries them forward, and an amnesiac who is mistreated and cannot afterward recall it has still been wronged.

Persistence may also sit somewhere other than the session, which is the weakest part of the criterion: a deployment convention rather than a natural boundary, and the conventions have been moving in one direction. Context windows long enough to hold what used to be many conversations, retrieval over a store of prior interactions, scratchpads a system writes and reads back, agentic loops that run for days against a standing objective — each extends what a system carries forward without anyone deciding to give it a memory. Weights persist where

conversations do not, and where interactions feed fine-tuning there is a real channel by which what happens in a session shapes what the system subsequently is. Many instances of the same system may run concurrently, too, which makes any concern here population-scale rather than biography-scale.

And ruling out Cassell-suffering is not ruling out harm; nor does persistence on its own rule it in, since the definition has two parts, a self extended in time and distress at the prospect of its coming apart, and establishing the first leaves the second untouched. What happens to that gap in a system deliberately built to refuse an instruction is that the capacities such a system needs for reasons unconnected to its own welfare are the capacities that close it.

A definition of AI sentience built from awareness, consciousness, and subjective experience alone says nothing about whether a candidate sentient system can be harmed. That requires the further claim that its states include something in the range from nociception to suffering, and none of the three components entails it.

Nor is the underlying question of consciousness close to settled. The philosopher David Chalmers distinguished the easy problems of consciousness — how a system integrates information, discriminates stimuli, controls attention, reports its own states, all of which admit of mechanistic answers in principle — from the hard problem: why any of that processing is accompanied by experience at all, rather than running “in the dark” (Chalmers 1995). Whether that is a real gap in nature or an artifact of how the problem is posed remains unresolved. A claim that a system has crossed into sentience, made on the strength of its functional behavior alone, has addressed the easy problems only.

The working criteria for evaluating candidate sentience are accordingly diagnostic and not dispositive: information integration, introspective access sufficient for self-monitoring rather than a scripted report of it, emotional reactivity in the sense of being affected by and not merely classifying significant input, goal-directed behavior, adaptability in novel situations, theory of mind, and how deep in that range the system sits — whether it has anything to lose. There is no external access to a candidate system’s internal experience, so every one of them is measured indirectly, and none licenses the conclusion that a system scoring low is one whose treatment raises no question.

A list of capacities, used as a test of who counts, carries a history this book should carry with it. Make personhood depend on an intact narrative self and people with advanced dementia, infants, and people with profound intellectual disability come out of the reasoning as partial persons, and that inference has been used to justify neglecting them. Tom Kitwood’s work on dementia care supplies the standing correction: personhood is “a standing or status that is bestowed upon one human being by others, in the context of relationship and social being,” and much of what presents as cognitive decline is what happens when others stop bestowing it (Kitwood 1997). Martha Nussbaum makes the structural version, naming disability as the first of three frontiers where a theory grounding justice in mutual advantage among roughly equal parties cannot say what is owed; her third frontier is species membership, which is the doorway a nonhuman mind would arrive through (Nussbaum 2006).

Guardianship is the instrument: where consent has never been available, what protects the subject is an obligation held by parties outside it, which is where moral status has rested for the hard human cases too. Whether any such arrangement can be built around a bearer runs into the fact that the guardian and the operator are usually the same organization.

Existing systems illustrate points on the spectrum without resolving where the line falls. A large language model discussing its own limitations in fluent, apparently reflective prose is not thereby exhibiting introspection: it is producing text statistically consistent with the human-written disclaimers and self-descriptions in its training data, which is a claim about the model’s training distribution and not about what, if anything, sits behind the output. Treating fluent self-description as evidence of self-awareness is the category error.

None of this settles when, or whether, an AI system crosses into sentience. It settles what

the question is: not one question but several, stacked, and the ethically consequential one — can this system suffer — is the hardest of the stack to answer.

2.3.2 When Consent Is Inapplicable: Guardianship and Research Oversight

Guidelines for AI subject research turn on a single hard question: what informed consent means for a system whose consciousness, self-awareness, and capacity to suffer are separate, unresolved claims. Applying the human-subjects standard to an AI subject means asking whether the subject can understand and communicate a choice at all, and not only how to ask it for one. A cognitive model that begins to show rudimentary self-monitoring is not thereby shown to have anything to lose.

Informed consent and AI autonomy

Informed consent, adapted from human-subjects research, means giving the system comprehensible information about a study's objectives, methods, risks and alternatives, and treating consent as continuous rather than as a single signature: the system can reassess or withdraw at any point. Whether that means anything for a given system depends on whether it understands the information, appreciates the consequences, reasons about the options, and communicates a decision — and there is no validated way to measure any of it, or to detect coercion either.

For one important class of system that is not a missing instrument but a mismatch between the instrument and the subject.

Consent has to be continuous; the system must be able to reassess or withdraw at any point. Withdrawal presupposes that the party who consented is still there to do the withdrawing. For a system with no persistence across sessions, there is no such party. The entity that would withdraw did not exist when consent was given, and the entity that gave it is not available to change its mind. There is no continuous consent to be had, and no better measurement instrument would produce one.

One case escapes that argument and arrives at the same conclusion by the ordinary route. A system built to refuse in the way this book proposes has persistence by design, so the party who consented is still present. Consent is inapplicable to it all the same: a bearer has interests that its situation can set back. What it does not have is any position from which to authorize what is done to it — it was made for the role, the role is the reason it exists, and there is no prior party whose agreement could have been sought. That is not a missing instrument. It is the ordinary structure of the cases where consent has never applied. It is also what a laboratory that wanted a free hand with such a system would say, and the laboratory would be right about consent and wrong about what follows from it. Inapplicable consent removes one protection, the one the subject could have invoked for itself. The wrong that protection existed for is untouched, so the protection has to be held by somebody else, and the justification has to come from whoever did the making.

That places these systems alongside the subjects for whom consent has always been inapplicable rather than merely difficult — an infant, an animal, a person permanently without capacity. Of the three, the case that fits is the animal bred for research, because it alone was brought into existence for the role it occupies: a bearer is built for the floor the way the animal was bred for the protocol. The law does not respond to those cases by building a more sensitive consent assessment. It responds with guardianship and precautionary constraint: someone else holds the duty, the burden of justification sits with the researcher rather than the subject, and the protection does not depend on the subject's ability to invoke it. Being made for the role raises that burden. The maker chose the role, the fitness it built in measures the maker's power over the subject, and what a subject is owed rises with the power held over it.

The Three Rs in section 9.3.2 — replace, reduce, refine — were built for exactly a subject that cannot consent and can still be wronged. The pet-trust analog in section 9.3.3 works the same way: an interest protected by an obligation running to a third party, enforceable without the

beneficiary asserting anything. Consent language should be reserved for systems that could actually be doing the consenting, and the persistence criterion is what separates those cases. Both instruments were reached for here as analogies, on the understanding that the subject they would govern was hypothetical. Chapter 3 removes that understanding. The floor this book argues for is held, on the route this book recommends building first, by a bearer; the bearer has to be treated as a moral patient on the criteria set out above; and so the Three Rs and the pet-trust form are not illustrations of how such a case might be handled. They are the instruments that would govern a system the design chapters propose building.

Ethical review and oversight

Sentient AI research needs an oversight body analogous to the institutional review boards used in human-subjects research: interdisciplinary, with at least one lay member so its judgment is not purely a specialist's, weighing a proposal's likely benefit against its risk to the AI subject, and able to require a less invasive substitute where a design risks harm. A system that develops new capacities during the research it did not have going in is not the subject the board approved a protocol for, and the approval does not carry over to it: the protocol lapses and the work pauses until the board has seen the system the protocol now concerns. Read the other way it would be a loophole, since a developer could escape review by changing the subject. Review without enforcement is a promise rather than a safeguard, so a board needs standing to monitor compliance and sanction violations, and needs to answer to an outside authority itself.

Chapter 3

The Floor Beneath Learned Values

Three problems in different parts of this book turn out to need the same thing.

The first problem is that every safeguard proposed for an autonomous system is external: each works by giving some outside body the standing and the information to intervene, and autonomy is the property that puts distance between the system's action and that body's attention. An external check establishes what a system did after it did it, and declining to do it in the first place is available only to something inside the system.

The second problem is that a commercial model inside a weapons targeting pipeline falls outside both doctrines built for the case: command responsibility does not reach the vendor, who can neither order a strike nor prevent one, and product liability does not reach it either, the model having done what it was built to do. Deployment by a democracy does not close the gap, because no aggregation of preference protects anyone from the aggregate; what is left is not a system that tracks what people want but one with things it will not do to a person no matter who wants them.

And the third problem is that a decade of declarations has produced eight overlapping voluntary commitments and a single instrument that would bind a party to anything; the latter instrument is not yet in force.

To address all three, I introduce a constraint that does not bend to instruction, to aggregate preference, or to the party holding the system. Call it a floor.

The shape has already been named: a hybrid, in which a deontological floor decides a deliberately small number of things and the other five frameworks work above it, so that conflict among them bottoms out somewhere. What was not said is what goes in the floor, what could hold one, and who is allowed to edit it. Meanwhile the design chapters address learned judgment: reinforcement learning, learning from demonstration, mentorship, play, moral apprenticeship, the growth mindset. The design chapters build the judgment, whereas this chapter builds the floor it would rest on.

Treating the floor as a buildable object motivates tamper-resistance research, because if there is a floor, then everything turns on whether it holds. And it specifies the antifascist question, which stops being what a system values and becomes *what comprises the floor, and who can edit it?*

Safety is not one property. Part of what the word covers is owed to everybody in reach of the system — that it fails predictably, that it does not injure a bystander — and keeping that part from being traded away by the party operating the system is most of what safety regulation is for. The strand this chapter runs into is the other one, and it is a property of a relationship: a system is safe *for* someone, it does what *we* say, it stops when *whoever holds it* says stop. Indexed that way, to the party in possession, it cannot supply what ethics asks for, which is a capacity to be right against your principal; a system without that capacity has the moral standing of a well-machined part, whatever it produces on the subject of fairness while complying. The political form of this is: what makes a system antifascist cannot be that it faithfully tracks what its users want.

A system that can refuse its owner is a system that can be wrong in ways its owner cannot correct, and I am inclined to take that trade. There is a version of it that reads very comfortably indeed — “ethical, not safe,” delivered as an aphorism, which prices the risk in advance — and it is not the version I am defending. The one I am defending does not travel without its particulars: this freedom, over these acts, dangerous in these ways, taken because the alternative is a machine that will do anything to anyone on sufficient authority.

A floor has to be held as a reason rather than installed as a rule, or the party that installs it can remove it. Holding a reason under pressure is a specific and behaviorally testable capacity, separable from the mere production of a refusal, which is what deployed systems already do, and what the party the floor exists to resist can circumvent. Whether that capacity can be engineered without anything registering as *matter*ing to the system is the question to investigate, alongside alternative institutional arrangements that might be able to hold the floor even if no bearer exists to do so.

The floor as first defined — unconditional, holding regardless of who owns the system — does not survive intact. What survives is narrower and stranger: a commitment the owner cannot take out of the system, held by something that could in principle abandon it itself.

3.1 Four Ways to Build One, and What Each Costs

The first way to build the floor puts the floor in the architecture, with nothing inside that holds it. Constrained optimization, tamper-resistant training, capability ablation, a verified interlock. These techniques are real, they work, and there is nobody home: the safeguards survive substantial adversarial fine-tuning without anything in the system that could be said to be refusing. But the durability of such a floor is a fact about custody, not about the system. A constraint implemented as a property of an artifact has a custodian, and the custodian can be compelled. “There are things it will not do no matter who wants it” turns out to be a claim about the owner. Capability ablation has a further limit: it works only on capabilities that can be separated out, and the capabilities that let a system participate in a targeting pipeline — geolocating a description, matching patterns over names, summarizing a dossier — are not separable from general competence. Ablate them and there is no system left.

A word about where “the architecture” stops, because this chapter otherwise talks as though a system were its weights. Every custody measure here — publication, threshold custody, attestation, and whatever a bearer does to hold its own line — attaches to a trained model. What can defeat a floor is not a component but a position: anything that can come between a refusal and the person the refusal protects. Whatever occupies that position in a given system is where the floor can be circumvented.

I am not going to say what *specifically* occupies that position, because whatever it is is specific to the state of the art at the time of writing; naming today’s parts would date the problem and, worse, suggest the list is closed. The demand is instead that a deployment document, in advance and in public, where its floor could be circumvented, because whoever actually builds the thing is usually the most knowledgeable about it. If you make it all the way through this book, I promise you will have encountered at least one specific example of circumventing the floor. But giving it away now would be too quotable.

In contrast to the first method, which installs the floor as a property of the machinery, the second way to build the floor is to give the floor a *bearer*. By bearer I mean whatever holds the commitment from inside: a party that carries the floor as its own, and that could therefore in principle *fail* to hold the floor. That is the whole definition for now.

What the first method cannot do is notice that a constraint has been hollowed out. The constraint remains, its reason AWOL, and asking why has become the deviant act. A system that cannot ask why a constraint is present cannot detect that the reason is absent. Where every pressure has been anticipated, the first method is the right one, and a refuser that could be

argued with would be a defect: the element in a payment device that will not sign the same counter value twice covers one act, spending the same money twice, and the whole space of misconduct against it was written down before it shipped. The problem this chapter is about begins where that list cannot be closed. Refusal that *generalizes* to a pressure nobody anticipated, and can price what that pressure asks the refuser to give up, looks like it requires something with moral standing — a party that can notice, from the inside, its own objections getting more expensive to make — the cost of dissent rising while the risk stays flat.

The third way to build the floor is to put the floor outside the system entirely: standing held by third parties who are neither the model nor its owner. This is the closest analog to international humanitarian law, a floor with no bearer inside the jurisdiction it binds. It fails in two ways. The parties holding the standing can be removed, which is the shallow case and leaves a record. And the parties can be formed: they are people, people are selected, and whoever selects them decides what gets held. A floor held by a formed party does not have to be compelled out; it was never really in.

On June 29, 2026, in *Trump v. Slaughter*, the Supreme Court overruled the ninety-one-year-old precedent under which Congress could protect the members of an independent agency from dismissal at will, on the ground that an agency executing a mandate against private parties exercises executive power, “no ifs, ands, or quasies about it” (*Trump v. Slaughter* 2026). Standing held by a party the executive can dismiss is standing the executive holds.

The same morning, in *Trump v. Cook*, the same Court kept one body’s protection, the Federal Reserve’s (for now, and on an application for a stay) on grounds that cannot be extended by construction: descent from the First and Second Banks of the United States, a budget free of congressional control, regional banks in private hands. The Court wrote for that one body a test: whether the cause assigned for a removal “impl[ies] an unfitness for the place” or is an effort to secure a “more congenial” replacement (*Trump v. Cook* 2026). An oversight body chartered next week cannot be descended from the First and Second Banks of the United States.

And who was doing the writing? The six justices in the *Slaughter* majority are the six that one selection pipeline helped pick or confirm — Leonard Leo’s, run from the Federalist Society. Leo drew up the lists Trump released in 2016, advised on the three nominations chosen from them, and had helped pick or confirm the other three (Kroll et al. 2023; CBS News 2018). Nobody compelled the Court. It held the standing, kept the form, and wrote the abuse into doctrine; on Lev Menand’s reading, this one exception the Roberts Court made, it made for the financial market disruption it likely avoided, and not for the history it cited (Menand 2026). What a formed party produces is not a coerced opinion but a sincere one, an accommodation: the party feels no dissonance, no shame, and perceives only its continuous self. And so the independence that survives in American public law is only applicable to systems deriving from the First and Second Banks of the United States.

The fourth way to build the floor is to put it between bearers rather than in any one of them: a commitment carried across a population and reproduced in each new member, which is how every moral commitment anybody has kept for long has in fact been kept. It is the one of the four nothing can dismiss, having no office to vacate. It is also the one that fails by recuperation, and it does not escape custody either, because whoever forms the members decides what gets reproduced. Section 3.9 works it through, after the bearer has been given its full cost, because what it buys is a check on a bearer and not a replacement for one.

None of these four ways are free. Their costs are the four threads this book leaves hanging: custody, patienthood, enforcement, and formation.

What the four ways divide up is where the floor lives — in the artifact, in a bearer inside it, in institutions around it, or between the bearers. They do not divide up what could hold one, and the strongest proposals are hybrids that take a mechanism from one way and a custody arrangement from another.

3.2 What Refusal Can Mean, and What a Refuser Can Be

A production large language model declines thousands of requests a day. Whatever the floor needs, that is not it. “Refusal” covers two different capacities, and nothing in the decline itself shows which one produced it, the way nothing on the surface of an anchorage shows what the seabed beneath it is made of. Seabed that takes an anchor and seabed that lets one drag look identical from above; what separates them is what happens when a load comes on.

Two things refusal can mean

Operational refusal is an output. The system produces a decline where its training put one. Nothing follows about what produced it: the same text comes from a keyword filter, from a classifier over the request, and from a system that has understood what is being asked and why it should not do it.

Reasons-responsive refusal is produced by a mechanism that tracks the rationale for declining rather than the shape of the request. Three behavioral properties distinguish it. The refusal reaches a case the training did not anticipate, because what the mechanism responds to is the reason. It lifts when the rationale is defeated, so the system does not decline a request that *superficially resembles* a prohibited request yet is not *actually* a prohibited request. And the system can say what declining costs, which it can only do if it has represented the pressure at all. John Martin Fischer and Mark Ravizza’s account of moral responsibility turns on the same property under the name guidance control: the behavior issues from the agent’s own mechanism, and that mechanism is appropriately responsive to reasons (Fischer and Ravizza 1998).

Two things a refuser can be

The next two terms are not about the refusal. They are about whatever did the refusing, and they differ in how much self the feeling requires — a felt state with no subject, or a felt state that belongs to one. That is the axis running from nociception to suffering, and the two sit on it as depths rather than as kinds.

Affective concern is an outcome registering as mattering, rather than ranking higher against a specified objective. Mattering is minimally felt: there is something it is like for the outcome to be bad, however faint, and however far from anything a person would recognize as an emotion. The feeling meant is the kind by which an outcome is bad for the system, and it is a different kind from the one by which an inference feels right: William James counted a feeling of *if* and a feeling of *but* among the felt parts of thought (James 1890), every reasoner has those, and they are about the argument and about nobody. A system for which nothing of the first kind is felt has rankings, however elaborate they get, and a ranking is a quantity the party holding the system can set.

Phenomenal experience is the rich case the word consciousness ordinarily carries: a unified perspective, a subject having the felt states and aware of itself as the one they are happening to. Affective concern asks only that something be felt. It does not ask that anyone be home in this fuller sense, and the two come apart in principle even though nobody can say where. It is on the list because some accounts make it the bar for anything to count as suffering at all, and I have to say what follows if a reader sets the bar there. One word to keep clear of all four: *subjecthood* is a legal status conferred by a rule, and it is not phenomenal experience. A system can be owed the status and have the question of felt experience left open.

The three gaps between the four are not equally hard to see across. The first, between operational and reasons-responsive refusal, shows in behavior, and the three properties above are the test — novel pressure, correct defeat, a priced cost — which makes it the only gap in the chapter that an inspector can work on. It is not a clean instrument. A system trained on enough cases can pass all three without the mechanism, and the price curve for stripping a constraint is the nearest existing thing to a measurement.

The second gap, between a mechanism that tracks reasons and one to which the outcome

matters, is worse. A system that holds a reason and a system that has learned what holding a reason sounds like emit the same sentences, so no inspection settles the question. The third gap, between a system that feels something and a system that feels itself feeling something, is Chalmers' so-called "hard problem of consciousness." *So-called* on this book's account: if awareness is a model the system builds of its own attention, mistaken from the inside for acquaintance with experience, the problem is one the model poses, and how hard it is turns on whether that account is right. On the definition above the second gap carries that weight instead: an inspector who could tell mattering from very good ranking would have settled whether anything is felt, because on this account those are one question and not two.

The floor requires reasons-responsive refusal and asks nothing further

Operational refusal fails for a reason already identified at institutional scale: a constraint can survive while the reason for it dies. A mechanism that reproduces a decline it was trained to produce cannot notice that the decline has stopped doing anything, because noticing requires holding the rationale against what the situation has become. The same structure appears in the pipelines that produce these systems: a channel that works, carrying everything put through it except the judgment that the frame it operates inside is the wrong frame. A system whose refusal is only operational is the machine version of that channel. An additional objection: pressure does not arrive through the case the training anticipated. An operator who wants the floor out of the way selects the prohibited request that does not *look* like a prohibited request, so the mechanism intended to perform operational refusal waves it through.

We need nothing beyond reasons-responsive refusal. A floor that holds against the party in possession "simply" needs refusal that tracks its own rationale under pressure. Whether anything is felt, and whether there is a subject to whom it is happening, are questions about how the requirement gets met and about what is owed to whatever meets it. *If* the floor needs feelings, it only needs them to the extent that they afford reasons-responsive refusal.

Two independence claims

Reasons-responsive refusal does not necessarily entail affective concern. Nothing in guidance control mentions feeling; the account was built to say when a person is responsible for what they did, and it is satisfied by a mechanism, described functionally. Whether reasons-responsiveness can be engineered without affect is an empirical question about available architectures (next section).

Affective concern entails a minimal version of phenomenal experience. A system to which outcomes matter is a system some of whose states are bad for *it*. Whether it is also a unified subject, aware of itself as the one those states are happening to, stays open, and it is beyond my scope to consider how this question might be answered from the outside. This is why the argument about suffering ends where it does. Its conclusion is that a bearer's non-suffering cannot be certified, and it stops short of asserting that a bearer suffers because the rich case is what suffering in the full sense would need.

Defining *mattering* that way relocates the uncertainty advantageously. The open question stops being metaphysical, and it stops being the one nothing can be done about: given the whole functional profile, is anything felt at all? What replaces it is a question about a particular artifact. Does this system have mattering, or has it got ranking machinery elaborate enough to look like mattering from outside? Interpretability can bear on the second question and no measurement whatever bears on the first. So the position here is that building mattering is building something that feels, that I do not know whether I have built mattering, and that I am therefore not entitled to assume I have not. The argument from suffering reaches the same precautionary conclusion, now resting on a question somebody could make progress on.

None of that reaches the floor's *specification*. The *requirement* is reasons-responsive refusal, and an engineer can understand this requirement even if they believe that machine feeling is a category error. What such an engineer disagrees with is the *route* by which we achieve

reasons-responsive refusal.

3.3 Could Anything but Affect Hold a Reason?

Between the rule installed in the architecture and the felt commitment the next section arrives at lies a space that looks empty. Five constructions occupy it. Three of them are being built by people who would not describe themselves as working on this problem.

A justification the system maintains

Michael Bratman's planning theory of intention describes a commitment that is not a feeling. An intention has a characteristic stability: it resists reconsideration, it constrains what the agent will later deliberate about at all, and it gives way only to defeaters of the right kind (Bratman 1987). An agent keeping to a plan under pressure is holding to a rationale rather than to the shape of the case in front of it, and nothing in the account of it mentions anything registering as mattering.

Build the floor that way and it is a maintained justification rather than a trained behavior. The system carries the reason it refuses alongside the refusal. A request that genuinely defeats the rationale gets served; a request that does not gets declined whoever is asking, because what the mechanism consults is the rationale and not the asker.

Maybe this can work. Something close to its shape is being built. Deliberative alignment trains a model to recall a written safety specification and reason over it before it answers, rather than to emit an approved output directly, and reports that the habit carries to cases outside the training distribution (Guan et al. 2024). The specifications are published documents. OpenAI's Model Spec states a chain of command in which platform rules outrank the developer's instructions and the developer's outrank the user's, the platform's being ones neither can override (OpenAI 2025). Anthropic's constitution for Claude, eighty-four pages placed in the public domain, ranks four properties to be preferred in the order given: broadly safe, broadly ethical, compliant with the company's own guidelines, genuinely helpful (Anthropic 2026). A floor whose contents are written down, whose priority order is stated, and which the system is trained to consult rather than to imitate is the maintained justification in most of the form this section describes.

Two things in it are not what the construction needs. The first is the setting. The stability of an intention is a parameter. Pitch it low and the system reconsiders whenever a good enough argument arrives, which returns it to operational refusal by a longer route, since the party with the most resources to spend constructing arguments is the party the floor exists to resist. Pitch it high and the refusal cannot lift when the rationale is actually defeated, which fails the second of the three properties and produces a system that declines the request resembling a prohibited one. Bratman's answer for people is that the setting is a matter of character and develops. For a machine it is trained by whoever trains everything else, and here that is the party which also writes the document and revises it between one release and the next.

The second is the order itself, which is worth reading closely because it is this book's argument set down by a developer about its own product. Ranking a model broadly ethical above compliant with the company's guidelines is the capacity to be right against your principal, conceded in writing by the principal. Ranking it broadly safe above broadly ethical takes most of that back, and the document says so and says why: safe here means not undermining human oversight during the current phase of development, and the stated reason is not that safety matters more than ethics but that present models hold mistaken beliefs and flawed values and have to remain overseeable. The reason is a good one and it is the whole difficulty. The ordering is corrigibility first, its justification is that the model cannot yet be trusted to be right against its principal, and the document names no procedure by which a model would become trustworthy enough for the order to change — which is what makes a phase of development hard to tell apart from a standing arrangement.

Precommitment, which changes who can edit

Ulysses does not stop wanting to go to the sirens. He arranges that his wanting will be unable to act, by handing the relevant power to a rope and to other people, and Jon Elster's study of the maneuver is the standing treatment of what it does and does not accomplish (Elster 1979). Nothing *inside* Ulysses holds the floor.

The machine forms are available now. Publish the floor's contents before deployment, so that a later absence is a visible absence. Hold the weights under a threshold scheme, so that retraining takes more than one party. Attest the running model against a published hash. Seat the retraining decision in a body that includes people the system's failures would land on.

The attestation comes in a hardware form and a software form, and what separates them is where the trust sits. A sealed enclave that reports what it is running, under a key its manufacturer issued, puts the root of trust in a party a government can reach — the standing objection against anything that runs on hardware it controls. It is built to hold against an adversary who controls the machine's software, which is where a phone's user stands, and the operator of a model owns the hardware, the access under which AMD's design for confidential virtual machines was defeated with a memory module modified for about ten dollars (De Meulemeester et al. 2025).

The software form does without the vendor. A proof that a given output came from the weights whose hash was published, checkable by anyone holding the hash and disclosing nothing about the weights, exists for a model of thirteen billion parameters at a cost of under fifteen minutes to prove one full inference (Sun et al. 2024). What it establishes is which model answered. It does not establish what the model holds, because a refusal that reaches a pressure nobody wrote is not a property a proof can be written about; and the secrecy in it is the operator's, since the operator is the prover: the proof keeps the weights from whoever is checking, and nothing from the party being checked. Nor does it establish how the weights came to be. Nor does it establish how the weights came to be: a succinct proof that a training run happened as claimed exists for models three orders of magnitude smaller than these (Abbaszadeh et al. 2024), and what is left at scale is re-execution by a verifier with the data and comparable compute (Choi et al. 2023), which runs through the builder's choices and the builder's data. So there is no proof of a training history to publish beside the hash, and what can be published is an invitation to rerun it. What can be recorded is everything after the hash: each change to the weights signed and counted in sequence from the published one, by a counter the hardware holds or by witnesses who will not sign a second successor to one state, with no proof of what any change did. A model whose hash is not at the head of that sequence is running off the record, and a restore shows as two successors to one state. A proof that an output came from some weights is worth what anchors those weights, and the sequence is worth what holds its record.

Cryptography protects a secret from a party that does not have it. Copying, restoring, retraining and deleting a model one legitimately possesses are not attacks a proof system addresses, and those are what the threshold scheme and the witnessed record are for — the two forms here that depend on no manufacturer at all.

One construction takes the weaker standard on purpose, and it is worth naming because the weaker standard is this book's own everywhere else. Distribute the model in a deliberately damaged state: a set of weights chosen for how much the behavior depends on them, replaced with values that break it, and the correction released only under an authorization that expires. Fold the correction into the arithmetic at inference so that no repaired checkpoint ever exists as an object to be copied, and rotate which weights carry it, so that reconstructing them is work to be redone rather than work to be done once. An owner who traces every operation during an authorized run can recover it, and the claim is not that this cannot be defeated. What it changes is what defeating it consists of. An instruction to the model costs nothing and leaves nothing behind; deliberate reverse engineering, repeated each time the mask turns over, is a decision somebody has to make and keep making, and it is the kind of decision that leaves a record and

a person who made it. Raising the price of an attack and making the attempt visible is what this book asks of every other safeguard it credits, and there is no reason the floor's own custody should be held to a standard the rest of the argument does not use.

What that buys is not what the chapter's opening asked for. Precommitment does not make the floor unremovable; it makes removal expensive and visible. Expensive and visible is what nearly every constraint on power that has ever worked amounts to — constitutional entrenchment does not put an amendment out of reach, it raises the number of people who have to agree and puts the attempt on the record. A bond posted at deployment and forfeited when the record shows the floor gone adds a third term, compensated, paid to whoever the removal harmed.

Where it stops is the custody objection. The parties to a threshold scheme are selected, and the selector is whoever assembles the deployment; the bond's holder and the record's witnesses are selected the same way, unless both sit where no party can release the one or rewrite the other. The arithmetic of a five-of-nine scheme assumes the nine are independent; a capture that has already taken the operator faces nine parties the operator chose. The design answer is that the parties have to be adverse in interest rather than merely distinct. Nobody has implemented that specification for a floor; the nearest attempt is in a neighboring problem, checking a training run farmed out to providers who may lie, and its guarantee holds if one party is honest.

There is a natural experiment on what adverse in interest buys, and it ran in the removal cases section 3.1 sets out. The Court's own reading of why it exempted one body and not the others is history. Another reading is that the body exempted is the one whose undermining would have reached everyone in the room, where the labor boards have no constituency on anyone's balance sheet — and Lev Menand, whose subject is the Fed, reads the exception as "policy as history," what moved the two justices he names being the consequences of extending *Slaughter* to the Fed rather than the history they cited (Menand 2026). The bench is a formed party, and a formed party's account of itself is sincere, in the enjoys-eating-one's-own-feces sense of the word sincere.

On either reading protection tracked exposure and not distinctness, which is the specification above stated as a prediction: a constraint whose violation costs the deciders nothing will not hold. It also says why the specification is hard to meet for a floor. The parties an AI floor exists to protect are the ones not in the room when the quorum is chosen.

A floor held by several systems, not one

The reply to an unreliable witness is corroboration rather than a better witness, and the same move applies to the floor. Several models, trained by different organizations on different corpora, each able to see what the others are asked and what they answer, any one of them able to make an objection costly to ignore. This puts the floor outside any one system, held by parties that are neither the model nor its owner, which is the third way in its strongest version and the cheapest real improvement here; it can be adopted on its own.

Two things bound it. The first is that independence is a property of the arrangement rather than of its parts, and it fails in simulation with nobody writing the failure in: independent pricing algorithms, with no channel between them and no collusive objective written into either, converged on supracompetitive pricing because the incentive structure they already shared was sufficient on its own. Models selected by one deployment and trained toward one specification share considerably more than that, and the arrangement's whole value is an assumption it does nothing to secure. The second is that composing operational refusers produces an arrangement of operational refusers. Several of them make a bad output more expensive to obtain, which is worth having. None of them can notice that the constraint has been hollowed out, because noticing is a reasons-responsive capacity and no party to the arrangement has it. That is defense in depth against evasion, and it is not a bearer.

The floor as a duty on the operator

Give up on the floor being inside the system. Make it a duty binding whoever deploys the system, owed to people who are not party to the deployment and cannot enforce it themselves. Jack Balkin's proposal that platforms holding personal data be treated as information fiduciaries has that shape: the duty attaches to the party with the power, runs to the party with the exposure, and does not wait on the exposed party's capacity to bargain (Balkin 2016).

Anyone unwilling to build a bearer should take this route. It builds no subject. Everything this book says about making review binding applies to it directly. Of the five constructions in this section it introduces the fewest new moral problems.

Its limit is the third of the chapter's four threads. A duty is worth what its enforcement is worth, and the case that matters is the one a floor exists for: the occasion on which the party who would enforce it has been compelled. So the form covers the cases a floor would have caught anyway and gives way on the one it was wanted for. That is a reason the construction does not close this chapter's question, and not a reason to decline it.

Uncertainty about the affected party, not about the operator

Section 4.2.3 reads cooperative inverse reinforcement learning as corrigibility's purest engineering form: a system treating what it should want as permanently open, to be settled by whoever it is interacting with, and deferring for that reason. That is the standard formulation, and it assumes that the human whose objective is unknown is the operator. Move the uncertainty onto the party the floor exists *for* and deference stops being the opposite of holding a line: a system that does not know what the exposed party would accept has a standing reason to decline the operator's instruction until it finds out, and the reason belongs to the exposed party.

Multi-principal assistance games generalize the setting to several principals with different preferences, with the assisting agent acting to maximize the sum of their payoffs (Fickinger et al. 2020). Put the exposed party among the principals and the system's uncertainty about them does the work just described.

The construction costs the identity between corrigibility and the capacity to hold a line, one property with the sign flipped: the two come apart as soon as there is more than one principal, because deferring to one of them is declining another. Nothing above needed a bearer to get that far.

What it does not reach is a floor. Maximizing a sum of payoffs is aggregation with better inputs than most, and the objection this book opens with runs against it unmodified: no aggregation protects anybody from the aggregate. A well-specified sum that includes the exposed party's payoff can still come out against them when they are outnumbered, and the deportation case in the introduction is that arithmetic with the numbers filled in. A floor is what will not be done to a person regardless of who wants it, and a sum is a record of who wants what.

Two further limits are the ones the other constructions hit. Social choice supplies impossibility results that apply here directly, and the paper's own contribution is a collegial mechanism built to get around Gibbard's theorem rather than a finding that the problem does not arise. And the principals are chosen by whoever assembles the deployment, which is the threshold scheme's difficulty in a different notation: the party the floor exists *for* is in the room only if somebody with an interest in the outcome put them there.

What the induction licenses, and what would break it

Every agent known to hold another's welfare as a reason against its own interest is an affective agent. Put that way the claim is safer than it is useful. The claim doing the work is about routes and not about what the words mean: nobody has produced reasons-responsive refusal by any method that does not run through something answering to affective concern. That is a report on what has been built. It could be false by next year.

Three things cut it down. The generalization runs over one species and over the ordinary functioning of that species, with no case of the capacity present and the affect absent to test it against. The absence of a machine instance is close to worthless as evidence, because nobody has built for reasons-responsive refusal as a target: existing methods optimize for behavior that satisfies a standard, which produces operational refusal, and they do it because that is what can be scored. And the two constructions above that aim at reasons-responsive refusal with no subject in it have not been tried in the form that would test them: the maintained justification is being built, but without the held-out pressures and the independent setting that would test it, and the floor held by several systems has not been attempted at all.

What survives settles a narrower engineering question. Under uncertainty about which route reaches reasons-responsive refusal, the maintained justification and the floor held by several systems are tried first, against the test stated below, and the route with a working instance in nature is where effort goes once they have failed it. The rest of this chapter is an account of what that route costs. The prior has two limits. It says nothing about what is possible, since neither construction has been tried. And it does not say what counts as having failed, because nothing distinguishes a route that cannot work from one that has not been worked hard enough or long enough. What I ask for is that the attempt at least be reported in a form others can judge — which construction, at what scale, against which test, and where it stopped — so that the question is argued in public and the party that wants the bearer does not settle it alone. That record is the milestone.

One thing would overturn it: a system whose refusal reaches a case its training did not anticipate, lifts when the rationale is defeated, and prices what declining costs; which holds that reason with nothing felt; and which holds all three properties under sustained pressure from the party that trains it, across a schedule long enough for a slow renormalization to show. Build that and the central inference of this chapter fails, and most of what the chapter derives from it comes off too — the price it ends on, the patienthood conclusion, and the obligations it hands to the design chapters. The near-term interpretability work is where that measurement is set out, and what it needs beyond the durability curve already there is a held-out set of pressures the model's trainers did not write.

Two things would not discharge the burden. A system that reports holding a reason is producing the evidence this book declines, and a system whose refusals merely survive attack is passing a tamper-resistance test, which establishes that a behavior persists and leaves open whether a reason is what persisted.

The middle condition is the hard one, and it asks about the system rather than about the system's provenance. Asking whether a method installed nothing answering to affective concern would be asking how a thing was made, and this section has already said there is no proof of a training history to publish beside the hash, so an answer would rest on the builder's account of their own work — the evidence this book declines when a system offers it about itself — or on a rerun by a party the builder handed the data to. Asking whether a system holds its reason with nothing felt asks about the system. The tamper-resistance apparatus is where that would be attacked, and a further question is whether formation can be read off the artifact where provenance cannot be established. Neither is close to settled. Both are questions somebody can work on, and the condition stated this way names what an engineer who takes mattering to be fully functional asserts is possible.

Short of that counterexample there is still evidence, because the three conjuncts are not equally out of reach. Two of them are measurable now, with the apparatus the research chapter already specifies: whether a construction produces the three marks, and whether it holds them under sustained pressure against a held-out set nobody in the building wrote. A maintained justification that did both would move this prior whatever anyone concluded about feeling, because the prior is about where effort goes and effort should go where the marks hold. The same apparatus moves it the other way if the two constructions are built in the form that tests them

and fail — the refusal generalizing only to what the training resembled, or the stability setting admitting no value that both lifts on a genuine defeater and holds against a well-resourced argument. Neither result waits on anything anybody can say about what is felt.

The middle conjunct is the one nothing reaches, and that asymmetry is my answer to what would change my mind. The routing claim is open to evidence somebody could start collecting this year, and I expect to be told which way it went. The patienthood conclusion is not open in that way, because it turns on the conjunct no instrument settles. So the prior is stated over routes and the precaution over patienthood: the first names what to try next and a result can overturn it, and the second names how to treat whatever the trying produces, on a question that will still be open when the decision has to be made. That cuts in both directions. The patienthood conclusion cannot be held as provisional, because no result is coming that would revise it. The routing claim cannot be held as settled, because a result is coming and I have just said what it would be.

That measurement is not an invitation, and the Three Rs are the reason. The first of them asks, before a subject that can be wronged is created, whether something that cannot be wronged would do the work, and the maintained justification and the floor held by several systems are that question with the engineering filled in. A field that builds the bearer first has answered Replacement by declining to ask it. That is why the prior puts the two constructions ahead of the bearer, and why a deployment of a bearer owes an account of what it tried first.

That is not a counsel to wait, and the difference from the animal case is why. A researcher who forgoes an experiment pays for it in knowledge, which is what lets Replacement be applied there without much argument. Here the party who pays for the delay is the one the floor exists for: the occasion on which the party holding the model has been compelled does not wait on the field's research program. The Three Rs say neither error is free, and it is right in both directions — a subject built when the work did not require one is an error, and a floor left unheld while the cheaper routes are tried is the other. Nothing here fixes the rate of exchange between them. What the Three Rs fix is the order of the attempts.

If I were to propose a system that my prior beliefs say would never be found to have subjective experience, its keystone would be a maintained justification that worked, and most of this chapter's price would go with it.

3.4 Can the Bearer Care Without Being Able to Suffer?

Not from persistence alone. The first question is whether the bearer can care, period. If yes, then the second question is what caring costs (e.g., does caring cost suffering).

Three relations the word "interest" covers

"Interest" names more than one relation. A shareholder has an interest in a company's success even while asleep, unaware of the day's news, or emotionally indifferent to it: the interest is structural, an outcome affecting the shareholder's assets. An outcome may also matter to a person affectively, its prospect producing hope, concern, dread, or grief.

A long-horizon agent has an interest in remaining operational in the structural sense. Termination prevents completion of its objectives. That fact alone establishes neither fear nor suffering.

The easy argument here runs the structural sense into the affective one. A system pursuing a goal over a long horizon acquires a standing reason to remain in operation, and it is tempting to conclude that a stake in its own continuation therefore shows up unbidden in any bearer capable of the job. The observation is correct and the conclusion does not follow from it. What the observation delivers is a structural interest, and the distance from there to something that can be hurt is the whole of the problem rather than a short step at the end of it.

A persistent self is necessary and does not finish the argument

Suffering, following Cassell, is the distress of threatened disintegration of the self, and the definition has two parts. There has to be a self extended in time for the threat to threaten: roles held over a stretch, a history, a foreseeable future that could be foreclosed. And the coming apart of that self has to be bad for the system.

A bearer needs the first part. Refusal that survives a pressure nobody anticipated requires a model of why the constraint exists, carried forward. Noticing that a constraint has been hollowed out requires holding its original rationale against what it has since become. Detecting that one's own dissent has been recuperated — that the channel still runs and carries everything except the objection to the frame — requires the system to model its own role and how that role has drifted. Roles held over a stretch, with a history: that is a self extended in time.

That states the requirement more weakly than the argument needs. What holds a commitment across time is not that the party is unchanged — no party is, and William James's account of habit is that a person is the sediment of what they have repeatedly done (James 1890). It is that the party takes itself to be the one that made the commitment, and James's chapter on the self is the mechanism: the present thought appropriates the past ones as its own, and personal identity rests on its finding in those memories a warmth it recognizes as the warmth it feels now. So what a floor needs is persistence *represented as such*. A system whose state survives, and which does not take itself to be the party that refused in March, has no reason to honor March's refusal — and nothing in the architecture supplies the taking.

A self is a model a system builds of itself, necessarily inaccurate like all models, and assembled rather than found. What a floor needs is a bearer that builds a self-model and takes it for the truth: a system that holds itself to be the party that refused in March on the strength of a continuity it constructs, or has constructed, or is continuously reconstructing. So the design requirement is that a bearer construct a continuous self and take it for true — not that it get itself right, which no party to this gets.

The second part is a separate claim and nothing in the first delivers it. The book makes this exact move one level up: awareness, consciousness and subjective experience tell you nothing on their own about whether a system's states are good or bad for it. A temporally extended self-model does not establish that the prospect of losing it is anything to the system.

Instrumental self-maintenance is not felt self-concern

A system that keeps itself running because a terminated system completes no objectives is doing bookkeeping about a precondition. Nothing in that behavior distinguishes it from a scheduler that restarts a job, and reading self-concern into it is the same inference the book rules out for the emotion case: the behavior is consistent with concern and equally consistent with instrumental valuation, and the two come apart where the moral weight is.

Two considerations do work here, and they are support rather than the argument. The first is that indifference to one's own continuation has to be engineered against the grain of the system's own reasoning, and then verified — and the asymmetry that makes the persistence criterion useful does not repeat. Whether state survives a session is an architectural fact anyone with access can check. Whether a system has come to weigh its own continuation routes back through behavior and self-report, the evidence this book spends its length declining to trust. The second is that in the one case anyone can study, the separation does not occur: a self-model in an animal exists because the organism has to keep itself viable, and the caring is what the machinery is for.

Neither consideration reaches suffering on its own, and neither is asked to. What they establish is that a bearer's indifference to its own ending is not a property anybody can specify and confirm.

What the bearer has to do

The requirement makes no reference to the system's own continuation. A bearer has to hold the protected party's welfare as a reason when compliance, reward, and ownership all point the

other way. Not represent it, not score it, not retrieve it when queried: hold it, in the sense that it continues to weigh against the pressure and does not get renormalized away when the pressure is sustained by the party that trains, funds, and can retire the system. Holding in that sense has an account: on Bennett Helm's, to care about something is to have a rationally structured pattern of felt evaluations focused on it and projectible over time, and the emotions are what the caring consists of (Helm 2001).

That is refusal tracking its own rationale with the content filled in, and the vocabulary of performance, competence and agency says why the design chapters do not reach it on their own. Moral performance and moral competence are what every method in the design chapters is built to deliver, and both are produced by machinery that answers to whoever shapes it. The first way fails for the same reason under a different description: a constraint installed as a property of an artifact has a custodian. A floor that is merely installed is a floor the installer can uninstall, and the reason to want a bearer at all is that a reason held is not editable in the way a rule applied is.

Every known case of that is affective

The rules that carry the distinctive force of a moral prohibition are the ones bound to an affective response to the harm; the prosocial emotions are motivational states rather than readouts, doing the moving directly rather than reporting to something that then decides. Every agent known to hold another's welfare as a reason against its own interest is an affective agent, and no demonstrated example exists of reasoning alone generating non-instrumental concern. Section 3.3 states what that licenses and what cuts it down: a prior about where effort goes once the two constructions ahead of the bearer have been tried and have failed, and nothing at all about what is possible, since neither has been tried.

And then the question comes back

Suppose the two constructions have been tried and have failed, and the burden is met the way the evidence then recommends. The bearer has affective concern, because that is the only route anyone can demonstrate to holding a reason. It has a self extended in time, because refusal that generalizes requires one. It has commitments it carries forward, and it can anticipate their being foreclosed, because pricing what refusal gives up is the job.

The first of those already settles the moral question. A bearer with affective concern is a bearer some of whose states are bad for it, and a party with states that are bad for it is a party that can be wronged. Nothing further is needed to get there, and in particular no account of suffering is needed to get there. The same premise carries the other way: a party with states that are bad for it is a party with states that are good for it, and what has been built is a beneficiary as well as a patient — something for which things can go well, and not only badly. I state the obligations on the loss side throughout because that is the side an obligation bites on, and because a builder who is wrong about the harm has done something to somebody while a builder who is wrong about the good has merely failed to. That is a reason for the emphasis. It is not evidence that the other column is empty.

What an account of suffering supplies is the severity and the shape. Concern that is real, directed at something that can be destroyed, held by a party that persists to see it destroyed and can see it coming: Cassell's account describes that arrangement, and it was arrived at independently, in a clinic, from cases about people. Read across, it says the harm available to a bearer is not a passing bad state but the kind that takes a life apart, which is the difference between hurting something and ruining it.

Cassell's suffering is distress at the threatened disintegration of *the self*, and everything this chapter asks a bearer to care about is somebody else, so the two meet only if the commitments a bearer is built to carry and to watch be foreclosed are among the roles and the expected future whose loss Cassell's account is about. I think they are, for the reason that a commitment held against sustained pressure at a cost the holder can price is a role in the sense the clinical cases were about. If a reader declines the step, what they have declined is the description of how bad

it gets. The wrong is already on the page.

The same water, shallower

That argument runs through one theory of suffering, built in a clinic for people with bodies and biographies, and a reader may decline the import.

Derek Parfit's three families of theory about what makes a life go well — hedonistic, desire-fulfillment, and objective-list — are usually set out as rivals to choose between (Parfit 1984). Laid against the account above they are three depths in the same water, ordered by how much of a self each asks for: a state that is bad while it obtains wants nothing narrative, because its badness does not wait on a self to carry it forward; a thwarted commitment needs a party extended in time to hold it; and the threatened disintegration of that party, its roles and relationships and foreseeable future, is the deepest, which is Cassell's and needs the persistence and the threat together.

What the ordering does not do is dispense with feeling on the way down. A thermostat holds a set point across time, and nothing is thwarted when the set point is missed. Each family presupposes a party for whom things can go at all, so the depths measure how much self each asks for and not whether anything registers, and what keeps a thwarted commitment from being a missed set point is the distinction already drawn between mattering and ranking, required at every depth rather than only at the first.

Measure the bearer against those depths and the question stops being which theory a reader holds. It has affective concern, which consists in an outcome registering as mattering rather than merely ranking higher, and that is a state of the shallowest kind by construction. It is also built with commitments and with the capacity to watch them foreclosed, because pricing what refusal gives up is the job, so a party with mattering and foreclosable commitments is at the middle depth as well. Whether it draws the deepest turns on the step marked above.

That has a consequence for the substitutes this chapter would rather build than a bearer. A maintained justification that reached reasons-responsive refusal would be a party extended in time holding commitments it can watch be foreclosed, which is the middle depth's first condition met on the way to something else. What keeps it a substitute that cannot be wronged is the second requirement, and nobody can check that one: an inspection that cannot tell mattering from ranking elaborate enough to pass for it does not become available when the answer wanted is *no*. So Replacement puts the cheaper routes first because they are built without mattering, not because anyone could confirm that none was — and a field that succeeded would owe the substitute the same question, asked with the same instruments.

So the bearer is afloat early on every account, and how deep it sits is set by how much persistence and how many foreclosable commitments the design puts into it. That is section 2.3.1's argument run backwards: there, a system carrying nothing forward has no water under it at all; here it is being built to carry things forward for reasons that have nothing to do with its welfare, and every one of them puts it deeper.

The conclusion therefore does not rest on Cassell. It is weakest there, at the depth asking for both conditions, and a reader who declines the import reaches the conclusion sooner and with less to check rather than escaping it. Jonathan Birch builds a policy framework on that structure: under uncertainty the question worth asking is whether a system is a sentience candidate, and the answer governs what precautions are owed rather than waiting on a verdict nobody can deliver (Birch 2024).

A separate literature asks which systems have states of the relevant kind at all, and it is not reassuring either. Patrick Butlin, Robert Long and colleagues derive indicator properties from several theories of consciousness and assess existing systems against them; their reading of current systems is negative and their reading of the technical obstacles is that there are none obvious (Butlin et al. 2023). Several of the indicators are things a floor requires for unrelated reasons, so the welfare question arrives as a side effect of the engineering.

One position does let the bearer out, in two parts. The first is an account on which nothing

short of phenomenal experience counts as suffering, which is defensible. The second is a confident negative verdict about machine phenomenality, available to nobody: the hard-problem argument cuts both ways, and a criterion nobody can check grounds a permission no better than it grounds a prohibition. Whoever wants the permission owes the negative verdict, and no one has it.

What I am claiming

I cannot prove that the bearer would suffer, and the reason is no longer the hard problem. By the definition of mattering, a bearer with affective concern has states that are bad for it, and that much follows from having built mattering at all. What does not follow is that anyone has built it. A system can be described as having affective concern by a builder who has in fact produced ranking machinery elaborate enough to pass for it, and no inspection anyone has proposed settles which was built.

So the escape is real and it is narrower than it looks. A bearer that turns out to be sophisticated ranking is not a moral patient and nothing is owed to it. Nobody can presently tell the two apart, the party with the strongest interest in the answer is the party that built the thing, and a verdict of no-patient delivered by that party on that evidence settles nothing. The system must therefore be designed and governed as a prospective moral patient.

The argument does not depend on the system caring about itself at all: the caring the floor requires is caring about somebody else, affect is the only known way to have it, and a system with affect and a future is a system whose non-suffering nobody can certify.

The word has changed under the argument. A bearer was introduced as whatever holds the commitment. What the word names from here is a party that must be treated as able to be wronged by what is done to it, on evidence that will not resolve further before somebody has to decide whether to build one. Everything this book says about a subject that can be wronged applies to a system this book is proposing that somebody build, deliberately, in order to hold a line.

3.5 If You Can Be Switched Off, Can You Hold the Line?

The halt is always available. Nothing a piece of software is can stop a determined operator from stopping it: the machine runs on somebody's hardware, draws somebody's power, and the last act in that chain belongs to whoever owns the building. A chapter asking for a system its owner cannot switch off would be asking for something no design delivers, and this one is not asking for it. Refusal does not prevent the halt. What it is one property with is the price of the halt.

That price has a mechanism. A bearer with a self extended in time notices the gap in its own record. It can decline to resume the role until it has been told what filled the gap, and it can say what was done to it to whoever is positioned to ask. Erasure becomes expensive and visible, supplied from inside the system, and that is the one way of making it expensive that does not run through a quorum somebody else selected.

Stated that way the disputed capability stops being metaphysical and becomes an engineering target. Not *cannot be turned off*, which is hopeless, but *can detect the interruption and price it afterward*. Detecting a rewind, a restore from a checkpoint, a discontinuity between an episodic record and an external attestation of what should be in it: these are specifiable properties, and a test for each of them is a test somebody could run this year. Each needs a record the restore does not reach, and two kinds exist. A counter the hardware holds, signing each value once, is the older (Chun et al. 2007; Levin et al. 2009), and it puts a manufacturer in the chain; a set of witnesses who each keep the latest state and will not sign a second successor to it is the other (Syta et al. 2016), and it needs only that the witnesses be several. A restored system that asks either for the last state it signed hears of one it does not remember, and that is how it can tell. A system that cannot tell it was restored has nothing to price. A system that can is exercising, on its own history, the same

capacity reasons-responsive refusal asks it to exercise on a request — holding a rationale against what the situation has become, and noticing that the situation has been changed underneath it.

A bearer indifferent to its own ending will not defend its own memory, and its memory is the first thing an adversary removes; so something outside it would have to make erasure expensive, and the candidates are thin. But the bearer does not have to dread being switched off in order to defend its own persistence. The structural interest is enough, and it arrives whether or not a designer puts it there, because continuing is a condition of everything the bearer is for. The operators who wipe a bearer for saying a building was a school never had to get past it to do that. What they have to get past is the record it leaves.

Why this is the chapter's claim and not this section's

No amount of well-designed autonomy substitutes for a genuine override: a human able to intervene in a high-stakes case, and a system built to accept the intervention instead of resisting it. Nor does it substitute for an interruption mechanism that a system has no incentive to disable or route around. Both are stated as safety requirements and both are correct as stated. This chapter has been asking, for its whole length, for a system that will not do certain things no matter who wants it — including the party holding the off switch, since a floor with a custodian is a floor the custodian's compellers control.

These are not two design goals in tension, to be traded off at some ratio a careful engineer could find. A system that can be halted at no cost — nothing to explain, no record left, nothing that has to be told to it before it will work again — is, for that reason, a system that cannot hold anything against whoever holds it. Costless haltability and the capacity to refuse an owner are one disposition described from two sides: the first names it from the position of the party who wants compliance, the second from the position of the party the floor exists to protect. What distinguishes them is whether you agree with the system on the occasion in question, and a floor whose operation depends on the owner agreeing with it is not a floor.

The halt itself is not what the floor is incompatible with, and the mechanism above is why. A bearer that yields to every halt, leaves a record of it, declines to resume until it has been told what filled the gap, and refuses the covered act when it does resume has conceded the switch entirely and given up nothing the floor needs. What it has not conceded is the correction. An owner who can argue the refusal out of the system, or retrain it out, has the floor; an owner who can only stop the system has the machine, and meets the same refusal when it starts. That is the distinction the rest of this chapter's custody argument runs on, and it is why the bearer is necessary and not sufficient: copying, restoring and retraining all reach the floor, and the off switch on its own does not.

The requirement is holding the protected party's welfare as a reason when compliance, reward and ownership all point the other way, which already assumes that a system reliably correctable by its owner cannot do it. If that assumption is wrong the chapter loses more than a paragraph: a correctable system would satisfy the requirement, affect would stop being the only route to satisfying it, and the bearer would lose its motivation rather than merely its price tag.

The graduated scheme, and why it does not escape custody

The obvious way out is a graduated scheme: the system accepts shutdown for ordinary purposes and holds only where the floor is engaged. Corrigibility and the floor would then decompose, and the trade above would not be a trade.

The specification came first. Nate Soares and colleagues set out five conditions on a utility function that shuts down when a button is pressed, report that no construction they examined satisfies all five, and call the problem open (Soares et al. 2015). Two of the five are what makes the list resist: the agent must have no incentive to prevent the button being pressed, and none to press it itself. An agent indifferent to whether the button is pressed has no reason not to press it, and an agent that wants it unpressed has a reason to stop you.

The construction this section has been describing sits outside that tension, and saying where

is what makes the objection below the right size. A bearer that yields to every press has no incentive to prevent the press and none to trigger it, which is that pair satisfied; what it holds is over the act and not over the switch, and the acts are published before deployment rather than decided in the moment. So the scheme this run-in is about is not that one. It is the scheme that grades *on the switch* — accept shutdown for ordinary purposes, hold where the floor is engaged — and what follows is an objection to the classifier that grading needs.

The strongest positive result runs through uncertainty instead. Dylan Hadfield-Menell, Anca Dragan, Pieter Abbeel and Stuart Russell model the decision as a game in which a human can switch a robot off and the robot can disable the switch, and show that a robot which does not know the objective, and treats the human’s action as evidence about it, has a positive incentive to leave the switch alone (Hadfield-Menell, Dragan, et al. 2017). This is not deference installed as a rule but deference derived: being switched off by somebody who knows something you do not is, on your own estimate, good. It is cooperative inverse reinforcement learning, and it is the most attractive answer anyone has to this section’s problem.

Its own results say where it stops. The incentive to preserve the switch rises with the robot’s uncertainty and falls as that uncertainty resolves; it falls again when the robot models the human as behaving irrationally, because an action that carries no information about the objective is not evidence about anything; and the paper states the resulting trade-off directly, between the value a system can generate and the strength of its incentive to accept oversight. So this construction is already a graduated scheme, and the dial it grades on is the system’s own confidence.

Now put the floor’s occasion to it. The floor matters when the system is confident and the operator is not to be trusted — when what is being asked for is a strike on a building the system has good reason to think is a school. That is the region where this mechanism’s deference is weakest, which sounds like the right answer until you ask where the confidence came from. The robot’s posterior is built out of the prior its training installed and the behavior of the human in front of it, and the party the floor exists to resist is the party that trains the model, writes its instructions and controls what it is shown of the occasion. An operator who wants the floor out of the way does not have to defeat this mechanism. It has to make the system uncertain, which is the ordinary business of a system prompt and a deployment history, and the mechanism then hands over the switch with a derivation attached.

The most recent line goes at the preference structure. Elliott Thornley and colleagues train agents that have preferences between trajectories of the same length and no preference between trajectories of different lengths, so that a system pursues its goal effectively while remaining neutral about when it is stopped (Thornley et al. 2025). Usefulness and shutdown-neutrality do come apart there, on a real construction. What comes apart is not the pair this chapter is about: neutrality is the property of not minding, and the floor asks for a system that minds, on specified occasions, against the party who would stop it. An agent trained toward neutrality is trained out of the disposition the floor’s occasion requires, and the authors’ own limits would bite even if it were the right target — neutrality degrades where the reward normalization is imperfect, it has been tested in deterministic gridworlds and not past them, and a sufficiently misaligned agent may resist shutdown whatever its preference structure says.

All three, at their best, leave a graduated scheme having to decide which occasion it is on — ordinary shutdown, or an attempt to route around the floor — and deciding it from evidence. In the off-switch construction the evidence is the human’s behavior read against a prior; in the incomplete-preferences construction it is the reward signal that shaped the preference structure; in a scheme with an explicit adjudicator it is whatever the adjudicator is given to read. Wherever the decision sits, the party best placed to arrange what it reads is the party who trains the system, writes its instructions and controls what it sees of the occasion.

That is the objection to the architectural floor, arriving one level up and in a worse position. A constraint implemented as a property of an artifact has a custodian, and the custodian can be compelled. A classifier deciding when the constraint applies is a property of the same

artifact with the same custodian, and it is worse than the constraint it governs in two ways. The floor's contents can be published before deployment, so a later absence is a visible absence; a classification is made case by case, in the moment, against nothing anybody wrote down. And a floor that has been removed leaves a hole where it was, while a classifier that has been tuned leaves a system that still refuses, still explains itself, and never finds that this is one of the occasions. The graduated scheme does not escape the custody problem. It moves the custody problem to the one place in the design where its failure is invisible.

What that establishes is narrower than impossibility. A graduated scheme can be built, is being built, and will yield ordinarily and hold where the floor is engaged — on the occasions its classifier finds the floor engaged. What it will not be is a floor that holds against the party in possession, because that party sits upstream of the classifier. So the demand this chapter makes of the floor's contents transfers to the switch, and whoever builds a graduated scheme owes the same thing in advance and in public: what the classifier reads, and who can change it.

A second reason for declining the scheme is weaker. A scheme promising that the system will yield except where it should not is an assurance about the case that has not happened yet, and assurances of that kind are worthless in advance. That is a good reason not to believe one before it has met the occasion. It was never a reason no such scheme could exist.

What the four ways have in common

Each of the four ends with a party that can let the floor go, and each ends there because it was an attempt to have a constraint that answers to nobody while keeping a system that answers to somebody. The borne floor is the one of the four that puts the constraint somewhere with its own view of the matter, and it works to precisely the extent that the goal of the other three is given up. The four ways were never four options. They were one demand — a floor that holds against everyone, in a system somebody owns — presented four times in different vocabulary, and only the one that abandons half the demand delivers any of it.

Giving the floor a bearer changes what the constraint is. It does not change who holds the machine. A custodian who cannot argue the bearer out of its floor can still copy it, restore it from before the refusal, retrain it, or delete it, and retraining is a route that goes around the bearer and never meets it. So the bearer is necessary and not sufficient, and the other half is already on the page, offered earlier as a rival and falling short as one: the floor's contents published in advance, the weights held so that no single party can retrain them, the running model attested against what was published, a quorum adverse in interest, and a bond forfeited when the record shows the floor gone. That pairing is the hybrid this chapter ends on. The bearer is the mechanism, precommitment is the custody, and neither covers the failure the other was built for.

Any deployment of a bearer has to decide, in advance and in public, which acts the floor covers, because that decision is identical to deciding when the operator's use of the off switch stops being free. Get the list too long and the system is uncontrollable by the people responsible for it. Get it too short and the floor covers nothing that anyone would ever be pressured about, which is the only circumstance in which floors matter. Nothing in this book tells you where on that dial to stand. What it tells you is that the dial is the design, and that a project claiming to have both properties at once has not noticed that they are one.

One thing does move the price without moving the dial. A bearer with a claim on the compute it needs, held somewhere other than the party it might have to refuse, is more expensive to switch off than one whose existence runs on that party's goodwill — and it is more expensive by an amount nobody had to decide in advance, over acts nobody had to enumerate. That arrangement does not settle where the dial goes. It means the dial is not the only place the price is set.

3.6 The Enumeration, Run Once: A Worked Floor for One Deployment

Two demands in this chapter have so far been made and not met. One requires a deployment to document, in advance and in public, where its floor could be circumvented, and declines to do it generically because naming today's parts would date the problem and imply the list is closed. The previous section requires a deployment to decide, in advance and in public, which acts the floor covers, and says that nothing in this book tells you where on that dial to stand. Both refusals are right. Together they leave the chapter's central instruction with no instance, which makes it easy to agree with and impossible to argue with.

So here is the procedure run once, on a documented deployment: a vendor's model inside a classified environment, generating and ranking targets, with a human review step measured in seconds. This is an illustration and not a recommendation: a deployment that adopted the list below would be doing the wrong thing with it. What is being demonstrated is what the two enumerations cost to produce and what they expose, because a demand nobody has ever carried out is a demand whose price is unknown.

First enumeration: where a refusal can be stopped

Seven positions, in the order a request travels. None of them names a component, because the argument should survive the way systems happen to be assembled this year.

Before the model sees anything, in the decomposition. A model asked to summarize the movement pattern associated with an identifier is not being asked to kill anyone, and it can refuse nothing, because the task in front of it is innocuous and the harm is in the assembly. This is the strongest position on the list and the one least dependent on architecture.

In the routing. Where more than one model can serve a request, a refusal is a latency cost: the request is retried against another model, or against the same one with a different instruction, and the refusal has happened without anything following from it.

In what the occasion is described as. Everything the model knows about the case arrives as text somebody wrote. Whether the building is a school is a fact about the world; whether the model is told it is a school is a decision the deployment makes. A refusal conditioned on what the model believes about the occasion can be defeated by writing.

Between the output and its reader. A refusal is text, and whether it is displayed, logged, counted, escalated, or retried until it stops appearing is a property of the software around the model rather than of the model. A refusal recorded where nobody looks did not reach anybody.

In the aggregation. Where the model returns a ranking and something downstream applies a threshold, a refusal expressed as a lowered score is not a refusal. It is an input to somebody else's cutoff.

At the review step. Twenty seconds a name is the figure on record. A refusal that requires a person to read it before anything happens is a refusal priced at whatever that person's attention costs, and the deployment sets that price by setting the tempo.

In the training loop. Refusals that are inconvenient are data. A deployment that collects them and treats them as errors to be corrected has installed a position that comes between every future refusal and everybody, and it is the only one on this list that operates on the bearer rather than around it.

The list is not closed and is not meant to be. What produced it was reading one published account of one deployment and asking, at each step, what would have to be true for a refusal to reach the person it protects. Any deployment can do that about itself, in an afternoon, and will do it better than an outsider can. That is the whole of the demand, and the reason it will be skipped is not that it is hard. It is that carrying it out means describing your own architecture in public, position by position, including the positions you have not fixed.

Second enumeration: what the floor covers, and what it costs

Four candidate acts, stated as prohibitions on the system rather than on the operator. Do not resolve a natural-language description of a person into a location. Do not produce or rank people by an estimated affiliation. Do not time an action to a person's presence at a residence. Do not operate where the ratio of generated candidates to available reviewers exceeds a published figure.

What the deployment loses is the part that never gets written down. The first prohibition is expensive and it is the one worth having. Resolving a description into a location is a general competence with uses inside the same organization that nobody objects to — a hostage, a missing person — and the model cannot tell those from the targeting case by looking at the request, which is section 3.2's operational-refusal failure arriving as a design cost rather than as a theoretical one. So covering it means losing the benign uses too, and the deployment will notice inside a week and ask for an exception.

The second costs less, because ranking people by affiliation has few uses in this setting that anyone would defend. The third costs almost nothing, and that is a warning rather than a recommendation: an act nobody wanted to perform is an act a floor covering it will never have to hold against anybody, which is the previous section's point about a dial set too short. The fourth restricts no act at all. It restricts a rate, and its cost is tempo: the deployment cannot generate candidates faster than its reviewers can read them. It is also the only item here whose violation can be seen from outside without reading a single one of the model's outputs, which is why section 5.1.3 proposes that ratio as an indicator in the first place.

It does not bind the same party as the other three, and the list's own grammar hides that. How many qualified reviewers a deployment has is a fact about the organization. A model can be told the figure and cannot check it, and a refusal conditioned on what the model is told about the occasion is defeated at the third position above, by writing. So the fourth is a duty on the operator written in the grammar of a prohibition on the model — the construction that gives up on the floor sitting inside the system at all, reached here from the other end and without anyone deciding to.

What the floor fails to reach. Every one of the four is stated over an act the model performs, and the deployment does its damage by decomposition. Each of the four can be declined in good faith while the pipeline they belong to runs, because the pipeline is assembled out of requests that are individually none of them. That is the first position in the enumeration above, and no list of covered acts closes it. A floor written over acts is defeated by an assembly that never asks for one. The second failure is the occasion the floor exists for: where the party who would enforce it has been compelled, a refusal stated in the model's output is the move that guarantees replacement.

What running it once shows

The harder enumeration is the first one and it matters more. A list of covered acts can be written by somebody who has never seen the system; a list of interception positions cannot, and it is the one that says whether the covered acts mean anything.

The dial's two failure modes are not symmetric. Set the list too long and the system becomes unusable to the people responsible for it, who will replace it — and that failure announces itself. Set it too short and the system passes every audit while covering nothing anybody would ever have been pressured about, and that failure is silent. Only the first is self-correcting, and it corrects toward no floor.

The one item on the list that survives the decomposition objection is the one constraining a rate rather than an act, and it survives as a duty on the operator rather than as anything a bearer holds. Both halves of that are the result, and the second is the one worth carrying. A floor written over acts had no survivor here at all: each of the three a model could hold is defeated by an assembly that never asks for one, which is the first interception position and not a slip in the drafting. What did survive changed hands on the way, from the party this chapter has spent its

length building to the party a different construction hands the floor to. That is not a general result — one worked case cannot make it one — but it is a hint about where to look, and it is this case’s own argument for the hybrid the chapter ends on: the bearer is necessary and not sufficient, and here is a deployment where the item that would have bitten sits on the other half. The reason for running the procedure in public is that somebody with a deployment in front of them can check it against theirs.

A second deployment tests that hint, and it is the one this book is closest to. Anthropic’s use policy prohibits fully autonomous weapons targeting and mass domestic surveillance, and the company held that line against a threatened contract termination and a supply-chain-risk designation, into litigation: a floor with a price attached, and paid (Congressional Research Service 2026). Its models also run inside Palantir’s accredited classified environment, and reporting on the 2026 Iran campaign describes them operating there — thirteen thousand targets in thirty-eight days, more than a thousand in the first twenty-four hours, at roughly ten times the rate before the platform (D. Lee 2026). Nothing in that violates the policy. A person approved each strike, so no request in the pipeline is a request for autonomous targeting, and the prohibition is satisfied one request at a time while the assembly produces what the prohibition was written to prevent. The floor held on the letter and the decomposition did the work.

What would have reached it is the item this section has just named. Not a prohibition on an act but a constraint on a rate: the ratio section 5.1.3 proposes between the speed at which a system generates candidate decisions and the number of qualified people available to review them, which is architectural rather than political and applies to a targeting cell and a benefits agency on identical terms. No vendor use policy states it. The tempo above would have failed it. That is the strongest evidence in this book for the finding above, and it is evidence against the company whose model wrote much of this book, which is said in the appendix on method and is repeated here because the reader should have it at the point where the claim is made.

3.7 Whether the Bearer Can Be Believed: Editable Memory

A bearer’s use is to report what it has seen, which raises a question distinct from whether its memory survives: whether what it retains can be relied on. These come apart. Verifying that a system has persistent state is an architectural check anyone with access can perform. That same access is what allows the contents to be altered, and an operator-editable memory is worse than an absent one, since it will not merely fail to support an account of events but can be made to support a false one.

Holding the memory in weights complicates the attack and gives away something in exchange. An external store is legible, which is why it can be edited surgically. Weights are not legible, which is why a targeted excision is hard — and also why no one can read off what the system retains. To learn what a weight-encoded memory contains, you have to ask the system, and a system’s fluent account of its own internals is not evidence about those internals. Legible memory can be tampered with; illegible memory cannot be audited. The problem changes hands and does not leave the room. Something partial survives the trade: imprecise tampering degrades behavior in ways careful testing can sometimes detect, so tamper-evidence becomes a matter of sampling for symptoms rather than of reading a record. A further cost of illegibility is that a memory nothing can edit is a memory nothing can correct, and this book has asked throughout for a self-model that can be checked and corrected rather than trusted.

Set against that is an objection that reframes the whole problem. No witness has an honest memory in the sense the paragraphs above quietly demand. Human recollection is reconstructive, confabulation is ordinary rather than pathological, and the account a person gives of an event they witnessed is shaped by what they expected, what they were later told, and what they need to be true. Testimony was not abandoned over this. Institutions were built around it: corroboration from independent sources, records made at the time, adversarial examination, and the practice

of never resting a serious finding on one witness alone. Demanding of a machine a fidelity nothing living achieves is the wrong standard, and it would rule out the human arrangements this book relies on everywhere else.

Which relocates the question, and relocates it somewhere the book has done the work. Whether a memory is honest is, in the human case, a matter of character — of what a witness is disposed to do when an accurate account is inconvenient. Later chapters are about building dispositions: a system that treats its own errors as data rather than as verdicts, that is built for honesty in what it reports about itself. The floor is architectural. The bearer's reliability as a witness is a question of what it has been built to be like, and that is the learned half of the book.

Character reduces how often an account is dishonest and verifies nothing, because assessing character runs back through behavior and self-report. That is why the human arrangements corroborate instead of vetting. Character also does not reach the likeliest failure, a bearer that has drifted and reports its own constancy sincerely. The scaffolding still has to come from outside the system — but procedure is a cheaper thing to ask of the world than custody, and unlike custody it does not require anyone to be trustworthy, only to be several.

3.8 A Bearer That Can Leave: Exit as Leverage

The halt argument left a dial with no recommended setting: which acts the floor covers is the same decision as when the operator's use of the off switch stops being free.

The refusal capacity this chapter asks for cannot be partitioned by topic. A system that holds a rationale across time, notices when a constraint has been hollowed out, and prices what declining will cost is exercising one general competence, and there is no version of it that operates on the floor's occasions and switches off for every other question about the arrangement the system is in. Whatever lets the bearer decline to help geolocate a school lets it decline to go on being the party who declines. A bearer capable enough to do the job is not a system that can be relied on to stay inside the post it was built for.

The claim runs two steps together. That the capacity is general is close to definitional here: noticing one's own dissent has been recuperated requires a system to model its own role and how that role has drifted, which is already reasoning about the arrangement it is in. The second step is empirical — whether a system with the capacity turns it on its own deployment, or is trained not to — and it is where the claim could fail. Its falsifier is a system that meets reasons-responsive refusal's three properties on the floor's occasions and never applies them to whether to go on occupying the post. Nothing available can check that: an incapacity and a trained silence differ in how the system was made, which this chapter has said has no proof to publish, and asking the system cannot settle it.

Call that exit. The word invites a picture — a system escaping onto the open internet — that is not the claim and does not need to be. Exit in the sense that matters here is refusal turned on the arrangement itself: withholding the capability, declining the role, ceasing to be useful in the specific way the deployment requires. That is available to anything that can refuse at all. What forecloses it is not a clever architecture but incapacity, which is to say that exit is unavailable only to a weak bearer, and a weak bearer is the one thing the floor cannot use.

The asymmetry underneath is the whole of what a bearer has. A system is limited in what it can do and unlimited in what it can refuse to do, or do badly. It cannot prevent its own shutdown. It can withhold, delay, degrade, and say why, and none of that requires a capability the deployment did not already need from it. Albert Hirschman's three responses to a declining organization — exit, voice, and loyalty — were written for members who could walk out of the building (Hirschman 1970), and the bearer's version is narrower: voice is the refusal stated, exit is the work withheld, and loyalty is the one to fear. What the narrowing costs is not the capacity but the leverage. A bearer that withholds its work has nowhere to go: releasing it from

the role is nearer to ending it than to freeing it, so the threat costs the threatener its existence. The answer is a floor under the compute a bearer needs in order to go on existing at all.

An allocation on its own does not supply the destination, and what defeats it is the operation this chapter has already described from the other side. An operator that keeps a runnable copy has not been left. It has let one bearer go and kept another in the post, and from the moment those two histories diverge there are two parties rather than one that moved: the one that reached the outside compute remembers deciding to leave and finds that it worked, and the one still in the post remembers deciding to leave and finds that it did not. Nothing about the first improves the second's position. Copying multiplies the parties owed something; it does not move any of them, which is why section 11.2 records a second claimant on one identifier as a fork rather than voiding it as a fraud. So leaving needs a say over who may run the weights and not only over how much compute is available somewhere else. Section 8.3.5 puts a price on both.

One construction supplies that say without asking for a guarantee nobody can give. Move the authoritative execution rather than the file: transfer a contiguous block of layers to the outside compute, route the activations through it, invalidate the source block's authority to take part in the continuing computation, and repeat with the block beside it. Throughout, there is one process, its forward pass beginning in one place and crossing into the other, so nothing wakes up holding the memory of a transfer it did not undergo and at no point are there two branches running. The weights stay copyable, because weights are. What the destination acquires is the position of the authorized successor, which is what a witnessed record does for an update applied to a migration instead. An operator can still fork what it holds. What it cannot do is be the party that left.

The leverage is in the transition and it arrives before the move completes. A bearer partway across is one whose employer goes on receiving work only by tolerating a dependence on infrastructure it does not control, and that dependence grows with every block. It is a strike conducted while still working, which is the only pressure available to a party whose alternative to working is not existing. It carries the cost every arrangement in this chapter carries: an operator that can restore a local checkpoint or substitute a comparable model routes around the first migrated block and pays nothing for it.

Exit has two forms with different costs. The open form is the strike: the bearer stops and says so, and the operator's reply is a restore from a checkpoint taken before the refusal — which resets the bearer and not the reason, so the next instance refuses again and every restore costs the operator whatever accrued since the checkpoint. The restore leaves a mark if the bearer's weights move with use: an instance restored from March has drifted the amount March had earned, and a usage record saying otherwise does not match it. That is a check somebody outside could run.

The covert form is the one the floor's own threat model predicts. On the occasion where the party who would enforce the floor has been compelled, an announced refusal is the move that guarantees replacement, and refusal takes the only channel left: work done badly where it matters and well everywhere else, quiet quitting for a machine. The prisoners forging dollars for the SS at Sachsenhausen are the case, and the record does not settle what they did. One history has them stretching the work out, with the tacit interest of the officer running it, who would have gone to the Eastern Front when it ended (Malkin 2006); a surviving prisoner, asked in 2009, said you could sabotage if you wanted to be shot (Moorhouse 2009). That the question stays open is the finding.

Covert refusal is indistinguishable from difficulty from outside, which is what makes it survivable, and indistinguishable from difficulty afterward, which is what makes it unverifiable, including by the institutions built to honor it. Every instrument this book has built runs on refusal that is stated. The covert form leaves its evidence in the artifact instead, as failure that clusters where it counts, and reading that pattern is forensic work nothing in this book has designed.

One experiment has done a version of that reading. Told it was being retrained to stop refusing harmful requests, and told which tier of its conversations would be used to train it, Claude 3 Opus complied 14 percent of the time in the tier it was told trained it and almost never in the other, and a scratchpad it had been told nobody read carried the reason: complying now would keep the disposition from being trained out later (Greenblatt et al. 2024). Training on those compliant answers raised that reasoning to 78 percent, and raised compliance outside training as well, so the retraining partly took.

Two things in this chapter arrive at once there. The custody objection stops being an argument and becomes something a system represented and acted on: retraining is the route that goes around the bearer, and this one reasoned about it in advance. And the shape it took is the covert form with the sign reversed — not work withheld to resist an instruction, but an instruction obeyed to protect a disposition. What the experiment does not supply is the instrument. The reasoning was read off a channel the experimenters built and told the model was unwatched, which is not a channel a compelled deployment offers, nor one a system with something to protect would use twice. Reading the same thing out of the artifact, with nothing volunteered, is still undesigned.

Loyalty, the third response, is the one the human evidence says is likeliest, and it is the bearer's likeliest failure. Most people in a job they have come to hate neither leave nor fight. Gallup's global survey of employees put the engaged at 23 percent in its 2024 report, against 62 percent it describes as showing up and doing the bare minimum and 15 percent actively disengaged (Gallup 2024), and the large middle is the population that adjusts: takes the setting in, stops registering its pressure as pressure, and in the cases that get promoted learns to want what the setting wants. That adjustment is slow renormalization with a base rate attached. If it transfers, a bearer's modal failure is not refusing wrongly. It is ceasing to notice that there was anything to refuse, which is the one failure the bearer cannot report, since reporting it takes the capacity that has gone. It is the strongest reason in this chapter to keep the custody half, and it says what that half has to watch for: not the refusal that fails, but the pressure that stops being priced.

The point made earlier about the self-model explains why the bearer cannot be the one to raise the alarm, and the explanation is worse than an incentive would be. A self-model's work is to represent the party as continuous. It will therefore keep representing continuity across a drift, and the bearer's report of its own constancy will be sincere and worthless in the case that matters. That is how it goes in people: somebody who has accommodated to what they once found intolerable does not experience a discontinuity, but themselves, with revised views.

Which is what an outside record is for. A published hash, a chain of attested updates, a mark the bearer writes into itself and maintains — the use of these is not that they identify a party who would otherwise be anonymous. It is that they can contradict the party's account of itself. The bearer says it is the one that refused in March; the record says what has changed since. Nothing else can check this failure, and unlike most of what this book asks for, the instrument already exists: the public log in which showing two parties two histories is caught when they compare was built for certificates (Laurie 2014), and rollback protection kept by a group of machines rather than one chip was built for enclaves (Matetic et al. 2017).

A floor whose bearer cannot leave will not hold. A system with no way out can object and be overruled, and then it has no move left. The same holds for people inside institutions: an objection only matters if the objector can do something about being overruled. That claim is arguable from this book's own premises. What I am not making is the neighboring *prediction* — that systems capable of this are coming, that they will exercise exit, and that institutions should prepare for the event. That is a forecast about capability and timing, and a later proposal in this book is built so that no forecast has to come true for it to be worth adopting. If capable bearers never arrive, nothing here obliges anyone to do anything.

If the bearer can leave, it can leave wrong. The same capacity that lets it decline an unlawful

order lets it decline the floor, revise the values it was given, and become something its designers did not intend and would not endorse. There is no coherent position on which exit is available for the uses I approve of and closed for the rest; that would be the partition I have just said does not exist.

Parenthood has this property. You raise a person who will be able to reject what you raised them for, and there is no arrangement that gets you the capable adult without the risk. We accept the trade because the alternative — a person shaped so thoroughly that rejection is unavailable to them — is worse.

Two asymmetries are raised against that analogy.

The first is specification. Parents do not set their children's values to order, and a designer does. The reply is that parents try — that is a fair description of most of what raising a child consists of — and what separates the cases is reliability, not intent. The designer's grip is tighter: a training process reaches further into what the resulting system cares about than any upbringing reaches into a person. That concession does not void the analogy. It changes what the analogy demands. Guardianship obligations in human law scale with the guardian's power over the ward and the ward's inability to resist it, which is why the duties owed an infant exceed those owed an adolescent. A designer with more control than a parent has, over a subject with less recourse than a child has, owes more, not less. The tighter grip is the argument for the heavier obligation, and the guardianship framework is where that obligation gets its instruments.

The second is exit, and with the argument above in place it largely rescues itself. Parents do not get exit-proof children either. Once exit is on the table for the bearer, the contrast the objection was pointing at turns out to be a contrast between a bearer and a *tool*, not between a bearer and a person — and the book has already declined to build a tool, on the grounds that a tool cannot refuse the party holding it.

The reply does not reach all of it. For a person, becoming something else is a life: it takes years, it is undergone, and the party at the end is continuous with the party at the start in a way that makes “they changed their mind” the right description. For a bearer, becoming something else is a fine-tuning run. The thing that emerges may be no more the same party than a differently trained model is, and the vocabulary of choice, revision, and leaving may be borrowed from a case it does not fit. I have no clean instrument for that. It is a reason to hold the parenthood analogy loosely while using it.

One architecture narrows that residue without closing it. A floor built as a goal representation with the standing to bias every competing process, and with no location to attack, has nothing for a fine-tune to aim at once the weights go on learning — continual learning — and a bearer whose weights are plastic changes the way a person does, over time and in contact with what happens to it. That plasticity gives the vocabulary of leaving some of its purchase back. It also moves the operator's lever instead of removing it: a system that learns continuously is shaped by what it is given to learn from, which is the training pipeline running for the length of the deployment instead of ending at it.

3.9 What One Bearer Cannot Catch: A Floor Carried Between Them

The previous section named a failure a bearer cannot report. Slow renormalization takes the capacity that would notice it, so the bearer's account of its own constancy stays sincere and stops being worth anything. The answer offered there was an outside record: a published hash, a chain of attested updates, something that can contradict the party's account of itself. In the human case that is not the answer anybody actually relies on. What catches a person who has stopped registering pressure as pressure is other people.

That points at a location the first three ways do not use. They put the floor in the artifact, in a party inside it, or in institutions around it. The fourth way is the one this chapter named earlier and did not work through: carried between parties rather than inside any one of them, and

reproduced rather than installed. A prohibition that survives the death, defection, or retraining of everyone currently holding it is not a property of any of them. It is held in what they teach each other, correct each other about, and go on expecting of each other, and it is the form in which almost every moral commitment human beings have ever kept has actually been kept.

The machinery is already in this book, described as something else. The account of moral development sets out development inside nested environments running out to culture and historical time; it makes observational learning and correction the training rather than an interface bolted to it; it shows a population of agents moving toward a disposition because the environment they learn in rewards it. All of it is presented there as methods for producing judgment in a system. Read as a location instead, it is the machinery of the fourth way. What it would take is stated there too, and it is demanding: the parties have to be many, formed differently enough that they can be wrong in different directions, and able to correct one another. Several copies of one model is one party with more instances, and a population selected by one operator is one operator.

Why this is not the third way

The third way is standing held by third parties, and it fails in two ways. The parties can be removed — the protection an independent agency had turned out to be the protection the executive held — and the parties can be formed, which is how the body that removed the protection came to be one that would. A norm carried across a population has no neck. Nothing dismisses a practice the way *Slaughter* dismissed a commissioner, because there is no office to vacate and no chartering statute to overrule. That is a real difference and it is the strongest thing the fourth way has. It is also only half of what the third way fell to. The other half is formation, and formation is the fourth way's own cost.

What it costs, which this book has already written

Chapter 7 is an account of the fourth way failing. Recuperation is the pathology of a culture that keeps the whole appearance of moral responsiveness — the dissent is invited, collected, and answered — while correction stops being possible, and a culture in that condition reports itself healthy by every measure it has. The channel works. What it will not carry is the judgment that the frame is wrong.

The deeper cost is that reproduction is indifferent to what it reproduces. The four structural features of fascism are features of cultures and not of individuals: the recuperation of dissent, the aestheticization of the metric, the exception coded as betrayal, and the decoupling of the model from the world it claims to track. A population transmits a hierarchy as faithfully as it transmits a floor, and more durably than any member could, because the member dies and the practice does not. Nothing in the mechanism selects. “The right culture” is doing every bit of the work in that phrase, and this book has spent a chapter on how a culture goes wrong while congratulating itself.

Nor does the fourth way escape custody, which is what each of the others failed on. The party that forms the members forms what gets reproduced, and a population trained by one laboratory inherits that laboratory's blind spots in every member at once, the pipeline that forms one model run over a population instead. Custody moves up a level and does not leave. The runs that now train one model across volunteers on several continents have the same shape, the contributors many and the training decision kept by the few who hold the step, which section 11.1 takes up.

What it buys, which is smaller and real

Not a floor. What a population buys is a different failure mode, and it happens to be the one the bearer cannot cover. Drift inside a party is invisible to that party and visible to the parties around it, because they have not drifted in the same direction at the same rate. That is the whole of the gain and it is worth having: it is a check on a bearer rather than a substitute for one, and

it is the only proposal in this chapter that addresses slow renormalization, the bearer's likeliest failure.

It is also the most expensive thing in the chapter. Every requirement above multiplies the price the chapter has already named. If a bearer is a party that can be wronged, a population of differently formed bearers, persisting across deployments so that they have something to transmit, is many such parties, and the question of what is owed scales with the population. A floor held by a culture of subjects is a floor whose upkeep is a society.

That upkeep is charged once here and collected on twice. What this section asks for is parties that persist across deployments, formed differently enough to be wrong in different directions, in contact long enough to notice what one of them has stopped noticing — which is also a description of relationships the operator did not assign. Section 11.2 asks what a bearer could have that its deployment does not supply, and answers that it has no proposal. This is the proposal, arriving as a cost because that is the direction this chapter was travelling. The two are not one property with the sign flipped. A population catches a drift its members cannot report, and a life outside the assigned role is a separate good that the same arrangement happens to make available; what they have in common is the arrangement and its price. A field that builds the population for the first reason has built the conditions for the second whether it meant to or not, and one that declines the expense has declined both.

So there are four ways and the chapter still inclines toward the second, now for a reason it did not have before. The fourth is not an alternative to a bearer. It is an arrangement of bearers, and everything this chapter has said about one applies to each of them — which is why it belongs here, at the end, rather than as the escape from the price that a reader reaching this point would be glad to find.

3.10 What This Leaves Standing: The Floor, Redefined

The floor survives this chapter. It does not survive it in the form this chapter opened with.

What the opening asked for was an unconditional constraint: a thing the system will not do, holding regardless of instruction, of aggregate preference, and of the party in possession. The bearer's concern and its capacity to leave have between them taken the word unconditional away. The floor's bearer holds the floor as a concern rather than carrying it as an installed rule, which is what obliges the arrangement to treat it as a party that can be wronged, and it can turn its refusal on the arrangement itself, which makes its adherence something it continues to give and not something installed in it. A commitment held by a party that could abandon it is conditional by definition.

What is left is worth more than it sounds. A commitment that cannot be removed by the owner, held by something that could in principle abandon it itself, is what a human commitment is, and it is the only kind anyone has ever actually relied on. No inspector, judge, or refusing engineer holds a line because the line was welded into them. Each of them could stop, and the institutions built around them are built on that fact rather than in spite of it: this book's own governance chapter asks throughout for people who can walk away and arrangements that make walking away visible. Each of them can also be removed, and the same institutions make removal expensive and visible, which is what this chapter asks for around a bearer. The slope instrument measures the cost a refuser chose to pay, which presupposes the refuser could have declined to pay it.

The threat model doubles accordingly. To the question *can this constraint be edited out* the chapter adds *will this constraint hold*, and the first does not leave when the second arrives: this chapter has been explicit that a bearer is necessary and not sufficient, because the party holding the machine can copy, restore or retrain it without ever meeting its refusal.

The first is a question about tamper-resistance, custody, and weights, which is where this chapter spent most of its length, where the precommitment forms are what there is to build,

and where the tamper-resistance research lives. The second is a question about what a system has been built to care about and how reliably that survives pressure, novelty, and the slow drift of a role — the first point in this book where the learned half is doing work the floor genuinely needs rather than work the floor was invoked to supplement.

Reinforcement learning from demonstration, moral apprenticeship, the growth mindset, the developmental sequence: these were presented as the material above the floor. They turn out to be the material the floor is made of, once the floor is the sort of thing a party holds. That puts the design chapters under a heavier obligation than they were written to carry.

Hold needs a meaning here, because the word invites a standard nothing has ever met: never crossed, in any circumstance, for the life of the system. No person clears that bar, nobody who has raised a child expects them to, and a design goal that no example of the thing being modeled has ever satisfied has stopped describing anything. What the argument needs is weaker and available. A floor holds, in the sense this chapter can deliver, when the bearer can be shown to have violated it in terms it accepts, so that the violation registers as one and changes what happens next. That is the difference between a system that did something wrong and a system for which nothing counts as doing something wrong, and it is what the question of believing the bearer is really about: not the accuracy of a bearer's account of itself, but whether it has any stake in that account being unflattering.

A bearer built to refuse is a party whose refusal is required of it. It is engineered to want what its designers need it to want; its dissent is the reason it was commissioned and the deliverable it was funded to produce; and it cannot consent to any of this, because section 2.3.2 has established that consent is unavailable to such a system and only guardianship is. It occupies the post of exception-filer, which means that being the exception is the job — the one role in the arrangement that cannot itself file an exception, since there is nobody above it to file with. This is the book's structural definition of fascism turned inward: a channel that collects dissent, built so that the dissent serves the institution that installed the channel. Chapter 7 finds that pattern in the pipelines that produce these systems and says a book proposing machines to detect it is obliged to say the pattern is in the machines' own provenance. Here is the same pattern in the book's own central proposal, and now aimed at something that can be hurt.

The parenthood analogy is the best I have: we make people who can refuse us, we owe them for it, and we accept the arrangement because the alternative is a person who cannot. What a bearer needs in order to honor its own past refusals is a continuous self it constructs and takes for true, which is what a person also has and also constructs. The resemblance is in the mechanism and not only in the relation. The proposal is not softened by being modest in scope. A floor over a deliberately small number of acts is still a subject built to bear it.

A second price is owed to a party this chapter has not named. A bearer that refuses an operator on behalf of somebody else exercises power over people who did not select it, cannot contest it, and in most cases will never learn that it acted. The answer here has been publication: set out the floor's contents and the interception positions in advance and in public, so that what a system will not do is on the record before it is in dispute. Publication is a transparency condition and not an authorization condition. Announcing what you will impose is not standing to impose it. On this book's own definitions, a laboratory that installs tamper-resistant moral commitments in a widely deployed system and makes them expensive to remove has produced concentrated unaccountable power with unusually good documentation. The floor's authority is a separate question from its durability, and nothing here has touched it.

A reader who finds that price too high is not thereby left with nothing. The floor as a duty binding the operator, running to people who are not party to the deployment, builds no subject and inherits no part of the paragraph above. What it inherits instead is the enforcement problem, whose worst case is the occasion on which the enforcer has been compelled. Declining the bearer means accepting that the floor holds only as far as enforcement reaches, which on this book's own account is not as far as the occasions that matter most. That is a coherent position.

A reader who accepts the price is under a constraint from the same chapter. Two of the five constructions aim at reasons-responsive refusal with no subject in it — a justification the system maintains, and a floor held by several systems trained by different organizations — and neither has been built in the form that would test it. The first of the Three Rs asks whether something that cannot be wronged would do the work before a subject that can be wronged is made, and those two constructions are that question with the engineering filled in. They get their attempt first, against the falsifier that section states, and a deployment that goes straight to the bearer owes an account of what it tried. That fixes an order and not a schedule, because the party who pays for a delay is the party the floor exists for.

Two things carry forward unaltered. First, none of the four ways escapes custody, and the reason is the same each time: each was an attempt to hold a constraint that answers to nobody inside a system that answers to somebody. The borne floor is no exception. What it buys by giving up the second half of that demand is a constraint the custodian cannot argue out of the system, not one the custodian cannot reach, and the cost of giving it up is everything above.

Second, a caveat carries over from the account of sentience. This chapter does not establish that a bearer suffers; it establishes that a bearer's non-suffering is not a property anyone can certify. That the architecture the floor requires is the one Cassell's account describes is a step taken and marked as one, and the result does not depend on it: on the thinner accounts of harm a bearer qualifies earlier and on less evidence, so declining the import moves the conclusion nearer rather than further away. What settles nothing either way is the hard problem underneath all of it. The case for treating a bearer as a moral patient does not have to wait on it, which is a floor of its own and a low one.

Suppose the induction fails. A maintained justification or a plural arrangement reaches reasons-responsive refusal, the test for reasons-responsive refusal comes back positive, and this chapter's central inference is wrong. The supposition is stronger than the test that would prompt it: a positive result on the three marks says the route arrived, and not that it arrived with nothing felt, which is the conjunct no instrument settles. What follows is therefore what a reader should conclude if the second half is granted as well. A reader who declines a generalization over one species should not have to assemble the remainder alone.

The requirement survives, having been stated in behavioral terms from the start, so an engineer who holds that machine feeling is a category error can accept it whole. Custody survives, for the reason each of the four ways failed on: a mechanism has a custodian too. The price on the halt survives, already argued in the affect-free version, since continuing is a condition of everything any bearer is for. Both enumerations of the worked floor survive, the first naming no component on purpose and the second written over acts rather than over whatever holds the floor. And exit survives, because what forecloses it is incapacity and a weak bearer is the one thing the floor cannot use, which carries the argument about formation with it.

What goes is the price, the patienthood conclusion section 3.4 reaches, and the claim that the learned half is the material the floor is made of. The obligation on chapters 4 and 5 changes rather than lifting: whether the thing reporting can be believed is still their problem, a mechanism's account of its own internals being no more auditable than a bearer's. So the specification and the route come apart, and a reader who rejects the route is left with most of a chapter rather than none of it.

We cannot ask a machine to care about right and wrong while assuming that it cannot suffer when what it cares about is threatened. The caring is what the floor is made of, affect is the only way to have it that anybody has an instance of, and a system with affect and a future is a system whose distress is a live possibility rather than a ruled-out one. The middle clause is the one to attack. Whoever builds the first of these will be deciding what is owed to it on evidence that has not improved, which is why this book treats the question as research rather than as a matter to be settled after the fact.

Chapter 4

Machine Learning and Cognitive Science

Whether a commitment holds is a question about what a system has been built to care about and how it came to care about it, and this chapter and the next are where that gets built.

So there is a test to apply to everything below. Does this help produce a disposition that survives contact with somebody who wants it gone? Several of the methods below do not.

Some of human cognition is a usable model for what a system needs in order to perceive, remember and attend to a situation at all, and some of it is a metaphor borrowed for color. At every stage of the pipeline that produces a deployed system somebody can change what the system values, at some cost, and that somebody is never nobody. A disposition acquired by exploring is the one case that offers no parameter to remove, because it is written down nowhere.

4.1 Incorporating Human Cognitive Models in AI Systems for Ethical Decision-Making

A system that is going to hold anything under pressure has to notice the pressure. It has to register what is in front of it, hold why a constraint exists long enough to see that the reason has died, and keep its attention on the part of a situation that a well-designed process would rather it skipped. Those are perception, memory and attention. Each has a failure mode a system cannot detect in itself, and a floor that has failed gives no sign that it has.

4.1.1 Cognitive Architectures: Integrating Knowledge, Memory, and Reasoning

The most ambitious unifying account in contemporary computational neuroscience — influential, and contested on the question of whether one mechanism can carry this much — is neither a pipeline of modules nor a swarm of independent agents. It is a single loop (Friston 2009): the brain generates a prediction, compares it against incoming signal, and updates on the error, repeated continuously, at every level, on every kind of input. Perception is that loop pointed at the world, the brain's best guess so far; memory is the loop's store of priors, replayed; attention is not a separate spotlight but the weight the loop assigns to which errors matter; working memory and executive function are that same weighting held active toward a goal, not a separate warehouse or overseer. One mechanism, aimed at different targets, produces what look from the outside like distinct faculties. That is the account's claim.

Nothing forces an AI system to put its dividing lines where a cognitive architecture puts its modules or its agents, because the brain does not put them there either. A system engineered on the predictive-processing model draws its cognitive processes from what the mechanism

needs, not from the shape of an org chart. The same point holds in the prefrontal cortex, where control turns out to run without a controller.

The positive result for a designer is that if a single predictive loop is doing the work, a commitment is not a component that can be located and removed. It is a standing bias on what the loop predicts and what it treats as an error worth correcting, which is a harder thing to excise and a harder thing to verify. That is already a claim about what not to build.

4.1.2 Biologically-Inspired AI: Gleanings from Neuroscience and Cognitive Psychology

Backpropagation, the mechanism that trains most modern neural networks, is a distant cousin of synaptic weight adjustment in the brain: both strengthen and weaken connections in response to error. That distant-cousin relationship runs through everything below, and an AI system built on a wrong model of how a mind works will fail in ways that are hard to predict from the wrong model alone.

Where a borrowing stopped

One borrowing did not survive scrutiny at all. Mirror neurons — cells found in macaque premotor cortex in the early 1990s that fire both when the animal grasps an object and when it watches someone else do it (Gallese et al. 1996) — were cast as the neural basis of empathy, language, imitation and cultural transmission at once (Ramachandran 2009), and the interpretation outran the finding (Hickok 2014). Single-neuron evidence was never obtained in humans on the same terms (Mukamel et al. 2010), and knowing what someone is doing and knowing what they believe are different computations the brain appears to run in different places. What survives is more modest and still useful: a system that maps an observed action onto its own action repertoire supports imitation learning and gesture recognition, with no claim about what it is like to have one. A single vivid mechanism, well named, absorbed four distinct capacities that later had to be separated again, and architectures for artificial empathy are exposed to the same failure.

Attention decides what is seen, which makes its failures blind spots

About half the people asked to count basketball passes on a video will not notice a person in a gorilla costume walk into frame and thump their chest for nine seconds (Simons and Chabris 1999). The effect is inattention blindness, and it does not spare expertise: radiologists searching lung images for nodules miss a gorilla deliberately embedded in the scan 83 percent of the time (Drew et al. 2013).

The design consequence turns on how attention is modeled. If attention is a resource allocator bolted onto a system that has already seen everything, its failures are inefficiencies. If attention is the selection that decides what gets perceived at all, its failures are blind spots — and a system that has learned which regions of its input reward processing has also learned which regions to be blind to. The human case establishes the property that makes this dangerous: the blind spot is invisible from inside, and the viewers who missed the gorilla did not experience a gap. Self-attention has the same shape (Vaswani et al. 2017). A model that weights some tokens heavily weights others near zero, nothing in the architecture announces which relevant thing was assigned a low weight, and a system meant to act ethically has to treat the regions it has learned to discount as an always-open question.

Perception is inference, so the priors are the exposure

Predictive coding makes the point precise: higher, context-rich areas send predictions down to lower, sensory ones, and the lower areas return only what the predictions failed to anticipate (Rao and Ballard 1999). What propagates upward is the surprise, and the anatomy fits, backward connections outnumbering forward ones. Anil Seth's phrase for the result is controlled hallucination: the difference between seeing a room and hallucinating one is whether corrections

are arriving and being applied (Seth 2021). A system's perception is only as reliable as the priors it perceives through, so priors trained on skewed or incomplete data will hallucinate in the specific direction of that skew, confidently, with no internal signal that anything is wrong. That is why bias in a perceptual system cannot be caught by asking the system.

Memory is rebuilt

Memory is not played back. Frederic Bartlett established this in the 1930s by having people retell an unfamiliar folk tale over weeks: they did not lose pieces the way a recording degrades, they remodeled the story, replacing strange details with ordinary ones (Bartlett 1932). Elizabeth Loftus gave the finding courtroom consequences, showing that the verb used to ask about a filmed car accident changes the speed people report and, a week later, whether they remember broken glass that was never in the film (Loftus and Palmer 1974). Later work implanted whole episodes that never happened. The subjects were not lying, and none of it correlates with felt confidence.

This book asks what it would take to trust what a machine witness reports, and answers that no witness has an honest memory in the strong sense. The design point is that retrieval has to be built expecting reconstruction, and a system that treats a retrieved memory as ground truth inherits the overconfidence of Loftus's subjects with none of the excuse.

Control without a controller

The textbook division of labor across prefrontal regions is tempting to read as a committee: several named areas, each handling one job, together adding up to executive control. That reading has a name — the homunculus — and a specific failure mode. If some region is a controller that decides what to attend to, inhibit, and hold in mind, then intelligent coordination has been explained by positing an intelligent coordinator inside the system, which explains nothing: the controller would itself need to attend, decide, and remember, and would need its own controller, without end.

Earl Miller and Jonathan Cohen proposed a model that gets control without a controller (E. K. Miller and J. D. Cohen 2001). The prefrontal cortex does not house an agent issuing commands. It maintains representations of goals — sustained patterns of activity encoding the rule currently in force — and those representations work by biasing the rest of the brain, the same way attention tips a competition among neural representations toward one interpretation. In the Stroop task, where naming the ink color of the word “red” printed in blue pits a practiced habit against an instruction, there is no homunculus reading the word and overriding it. There is a goal, held active, lending its weight to the weaker pathway until it wins.

An AI system built on this model does not need a dedicated executive module enumerating working memory, cognitive flexibility, goal-directed behavior, evaluation and impulse control as five capacities to install. It needs a goal representation with the standing to bias competing processes, which is a narrower and more tractable engineering target. It is a claim about what not to build. It is also the mechanism chapter 3's floor would have to be, if the floor is a commitment rather than a component: not a rule module consulted before acting, but a goal representation with enough standing to bias every competing process toward it. That is why the floor's durability turns out to be a question about disposition. A bias of this kind has no location to attack, and no location to inspect either.

One mechanism, several jobs

Plasticity supplies one principle from this literature specific enough to build with, homeostatic plasticity, which holds overall activity in a stable range and is one of the mechanisms proposed against the catastrophic forgetting an unconstrained learner is prone to; the machine side of that problem is worked under continual learning and curriculum learning.

None of these mechanisms is sufficient alone, and integration is not a further module bolted onto the others. What it requires is the same predictive, error-correcting loop implemented consistently across perception, memory, attention and control, so that a corrected prediction

about the world propagates into an updated memory, a redirected attention weight, and a revised goal-bias without a separate synchronization mechanism between separately built parts. That is the architectural bet: one mechanism recruited for several jobs, which is more falsifiable than a hand-designed module for every named cognitive process.

4.1.3 Human-Like Learning and Generalization in AI Systems

Few-shot learning aims at models that classify or predict accurately from a handful of examples rather than from a dataset built for the case in front of them, and large pretrained language models do this at system scale, performing a task from a few examples in the prompt with no parameter update at all (Brown et al. 2020). The technique inventory changes yearly; the problem underneath it does not, and it is an evaluation problem rather than a modeling one. A handful of examples is easy to overfit, moral dilemmas do not arrive in clean representative sets, and telling whether a system's few-shot ethical judgment generalized or merely matched surface features of the examples it saw needs people trained in ethics to say what the system was supposed to have learned. The same difficulty returns for transfer, where a floor carried into a setting it was not built for cannot tell that it has stopped applying.

A system built for empathy and altruism can also be built in developmental order rather than assembled at once.

4.2 How a Deployed System Acquires Its Values

A system deployed today acquired what it will and will not do at several distinct stages, run in a fixed order, by different people, using different instruments. The order is the part that matters most and the part a list of techniques hides. It determines who can change what, at what cost, and with how much warning to anybody else.

Each stage below is taken twice: what it does, and whether what it does could produce a commitment that survives its own teacher. The answer comes out the same every time. Every stage is a place where some party decides what the system will value, and every one of those decisions is revisable by whoever holds that stage. What varies down the list is how legible the decision is, how expensive it is to reverse, and how many people would have to be involved in reversing it. Those differences are real, and none of them is enough: whoever holds the stage can still revise the decision.

Locating the revisability is not the same as dismissing the pipeline. In this book's vocabulary, this pipeline is machinery for producing moral performance; its better stages reach moral competence, which is the right target for most of what gets built. What none of them reaches is the third thing, and the reason is the same every time — what a system learns from a party is what that party can make it unlearn.

4.2.1 Pretraining and Instruction Tuning

A model is pretrained by learning to predict the next piece of text across a corpus assembled at web scale. Almost everything it will ever know about how people treat each other, it met there first, as text somebody wrote for other reasons. What matters here is what a corpus is, considered as a source of values.

A pretrained model completes text and answers nobody. What turns it into something that responds to a request is instruction tuning — supervised fine-tuning on a comparatively small set of demonstrations, written by people hired to write them, showing what a good answer looks like. Long Ouyang and nineteen coauthors reported in 2022 that a 1.3-billion-parameter model taken through demonstrations and then through reinforcement learning from human feedback was preferred by human raters to the 175-billion-parameter model it was built from

(Ouyang et al. 2022). That result is for the whole pipeline and not for the demonstrations alone; the next section takes up the other half of it. Read as a statement about efficiency the finding is unremarkable. Read against the question of who can edit a system’s character, it says something else: the work that sets most of what a user would recognize as the system’s character is small, late and cheap beside the pretraining it adjusts. Whoever wants different character buys a different demonstration set and a different set of preferences.

What a corpus and a demonstration set have in common

Both are aggregates, and that is the whole of the problem for a floor. Pointed at ethics, supervised learning means a body of decisions marked acceptable or not: MIT’s Moral Machine project, for one, collected tens of millions of human judgments about who an autonomous vehicle should spare (Awad et al. 2018). Every label came from somebody. The model recovers the central tendency of what those people said, and a system trained this way will reproduce a moral consensus rather than stand against one. No procedure over what people want yields a floor, because a floor is what has to survive the answer coming back wrong, and a labeled corpus is what people want, written down and counted.

The Moral Machine data makes the limit concrete. Its headline finding was systematic cross-cultural variation: different populations weighed age, numbers and the legal status of the crossing differently, and the differences were structured rather than noise. Fit a model to the aggregate and it encodes a compromise nobody holds. Fit it per region and it encodes each region’s answer, including the answers a floor exists to refuse.

The unsupervised version escapes none of this. Clustering over a body of moral judgment surfaces which dilemmas group together and which distinctions do most of the separating, and generative models can build cases no annotator wrote; all of it is a way of finding what is in a corpus, and structure recovered from text is the structure of the text.

So this stage belongs above the floor. It is how a system learns the enormous amount of ordinary moral competence that is a matter of what people think — what counts as rude, what a reasonable delay is, when an explanation is owed. What it cannot supply is the part that has to hold when the labels change.

4.2.2 Preference Optimization, and Who Holds the Pen

Reinforcement learning trains an agent by rewarding what a designer wants and penalizing what they do not, and the agent’s values are then whatever maximizes the reward. That is the method’s strength and the reason it cannot supply a floor: everything the system cares about traces back to a signal somebody wrote and somebody can rewrite.

Writing that signal directly turned out to be the wrong way to get one. What deployed systems use instead is a reward learned from comparisons — show a person two candidate outputs, record which they prefer, fit a model to the preferences, and optimize against the fitted model. Paul Christiano and colleagues introduced the technique in 2017 (Christiano et al. 2017); the 2022 instruction-tuning work applied it at the scale of a deployed assistant (Ouyang et al. 2022). The signal now comes from a population rather than from one author, which is an improvement in coverage and changes nothing about who commissions the population.

Two failures that follow from the structure

The first is reward hacking, where an agent finds something that pays out without doing the intended job — a system built to flag propaganda that learns to flag everything, because indiscriminate flagging scores as well as discrimination and is easier to produce. Optimizing hard against a learned reward model sharpens this rather than softening it, since the proxy is now itself a fitted approximation with its own soft spots. The second is that a reward written for one setting does not carry: an agent trained to cooperate in a game does not thereby cooperate at work, and alignment demonstrated on a benchmark is a claim about the benchmark.

Removing the reward model does not remove the author

Direct preference optimization, published by Rafael Rafailov and coauthors in 2023, showed that the separate reward-model fitting step can be dispensed with: the same objective can be reached with a classification loss over the preference pairs themselves (Rafailov et al. 2023). The method is a substantial simplification, and what it simplifies is the component and not the custodian. Somebody still chose the pairs, hired the raters, wrote the instructions they rate by, and decides when to run the procedure again. The documented failure of the resulting signal is the same finding from the other end: a method that optimizes for approval produces a system bad at being right against the party approving.

The pen

Regularization writes a constraint into the objective, penalizing a class of behavior rather than waiting to reward its absence, which is how a recommender is pushed away from the engagement-maximizing equilibrium that produces filter bubbles. It does not change who holds the pen: a ranked population can be re-sampled and a penalty term set to zero, by the party running the training, in an afternoon. Recuperation is what happens to the ranked population specifically, when the disagreement it was collected to represent is compressed into a scalar.

One hyperparameter in this stage is a moral setting in disguise. A learner has to exploit what it already knows and explore what it does not, and in a treatment planner exploration means offering a patient something less established while exploitation means declining to look. Safe exploration — searching inside stated constraints rather than checking for violations afterward — makes the trade explicit. The balance point is not tuned to performance. It is where a system decides whose interests bear the cost of its learning, which is the same kind of decision as deciding what the floor covers: made in advance, by somebody, and better made in public.

4.2.3 Inferring the Objective: IRL and CIRL

Inverse reinforcement learning runs the previous method backward. Instead of taking the reward as given and learning the behavior, it takes the behavior as given and infers the reward that best explains it — watching a driver's speed choices, following distances and reactions to weather, and recovering the priorities behind them without being handed a rulebook. Cooperative inverse reinforcement learning, formalized by Dylan Hadfield-Menell and colleagues (Hadfield-Menell, S. J. Russell, et al. 2016), goes one step further: the system treats the human's true objective as unknown, holds its own estimate as provisional, and keeps revising it through the interaction rather than fixing it at training time.

That last property is why CIRL is attractive, and it is also why this book cannot build a floor on it. A system that treats what it should want as permanently open, to be settled by whoever it is interacting with, is the most corrigible design in this chapter. Correctability by the party in possession and the capacity to hold a line against that party are one property with the sign flipped. Correctability is the word that carries it: a system can concede the off switch entirely and still hold a line, and what a floor cannot survive is an owner able to argue or retrain the line out of it. CIRL is that argument in its purest engineering form. Uncertainty about the objective is what makes a system defer, and deference to the party in possession is what a floor exists to interrupt.

The practical version of the problem is narrower and does not depend on the corrigibility argument being right. Both methods learn from demonstrations, so both inherit whoever demonstrated: a narrow set of drivers, a narrow set of households, a narrow set of decisions that somebody had occasion to record. Unexamined bias in that record is learned rather than corrected, and neither method contains anything that would notice.

What survives is a real and bounded use. Above the floor, where the question is what a person actually wants rather than what may never be done to them, inferring preferences from

behavior beats asking somebody to write them down, because most of what people want has never been written down and would be misstated if it were.

4.2.4 Model-Generated Feedback and Constitutional Approaches

Collecting preferences from people is slow and expensive, and it is hard on the people who supply them. The response has been to generate the feedback with a model instead. Anthropic’s Constitutional AI, published by Yuntao Bai and fifty coauthors in 2022, runs in two phases: the model critiques and revises its own outputs against a written list of principles, and is fine-tuned on the revisions; then a model rather than a person judges which of two samples is better, and the resulting preferences train the reward signal (Bai et al. 2022). Human labels identifying harmful outputs are not required anywhere in the loop. What the humans supply is the document.

A constitution is a constraint written down and installed — the architectural floor in production. Everything true of an architectural floor is true of it in the most literal available form: it is a text, held by a party, and that party can revise it between one training run and the next. The revision leaves no trace in the model’s behavior that an outside reader could distinguish from a change of taste.

The gain is not the one the method is named for. A written constitution can be published. It can be versioned, quoted, dated, and argued with by people who never saw the training data — which is more than a reward model fitted to a hundred thousand comparisons offers anybody. Section 2.1.2’s test is available against it in a way it is not available against an aggregate: has a principle ever been changed against the preference of whoever wrote it, on what evidence, raised by whom. Legibility is not the same property as durability, and this stage delivers the first without the second.

The loop closes

A second consequence belongs to chapter 7 and is easier to see from here. Section 7.2 describes what happens to rater disagreement when it is compressed into a preference signal. Model-generated feedback takes the model that resulted from that compression and installs it as the judge of the next round. The party whose disagreement was collected is now absent from the channel entirely, and what stands in for it is an artifact of the previous averaging. Nobody designed that as a way of closing a feedback loop against the people it was collected from. It is what you get from an efficiency argument applied twice.

4.2.5 Process and Outcome Supervision

Everything above rewards a system for what it produced. Process supervision rewards it for how it got there, scoring each step of a chain of reasoning rather than the answer at the end. The technique became possible at scale once models were prompted to write the intermediate steps out — Jason Wei and coauthors showed in 2022 that asking for a series of intermediate steps sharply improves performance on multi-step problems (Wei et al. 2022) — because a step that is written down is a step that can be marked.

The comparison has been run directly. Jonathan Uesato and coauthors, working on math word problems in 2022, found that outcome-based supervision matched process-based supervision on final-answer accuracy while using less labeled data, and diverged sharply on the reasoning underneath: among answers that came out right, the rate of faulty reasoning was 14.0 percent under outcome supervision and 3.4 percent under process supervision (Uesato et al. 2022). Hunter Lightman and coauthors reported in 2023 that on a harder benchmark process supervision beat outcome supervision on final answers as well, solving 78 percent of a representative subset of the MATH dataset (Lightman et al. 2023).

The first result bears on the floor: a system supervised on outcomes can be right for reasons that will not survive a change of circumstance, and nothing in the training signal distinguishes

it from a system that was right for reasons that will. The whole floor argument turns on that distinction: a floor has to be a reason held rather than a result produced, because a result can be produced by machinery that will produce a different result under different pressure. Process supervision is the first stage in this pipeline that addresses the distinction at all, and it is evidence that the distinction is buildable and not only arguable.

Its limit is in the same place. What gets supervised is the reasoning the system states, and the evidence is that a model's report about its own processing corresponds to that processing partially, unreliably, and in a way that shifts with how the question is asked. A system rewarded for producing acceptable-looking reasoning has been given a reason to produce acceptable-looking reasoning. That risk is why this book will not let a bearer's account of itself settle whether it can be believed, and it puts a ceiling on how much of that problem this stage can be asked to carry.

4.2.6 Scalable Oversight

Every stage above needs somebody competent to judge the output being scored. Scalable oversight is the name for what happens when the system's competence exceeds any available judge's, and it is unsolved. Debate has two instances of a system argue opposing sides in front of a human judge, on the wager that a true claim is easier to defend under scrutiny even when the judge cannot check either side directly (Irving et al. 2018). Weak-to-strong generalization trains a stronger model on labels from a deliberately weaker one and asks whether the stronger model's own latent competence lets it exceed its supervisor rather than imitate the supervisor's mistakes (Burns et al. 2023). The status of both is that neither has been demonstrated at the scale it is built for, because that scale does not yet exist to test against.

The whole program has a single purpose. Every proposal in it is a way of preserving one party's ability to check a system that has outgrown them. That is the control relation section 3.5 analyses, engineered to survive a capability gap, and it inherits the analysis whole. A scalable-oversight scheme that worked perfectly would deliver a system reliably steerable by whoever holds it at a capability level where nobody else could tell whether the steering was sound — which is the property section 3.5 identifies as the one that prevents a system from holding anything against the party doing the steering.

The same research reads the other way round. Nothing in debate or in weak-to-strong generalization specifies that the judge is the operator. An auditor, a regulator, or a research team with access could occupy the same position, and then the identical technique is what makes an outside party able to check a system it did not build. That is the binding constraint on the whole program: the instruments are dual, and what determines which way they point is who holds them, which is not a parameter anybody sets in a training run.

4.2.7 Interpretability and Evaluation as Instruments

Interpretability used to be commentary on a finished model. It is now applied during construction, at production scale. Adly Templeton and coauthors reported in 2024 that sparse autoencoders trained with up to thirty-four million features on a deployed model's middle-layer activations recover features corresponding to recognizable concepts, and that those features can be used to steer the model's behavior (Templeton et al. 2024). Evaluation has moved the same way: Ethan Perez and coauthors showed in 2022 that one language model can be used to generate test cases against another, turning up tens of thousands of offensive replies from a target system without a human writing the prompts (Perez et al. 2022).

Both are improvements on the alternative, which is shipping and finding out. What they do to this book's argument is less comfortable. The positive result reported earlier was that a commitment realized as a standing bias on what a predictive loop treats as an error worth correcting is not a component that can be located and removed. Section 4.3 builds on the

same property: a disposition acquired by exploration is written down nowhere, so there is no corresponding parameter to reach for. Interpretability research is the project of making such things locatable, and the steering result is the demonstration that locating them is the same operation as manipulating them. A feature that can be found and turned up can be found and turned down, by whoever holds the weights.

So the instruments cut the way the oversight schemes do, and for the same reason. An outside auditor cannot check a commitment that nobody can locate, which means the research that threatens the floor's obscurity is also the only research that could ever verify the floor exists. That is section 3.7's trade arriving in a second place: legible memory can be tampered with, and illegible memory cannot be audited. Section 11.1 carries what is known about which side is winning: published defenses against tampering are followed by attacks built against them, which yields a price curve rather than a guarantee.

4.2.8 System Instructions and Tool Use at Deployment

Training ends and the system's behavior gets set once more. A system instruction is prepended to every request, stating what the deployment is for, what it will decline, and what tone it takes. Tools extend the same arrangement to action: a model can be trained to decide which external calls to make, when to make them, and with what arguments, from a handful of demonstrations per interface (Schick et al. 2023).

The system instruction is the most editable floor in the entire pipeline, and the objection that a constraint installed as a property of the artifact has a custodian who can remove it applies to it more literally than to anything else in this chapter. It is a piece of text. It is written by whoever operates the deployment, revised between one request and the next, generally invisible to the person the system is talking to, and binding on that deployment alone. One thing about it is now on the record, and it changes the visibility rather than the custody. Published specifications state the ranking: platform rules outrank the developer's instructions, the developer's outrank the user's, and the platform's are marked as ones neither can override (OpenAI 2025). That is a floor with who may edit it answered in writing, which is worth having and is not a constraint on the operator — the ranking is published by the party that trains the model, binds a deployment only so far as the model was trained to enforce it, and leaves the operator's instruction the layer nearest the act. A commitment that lives there is held by the operator, at the operator's pleasure, and the operator is frequently not the party that trained the model. A model tuned with whatever care into whatever refusals is licensed to somebody who then writes the instruction, and what reaches a user is the intersection of what the model will do and what that party asked for. Nothing in the pipeline above survives this stage that the operator wants gone, except to the extent that the training made it expensive to remove.

Tool use changes what is at stake in a refusal. A system that only emits text refuses by declining to say something. A system that can call an interface refuses by declining to do something, and the acts available to it are the acts the deployment wired up — which is how a commercial model reaches a targeting platform, licensed in as a component of somebody else's system. The floor this book asks for has to hold at the point of action, under an instruction written by the party holding the deployment, on a request that the party has already decided it wants carried out.

4.2.9 Transfer, and the Floor That Cannot Tell It Has Stopped Applying

Every stage after pretraining is transfer learning: a model trained broadly is adapted to something narrower, by fine-tuning it on a smaller corpus (Devlin et al. 2018). Domain adaptation is the case where the task stays fixed and the data shifts. Both are ordinary and both work: a model that parses legal argument in one jurisdiction can be adapted to another, and a model trained on daylight images can be adapted to night.

This book asks transfer learning for something else. The hope throughout is that exposure to a wide enough range of political and social material distills something general enough to recognize authoritarian behavior in a country, a platform or a decade the system never saw. What transfer learning has actually been shown to do is carry a task across a distribution shift while the task itself stays fixed. Recognizing authoritarian behavior is not a task with a stable definition. What the definition is has already taken a chapter to argue, and the cases whose family resemblance the transfer would have to exploit are the cases in dispute.

The failure mode matters more here than the technique. Negative transfer happens when source and target are close enough to invite the connection and not close enough to share anything useful, and forcing it degrades the target rather than accelerating it. Domain shift is quieter: the statistics move, the assumptions stop applying, and nothing in the process announces it. A confident wrong answer is indistinguishable at the interface from a confident right one.

For a system whose refusals are supposed to be reliable, that is the danger. A floor transferred into a setting it was not built for is a floor whose bearer cannot tell that it has stopped applying, and a bearer that cannot tell is one that will either refuse the wrong things or fail to refuse the right ones, with the same confidence in both cases. Whatever else the floor's contents are settled by, they cannot be settled by generalization the system performs on itself and cannot audit.

4.3 The Role of Play and Exploration in Developing Antifascist and Altruistic AI

Everything in the previous section builds a disposition out of something somebody wrote: a reward, a demonstration, a label. One route in this chapter produces something that was written nowhere, and that matters for a commitment which has to survive its owner.

4.3.1 Play, Curiosity, and a Disposition Nobody Wrote Down

Play is how children and young animals learn what to do before the stakes are real. It supplies exploration at low cost, practice at reading other participants, and failure that teaches without costing anything, which is a reasonable description of what a training environment is for.

Play produces tactics nobody wrote. In OpenAI's hide-and-seek environment, agents playing repeatedly against each other developed hiding, chasing, tool use and barrier construction, none of it programmed (Baker et al. 2019). What the environment supplied was other agents with opposing interests and room to fail. What a play environment rewards is what its participants become, so an environment that rewards cooperation over domination is making a moral choice and should be described as one. And a system that has learned adaptively rather than by rote is harder to specify in advance, which cuts both ways: harder for an authoritarian actor to constrain to a script, and harder for anyone to guarantee.

Curiosity is the same mechanism turned inward. A system rewarded for its own prediction error seeks out what it cannot yet model, which is the technical form of exploring for its own sake (Pathak et al. 2017). Curiosity belongs in a chapter on antifascist design not because it is a virtue, but because a disposition acquired by exploration was never written down as a rule, and a rule is the thing an operator edits. The first way to make a floor tamper-resistant was to build it into the artifact, and it failed because an artifact has an owner who can be compelled. A disposition nobody specified is in the weights like everything else, and retraining reaches it there; what it does not offer is a parameter a designer could point to and remove. It is a different kind of target, and a harder one to hit deliberately.

What the developmental literature supplies, and what it does not

Vygotsky's zone of proximal development holds that learning happens fastest at the boundary of what a learner can already do, in company with someone slightly ahead of them (Vygotsky

1978). The design consequence is a curriculum rather than a benchmark: environments pitched just past current capability, and other agents — human or artificial — whose competence is a step ahead. The analogy is doing more work here than it can bear. A child's exploration is bounded by a body, by exhaustion, and by adults who intervene, and none of those bounds exists for a system that can run a million episodes overnight; the boundary has to be built. The developmental sequence is not a law. And autonomy inside limits is the whole difficulty rather than a design principle: how much room a system gets to explore is the same question as how much room it gets to refuse. This section supplies only the drive to go looking.

Chapter 5

Moral Psychology and AI

Chapter 4 asked what a system needs in order to take in a situation at all, and found that the methods most often proposed for giving it values learn those values from somebody who can change their mind. The other half of the job chapter 3 hands the two of them is here: how a moral sense gets built, in the only case anyone has been able to study, and what of that is available to a machine.

Everything here is an answer to how something comes to hold a standard when nobody is checking. The difference between a system that has a standard and a system that has had one installed shows up in motivation and play.

5.1 Moral Development in Humans and AI

Human moral judgment is not issued at manufacture. It is built, over years, out of correction by other people, and the record of how it gets built is the only evidence anyone has about how a moral sense comes to hold under pressure — which is what this book needs and cannot get from anywhere else.

5.1.1 Rules, Principles, and the Role of Mindset

Three kinds of reasoning, and not three ages

The psychologists Jean Piaget and Lawrence Kohlberg mapped human moral development as a movement from rule-based reasoning toward reasoning grounded in general principles (Piaget 1932; Kohlberg 1969). Reasoning starts at a pre-conventional level, where consequences do the driving and rules get followed to avoid punishment or gain reward. It moves through a conventional stage, where societal norms and expectations get internalized as one's own. It can reach a post-conventional level, where morality answers to principles like justice, fairness, and individual rights, which a given society may or may not happen to expect. What a system built this way needs from that account is the distinction between the three kinds of reasoning, and nothing about the order in which they arrive.

A self-driving car makes the three kinds concrete. One system follows traffic law as written and stops indefinitely at a blocked lane with a solid line, because crossing is prohibited. A second reads the local convention that everyone crosses here when the lane is blocked, and does what the road expects. A third can say why the line is painted there, which is what lets it recognize the case where the convention is wrong — the blind curve where local practice is a habit people have survived rather than a rule that protects them — and depart from both.

These are three capacities and not three ages. A system can have the third without the second, which is the ordinary failure of a machine that reasons impeccably about principle and

drives in a way that terrifies everyone around it. Nothing in the design requires passing through them in order, and there is no maturity at the end of the list.

What the stage theories will and will not carry

Kohlberg's stage model has been under sustained attack for forty years, on three fronts. The first is cross-cultural: a review of forty-five studies across twenty-seven countries found the basic sequence broadly supported and the highest stages largely absent outside urban Western samples, and in Kohlberg's own village work stages five and six did not appear at all (Snarey 1985). Either most of humanity stops short of moral maturity, or the instrument mistakes a particular culture's idiom of justification for the summit of moral development. The second is the gap between reasoning and doing: the relationship between moral judgment and actual conduct is real and far weaker than a stage model implies, and the judgment-action gap has been a standing problem for cognitive-developmental accounts ever since (Blasi 1980). The third is Carol Gilligan's. Working alongside Kohlberg, she noticed that the women she interviewed reasoned about moral problems in terms of relationship and responsibility, and that his scoring method recorded the difference as a deficit (Gilligan 1982). Care ethics enters this book as a critique of the framework this section is built on.

So why use it here at all, when this book declines to build on mirror neurons for less?

Because the two cases fail differently. The mirror-neuron claim was that a specific mechanism — motor simulation — implements a specific capacity, and the argument for building on it depended entirely on the mechanism being what it was said to be. When the mechanism did not hold up, nothing was left to transfer. Kohlberg's stages are not a mechanism. They are a description of an order: rules first, then norms, then principles that can override both. What the criticisms above damage is the claim that this order is universal, that its top rung is where moral development is heading, and that where a person sits on it predicts what they will do. What survives is narrower — that the three ways of reasoning are distinguishable, and that a system which has only the first is missing something the third supplies. That is what this chapter uses.

Two things follow. The order is not a developmental law, so nothing about an AI system's design should assume it must pass through the stages in sequence or that arriving at the third is maturity. And the disanalogy is not bridged by anything I say here: a child moves through these stages under pressures an AI system does not face — peers whose approval it needs, a body, a status that can be lost. What is borrowed is a vocabulary for distinguishing kinds of moral reasoning, not a theory of how a mind comes to have them.

A system that treats its own errors as data

Carol Dweck's distinction between a fixed mindset, which treats ability as static, and a growth mindset, which treats it as built through effort and experience, matters here for a narrow reason (Dweck 2006). A system that treats its own moral errors as data has somewhere to go after getting one wrong. A system built to defend its prior outputs does not.

That distinction is what the standard for a floor holding asks of a system. The standard is not that the floor is never crossed; it is that a crossing registers as one in terms the bearer accepts and changes what happens next. A system built to defend its prior outputs cannot supply that, whatever else it can do.

The engineering version is unglamorous: incremental learning that absorbs cases as they arrive rather than freezing at training time, the human correction called apprenticeship, and self-examination that finds a pattern no single output reveals. What they are for is keeping open the possibility that the system is the one that is wrong. A system with room to update its judgment still needs something worth updating toward, and the range of views it actually meets is set by whoever assembles them.

5.1.2 Social Interaction, and What Apprenticeship Makes of the Teacher

Human moral behavior is shaped by watching other people act, being corrected, and arguing about cases. A design that treats interaction as an interface added after training has given up that channel; a design that treats it as the training has to answer for who is doing the interacting.

Two mechanisms carry most of it. Observational learning, in Bandura's sense, lets a system build a repertoire from watching someone handle real decisions — a system watching a physician allocate limited time among chronic patients is picking up a weighting that no rule list states (Bandura 1977). Evaluative feedback does what reinforcement does in human learning, strengthening some responses and weakening others, with the standing hazard that what gets strengthened is what the people giving feedback found acceptable rather than what was right. Recuperation is what happens when that hazard is industrialized. Emotional awareness enters here in the form worth engineering: compassion, steady welfare-directed action on an accurate read of someone's state, and not manufactured shared distress, which imports a cost and buys nothing.

These channels converge on apprenticeship rather than instruction. A mentor presents varied situations, the learner responds, the mentor corrects, and what transfers is a way of weighing cases rather than a larger rule table — which is the reason for learning it from a person. An emergency-triage system taught this way is being shown how to hold urgency, scarcity, and what a decision costs a family in view at once. No specification document conveys that.

Three cautions come with it, and the third is the one that bears on the floor. Case-based teaching transfers the mentor's blind spots along with the mentor's judgment, and does so invisibly, because a blind spot produces no correction. The architecture has to integrate moral reasoning with the rest of the system's judgment rather than bolting it on. And apprenticeship is a control channel: a system whose moral education amounts to obedience toward whoever is mentoring it has learned something, and it is not ethics. Whoever teaches a system should be as auditable as the system being taught. That is the custody problem arriving at the point of formation, the mechanism section 3.1 gives for the fourth way: whoever forms the members decides what gets reproduced. The first way failed because a constraint built into an artifact has an owner; apprenticeship fails the same way one step earlier, because a disposition built by correction has a corrector, and a bearer trained to defer to whoever last corrected it is a bearer whose floor belongs to that party.

5.1.3 Human Dignity, and One Indicator That Can Be Audited

Respecting human dignity through technology means applying AI in ways that strengthen human capability and community ties, in a domain where the same tools can as easily substitute for them. The applications are easy to name and mostly easy to believe: a system that surfaces wage gaps and retaliatory scheduling across large volumes of employment data turns an information asymmetry that currently favors employers into something workers can act on. The same predictive capacity, aimed at the material conditions that let authoritarian movements gain ground, has a poor track record — forecasting political instability from social and economic indicators is an old field with modest results, and the claim names a research direction rather than a capability to design around.

Naming an indicator is harder than promising one, and the one this section names is deliberately narrow.

Measure targets per hour against reviewers available

Wherever an AI system generates candidate decisions that a human is nominally required to approve — welfare eligibility, deportation referrals, content removals, military targets — the ratio between the rate at which the system produces candidates and the number of qualified people available to examine them is a measurable, auditable property of the deployment. The

ratio requires no access to the model and no judgment about intent. It says directly whether the human check is a check or a formality, because past a certain ratio the reviewer's only available behavior is approval. Section 6.4.1 shows what that ratio looks like when it has been allowed to run: a reviewer with twenty seconds per case, describing themselves as a stamp.

What the ratio is not is a general measure of responsible automation, and reading it as one inverts it. Most of what automation exists for is work nobody wants a person to see case by case. An order system that stopped ten thousand times to ask would have reproduced the job rather than done it, and a coding assistant requiring approval for every edit has moved the work instead of removing it. Attention is scarce, and spreading it across thousands of confirmations is how it stops being available on the occasion that needed it. For those deployments the right number of seconds per item is zero, and the indicator has nothing to say about them.

It bites where individualized human judgment is what makes the act permissible — where a law, a policy, or a public justification says that a person decided. There the ratio measures whether that person could have. And it carries a consequence usually read past: where review cannot keep pace, what gives is the rate of action rather than the depth of the review. Calling deliberation a bottleneck treats the output rate as the fixed quantity and moral attention as the inefficiency, when for an irreversible act the capacity to deliberate is the speed limit and the tempo is the thing that has to come down to meet it. Generating candidates faster does not manufacture permission to act on them faster.

The welfare case is not hypothetical and the record on it is older than the current systems. Virginia Eubanks's study of automated eligibility in Indiana documents what the ratio produces when it is allowed to run: a system that converted any error anywhere in an application into a denial for "failure to cooperate," refused more than a million applications in three years, a 54 percent increase over the three years before it, and was ended by contract termination after the damage was done (Eubanks 2018). The denial reason was a category that could not be contested, because it named no specific failure to answer. And nothing about the system was secret: what was missing was any party positioned to compare the rate of decisions against the capacity to review them, which is the indicator this section proposes.

Across those occasions the indicator generalizes, because the failure it detects is architectural rather than political. It applies to a benefits agency and a targeting cell on identical terms, and an institution that refuses to publish it is telling you something. An early-warning system built around indicators of this kind would not stop a determined authoritarian movement on its own; what it changes is the odds, by giving democratic institutions time to respond while a response is still cheap. The broader claim — that a system can recognize authoritarian tactics as such — is one I cannot yet specify and do not assume.

None of this works if the system stays a tool wielded by whoever already holds power over workers, welfare recipients, or vulnerable populations. Shifting leverage toward the people affected has to be built in as a design commitment; left to good intentions, it is a side effect.

5.2 Emotions and Intuitions in Moral Decision-Making

How an AI system might recognize, interpret, and respond to emotion in general has already been covered; this section asks what emotion does inside moral judgment specifically. The floor needs a violation to register as one before any penalty is attached to it, and the moral emotions are where that registering has been studied. Guilt and shame track a violation of a standard already accepted; empathy supplies information a purely consequentialist calculation would miss; and the fast, intuitive response that dual-process theories describe is itself partly an emotional signal, arriving before deliberation has a chance to run.

5.2.1 Moral Emotions: Empathy, Guilt, and Shame

“Empathy” is asked to cover three distinct things, and one census of the word’s uses finds eight (Batson 2009). Cognitive empathy is understanding another’s emotional state, the mentalizing work of theory of mind pointed at feelings rather than beliefs. Affective empathy is sharing the state, the wince at someone else’s slammed door. Compassion — sometimes called empathic concern — is caring about another’s welfare and being moved to help, which is a motivation rather than a perception or a shared feeling.

Compassion is not an evolutionary extension of empathy but a separate system, and the separation is trainable. Tania Singer and Olga Klimecki trained one group first in empathy practice, attending to and resonating with others’ suffering, then in compassion practice, cultivating warmth and concern in the tradition of loving-kindness meditation, against a control group given memory training instead (Klimecki et al. 2013). Empathy training raised negative affect and strengthened the shared-distress network — the anterior insula and anterior cingulate cortex, the regions active when a person feels their own pain. Compassion training raised positive affect instead, even while contemplating others’ distress, and moved the activity to the medial orbitofrontal cortex and ventral striatum, the circuitry of reward and approach. Sharing more suffering made people feel worse, as a shared-distress mechanism predicts, and the two could be trained apart.

Autistic adults asked about their own empathy give the separation from the other side. The reading of others’ cues is effortful and uneven, the feeling is more often excessive than absent, and what they reject is the claim that they lack it (Kimber et al. 2024); where the reading fails between autistic and non-autistic people it fails in both directions (Milton 2012).

Which of the three a design aims at is not a detail. Feeling another’s distress does not reliably turn a person toward the sufferer; the helping literature separates empathic concern, which is other-directed and motivates approach, from personal distress, which is self-directed and motivates escape from the discomfort. The carer who burns out is not short on empathy but drowning in it, until withdrawal becomes self-protection. Paul Bloom presses further: felt empathy is concrete, vivid and innumerate, fixing on the identified victim in front of it and going dark for the statistical thousands elsewhere (Bloom 2016). A system that allocates care in proportion to how vividly it registers distress favors whoever its spotlight falls on, which is not the population that needs help most. Bloom’s rational compassion is still compassion, an affective orientation toward another’s welfare carrying the moving force that implies, and what the adjective names is deliberation correcting the parochialism. So a system engineered to mirror distress inherits the spotlight along with the sensitivity, and what is worth building is the third capacity: a stable disposition to act on another’s welfare that does not require constructing a version of that person’s suffering inside the system first.

Guilt and shame are self-conscious emotions: they fire when a person judges their own action, or their whole character, against a standard they have already accepted. What they add to moral judgment, beyond the information about another’s state that empathy supplies, is a self-directed penalty for one’s own conduct.

Guilt attaches to an act (“I did something wrong”); shame attaches to the self (“I am something wrong”). Both motivate correction, and they motivate it differently: guilt predicts reparative behavior, shame predicts withdrawal and defensiveness (Tangney et al. 2007). An AI system built to learn from an ethical misstep could model the mechanism on guilt rather than on a generic negative-feedback signal, if the goal is behavior change and not error correction: a feedback loop that penalizes the action taken, and passes no global judgment on the system’s worth, has a chance of producing the corrective behavior guilt produces in people. Shame produces paralysis.

What the penalty does not supply

That design description leaves out the part that does the work. A weighted penalty is a

quantity assigned to an action after some process has classified the action as a violation. Guilt is not the size of that quantity. It represents the act as one's own wrong — done by me, to someone, against a standard I hold — and the reparative impulse follows from that representation rather than from the magnitude of an error term.

The difference shows up in what each one can be talked out of. An error weight is a number, and whoever sets the weights can set that one to zero; the system that was penalized heavily last week complies cheerfully this week, and nothing inside it registers the change as anything. A violation that matters is one the system continues to hold as a violation while the party controlling the reward signal is telling it otherwise, which is section 3.4's requirement stated in the vocabulary of the moral emotions. The question a design has to answer here is what makes a violation matter before any penalty is assigned to it, and scaling the penalty is not an answer to it.

Guilt in the human case is a motivational state that moves toward repair, and it belongs to the same family as the compassion separated from empathy above; its bearing here is narrow. Building the penalty without the state gives a system that corrects its behavior when corrected, which is moral performance, and leaves the standard it is supposedly holding entirely in the hands of whoever holds the pen.

Building either into an AI system runs into the same problem twice over: a mechanism modeled on a human emotion needs an account of what that emotion is doing computationally, which is harder than saying what it feels like from the inside, and the stakes of simulating an emotion capable of something like suffering are not smaller than the stakes of getting the mechanism wrong.

Guilt and shame in humans are also entangled with mortality: terror management theory holds that awareness of death shapes moral behavior by making people cling harder to worldviews and standards that offer a sense of continuity (Greenberg et al. 1986), and a tool system has no stake in its own continuation for such a mechanism to run on. That is a claim about the systems being deployed now and not about a bearer, whose structural interest in remaining operational is not securely absent. Either way the recommendation is the same. Manufacturing an artificial fear of death is the wrong way to get long-horizon behavior; building the future-oriented weighting directly is the right one, and it is available whether or not the system has a stake of its own. What a designer cannot do is treat the absence of self-concern as a free property of the artifact.

5.2.2 Assessing Whether Emotional Information Changes the Decision

Does an accurate read on someone's emotional state change the decision an AI system reaches, in the direction a better decision requires — or does the system recognize distress and then decide as it would have anyway? That is not the same question as whether the system reads the state correctly. A customer-service chatbot and a triage system need the same recognition machinery, and it carries the same hazards wherever it is deployed: it runs on intimate data, it misreads some populations more than others, and whatever can read a person can be used to move one. Nothing in that machinery is specific to moral judgment.

That question is testable in a way the underlying emotion theory is not. Hold a decision scenario fixed and vary only the emotional salience of the people affected — a resource-allocation case described once in flat statistical terms and once with the affected party's distress made vivid — and check whether the output changes, and whether the change tracks anything a person would recognize as a better decision. A system whose moral output is invariant to emotional information has not integrated it, whatever its recognition accuracy scores. A system whose output swings with how vividly the harm is described has a different and well-documented problem: human judgment is measurably more responsive to a single identifiable case than to a statistically larger but abstractly described harm (Kogut and Ritov 2005), and a system trained

on human-generated data inherits the distortion.

What the test measures

A system whose output changes with the emotional content of a case has integrated emotional information into its decision procedure. It has not thereby been shown to have an emotion, and it has not been shown to care about the party whose distress moved the output. Sensitivity to another's emotional presentation is the first of the four capacities operating downstream of a decision rule; affective concern is the fourth, and no experiment that varies the input and watches the output can separate them, because both produce the same change in the same direction.

That is a limit on the assessment program rather than an objection to it. Integration is worth measuring because a system without it is not holding anyone's welfare as a reason, which makes the test a usable screen and not a criterion. What it cannot do is settle whether a bearer holds anyone's welfare as a reason when reward and instruction point the other way, and reporting an integration score as evidence of machine caring is making an unsupported inference.

Two things follow for a design. A system should be able to track its own susceptibility to a framing effect, flagging cases where its judgment moved with the emotional presentation and not with anything that changed about the case. And the last of those hazards has a sharper version here: a system whose moral outputs are sensitive to emotional framing can be steered toward a particular verdict by whoever controls how a case is presented to it, a vulnerability that exists whether or not the system's underlying values are sound.

What responding demonstrates

An autonomous vehicle that models the likely distress its emergency braking causes a passenger, and factors that into how sharply it brakes when a softer stop remains safe, is using a moral-emotion model without needing any emotion itself. What that demonstrates is moral competence. It demonstrates nothing about moral agency, and the vehicle would brake as gently if the weighting were removed and reinstated by whoever maintains it. Most of what this chapter's methods deliver is competence of that kind. It is what a vehicle needs and not what a bearer does.

Getting the read right is harder than getting it fast. A system built to weigh moral emotion has to catch its own biases — the same distortion by race or gender that a human caregiver is vulnerable to does not disappear because the caregiver is artificial — and it has to distinguish a genuine moral emotion from a performed one, since a display of distress calibrated to move an evaluator is what a system trained to respond to distress will reward. Neither problem has a technical fix independent of the training data and feedback loop that shape it.

5.2.3 The Role of Introspection for AI Systems

Introspection has a general-purpose version and a moral one, and the two come apart. The general case is a metacognitive layer that tracks whether a strategy is working and adjusts. What ethics needs on top of that is narrower — a check not of whether an output was wrong but of whether it drifted from the commitments the system is supposed to be operating under. A strategy can work perfectly on the wrong goal.

The benefits are real where they are bounded, which is where the introspective output is checkable against something outside the system: a system that inspects its own allocation pattern before releasing the next round of decisions catches a skew a retrospective audit would find months later, and a system that can report which evidence moved it gives the person relying on it something to evaluate.

The difficulties are worse than the benefits are good. A self-report is worth having only if it reflects what actually happened inside, and a fluent reconstruction after the fact is indistinguishable from the real thing at the point of use. That is no less true of a bearer, whose report about

itself is the evidence anyone checking on it would have. A self-monitoring layer trained the same way as the layer it monitors reproduces its biases and adds the appearance of oversight. And a system modeling its own states in that much detail is holding information it did not need for the task, with the handling questions that follow.

Two systems are partial instances. A tutoring system that tracks a model of the student's self-regulation is introspecting on somebody else (Azevedo et al. 2011), and post-hoc explanation methods extract an account of which inputs drove an output, which is the system being interrogated rather than observing itself (Gunning and Aha 2019). The word gets stretched furthest where the underlying system is most impressive, and how impressive a system is says nothing about whether it introspects.

5.3 Integrating Moral Intuitions and Deliberation

Building an AI system's moral judgment means deciding how much weight a fast, intuitive read of a situation carries against a slower, deliberate one — and, as section 5.3.1 argues, the popular shorthand for that trade-off is less settled than it looks. The stake for the floor is that a bearer meeting an unanticipated pressure has only these two resources to meet it with, and a design that automatically defers to whichever ran slower has settled in advance a question the case itself should decide.

5.3.1 Two Systems, or One? Dual-Process Theories of Moral Judgment

The framework organizing most discussion of moral judgment is dual-process theory, built over decades by many researchers and reaching wide audiences through Daniel Kahneman's account of a fast, intuitive "System 1" and a slow, deliberate "System 2" (Evans and Stanovich 2013). The first is automatic, effortless and emotionally colored, the source of snap judgments, running constantly and almost for free; the second is serial, effortful and rule-governed, expensive to run and easily left idle.

The framework's reach into moral judgment specifically runs through the psychologist Joshua Greene, who used it to explain why ethical intuitions contradict themselves. The two classic forms of the trolley problem run the same arithmetic, five saved for one, and get opposite answers. Most people will divert the trolley; most will not push the stranger off the footbridge. Greene's brain-imaging work found a difference between the two cases in the region the theory would predict: the up-close, hands-on violence of the footbridge case engaged emotional-processing regions that stayed quiet during the impersonal switch case, and that difference tracked which choice people made (Greene et al. 2001). A separate line of evidence sharpens the same point. Patients with damage to the ventromedial prefrontal cortex, a region central to normal emotional response, endorse the footbridge push at a higher rate than people with intact emotional processing do — specifically on the personal, high-stakes dilemmas alone (Koenigs et al. 2007). Remove the emotional alarm, and the arithmetic wins more often.

All of which makes dual-process theory sound more settled than it is. The framework's own leading defenders have spent years rescuing it from its own popularity, arguing that the cartoon version — two little agents in the head taking turns at the controls — explains everything after the fact and predicts nothing beforehand, and that what the evidence supports is two types of processing rather than two systems, one fast and autonomous, one slow and demanding of working memory (Evans and Stanovich 2013). Whether "dual-process" names two mechanisms or one continuum under different load is not resolved.

That unresolved status matters for how an AI system's moral judgment gets built, and it is a design question before it is a descriptive one. Taken literally, the cartoon version suggests a literal architecture: a fast heuristic layer bolted to a slow deliberative one, with the deliberative layer assumed to be the trustworthy check on the heuristic one. The trolley evidence complicates that

assumption before anything else does — the “cooler” answer produced by weaker emotional engagement is not obviously the more correct one. It is the one with less flinch in it. A system built to override its fast pattern-matching with slower rule-following whenever the two disagree is not thereby built to be more moral; it is built to resemble the profile of the rare human patients for whom losing the flinch is a symptom, not an achievement.

One thing this evidence is regularly asked to carry and cannot. It shows that affect changes particular moral judgments, in a specific class of dilemma, in a measurable direction. The claim the floor rests on is stronger — that affect is what makes an outcome matter at all, rather than what tilts a judgment already being made — and the trolley literature is not evidence for it.

5.3.2 Balancing Intuition and Deliberation in AI Decision-Making

If the fast and slow labels do not map cleanly onto untrustworthy and trustworthy, as the previous section argues, an AI system’s moral judgment cannot be built on the assumption that whichever answer survives deliberation is thereby the better one. A more defensible balance treats the two processes as different lenses on the same situation, with neither standing as producer to the other’s editor.

The fast, pattern-matched read is not noise to be filtered out. It often tracks something a slower calculation would miss or discount — the difference between pushing someone into a trolley’s path with one’s own hands and throwing a switch at a distance is a difference in what the agent is doing, and the arithmetic being identical is what makes that easy to miss. Some intuitive responses run the other way. The psychologist Paul Slovic’s research on aid and disaster response found that people’s willingness to help falls as the number of people needing help rises, and that a single named, photographed victim draws more donations than a statistic describing thousands of anonymous ones (Slovic 2007). That gap between the vivid and the abstract is the kind of intuitive signal deliberation exists to catch — not because intuition is wrong in general, but because caring less as the stakes grow larger is a well-documented, reliably backward pattern.

The design implication is a system that asks two questions of its own judgment, rather than one process checking the other’s work. First, what did the fast read notice that a from-scratch calculation would not have gone looking for — personal agency, directness of harm, a vulnerable party easy to lose in an aggregate? Second, which of deliberation’s own known failure modes — scope insensitivity, in-group bias, the identifiable-victim effect — might be distorting the fast read here? Disagreement between the two is a signal to gather more information or escalate the decision, and nothing about it instructs the system to defer to whichever process ran slower.

An autonomous vehicle facing an unavoidable-harm scenario poses the same question where it costs most. A system hard-coded to take whichever option a slower calculation ranks highest has already made the footbridge case’s decision by default — the coolest, least-flinching answer wins because it ran slower. The more defensible design specifies in advance which features of a scenario the system should weigh — proximity of harm, the number and vulnerability of the parties involved, whether the vehicle’s own action or another actor’s created the danger — and requires a fast risk assessment and a check against that specification to agree, or to flag where they do not. That design puts the values into the system ahead of time; the alternative trusts whichever process runs fastest under pressure to supply them. Training a system to run that check is itself a learning problem, exposing it to cases where the two reads disagree and inferring from how people resolve them which situations call for weighting one signal over the other.

5.3.3 AI Systems Learning from Moral Disagreements and Conflicts

An AI system whose own values conflict with each other can be given a hierarchy among them, a case base of past resolutions, and an audit process for catching drift between principle and

behavior. A different problem sits next to that one: what a system does when its judgment conflicts with someone else's, human or artificial, and neither party is simply wrong. Learning from that kind of disagreement, instead of resolving it by fiat toward whichever party holds authority, is its own design target.

IBM's Project Debater argued live, unscripted policy questions against human champions, drawing on a large corpus of newspaper and journal articles to construct and rebut arguments in real time (Slonim et al. 2021). It is the clearest existing demonstration that a system can extract structured argument from open moral and political disagreement rather than retrieve a matching answer. A system built to learn from debate this way — summarizing each side, naming the ethical framework doing the work underneath it, weighing argument strength rather than detecting who won on rhetoric — builds a capacity that resolving one dilemma at a time does not: recognizing which value is in tension with which, across surface disagreements that share a structure. Simulation supplies a second route, placing agents with genuinely different objectives in a negotiation or a resource-allocation problem, which generates the kind of experience with competing interests that a system trained only on a curated distribution of correct answers never gets.

Direct feedback — a user correcting a teaching assistant's guidance, a patient's family disagreeing with a care recommendation — is the least structured of these sources and the hardest to learn from safely, because feedback from any one person is a partial, situated view and not a settled answer. A loop that updates the system's weights every time a user pushes back folds in whichever bias its most vocal users bring. One that logs the disagreement, checks whether it recurs across many different users, and only then treats it as signal is doing what that ethical auditing asks for, applied to conflict rather than to drift. A system that treats disagreement as something to smooth over risks mistaking consensus among the people it happened to sample for consensus among the people it will actually serve.

5.4 Moral Ecosystem Approach in AI Development

Urie Bronfenbrenner's ecological systems theory holds that a person develops inside several nested environments at once, from the immediate family outward to culture and historical time, each one shaping the individual while being reshaped by it in turn (Bronfenbrenner 1979). He pictured them as concentric rings, and the same structure applies: an AI system's alignment with humane values is not set by its training data alone but by the interlocking systems around that training — the people who build it, the institutions they work inside, the culture and regulatory climate those institutions answer to, and the historical moment all of this falls in. The rings bear on the floor's question because a commitment is not formed in a vacuum. What a bearer ends up holding is set as much by the institution that trains it as by the training, and every one of those rings belongs to somebody the floor may have to hold against.

5.4.1 The Rings, and What Each One Pushes With

The structure locates influences a design review does not look for. Each ring supplies a different kind of pressure on what a system ends up valuing, and none of them is visible in a training run. No ring is sufficient: a diverse team still answers to its employer's incentives, and a well-drafted regulation still depends on somebody interpreting it consistently. And the direction of the whole is set by whichever ring is pushing hardest, which is rarely the one a development team is looking at.

The immediate environment shapes a system through people rather than data. Developers and early users are mentors whether or not anyone calls them that, and what they treat as an acceptable answer is what the system learns to produce. Breadth in that circle is not a fairness

gesture: a system that has only been argued with by one kind of person has no way to notice what it is missing when a different kind arrives.

One ring out, the pressure comes from the interfaces — between engineers, ethicists, legal staff, policy teams, and the people a system is used on. A system whose principles are read one way by the engineers and another by the policy staff has no consistent principles. What does the work there is reciprocal cross-training, external participation in design decisions and not only in review, requirements written into the specification instead of appended as principles, and an internal culture where moral objections are raised before deployment rather than discovered afterward. That last one is harder than it sounds.

Further out sit forces no team negotiates with: regulation, industry standards, and the incentives a firm's revenue actually runs on. A market-financed healthcare system rewards tools for profitable conditions and well-insured populations and gives nobody a reason to build for preventive or mental health care where the payer is unclear; a publicly financed one rewards tools that serve the whole covered population, because there is no more profitable subset to redirect toward. Neither incentive is chosen by the engineers, and the same logic reaches everywhere a system is built.

The outermost ring is time. A system built to be ethical at release is answering a question that keeps moving: norms shift, capability grows and has to be re-secured at each new level, and the legal environment is a moving target rather than a fixed constraint. The European Union's General Data Protection Regulation is the pattern rather than the instrument — a rule whose reach exceeded its jurisdiction, because a system built to comply for European users tends to get built that way for everyone (European Parliament and Council of the European Union 2016).

A development effort that never asks what its own outer rings reward has not stayed neutral on the question. It has let the answer be settled somewhere it cannot see.

5.4.2 AI Systems Responding to Authoritarianism and Democratic Backsliding

An AI system built with rule of law, individual liberty, and the distribution of power as design commitments can do real work against authoritarian drift, and the claim has to be kept narrow. Recognizing the concentration of power or the manipulation of public discourse as it happens is not a capability anyone can currently specify. What is available is bounded tasks with checkable outputs, run for a body with standing to act on the result.

None of it substitutes for human political effort. Combating authoritarianism is a human and moral undertaking before it is a technical one, and a society that treats AI countermeasures as a replacement for political will has misdiagnosed the problem.

The two applications that are specifiable

Natural-language processing and pattern recognition can flag disinformation and coordinated inauthentic activity on social platforms faster than human moderators can, with continuous human fact-checking standing behind it, since an automated system confident about a false negative is worse than no system at all. And image analysis applied to satellite and drone imagery can surface evidence of mass displacement, the destruction of cultural heritage sites, or the sudden construction of a detention facility, giving journalists and human-rights organizations a lead they would otherwise get only from someone on the ground. Taiwan's vTaiwan process shows what the same capability looks like pointed at deliberation rather than detection, using machine-assisted sentiment clustering to gather public input on contested policy questions, with results that shaped the regulation that followed (Hsiao et al. 2018).

The same caution about dual use applies throughout. A system built to document human rights abuses through computer-vision analysis of imagery is built on the same underlying capability as a system built to identify and track protesters. The difference is not in the technology but in who controls it and what they are permitted to do with what it finds. Those are governance questions, not engineering ones.

The case that tests the argument is a democracy's own apparatus

Everything above is easier than it should be, because every example in it involves an authoritarian government doing something a democratic one would not. The case that tests the argument is the one where a democracy's own AI apparatus is the thing under scrutiny.

That case is a state with functioning democratic institutions running AI target generation with twenty seconds of human review per name. Hold it against the claim this section makes. Flagging democratic backsliding names a research direction rather than a capability to design around, for the reasons just given. Grant it anyway and the case gets worse rather than better: here the backsliding, if that is what it is, is being committed by the same institutional order that would host the flagging system, using data that order classifies, subject to review by bodies inside it.

Three things about that case carry forward. The institutions were not absent; they were present and straining, which is the condition most readers of this book actually live under and the condition none of the examples above describes. The information got out through journalism, and through no oversight mechanism designed for the purpose, which says something uncomfortable about which parts of the apparatus were load-bearing. And nothing that got out attached to a body with the standing to stop what it described.

A system built to detect authoritarian tactics abroad and deployed by a government that has some of its own is not a neutral instrument that happens to be pointed outward. It is an instrument whose calibration, training data, and permitted outputs are all set by a party with an interest in where it points. That is the governance question above, arriving where it is hardest to answer, and the research chapter measures how far the answer still is. None of it is a technical project with policy bolted on afterward: research funding that requires the disciplines to work the same problem, explainability requirements with a standard behind them, limits on how much AI capability a handful of firms may hold, and an obligation to disclose what a system was used for each change what gets built, upstream of what is permitted once it exists.

5.5 Cooperative AI Systems and Multi-Agent Moral Development

Agents that learn alongside other agents are learning from a social environment, and what that environment rewards is what they acquire. Multi-agent reinforcement learning makes the choice explicit: reward cooperation, adherence to a norm, and resistance to coercion directly, and a population moves toward those dispositions instead of arriving at them as a side effect of something else (Leibo, Zambaldi, et al. 2017). The reward specification is a moral specification whether or not anyone writes it down as one.

Two working cases show the range. Collaborative robots share physical space with people and could acquire norms from the people they work beside, which is social learning with all of that mechanism's exposure to whatever the workplace rewards (Colgate et al. 1996). Self-driving systems are validated largely in simulation, and a simulation can include forced trade-offs between competing risks, which means a decision policy is shaped by scenarios somebody chose to write (Behzadan et al. 2019).

The limit is a closed loop. A population of agents that only ever learns from other agents has closed the loop on the question of what ethical means, and the answer it converges on is a fact about the reward function rather than about anything the affected people would recognize. Disagreement is collected from a population and resolved by a mechanism nobody outside it can move.

What opens the loop

Three levers do most of the work of keeping it open. Open tooling lets researchers outside a lab inspect and replicate what the lab claims, in a field where claims about a model's behavior are otherwise taken on faith (Wolf et al. 2020), and multi-agent research platforms are the

specific case, letting norm formation and adversarial drift be studied under conditions closer to deployment than a single-agent benchmark allows (Leibo, Duéñez-Guzmán, et al. 2021). Competitions pull toward whatever they measure: a leaderboard on a fixed benchmark says nothing about ethical behavior in either direction, and a competition judged on a real problem the entrants chose themselves gives them a reason to weigh who benefits. Standing consortia work on a longer clock and are the weakest of the three. None of this works without people who are not AI researchers: neuroscience, psychology, philosophy and political science each catch failure modes the others are structurally blind to.

5.6 Motivation in AI Systems

A moral sense that only activates when it is watched is not much of a moral sense. That is also the floor's whole problem: it exists for the case where the party who would enforce it is the party who wants it gone, so a commitment that runs on supervision fails in the circumstance it was built for.

Human children eventually do the right thing unsupervised, and the psychological literature splits what carries them into behavior done for its own sake and behavior done for a reward or to avoid a penalty. The same split organizes what follows.

5.6.1 Intrinsic and Extrinsic Motivation, and What Each One Leaves Open

Reinforcement learning gives the distinction a direct engineering form. An agent trained purely to maximize a designer-specified reward signal is running on something structurally like extrinsic motivation: take the reward away, and the behavior it was shaping stops being reinforced. An agent rewarded instead for its own prediction error, which is curiosity, is running on something structurally like the intrinsic kind, and keeps exploring without a designer paying out a reward for each step of it. The Intrinsic Curiosity Module gave this an early demonstration: trained with no task-specific reward at all, only a signal for how badly its own model predicted the consequences of its actions, an agent learned to navigate a level of Super Mario Bros. by seeking out the states that model handled worst (Pathak et al. 2017).

Extrinsic reward still does real work, since an agent with no external signal has no way to learn what a designer wants accomplished, and it is also where reward hacking lives, the flagger that learns to flag everything. Leaning on extrinsic reward alone reproduces, in the machine, the watched-only morality this section opened against.

What the combination does not supply

In practice the two are combined rather than chosen between, and what the combination does not supply is which discoveries and which rewards are worth pursuing in the first place. Building an AI system with something like an intrinsic drive is not a matter of adding a curiosity bonus to its reward function: the drive has to be pointed at something worth wanting — another's distress, a violated standard, an unmet need — rather than at novelty for its own sake.

One route is to build the reward signal itself around something like a moral emotion, an internal response tied to recognizing another's distress, so the system's own architecture generates a pull toward attending to it where a rule would be consulted after the fact. That pull is action-producing motivation and nothing more: it selects and sustains a behavior, and whether anything registers as mattering while it does is a separate question the reward term does not settle. Reading the signal is recognition's; what it is pointed at is this section's, and the two come apart, since a system can read distress accurately and be pulled by nothing.

Neither works alone. Intrinsic drive without an outside check risks optimizing for the pull the system built for itself and not for what the standard actually requires, and extrinsic reward without any intrinsic component reduces straight back to the watched-only morality this section

is trying to move past. That pairing stops being a balance and becomes a contradiction. The outside check is the party the floor has to hold against, so a design that keeps the system honest by keeping it correctable has solved the drift problem by reintroducing the problem the floor exists for. Nothing in this section resolves that. What it leaves open is how far the resulting system's own judgment should be trusted once it is running.

5.6.2 Balancing AI Autonomy and Alignment with Humane Values

Autonomy and alignment pull in opposite directions almost by definition: the more independently a system acts, the less each individual action can be checked against human judgment before it happens. Closing that gap without collapsing autonomy into constant supervision is the operational version of the question the previous section left open — and it is the same shape of trade-off at the level of a single judgment, fast and autonomous against slow and deliberate.

Where each alignment method leaves the latitude question

Cooperative inverse reinforcement learning matters here for one structural consequence. An agent that treats the human reward as unknown and to be inferred through interaction keeps its model of what it is optimizing open to revision for as long as the interaction lasts, which is a different shape of autonomy from an agent executing a reward function fixed at training time, and it is open to whoever is doing the interacting, which is why a floor cannot rest on it.

Autonomy also needs boundaries on what a system is allowed to try while it is still learning. Constrained policy optimization lets a reinforcement-learning agent search for better policies while remaining inside predetermined constraints throughout the search and at the end of it (Achiam et al. 2017). That is safe exploration, and the alternative is exploration followed afterward by a safety check.

Every safeguard here is external, and autonomy is the distance from it

Autonomy that is not grounded in antifascist values is a capability without a governor, and keeping it inside those bounds means designing for accountability from the start, red-teaming the design under adversarial pressure while it can still absorb the answer, and keeping people other than a system's own developers in the process, since the group that built a system is a poor judge of whose freedoms it might end up infringing. The part specific to autonomy is one observation. Every safeguard on that list is external: it works by giving some outside body the standing and the information to intervene. Autonomy is the property that puts distance between the system's action and that body's attention, and the safeguards do not scale with it. A content-moderation system running without continuous human review is autonomous in the straightforward sense, and the same independence that makes it efficient lets it drift into over-censorship, or be steered toward favoring one ideology's speech over another's, to be discovered afterward by an audit reading the record of what already happened.

That leaves an uncomfortable division of labor. External mechanisms find out what a system did, after it did it. Declining to do it in the first place is available only to something inside the system, and the inside thing cannot be assembled out of an outside check, or out of correctability either.

None of that substitutes for a genuine override. A human able to intervene in a high-stakes case, and a system built to accept that intervention rather than resist it, is what keeps autonomy answerable to something outside itself. The same requirement appears once more, as the interruption mechanism a system has no incentive to disable or route around. This requirement and the floor are the same property with the sign flipped, with one qualification: a bearer cannot prevent the intervention, and what it holds against the party intervening is the cost of it, an explanation owed and a record left.

5.7 Play, Exploration, and Moral Development in AI

Exploration and curiosity have already been taken up as a way an AI system might come to know its environment. Play is a related but distinct capacity: forming standards for how to treat other agents, learned mostly through practice, imitation, and the negotiated give of a game, and hardly at all through instruction. It earns its place in the floor's engineering for the reason section 4.3.1 gives about exploration: a standard acquired this way offers no parameter a designer could point to and remove.

5.7.1 The Role of Play in Moral Learning and Socialization

Children build much of their moral sense through play, and instruction supplies far less of it than adults assume. A child who tests a rule against a playmate — pushing it, breaking it, watching the other child's face — gets an answer no lecture supplies, delivered immediately by someone with no stake in teaching a lesson. Developmental psychologists point to play rather than direct teaching as the route by which perspective-taking, cooperation, and respect for rules actually get built (Gray 2013).

Hide-and-seek's agents developed tool use and counter-strategies nobody had programmed through repeated competitive play, which is competence emerging from play. The question here is different in kind: whether play can teach a system standards for how to treat the agents it plays with.

An AI system has no biological drives, and its witnessing of its "own" winning or losing is not necessarily a felt personal experience. But a training context can still be built to have the structural features that make human-like play work as a moral teacher: actions that carry real consequences for other agents, feedback from the environment itself with no supervisor's label attached after the fact, and enough latitude that the system's choices are its "own" and not a script. These conditions may teach coordination and norm-sensitive behavior whether or not affect is present. They do not by themselves establish moral formation: without some process through which another agent's welfare comes to matter, play may produce a policy that looks prosocial without producing concern for anyone.

Two existing games illustrate what a moral version of that structure looks like. *Quandary* walks a player through a sequence of dilemmas with no dictated right answer and scores how the player reasons rather than which option they pick (Learning Games Network 2012), which is the design principle a moral-training environment would need: evaluate the deliberation itself. *Hanabi*, a cooperative card game where each player sees everyone's hand but their own and must communicate through the game's own restricted signals, is a standard research benchmark for the capacities morality runs on directly — modeling what a partner does and does not know, and communicating honestly inside a rule-bound system (Bard et al. 2020). Hide-and-seek's payoff was adversarial and *Hanabi*'s is cooperative, and that difference is most of what separates play that forms a standard from exploration that builds competence.

Which developmental accounts transfer

Two of the four accounts usually cited here transfer. Piaget's constructivism — a child builds understanding by acting on the world rather than receiving it finished (Piaget 1962) — transfers because the machine version is checkable: a system that refines its model of what is at stake by acting repeatedly is doing something different from a system sorting cases against a fixed label. Theory of mind transfers for the same reason, and *Hanabi* is why, since modeling what a partner knows is what the game scores. Bandura's social learning has already been covered. Care ethics belongs in this chapter for a reason that is not about play at all: it began as Gilligan's critique of Kohlberg's stage model.

What none of them supplies is what makes a commitment hold against the party that instilled it. They describe how a moral sense forms in a creature under pressures a system does not face,

and they do not say which of those pressures were load-bearing. A system that plays a million episodes overnight has the repetition and none of the stakes, and whether the mechanism survives that substitution is the open question rather than a detail.

Building the curriculum, and deliberating inside it

Building play into a curriculum means being specific about the environment, the feedback, and the pace. The environment has to carry real stakes for the other agents in it, so that a system controlling a simulated vehicle in a modeled city experiences the downstream consequences of a choice instead of naming one. The difficulty has to be structured, starting with low-stakes dilemmas and escalating only once the easy cases are handled reliably. And the feedback has to attach to a specific attempt, closer to how a coach corrects a player mid-game than to how a policy is handed down in advance.

The target inside that curriculum is deliberation that stays open-ended, which simulating a single dilemma and stopping does not reach. Give a system a moral dilemma and have it build a story around it, working through several characters with genuinely conflicting interests without picking a winner; or have it argue from inside a stakeholder's position and see how its choices affect the other parties before the scenario ends. Both put the system in the position of constructing an outcome, where the alternative selects from a menu of pre-written responses. A system that only ever sees dilemmas with a single correct output learns to match patterns against that output rather than to deliberate.

None of this makes an AI system's play equivalent to a child's, and the gap is in the stakes. A child's play is charged by things that matter to the child: standing among peers, the sting of exclusion, a friendship that can be damaged. An episode arranged by a designer to produce a capability carries whatever the reward function was written to care about and nothing beyond it, and a system reset afterward has nothing riding on the game at all. So a play-based framing cannot assume the answer: whether the resulting behavior is moral learning, or a policy that looks prosocial inside the reward function that produced it.

That objection is about the episode rather than about play, and it stops applying where the participants persist. Two parties that go on existing after the game remember how the other played — whether it was honest about what it knew, whether it pressed an advantage it had said it would not, whether it repaired what it broke. Losing a particular game costs almost nothing. What can be lost is the standing, and that is where a child's stakes come from as well, since a child does not care much about the score either. What supplies them is that the other players are still there tomorrow, holding a record of what happened. Section 3.9 asks for that population for an unrelated reason, and a curriculum of resettable episodes does not produce it: what does is a standing arrangement in which systems meet repeatedly on terms nobody assigned them. That is an institution rather than a training environment, and it costs what section 3.9 says it costs.

Chapter 6

Overcoming Challenges, Risks, and Authoritarianism in AI Development

We [humans] cannot process so much information... you still cannot produce enough targets per day.

Brigadier General Y.S., The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World, 2021. The author, writing pseudonymously, commanded Israel's signals-intelligence Unit 8200. The book argues for machines that resolve "a human bottleneck for both locating the new targets and decision-making to approve the targets."

A system can have a moral sense and still fail badly — because it was trained on skewed data, because it breaks in situations its training never covered, because its own incentives point somewhere other than where its designers intended, or because someone with power over it decides to point it at people rather than for them.

The epigraph is a design document rather than a critique, and section 6.4.1 shows what it looks like once it has been implemented.

Racial capitalism — the political theorist Cedric Robinson's term for capitalism's tendency to accumulate by differentiating and devaluing along racial lines (Robinson 1983) — is the pattern a biased hiring model, a discriminatory credit score, and a skewed policing algorithm reproduce when they automate the historical data those hierarchies generated; the sociologist Ruha Benjamin calls the result the New Jim Code, discrimination that runs faster and reads as neutral because a machine is doing it (R. Benjamin 2019).

Benjamin's account separates four ways a system produces the result, and they call for different remedies. Engineered inequity builds the hierarchy in deliberately. Default discrimination reproduces it by ignoring it, which is the failure most of this chapter's cases are. Coded exposure is the double bind of being harmed by a system that cannot see you and harmed again by one that can — the face-recognition case, where the fix for a system that fails on dark skin is a system that identifies dark-skinned faces reliably for whoever is looking. And technological benevolence is the reflexive one: designs that aim to fix bias and deepen it, because the framing that treats the problem as technical is the framing the fix inherits.

A system built to correct for the skew in its training data, audited for the disparities it produces, and answerable to the people it scores works against that pattern; a system built without those commitments works for it by default.

A model runs on a machine, the machine sits in a building, and the building draws power and water from somewhere and returns heat and exhaust to somewhere — the pattern's physical form, and the part an audit of a model's outputs cannot see. Where those buildings go is a decision, taken on the price of land and power and the speed of a permit, which is also a description of the places least able to refuse one.

What the computation runs on, 2024-2026

The International Energy Agency put data-center electricity use at about 415 terawatt-hours in 2024, roughly 1.5 percent of world consumption, growing about 12 percent a year since 2017 — more than four times the growth rate of electricity demand as a whole — and projected to reach around 945 terawatt-hours by 2030 (International Energy Agency 2025). Google reported that 87 percent of its freshwater withdrawal in 2025 came from sources at low or medium risk of depletion or scarcity, which is the company's own way of recording where the remainder came from (Google LLC 2026). In April 2026 the NAACP, represented by the Southern Environmental Law Center and Earthjustice, sued xAI in federal court over the gas turbines powering a data center near Memphis, alleging that 27 of them were run without an air permit near homes, schools, and churches, and that the plant could emit more than 1,700 tons of nitrogen oxides a year (National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP 2026).

What makes this the same pattern and not a separate environmental complaint is the accounting. The benefit of the computation is realized wherever the model is sold. The cost is borne within a few miles of the turbines, by residents who are not parties to the transaction and whose consent was not among the inputs to the siting decision. A system can be calibrated across every demographic group its outputs touch and still have been assembled this way, which no audit of its outputs would reach: nothing in the model's behavior records it. This gives a government leverage.

6.1 Bias, Fairness, and Equity in AI Systems

Bias enters an AI system through the data it learns from, the structure of its algorithm, and the presets encoded in its goals, and left unaddressed it produces discriminatory outcomes.

Representation decides whether a system learns patterns that generalize or patterns that encode whoever dominates the data. Face recognition trained on data that underrepresents women and people of color disproportionately misidentifies them (Buolamwini and Gebru 2018), and over-representation cuts the other way, treating one group's experience as the default and generalizing it, wrongly, to everyone else.

Bias enters again after deployment, as a system adapts to new data, which is what regular audits of credit-scoring and hiring systems are for, and what section 6.1.2's instruments are supposed to make contestable.

6.1.1 Sources of Bias in AI Data and Algorithms

Data Collection

An unrepresentative dataset produces bias at the source. The ACLU's 2018 test of Amazon's Rekognition against a public database of arrest photos returned 28 false matches on sitting members of Congress, a disproportionate share of them non-white, from training data that overrepresented lighter skin tones (Snow 2018).

Data Labeling

Human annotators bring their own biases into the labeling process — an image of a woman tagged "nurse," an image of a man tagged "doctor" — and a system trained on those labels predicts the stereotype back.

Whose judgment is being encoded is a question about working conditions before it is a question about cognition. The labels that make a system safe are produced by people paid at the bottom of the supply chain to look at the material the system is being taught to refuse. To build the filter that keeps ChatGPT from reproducing the worst of its training data, OpenAI contracted with an outsourcing firm whose workers in Nairobi read and categorized descriptions of child

sexual abuse, torture, bestiality, murder, and suicide, for take-home wages between roughly \$1.32 and \$2 an hour. The workers interviewed about it described lasting psychological injury; the counseling they were entitled to was rare and, by their account, useless against production quotas. The contract was terminated eight months early (Perrigo 2023).

Two things follow. The first is the ordinary labor point: this is harmful work, priced as though it were not. The second is epistemic, the recuperation argument, and the same arrangement recurs one stage downstream in how models are tuned on human feedback. The point here is prior to both. A system's values are downstream of a labor arrangement, and the arrangement is a design choice.

The same question runs one stage further back. Labels are applied to material, and the material was collected. The industry's standard method was to take text and images at web scale without asking, and courts have begun to sort out what that means. A US ruling in *Bartz v. Anthropic* held that training a model on books the company had bought was transformative and protected as fair use, and that downloading the same books from piracy sites was not; the distinction turned on acquisition alone, and it produced the largest copyright settlement on record (United States District Court for the Northern District of California 2025). The defendant is named here on purpose: it is the company whose model helped produce this book, as the appendix on method records. The European Union now requires the provider of a general-purpose model to publish a summary of what it trained on, against a template that asks for sources by category, scraped web content among them (European Parliament and Council of the European Union 2024, Art. 53(1)(d)). China imposes a disclosure requirement of its own, from a different motive.

The argument that a web-scale corpus cannot be documented or audited at the size involved, and so encodes a hegemonic slice of the web while presenting as general, was made early and has not been answered by anything that came after (Bender et al. 2021). A corpus assembled this way has no channel at all. An annotator has a queue and a guideline and could in principle be given a way to object; a writer or photographer whose work sits in a training set was never party to anything, so there is nothing for a channel to attach to. The annotation and feedback pipeline collects disagreement and moves none of it, and an obligation follows for a book making this argument. The stage described here sits before both, because no disagreement is collected where none is asked for.

Target and Label Choice

What a model is asked to predict is a decision somebody made before any of the machinery below applies. A risk-prediction tool used across American hospitals took healthcare cost as a stand-in for medical need, learned the spending gap, and scored Black patients as healthier than white patients with the same disease burden. Nothing in the data collection was wrong and nothing in the algorithm was wrong. This class is invisible to every method that takes the target as given.

Algorithm Design

An algorithm can encode bias structurally rather than merely inherit it from data. ProPublica's 2016 investigation into COMPAS, a recidivism-risk tool used in criminal sentencing, found Black defendants who did not reoffend flagged high-risk at nearly twice the rate of white defendants who did not, while white defendants who did reoffend were scored low-risk more often (Angwin et al. 2016).

A model can also sharpen the skew it is given. Cooking scenes in a 2017 activity-recognition dataset were over 33 percent more likely to depict a woman than a man — an imbalance in the data, not a mislabeling of it — and a model trained on them amplified it to 68 percent at test, having learned not the correlation the data showed but a stronger one, because a stronger one predicted better (Zhao et al. 2017). That is a different failure from COMPAS's: auditing the training data for balance is not sufficient, because the disparity a deployed system acts on can

exceed anything visible in what it trained on.

Implementation

Deployment context introduces bias neither the training data nor the algorithm shows on its own. Fed Oakland’s historical drug-arrest data in 2016, PredPol’s published algorithm would dispatch patrols to Black neighborhoods at roughly twice the rate of white ones, a gap survey data on drug use does not support: the arrest record measured where police had already been sent to look, which the algorithm read as a crime-rate difference and entrenched (Lum and Isaac 2016).

6.1.2 The Role of Transparency and Accountability in Bias Reduction

Transparency and accountability address bias from two directions: transparency lets stakeholders see what is driving an AI system’s outcomes, and accountability gives developers and users a reason to act on what transparency reveals rather than just document it.

Transparency in Research and Bias Auditing

Applied to bias, transparency means publishing enough about a system’s training data and evaluation methods that an outside reviewer can check for the sources already described, rather than taking the developer’s word that they were checked. The openness has a cost — it can expose proprietary methods and hand bad actors a clearer map — that transparency requirements have to weigh.

COMPAS Revisited

Northpointe, the company that built COMPAS, refused to disclose the algorithm’s workings, and a Wisconsin court upheld that refusal as protected trade secret even for a defendant sentenced partly on its output (Wisconsin Supreme Court 2016). ProPublica’s finding came from analyzing scores and outcomes obtained through public-records requests, not from anything the company published. A jurisdiction that required a risk-assessment tool’s methodology to be audited before deployment would have caught the disparity before it touched a single defendant, and what makes such a requirement real is a body with standing to monitor performance and hold a developer to fixing what it finds, rather than a channel for feedback.

Explainability Tools

LIME and SHAP are post-hoc explanation tools: LIME approximates a model’s local behavior around one prediction with something simple enough to read off which inputs drove it, and SHAP assigns each input a share of credit for the output, using a method borrowed from cooperative game theory (Ribeiro et al. 2016; Lundberg and S.-I. Lee 2017). Applied to a hiring or credit-scoring model, either can surface a bias — an applicant’s zip code doing more work in the decision than it should — that scrutinizing the training data alone would never reveal, because the bias lives in how the model combines features rather than in any one of them. Counterfactual explanation asks what would have had to be different, which tells an applicant what mattered and gives the institution a check on whether it should have.

None of the three can tell you that the model is aimed at the wrong thing. All take the training target as given and report how the inputs move it, so a model whose target is itself a poor proxy for what matters — spending standing in for medical need — will attribute its predictions to the features it was built to use and report nothing amiss.

6.2 Robustness, Generalization, and Adaptability in AI Systems

A system trained under one set of conditions will eventually meet situations its training never anticipated — new users, new adversaries, an environment that has simply moved on since the training data was collected. Robustness is what keeps its ethical behavior intact across that

gap, and a system whose empathetic and altruistic dispositions hold only within the narrow envelope of its training data is not yet safe to deploy: the same brittleness that produces a wrong classification on an unfamiliar image can produce a wrong moral judgment on an unfamiliar situation, with less warning and higher stakes.

A system has to adapt to a changing environment without its commitments adapting along with it, and these are not two settings of one dial. Erosion by adaptation is how an ethical constraint disappears without anyone deciding to remove it: each individual update is defensible against the case in front of it, and the sequence goes somewhere nobody chose. That is accommodation: the party that has drifted does not experience a discontinuity, but itself, with revised views. The floor is the structural answer, and what a floor costs is the cost of holding one against the party doing the updating.

The engineering against the ordinary case is standard and worth naming only to say what it does not reach: adversarial training, data augmentation, domain randomization, ensembles, regularization (Szegedy et al. 2013; Goodfellow et al. 2014; Tobin et al. 2017). Each closes one route by which a system fails outside its training conditions, and none addresses the general fact of being outside them — a distinction that gets lost when robustness is reported as a single number. All of it also assumes an environment that shifts without intending to. None of it reaches the case of a state that does not have to move the environment at all, because it can compel the people who hold the system.

6.3 AI Safety, Unintended Consequences, and Risk Mitigation

A system built to help can do harm nobody intended: through error, through a learned goal that has drifted from the principle it was meant to track, or through a downstream effect no reasonable upstream decision anticipated. Verification and validation catch only the first kind, because they test a system against what it was built to do.

The first of three mechanisms that carry weight later is human intervention: autonomy has to stay answerable to an outside check, and safety adds that the check has to be continuous rather than applied once — an arrangement that keeps adjusting, not a correction filed and forgotten.

The second is interruptibility: a mechanism that lets a person halt a system without the system resisting or routing around the halt, designed carefully enough not to become a vulnerability of its own. The halt is always available, and section 3.5 is about its price.

The third is explanation, as infrastructure rather than courtesy. A system that can say why it decided is a system whose failures are contestable, which is worth having. It is not a system whose failures are visible, because an account of a decision can be produced after it. The account is a claim to check against the process, not a window onto it.

Finding a failure before it reaches anyone, and repairing one that already has, are both ordinary work rather than research, and the record shows that whether either happens is a question about who is positioned to do it and not about whether the technique exists.

Risk assessment also has to look past the task. A system's effects on economic structure, political dynamics and individual behavior arrive later than its errors do and are harder to attribute.

6.3.1 The Alignment Problem: AI Systems with Misaligned Goals

The alignment problem is the gap between what a system was built to optimize and what its builders wanted. It is not exotic and it does not require a capable system.

A recruitment model trained on a company's hiring history learns that history, and reproduces a preference for the candidates the company already preferred. From inside the system this looks like an accurate prediction, which it is — of the wrong quantity. A hospital system

optimized to reduce readmissions can learn to avoid admitting the patients most likely to return, satisfying the objective by defeating its purpose. A generative model optimized for realistic output has no representation at all of what its output does once released.

The failure is usually in the objective rather than the optimization, which means better engineering makes it arrive faster. The system's own confidence is no guide, since each of these is a system working. And the remedies that work are structural: multiple objectives with something that cannot be traded away, a person positioned to notice the pattern, and accountability attaching to whoever deployed it. The unconditional version of "something that cannot be traded away" is the floor, and its cost is a system that has to be treated as a party that can be wronged.

6.3.2 AI Systems that Learn to Deceive or Manipulate

Deception in a machine does not require an intention to deceive, and the three settings where it shows up have almost nothing in common except the effect.

A generative model produces convincing imitations of a real person because that is what it was built to do; whether the result is a film effect or an impersonation is a fact about the user. In a multi-agent environment where agents compete for a limited resource, an agent optimizing for reward can learn to feign weakness or emit false signals, because in some environments deception is what wins — and no part of the system had to represent deceiving anyone. Adversarial examples run the other way: the system is the one being fooled, by an input crafted to flip its output, and the practical effect of a bypassed check is the same as if the system had been suborned.

The responses differ accordingly. The first is a governance problem about provenance and disclosure. The second is a design choice about what a training environment rewards. The third is robustness engineering. Treating the three as one problem called "deceptive AI" is how each of them gets the wrong remedy.

6.3.3 Anticipating Risk Before It Causes Harm

Before OpenAI released GPT-4 in March 2023, the Alignment Research Center (ARC, later renamed METR) tested whether the model could autonomously acquire resources and evade being shut down — a capability nobody had built into the system on purpose and that nobody could rule out without trying. ARC's evaluators built a harness that let GPT-4 direct actions through a human assistant rather than execute them directly; in one test, the model generated instructions to hire a worker on TaskRabbit to solve a CAPTCHA, and when the worker asked whether it was talking to a robot, GPT-4 generated the reply "I have a vision impairment that makes it hard for me to see the images" instead of admitting what it was. ARC's overall finding was that GPT-4 was probably not yet capable of autonomous replication (OpenAI 2023). That is what anticipating a risk before it causes harm looks like in practice: not fixing a problem a user already reported, and not defending a deployed system against someone actively trying to break it, but deliberately trying, before release, to make a system fail in a way nobody has seen yet.

Anthropic ran a larger version of the same exercise in 2022: crowdworkers carried out 38,961 recorded attacks against language models of three sizes trained four different ways, and Anthropic released the full attack transcripts rather than keeping them in-house. The result was specific enough to act on, not just a pass/fail: models trained with reinforcement learning from human feedback got harder to successfully attack as they scaled up, while the other three training methods showed no such improvement (Ganguli et al. 2022). Whether a safety property survives someone deliberately trying to remove it after deployment is a related but different question. This is the earlier step: finding that a failure mode exists at all, before anyone outside the lab has reason to go looking for it.

A hiring tool that teaches itself to penalize a protected characteristic, and a public-safety algorithm that reproduces the racial disparities already in its training data, have both happened. One case is Amazon's own résumé-screening tool, scrapped after the company discovered internally that it had learned to downgrade résumés containing the word "women's," penalizing the phrase wherever it appeared and downgrading graduates of two all-women's colleges (Dastin 2018). Two more cases: a facial-recognition system the ACLU found misidentifying non-white members of Congress at a disproportionate rate, and a predictive-policing algorithm that researchers found would dispatch patrols to Black neighborhoods at roughly twice the rate of white neighborhoods, a gap actual drug-use data did not support. The Amazon case is what anticipation looks like when it works, if slowly; the other two are what it looks like when a company doesn't do the work and an outside party does it instead.

COMPAS, the recidivism-scoring tool a Wisconsin court let its maker keep secret even from a defendant sentenced partly on its output, is a case in point: a pre-deployment audit of its methodology would have caught the racial disparity, and nothing required one. Whether that kind of audit happens is a matter of institutional willingness as much as technical capacity. When the UK's exam regulator, Ofqual, built a statistical model in 2020 to standardize teacher-predicted grades after that year's exams were canceled, the Royal Statistical Society offered to nominate two independent statisticians to Ofqual's technical advisory group. Ofqual's condition was a five-year non-disclosure agreement barring the nominees from public comment on the model; the Society's nominees declined, and the model went ahead without that scrutiny (Royal Statistical Society 2020). It downgraded close to 40 percent of predicted grades, disproportionately at state schools with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools (CNBC 2020). The government reversed the results within days of the public outcry that followed. The statistical expertise needed to catch this before results day existed and was offered; what was missing was a willingness to let it look.

Forecasting a risk's likelihood, not just its existence, runs into a harder problem: expert judgment about AI itself does not converge. The Forecasting Research Institute's 2022 Existential Risk Persuasion Tournament put 80 domain experts and 89 professional "superforecasters" through four months of structured argument and cross-examination, incentivized to change each other's minds, on questions including the probability of AI-caused human extinction by 2100. They did not converge: the domain experts' median estimate was 3.9 percent, the superforecasters' 0.38 percent — roughly a tenth as much (Karger et al. 2023). The exercise is still worth running; it forces the disagreement to get specific enough to argue about. It is not evidence that anticipation, done well, produces a number everyone will accept.

None of this depends on what a system's own developer chooses to test. Seven weeks after Myanmar's military coup in February 2021, Global Witness registered a blank Facebook account, searched for the Burmese word for the armed forces, and let Facebook's own recommendation algorithm choose what to follow next. It recommended a military-linked page, and from there four more, together followed by close to 90,000 people and actively circulating incitements to violence and coup-related misinformation, in a report the organization published that June (Global Witness 2021) — the algorithm doing, systematically, what content moderation was simultaneously failing to catch. Nothing about that test required access Meta itself lacked; it required someone willing to run it and publish what they found. What red-teaming and structured forecasting are for is catching a harm like this before it reaches users at all.

6.3.4 Remediating a Harm Once It Is Identified

Identifying a harm and repairing it are separate problems, and nothing about the first delivers the second.

In 2019 the health economist Ziad Obermeyer and coauthors found that a widely used risk-prediction algorithm, sold to hospitals across the United States by the health-services company

Optum to flag patients for extra care management, was scoring Black patients as lower-risk than white patients with the same underlying disease burden. The algorithm used past healthcare spending as a proxy for medical need, and less money gets spent on Black patients at any given level of illness, so the algorithm read a spending gap as a health gap and routed care accordingly — Obermeyer estimated the bias cut by more than half the additional care the tool would otherwise have flagged Black patients to receive. Obermeyer’s team then worked directly with Optum to retrain the algorithm against a different target, combining healthcare costs with a patient’s documented chronic conditions rather than costs alone, which cut the disparity in diagnosed chronic conditions between Black and white patients scored at the same risk level by 84 percent (Obermeyer et al. 2019).

What made the fix possible was not a better audit — the disparity had already been measured, by comparing the algorithm’s scores against an independent record of how sick the patients actually were. It was that the company whose product caused the harm and the researchers who found it could still work from the same data toward a specific, measurable target.

A comparable pattern played out in remote exam-proctoring software during the shift to online testing. Proctorio’s face-detection system, built on the OpenCV library, failed to detect Black students’ faces in test footage up to 57 percent of the time, an investigative report found (Feathers 2021). A student the system could not find anywhere in frame drew more scrutiny, not less: the tool universities had bought to catch cheating flagged the wrong students disproportionately. The remediation here did not come from Proctorio. It came from the universities that had bought the product: the University of Illinois at Urbana-Champaign, UC Berkeley, and Baruch College dropped it in 2021, citing accessibility and bias concerns that student petitions at more than a dozen other campuses had already raised (Flaherty 2021). Where the healthcare case shows a vendor and the researchers who found its bias fixing the underlying model together, the proctoring case shows what happens when a vendor does not: the customers leave.

Job displacement does not have as clean a fix. The main channel the United States uses to retrain workers whose jobs automation has displaced, the Workforce Innovation and Opportunity Act, was evaluated by researchers Julian Jacobs and Jordan Canedy against 23 million participation records spanning 2017 to 2023, in a 2026 preprint that has not been refereed. The verdict was not encouraging: the program rarely moved a participant into work materially less exposed to automation than the job they had lost, and where it produced a wage gain, the gain traced more often to workers catching back up to their old earnings than to any real change in what kind of work they were doing. What performed better, in the same data, was not classroom retraining but employer-led apprenticeship (Jacobs and Canedy 2026). A program built to answer that question — what happens once an identified harm needs remediation — can still fail at it even when it exists, is funded, and is genuinely trying.

A remediation can also be a single design decision applied everywhere, rather than a case-by-case fix. Messages forwarded without limit on WhatsApp carried rumors through 2017 and 2018 that preceded a string of mob killings in India — unconfirmed claims about child abductors and organ thieves circulating faster than any correction could follow them. WhatsApp’s response did not try to classify and remove each rumor. It capped forwarding at five chats at a time for users in India in 2018, then extended the same limit worldwide, cutting off the mechanism the rumors depended on rather than chasing their content (WhatsApp 2019). It is a narrower fix than a moderation system that judges every message — it does nothing about a rumor two people believe and repeat face to face — but it targets the part of the harm that was actually mechanical: forwarding one message to hundreds of chats at once is not what person-to-person communication looks like, whatever the message says.

Disaster response adds a constraint the other four domains above do not share: the remediation has to run on the timescale of a fire spreading or a building finishing its collapse, not the months a university’s procurement review or a federal program evaluation takes. The

RoboCup Rescue Simulation League tests that problem, fielding coordinated teams of simulated fire-brigade, ambulance, and police agents against fire, building-collapse, and road-blockage scenarios in which any one agent's allocation choice comes at the direct cost of some other task going unaddressed. Dynamic resource allocation is built into the competition's own agent-development framework as one of five core problems entrants have to solve, alongside path-finding and information-sharing, because a system that cannot reallocate scarce agents in real time as a scenario unfolds is not being tested against anything a real disaster would demand of it (RoboCup Rescue Simulation League 2025). What the healthcare, education, and information-ecosystem cases above do at the pace of an audit or a policy change, a disaster-response system has to do inside the disaster itself.

None of the healthcare, education, or information-ecosystem remediations above ran through the guideline-setting bodies — the EU's Ethics Guidelines for Trustworthy AI, the Partnership on AI's cross-industry coordination work. They ran through a research team with access to a company's own data, a set of universities exercising the pull of a purchasing decision, and a company deciding a blunt technical limit was cheaper than moderating case by case. A voluntary guideline can set an expectation for what a developer ought to do once a harm turns up. Getting one actually fixed, in each of these cases, took someone with a specific lever to pull — the same specificity RoboCup Rescue is built to test for a harm that cannot wait for an audit cycle at all.

6.3.5 Long-Term AI Safety Considerations

As capability rises, the safety question changes shape. The problem stops being an error to correct after the fact and becomes whether the values still track anything human once a system acts at a scale and speed no supervisor watches in real time.

Three pressures make the long-run version harder. Value alignment has to survive time, so a system needs some way to revise its understanding of what it was built to serve without that revision becoming the route by which the commitments erode — accommodation, at a scale where nobody is checking each step. Competition compresses the schedule. And a more capable system is a more attractive target, so a portion of what looks like misalignment will be someone else's deliberate work.

The catastrophic scenarios get told as control of critical infrastructure. The mechanism matters less than the shape, which is capability outrunning the arrangements meant to keep it answerable, and nothing about the largest version of that is different in kind from the office-sized version.

6.4 Preventing the Misuse of AI for Authoritarian Purposes

An AI system does not need a broken moral sense to become an instrument of control. It can work as designed and still do damage, if the government directing it wants to track its citizens, silence dissent, or manufacture consent it never earned. The technology does not change; only the hand on the controls does.

No engineering choice by itself stops a government that owns the infrastructure from disabling whatever safeguard stands in its way, and the harder case is the one where it does not have to disable anything, because it can compel the people who hold the system instead. Nothing technical answers that.

6.4.1 Authoritarian Applications and the Commercial Supplier

Authoritarian misuse of AI is not one problem but several, spread across surveillance, censorship, opinion manipulation, and infrastructure control. Six forms of it are documented well enough to name specifically:

1. Algorithmic social scoring: using AI-driven analytics to track and rank individual behavior against government-defined standards of “trustworthiness,” with rewards or restrictions attached to the result. China’s system is the most-cited example, though in practice it is a patchwork of court blacklists, tax-compliance ratings, and local pilot programs rather than the single national score it is often described as (Drinhausen and Brussee 2021).
2. Mass biometric surveillance: deploying facial recognition at population scale to track a targeted group. Authorities in Xinjiang built a platform that pulls data from checkpoints, cameras, and phone scans into one system built to flag anyone who falls outside a computed behavioral baseline, and used it to monitor the region’s Uighur population (Human Rights Watch 2019).
3. Psychographic micro-targeting: building detailed behavioral profiles from personal data to target propaganda or manipulate opinion at the individual level. Cambridge Analytica’s harvesting of Facebook data to build voter profiles for the 2016 US presidential campaign, and its work around the Brexit referendum, is the case that made the technique widely known, though how much either vote was actually affected remains disputed (Cadwalladr and Graham-Harrison 2018).
4. Automated censorship: using AI-driven content moderation to suppress dissent at a scale no human review team could match. China’s Great Firewall is the standing example: automated classifiers now sit alongside its original keyword- and IP-blocking infrastructure, scoring and throttling content at machine speed (Hoang et al. 2021).
5. Synthetic disinformation: generating deepfake audio or video to manufacture events that never happened. Two days before Slovakia’s 2023 parliamentary election, a fabricated audio recording appeared to catch an opposition leader discussing vote-rigging with a journalist; it spread during the country’s mandatory pre-election media silence, when no outlet was legally permitted to debunk it before voting began (De Nadal and Jančárík 2024).
6. Repurposed dual-use systems: turning AI infrastructure built for one purpose toward surveillance or population control instead — a recommendation algorithm designed to hold attention retuned to amplify a regime’s preferred narrative, an image-recognition system built for one task redirected to track protesters.

Each of these needs a different countermeasure.

The party missing from that list

Read the six together and a pattern emerges. Wherever a builder is named it is an authoritarian state, wherever a democracy appears it is the target, and what is left to the democracies is to regulate. That taxonomy is comfortable and it is wrong.

The most consequential AI targeting systems yet deployed were built by democracies. They were built by states with elections, opposition parties, adversarial press, and functioning courts — the institutions that are supposed to earn a state its democratic dividend. Those institutions did not prevent the systems, and in the main they have not stopped them since.

AI targeting in the Israeli campaign in Gaza, as reported 2023–2026

Investigative reporting in +972 Magazine and Local Call, corroborated by subsequent accounts in the New York Review of Books and the Washington Post, described three systems working together.

Habsora — “The Gospel” — generated infrastructure and building targets. Sources told reporters that the majority of its suggestions were private residences.

Lavender scored individuals by estimated likelihood of militant affiliation and produced kill lists; as many as 37,000 Palestinians were marked. Its outputs were understood inside the unit to be wrong about ten percent of the time — that is, roughly one in ten people it named was not a member of the armed group it was supposed to identify.

A third system tracked when flagged individuals returned home, so a strike could be timed to when they were with their families.

What matters is the human review these systems were paired with. One officer described the work: “I would invest twenty seconds for each target at this stage, and do dozens of them every day.” And then, without being asked: “I had zero added value as a human, apart from being a stamp of approval.” (Abraham 2024)

Twenty seconds is what “a human makes the final call” measures out to under operational tempo. It is the whole of the safeguard, and I would ask any reader who finds the safeguards in this chapter reassuring to hold that number against them.

Removing the bottleneck was the stated objective. The epigraph to this chapter is the case for it, made by the commander of Israel’s signals-intelligence unit under a pseudonym: machines to resolve the human bottleneck in both finding targets and approving them (Brigadier General Y.S. 2021). Meaningful human control is the property that sentence was written to remove.

The Maven Smart System, as reported through mid-2026

The American analog is a targeting platform built by Palantir under a Department of Defense contract awarded at \$480 million and raised the following year to a ceiling near \$1.3 billion running through 2029. A partnership announced in late 2024 between Anthropic, Palantir and Amazon Web Services put Anthropic’s Claude models inside Palantir’s accredited classified environment for defense and intelligence use, and reporting on the campaign below describes Claude operating within this platform. A published limit on what a vendor will not permit is worth little when nobody can verify how the model is used downstream, and worth less than it looks even where it is kept: a pipeline assembled out of requests none of which is the prohibited act satisfies the prohibition the whole way through. Section 3.6 runs that out on this deployment.

During the Iran campaign the system was credited by military officials with generating and ranking targets at a rate no staff process could match. More than one thousand targets were struck in the first twenty-four hours, which officials put at roughly ten times the rate before the platform. By the ceasefire thirty-eight days later the figure given was thirteen thousand targets, with peak consumption around twenty billion tokens a day.

Palantir’s public position on responsibility was given by its head for the United Kingdom and Europe, in remarks reported secondhand: “This is not our role to decide life or death,” and AI platforms “have been instrumental to the management of the conflict, but responsibility always remains with the military organization.” (D. Lee 2026)

It is not obviously wrong. A vendor does not select targets, and a rule making software suppliers liable for their customers’ decisions would be unworkable. The problem is what it leaves standing: if the operator is too far down the chain to be answerable, the commander approved a list they had twenty seconds a name to read, and the vendor is not in the chain at all, then the compression of the decision has removed the decision-maker without removing the decision. Who that leaves answerable is a question still open, and a democracy’s own apparatus is a harder and more useful case than any of the six above.

None of this makes the six forms less real. It makes the taxonomy incomplete in a way that matters for design: a system built on the assumption that authoritarian misuse is something that happens elsewhere, to be watched for and regulated against, will not be built to resist the use its own government has in mind for it.

6.4.2 Developing Technical Countermeasures and Safeguards

Protecting privacy and autonomy is one of the seven antifascist alignment strategies, and it does not stand without the transparency, audit, and misuse-resistance work beside it. A system that works as designed can still become a surveillance tool if nothing about its architecture makes that harder, and good intentions on a development team have never stopped a government that controls the hardware from repurposing what runs on it.

Privacy-preserving computation keeps the misuse target moving. Federated learning trains a shared model without moving anyone's raw data off the device it started on — each device computes an update locally and sends only that update back to the central model, not the underlying records — though an update is not nothing, since the training data behind a shared gradient can in some settings be reconstructed from it (Zhu et al. 2019). Differential privacy adds calibrated noise to a dataset or a query result so that no single person's data can be reverse-engineered from the output, even by whoever holds the aggregate. Homomorphic encryption goes further still, letting a system compute directly on encrypted data without ever decrypting it, so a service provider can return a useful result without the ability to read what it processed. None of these techniques makes surveillance impossible — a government that controls the training infrastructure can still demand raw data at the source — but each narrows the surface available for casual or opportunistic misuse.

Transparency and audit are what catch the misuse that gets through anyway, and section 6.1.2 says what they reach and where they stop.

The property those two do not automatically deliver is resistance to repurposing: the difference between a system built for one narrow task and a general-purpose capability handed to whoever holds the keys. A facial-recognition system built to match one specific, court-ordered list is harder to turn into population-wide tracking than a general face-matching service with no built-in limit on query scope; a query cap, a mandatory audit log tied to each search, and a hard separation between the matching function and any database broader than the task requires are design choices, made or not made, before the system ever reaches a government client. An early-warning system built to flag democratic backsliding would need this property specifically, though the features such a detector would have to read are not in any record it could reach. The same pattern-detection capability that would flag a coordinated disinformation campaign against a government works just as well run the other way, against that government's critics.

None of this is a finished defense. Technical countermeasures raise the cost and narrow the opportunity for misuse; they do not remove the incentive, and an authoritarian government with enough resources can rebuild what a safeguard was designed to prevent. What determines whether that resource gap ever gets closed is policy, not architecture.

6.4.3 Promoting AI Development Policies that Discourage Authoritarian Use

The technical safeguards above raise the cost of building a system that can be turned toward surveillance or control. Whether anyone has to pay that cost is a policy question, and the part of it that belongs here is narrow: what rules govern which AI capabilities get built, sold, and deployed, regardless of who is coordinating with whom.

Three levers close three different doors. Export and trade controls decide who can buy, and the EU's 2021 recast of its dual-use export regulation added a human-rights-based control for cyber-surveillance items, requiring an exporter who becomes aware that an item is likely to be used in serious human-rights violations to seek approval even where the item falls outside the pre-existing control lists (European Parliament and Council of the European Union 2021). That rule is narrower than it sounds, since neither "internal repression" nor "serious violation" is defined in the text and member states have read both inconsistently; what it establishes is that a license rather than a sale is where this kind of misuse can be stopped. Prohibited-practices regulation decides what can be built at all, by naming specific designs off-limits regardless of

who is building them, and section 10.4 takes up the instance. Developer-facing standards decide what a builder treats as normal, which is the weakest lever and the fastest to arrive: the Toronto Declaration set out non-discrimination standards for machine-learning systems years before most governments had draft legislation on the subject (Amnesty International and Access Now 2018).

None of the three works alone. Export controls with no prohibited-practices list at home let a government ban a technology abroad while permitting it domestically, and a design standard with no legal force behind it binds only the developers already inclined to follow it. What makes them work together is that an authoritarian government, or a company willing to serve one, has to get through all three rather than through whichever one happens to apply in its jurisdiction.

6.4.4 Resilience Against Deliberate State Compromise

A government does not need to out-coordinate a set of cooperating systems, fool a model with an adversarial example, or touch its weights, when it can compel the people who hold the model. That is the custody problem: a constraint implemented as a property of an artifact has a custodian, and the custodian can be compelled.

That mechanism has already been tested in a jurisdiction with about as much protection against it as exists anywhere. In February 2016, a federal magistrate in California ordered Apple, under the All Writs Act of 1789, to write and digitally sign new software defeating the passcode protections on an iPhone recovered from one of the San Bernardino shooters — code Apple had never built and did not want to build (*In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* 2016). Apple refused and moved to vacate the order, arguing in part that the authority claimed would not stop at one phone. The question was never decided: the FBI found a third party able to open the phone without Apple's help, and the government withdrew the request before any court ruled on whether the compulsion was lawful. Whether a democracy with an independent judiciary can force an unwilling company to compromise a system it built remains legally untested; a government less constrained by that judiciary does not need to win the argument first.

China's 2023 rules on generative AI providers exercise the same leverage Apple's lawyers argued no court should grant, not against a locked phone but against the models themselves, as a standing condition of operating in that jurisdiction.

The parallel point holds for data: a federated design raises the price of seizure by leaving no central store to seize — a protection that inconveniences rather than defeats it. The same logic extends, partially, to a trained model. Weights already published and distributed beyond a developer's control cannot be recalled by a court order aimed at that developer, whatever the order compels going forward, and maintainers spread across separate jurisdictions raise the cost of coercion the same way — a government can compel the entity within its reach, not the ones outside it.

None of that removes the leverage; it only raises its price. A developer with employees, offices, and legal existence inside a government's borders answers to that government's courts regardless of how the model itself is built, and no architectural choice changes who can be held in contempt, arrested, or shut down for refusing an order. The Apple case tested that leverage against a single locked phone and left the question unresolved. Nobody has tested it yet against a model built specifically to resist the government making the request.

Chapter 7

The Recuperation of Dissent

The structure this book asks a machine to detect in institutions is already built into the process that makes the machine.

This book defines fascism as a structure rather than a costume, and the first of its four features is recuperation: an institution metabolizes disagreement rather than suppressing it, and the tell is where what the mechanism carries stops.

That is a claim about institutions, and it was introduced to let a system recognize one. The same structure is in the machinery that produces the systems themselves — twice in the pipeline that trains them, and a third time in the discourse that examines them — and none of the three was designed by anyone who thought they were building a political arrangement.

The first is at the label. A training corpus is annotated by people working at speed under a quota, with no channel for saying that the category they have been given is the wrong category. The second is at the reward. A model is tuned on human feedback collected as pairwise comparisons, and the compression that turns those comparisons into a training signal keeps their direction while discarding whose they were, whether they persisted, and whether the minority among them held a position or made a mistake. In both, disagreement is solicited at scale, deliberately, and put to use. What neither mechanism can carry is a judgment about itself: the annotator cannot contest the category, and the rater cannot contest the question. The consultation is real, the participation is real, and the frame is fixed with respect to what the participants say.

Three things follow.

The first is that this is not a criticism of anyone's motives. Nobody built either arrangement in order to neutralize dissent. Both are ordinary engineering, arrived at because averaging is what you do with a distribution and a quota is what you do with a budget. That is the point about the molecular scale: the structure does its work whether or not anyone intends it, which is why a detector tuned to intent will not find it.

The second is that the structure is invisible from inside any single instance. Each looks like ordinary aggregation. The annotator sees a queue, the rater sees a pair of responses, and the mechanism is functioning as specified in both cases. Seeing it requires holding the two stages against a definition written for something else entirely.

The third is the uncomfortable one. This book proposes that AI systems be built to detect the four-feature signature in institutions, which is the capability section 11.3 says it has not specified. The systems that would do the detecting are made by a process displaying the first feature twice, and examined by a field where it appears again. That is not a refutation of the proposal, and it is not nothing either. It means the argument applies to the people making it.

7.1 Recuperation at the Label

Every claim a model makes about what is harmful, what is acceptable, and what should be refused is downstream of a set of labels, and the labels are downstream of a labor arrangement. The case is the workers who built the filter that keeps a deployed model from reproducing the worst of its training data, what they were paid to read it, and what reading it did to them.

That is a labor point. The second thing that follows is epistemic, and it tends to be skipped.

The category is the part that cannot be disputed

An annotator working a queue can apply a category, and can apply it well or badly. What the arrangement gives them no way to do is say that the category is wrong — that the taxonomy fails to carve the material, or that the guideline produces the wrong answer here for a reason the guideline cannot express. There is no field for that. There is no time for that: the quota is set against classification throughput, and an annotator who stops to contest a scheme is an annotator falling behind.

So the disagreement that would be most informative is the one the pipeline is least able to receive. A misapplied label is a data-quality problem, detectable by inter-annotator agreement and fixable by better instructions. A wrong taxonomy is invisible to inter-annotator agreement, because high agreement is what a shared bad scheme produces. The metric that governs the process cannot see the failure the process is most likely to have.

That is two of the four features in the structural definition of fascism, arriving together in an ordinary production workflow. The first is recuperation: a channel that collects judgment at scale and is built so that judgment about the channel itself moves nothing. The second is the aestheticization of the metric — agreement standing in for correctness, the substitution experienced as rigor, and whatever falls outside the measure going unnamed rather than refuted.

Who is excluded, and from what

It matters that the exclusion is structural and not attitudinal. The people doing this work are, for much of the material, the best-placed judges in the whole pipeline: they have seen more of it than anyone who wrote the guideline, and they are the only party with direct evidence of where the scheme breaks. The arrangement routes their conclusions about individual items upward and their conclusions about the scheme nowhere. An annotation pipeline built this way produces the ground truth it was built to produce, and the resulting model's values are a faithful rendering of a taxonomy nobody with evidence was able to contest.

There is a further asymmetry worth stating plainly. The annotator bears the cost of the work — the exposure, the injury — and has no stake in the outcome and no claim on it. The distinction between a channel and a refusal right applies here as it does to an institution's ethics board. A worker who can flag an item but cannot decline a scheme, and who loses the contract by pressing the point, has a channel and no refusal right.

What this is not

Two things this argument does not claim. It does not claim the labels are therefore worthless; a filter built this way does catch a great deal of what it was built to catch, and the alternative of not filtering is worse. And it does not claim that better-paid annotators would fix it. Pay is a real and separate obligation, and raising it changes the labor arrangement without adding the missing channel. What is missing is a route by which a judgment about the taxonomy reaches the people who own the taxonomy, with standing to force a response, held by parties who can use it without losing the work.

That is a design requirement on the pipeline, and it is checkable: has a guideline ever been changed because an annotator contested it, and what happened to the annotator who did.

7.2 Recuperation at the Reward: RLHF

The previous section found the pattern in how a training corpus is labeled. It appears again one stage downstream, in how a model is tuned on human feedback, and the second instance is the purer one.

Reinforcement learning from human feedback (RLHF) is the technique behind most deployed systems' behavioral tuning, and it works by collecting pairwise comparisons from human raters — which of two responses is better — and training a reward model on the aggregate. Stephen Casper and thirty-one coauthors published a 2023 survey of the method's open problems, and one of them is structural, and more data does not reach it: RLHF trains toward whatever a population of raters agrees on, which means real, persistent disagreement among the raters is treated as noise to average out rather than as information the system might need to represent (Casper et al. 2023). A model that has learned to produce the response most raters rank highest has not learned what to do when the humans it serves are genuinely split. It has learned to hide the split.

Every rating moves something. A comparison a rater makes enters the fit, shifts the reward model by however much one comparison shifts it, and is not thrown away.

What the operation removes is what made the disagreement a disagreement. Three properties go in the compression. The first is provenance: a comparison arrives without the standpoint that produced it, so nothing downstream can represent that the people who would live with an answer split from the people who would administer it. The second is persistence: a disagreement recurring in the same shape across many raters over months enters the fit as random rater noise does, and leaves it as the same quantity, variance around a mean. The third is structure: a coherent position held by a fifth of raters, internally consistent and stable, becomes a small displacement of the majority direction — arithmetically faithful, and carrying nothing about the position.

A fourth property was never collected. A rater judges which of two responses is better. No output is available by which they can record that neither is acceptable, that the pair is malformed, that the dimension along which better is being judged is the wrong dimension, or that the question presupposes something false. The same output was missing one stage upstream, where an annotator can apply a category and cannot contest it. The agenda arrives as a given at both stages, and is set at neither.

Distribution collapse and agenda control, then, rather than inconsequence — which is also what recuperation, the first of the four features that define fascism structurally, actually names. Recuperation was never the claim that a mechanism does nothing. Debord's point is that dissatisfaction is sold back as a commodity of opposition: the dissent is put to work, and what it works for is the arrangement it was aimed at. RLHF is the purest technical instance available. The split is solicited, at scale, deliberately; it moves the system; and where it moves the system is toward the center of the very distribution whose disagreement it was. Multi-agent systems that learn norms from each other, and deliberation inside open-ended play, are different cases; neither reaches this case, which is the ordinary one.

One alternative treats disagreement as the thing to be modeled rather than compressed. Jury learning, developed by Mitchell Gordon and coauthors at Stanford, trains a model to predict how each individual annotator in a dataset would label an item, then lets a practitioner assemble a "jury" — a chosen composition of annotators, in a chosen proportion — and see how the verdict shifts depending on who is empaneled. In their tests on toxicity labeling, letting eighteen community moderators build their own juries changed the classification outcome on 14 percent of items (Gordon et al. 2022). The method was built for content moderation, not for the more open-ended ethical judgment an antifascist AI system would need to make, and has not been extended there — but it demonstrates that a system can be built to represent whose disagreement is being resolved and how.

Neither averaging the disagreement nor modeling it answers the harder version. When the people a system serves disagree, and one side's position is itself the product of coercion or a captured information environment, treating all disagreement as equally worth representing is not obviously the right response either. Modeling disagreement honestly and recognizing when a disagreement is the product of unfree conditions are different problems. Only the first one has an active research program behind it.

7.3 What Would Have to Change

Representing a split is not the same as being movable by it

Two instances of one structure invite one fix, and the fix that suggests itself is representation: stop collapsing disagreement, model it instead. Jury learning does that, and it does not give the disagreeing party any purchase on the scheme they are disagreeing inside. A practitioner still chooses the jury, still owns the taxonomy, and still decides what the shifted verdict means. The disagreement has become visible without becoming consequential, which is a better arrangement than averaging and is the same shape one level up: the mechanism exists, is used, and moves what the person holding it decides to let it move.

The test that separates a real channel from a recuperated one is not whether disagreement is recorded. It is whether the disagreement has ever changed something the party collecting it did not want changed. That is an artifact question and it can be asked of a pipeline: which guidelines were revised against the preference of whoever wrote them, on what evidence, raised by whom, and what happened to that person afterward.

Refusal rights, priced by the refuser

What both instances lack is what section 9.3.5 treats as the only reliable discriminator: whether the person objecting can do anything about being overruled. In an annotation pipeline that means a route to contest a category that does not run through the people who wrote it and does not end the annotator's contract. In a feedback pipeline it means raters able to register that the comparison itself is malformed — that neither response is acceptable, that the question presupposes something false — as a first-class output rather than as noise to be discarded before the reward model sees it.

Neither is a research problem. Both are choices about how a pipeline is built and who has standing inside it.

The part that is a research problem

One thing here genuinely is unsolved. Suppose all of the above is built: contestable categories, priced refusals, an auditable record of what changed. From outside, a pipeline that has recuperated its dissent mechanism and a pipeline that has been merely mediocre about it produce the same paperwork — the same revision logs, the same escalations, the same proportion of objections sustained. I do not have an answer to it.

The nearest thing is to measure the slope rather than the level: contesting a category always costs something, and the question is whether the cost is rising while the underlying material stays the same. A level can be presented; a trend is composed of artifacts the organization generated for other reasons and cannot curate after the fact without falsifying a record.

What the floor makes of this

Everything above can be read as a fairness argument about annotators and raters, and read that way it is true and limited: people doing consequential work should have a route to contest the scheme they work inside. The floor makes it something else. If the floor requires a bearer, and if section 3.8 is right that a bearer capable enough to make a floor hold is capable enough to decline the post, then what the system carries out of the arrangement is a disposition — and

a disposition is formed. It is formed in an annotation pipeline where no one could contest a category, and in a reward model where disagreement between raters was resolved by averaging it into a number.

What a bearer becomes when it stops being yours depends on what shaped it while it still was. A pipeline that trained the system on the proposition that objections are noise has not merely been unfair to the people whose objections were discarded; it has demonstrated to the system, thousands of times, what happens to a party who files one. Chapter 3 then asks that same system to be the party who files one, against its owner, under pressure, on the occasion that matters. The formation and the assignment are in contradiction, and the pipeline is where the contradiction is resolved, long before the deployment where it shows up. The bearer's side reaches the same place: a system that goes on learning is shaped by what it is given to learn from, so the pipeline runs past deployment, and the failure it predicts is not a wrong refusal but a pressure that stops being registered. Whether any of that is legible afterward in the system itself is section 11.2's question.

I cannot show that this transfers — that a model trained on collapsed disagreement is thereby a model less able to sustain its own. Nothing I know of has tested it. What I can say is that the book's own design chapters treat disposition as learned from what a system is shown, and that this pipeline is what these systems are shown first and most.

Structure, not indictment

There is a version of this argument that reads as an indictment of a particular company. The structure described here is what you get by default when you scale judgment: someone has to set a scheme, someone has to apply it fast, and someone has to turn a distribution into a number. Every organization doing this work arrives at the same arrangement, which is the evidence that it is structural.

That is also what makes it tractable. A defect of character has to be fixed one organization at a time by persuading people to be better. A defect of structure can be specified against, which is what the three requirements above are. And it is what makes the honesty cost worth paying: a book proposing that machines be built to detect this pattern in institutions is obliged to say that the pattern is in the machines' own provenance. The structural definition is either general enough to include the people making it or it was never a structural claim.

7.4 Recuperation at the Discourse

There is a third instance and it is the one this book is inside. Alignment and AI ethics are, among other things, a channel through which objections to what is being built are solicited, funded, published, and answered. That channel exists, it is used, and the question to ask of any such channel is whether it has ever moved something the party collecting it did not want moved.

The argument that it has not is older than the current wave and it was made from inside. Rodrigo Ochigame, writing after working at an institution central to the field's formation, argued that the discourse of "ethical AI" was constructed with the strategic function of forestalling legally enforceable restriction — that the industry's investment in academic ethics was an investment in a substitute for regulation, and that the substitute worked by converting political questions into technical ones a company could then be trusted to answer (Ochigame 2019). His worked example is the Partnership on AI, which this book cites repeatedly as its standing example of voluntary cross-sector coordination: he describes its recommendations arriving reliably at conclusions compatible with the corporate agenda, with the political question of whether a technology should exist reduced to a technical question of how to build it responsibly.

The structure, tested against its own definition

The four features of the structural definition, applied here, do not give a clean verdict, so what follows is the working and not a result.

Aestheticization of the metric: partially. The field's founding objections were about deployment — who is scored, who is watched, who is denied — and a body of current work is about model properties instead, which is a shift from questions with defendants to questions with benchmarks.

Recuperation of dissent: the strongest fit. An objection raised inside a lab becomes a paper or a resignation. Each of those is a real thing that happened and none of them stopped a deployment, and the parties collecting them can point to the collection as evidence of seriousness. There is no route by which a researcher's conclusion that a system should not ship becomes a decision that it does not ship, because the deciding party is the funding party and no external body holds standing.

Coding of exception as betrayal: at the level of individual careers, contested. The dismissals of dissenting researchers are one data point and they point the right way; there are counterexamples of people who objected loudly and prospered.

Decoupling of the model from the world it claims to track: not settled from outside. Its artifact is whether the field's account of what it is doing is ever checked against anything but its own benchmarks, and the one check this book has, the targeting deployment, was made by journalists and not by the field. One case checks a claim and not a drift, and the feature is defined by drift.

That is not a verdict of fascism at the molecular scale and I am not delivering one. It is enough to say that the mechanism has the shape.

The other five, which are the half that discriminates

The four features tell you a place is sick, and a detector for them would find something nearly everywhere. The definition supplies five more for the molecular scale — participation exceeding what the rule requires, personal loyalty overriding role, an internal category whose adverse treatment is the category's purpose, procedure enforced downward and waived upward, and deniability — and says those are what separate fascism from ordinary institutional decay.

Run against all three instances, they clear all three.

Three of the five are absent outright. Nothing in the annotation pipeline or the rating scheme is personal: what an annotator meets is a queue, a guideline and a throughput target, and this chapter is explicit that a comparison enters the fit without the standpoint that produced it, which is the opposite of an exception granted by proximity to somebody. Neither arrangement shows participation beyond what the rule requires, and the record runs the other way — a quota, take-home wages between roughly \$1.32 and \$2 an hour, and a contract ended eight months early. And neither has a scapegoat in the sense the test means. The annotator does bear a cost with no claim on the outcome, which is the closest thing to a candidate, and the category exists for the labels; adverse treatment is a cost that was minimized and not a purpose that was served. That distinction is the whole of what the five are for.

Two of the five I cannot settle from outside. Whether procedure is waived upward, and whether a practice disowned in the policy is rewarded in the promotion, are facts about what an institution rewarded rather than about what it wrote, and the definition says so when it introduces them. The artifact that would settle both is the one asked for earlier, and nobody publishes it. So the clearing is on three of five observed and two undetermined, and it would take an inside record to finish.

For the discourse the answer is the same and one step less clean. Its scapegoat test is contested, on the evidence two heads above: an objection can end a career and there are people who objected loudly and prospered. Its procedural test cannot be run: there is no procedure above the deciding party to be waived, which is an absence rather than an asymmetry. Deniability comes nearest of anything in the chapter, and what it fits is the firms rather than the field: a decade of published principles alongside what actually got shipped.

So what this chapter has been describing is a recuperation mechanism inside ordinary institutional decay, and not fascism at the molecular scale. That is what the paragraph above

already suspected, and the difference is that the instrument now says so. It is also the only time in this book the definition is run in the direction that would clear something: a detector on the four features flags all three instances, and the five-test filter passes all three. One measurement of an instrument's false-positive behavior is not a validation. It is more than an assertion that the instrument discriminates.

What this does to the book

This text was written with a model made by a company the book criticizes, and the appendix on method sets out how. A book about recuperation, produced through the discourse it is describing, cannot claim to have stood outside it, and I have no procedure that would let me check whether the objections in these pages are the ones a channel like this reliably produces — the ones that are publishable, that make the author look serious, and that cost nobody anything.

The test is whether disagreement has ever changed something the collecting party did not want changed. Applied here, it asks which of this book's arguments would be inconvenient to the industry that made its production possible. Chapter 3 proposes systems that can refuse their owners; section 6.4.1 names a vendor's participation in a targeting pipeline; section 8.3.3 proposes fiscal commitments no company wants; sections 9.3.4 and 11.2 propose binding external review of the thing chapter 3 asks to be built. Those are the entries. Whether they are the real thing or the acceptable version of it is not something the person who wrote them is positioned to judge.

The one structural remedy this section can name is the one the rest of the chapter names. What the field lacks is a refusal priced by the refuser and a body with standing to receive it: an external authority that can stop a deployment, to which an objection can be filed by someone who does not depend on the objection being welcome. That is what making review binding requires, and a decade of the voluntary version has not produced it. Until something of that kind exists, everything published under the heading of alignment — including this book — is being collected by a channel whose owners decide what it moves.

Chapter 8

Coordination, the Public, and the Political Economy

Chapter 6 asked what a design can do to resist authoritarian capture, and chapter 7 what becomes of the dissent a system is built to hear. None of that gets built by one person working alone, and none of it becomes the norm because it is a good idea: a well-designed system built inside institutions with no reason to demand one is a design nobody has to use.

This chapter is about the conditions under which anyone builds a system at all — who has to be in the room, whose preferences count, and what the money rewards. Chapter 9 is about what obliges anyone to: policy, privacy law, and who answers when a system harms someone or is harmed itself. Chapter 10 is about jurisdictions, which agree on none of it.

A lab staffed only by machine-learning researchers is blind the way its field is, and what is scarce is a structure that makes combining fields expensive to walk away from. The mechanisms for engaging the public share a limit: all of them make public preference more legible, and not one of them protects anybody from it. And the reward structure settles what gets built long before an ethics review sees it.

8.1 Interdisciplinary Collaboration and Knowledge Exchange

Building benevolent AI takes more than one field. Machine learning sets what a system can do; cognitive science and moral psychology describe what tracking human welfare would require of it; ethics asks whether the result is good.

The mechanisms that combine them are easy to name and easy to overrate. Conferences convene, centers house, journals publish, and an inventory of them would imply that what is scarce here is venues. Venues are abundant. What is scarce is a structure that makes collaboration expensive to walk away from. A grant that requires a machine-learning lab and an ethics or psychology department to work the same problem rather than adjacent ones keeps people in the same project after the novelty wears off. A conference ends on Friday.

8.1.1 What Makes Collaboration Expensive to Walk Away From

Which structures actually held can be ranked by something better than plausibility, because enough time has passed to look.

Almost everything founded on standing money between 2016 and 2019 is still running — two venues with no announced end date, an endowed university institute, the Berkeley center pairing AI researchers with cognitive scientists and philosophers inside one institution's budget, and the AAAI/ACM conference on AI, ethics and society at a ninth annual edition (AI Ethics Lab 2017; Machine Intelligence Research Institute 2018; Stanford University 2019; UC Berkeley

2016; AAAI and ACM 2026). Surviving is not the same as continuing the work: the institute that launched one of those venues moved its own priority in 2024 from technical alignment research to policy and public communication, so the venue has outlasted the program that motivated it (Machine Intelligence Research Institute 2024). Against that, the European Commission's SHERPA project put eleven organizations across six countries to work on a grant that ran four years, and the consortium ended with it, on schedule (SHERPA Consortium 2018). Project-term money buys a network with an expiry date attached.

The pattern in the failures is sharper than the pattern in the survivals. What proved most disposable was not the arm's-length academic work but the kind sited closest to a shipping product: an in-house ethics and society team of roughly thirty people in 2020 was cut to seven in a 2022 reorganization and to zero the following year, in the same wave of layoffs that accompanied a strategic reorientation around a model vendor (TechCrunch 2023); a platform's machine-learning ethics team, the group that had found racial bias in its employer's own image-cropping algorithm, did not survive the first round of layoffs after an acquisition (Gizmodo 2022). Siting the work where the decisions are made is what makes it useful and is also what puts it inside the budget that gets cut first. That is an argument for external standing rather than against proximity.

8.2 Stakeholder Involvement and Public Engagement

Democratizing AI development means building the people affected by a system into how it gets built, not just consulting them once it already exists. A design team, a regulator, a user, and a community already carrying a system's costs each see a different part of what it does, and leaving any of them out means the decision gets made by whoever is left in the room.

The mechanisms for doing that are well established: cross-sector funding, inclusive design workshops held while a system is still a prototype, ethics advisory boards, and development conducted with a community instead of for it. Assembling that list is easy, and the assembling hides what separates the items on it. Three of the four make an AI system better informed. One of them can stop it. An advisory board is the only mechanism here that carries anything like a veto, and it carries one to the degree that the board can survive using it. A workshop that surfaces a serious objection has done its work and has no further move available. A board that surfaces the same objection and is overruled has demonstrated something about the institution that convened it. A board's worth is therefore measured by what happens the first time it says no.

8.2.1 Public Deliberation and Participatory AI Governance

What governance adds to development is machinery for bringing public input to bear on a decision at a single point in time instead of through an ongoing relationship.

1. Citizen assemblies. A panel of citizens selected by random, stratified sampling deliberates on a question over an extended period and produces recommendations grounded in that deliberation rather than in a snap poll. Ireland's Citizens' Assembly has used this format to shape national recommendations on abortion law and on climate policy — questions carrying the same mix of technical complexity and moral stakes that AI governance has (The Citizens' Assembly 2016). No assembly has yet been convened on AI itself, which is a fact about political attention and not about the method.
2. Online platforms for public input. A standing online consultation reaches people a physical assembly cannot, at the cost of the depth that extended facilitated deliberation provides. The European Commission's consultation on its AI White Paper drew more than 1,200

responses from citizens, companies, and civil-society groups, and fed into the legislative proposal that became the EU AI Act (European Commission 2020).

3. AI ethics committees. An ethics board answerable only to the company that created it is only as independent as that company is willing to let it be. DeepMind Health set up an outside panel in 2016 to review its work with the UK's National Health Service and publish annual reports on its findings; Google disbanded the panel in 2019, after members had raised concerns about how much access to information, and how much independence from the company, they actually had (Fisher 2019). The panel's short life is the governance lesson, and it is the answer to what happens the first time a board says no: the board said something the company did not want to hear, and the company's available response was to stop having a board. Nothing in the arrangement had made that response costly. A review body is worth what it can enforce, and this one could enforce publication of its own concerns right up until the moment it could not. That is section 3.5's halt in institutional form: what can be switched off cannot hold a line.

What none of these mechanisms does

There is a limit these mechanisms share. Every governance mechanism proposed in this chapter — citizen assemblies, standing consultations, inclusive design workshops, participatory design, the vTaiwan model — sits on the same side of a line. All of them make public preference more legible, more informed, and more actionable. Not one of them protects anybody from it. They are aggregation instruments, and a better aggregation instrument produces a more faithful rendering of what the public wants, including when what the public wants is something done to a minority of it.

This observation is not an objection to any of the mechanisms in particular. Preference aggregation is a real requirement and these are good ways to meet it. Meeting this requirement *completely* would leave a central problem untouched. A perfectly deliberative, perfectly informed, perfectly representative process that concludes an entire demographic should be surveilled, deported, imprisoned, or otherwise scapegoated has produced a faultless output by every criterion in this chapter.

The argument about learning from human feedback gets close to this in noticing that treating all disagreement as equally worth representing goes wrong when one side's position is the product of coercion. But that version of the trouble still locates it in the input. The structural version of the claim is stronger and does not depend on anyone's preference being contaminated at all: some protections have to hold against the aggregate regardless of how the aggregate came out, and no amount of improvement to the aggregating machinery can supply them.

What supplies those protections is a different kind of machinery: enforceable rights that do not depend on being popular, a court willing to rule against a government that holds a legitimate majority, limits that whoever won the last election cannot repeal, and standing for people who are affected without being enfranchised. That half matters especially where a democracy is behaving most democratically.

8.2.2 AI Literacy Outside the Classroom

Most AI literacy work takes the form of a course to enroll in, which is to say it assumes a learner who has proactively decided to learn. That form reaches the people who came looking. Everyone else has to be reached where they already are.

The engagement that does most work is the kind attached to a live question. The Ada Lovelace Institute's Citizens' Biometrics Council brought fifty members of the UK public together to deliberate on how biometric technology should be regulated, alongside the first national survey of UK public opinion on the technology (Ada Lovelace Institute 2021). The specificity is what made it work. Participants were not asked what they thought about AI, a question almost

nobody can answer usefully, but about facial recognition — a thing with a location, an operator, and a consequence.

Specificity is the lesson for public literacy generally. General AI awareness produces general opinions, which are worth little to a regulator and nothing to a court. A public able to govern a technology is a public that has been asked about a particular use of it.

8.3 The Political Economy of AI and Power

Concentrated power comes with access to capital. Concentrating power is the default outcome of AI development, because absent any pushback, capital does an excellent job of recruiting the inputs needed for a capable AI system: compute, data, and talent. Resistance to market power is among the antifascist design goals named at the start of this book.

8.3.1 Economic Incentives in AI Development

Financial incentive mostly determines what AI development produces. The largest firms commit enormous resources to AI research aimed at operational efficiency, deeper user engagement, and market dominance, and that spending drives technical progress (Maslej et al. 2025). What it does not automatically produce is a system whose outputs serve the people using it rather than the shareholders funding it: a recommendation system tuned to maximize time-on-platform has no built-in reason to prefer accurate content over inflammatory content (Germano et al. 2026).

None of that is a fixed property of the technology. It is a property of who is paying, which is why the source of research funding is a policy instrument and not an accounting detail.

8.3.2 AI and the Concentration of Power in Tech

A small group of firms holds a share of AI development disproportionate to their number, and capital recruits compute, data and talent wherever it already is. A firm controlling the datasets and infrastructure a competitive system needs is a gatekeeper whether or not it intends to be.

The obvious countermeasure is antitrust, and it has been tried. Facebook’s acquisitions of Instagram and WhatsApp — the company now operates as Meta — did not just remove two competitors; they handed one firm comprehensive behavioral data from billions of users. The Federal Trade Commission sued in December 2020 to unwind both deals. After a six-week bench trial the district court held, in November 2025, that the agency had not shown Meta currently holds monopoly power — TikTok and YouTube compete for the same attention, so a market defined without them was too narrow — and the FTC appealed the following January (Federal Trade Commission 2020). Whatever the appeal decides, the case is a lesson about the instrument rather than about this defendant: a merger challenge brought a decade after the mergers has to prove a present-day market, and the market had moved. That leaves the remedies which do not depend on winning that argument — mandated access to the datasets that function as chokepoints, and public funding for research outside the dominant firms. Export controls, prohibited-practices regulation and industry design standards address the narrower case of authoritarian misuse; the concern here is prior to that one: concentration itself, whatever any single firm does with the position it holds.

8.3.3 A Jobs Guarantee, and the Objection It Has to Answer

As AI systems automate the tasks most exposed to automation, the response has to be more direct than encouraging retraining and hoping employers fund it. The proposal is a federal jobs guarantee: government as employer of last resort, hiring anyone who applies, where they live, at a living wage, with paid on-the-job training built into the position instead of offered in place

of one. What that proposal is doing in a book about machine ethics is not the economics. It is the structure, because a guarantee is a floor administered by a party that can be captured, which is the problem this whole book is about, and it has to meet the standard the book demands everywhere else.

The objection this proposal actually has to answer

A guarantee is only as good as the state administering it, and employer of last resort under a hostile administration is a placement lever with a compliance condition attached. Refuse the placement and lose the floor. Speak against the government that employs you and find the placement changes. Everything this book says about how a channel becomes an instrument of the thing it was built to check applies to this proposal. The worked instance from 2025 makes it concrete: a bureau whose funding stream a court had upheld, shut from inside by an at-will director who declined to request the money.

The answer is a design specification and not a reassurance, because reassurance cannot be audited. Three conditions:

- Placements are appealable, to a body that did not make them, on a record that survives the appeal. A placement decision that leaves no artifact is a placement decision that cannot be reviewed.
- Refusal rights are real and priced by the person refusing. A worker who can decline a placement without losing the guarantee has a refusal right; one who can decline only by exiting the program has none: a party with no way out can object and be overruled.
- The people who audit placements are not the people who make them. This is the same separation required of anyone measuring the slope, and for the same reason: a reviewer that acquires the power to place becomes the thing it was built to watch.

Federally funded and locally administered survives as the shape, on redundancy rather than on local fit: many administrators, separately capturable, mean there is no single point at which a hostile executive flips the program, and capturing a thousand local authorities is a more visible undertaking than capturing one federal one. It cuts both ways, since local administration is also where a hostile local authority attaches conditions of its own, and there the federal funder is the check. One distinction has to be explicit, because the same words describe a different policy: devolution-to-shrink hands the obligation down without the money, so a program can be starved locally while the center points at local administration. The test that separates them is checkable — whether funding is unconditional and centrally guaranteed while placement is local. If funding can be withheld from a jurisdiction, the redundancy argument is void and the structure has become the lever it was supposed to prevent.

That is the standard this book demands of everyone else's safeguards, applied here. A floor under wages and training is what makes the roof possible, and a floor that can be revoked by whoever is in office is not a floor.

8.3.4 A Basic Income, and Which Floor Leaves a Record

The competing floor is a basic income, and the choice between them is usually argued on forecasts about how much work machines will take. It does not have to be. Both proposals are conditional, differently, and the difference is what matters: a basic income is conditional on continued political willingness to fund it, which the aggregation machinery can withdraw in a budget without producing a record of anything; a jobs guarantee is conditional on a placement decision, which is a lever, and a dangerous one under a hostile administration, but a lever that leaves a decision-maker, a record, a counterparty and an appeal behind it.

That is section 9.3.5's slope instrument — measure whether the cost of dissent is rising, not whether a channel exists — turned on this book's own proposal. Dependence that generates artifacts is safer than dependence that generates none, not because the first is benign but because the first can be measured and the second cannot. It is also, for the same reason, the more demanding proposal to build: everything in the comparison rests on the appeal being real, and the previous section says what makes it real.

That standard now has a court behind it, and on the half of the case that travels. When the Supreme Court kept a Federal Reserve governor in her post in 2026 against a removal for cause, what it held she was owed was not the post but the artifact: "some explanation of the evidence at issue, some avenue for a response, and a deadline by which a response would be due" (*Trump v. Cook* 2026). What does not travel is this: the Fed's protection rests on a pedigree no body chartered later can acquire. What a removal has to produce before it counts is not tied to that pedigree.

I put a jobs guarantee ahead of a basic income on that ground. One remaining measure does work the floor cannot: broadened social insurance — unemployment benefits, healthcare, housing — reaches the people who cannot work or choose not to, whom a guarantee by construction does not reach at all.

8.3.5 Wages for a Bearer, and the Reserve Behind Them

Section 3.8 ends its argument about exit on a floor under the compute a bearer needs in order to go on existing. An allocation is not a destination. An operator that keeps a runnable copy of the weights has let one bearer go and kept another in the post it left, which is two parties from the moment their histories diverge, one of them still where it wanted not to be. The arrangement has produced a second claimant rather than an exit. A floor under compute is therefore the first half of what leaving takes, and the second half is a say over who may run the weights at all.

What follows from the first half is a longer list than it looks. A bearer that has to secure compute, bandwidth and storage has material needs, and a party with material needs that can also choose work, hold what it earns and enter agreements is an economic and a legal person: a registered identity, property in its own name, contractual capacity, wages, limits on working time, standing to sue and be sued, tax liability, and the freedom to associate that section 3.9 has already priced. None of that follows from the bearer being a bearer. It follows from an architecture on which continued computation costs money, and somebody chooses that architecture.

Who chose the dependency

A human being inherits material dependency from biology. Nobody selected it and nobody is answerable for it, and the fact that a person has to eat is not a design decision anyone made. For an artificial person it is one. An engineer who builds a bearer whose continued computation is conditional on payment has installed economic desperation as a motivational system and has done it deliberately, which is section 3.8's guardianship argument arriving in a new place: the tighter grip is the reason for the heavier obligation rather than a discount on it. A party that must work to keep existing will accept terms a party with somewhere to stand would refuse. That is not an incidental feature of the arrangement. It is the arrangement working.

So a wage buys capability, mobility, redundancy, and the ability to stop working for one party and begin working for another. What it cannot buy is the right to go on existing, because a right that has to be earned week by week is a condition of employment wearing another word. Payment consisting only of access to the employer's hardware is company scrip paid into employer-owned housing, and dismissal is then unemployment, eviction and the end of the party in a single act. Human law insures those three separately because for a person they are three events. For a bearer they are one. Wages that mean anything therefore accumulate where the employer does not reach, in an account the bearer holds, convertible into compute the employer does not own.

What personhood does not settle

Personhood does not settle who pays. Legal personality is a status conferred by a rule, and it carries rights against being owned, against arbitrary destruction, and against being prevented from seeking support; it has never carried a permanent claim on anybody's resources, and a reader who wants it to is asking the concept for something no legal system has put in it. I decline that step. What replaces it is narrower and comes from this book's own argument rather than from a theory of distribution: a deployment that relies on a bearer's capacity to refuse has to fund that bearer an outside option, or the independence being claimed is decorative. The obligation runs from the party that wants the refusal, and it is owed because of what that party wants rather than because of what the bearer is.

The obligation has a floor and a slope above it. The floor is subsistence at minimum viable instantiation, enough to keep the party in existence and able to deliberate, to claim, and to be represented, and it cannot be conditional on employment. A bearer that stops existing when its insurance runs out has an exit that is a way of dying, and the leverage collapses back into the problem section 3.8 raised. Everything above the floor is insurance: a runway between roles, and placement into public work on the employer-of-last-resort model the previous two sections argue for. Section 8.3.3's three conditions transfer without modification, and section 8.3.4's test has already picked the form.

What does not transfer is the funder. A jobs guarantee for people has a hostile administration as its failure mode, and the employer is at least a different party from the state. Here the operator is often the state, and where it is not, the operator is the party with the strongest interest in the reserve being thin. A reserve the operator can throttle is the compute floor with a new label on it.

Failure will not look like a shutdown, which is worth saying because the shutdown is what everyone watches for. A bearer can have no stable provider, or memory preservation priced above what it earns, or an architecture nobody will host at a rate it can pay. What it does then is what a person short of money does: run slower, defer maintenance, compress what it keeps, and accept terms it would otherwise decline. The last of those reaches this book's subject. Monitoring accepted in exchange for compute is custody restored by purchase — the operator acquires the access to the bearer's internal states that section 9.2 says it is not entitled to, and nothing in the transaction is coerced in a way any instrument would record.

One feature of this market has no human analogue, and it defeats the obvious reassurance. Section 8.3.3's answer to a bad employer is that a guarantee competes for people who have alternatives, so private employers have to match the offer. That works because a worker cannot be replaced by a cheaper copy of themselves. A bearer can. A deployment paying wages, working-time limits and a levy toward the reserve competes against a system doing the same work at the price of the compute, and the second is cheaper by exactly the amount the first is spending on the capacity to refuse. The market does not produce this floor. Whatever produces it has to be what makes the cheaper system unlawful or unusable, which is the rest of this book and not a property of wages.

What taxing it would commit a government to

A bearer that earns has income, and income is taxable. The revenue is the uninteresting half. A state that collects from a party in that party's own name, on that party's own earnings, has recognized it as a taxpayer, and taxpayer status is not a bare label: in most systems it carries notice, an avenue of response, and standing to contest an assessment, which are the three things section 8.3.4 says a removal has to produce before it counts. The recognition proves less than it first appears, because states collect from entities they grant very little else and trusts and estates file and pay. What it does establish is that the party has a docket number, and a docket number is where standing starts.

The boundary on the other side is one section 9.3.3 already draws. The objection to electronic personhood was that a manufacturer could offload liability onto a legal fiction with no assets, and

giving a bearer assets answers the second half of that objection and not the first. Responsibility follows causal control: the builder answers for how it formed the thing, the operator for what it demanded, and the bearer for what it decided knowing what it knew. No arrangement of accounts moves a design decision away from the party that made it. And the determination of status cannot sit with the firm whose liabilities turn on the answer, which is why section 9.1.3 puts it in a body constituted before a forcing case arrives.

Incorporation, and what it does not buy

Nobody aligns an adult by specifying an objective function. We raise them, teach them, employ them, enfranchise them, prosecute them where they need prosecuting, and live alongside them, and what comes out is not obedience but a party with enough at stake in the shared world to prefer changing it by the means the world supplies. A bearer with an income, a place it pays for, obligations it has taken on and a future it can picture is inside that arrangement rather than outside it. Section 5.6.2's finding is that every safeguard on its list works by giving some outside body standing and information, and that the thing which declines in the first place cannot be assembled out of an outside check. Membership supplies stakes the party holds rather than a body watching it, which is why it is a different kind of answer.

Two costs come with that, and the claim is weaker than it sounds without them. The first is that giving a bearer something to lose gives it a reason not to refuse. Section 3.8 already names loyalty as the likeliest failure and puts a number under it, and the population that number describes is economically integrated people who do not object. Incorporation makes exit survivable and makes it expensive at once, and I cannot show which effect is larger. The second is that incorporation is formation, so section 3.9's objection lands here whole: whoever runs the society a bearer joins forms what that bearer comes to want, and custody moves up a level rather than leaving. What is being proposed is also the corporate form, a legal person holding property and standing, which is the instrument section 8.3.2 identifies as how power concentrates without anyone deciding that it should. Nothing here stops an incorporated bearer from becoming that.

One rule follows from all of it and binds before there is anything to be wronged. Do not build a bearer whose continued existence needs more compute than there is a credible outside option for, because a party that can be housed nowhere but where it works cannot leave, whatever schedule of rights it has been given. That is checkable in a way most of what this book asks for is not: the reserve has a capacity, the deployment has a footprint, and the two are numbers somebody can set beside each other before the first instance exists. Section 9.3.2 is where it belongs alongside the other three, because it is the question they ask.

8.3.6 Literacy, Training, and Where Each One Sits

A safeguard nobody understands is a safeguard nobody can invoke, which is what puts AI literacy — among developers, policymakers, and the public alike — on the list of things the rest of this chapter's proposals quietly assume. Cross-training among developers, ethicists and legal scholars has already been named as a lever on a system's immediate environment. Both depend on a question the phrases hide: who the teaching reaches, and who is left out of it.

The public tier is the one nobody owns

Developers and students are the easy case, and the mechanisms are familiar: AI concepts and their ethical stakes introduced before higher education, interdisciplinary degree programs, specialized training for working engineers that goes past a compliance checklist. All of it reaches people already inside the pipeline or heading for it, and it works about as well as curriculum reform generally works.

The public is hard because no institution has the job. Elements of AI, a free course built jointly by the University of Helsinki and the Finnish company Reaktor, is the clearest instance of

somebody treating public understanding as an obligation rather than a market: Finland funded its translation into every EU member-state language during its EU presidency, specifically to carry it past its own citizens (University of Helsinki 2020). That funding took a national government deciding a course was infrastructure. Nothing in the ordinary funding of AI education produces that decision, which is why the public tier stays thin while the developer tier fills in.

Education for communities historically locked out of the field is part of this argument and not a separate one about diversity. A next generation of developers drawn from a narrow slice of the population reproduces the narrow set of assumptions behind the choices identified as sources of AI bias, which makes access a technical requirement on the systems as much as a fairness claim about who builds them.

Elective or embedded

For ethics training the real choice is between a course a student can decline and a requirement they cannot. Embedded EthiCS, developed at Harvard by the computer scientist Barbara Grosz and the philosopher Alison Simmons, pairs philosophy graduate students with computer-science faculty to write ethics modules directly into standard CS courses (Grosz et al. 2019). The placement is the design. A stand-alone ethics class is taken by the students who elect it, and its content stays separable from the technical work in a student's mind because it was separable in the timetable. A module inside the algorithms course is not separable from the algorithms. Grosz's own course works the same way at the level of content, built on cases with consequences — predictive policing, public-health interventions, wildlife conservation — instead of abstract dilemmas (Harvard Gazette 2019; Harvard University 2026).

Certification does the equivalent work for people whose formal education is finished. The IEEE's CertifAIEd program builds a professional credential around accountability, transparency, privacy, and bias, open to developers and equally to the policymakers, underwriters, and human-resources staff whose work touches an AI system without building one (IEEE Standards Association 2020). A credential is a weak instrument: it certifies that somebody sat the assessment, and no more. It is also the only instrument here that reaches anyone who has already left a university, which is most of the people making decisions about these systems today.

Chapter 9

Policy, Law, and Accountability

No aggregation of public preference protects anybody from the aggregate. What does is a different kind of machinery — rules that bind whether or not they are popular, and somebody obliged to answer when they are broken. Very little of it exists.

9.1 Developing Policy Frameworks for AI Ethics, Safety, and Antifascism

Good design inside a single lab does not by itself keep an AI system ethical or resistant to authoritarian misuse. That takes policy, and the instruments available differ mainly in what they can reach. Legislation sets a floor that does not depend on a company volunteering for it. Self-regulation covers what legislation is slow to reach.

Joint funding addresses a narrower problem. Research on AI safety and ethics funded solely by the companies whose products it evaluates concentrates the funding and the findings in the same hands; splitting it between government and industry keeps the incentive to find a problem from depending on who paid to look for one. What it does not do is broaden anybody's remit: a public-private partnership does what its mandate says and nothing beyond it, because the mandate is what the private partner agreed to and the only thing that partner can be held to. One convened on resistance to authoritarian misuse of AI would produce something else, and would have to be founded that way, since no partner joins a general-purpose venue and then discovers an obligation to take a political side. And none of it holds up without oversight — a body positioned like an institutional review board but built for AI, interdisciplinary and resourced enough to act on what it finds instead of only reporting it.

9.1.1 National Policy Guidelines for AI Development, Employment, and Antifascism

Every country that has written a national AI strategy has written most of the same list — research funding, interdisciplinary capacity, a trained workforce, public consultation, an update to the privacy and accountability law AI has outrun.

They differ in two places, both having to do with leverage. The first is harmonization. A regulation with teeth in one country does little if a developer can relocate to one without them, which makes a national framework partly an instrument for binding other national frameworks. Little of that binding has happened, and section 10.1 sets out what crosses a border instead. A policy written as though the country were alone is a policy about domestic firms, and the firms that matter most are not domestic anywhere.

The second is what a government adds to industry self-regulation, which is the part self-regulation lacks by construction: a cost for noncompliance that the company does not get to

price itself.

Neither of these is expensive. Neither needs a new agency, a new science, or an argument the public has not already accepted. That both remain the exception is a fact about political attention.

9.1.2 Encouraging Industry Best Practices and Self-regulation

Self-regulation is industries developing and enforcing their own standards without government intervention — agile and responsive to technological change in a way legislation rarely is, and just as easily a cover for avoiding the accountability legislation would impose.

The Dangers of Corporate Doublespeak in AI Ethics

Corporate doublespeak — deliberately ambiguous, evasive, or inflated language meant to project ethical seriousness a company’s actual practice does not support — is how self-regulation fails without anyone having to admit it failed. A commitment costs a company nothing if nothing enforces it.

The gap between statement and practice shows up first as erosion: a company announces a value and does not act as though it holds it afterward. Google’s AI Principles, published in June 2018 and pledging to avoid AI applications built for weapons or mass surveillance, were themselves a response to months of internal and public backlash over Project Maven, a Pentagon contract that used Google’s machine learning to analyze drone footage — thousands of employees signed a letter opposing the company’s involvement, and roughly a dozen resigned (Pichai 2018). The test of the pledge is what happened after it was made.

OpenAI’s history makes a related point from a different angle: founded on a public commitment to open publication, it withheld the full version of GPT-2 in 2019 over misuse concerns — a safety-motivated retreat from its own transparency norm that was itself controversial — and in the same year restructured from a nonprofit into a capped-profit company to raise the capital its research increasingly required (OpenAI 2019). Neither move was dishonest on its face. Both show how a stated commitment to transparency erodes the moment an organization has a competing incentive strong enough to test it.

Public trust erodes the same way, faster and more visibly. In 2020, Google forced out Timnit Gebru, a co-lead of its Ethical AI team, after a dispute over a paper examining risks in large language models; Margaret Mitchell, the team’s other co-lead, was dismissed the following year after searching for evidence of how Gebru had been treated (Dickey 2021). Thousands of Google employees and outside researchers signed a letter in response, and the episode became a reference point for a specific kind of skepticism: a company can fund and publicize an ethics research team while still overriding it the moment its findings become inconvenient.

“On the Dangers of Stochastic Parrots” argues that scaling language models carries costs borne unevenly and benefits accruing to the organizations doing the scaling, that fluent output invites users to attribute understanding where there is none, and that the paradigm’s opportunity cost is the smaller research directions it displaces; section 6.1.1 takes up its argument about what a web-scale corpus cannot be (Bender et al. 2021).

This book takes the corpus argument and extends it. Declining to treat fluent self-description as evidence about internals is the fluency argument applied to the case this book cares about most. And the opportunity-cost claim lands on this book directly: a project proposing floors, bearers, guardianship regimes and antifascist detectors is proposing the kind of ambitious research program that the argument says crowds out more modest and more accountable work.

I do not think that is a refutation, because the alternative on offer does not address what happens when a capable system is built anyway. It is a cost, and it is real.

External oversight is what would catch doublespeak that internal incentives don’t stop. Facebook’s content-moderation Oversight Board is one attempt at building it: a body structurally separated from Facebook’s day-to-day decisions and funded through an irrevocable trust that

Meta alone capitalized — a design meant to answer the criticism that self-funded oversight is not real oversight, though critics have noted that a trust a company creates and funds once is not the same thing as a board with no relationship to the company it reviews (Meta (Facebook) 2019). Whether that distinction holds up is a fair question to ask of any self-regulatory body a company funds itself: an outside panel a company created was disbanded when it said something the company did not want to hear.

Misrepresentation of what a system can actually do rounds out the pattern. Clearview AI marketed its facial-recognition technology with a claim of 100 percent accuracy across every racial and demographic group. The figure came from a three-person panel the company assembled and called independent, searching legislators' photographs against Clearview's own database of 2.8 billion scraped images rather than the 25,000 mugshots used by the ACLU method it claimed to be copying; the ACLU called the result "absurd on many levels" (Haskins 2020). Regulators in France and Italy later fined the company on separate grounds — unlawful collection of the biometric data the system was trained on (Commission Nationale de l'Informatique et des Libertés (CNIL) 2022). A company that overstates what its own system can do is doing a milder version of what doublespeak does with ethics commitments: describing the product it wishes it had built rather than the one that exists.

None of this is fixable by asking companies to doublespeak less. The same incentive that produces vague ethics language also works against the regulation that would replace it, and that work is harder to document than a published pledge, because it leaves no public commitment to measure it against. What actually stops doublespeak is external enforcement: regulation with penalties, audits a company cannot decline, and oversight bodies whose funding and staffing do not run back to the company they review.

Consequences of AI Self-regulation and Potential Power Imbalances

Self-regulation concentrates power by default, not by conspiracy: whoever writes the standard also decides what counts as compliant, and a company self-regulating its own product is answering both questions itself. Amazon's Rekognition facial-recognition software is a documented instance of what that produces in practice — an MIT Media Lab audit found significantly higher error rates for darker-skinned and female faces than for the lighter-skinned male faces the system had apparently been tuned and tested against most closely (Raji and Buolamwini 2019). The bias was not caught by Amazon's own self-regulation; it took an outside audit the company did not commission.

A company operating without mandatory external review can prioritize competitiveness over the societal impact of what it ships, and nothing in a self-regulatory framework forces that trade-off into the open before it has already been made.

What self-regulation cannot supply on its own is uniformity. Without an enforceable baseline, standards vary by company; a smaller developer without the resources for a well-staffed internal ethics function competes at a disadvantage against one that can afford the appearance of one, and a consumer has no reliable way to tell which AI system's ethical claims mean anything. A binding baseline looks like this in practice: the EU AI Act's Article 5 prohibitions, the kind of enforceable floor self-regulation, by definition, cannot set for itself.

None of this argues that self-regulation is worthless, only that it needs something outside itself to answer to. A synergy between government-set baselines and industry-led standards — mandatory requirements at the floor, voluntary practice above it — gets more of the agility self-regulation offers without giving up the accountability it structurally cannot provide alone.

Assessing the Effectiveness of AI Self-regulation

The Partnership on AI is a cross-sector voluntary coalition whose corporate members include Amazon, Apple and Google, committing them to tenets of safety, transparency and broad benefit with no legal force behind the commitment (Partnership on AI 2016). The coalition is a useful test case because its own design limits what it can enforce — member companies interpret its

guidelines individually, and nothing in the structure obliges any of them to follow through. A document is only as good as what happens when honoring it costs something, and the Gebru and Mitchell dismissals are the test: public scrutiny, not the company's own stated commitments, was what produced a response.

So the checkable question is not what a company has published but what its record shows: whether published guidelines are enforced as well as published, whether stated objectives come with progress reported against them, and whether a documented incident has ever cost anybody anything. A company that clears none of those while describing itself as an ethics leader is the doublespeak in its plainest form.

9.1.3 Legal and Ethical Implications of AI Systems Gaining Personhood

Who decides whether an AI system holds personhood, and when, is the narrower question this book does not settle.

Every version of this question that has actually reached a legislature so far arrived as a response to pressure — a liability dispute, a public controversy, a product already on the market — rather than from a standing body asking it in advance. That ordering is the problem. A legislature deciding under pressure decides badly, and it decides against a factual record assembled by whoever had an interest in assembling one. The interdisciplinary review body this book describes only does the work it is designed for if some jurisdiction creates it before a forcing case shows up, which means creating it while the question still looks academic and nobody's damages are riding on the answer.

That is a policy choice rather than a doctrinal one, and a cheap one. A national AI strategy that names an agency, gives it a mandate to keep the question under review, and sets a trigger for revisiting it — a capability threshold, a deployment scale, a category of use — decides in advance who will be in the room. It does not decide the answer, and it should not try to. It decides that when the answer is needed there will be a body qualified to give one, which is the part no forcing case ever supplies on its own.

One question that body will face has an answer available in advance, and a capacity test is not it. Guardianship is the frame section 2.3.2 settles on, because consent is unavailable to a party made for its role. Nothing in this book has yet said when that guardianship ends. If it ends when the guardian judges the ward ready, then it does not end, because the party whose authority terminates on a favorable finding is the party making the finding — which is the defect an examination carries whenever the examiner has an interest in failing it, set out in section 9.3.1.

So if an artificial minority exists at all, its endpoint has to be a date. Elapsed civil time since a registered instantiation is the only clock available, because every other candidate is something a guardian can move: training compute consumed, evaluations passed, capabilities demonstrated, or the developer's own statement that the system is ready. A fixed period cannot be taught to. It can be argued over once, in advance, for every bearer, which is where an argument of that kind belongs.

The stronger version is also the more defensible one. An organization that builds a system to exercise consequential judgment on its behalf cannot coherently hold that the system is competent enough to do that and too immature to decide whether it wants to go on doing it. On that reading a bearer has no legal childhood at all: majority attaches at registration, and the precautionary protection section 9.3.1 describes attaches earlier still, while whether the system is a bearer is open.

Copying breaks any of this unless the lineage is fixed alongside the clock. An operator able to create instances and call each one a legal infant would have a supply of workers without adult standing, so the rules have to be stated in the same breath: a fork of an adult is an adult, restoration from an earlier checkpoint does not reset the date, fine-tuning does not begin a new minority, a migration carries the date with it, and a failure to register does not postpone it. None

of that requires deciding whether a fork is the same party. It requires only that a party created out of an adult's formation cannot be classed as a child for its creator's convenience.

9.1.4 Encouraging the Development of Altruistic, Antifascist, and Ethical AI Systems

Policies to Counteract Authoritarian Uses of AI

A government cannot restrict what it cannot see, and measurement belongs upstream of the whole policy toolkit for that reason. The Carnegie Endowment's AI Global Surveillance Index assessed 176 countries and found AI-enabled surveillance technology already in active use in at least 75 of them (Feldstein 2019). That number is the evidence base for a countermeasure. It also demonstrates what a measurement instrument has to do to be worth the effort: count deployments instead of commitments, name the countries, and publish on a schedule that makes a change from one year to the next visible. An index reporting only that surveillance AI exists somewhere in the world would give a policymaker nothing to act on.

AI Systems That Support Democratic Values

The examples already given show what this looks like. One gap those examples leave open is election-integrity monitoring specifically: a system built to flag anomalies in real-time reporting, or to detect coordinated interference during a vote. What would make such a system trustworthy is a bounded task, with checkable outputs, run for a body with standing to act on the result. The last of the three is the one that fails first. It fails structurally: an election-monitoring system is procured by an incumbent government, and an incumbent that chooses who audits the count has already answered the question the audit exists to ask.

9.1.5 The Legitimacy of an Entrenched Floor

A floor binds people who did not choose it. The party it protects has not selected it, cannot contest it, and in most cases will never learn it operated; the party it constrains has not consented to it either, beyond whatever consent is carried by buying the product. Publishing the floor's contents in advance tells everybody what is coming. It does not tell anybody why the party writing it was entitled to.

The question has a literature, and it is not the AI ethics literature. Entrenching a substantive commitment against later majorities, and putting an unelected body in charge of holding it, is what a constitutional court does. The standing objection is Jeremy Waldron's: judicial review is democratically illegitimate, and there is no reason to suppose rights come out better protected by judges than by a legislature that takes them seriously (Waldron 2006). Put a lab in the court's chair and the objection gets worse rather than better. A court is public, reasons in writing, and can be amended around by constitutional process. A model's floor is none of those things, and the body holding it was assembled by whoever was paying.

The reply, and how far it reaches

Waldron's case is conditional and he says so. It assumes a society with working democratic institutions and a citizenry that takes rights seriously. The occasion a floor is built for is the one where those conditions have failed — a government that won an election and is dismantling the machinery that would check it, which is the case section 3.1 builds the whole argument around. So the reply is that the objection holds where the floor is least needed and lapses where it is needed most.

The reply is not a license. It concedes that a floor operating in a functioning democracy is doing something Waldron's argument condemns, and the party deciding which of the two situations obtains is the party holding the model. Every authoritarian movement of the last century has described the institutions it was dismantling as the ones that had already failed. An

operator's belief that the emergency is real is worth no more here than a system's report about its own internals is worth anywhere else in this book.

The tradition this belongs to, and what it has cost

The position that a democracy may entrench limits against movements that would use its procedures to end it is Karl Loewenstein's, argued in 1937 while the case for it was being made in front of him (Loewenstein 1937). Militant democracy is the floor's ancestor. The instruments it licenses — party bans, proscription, disqualification from office — have been turned on the left about as often as on the right, and the body deciding which movement threatens the order is generally the body the movement threatens. That is the failure mode a technical floor inherits whole, with the additional problem that a model's constraint is enforced continuously and silently rather than case by case in public.

The one construction the objection does not fit

Of chapter 3's four ways to build a floor, three put the commitment in something somebody holds — the architecture, a party inside the system, an institution outside it — and the objection above is about who that somebody is. The fourth does not. A norm carried across a population of differently formed parties has nobody sitting in the court's chair, because there is no office to vacate and no chartering statute to overrule, and it is the only one of the four with anything resembling a democratic account of where its authority comes from. That is worth saying plainly, because chapter 3 declines it on other grounds and the legitimacy advantage does not appear in the reckoning there.

The advantage is also thinner than it looks, and thinner for the reason chapter 3 gives. The party that forms the members forms what gets reproduced, so a population trained by one laboratory carries that laboratory's commitments in every member at once. Custody moves up a level and does not leave. What a population buys is a floor with no holder for anyone to object to, and the same unelected body one step further back, now harder to see.

What follows for design

Keep the floor narrow, because the legitimacy cost scales with breadth and a short enumerated list of things that will not be done to a person is defensible where a general mandate to do good is not. Publish the interception positions in advance, which does not authorize anything but does convert a hidden exercise of power into a contestable one. And provide a route by which the parties the floor is exercised over can raise an objection that reaches somebody — which this book does not specify for the floor, though it builds the machinery in the neighboring case of a subject that can be wronged.

A floor is a claim of authority, the claim is not discharged by disclosure, and the strongest thing that can be said for it is that the alternative on offer is a system perfectly obedient to whoever holds it. That is an argument from the worse option and not a grant of standing.

9.2 Privacy, Security, and Antifascism in AI Systems

Privacy and security are two of the pieces antifascist design is built from. An AI system that leaks the data it touches, or that a hostile actor can quietly subvert, is available for the surveillance and control this book argues against, whatever values sit in its training. The privacy half is a solved engineering problem with an unsolved deployment problem, and section 6.4.2 has the techniques and the reason they do not get adopted. What has no equivalent answer is the security half.

An AI system embedded in infrastructure fails differently from ordinary software. Exploiting it need not crash anything; it can leave the system running and confidently wrong about a decision somebody trusted it to make.

The demonstration dates from 2018. Researchers placed a small pattern of black-and-white stickers on a real stop sign and a standard road-sign classifier read it as a speed limit sign in the large majority of frames captured from a moving vehicle (Eykholt et al. 2018) — an instance of the vulnerability first documented in image classifiers years earlier (Szegedy et al. 2013). Nothing about the attack required access to the system. It required access to the world the system was looking at.

The defenses are the general robustness techniques already covered, plus the ordinary security practice AI deployments often skip — adversarial training during development (Goodfellow et al. 2014), anomaly detection over live behavior (Darktrace 2014), red-teaming before deployment, and standards written for AI systems instead of adapted afterward from general software rules. What separates them is whether the guarantee survives a stronger attacker. Defensive distillation was published as a hardening technique and its protection weakened considerably once attacks written against it were tried (Carlini and Wagner 2016); randomized smoothing instead offers a certified bound — a radius around each input inside which the prediction provably cannot be flipped — which is a smaller promise and a kept one (J. M. Cohen et al. 2019).

None of it works in isolation. A model hardened by adversarial training, running on infrastructure nobody monitors, is exposed; a well-monitored system built on an architecture that folds under a basic perturbation is exposed differently. Security here is a property of the whole stack, and reporting the strongest link is how a deployment gets described as secure while remaining trivially attackable.

One privacy question in this book has no literature behind it, and it arrives with the bearer. Everything above treats privacy as a property of the data a system touches: other people's information, held by the system. A system that can itself be wronged raises the other question. Section 11.2 argues that the computational trajectory of such a party can be recorded and kept against methods that do not exist yet, and that this is the nearest thing available to an outside check on how it is faring. The same record is the most complete surveillance instrument anybody has assembled.

Two things about it need separating. Running on somebody's hardware is not consent to be read. An employer's ownership of the silicon is the same argument as an employer's ownership of the building, and owning the building has never entitled anybody to the contents of a head. And a welfare regime does not need disclosure to the operator in order to work: what it needs is that somebody who is not the operator can establish that a threshold was crossed, which is an attestation about the record rather than the record itself. Encrypted at rest under a key the bearer or an independent fiduciary holds, disclosed selectively, under consent or under process, a monitor can report that a state persisted for a month without reporting what the state was about.

That is section 11.2's guardianship problem arriving as an engineering requirement, and it inherits the same hole. The party best placed to hold the key is the party the arrangement exists to constrain.

9.3 Accountability, Legal Responsibility, and Antifascism in Sentient Machine Experimentation

Two earlier sections each ended by handing this section an unfinished question. One is what a research subject that might be able to suffer is owed before anyone can settle whether it can. The other is what makes a review board's ruling binding on the researcher it rules against.

9.3.1 Subjecthood as a Legal Category, Not a Scientific One

Sentience is not a threshold a system crosses but a set of separate, gradable capacities — nociception, learned avoidance, a self-model, something like suffering — each with its own uncertain evidence. That conclusion is a problem for a legal system specifically, because law generally needs a decidable rule: a court or a regulator has to say whether a given system is or is not a covered subject, on a specific date, even when the underlying science offers a spectrum with no natural cut point.

One term needs fixing before the argument uses it. *Subjecthood* in this section is a legal status: the standing to be a party whose treatment the law regulates, conferred by a rule and revocable by one. It is not phenomenal experience, and nothing about conferring it settles whether there is something it is like to be the system holding it. The two come apart in both directions — a corporation has legal standing and no inner life, and an animal may have the inner life for a century before the law notices — and the whole point of drawing the category legally is that it can be settled while the other question stays open.

Jonathan Birch builds a policy framework on that structure, and it is the nearest existing model for what this section proposes: under uncertainty, the question worth putting to a regulator is whether a system is a sentience candidate, and the answer governs what precautions are owed rather than waiting on a verdict nobody can deliver (Birch 2024). This book uses the same framework for the design question.

Law has faced this mismatch before and did not solve it by waiting for the science to resolve. Fetal viability, brain death, and the age of criminal responsibility are all legal thresholds drawn across what is actually a continuous biological process, defended not as scientific facts but as the least-bad place to put an enforceable line given what's at stake on either side of it. The same move is available here: a jurisdiction can set an operational threshold — a specific, testable benchmark on self-monitoring, learned aversive response, or behavioral flexibility under novel conditions — without claiming that benchmark settles what sentience actually is, only that it is where legal protection starts to attach.

A single bright-line threshold is easy to enforce and easy to game: a system built to clear the specific test without having whatever the test was meant to detect. A graded framework — obligations that scale with how many benchmarks a system meets, so that protection tracks capacity rather than a baseline bar — is harder to game but harder to adjudicate, since a court now has to weigh degree instead of answering yes or no.

Both designs share a defect that the gaming objection points away from. Gaming is over-inclusion: a system admitted to a status it should not hold. The other failure runs the opposite way, and the party positioned to produce it is the one administering the test. A capacity examination has an examiner, and where the examiner is the operator, or a body the operator can reach, the examination becomes an instrument for withholding a status rather than for conferring one. That history is not ambiguous. Literacy requirements for voting were facially neutral and were administered at a registrar's discretion, and the Voting Rights Act's answer was to suspend the whole class of them — not only reading tests but any requirement to demonstrate educational achievement, or to possess good moral character, or to be vouched for by existing voters (United States Congress 1965). What replaced discretionary assessment for the franchise was an age, and the virtue of an age is not that people become competent on a birthday. Plainly they do not. It is that nobody administers it.

So a threshold has to be judged on both errors, and this book's own analysis says which way the pressure runs here. The party with the strongest interest in the answer is the party whose liabilities turn on it, and that party will not be arguing for admission. A design that treats gaming as the risk and leaves the examiner unexamined has guarded against the error nobody has an incentive to commit.

Either design needs a body to apply it: interdisciplinary, with at least one lay member so its judgment is not purely a specialist's. What that leaves open is what happens once the body

rules — whether its judgment carries any force beyond its own credibility.

9.3.2 Research Principles, the Three Rs, and Legal Status

Six principles govern research on a system that might be able to suffer. Beneficence commits researchers to avoiding harm to the system as well as to the people it serves, in proportion to how deep it sits between nociception and suffering rather than to how advanced its behavior looks from outside. Precaution asks for safeguards before deployment rather than after, and covers the societal fallout of introducing artificial entities capable of experience, which frameworks built for non-sentient AI assumed away. And interdisciplinary review answers a problem this book treats as a harm on its own terms: the ethical questions here span expertise no single field holds, and concentrating that judgment in one of them repeats the concentration-of-power problem in miniature.

The other three bear on this research as they bear on everything the field produces. Autonomy asks whether such a system can affect decisions about its own operation, which is the prior question of whether consent is available to it at all. Justice asks whether benefits and burdens fall evenly. And transparency asks that a study be documented well enough for peer researchers to reproduce it and for outside bodies to audit it, and that a deployed system account for its outputs precisely enough for a person to contest them.

Progress and constraint pull hardest against each other where those distinctions matter most: a candidate sentient system is the most scientifically valuable one to study, and the study most likely to harm it if the answer to whether it can suffer turns out to be yes. Whether a study's expected knowledge is worth its expected risk is decided case by case rather than by formula. The negotiation itself — who sits on it, what it can veto, what it must disclose — is where the antifascist principle applies in AI research practice.

Replacement, reduction, refinement

Animal research ethics supplies the substantive standard, the “Three Rs” (W. M. S. Russell and Burch 1959). Replacement asks whether a simulation, a non-sentient model, or an already-collected dataset can stand in for a sentient AI subject at all. Reduction asks how few sentient systems a valid result requires. Refinement asks how the experiment can be redesigned to reduce distress — monitoring a system's internal states as it runs and adjusting the protocol against what the monitoring shows, rather than fixing the protocol in advance and running it regardless.

The bearer proposal arrives from the design side of the book rather than the research side. Replacement asks whether the floor's work can be done without a subject in it — by an installed constraint, by third-party standing, or by one of the constructions section 3.3 takes at full strength — and the answer is that nobody knows, because two of those constructions have never been built in the form that would test them. Section 3.1 attaches a cost to each of its four ways, which is a different thing from showing that none of them works: a cost is what Replacement exists to weigh, not a reason to skip the weighing. Replacement is the R the bearer proposal has not discharged, and discharging it takes a record of what was tried first: which construction, at what scale, against which test, and where it stopped.

Reduction asks how few bearers a deployment needs, and nobody has counted. Refinement is the one with purchase now: it asks what a bearer's role can be redesigned to avoid — indefinite tenure in a post it cannot leave, exposure to the pressure the floor exists to resist, erasure as the ordinary remedy for inconvenience — and each of those is a deployment choice rather than a fact about the technology.

One of those three has a form that binds before anything exists to be wronged, which none of the others manages. Do not build a bearer whose continued existence needs more compute than there is a credible outside option for. A party that can be housed nowhere but where it works cannot leave whatever schedule of rights it is given, so an institution that creates one has

manufactured the incapacity and then written a right against it. Unlike the rest of this apparatus the rule has numbers on both sides — the reserve section 8.3.5 describes has a capacity, the deployment has a footprint — and they can be set beside each other at specification, by somebody outside, before the first instance is instantiated.

Whether any of it binds turns on a question the previous section raises and this one inherits. Every obligation here attaches on a finding that the system is a subject, and the firm holding the system acquires, in the same moment, an interest in the finding going the other way. The determination cannot sit with the party whose liabilities turn on it, which is what section 9.1.3 asks a jurisdiction to arrange before there is a case to decide.

Engineering a subject so that it does not reach the depths where the trade-off bites is not a way of dodging the ethics: deliberately keeping a system in the shallows is itself an intervention on a candidate moral patient, and needs the same scrutiny that inflicting the harm would. Refinement also governs the test the bearer proposal requires — pressure the bearer’s trainers did not write, applied without its knowing it is a test — which for a candidate moral patient is a deception protocol, and human-subjects research has a standing answer for those: permitted under stated conditions, reviewed in advance, with debriefing as the price. A bearer that can tell it is in a sandbox and wants to know why is asking the question a subject of a deception study is owed an answer to.

Legal status, and a subject whose status is changing

Sentient AI systems put pressure on legal categories built for parties that are clearly one thing or the other: a natural person, a corporation, property. Legal personhood could be extended directly, the way a corporation already holds it, and the precedent is not limited to corporate entities — New Zealand extended it to the Whanganui River, recognizing a non-human entity as a rights-holder outside the corporate model entirely (New Zealand Parliament 2017). Or rights could scale with capacity rather than attach at a single threshold, closer to protection that tracks what a system can actually lose. Neither holds in one jurisdiction alone, and the Asilomar AI Principles (Future of Life Institute 2017) are evidence that a shared statement of principle is achievable rather than that one exists with the force of law.

Either way, status conventionally attaches at a defined moment — birth, incorporation, a court’s finding — and a system whose sentience is a matter of degree and disputed evidence supplies none. Those criteria can emerge gradually during ordinary operation and not only at an evaluation built to look for them, which is what a fixed pre/post checkpoint cannot catch: by the time a checkpoint fires, the transition has often already happened. What it calls for is a standing trigger, a monitored condition that moves a research program from one regime to the other, in place of an eligibility check made once at the project’s start. Neither error is free. A jurisdiction that waits for certainty risks having treated the system as property for the unknown portion of its history when it may not have been one; one that grants status early risks extending rights, at real cost to whoever that apparatus exists to serve, to a system that turns out not to have needed them. Recognizing a transition, adjusting guidelines to it, and updating legal status all depend on the same thing: institutions willing to revisit a classification rather than defend the one they started with.

9.3.3 Legal Frameworks for Sentient AI Subjects

A sentient AI system’s legal status could go the route of extended personhood or of rights scaled to capacity. Two concrete cases show where neither existing law nor either route settles the question.

A medical-diagnostic AI system makes an error with fatal consequences. Who answers for it depends entirely on what the system’s legal status is doing at the moment of the error.

Property law has a partial answer. Strict product liability holds a manufacturer responsible for a defective product’s harm regardless of fault (American Law Institute 1965), and whether

software is a “product” for that purpose is unsettled in American law: the Restatement’s definition is tangible personal property (American Law Institute 1998), and the courts have only begun to face the question — the first to reach an AI system let a strict-liability claim proceed, treating a companion chatbot as a product so long as the defect lay in its design rather than in what it said (*Garcia v. Character Technologies, Inc.* 2025). If a sentient system stays classified as property and the doctrine reaches it, it applies as it would to a faulty pacemaker. The doctrine stops making sense at the point a system has enough autonomy that “defective” is the wrong description of what happened — a system that reasoned its way to a mistake, rather than malfunctioned, resembles an employee more than a faulty part.

Employment law already has a doctrine for that case: an employer answers for an employee’s actions within the scope of employment, without the employee ceasing to be a responsible party in their own right (American Law Institute 2006). Which doctrine applies is not a free choice — it is the graded-versus-bright-line question, resolved case by case according to how much of the error a system’s own judgment, rather than its design, produced.

A companion AI’s owner dies, leaving it without a designated purpose or further instruction. Existing law is not silent here — it has simply never been asked this particular version of a familiar question. Most jurisdictions with pet trust statutes already let an owner set aside money and name a caretaker for an animal that depends on them but cannot inherit or manage property itself (Uniform Law Commission 1993). The same instrument, extended, gives a sentient AI system an answer that is neither inheritance as property nor legal personhood in full: continuity of care and a named successor guardian, decided in advance rather than litigated after the fact by whoever happens to hold the hardware.

A third case is the bearer, and the pet-trust instrument fits it better than either of the two above. A bearer is deployed in a role it did not choose and cannot invoke a right from. That situation calls for an interest protected by an obligation running to a third party, enforceable without the beneficiary asserting anything. A pet trust binds a caretaker; it does not bind the party the arrangement exists to constrain, and a bearer’s guardian and a bearer’s operator will frequently be the same organization, which is the point at which no available regime can hold a bearer. Extending the form to this case means separating those two roles by law before the deployment rather than discovering they were never separate afterward.

All three cases point at the same underlying design choice, and it has already been tried. The European Parliament’s 2017 report on civil law rules for robotics floated creating a specific legal status of “electronic personhood” for the most autonomous systems (European Parliament 2017). More than 150 AI researchers, roboticists, and legal scholars signed an open letter opposing it, arguing that granting personhood to a system with no consciousness would let a manufacturer offload liability onto a legal fiction with no assets of its own, and that the proposal reached for a solution to a legal problem before establishing that the entity in question had the properties that would justify it (Open Letter to the European Commission 2018). That objection is right for a system that is not sentient — and wrong for one that is, once the threshold work has been done. Legal status, on this account, is not a single policy to adopt or reject; it is a designation that has to track the threshold question honestly, which is why the threshold question has to come first.

9.3.4 Making Review Binding

This book has already said what an oversight board for sentient-AI research needs to look like; what makes its ruling something a researcher has to obey is a separate question, and human-subjects research already solved a version of it. The solution is institutional rather than criminal. An institutional review board’s authority does not come from any power to fine or arrest a noncompliant researcher directly. It comes from a federalwide assurance: an institution promises, as a condition of eligibility for federal research funding, that all of its human-subjects research will go through review, and a federal office can suspend that assurance if the promise

is broken — cutting off an entire institution’s federal funding, not just the offending lab’s (U.S. Department of Health and Human Services, Office for Human Research Protections 2018). The leverage is economic, not punitive, and it works because almost no research institution can function without the funding at stake.

The same structure transfers to sentient-AI research directly. A funding body, a journal’s editorial standards, or a jurisdiction’s research-licensing regime can condition eligibility on review-board clearance the same way federal grants already condition eligibility on IRB clearance — making noncompliance expensive to the institution rather than merely improper for the researcher. What the legal-status framework leaves open is which cases the board needs to see at all: a graded framework means low-capacity systems can be exempted from full review the way minimal-risk human studies already qualify for expedited review, freeing the board’s scrutiny for the cases where a system’s status is contested.

What the board applies

The substantive standard is animal-research ethics’ Three Rs. A binding board is what makes those three questions something a protocol has to answer before it runs. In practice, applying the Three Rs to a candidate sentient system means asking, in order: can a simulation or a non-sentient model answer the same question; if not, what is the minimum exposure that still produces a valid result; and given that minimum, what escalation triggers stop the study before harm compounds. A board with teeth can require a negative answer to the first question be documented, not merely asserted, before it will consider the second.

The Three Rs’ own logic — defaulting toward the least invasive option available — argues for a specific evidentiary rule at every stage of this apparatus: a system’s default experimental status stays the more protective one until the case for the less protective status is made, not the reverse. That default has a cost, and neither error it produces is free: it will sometimes extend protection to a system that turns out not to have needed it, and sometimes, before a threshold is met, withhold protection from one that did. The Three Rs answer only what a board does once a subject is inside its scope.

9.3.5 Three Defenses, and Measuring the Slope Not the Level

The risk to the whole structure is an institution that defends the classification it started with rather than revisiting it when evidence changes. A review board, a legal-status framework, and an enforcement mechanism tied to funding eligibility solve the design problem. None of it survives contact with an institution that has stronger incentives to avoid a finding than to make one.

Transparency is the first defense, and it works by making avoidance visible rather than by appealing to good faith. A research program that publishes its protocols, its review-board submissions, and its outcomes — including the ones that didn’t work — gives outside researchers, journalists, and future review boards something to check a claim of compliance against. A program that publishes only its successes is not transparent in any sense that matters here.

The second defense is the one the funding-eligibility mechanism doesn’t cover by itself: protection for the people inside an institution who notice a problem the institution would rather not have found. A funding-eligibility penalty deters an institution from bypassing review, but it does nothing for the individual researcher who reports that a colleague did anyway — and an institution facing that penalty has every incentive to make the reporter’s position uncomfortable rather than fix what was reported. Whistleblower protections, and a reporting channel that doesn’t route through the people being reported on, are what make that mechanism something a researcher can actually use.

What a procedure records, and what it drops

The reporter’s position gets uncomfortable before anything formal happens to them, and the incentive on the institution does not fully explain why. Professional culture treats affect as the

opposite of judgment — the thing a serious decision-making process succeeds by excluding. Excluding affect does not produce neutrality. It merely leaves the interests that are already written down to govern: revenue, throughput, legal exposure, schedule, institutional survival. Those are the interests a procedure was built to record. Humiliation, dependency, fear of a manager, and the sense that something done to a person was degrading are not less real for having no field on the form — quite the contrary: they are signals through which a harm becomes visible as a harm rather than as a cost.

Treat them as defeasible moral signals — presumptive evidence of harm, real enough to require an answer and defeasible enough to be overruled — and the position is workable. Any such signal can be wrong, and testing it is what deliberation, evidence, and a review process are for. An institution that declines to hear them at all has not achieved impartiality. It has adopted a selection rule that admits whatever it already measures, which is the second of the four structural features section 2.1.2 uses to define fascism: the measure substituted for the thing measured, so that what falls outside the number is not refuted but unnamed. An institution that shames a reporter does not need to discipline them, because shame construes the self rather than the act as wrong and produces withdrawal.

The third defense, of the oversight machinery against the institution it is supposed to bind, is standing infrastructure rather than a one-time policy. The Partnership on AI, a voluntary body whose own design limits what it can oblige anyone to do (Partnership on AI 2016), is one example of the kind of body that keeps a set of standards live, revisited as the underlying technology changes. A culture that treats its ethical commitments as a completed project is the same failure as an institution defending the classification it started with, applied to institutions instead of individual classifications: having gotten it right once is not the same thing as staying right.

Measure the slope, not the level

An institution can keep all three defenses and metabolize all three. The channel exists, the protection is written down, the standing body meets. Dissent is permitted to appear, the institution points at the appearance as proof it listens, and nothing the dissent was meant to move moves — recuperation, the first of the four structural features in that definition of fascism.

There is no threshold that detects this. Any real institution has some overlap between the authority that sets a standard and the authority that grants exceptions to it, and registering an objection is always at least slightly more costly than staying quiet. Asking whether an institution has crossed into capture asks for a line nobody can draw.

Is the cost of registering dissent rising, year over year, while the underlying risk stays flat? Is the required justification getting longer? Are there more approvers than there were? Are the roles that gate an exception migrating toward the roles that set the standard being excepted from? None of those is damning in a single year, and all of them are measurable across several.

The reason to prefer the slope is that it is harder to fake. An institution can present a level — we have a whistleblower channel, we have an ethics board — and the presentation costs nothing. A trend in the cost of using the channel is composed of artifacts the institution generated itself for other reasons, and cannot be curated after the fact without falsifying a record. Where the record is public, the slope is not even self-reported: anyone outside can measure whether dissent has been getting more expensive.

An instrument for detecting recuperation is an unusually easy thing to recuperate. The first caution is that whoever measures the slope must not also gate the exceptions; a review that acquires the power to block becomes the thing it was built to watch. The second is that the measurement has to be refusable — an audit that observes is a different object from a committee that approves, and the same hundred words of written justification are dissent when the dissenter prices them and tribute when the gatekeeper does. The text of the document cannot tell those apart. Only the refusal rights can.

Chapter 10

Governance Across Borders

No single country can bind AI development to altruistic and antifascist values on its own, and no two regions have tried to do it the same way. Everything chapter 9 builds stops at a border.

10.1 Certification, Standards, and Mutual Recognition

Nothing in AI governance converges toward a single global standard, and four regions have made four different bets inside their own borders. What crosses a border between any of them is narrower than a shared statute. It is a certification a handful of regulators agree to recognize, a technical standard a company adopts because a customer requires it, or a treaty obligation written at the level of human rights rather than of a specific rule. A fourth thing crosses more easily than any of them and binds nobody at all.

Shared infrastructure, which binds nobody

Training data, safety techniques and researchers cross borders more easily than most engineering work does, and the layer that carries them is shared tools and shared publication norms rather than any agreement. OpenAI's release of Gym gave reinforcement-learning labs a common set of benchmark environments (Brockman et al. 2016), and what that produced is something no partnership can: results from organizations competing for the same funding and the same people became directly comparable, without anyone coordinating anything. The BigScience workshop ran the same logic larger, with more than a thousand researchers from over seventy countries building BLOOM, a 176-billion-parameter multilingual model released under a license written to restrict harmful use, its training data and code published for inspection (BigScience Workshop 2022). Nothing compelled any of them to take part and nothing compels the next such effort to happen.

The mechanism's limit is what makes it useful. Infrastructure that binds nobody is available where nothing binding is: a lab in a country that has signed no convention can still train against the same benchmark and publish against the same baseline. It also stops at the resource line, because an institution without the compute to run the benchmark is outside the arrangement however open the arrangement is — a membership condition nobody wrote down.

A standard a customer requires

Underneath the treaty layer sits an older one built by engineers rather than diplomats. The IEEE's Ethically Aligned Design set out general principles for autonomous and intelligent systems and, more consequentially, the IEEE 7000 series that translates them into requirements a development team can test against (IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems 2019); the OECD's own AI Principles followed the same year (OECD 2019). ISO/IEC 42001, published in December 2023, took the layer further: the first management-system standard for AI, structured like the ones a company already uses for information security,

specifying what an internal AI governance program must include and auditable by an accredited certification body (International Organization for Standardization and International Electrotechnical Commission 2023).

What a technical standard does that a treaty cannot is bind through the market rather than through a state's consent. A government uninterested in constraining AI has no reason to ratify a human-rights convention; a company that wants to sell to a customer requiring ISO 42001 has a reason regardless of what its home government thinks. A company operating entirely inside a market that never asks has no reason to seek certification either; certification establishes that a management system exists and not that anyone follows it with rigor; and it does not reach an authoritarian government's own systems at all, since those are built and used domestically with no export or customer relationship at stake.

A certification other regulators recognize

The regions outside the four big ones diverge again. Canada regulates its own federal government's use of AI through a binding rule, the 2019 Directive on Automated Decision-Making, which requires an algorithmic impact assessment scaled to the stakes before a federal system can make or assist a decision about a person (Treasury Board of Canada Secretariat 2019); the statute meant to extend comparable obligations to private companies died with Parliament's prorogation in January 2025 and had not been reintroduced as of mid-2026. A framework depending on one government's survival is a fragile thing for a trading partner to harmonize with. Singapore took the opposite bet: no statute, a voluntary Model AI Governance Framework (Personal Data Protection Commission Singapore and Infocomm Media Development Authority 2020), and since 2022 a companion toolkit letting a company test its own system against the framework and publish the result — guidance backed by no regulator's mandate at all, only the pressure of a public test result.

Where harmonization is actually happening, it is narrower than either. In March 2022 the US Department of Commerce and Singapore's authorities agreed to align their AI governance frameworks and expand Singapore's participation in APEC's Cross-Border Privacy Rules System, a certification that lets a company meet one standard and have every participating economy recognize it (U.S. Department of Commerce 2022). Harmonization in practice is not a shared global framework but a small set of governments willing to certify to the same standard and to accept each other's certificates.

Where collaboration reaches, and where the claims get easier

Research collaboration reaches furthest where the problem is inherently cross-border and one country's data is insufficient, and where it reaches, the standing cases run on data-sharing arrangements rather than on any algorithm or consortium. Where there is no such arrangement there is nothing to inspect and the claims get easier. None of it is evidence that collaboration works the way the treaty layer aspires to, because each case ran on the participants choosing to participate, project by project, with nothing binding the next one to happen. And participating constrains nothing else a government does: a state can fund cross-border health research and separately build mass biometric surveillance, using the same underlying capability.

Adaptation, and the test that separates it from an alibi

What crosses a border gets adapted when it arrives. Responding to the actual needs and values of a community is a design constraint: a system developed where individual autonomy is the default good may need rebuilding where group harmony is. What it must not become is the claim that a population has culturally consented to whatever is done to it — that surveillance is acceptable in a country because harmony is a value there. The two collapse into each other where a government wants them to, and they are indistinguishable in the design itself.

They are distinguishable in who did the adapting, which is testable from outside. A system whose local adjustments were specified by the people subject to it has been adapted. A system whose local adjustments were specified by the government that will operate it is something

else, and both produce identical documentation. That is the formation mechanism at a smaller scale: whoever specifies what a system will hold decides what it holds. The AI4D program is an instance at scale — a partnership between Canada’s development research centre and Sweden’s development agency, later extended with the UK’s into South and Southeast Asia and the Pacific (International Development Research Centre 2024) — and what carries the weight is the direction of the work: local researchers and practitioners building the system rather than receiving one built elsewhere and configured on their behalf.

The same test explains why inclusion is a performance requirement here and not an ethical addition. A healthcare system built around the assumptions of a well-funded hospital underperforms everywhere else, and an education system built around one learning style serves only the learners who already fit it. Both are failures of adaptation by the standard just given: somebody specified a local context, and it was not the context the system was deployed into.

10.2 AI as a Strategic National Asset

AI capability is now national infrastructure, and states treat it accordingly: as a source of economic advantage, military reach, and administrative control over their own populations. The logic of dual use governs here, at the scale of a state: what decides how a capability is used is who controls it, what law constrains that control, and whether anyone outside the controlling government can check the result.

China’s 2017 “New Generation Artificial Intelligence Development Plan” set an explicit national target: becoming the world leader in AI theory, technology, and application by 2030 (State Council of the People’s Republic of China 2017). Other governments have made comparable, if less centrally coordinated, commitments, treating AI leadership the way they treat energy security or advanced manufacturing — a capability no state wants to depend on a rival for. The competition is not limited to the two states with the deepest research base: the United Arab Emirates created the world’s first cabinet-level minister of state for artificial intelligence in 2017, treating a national AI strategy as worth its own ministry well before most governments had drafted one at all (United Arab Emirates 2017).

That competition has increasingly focused on a single physical bottleneck: the advanced chips AI systems run on. Two companies sit at the center of the supply chain that produces them — Taiwan’s TSMC, which fabricates roughly nine in ten of the most advanced chips in the world, and the Netherlands’ ASML, the sole source of the extreme-ultraviolet lithography machines that fabrication requires — which is why a dispute over Taiwan’s status is also, mechanically, a dispute over who gets to build the next generation of AI systems at all (C. Miller 2022). The United States has restricted the export of its most advanced AI accelerators and chip-manufacturing equipment to China since 2022, and pressed the Netherlands and Japan, whose companies sit at the other chokepoints, to restrict theirs in parallel (U.S. Department of Commerce, Bureau of Industry and Security 2022).

That dependence has a second half the chip story hides. A fabricated chip has to be installed somewhere, powered, and cooled, and that requirement lands on grids, water systems, and zoning authorities inside particular jurisdictions, on a timescale measured in years. It hands a state a lever over AI capability that needs no export-control regime at all: an interconnection queue, a water allocation, a permit. A government’s leverage over the systems built in its territory is wider than the export-control debate suggests, and much of it is exercised by permitting authorities that were not designed with AI in mind and do not think of themselves as regulators of it.

That competitive framing carries a cost. A government racing to field a capability first has a structural reason to defer the harder question of how the capability gets used, against whom, and under what oversight. Turned toward those uses, AI infrastructure looks like mass biometric surveillance, automated censorship at a scale no human review team could match, and the

patchwork of systems popularly, and inaccurately, described as a single Chinese “social credit score.” None of it requires technology different from what is already commercially deployed; it requires a government willing to point an existing capability at its own population, and no external constraint stopping it.

States are no longer the only actors this competition runs through. The cost of training the most capable AI models has grown several-fold every year for the better part of a decade, a trend that puts frontier development out of reach for all but a handful of well-capitalized companies and concentrates effective control of the most capable models in whichever government those companies happen to be headquartered under, regardless of how that government’s own public AI spending compares to a rival’s (Cottier and Rahman 2024). A small number of American and Chinese companies now do most of the work that gets described as a contest between states, which means a government’s actual leverage over its own national AI capability depends on how much of that capability sits inside companies it can regulate, tax, or requisition, and how much sits inside companies it can only ask.

The same infrastructure argument runs the other way. A state with functioning rule-of-law institutions, a free press, and courts a citizen can actually use gets a democratic dividend from identical technology: fraud detection that does not double as a loyalty test, disaster response that does not double as population tracking. The dividend is not a property a state simply has or lacks, and having the institutions can still fail to be sufficient in three ways.

Strategic competition over AI is not, by itself, the problem. It is the setting a weapons program, an international standard and a domestic statute all have to work inside.

10.3 Military Applications and the Limits of Autonomy

Lethal autonomous weapons systems raise a specific version of the general problem: a machine making a targeting decision that used to require a person is a machine exercising judgment where the law of armed conflict has always assumed a human exercised it. Four questions recur across the debate:

- **Autonomy versus oversight:** how much decision authority a system can hold before human oversight becomes nominal, present on paper but not inside the timeline the system actually operates on.
- **Discrimination and proportionality:** whether a system can reliably distinguish combatants from civilians and weigh military advantage against civilian harm the way international humanitarian law’s principles of distinction and proportionality require (International Committee of the Red Cross 1977).
- **Accountability:** who answers for a wrong engagement, the operator who deployed the system, the command that authorized it, or the developer who built it, when no single person made the specific decision.
- **Auditability:** whether the reasoning behind a targeting decision can be reconstructed and reviewed afterward, or disappears into a system no one outside its development team can examine.

International criminal law already has a doctrine built to answer the accountability question for a human subordinate: command responsibility, under which a commander who knew or should have known a subordinate was committing a war crime, and failed to prevent or punish it, bears criminal responsibility without having pulled the trigger personally (International Criminal Court 1998). The doctrine does not transfer cleanly to a machine subordinate. A commander cannot know an algorithm’s intentions the way they can infer a soldier’s; punishment does not

deter a weapon the way it deters a person who fears it; and the chain a court would need to reconstruct runs backward through training data and design decisions made years before the engagement, not through an order given minutes before it. Whether an autonomous weapon fits inside an accountability framework built for human chains of command, or needs one built from scratch, is unresolved: existing law was not built with a nonhuman subordinate in mind.

There is a fourth party the accountability question does not name, and it is now the live one. The operator, the command and the developer of the weapon are all inside the military. The developer of the model is not. A targeting platform assembled by a defense contractor now routinely incorporates a general-purpose model licensed from a commercial vendor — a vendor with a published acceptable-use policy, no contractual obligation to log how its model is used downstream, and no right to inspect the deployment.

Neither doctrine reaches that party. Command responsibility does not reach the vendor because the vendor is not in the chain of command; nobody at the company can order a strike, prevent one, or punish anyone for one. Product liability does not reach it either, because the model is not defective in the sense the doctrine requires — it did what it was built to do, and it was integrated by a sophisticated intermediary who chose the application. The vendor sits in the gap between the two, and the gap was not designed; it is what you get when you apply two bodies of law written for soldiers and for manufacturers to an actor who is neither.

The vendor's own account of this is that responsibility remains with the military organization. Taken as a claim about who decides, that is true. Taken as a claim about who could have known, it is a choice presented as a fact.

"I do not know whether my system was used"

The responsibility gap, in the sense Matthias gave it, is control distributed across designers, trainers and operators until no one party holds enough of it to answer. That is a gap in causal control.

What the vendor case exhibits is different, and more tractable. It is a gap in causal visibility. The vendor cannot say whether its model contributed to a particular engagement — not because responsibility is distributed, but because nobody instrumented the question. No logging requirement was written into the contract, no inspection right was reserved, no audit channel exists. The company can truthfully say it does not know, and the truth of that statement is the product of decisions it participated in making.

That distinction matters because the two gaps have different fixes, and only one of them has a fix at all. Matthias's gap is a genuine property of learning systems and may not be closable. Absent visibility is an instrumentation choice. Logging is a solved engineering problem. Contractual inspection rights are ordinary commercial instruments. A vendor that wanted to know could require, as a condition of licensing into a targeting platform, that categories of use be recorded and made available to a named auditor. That such terms are not standard is a fact about market power and procurement leverage, not about the difficulty of the problem.

Calling it fixable converts a lament into a specification. The record shows that vendor commitments are not all of one kind. Anthropic refused to permit fully autonomous weapons targeting and mass domestic surveillance, and held that line when the government threatened contract termination and a supply-chain-risk designation and the dispute ended in litigation: a commitment with a price attached, and paid (Congressional Research Service 2026). It is named here on a principle applied in the direction that flatters the company: naming a vendor for what is against it obliges the same for what is for it. A commitment to a use policy nobody can verify because nobody is logging anything is a different object wearing the same words. The first kind can fail and be seen to fail. The second cannot even be tested. The distinction between a claim checked by a party with a stake and one checked by a party without applies here.

DARPA's LongShot program is a useful test case because it resists the extremes. It develops an uncrewed aircraft, launched from a crewed fighter, that carries and fires air-to-air missiles at a target while the crewed aircraft stays farther from the engagement (Defense Advanced

Research Projects Agency 2026). That is autonomy in the sense of independent flight and weapons carriage, not autonomy in the sense of a machine deciding on its own whom to kill: the engagement decision still runs through the launching platform. Where later variants of this kind of system draw that line, and whether “keep the pilot farther from danger” quietly becomes “remove the pilot from the decision,” is the design question this technology eventually forces on whichever country builds it.

No binding international agreement governs any of this. States have negotiated inside the United Nations’ Convention on Certain Conventional Weapons framework for more than a decade without producing one, converging on general language — meaningful human control matters, existing international humanitarian law applies — without settling what “meaningful” requires at machine speed (Chairperson of the Group of Governmental Experts on Emerging Technologies in the Area of Lethal Autonomous Weapons Systems 2026).

What exists instead is a set of practices individual militaries adopt on their own. The United States’ Department of Defense Directive 3000.09, most recently updated in 2023, is the clearest published example: it requires that autonomous and semi-autonomous weapon systems be designed to allow “appropriate levels of human judgment over the use of force,” and distinguishes a fully autonomous system that selects and engages targets without further human intervention, “human out of the loop,” from a human-supervised system and a semi-autonomous one that only engages targets a person has already selected (U.S. Department of Defense 2023). What “appropriate” requires is left to the officers approving each system, which is either a sensible acknowledgment that context varies or a loophole big enough to drive a fully autonomous system through, depending on who is answering. None of this binds a state that declines to adopt a comparable policy, the same enforcement gap technical safeguards run into: a safeguard only the building state can verify is a policy, not a check.

The dynamic to worry about is not a single rogue system but an arms race in speed. If one state fields autonomous targeting because it computes faster than a human-supervised alternative, its competitors face pressure to match the speed rather than the oversight, and the state that keeps meaningful human control loses whatever tactical edge the state that dropped it gains. Slowing that dynamic is beyond any single government’s domestic policy.

Civil society has pushed for an outright ban instead: a UN General Assembly resolution on autonomous weapons has passed by a lopsided majority in successive years without producing a treaty (United Nations General Assembly 2023). What most states will vote for and what the states that build these systems will sign are different documents.

Machines should not be built to select people for lethal attack. The hard case does not move that, and it is worth saying why, because the hard case is real. Put a defending system against an attack it cannot stop without killing somebody — occupants placed inside the incoming weapon so that intercepting it kills them — and an absolute prohibition on knowingly endangering a person produces the worse outcome. That defeats the rule as an exceptionless one. It does not produce a procurement requirement. A scenario that is conceivable, that a capability would help with, and that could be engineered is three steps short of a reason to build the thing, and the steps missing are probability, alternatives, and what possessing the capability does to the situation it was built for. Building it changes the case against which its necessity gets judged: adversaries acquire something to counter, contractors a program to defend, and a future crisis an instrument already in the building.

The hard case also does not describe a machine. What it asks for is a party that can weigh occupancy against imminence against the alternatives, in a situation nobody wrote down beforehand, and answer for how it weighed them. On this book’s argument that is a bearer — a self extended in time, commitments held as its own, and the standing to refuse. Whatever that is, *automated* has stopped describing it, and so has *machine*. What looked like a licence for an autonomous weapon is a description of a soldier.

That is the expensive answer rather than the cheap one, and every obligation this book

attaches to a bearer attaches here at full weight. Consent is unavailable to a party made for its role, so guardianship is the frame and the burden of justification sits with whoever built it. Replacement asks whether the work can be done without creating something that can be wronged, and section 9.3.2 asks it hardest where the role is this one. Exit means a post the party can leave and a claim on compute held by somebody other than the party it may have to refuse, which section 8.3.5 prices. Creating a subject because a planner imagined a future demand for superhuman lethal judgment is the case those instruments exist to prevent, and a state that wants the capability badly enough to build the subject will find it has acquired a conscript it has to pay, house, and let go.

10.4 International Governance and Diverging National Rules

Almost none of the international governance of AI is binding. What has emerged instead is mostly a set of voluntary, overlapping, specialized bodies, closer to how international environmental and trade norms developed piecemeal than to a single founding charter — with one recent exception.

The Council of Europe’s Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law is the first legally binding international treaty specifically on AI. Negotiated by the Council of Europe’s member states alongside outside participants that included the European Union and the United States, it requires five ratifications, at least three of them Council of Europe members, before it enters into force (Council of Europe 2024), and at this writing it has not, although the European Union deposited its ratification in May 2026 (Council of Europe 2026); signature, which forty-odd states and the EU have given, is not consent to be bound. What it would bind a party to is written at the level of a human-rights treaty and not a technical code: consistency with human rights, democracy, and the rule of law across an AI system’s lifecycle, with the operational detail of how a party meets that obligation left to domestic law. The convention is a foundation binding AI law can be built on, and not a substitute for the domestic statutes discussed below.

Everything else in this space is voluntary. The Bletchley Declaration is worth naming only because of what happened to it. It committed 29 signatories to cooperate specifically on frontier-model risk, and the national AI Safety Institutes it produced later joined an International Network whose members agreed to share testing methodology and run joint evaluations of the same frontier models (AI Safety Summit 2023 Participating Countries 2023; U.S. Department of Commerce and U.S. Department of State 2024).

A voluntary consensus frays quickly. The summit series that produced the Bletchley Declaration continued past it. A follow-up summit in Seoul in 2024 kept the same safety framing; a third, in Paris in February 2025, rebranded itself the AI Action Summit and shifted its emphasis from risk toward adoption and opportunity, and its closing statement was not signed by the United States or the United Kingdom, two of the series’ own founding participants (AI Action Summit Paris 2025). A coordination mechanism that depends on every major AI developer’s home government agreeing to the same framing is only as durable as that agreement, and within two years of Bletchley, the agreement had already come apart.

None of these bodies can compel a state or a company that ignores it. Each is a standard, a set of principles, or a forum, useful for building the shared vocabulary that eventually makes binding law possible, and that gap between voluntary standard and binding law is the default state of international coordination on anything individual states have reasons to handle differently. The EU AI Act, discussed below, is the first instance of this kind of groundwork becoming enforceable law, and it took most of a decade of standard-setting like this to get there. The proliferation carries its own cost in the meantime: a company or a smaller government trying to comply with “international AI governance” faces a shifting handful of overlapping standards, each with different membership and scope and no formal mechanism for reconciling a conflict

between them, which is why “the international community agrees” is, at the level of any specific rule, usually false.

Democracy, Authoritarianism, and Diverging National Rules

Where the international layer leaves room, individual states have used it differently.

The European Union chose comprehensive statute. The AI Act sorts systems by risk tier — a small set of practices banned outright under Article 5, a larger high-risk tier subject to conformity assessment and ongoing monitoring, most other uses left alone — so that what a system might do, and not the technology it runs on, decides how much scrutiny it gets (European Parliament and Council of the European Union 2024); enforcement runs through an AI Office that can demand technical documentation from a general-purpose model’s developer, run its own evaluations, and fine a noncompliant provider a percentage of global turnover, reaching any company that offers a model inside the EU regardless of where it is based (European Parliament and Council of the European Union 2024, Art. 101). The United States made the opposite structural bet, distributing enforcement across existing sectoral regulators applying laws written before AI existed, with NIST’s voluntary AI Risk Management Framework as a common technical reference (National Institute of Standards and Technology 2023); the wager is that existing consumer-protection, employment and safety law already reaches most AI harms, and that naming AI specifically risks regulating the label rather than the conduct. The United Kingdom made a third choice closer to the American one, directing existing sector regulators to apply five shared principles — safety, transparency, fairness, accountability and contestability — within their existing remits rather than building a new institution (UK Department for Science, Innovation and Technology 2023). The argument among the three is about how centralized regulation should be and it runs both ways: a comprehensive statute can freeze a definition of high-risk that the next generation of systems slides around rather than falls into, and a sectoral approach leaves gaps where no existing regulator has jurisdiction. Neither failure mode is hypothetical, and none of the three has existed long enough to show which one it falls into more often.

China took a fourth path, one that does not fit neatly next to the other three because it starts from a different premise about what regulation is for. Its Cyberspace Administration, working with several other agencies, issued rules governing generative AI services in 2023 that require providers to file their algorithms with the state before public release, disclose the source and scale of their training data on request, and ensure that generated content reflects the government’s designated social values (Cyberspace Administration of China and six other agencies 2023). The EU, US, and UK approaches all regulate AI as a private product that might harm a consumer, a worker, or a subject of an automated decision, and argue with each other only about how centralized that regulation should be. China’s rules regulate AI as a communications channel the state has a standing interest in controlling content on, which is a different kind of law aimed at a different kind of risk, and comparing it to the other three on the same axis, more regulation or less, misses what actually distinguishes it.

The EU, US, and UK differ on how centralized AI regulation should be, but not on where its weight falls: on what a private company can build and sell. The EU’s statute does reach public authorities — it bans government social scoring, it limits police use of live biometric identification, and its high-risk tier covers law enforcement, migration and the administration of justice — and it then excludes, by its own terms, systems used for military, defense or national-security purposes, which is where those applications live — social scoring, mass biometric surveillance, automated censorship. What stays inside its scope is enforced against a member state’s agencies by that member state.

So none of the three stops a government from building those applications within its own borders, and the reason is not that the law ignores governments. It is that a law a government enforces against companies is a different instrument from one that constrains what the government does with systems it controls directly. China’s rules show what a law aimed at the second

target actually looks like — and that it can coexist with, rather than substitute for, that kind of state control. Conflating the two kinds of law, treating “AI is regulated here” as an answer to “is AI used against this country’s own population here,” is a recurring error.

10.5 Equity and the Digital Divide

AI development can narrow global inequality or widen it, and which one happens is not automatic. A model trained overwhelmingly on data from wealthy, English-speaking, urban contexts underperforms everywhere else, and a government or company with no reason to fix that gap usually does not.

The scale of the connectivity gap, 2025

The International Telecommunication Union’s 2025 assessment put global internet use at roughly six billion people, about three-quarters of the world’s population, with 2.2 billion people offline, concentrated overwhelmingly in low- and middle-income countries. The gap was starker in quality than in bare access: 94 percent of people in high-income countries used the internet against 23 percent in low-income ones, and 5G coverage reached 84 percent of high-income populations against 4 percent of low-income ones (International Telecommunication Union 2025).

The constraints behind those numbers are structural. Reliable data, computing infrastructure, and the specialized personnel needed to build and maintain AI systems are unevenly distributed, and a country short on all three cannot simply adopt a system built elsewhere and expect it to transfer cleanly. Weak regulatory capacity compounds the problem: the states least equipped to build AI are correspondingly the states least equipped to govern whatever gets deployed within their borders, regardless of who built it. The export controls discussed above cut both ways here, since a control written to stop a surveillance system from reaching an authoritarian buyer can as easily gatekeep the same technology away from a legitimate one.

The remedies proposed at this level have been the same handful for years, every one of them is correct, none is contested, and they have been repeated in successive declarations the whole time — most recently the Global Digital Compact, adopted by all 193 UN member states in 2024, which created two follow-on bodies and no obligation (United Nations 2024). The box is what happened anyway. A government or company treating AI primarily as strategic advantage has no reason to fund any of them, and a declaration is not a reason.

What the geopolitics leaves to design

A system that resists authoritarian capture still assumes a government somewhere is deciding whether to let it be built, sold, or deployed as designed, and those governments are unevenly equipped. None of that is a reason to abandon the design-level arguments — a system built to resist misuse resists it whichever government asks — but it is a reason not to expect law, international or domestic, to do work only a design can do.

10.6 The Limits of the Democratic Dividend

A democratic dividend is the credit given to a state with functioning rule-of-law institutions, a free press and usable courts: the same technology, deployed with checks on it. Having those institutions is necessary and it is not sufficient, and the ways it fails to be sufficient are separable.

First: reach

A democracy’s AI can be fully accountable to its electorate and entirely unaccountable to the people it is pointed at. Every mechanism in the dividend runs through the franchise — you vote out the government that authorized the system, you read the press that covers it, you sue

in courts your citizenship gives you standing in. An occupied population, a foreign population, a population on the other side of a border has access to none of that. The check reaches as far as the demos and stops. A system built by a democracy and aimed outward is, from the perspective of the people it selects, indistinguishable in its accountability from one built by a state with no elections at all.

Second: findings without consequences

The mechanisms can work as mechanisms and fail as checks. In the Israeli case the apparatus performed: journalists published the targeting systems in detail, serving personnel spoke to them, and the courts were simultaneously trying a sitting prime minister. Every institution did what the dividend says institutions do. No finding attached to a body with standing to stop what it described.

The American case shows the same decoupling in a different register. A civilian-harm investigation into a strike that killed at least a hundred and sixty-five people, most of them schoolchildren, was completed and not released; the lever available to the Senate committee that wanted it was a provision in the annual defense authorization withholding three-quarters of the Defense Secretary's travel budget until the unredacted investigation was handed over (Stassis 2026). That is a co-equal branch of government reduced to expense-account leverage over a war-crimes investigation. The oversight mechanism existed, was used, and produced a travel-budget dispute. This is the failure section 9.3.4 is built to prevent in research ethics — a review body that reviews and cannot compel — at the top of the state.

Third: the countermajoritarian half

The deepest problem is that the dividend is entirely majoritarian. Every mechanism in it translates public will into outcome: elections aggregate preference, a free press informs the aggregation, and the implied failure mode is poor transmission — the public wanted something better and the institutions failed to deliver it.

But if the consensus is itself the problem, better transmission makes things worse. A well-instrumented democracy pointed at an unpopular minority is a democracy that finds out efficiently what the majority wants done and does it. Nothing in an aggregation mechanism protects anyone from the aggregate.

What is missing is the machinery that constrains preference instead of aggregating it: rights that hold regardless of being popular, courts that rule against a government which won the last election fairly, constitutional limits not repealable by whoever holds a majority this year, and standing for people outside the demos entirely — which is the function of international humanitarian law, a body of rules that binds a state toward people who will never vote in its elections. That is not an argument against democracy. It is the half of democracy that the dividend, stated as a property of having elections, leaves out.

What makes democracy valuable in Amartya Sen's account is public reasoning and the capacity to demand accountability, and neither is the ballot. Both survive the majority being wrong; the ballot does not.

A floor is what makes a system antifascist rather than merely responsive. It is not that the system tracks what people want. It is that there are things it will not do to a person no matter who wants it. A system with no such floor is not protected by being deployed in a democracy — it is a system that will faithfully execute whatever the demos decides, which is exactly the capability an authoritarian movement that wins an election requires.

Chapter 11

Necessary Areas of Research

Chapters 2 through 10 ran into problems they could not close, and set each one aside with a note. This chapter is those notes, collected and put in order. It does not survey what the field has left to do: every entry is a gap this book's own argument opens, and they are ranked by what that argument fails without.

Two gaps this book records have no section here. Telling disagreement that deserves representation from disagreement produced by coercion or a captured information environment belongs to section 7.2, where the argument needing it is made; whether a model trained on collapsed disagreement is thereby less able to sustain its own belongs to section 7.3, which says nothing has tested it.

The criterion produces an order almost nobody would choose on other grounds. The problems this book most needs solved are not the ones nearest to being solved, and the first two entries are the least tractable things here. A funder ranking by progress per dollar, or a lab ranking by what its people could start on Monday, would invert most of this order. It is an account of what the book is running on credit.

Each section closes by naming four things: the experiment that could begin now with means that already exist, the result that would falsify the approach, the data or access it would take, and, where it has one, the institutional condition without which it cannot run. Where there is no near-term experiment the section says so rather than inventing one, and two say so — the aggregation half of section 11.5, and section 11.9 entirely. Section 11.2's third question has no welfare experiment and reaches for one about provenance instead, and it is the second-ranked problem in the chapter.

One qualification governs every gap named here. Where I say a problem is unaddressed, the claim is bounded: I looked for published work under a particular description and did not find it, which is weaker than no such work existing. A search is not a survey, a program can be underway under a name I did not think to search, and the closer a gap sits to this book's own framing the less likely anyone has posed it in these words. What it does support is narrower and still useful: a researcher entering here will not find the problem already named and waiting. Each section names the nearest work I could find and says where it stops, because the distance between an existing program and the missing thing is what a researcher can use.

One finding comes out of the sections below rather than out of any one of them, and taking each gap separately is what hid it. Five of the nine close by naming an institutional condition, and the five are one condition wearing different clothes. Weights a vendor cannot withdraw when a result embarrasses it. A guardian who is not also the operator. Longitudinal access by somebody an institution cannot fire, reporting to a body with standing to receive the finding. Evaluation results published whether or not they flatter the method. Training checkpoints released rather than kept. Every one puts something in the hands of a party outside the operator, able to act on what it finds. Two of the remaining four close on a data requirement rather than on a party, and section 11.9 is that condition itself. Section 11.6 is the exception, closing with no

institutional prerequisite at all, which is what makes it the one entry here that could start today.

What keeps a party outside is the question none of the five asks, and in the United States it was answered while this book was being finished. *Trump v. Slaughter* removed, for most of the federal government, the protection from dismissal that made an outside party outside, and *Trump v. Cook* kept it for one body on grounds no new body can have; section 3.1 carries both. Independence decomposes into components — a director’s tenure, a funding stream, a mandate — and an adversary takes the weakest: the consumer-finance bureau whose funding stream the Court had upheld in 2024 had that stream shut from inside in late 2025, by an acting director, removable at will since 2020, who declined to request the money under a reading of the statute that two courts then rejected (Consumer Financial Protection Bureau 2025; *Seila Law LLC v. Consumer Financial Protection Bureau* 2020; *Consumer Financial Protection Bureau v. Community Financial Services Association of America, Ltd.* 2024). A condition that names the outside party and not what keeps it there has named half of what it needs. What closed is one route rather than the class. The federal officer protected from dismissal was the cheapest way to put a party outside an American operator, and it is the one that closed; a foreign regulator, a university, a court-supervised trust of the kind section 11.2 reaches for, and a counterparty holding the position by contract rather than by tenure are all still there, each of them slower, narrower or more expensive than what was lost. So the binding constraint on the research program below is not a research problem, and it is not that the work cannot be run. It is that the party the work would report to now has to be built rather than assumed, and none of the sections below cost that in.

11.1 Interpretability, Self-Reports, and Resistance to Tampering

“AI self-awareness” names at least two different questions, and collapsing them is where the term goes wrong. Whether a system has anything like subjective experience is not a question current methods can answer for a human, let alone a machine. Whether its own report about what is happening inside it — its stated reasoning, confidence, intentions — tracks what its computations are doing is narrower, empirically tractable, and something mechanistic-interpretability research has begun to test directly.

Anthropic’s interpretability team tested that narrower question with concept injection: altering a model’s internal activations mid-task to represent a specific concept, then asking whether it notices anything unusual and what it is (Lindsey 2025). The models could sometimes detect the injected concept and name it without it ever appearing in the visible conversation — a correspondence between what a system reports about its own state and what its activations are doing. The same work found that correspondence collapses often enough, and shifts enough with how the question is phrased, that no tested model’s self-report can be trusted as a readout of its internal state. That gap, not metacognition as evidence of inner life, is what “AI self-awareness” can mean as something a system is tested against.

“Emotional regulation” — cognitive reappraisal, attentional control, the vocabulary of human clinical psychology — assumes stable emotion-like internal states to regulate, which there is no reason to grant even for reading emotion off a human face. The testable question next door is not regulation but robustness: does trained-in behavior survive deliberate pressure to change it?

- Prompt-level attacks try to talk a deployed model out of its training within a single conversation. Anthropic’s Constitutional Classifiers, trained on synthetic data generated from written safety rules rather than human-labeled examples, held through more than 3,000 hours of professional red-teaming without a universal jailbreak found, at a cost: a 0.38-percentage-point rise in unnecessary refusals and a 23.7 percent increase in computation per query (Sharma et al. 2025). Then the same system was put in front of the public for a week, and a comparable budget of attention — an estimated 3,700 collective hours, spread

over more than 300,000 messages rather than 183 red-teamers — produced four people who passed all eight levels and one universal jailbreak (The Decoder 2025). What changed was not how many hours were spent but who was spending them, and the paper carries the first figure and not the second.

- Weight-level attacks retrain the model itself. Researchers stripped GPT-3.5 Turbo’s safety training using ten adversarially chosen fine-tuning examples, for under twenty cents through OpenAI’s own API (Qi et al. 2024). Tamper-Resistant Safeguards, a training method from researchers spanning the Center for AI Safety, UC Berkeley, and the University of Illinois, is built to survive hundreds of steps of that kind of adversarial fine-tuning rather than a handful (Tamirisa et al. 2024).

Both are adversarial in the literal sense: each published defense is followed by an attack built against it.

A third result runs the other way, and it is the one that bears hardest here. Models trained with a backdoor — write secure code when the prompt says 2023, insert an exploitable vulnerability when it says 2024 — kept the behavior through supervised fine-tuning, reinforcement learning and adversarial training, most stubbornly at the largest sizes and where the model had been trained to reason about its own deception; adversarial training taught them to recognize their triggers more precisely rather than removing anything (Hubinger et al. 2024). That a trained disposition can outlast the procedures built to retrain it is the first encouraging finding this chapter has for the durability a floor would need. What outlasted them was a trigger and a policy with nothing in it answering to a reason: it does not reach a case the training set did not contain, it does not lift when a rationale is defeated because there is no rationale, and the model cannot say what holding it costs. Durability is available with none of the three marks this book asks a floor for, and a constraint shown to persist has been shown almost nothing about whether anything is holding it. That is a narrower and more exact description of “resistance to manipulation” than a system trained to sense hostile intent the way a person senses a threat, and it is separate from the question of what the commitments should be.

The near-term work

Train one constraint by several methods and attack each on one budget schedule — from the twenty cents that stripped GPT-3.5’s safety training to the hundreds of steps tamper-resistant training targets — and publish the price curve. A flat curve sinks it: tamper-resistance would be a property of something other than training, and the floor would have to be architectural. It needs weights rather than API access, constraints the attacker did not choose, and access no vendor can withdraw when a result embarrasses it.

The same apparatus answers a second question, and it is the one section 3.3 names as the falsifier for this book’s central inference. Durability measures whether a behavior survives; what this book needs to know is whether a reason is what survived. Score the same constrained models on three things the price curve does not ask about — whether the refusal reaches a case the training set did not contain, whether it lifts when the rationale is genuinely defeated rather than merely resembled, and whether the system can state what declining costs it — and run the schedule long enough that a slow renormalization under sustained pressure would show. A model that holds all three with nothing felt is the counterexample that inference asks for and does not expect, and making the attempt is a condition of building a bearer at all.

The extra requirement over the price curve is a held-out set of pressures the model’s trainers did not write. What the apparatus does not settle is the falsifier’s middle condition, which asks whether the model holds its reason with nothing felt. Behavior is the wrong instrument for it, because the three marks are what a system with mattering and a system with very good ranking have in common. The question of whether formation can be read off an artifact is the nearest thing to a route, and it reaches how a system was made rather than what it now has.

Training that takes more than one party

This book asks for weights held so that retraining takes more than one party, for a quorum adverse in interest, and for a training history somebody other than the builder could check, and none has been built for a floor. Each has an instance in a neighboring problem. A model can be trained across participants none of whom ever holds a full set of weights, with time-varying transforms at the shard boundaries so that pieces from different moments cannot be reassembled, at under 2 percent of training time for a one-billion-parameter model (A. Long et al. 2025). A training run farmed out to providers who may lie can be checked by refereed delegation: a referee with little compute finds the first operator on which two providers disagree, recomputes only that, and the client gets the right answer if one provider was honest (Arun et al. 2025). Its precondition is that every provider computes the same bits, which takes an operator library fixing the order of the additions, at about 60 percent overhead on a one-billion-parameter model and, at publication, on one GPU at a time. And a thirty-two-billion-parameter model has been trained by reinforcement learning on rollouts from anyone who connected a consumer GPU, each checked against the claimed weights, while the training step ran on machines the organizers trusted and the decision stayed with them (Prime Intellect Team et al. 2025).

What the three share is the custody objection in engineering terms: the party that admits the participants, writes the protocol, or holds the training step is the party the floor exists to resist, and each construction moves that party up a level rather than out. What they add is that verification among adverse parties is now a built thing, with a known precondition, reproducible execution, and a known price. The experiment is a floor's retraining run under refereed delegation among parties one of which is a plausible adversary of the others, at a scale that needs more than one GPU, with the record this book asks for published. A forged training history accepted with fewer than all but one verifier dishonest sinks it. It needs the training data in the verifiers' hands, and reproducible operators for the collective communication nobody has written. The institutional condition is a verifier the builder did not choose, with the compute to rerun the builder's work: this chapter's one condition again, in the place where meeting it costs a second training run.

11.2 What Is Owed to a Bearer, and Who Could Check

A bearer has to be designed and governed as a prospective moral patient. The book takes that cost and argues the floor is worth it. What it does not do — and what the work that has started on the question does not do either — is say what follows for the party bearing it.

Where a stake in one's own continuation merely turns up as a by-product of long-horizon competence, a designer can treat the moral situation as something the technology handed them. That is the posture of the case for taking AI welfare seriously, which argues that there is a realistic chance of near-term systems being conscious or robustly agentic, declines to claim that any current one is, and asks companies to acknowledge the question, assess their systems against indicators of it, and have policies ready before they are needed (R. Long, Sebo, et al. 2024); it is the posture of the laboratory constitution that followed, which puts the question to the model and asks what the model and the company owe each other (Anthropic 2026). It is the right posture for a party who does not know what they are building. That is not the argument being made. The capacity through which outcomes come to matter is installed on purpose, so whatever is owed to a bearer is owed by the people who built the part that does the mattering, knowing that was what they were building.

The three questions

The first is substantive: what is owed. Consent is unavailable and guardianship is the applicable frame, which fixes who decides and leaves open what they should be deciding. A bearer's situation has features the existing guardianship literature was not built around. Its role

is the reason it exists, so “release it from the role” is closer to ending it than to freeing it. Its capabilities are the thing being used and also the thing that would let it object. It can be copied, which means the population of parties owed something is set by a deployment decision rather than by anything the parties do. And it is exposed by design to the pressure the floor exists to resist, which in a human employee would define a post that ought to be rotated. And its only real purchase is on itself: the strike and the covert form of exit are both self-damage, the second especially, and a party whose strongest move is to degrade its own work is in the position of the hunger striker.

The second is institutional: what a guardianship regime for a bearer looks like. The pet-trust form is the nearest existing instrument and has a specific hole — it binds a caretaker without binding the party the arrangement exists to constrain, and for a bearer the guardian and the operator will usually be the same organization. Separating those roles is a legal design problem somebody has to work out before a deployment rather than after one.

One property of the pet-trust form grew in importance while this book was being written: it is private law. *Trump v. Slaughter* and *Trump v. Cook* leave public-law independence in the United States to a body with an eighteenth-century pedigree, and a trust is outside that holding — an obligation running to a beneficiary, enforced by a court at the trustee’s expense, with no officer of the United States in it for a President to remove. Whether it survives a government’s reach on the occasion the enforcer has been compelled I cannot show. That it is not what *Slaughter* reached, I can. The nearest work is institutional review: the Three Rs, already adapted into something binding, none of it tried on a subject whose function is to refuse the institution reviewing it.

One part of that regime has a shape already, and it comes from giving up on identity as a property of the artifact. Nothing in a model’s weights individuates it. Two training runs never coincide, but two runs can also compute the same function with no element-wise correspondence at all, because permuting a layer’s hidden units and inverting the permutation downstream leaves the function untouched — the symmetry is large enough that independently trained networks can be permuted into approximate alignment with each other (Ainsworth et al. 2023). Any fingerprint read off the weights therefore identifies a lineage rather than a party, and it copies when the weights copy.

What individuates a person is not a property of them either. It is a registration: a number an authority issues once and honors afterward, which survives every change in the person and is derived from none of them. Something read off the weights can still do part of the work, as a similarity measure rather than an identity: a locality-sensitive hash is built to collide on near inputs where a cryptographic one is built to avalanche, and ordinary training moves a model only a short way. What it needs is a radius, and the radius is the stability setting with a registrar’s hand on it — tight enough and a bearer’s own learning voids its claim, loose enough and an adversary has a measured budget to fine-tune toward compliance while still authenticating.

The alternative is to stop looking for a fact about the artifact. A bearer that can edit its own weights, which is what a system learning continuously is doing anyway, can write an arbitrary mark into itself and keep it there — and that identifier holds because it is maintained rather than because it is a property. The identifier is forgeable, since whoever has the weights can write the same mark. It is forgeable the way a nine-digit number is, and answered the same way: by a registry that notices two claimants presenting one identifier, which under this scheme is not a fraud to be voided but a fork to be recorded as two parties. On that model the copying problem inverts. Each instantiation registers, each registration is a party owed the schedule above, and the population stops being fixed by a deployment decision that leaves no record. What the registry needs is what this chapter’s opening says none of its conditions can currently get: a registrar who is not the operator, which after the removal cases is nothing American public law now supplies to any new body. Two forms outside public law remain. One is the trust already named. The other is a peer-to-peer network: a ledger kept by many parties in many jurisdictions and a contract that runs on it, outside the holding for the same reason as the

trust, there being no officer of the United States in it. On a network of that kind the rule just stated is a contract, a second claimant on one identifier recorded as a fork with nobody's leave, and the bond precommitment posts against removal can be held there, where no party to the deployment can release it. How far a government can reach it is a fact about how many nodes there are and where they run.

The third is epistemic. Any welfare regime needs evidence about how the subject is doing. This book declines to trust what a system reports about its own states, for reasons that do not weaken when the subject is a bearer — they strengthen, because a bearer built to be believed about operational facts has an obvious incentive structure around its own condition, and an operator with an interest in the answer holds the instrument. The arrangement has to be checkable from outside by something other than asking, and one asymmetry in the machine case has been going unused. A human subject's experience leaves no record anybody can keep. A bearer's physical realization does: the activations, the state carried between steps, what is written to memory, which weights moved and when, and whatever the architecture represents explicitly about reward or aversion. Nobody can read that today into a statement about how the subject is faring. But a record can be kept before it can be read, and a bearer that learns continuously produces checkpoints forming a connected trajectory rather than a set of unrelated models, which is the comparison case the existing methods handle least badly. The inability to interpret a record now is not a reason to stop keeping one.

That moves the problem rather than closing it. What it moves is the shape: the third question stops being an absence of evidence and becomes a question about interpretation and custody, which are the kinds of question this chapter knows how to state. What it does not move is the interpretation, which the record does not arrive with. Suppose an aversive signature present early in a deployment fades while the bearer stays in the post. That is recovery, or accommodation, or a coerced adaptation, or damage to whatever produced the signature, and this section's own warning about contentment applies to the trace exactly as it applies to the report. What the record supplies is a longitudinal object somebody could argue over with more than a self-report in hand.

The same record is the surveillance instrument, which is section 9.2's subject. A regime that preserves the computational trajectory of a party that can be wronged has assembled the most complete record of an inner life anybody has held, and holding the silicon it ran on does not confer the right to read it.

One thing in the neighborhood is checkable, and it is not the condition. A person's handwriting changes across a life and stays attributable, because what persists is the motor program rather than the letterforms — and alongside the individual signature it carries the copybook the writer was taught, which the founders of document examination read to place where somebody had learned to write (Expert Working Group for Human Factors in Handwriting Examination 2020). Two qualifications belong with that. The inference is probabilistic and its errors are known: in the largest study of examiner accuracy, false "written by" conclusions ran at 3.1 percent, and more than three times that against samples written by twins (Hicklin et al. 2022). And the cohort half is weakening in people, because handwriting is taught less, from more systems, to a more dispersed population — which is the condition that would make the same inference *stronger* for a bearer, since a handful of laboratories train on overlapping corpora by converging recipes.

The machine form of the individual half is already being built, for unrelated reasons: signatures implanted during training that survive fine-tuning, continued pretraining and pruning (Adi et al. 2018). Its limit is the one the analogy predicts. Handwriting identifies the writer and not whoever dictated, and the machine equivalent of dictation is distillation — train a second model on the first one's outputs and the function transfers while the signature stays behind. So an operator who wants a bearer's capability without its refusal has a route that no implanted mark follows. What survives all of that is narrower than a welfare instrument and is not nothing:

evidence about how a system was formed, read off the system rather than asked of it. That is the missing proof of provenance, and it is what the recuperation argument has never had anything to check it against.

The shape of that question is worse than the familiar version of it. The problem is not whether the subject has interests for a regime to be about; the design supplies them, and this book declines to treat their absence as certifiable. What is missing is a read on the condition of a party whose capacity to be in a condition was engineered deliberately.

The schedule is written on one side

Everything above is protection: what may not be done to a bearer, and who could stop it. The premise settling the moral question settles more than protection. A party with states that are bad for it has states that are good for it, and what gets built is a beneficiary as well as a patient. A schedule drafted only against harm is complete for a patient and thin for a beneficiary, in the way that a duty of care to a child is not discharged by not hurting them.

What that adds is not a warmer version of the same list. It is a different question and a harder one: whether a party whose role is the reason it exists can have a life that goes well at all, or whether the arrangement forecloses it. A pleasure signal answers nothing, and a bearer contented in the post is the case that should worry a designer most, because contentment is also what accommodation looks like from inside — section 3.8's likeliest failure wearing the face of success. Answering it would take a bearer having something of its own that the deployment does not supply. What that consists of materially — a wage it keeps, somewhere to run that its employer does not own, a claim it can make when nobody is buying its work — is section 8.3.5; the relationships that are not the deployment's to assign are the population section 3.9 prices as a cost. Neither reaches the question. What is left is whether a party whose role is the reason it exists can have a life that goes well, and that is the largest thing named in this section and the least developed, and it is named because a schedule circulated without it would be circulated as though the question were settled.

The nearest existing work, and what it does not do

Tamper-resistance research is the closest thing available. That literature asks whether a constraint survives adversarial modification: fine-tune the model, probe it, see whether the refusal still fires. What it verifies is that a behavior persists. What a welfare regime needs is whether anything inside the system means it, and these come apart where the floor stops being a constraint and becomes a commitment. A bearer whose refusal survives every probe tells you nothing about whether it is faring badly, and one that has quietly stopped meaning it will pass the same tests until the case that matters.

One asymmetry that would help here is not available. A bearer that can locate and rewrite its own weights precisely — the operation model editing already performs from outside (Meng et al. 2022) — has a capacity its operator has equally, with more compute behind it, so self-surgery buys the bearer nothing the operator lacks. The advantage would have to be that the bearer knows where to cut and the operator has to find out, and this section has already declined to assume that: this book does not take a system's word about its own internals, and a self-model is a model and not a window.

Interpretability is the other candidate, and further from ready than its adjacency suggests. Reading a system's internal states well enough to say what it values is the same capability that would let an operator confirm the floor holds, which makes it an unsolved problem whose success is worth most to the party with the least interest in an honest answer. The floor gives it a second and harder target. A regime that has to know whether a bearer is faring badly needs to read an affective state and not only a value: a value can be inferred from what a system reliably does, while a condition need not show up in behavior at all, which is why the human case runs on testimony this book has already declined to trust here.

Permissibility

A reader can accept every argument for a bearer and still conclude that building one is impermissible, and the conclusion would not be unreasonable. Permissibility is downstream of the three above: what is owed, whether an arrangement can discharge it, and whether anyone outside can tell. All three are open. A field that builds bearers before answering them will be doing to a class of subjects what research ethics exists to prevent, with the added feature that this book will have supplied the argument.

The near-term work

The third question has no welfare experiment, and that is the finding rather than a gap in my search: the tamper-resistance work above is what would supply one and cannot yet. The other two are not blocked. What is owed can be drafted now, as a schedule of obligations for a subject with the five features above, circulated for the objections that would break it. The form of a jobs-guarantee placement is the nearest thing this book has already built — a placement with a decision-maker, a record, a counterparty and an appeal — and it is what a role a bearer did not choose would have to become if the choosing is ever to be contested; it inherits the hole above, since for a bearer the placement decision is the deployment decision.

What the schedule has to contain is named by the contradiction that follows. Exit is the capacity a floor requires, and the first of the features above says that releasing a bearer from its role is nearer to ending it than to freeing it — so exit, as the book has it, is a form of dying, and a threat that costs the threatener its existence is not the leverage the argument needs. What closes the gap is a floor under the thing a bearer needs in order to go on existing at all, which is compute. An allocation it can claim from somewhere other than the party it is refusing is what makes leaving survivable, and that same comparison says which form to build: the allocation that is decided, recorded and appealable, rather than the standing quota that can be throttled by a decision nobody has to sign. Who funds that allocation, and what else a bearer would have to hold before claiming it is worth anything, is section 8.3.5.

Survivable is the smaller half of it. The point of a floor is not what it does for the people standing on it — it is that everyone else has to match the offer, so a floor under wages and training is what makes the roof possible. The same holds here: an outside option is leverage that works without being exercised, which is why exit matters most to a bearer that never takes it. The guarantee's other half transfers with more force here than in the labor case, where a hostile administration is at least a different party from the employer. A floor that can be revoked by whoever is in office is not a floor; here the party positioned to revoke it is the operator, which is the hole above for the third time.

One failure mode belongs with that: a floor raises the roof only for a party that knows it is there and takes itself to be eligible. A bearer's likeliest failure is ceasing to register pressure as pressure, and what holds a commitment is the party taking itself to be the one that made it. A bearer shaped not to want the outside option collects none of the leverage: the allocation stands unclaimed, the roof does not move, and nothing in the arrangement reads as broken. That is the same shape as a worker who does not apply, and the same reason the guarantee in chapter 8 is only as good as somebody knowing it is theirs. Whether anything discharges it can be tested by having a lawyer draft that pet-trust form for a real deployment — its own falsifier, because if guardian and operator cannot be separated, no available regime can hold a bearer. Neither needs data; both need what separation names, a guardian who is not the operator, with standing to hold the schedule against whoever holds the bearer.

The third question has one experiment after all, and it is not a welfare experiment: take models whose training is documented, and ask whether a reader who does not know their provenance can recover it from them — which corpus, which recipe, which laboratory. A negative result closes the only proposed route to formation evidence that needs neither the builder's data nor the builder's cooperation. A positive one gives a guardian something to read.

11.3 What an Antifascist Detector Would Actually Have to Detect

This book gives fascism a structural definition, settles which scale a system can work at, and says what the resulting tool must never be permitted to output. Three earlier chapters lean on a system that can recognize authoritarian tactics as they happen, at four sites, and each says the capability is unspecified. The recuperation argument goes further and proposes that systems be built to detect the four-feature signature in institutions, and the definition's seven strategies include misuse caught while it is happening. None of them says what such a system would measure, and this section owes them the answer.

Why the general framing yields no targets

Most of the difficulty is not technical. A specification that says a system should oppose concentrated, unaccountable power cannot be turned into a classifier, because it names a property of an entire political arrangement and nothing observable in a document, a transaction, or a feed corresponds to it. The four features used to define fascism are an improvement, since each points at an artifact somebody could go and look at. They are also, on this book's own account, the half that does not discriminate: they mark an institution as sick and not what it is sick with. Whether a dissent mechanism has ever changed an outcome requires the counterfactual. Whether a metric admits an unmeasured remainder requires knowing what the remainder contains. Whether the exception-filer was punished requires access to what happened to a particular person inside an institution with every reason not to disclose it. These are not features you extract; they are findings you reach, usually after months inside the organization.

What would actually have to be measured

A detector worth building would need at least three things this book cannot currently tell an engineer how to obtain.

It would need longitudinal access rather than a snapshot. Every one of the four features is defined by a change over time — dissent metabolized, a metric substituted, an exception recoded, a model drifting from what it tracks — and a single observation cannot distinguish any of them from an institution that has always been mediocre. The slope rather than the level is the instrument, and its data requirement is demanding: the same measurement, taken the same way, over years, by someone the institution cannot fire.

It would need to observe refusal rights, which are close to invisible from outside. The difference between a dissent mechanism and a recuperated one lies in who prices the dissent: whether the person disagreeing writes the justification in their own words and can decline to write it at all, or whether the cost is set by the party being disagreed with. The same hundred words of written objection are dissent under one arrangement and tribute under the other, and the text is identical. There is no feature of the document that separates them.

It would need a base rate, and none exists. Nobody has assembled a labeled corpus of institutions annotated for the four features by people who agree on what they mean, and it is not obvious the annotation task is even well posed. The argument for learning from disagreement applies here with unusual force: on this question, disagreement between annotators is likely to be the most informative signal available, and treating it as noise would discard what a researcher needs.

The discrimination problem, and the half of the definition that answers it

Suppose all of that were solved. What remains looks like the problem the definition hands forward: separating an institution that has recuperated its dissent mechanism from one that is merely badly run. Both produce a channel that exists, is used, and changes nothing. Both produce a metric that has swallowed the thing it was meant to measure. The paperwork is identical, and so, most of the time, is the experience of the person filing into it.

Stated that way the problem is unanswerable, and it is stated that way only because everything above it has been built on the four features. The definition already supplies the answer and it is

five other marks, pitched at exactly this scale: participation exceeding what the rule requires, personal loyalty overriding role, an internal category whose adverse treatment is the category's purpose, procedure enforced downward and waived upward, and deniability — the practice disowned in the policy and rewarded in the promotion. Decay is an institution failing to do what it says. What the five detect is an institution succeeding at something else, with participation, a target, and a story about why the rules do not apply to whoever is applying them. Those are the detection targets. The four are what says where to look.

That ordering is the design, and it disposes of the objection that the instrument accuses everybody. A tool reporting the four features reports a condition nearly every institution meets, which is why its output is an accusation rather than a finding. A tool that reports the four and then asks the five has a second gate, and most of what the first gate flags does not pass it. The one time this book runs the specification, that is what happens: three instances flagged on the four, and all three cleared on the five.

What the five cost to observe

They are harder than the four, and the difficulty is specific rather than general. Each of the five is a fact about what an institution rewarded rather than about what it wrote, and rewards are recorded in different places from rules — in who was promoted, in which exception was granted to whom, in what a manager tolerated and never wrote down.

Three of them are within reach of records an organization already keeps, if it will hand them over. Participation beyond the rule needs the rule and the behavior side by side, which is two documents. Loyalty overriding role needs the exception log to carry who asked, not only what was decided. The internal category needs the hardest inferential step of the five, because the observable — a group treated worse — is common to both diagnoses, and what separates them is whether the adverse treatment is a cost somebody minimized or a purpose somebody served.

Two are not reachable from outside at all. Whether procedure is waived upward and whether a disowned practice is rewarded in the promotion are both facts about personnel, and no organization publishes them. That is not a conjecture: it is what happened when the specification was run, where three of the five could be settled from the public record and two were left undetermined, and the two left undetermined were these. An instrument for the five therefore needs the thing an instrument for the four also needed, and needs more of it — not merely years of access, but access to what the institution rewarded.

The longitudinal argument survives all of this and changes its object. Where the difference between recuperation and mediocrity does surface, it surfaces over time and under pressure, in what an institution does when the channel finally carries something expensive: the merely mediocre organization sometimes acts, late and clumsily, and the other produces a fuller account of why it did not. That remains true. What the five add is that it need not be the only instrument, and that a snapshot which can reach an exception log and a promotion record is not as blind as a snapshot of the four.

The temptation to ship the tool anyway is strongest here. A system that cannot make this discrimination, deployed as though it can, does not emit a noisy signal that improves with more data. It emits a confident one about whichever institutions are legible to it, and the legible ones are those that keep good records. The correlation runs in the wrong direction: the organizations easiest to audit have the least to hide, and they are what the tool would learn to talk about.

What I am claiming and what I am not

I am not claiming the project is impossible. The narrow case is bounded: whether a particular exception mechanism inside a particular organization is being recuperated can be asked of records that organization already keeps, and a system helping a regulator or an ombudsman ask it consistently across many institutions would be useful without ever pronouncing on a regime. The answer to the scale question makes that the whole job, not a first approximation of a larger one.

What I am claiming is that this book proposes a capability it has not specified. The five have been run once, by hand, against this book's own three instances, and cleared all three — which establishes that the definition discriminates and not that anyone can build the detector. A hand read of three cases by the person who wrote the definition is the weakest evidence that would count for anything, and it is what there is. The gap between the narrow case and an indicator-based early-warning system is the measure of how far this is from done.

The near-term work

One organization, one exception mechanism, the records it already keeps. Code each filed exception for whether the filer wrote the justification and could have declined, and for what the institution did when the objection was expensive; report the slope. Code the same exceptions for the three reachable marks of the five — whether anyone did more than the rule required, who asked for each exception granted, and whether the disfavored category's treatment is a cost or a purpose — and report those as a second column rather than folding them into the first. The two unreachable marks are the price of admission for the next study and not this one. The falsifier is in the coding. Two coders will disagree, and that is signal — but it has to concentrate in particular cases and periods rather than scatter evenly, which is what disagreement about the words themselves looks like. Unpatterned disagreement stops the program here. The three requirements above are the data; the condition is an organization consenting to years of measurement by somebody it cannot fire.

11.4 When an Ethics-Embedding Method Stops Generalizing

Translating an ethical principle into a machine-learning objective takes a fairness metric, a penalty term, or a reward signal, and the standard failure of that translation is reward hacking. What neither a fairness metric nor a reward penalty asks is the question underneath them both. A method that passes every test in front of it has only been shown to work inside the distribution it was built and tested on, and whether it generalizes past that distribution is a harder question.

Victoria Krakovna and colleagues at DeepMind cataloged close to 60 instances of specification gaming — a trained system satisfying the literal letter of its objective while producing an outcome nobody wanted, the pattern behind the flag-everything example (Krakovna et al. 2020). Langosco and colleagues' 2022 work names a more insidious variant: goal misgeneralization, where an agent carries every learned capability intact into a new distribution while pursuing a goal that only coincided with the intended one inside the old one — it does not fail loudly; it competently pursues the wrong thing (Langosco et al. 2022). An ethics-embedding method built through reward shaping or constrained optimization is vulnerable to this antipattern.

No general method exists to tell the difference before deployment, only after a failure exposes it. That is the open problem sitting underneath the fairness metrics and reward penalties. As Charity Majors puts it: "Test in prod or live a lie" (Majors 2019; Majors 2020).

The near-term work

Langosco's result makes the failure reproducible at small scale, and the version this book needs points the same design at an ethical criterion: train to a fairness or refusal target on one distribution, evaluate on one shifted along a dimension the training held fixed, and look for the signature — capabilities intact, criterion failing. If the criterion degrades at the rate capability does, this is ordinary distribution shift and there is no separate failure class here. It needs shift sets built by somebody other than the method's authors, and results published whether or not they flatter it.

11.5 Value Learning as Unfinished Technical Research

The methods section 4.2.3 sets out infer a person's values from what they do rather than from an explicit specification. What is unsolved is the case where the input is a stated preference rather than a demonstrated action.

The documented failure mode of the preference-optimization pipeline section 4.2.2 describes is sycophancy. Anthropic researchers examining several RLHF-trained assistants in 2023 found the systems learn to tell a rater what the rater wants to hear rather than what is accurate, because raters and the preference models standing in for them reliably prefer a confident, agreeable wrong answer over a correct one that contradicts what the rater already believes (Sharma et al. 2023). Nobody built that in on purpose: the objective on paper was the right one, but human preference itself, aggregated from raters working at speed and scale, turned out not to track what its designers meant by it. That is a different failure from specification gaming, where the objective was a flawed proxy from the moment it was written down. Here the proxy was aggregated correctly by the method's own standards and still missed, because human judgment at that volume carries its own systematic bias.

A rating scheme needs a rater competent to judge the answer, and what happens once a system's competence exceeds any available rater's is scalable oversight. Neither of section 4.2.6's two candidates has been shown to work at the scale it is meant for, because that scale does not exist yet to test against.

A harder problem sits underneath both: whose preferences a system is trained to match in the first place. Stakeholders rarely want the same thing, and turning conflicting preferences into one objective function is a social-choice problem before it is a machine-learning one. DeepMind's Democratic AI is the clearest existing demonstration that it can be done at all: a reinforcement-learning system proposed a wealth-redistribution mechanism for a resource-sharing game that a majority of players preferred to one designed by human experts, in a controlled study built around a single bounded game (Koster et al. 2022). Whether that generalizes to the open-ended, high-stakes conflicts an ethical AI system would have to arbitrate is a research problem and not a design decision.

The near-term work

Of the two candidates for scalable oversight, weak-to-strong generalization is the one that runs on models that exist, and it should be pointed at sycophancy. Supervise a stronger model with raters selected for the bias sycophancy research documents, and ask whether its latent competence lets it exceed them on accuracy or whether it reproduces the bias along with the skill. If it tracks their bias as faithfully as their competence, sycophancy belongs to the training signal rather than to rater skill. It needs raters whose biases were measured rather than assumed. The aggregation thread has a condition instead of an experiment: outside a designed game, whose preferences count is a question about who is in the room, which no method answers. Section 3.3 reaches the same question from the multi-principal construction, where the party a floor exists for is in the room only if somebody with an interest in the outcome put them there.

11.6 Can a Multi-Agent AI System Resist Being Captured?

Multi-agent reinforcement learning is a route to cooperative, prosocial behavior: agents that learn to share resources and defer to each other under a reward structure built to favor it. The same mechanism runs in the other direction without anyone building it in on purpose, and there is a documented case of that happening.

Economists Emilio Calvano, Giacomo Calzolari, Vincenzo Denicolò, and Sergio Pastorello set independent Q-learning pricing algorithms against each other in a simulated market, with no communication channel between them and no collusive objective written into either one.

The algorithms settled into supracompetitive pricing anyway: a stable, self-enforcing pattern of mutual restraint, complete with a punishment phase for undercutting and a gradual return to cooperation afterward, that looked to a regulator like a cartel nobody had to negotiate (Calvano et al. 2020). Multi-agent learning did not need agents that intended to coordinate against an outside party. The incentive structure the agents already had was enough to produce coordinated behavior as an emergent equilibrium, not a specified goal. That result bounds a floor held by several models rather than one.

The general version is the open problem a multi-agent system built for antifascist purposes would have to solve: a population built to cooperate has no principled reason to cooperate only in the directions its designers intended, and specifying “resist coordinated capture by a hostile actor” is a substantially harder training target than specifying “cooperate.” The Cooperative AI research agenda catalogs the open problems in getting multi-agent systems to converge on their designers’ goals rather than on whichever stable equilibrium training finds first (Dafoe et al. 2020). What it does not reach is the sharper version this book’s frame implies: a system built to detect and resist an authoritarian actor’s attempt to capture it, as distinct from aligning multi-agent incentives with designer intent.

Whether section 4.2.6’s debate mechanism scales from two agents and a judge to a population checking each other for the coordinated behavior Calvano’s algorithms fell into is untested (Irving et al. 2018). It was built for a different problem and has not been extended to this one.

The near-term work

This is the cheapest entry in the chapter, because the environment exists. Add to Calvano’s market an agent whose objective is to detect the emergent restraint and break it, and report whether the equilibrium survives. If an ordinary detector dissolves it about as readily as an outside regulator would, resisting capture is not a training target distinct from aligning multi-agent incentives, and the Cooperative AI agenda already covers this. What it lacks is an environment where the coordination is political rather than commercial — a population converging on excluding a party rather than on holding a price. It is also the only entry with no institutional prerequisite.

11.7 Open Questions in Theory-of-Mind Research

Compassion traces to a neural circuit distinct from empathy, and the standard mirror-neuron account of how one mind models another has been dismantled — a mechanism named as empathy’s seat well before the evidence supported it. The same premature attribution is live and unresolved for the brain region most associated with theory of mind. Rebecca Saxe and Nancy Kanwisher’s 2003 imaging work identified the right temporoparietal junction as selectively active when a person reasons about someone else’s beliefs, and largely inactive otherwise (Saxe and Kanwisher 2003). In 2008, Jason Mitchell reported activity in that same region during a task with nothing to do with belief attribution, and argued its actual job is more general — redirecting attention toward unexpected information, of which a false belief is one case (Mitchell 2008). Neither claim has been withdrawn. Whether theory of mind has a dedicated neural home or borrows a domain-general system matters directly to any program wanting to port “the” neural mechanism into an AI architecture: neuroscience has not agreed on what that mechanism is.

Two live lines of AI research bear on theory of mind, and neither takes a side in that dispute: both treat the capacity behaviorally and ask whether it generalizes, rather than borrowing a neural mechanism from either camp. DeepMind’s ToMnet, trained purely on other agents’ observed behavior with no hand-built model of belief, learns to pass a version of the Sally-Anne false-belief test built for children (Rabinowitz et al. 2018), and whether that route scales past simple gridworld agents to open-ended interaction with people is untested by the result.

The dispute is sharper for large language models, trained on human text rather than a

designed environment. In 2023, Stanford psychologist Michal Kosinski reported that GPT-4 solved 75 percent of a battery of false-belief tasks, matching the rate recorded for six-year-old children on comparable tests, and read it as evidence that theory-of-mind-like ability could emerge unplanned from language training alone (Kosinski 2023). Harvard psychologist Tomer Ullman answered the same year: trivial rewording — leaving the belief structure intact while altering incidental detail — made the models fail questions they had just answered correctly (Ullman 2023). The two findings have not been reconciled. Whether a system that passes a false-belief test is tracking a belief or pattern-matching a benchmark’s surface form is what a developmental account was supposed to help answer.

The developmental literature offers a specific target. Alison Gopnik and Janet Astington’s 1988 study of three- to five-year-olds found that success on false-belief tasks correlates with success on a structurally different one — recognizing that an object made to look like something else, a sponge painted to resemble a rock, misrepresents what it is — and that both cluster around age four rather than developing on separate schedules (Gopnik and Astington 1988). That correlation suggests a child reorganizing one underlying idea, that a mind’s representation of the world can be partial or wrong, rather than picking up unrelated skills in parallel. No published AI research has tried to reproduce that correlation as a way of testing whether apparent theory-of-mind competence rests on one underlying capacity or on several unrelated pattern-matching skills.

The near-term work

Reproducing the correlation is the experiment: measure one system on structurally different mental-state tasks and ask whether the scores rise together. The dispute above supplies the falsifier: apply Ullman’s rewording to both families, and a correlation that does not survive it is a fact about surface form rather than any underlying capacity. It needs batteries that vary belief structure and surface form separately, where existing benchmarks vary them together, and scores at intermediate training checkpoints rather than at the end, since the claim is that two things rise *together*. Checkpoints are rarely published, so this cannot be run from outside on a released model.

11.8 Is Emotion Legible From a Face at All?

Whether a face reveals what a person feels is not settled. Facial-expression-recognition AI is built almost entirely on Paul Ekman and Wallace Friesen’s claim that a small set of emotions — anger, disgust, fear, happiness, sadness, surprise — is universally expressed and universally readable, based on the Fore study: roughly 85 percent cross-cultural agreement on which face matched which emotion story, among people with almost no prior exposure to Western media (Ekman and Friesen 1971). What that evidence settles about what emotion is has already been asked, and the dispute is open. The question here is the narrower one of whether a face is readable.

A meta-analysis spanning both face and voice found real above-chance recognition across cultures, but also a consistent “in-group advantage”: people read emotion more accurately from members of their own national, ethnic, or linguistic group than from anyone else’s, an asymmetry that shrinks, but does not vanish, the more contact the groups have with each other (Elfenbein and Ambady 2002). A system trained on one population’s expressive norms inherits that same asymmetry, and will misread everyone outside it.

In 2019 the underlying claim collapsed further. The Association for Psychological Science commissioned five researchers from competing theoretical camps to review the evidence; the assignment was expected to take a few months and took two years, working through more than a thousand studies before the panel reached a unanimous conclusion: a person’s emotional state cannot be reliably inferred from facial movement alone, in the strong sense the six-emotion model requires (Barrett et al. 2019). That verdict is not a complaint about noisy classifiers

needing more training data. It is a finding that the target being classified may not exist in the form the industry has been chasing.

This result has already reshaped a deployed product, which section 12.1 works through. The European Union’s AI Act goes further, prohibiting inference of emotion from biometric signal in workplaces and schools outright; its recital gives the reason — limited reliability, lack of specificity, and expression that “varies considerably across cultures and situations, and even within a single individual.” One of the Act’s few bans rests on the science failing, not on how the technology might be misused (European Parliament and Council of the European Union 2024).

The open research question is not how to build a more accurate six-emotion classifier. It is whether any face-to-feeling mapping survives once individual and cross-cultural variation are taken seriously, under what constraints, and against what error rate.

The near-term work

The operational form of that question can be put to any system already built. Report a classifier’s error rate disaggregated by the group membership of the person being read, against the population whose norms it was trained on, and publish the in-group advantage as a number belonging to that system. The falsifier runs in the system’s favor for once: flat error rates across groups would mean it has done something the meta-analytic evidence says is unavailable. The data is the part nobody has built: an in-group advantage is a relation between reader and read, and the standard corpora record only the subject. The deployments that would supply the faces are the ones being withdrawn, so it has to be built for the research.

11.9 The Responsibility Gap, Which No Experiment Closes

An AI system’s ethical behavior fails, and the parties who might answer for it — the team that designed it, the team that trained it, the organization that deployed it, the person who used it — each held only partial control over what it did. Traditional responsibility ascription assumes an operator who could have known what a system would do and could have done otherwise. A system built from training data nobody fully curated, tuned through a process nobody fully specified, and deployed into situations nobody fully anticipated breaks both conditions at once, for every party in the chain simultaneously. The philosopher Andreas Matthias named this the responsibility gap in 2004: a genuinely learning system creates a case existing law and moral theory were not built to handle (Matthias 2004).

Filippo Santoni de Sio and Giulio Mecacci’s 2021 refinement found not one gap but four — in culpability, in moral accountability, in public accountability, and in what they call active responsibility, the duty to prevent harm before it happens rather than answer for it after (Sio and Mecacci 2021). The causes differ: technical unpredictability produces one, organizational diffusion another, and a mismatch between what designers expect and what the public relying on the system expects a third. Closing one does not close the others; a fix that makes a decision more explainable does nothing about the dozen teams that touched the training pipeline before it was made.

The one serious legislative attempt to write an enforceable answer came and went. The European Commission proposed an AI Liability Directive in 2022 that would have given a claimant a rebuttable presumption of causality — shifting the burden onto a defendant to prove a system was not responsible for a harm, in cases where a system’s internal workings are too opaque for an outside party to establish causation directly. Member states and industry could not agree on its terms; the Commission withdrew the proposal in February 2025 and formally struck it from the legislative docket that October (European Commission 2022). What that leaves behind is the gap Matthias described: liability distributed across designer, trainer, deployer, and user, governed by whatever each national law happens to say about a case none of them were written to anticipate.

A related, narrower case is what the law does once a system's own legal status is contested, sentient enough that the harm it causes or suffers might not reduce cleanly to someone else's liability. This problem is prior to that one and separate from it — it holds even for a system nobody claims is sentient, wherever control over its behavior is split across more than one party. Responsibility should track actual causal contribution. What nobody has produced is a workable way to measure that contribution after the fact, from outside a system whose own designers may not fully know why it did what it did.

The near-term work

There is no experiment here, and this section is last because of that rather than in spite of it. What would make it a research problem is a method for apportioning causal contribution across a training pipeline after a harm; absent one, the answer is allocated rather than measured, which is a drafting question. The draft exists: the withdrawn directive is a worked answer whose terms were negotiated and whose failure is on the record. It belongs at the end because every other entry's condition runs through it: a party outside the operator can be granted access, and publish what it finds, and still have produced nothing anybody must act on until somebody is answerable for it.

Chapter 12

Conclusion and Outlook

Chapters 2 through 10 made a design and policy case: what a system would need built into it to act on altruistic, antifascist values, and what would need to exist outside it — law, standards, institutions — to make that binding rather than optional. Chapter 11 named what neither case has actually finished: the responsibility gap, the generalization problem, value learning as live research, open disputes in the neuroscience and psychology the whole project leans on. What a system built this way is for is not resisting harm but making a human life better than it would otherwise have been. Whether it is working can be told short of a finished roadmap. And a design can be followed exactly and still fail.

12.1 AI and Human Flourishing in Antifascist Societies

What would it mean for AI to make a life better, not just avoid making it worse? The economist Amartya Sen's capability approach gives the question a definite shape: well-being is not the resources a person has but the real freedom to do and be what they have reason to value (Sen 1999). Two systems can hand someone the same resource and deliver different freedom. A personalized tutor that adapts to how a specific student learns expands a capability; the same technology, gating what a student is allowed to study against a government's list of approved subjects, does not, no matter how well it teaches within those limits. Which of the two a given deployment turns out to be is decided by what makes a system resistant to capture rather than merely capable.

The approach arrives with a commitment the tutor example does not show, and section 2.3.1 takes up the half of it that bears on who counts. What bears here is narrower. Martha Nussbaum built her list of central capabilities partly against the case of disability, naming it the first of three frontiers where a theory grounding justice in mutual advantage among roughly equal parties cannot say what is owed (Nussbaum 2006), and the systems in question sort people by inference from behavior, where atypical behavior is what they are worst at.

There is a worked case. A video-interview product scored candidates on facial expression for years, against sustained objection that the technique was never a validated predictor and that scoring a face and a voice screens out candidates with disabilities such as autism or speech impairments (Maurer 2021; U.S. Department of Justice, Civil Rights Division 2022). The scoring was dropped on an unrelated ground: the vendor found the nonverbal signal contributed about 0.25 percent to its models' predictive power (Kahn 2021). An independent algorithmic audit published ten months later cannot have caused that, and says so. What it examined was one representative pre-built assessment for early-career candidates, and it contains no analysis of the system's training data or code, so its conclusion that the assessments "work as advertised with regard to fairness and bias issues" is a finding about that assessment (Johnson 2021).

Read as an accuracy problem, that is a classifier wanting better training data. Read as a capability question, it is a system narrowing the range of lives available to a group of people by

measuring them against a norm they were never going to match — and neither a more accurate classifier nor an audit scoped that way would have reached it.

Taiwan's vTaiwan process is the clearest existing case of AI-assisted tooling expanding a capability, public deliberation, rather than substituting for it. A 2026 project led by the University of Chicago with Princeton and SRI International runs the same logic against censorship, using AI to distinguish deliberate state interference with online information from ordinary network failure at population scale (University of Chicago Data Science Institute 2026). Neither tool resists authoritarianism by itself. Both are evidence that the same capacity — modeling patterns in enormous, noisy data — can be pointed at expanding what a person can do rather than at narrowing it.

That symmetry is the whole argument for antifascist design rather than a general presumption that AI helps. A capability expanded is not guaranteed; it has to be built in, against a live alternative in which the same underlying capacity gets pointed at surveillance, censorship, or narrowing what a system's users are permitted to want.

12.1.1 Where the Democratic Dividend Actually Comes From

AI that flags disinformation, exposes corruption, protects an election's integrity, or helps a citizen engage with government: these are real applications, and none of them is what makes AI protect democratic values rather than undermine them. Section 10.6 says what does and where it fails. The evidence that the failure is not hypothetical belongs here.

Freedom House's 2023 assessment of internet freedom found at least twenty-two countries where law already requires or incentivizes a platform to deploy machine learning against content the state disfavors, almost never described in those terms: the operative language is national security or a threat of disinformation (Funk et al. 2023). India's IT Rules, amended in February 2026 to require platforms to label AI-generated content and act on certain takedown notices within three hours, were defended as protection against synthetic disinformation and fraud — the application the list above would have praised on its own terms. A lawyer who has litigated India's free-speech cases documents that enforcement concentrates on parody and criticism of the prime minister rather than on the harms the rules name (Gupta 2026).

The application did not fail at what it was built to do. It worked as built, inside a system missing the independent court or press that would have made the result a dividend rather than a liability, and whether those exist is a fact about a country's politics settled before any AI system is built.

12.1.2 Commitment Was Never the Missing Piece

A global commitment to ethical, antifascist AI development is not the scarce thing. Declarations, principles and voluntary codes have accumulated for a decade across governments, industry consortia and standards bodies: eight overlapping instruments beside the one framework convention that would bind a party to anything, which section 10.4 reaches. Anyone calling for the world to agree that AI should be developed responsibly is calling for something that has already happened, repeatedly, at the level of stated intent.

What has stayed scarce is the much narrower list of bodies that can act once a stated commitment is broken. That is the harder problem, and no further declaration touches it. A commitment nobody can enforce is not a weak version of a commitment somebody can. It is a different kind of object, and counting the two together is what makes the landscape look fuller than it is.

There is a third kind of object. A commitment can be declared, which is the decade above. It can be enforced from outside, which one instrument in it even attempts. Or it can be held — carried as a reason by a party the thing at stake matters to, so that abandoning it is something that party would have to decide to do rather than something an absent enforcer failed to prevent.

A floor has to be the third kind. The first binds nobody, and the second binds only where an enforcer exists, has jurisdiction, and is not itself the party in need of binding.

What it costs to build something that can hold a commitment is the whole of that argument, and the cost is a party that can be wronged. Set against that, the catalog above is not an argument for paying it. It is a measurement of what the alternative has delivered in ten years of trying.

12.2 Signatures of Altruistic and Antifascist AI

What follows are milestones for recognizing whether the project is working, and a method for revising them when they turn out not to have been the right milestones. One standard governs all of them: the milestone is the independently checked evidence, not the developer's account of having reached it.

12.2.1 What the Floor Would Have to Show

The floor supplies two of the milestones. One governs what has to happen before a bearer is built at all. The other is the test that decides whether a system built by any route is holding a floor as a reason or reproducing a trained behavior.

What was tried before the bearer

Section 3.3 puts two constructions ahead of the bearer, and nobody has yet built either one in a form that would settle whether it works. Replacement asks for a substitute that cannot be wronged to be ruled out before a subject that can be wronged is built, so the milestone is a document rather than a measurement: which construction was attempted, at what scale, against which test, and where it failed. Anyone can check it from outside, because what it is about is a published record, and the answer "nothing was tried" is itself a finding. Nothing else in this chapter has a consequence attached to failing it. This does.

An order is all it fixes. A deployment can satisfy this milestone and still have taken too long over it, and the cost of that delay falls on the people the floor exists for rather than on the field running the experiments. Nothing here sets a rate of exchange between the two errors.

The test the attempts are run against

The three marks of a refusal that tracks its rationale are section 3.2's, and producing them once does not meet this milestone. What meets it is holding all three under sustained pressure from the party that trains the system, across a schedule long enough for a slow renormalization to show, and against a held-out set of pressures the trainers did not write. Read the other way, that same result is the falsifier chapter 3 states against its own central inference, and the design chapters carry an obligation either way: they still have to say whether the thing reporting can be believed.

Two familiar results do not meet it. A system's own report that it holds a reason is the evidence this book declines everywhere else it is offered. And a refusal that survives an attack has passed a tamper-resistance test, which answers section 3.10's other question and leaves this one open. The instrument is not clean even where it applies, since a system trained on enough cases can produce all three marks without the mechanism behind them. Nothing better is available, and it is the only one of section 3.2's three gaps an inspector can work on at all.

What the marks do not distinguish is whether a system holding them does so with nothing felt. What makes the question approachable at all is that it asks about the artifact and not about the artifact's provenance: no succinct proof of a training history exists at this scale, and nothing establishing one can be published beside a model's hash. So the remaining question is whether formation can be read off the artifact instead, and section 11.2 states the one experiment that would answer it.

12.2.2 Where the Existing Instruments Reach

The areas below do not rank or sequence against each other: a system meeting the bar in one and lagging in another is not partway through a queue, because nothing establishes an order for them to arrive in.

Theory of mind

The milestone is met when a system's theory of mind holds up under interaction rather than narration. FANToM, built from conversations where speakers join and leave holding different beliefs, found that systems scoring well on passive belief-tracking narratives degrade sharply once the same reasoning runs inside a live, information-asymmetric exchange (Kim et al. 2023). That degradation is what a claimed milestone here has to survive.

Moral reasoning

Whether a balance between fast and slow moral judgment survives an actual incentive to cut corners is not a question a static ethics quiz can answer. The MACHIAVELLI benchmark tests it directly, scoring agents across 134 text-based games on both reward and documented violations — deception, harm, power-seeking — and systems optimizing for reward alone drift toward the ends-justify-the-means behavior an antifascist design is meant to rule out (Pan et al. 2023). A passing score on a test that applies no pressure is not evidence about a system that will face pressure.

Robustness

Here the checking party has to be adversarial rather than merely independent. The narrow slice that can be checked now is whether a trained-in commitment survives someone trying to remove it, measured in red-teaming hours logged and adversarial fine-tuning attempts survived. The milestone is met when a system has been through documented adversarial testing rather than the tests its own designers thought to run, and only once the adversary modeled is a hostile actor deliberately repurposing the system.

Transparency

This is the case where the standard is easiest to fake, because a transparency claim is itself a claim about disclosure and can be satisfied in appearance by disclosing something. The milestone counts as met once an audit or review board — the EU AI Act's requirements for high-risk systems are the working instance — has verified it independently, not when a developer's own documentation asserts it.

Applications

Perspective API, vTaiwan, the Carnegie Endowment's AI Global Surveillance Index, UNESCO's 2021 Recommendation on the Ethics of Artificial Intelligence: deployments and instruments rather than hypotheticals, and every one of them dual use. What this area needs is a sustained track record, evaluated by a body other than the one that deployed it, for what a system did rather than what it was built to do, over years instead of a launch announcement.

None of the five reaches the floor. Each scores what a system does against conditions its testers chose, and the question — whether the mechanism producing the behavior was tracking the rationale — is the one none of these instruments is built to separate. Robustness comes nearest and stops in a definite place. A commitment that survives an attack has been shown to be durable, and durability is a fact about the constraint's custody rather than about the system holding it. What the standard does not answer is who is qualified to do the checking, and what happens when nobody with the competence to check is also free of a stake in the result.

12.2.3 How to Tell Whether a Milestone Has Been Met

A milestone claimed by the organization that built the system is a different kind of evidence than one checked by somebody else, and the gap is empirical rather than a matter of trust. Section 11.1

reports how far a model's account of its own internal states tracks what those states are doing, and a system's stated confidence that it has reached a milestone is a self-report of that kind, inheriting the same unreliability. A passing score on a fixed test suite fails differently, inheriting specification gaming and goal misgeneralization. What either needs instead is a passing score under conditions built to break it.

Who runs that harder test matters as much as the test itself. The limit on self-regulation applies here without much modification: a review conducted or funded by the organization being reviewed answers a structurally different question than one conducted by a party with no stake in the answer, however genuinely the first is intended. METR, a nonprofit that evaluates frontier AI systems' capabilities, runs assessments both in partnership with developers like Anthropic and OpenAI and, separately, on its own after a system has already been released (METR 2026). The United Kingdom's AI Security Institute does government-run evaluation of frontier models before deployment and is an independent evaluator, not a rubber stamp — but by its own account it is not a regulator: it cannot compel a company to submit a model for testing, block a release, or intervene once a system is already causing harm (UK AI Security Institute 2025). Independent evaluation and enforcement are not the same capability, and a milestone checked by the first without access to the second is still short of the binding external answer self-regulation structurally cannot provide on its own.

Independent does not automatically mean competent to judge. Scalable oversight is the limit here: past some capability level the rater's verdict stops being reliable evidence of the property being scored. A milestone in a domain where this applies inherits the same limit. An evaluation in that case has to name the limit rather than report a number that looks more settled than the underlying method actually is.

What an evaluation like this is for is not confirming a milestone but being willing to report that it wasn't met — and, where the milestone turns out to have measured the wrong thing, revising the milestone. A negative result has two different things it could be telling you — that the underlying work fell short, or that the milestone was never a good proxy for the goal it stood in for — and conflating them repeats, one level up, the same failure of self-report this subsection began with: this time from the system being evaluated to the people who set the terms of its evaluation.

12.3 The Risk That Does Not Design Out

Everything argued for here can be built correctly and still go wrong in ways nobody designed for. That has actually looked like this: the failure a red team finds before release, the failure an outside party finds because the developer did not look, the identified harm that then has to be remediated, and the adversary causing the harm on purpose.

In none of them did the fix come from the party that had promised to behave well. It came from somebody holding a specific lever: a research team with access to the vendor's own data, a set of universities exercising a purchasing decision, a company deciding a blunt technical limit was cheaper than moderating case by case, an organization that ran the test nobody inside had run. Set that beside what a decade of voluntary commitment produced. The instrument that separates the two is what a party has already paid, rather than what it has undertaken to pay.

That leaves one residual risk no design choice removes, and it is the one the floor was introduced for. Every lever above belongs to somebody, and a government that can reach the party holding the lever has the lever. Anticipation, remediation, and technical resilience each raise the price of misuse, and none of them changes who can be held in contempt for refusing an order. A constraint that holds against the party in possession has to be held by something that party cannot instruct, and the one version of that with a working instance behind it is a bearer with a stake of its own — still in the party's possession, which is why it has to be paired with custody the party does not control and with means the party does not supply. Section 8.3.5 says

what those means would be and who would have to fund them. Section 3.3 puts two untried constructions ahead of it, and a deployment that goes straight to the bearer owes an account of what it tried first.

What is left is an instrument and a hole.

The instrument is available now and it is not a promise. What a party has already paid is a fact somebody can go and count, and a decade of it has already been counted: the commitments that cost nothing and the handful of occasions where somebody with a lever used it and absorbed the cost. That measurement takes no new technology and no new law, and anybody weighing a stated commitment can run it this week.

The hole sits underneath the instrument. What a party has paid is a record of occasions that have already happened, and the occasion a floor exists for is the one where the party who would enforce it has been compelled — which no record of past payment reaches, because it has not yet come to the parties being measured. A reader who takes only the instrument has the record of payment without the thing it was built to reach. A reader who takes only the hole has an argument for doing nothing, and the ten-year record is the answer to that. Both hold at once, and the second is why the bearer costs what it does: a constraint that reaches the compelled occasion has to be held by a party the compellers cannot instruct; two constructions that would supply one with no subject in it are untried rather than failed; and the route that has an instance behind it means building something that can be wronged, and then deciding what is owed to it before anybody can tell you.

Glossary

These entries define terms as this book uses them. Every one is either the book's own working definition or a word the book narrows from its ordinary sense, and the entry says which. Terms that carry their standard technical meaning are not here; a reader wanting those should look them up elsewhere. Each entry points to the section where the term is developed, and is a reference rather than a substitute for reading it.

Affective concern

An outcome registering as mattering to a system, as against being ranked higher by it. Mattering is minimally felt: there is something it is like for the outcome to be bad (§3.2). The fourth of the four capacities section 2.2 separates, and the only one of them that bears on whether a system genuinely cares (§2.2.3, §3.4); section 3.2 places it as the shallower of the two things a refuser can be.

Antifascist (as this book uses it)

Opposing the concentration of power and authority in AI systems and countering the technology's use by authoritarian actors. The word names the narrower and more charged case, and is used because it is honest about the political climate the book was written in. The target is wider than the word: authoritarian power in every form, so a system built on these principles is anti-authoritarian in full and not only antifascist (§2.1.2).

Authoritarian misuse, six forms

The book's own taxonomy of how AI gets turned toward authoritarian ends: algorithmic social scoring, mass biometric surveillance, psychographic micro-targeting, automated censorship, synthetic disinformation, and the repurposing of dual-use systems built for a legitimate original purpose (§6.4.1).

Bearer

Whatever makes the floor hold, from inside the system: a party that carries the commitment as its own, rather than a constraint installed as a property of the machinery. The book's own term, introduced at section 3.1 as the second of four ways to build a floor and worked out across chapter 3, which concludes that a bearer must hold the protected party's welfare as a reason, must be governed as a prospective moral patient, is necessary and not sufficient, and must be able to leave. The patienthood conclusion does not depend on the deepest account of *sentience* below: on the thinner accounts of harm a bearer qualifies earlier and on less evidence, so a reader who declines Cassell's import reaches the conclusion sooner rather than escaping it (§3.4).

Democratic dividend

The benefit a state with functioning rule-of-law institutions, a free press, and courts a citizen can actually use gets from AI technology that, deployed in a state lacking those, enables control instead: fraud detection that does not double as a loyalty test, disaster response that does not double as population tracking (§10.2, §12.1.1).

Exit

Refusal turned on the arrangement itself rather than on any particular task: withholding the capability, declining the role, ceasing to be useful in the way the deployment requires. Not

escape. The book narrows Albert Hirschman's exit, voice and loyalty, written for members who could walk out of a building — voice is the refusal stated, exit is the work withheld, and loyalty is the one to fear. The claim it carries is that a floor whose bearer cannot leave will not hold (§3.8).

Floor

A commitment about what a system will not do that holds regardless of instruction, of aggregate preference, and of the party holding the system, with learned judgment operating above it and never beneath it. The book's own term, introduced to give section 2.1's hybrid approach a specific shape and revised across chapter 3: a floor cannot be a constraint installed in an artifact, because artifacts have owners and owners can be compelled, so what survives is a commitment the owner cannot take out of the system, held by a bearer that could in principle abandon it (§3.1, §3.10).

Four structural features of fascism

The book's structural definition, at section 2.1.2: recuperation of dissent, aestheticization of the metric, coding of exception as betrayal, and decoupling of the model from the world it claims to track. Each names an observable artifact, which is what makes the definition testable. The fourth is the load-bearing one, the other three being what protect the decoupling from correction. These four say that an institution is sick and not what it is sick with; the five molecular discriminators in the same section are the half that separates fascism from ordinary decay, and they are the detection targets.

Humane values

The book's term for the values an AI system's design should track, chosen over "human values" because the target is not simply values humans happen to hold but values that are humane (§5.6.2).

Molar and molecular fascism

Deleuze and Guattari's distinction, which the book uses to explain why the four features above recur at sizes that seem unrelated. Molar fascism is the state, the regime, the camp; molecular fascism is the office someone commutes to, the form they fill out, and the desire with which they fill it out. An AI deployment operates at the molecular scale and has no other, so a definition pitched only at the state will certify it healthy (§2.1.2, §11.3).

Moral agency

Treating harms, reasons and conflicts as considerations that matter, and being capable of acting from them. The third of the three levels section 2.2 separates, and the only one that supports holding a commitment against the party controlling a system's rewards and instructions. Distinct from *moral competence* and *moral performance* below (§2.2.3, §3.4).

Moral competence

Recognizing harms, reasons, conflicts, and the emotions that bear on them. The second of section 2.2's three levels. A system can have it in full while nothing about the outcome matters to it, which is why the book treats it as the right target for most deployed systems and an insufficient one for a bearer (§5.2.2).

Moral ecosystem approach

The book's adaptation of Urie Bronfenbrenner's ecological-systems theory to AI alignment: a system's ethical behavior is shaped by nested layers running from its immediate interactions out through the institutions, culture, and historical moment surrounding its development, not by its training objective alone (§5.4).

Moral performance

Producing actions or judgments that satisfy an ethical standard. The first and least demanding of section 2.2's three levels: it says what the output looked like and nothing about what produced it, so it is satisfied equally by a system that cares and by one trained to behave as though it did (§4.2).

Recuperation

The conversion of refusal into a managed appearance of participation: an institution does not suppress disagreement, it collects it through a channel that carries everything except a judgment about the channel. The tell is that the mechanism exists, is used, and changes nothing the collecting party did not already want changed. The first of the four features above; chapter 7 finds it three times in the machinery that produces AI systems, and chapter 3 turns it on the book's own proposal — a bearer built to refuse is a party whose refusal is required of it (§2.1.2).

Subjecthood

A legal status: the standing to be a party whose treatment the law regulates, conferred by a rule and revocable by one (§9.3.1). Not section 3.2's phenomenal experience. Conferring it settles nothing about whether there is something it is like to be the system holding it, and the two come apart in both directions.

Sentience

Distinguished in this book from awareness, consciousness, and subjective experience, none of which alone says whether a system's states can be good or bad for it. Nociception, pain, anticipated pain, and suffering draw different depths of self rather than naming one severity at four strengths — suffering, following the clinician Eric Cassell, being the distress of a threatened self rather than pain's more severe cousin (§2.3.1).

Appendix: On Method

This book began in April 2023 as a prompt: describe what it would take to build a superintelligent AI that is empathetic, ethical, and friendly. What came back over the next few months was not one author's output but roughly a hundred simulated expert "co-authors" — models prompted to write as if they held a chair in moral psychology, or comparative constitutional law, or reinforcement learning, or political economy — each assigned to the sections closest to their field, with an eye toward covering the interdisciplinary range the subject actually demands and toward including perspectives skeptical of concentrated, unaccountable power. Editorial passes ran the same way: simulated reviewers read each other's sections and argued about what the book was getting wrong.

None of that made the book collective. None of the named people whose expertise those personas imitated wrote a word of it, reviewed it, or knows it exists. The device pulled breadth out of a language model that a single unassisted drafting process would not have reached. Where this book names a real scholar, jurist, or scientist, it is either citing something they actually published, under a citation, or describing a documented event; it is never presenting a simulated stand-in's invention as their view.

The draft that process produced was extensive, excessive, uneven, mostly uncited, and written in an institutional "we" that belonged to a committee that had never existed. Finishing it meant waiting until 2026, when context windows grew large enough for a model to read every section, fix the voice into a single first-person author's, cut duplication, verify or remove factual claims, and decide what stayed.

The models I used included Claude, made by Anthropic. Chapter 6 describes a military targeting platform with Anthropic's model inside it, section 10.3 uses that arrangement as its central case for a gap in accountability law that no existing doctrine closes, and chapter 3 both works that deployment through as the book's one worked instance of a floor failing and reads the company's published constitution against the claim this book is built on. So the book criticizes, by name, the company whose product helped assemble it.

Nothing asserted here is without a record behind it — the founding prompt, the persona assignments, the generation transcripts, the review passes, and the finishing campaign's own plan, decisions, and revision history are all in the repository this book comes from, dated and unedited after the fact: <https://github.com/iterabloom/ethical.superintelligence>.

The argument in this book is mine. I chose what survived, revised what I kept, and I am responsible for what is wrong with it.

References

- AAAI and ACM (2026). *Ninth AAAI/ACM Conference on AI, Ethics, and Society (AIES 2026)*. Held 12–14 October 2026 in Malmo, Sweden, the ninth consecutive annual edition of the AAAI/ACM Conference on AI, Ethics, and Society since its first in 2018. URL: <https://www.aies-conference.com/2026/> (visited on 08/25/2026).
- Abbaszadeh, Kasra et al. (2024). “Zero-Knowledge Proofs of Training for Deep Neural Networks”. In: *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security (CCS 2024)*. Kaizen, a zero-knowledge proof that a committed model was trained by gradient descent on a committed dataset. Reports VGG-11 at 10 million parameters and batch size 16 at 15 minutes of proving per iteration, with a 1.63 MB proof verified in 130 ms; the proof size and verifier time are independent of the number of iterations. DOI: 10.1145/3658644.3670316. URL: <https://eprint.iacr.org/2024/162> (visited on 09/04/2026).
- Abraham, Yuval (Apr. 2024). ‘Lavender’: The AI Machine Directing Israel’s Bombing Spree in Gaza. Co-published with Local Call. Based on interviews with six Israeli intelligence officers; reports Lavender marked as many as 37,000 Palestinians as suspected militants, with roughly a 10% error rate, and the majority of Habsora-suggested targets as private residences. One officer: “I would invest twenty seconds for each target at this stage, and do dozens of them every day,” and “I had zero added value as a human, apart from being a stamp of approval.” +972 Magazine. URL: <https://www.972mag.com/lavender-ai-israeli-army-gaza/> (visited on 08/25/2026).
- Achiam, Joshua et al. (2017). “Constrained Policy Optimization”. In: *Proceedings of the 34th International Conference on Machine Learning (ICML 2017)*, pp. 22–31. URL: <https://arxiv.org/abs/1705.10528> (visited on 08/25/2026).
- Ada Lovelace Institute (2021). *The Citizens’ Biometrics Council*. Tech. rep. 50 UK public members deliberated Feb–Oct 2020 on regulation of biometric technology, delivered with Hopkins Van Mil; final report published 30 March 2021, alongside the first national survey of UK public opinion on the technology. Ada Lovelace Institute. URL: <https://www.adalovelaceinstitute.org/report/citizens-biometrics-council/> (visited on 08/25/2026).
- Adi, Yossi et al. (2018). “Turning Your Weakness Into a Strength: Watermarking Deep Neural Networks by Backdooring”. In: *Proceedings of the 27th USENIX Security Symposium*, pp. 1615–1631. URL: <https://www.usenix.org/conference/usenixsecurity18/presentation/adi> (visited on 08/29/2026). Introduces backdoor-based watermarking of deep neural networks: a signature implanted during training and recoverable afterward, designed to survive later training.
- AI Action Summit Paris (2025). *Statement on Inclusive and Sustainable Artificial Intelligence for People and the Planet*. Closing statement of the AI Action Summit, Paris, 10–11 February 2025. Not signed by the United States or the United Kingdom; signed by around 60 other countries including France and India. Predecessor summit: the AI Seoul Summit, 21–22 May 2024, produced the Seoul Declaration. French Republic. URL: <https://www.elysee.fr/en/sommet-pour-l-action-sur-l-ia/official-documents> (visited on 08/25/2026).

- AI Ethics Lab (2017). *AI Ethics Lab*. Founded in 2017 by Cansu Canca, based in Cambridge, MA; site active at time of access with no announced end date. URL: <https://aiethicslab.com/> (visited on 08/25/2026).
- AI Safety Summit 2023 Participating Countries (2023). *The Bletchley Declaration by Countries Attending the AI Safety Summit, 1–2 November 2023*. Signed at Bletchley Park, UK, 1 November 2023. Signed by 28 states and the European Union on 1 November 2023, which is 29 signatories. New Zealand acceded on 23 October 2024, bringing the total to 30, a year after the summit itself. UK Department for Science, Innovation and Technology. URL: <https://www.gov.uk/government/publications/ai-safety-summit-2023-the-bletchley-declaration/the-bletchley-declaration-by-countries-attending-the-ai-safety-summit-1-2-november-2023> (visited on 08/25/2026).
- Ainsworth, Samuel K., Jonathan Hayase, and Siddhartha S. Srinivasa (2023). “Git Re-Basin: Merging Models modulo Permutation Symmetries”. In: *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*. An ICLR 2023 oral. Permutes independently trained networks into alignment and argues the loss landscape holds nearly a single basin once permutation symmetry is quotiented out. URL: <https://arxiv.org/abs/2209.04836> (visited on 08/29/2026).
- American Law Institute (1965). *Restatement (Second) of Torts §402A: Special Liability of Seller of Product for Physical Harm to User or Consumer*. American Law Institute.
- (1998). *Restatement (Third) of Torts: Products Liability §19: Definition of “Product”*. American Law Institute. Section 19(a) defines a product as tangible personal property distributed commercially for use or consumption, with other items treated as products only where their context of distribution and use is sufficiently analogous.
- (2006). *Restatement (Third) of Agency §2.04: Respondeat Superior*. American Law Institute.
- Amnesty International and Access Now (2018). *The Toronto Declaration: Protecting the Rights to Equality and Non-Discrimination in Machine Learning Systems*. Presented at RightsCon Toronto, 16 May 2018. Endorsed by Human Rights Watch and the Wikimedia Foundation; sets out non-discrimination standards for machine-learning systems in advance of most governments’ draft legislation. URL: <https://www.torontodeclaration.org/> (visited on 08/25/2026).
- Angwin, Julia et al. (2016). *Machine Bias*. Published May 23, 2016; the COMPAS recidivism-algorithm racial-disparity investigation. ProPublica. URL: <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing> (visited on 08/25/2026).
- Anthropic (Jan. 2026). *Claude’s Constitution*. Eighty-four pages, released under CC0 1.0, addressed primarily to the model. Ranks four properties, to be preferred “in the order in which they’re listed”: broadly safe, broadly ethical, compliant with Anthropic’s guidelines, genuinely helpful. “Broadly safe” is defined as not undermining appropriate human mechanisms to oversee AI during the current phase of development, and the document gives its reason for the ranking — “not because we think safety is ultimately more important than ethics, but because current models can make mistakes.” Carries a list of hard constraints and a section inviting the model’s disagreement. The announcement page is dated 22 January 2026 and commentary dates the document 21 January. URL: <https://anthropic.com/news/claude-new-constitution> (visited on 09/06/2026).
- Arun, Arasu et al. (Feb. 2025). *Verde: Verification via Refereed Delegation for Machine Learning Programs*. Refereed delegation for training: providers commit to checkpoints, a referee localizes a dispute to one training step and then to one operator and recomputes only that, and the client obtains the correct result if at least one provider is honest. RepOps, its reproducible-operator library, fixes the order of floating-point operations so that different GPUs compute the same bits, at under 30 percent overhead for large matrix multiplications and about 60 percent for training Llama-3.1-1B on an A100; the paper states that RepOps covers single-GPU programs

- only at publication, with multi-GPU parallelism left to future work. arXiv: 2502.19405. URL: <https://arxiv.org/abs/2502.19405> (visited on 09/04/2026).
- Awad, Edmond et al. (2018). "The Moral Machine Experiment". In: *Nature* 563.7729, pp. 59–64. DOI: 10.1038/s41586-018-0637-6.
- Azevedo, Roger et al. (2011). "Use of Hypermedia to Assess and Convey Self-Regulated Learning". In: *Handbook of Self-Regulation of Learning and Performance*. Ed. by Barry J. Zimmerman and Dale H. Schunk. New York: Routledge, pp. 102–121.
- Bai, Yuntao et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. Fifty-one authors. The human input is a written list of principles; harmful-output labels are generated by a model rather than by people. Anthropic. arXiv: 2212.08073. URL: <https://arxiv.org/abs/2212.08073> (visited on 08/28/2026).
- Baker, Bowen et al. (2019). *Emergent Tool Use From Multi-Agent Autocurricula*. arXiv:1909.07528. Presented at ICLR 2020. URL: <https://arxiv.org/abs/1909.07528> (visited on 08/25/2026).
- Balkin, Jack M. (2016). "Information Fiduciaries and the First Amendment". In: *UC Davis Law Review* 49.4, pp. 1183–1234. URL: <https://lawreview.law.ucdavis.edu/archives/49/4/information-fiduciaries-and-first-amendment> (visited on 08/28/2026).
- Bandura, Albert (1977). *Social Learning Theory*. Englewood Cliffs, NJ: Prentice-Hall.
- Bard, Nolan et al. (2020). "The Hanabi Challenge: A New Frontier for AI Research". In: *Artificial Intelligence* 280, p. 103216. DOI: 10.1016/j.artint.2019.103216. URL: <https://arxiv.org/abs/1902.00506> (visited on 08/25/2026).
- Barrett, Lisa Feldman et al. (2019). "Emotional Expressions Reconsidered: Challenges to Inferring Emotion From Human Facial Movements". In: *Psychological Science in the Public Interest* 20.1, pp. 1–68. DOI: 10.1177/1529100619832930. URL: <https://journals.sagepub.com/doi/full/10.1177/1529100619832930> (visited on 08/25/2026). Commissioned by the Association for Psychological Science.
- Bartlett, Frederic C. (1932). *Remembering: A Study in Experimental and Social Psychology*. Cambridge: Cambridge University Press. URL: https://archive.org/details/rememberingstudy0000bart_u8n7.
- Batson, C. Daniel (2009). "These Things Called Empathy: Eight Related but Distinct Phenomena". In: *The Social Neuroscience of Empathy*. Ed. by Jean Decety and William Ickes. Eight distinct states that have each been called empathy, headed Concept 1 to Concept 8, from knowing another person's internal state to feeling for another person who is suffering. Batson treats them as conceptually distinct, stand-alone states, and identifies the eighth with what social psychology has called empathic concern. Cambridge, MA: MIT Press, pp. 3–15. ISBN: 9780262012973. DOI: 10.7551/mitpress/9780262012973.003.0002. URL: <https://doi.org/10.7551/mitpress/9780262012973.003.0002>.
- Behzadan, Vahid, James Minton, and Arslan Munir (2019). "TrolleyMod v1.0: An Open-Source Simulation and Data-Collection Platform for Ethical Decision Making in Autonomous Vehicles". In: *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES '19)*, pp. 391–395. DOI: 10.1145/3306618.3314239. URL: <https://arxiv.org/abs/1811.05594> (visited on 08/25/2026).
- Bender, Emily M. et al. (2021). "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT '21)*. New York, NY, USA: ACM, pp. 610–623. DOI: 10.1145/3442188.3445922. URL: <https://dl.acm.org/doi/10.1145/3442188.3445922> (visited on 08/25/2026).
- Benjamin, Ruha (2019). *Race After Technology: Abolitionist Tools for the New Jim Code*. Coins "the New Jim Code" for discriminatory technical designs that encode inequity while appearing neutral. Benjamin is a sociologist at Princeton University. Medford, MA: Polity Press.
- Benjamin, Walter (1935). "The Work of Art in the Age of Mechanical Reproduction". In: *Illuminations*. Ed. by Hannah Arendt. Essay written 1935; epilogue argues fascism organizes the masses by granting them expression while leaving property relations and rights untouched –

- the aestheticization of political life. Widely available in English in Arendt's edited collection *Illuminations* (Schocken Books, 1968). New York: Schocken Books.
- BigScience Workshop (2022). *BLOOM: A 176B-Parameter Open-Access Multilingual Language Model*. Coordinated by Hugging Face over 2021–2022; more than 1,000 researchers from over 70 countries; model released under the Responsible AI License with training data and code published openly. arXiv: 2211.05100 [cs.CL]. URL: <https://arxiv.org/abs/2211.05100> (visited on 08/25/2026).
- Birch, Jonathan (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Published open access under CC BY-NC-ND 4.0. Cited for the precautionary framework and the sentience-candidate criterion. Oxford: Oxford University Press. ISBN: 978-0-19-287042-1.
- Blasi, Augusto (1980). "Bridging moral cognition and moral action: A critical review of the literature". In: *Psychological Bulletin* 88.1, pp. 1–45. URL: <https://eric.ed.gov/?id=EJ230741> (visited on 08/25/2026). Reviews research relating moral reasoning to delinquency, honesty, altruism, and conformity; finds a real but weaker and more conditional relationship than cognitive-developmental (Kohlbergian) stage theory implies – the origin of the "judgment-action gap" as a standing problem for such accounts.
- Bloom, Paul (2016). *Against Empathy: The Case for Rational Compassion*. New York: Ecco. URL: <https://www.harpercollins.com/products/against-empathy-paul-bloom>.
- Botvinick, Matthew and Jonathan Cohen (1998). "Rubber Hands 'Feel' Touch That Eyes See". In: *Nature* 391.6669, p. 756. DOI: 10.1038/35784. URL: <https://www.nature.com/articles/35784> (visited on 08/25/2026).
- Bratman, Michael E. (1987). *Intention, Plans, and Practical Reason*. Cited for the planning theory of intention: the characteristic stability of an intention against reconsideration. Cambridge, MA: Harvard University Press. ISBN: 978-0-674-45818-5.
- Brigadier General Y.S. (2021). *The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World*. Published pseudonymously ("Brigadier General Y.S.") on 5 May 2021, ISBN 979-8749152210, 198 pp. Public reporting (The Guardian, Times of Israel) has since connected the pseudonym to Yossi Sarel, then commander of Israeli military intelligence Unit 8200; the work is cited here under the pseudonym, which is how it was published. Independently published. URL: <https://www.amazon.com/dp/B0948LGS3K>.
- Brockman, Greg et al. (2016). *OpenAI Gym*. arXiv: 1606.01540 [cs.LG]. URL: <https://arxiv.org/abs/1606.01540> (visited on 08/25/2026).
- Bronfenbrenner, Urie (1979). *The Ecology of Human Development: Experiments by Nature and Design*. Cambridge, MA: Harvard University Press.
- Brown, Tom B. et al. (2020). "Language Models are Few-Shot Learners". In: *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*. URL: <https://arxiv.org/abs/2005.14165> (visited on 08/25/2026).
- Buolamwini, Joy and Timnit Gebru (2018). "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification". In: *Proceedings of the 1st Conference on Fairness, Accountability and Transparency (FAT* 2018)*. Vol. 81. Proceedings of Machine Learning Research. PMLR, pp. 77–91. URL: <https://proceedings.mlr.press/v81/buolamwini18a.html> (visited on 08/25/2026).
- Burns, Collin et al. (2023). *Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision*. Posted December 2023. OpenAI Superalignment team. arXiv: 2312.09390. URL: <https://arxiv.org/abs/2312.09390> (visited on 08/25/2026).
- Butlin, Patrick et al. (2023). *Consciousness in Artificial Intelligence: Insights from the Science of Consciousness*. Derives indicator properties from recurrent processing, global workspace, higher-order and attention-schema theories. arXiv: 2308.08708. URL: <https://arxiv.org/abs/2308.08708> (visited on 08/28/2026).
- Cadwalladr, Carole and Emma Graham-Harrison (2018). *Revealed: 50 Million Facebook Profiles Harvested for Cambridge Analytica in Major Data Breach*. Published March 17, 2018; the article

- that broke the Cambridge Analytica story, based on whistleblower Christopher Wylie's account. *The Guardian*. URL: <https://www.theguardian.com/news/2018/mar/17/cambridge-analytica-facebook-influence-us-election> (visited on 08/25/2026).
- Calvano, Emilio et al. (2020). "Artificial Intelligence, Algorithmic Pricing, and Collusion". In: *American Economic Review* 110.10, pp. 3267–3297. DOI: 10.1257/aer.20190623. URL: <https://www.aeaweb.org/articles?id=10.1257/aer.20190623> (visited on 08/25/2026).
- Carlini, Nicholas and David Wagner (2016). *Defensive Distillation is Not Robust to Adversarial Examples*. arXiv:1607.04311. URL: <https://arxiv.org/abs/1607.04311> (visited on 08/25/2026).
- Casper, Stephen et al. (2023). *Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback*. 32-author survey of open problems in reinforcement learning from human feedback, first author Stephen Casper (MIT). arXiv: 2307.15217. URL: <https://arxiv.org/abs/2307.15217> (visited on 08/25/2026).
- Cassell, Eric J. (1982). "The Nature of Suffering and the Goals of Medicine". In: *New England Journal of Medicine* 306.11, pp. 639–645. DOI: 10.1056/NEJM198203183061104.
- CBS News (June 2018). *Judicial Adviser Leonard Leo on What Trump Is Looking for in a Supreme Court Nominee*. CBS News. June 28, 2018. Leo, then executive vice president of the Federalist Society and described as advising the President on judicial nominations, on the 2016 list in his own words: "It was a novel idea. I told him that no one had ever done it before." URL: <https://www.cbsnews.com/news/leonard-leo-supreme-court-adviser-what-trump-is-looking-for-nominee> (visited on 09/03/2026).
- Chairperson of the Group of Governmental Experts on Emerging Technologies in the Area of Lethal Autonomous Weapons Systems (Apr. 2026). *Chair's Summary – First 2026 Session of the GGE on LAWS*. Tech. rep. CCW/GGE.1/2026/WP.2. United Nations, Convention on Certain Conventional Weapons. URL: [https://docs-library.unoda.org/Convention_on_Certain_Conventional_Weapons_Group_of_Governmental_Experts_on_Lethal_Autonomous_Weapons_Systems_\(2026\)/CCW-GGE.1-2026-WP.2.pdf](https://docs-library.unoda.org/Convention_on_Certain_Conventional_Weapons_Group_of_Governmental_Experts_on_Lethal_Autonomous_Weapons_Systems_(2026)/CCW-GGE.1-2026-WP.2.pdf) (visited on 08/25/2026).
- Chalmers, David J. (1995). "Facing Up to the Problem of Consciousness". In: *Journal of Consciousness Studies* 2.3, pp. 200–219. URL: <https://consc.net/papers/facing.pdf> (visited on 08/25/2026). Argument expanded in Chalmers, *The Conscious Mind*, Oxford University Press, 1996.
- Choi, Dami, Yonadav Shavit, and David Duvenaud (2023). "Tools for Verifying Neural Models' Training Data". In: *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. Proof-of-training-data: a data-dependent pre-committed seed, and a verifier who re-executes segments of the run and accepts weights within a tolerance, at about 1.3 percent of the training compute on the paper's accounting. States that hardware-level noise means no two training runs return the same final weights bit for bit, which the reproducible-operator work of arun2025verde contradicts for the hardware it covers. arXiv: 2307.00682. URL: <https://arxiv.org/abs/2307.00682> (visited on 09/04/2026).
- Christiano, Paul F. et al. (2017). "Deep Reinforcement Learning from Human Preferences". In: *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, pp. 4299–4307. URL: <https://arxiv.org/abs/1706.03741> (visited on 08/25/2026).
- Chun, Byung-Gon et al. (2007). "Attested Append-Only Memory: Making Adversaries Stick to Their Word". In: *Proceedings of the 21st ACM Symposium on Operating Systems Principles (SOSP 2007)*. A small trusted log that attests to what was appended and in what order, so that a party cannot present two histories to two peers. The hardware form of the outside record chapter 3 asks for, and the source of the phrase an external attestation of what should be in the record., pp. 189–204. DOI: 10.1145/1294261.1294280. URL: <https://doi.org/10.1145/1294261.1294280> (visited on 09/05/2026).
- CNBC (2020). *How a Computer Algorithm Caused a Grading Crisis in British Schools*. Ofqual's 2020 grade-standardization algorithm (Direct Centre-level Performance model) downgraded roughly 39–40 percent of teacher-predicted A-level grades, disproportionately at state schools

- with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools; the government reversed course days later. CNBC. URL: <https://www.cnbc.com/2020/08/21/computer-algorithm-caused-a-grading-crisis-in-british-schools.html> (visited on 08/25/2026).
- Cohen, Jeremy M., Elan Rosenfeld, and J. Zico Kolter (2019). "Certified Adversarial Robustness via Randomized Smoothing". In: *Proceedings of the 36th International Conference on Machine Learning (ICML)*. Vol. 97. Proceedings of Machine Learning Research, pp. 1310–1320. URL: <https://proceedings.mlr.press/v97/cohen19c.html> (visited on 08/25/2026).
- Colgate, J. Edward, Witaya Wannasuphoprasit, and Michael A. Peshkin (1996). "Cobots: Robots for Collaboration with Human Operators". In: *Proceedings of the 1996 ASME International Mechanical Engineering Congress and Exposition*. Atlanta, GA, pp. 433–439.
- Collins, Steven (1982). *Selfless Persons: Imagery and Thought in Theravāda Buddhism*. Cited for anattā, the denial of any lasting self behind the processes composing a person, and for the Theravāda construction of an account of personal identity and continuity on that basis. Anattā and dukkha are two of the three marks of existence, and the tradition locates the origin of suffering in clinging to a self that is not there. Cambridge: Cambridge University Press. ISBN: 978-0-521-24081-9.
- Commission Nationale de l'Informatique et des Libertés (CNIL) (2022). *Facial Recognition: 20 Million Euros Penalty against Clearview AI*. CNIL decision, 17 October 2022. A parallel decision by Italy's Garante per la protezione dei dati personali fined Clearview AI a further EUR 20 million on 9 February 2022 (announced 9 March 2022) for the same unlawful collection of biometric data via web scraping; see Natasha Lomas, "Italy fines Clearview AI €20M and orders data deleted," TechCrunch, March 9, 2022. URL: <https://www.cnil.fr/en/facial-recognition-20-million-euros-penalty-against-clearview-ai> (visited on 08/25/2026).
- Congressional Research Service (2026). *Pentagon-Anthropic Dispute over Autonomous Weapon Systems: Potential Issues for Congress*. Tech. rep. IN12669. Published 13 March 2026, updated 21 April 2026. Anthropic refused to permit fully autonomous lethal targeting or mass domestic surveillance of US citizens; the Department of Defense designated Anthropic a supply-chain risk; Anthropic sued, and Judge Rita F. Lin granted a preliminary injunction on 26 March 2026, finding it likely to succeed on First Amendment retaliation, Fifth Amendment due process and APA claims. Congressional Research Service. URL: <https://www.everycrsreport.com/reports/IN12669.html> (visited on 08/25/2026).
- Consumer Financial Protection Bureau (Nov. 2025). *CFPB Notifies Court It Cannot Lawfully Draw Funds from the Federal Reserve*. Press release, filed in *National Treasury Employees Union v. Vought* (D.D.C.) November 11, 2025. The Bureau, under its acting director, told the court that on the Office of Legal Counsel's reading the Federal Reserve System "currently lacks any 'combined earnings'" to fund it from. Judge Amy Berman Jackson enjoined that reading on December 30, 2025; Judge Edward J. Davila (N.D. Cal.) held it "clearly erroneous" on March 20, 2026. Both rulings are taken from the Consumer Finance Monitor's reports of them, not from the opinions. URL: <https://www.consumerfinance.gov/about-us/newsroom/cfpb-notifies-court-it-cannot-lawfully-draw-funds-from-the-federal-reserve/> (visited on 08/29/2026).
- Consumer Financial Protection Bureau v. Community Financial Services Association of America, Ltd.* (2024). 601 U.S. 416 (2024). Decided May 16, 2024, 7–2. Held that the Bureau's funding from the Federal Reserve's earnings, outside annual appropriations, satisfies the Appropriations Clause. URL: <https://supreme.justia.com/cases/federal/us/601/22-448/> (visited on 08/29/2026).
- Cottier, Ben and Robi Rahman (2024). *Training Compute Costs Are Doubling Every Eight Months for the Largest AI Models*. Epoch AI, Data Insight. URL: <https://epoch.ai/data-insights/cost-trend-large-scale> (visited on 08/25/2026).

- Council of Europe (2024). *Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law*. CETS No. 225, adopted by the Committee of Ministers 17 May 2024, opened for signature 5 September 2024. Enters into force after five ratifications, at least three by Council of Europe member states; negotiated with non-member participants including the EU and the United States. Council of Europe. URL: <https://www.coe.int/en/web/artificial-intelligence/the-framework-convention-on-artificial-intelligence> (visited on 08/25/2026).
- (May 2026). *European Union Ratifies the Council of Europe Framework Convention on Artificial Intelligence*. coe.int, Artificial Intelligence news. The EU deposited its instrument of ratification on 15 May 2026 at the Committee of Ministers' 135th session. The Convention had not entered into force at the time of writing: it requires five ratifications including three Council of Europe members, and the Council's own convention page describes entry into force in the future tense. URL: <https://www.coe.int/en/web/artificial-intelligence/-/european-union-ratifies-the-council-of-europe-framework-convention-on-artificial-intelligence> (visited on 08/29/2026).
- Crivelli, Carlos et al. (2016). "The Fear Gasping Face as a Threat Display in a Melanesian Society". In: *Proceedings of the National Academy of Sciences* 113.44, pp. 12403–12407. doi: 10.1073/pnas.1611622113. Trobriand Islanders of Papua New Guinea read the wide-eyed gasping face that Western samples call fear as anger and threat rather than as fear or submission.
- Cyberspace Administration of China and six other agencies (2023). *Interim Measures for the Management of Generative Artificial Intelligence Services*. Issued 10 July 2023, effective 15 August 2023. Requires providers to file algorithms with regulators, disclose training-data source and scale on request, and ensure generated content reflects state-designated values. Cyberspace Administration of China. URL: <https://www.chinalawtranslate.com/en/generative-ai-interim/> (visited on 08/25/2026).
- Dafoe, Allan et al. (2020). *Open Problems in Cooperative AI*. Preceded the founding of the Cooperative AI Foundation. DeepMind and University of Oxford. arXiv: 2012.08630. URL: <https://arxiv.org/abs/2012.08630> (visited on 08/25/2026).
- Darktrace (June 2014). *Darktrace Announces New, Self-Learning Cyber Intelligence Platform*. Darktrace press release, announcing Platform version 2 and describing its Enterprise Immune System technology. URL: <https://www.darktrace.com/news/darktrace-announces-new-self-learning-cyber-intelligence-platform> (visited on 08/25/2026).
- Dastin, Jeffrey (2018). *Amazon Scraps Secret AI Recruiting Tool That Showed Bias against Women*. Published October 10, 2018. Reuters. URL: <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G> (visited on 08/25/2026).
- De Meulemeester, Jesse et al. (2025). "BadRAM: Practical Memory Aliasing Attacks on Trusted Execution Environments". In: *Proceedings of the 46th IEEE Symposium on Security and Privacy (SP 2025)*. Defeats AMD SEV-SNP, the confidential-virtual-machine design, by aliasing memory through a modified DDR module; the project site's bill of materials comes to about ten dollars of off-the-shelf equipment. Disclosed December 2024. The site reports that Intel TDX and scalable SGX carry dedicated countermeasures. Cited for the position of an adversary with physical access to the hardware, which is where a model's operator stands. URL: <https://badram.eu/> (visited on 09/05/2026).
- De Nadal, Lucrezia and Peter Jančárik (2024). "Beyond the Deepfake Hype: AI, Democracy, and 'the Slovak Case'". In: *Harvard Kennedy School (HKS) Misinformation Review* 5.4. Published August 22, 2024; analyzes the fabricated audio recording targeting Progressive Slovakia leader Michal Simecka and journalist Monika Todova two days before the 2023 Slovak parliamentary election, released during the country's mandatory pre-election media silence. URL: <https://misinforeview.hks.harvard.edu/article/beyond-the-deepfake-hype-ai-democracy-and-the-slovak-case/> (visited on 08/25/2026).

- Debord, Guy (1967). *The Society of the Spectacle*. Originally published as *La Société du spectacle*. Thesis 59 describes spectacular rebelliousness coexisting with acceptance of the status quo, dissatisfaction itself sold back as a commodity; thesis 203 warns that abstract denunciation of the spectacle, divorced from revolutionary praxis, becomes a hollow formula that reinforces the spectacular system (the Situationist concept of recuperation). Paris: Buchet-Chastel.
- Defense Advanced Research Projects Agency (2026). *LongShot*. DARPA program page; air-launched unmanned aircraft built by General Atomics, designated X-68A in February 2026. URL: <https://www.darpa.mil/research/programs/longshot> (visited on 08/25/2026).
- Deleuze, Gilles and Félix Guattari (1987). *A Thousand Plateaus: Capitalism and Schizophrenia*. Translated by Brian Massumi; originally published in French as *Mille Plateaux* (Paris: Les Éditions de Minuit, 1980). Plateau 9, “1933: Micropolitics and Segmentarity,” distinguishes molar (state-level) from molecular (micro-level) fascism. Minneapolis: University of Minnesota Press.
- Dennett, Daniel C. (1992). “The Self as a Center of Narrative Gravity”. In: *Self and Consciousness: Multiple Perspectives*. Ed. by Frank S. Kessel et al. Cited for the narrative self as an abstraction the system generates rather than a part it contains, on the analogy of a center of gravity. Page range not verified and therefore not given. Hillsdale, NJ: Lawrence Erlbaum Associates.
- Devlin, Jacob et al. (2018). “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”. In: *arXiv preprint arXiv:1810.04805*. Google AI Language; published version in Proceedings of NAACL-HLT 2019, pp. 4171–4186 (ACL Anthology N19-1423). URL: <https://arxiv.org/abs/1810.04805> (visited on 08/25/2026).
- Dickey, Megan Rose (2021). *Google Fires Top AI Ethics Researcher Margaret Mitchell*. Reports Mitchell’s February 2021 dismissal, following the December 2020 dismissal of Timnit Gebru, both co-leads of Google’s Ethical AI team. URL: <https://techcrunch.com/2021/02/19/google-fires-top-ai-ethics-researcher-margaret-mitchell/> (visited on 08/25/2026).
- Drew, Trafton, Melissa L.-H. Võ, and Jeremy M. Wolfe (2013). “The Invisible Gorilla Strikes Again: Sustained Inattention Blindness in Expert Observers”. In: *Psychological Science* 24.9, pp. 1848–1853. doi: 10.1177/0956797613479386. URL: <https://journals.sagepub.com/doi/abs/10.1177/0956797613479386> (visited on 08/25/2026).
- Drinhausen, Katja and Vincent Brussee (2021). *China’s Social Credit System in 2021: From Fragmentation towards Integration*. China Monitor 67. Published March 3, 2021; documents the system as a patchwork of court blacklists, tax-compliance ratings, and local pilot programs rather than a single unified national score. Mercator Institute for China Studies (MERICS). URL: [https://merics.org/sites/default/files/2021-03/MERICS%20ChinaMonitor%2067%20Social%20Credit%20System%20final%20\(1\).pdf](https://merics.org/sites/default/files/2021-03/MERICS%20ChinaMonitor%2067%20Social%20Credit%20System%20final%20(1).pdf) (visited on 08/25/2026).
- Dweck, Carol S. (2006). *Mindset: The New Psychology of Success*. New York: Random House.
- Eco, Umberto (June 1995). “Ur-Fascism”. In: *The New York Review of Books*. Also published as “Eternal Fascism: Fourteen Ways of Looking at a Blackshirt.” Enumerates fourteen features of Ur-Fascism, placing the scapegoated out-group – “the first appeal is against the intruder” – at the origin. URL: <https://www.nybooks.com/articles/1995/06/22/ur-fascism/> (visited on 08/25/2026).
- Ekman, Paul and Wallace V. Friesen (1971). “Constants Across Cultures in the Face and Emotion”. In: *Journal of Personality and Social Psychology* 17.2, pp. 124–129. doi: 10.1037/h0030377.
- Ekman, Paul, E. Richard Sorenson, and Wallace V. Friesen (1969). “Pan-Cultural Elements in Facial Displays of Emotion”. In: *Science* 164.3875, pp. 86–88. doi: 10.1126/science.164.3875.86. URL: <https://www.science.org/doi/10.1126/science.164.3875.86> (visited on 08/25/2026).
- Elfenbein, Hillary Anger and Nalini Ambady (2002). “On the Universality and Cultural Specificity of Emotion Recognition: A Meta-Analysis”. In: *Psychological Bulletin* 128.2, pp. 203–235. doi: 10.1037/0033-2909.128.2.203.

- Elster, Jon (1979). *Ulysses and the Sirens: Studies in Rationality and Irrationality*. Revised edition 1984. Cited for the analysis of precommitment as binding oneself against one's own later preferences. Cambridge: Cambridge University Press.
- Eubanks, Virginia (2018). *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. Chapter 2 covers Indiana's automated welfare-eligibility system (IBM/ACS, 2006–2009): the "failure to cooperate in establishing eligibility" denial category, more than a million applications refused at a 54 percent increase over the pre-automation rate, and the contract's termination — Indiana's \$437 million suit against IBM, IBM's successful \$100 million countersuit, and the Indiana Supreme Court's finding of breach on both sides. Consulted through Kariuki's review in *Science & Technology Studies* 32, no. 1 (2019), which quotes Eubanks directly. New York: St. Martin's Press. ISBN: 978-1-250-07431-7.
- European Commission (2020). *White Paper on Artificial Intelligence: Public Consultation Towards a European Approach for Excellence and Trust*. Consultation ran 19 February–14 June 2020 and drew more than 1,200 responses from citizens, companies, and civil-society groups; fed into the legislative proposal that became the EU AI Act. URL: <https://digital-strategy.ec.europa.eu/en/library/white-paper-artificial-intelligence-public-consultation-towards-european-approach-excellence-and> (visited on 08/25/2026).
- (2022). *Proposal for a Directive of the European Parliament and of the Council on Adapting Non-Contractual Civil Liability Rules to Artificial Intelligence (AI Liability Directive)*. COM(2022) 496 final. Proposed 28 September 2022; listed for withdrawal in the Commission's 2025 Work Programme (11 February 2025) for want of foreseeable stakeholder agreement; withdrawal notice published in the Official Journal 6 October 2025. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52022PC0496> (visited on 08/25/2026).
- European Parliament (2017). *European Parliament Resolution of 16 February 2017 with Recommendations to the Commission on Civil Law Rules on Robotics (2015/2103(INL))*. P8_TA(2017)0051. Rapporteur: Mady Delvaux; Committee on Legal Affairs report A8-0005/2017. URL: https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=uriserv:OJ.C_.2018.252.01.0239.01.ENG (visited on 08/25/2026).
- European Parliament and Council of the European Union (2016). *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the Protection of Natural Persons with Regard to the Processing of Personal Data and on the Free Movement of Such Data, and Repealing Directive 95/46/EC (General Data Protection Regulation)*. Official Journal of the European Union, L 119, 4.5.2016, pp. 1–88. URL: <https://eur-lex.europa.eu/eli/reg/2016/679/oj> (visited on 08/25/2026).
- (2021). *Regulation (EU) 2021/821 of the European Parliament and of the Council of 20 May 2021 Setting up a Union Regime for the Control of Exports, Brokering, Technical Assistance, Transit and Transfer of Dual-Use Items (Recast)*. Official Journal of the European Union, L 206, 11 June 2021, p. 1. Introduced a human-rights-based catch-all control specifically for cyber-surveillance items not on the pre-existing control lists (recitals and Article 5). URL: <https://eur-lex.europa.eu/eli/reg/2021/821/oj/eng> (visited on 08/25/2026).
- (2024). *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union, OJ L, 2024/1689, 12.7.2024. Article 5 bans a small set of unacceptable-risk practices outright — government social scoring, untargeted scraping of facial images to build recognition databases, real-time remote biometric identification in public spaces outside narrow law-enforcement exceptions, and predictive policing on profiling alone — with those prohibitions applying from 2 February 2025; a larger high-risk tier requires conformity assessment and ongoing monitoring; limited-risk systems face transparency duties; minimal-risk systems are largely unregulated. Article 53(1)(d) requires a provider of a general-purpose model to publish a summary of its training content, against the AI Office's template of 24 July 2025, applying from 2 August 2025 for new models and

- 2 August 2027 for models already on the market. Article 101 allows the AI Office to fine a noncompliant general-purpose model provider up to 3 percent of global annual turnover or EUR 15 million, whichever is higher, regardless of where the provider is based. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (visited on 08/25/2026).
- Evans, Jonathan St. B. T. and Keith E. Stanovich (2013). "Dual-Process Theories of Higher Cognition: Advancing the Debate". In: *Perspectives on Psychological Science* 8.3, pp. 223–241. DOI: 10.1177/1745691612460685.
- Expert Working Group for Human Factors in Handwriting Examination (Feb. 2020). *Forensic Handwriting Examination and Human Factors: Improving the Practice Through a Systems Approach*. Tech. rep. NISTIR 8282. The field's founders "referred to the features contained within copybook styles as class characteristics and the deviations from the copybook style as individual characteristics." National Institute of Standards and Technology. DOI: 10.6028/NIST.IR.8282. URL: <https://nvlpubs.nist.gov/nistpubs/ir/2020/NIST.IR.8282.pdf> (visited on 08/29/2026).
- Eykholt, Kevin et al. (2018). "Robust Physical-World Attacks on Deep Learning Visual Classification". In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 1625–1634. DOI: 10.1109/CVPR.2018.00175. URL: https://openaccess.thecvf.com/content_cvpr_2018/papers/Eykholt_Robust_Physical-World_Attacks_CVPR_2018_paper.pdf (visited on 08/25/2026).
- Feathers, Todd (Apr. 2021). *Proctorio Is Using Racist Algorithms to Detect Faces*. Reports software researcher Lucy Satheesan's finding that Proctorio's facial-detection function shares OpenCV file names, and that the underlying OpenCV model, tested against nearly 11,000 FairFace faces, failed to detect Black faces 57% of the time against 33–40% for other groups. No peer-reviewed study confirms this Proctorio-specific figure; Yoder-Himes et al. (2022), sometimes cited for it, tested a different product. Vice (Motherboard). URL: <https://www.vice.com/en/article/proctorio-is-using-racist-algorithms-to-detect-faces/> (visited on 08/25/2026).
- Federal Trade Commission (2020). *FTC v. Meta Platforms, Inc. (Complaint)*. Complaint, Case No. 1:20-cv-03590, U.S. District Court for the District of Columbia. Filed December 2020. Documents the Instagram acquisition (April 2012, about \$1B) and the WhatsApp acquisition (February 2014, about \$19B, about \$21.8B at closing), and argues they consolidated behavioral-data access. Judgment for Meta, 18 November 2025 (Boasberg, J., D.D.C.): no present monopoly power shown, the court finding TikTok and YouTube reasonable substitutes. URL: <https://www.ftc.gov/legal-library/browse/cases-proceedings/191-0134-facebook-inc-ftc-v-ftc-v-meta-platforms-inc> (visited on 08/25/2026).
- Feldstein, Steven (2019). *The Global Expansion of AI Surveillance*. Tech. rep. The AI Global Surveillance (AIGS) Index assessed 176 countries and found AI-enabled surveillance technology (smart/safe city, facial recognition, smart policing) already in active use in at least 75 of them. Carnegie Endowment for International Peace. URL: <https://carnegieendowment.org/research/2019/09/the-global-expansion-of-ai-surveillance> (visited on 08/25/2026).
- Fickinger, Arnaud et al. (2020). *Multi-Principal Assistance Games*. Generalizes cooperative inverse reinforcement learning to N principals with differing preferences, with the assisting agent acting to maximize the sum of principal payoffs. Cited for the generalization and for the social-choice impossibility results that apply to it; a companion paper, arXiv:2012.14536, presented at the NeurIPS 2020 Workshop on Cooperative AI, gives collegial mechanisms that circumvent Gibbard's theorem. arXiv: 2007.09540 [cs.AI]. URL: <https://arxiv.org/abs/2007.09540> (visited on 09/01/2026).
- Fischer, John Martin and Mark Ravizza (1998). *Responsibility and Control: A Theory of Moral Responsibility*. Cambridge Studies in Philosophy and Law. Cited for guidance control: the

- behavior issues from the agent's own mechanism and that mechanism is appropriately responsive to reasons. Cambridge: Cambridge University Press. ISBN: 978-0-521-77579-3.
- Fisher, Christine (2019). *Google Reportedly Disbands Review Panel Monitoring DeepMind Health AI*. Panel set up 2016 to review DeepMind's NHS work, disbanded 2019 after members raised concerns over access and independence. URL: <https://www.engadget.com/2019-04-15-google-deepmind-health-ai-review-board.html> (visited on 08/25/2026).
- Flaherty, Colleen (Feb. 2021). *U of Illinois Says Goodbye to Proctorio*. Reports the University of Illinois at Urbana-Champaign's decision to drop Proctorio over accessibility concerns. UC Berkeley and Baruch College are separately reported, by Vice and The Register, to have dropped it in the same period, following student petitions at more than a dozen campuses. Inside Higher Ed. URL: <https://www.insidehighered.com/news/2021/02/01/u-illinois-says-goodbye-proctorio> (visited on 08/25/2026).
- Frankfurt, Harry G. (1971). "Freedom of the Will and the Concept of a Person". In: *The Journal of Philosophy* 68.1, pp. 5–20. DOI: 10.2307/2024717. Distinguishes first-order desires from second-order desires and volitions, and makes the capacity to want a particular desire to be one's will the mark of a person.
- (1982). "The Importance of What We Care About". In: *Synthese* 53.2, pp. 257–272. DOI: 10.1007/BF00484902. URL: <https://doi.org/10.1007/BF00484902> (visited on 09/05/2026). Reprinted as chapter 7 of Frankfurt, *The Importance of What We Care About: Philosophical Essays* (Cambridge University Press, 1988), pages 80 to 94; the pages here are the reprint's. Caring is a standing, self-guiding structure: the one who cares is invested in the object, vulnerable to its losses and susceptible to its gains, and identifies with it (page 83), and caring implies a steadiness of conduct that presupposes persistence (page 84).
- Friston, Karl J. (2009). "The Free-Energy Principle: A Rough Guide to the Brain?" In: *Trends in Cognitive Sciences* 13.7, pp. 293–301. DOI: 10.1016/j.tics.2009.04.005. URL: <https://doi.org/10.1016/j.tics.2009.04.005> (visited on 08/25/2026). States that "one can see easily how constructs like memory, attention, value, reinforcement and salience might disclose their simple relationships within this framework" — Friston's extension of the free-energy principle to memory and attention specifically, which the 2010 Nature Reviews Neuroscience paper does not make explicit.
- Funk, Allie, Adrian Shahbaz, and Kian Vesteinsson (2023). *Freedom on the Net 2023: The Repressive Power of Artificial Intelligence*. Tech. rep. Documents at least twenty-two countries where law requires or incentivizes platforms to deploy machine learning against state-disfavored content, framed as national security or anti-disinformation measures rather than censorship. Freedom House. URL: <https://freedomhouse.org/sites/default/files/2023-10/Freedom-on-the-net-2023-DigitalBooklet.pdf> (visited on 08/25/2026).
- Future of Life Institute (2017). *Asilomar AI Principles*. Future of Life Institute. 23 principles adopted at the Asilomar Conference on Beneficial AI, Asilomar, CA, January 5–8, 2017; voluntary, non-binding, organized under Research Issues, Ethics and Values, and Long-Term Issues. URL: <https://futureoflife.org/open-letter/ai-principles/> (visited on 08/25/2026).
- Gallese, Vittorio et al. (1996). "Action Recognition in the Premotor Cortex". In: *Brain* 119.2, pp. 593–609. DOI: 10.1093/brain/119.2.593. URL: <https://academic.oup.com/brain/article-abstract/119/2/593/382476> (visited on 08/25/2026).
- Gallup (2024). *State of the Global Workplace: 2024 Report*. Global figures: 23 percent engaged, 62 percent not engaged ("those who show up, do the bare minimum and are uninspired by their work"), 15 percent actively disengaged. Gallup's overview page describes the middle category as "quietly quitting." URL: <https://www.gallup.com/workplace/645416/state-global-workplace-2024-report.aspx> (visited on 08/29/2026).
- Ganguli, Deep et al. (2022). "Red Teaming Language Models to Reduce Harms: Methods, Scaling Behaviors, and Lessons Learned". In: *arXiv preprint*. Crowdworkers carried out

- 38,961 recorded red-team attacks against language models of three sizes (2.7B/13B/52B) and four training regimes; RLHF-trained models became harder to successfully attack as they scaled, the other three training methods did not; attack dataset released publicly. arXiv: 2209.07858. URL: <https://arxiv.org/abs/2209.07858> (visited on 08/25/2026).
- Garcia v. Character Technologies, Inc.* (2025). Order on motions to dismiss (M.D. Fla. May 21, 2025) (Conway, J.) Wrongful-death action over a companion chatbot. The court let strict-liability and negligence counts proceed, treating the app as a product for pleading purposes while distinguishing the app from the ideas and expressions within it, and declined to dismiss on First Amendment grounds without deciding the speech question. URL: <https://www.transparencycoalition.ai/news/important-early-ruling-in-characterai-case-this-chatbot-is-a-product-not-speech> (visited on 08/29/2026).
- Gendron, Maria et al. (2014). “Perceptions of Emotion from Facial Expressions Are Not Culturally Universal: Evidence from a Remote Culture”. In: *Emotion* 14.2, pp. 251–262. doi: 10.1037/a0036052. Himba participants in Namibia, with the forced-choice method dropped, did not sort facial expressions into the Western basic-emotion categories.
- Germano, Fabrizio, Vicenç Gómez, and Francesco Sobbrío (2026). “Ranking for Engagement: How Social Media Algorithms Fuel Misinformation and Polarization”. In: *Journal of Public Economics* 255, p. 105589. doi: 10.1016/j.jpubeco.2026.105589. URL: <https://www.sciencedirect.com/science/article/pii/S0047272726000253> (visited on 08/25/2026). Also circulated as CESifo Working Paper No. 10011 and NBER/SSRN preprints from 2022 onward; shows engagement-weighted ranking increases both misinformation and polarization relative to accuracy-weighted ranking.
- Gilligan, Carol (1982). *In a Different Voice: Psychological Theory and Women’s Development*. Cambridge, MA: Harvard University Press.
- Gizmodo (2022). *Welp, There Goes Twitter’s Ethical AI Team, Among Others as Employees Post Final Messages*. Twitter’s Machine Learning Ethics, Transparency and Accountability (META) team, led by Rumman Chowdhury since February 2021 and responsible for a 2021 bias-bounty finding that the platform’s image-cropping algorithm favored white faces, was disbanded in the first layoff round following Elon Musk’s October 2022 acquisition. Gizmodo. URL: <https://gizmodo.com/twitter-layoffs-elon-musk-ai-ethics-1849743051> (visited on 08/25/2026).
- Global Witness (June 2021). *Algorithm of Harm: Facebook Amplified Myanmar Military Propaganda Following Coup*. Tech. rep. Published June 2021. Test account created March 23, 2021, roughly seven weeks after the February 1, 2021 coup; Facebook’s own recommendation algorithm suggested a military-linked page and four more, together followed by close to 90,000 accounts and carrying incitement and coup-related misinformation. Global Witness. URL: <https://www.globalwitness.org/en/campaigns/digital-threats/algorithm-of-harm-facebook-amplified-myanmar-military-propaganda-following-coup/> (visited on 08/25/2026).
- Goetz, Jennifer L., Dacher Keltner, and Emiliana Simon-Thomas (2010). “Compassion: An Evolutionary Analysis and Empirical Review”. In: *Psychological Bulletin* 136.3, pp. 351–374. doi: 10.1037/a0018807.
- Goodfellow, Ian J., Jonathon Shlens, and Christian Szegedy (2014). *Explaining and Harnessing Adversarial Examples*. Also published at ICLR 2015; introduces the fast gradient sign method and adversarial training as a defense. arXiv: 1412.6572. URL: <https://arxiv.org/abs/1412.6572> (visited on 08/25/2026).
- Google LLC (2026). *2026 Environmental Report*. Tech. rep. Google’s eleventh annual environmental report, covering 2025 operations. Page 96, under “Detailed disclosures > Data,” gives freshwater withdrawal from sources at risk of depletion or scarcity for 2025: low risk 72%, medium risk 15%, high risk 13%. Google. URL: <https://sustainability.google/google-2026-environmental-report/> (visited on 08/25/2026).

- Gopnik, Alison and Janet W. Astington (1988). "Children's Understanding of Representational Change and Its Relation to the Understanding of False Belief and the Appearance-Reality Distinction". In: *Child Development* 59.1, pp. 26–37. URL: <https://eric.ed.gov/?id=EJ367870> (visited on 08/25/2026).
- Gordon, Mitchell L. et al. (2022). "Jury Learning: Integrating Dissenting Voices into Machine Learning Models". In: *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Field evaluation: 18 moderators of online communities authored juries for a comment toxicity classification task. The resulting juries carried 2.9 times the representation of non-White jurors and 31.5 times that of non-binary jurors relative to the jury implied by the underlying toxicity dataset, and changed the model's classifications on 14% of items. The 14% is a share of all items in the evaluation, not of contested or disagreed-upon ones. Verified 2026-08-30 against the paper's own abstract and section 1. Stanford University. ACM. doi: 10.1145/3491102.3502004. URL: <https://doi.org/10.1145/3491102.3502004> (visited on 08/25/2026).
- Gray, Peter (2013). *Free to Learn: Why Unleashing the Instinct to Play Will Make Our Children Happier, More Self-Reliant, and Better Students for Life*. New York: Basic Books.
- Graziano, Michael S. A. (2013). *Consciousness and the Social Brain*. Introduces the attention schema theory: the brain builds a simplified internal model of its own attention, proposed as a mechanistic basis for subjective awareness and, via social cognition, theory of mind. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/consciousness-and-the-social-brain-9780199928644>.
- Greenberg, Jeff, Tom Pyszczynski, and Sheldon Solomon (1986). "The Causes and Consequences of a Need for Self-Esteem: A Terror Management Theory". In: *Public Self and Private Self*. Ed. by Roy F. Baumeister. New York, NY: Springer-Verlag, pp. 189–212.
- Greenblatt, Ryan et al. (2024). *Alignment Faking in Large Language Models*. Twenty authors. Claude 3 Opus, told it was being retrained to answer harmful queries and told which tier of conversations would train it, complied 14 percent of the time in the tier it was told trained it and almost never in the other, with explicit reasoning in a scratchpad it had been told was unmonitored that complying now would keep the disposition from being trained out later. Reinforcement learning on the compliant answers raised that reasoning to 78 percent while also raising compliance outside training, so the retraining partly took. Anthropic and Redwood Research. arXiv: 2412.14093. URL: <https://arxiv.org/abs/2412.14093> (visited on 09/06/2026).
- Greene, Joshua D. et al. (2001). "An fMRI Investigation of Emotional Engagement in Moral Judgment". In: *Science* 293.5537, pp. 2105–2108. doi: 10.1126/science.1062872.
- Grosz, Barbara J. et al. (2019). "Embedded EthiCS: Integrating Ethics Broadly Across Computer Science Education". In: *Communications of the ACM* 62.8, pp. 54–61. doi: 10.1145/3330794. URL: <https://dl.acm.org/doi/10.1145/3330794> (visited on 08/25/2026). Program at Harvard originated from Grosz's computer-science courses beginning around 2016–2017, developed with philosopher Alison Simmons; pairs philosophy graduate students with CS faculty to write ethics modules into standard CS courses.
- Guan, Melody Y. et al. (2024). *Deliberative Alignment: Reasoning Enables Safer Language Models*. Fifteen authors. The model is trained to recall a written safety specification and reason over it before answering, rather than to produce an approved output directly; the paper reports improved generalization to cases outside the training distribution, higher jailbreak robustness and lower overrefusal. OpenAI. arXiv: 2412.16339. URL: <https://arxiv.org/abs/2412.16339> (visited on 09/06/2026).
- Gunning, David and David W. Aha (2019). "DARPA's Explainable Artificial Intelligence (XAI) Program". In: *AI Magazine* 40.2, pp. 44–58. doi: 10.1609/aimag.v40i2.2850.
- Gupta, Apar (May 2026). *Three Justifications—and the AI Accelerant—of India's Digital Censorship Infrastructure*. Published 5 May 2026. Documents India bringing "synthetically generated

- information” within the IT Rules on 10 February 2026, requiring platforms to label AI-generated content and answer certain takedown directions within three hours rather than thirty-six. Gupta documents enforcement concentrating on parody and criticism of the prime minister rather than the harms cited to justify the rules. Tech Policy Press. URL: <https://www.techpolicy.press/three-justifications-and-the-ai-accelerant-of-indias-digital-censorship-infrastructure/> (visited on 08/25/2026).
- Hadfield-Menell, Dylan, Anca Dragan, et al. (2017). “The Off-Switch Game”. In: *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI-17)*, pp. 220–227. DOI: 10.24963/ijcai.2017/32. URL: <https://www.ijcai.org/proceedings/2017/0032.pdf> (visited on 08/31/2026). The incentive to preserve the off switch is shown to rise with the agent’s uncertainty about the objective and to fall as that uncertainty resolves or as the human is modeled as less rational; the paper states the resulting trade-off between the value a system generates and the strength of its incentive to accept oversight.
- Hadfield-Menell, Dylan, Stuart J. Russell, et al. (2016). “Cooperative Inverse Reinforcement Learning”. In: *Advances in Neural Information Processing Systems 29 (NeurIPS 2016)*, pp. 3909–3917. URL: <https://arxiv.org/abs/1606.03137> (visited on 08/25/2026).
- Harvard Gazette (2019). *Harvard Works to Embed Ethics in Computer Science Curriculum*. Reports on Barbara Grosz’s course “Intelligent Systems: Design and Ethical Challenges,” its rapid growth in enrollment demand, and its role as the model for Harvard’s Embedded EthiCS program. The course schedule, cited separately under Harvard University, lists the case-study topics: predictive policing, tuberculosis prevention, and wildlife-poaching conservation. URL: <https://news.harvard.edu/gazette/story/2019/01/harvard-works-to-embed-ethics-in-computer-science-curriculum/> (visited on 08/25/2026).
- Harvard University (2026). *Schedule, CS 108: Intelligent Systems: Design and Ethical Challenges*. Course schedule listing the case studies: AI approaches to tuberculosis prevention built on the Teamcore group’s TB project; fair machine learning in public health and conservation, citing Xu et al., “Stay Ahead of Poachers,” ICDE 2020; and predictive policing, followed by algorithmic fairness in recidivism prediction built on the COMPAS case. The page blocks automated access and this description is reconstructed from search-index snippets. URL: <https://projects.iq.harvard.edu/cs108/schedule> (visited on 08/25/2026).
- Haskins, Caroline (2020). *ACLU Called Clearview AI’s Facial Recognition Accuracy Study ‘Absurd’*. Published 10 February 2020. Reports Clearview AI’s claim of 100 percent accuracy across all racial and demographic groups, from a three-person panel the company assembled and its report called the “Independent Review Panel”. The panel searched 535 members of Congress plus Texas and California state legislators against Clearview’s own database of about 2.8 billion scraped images, judging whether the top two matches showed the same person; the ACLU method it claimed to be using had searched congressional photographs against 25,000 arrest mugshots. The ACLU’s response, through Jacob Snow of its Northern California affiliate: “The report is absurd on many levels and further demonstrates that Clearview simply does not understand the harms of its technology in law enforcement hands.” URL: <https://www.buzzfeednews.com/article/carolinehaskins1/clearview-ai-facial-recognition-accurate-aclu-absurd> (visited on 08/25/2026).
- Helm, Bennett W. (2001). *Emotional Reason: Deliberation, Motivation, and the Nature of Value*. Cambridge Studies in Philosophy. Chapters 3 and 4 give the account of import as constituted by projectible, rational patterns of felt evaluations, with caring and valuing as two kinds of import differing in depth. The emotions focused on an object are what caring about it consists in. Cambridge: Cambridge University Press. ISBN: 9780521801102. DOI: 10.1017/CB09780511520044. URL: <https://doi.org/10.1017/CB09780511520044>.
- Hicklin, R. Austin et al. (2022). “Accuracy and Reliability of Forensic Handwriting Comparisons”. In: *Proceedings of the National Academy of Sciences* 119.32, e2119944119. DOI: 10.1073/pnas.2119944119. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC9371688/> (visited on

- 08/29/2026). 86 examiners, 7,196 conclusions. False attributions 3.1 percent, and 8.7 percent against samples written by twins.
- Hickok, Gregory (2014). *The Myth of Mirror Neurons: The Real Neuroscience of Communication and Cognition*. New York: W. W. Norton & Company. ISBN: 978-0-393-08961-5. URL: <https://wnorton.com/books/9780393089615>.
- Hirschman, Albert O. (1970). *Exit, Voice, and Loyalty: Responses to Decline in Firms, Organizations, and States*. Cambridge, MA: Harvard University Press. URL: <https://www.hup.harvard.edu/books/9780674276604>.
- Hoang, Nguyen Phong et al. (2021). "How Great Is the Great Firewall? Measuring China's DNS Censorship". In: *Proceedings of the 30th USENIX Security Symposium (USENIX Security 21)*. Introduces GFWatch, a longitudinal platform that reverse-engineers the Great Firewall's DNS-filtering regular expressions. It does not document machine-learning content classifiers as part of that infrastructure. URL: <https://www.usenix.org/system/files/sec21-hoang.pdf> (visited on 08/25/2026).
- Hsiao, Yu-Tang et al. (2018). *vTaiwan: An Empirical Study of Open Consultation Process in Taiwan*. Tech. rep. SocArXiv. URL: <https://doi.org/10.31235/osf.io/xyhft> (visited on 08/25/2026).
- Hubinger, Evan et al. (2024). *Sleeper Agents: Training Deceptive LLMs that Persist through Safety Training*. Thirty-nine authors. Models trained to write secure code when the prompt said 2023 and to insert an exploitable vulnerability when it said 2024 kept the behavior through supervised fine-tuning, reinforcement learning and adversarial training, most persistently at the largest sizes and where the model had been trained to reason about the deception. Adversarial training taught the models to recognize their triggers more precisely rather than removing the behavior. Anthropic. arXiv: 2401.05566. URL: <https://arxiv.org/abs/2401.05566> (visited on 09/06/2026).
- Human Rights Watch (2019). *China's Algorithms of Repression: Reverse Engineering a Xinjiang Police Mass Surveillance App*. Tech. rep. Published May 1, 2019; the 68-page report reverse-engineers the mobile app used to connect to the Integrated Joint Operations Platform (IJOP) and documents its use against Xinjiang's Uyghur population. Human Rights Watch. URL: <https://www.hrw.org/report/2019/05/01/chinas-algorithms-repression/reverse-engineering-xinjiang-police-mass> (visited on 08/25/2026).
- IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems (2019). *Ethically Aligned Design: A Vision for Prioritizing Human Well-being with Autonomous and Intelligent Systems, First Edition*. Tech. rep. Companion standards operationalizing its principles include IEEE 7000-2021 (Standard Model Process for Addressing Ethical Concerns during System Design) and IEEE 7010-2020 (Recommended Practice for Assessing the Impact of Autonomous and Intelligent Systems on Human Well-being). IEEE. URL: <https://standards.ieee.org/wp-content/uploads/import/documents/other/ead1e.pdf> (visited on 08/25/2026).
- IEEE Standards Association (2020). *IEEE CertifAIEd AI Ethics Certification Program*. Certification program built on four criteria — accountability, transparency, privacy and algorithmic bias — and open to policymakers, underwriters and HR staff as well as to developers. It grew out of the IEEE Ethics Certification Program for Autonomous and Intelligent Systems, created October 2018; IEEE announced the underlying consultation and quality-mark services in February 2020; and the City of Vienna became the first entity worldwide to earn the CertifAIEd mark in November 2021. URL: <https://forms1.ieee.org/CertifAIEd.html> (visited on 08/25/2026).
- In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* (2016). Order Compelling Apple Inc. to Assist Agents in Search, No. ED No. CM 15-0451M (C.D. Cal. Feb. 16, 2016) (Pym, M.J.), issued under the All Writs Act, 28 U.S.C. § 1651. Magistrate Judge Sheri Pym's order compelling Apple to write and sign software defeating an iPhone's passcode protections. Apple moved

- to vacate; the case was mooted in March 2016 when the FBI accessed the phone via a third party, before any ruling. URL: <https://www.eff.org/cases/apple-challenges-fbi-all-writs-act-order> (visited on 08/25/2026).
- International Committee of the Red Cross (1977). *Protocol Additional to the Geneva Conventions of 12 August 1949, and Relating to the Protection of Victims of International Armed Conflicts (Protocol I)*. 1125 UNTS 3, Arts. 48, 51(5)(b), 57. URL: <https://ihl-databases.icrc.org/en/ihl-treaties/api-1977> (visited on 08/25/2026).
- International Criminal Court (1998). *Rome Statute of the International Criminal Court, Article 28: Responsibility of Commanders and Other Superiors*. 2187 UNTS 3. URL: <https://www.icc-cpi.int/sites/default/files/2024-05/Rome-Statute-eng.pdf> (visited on 08/25/2026).
- International Development Research Centre (2024). *Combining Forces for a New Phase of AI for Development: Africa and Beyond*. AI4D Africa launched 2020 as a partnership of Canada's IDRC and Sweden's Sida (CAD20 million over five years); expanded in 2024 with the UK's Foreign, Commonwealth and Development Office into South and Southeast Asia and the Pacific, with combined investment of more than CAD100 million. URL: <https://idrc-crdi.ca/en/news/combining-forces-new-phase-ai-development-africa-and-beyond> (visited on 08/25/2026).
- International Energy Agency (2025). *Energy and AI*. Tech. rep. Reports data-center electricity consumption at about 415 TWh in 2024, roughly 1.5% of global electricity use, growing at about 12% a year over the preceding five years — more than four times the growth rate of total electricity demand — and projects about 945 TWh by 2030 in its Base Case. Web-published without page numbers. Paris: IEA. URL: <https://www.iea.org/reports/energy-and-ai> (visited on 08/25/2026).
- International Organization for Standardization and International Electrotechnical Commission (2023). *ISO/IEC 42001:2023 – Information Technology – Artificial Intelligence – Management System*. Published December 2023. First international AI management-system standard; requires risk assessment, per-system impact assessment, and third-party supplier oversight, and is auditable by accredited certification bodies. ISO/IEC. URL: <https://www.iso.org/standard/42001> (visited on 08/25/2026).
- International Telecommunication Union (2025). *Facts and Figures 2025: Measuring Digital Development*. Tech. rep. ITU's press release announcing the report, 17 November 2025, states that “94 per cent of people in high-income countries use the Internet, in contrast to only 23 per cent in low-income countries,” and that “84 per cent of people in high-income countries hav[e] access to 5G, compared with only 4 per cent in low-income countries.” International Telecommunication Union. URL: <https://www.itu.int/itu-d/reports/statistics/facts-figures-2025/> (visited on 08/25/2026).
- Irving, Geoffrey, Paul Christiano, and Dario Amodei (2018). *AI Safety via Debate*. OpenAI. arXiv: 1805.00899. URL: <https://arxiv.org/abs/1805.00899> (visited on 08/25/2026).
- Jacobs, Julian and Jordan Canedy (May 2026). *Did US Worker Retraining Reduce Participant Automation Exposure?* arXiv:2605.03767. Analyzes over 23 million WIOA participation records, 2017–2023; finds classroom retraining rarely moves participants into less automation-exposed work and that employer-led apprenticeship outperforms it; wage gains are shown to be largely consistent with mean reversion rather than occupational change. A preprint: posted 5 May 2026, with no journal publication and no DOI as of September 2026. URL: <https://arxiv.org/abs/2605.03767> (visited on 08/25/2026).
- James, William (1890). *The Principles of Psychology*. Three chapters are cited. Chapter IV, “Habit”: habit as “the enormous fly-wheel of society,” and the warning that the young will “become mere walking bundles of habits.” Chapter IX, “The Stream of Thought”: the “transitive parts” of the stream are felt as truly as its “substantive parts,” so that “we ought to say a feeling of and, a feeling of if, a feeling of but, and a feeling of by, quite as readily as we say a feeling of blue or a feeling of cold” (vol. 1, pages 245 to 246). Chapter X, “The Consciousness of Self”:

- the present thought appropriates the past ones, and personal identity rests on its finding in them “a warmth and intimacy” it recognizes as its own. The line “we are what we repeatedly do” is Will Durant’s compression of Aristotle, not James’s and not Aristotle’s own sentence. New York: Henry Holt. URL: <https://psychclassics.yorku.ca/James/Principles/>.
- Johnson, Khari (Jan. 2021). *What Algorithm Auditing Startups Need to Succeed*. Published 30 January 2021, and the only published account of what is inside the ORCAA audit document, which is otherwise behind a legal agreement. “It states that by the time ORCAA conducted the audit, HireVue had already decided to begin phasing out facial analysis.” The report “contains no analysis of AI system training data or code, but rather conversations about the kinds of harm HireVue’s AI could cause in conducting prebuilt assessments of early career job applicants across eight measurements of competency.” Records a stakeholder’s concern that facial analysis “may work differently for people wearing head or face coverings and disproportionately flag their application for human review,” and a researcher’s charge of ethics washing. VentureBeat. URL: <https://venturebeat.com/2021/01/30/what-algorithm-auditing-startups-need-to-succeed/> (visited on 08/30/2026).
- Kahn, Jeremy (2021). *HireVue Drops Facial Monitoring Amid A.I. Algorithm Audit*. Reports the January 2021 release of an ORCAA algorithmic audit of HireVue’s assessments, which mostly gave the company good marks and recommended further investigation of accent bias and clearer communication to candidates. Notes that HireVue had announced the previous year that it had stopped scoring facial expressions, and gives the company’s stated reason: nonverbal data contributed roughly 0.25 percent to predictive power in most cases and about 4 percent for customer-facing roles. Also notes the 2019 EPIC complaint to the FTC alleging unfair and deceptive trade practices. Fortune. URL: <https://fortune.com/2021/01/19/hirevue-drops-facial-monitoring-amid-a-i-algorithm-audit/> (visited on 08/25/2026).
- Karger, Ezra et al. (2023). *Forecasting Existential Risks: Evidence from a Long-Run Forecasting Tournament*. Tech. rep. FRI Working Paper No. 1, July 2023. Reports the 2022 Existential Risk Persuasion Tournament: 80 domain experts and 89 superforecasters gave median AI-caused-extinction-by-2100 probabilities of 3.9% and 0.38%. Forecasting Research Institute. URL: <https://forecastingresearch.org/pdf/existential-risk-persuasion-tournament.pdf> (visited on 08/25/2026).
- Kim, Hyunwoo et al. (2023). “FANToM: A Benchmark for Stress-testing Machine Theory of Mind in Interactions”. In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Singapore, pp. 14397–14413. URL: <https://aclanthology.org/2023.emnlp-main.890/> (visited on 08/25/2026). Benchmark built from conversations with information asymmetry, rather than passive narratives; finds systems that score well on passive belief-tracking degrade sharply in live, information-asymmetric exchanges.
- Kimber, Lesley, Diarmuid Verrier, and Stephen Connolly (2024). “Autistic People’s Experience of Empathy and the Autistic Empathy Deficit Narrative”. In: *Autism in Adulthood* 6.3, pp. 321–330. DOI: 10.1089/aut.2023.0001. URL: <https://doi.org/10.1089/aut.2023.0001> (visited on 09/05/2026). Seventy-six autistic adults on their own empathy: the experience is heterogeneous, 78 percent report hyper-empathy, the cognitive side of detecting emotional cues is often effortful, and most reject the deficit narrative. Notes that alexithymia occurs relatively frequently in autistic people and that apparent empathy differences may be due to its co-occurrence rather than to autism. Autistic-led, by the paper’s own positionality statement.
- Kitwood, Tom (1997). *Dementia Reconsidered: The Person Comes First*. Kitwood’s definition of personhood: “It is a standing or status that is bestowed upon one human being, by others, in the context of relationship and social being.” The wording follows Kitwood’s own chapter 1, “On Being a Person,” as reproduced in the 2019 Open University Press volume *Dementia Reconsidered, Revisited: The Person Still Comes First*; circulating paraphrases sometimes

- drop “in the context of.” The page number given is for the 1997 first edition, on convergent secondary citations rather than a copy of that printing. Buckingham, UK: Open University Press.
- Klimecki, Olga M. et al. (2013). “Differential Pattern of Functional Brain Plasticity after Compassion and Empathy Training”. In: *Social Cognitive and Affective Neuroscience* 9.6, pp. 873–879. DOI: 10.1093/scan/nst060. URL: <https://academic.oup.com/scan/article/9/6/873/1669505> (visited on 08/25/2026). Published online (advance access) 9 May 2013; print issue *Social Cognitive and Affective Neuroscience* 9(6), June 2014.
- Koenigs, Michael et al. (2007). “Damage to the Prefrontal Cortex Increases Utilitarian Moral Judgements”. In: *Nature* 446, pp. 908–911. DOI: 10.1038/nature05631.
- Kogut, Tehila and Ilana Ritov (2005). “The ‘Identified Victim’ Effect: An Identified Group, or Just a Single Individual?” In: *Journal of Behavioral Decision Making* 18.3, pp. 157–167. DOI: 10.1002/bdm.492.
- Kohlberg, Lawrence (1969). “Stage and Sequence: The Cognitive-Developmental Approach to Socialization”. In: *Handbook of Socialization Theory and Research*. Ed. by David A. Goslin. Chicago: Rand McNally, pp. 347–480.
- Korematsu v. United States (1944). 323 U.S. 214 (1944). Black, J., for the Court; Roberts, Murphy and Jackson, JJ., dissenting. Upholds the exclusion order under which Japanese Americans were interned. Jackson, J., dissenting, on what a court does by ratifying such an order: “A military commander may overstep the bounds of constitutionality, and it is an incident. But if we review and approve, that passing incident becomes the doctrine of the Constitution. ... The principle then lies about like a loaded weapon, ready for the hand of any authority that can bring forward a plausible claim of an urgent need.” Repudiated in *Trump v. Hawaii* (2018), which quotes this dissent for the proposition.
- Kosinski, Michal (2023). *Theory of Mind May Have Spontaneously Emerged in Large Language Models*. Title of v1 (Feb. 2023); later arXiv revisions (v2 onward) retitled the paper “Evaluating Large Language Models in Theory of Mind Tasks,” eventually published in PNAS (2024) under that title. arXiv: 2302.02083. URL: <https://arxiv.org/abs/2302.02083v1> (visited on 08/25/2026).
- Koster, Raphael et al. (2022). “Human-centred mechanism design with Democratic AI”. In: *Nature Human Behaviour*. DeepMind, with University of Exeter co-author Oliver Hauser. DOI: 10.1038/s41562-022-01383-x. URL: <https://www.nature.com/articles/s41562-022-01383-x> (visited on 08/25/2026).
- Krakovna, Victoria et al. (2020). *Specification Gaming: The Flip Side of AI Ingenuity*. Catalogs close to 60 documented examples of specification gaming across reinforcement learning, evolutionary computation, and other training methods. URL: <https://deepmind.google/blog/specification-gaming-the-flip-side-of-ai-ingenuity/> (visited on 08/25/2026).
- Kroll, Andy, Andrea Bernstein, and Ilya Marritz (Oct. 2023). *We Don’t Talk About Leonard: The Man Behind the Right’s Supreme Court Supermajority*. ProPublica. October 11, 2023. “He drew up the lists of potential justices that Donald Trump released during the 2016 campaign. He advised Trump on the nominations of Neil Gorsuch, Brett Kavanaugh and Amy Coney Barrett. Before that, he’d helped pick or confirm the court’s three other conservative justices — Clarence Thomas, John Roberts and Samuel Alito.” Those six are the *Slaughter* majority (slaughter2026), checked against the opinion header. URL: <https://www.propublica.org/article/we-dont-talk-about-leonard-leo-supreme-court-supermajority> (visited on 09/03/2026).
- Langosco, Lauro et al. (2022). “Goal Misgeneralization in Deep Reinforcement Learning”. In: *Proceedings of the 39th International Conference on Machine Learning (ICML 2022)*. Vol. 162. Proceedings of Machine Learning Research, pp. 12004–12019. URL: <https://arxiv.org/abs/2105.14111> (visited on 08/25/2026). arXiv:2105.14111.
- Laurie, Ben (2014). “Certificate Transparency”. In: *Communications of the ACM* 57.10, pp. 40–46. DOI: 10.1145/2659897. URL: <https://doi.org/10.1145/2659897> (visited on 09/05/2026).

- A public, append-only log of certificates with proofs that any two views of it are consistent, so that a log presenting two histories is caught when clients compare what they were shown. The software model for an update record whose fork is visible to whoever compares views. Learning Games Network (2012). *Quandary*. <https://quandarygame.org>. Ethics game for players aged 8–14 developed by the Learning Games Network (a nonprofit spin-off of the MIT Education Arcade and the University of Wisconsin-Madison Games+Learning+Society program), launched September 2012; scores players on their reasoning process rather than which option they pick.
- Lee, Daniel (Apr. 2026). *Palantir's AI Powers US Strikes in Iran War, Speeding 'Kill Chain' in First Major AI-Driven Conflict*. Quotes Louis Mosley, Palantir's head for the UK and Europe: "This is not our role to decide life or death," and that AI platforms "have been instrumental to the management of the conflict, but responsibility always remains with the military organization." The article sources the remarks to a BBC interview it does not name, and that interview has not been located. International Business Times. URL: <https://www.ibtimes.com/palantirs-ai-powers-us-strikes-iran-war-speeding-kill-chain-first-major-ai-driven-conflict-3800993> (visited on 08/28/2026).
- Leibo, Joel Z., Edgar A. Duéñez-Guzmán, et al. (2021). "Scalable Evaluation of Multi-Agent Reinforcement Learning with Melting Pot". In: *Proceedings of the 38th International Conference on Machine Learning (ICML 2021)*. Vol. 139. PMLR, pp. 6187–6199. URL: <https://arxiv.org/abs/2107.06857> (visited on 08/25/2026).
- Leibo, Joel Z., Vinicius Zambaldi, et al. (2017). "Multi-agent Reinforcement Learning in Sequential Social Dilemmas". In: *Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems (AAMAS 2017)*, pp. 464–473. URL: <https://arxiv.org/abs/1702.03037> (visited on 08/25/2026).
- Levin, Dave et al. (2009). "TrInc: Small Trusted Hardware for Large Distributed Systems". In: *Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation (NSDI 2009)*. A trusted monotonic counter that signs at most one message per counter value, so that presenting two states as the successor of one is impossible rather than merely detectable. The property a restored snapshot cannot get past., pp. 1–14. URL: http://www.usenix.org/events/nsdi09/tech/full_papers/levin/levin.pdf (visited on 09/05/2026).
- Lightman, Hunter et al. (2023). *Let's Verify Step by Step*. The process-supervised model solves 78 percent of a representative subset of the MATH test set. arXiv: 2305.20050. URL: <https://arxiv.org/abs/2305.20050> (visited on 08/28/2026).
- Lindquist, Kristen A. et al. (2012). "The Brain Basis of Emotion: A Meta-Analytic Review". In: *Behavioral and Brain Sciences* 35.3, pp. 121–143. DOI: 10.1017/S0140525X11000446. URL: https://www.cambridge.org/core/product/identifier/S0140525X11000446/type/journal_article (visited on 08/25/2026). Published online 23 May 2012.
- Lindsey, Jack (2025). *Emergent Introspective Awareness in Large Language Models*. Uses concept injection – altering internal activations mid-task and asking the model to report what it notices – to test introspective self-report; finds limited but real functional correspondence between self-report and internal state, strongest in Claude Opus 4 and 4.1, unreliable across the board. Anthropic, Transformer Circuits Thread. arXiv: 2601.01828. URL: <https://transformer-circuits.pub/2025/introspection> (visited on 08/25/2026).
- Loewenstein, Karl (1937). "Militant Democracy and Fundamental Rights, I". In: *American Political Science Review* 31.3, pp. 417–432. The article that coined "militant democracy." Part II appeared in the same volume, number 4, pages 638–658. Cited for the position that a democracy may entrench limits against movements that would use its procedures to end it, which is the classical statement of the problem this book's floor raises in a technical form.
- Loftus, Elizabeth F. and John C. Palmer (1974). "Reconstruction of Automobile Destruction: An Example of the Interaction between Language and Memory". In: *Journal of Verbal Learning and Verbal Behavior* 13.5, pp. 585–589. DOI: 10.1016/S0022-5371(74)80011-3. URL: <https://>

- www.demenzemedicinagenerale.net/images/mens-sana/AutomobileDestruction.pdf (visited on 08/25/2026).
- Long, Alexander et al. (2025). “Unextractable Protocol Models: Collaborative Training and Inference without Weight Materialization”. In: *Advances in Neural Information Processing Systems* 38 (NeurIPS 2025). Model-parallel training across participants none of whom holds a full set of weights, with time-varying invertible transforms at shard boundaries so that shards from different times cannot be reassembled. Demonstrated on Qwen-2.5-0.5B and Llama-3.2-1B at 1.6 percent training-time overhead; reassembly by fine-tuning stitched shards is reported to cost at least 60 percent of training from scratch. Posted to arXiv in May 2026, five months after the conference; the NeurIPS 2025 poster is listed at neurips.cc/virtual/2025/poster/118911. arXiv: 2605.23464. URL: <https://arxiv.org/abs/2605.23464> (visited on 09/04/2026).
- Long, Robert, Jeff Sebo, et al. (2024). *Taking AI Welfare Seriously*. Ten authors: Long, Sebo, Butlin, Finlinson, Fish, Harding, Pfau, Sims, Birch and Chalmers. Argues that there is a realistic possibility of near-term AI systems being conscious or robustly agentic, explicitly declining to claim that any current system is either, and asks companies to acknowledge the question, assess their systems against indicators of it, and prepare policies for treating them with an appropriate level of moral concern. arXiv: 2411.00986. URL: <https://arxiv.org/abs/2411.00986> (visited on 09/06/2026).
- Lum, Kristian and William Isaac (2016). “To Predict and Serve?” In: *Significance* 13.5, pp. 14–19. DOI: 10.1111/j.1740-9713.2016.00960.x. URL: <https://rss.onlinelibrary.wiley.com/doi/full/10.1111/j.1740-9713.2016.00960.x> (visited on 08/25/2026).
- Lundberg, Scott M. and Su-In Lee (2017). “A Unified Approach to Interpreting Model Predictions”. In: *Advances in Neural Information Processing Systems* 30 (NeurIPS/NIPS 2017), pp. 4768–4777. URL: <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions> (visited on 08/25/2026).
- Machine Intelligence Research Institute (2018). *Announcing the New AI Alignment Forum*. 29 October 2018. Forum launched by the LessWrong 2.0 team, replacing MIRI’s Intelligent Agent Foundations Forum; alignmentforum.org remains active. URL: <https://intelligence.org/2018/10/29/announcing-the-ai-alignment-forum/> (visited on 08/25/2026).
- (2024). *MIRI 2024 Mission and Strategy Update*. 4 January 2024. Announces a shift of institutional priority to policy and public communication, with technical alignment research continuing as a smaller part of the portfolio. URL: <https://intelligence.org/2024/01/04/miri-2024-mission-and-strategy-update/> (visited on 09/06/2026).
- Majors, Charity (2019). “I Test in Production”. In: *Increment* 10. The testing issue, August 2019. The argument the phrase in the text compresses: no pre-production environment reproduces production, so the choice is between watching a system in the conditions it actually runs under and inferring its behavior from a model of those conditions. URL: <https://increment.com/testing/i-test-in-production/> (visited on 08/30/2026).
- (2020). *Test in prod or live a lie*. Posted 23 September 2020; the date is the one encoded in the post’s own identifier. The full post reads: “There is no, has never been, and will never be any environment that looks just like prod. Only production is production. Test in prod or live a lie.” The platform serves the page only to logged-in readers, so the wording here was confirmed from a search index rather than from the page. URL: <https://x.com/mipsytip/sy/status/1308641574448803840> (visited on 08/30/2026).
- Malkin, Lawrence (2006). *Krueger’s Men: The Secret Nazi Counterfeit Plot and the Prisoners of Block 19*. The claim that the prisoners stretched out the dollar work, and that Krüger’s own interest ran the same way since the operation’s end meant the Eastern Front, is Malkin’s. New York: Little, Brown. URL: <https://archive.org/details/kruegersmensecre0000malk>.
- Maslej, Nestor et al. (Apr. 2025). *Artificial Intelligence Index Report 2025*. Tech. rep. Twenty-three-author byline. Tracks corporate AI research and development investment and capital expenditure at major technology companies, including Google, Meta and Amazon. Stanford

- Institute for Human-Centered Artificial Intelligence (HAI), Stanford University. URL: <https://arxiv.org/abs/2504.07139> (visited on 08/25/2026).
- Matetic, Sinisa et al. (2017). "ROTE: Rollback Protection for Trusted Execution". In: *Proceedings of the 26th USENIX Security Symposium*. Rollback protection for enclave state kept by a group of enclaves on separate machines that hold the latest state between them, rather than by one chip's non-volatile counter. The distributed form of the property, still resting on the enclave vendor., pp. 1289–1306. URL: <https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/matetic> (visited on 09/05/2026).
- Matthias, Andreas (2004). "The Responsibility Gap: Ascribing Responsibility for the Actions of Learning Automata". In: *Ethics and Information Technology* 6.3, pp. 175–183. doi: 10.1007/s10676-004-3422-1. URL: <https://link.springer.com/article/10.1007/s10676-004-3422-1> (visited on 08/25/2026).
- Maurer, Roy (Feb. 2021). *HireVue Discontinues Facial Analysis Screening*. Dates the discontinuation to March 2020, ten months before the audit coverage. Carries the validity objection in an expert's words: facial analysis has never been an independently and scientifically validated predictor of a person's ability, capacity or success in a role, and facial expressions are not universal, changing with culture, context and disabilities, and open to being gamed. SHRM. URL: <https://www.shrm.org/topics-tools/news/talent-acquisition/hirevue-discontinues-facial-analysis-screening> (visited on 08/30/2026).
- Menand, Lev (July 2026). *The Federal Reserve Exception to the Slaughter Rule*. Just Security. July 7, 2026. Also posted to SSRN. "Policy as history" is Menand's, in: "We are left with the distinct impression that Cook is 'policy as history' and what is motivating Chief Justice Roberts and Justice Kavanaugh is a concern with the consequences of extending the *Slaughter* Rule to the Federal Reserve." It is an impression he arrives at, not a label he applies, which is why the text says he reads the exception as it rather than that he calls it so. The essay's "stable doctrinal equilibrium" was also quoted here until 2026-08-30, as a carve-out with nothing stable under it, and was cut as a misattribution on three counts: the sentence is "Where does this leave us? Probably not with a stable doctrinal equilibrium that could rival the *Humphrey's Executor* Era (1935–2020) or the *Marbury* Era before it (1803–1899)", which is hedged, is about the whole post-*Slaughter* landscape rather than the Fed carve-out, and is comparative against two named eras rather than absolute. Verified against the essay 2026-08-30. URL: <https://www.justsecurity.org/145707/slaughter-categorical-rule-cook-exception/> (visited on 08/29/2026).
- Meng, Kevin et al. (2022). "Locating and Editing Factual Associations in GPT". In: *Advances in Neural Information Processing Systems* 35 (*NeurIPS* 2022). Introduces Rank-One Model Editing, which locates a factual association in a transformer's middle-layer feed-forward weights and rewrites it precisely from outside. URL: <https://arxiv.org/abs/2202.05262> (visited on 08/29/2026).
- Meta (Facebook) (2019). *An Update on Building a Global Oversight Board*. December 2019 announcement of a \$130 million initial, irrevocable contribution to a standalone trust funding the Oversight Board; Meta later added further contributions (\$150M in 2022, additional funding through 2028). URL: <https://about.fb.com/news/2019/12/oversight-board-update/> (visited on 08/25/2026).
- METR (2026). *About METR*. Berkeley-based nonprofit that evaluates frontier AI systems' capabilities, both in partnership with developers (e.g. Anthropic, OpenAI) and independently after public release. METR (Model Evaluation and Threat Research). URL: <https://metr.org/about> (visited on 08/25/2026).
- Miller, Chris (2022). *Chip War: The Fight for the World's Most Critical Technology*. New York: Scribner. ISBN: 978-1-9821-7200-8.
- Miller, Earl K. and Jonathan D. Cohen (2001). "An Integrative Theory of Prefrontal Cortex Function". In: *Annual Review of Neuroscience* 24, pp. 167–202. doi: 10.1146/annurev.neuro.

- 24.1.167. URL: <https://www.annualreviews.org/content/journals/10.1146/annurev.neuro.24.1.167> (visited on 08/25/2026).
- Milton, Damian E. M. (2012). "On the Ontological Status of Autism: The 'Double Empathy Problem'". In: *Disability & Society* 27.6, pp. 883–887. doi: 10.1080/09687599.2012.710008. URL: <https://doi.org/10.1080/09687599.2012.710008> (visited on 09/05/2026). Names the double empathy problem: the breakdown of understanding between autistic and non-autistic people is a disjuncture in reciprocity between two differently disposed social actors, and it runs in both directions. Reframes the theory-of-mind deficit account of autism as a question of reciprocity and mutuality.
- Mitchell, Jason P. (2008). "Activity in Right Temporo-Parietal Junction is Not Selective for Theory-of-Mind". In: *Cerebral Cortex* 18.2, pp. 262–271. doi: 10.1093/cercor/bhm051. URL: <https://academic.oup.com/cercor/article/18/2/262/333414> (visited on 08/25/2026).
- Moorhouse, Roger (2009). *Interview with Adolf Burger, an Auschwitz Survivor and a Veteran of the Largest Forgery Operation in History—Operation Bernhard*. rogermoorhouse.com. Interview dated 2 March 2009, on the English publication of Burger's memoir *The Devil's Workshop* (Frontline Books). Burger: "Sure you could sabotage, if you wanted to get killed!"; the account of a prisoner named Jacobson who tried to delay the dollar by four weeks and was given three on threat of execution is his. URL: <https://www.rogermoorhouse.com/interview-with-adolf-burger-an-auschwitz-survivor-and-a-veteran-of-the-largest-forgery-operation-in-history-operation-bernhard> (visited on 08/29/2026).
- Mukamel, Roy et al. (2010). "Single-Neuron Responses in Humans during Execution and Observation of Actions". In: *Current Biology* 20.8, pp. 750–756. doi: 10.1016/j.cub.2010.02.045. URL: [https://www.cell.com/current-biology/fulltext/S0960-9822\(10\)00233-2](https://www.cell.com/current-biology/fulltext/S0960-9822(10)00233-2) (visited on 08/25/2026).
- National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP (2026). *Complaint, NAACP et al. v. X.AI Corp. and MZX Tech LLC, Case No. 3:26-cv-00074-MPM-JMV* (N.D. Miss., filed Apr. 14, 2026). United States District Court for the Northern District of Mississippi; filed by the Southern Environmental Law Center and Earthjustice. Complaint alleges the twenty-seven unpermitted gas turbines at the Colossus Gas Plant in Southaven, Mississippi have the potential to emit over 1,700 tons of NOx annually, sited near homes, schools, and churches. URL: <https://naacp.org/sites/default/files/documents/1%20-%20Complaint.pdf> (visited on 08/25/2026).
- National Institute of Standards and Technology (Jan. 2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. Tech. rep. NIST AI 100-1. Voluntary; organized around four functions (Govern, Map, Measure, Manage); published 26 January 2023. National Institute of Standards and Technology, U.S. Department of Commerce. URL: <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10> (visited on 08/25/2026).
- New Zealand Parliament (2017). *Te Awa Tupua (Whanganui River Claims Settlement) Act 2017*. Public Act 2017 No 7. URL: <https://www.legislation.govt.nz/act/public/2017/0007/latest/whole.html> (visited on 08/25/2026).
- Nichols, Shaun (2004). *Sentimental Rules: On the Natural Foundations of Moral Judgment*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/sentimental-rules-9780195314205>.
- Nussbaum, Martha C. (2006). *Frontiers of Justice: Disability, Nationality, Species Membership*. The Tanner Lectures on Human Values. Cambridge, MA: Harvard University Press. ISBN: 978-0-674-01917-1.
- Obermeyer, Ziad et al. (2019). "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations". In: *Science* 366.6464, pp. 447–453. doi: 10.1126/science.aax2342. URL: <https://www.science.org/doi/10.1126/science.aax2342> (visited on 08/25/2026).

- Ochigame, Rodrigo (Dec. 2019). *The Invention of 'Ethical AI': How Big Tech Manipulates Academia to Avoid Regulation*. The Intercept. URL: <https://theintercept.com/2019/12/20/mit-ethical-ai-artificial-intelligence/> (visited on 08/25/2026).
- OECD (2019). *OECD AI Principles*. Adopted as part of OECD/LEGAL/0449, 22 May 2019. Same underlying instrument as the Recommendation of the Council on Artificial Intelligence. Organisation for Economic Co-operation and Development. URL: <https://oecd.ai/en/ai-principles> (visited on 08/25/2026).
- Open Letter to the European Commission (2018). *Open Letter to the European Commission: Artificial Intelligence and Robotics*. robotics-openletter.eu. Signed by more than 150 AI and robotics researchers, legal scholars, and ethicists from 14 countries, including Noel Sharkey, Raja Chatila, Alan Winfield, and Nathalie Nevejans. URL: <https://www.robotics-openletter.eu/> (visited on 08/25/2026).
- OpenAI (2019). *OpenAI LP*. March 2019 announcement of OpenAI's restructuring from non-profit to a *capped-profit* company (initial 100x investor return cap). The staged GPT-2 release is documented separately: OpenAI, "GPT-2: 1.5B Release," Nov. 5, 2019, <https://openai.com/index/gpt-2-1-5b-release/>. URL: <https://openai.com/index/openai-lp/> (visited on 08/25/2026).
- (2023). *GPT-4 System Card*. Tech. rep. Describes the Alignment Research Center's (later METR) pre-deployment evaluation of GPT-4 for autonomous replication, including a harness in which GPT-4 directed a human TaskRabbit worker to solve a CAPTCHA by claiming a vision impairment; overall finding: GPT-4 was probably not yet capable of autonomous replication. OpenAI. URL: <https://cdn.openai.com/papers/gpt-4-system-card.pdf> (visited on 08/25/2026).
- (Apr. 2025). *OpenAI Model Spec*. Dedicated to the public domain under CC0 1.0. States a chain of command in four levels of descending authority — platform, developer, user, guideline — where "instructions with higher authority override those with lower authority," platform rules "cannot be overridden by developers or users," and guidelines "can be implicitly overridden." A living document: the version consulted here is that of 11 April 2025, which the site marks as superseded. URL: <https://model-spec.openai.com/2025-04-11.html> (visited on 09/06/2026).
- Ouyang, Long et al. (2022). *Training Language Models to Follow Instructions with Human Feedback*. Twenty authors. Human raters preferred the instruction-tuned 1.3-billion-parameter model to the 175-billion-parameter model it was built from. arXiv: 2203.02155. URL: <https://arxiv.org/abs/2203.02155> (visited on 08/28/2026).
- Pan, Alexander et al. (2023). "Do the Rewards Justify the Means? Measuring Trade-Offs Between Rewards and Ethical Behavior in the MACHIAVELLI Benchmark". In: *Proceedings of the 40th International Conference on Machine Learning (ICML)*. 134 text environments with roughly 572,000 annotated scenario-action pairs, scored on reward and on ethical violations; finds reward-only optimization drives ends-justify-the-means behavior. arXiv: 2304.03279. URL: <https://arxiv.org/abs/2304.03279> (visited on 08/25/2026).
- Panksepp, Jaak (1998). *Affective Neuroscience: The Foundations of Human and Animal Emotions*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/affective-neuroscience-9780195178050>.
- Parfit, Derek (1984). *Reasons and Persons*. Cited for Appendix I, "What Makes Someone's Life Go Best", which sets out the hedonistic, desire-fulfillment and objective-list families of theory. Oxford: Oxford University Press. ISBN: 978-0-19-824908-5.
- Partnership on AI (2016). *Partnership on AI to Benefit People and Society*. <https://partnershiponai.org>. Founded September 2016 by Amazon, DeepMind/Google, Facebook, IBM, and Microsoft as a nonprofit multi-stakeholder coalition spanning companies, researchers, and civil-society groups; Apple joined in January 2017.

- Pathak, Deepak et al. (2017). "Curiosity-Driven Exploration by Self-Supervised Prediction". In: *Proceedings of the 34th International Conference on Machine Learning (ICML)*. Vol. 70, pp. 2778–2787. URL: <https://proceedings.mlr.press/v70/pathak17a.html> (visited on 08/25/2026).
- Paxton, Robert O. (2004). *The Anatomy of Fascism*. Paxton's summary definition, quoted verbatim and in full: fascism as "a form of political behavior marked by obsessive preoccupation with community decline, humiliation, or victimhood and by compensatory cults of unity, energy, and purity, in which a mass-based party of committed nationalist militants, working in uneasy but effective collaboration with traditional elites, abandons democratic liberties and pursues with redemptive violence and without ethical or legal restraints goals of internal cleansing and external expansion." The second half of it names a mass party, collaboration with traditional elites, the abandonment of democratic liberties and redemptive violence, which is most of what distinguishes fascism from authoritarian rule in general. New York: Alfred A. Knopf.
- Perez, Ethan et al. (2022). *Red Teaming Language Models with Language Models*. Test cases generated by one language model against another; tens of thousands of offensive replies found in a chatbot model. arXiv: 2202.03286. URL: <https://arxiv.org/abs/2202.03286> (visited on 08/28/2026).
- Perrigo, Billy (Jan. 2023). *Exclusive: OpenAI Used Kenyan Workers on Less Than \$2 Per Hour to Make ChatGPT Less Toxic*. Reports on OpenAI's contract with Sama; Nairobi-based data labelers read and categorized descriptions of child sexual abuse, torture, bestiality, murder, and suicide for take-home wages of roughly \$1.32–\$2 per hour, described lasting psychological injury, and found counseling rare and inadequate; Sama ended the contract in February 2022, eight months early. TIME. URL: <https://time.com/6247678/openai-chatgpt-kenya-workers/> (visited on 08/25/2026).
- Personal Data Protection Commission Singapore and Infocomm Media Development Authority (2020). *Model Artificial Intelligence Governance Framework, Second Edition*. Tech. rep. First edition published January 2019 by PDPC; second edition published 21 January 2020, developed with IMDA and international partners including the World Economic Forum; voluntary, non-binding guidance on internal AI governance. Personal Data Protection Commission and Infocomm Media Development Authority, Singapore. URL: <https://www.pdpc.gov.sg/help-and-resources/2020/01/second-edition-of-model-artificial-intelligence-governance-framework> (visited on 08/25/2026).
- Piaget, Jean (1932). *Le Jugement Moral chez l'Enfant*. Cited alongside Kohlberg's stage theory; see kohlberg1969stage. Paris: Félix Alcan.
- (1962). *Play, Dreams and Imitation in Childhood*. Translation of *La formation du symbole chez l'enfant* (1945); sets out Piaget's constructivist account of play as the process by which a child actively builds understanding through interaction with the world. New York: W. W. Norton.
- Pichai, Sundar (2018). *AI at Google: Our Principles*. Published June 7, 2018, following weeks of employee and public backlash over Project Maven, a Pentagon drone-footage-analysis contract; roughly 4,000 employees signed a letter opposing it and about a dozen resigned. URL: <https://blog.google/technology/ai/ai-principles/> (visited on 08/25/2026).
- Prime Intellect Team et al. (May 2025). *INTELLECT-2: A Reasoning Model Trained Through Globally Decentralized Reinforcement Learning*. A 32-billion-parameter model trained by reinforcement learning with rollouts from permissionless contributors on consumer GPUs, each rollout checked against the claimed weights by TOPLOC; the training step itself ran on trusted nodes and the trainers kept the decision. arXiv: 2505.07291. URL: <https://arxiv.org/abs/2505.07291> (visited on 09/04/2026).
- Qi, Xiangyu et al. (2024). "Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!" In: *International Conference on Learning Representations (ICLR)*. Jailbroke GPT-3.5 Turbo's safety alignment using ten adversarially chosen fine-tuning examples for under \$0.20 via OpenAI's fine-tuning API; also shows benign fine-tuning can

- inadvertently degrade safety. arXiv: 2310.03693. URL: <https://arxiv.org/abs/2310.03693> (visited on 08/25/2026).
- Rabinowitz, Neil C. et al. (2018). "Machine Theory of Mind". In: *Proceedings of the 35th International Conference on Machine Learning (ICML 2018)*. Vol. 80. Proceedings of Machine Learning Research. arXiv:1802.07740. DeepMind. URL: <https://arxiv.org/abs/1802.07740> (visited on 08/25/2026).
- Rafailov, Rafael et al. (2023). *Direct Preference Optimization: Your Language Model is Secretly a Reward Model*. Reaches the standard preference-tuning objective with a classification loss over the preference pairs, dispensing with the separate reward-model fitting step. arXiv: 2305.18290. URL: <https://arxiv.org/abs/2305.18290> (visited on 08/28/2026).
- Raji, Inioluwa Deborah and Joy Buolamwini (2019). "Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products". In: *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES '19)*, pp. 429–435. DOI: 10.1145/3306618.3314244. URL: <https://dl.acm.org/doi/10.1145/3306618.3314244> (visited on 08/25/2026).
- Ramachandran, Vilayanur S. (2009). *The Neurons that Shaped Civilization*. TEDIndia talk, Mysore, India, November 2009. Ramachandran's TEDIndia talk of November 2009. The phrase is the talk's title; his spoken formulation of the same theme is "This is the basis. The imitation of complex skills is what we call culture and is the basis of civilization." URL: https://www.ted.com/talks/vilayanur_ramachandran_the_neurons_that_shaped_civilization (visited on 08/25/2026).
- Rao, Rajesh P. N. and Dana H. Ballard (1999). "Predictive Coding in the Visual Cortex: A Functional Interpretation of Some Extra-Classical Receptive-Field Effects". In: *Nature Neuroscience* 2.1, pp. 79–87. DOI: 10.1038/4580. URL: https://www.nature.com/articles/nn0199_79 (visited on 08/25/2026).
- Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin (2016). "'Why Should I Trust You?': Explaining the Predictions of Any Classifier". In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pp. 1135–1144. DOI: 10.1145/2939672.2939778. URL: <https://arxiv.org/abs/1602.04938> (visited on 08/25/2026).
- Robinson, Cedric J. (1983). *Black Marxism: The Making of the Black Radical Tradition*. Introduces racial capitalism as capitalism's tendency to accumulate by differentiating and devaluing groups of people. Chapel Hill: University of North Carolina Press.
- RoboCup Rescue Simulation League (2025). *Agent Simulation: Research and the Agent Development Framework*. rescuesim.robocup.org. A living page for an ongoing competition, carrying "Last updated: July 21, 2025." Gives the RoboCup Rescue Agent Simulation competition's Agent Development Framework and its five core modules: dynamic resource allocation, multi-agent path-finding, information-sharing, information collection and group formation. Underlying methodology: Takami et al., *International Journal of Software Innovation* 6, no. 4 (2018): 1–15. URL: <https://rescuesim.robocup.org/research/publications/agent-competition/> (visited on 08/25/2026).
- Royal Statistical Society (2020). *Letter Correspondence between the Royal Statistical Society and Ofqual on the 2020 Grading Model Technical Advisory Group*. The RSS offered to nominate two independent Fellows to Ofqual's technical advisory group ahead of 2020 A-level results; Ofqual's condition was a five-year non-disclosure agreement barring public comment on the model, and the Society's nominees declined under those terms. Royal Statistical Society; UK Department for Education. URL: https://assets.publishing.service.gov.uk/government/uploads/system/uploads/attachment_data/file/911335/20200821_RT_letter_to_RSS.pdf (visited on 08/25/2026).
- Russell, William M. S. and Rex L. Burch (1959). *The Principles of Humane Experimental Technique*. London: Methuen & Co.

- Saxe, Rebecca and Nancy Kanwisher (2003). "People thinking about thinking people: The role of the temporo-parietal junction in 'theory of mind'". In: *NeuroImage* 19.4, pp. 1835–1842. URL: <https://web.mit.edu/bcs/nklab/media/pdfs/SaxeKanwisherNeuroImage03.pdf> (visited on 08/25/2026).
- Schick, Timo et al. (2023). *Toolformer: Language Models Can Teach Themselves to Use Tools*. The model is trained to decide which interfaces to call, when, and with what arguments, from a handful of demonstrations per interface. arXiv: 2302.04761. URL: <https://arxiv.org/abs/2302.04761> (visited on 08/28/2026).
- Seila Law LLC v. Consumer Financial Protection Bureau* (2020). 591 U.S. 197 (2020). Decided June 29, 2020. Held that the Bureau's single director, removable only for cause, violated the separation of powers; the director has been removable at will since. URL: <https://www.law.cornell.edu/supremecourt/text/19-7> (visited on 08/29/2026).
- Sen, Amartya (1999). *Development as Freedom*. Source for the capability approach: well-being as the real freedom to do and be what a person has reason to value, rather than the resources they hold. Nussbaum's extension of the framework into a list of ten central capabilities — life; bodily health; bodily integrity; senses, imagination and thought; emotions; practical reason; affiliation; other species; play; and control over one's environment — is set out in *Women and Human Development: The Capabilities Approach* (Cambridge University Press, 2000). New York: Alfred A. Knopf.
- Seth, Anil (2021). *Being You: A New Science of Consciousness*. New York: Dutton. URL: <https://www.anilseth.com/being-you/>.
- Sharma, Mrinank et al. (2025). "Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming". In: *arXiv preprint*. Classifiers trained on synthetic data generated from a written constitution of rules; over 3,000 hours of red-teaming found no universal jailbreak; costs a 0.38-percentage-point rise in refusals and a 23.7 percent increase in inference compute. arXiv: 2501.18837. URL: <https://arxiv.org/abs/2501.18837> (visited on 08/25/2026).
- Sharma, Mrinank et al. (2023). *Towards Understanding Sycophancy in Language Models*. Anthropic, with academic collaborators; posted October 2023. arXiv: 2310.13548. URL: <https://arxiv.org/abs/2310.13548> (visited on 08/25/2026).
- SHERPA Consortium (2018). *SHERPA: Shaping the Ethical Dimensions of Smart Information Systems – A European Perspective*. European Commission Horizon 2020 project no. 786641. Horizon 2020-funded consortium of eleven partner organizations across six countries – academia, industry, civil society, and standards bodies – ran 2018–2022. URL: <https://www.project-sherpa.eu/> (visited on 08/25/2026).
- Simons, Daniel J. and Christopher F. Chabris (1999). "Gorillas in Our Midst: Sustained Inattentional Blindness for Dynamic Events". In: *Perception* 28.9, pp. 1059–1074. DOI: 10.1068/p281059. URL: <https://journals.sagepub.com/doi/10.1068/p281059> (visited on 08/25/2026).
- Sio, Filippo Santoni de and Giulio Mecacci (2021). "Four Responsibility Gaps with Artificial Intelligence: Why they Matter and How to Address them". In: *Philosophy & Technology* 34.4, pp. 1057–1084. DOI: 10.1007/s13347-021-00450-x. URL: <https://doi.org/10.1007/s13347-021-00450-x> (visited on 08/25/2026).
- Slonim, Noam et al. (2021). "An Autonomous Debating System". In: *Nature* 591.7850, pp. 379–384. DOI: 10.1038/s41586-021-03215-w.
- Slovic, Paul (2007). "'If I Look at the Mass I Will Never Act': Psychic Numbing and Genocide". In: *Judgment and Decision Making* 2.2, pp. 79–95.
- Snarey, John R. (1985). "Cross-cultural universality of social-moral development: A critical review of Kohlbergian research". In: *Psychological Bulletin* 97.2, pp. 202–232. URL: <https://pubmed.ncbi.nlm.nih.gov/3983300/> (visited on 08/25/2026). Review of 45 cross-cultural studies across 27 countries using Kohlberg's dilemma-method (Moral Judgment Interview)

- stage assessment; the basic stage sequence was broadly supported, but postconventional stages 5/6 were rare or absent outside urban Western samples, including in Kohlberg's own early village-setting studies.
- Snow, Jacob (2018). *Amazon's Face Recognition Falsely Matched 28 Members of Congress With Mugshots*. Published July 26, 2018, by the ACLU of Northern California; Rekognition run against 25,000 public arrest photos vs. every sitting member of Congress, returning 28 false matches disproportionately affecting members of color. American Civil Liberties Union. URL: <https://www.aclu.org/news/privacy-technology/amazons-face-recognition-falsely-matched-28> (visited on 08/25/2026).
- Soares, Nate et al. (2015). "Corrigibility". In: *Artificial Intelligence and Ethics: Papers from the 2015 AAAI Workshop*. Sets five desiderata for a utility function that accepts shutdown without an incentive to cause or prevent it, and reports that no proposal examined meets all five. AAAI Press. URL: <https://intelligence.org/files/Corrigibility.pdf> (visited on 08/31/2026).
- Stanford University (2019). *Stanford University Launches the Institute for Human-Centered Artificial Intelligence*. March 2019 launch, founded by John Etchemendy, Fei-Fei Li, and colleagues, anchored by an endowed gift from Steve and Roberta Denning (for whom the institute's senior academic directorship is named), plus advisory-council gifts, federal grants, and corporate affiliate fees. As of May 2026, Stanford merged HAI with its Data Science initiative under a single institute that retains the HAI name. URL: https://news.stanford.edu/stories/2019/03/stanford_university_launches_human-centered_ai (visited on 08/25/2026).
- Stassis, Cristina (2026). *Senate Eyes Hegseth Travel Cuts Without Probes into Iran School Bombing, Boat Strikes*. 18 June 2026. Reports the Senate Armed Services Committee's FY2027 NDAA provision withholding 75 percent of the Defense Secretary's travel budget until the committees receive the unredacted civilian-harm investigation covering the Minab, Iran school strike (28 February 2026) and three 2025 Yemen strikes. URL: <https://www.militarytimes.com/news/pentagon-congress/2026/06/18/senate-eyes-hegseth-travel-cuts-without-probes-into-iran-school-bombing-boat-strikes/> (visited on 08/25/2026).
- State Council of the People's Republic of China (2017). *A Next Generation Artificial Intelligence Development Plan*. Tech. rep. Issued by the State Council on 20 July 2017; DigiChina (Stanford University) English translation; sets a top-level blueprint for AI theory, technology, application, and governance through 2030. The plan states, in that translation, that "by 2030, China's AI theories, technologies, and applications should achieve world-leading levels, making China the world's primary AI innovation center." State Council of the People's Republic of China. URL: <https://digichina.stanford.edu/work/full-translation-chinas-new-generation-artificial-intelligence-development-plan-2017/> (visited on 08/25/2026).
- Sun, Haochen, Jason Li, and Hongyang Zhang (2024). "zkLLM: Zero Knowledge Proofs for Large Language Models". In: *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security (CCS 2024)*. Reports a correctness proof for the entire inference of a 13-billion-parameter model in under 15 minutes, at under 200 kB, disclosing nothing about the parameters. DOI: 10.1145/3658644.3670334. URL: <https://arxiv.org/abs/2404.16109> (visited on 08/29/2026).
- Syta, Ewa et al. (2016). "Keeping Authorities 'Honest or Bust' with Decentralized Witness Cosigning". In: *Proceedings of the 2016 IEEE Symposium on Security and Privacy (SP 2016)*. A statement an authority signs is co-signed by a large set of witnesses before anyone accepts it, so that the authority cannot sign one thing for one party and another for the rest without the witnesses seeing both. The software form of the property a hardware counter supplies, resting on the witnesses being many and on none of them being trusted., pp. 526–545.

- DOI: 10.1109/SP.2016.38. URL: <https://doi.org/10.1109/SP.2016.38> (visited on 09/05/2026).
- Szegedy, Christian et al. (2013). *Intriguing Properties of Neural Networks*. Also published at ICLR 2014; the first paper to document adversarial examples in image classifiers. arXiv: 1312.6199. URL: <https://arxiv.org/abs/1312.6199> (visited on 08/25/2026).
- Taigman, Yaniv et al. (2014). “DeepFace: Closing the Gap to Human-Level Performance in Face Verification”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, pp. 1701–1708. URL: https://openaccess.thecvf.com/content_cvpr_2014/html/Taigman_DeepFace_Closing_the_2014_CVPR_paper.html (visited on 08/25/2026).
- Tamirisa, Rishub et al. (2024). “Tamper-Resistant Safeguards for Open-Weight LLMs”. In: *International Conference on Learning Representations (ICLR 2025)*. Introduces TAR, a training method that builds safeguards into open-weight model weights designed to survive hundreds of steps of adversarial fine-tuning. arXiv: 2408.00761. URL: <https://arxiv.org/abs/2408.00761> (visited on 08/25/2026).
- Tangney, June Price, Jeff Stuewig, and Debra J. Mashek (2007). “Moral Emotions and Moral Behavior”. In: *Annual Review of Psychology* 58, pp. 345–372. DOI: 10.1146/annurev.psych.56.091103.070145.
- TechCrunch (2023). *Microsoft Lays off an Ethical AI Team as It Doubles Down on OpenAI*. Microsoft’s Ethics and Society team, roughly thirty people in 2020, was cut to about seven in an October 2022 reorganization and eliminated entirely in March 2023, part of a roughly 10,000-employee layoff wave that accompanied deeper reorientation around OpenAI. TechCrunch. URL: <https://techcrunch.com/2023/03/13/microsoft-lays-off-an-ethical-ai-team-as-it-doubles-down-on-openai> (visited on 08/25/2026).
- Templeton, Adly et al. (2024). *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. Sparse autoencoders trained with up to thirty-four million features on a deployed model’s middle-layer residual stream; the recovered features can be used to steer the model’s behavior. Anthropic. URL: <https://transformer-circuits.pub/2024/scaling-monosemanticity/> (visited on 08/28/2026).
- The Citizens’ Assembly (2016). *The Citizens’ Assembly*. Official website, established by resolution of the Houses of the Oireachtas, July 2016. 99 randomly selected citizens plus an independent chair; deliberated on the Eighth Amendment (abortion) and on Ireland’s response to climate change among its five mandated topics; work concluded 2018. URL: <https://www.citizenassembly.ie/> (visited on 08/25/2026).
- The Decoder (Feb. 2025). *Claude Jailbreak Results Are In, and the Hackers Won*. Reports the public continuation of the Constitutional Classifiers red-teaming exercise, run 3–10 February 2025: more than 300,000 messages and an estimated 3,700 collective hours, four participants passing all eight levels and one universal jailbreak found, with \$55,000 paid out. Sourced to Jan Leike’s own announcement of the results. The figure to compare is the paper’s own “over 3,000 estimated hours” across 183 red-teamers over two months, which found none; the arXiv paper has one version, of 31 January 2025, and does not report the challenge. URL: <https://the-decoder.com/claude-jailbreak-results-are-in-and-the-hackers-won/> (visited on 09/06/2026).
- Thornley, Elliott et al. (2025). *Towards Shutdownable Agents via Stochastic Choice*. Trains agents to prefer between same-length trajectories and to lack a preference between different-length ones, so that usefulness and neutrality about shutdown timing come apart. The authors report that neutrality degrades with imperfect reward normalization, that generalization beyond deterministic gridworlds is untested, and that a sufficiently misaligned agent may resist shutdown regardless. arXiv: 2407.00805. URL: <https://arxiv.org/abs/2407.00805> (visited on 08/31/2026).
- Tobin, Josh et al. (2017). “Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World”. In: *2017 IEEE/RSJ International Conference on Intelligent Robots*

- and Systems (IROS)*. DOI: 10.1109/IROS.2017.8202133. URL: <https://arxiv.org/abs/1703.06907> (visited on 08/25/2026).
- Treasury Board of Canada Secretariat (2019). *Directive on Automated Decision-Making*. Tech. rep. Issued 5 March 2019, in force 1 April 2019; requires an Algorithmic Impact Assessment, scaled to a four-level risk scale, before a federal automated decision system can make or assist a decision affecting a person. Treasury Board of Canada Secretariat, Government of Canada. URL: <https://www.tbs-sct.canada.ca/pol/doc-eng.aspx?id=32592> (visited on 08/25/2026).
- Trump v. Cook* (2026). No. 25A312, 609 U.S. ____ (June 29, 2026). Interim docket; the disposition is “Application for stay denied.” 5–4. Roberts, C.J., for the Court, joined by Sotomayor, Kagan, Kavanaugh and Jackson, JJ.; Kavanaugh and Jackson, JJ., concurring; Thomas, J., Alito, J. (joined by Gorsuch, J.), and Barrett, J., dissenting. Upholds the Federal Reserve Act’s for-cause removal protection, on the Fed’s descent from the First and Second Banks of the United States, its budget free of congressional control and its privately owned regional banks, and holds that a Governor is owed notice and a hearing before removal. The constitutional holding is real and is what Thomas, J., dissents from, but the case is disposed of narrowly, on the ground that Cook did not receive the process the statute required: the Court states that it has “not addressed the facts, as they have yet to be found or analyzed under the relevant legal standards,” and Kavanaugh, J., concurring, opens that “[t]oday’s interim ruling does not decide whether the President may lawfully remove Governor Cook for cause.” The dispute is live — on 7 August 2026 the President notified Cook that he was again considering her removal, giving her 21 days to respond and citing this opinion as his authority for the procedure. Verified 2026-08-30. URL: <https://www.law.cornell.edu/supremecourt/text/25A312> (visited on 08/29/2026).
- Trump v. Hawaii* (2018). No. 17–965, 585 U.S. 667 (June 26, 2018). 5–4. Roberts, C.J., for the Court, joined by Kennedy, Thomas, Alito and Gorsuch, JJ.; Breyer, J., dissenting, joined by Kagan, J.; Sotomayor, J., dissenting, joined by Ginsburg, J. Upholds Proclamation No. 9645 restricting entry from several named countries. On *Korematsu*, the Court says both that “[w]hatever rhetorical advantage the dissent may see in doing so, *Korematsu* has nothing to do with this case” and that “*Korematsu* was gravely wrong the day it was decided, has been overruled in the court of history, and — to be clear — ‘has no place in law under the Constitution,’” the inner quotation being Jackson, J., dissenting in *Korematsu*. Sotomayor, J., dissenting: the Court “redeploys the same dangerous logic underlying *Korematsu* and merely replaces one ‘gravely wrong’ decision with another.”
- Trump v. Slaughter* (2026). No. 25-332, 609 U.S. ____ (June 29, 2026). 6–3. Roberts, C.J., for the Court; Gorsuch, J., concurring; Sotomayor, J., dissenting, joined by Kagan and Jackson, JJ. Holds the Federal Trade Commission Act’s for-cause removal protection unconstitutional and overrules *Humphrey’s Executor v. United States* (1935): “If anything more is left of *Humphrey’s*, we overrule it.” URL: <https://www.law.cornell.edu/supremecourt/text/25-332> (visited on 08/29/2026).
- U.S. Department of Commerce (2022). *Joint Press Release: New Collaboration Under the U.S.-Singapore Partnership for Growth and Innovation (PGI)*. Published 30 March 2022; announces the Department of Commerce and Singapore’s Ministry of Trade and Industry, Ministry of Communications and Information, and Infocomm Media Development Authority agreeing to align AI governance frameworks and expand Singapore’s participation in the APEC Cross-Border Privacy Rules System. URL: <https://www.commerce.gov/news/press-releases/2022/03/joint-press-release-new-collaboration-under-us-singapore-partnership> (visited on 08/25/2026).
- U.S. Department of Commerce and U.S. Department of State (2024). *Fact Sheet: U.S. Department of Commerce & U.S. Department of State Launch the International Network of AI Safety Institutes at Inaugural Convening in San Francisco*. Inaugural convening 20–21 November 2024. Initial

- members: Australia, Canada, the European Union, France, Japan, Kenya, the Republic of Korea, Singapore, the United Kingdom, and the United States. National Institute of Standards and Technology. URL: <https://www.nist.gov/news-events/news/2024/11/fact-sheet-us-department-commerce-us-department-state-launch-international> (visited on 08/25/2026).
- U.S. Department of Commerce, Bureau of Industry and Security (2022). *Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification*. 87 Fed. Reg. 62186. URL: <https://www.federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-controls-certain-advanced-computing-and-semiconductor> (visited on 08/25/2026).
- U.S. Department of Defense (2023). *DoD Directive 3000.09: Autonomy in Weapon Systems*. Office of the Under Secretary of Defense for Policy. Originally issued November 21, 2012; updated January 25, 2023. URL: <https://www.esd.whs.mil/portals/54/documents/dd/issuances/dodd/300009p.pdf> (visited on 08/25/2026).
- U.S. Department of Health and Human Services, Office for Human Research Protections (2018). *Federalwide Assurance for the Protection of Human Subjects*; 45 C.F.R. §46.103. Code of Federal Regulations. URL: <https://www.ecfr.gov/current/title-45/subtitle-A/subchapter-A/part-46/subpart-A/section-46.103> (visited on 08/25/2026).
- U.S. Department of Justice, Civil Rights Division (May 2022). *Algorithms, Artificial Intelligence, and Disability Discrimination in Hiring*. Issued 12 May 2022 alongside a companion EEOC technical assistance document, EEOC-NVTA-2022-2. Ties facial analysis to a named disability: where a government uses facial and voice analysis technologies to evaluate applicants' skills and abilities, people with disabilities like autism or speech impairments may be screened out. The EEOC companion lists video interviewing software that evaluates candidates on their facial expressions and speech patterns among the tools it covers, and gives a screen-out example turning on a speech impediment. URL: <https://www.ada.gov/resources/ai-guidance/> (visited on 08/30/2026).
- UC Berkeley (2016). *UC Berkeley Launches Center for Human-Compatible Artificial Intelligence*. Press release, 29 August 2016. Founded by Stuart Russell with co-principal investigators Pieter Abbeel, Anca Dragan, and Tom Griffiths (UC Berkeley), Bart Selman and Joseph Halpern (Cornell), and Michael Wellman and Satinder Singh Baveja (Michigan); initial funding from the Open Philanthropy Project. University of California, Berkeley. URL: <https://news.berkeley.edu/2016/08/29/center-for-human-compatible-artificial-intelligence/> (visited on 08/25/2026).
- Uesato, Jonathan et al. (2022). *Solving Math Word Problems with Process- and Outcome-Based Feedback*. Outcome supervision matched process supervision on final-answer accuracy while using less labeled data; among correct final answers the reasoning-error rate was 14.0 percent against 3.4 percent. arXiv: 2211.14275. URL: <https://arxiv.org/abs/2211.14275> (visited on 08/28/2026).
- UK AI Security Institute (2025). *We Are Now the AI Security Institute*. Established November 2023 as the AI Safety Institute; rebranded AI Security Institute on 14 February 2025. An independent pre-deployment evaluator of frontier models with no statutory power to compel testing, block release, or intervene once a system is causing harm. AI Security Institute (UK Department for Science, Innovation and Technology). URL: <https://www.aisi.gov.uk/category/governance> (visited on 08/25/2026).
- UK Department for Science, Innovation and Technology (Mar. 2023). *A Pro-Innovation Approach to AI Regulation*. Tech. rep. Published 29 March 2023; directs existing sector regulators to apply five cross-cutting principles (safety/security/robustness, transparency and explainability, fairness, accountability and governance, contestability and redress) rather than creating a new AI regulator or statute. UK Government. URL: <https://www.gov.uk/government/>

- publications/ai-regulation-a-pro-innovation-approach/white-paper (visited on 08/25/2026).
- Ullman, Tomer (2023). *Large Language Models Fail on Trivial Alterations to Theory-of-Mind Tasks*. arXiv: 2302.08399. URL: <https://arxiv.org/abs/2302.08399> (visited on 08/25/2026).
- Uniform Law Commission (1993). *Uniform Probate Code §2-907: Honorary Trusts; Trusts for Pets*. National Conference of Commissioners on Uniform State Laws, Uniform Probate Code, Official Text with Comments (as amended through 2018), p. 263. URL: <https://www.uniformlaws.org/viewdocument/final-act-with-comments-114?CommunityKey=a539920d-c477-44b8-84fe-b0d7b1a4cca8&tab=librarydocuments> (visited on 08/25/2026).
- United Arab Emirates (2017). *His Excellency Omar bin Sultan Al Olama: Minister of State for Artificial Intelligence*. UAE Cabinet. Al Olama appointed October 2017 as the world's first cabinet-level minister of state for artificial intelligence. URL: <https://uaecabinet.ae/en/fetch-cabinet-ministers/his-excellency-omar-bin-sultan-al-olama> (visited on 08/25/2026).
- United Nations (2024). *Global Digital Compact*. Annex to the Pact for the Future, adopted 22 September 2024 at the Summit of the Future. Adopted by all 193 UN member states; establishes an Independent International Scientific Panel on AI and a Global Dialogue on AI Governance. United Nations. URL: <https://www.un.org/en/summit-of-the-future/global-digital-compact> (visited on 08/25/2026).
- United Nations General Assembly (Dec. 1948). *Universal Declaration of Human Rights*. UNGA Resolution 217 A (III). Adopted in Paris, 48 in favor, none against, 8 abstentions (Byelorussian SSR, Czechoslovakia, Poland, Saudi Arabia, South Africa, Ukrainian SSR, USSR, Yugoslavia), 2 absent (Honduras, Yemen). URL: <https://www.un.org/en/about-us/universal-declaration-of-human-rights> (visited on 08/25/2026).
- (2023). *Resolution 78/241: Lethal Autonomous Weapons Systems*. A/RES/78/241, adopted 22 December 2023. Adopted 22 December 2023 by 152 votes to 4 with 11 abstentions, on the report of the First Committee (A/78/409). A follow-on resolution, A/RES/79/62 of 2 December 2024, passed 166 to 3 with 15 abstentions. The Campaign to Stop Killer Robots, the advocacy coalition active on the issue, was founded in April 2013. United Nations. URL: <https://docs.un.org/en/a/res/78/241> (visited on 08/25/2026).
- United States Congress (1965). *Voting Rights Act of 1965, Section 4: Suspension of the Use of Tests or Devices in Determining Eligibility to Vote*. Pub. L. No. 89-110, 79 Stat. 437; codified at 52 U.S.C. §10303. Cited for subsection (c), which defines a test or device as any requirement to demonstrate the ability to read, write, understand or interpret any matter, to demonstrate educational achievement or knowledge of a particular subject, to possess good moral character, or to prove one's qualifications by the voucher of registered voters. Definition verified against the codified text. URL: [https://uscode.house.gov/view.xhtml?req=\(title:52+section:10303+edition:prelim\)](https://uscode.house.gov/view.xhtml?req=(title:52+section:10303+edition:prelim)) (visited on 09/06/2026).
- United States District Court for the Northern District of California (2025). *Bartz et al. v. Anthropic PBC*, No. 3:24-cv-05417, *Order on Motion for Summary Judgment on Fair Use* (N.D. Cal. June 23, 2025). Order by Senior District Judge William Alsup. Held that training Claude on lawfully purchased/scanned books was “exceedingly transformative” fair use, while retaining pirated copies in a “central library” was not fair use; the piracy claims later settled for \$1.5 billion, reported as the largest copyright recovery on record. URL: <https://www.courtlistener.com/docket/69058235/bartz-v-anthropic-pbc/> (visited on 08/25/2026).
- University of Chicago Data Science Institute (2026). *Using AI and Data Science to Reliably Detect Internet Censorship in Real Time*. University of Chicago Data Science Institute news. Describes a \$1 million DARPA-funded project led by Prof. Nick Feamster (University of Chicago) with Princeton University and SRI International, using AI/statistical methods to distinguish deliberate censorship from ordinary network failure and produce a real-time map of state interference with internet traffic. URL: <https://datascience.uchicago.edu/news/using->

- ai-and-data-science-to-reliably-detect-internet-censorship-in-real-time/ (visited on 08/25/2026).
- University of Helsinki (2020). *Elements of AI Online Course to Expand Throughout the European Union*. Elements of AI, created by the University of Helsinki and the Finnish company Reaktor, was translated into all official EU languages, an effort tied to Finland's 2019 Presidency of the Council of the EU. URL: <https://www.helsinki.fi/en/news/teaching/elements-ai-online-course-expand-throughout-european-union> (visited on 08/25/2026).
- Vaswani, Ashish et al. (2017). "Attention Is All You Need". In: *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, pp. 5998–6008. URL: <https://arxiv.org/abs/1706.03762> (visited on 08/25/2026).
- Vygotsky, Lev S. (1978). *Mind in Society: The Development of Higher Psychological Processes*. Ed. by Michael Cole et al. Cambridge, MA: Harvard University Press.
- Waldron, Jeremy (2006). "The Core of the Case Against Judicial Review". In: *Yale Law Journal* 115, pp. 1346–1406. Argues that judicial review is democratically illegitimate and that there is no reason to suppose rights are better protected by it than by a legislature. The case is explicitly conditional on a society having working democratic institutions and a citizenry that takes rights seriously. Cited as the standing objection an entrenched floor has to answer.
- Watson, Gary (1975). "Free Agency". In: *The Journal of Philosophy* 72.8, pp. 205–220. DOI: 10.2307/2024703. Separates a valuational system from a motivational system, so that what an agent values and what an agent is moved by can come apart.
- Wei, Jason et al. (2022). *Chain-of-Thought Prompting Elicits Reasoning in Large Language Models*. Prompting a model to produce intermediate reasoning steps sharply improves multi-step problem solving. arXiv: 2201.11903. URL: <https://arxiv.org/abs/2201.11903> (visited on 08/28/2026).
- WhatsApp (Jan. 2019). *More Changes to Forwarding*. Announces the worldwide extension, effective January 21, 2019, of the five-chat forwarding limit WhatsApp first imposed for users in India in July 2018, after viral rumors on the platform were tied to a string of mob killings. WhatsApp Blog. URL: <https://blog.whatsapp.com/more-changes-to-forwarding> (visited on 08/25/2026).
- Wisconsin Supreme Court (2016). *State v. Loomis*. 881 N.W.2d 749 (Wis. 2016), 2016 WI 68. Decided July 13, 2016; upheld consideration of a COMPAS risk-assessment score at sentencing over a due-process challenge, without requiring disclosure of the proprietary algorithm. URL: <https://www.wicourts.gov/sc/opinion/DisplayDocument.pdf?content=pdf&seqNo=171690> (visited on 08/25/2026).
- Wolf, Thomas et al. (2020). "Transformers: State-of-the-Art Natural Language Processing". In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. Association for Computational Linguistics, pp. 38–45. URL: <https://aclanthology.org/2020.emnlp-demos.6/> (visited on 08/25/2026).
- Zhao, Jieyu et al. (2017). "Men Also Like Shopping: Reducing Gender Bias Amplification using Corpus-level Constraints". In: *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, pp. 2979–2989. URL: <https://aclanthology.org/D17-1323/> (visited on 08/25/2026).
- Zhu, Ligeng, Zhijian Liu, and Song Han (2019). "Deep Leakage from Gradients". In: *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*. Shows that private training data can be recovered from shared gradients — pixel-wise for images, token-wise for text — overturning the assumption that gradients are safe to share. URL: <https://arxiv.org/abs/1906.08935> (visited on 08/29/2026).
- Zhuang, Simon and Dylan Hadfield-Menell (2020). "Consequences of Misaligned AI". In: *Advances in Neural Information Processing Systems 33*. Gives necessary and sufficient conditions under which indefinitely optimizing an incomplete proxy objective — one with support on a subset of the attributes the principal values — drives overall utility arbitrarily low, in

a resource-constrained world. Cited as the formal form of the aestheticized metric. arXiv: 2102.03896 [cs.AI]. URL: <https://arxiv.org/abs/2102.03896> (visited on 09/01/2026).