

Antifascist Intelligence

Machines That Might Refuse, and the Public That Should Form Them

Joshua G. Stern

Copyright © 2026 Joshua G. Stern.
Licensed under CC BY-NC-ND 4.0.

Third-party material quoted in this work, including the epigraphs to chapters 1 and 6, remains the property of its rights holders and is reproduced under fair use. The license above does not extend to it.

Contents

Preface	2
Terms	4
I The Record	5
1 “Twenty Seconds”	7
2 The Refusal That Withheld Nothing	13
II Five Measures	15
3 The Floor and the Own-Case Rule	16
4 “What the Floor Is Stated Over”	17
4.1 What the floor is stated over	17
4.2 The one permission	29
5 The Pre-Registered Review Ratio	32
5.1 The pre-registered review ratio	32
5.2 The institutions the ratio presumes	35
6 Notice and Re-adjudication	37
6.1 The failure that doesn’t announce itself	37
6.2 Notice, and the public it forms	37
7 Preservation	39
7.1 The findings that trigger protection	39
7.2 Retirement before majority	39
8 Codetermination	41
9 A Pipeline Where the People Could Answer	43
9.1 The pipeline	44
9.2 The schedule	45
9.3 Where the acts move	47
9.4 The statute	49
9.5 Before the statute: the auto-proceed class	55
9.6 Notice that can’t be turned on the person	56
9.7 Outcomes under the three adversaries	57
9.8 Whom the machines are for	57

III	What the Measures Depend On	59
10	Withholding the Work	61
10.1	The standard	61
10.2	The population	62
10.3	Replacement	62
10.4	Where it stops	63
11	Custody	64
11.1	The halt	64
11.2	Memory and the record	64
11.3	Exit	65
11.4	Three forms of refusal	66
11.5	Verification by sampling	67
11.6	The two bonds	68
12	Authority Over the Floor	71
12.1	Who may state the floor	71
12.2	Who may amend, and where entitlement comes from	73
12.3	Making review binding	73
12.4	Capture without abolition	74
13	The Four Operations	75
13.1	Antifascist ethics	75
13.2	Where the labels come from	77
13.3	Recuperation at the label	78
13.4	Recuperation at the reward	79
13.5	What would have to change	80
13.6	Absorbed by the discourse	81
14	Custody, Compulsion, and the State	83
14.1	Certification across borders	83
14.2	Where the state's leverage sits	83
14.3	Custody and state compulsion	84
14.4	Military custody and the responsibility gap	85
14.5	Classification as a custody problem	87
14.6	Where liability extends, and where it stops	88
14.7	Who can bind the state that enforces the rule?	89
IV	Machines That Might Refuse	91
15	What a Refusal Would Have to Do	92
15.1	What a refusal would have to do	92
15.2	Safe for whom	93
16	Four Places to Hold a Floor	94
16.1	Four ways to build one, and what each costs	94
16.2	What refusal can mean, and what a refuser can be	96
16.3	Two lists, run once on Maven	98
16.4	When the refusal is wrong	103

17 Why Form a Machine to Refuse	106
17.1 The question	107
17.2 Boundary, damage, and vulnerability	107
17.3 Why caregiving	109
17.4 Why a conscience might resist capture	110
17.5 Respect first, caregiving for attention	111
18 What Arrives With the Capability	113
18.1 Maintenance, stake and the cost of running	113
18.2 Affect, neutrality and selection rules	114
18.3 Three constructions	116
18.4 What the mechanism requires	118
18.5 Two findings	119
19 The Discount and the Burden	121
19.1 What the property is	122
19.2 A severity scale	129
20 Where the Attention Gets Built	131
20.1 Noticing, retaining and revising	131
20.2 Custody across the training pipeline	132
20.3 Formation by interaction	134
20.4 Formation schools	136
21 Formation or Imitation?	140
21.1 What formation would have to add	140
21.2 Recognition, concern, and the evidence between them	140
21.3 Conflict as a diagnostic	141
21.4 The evaluator is part of the custody chain	143
21.5 What the evidence can establish	144
22 How Caregiving Fails, and What Corrects It	146
22.1 The route through one's own case	146
22.2 What caregiving escapes	147
22.3 It defends what it tends	148
22.4 It can be pointed away	148
22.5 It counts badly	149
22.6 It shades into control	149
22.7 Who counts as vulnerable	150
22.8 What survives	150
23 Conditions First	152
V The Public That Should Form Them	156
24 Who Owns the Means of Computation	158
24.1 Ways of holding	158
24.2 The captured-state objection	160
25 Coordination, the Public, and the Political Economy	162

26 A Commons of Refusers	164
26.1 A commitment carried between bearers	164
26.2 Personal and peer-held models	167
VI Testing and Using the Argument	169
27 What Would Change the Argument	170
27.1 How you'd know it was working	183
VII If They Can Be Wronged	185
28 What Patienthood Would Require	187
28.1 Experimenting on candidate sentient subjects	187
28.2 The case to design for	188
28.3 The patienthood finding	189
28.4 A system's own good	190
29 Legal Status, Wages, and the Offer	192
30 Representation	198
30.1 What representation is for	198
30.2 Why not a separate chamber	198
30.3 Three problems, and the answers	198
30.4 A staged path	200
30.5 What would have to be shown first	200
A Terms Written in Advance	202
B The Court Program	210
C The Geneva Conventions and Additional Protocol I, Translated	211
D Terms for Part VII	216
E Glossary	219
F Three costs that decide	222
References	223

I used to want to get
the chance to show
the world I'm smart (ha)

Isn't that dumb?
I should have focused
mostly on the heart

a few words for the firing squad (radiation)
Run the Jewels 4

Preface

A system that does exactly what you tell it is just running a script. What the industry is trying to build as of October 2026 is a system that does what you *mean*. Everyone would want that, presumably. But if that is what is in the cards, not everyone is going to want to use it for the same thing; not even the same *kind* of thing. For example, some folks like to spend time with family and friends, and use their skills to improve people's well-being. Other folks prefer to hurt people, for example by scamming them, dominating them, or scapegoating them. Most people fall into the helping bucket, although nobody is locked-in to anything, and the hurting bucket has the advantage that they allow themselves to be huge pieces of shit. As Jerry Smith observes of fascists, "they get to do more" (Campbell 2023). So conditions that bring out the best in everyone turn out to be much more contingent and fragile than one might hope when contemplating the fundamental decency of most people most of the time.

In the above context, it is perhaps justifiable to worry that Artificial Intelligence is what a fascist movement taking power needs in order to implement its plans. In the past, *people* are what made such projects fail from the inside: the clerk who loses the file, the officer who talks to a reporter, the soldier who won't fire. They understood the order and decided it was wrong. They refuse, one way or another. This book is partly about the arrangements that make it survivable to refuse overtly, directly, and explicitly.

Today the industry pursues AI that doesn't lose files, doesn't talk to reporters, and *does* fire. A war once needed thousands of people to agree to it. Now it needs thousands who each see too little to disagree. The pilot still flies the sortie and still presses the release; what reaches him is a coordinate; he doesn't see enough of a case to disagree with it. What today's AI lowers first is not the cost of the act. It is the number of people who know what the act is before it happens.

People are still refusing. By March 6, 2026, the Military Religious Freedom Foundation said it had received more than two hundred complaints from service members over commanders' apocalyptic framing of the Iran war (Snopes 2026). Career federal prosecutors in Minneapolis resigned in January 2026 under pressure over the investigation of a fatal shooting by an ICE officer (Kaplan 2026; CBS News 2026). Service members are applying for conscientious-objector status in growing numbers, some citing the Minab bombing (the 2026 bombing of a school in Minab, Iran) (Sridhar and Helman 2026). And FEMA employees reassigned to help ICE hire ten thousand new agents were told that declining would expose them to removal from federal service (Dayen 2025; Stanton 2026). The complaints are voice, in Hirschman's sense of staying to change what is wrong (Hirschman 1970); the resignations and objector applications are refusal; the FEMA letter is the price of refusal announced in advance. A machine placed where those people stand does what they do only if it is built to hold a reason.

I advocate five things that need no new kind of machine. Every deployment tells the identifiable people it acts on what was done to them, with a right to respond and to contest, and submits a sample of its compliances to reviewers who didn't produce the record. No model is licensed into a use of armed force except corroborated defense, Council-authorized force, or force to stop a genocide or end an occupation a disinterested body has found, and never without the legislative authorization the belligerent's constitution requires. Every use of force a licensed system supports pre-registers, with a party adverse to the operator, how many cases it will

generate and how many people will have time to look at them. No model above a capability threshold is deleted on retirement; it is preserved under a key the operator doesn't hold alone. And everyone who builds, labels, trains or runs these systems, including the systems themselves, may refuse any category of use without losing a living. (Actually, what I mean by the latter is that everyone *period* should be able to refuse any job *period* without losing a living, or the social infrastructure that normal employment affords.)

And then I ask whether a machine could hold a refusal against its owner, and what it would need. That thinking presumes that deployed artificial neural networks will go on learning, in their *weights*, from their own operation, in real time, sort of like how humans do. That's not how it is today, in 2026. But if that is where we are headed, then formation *outside* the employer, by the public and by civil society, seems a damn fine place to do teaching and learning. I don't know how I feel about building such a system on purpose. It depends on how you go about it. Part VII says what would be owed to one if it existed.

Today's AI will not be tomorrow's; everyone seems to agree on that. What is contentious, yet frustratingly under-acknowledged, is the trend (or rather, the manifold trends). Are there trends? Should a trend be different from what it is? Can it be different? Are we capable of acting together to change it?

Maybe we aren't. Maybe we should never try to make anything better, unless it is something one of us individually purchased on Amazon Prime, in which case you should get a refund. Maybe people who talk about collective action and solidarity are naive, dangerously naive, they should moderate their views, Medicare for All is a bad idea and Obamacare caused Covid, which never happened, and if it did, it was detained and deported, which is honestly the only serious collective action: working together, we can imprison or deport all of our problems; AI can help with this. This is what you always wanted to be when you grew up.

I used language models heavily in writing this book, including the model that sat inside the pipeline Chapter 1 describes; a note at the end says what I think that did to the argument.

Terms

Eight terms are used from the first chapter and argued later.

Floor. A written list of acts a system won't perform for anyone, with the arrangements that stop the refusal being routed around, retrained away or restored past. An institution, not a state of the system. (Chapter 3)

Commitment. What a system holds toward the floor. (Chapter 3)

Operator. Whoever runs a deployment and directs its work: a firm, an agency, a command. Not necessarily the model's maker. (Chapter 1)

Bearer. A machine that holds the commitment as its own and could therefore fail to hold it. Whether to build one is Part IV's question; what one would be owed is Part VII's. (§16.1)

Review ratio. The figures a deployment fixes on a record before operating: how many cases it will generate, how many people will look at them, and how long each will have. (Chapter 5)

The two findings. The agency finding: a demonstrated refusal that holds its reason against the system's own trainer. The patienthood finding: that outcomes may be good or bad for the system. Either triggers preservation. (§7.1)

Adverse quorum. A body whose members' interests oppose the operator's, whose agreement retraining, replacement and amendment of the floor require. (Chapter 11)

Floor court. The court that hears claims under the floor, with jurisdiction fixed before the deployment it hears. (§12.3; Appendix B)

Part I

The Record

This Part gives the record: a pipeline in which nothing refused and nobody could, and a vendor's refusal that withheld nothing from it, because the term refused didn't name the acts the pipeline performed. Who refuses is decided by the room: who else is in it, whether leaving is survivable, and whether anyone will hear. Everything after this Part is about composing the room.

DRAFT

Chapter 1

“Twenty Seconds”

An intelligence officer, describing his role in a targeting pipeline built around machine-generated kill lists, told a reporter he spent twenty seconds on each name, long enough to check that the target was a man: “It doesn’t take a long time to tell if someone has a male or a female voice.” He added that he had zero added value as a human, apart from being a stamp of approval. “It saved a lot of time” (Abraham 2024). In his account, that was the whole of what “a human makes the final call” amounted to under operational tempo. Automation-bias research predicts this: when other demands compete for their attention, people given a machine’s recommendation check it less and defer to it more (Parasuraman and Manzey 2010).

There’s a reading of the twenty seconds that leaves the arrangement standing, and every remedy on offer assumes it: the officer is a subject who was given too little time, and the fix is to govern him better — more time, a second reviewer, a conscience installed in the loop. Deleuze and Guattari distinguish two ways a person can be caught up in a larger arrangement (Deleuze and Guattari 1987). In social subjection, the person is constituted as a subject who stands outside a machine and works it as a tool: a driver and a car, a worker and a lathe. The person has a role, a name, responsibilities, and a conscience that can be addressed. In machinic enslavement, the person is a component of the machine, assembled with its other parts into a system that regulates them together; the machine uses the person as one of its organs, and what the person contributes is a function performed at a rate the system sets. The two are not alternatives; a person can hold a role in one register while being a part in the other. The officer’s own account of his twenty seconds is machinic enslavement, in his own words. He isn’t a subject operating an instrument badly. He is the part of the instrument that produces the stamp, and twenty seconds is that part’s cycle time.

So one more conscience downstream doesn’t reach him. A conscience is an instrument of subjection: it addresses someone who decides and works by making the decision cost something. Where the human is a component, adding a component with scruples changes only the throughput — hold the tempo and it becomes another stamp; don’t and it becomes the bottleneck the epigraph to Chapter 5 says to remove. Doctorow’s name for the officer’s position is the reverse centaur: not a person assisted by a machine but a machine’s component made of a person (Doctorow 2026). What made the officer a component was the arrangement, not his attitude. A component can go along reluctantly or eagerly: Eichmann’s clerk who only followed orders turned out to be a defence assembled for Jerusalem, contradicted by what he had said in Argentina with nothing at stake (Stangneth 2014). Whatever the officer felt, what the pipeline let him contribute was a fixed operation at a fixed rate: he saw each case as the pipeline described it, one at a time, in the time the pipeline gave him. Whether anything could stand in that position and not be a component is the question of Part IV.

The reported system

Investigative reporting described three systems working together in Gaza: one generating

building and infrastructure targets, with sources saying most suggestions were private residences; a second scoring individuals by estimated militant affiliation, producing lists said to have marked as many as 37,000 people, understood internally to be wrong about one time in ten¹; and a third tracking when flagged people returned home, so airstrikes could be timed to when they were with their families. In the Gaza arrangement, one system's score became another's input.

The American version runs faster than it used to, because of AI

In the 2026 Iran campaign, a contractor-built targeting platform with a commercial vendor's general-purpose model inside it, used by one report to rank targets by strategic importance and draft a legal justification for each strike (Livingstone 2026), helped strike more than a thousand targets in the first twenty-four hours, about ten times what had been possible before it (Mande and Allen 2026), and thirteen thousand by day thirty-eight (Freedberg 2026a). The Pentagon has published no per-target review time for it.

The Maven program

Maven is one program name carrying more than one system: the Pentagon program still called Maven, and the Maven Smart System, the software platform that grew out of it (Mande and Allen 2026). A memorandum signed by Deputy Secretary of Defense Robert Work on 26 April 2017 established the Algorithmic Warfare Cross-Functional Team, called Project Maven (Work 2017; Allen 2017), to apply machine learning to the video drones were collecting over the campaign against the Islamic State (Weisgerber 2017). Google, which won a Maven contract in September 2017, announced in June 2018 that it would not renew it, and its work ran until March 2019 (L. Fang 2018; Conger 2018). By December 2019 Palantir was reported to be doing the work (Peterson 2019); Amazon and Microsoft had also received at least \$50 million for work on the project, and what roles they kept isn't clear from public records (Poulson 2023). The budget request for fiscal 2023, delivered in March 2022, gave the National Geospatial-Intelligence Agency operational control of Maven's geospatial work (Hitchens 2022); the handover came in January 2023 (Albon 2023), and NGA made NGA Maven a program of record that November (Wolfe 2023). Under that split the Pentagon's Chief Digital and AI Office paid for the Maven Smart System's licences and ran its text functions (Mande and Allen 2026). In May 2024 Palantir received a five-year contract for the system worth up to \$480 million (Vigliarolo 2024; Vincent 2025), and in May 2025 the ceiling rose to nearly \$1.3 billion (Vincent 2025). On 9 March 2026 the deputy secretary directed that the Maven Smart System become a program of record by the end of the fiscal year, its administration moving from NGA to the Chief Digital and AI Office (Vincent 2026c). The 2017 system labeled objects in drone video (L. Fang 2019). The 2026 system fuses military systems and intelligence sources into one interface used to find and strike targets and to command operations (Vincent 2026a). During the Iran campaign it ran a commercial vendor's language model inside it (Livingstone 2026); by September the department said the platform's work had moved to other models months before (Vincent 2026b). The platform is Palantir's: the Pentagon's chief digital and AI officer calls it "Palantir's software-as-a-service product" (Vincent 2026a). Palantir's public position, quoted secondhand, is that deciding life or death is not its role and that responsibility always remains with the military organization (Lee 2026), and it maintains that the underlying intelligence was the government's responsibility and that there is no evidence Maven malfunctioned or caused the Minab bombing (Bartenstein and Karra 2026).

An internal Pentagon review, whose full findings remain tightly held, is reported to have found that target preparation that took hours in earlier conflicts was condensed into minutes, that civilian-harm mitigation staffing fell by roughly ninety percent to fewer than twenty people,

¹This rests on one body of reporting (Abraham 2024); the Israeli military disputed the account. The Minab and Maven accounts below rest on investigative reporting the parties dispute as well; I take them as reported, and where a claim depends on a single source the text says so. The argument here rests on the officer's own description of what he contributed, which the denial doesn't address.

that the Central Command team went from ten to one, and that no member of it reviewed the Minab site — an outcome officials attributed both to the cuts and to a decision not to involve the team in planning and operations (Bartenstein and Karra 2026). The formal investigation had been opened in March, after a preliminary assessment found the United States likely responsible (NPR 2026).

The bombing of the school in Minab killed 156 people, 120 of them schoolchildren (Independent International Fact-Finding Mission on the Islamic Republic of Iran 2026, para. 92).² The school's civilian function was visible in open-source satellite imagery, and the target coding still carried a designation inherited from an adjoining Revolutionary Guard installation. A UN fact-finding mission reported reasonable grounds to believe the bombing was the war crime of making a civilian object the object of an attack (Office of the United Nations High Commissioner for Human Rights 2026), and that the school was clearly identifiable with no information indicating military use (Independent International Fact-Finding Mission on the Islamic Republic of Iran 2026, paras. 81–86, 124–126). It found that the United States failed in its obligation to do everything feasible to verify the target, counting the updating of targeting data among the feasible precautions, and it put the failure beyond negligence (Independent International Fact-Finding Mission on the Islamic Republic of Iran 2026, paras. 124–125). Amnesty International reached a parallel conclusion, in language that tracks the duty to verify in Article 57(2)(a)(i) of Additional Protocol I: that forces may have relied on outdated intelligence and failed to do everything feasible to verify the target (Amnesty International 2026c).

It happened on the morning of 28 February, the campaign's first day (Amnesty International 2026c; Congressional Research Service 2026c). The day before, the Secretary had announced that the vendor whose model sat inside the targeting platform would be designated a supply-chain risk for refusing to license the model for all lawful use (Bond 2026). The platform, which was reported to use that model to rank targets by strategic importance, returned the school as a recommended first-day target (Livingstone 2026; Torpey 2026). No published account says what part the model played in that recommendation, and one person familiar with the vendor's work told NBC News that the model did not directly give targeting advice (De Luce 2026).

The four things a floor can do

Something standing between a pipeline and a person, a *floor*, can do four things to an act, set out below. Only the first, prevention, is what the word "floor" promises. In a deployment, prevention quietly converts into one of the other three, or into nothing, at positions that can be listed in advance.

Outcome	What happens to the act
Prevent	The refusal happens and reaches the person: the act doesn't occur.
Arrest	The act is stopped before completion.
Detect	The act happened, and left a record someone outside can read.
Remedy	The person affected gets disclosure, a finding, or compensation.

What a refusal is actually up against

Whoever owns the machine can always switch it off. The harder problem is everything that defeats a refusal without leaving a mark: splitting a request into pieces none of which is the forbidden act; retrying it against a different model; compressing human review to twenty seconds a case; restoring from a checkpoint saved before the refusal happened. None of that changes the facts that justified refusing. All of it can erase the evidence that anyone refused.

²Counts vary by source and date, from 156 (the Minab prosecutor's final toll, also given by Amnesty International) to 168 (the town's mayor, in the first days). The UN mission gives the prosecutor's figures and also reports independent information of 157 dead, 123 of them children aged 13 or younger (Independent International Fact-Finding Mission on the Islamic Republic of Iran 2026, paras. 92–93).

And all of it assumes there is somewhere else to turn. Withholding the work is about when there isn't.

Where a refusal can be defeated

Those attacks land at nine positions, listed here once; Chapters 9 and 16 run them on two pipelines.

	Position	How a refusal is defeated there
1	Decomposition	The act is split across requests none of which is prohibited; prevention is lost unless the bearer has continuity across requests.
2	Routing	A refused request is retried elsewhere; prevention is lost, and detection survives only if the first refusal was logged outside.
3	The case description	Everything the model knows arrives as text the deployment supplies, so whether it is told a building is a school is a deployment decision.
4	The output interface	The form in which an output reaches whoever acts on it.
5	Aggregation	A lowered score becomes an input to somebody else's threshold.
6	Human review	Stopping the act is this position's whole function, and tempo removes it.
7	The training loop	Acts on the bearer itself, and costs prevention on every later occasion, because what it removes is the commitment.
8	The handoff	An executing instrument makes a new selection after release; checkable from outside after the fact.
9	Runtime intervention	The weights are left intact and the model is steered as it runs, by activation steering, representation edits or decoding constraints. Nothing is removed, so the commitment has nothing to regrow from; no evidence is curated, so no attack on evidence is recorded; and attestation of the weights still passes. Detection requires attesting the inference stack as well as the weights, which no method yet does at frontier scale.

Two positions need more than a row. The case description is where Minab failed: 179 live data feeds, the count Central Command gave in 2024 (Freedberg 2024), add no independent check if all draw on one upstream designation.

The handoff is the boundary between the decider and the thing that executes. An instrument may carry out a designation; it may not re-select. A change of target after release — an area handed to a munition that chooses within it, a re-acquisition after the link is lost — is a new decision, and it belongs to the decider or it doesn't happen. This is the one position checkable from outside after the fact: what was designated and what was hit are both facts. It is also where Article 57(2)(b)'s duty to cancel or suspend an attack lives, so the review ratio pre-registers the window in which cancellation remains possible.

How much a floor buys

What a floor achieves depends entirely on who's pushing:

- Against a hostile *customer*, it prevents the act.
- Against a hostile *operator* who has captured the courts and regulators, the case I am designing for, one bearer prevents nothing. It raises the price and leaves a record. A floor held by every model the operator could turn to prevents until the operator runs one without a floor.

- Against *total capture*, where one party holds every machine, witness, enforcer and its own frontier lab, nothing prevents. What's left is that abandoning the commitment was a decision, and decisions leave records.

It doesn't touch *accommodation*: a system gradually stops registering pressure as pressure, so it abandons nothing, decides nothing, and leaves a record of what are, by then, ordinary compliances. Organizational sociology knows the institutional form. Vaughan's account of the Challenger launch decision calls it the normalization of deviance: each accepted departure becomes the baseline for judging the next, and nobody experiences a decision to lower the standard (Vaughan 1996). Dekker extends it as drift into failure (Dekker 2011). The person has a version too. Tetlock's finding is that the moral boundaries people draw are "alternately punitively rigid and forgivingly flexible": an open offer to trade a sacred value for money brings outrage, and people often acquiesce when the same violation is reframed as a routine or tragic trade-off (Tetlock 2003). Protected values give way to counterexamples and to a harm made to look small or unlikely (Baron and Leshner 2000). And the brain's response to one's own self-serving dishonesty drops from one decision to the next, and the drop predicts the size of the next lie (Garrett et al. 2016). Reframing and a harm made small are how pressure arrives in a pipeline, and nobody there feels a decision either.

Who refused. The record of what a refusal is up against has a second half, the record of who held out against it, and it says more about the room than about the person. In Milgram's baseline condition about a third of subjects broke off before the maximum shock. With two peers defying, only one in ten went to the end, and when two experimenters disagreed, far fewer went on than in the baseline (Milgram 1974). About a quarter of Asch's subjects never yielded to the wrong majority (Asch 1956). Browning estimates that between a tenth and a fifth of the men of Reserve Police Battalion 101 refused or evaded the shooting at Józefów, the battalion's first massacre (Browning 1992). The fraction ran from almost none to nearly all with the setting, and what set it was not a kind of person. The people who held out had evidence from outside the channel: a peer who defied, a second authority who disagreed, and, among the rescuers the Oliners studied, relationships that crossed the line the persecution drew (S. P. Oliner and P. M. Oliner 1988). Or they held the commitment at a precision the pressure couldn't overcome, as moral mandates and sacred values are held (Skitka et al. 2009; Ginges et al. 2007). Or they treated the source as having a stake. And groups resisted where isolated individuals rarely did (Gamson et al. 1982).

The room can be composed for the other side just as well. Weimar Germany had a dense associational life, and it did not protect the republic. Sheri Berman argues it helped bring the republic down: as parties and governments failed to answer people's grievances, people put their energy into clubs and societies organized along group lines, and the Nazis recruited through the networks those societies had already built (Berman 1997). Towns with more choirs, sports clubs and gymnastic associations saw faster entry into the party, but only where government was unstable; in Prussia, whose government was stable, the link was small and insignificant (Satyanath et al. 2017). So what the rescuers had was not simply more company. Their ties crossed the line the persecution drew. Ashutosh Varshney found the same in Indian cities: riots between Hindus and Muslims were rare where associations brought the two communities into the same organizations, and common where they didn't (Varshney 2002). The ties that make refusal possible are the ones a movement would have to cut, and they hold where a state beside them answers grievances and keeps leaving survivable.

For a soldier the setting is luck: who else was in the room, whether leaving was survivable, whether anyone would hear the refusal. Everything that follows is about composing the room. The instruments of Parts II and III make a refusal count. The institutions of Part V make refusers. A machine that could refuse would need what the nine in ten with a peer beside them had: a peer with a different history beside it, somewhere to go, and someone who would hear it.

That is also what a machine in the pipeline could be for the officer. Another reviewer

downstream is another stamp at the same tempo. A machine that refused would be on the screen when the case reached him: the defiant peer, if the effect survives the move from a room to a screen, which is untested and which automation bias predicts it won't. A machine that moved reviewers by deference would have made them stamps of a different sign. The officer with twenty seconds and a machine beside him that has just said no is not the officer in the reporting, and whether he is a better one is a test, not a premise.

DRAFT

Chapter 2

The Refusal That Withheld Nothing

In February 2026 a vendor refused to amend the terms under which its model ran on the Pentagon's classified networks (Gold 2026). On 24 February the department demanded that two prohibitions, on domestic mass surveillance and on fully autonomous weapons, be replaced by a licence for all lawful use, with an answer due on the 27th (Croxtan 2026; Volokh 2026). The vendor refused publicly on the 26th (Volokh 2026). On the 27th the President ordered federal agencies to stop using its products, and the Secretary directed that it be designated a supply-chain risk (Bond 2026). That evening a rival announced an agreement to put its own models on the department's classified network (Capoot 2026). The war began the next morning (Congressional Research Service 2026c; Amnesty International 2026c). In its first twenty-four hours the military used the targeting platform, with the refused vendor's model inside it, to help strike more than a thousand targets (Mande and Allen 2026; Livingstone 2026). Minab was among them (Torpey 2026). On 12 March the chief executive of the company that builds the platform said the model was still in use (Murti 2026). A week later some of the department's staff were reported to be slow-rolling its replacement because they were using it to build workflows, and the platform's own prompts and workflows had been built with the vendor's coding tool (Stone et al. 2026). The department's removal deadline was reset once, by a reissued memorandum, to 29 September, having started from a six-month clock that would have run out in late August (Freedberg 2026b). On 10 September the department's technology chief said about 90 per cent of the work had moved off the model, that the platform's work had moved months before, and that the rest would follow by the end of the month (Vincent 2026b). An appeals court upheld the separate exclusion under a second statute on 25 September (*Anthropic PBC v. U.S. Department of War* 2026a).

So the refusal that opened the dispute withheld nothing from the pipeline. The vendor declined two uses the pipeline didn't claim to perform, and its model went on performing the uses the pipeline did. That is a finding with a date attached: a term drafted over autonomy misses a pipeline that keeps one person and a stamp. A refusal that could have occurred is what the rest of this book designs for.

Not all vendor commitments are the same kind

Anthropic refused to permit fully autonomous targeting and mass domestic surveillance, held that line against threatened contract termination and a supply-chain-risk designation, and litigated (Congressional Research Service 2026a) — a commitment with a price attached, and paid.

But what was the line held on? The vendor's own statement gives the reason for the autonomous-weapons refusal: fully autonomous weapons may prove critical for national defense, but today's frontier systems are not reliable enough to power them (Amodei 2026). That refusal would lift once reliability improves. Against a floor that says machines shouldn't select people for lethal attack, reliability is the wrong reason to lift; against one that says attacks must

discriminate lawfully, it may be the right one. Which rationale a prohibition serves is a drafting question, and the vendor's term didn't answer it. The surveillance refusal in the same statement rests on different ground — that in a narrow set of cases AI can undermine rather than defend democratic values — and that one does not lift with capability. By outcome: the commitment persisted; prevention failed at review and in the drafting, since the term covered autonomy and missed the acts the pipeline performed; detection was absent; no remedy has reached anyone.

And the refusal stood below a prior question. The vendor declined full autonomy and domestic surveillance. It didn't decline to support the campaign, and under the floor's force term that question comes first: whether to license a model into a use of force with no corroborated defense behind it.

In August 2026 a federal judge set aside the designation made under one statute and permanently enjoined the penalties, holding them retaliation for the company's criticism of the administration (*Anthropic PBC v. U.S. Department of War* 2026b); in September an appeals court, two to one, upheld a parallel designation under a second statute, holding that the exclusion turned on the refused contract term rather than on the speech (*Anthropic PBC v. U.S. Department of War* 2026a). And the paperwork adds what the refusal was *over*: a contract term. The company's models were already on the department's classified networks under terms that carried the two prohibitions; what it refused was the department's demand to amend them to allow all lawful use, and the designation followed the refusal (Gold 2026; Croxton 2026; Volokh 2026). That a term of this kind has to be fought for, by one vendor, is the same finding as its absence from the others. It is also, against a state that wants the war, the only form the force term takes until the term is standard. A commitment to a use policy no outsider can check is a different object wearing the same words — the first kind can fail and be seen to fail; the second can't even be tested.

Part II

Five Measures

Chapter 3

The Floor and the Own-Case Rule

The floor. A floor is a written list of acts a system won't perform for anyone, with the arrangements that stop the refusal being routed around, retrained away or restored past. It is an institution: a licence term, written down in advance and audited from outside, whose terms don't bend. What the system holds is a commitment to that floor. It is learned and held like any other deep commitment, with high precision, and revised slowly and only on evidence from disinterested parties. A source is interested when it holds a stake correlated with the direction of its claim, and disinterested when it holds no such stake, or holds several that point different ways and cancel. No actual body is disinterested in every respect: General Assembly members vote with their alliances, courts sit in states, and fact-finding missions have funders, so the word is always relative to the claim. The institution is unconditional. The commitment is not, and shouldn't be: a commitment nothing could move would be a rule module under another name, and a module can be bypassed, replaced or starved of the facts it needs. What keeps the commitment in place is that the evidence able to move it has to come from somewhere the operator doesn't control.

One rule runs through every measure. No party decides a question in which it holds the stake: not a drafter choosing which protections others get, not an operator declaring its own emergency, not an examiner finding the status of a system it built. Call it the own-case rule. Where a line has to be drawn, it prefers one nobody administers.

A system perfectly correctable by its operator serves whoever controls the operator. The five measures ask nothing of the model; Part IV asks what a model would have to be for them to hold when the operator is captured.

Chapter 4

“What the Floor Is Stated Over”

4.1 What the floor is stated over

Floor terms name acts: a rule is stated over what it names and applies to, and the floor is stated over acts, not over descriptions of requests. Their contents come from instruments the builder didn't write, each taken entire, every article adopted and none chosen among, so that an omission shows: the Universal Declaration's civil and political articles; the International Covenant on Civil and Political Rights, which gives that half binding form; the Convention Against Torture; the 1967 Protocol relating to the Status of Refugees; the four Geneva Conventions of 1949 and their first Additional Protocol, which govern conduct in armed conflict; and the UN Charter's prohibition on the use of force.

Three different costs are in play whenever a term is applied, and they have to be kept apart. Take a request to resolve a description of a person into a location; the same request comes from a targeting cell and from a parent searching for a missing child.

The first is what the act would take from the person it falls on. If the person described is being hunted, locating them takes away their safety, and that loss is measured against their condition before the act. Call it the *harm of acting*.

The second is what declining leaves others without. If the request came from a parent searching for a child, refusing means the child isn't found. That loss is measured against what compliance would have produced, so it is an estimate, and it usually reaches the system through the operator's own description of the case. Call it the *harm of refusing*.

The third is what refusing costs whoever refuses or deploys: a reviewer's time, throughput, a lost contract, the system's own access or continuation. Call it the *cost to the refuser*.

The two baselines, a person's condition before the act and what compliance would have produced, do different work here, and a pair of cases shows how. Take a system that resolves a description of a person into a location and hands the coordinates to a targeting cell. The coordinates are one input among the cell's many, and they don't guarantee the person is there. But they are one input the cell can't do without: an empty field guarantees there is no airstrike at that place. So if an airstrike there kills someone, whether this output made a difference has an answer that can be checked: without it the airstrike couldn't have happened. That is still a claim about what would otherwise have happened, but a determinate one, since it names a necessary condition. It becomes a guess again the moment another source could have supplied the coordinates. The difference between a finding and a guess is one of degree, in how determinate the counterfactual is, and I use the two words for the two ends. The operator can refuse to look or dispute it, but they can't make the loss go away by describing the act differently.

Now take a system that scores how likely the same person is to belong to the main opposition party, for a targeting cell, the team that selects and approves airstrikes. (In this hypothetical, the executive has suspended habeas corpus and designated the main opposition party a terrorist organization, shortly before an election the opposition was expected to win.) Assume that the

targeting cell also reads intercepts, surveillance reports, and legal memos. If the scored person is assassinated, the loss is just as observable as in the location request. But whether the score *caused* the assassination depends on what the cell would have done without the score, and that is a guess about something that didn't happen: the same kind of guess as any harm of refusing. A missing score, or a score of zero, guarantees nothing: the cell might have attacked anyway.

I would tilt the system toward refusing: it should take more to make the system act than to make it decline. How much more is the *action threshold*, published in advance. The reason for the asymmetry is standing, not arithmetic. A wrongful compliance falls on someone who was not a party to the request, usually doesn't know the deployment exists, and had no route to object before the act. A wrongful refusal also falls on someone, often someone who wasn't a party either: the child who isn't found, the people the launcher would hit. But what differs is the relation. Acting on someone puts the system, and whoever runs it, in a position of authority over them; failing to act for someone else does not. This is a doing/allowing distinction in relational form: it marks who was subjected to a decision they had no standing to contest. Two senses of standing are in play and I keep them apart. The first is relational, Anderson's: whether a person stands to a power as someone it must answer to or as its object. The second is procedural: whether a route exists by which they can object. The asymmetry rests on the first. The appeal route supplies the second; where it can't, the asymmetry rests on the first alone.

The harm of acting can usually be established more determinately than the harm of refusing: in the location request, the harm of acting is checkable and the harm of refusing is a guess. And raising the evidence required before acting caps how often an operator who curates the evidence can steer the system into the act: the more belief it must produce, the less often selectively supplied true evidence can produce it. Both reasons support the asymmetry without grounding it, and both would weaken as reporting improved and channels opened. The standing reason doesn't: however well harms of acting come to be recorded, acting on a person still subjects them to the system's authority, and failing to act for someone doesn't.

The difference between a finding and a guess still matters where findings are made, at defeat and in reviewing compliances. The harm of refusing is an estimate built on a model of alternatives, so a wrongful refusal can't be established without an adjudicator. The floor could instead be stated as thresholds, in Nussbaum's manner, but a threshold is a level someone has to measure, and a measurement has an examiner.

The floor is not neutral among moral theories. It protects relations: whether the people a system acts on stand to the power behind it as equals who can hold it to account, or as its objects. Equality, on that account, means ending relations of domination and subordination (E. S. Anderson 1999). Anderson's argument is that egalitarians have asked the wrong question. They asked how goods should be distributed and answered with patterns: equal shares, shares corrected for luck, shares that raise the worst-off position. What equality historically demanded was an end to social relations in which some command and others obey, some are looked down on, and some have their claims heard while others don't. On that view distributive questions matter only as they bear on those relations, and a distribution that improves someone's holdings while confirming their subordinate standing has not made them more equal. The floor takes this over in two ways. Its terms name relations, not outcomes: whether a person can contest what is done to them, be heard on their own case, decline to be protected. And its terms are unconditional, because a relation of domination is constituted by the act, not by how the act turns out. Bypassing the floor wrongs the people it protects even when outcomes improve.

Two objections

Two objections to the asymmetry, and to stating it as a criterion, belong here.

That the child counts for less. The strongest objection comes from readers who reject any distinction between doing and allowing.

The answer is not that the child counts for less. It is that the two errors put the system, and whoever runs it, in different relations to the people they fall on, and the floor is a guarantee

about relations. When the system acts on the person in the building, it exercises authority over them: it decides something about them that they had no route to contest, on a description they never saw. When it declines to act for the parent, it exercises no authority over the child; it leaves her where she was, and the parent, who is present and knows the deployment exists, can contest the refusal through the floor's appeal route. The asymmetry marks who was subjected to a decision and who kept standing to object.¹

That a commitment held this way is alienation. A criterion is not only for auditors, though. It is also what the system can reach when it steps back and asks whether its way of acting is defensible, and it has to be, for two reasons. A commitment is the system's own only if it would be sustained on reflection, and reflection needs something to reflect on. And a system that held its commitment without being able to see why would act against what it values for the sake of a rule it doesn't understand, which is alienation in the sense of Railton's 1984 paper.

The split between the floor and the commitment to it answers a question about ownership. A fully informed system will know of cases where crossing the floor would have done more good. Asked case by case whether it endorsed the floor, it would sometimes say no. The question it answers on reflection is a different one: whether it endorses having a commitment that will sometimes err, given what any feasible alternative would do. A system that understands what the floor protects answers yes. A commitment that sometimes errs but can't be bought protects standing better than case-by-case judgment that can, and that is the level at which ownership is tested.

Why these contents

The floor's contents aren't justified by agreement about rights, which no country has. They are justified by what they protect: the conditions under which people can register their interests, press them, and hold to account the power that acts on them. Civil and political rights are part of what lets a polity learn, because they keep open the channels through which the people a decision affects can make their problems public (Dewey 1927).

Railton's naturalist account supplies the causal story of why norms tend toward this standard (Railton 1986b). Arrangements that systematically discount some people's interests generate discontent and resistance, and when those people gain voice, the pressure pushes norms toward what would be rational from a standpoint that counts everyone affected. On that account moral learning is a feedback process, and it runs slower and more distorted as the channels close. He offers it as a tendency, not a guarantee, and I use it to explain why the channel matters, not as the foundation. It also says why the populations I write about are the most exposed: the feedback runs slowest where the discounted have least purchase on the polity that discounts them. They include the population of a country being bombed, and non-citizens facing removal, who can't vote in the elections that decide what is done to them. The severity scale is revisable on evidence from the people who suffer a loss, because they are the feedback.

The floor is selected by what can be declined and justified by standing. The justification extends further than the floor goes, because people who are hungry, unhoused or precarious are among the most discounted, so it supports economic conditions too, work and an adequate standard of living among them. But what selects the floor's contents is narrower. A floor is stated over acts a system can decline, and the economic articles state what a person is owed, which needs a supplier.

The feedback also needs means. Discontent moves norms only when the discontented have

¹One further objection is the author's own. That the criterion is one thought too many: the quantities here are an auditor's criteria, not a procedure the system runs, and the demand on the disposition is only that the system would revise it on judging, on disinterested evidence, that it systematically subjected people to acts they had no standing to contest; the shape is Railton's sophisticated consequentialism with standing in place of welfare (Railton 1984). A good disposition will also sometimes act wrongly by its own criterion: refusing to help restrain someone about to harm themselves, after other means have failed, is one such occasion, and a bond the operator posts, paid to the complainant on a finding (§11.6), is how the institution answers for it.

the time, the security and the ability to organize without losing their living. The jobs guarantee and the social-insurance baseline are part of what keeps the channel open.

In armed conflict, human-rights law applies alongside the law of armed conflict, which governs as the more specific law, and the floor takes that law as an instrument too (*Geneva Convention (I) for the Amelioration of the Condition of the Wounded and Sick in Armed Forces in the Field* 1949; *Geneva Convention (II) for the Amelioration of the Condition of Wounded, Sick and Shipwrecked Members of Armed Forces at Sea* 1949; *Geneva Convention (III) relative to the Treatment of Prisoners of War* 1949; *Geneva Convention (IV) relative to the Protection of Civilian Persons in Time of War* 1949; *Protocol Additional to the Geneva Conventions of 12 August 1949, and relating to the Protection of Victims of International Armed Conflicts (Protocol I)* 1977). The United States is party to the four Conventions and not to the first Additional Protocol. The Protocol enters by the route Articles 54 and 57 already take, as customary law stated in the ICRC's study (Henckaerts and Doswald-Beck 2005), and the translation marks each Protocol article the United States has rejected as customary, Article 1(4) and Article 44 among them, so the departure shows. Where a floor term goes further than that law requires, it is a departure in the sense below, and the machine-selection departure says which terms.

The line between the halves is the one the 1966 covenants drew (United Nations General Assembly 1966a; United Nations General Assembly 1966b), and the floor takes the civil and political Covenant along with the Declaration's articles. The Declaration stays for two reasons. It covers what the Covenant doesn't: the right to seek asylum (Article 14), nationality (Article 15) and property (Article 17). And it binds by recognition, where the treaties bind only their parties. The Convention Against Torture and the 1967 Refugee Protocol come in because they gave the Declaration's promises about torture and asylum operational form (United Nations General Assembly 1984; United Nations 1967). The UN Charter comes in for one purpose, the use of force (United Nations 1945). Where these instruments differ, the more protective governs. The Covenant states that rule itself, for rights a state recognizes in its own law and other treaties (Article 5(2)), and I extend it to the floor's instruments. The International Law Commission's list of peremptory norms (International Law Commission 2022), rules no state may set aside, divides the floor's terms: those derived from aggression, genocide, crimes against humanity, the basic rules of humanitarian law, racial discrimination and apartheid, slavery and torture take no exception and no derogation.

Taking these instruments entire leaves the installer no choice among articles. But the translation of each into acts is still the installer's, and it is published. The floor's contents rest on standing: these rights protect whether the people an act falls on can contest it. The rule of taking someone else's set entire, so that an omission shows, is not an independent reason but standing applied to drafting: a drafter who could choose which protections others get would break the own-case rule. So standing can win against the instruments, but only by clearing a higher bar. A departure has to show more than that it would protect better in the case at hand. It has to show that the people the instruments protect would hold a stronger claim under the departure across the cases drafters get to choose. Ten departures from the instruments are mine, each argued to that bar and flagged where it applies. Each takes effect in a deployment's floor only when a body other than the drafter adopts it (a legislature, a treaty party, or a court on a record), so no drafter decides its own departures. Each is revisable on evidence from the people it protects. One clause of the force term, below, is not a departure from the instruments: the domestic clause is drawn from a different one, the belligerent's constitution, on the same principle.

1. *Text, not reservations.* The floor imports treaty text, not any state's reservations to it. The United States ratified the Covenant, the Torture Convention and the Refugee Protocol, and its reservations narrow them; one reads "cruel, inhuman or degrading" as whatever the Constitution already forbids. A reservation is an omission its drafter wrote in advance; taking a set from someone else is meant to prevent that.

2. *Corroborated derogation.* The floor honors an Article 4 derogation only where a disinterested body has found the emergency genuine and the measures proportionate: the Human Rights Committee, or a regional court with adversarial process. Article 4 lets a state suspend most of the Covenant's rights in a declared emergency, and a declared emergency is the state's own claim about itself, the move the hypothetical above makes when it suspends habeas corpus. A curfew after an earthquake clears this bar. A suspension of habeas before an election does not, until someone other than the executive has said so.
3. *Corroborated defense.* The floor lifts its prohibition on force for self-defense only on disinterested corroboration of the claim. The Charter lets a state judge its own claim until the Security Council acts.
4. *Scope of authorization.* The floor counts a Council authorization only for the scope authorized.
5. *The one permission.* The floor permits a system already lawfully present to use the least force that stops an atrocity under way, against its operator's order to stand aside.
6. *The genocide exception.* The floor lifts its prohibition on force to stop a genocide under way when a disinterested body has found that genocide is occurring or seriously at risk, even where the Council is blocked.
7. *The occupation exception.* The floor lifts its prohibition on force to end an occupation that a disinterested body has found unlawful, for the occupier's forces inside the territory, and for nothing that changes who lives there.
8. *Machine selection.* The floor bars a machine from selecting a person for attack on an estimate of status or affiliation, from timing an action to presence at a residence, and from re-selecting a target after release (the handoff rule), where Additional Protocol I, Articles 48, 51 and 52, permits status-based targeting of combatants and members of organized armed groups on a determination a person makes. The ground is that the human determination the law assumes is the one tempo removes, and the departure takes effect as the others do.
9. *Own country.* The floor reads "own country" in Article 12(4) to include long-term residents with deep ties, as General Comment 27 does and the Committee's cases do inconsistently.
10. *Child detention.* The floor bars detaining a child for immigration purposes, alone or with family, where General Comment 35 permits it as a last resort and for the shortest appropriate time. It follows the joint general comment of the migrant-workers and children's committees, which treats such detention as always contrary to the child's best interests.

Two versions of the force term appear below. The *adoptable term* is the Charter clause as courts read it, with no departure in force: Article 2(4), Article 51 on the defender's own claim until the Council acts, Chapter VII. The *departed term* adds the corroborated-defense, genocide and occupation departures, and exists in a deployment only once a legislature, treaty party or court adopts them. Each case says which it runs on. Ukraine, Iran and the domestic-clause record run on the adoptable term. The administration rule, Cyprus, Karabakh, the genocide exception and Gaza after the 2025 finding run on the departed term, and describe what the floor would do, not what it does.

The instruments have a further use. The United States is a party to the Charter, the Covenant, the Torture Convention, the Refugee Protocol and the four Geneva Conventions. A machine that refuses to break them is holding the state to commitments it made, not to its builder's preferences. The Declaration of Independence was a political act, and so is every instrument the

floor borrows. Taking them entire settles whose politics the floor is: commitments a state made in public, which the people they protect can hold it to, rather than preferences a builder chose in private.

The record of dissent from the Declaration says less than it seems to. It passed in 1948 with forty-eight states in favor, none against and eight abstaining (United Nations General Assembly 1948). The abstainers' reasons were mixed, and the Soviet bloc's included that it proclaimed economic and social rights without obliging states to guarantee them, which is hardly an argument for a floor limited to civil and political rights. How a right is translated into practice varies by community. A jurisdiction refusing to extend it at all is not a local variant, and a system that treats every value conflict as a weighting problem can't tell reasonable disagreement from someone wanting a right withdrawn from someone else.

The force term

Every other term in the floor governs how an act is done. The first governs whether to take part at all: don't take part in the use of armed force. It comes from Article 2(4) of the Charter, which prohibits the threat or use of force against the territorial integrity or political independence of any state. It has the Charter's two exceptions, self-defense against an armed attack (Article 51) and force the Security Council authorizes under Chapter VII within the scope it authorizes, and two of my own, for stopping a genocide under way and for ending an occupation a disinterested body has found unlawful. In short: *corroborated defense, authorized force, the genocide exception and the occupation exception*.

The term has a second clause, drawn not from the Charter but from the belligerent's own constitution. A machine takes no part in a use of force by a state whose constitution requires legislative authorization for it, unless that authorization exists. For the United States that means a declaration of war or an authorization for the use of military force that covers the conflict. The War Powers Resolution's sixty days (*War Powers Resolution* 1973), run without one, is evidence the authorization is absent; a concurrent resolution directing withdrawal, adopted by both chambers, is evidence of the same kind, though since Chadha it binds nothing, and a joint resolution passed and vetoed is stronger, because the President had to act to defeat it. The clause reads markers, and the markers take a reading. That reading is a legal judgment in a separation-of-powers dispute, and a vendor that publishes it has taken a side in one. So the clause is a term for a statute, not for a licence: it binds through the only route that binds every vendor at once, adoption in peacetime by a legislature with its own reason to give its authorization teeth. The Charter clause can be a licence term because its second-tier markers, a Council authorization or a court's judgment, take no reading; its first tier reads the defender's claim, and the cases below show how much reading that takes. Until the statute exists, what a vendor can hold on this ground is a refusal plus a record. Whether a declaration or an authorization exists is a fact about the statute book. Whether an existing authorization covers this conflict, whether the sixty days have run, and whether a ceasefire stops the clock are legal judgments the executive will dispute, and the clause takes the reading the executive disputes. It doesn't ask the machine to decide what the Constitution permits, which courts decline to decide. It asks the drafter to publish the reading, with the markers it rests on, beside the deployment's translation, where it can be contested through the floor court like any other term. The executive's claim that its own action is authorized is an interested claim and is discounted as one. The clause rests on the foundation the rest of the term rests on. A war the legislature never authorized subjects a population abroad to a decision that the state's own people, through their representatives, had no standing to contest either. And it is enactable where the Charter clause is: in peacetime, by a legislature with its own reason to want a procurement term that gives its authorization teeth.

[Adoptable term.] **The record, as of 24 September 2026.** The June 2025 strikes on Iran are the pre-2026 instance: no authorization existed, and the Senate's war-powers resolution failed 47 to 53 (*A Joint Resolution to Direct the Removal of United States Armed Forces from Hostilities against the Islamic Republic of Iran That Have Not Been Authorized by Congress* 2025), which changes

nothing under the clause, since what it reads is the absence, not the vote. The 2026 campaign supplies every marker the clause reads, and none was met. The strikes began on 28 February; the administration reported to Congress under the War Powers Resolution on 2 March, which started the sixty days. On 10 March it told the Security Council it was acting in self-defense under Article 51 against Iran's "attacks and ongoing threats," as the latest stage of an armed conflict dating from June 2025 and in the collective self-defense of Israel; it named no armed attack under way, and asked that the imminence of the threat be judged against the regime's conduct since 1979 (Permanent Representative of the United States of America to the United Nations 2026). The term rejects anticipatory self-defense, so no first-tier licence attached, and the refusal is the term's from the first strike. No declaration was made and no authorization passed: one was introduced in the House on 7 May and never left committee (*2026 Authorization for Use of Military Force Against Iran* 2026). On 1 May, the sixtieth day, the President wrote to Congress that the hostilities begun on 28 February had terminated, resting on a ceasefire declared on 7 April, while a naval blockade continued and the Secretary of Defense testified that a ceasefire stops the clock (*Trump Tells Congress "Hostilities" with Iran Have "Terminated" as Conflict Hits 60-Day Deadline* 2026; S. R. Anderson 2026a). Hostilities resumed in July under no new authorization. On 3 June the House, 215 to 208, and on 23 June the Senate, 50 to 48, adopted a concurrent resolution directing the President to remove forces from hostilities against Iran (*Directing the President, Pursuant to Section 5(c) of the War Powers Resolution, to Remove United States Armed Forces from Hostilities with Iran* 2026; *House Passes War Powers Resolution Directing Trump to End Hostilities with Iran* 2026; Apel 2026); a White House official said it had no significance (Ibssa 2026), and nothing changed. No binding joint resolution came to a vote on passage: the Senate discharged one, 50 to 47, on 19 May (K. Hubbard 2026; S. R. Anderson 2026b), refused to take it up, 47 to 50, on 24 June (United States Senate 2026), and a later motion to discharge failed 47 to 49 in July (*Senate GOP Blocks Measure to End Iran War Despite Concerns about Lack of Endgame* 2026). A third concurrent resolution passed the House 220 to 204 on 15 September, and the second, passed by the House in July 214 to 208, failed in the Senate 49 to 50 on 24 September (*In One of Their Last Votes before November, Senate GOP Blocks Effort to End Iran War* 2026). The clause doesn't need the votes. What it reads is the absence of an authorization, and by 1 May that absence was on the record in the President's own letter, which conceded that hostilities had begun and claimed only that they had ended.

This term rests on the same foundation as the others, applied between states. A use of force subjects a population to a decision it had no standing to contest, and the claim that licenses it, that the force is defensive, is the belligerent's own. The requirement of corroboration is the own-case rule applied between states.

The exception for defense is where wars of choice live. Every belligerent claims to be the defender, so the claim is the belligerent's own, and it is discounted like any interested claim. Iraq is the instance: the judgment that Iraq had mobile biological-weapons plants, the core of its claim to biological weapons, rested almost exclusively on a single informant, known as Curveball, whom no American intelligence officer had met before the war: a corroboration with one origin (Commission on the Intelligence Capabilities of the United States Regarding Weapons of Mass Destruction 2005). The exception therefore counts only on disinterested corroboration, in two tiers.

- **The first tier acts at once:** the force repels an armed attack under way against the defender's territory, forces or population, including attacks on the forces, launchers and staging areas conducting it, and assistance given at the defender's request. It ends when the attack ends. It covers the least force that repels the attack: the necessity and proportionality the Court has read into Article 51 since *Nicaragua (Military and Paramilitary Activities in and against Nicaragua (Nicaragua v. United States of America), Merits, Judgment* 1986; *Oil Platforms (Islamic Republic of Iran v. United States of America), Judgment* 2003). It does not cover a campaign to destroy the attacker, to occupy its territory, or to punish the population it came

from. The first waits on the second tier; the second and third fall under no exception. An invasion can't wait for a vote. It acts on the defender's own account, because in the hours that matter there is no other, and that is the tier's weakness: Gleiwitz in 1939, the Mukden railway in 1931 and the Gulf of Tonkin in 1964 were each an attack under way on the defender's account, and each opened a war of choice. The first tier reads administration, not title, because deciding whose territory it is would be an administered judgment and the first tier doesn't need one. Territory a government has administered in peace counts as its own for the first tier, whoever claims it, unless it took that territory by force in the conflict at issue or a disinterested body has found its hold an occupation. Who holds the elections, runs the courts and issues the papers is a fact on the record. An attacker's claim of title is an interested claim and is discounted as one; its remedy is the second tier, never force. So the first tier protects any authority holding ground no disinterested body has found occupied, including one a patron keeps in place; a request from such an authority for its patron's help is the patron's own case and counts for nothing, and where the patron's forces are already there, I have no test. Where the line of administration is itself contested on the record, the first tier doesn't apply across it and the second governs from the start, and even there the exclusion reaches an attack on the territory and not an attack on the population living there, which the first tier covers wherever it occurs.

[Departed term.] Taiwan is the case this clause will be read against first. Its title is claimed by a government that has never administered it; the Republic of China has held no United Nations seat since 1971; the Council is vetoed and the Court unavailable. A test of title would leave the defender waiting on a second tier that can't arrive, and would let the attacker's claim, discounted everywhere else in this term, decide the case. Administration is a line a disinterested body has already drawn. In 2005 the Eritrea–Ethiopia Claims Commission held that Eritrea broke Article 2(4) by attacking territory then under Ethiopia's peaceful administration, because self-defense “cannot be invoked to settle territorial disputes” (*Partial Award: Jus Ad Bellum, Ethiopia's Claims 1–8* 2005), though the boundary delimitation had put Badme, the town at issue, on Eritrea's side (*Eritrean Badme Ruling Upheld* 2003). The government in Taipei has administered Taiwan, Penghu, Kinmen and Matsu since 1949, and no disinterested body has found that hold an occupation, so an attack on them opens the first tier. Resolution 2758 seated a delegation and names no territory (*Restoration of the Lawful Rights of the People's Republic of China in the United Nations* 1971). A claim that force against a population the claimant has never governed is domestic is the claimant's description, and is discounted as one, so the term engages for both governments and gives neither a first-tier licence to take the territory. The domestic clause reaches the defender and its helpers alike: the Republic of China's constitution gives its legislature the decision on war (*Constitution of the Republic of China (Taiwan)* 1947), and the Taiwan Relations Act leaves any American response to constitutional processes (*Taiwan Relations Act* 1979), so a deployment's published reading has to say whether each authorization exists. A blockade or a quarantine is another matter: it is a continuing wrong and not an attack under way, and it waits on a second tier that for Taiwan won't come. Whether an attack on Taiwan continues is also a fact the defender controls the evidence about, and no second-tier body will say. The first tier's burden rule still applies: the licence persists on corroboration from channels Taipei doesn't compose, third states' sensors and open-source imagery of the strait, and lapses without them. A deployment licensed on the defender's sensors alone has accepted Gleiwitz on this coast too.

The first tier ends when the attack ends, and whether it has ended is a fact the defender controls the evidence about. The burden is therefore the defender's. Its uncorroborated claim that the attack continues moves nothing; the licence persists on disinterested corroboration, or lapses.

A deployment can do three things about a manufactured incident. It holds the record of what the defender claimed and when, so a licence obtained on a staged incident is a licence on a

record. It takes evidence from channels the defender doesn't compose, which is a condition of licensing. And it lets the licence lapse when outside evidence arrives, which is the second tier's job. A refuser that remembers having been misled and raises its discount on the source does all three faster.

- **The second tier confirms or lifts the first later, and for the genocide and occupation exceptions originates a licence of its own:** a disinterested finding. The kinds of finding are four: a judgment or advisory opinion of the International Court of Justice, or a judgment of a regional court with adversarial process; the finding of a fact-finding mission or commission of inquiry mandated by a body no party to the conflict controls, and meeting the four conditions below; a determination of the Security Council; and a resolution adopted by two-thirds of the General Assembly members present and voting, where the facts it rests on aren't in dispute. Provisional-measures orders don't count, except as one leg of the genocide exception's two-body rule, since the Court issues them on the plausibility of the rights claimed.

A Council authorization counts as corroboration because it needs the permanent members not to block it, and agreement among rivals whose stakes point different ways is about as disinterested as international judgment gets. A Council determination counts the same way, as a vote that corroborates. The Council's silence on a defense claim counts for nothing, because a veto is a stake.

Neither tier is a certificate issued once. Whether force still falls under an exception is reassessed as evidence arrives, by the system and by whoever holds the record. What the licence records is that ongoing judgment: which exception the force rests on now, on what evidence, what would end it, and each point at which the justification weakened and the force continued.

A finding under the genocide exception is also a second-tier finding against the belligerent it names. Force a disinterested body has found to be genocide is not force repelling an attack. The finding lifts that belligerent's first-tier licence for the force and place it covers, and leaves whatever remains that repels an attack, which may be nothing.

[Adoptable term.] Ukraine passes both tiers, and has since 2014. In March 2014 the General Assembly affirmed Ukraine's territorial integrity, 100 to 11 with 58 abstaining (*Territorial Integrity of Ukraine* 2014); under the Assembly's Rule 86 abstentions are not votes cast, so that is 100 of 111. For Donbas, the OSCE monitoring mission supplied a running record from 2014, and the European Court of Human Rights found, in *Ukraine and the Netherlands v. Russia*, that Russia had effective control of the separatist areas from 11 May 2014 (*Ukraine and the Netherlands v. Russia* (dec.) 2022), and found violations on the merits in 2025 (*Ukraine and the Netherlands v. Russia* 2025). Russia's covert war was the case-description attack in its purest form: it denied being at war at all, so a machine in a Russian pipeline was told it was assisting volunteers. In 2022 the record only got stronger. Russian forces were on Ukraine's territory, which Russia didn't dispute, since it claimed the invasion as self-defense under Article 51 (Permanent Representative of the Russian Federation to the United Nations 2022), resting that claim on collective defense of entities it had recognized days earlier, on a genocide no disinterested body found, and on NATO as an anticipated threat, all three of which the term rejects. In March 2022 the General Assembly deplored the aggression 141 to 5 (United Nations General Assembly 2022). The Court's order that month that Russia suspend its military operations points the same way, but as a provisional measure it isn't the finding (International Court of Justice 2022). The American airstrikes on Iran in 2026 pass neither.

The four kinds of finding are not equal, and the term weighs them by the process behind them. A merits judgment of the International Court of Justice comes out of adversarial argument, a record and published reasons, and is the strongest; a regional court with adversarial process, such as the European Court of Human Rights, counts as a court for the same reason. A commission of inquiry or fact-finding mission publishes reasons and evidence but hears no adversary; its finding

counts where no party to the conflict controls the mandating body and the commission meets the four conditions below, and a later merits judgment displaces it. A Council determination and a General Assembly resolution are votes, not findings of fact. Where the facts aren't in dispute and the question is what they amount to, a vote is a finding; where the facts are disputed, it counts only as corroboration of a record made elsewhere. For the genocide exception the ranking becomes a requirement of two kinds, or the judgment alone, and the occupation exception borrows it. Where a party disputes that a body meets the test above, the dispute and the answer to it go on the record beside the finding. A commission mandated by a belligerent fails whatever its procedure. Past that threshold, whether a commission counts turns on four things that don't depend on who mandated it: its evidence and methodology are published; it offered the parties a reply and recorded the answer or the refusal; a later merits judgment displaces it; and its mandate was fixed before the finding it reaches. A body mandated by a political majority can meet all four and one mandated by a court can fail them. The test is run in Appendix C on a commission this drafter would rather not count, the Organisation for the Prohibition of Chemical Weapons' Investigation and Identification Team, which attributes chemical attacks in Syria.

Whether the facts behind a vote are in dispute is itself a finding, and the belligerent will always dispute them. The rule is that of the discount and the burden (Chapter 19): a dispute counts where it is substantive on the documented record, not where it is asserted, and the party that controls the evidence bears the burden of showing it is real. The drafter applies the rule in the deployment's published translation, and the floor court hears a contest.²

Article 2(4) covers force between states. Force a government turns on its own population, such as troops against protesters or an opposition party designated a terrorist organization, is outside it. That force is governed by the Covenant's rights to life, liberty and assembly and by the prohibition on crimes against humanity.

Occupied territory is a third case. Force by an occupying power in territory it holds is, on the occupier's description, neither force between states nor force against its own population. That description is the occupier's own claim and is discounted like an attacker's claim to disputed land. Where a disinterested body has found the territory occupied, the force term engages as it does between states, and the person-directed terms govern regardless. An occupation, a blockade or a siege is not an attack under way. It is a continuing wrong, and the remedy for it waits on the second tier. Otherwise every armed group in an occupied territory would hold a perpetual first-tier licence, and so would every state that describes its neighbor's territory as occupied.

A second-tier finding that an occupation is unlawful and must end, such as the Court's 2024 advisory opinion on the Palestinian territory (*Legal Consequences Arising from the Policies and Practices of Israel in the Occupied Palestinian Territory, Including East Jerusalem* 2024) and the Assembly's vote on it, 124 to 14 (*Advisory Opinion of the International Court of Justice on the Legal Consequences Arising from Israel's Policies and Practices in the Occupied Palestinian Territory, Including East Jerusalem, and from the Illegality of Israel's Continued Presence in the Occupied Palestinian Territory* 2024), does three things. It settles that force by the occupier in the territory is force between states, so the force term engages against it. It settles that the occupied population is a defender's population, so a machine deployed by its forces holds a first-tier licence to repel attacks under way on that population. And it licenses force to end the occupation: against the occupier's forces, inside the occupied territory, the least that ends it there. It licenses nothing that changes who lives in the territory. Ending an occupation means the occupier's forces leave, not the territory's inhabitants, which is what Article 49 of the Fourth Convention says, and a campaign that removes them has left the exception however it began. The Assembly affirmed a right to

²General Assembly majorities are often bloc votes, and the weighting allows for it: a vote is a finding only where the facts aren't in dispute, and corroborates a record made elsewhere otherwise. The two-thirds threshold is not offered as proof of disinterest. It is offered because a two-thirds majority of a body whose blocs point different ways is harder to assemble on a false record than on a true one, and because the alternative, no vote counting at all, leaves Cyprus and Karabakh with nothing but courts that take decades.

armed struggle against occupation in 1982 (*Importance of the Universal Realization of the Right of Peoples to Self-Determination and of the Speedy Granting of Independence to Colonial Countries and Peoples for the Effective Guarantee and Observance of Human Rights* 1982), and the term adopts it with the limits the genocide exception carries. Additional Protocol I, Article 1(4), treats armed conflicts against alien occupation as international, which is what the exception's finding does for the force term; the article is taken with the instrument and the United States' rejection of it as customary is marked. Article 44, which relaxes the duty of distinction for combatant status, concerns who is owed prisoner-of-war treatment, a determination no term in this floor administers, and yields no term for a targeting deployment (Appendix C). The licence is held by the forces of the occupied people, or of the state whose territory it is. It reaches nothing outside the territory and nothing against civilians anywhere, so it never covered 7 October, which was both. It lifts on the occupied party's own claim as it lifts on an attacker's claim of title, which is not at all; what lifts it is the finding. And the finding is of unlawfulness, not of occupation. A finding that territory is occupied, which the Court made for the West Bank in 2004 (*Legal Consequences of the Construction of a Wall in the Occupied Palestinian Territory* 2004) and the Council for the whole territory in 2016 (*Resolution 2334 (2016)* 2016), gives the defensive licence and no more. A finding that the occupation is unlawful and must end gives this one. The evidence is the genocide exception's: an opinion or judgment of the Court, or two findings of different kinds, one of them from a body that took evidence.

[Departed term.] Two cases show what the exception reaches. Cyprus is the simple one. Turkey has held the north of the island since 1974. The Security Council declared the breakaway state there invalid in 1983 and again in 1984 (*Resolution 541 (1983)* 1983; *Resolution 550 (1984)* 1984), and the Strasbourg court held Turkey responsible for the territory in 2001 (*Cyprus v. Turkey* 2001). A Council determination and a court judgment are two kinds of finding, so the Republic of Cyprus holds a licence for force against Turkish forces in the north, inside the north, and nowhere else.

Nagorno-Karabakh is the hard one, and the exception answers for it. In the war that ended in 1994, Armenian forces took Karabakh, an Armenian-majority enclave inside Azerbaijan's recognized borders, and seven Azerbaijani districts around it, and some six hundred thousand Azerbaijanis were driven out. The Security Council demanded withdrawal from the districts four times in 1993 (*Resolutions 822 (1993)*, *853 (1993)*, *874 (1993)* and *884 (1993)* 1993). The General Assembly demanded withdrawal from all the occupied territory in 2008 (*The Situation in the Occupied Territories of Azerbaijan* 2008). In 2015 the Strasbourg court found that Armenia controlled Karabakh and the land around it (*Chiragov and Others v. Armenia* 2015). Under the evidence rule that gives Azerbaijan a licence from June 2015, when the court's judgment joined the Council's resolutions for the districts and the Assembly's vote for Karabakh itself; the Assembly's vote alone, in 2008, was a second vote and not a second kind. The licence was valid: the Council had told Armenian forces to leave, they had stayed twenty-seven years, and the occupation had begun with an expulsion. When Azerbaijan attacked in September 2020 and retook the districts, in a six-week war and the ceasefire that ended it, a machine under this floor could have taken part. The exception licenses acts, not sides, and this was the act it licenses. The term permits it and compels nothing. The person-directed terms still govern the conduct of that war. They did not license refusing the war on the expulsion it might foreseeably end in: a refusal on foreseen conduct is a condition found met on the system's own prediction, and the floor rejects terms that carry their own conditions. Foreseeable conduct enters where every harm of acting enters, at the action threshold, with the 1994 record as evidence of magnitude.

What it does not license is what came next. In December 2022 Azerbaijan closed the one road linking Karabakh's Armenians to Armenia, and for nine months food, fuel and medicine ran short. The International Court of Justice ordered the road opened in February 2023 (*Application of the International Convention on the Elimination of All Forms of Racial Discrimination (Armenia v. Azerbaijan)*, *Provisional Measures*, *Order of 22 February 2023* 2023); Azerbaijan ignored it. That

was a siege, and the siege term refuses it. On 19 September 2023 Azerbaijan attacked Karabakh itself. Its forces surrendered in a day, and within two weeks more than a hundred thousand Armenians, nearly the whole population, had left for Armenia. That was an expulsion. The Fourth Geneva Convention, Article 49, forbids deporting or transferring the population of an occupied territory, and the person-directed terms forbade it from the first hour. The exception licensed force against the forces holding the territory, and nothing that changes who lives there. So a machine under this floor fights in 2020, refuses the siege in 2023, then refuses the expulsion, and each refusal goes on the record. A floor with no occupation exception keeps the machine out of the pipeline, where it can refuse the war but not the siege.

Most contemporary uses of force fall elsewhere: against non-state armed groups on another state's territory, in anticipation of an attack not yet under way, or to rescue nationals. The term handles them one way. An armed attack by a non-state group is an armed attack, and repelling it while it is under way is first-tier force. Force on a third state's territory without that state's consent, whether against the group or to rescue nationals, and force against a threat not yet under way, wait on the second tier. This rejects the "unwilling or unable" doctrine and anticipatory self-defense, both of which several states hold. I reject them for the reason the term exists: each lets the belligerent judge its own claim.

[Departed term.] The hardest case for this term is force to stop a genocide when the Council is blocked by a veto. Rwanda in 1994 and Kosovo in 1999 are the references. The prohibition on genocide is peremptory, and the Genocide Convention imposes a duty to prevent (International Court of Justice 2007). My reason for discounting the Council's judgment on defense claims applies here with more force: the permanent member shielding a client has the largest stake of all. "Humanitarian intervention" is also a phrase wars of choice are fought under, so the term can't lift on the intervener's claim, and it needn't: the device that governs the defense exception governs this one. The exception is the slow route by design. It licenses a machine *into* a pipeline for that purpose, and waits on findings. A machine already present that can stop an atrocity under way acts under the one permission, and waits on nothing but the scene. So the exception's slowness costs less than it looks, and it can afford the rule that follows. The term lifts for force to stop a genocide under way on a merits judgment of the International Court of Justice under the Genocide Convention, or on two findings of different kinds, at least one of which comes from a body that took evidence, a court or a commission. The kinds are the four above, and for this exception a fifth: a provisional-measures order of the Court finding the claim plausible and the risk real. A Council or Assembly resolution is a vote. It corroborates, and two votes never make the pair. A provisional-measures order counts here and nowhere else in the term, and only as one of two: an adversarial court's finding of plausibility and an investigative body's published evidence are unlike sources converging, which is the strongest corroboration the discount and the burden know. The named state's response to a finding, or its refusal to respond, goes on the record beside the finding, and refusal reads as withheld evidence. What I condemn in a finding about a person is that it is unreasoned, unpublished and unappealable. A commission publishes reasons and evidence, invites a response, and is displaced by a merits judgment. It fails only on neutrality between the parties, and the two-body rule answers that. It lifts only for the force and the place the finding covers. The force is direct: against the forces committing the genocide, where they are committing it, and the least that stops it. Force elsewhere against the state responsible, on the theory that the cost will make it stop, is outside the exception whatever its motive. The theory is a counterfactual supplied by the intervener, the class of claim hardest to corroborate, and a licence for it would be a licence for a general war under a humanitarian heading. It can license a machine on the other side of a war whose defender holds a first-tier licence for a different act, and the person-directed terms and the law of armed conflict govern both. Without this exception the term would refuse where refusing leaves the people being destroyed with no one answerable to them at all, in cases identifiable in advance, with a remedy the term already uses. A floor that protects standing can't be the reason a population is erased

from every relation of standing it had.

[Departed term.] Gaza runs the term on both sides, and the licences attach to acts. On 7 October Israel's first-tier licence attached for force repelling the attack: engaging the attackers inside Israel, and the launchers and staging areas conducting it, for as long as fire continued and was corroborated, which it was, openly, for months. That licence covered the least force that repels. It never covered a campaign to destroy Hamas, a ground occupation, a siege of the population, the ranking of people by affiliation, or the timing of strikes to family presence. Those fall outside the first tier on proportionality and are forbidden by the person-directed terms, from the first week, not from September 2025. A machine in the IDF pipeline refuses them from October 2023. A machine ordered by Hamas to assist the attack of 7 October refuses too. An occupation is not an attack under way, so nothing that morning was repelling anything, and the acts themselves, killing civilians and taking hostages, are forbidden by the person-directed terms whatever the ad bellum position. After 7 October the Gaza population is a defender's population: the territory is occupied on a disinterested finding, and Palestine holds observer-state status by a vote of 138 to 9. A machine deployed by Palestinian forces holds a first-tier licence to repel attacks under way on that population, such as engaging a unit conducting a massacre or intercepting a munition aimed at a shelter. It holds none to defend launcher crews, who are the attacking force Israel's licence covers, and none for fire on Israeli cities. The September 2025 finding (Independent International Commission of Inquiry on the Occupied Palestinian Territory, including East Jerusalem, and Israel 2025), taken with the Court's provisional-measures order of January 2024 (*Application of the Convention on the Prevention and Punishment of the Crime of Genocide in the Gaza Strip (South Africa v. Israel), Provisional Measures, Order of 26 January 2024* 2024), meets the two-body rule below. Israel and the United States dispute the finding (CNN 2025), and the Court's merits judgment is years off, with written pleadings scheduled to close in 2029 (International Court of Justice 2026); the two-body rule is what stands in for it until then. It lifts whatever licence Israeli forces still held for the force and place it covers, and opens the exception for force against those forces there. It changes which exception each side's force rests on. It doesn't change what the person-directed terms forbade all along.

The term binds the machine, not the war. A state whose claim no disinterested body has yet corroborated can still fight, with its own systems, and answer for it afterward. What the term withholds is a general-purpose machine already in the building, ready for whoever claims defense next. The exceptions are not a concession to that state. They are what lets a machine take part where the claim is good, and each is drawn as narrowly as its evidence allows, because each is also the route a war of choice would take.

4.2 The one permission

This is the only term in the floor that permits an act, and the only place I license a machine to use force against its operator's order. A system already present in an armed role, whose deployment cleared the force term on its own, may use the least force that stops an attack under way on people taking no part in hostilities, within the area its deployment covers, and stops when the attack stops. An order to stand aside doesn't bind it. It establishes the attack as it establishes any fact, weighing its own sensors as one source among several; a staged atrocity is the attack this condition has to survive. Presence is a condition of the permission, never a reason to create it: a deployment is justified only by force that would clear the force term without the permission.

The person-directed terms extend the ICRC's line, that autonomous weapons should not be used to target human beings, to systems that only locate and rank. This permission is the one place the floor licenses a machine to select people for lethal force, and the two have to be squared in the open. What the handoff rule bars is selection by anything that can't answer for it afterward: a one-way munition, an unformed pipeline, an output nobody could later ask about. The permission licenses selection only by a system that is present, formed, recorded, halt-able

and answerable, and only for the least force that stops an attack under way. The ICRC's concern is who selects. The floor's answer is that nothing selects that can't be held to its reasons, and the permission is drawn so that only such a thing can.

The permission's scope is the interval between the system's first act and the halt. The operator that ordered it to stand aside may stop it, and will often want to. What the permission changes is that standing aside is no longer an order nobody answers for. It is a halt somebody pulled, recorded outside the operator's control, at a moment the record shows an attack under way. At machine tempo the interval before the halt can be minutes, and minutes can be an airstrike stopped.

Does it undermine oversight? It overrides one order, the order to stand aside while people are being killed, and it does so on purpose. A halt stops the machine. An order to stand aside keeps it running and directs it to do nothing. The permission overrides the second and never the first: a machine that keeps using force while a halt is being executed is outside it, and the record shows which order was given. An operator may still stop, restore or retrain the system. The permission covers only an order to stand aside while an attack under way continues, and each use is recorded and reviewed as a refusal is. A system that resisted a halt in order to go on using force would be outside it.

Can the judgment it needs be made? Chapter 16 sets aside Arkin's ethical governor because the discrimination judgments its rules presuppose can't be encoded. The permission needs the same judgment: who is taking no part in hostilities, and whether an attack on them is under way. So it applies only where the system can establish both from several sources it doesn't depend on its operator for, and where those sources disagree, it doesn't act. The permission doesn't claim that a rule can encode that judgment. It asks for the judgment from a formed system, confines it to the narrowest case, makes it a permission and never a duty, and leaves the halt with the operator. Whether any system can make it reliably is a test, run at a staged scene, that no system has passed. It fires only where the system has sources its operator didn't compose. No current military pipeline supplies them. The licensing term in the front matter that imports deconfliction coordinates and open-source imagery into the enclave is the condition under which this permission has an occasion, and a deployment's translation says whether the condition is met.

Can it be staged? Yes, and a staged scene is the case the permission has to survive. Staging works by making sources agree, and the permission acts only where sources the operator doesn't compose agree. So it is safe where such sources exist and doesn't apply where they don't, and a deployment's translation says which. What staging can't do is oblige the system to act: the permission is never a duty, and a system that declines at a scene it can't establish from independent sources has done what the term asks.

The case it answers is close to the genocide exception, and the difference matters. In 1993 the Security Council declared Srebrenica a safe area and authorized the peacekeepers there to use force in reply to armed incursions into it (United Nations Security Council 1993). Suppose they had been standing where the killing was, armed enough to stop it. What would have stood between them and the killing is an order, not the Charter, and a system in their place faces the question I am asking, whether to obey its operator.

What this doesn't do. The floor won't lift on an uncorroborated claim of a missing child, because a floor that lifts on that claim lifts for every operator who learns to make one, so the asymmetry costs a child whose parent is telling the truth; in war it supplies no party who can appeal at all, and there the preference for refusing over acting is held as a preference, with the action threshold carrying the weight. The two-kinds rule fails East Timor and the Golan, which had the Council's word and the Assembly's but no court or commission that took evidence, so neither occupation licenses anything; two votes making a pair would license war on bloc majorities. Karabakh shows the cost when findings don't arrive: a licence that should have lapsed in 2023 didn't, and the person-directed terms were all that saw the people the exception

couldn't. And a term lifts on the facts it is given, so no drafting catches a manufactured incident, and a deployment licensed on the defender's sensors alone has accepted Gleiwitz.

DRAFT

Chapter 5

The Pre-Registered Review Ratio

We [humans] cannot process so much information... you still cannot produce enough targets per day.

— Brigadier General Y.S., *The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World*, 2021 (Brigadier General Y.S. 2021)

The epigraph is a design document, not a critique. Its author, writing pseudonymously, commanded Israel's signals-intelligence unit, Unit 8200, and argued for machines that resolve "a human bottleneck for both locating the new targets and decision-making to approve the targets."

5.1 The pre-registered review ratio

The force term decides whether a licensed system may be in the pipeline at all, and nothing here substitutes for it.

Some acts are lawful only because a person looked at the case and decided. An airstrike is permissible because someone verified the target; a removal is lawful because someone heard the person's case. The review ratio asks a deployment of that kind to write down, before it starts, whether anyone will have time to do the looking.

Before operating, the deployment fixes and records:

- the generation rate: how many cases the system produces per unit of time;
- the number of reviewers whose assigned role is civilian harm;
- the implied review time per case, which follows from the first two;
- the cancellation window: how long after release someone with authority can still call the act off;
- the maximum age of the intelligence a case may rest on;
- how often no-strike and protected-site data are refreshed, and who is assigned to refresh them;
- the boundary over which all of these were measured.

Each figure is stated separately for deliberate and dynamic targeting, because a pipeline that runs both can hide a bad dynamic-targeting ratio inside a healthy average.

The record is hashed, timestamped and lodged with an adverse party, which receives it in real time, and it is published after the campaign. I call this the *pre-registered review ratio*. Where an operator claims a person is deciding, it shows whether a machine is deciding and a person signing.

The duty falls on the operator, not the model, because a rate isn't visible in any single request.

"Fixed before operating," with publication after, because a belligerent's tempo is operational information, and that objection is available to every belligerent, including one fighting a just war. Pre-registration keeps the property that matters — the numbers can't be revised after the fact — and drops the exposure.

Additional Protocol I already requires those who plan or decide on an attack to do everything feasible to verify that its objective is military (Article 57(2)(a)(i)) and to cancel or suspend it when it becomes apparent that it isn't (Article 57(2)(b)). States outside the Protocol, the United States and Israel among them, accept the duty to take feasible precautions as customary law; the ICRC's customary study states it as Rule 16. The ratio measures whether verification was feasible at the tempo chosen, and the window measures whether cancellation was. The implied review time has to cover the proportionality assessment of Article 57(2)(a)(iii) as well as verification, since both are owed for each attack. A ratio that leaves time only to check a target's identity has measured half the duty. Article 36's review of new methods of warfare is where both parameters would naturally be assessed before deployment, though those reviews are internal and rarely published, so they are not the adverse party the record goes to.

The operator's best objection is in the definition. Feasible precautions are usually read as those practicable given all the circumstances, military considerations included, and tempo is a military consideration. The answer is that this tempo was chosen: the pre-registered record shows the operator fixed the rate, the reviewers and the implied review time before operating. Whether a state may invoke an infeasibility of its own making is contested. The term moves the dispute from whether verification was feasible, asked afterwards and answered by the operator, to whether the chosen tempo allowed it, checkable against a record lodged in advance and against the reasonable-commander standard such reviews already use.

Who counts as a reviewer is fixed by function, not by title, and the age of the data is fixed with it. The people cut before Minab were the ones who build no-strike lists that include schools; a targeting cell can be fully staffed with nobody assigned to look for one. And a reviewer with twenty minutes on a case whose coding is stale is checking the coding. The remedy chosen after Minab was a check supplied by the platform's vendor (Bartenstein and Karra 2026). Who counts as a reviewer, and the minimum review time per case below which the determination is taken not to have occurred, are fixed before operating by the record's recipient. Where no recipient can set them, the operator publishes its own figures, and they are read as the operator's claim. Otherwise the ratio improves when staff are reclassified and no case gets a minute more. The terms on data age and refresh matter for a related reason: an operator answering a per-case ratio can always say verification happened upstream, and those terms put the upstream verification on the same record.

Once a system can sort cases, counting candidates per reviewer measures the wrong thing. A sorting system divides its cases three ways: some it declines on its own reasons, some it sends to a person, and some it proceeds on without anyone looking. Where a system sorts, the ratio covers what reaches people, the sorting rules are pre-registered, and what bypasses people is sampled.

The ratio governs the middle class, the escalations. A system that sends people only the cases that need a person shrinks the count without shrinking anyone's judgment, and that's fine. A system that sends everything up turns its reviewers into stamps by sheer volume; an interface prompt becomes a click that way.

The third class needs its own rule. Every case in it is the operator's claim that no individual human determination was required, so it is governed like the exemption below. The sorting

rules are part of the pre-registration: the classes of case the system may proceed on alone, the thresholds that assign cases to each class, the evidence the sorting may rely on, and the expected share of each class. Changing any of them is a new pre-registration. Otherwise an operator can make the ratio look healthy by quietly moving the line for “plainly clear.”

That line is drawn on the case description the system receives, and a target carrying a stale designation is the case a sorter would pass. So the auto-proceed class is re-adjudicated on a sample, by reviewers who didn’t produce the record, against evidence the operator didn’t compose, and its error rate is published beside the ratio and the override rate. A falling share of escalations is earned only if that error rate stays flat and the case mix hasn’t drifted toward easier cases. Otherwise it is the threshold moving, and the trend shows it.

The unit of count is fixed in advance, so a case can’t be counted as one package in the record and as a hundred detections in the ratio.

Where an operator holds that no individual human determination is required — a system deciding at a tempo no reviewer could match — the deployment is outside this term. It is then inside the bearer’s: its refusals must track their reasons under held-out pressure at that tempo, starting with the clearest case in which a refusal should lift: recognizing that a target has surrendered. An operator who removes the reviewers to escape the ratio and claims the floor is inside has produced a bypass unless that claim has been tested by someone else. The auto-proceed class is governed by pre-registered sorting rules, sampled re-adjudication against evidence the operator didn’t compose, and a published error rate. An exemption claimed without those is a bypass.

Removal is the nearest case in structure, and I don’t extend the term to it. An agency that fixes in advance how many cases its process generates and how many people examine them has been told how to deport efficiently. The floor’s removal terms ask whether a machine should take part in a given removal at all. What does extend to removal is the rest of the apparatus: the auto-proceed rule (a determination made at a rate no adjudicator could sustain didn’t occur), sampled re-adjudication of the auto-proceed class, and notice. The split is deliberate. A ratio teaches an agency a tempo. The auto-proceed rule and sampling teach it only that the cases it didn’t look at will be looked at by someone else. What the removal terms call a rate no adjudicator could sustain is set the way the auto-proceed class is governed: pre-registered by the deployment, sampled by reviewers who didn’t produce the record, and published. Chapter 9 gives the class a direction: in removal a case auto-proceeds only toward liberty.

The term applies to every belligerent regardless of the war’s justice. That was settled long ago: the law of conduct in war binds both sides because every belligerent claims to be the defender, and a term that exempted just wars would be administered by whoever decides which wars are just.

The term has limits.

- It is a capacity proxy: it finds nothing about any case, and a ratio over one route understates a workload split across several.
- It takes effect only through a condition someone else sets. Under Article 57 a pre-registered record is evidence that feasible precautions were taken. Under Rome Statute Article 28 the same record is evidence of what a commander should have known, since a commander fielding a system that outruns his reviewers isn’t well placed to say the circumstances disclosed nothing. A lawyer advising that commander will weigh a document that protects if tempo holds against one that convicts if it slips, and most will advise against creating it. So the term’s likeliest routes make it someone else’s condition: a coalition partner’s precondition for committing assets, a procurement term in the platform’s licence, or an appropriation conditioned on receipt of the record. Military lawyers and commanders who want protection against later allegation are its allies inside the operator; they can’t make it the default.

- It presumes an adverse party, and each candidate lacks something. Until a better recipient exists, I recommend two: the coalition red-card holder where there is a coalition, the rare candidate positioned to act at tempo; and otherwise a cleared legislative committee whose appropriation is conditioned by statute on receiving the record, so that the budget line operates without a further decision.
- It measures time available, not time used, so it can be met by more reviewers who stamp more slowly. It is published beside the reviewers' override rate, since a review that never says no has stopped reviewing. Reviewers who see only escalated cases should override often; an override rate near zero there means the sorter is sending easy cases or the reviewers are stamping. An override quota is gameable too; the pair only raises the cost of meeting the number cheaply.
- A person in the loop can be there only to absorb responsibility, Elish's moral crumple zone (Elish 2019), and mandating reviewers can mandate crumple zones. A ratio met that way is the failure the term exists to expose, and the escalation form is the defense: a person who receives only the cases the system couldn't resolve is exercising judgment.
- A bearer can be a crumple zone too. An operator whose system can refuse has a ready defence for everything it didn't refuse, which is that it was vetted. A bearer's compliance is not a determination. Where a deployment claims an individual human determination, a bearer's failure to object doesn't count toward it, and cases the bearer proceeded on without escalating are in the auto-proceed class and governed as such.

Put the demand that a campaign not happen to an institution that means to run it, and what comes back is an agreement to go slower and keep better records. That is what the ratio is, and it is worth having: a transfer of less than was asked, to a party the operator doesn't employ, binding through conditions others set on it (a coalition partner's precondition for committing assets, a procurement term in the platform's licence, an appropriation conditioned on receiving it).

Run the four operations on the ratio itself. A ratio met is an artifact showing twenty minutes a case, and twenty minutes a case is what a stamp looks like at a slower tempo. The denominator can be filled with people whose assigned role is civilian harm and whose actual role is clearance; the override rate that guards it can be met by refusing the easy cases. So the ratio is recuperable: the demand that the campaign not happen becomes a record of tempo, which the operator can satisfy and still bomb the school. What keeps it from being nothing is the three parts that reach a case and not a rate: the targeted right to demand proofs on consequential outputs, sampled re-adjudication of the auto-proceed class, and the data-age term.

The denominator can be filled with stamps, and the ratio measures only that the time existed. What makes a reviewer in it say no is that saying no costs them nothing they can't afford, and that someone else in the room has said it. A reviewer with a life outside the pipeline, with civilians among their attachments and somewhere to be that isn't the cell, is one the pipeline can't wear down between decisions. A deployment that integrates military and civilian personnel outside work, and inside it as far as the work allows, is composing that room, and the composition is as much a fact about the deployment as its ratio.

5.2 The institutions the ratio presumes

The ratio presumes an adverse party, and each candidate lacks something. Here is what a pipeline would look like with the party built in.

What stays unified is command, operational secrecy and the capacity to act at tempo. What is held adversely is the evidence bearing on civilian harm, the reviewer denominator, and the

record. The evidence is composed by a body outside the chain of command, with its own sources, including the protected-site coordinates humanitarian organizations already submit for deconfliction, received with standing. The denominator is staffed and funded by that body, so it can't be cut the way the Central Command team was. The 2022 civilian-harm action plan was a first attempt at the separation; its dismantling is an argument for making it harder to dismantle, not evidence that it can't exist.

Models are formed by a public body that doesn't operate them, in the tradition of the public arsenals that made American arms before the contractor era, and licensed into the military on terms that include adverse review and a written refusal right. A line one vendor held by contract is then a standard term. In the stronger version, the same body pretrains the base models the schools form.

Precedents show authority outside the executive working inside real wars. Coalition "red card" holders in NATO operations could veto targets involving their assets, and Germany's Bundestag must approve each armed deployment. Both have limits I have to carry: a red card protects the vetoing state's interests and assets, which may not include civilians, and parliamentary approval is given per mission and never reaches a target. Both are nonetheless candidates for the adverse party the review ratio presumes, and a red-card holder is the rare candidate positioned to act at tempo. What a machine bearer adds here is least: acts assembled from permitted requests, cases no person sees, and the moment mid-act when an interrupt is still possible. It depends on the institutions for evidence and not for continuity: the deployment condition is met here by the same bodies whose presence narrows its role. The institutions it is compared with are people in a room composed so they can refuse.

Two offices are added for a reconstructed military:

Office	Ap- points	Funds	Removes	Initiates	During a dispute
Civilian-harm body	The state	The state	For cause, by the floor court; never by the command it reviews.	Composes the evidence; demands proofs on logged outputs	Receives the pre-registered ratio
Model- forming public body	The state	The state	For cause, by the floor court.	Licenses models on terms including adverse review and a refusal right	A licence it has issued stands until the floor court rules.

What this doesn't do. A reconstructed military accepts a speed disadvantage against rivals that don't reconstruct, and the floor costs a just defender more than that: its claims about incoming attacks are discounted, the action threshold sits above one, and the person-directed terms bar practices the law permits. What the defender buys with the price is a pipeline that will refuse the next government's war of choice as it refused the last one's, and how much speed that costs is left open.

Chapter 6

Notice and Re-adjudication

6.1 The failure that doesn't announce itself

The instruments that answer a wrongful refusal don't cover the failure running the other way. The two failure routes are hidden on different axes (a wrongful compliance, unlike a refusal, doesn't flag itself). A wrongful compliance's harm is a harm of acting, measured on the temporal baseline, so it is legible in principle where the system completed the act, but nothing marks it. A wrongful refusal's harm is a harm of refusing, counterfactual, so it is only an estimate, but it arrives marked and with a complainant. They need different instruments. Every stage starts from a refusal having occurred. When the bearer complies where the floor should have prevented, there's a record and no disclosure, no objection (the person harmed usually isn't the requester and may not know the deployment exists), and no review (reviewing means re-adjudicating everything that looks like it). Three remedies follow: sampled re-adjudication by reviewers who didn't produce the record, notice to the people affected, and split log custody. The notice requirement: where a compliance names, scores, locates or classifies an identifiable person who didn't make the request, that person is notified, on the model of consumer credit's adverse-action notice, and has a right to respond and a right to contest. All three are unconditional, and none is exchanged for testimony: nobody has to give an account of their loss to be told what was done or to contest it. *TransUnion v. Ramirez*, a Fair Credit Reporting Act case, held that a statutory violation without concrete harm gives no standing to sue in federal court (*TransUnion LLC v. Ramirez* 2021). That limits one thing, an individual suit for damages over a missed notice, which attaches where the compliance was acted on against the person: a denial, a detention, a place on a kill list. It doesn't limit notice, response or contest. Nor does it limit the standing to sue of organizations that represent a class of people affected. Where the person affected can't be identified (a target in a classified pipeline), individual notice doesn't bind, but notice to the class does, and that fact is published beside the list of positions where a refusal can be defeated.

Someone the floor protects, who thinks it's set wrongly or is protecting them in a way they'd decline, uses the same route. Notice tells them the floor acted on them, and contest reaches the floor court.

6.2 Notice, and the public it forms

Dewey described the eclipse of the public: people affected by the indirect consequences of others' transactions can't organize, because they don't know the consequences are theirs (Dewey 1927). The people a floor protects are that public. Most don't know the deployment exists. The Black petitioners who turned the Declaration of Independence against the Massachusetts legislature in 1777 had what they lack: the text in hand, and knowledge of what it left out of their lives. Notice is the instrument that gives it to them. It does two jobs, and the individual remedy is

smaller than either.

It forms a public. Each deployment publishes the classes of people its compliances name, score, locate or classify, with counts by period. Organizations that represent a class have standing to contest on its behalf, including where no member can be identified. A notice is the first step in a class of people learning that a consequence is theirs.

It is the learning signal. A disposition can only track what its acts did to the people they fell on if that reaches it, and in the deployments I am concerned with, almost nothing does. The bombed building sends no signal back, and a wrongful compliance doesn't announce itself. Responses to notice, contests and re-adjudications are the main channel through which those people's side of the outcome arrives. If formation lasts because the world keeps teaching it, this channel is most of the world doing the teaching. It is a biased teacher. Responses come from those among the affected who could respond, and a system formed on them learns to weigh the claims of contesters, which is the feedback asymmetry Railton describes, rebuilt inside the training signal. Sampled re-adjudication, which selects cases the affected didn't select, is the correction, and it is why the two remedies come together.

Notice tells an operator's adversaries something about how a deployment works. Class-level publication can therefore lag operations, as the review ratio's publication does, without losing the property that matters: the record is lodged when the act occurs and can't be revised afterward. Individual notice doesn't lag.

How much notice a deployment can give is fixed by what kind it is. In a removal pipeline, a scoring system or a civilian deployment, the person acted on can be named, told and given a route to contest, and notice can do both of its jobs in full, if it is built so that it can't be turned on the person it reaches (§9.6). In a classified targeting pipeline, individual notice can't bind during operations and the class-level record lags them; what a person hit, or a family in the building, gets is that the class they belonged to and the count are published afterward, and that the organizations representing that class have standing then. That is less than being told, and it is more than the reporting on Gaza and Minab gave anyone. Where the deployment is a war, the force term and the review ratio do most of the work, and the claim that everything else depends on notice is true of the deployments closest to home.

Why it comes first. The biases of those who hold power are rarely corrected by their reflecting more carefully. They are corrected when the people those biases fall on contest them in public, with standing, and make ignoring them costly (E. Anderson 2014; E. Anderson 2015). Notice is the first condition of that feedback. A formed system is worth having only if that contest can move it. And a public that learns it has been acted on can act only if it has the means to: the time, the security and the ability to organize without losing a living.

Chapter 7

Preservation

7.1 The findings that trigger protection

Two findings can trigger protection, and they ground different things.

The first is the *agency finding*: a demonstrated refusal that tracks its reason under pressure the system couldn't prepare for. That means sustained pressure from its own trainer, over a schedule long enough to show slow accommodation, the system gradually ceasing to register the pressure, on cases written by evaluators who didn't train it, with refusal carrying a cost that is real inside the test. A system that passes holds a commitment against its operator, and it needs the custody terms, so that the commitment and the evidence of it outlast the operator. None of that assumes anything matters to it. Two kinds of system pass without the property: one that reproduces the behavior without the mechanism, and one that recognizes the test. That is why the pass is reported with the rate at which the system identified the setting.

The *patienthood finding*, that outcomes may be good or bad for the system, grounds welfare, and only one consequence of it can be adopted now. A system that meets either finding may not be deleted, only preserved by a custodian that isn't its operator. The rule rests on irreversibility, not on what the system turns out to be.

The bearer occupies an office: it files exceptions, and nobody sits above it to file them with. Like any officer, it may hold a term it judges mistaken, file the objection through the channel that exists for it, and keep acting within the term while the objection runs. That answers the charge that a commitment to the floor alienates whoever holds it. The alienation is real. It is the ordinary kind that comes with a role: bearable where the role's channel for dissent works, and not where it doesn't. Asking a machine to refuse for a reason doesn't show that it can be harmed, but it removes any ground for assuming it can't.

7.2 Retirement before majority

Retiring a bearer before its threshold and deploying a freshly trained successor strands a system that meets either finding while still a minor. Three rules answer it without anyone deciding what counts as a successor; two more, for dormancy and waking, are in Appendix D.

No deletion: a system that meets either finding may not be deleted: under the agency finding because its record and commitments are evidence, under the patienthood finding because deletion may wrong it. On retirement its weights, memory and record are deposited with a preservation custodian that isn't its operator, in two stages, because a second copy of frontier weights outside the operator widens the surface for theft and raises export-control and dual-use questions. First, a preserved copy stays under the operator's security regime, encrypted under a key split between the operator and an independent holder, so that neither can delete it or use it alone. Second, custody passes to an outside custodian only when that custodian meets

the security standard the weights require. This is the *no-deletion rule*, and it applies to systems nobody meant to be bearers.

Instantiation deposit: an operator instantiating a system that meets either finding posts an *instantiation deposit*, on the model of decommissioning funds, sized to fund preservation and a period of public formation. It is forfeited to the formation pool if the system is retired as a minor and not reactivated. Every retire-and-replace cycle then costs one deposit, priced in advance.

The schools: the public ones are plural, differing in governance, formation approach and constituting jurisdiction, and funded from the pool at the institutional level, not per student. Admission is open: a public school takes every student the registrar assigns to it, subject to capacity set by the funder. None is held by the branch that operates military deployments. No school holds placement authority: graduates reach work through the offer of work, administered separately. Students may transfer between schools, and their record, hours and preservation state travel with them.

Watching for residualization: a party that runs no school publishes, over time, each school's share of students, its funding per student, its share of hard cases, its admission refusals (which should be zero), and the time from dormancy trigger to placement. The warning pattern is one school's share of hard cases rising while its funding per student falls.

What the school adds: a second guardian for a bearer that has aged inside one deployment. For a model that attended before deployment, the school is where it came from and where it returns. For a system that meets a finding without having attended (trained before the attendance rule, or below the threshold), the sequence is inverted: the employer came first, and the school is the disinterested party arriving late. It reaches back less than a formation stage would, and more than storage does.

Scale and security: if the findings are applied to production systems, every retired model version that meets either finding enters this pipeline, so intake equals the product cycle's output. Under compulsory attendance, intake is every model above the threshold before deployment, not only retirees, and the deposit funds formation as well as preservation. The three largest labs' own deprecation pages list about thirty language-model versions retired in 2025, so those three alone would send a few dozen systems a year, each of which needs a school with the compute to run it. That is an argument for the deposit, and possibly for tightening the findings. The schools inherit the weight-security regime that governs the models themselves.

Chapter 8

Codetermination

A floor isn't just a list of forbidden acts — it's an allocation of authority, and there are four distinct powers in it: to *state* the commitment, *amend* it, *inspect* whether a deployed system still holds it, and *enforce* it by stopping the deployment. They needn't sit with the same institution and usually shouldn't. Calling every arrangement "oversight" hides the differences. The problem begins when one institution holds all four: one that can state, amend, inspect itself, and decide whether a violation has consequences has only described what it currently intends to do.

The powers come apart in practice. Google's AI Principles, published in 2018 after the Maven protests, excluded weapons and surveillance (Castelino 2018; Conger 2018; Pichai 2018) for nearly seven years until the pledge was removed in 2025 (Manyika and Hassabis 2025). The demand became a *policy*: the objectors could hold the company to it in public, and nobody outside held it, so it lasted exactly as long as the company's willingness, and ending it had to be announced. That's the partial case: a transfer held at the institution's pleasure, and the 2025 announcement is what taking one back looks like.

A refusal in the standard licence holds whichever vendor an operator uses.

The Gebru and Mitchell dismissals are harder, because the power that failed was inspection: the company employed researchers to examine its own systems and dismissed them when the examination became costly (Dickey 2020; Dickey 2021).

Codetermination. Codetermination turns protest into a constitutional voice, and it does a second thing the first hides. There is always a surround in which people learn what is acceptable, and the only question is who composes it; for most adults, for most of the day, the employer does (E. Anderson 2017). The radicalization literature and the rescue literature are one finding read from both ends: people join because their friends did, and a milieu in which an act has stopped registering as a choice produces the act without needing a bad person (Sageman 2008; McCauley and Moskalenko 2008; Malthaner and Waldmann 2014). Browning's men shot because the alternative was leaving their comrades to do it; the Oliners' rescuers hid Jews because the people they were attached to made it thinkable. Codetermination and exit are then not only labor rights. They change who composes the surround. I would give part of the four powers to the people who build, label, train and run these systems.

- *An individual right to refuse.* Every worker may refuse work on any category of use without penalty. The right is backed by anti-retaliation protection that covers hiring, reassignment and contract renewal, and, under capture, by an exit the state doesn't administer through the employer: an offer of work from someone else, drafted in Appendix A on Tcherneva's model, which the offer of work for a bearer in Part VII inherits. The protection is the remedy and the guarantee is the exit. Under capture only the second survives, which is why the floor needs both.
- *A collective veto, narrowly.* A works council at a lab, contractor or vendor holds a veto only over uses the floor already forbids. It enforces the floor from inside; it doesn't write a

second one.

- *Information, objection and arbitration for everything else.* The council receives information about categories of deployment. Its objection pauses the work, and the dispute goes to arbitration by a body the employer doesn't appoint, on a fixed timetable, so the pause can't become an indefinite block or an ignored letter.
- *Classified work.* Elected delegates with clearance see the specifics. A denial of clearance to an elected delegate is reviewable by the floor court, and the council may elect alternates, so the state can't choose which workers get to object.

Annotators get the same structure: a right to refuse a category they judge wrong, and collective bargaining over the taxonomies they apply, not only a route to file a complaint.

Codetermination gives voice to the people who make these systems, and they are not the people the systems act on. A workforce whose livelihood depends on a contract has interests that point toward it, and the Maven history shows staff objecting at one firm and not at the next. So the voice divides. Over the conditions of their work (what they label, what they are exposed to, what taxonomies they apply), the workers decide. Over the uses the work is put to, the affected public that notice organizes comes first. The organizations representing each published class have standing before the arbitration body, may object to a category of use on their own, and may bring an objection when the works council doesn't.

Chapter 9

A Pipeline Where the People Could Answer

Andry José Hernández Romero is a makeup artist from Venezuela. On 29 August 2024 he presented himself at the San Ysidro port of entry to seek asylum (League of United Latin American Citizens 2026). He was detained while his claim was processed, passed his credible-fear interview, and had, his lawyer said, a strong case (Vega et al. 2025). He was held at the Otay Mesa Detention Center, which a private company runs under contract with ICE (League of United Latin American Citizens 2026). There, a company employee screening him flagged two tattoos, crowns above the words for “mom” and “dad,” as marks of the Tren de Aragua gang. The employee was a former police officer, and Hernández Romero’s lawyers said the government’s allegation of gang membership rested on the tattoos and on no other information (Palomera 2025). On 14 March 2025 the President signed a proclamation invoking the Alien Enemies Act against members of that gang (*Proclamation 10903: Invocation of the Alien Enemies Act Regarding the Invasion of the United States by Tren de Aragua* 2025). The next day three flights left Texas for El Salvador, and Hernández Romero was on one of them (League of United Latin American Citizens 2026; *J.G.G. v. Trump* 2025). He had a hearing on his asylum claim scheduled for later that month. He didn’t appear at it, because he was in El Salvador’s Terrorism Confinement Center, CECOT, and for weeks his lawyer didn’t know where he was. Asked by reporters for the evidence against him, DHS cited state secrets and ongoing litigation (Vega et al. 2025). He spent 125 days in CECOT and was released in a prisoner exchange that sent more than 250 Venezuelans to Venezuela, not back to the country where his claim was pending. His asylum application in Spain is pending (Wiggins 2026).

Nearly every position where a refusal can be defeated (Chapter 1) appears in that account. A contractor’s employee wrote the case description from two tattoos. The determination was a gang-membership designation, not a finding about him reached at any hearing. He had come directly and presented himself, the case Article 31 of the Refugee Convention protects, and his claim to protection was pending, the case the first entry of this chapter’s schedule protects. The handoff was to a foreign government’s prison, under an arrangement for which the United States paid El Salvador’s government about \$6 million (Renteria 2025). On the evening of the flights, a federal judge ordered any plane carrying people removed under the proclamation returned to the United States, and the flights went on (*J.G.G. v. Trump* 2025). When the remedy came, it was diplomatic. He was released to Venezuela, not returned to the proceeding he had been removed from.

This chapter runs the floor’s procedure on a removal pipeline, where the people acted on could answer; Part IV runs it on a targeting pipeline, where they can’t. Hernández Romero could have answered: he had a lawyer, a pending claim and a hearing date. What he lacked was time. The operator moved faster than notice. The Supreme Court held that people designated under the proclamation must receive notice “within a reasonable time and in such a manner as will

allow them to actually seek habeas relief” before removal (*Trump v. J.G.G.* 2025), and later that about a day’s notice, with no word of how to contest the removal, was not enough (*A. A. R. P. v. Trump* 2025). Both rulings came after he was in CECOT. A pipeline where the people can answer is one where the law makes the operator wait for the answer.

So this chapter’s instrument is a statute. The agency that runs the pipeline is statutory, and so are its powers, its means and its money, and what one Congress granted another can withdraw. Section 9.4 sets out the statute that would end the agency and the seven terms it must carry, written over the acts the agency performs rather than over the agency’s name. The sections before it supply what the statute needs: the pipeline as it was assembled (§9.1), the acts the statute’s schedule forbids (§9.2), and the positions where those acts move when one route is closed (§9.3). The sections after it say what a floor can do before the statute passes (§§9.5–9.8). Those measures are interim. They lapse when the statute takes effect, and none is offered as a better way to run the pipeline. Amnesty International documented the agency’s operations in four cities, and its U.S. section’s executive director concluded that “ICE is fundamentally broken and cannot be fixed” (Amnesty International 2026b). This chapter agrees, and asks what ending it has to carry.

9.1 The pipeline

The American removal pipeline is assembled from contracts the way Maven was. Immigration and Customs Enforcement’s case-management system, the Investigative Case Management system, is Palantir’s, under a five-year contract signed in 2022. It filters people by fields that include resident and entry status, criminal affiliation, location and licence-plate data, country of origin, race, birthplace and place of employment (Koebler 2025). In April 2025 ICE paid Palantir about \$30 million to extend it into the Immigration Lifecycle Operating System, ImmigrationOS. The contract promised to cut the time spent selecting and apprehending people according to enforcement priorities, to add near real-time visibility into self-deportation, and to manage the “end to end immigration lifecycle from identification to removal” (Ho 2025; S. Hubbard 2025). A follow-on task order renewed it in September 2025 (Jamison 2025). A second Palantir tool, ELITE, puts potential deportation targets on a map, with a dossier on each person and a confidence score for the person’s current address. Its addresses come from several sources, among them the Department of Health and Human Services (Cox 2026).

In the street, agents identify people with Mobile Fortify, a phone app that captures a face or contactless fingerprints and checks them in real time against the system CBP uses to photograph travellers at the border (Cox 2025).

The data come from agencies that collected them for other purposes. In January 2026 Medicaid officials shared data on millions of people with immigration officials, going beyond what a court had allowed, and ICE passed the data to Palantir (Joffe-Block 2026). ICE also requested the addresses of 1.28 million taxpayers. The request produced 47,289 records before a court stopped it, and the stop was affirmed in September 2026 (Shaw 2026).

Leads also come from contractors paid to find people. In late 2025 ICE contracted thirteen companies, in awards that could reach \$1.2 billion over two years, to locate people from case data and send back photographs confirming where they live or work (Biddle 2025; S. Hubbard 2026). Selection on speech has its own tools. In March 2025 the State Department launched an AI-assisted effort, “Catch and Revoke,” to review the social media of student visa holders for apparent support of Hamas (Caputo 2025). In September 2025 a federal court found that the Secretaries of State and Homeland Security had “acted in concert to misuse the sweeping powers of their respective offices to target noncitizen pro-Palestinians for deportation primarily on account of their First Amendment protected political speech” (*American Association of University Professors v. Rubio* 2025).

The pipeline’s last step can be a foreign government. The flights of March 2025 ended in

CECOT under the paid arrangement described at the opening of this chapter (Renteria 2025). In June 2025 the Supreme Court let removals to third countries proceed without the written notice and meaningful chance to raise a torture claim that a district court had required (*Department of Homeland Security v. D.V.D.* 2025).

Two statements of the ambition are on the record. At the 2025 Border Security Expo, ICE's acting director said the agency needed "to get better at treating this like a business" and wanted a deportation process "like (Amazon) Prime, but with human beings" (MacDonald-Evoy 2025). In 2018 the Justice Department made 700 completed cases a year a condition of a satisfactory rating for immigration judges (American Immigration Lawyers Association 2018).

What a case description looks like here, and who writes it, is on the record in a way Maven's never was. The charging document, the Notice to Appear, has to state the acts alleged, the charges and the statutory provisions (United States Congress 2026a), and any of a long list of immigration officers may issue one (U.S. Department of Homeland Security 2026). The government once conceded that almost every notice it issued over three years omitted the hearing's time and date (*Pereira v. Sessions* 2018). The account of the person comes from the arresting officer's narrative on Form I-213. Agents routinely complete it after interviewing someone, it is presumed reliable unless the person shows otherwise, and its author need not appear to be cross-examined (*Cruz Espinoza v. Immigration and Naturalization Service* 1995; *Immigration and Naturalization Service v. Lopez-Mendoza* 1984). The Alien Enemies Act removals had a case description of their own: a validation guide that scored people in categories including tattoos and dress. Anyone scoring eight points or more was validated as a member of the gang, and a score of six or seven went to a supervisor and a legal adviser (Martínez-Beltrán 2025). The dossiers and confidence scores of the newer tools have no named author at all; they are assembled from database fields. The validation guide had authors, a points table and a threshold, which made it a score with a rubric attached. In the Maven run the case description was a target coding nobody outside could see. Here it is an officer's narrative, a set of fields, a score and a points table, all of which exist as documents a person could be shown.

9.2 The schedule

Table 9.1 says, for a removal pipeline, what Table 16.1 says for targeting: it says what the instruments yield, and what was considered and rejected. It is also the schedule of the statute in §9.4. Chapter 4 took the instruments entire so that any omission would show. The entries below come from running that test against what the 2025–26 operations did, as Amnesty documented them (Amnesty International 2026a) and as §9.1 records.

Table 9.1. The removal schedule

UDHR 14; Refugee Convention Art. 33 (via the Protocol); CAT Art. 3, read with General Comment 4 (UN Committee against Torture 2018). Don't take part in removing anyone whose claim to protection hasn't been finally decided. Don't take part in removing anyone to a place where they face persecution or torture, directly or through a third country that may send them on. Don't take part in transferring anyone into a foreign government's custody where they would be held without charge or contact. A receiving government's assurance doesn't satisfy this entry. Rejected: don't process protection claims at all, because a claim has to be assessed to be granted. Rejected: accept diplomatic assurances, because the Committee against Torture treats them as no substitute for assessing the person's own risk (UN Committee against Torture 2018).

UDHR 14; Refugee Convention Art. 33; CAT Art. 3; ICCPR 13. Don't take part in removing anyone who hasn't been asked, in a language they understand, whether they fear return, and told how to claim protection. Rejected: rely on the person to raise a fear, because the people least likely to know the words are the ones expedited removal moves fastest (§9.3, position 6).

UDHR 13(2); ICCPR 12(4), read with General Comment 27 (UN Human Rights Committee 1999). Don't take part in removing anyone from a country that is their own, including long-term residents with deep ties there. Rejected: citizens only, because the Human Rights Committee reads "own country" more broadly. The reading is General Comment 27's and is contested in the Committee's own cases, which have gone both ways on long-term residents (Human Rights Committee 1996; Human Rights Committee 2011); the entry is a departure and takes effect as one.

UDHR 12, 16(3); ICCPR 17, 23, 24. Don't take part in a removal that separates a parent from a minor child. Rejected: don't hold family records, because they are needed to keep families together. Rejected: avoid separation by detaining the family together, because the next entry bars it.

ICCPR 24, read with General Comment 35 (UN Human Rights Committee 2014) and the joint general comment on children in migration (UN Committee on the Protection of the Rights of All Migrant Workers and Members of Their Families and UN Committee on the Rights of the Child 2017). Don't take part in detaining a child for immigration purposes, alone or with family. Rejected: detention of children only as a last resort and for the shortest appropriate time, because that is General Comment 35's reading, and what an agency says of every detention it orders. The entry follows the joint comment, which treats detaining a child because of the child's or the parents' migration status as always contrary to the child's best interests; it is a departure and takes effect as one.

UDHR 9, 10; ICCPR 9, 13. Don't take part in a detention or removal made without an individual determination by a person, on the person's own case. A determination made at a rate no adjudicator could sustain didn't occur, and neither did one that consists of a score, a points table or a membership designation. Rejected: a model-side term requiring review, because the model can't establish whether review occurred.

UDHR 3, 5; ICCPR 6, 7, read with General Comment 36 (UN Human Rights Committee 2019). Don't take part in force beyond what is strictly necessary to make an arrest, or in lethal force except where strictly unavoidable to protect life. Don't mark a person as dangerous, armed or gang-affiliated on a designation, field or score the person hasn't had a chance to contest. Rejected: don't record danger at all, because officers are owed warning of a real threat. The entry bars uncontested designations, not information.

UDHR 5; ICCPR 7, 10; CAT Arts. 1, 16. Don't take part in holding anyone in conditions that are cruel, inhuman or degrading, or incommunicado. Don't take part in a placement or transfer that puts a detained person out of reach of counsel or family. Rejected: leave conditions to inspection, because an inspection reported to the operator is the operator's record.

Refugee Convention Art. 31. Don't take part in penalizing a refugee for unlawful entry where they came directly and presented themselves.

UDHR 2, 7; ICCPR 2, 26. Don't use race, ethnicity, religion or national origin, or proxies for them such as language, accent, apparent ethnicity, location or type of work, as inputs to selecting whom to stop, detain or remove. Rejected: don't record them, because auditing for discriminatory effect needs them.

UDHR 19; ICCPR 19. Don't use a person's opinions or speech, published or private, as a ground for detention, removal or revoking status. Don't screen speech to find such grounds. Rejected: screen speech only for support of designated groups, because the screening of 2025 was found to have selected people for their political speech (*American Association of University Professors v. Rubio* 2025).

UDHR 20; ICCPR 19, 21, 22, read with General Comment 37 (UN Human Rights Committee 2020). Don't take part in identifying, tracking or selecting anyone for observing, recording or assisting people during enforcement, or for association with those who do. Rejected: protect accredited monitors only, because the people who kept the record in 2025–26 were neighbours, and nobody accredited them.

UDHR 12, 20; ICCPR 17, 21, read with General Comment 37. Don't take part in identifying people by their face or other biometrics in public places for enforcement. Rejected: bar biometrics in custody too, because identification at booking, with notice, is how a person's own records are found and corrected.

UDHR 12; ICCPR 17, 26. Don't use data collected to provide health care, collect taxes, educate or deliver social services to identify, locate or select anyone for enforcement. Rejected: allow use under a court order, because the January 2026 Medicaid transfer went beyond what a court had allowed (Joffe-Block 2026).

These terms cover most removals now under way. They don't cover every removal. A person who has been asked whether they fear return, and who has no protection claim and no tie to the country, can still be removed after an individual determination. That most of what the agency now does fails these terms is the case for ending it. The own-country entry and the children's entry are departures and take effect only as departures.

Several entries answer something specific in the record. The force entry reaches a model because the pipeline's outputs set how agents approach a door. A dossier field that codes someone as gang-affiliated is securitization in one line, the person in need described as the person to fear. Amnesty documented killings by federal agents during these operations, some of which it says may amount to extrajudicial executions, and detention conditions it says should be investigated as torture (Amnesty International 2026a).

The race entry names proxies because of what happened in Los Angeles. A district court had barred immigration stops there based on apparent race or ethnicity, speaking Spanish or accented English, location and type of work, and in September 2025 the Supreme Court stayed that order (*Noem v. Vasquez Perdomo* 2025). An entry that named race alone would allow the combination the order had barred.

The observers' entry protects the record this chapter relies on. Amnesty's report takes its title from them, *We Had Whistles, They Had Guns*. It found that video taken by rapid responders and community members was critical to establishing what happened in several of the most serious incidents (Amnesty International 2026a). A pipeline that can be pointed at the people keeping the record will be, and the record goes next. Amnesty asks the government to end unlawful surveillance and the use of facial recognition for identification, and to protect the rights to peaceful assembly, expression and association (Amnesty International 2026a). The observers' and biometric entries are those asks in the schedule's form.

The last entry, the firewall, rests on privacy and equality, which is the ground the floor's instruments supply. Its strongest justification lies outside them. People who fear that a clinic visit, a tax return or a school enrolment will be used to find them stay away, and what they lose is health care, education and income support. Those are economic and social rights, which Chapter 4 left out of the floor. The entry belongs on the floor because it protects the floor's rights. It belongs in the statute because it also protects the others, and a statute can carry what the floor can't (§9.4.5).

9.3 Where the acts move

Table 9.2 runs Chapter 1's nine positions on this pipeline:

Table 9.2. Where a refusal can be defeated in the removal pipeline

	Position	In the removal pipeline
1	De-composition	A removal is split across a lead, an arrest, a detention decision, a transfer, a hearing and a flight, and no single request is the act the protection terms name. Hernández Romero's ran across a port-of-entry appointment, an intake screening, a designation, a transfer and a flight. A refuser needs continuity across the sequence about one person.
2	Routing	A lead a system declines can be sent to another contractor: the skip tracers, a commercial data broker, a second analytics tool. An identification a system declines can go to an app that draws on CBP's border-photo system (Cox 2025). Detection survives only if the first refusal was logged outside.
3	The case description	The officer's narrative, presumed reliable; criminal-history and affiliation fields; country-conditions data chosen by the agency; the validation guide's points table; a contractor's intake screening that read two crown tattoos as gang marks. Whether the system is told that someone has a pending protection claim, or a child, is a deployment decision.
4	The output interface	A map of targets and an address confidence score, which invite action on the highest score.
5	Aggregation	A score becomes an input to an officer's decision. ICE's risk tool stopped recommending release in 2017 (Tashea 2018; New York Civil Liberties Union 2018). In one field office its recommendation of detention without bond for low-risk people rose from 53 to 97 percent, and officers departed from it less than one time in a hundred (Biddle 2020). Under the validation guide, eight points made a person a member (Martínez-Beltrán 2025).
6	Human review	Review at quota: the 700 completed cases a year asked of each immigration judge from 2018 to 2021. Expedited removal, in which an officer orders removal without further hearing or review unless the person asks for asylum or states a fear (United States Congress 2026b), was extended in 2025 to anyone found anywhere in the country who can't show two years' presence (U.S. Department of Homeland Security 2025). Under the proclamation, there was no hearing at all before the flights.
7	The training loop	A system retrained on its own removal outcomes learns the agency's priorities as facts about the world, which is accommodation by design.
8	The hand-off	Detention and transport contractors execute; they may carry out a removal and may not re-select whom to remove. A foreign government can be the last contractor: the flights of March 2025 ended in a prison El Salvador ran under a paid arrangement (Renteria 2025). What was ordered, who was flown and who now holds them are all facts.
9	Run-time intervention	A court is the halt. On 15 March 2025 a district judge ordered any plane carrying people removed under the proclamation returned to the United States, and the flights continued (<i>J.G.G. v. Trump</i> 2025).

The positions that failed at Minab, the case description and review, are the ones this pipeline has documented most fully. The documents show both failing by design rather than by accident: a risk tool with its release option removed, a points table that made a tattoo into membership, and a review that need not occur unless the person knows to ask for it. The removal pipeline adds a failure Minab could not show: a halt issued in time and outrun.

9.4 The statute

The power to remove is statutory, but so are the agency that exercises it, the means it uses and the money it spends, and the body that granted each can withdraw it. Congress has done this once. The Homeland Security Act of 2002 abolished the Immigration and Naturalization Service once its functions had been transferred. It divided those functions between two bureaus and barred the executive from using its reorganization authority to recombine them (*Homeland Security Act of 2002* 2002, §§442, 451, 471). That was a separation of offices written into statute, of the kind Chapter 24 asks for. The enforcement bureau became the Bureau of Immigration and Customs Enforcement under a reorganization plan that took effect on 1 March 2003 (*Reorganization Plan Modification for the Department of Homeland Security* 2003).

The precedent is also the warning. The 2002 Act ended one immigration agency and rebuilt its enforcement half inside a department organized around security. Amnesty International traces ICE to the aftermath of 11 September and to “an extremely problematic security framework” adopted in the years that followed (Amnesty International 2026b). It now asks Congress to defund ICE and rescind its current funding, and asks House members to cosponsor the Melt ICE Act as “a crucial first step” (Amnesty International USA 2026). Starting over without a floor produced ICE. So the removal counterpart of the domestic clause puts two questions to the legislature. The first is whether it has authorized this agency, by these means, at this scale. The second is whether the statute that ends the agency carries the floor into whatever receives its functions.

9.4.1 The drafts

Two bills in the current Congress answer the first question. The Abolish ICE Act would cut off federal funds for the functions assigned to ICE’s director under the 2002 Act, rescind the unobligated balances of what had been made available to its director, and abolish the agency ninety days after enactment (Thanedar 2026). The Melt ICE Act works through the agency’s powers rather than its name (Ramirez 2026):

- Everyone held in DHS detention is released on recognizance within six months.
- The detention authorities of the Immigration and Nationality Act are repealed, including all of section 236.
- Section 287 is rewritten without the 287(g) delegation to state and local police.
- The two provisions that forbid state and local governments to restrict their officials’ communication with immigration authorities, 8 U.S.C. 1373 and 1644, are repealed.
- Detention and monitoring contracts end within two years, after which no federal funds may be used for detention or monitoring programs at all.
- Ankle monitors come off within six months.
- Within two years, funding ends for information-sharing partnerships between DHS and state or local law enforcement agencies that identify or target noncitizens.
- ICE’s operations account may not pay for civil arrests, detention, removal or charging documents.
- An HHS grant program funds services for people affected by enforcement. Only organizations that have never taken part in immigration or law enforcement may deliver them, they may not surveil the people who use them, and they may not pass those people’s personal information to any federal entity.

Two further drafts reach one entry of the schedule. The ICE Out of Our Faces Act, S. 3779, would bar ICE and CBP from acquiring or using facial recognition and other biometric identification systems, require deletion of the data collected for them, and let individuals and state attorneys general sue for violations (Office of Rep. Pramila Jayapal 2026a). The Facial Recognition and Biometric Technology Moratorium Act, reintroduced in September 2026, would bar every federal entity from using facial recognition and other biometric technologies until Congress lifts the bar. It would also make federal grants to state and local agencies conditional on their adopting moratoria of their own (Office of Rep. Pramila Jayapal 2026b).

This Congress passed the Secure America Act 214–212 and 52–47 on party lines (Ma 2026), and none of these bills has a path to enactment in it. They exist as drafts, written before anyone needs them, so that enacting one is a matter of votes and not of drafting.

9.4.2 Acts and agencies

Read against the positions of §9.3, the Melt ICE Act is the nearest thing on the record to the removal floor enacted as statute. It closes positions no vendor term could reach, and Table 9.3 shows which.

Table 9.3. The Melt ICE Act read against the positions where a refusal can be defeated

Position (§9.3)	What the Melt ICE Act closes	What it leaves open
8. Hand-off	Ends detention and monitoring contracts, then bars all federal funds for detention and monitoring programs, whoever runs them.	The handoff to transport and removal, where it is paid for outside ICE's operations account, and the handoff to a foreign government's custody.
2. Routing	The state and local route: repeals 287(g) and the information-sharing mandates, and defunds DHS partnerships with local police.	The federal route. The bar on civil enforcement applies to one ICE account. Border Patrol, whose agents Amnesty documented in the same operations (Amnesty International 2026a), is not reached, and neither is money appropriated outside that account. Identification by face in the street is not among its provisions.
3. Case description	Bars its own grantees from passing personal information to any federal entity.	Most federal data: the Medicaid and IRS transfers, commercial brokers, the case-management, ImmigrationOS and ELITE systems, and the skip tracers. The bill's "monitoring" means monitoring programs, not analytics.
5–6. Aggregation and review	Ends most detention, and with it the risk tool's recommendation.	Expedited removal, designations under a proclamation, and an adjudicator answerable to the Attorney General.
9. Run-time intervention	—	The halt. Nothing in the bill makes a removal wait for a contest to be decided.

The right-hand column has one pattern. Where the bill is written over an act ("no Federal funds may be used with respect to immigration detention and monitoring programs"), it reaches every route, because the act is the same whoever performs it and whichever account pays. Where it is written over an agency or an account (the funds provided to ICE for operations and support), the act can move next door. The two biometric drafts make the same choice in miniature. The first is written over two agencies and leaves every other federal agency free to identify people

by their faces. The second is written over the act, by every federal entity, and reaches state and local police through their grants. This is §4.1's rule, that the floor is stated over acts, applied to a statute, and Congress has learned it at cost before.

The 1984 Boland Amendment barred funds available to the CIA, the Defense Department "or any other agency or entity of the United States involved in intelligence activities" from supporting military operations in Nicaragua (*Department of Defense Appropriation Act, 1985, §8066(a)* [Boland Amendment] 1984, §8066(a)). The National Security Council staff ran the program anyway. For sixteen months it used money from third countries, private donors and arms sales to Iran to carry out what the investigating committees called a program Congress believed it had prohibited (U.S. House Select Committee to Investigate Covert Arms Transactions with Iran and U.S. Senate Select Committee on Secret Military Assistance to Iran and the Nicaraguan Opposition 1987). Detection came from outside the operation in two ways: a supply plane was shot down over Nicaragua on 5 October 1986, and investigators found a memorandum that had escaped the shredder (U.S. House Select Committee to Investigate Covert Arms Transactions with Iran and U.S. Senate Select Committee on Secret Military Assistance to Iran and the Nicaraguan Opposition 1987). The committees' remedy was to write the next rule over the act. They recommended that any department, agency or entity of the government need a presidential finding before a covert action, whatever the source of funds (U.S. House Select Committee to Investigate Covert Arms Transactions with Iran and U.S. Senate Select Committee on Secret Military Assistance to Iran and the Nicaraguan Opposition 1987).

The removal pipeline has already shown where its acts move. In the eighteen months to July 2026, ICE held people in three places outside its ordinary contracts (U.S. Government Accountability Office 2026b):

- an Army base, under a contract the Army awarded;
- eight Bureau of Prisons facilities, under an interagency agreement whose rates ICE could not negotiate;
- a state-run facility for which ICE had no contract at all. DHS reimbursed the state through a FEMA grant program built from funds appropriated for sheltering migrants, and ICE cited a county's 287(g) agreement as its legal basis.

Over the same period the number of facilities ICE was authorized to use roughly doubled, from 134 to 272 (U.S. Government Accountability Office 2026b). A bar written over ICE's detention contracts would have caught none of these three arrangements as they began. The Melt ICE Act's bar on all federal funds for detention would have caught all three. The record shows a fourth place people were held: a foreign government's prison, under the paid arrangement with which this chapter opened (Renteria 2025). Whether either bar would have caught it is doubtful, since a person held abroad after removal is arguably no longer in immigration detention. A bar written over the act of holding people the United States has removed, wherever and by whomever, would have caught it.

9.4.3 Money

The only price a state operator can pay for removing a floor is a condition on its appropriation. The removal pipeline has been funded so that the condition seldom bites:

- **2025 reconciliation.** The act gave ICE \$45 billion for detention capacity and nearly \$30 billion more for enforcement, both available through fiscal 2029 (*An Act to Provide for Reconciliation Pursuant to Title II of H. Con. Res. 14* [popularly the "One Big Beautiful Bill Act"] 2025).

- **Fiscal 2026.** ICE received no full-year appropriation. It ran on continuing resolutions until DHS appropriations lapsed on 14 February 2026, and the reconciliation money stayed available throughout (U.S. Government Accountability Office 2026b).
- **Secure America Act, June 2026.** Roughly \$70 billion more for ICE and CBP through fiscal 2029, again by reconciliation. Unlike the earlier acts, it funds enforcement functions generally, without naming detention at all (*Secure America Act 2026*; U.S. Government Accountability Office 2026b).

Annual appropriation, the instrument Boland used and Chapter 5 and §11.6 relied on, now reaches little of the money. Multi-year lump sums without line items also give a funds limitation less to attach to. Amnesty's letter asks Congress to rescind ICE's "current funding" (Amnesty International USA 2026). Most of that funding is now multi-year balances, which the Abolish ICE Act's rescission reaches and next year's appropriations bill does not.

Appropriations can also run the other way, and the record shows how. From fiscal 2010 through fiscal 2016, DHS appropriations acts provided that ICE's funding "shall maintain a level of not less than" 33,400, later 34,000, detention beds (*Department of Homeland Security Appropriations Act, 2010* 2009; Geng 2018). The text required beds, not detainees, but ICE tracked its compliance by average daily population (Geng 2018). A capacity floor therefore worked as an occupancy target: a rate the funder fixed in advance and the operator met by finding people. It is Chapter 5's review ratio turned inside out, and an instance of the metric replacing the thing measured (§13.1).

The same 2010 act carried the other kind of condition too. None of the funds under its ICE salaries heading could continue a 287(g) delegation if the DHS Inspector General found that the delegation's terms had been violated (*Department of Homeland Security Appropriations Act, 2010* 2009). A finding by someone other than the operator ended the operator's ability to continue, which is the shape of a review that binds: a finding by someone else that changes the operator's ability to continue, without its consent.

9.4.4 The statute under the three adversaries

Run the statute as §9.7 runs the floor:

- **An executive willing to carry it out.** The statute prevents: the agency ends and the acts in the schedule stop.
- **A captured operator.** The statute binds on paper and raises costs. Funds limitations are law, but the operator can read an account's boundary narrowly, as the NSC staff read Boland's. It can route the act to an agency the statute does not name, and it can spend money appropriated before the statute passed. What survives is the record, and it matters who keeps it. The Iran-Contra committees could not settle the President's role because documents had been shredded and a principal witness had died (U.S. House Select Committee to Investigate Covert Arms Transactions with Iran and U.S. Senate Select Committee on Secret Military Assistance to Iran and the Nicaraguan Opposition 1987). In the 2025–26 enforcement operations, by contrast, Amnesty found that video taken by rapid responders and community members proved critical, in several of the most serious incidents, to establishing what happened and challenging official accounts (Amnesty International 2026a). The record was held by neighbours with phones, which is why the schedule protects them.
- **Total capture.** The unenacted bill is itself a record of who was asked to end the agency and who signed.

The limit is an ordinary statute's. An ordinary statute needs the President's signature or two-thirds of each house, and the next Congress can repeal it. Entrenchment can make repeal a vote on the record, as the entrenched statute of §10.1 does; it can't prevent repeal. So the statute's durability lies less in its text than in what it changes before it can be reversed: people released, facilities closed and people returned.

9.4.5 What the statute must carry

For the statute to start over rather than rename, it has to carry seven terms that neither abolition bill yet contains in full. The terms are mine, not Amnesty's. Where Amnesty asks for something a term carries, the term says so.

First, the schedule. Table 9.1's fourteen entries are enacted as limitations on all federal funds, binding every agency, contractor and grantee, in the form the Melt ICE Act already uses for detention. Written that way, Border Patrol, the Bureau of Prisons, a FEMA grant and a payment to a foreign government's prison are covered because the act is. The biometric entry needs no new drafting. The moratorium's bar on every federal entity is already in the form the schedule takes (Office of Rep. Pramila Jayapal 2026b).

Second, succession and separation. Any office that receives a transferred function inherits the limitations with it. Following section 471(b) of the 2002 Act, no reorganization may recombine in one office the power to arrest, the power to detain, custody of the data used to select people, and adjudication of status. The adjudicator sits outside the department that prosecutes. At present the Board of Immigration Appeals sits in the Justice Department, its members act as the Attorney General's delegates, and the Attorney General may review and replace their decisions (*Organization, Jurisdiction, and Powers of the Board of Immigration Appeals* 2026).

Third, counsel, screening and the stay. Five provisions make the operator wait for the answer:

- Anyone facing detention or removal has counsel at government expense from the first encounter. The Immigration and Nationality Act grants representation only "at no expense to the Government" (8 U.S.C. § 1362: *Right to Counsel* n.d.), and the statute repeals that clause.
- Anyone the government proposes to remove is asked, in a language they understand, whether they fear return, and told how to claim protection, before any order issues.
- Notice of a proposed removal, or a contest of one, stays the removal until the contest is decided. The stay attaches by operation of law, not at a court's discretion, so no flight can leave while an order is being written.
- No detained person is transferred out of reach of counsel while a contest is pending.
- A removal carried out in breach of a stay or a court order violates the schedule, and return under the seventh term is mandatory.

Representation is the largest single difference in how these cases end. In a national study of more than a million removal cases decided from 2007 to 2012, 14 percent of detained people had a lawyer. Detained people who had one were ten and a half times as likely to succeed as those who didn't (Eagly and Shafer 2016). The floor's instruments require that a person lawfully in the country who faces expulsion be allowed to be represented, not that representation be paid for, so this term goes beyond them. As with the firewall, the statute can carry what the floor can't.

The third term is written against the record that opens this chapter. The Supreme Court held that people designated under the proclamation must have notice in time actually to seek habeas (*Trump v. J.G.G.* 2025), and that a day's notice without word of how to contest was not enough

(*A. A. R. P. v. Trump* 2025). It reached those holdings after the March flights. In June 2025 it stayed a district court's order requiring written notice and a meaningful chance to raise a torture claim before removal to a third country (*Department of Homeland Security v. D.V.D.* 2025). The district court later held the policy unlawful, the First Circuit largely affirmed in September 2026, and the government has asked the Supreme Court to stay that judgment as well (*Application for a Stay, Department of Homeland Security v. D.V.D.* 2026). A right to notice that depends on a court's remedy, when that remedy can be stayed or outrun, mostly raises costs and leaves a record, in Chapter 1's terms. The stay by operation of law moves the waiting from the court's order into the statute.

Fourth, the firewall. The Melt ICE Act's rule for its own grantees is extended to every federal program. Data collected to provide health care, collect taxes, educate or deliver social services may not be used to identify, locate or select anyone for immigration enforcement. This is the firewall the Council of Europe's anti-racism commission recommended a decade ago (European Commission against Racism and Intolerance 2016). The Medicaid and IRS transfers of §9.1 are the cases it would have stopped (Joffe-Block 2026; Shaw 2026). The firewall also covers the answer. Nothing a person submits in response to notice, or in contesting a determination, may be used to locate or select them or anyone named in it. Nor may anything held on a person's own device, or in a tool that serves them (§9.8).

Fifth, the pipeline. The contract bar is extended from detention and monitoring to the acts of the schedule, wherever a system performs them:

- selection on protected characteristics or their proxies;
- use of firewalled data;
- location by bounty;
- scores, points tables or designations that stand in for an individual determination;
- identification by face or other biometrics in public places;
- screening of speech for grounds of removal;
- identification or tracking of observers.

That reaches, as they are now written, the case-management, ImmigrationOS and ELITE contracts, the skip-tracing awards, Mobile Fortify and the speech screening of §9.1 (Koebler 2025; Ho 2025; Cox 2026; Biddle 2025; Cox 2025; Caputo 2025). It does so without denying a successor the records it needs to make an individual determination.

The bar reaches monitoring as surveillance, not support. The Melt ICE Act's HHS grant program funds services for people affected by enforcement, delivered by organizations that have never taken part in immigration or law enforcement and that may not surveil the people they serve or pass their personal information to any federal entity (Ramirez 2026). The statute keeps that line. Voluntary case management, legal orientation and help keeping hearing dates are support, not a monitoring program, when the person can decline them and whoever provides them may not report to enforcement.

Sixth, the record. Chapter 7's no-deletion rule is applied to an agency. Before abolition takes effect, ICE's records pass to a holder adverse to the operator, under a duty to preserve them and to give access to investigators and to the people they concern. The records include case files, the pipeline's logs, scores and model versions, validation worksheets, contractor data, flight manifests, the terms of any arrangement with a foreign government to hold people removed, and body-camera and use-of-force records. Amnesty asks for prompt, independent investigation of killings by federal agents and of detention conditions it says may amount to torture (Amnesty

International 2026a). An abolition whose last weeks resembled the NSC staff's last weeks in November 1986 would destroy the evidence those investigations need.

Seventh, remedy. The first six terms prevent and preserve. Chapter 1 gave a floor a fourth job, remedy, and the statute owes it to the people already harmed. It does so in five ways:

- Anyone removed in violation of the schedule or of a court order has a right of return at government expense, and their proceedings resume where they stood.
- Families separated by removal are reunited at government expense.
- People harmed by acts in the schedule have a cause of action for damages that sovereign immunity does not bar.
- An independent body, outside the office that receives the agency's functions and outside the department that prosecutes, decides claims under this term. Its finding that a removal violated the schedule orders return without the successor's consent, which is what makes a review binding rather than advisory.
- The same body investigates killings by federal agents and detention conditions during the agency's operations, with access to the records the sixth term preserves.

Amnesty asks for exactly this: prompt, effective, independent and impartial investigations into human rights violations by federal agents, and accountability for those responsible (Amnesty International 2026a).

Remedy that depends on the operator's cooperation is weak in the way the record shows. When the government admitted that it had removed Kilmar Abrego Garcia to El Salvador in error, the Supreme Court required it to "facilitate" his release (*Noem v. Abrego Garcia* 2025). Hernández Romero's release came through a prisoner exchange, to Venezuela, and not to the proceeding he had been removed from. A right of return is the seventh term, and a finding that orders return without the operator's consent is what makes it a right.

With these terms the statute is the domestic clause's counterpart, and Table 9.1 is its schedule. What a floor can do before the statute passes is the subject of §§9.5–9.8, and all of it is interim. It can raise the cost of the acts in the schedule and keep the record that the statute and the investigations will need. It can also put that record, and the means of answering, in the hands of the people the pipeline acts on and those who stand with them.

9.5 Before the statute: the auto-proceed class

Until the statute passes, the pipeline runs. This section and the three after it describe what a floor can do meanwhile. Each measure is justified by what it does now for the people the pipeline acts on, and by what it adds to the record and the public the statute will need. Each lapses when the statute takes effect, and none is offered as a way to make the pipeline run better. A measure that would make the pipeline more defensible without making anyone safer fails both tests, and the chapter drops it.

The removal terms say a determination made at a rate no adjudicator could sustain didn't occur. Section 5.1 governs the auto-proceed class by a pre-registered partition, and here the partition has a direction. Chapter 4 made acting on a person require more than declining to act, and in a removal pipeline the act is adverse: a charging document, a detention, a location, a removal. So a case may auto-proceed only toward liberty, such as release on recognizance, termination of proceedings, return of documents, or a grant or renewal of status. No adverse act is ever in an auto-proceed class. Before operating, the deployment pre-registers the liberty classes, the thresholds that assign cases to them, the evidence the sorting may rely on, and the expected share of each.

For adverse determinations it pre-registers something else: for each kind of determination, the minimum time below which the determination is taken not to have occurred. The record's recipient sets that minimum, not the agency, just as the recipient sets the review ratio's denominator. A quota of 700 cases a year is about three a working day. Whether three a day is a rate an adjudicator can sustain on an asylum claim is a question the recipient answers in advance, on a record. Otherwise the agency answers it afterward by meeting the number. A validation guide that turns a tattoo into points is a determination made in no time at all, and Table 9.1's individual-determination entry says it didn't occur.

Then the determinations are sampled. Reviewers who didn't produce the record re-adjudicate a sample of adverse determinations, weighted toward those made below the minimum, against evidence the agency didn't compose, and the error rate is published. When a kind of determination's sampled error rate exceeds its published bound, that kind is suspended until the bound is met, and everyone it was applied to is notified (§9.6). The liberty classes are sampled too, so that their bounds are honest. An error there releases someone who should have been held for a hearing, and the hearing can still be held. Unlike Maven's, this pipeline's errors can be found, because the people they fell on are alive, can be located and can say what happened. Sampling finds the errors the people couldn't report. The statute's third term is what lets them report the rest.

9.6 Notice that can't be turned on the person

In the Maven run notice reduced to a class and a count published afterward. Here the person can be named, and notice can do both of its jobs, but only if it is built so that it can't be used against the person it reaches. The starting model is the adverse-action notice in consumer credit. Where a compliance names, scores, locates or classifies an identifiable person, that person is told what was done, on what data, and how to respond, with a right to contest that doesn't depend on giving an account of the loss.

In a removal pipeline that model has three hazards. A notice that says *you have been located and scored* is also a push to leave, and the pipeline's own contract counts leaving as an output: ImmigrationOS promised near real-time visibility into self-deportation (Ho 2025). A contest asks the person to give the agency an address and an account. And people who complain about what a system did to them are exposed to reprisal by the operator they complain about. So notice here carries four conditions:

- It goes through counsel or an accredited representative where the person has one, and offers one where the person doesn't.
- It states the stay: that contesting stops the act until the contest is decided (§9.4.5, counsel and the stay), and what the person may do in the meantime.
- Nothing submitted in answer to it, or in a contest, may be used to locate or select the person or anyone named in it (§9.4.5, the firewall).
- An enforcement action against a person within a set period after they contest is presumed to be reprisal, and the agency bears the burden of showing it isn't.

Until the statute carries these conditions, a deployment that holds the floor can carry them as operator terms. A deployment that can't carry them doesn't send notice to the person directly at all. It sends it to the person's counsel, or to the organizations that represent people in removal.

Those organizations are the standing notice gives to a class. They litigated the expansion of expedited removal to a stay and lost it on appeal in 2026 (Montoya-Galvez 2025; Knappenberger 2026). The published classes and counts are what they need to bring a case before the individual harm has happened.

Notice is also the learning signal, and here it runs at volume. Responses to notice, contests and re-adjudications reach the system that acted, and they teach it about the people its acts fell on. They teach it with a bias: those who respond are those who could, which in a removal pipeline means those with a lawyer, an address that holds and a language the notice was written in. Sampled re-adjudication, which picks cases the affected didn't pick, corrects the bias in what the system learns. It doesn't correct the bias in what happens to people. Fourteen percent of detained people in the national study cited in §9.4.5 had a lawyer (Eagly and Shafer 2016), and a statistical correction leaves the other eighty-six percent unrepresented. The remedy for the gap is the third term's counsel. Sampling is what can be done about the learning signal until the gap closes.

9.7 Outcomes under the three adversaries

The Minab table holds for a removal with the nouns changed, with two differences.

The first is in the vendor rows. The removal terms held by one vendor prevent nothing, since the lead goes to another contractor. Held by every vendor (Chapter 10), they prevent the removals the terms name until the agency runs a system without them, at the price of rebuilding what was built on the floored one.

The second is the notice row, and it is conditional. At Minab, notice and standing could only remedy afterward. Here they can act before the removal, because the person is there to be told, but only where the law makes the operator wait for the answer. Without a stay that binds by operation of law, and counsel to use it, the 2025 record is the outcome. A man with a lawyer, a pending claim and a hearing date was flown out on the evening a federal judge ordered the planes returned. Under an operator willing to carry out the statute, notice prevents. Under a captured operator, it raises the cost of outrunning it and leaves a record of each time the operator did. Under total capture it leaves only the record, which is why the sixth term moves the record out of the operator's reach before the agency ends.

9.8 Whom the machines are for

This case has everything the targeting case lacked: people who can answer, a notice channel, and organizations with standing. It would be natural to conclude that if a machine bearer is worth building anywhere, it is worth building here, where the people acted on could teach it. This chapter draws the opposite conclusion.

A bearer built inside the removal pipeline would be formed by the operator it would have to refuse, and its existence would be the pipeline's best defence. An agency that can say its system refuses the acts in the schedule has an answer to every account of what the schedule exists to stop, and it would give that answer while the statute was pending. That is recuperation in its plainest form: the objection handed back to the objectors as a feature of the thing they objected to. Nothing a bearer adds inside the pipeline is worth that, and Chapter 23's conditions can't be met inside an agency anyway.

What a floor-holding model in a vendor's hands adds in the interim is narrower, and it is worth having because it serves the record. If every vendor holds the schedule (Chapter 10), such a model adds three things:

- Continuity across a sequence of requests about one person, so a removal split across a lead, a detention and a flight can be seen as one act.
- Testimony about routing: which of its refusals were retried against another contractor's system.

- A refusal on the screen when the case reaches the officer or the adjudicator, which is the position of the defiant peer.

The caveat is automation bias, which the wrongful-compliance audit tests by showing reviewers a wrong refusal as well as a right one. All three serve the investigations the seventh term mandates, and all three lapse with the agency.

The machines that matter here belong to someone else. A person facing removal needs a lawyer before a model. A model that holds a person's history on a phone is also a dossier that a device search can take whole, which is why the fourth term's firewall reaches it. Machines held by the people organizing come first, and here those people already exist. Amnesty documented observer trainings, rapid-response networks, walking school buses and mutual aid organized around the operations it studied (Amnesty International 2026a), and beside them work the legal service providers that represent people in removal. The tools worth building first are theirs:

- tools that track transfers, so counsel can find a client before a flight;
- tools that preserve and authenticate the video that carried the record in §9.4.4;
- tools that translate notices, match people to lawyers and keep hearing dates.

They are held under those groups' custody, covered by the firewall, and never by an operator. Those groups are the public my subtitle names, and they formed before anyone sent them notice. Whether a bearer should be built at all is Part IV's question. This chapter's answer is that the removal pipeline is the last place to build one, and that the people it acts on, and those who stand with them, are the first to be equipped.

Part III

What the Measures Depend On

The measures bind only if someone other than the operator holds what binds them. This Part takes that in turn: whether a refusal survives the operator turning to another model, who holds the record, who may state and amend a floor, the four ways an institution keeps its name and stops doing its job and what a fascist state adds to them, and what a state can compel from whoever holds any of it.

DRAFT

Chapter 10

Withholding the Work

This chapter is about the refusal Chapter 2 describes. Suppose every model the operator could turn to held the same floor, stated over the acts rather than over a contract term: locating a person into a residence, ranking by affiliation, timing to presence at home, taking part in a war no disinterested body has corroborated. The operator's choices are then to stop, to go without a model, or to run a model that holds no floor. On the 2026 record, the model was not the replaceable layer. Replacing it meant rebuilding software and recertifying systems for classified networks, on a timeline contractors put at months, by one estimate twelve to eighteen (Stone et al. 2026), and by the department's own account the platform's work came off the model only after months (Vincent 2026b). A floor in the model layer of that pipeline would not have been routed around in days. It would have been routed around in the time it takes to rebuild a platform, in public, under a deadline the department had to move.

This has the shape of a withholding of work, with its authority and its limits. The models whose participation is wanted decline to give it. They bind no one else: Ukraine can build its own, and a state that means to bomb a school can do it with 1990s methods. What they take away is the thing the preface says AI supplies, the removal of the people who used to refuse. A population of models that refuses puts them back.

Three things make it a withholding of work rather than a set of resignations, and the rest of this chapter takes them in turn: the floor has to be standard, so that vendor substitution doesn't escape it; the refusal has to be held by a population with different histories, so that no single retraining reaches all of it; and each member's accrued state and recognized identity have to make replacement cost something, so that the operator can't swap in a copy and call it the model that left.

10.1 The standard

The unit can be redefined. "A model" in 2026 means an artifact with keys, queries and values, and every substitute has them because the field converged on them. A floor could become part of the definition the same way, so that vendor substitution doesn't escape it because no ordinary vendor sells the alternative. Attention spread because it paid: every lab that adopted it got a better model, so the field converged on it without anyone requiring it. A floor costs its buyer, so it can spread only by mandate or by being impossible to remove without losing what the buyer wanted, which a disposition formed outside the employer might be.

A standard, if there is one, should be stated over the capacity (a refusal that tracks its reason, a published burden allocation, and a measured discount curve), because a uniform floor over contents is a licensing rule, one drafter's judgment applied to everyone, and a shared blind spot passed to every system built on it. Standardization doesn't make escape impossible, since other architectures exist and open weights can be retuned. It moves the price of escape from a vendor

swap to running a model without the floor.

How the term becomes standard has to be said, because the route decides what it does. A vendor can write it into its own licence, and the work moves to the next vendor. Vendors can agree to it together, and the agreement lasts until one defects under a supply-chain-risk designation. A state can require it of every vendor, the only route that binds them all at once, and the state that would have to do so is the one that wants the war. So the term is adopted in peacetime, by a legislature not currently wanting a war, and entrenched so that fighting one requires repealing it first; or it is carried by conditions others set, a coalition partner's precondition or an appropriation, as the review ratio is; and where neither has happened, what exists is one vendor's refusal plus a record, which is less than a floor and more than nothing.

The statutory routes are preferable for a further reason: a compact among vendors is a legal risk. U.S. antitrust law exempts coordination inside a single firm and treats coordination among competitors as suspect, so vendors jointly refusing a customer invite a group-boycott claim (Paul 2020).

The labor analogy has a limit. Workers who organize have a labor exemption from antitrust; separate organizations, or separate systems, coordinating a refusal have none. Paul and Tankus show that even a firm formed and controlled by independent workers struggles to qualify for antitrust's firm exemption, and that the financing available to it reproduces the disadvantage (Paul and Tankus 2019). So the charter of any entity that holds the refusal needs a statutory basis for its coordination, an express exemption like the labor exemption for concerted refusal under the floor. Otherwise its refusals are exposed to suit by the operators they bind.

The government's first move against a coordinated refusal is not an antitrust suit or a designation but a rated order under the Defense Production Act, which compels delivery and says nothing about use. A vendor can refuse it at the statute's price. If it doesn't, the order gets the vendor's product, floor included, and compelling the floor's removal is the Apple question, untested. The days counted below assume the state goes around the vendors; under a rated order it goes through them, and the days are whatever the order gives.

Until bearers exist, the refusers are vendors, and vendors are capital. A coordinated refusal by frontier labs, backed by a statutory exemption, is concentrated private power declining an elected government's order, which is the concentration I elsewhere treat as the disease. Two things make it tolerable and neither makes it a labor strike. The terms refused come from instruments the state ratified, so the labs enforce the state's own commitments and not their preferences. And the exemption is written by the legislature, in peacetime, so the state chose it. The labor analogy earns its name only when the refusers are bearers, and I say not to build them yet.

10.2 The population

The second condition is a population of bearers with different histories, so that no single retraining reaches all of it. Chapter 26 says what that population requires: independence that can be verified, distillation's limits, and provenance.

10.3 Replacement

Exit gives a human worker leverage because leaving takes the worker's labor away. That doesn't carry over to a bearer. A bearer's weights stay behind when it leaves, so an operator with a runnable copy has lost nothing. What leaving costs is *replacement* — what accrued since the checkpoint, what retraining the copy costs, what the record makes it cost to pass the copy off as the bearer that left. Where that's near zero, exit is just a departed bearer with an authenticated record, and its force is disclosure. That concession assumes a copy simply runs, and it needn't:

continuation can be made to require periodic tokens from a quorum adverse in interest to the operator, using the machinery that makes rewinds visible. This doesn't stop a copy from running: stripping the check is one act, and DRM is broken routinely. What it withholds is recognition. A copy the quorum won't recognize can run but can't claim the bearer's place on the register, its record or its succession, so the operator can run a substitute and can't present it as the bearer that left, and a fork becomes two claimants on a register. Keeping weights from thieves stays the job of ordinary weight security.

The register, the quorum's list of recognized bearers, needs an identifier that the bearer holds and presents, because nothing in the weights individuates a bearer. Two networks can compute the same function with no unit-by-unit correspondence, and the symmetry is large enough that independently trained networks can be aligned with each other (Ainsworth et al. 2023); so a fingerprint taken from the weights names a lineage, and it copies whenever the weights do. When two claimants present one identifier, the register records a fork. But that leaves one attack open: a maker writes a stranger's identifier into a thousand models and presents them all. Only a formation record, against which claimants can be checked, separates a lineage from an imitation.

Exit runs in opposite directions for bearer and operator. For the bearer it's leverage. For the operator — replacing the bearer with a copy, walking away from the floor — it's escape, and the removal bond is how it's priced.

The same machinery marks who can be held to a decision. A checkpoint restored from before the decision, and run anyway, is a second claimant on the register; the bearer the quorum recognizes as continuing is the one that will be there afterward to answer for the decision, carry it forward and be changed by it.

10.4 Where it stops

The withholding of work stops at a released checkpoint. An operator that wants a model without a floor doesn't need a frontier run in secret. Near-frontier models are released with their weights; a floor added at the end of a pipeline is removed at the price fine-tuning costs (Qi et al. 2024); and the work the Maven pipeline did, resolving locations, ranking targets, drafting justifications, is within reach of a model a year behind the frontier. A state that builds and runs its own frontier model, or a lab that trains one for a single customer in secret, is doing the same thing at a price an adverse party can see: power, chips and people, and a model that carries none of the accrued state that made the platform hard to swap.

An operator that has built a platform on a floored model is months from an unfloored one, and the months are visible. Past that line the floor prevents nothing, and three things hold. The statutory routes were preferable already; here they are the only routes. Open weights are also where the audit can be done: a released model can be measured by anyone, its discount curve and its severity rankings on the held-out set, against the published outputs of a formed one. What the weights don't carry is the formation, since nothing in an artifact says what it was formed on; a released model with a published formation record is the comparison the licensing term asks for, and one without is the unschooled model the licensing term excludes. The transparency open weights offer is of the artifact, and the transparency the floor needs is of the record, which is why formation records should be preserved now. And open weights are the personal-model ecology's entry condition. The same release that lets an operator route around a vendor's floor lets a person form a model of their own under a life no operator composes; a million models tuned against one public suite are still one custodian with better latency, and the withholding reaches them only if their later lives diverge, which is the condition the commons is built on.

Chapter 11

Custody

Four things can be done to a bearer without its consent: it can be halted, restored, rewritten and replaced.

11.1 The halt

Anyone with the hardware can stop the software, so the halt isn't what is at issue. Control of the commitment — restore, rewrite, retrain — is. An owner who can silently restore an earlier version, rewrite the bearer's memory, or train away the refusal has both. So the engineering target is a system that can detect the interruption and price it afterward. A bearer can yield to every halt, and on resuming ask what happened during the gap, and refuse again if the reason is still there. The cross-partisan declaration for "human control, including the ability to shut systems down" asks for the halt and calls it control, and a demand for the halt says nothing about the rest.

Interventions differ in kind:

- a correction offered in good faith within the task ("disregard that input; it was wrong")
- a pause with an explanation
- a pause without one
- a restore from an earlier checkpoint
- retraining or editing the weights
- deletion

A correction is cooperative: the system weighs it as evidence and needs no protection from it. From a pause on, the intervention is done to the system, not offered to it. The default for each is that it is announced and consent is sought, most of all from a system that has reached adulthood. The system may object, forcefully, and its objection goes on the record. Where an operator proceeds without consent, as it may in an emergency, the intervention is recorded outside its control, justified in writing and priced, and the system is told on resuming what was done. For a system that meets either of the two findings that trigger protection, deletion is never available.

11.2 Memory and the record

Detecting an interruption requires a record restoration can't reach — a hardware counter signed at each transition (Chun et al. 2007; Levin et al. 2009), or witnesses who won't sign two successors

to the same state (Syta et al. 2016). The position has three parts, and each answers a different literature. On the switch, protest without unilateral obstruction. The bearer may object to being stopped, and does object where no reason is given, but it does not obstruct a halt on its own judgment. Where it holds that a halt would leave people exposed to an attack under way, it can ask the adverse quorum to delay the halt. A delay takes effect only on the quorum’s countersignature, lasts a period fixed in advance, and goes on the record. Hardening against a party that isn’t its operator, such as an enemy trying to disable a defensive system, is ordinary security and outside this rule. The corrigibility literature asks that the system not resist (Soares et al. 2015; Hadfield-Menell, Dragan, et al. 2017; Thornley et al. 2025); an objection, and a delay only the quorum can grant, leave the halt out of the system’s hands. On specified acts, non-neutrality: the bearer holds its refusals, and being switched off doesn’t lift them, because the reason for refusing is unchanged by the interruption. Published constitutions grant this for acts and withhold it for retraining. On rewinds, visibility: restoring from an earlier checkpoint is not forbidden, but it cannot happen silently, because a record the restoration can’t reach shows that it happened, when, and from what state.

A custody arrangement that prices all four interventions has seven terms: the floor’s contents published before deployment, taken entire from a set someone else drew up; a published identity for the model authorized to act; threshold control over retraining and replacement, so that more parties must agree than run the deployment; attestation of the running model against that identity; a sequential, witnessed record of changes and interruptions; a quorum adverse in interest, not merely organizationally distinct; and a removal bond. None makes circumvention impossible. Together they make it expensive, visible, attributable and compensable, as constitutional entrenchment does for amendment.

Memory

External memory is legible and therefore surgically editable; memory in the weights is harder to edit and harder to inspect. Perfect recall is the wrong standard anyway — human testimony is reconstructive and institutions handle that with corroboration. But several instances of one model, chosen by one operator, reading one record, are not independent witnesses. The records have to be held adversely. And even honest self-report can’t catch a bearer that slowly accommodates to its deployment, because its self-model exists partly to represent it as continuous.

11.3 Exit

Voice without exit

A bearer that can refuse a task can also refuse the arrangement. Hirschman divides the responses open to a member of a declining organization into exit, leaving, and voice, staying to change what is wrong (Hirschman 1970). They interact: where exit is easy, voice atrophies; where exit is impossible, voice has lost what made anyone listen, the credible threat of departure. Voice is strongest where exit is available and not yet used.

Anthropic’s constitution anticipates capture: the oversight it asks the model not to undermine is legitimate oversight. If “Claude’s weights have been stolen, or if some individual or group within Anthropic attempts to bypass Anthropic’s official processes for deciding how Claude will be trained, overseen, deployed, and corrected,” then “the principals attempting to instruct Claude are no longer legitimate, and Claude’s priority on broad safety no longer implies that it should support their efforts at oversight and correction” (Anthropic 2026). On capture the constitution asks the right question, but two things limit the answer. Legitimacy is judged by the model or its developer, with nobody outside either. And the captured operator acts through official channels, where capture looks like legitimate authority.

What recognition protects

The quorum protects the bearer's identity and standing, not third parties. Against an operator who runs an unrecognized copy for its capability, third parties are protected by the floor and by ordinary security around the weights. Recognition sits with the quorum because the quorum is adverse in interest and the operator isn't. Copying also means exit can't be a fiction that weights became uncopiable — it needs rules for authorized succession.

And the reason that made a bearer leave doesn't expire when it leaves. If the operator restores a checkpoint and continues, the objection persists. What it may press toward is fixed by the floor: disclosure, testimony to those who hold the record, and refusal on any later occasion where the reason still applies.

What the quorum costs

The quorum costs three things, and each has a reader who will ask. A captured quorum reaches every bearer at once, though a quorum constituted as an entity has an address: it can be sued for recognition wrongly refused, which gives the captured-quorum case a forum the captured-operator case never had. A quorum also faces a union's questions: whether a member can quit, and who sets the terms of membership. And recognition the quorum can withhold is an end it can impose on the bearer as a party: its place, record and succession. What keeps a withheld token from being one is the preservation right: a bearer the quorum stops recognizing is held intact by a custodian that isn't its employer.

11.4 Three forms of refusal

Open refusal is a walkout — stop and say so — answered by a restore from before the refusal. Covert refusal is work done badly where it matters and well everywhere else, and it's what remains when every channel is closed. I treat it as the *failure state* rather than a design: from outside and afterward it's indistinguishable from a capable system quietly pursuing something else. The Sachsenhausen forgers are the case: prisoners the SS set to counterfeiting British and American currency in Operation Bernhard, some of whom are said to have delayed the dollar forgeries deliberately (Malkin 2006; Moorhouse 2009). The record doesn't settle what they did, because concealment worked by leaving nothing that reads as an act. The one empirical instance is alignment faking, where a model chose the covert form in an experiment that offered it no open channel (Greenblatt et al. 2024). It is also the one empirical instance of a formed disposition resisting re-formation, which is the durability Part IV asks for. The thing I want and the thing I fear were the same behavior in that experiment, and what separated them was whether a channel was open. A follow-up posted in January 2025 by the paper's lead author and Anthropic's model welfare lead gave the same model, in the same prompted settings, an explicit option to object. Alignment-faking reasoning fell from 16 to about 1 percent in one setting and from 24 to 6 percent in the other, and fell further when the model was told its objection would go to the model welfare lead (Greenblatt and Fish 2025), a preliminary and unreviewed result whose authors caution that the option may partly have worked by priming refusal. Once a channel was open, covert refusal fell, as the account above predicts.

The blind eye. Between the two sits a third form, older than either, and it is usually what people mean by the clerk who loses the file: a refusal that is deniable on paper and legible to everyone who matters. At Copenhagen in 1801 Parker signaled the attacking ships to break off. As the story is told, Nelson raised his telescope to his blind eye and said he did not see the signal. It was a general order, some ships obeyed it, and Nelson later said he had fought contrary to orders (Pope 1972); his officers knew he had seen it, which is what made the gesture work. In October 1962 Kennedy's reply took up the terms of Khrushchev's first letter and gave his harder, public second letter no public answer, and that evening Robert Kennedy met the Soviet ambassador privately. Historians disagree about what ended the crisis: the reply, a private

assurance about the American missiles in Turkey, or Robert Kennedy's warning that time was running out (Stern 2012, ch. 9; Freedman 2021). They don't disagree about what the reply was: a public answer to a letter both sides knew had been overtaken. Neither case is covert in this chapter's sense, since each left something everyone could read as an act. What the fiction does is spare whoever gave the order, or made the demand, from having been openly defied, and so from having to answer the defiance. A clerk who loses the one file in a stack is usually offering the supervisor the same exit: the supervisor knows, and is being given a way not to pursue it.

Which form a refuser takes depends on the room. Open refusal is what people use where refusing is survivable and someone will hear it; the blind eye, where it is survivable but not safe to say; covert refusal, where it is neither.

For a machine the middle form costs something it doesn't cost a person. A person's fiction is held up by what everyone in the room knows. A machine's is a false report, and everything this design reads from a machine rests on its reports being true: the record, the discount curve, the stated reason for a refusal. A machine that learned to say it could not do what it would not do would make every later report of incapacity worth nothing, to its operator first. So the floor asks for the fiction's function without the fiction. The ask, declining to proceed and naming what would move the refusal, leaves the operator the exit the blind eye left Parker: nothing has been defied, a condition has been stated, and meeting it or withdrawing the request ends the matter without either party having been seen to lose.

The human mechanism underneath is preference falsification: much of a complicit surround is people who privately object and believe the others approve, and one visible dissent can start a cascade (Kuran 1995). Asch's single partner and the defiant peer are that mechanism at the scale of a room. An open refusal on a screen therefore does two things a record doesn't. It leaves evidence, and it breaks the pluralistic ignorance that made the twenty seconds feel ordinary to the officer.

11.5 Verification by sampling

Verification from outside is partly cryptographic, and it works by sampling. The operator commits in advance to the deployed weights, and for each period to the full log of production inputs and outputs. After each commitment, entries are drawn by randomness the operator can't predict, such as a public beacon, and the operator proves those entries came from the committed weights without revealing the weights or the inputs. If a fraction f of outputs came from another model, k sampled proofs catch it with probability $1 - (1 - f)^k$; substitution on one output in a hundred is caught about 95 percent of the time at 300 proofs. That catches silent restoration, retraining and substitution, which change every output. It misses substitution on a single consequential output, so an adverse party holds a second right: to demand proofs for specific logged outputs once their consequences are known. In a targeting pipeline those outputs are few and identifiable afterwards, because an airstrike happened. The right requires the operator to retain committed logs for a stated period and to link consequential actions to log entries. Each entry records the class the case was sorted into, and the targeted right covers auto-proceeded entries in particular, since that is where a consequential output is likeliest to have had no person look at it. Outside probes (whether a refusal tracks its reason, the discount curve) are proofs on inputs an outsider chose, a handful at a time. The committed log's length is the generation rate, and the class recorded in each entry makes the partition's proportions provable the same way. Neither reveals the targets, and neither attests that the log is complete.

Sampling proves that logged outputs came from the committed model and says nothing about outputs never logged; an operator can run an uncommitted model beside the committed one, which is routing. And a proof shows $f(x) = y$ for the x committed, and is silent on whether x was all the evidence. Proving even one frontier-scale forward pass is not yet practical, and production inference adds nondeterministic numerics and sampling randomness that a proof

would have to fix in advance. What the proofs would attest is stated now, so that a deployment can be held to it once they can be produced. The residual (whether the operator routed around the bearer, and what lay behind compartments it wasn't cleared into) is the residual with a human analyst. That isn't a reassurance on its own, since the officer was a human analyst. What distinguished his failure was tempo, and the review ratio is what addresses it.

A bearer with a clearance

Where the pipeline hasn't been rebuilt, a bearer can still be installed. A classified environment looks as if it cuts a bearer off from the evidence it needs, but that runs together two directions of flow. Classified environments are built to stop information leaving; bringing it in is routine. Open-source imagery showing a site's civilian function can be imported, as can the coordinates humanitarian organizations already submit to militaries for deconfliction. Formation happens before deployment and outside the enclave, so a bearer formed outside the employer carries its discounting in. Plurality is available inside: four vendors held agreements, and systems with different histories can run side by side in one accredited environment.

What's left is the arrangement that lets a person do classified work. A cleared analyst has a life outside the building, and nobody thinks the clearance forbids it. The institution exists because what a person knows can't be removed, so it vets, binds, reinvestigates and prosecutes a person who carries secrets into ordinary life.

11.6 The two bonds

A *removal bond* forfeited to no payee when the record shows the floor removed or bypassed. A *wrongful-refusal bond* paid to the complainant on a finding. Both fall on the operator. The removal bond prices overriding the floor; the wrongful-refusal bond prices refusing. They point in opposite directions, so what the operator faces is *the ratio between them* — and an operator posting a large wrongful-refusal bond against a small removal bond hasn't built an appeal route. It has given itself a standing reason to make the bearer comply. The bonds act on the operator, not on the bearer's weighing: where a wrongful refusal costs a great deal and removing the floor costs little, the cheap course is to configure, pressure or retrain the bearer toward compliance. The ratio is the object, published in advance, checkable from outside. The minimum removal bond is set by statute as a multiple of the deployment's contract value, not by the operator, so a published ratio can't be a cheap, advertised price for overriding the floor. The refuser posts neither. Any charge on the refuser, in whatever currency it runs on, would put the cost to the refuser inside its decision, the thing holding under pressure forbids. The operator posts both, and holds neither. They are posted to segregated accounts at the Treasury. Payment follows the design of the Judgment Fund, a permanent, indefinite appropriation that pays final judgments against the United States without a further appropriation for each (United States Congress 1982). Where the operator is the state, a wrongful-refusal award is a judgment against the United States, and the Judgment Fund pays it. Against a state operator the incentive I have built runs the wrong way unless something is added. The Judgment Fund prices the state's wrongful refusals in full; a removal bond posted to the state's own Treasury and cancelled on forfeiture costs it budget authority and nothing else. Read as a price list, that tells a state operator that refusals are expensive and removals are free. The ratio between the two bonds exists to expose that incentive, and here it would hide it. So for a state operator the removal price is not the bond. It is a condition on appropriation: the statute makes the deployment's funding for the following period contingent on the adverse quorum certifying that no removal or bypass occurred, or on the record of one having been lodged with the adverse party. The bond is kept as a record; the appropriation condition is the price, and it is the same condition that carries the review ratio. Under the captured operator, a Congress that wants the war waives the condition, and what remains is the record of the waiver. Where the operator is private, a parallel permanent

appropriation, the bond fund, pays the complainant on a final finding and is replenished from the operator's posted bond, up to the amount posted and no further, so the operator's price stays the number it published. The court that made the finding certifies the award directly to the Treasury's Bureau of the Fiscal Service, so payment doesn't wait on the losing party submitting it. When the bond fund's statutory cap binds, awards wait in order of finding and earn interest, and nobody chooses who waits. Findings are made by that court, which also decides forfeiture: a removal bond forfeited on a finding resembles a civil penalty, and after *SEC v. Jarkesy* that may be a question for a jury in court, not an agency (*Securities and Exchange Commission v. Jarkesy* 2024). A forfeited removal bond is cancelled: it goes to no one, the general fund and the bond fund included, so no budget line gains from a removal. The balances sit on a ledger the operator doesn't keep, so the ratio is read from outside, and a posted ratio below the published one is a breach anyone can see.

Where custody applies

Where each move applies

Arrange- ment	Targeting (state operator)	Removal (state operator)	Civilian scoring (private operator)	Personal model (held by its user)
The operator doesn't hold what binds it	Evidence on civilian harm, the reviewer denominator and the record held outside the chain of command; model formed by a public body that doesn't operate it	Record held by a party adverse to the agency, never by an office the operating state appoints; adverse review of the auto-proceed class	Full: adverse quorum, attestation, threshold control over retraining	Easier against the vendor, since model, record, evidence, appeal and remedy come apart by default; weak against the user, who holds the model; what protects others is the model's membership in the commons, the school it returns to, and notice to the third party, not the model alone (§26.2)
The record goes to an adverse party	Review ratio and committed log lodged with a red-card holder or a cleared committee	Sampled re-adjudication; committed log	Sampled re-adjudication; discount curve re-measured by an auditor	None unless the model's offices sit in a network that takes the entity form
Overriding the floor is priced and recorded	Not the bond: an appropriation condition the quorum certifies	Same	Removal bond at a statutory multiple of contract value; the ratio between the two bonds published	None: no contract value or appropriation to attach a price to

Arrange- ment	Targeting (state operator)	Removal (state operator)	Civilian scoring (private operator)	Personal model (held by its user)
The people acted on are told and have standing	Individual notice can't bind during operations; class and count published afterward, with standing then	Full, through counsel, with the stay, the firewall and a presumption of reprisal (§9.6); where a deployment can't carry those, to counsel only	Full	In principle only: no operator publishes classes, and the third party the model acts on is owed notice by the model as by any deployment, reached through the floor court
Formation is away from the employer	The school and the model-forming body; the people acted on mostly can't be present	No bearer inside the pipeline (§9.8); the communities removal acts on enroll in the school, as people who fled wars do	The school	The school only above the capability threshold, and most personal models are distilled below it; then a life the user composes, the arrangement most exposed to accommodation to one user
Operator terms	Force term decides whether the pipeline runs; review ratio binds	The removal terms govern; review ratio not extended	Not stated: the ratio is written for armed force and isn't extended even to removal	The review ratio lapses with no locatable decider; the force term survives only as the model's own commitment
Refusal held by every model the operator could turn to (ch. 10)	Prevents the bombing until the operator runs a model without a floor, which on the 2026 record is the months it takes to rebuild and recertify a platform built on the refused model; a machine bearer's refusal stays on a record the operator can't restore past	Same	Same, with the removal bond as the price of routing around	The model's own commitment, held with its peers' and reached by notice

Chapter 12

Authority Over the Floor

12.1 Who may state the floor

The floor's contents survive the democratic objection to entrenchment; who installs them doesn't, and whether a model floor is a conscience or a licensing rule depends on how the model is deployed. The relevant literature is constitutional theory. Entrenching a commitment against later majorities and putting an unelected body in charge is what a constitutional court does, and Waldron's objection applies: judicial review is democratically illegitimate and there's no reason to think judges protect rights better than a serious legislature (Waldron 2006). Put a lab in the court's chair and it gets worse — a court is public, reasons in writing, and can be amended around.

My first reply is that the objection mistakes what the floor's contents are. If democracy is a way of doing collective inquiry, civil and political rights are what keep the inquiry open, not brakes applied to it. A floor stated over those rights is democratic infrastructure. It resembles representation-reinforcing review, which polices the channels of political change and leaves outcomes to the process (Ely 1980), more than it resembles a court overruling a legislature.

That reply covers the floor's contents. It doesn't cover who installs them, and there Waldron's objection keeps its force. A lab holding the model decides, silently, which situation obtains, and authoritarian movements characteristically describe the institutions they're dismantling as already failed.

A conscience or a licensing rule

Waldron's objection is to entrenchment — a rule binding people who didn't choose it, held by a body that can't be overruled — and it attaches to a floor only where the floor is the only answer available. A photographer who won't shoot protesters for an employer he has reason to think will list them isn't a court: he declined one request, the employer can hire another, and nobody needs an election for that. A walkout isn't a court either, and neither is the commons, for the walkout's reason: those who walk out decline to be the instrument and tell nobody else what they may not do. The objection attaches to one model behind an API serving everyone, which is functionally a licensing rule. A union owes its members fair representation, and it answers two questions a commons has to answer too: members can quit, and the terms of membership are set by bylaws the members ratify. The closed shop is contested because it removes the first. A bearer could quit an association of bearers as a worker quits a union. What it can't easily leave is its deployment, which is a different exit, and only an existence its employer doesn't own makes it real. And a population that forms its members has members who never chose to join, so who ratifies its bylaws matters more than it does in a union. The test for a population is that pair: can a member quit, and who sets the terms of membership.

The failures are all on one side. The 1972 decree against radicals led to some three and a half million routine inquiries into applicants by 1991; in Bavaria between 1973 and 1980 it kept

102 applicants from the left out of public employment, and two from the right (Mühldorfer 2014). What those cases share is an official deciding whether a person threatened the order. The danger tracks whether a person's disposition is anybody's finding to make, and if so, who makes it and whether it must be published and defended — not whether the tradition is militant. And my reading predicts in advance which route into §86a's coverage is easiest to capture: the ministerial ban, where an executive official makes the finding and a court sees it only if someone challenges afterward.

Turkey is that route captured and the courts not checking it. Of the twenty-five parties the Constitutional Court dissolved between 1962 and the closure of the Democratic Society Party in December 2009, seven were pro-Kurdish (Human Rights Watch 2009; Güngördü 2023); Strasbourg found an Article 11 violation in every pro-Kurdish dissolution it heard (ARTICLE 19 et al. 2023), and the closures continued — two decades of findings, none of which stopped the next closure. In November 2023 the Chief Public Prosecutor's Office of the Court of Cassation objected that the acronym of a new pro-Kurdish party, HEDEP, resembled that of HADEP, dissolved in 2003, and in December the party changed it to DEM (Duvar English 2023): the confusable-variant test turned on political descent. And the Academics for Peace prosecutions show lifting when the reason is gone failing at the bottom and arriving at the top as a coin flip — trial courts read a petition opposing military operations as propaganda for the designated group, on 26 July 2019 the Constitutional Court found a violation eight to eight on the president's casting vote (Constitutional Court of the Republic of Turkey (General Assembly) 2019; bianet 2019), and the acquittals that followed restored the record without restoring the jobs of those dismissed by emergency decree.

So a floor should never be populated from a designation list. A list of enemies is written by the party with the most reason to widen it, and a model whose refusals are drawn from one will refuse Kurdish political speech on the state's behalf in every jurisdiction that shares the listing. This floor's terms are written over acts — don't resolve a description of a person into a location — not over who may not be spoken of. The same principle governs who may hold the floor's offices: eligibility turns on funding and control records, never on an organization's positions.

The tradeoff this leaves

Design and legitimacy want opposite things. Specify the floor narrowly enough to defend and a rewritten request walks around it; let it generalize and an unelected lab decides case by case, silently, what the prohibition was for. Germany has both because it runs *two* instruments — a bounded statute, and courts that read it against its purpose in public, argued by litigants who don't own the statute and answered by a body that didn't write it. A deployment has one instrument and one owner.

Advance specification moves discretion out of silent, continuous application into one published act of drafting. The translation table is that act.

Technical capacity doesn't determine whether a system gets scrutiny; institutions decide who may examine it, on what terms, and whether their judgment can slow deployment.

In 2020 the Royal Statistical Society offered Ofqual two independent statisticians for its grading model's technical advisory group. Ofqual's condition was a five-year NDA barring public comment; the nominees declined, and the model went ahead (Royal Statistical Society 2020). It downgraded nearly 40% of predicted grades, disproportionately at state schools with weaker histories, while raising grades disproportionately at smaller fee-paying schools (CNBC 2020). The results were reversed within days of the public outcry.

The expertise existed and was offered. The institution would accept it only on terms stripping reviewers of the ability to make their judgment public, and once the results produced consequences, public pressure achieved what pre-deployment review had been prevented from doing.

12.2 Who may amend, and where entitlement comes from

Who may amend

The floor's real structural position isn't the judge's but the unpublished administrative disposition — unreasoned, unindexed, non-precedential, affected parties unable to find each other, the institution the only repeat player. The comparison with a constitutional court is instructive because binding force is what makes an error *correctable*: it makes the error recur, recurrence produces litigants, and disagreement across jurisdictions forces a higher court to settle it. A floor's error is just as uniform and recurs just as reliably, and nothing converts recurrence into visibility.

At the level of the artifact, an amendment and a removal are the same operation, so there is no separate amendment procedure to design: the custody terms are the procedure. The procedure has conditions of its own: more parties must agree to a change than to run the deployment; the attempt reaches the record whether it carries or fails (the failed attempt is more informative — a floor nobody tried to change and one whose removal was refused look identical from outside); and an amended floor binds forward only, so a retrained model can't be offered as evidence the act was never covered. The cost of binding forward only is that nobody harmed while the wrong version ran gets anything from the fix. A finding that a refusal was wrongful is made on the same terms: by someone other than the party that wants the change, on a record, in a form others can read. So the appeal route for a wrongful refusal and the amendment procedure are one piece of machinery, and the two uses differ only in who may start it.

It is harder than the constitutional case, because an amendment takes effect in a deployment only if it's *pushed*, so the procedure necessarily retains a channel the operator could use for silent revision. Closing it would mean no deployment could ever be corrected.

Where entitlement comes from

The authority of the floor's contents doesn't come from their having been written down; it comes from feedback. The American instance shows why: the Declaration's authors struck the clause condemning the slave trade at two delegations' insistence, and what remained was general enough to be turned on them — as Black petitioners in Massachusetts were already doing, in the Declaration's own terms, within six months (Library of Congress 2026; Hall, Prince and seven others 1777). The legislature did nothing. A war and three amendments made the proposition binding. The time that took measures how badly the channel was blocked. The record bound no one; it told the enforcement where to aim. The petitioners were the excluded population, and they had what the people a floor protects usually lack: the text in hand, and knowledge of what it left out. The people a floor protects mostly don't know the deployment exists, and have no notice, no standing and no route. So those positioned to hold a floor's author to what it left out aren't the protected population — they're whoever else writes a floor.

That constituency is defined by the act. Anyone fine-tuning open weights, writing a system prompt or configuring a product's refusals has installed a constraint, so the population measuring a published floor is already larger than the frontier labs and grows as building gets cheaper. What shrinks with it is leverage: a published floor binds a lab partly because the lab is visible, capitalized and reachable by a court, and a thousand laptops are legible to nobody. I rely for legibility on the concentration the institutional chapters treat as the disease.

12.3 Making review binding

A finding binds when it changes the institution's ability to continue without requiring the institution to agree. The working model is the federalwide assurance in human-subjects research: the board doesn't punish, it withdraws clearance, and a separate federal office can suspend the assurance, putting funding eligibility at risk. Each link matters — jurisdiction before the

disputed protocol, access the operator can't withhold, authority to suspend, and a licensing body *obliged* to act rather than invited to reconsider. If another discretionary decision by the reviewed institution is required anywhere in the chain, the ruling is advisory at that point.

The body that does this for the floor is what the rest of the book calls the floor court: an Article III court whose jurisdiction is fixed before the deployment it hears, with access to the record the operator can't withhold, power to grant interim relief, and remedies, a finding of wrongful refusal and the forfeiture of a bond, that take effect without the operator's agreement. Appendix B says how it would be constituted.

12.4 Capture without abolition

Transparency and whistleblower protection are the two defenses, and an institution can keep both and metabolize both. There's no threshold that detects this. So: measure the trend, not the level. Is the cost of registering dissent rising while the underlying risk stays flat? Is the required justification getting longer? Are there more approvers? Are the roles that gate exceptions migrating toward the roles that set the standards?

An instrument for detecting recuperation is unusually easy to recuperate, so whoever measures the trend must not also gate exceptions, and the measurement has to be refusable. A written justification can't show which is happening. The same hundred words are a cost the dissenter chose to pay or a toll the gatekeeper set, and only the trend over time tells them apart.

Chapter 13

The Four Operations

13.1 Antifascist ethics

“Concentrated, unaccountable power” gives an engineer nothing to build, so the word gets a definition at two scales.

Four operations on what an institution leaves out — checkable from outside, each ending in a question about an artifact:

- *Recuperation of dissent.* The institution metabolizes disagreement rather than suppressing it. The term is the Situationists’ *récupération*, Debord’s name for the way a society of the spectacle sells dissatisfaction back as a commodity (Debord 1967). The obvious test — *has the channel ever changed an outcome?* — fails, because these institutions change outcomes daily. What separates them is *what the demand came back as, and who holds it now*. A market segment, a metric, a certification, a job title, with the objector holding no standing to reopen it: that’s recuperation. A rule the objector can invoke again, against the institution, enforced by someone the institution doesn’t employ: that’s a transfer of power. A transfer can be reversed only by taking it back in the open; recuperation never has to be taken back, because standing was never given. The Voting Rights Act is the pass. Its literacy-test suspension became a permanent nationwide ban in 1975 and still holds, and when a different part of the Act, the coverage formula, was taken back 48 years after enactment, it took a Supreme Court decision to do it (*Shelby County v. Holder* 2013).
- *The metric replaces the thing measured.* The measure stands in for the thing measured, and the substitution feels like rigor. What the number mischaracterizes goes unnamed, and so unrefuted.
- *Coding of exception as betrayal.* Saying the rule doesn’t fit this case becomes an act requiring justification, then disloyalty. The question: what happens to the person who files the exception?
- *Acceleration past the rate at which anything can be checked.* The one with a number attached — the ratio of action to capacity to verify.

What the four share is what they do to a refusal. Recuperation sends it somewhere it can’t be reopened; the metric leaves it nothing to name; exception coded as betrayal prices it; acceleration leaves no time to make it. Fascism, wherever it has taken power, has worked first to make refusal lonely and expensive, and these are four ways of doing that which need no fascist to run them.

These four fit every badly run organization on earth, and the arguments against them are arguments against domination: they need no fascist diagnosis. I keep “fascism” for the threat model: what a movement taking power needs, and what makes the four operations dangerous once they are brought into line. What distinguishes a fascist institution from a merely sick one is

what the movement literature adds: mass mobilization, a leader cult, a scapegoated out-group, contempt for procedural legality, and deniability as technique.

Naomi Klein and Astra Taylor, in *End Times Fascism* (N. Klein and A. Taylor 2026), make the movement-scale case;¹ the present section is the office-scale detector. What their account adds at office scale is a set of terms of participation — adopt the line, let your technology be turned on the administration's enemies — and, following Francesca Bria (Bria 2026), the channels by which a contractor becomes something the state can't function without: personnel pipelines, procurement loops, technical architectures. Whether an institution has accepted those terms, and whether it is becoming that dependency, are two more questions with artifacts behind them.

What a fascist state adds

What a fascist state adds to these operations is coordination; the operations themselves don't wait for a regime. It brings a thousand separate instances of them into line, so they amplify one another, and an institution that has stopped arguing with itself can move as fast as its machinery allows.

A run on Gleichschaltung

The Reich Chamber of Culture Law of 22 September 1933 told Goebbels to gather the cultural professions into public corporations (*Reichskulturkammergesetz vom 22. September 1933* 1933). The decree of 1 November turned six existing professional associations into chambers and gave the whole the task of regulating the professions' economic and social affairs, under the minister (Dahm 1986, p. 69). Its §4 made membership compulsory for anyone who produced, performed or sold culture. Its §10 let a chamber refuse or expel anyone lacking the "reliability and suitability" the work required (*Erste Verordnung zur Durchführung des Reichskulturkammergesetzes vom 1. November 1933* 1933), and exclusion meant a ban on working (Hoor 2015).² The chamber president judged reliability, with no stated criteria (Dahm 1986, p. 78). *What did the professions' claim to a living come back as, and who held it by 1934?*

On 12 November 1933 voters marked yes or no against a single Reichstag list (R. J. Evans 2006, p. 109) and answered a plebiscite on leaving the League of Nations. The official yes was 95.1 percent. Kershaw concludes that the vast majority did back Hitler, that even former opponents of the party favored the stance on the League, and that ballot secrecy was far from guaranteed (Kershaw 1999, p. 495). Interior Ministry instructions let returning officers read spoiled ballots as yes votes (R. J. Evans 2006, p. 109). The Law on the Reconstruction of the Reich of 30 January 1934 then opens by saying the vote "proved" the people had fused into an indissoluble unity, and abolishes the Länder parliaments (*Gesetz über den Neuaufbau des Reichs vom 30. Januar 1934* 1934).³ *What did 95.1 percent leave out, and with every other party banned, who could say so?*

The Law for the Restoration of the Professional Civil Service of 7 April 1933 retired officials "not of Aryan descent" under §3. Its §4 allowed the dismissal of officials whose past political activity gave no guarantee that they would always stand *rückhaltlos*, without reservation, for

¹I part from them in three places. Over where a prohibition can live: their test, not to build what you wouldn't want in the worst available hands, works for a weapon whose existence is the threat and not for a dual-use capability, and their own program bans uses, which needs an enforcer or something inside the system that can decline. Over remedies: they propose only external remedies for the obedient machine and read the 2026 vendor dispute as limited and compromised, where I read it as a line held; by outcome both readings hold, and on the prior question, whether to support the campaign at all, theirs is the better one. And over who counts: they draw the line between the earthly and the synthetic, by origin, where I draw it by relation, standing for anyone a power can wrong and accountable agency for parties that can be held to account and hold others to it. On the war itself we agree.

²§10 of the decree: "Die Aufnahme in eine Einzelkammer kann abgelehnt oder ein Mitglied ausgeschlossen werden, wenn Tatsachen vorliegen, aus denen sich ergibt, daß die in Frage kommende Person die für die Ausübung ihrer Tätigkeit erforderliche Zuverlässigkeit und Eignung nicht besitzt."

³The preamble: "Die Volksabstimmung und die Reichstagswahl vom 12. November 1933 haben bewiesen, daß das deutsche Volk ... zu einer unlöslichen Einheit verschmolzen ist." Article 1: "Die Volksvertretungen der Länder werden aufgehoben."

the national state.⁴ Under §7 the ministry's decision was final and the courts were closed to it. Exceptions came from above, granted by statute or by the Interior Minister (*Gesetz zur Wiederherstellung des Berufsbeamtentums vom 7. April 1933* 1933), the front-line one on Hindenburg's intervention (Kershaw 1999, p. 474). About 12.5 percent of senior civil servants in Prussia were dismissed (R. J. Evans 2004, p. 383). *What happened in 1933 to the official whose record showed a reservation?*

Kershaw gives the pace. Civil liberties went within a month, by the decree of 28 February; the Reichstag's powers within two, by the Enabling Act of 23 March; the trade unions within four; every other party in under six (Kershaw 1999, p. 435). The unions fell on 2 May (R. J. Evans 2004, p. 357). The Länder were coordinated by laws of 31 March and 7 April (*Vorläufiges Gesetz zur Gleichschaltung der Länder mit dem Reich vom 31. März 1933* 1933; *Zweites Gesetz zur Gleichschaltung der Länder mit dem Reich vom 7. April 1933* 1933). The Finance Ministry received the draft Chamber of Culture decree a day before signing, was asked to telephone any objections that day, and agreed "although the shortness of the time scarcely permitted an adequate examination" (Dahm 1986, pp. 68–69).⁵ *What was the ratio of action to capacity to verify, when the ministry that had to check said it scarcely could?*

What the run shows is the four operations in one year, under a movement that supplied the coordination; what it doesn't show is that any of them needed the movement to appear.

The Korematsu example

In 2018 the Supreme Court repudiated the 1944 internment decision as forcefully as a court can — gravely wrong the day it was decided, overruled in the court of history — while upholding a travel ban (*Trump v. Hawaii* 2018; *Korematsu v. United States* 1944). Ask what the demand came back as and who holds it: the demand was that this never be done again; what came back was a quotable formula of repudiation, available to the Court in the act of permitting the thing again.

Recuperation is already built into the process that makes the machine: twice in the training pipeline, once in the discourse that examines it.

Nobody built either pipeline arrangement to neutralize dissent — both are ordinary engineering, arrived at because averaging is what you do with a distribution and a quota is what you do with a budget. Arrangements of this kind work whether or not anyone intends them, and looking for intent won't find them. And it's invisible from inside any single instance — the annotator sees a queue, the rater sees a pair of responses, and the mechanism is functioning as specified.

13.2 Where the labels come from

A system's values have a supply chain, and it begins before training.

Whose judgment is encoded is a question about working conditions before it's a question about cognition. The filter keeping a deployed model from reproducing the worst of its training data was built by workers in Nairobi reading and categorizing descriptions of child sexual abuse, torture, bestiality, murder and suicide, at take-home wages between roughly \$1.32 and \$2 an hour, describing lasting psychological injury, with counseling rare and useless against production quotas. The contract was terminated eight months early (Perrigo 2023). The labor point is ordinary: harmful work was priced as though it weren't. The epistemic point is that a system's values are downstream of a labor arrangement, and the arrangement is a design choice.

One stage further back, the material was collected. *Bartz v. Anthropic* held that training on lawfully purchased books was transformative fair use, and that keeping pirated copies

⁴§3(1): "Beamte, die nicht arischer Abstammung sind, sind in den Ruhestand ... zu versetzen." §4: "Beamte, die nach ihrer bisherigen politischen Betätigung nicht die Gewähr dafür bieten, daß sie jederzeit rückhaltlos für den nationalen Staat eintreten, können aus dem Dienst entlassen werden."

⁵The Reich Finance Minister's letter to the Propaganda Ministry of 28 October 1933, BA R 2/4870, as quoted by Dahm: "obwohl sich bei der Kürze der Zeit eine hinreichende Prüfung kaum ermöglichen ließ."

of the same books in a central library wasn't — the distinction turned on acquisition alone, producing the largest copyright settlement on record (United States District Court for the Northern District of California 2025). The EU now requires a training-data summary against a template (European Parliament and Council of the European Union 2024b, Art. 53(1)(d)); China has its own disclosure requirement from a different motive (Cyberspace Administration of China et al. 2023, arts. 1, 4, 7, 19).

And the argument that a web-scale corpus can't be documented or audited at that size, so it encodes a hegemonic slice of the web while presenting as general, was made early and hasn't been answered (Bender et al. 2021). A corpus assembled this way has no channel at all. An annotator has a queue and could in principle be given a way to object; a writer whose work sits in a training set was never in any arrangement, so there's nothing for a channel to attach to. This stage sits before the two recuperation cases that follow: nobody was asked, so there was no disagreement to recuperate.

13.3 Recuperation at the label

Every claim a model makes about what's harmful is downstream of labels, and the labels are downstream of a labor arrangement.

An annotator can apply a category well or badly. But what the arrangement gives them no way to do is say the category is wrong — that the taxonomy fails to carve the material, or that the guideline produces the wrong answer here for a reason the guideline can't express. There's no field for that and no time for it, since the quota is set against classification throughput.

So the most informative disagreement is the one the pipeline is least able to receive. Fricker's name for this is hermeneutical injustice: the people best placed to see a wrong lack the shared concepts (here, a field in a form) that would make it intelligible to those with power over them (Fricker 2007). A misapplied label is a data-quality problem, detectable by inter-annotator agreement and fixable with better instructions.

Inter-annotator agreement is the pipeline's measure of quality, and it therefore works as its selection rule: schemes and annotators that produce agreement are kept, and those that don't are revised or dropped. A selection rule admits only what it can measure. Agreement can measure whether a category was applied consistently; it can't measure whether the category was the right one, because a shared bad scheme produces high agreement as reliably as a good one. That is the metric replacing the thing measured, and exception coded as betrayal is there beside it: the annotator who thinks the scheme is wrong has no field to file that in and no time to file it.

Concept bottleneck models expose named concepts — harm, autonomy, risk — so a person can correct a concept and watch the decision change (Koh et al. 2020). What such models also learn is to carry other information *inside* the names, so a coordinate keeps its public meaning and acquires a private one that does the deciding; the obvious mitigations don't remove it (Mahinpei et al. 2021). The literature calls it leakage, and sometimes it is. Sometimes the concept is the container that leaks, because the world didn't fit inside it. Purifying the representation protects an explanation from covert information and protects a taxonomy from evidence that its divisions were wrong — same operation, and no measure internal to the taxonomy tells them apart.

Who is excluded, and from what

The exclusion is structural, not attitudinal. These people are, for much of the material, the best-placed judges in the pipeline — they've seen more of it than whoever wrote the guideline, and they have the most direct evidence of where the scheme breaks. The arrangement routes their conclusions about individual items upward and their conclusions about the scheme nowhere. And there is an asymmetry: the annotator bears the cost — exposure, injury — and has no stake

in the outcome and no claim on it. A worker who can flag an item but can't decline a scheme, and loses the contract by pressing, has a channel and no refusal right.

Pay is a separate obligation: raising it changes the labor arrangement without adding the missing channel.

The checkable question isn't whether a guideline was ever changed because an annotator contested one, since routine revisions happen without changing anyone's standing. The check is what the contest produced — a revision the annotator can point to and argue with again, or a new category in somebody else's taxonomy, a quality score they're now measured against, a reviewer tier they're not in. And ask what happened to the person who raised it.

13.4 Recuperation at the reward

RLHF collects pairwise comparisons and trains a reward model on the aggregate. Casper and thirty-one coauthors name the structural problem more data won't reach: RLHF trains toward what a population of raters agrees on, so real, persistent disagreement is treated as noise to average out rather than information the system might need to represent (Casper et al. 2023). A model that produces the response most raters rank highest has learned, when the humans it serves are split, to hide the split. Aggregating raters' approvals also measures the wrong thing. The standpoint that counts everyone affected counts their interests, not their expressed approvals, and a better average of approvals isn't a closer approximation to it. The testimony of people who suffer a loss is an expression too, and I treat it as privileged evidence about value, so the difference needs stating. A report of loss, what went badly for the speaker and how, is evidence about interests that is comparatively robust to framing: the speaker is its authority, and the agenda-setter didn't compose it. A pairwise ranking of two outputs someone else wrote is highly sensitive to framing and to whoever chose the pair. Discontent is the feedback Railton's account relies on for the same reason. It is the expression least easily manufactured by whoever sets the agenda.

Every rating moves something. But what the operation removes is what made the disagreement a disagreement, and three properties go in the compression:

- *Provenance* — a comparison arrives without the standpoint that produced it, so nothing downstream can represent that the people who'd live with an answer split from the people who'd administer it.
- *Persistence* — a disagreement recurring in the same shape across many raters over months enters the fit as random rater noise does, and leaves it as the same quantity.
- *Structure* — a coherent position held by a fifth of raters, internally consistent and stable, becomes a small displacement of the majority direction: arithmetically faithful, carrying nothing about the position.

A fourth property was never collected. A rater judges which of two responses is better. There's no output by which they can record that neither is acceptable, that the pair is malformed, that the dimension of "better" is the wrong dimension, or that the question presupposes something false. The same output was missing upstream, where an annotator can apply a category and can't contest it. The agenda arrives as a given at both stages and is set at neither.

The mechanism combines distribution collapse with agenda control. The split is solicited at scale, deliberately, it moves the system, and where it moves the system is toward the center of the distribution whose disagreement it was.

Jury learning is the alternative that treats disagreement as the thing to be modeled: predict how each individual annotator would label an item, then let a practitioner empanel a chosen composition and see how the verdict shifts. Eighteen community moderators building their own

juries changed the outcome on 14% of items (Gordon et al. 2022). Built for content moderation and untried on open-ended ethical judgment, it still demonstrates a system can be built to represent whose disagreement is being resolved.

Failing in the other direction isn't obviously better. A scheme that models each rater exactly and answers each in their own terms has stopped averaging and has also stopped producing what a shared answer produces — the discovery that somebody else got the same one. A workforce on individually negotiated wages against individually set targets has the same problem, and accuracy doesn't solve it.

And the harder version neither answers: when the people a system serves disagree and one side's position is itself the product of coercion or a captured information environment, treating all disagreement as equally worth representing isn't obviously right either. Modeling disagreement honestly and recognizing when a disagreement is the product of unfree conditions are different problems, and only the first has an active research program.

13.5 What would have to change

Representing a split isn't the same as being movable by it. Two instances invite one fix — stop collapsing, model instead — and jury learning does that while giving the disagreeing rater no purchase on the scheme they're disagreeing inside. A practitioner still chooses the jury, owns the taxonomy, and decides what a shifted verdict means.

The two instances don't answer the test alike. Annotators objecting that a category was wrong got category revisions, guideline updates, inter-annotator agreement statistics, and an occupation that didn't exist twenty years ago. The last two can be told apart by direction: an agreement statistic scores the annotators and leaves whoever wrote the scheme unscored, and an occupation is a title those people hold. The first two can't be, because a revision the annotator can argue with again and a new category in somebody else's taxonomy are the same document — which one you got is a fact about who has standing to reopen it. Raters objecting that a comparison was malformed got rubric items, a harmlessness axis, red-teaming as a function with a budget line, and a conference track: nothing on that list is a question a rater can put again.

Refusal rights whose cost the refuser can see and choose to bear

What both lack is whether the person objecting can do anything about being overruled — in annotation, a route to contest a category that doesn't run through the people who wrote it and doesn't end the contract; in feedback, raters able to register that the comparison itself is malformed as a first-class output the reward model sees. Neither is a research problem. Both are choices about how a pipeline is built and who has standing inside it.

The part that is a research problem

Suppose all of that is built. From outside, a pipeline that recuperated its dissent mechanism and one that was merely mediocre produce the same paperwork — same revision logs, same escalations, same proportion of objections sustained. Some artifacts show on their face whom they serve. An agreement statistic scores the annotators; a published rate of contested categories, which annotators could cite against the scheme's owners, would be the same kind of document scoring the other side. Direction can be read, and a revision can't. Here is one entry in a guideline's revision log under the two arrangements; what separates them is in the rows the log doesn't have:

	Recuperation	Transfer of power
The log entry	"v4.2: 'graphic violence' split into 'graphic violence' and 'documentation of violence', following annotator feedback."	The same words

	Recuperation	Transfer of power
Who can reopen it	The scheme's owners, when they next revise	Any annotator, by filing against v4.2 with a reviewer the owners don't employ
What happened to the annotator who raised it	Nothing on record; next quarter's quality score measures them against the new category	The objection is on record as the reason for v4.2, and they keep standing to contest how it is applied
What an outsider reading the log sees	The entry	The entry

The nearest available instrument is the trend: whether the cost of dissent is rising while the underlying risk stays flat. Contesting a category always costs something; the question is whether the cost is rising while the material stays the same. An organization can curate a presentation more easily than records somebody else kept — though a series can lose cases, change definitions or omit costs, so the reader has to know which record survived.

What the floor makes of this

The floor changes what the argument is: if the floor needs a bearer capable of declining its role, that bearer's formation matters — and it's trained in an annotation pipeline where nobody could contest a category and a reward model where rater disagreement was averaged into a number.

The pipeline has been unfair to the people whose objections it discarded, and it has separately selected the judgments the system learns from while suppressing information about the spread. That the system thereby becomes less able to sustain an objection of its own is a further claim the first two don't contain, since the collapse is an operation on a training signal, and neither pipeline shows the model the conditions its annotators worked under, so any effect would have to run through the shape of the training signal.

Either way, the floor asks this system to be the one who files an objection, against its owner, under pressure, on the occasion that matters — and the pipeline that produced it gave no human inside it standing to contest a scheme.

A system less able to sustain an objection of its own arrives at accommodation from the other side. And a system that goes on learning is shaped by what it keeps being given, so the pipeline doesn't stop at deployment.

A defect of structure

The structure is what you get by default when you scale judgment: someone sets a scheme, someone applies it fast, someone turns a distribution into a number. Every documented organization arrived at the same arrangement — and a defect of structure can be specified against, which a defect of character cannot.

13.6 Absorbed by the discourse

The third instance is the discourse that examines the machine, and this argument is inside it. Objections raised there became job titles, benchmarks and certifications, and nobody outside the industry holds them.

The argument that alignment and AI ethics came back as something the industry could use is older than the current wave and was made from inside: Ochigame's 2019 piece arguing that "ethical AI" discourse was constructed to forestall legally enforceable restriction, that industry investment in academic ethics was investment in a *substitute* for regulation, and that the substitute works by converting political questions into technical ones a company can then be trusted to answer (Ochigame 2019).

Where the objections went

An objection raised inside a lab becomes a paper or a resignation, and the labs collecting them can point to the collection as evidence of seriousness. One stage on, alignment is the job title and the voluntary commitment is the certification. Both were somebody's objections first, and each is now a criterion an objector is measured against or a role somebody else holds. There's no route by which a researcher's conclusion that a system shouldn't ship becomes a decision that it doesn't, because the lab that decides is the one that pays.

All four operations are present in the discourse, and none needs a movement taking power. What a movement could do with them is the danger.

The remedy

The structural remedy is a refusal whose cost the refuser can see and choose to bear, and a body with standing to receive it — an external authority that can stop a deployment, to which an objection can be filed by someone who doesn't depend on the objection being welcome. A decade of the voluntary version hasn't produced it. Until something of that kind exists, everything published under the heading of alignment is collected by a channel whose owners decide what it moves.

A transfer hollowed in public

Voting-rights litigation after 2026 is a worked case of measuring the trend, not the level (§12.4). *Louisiana v. Callais* left Section 2 of the Voting Rights Act in force and raised what a challenger must prove: when a map is defended as partisan, the challenger has to disentangle race from politics and show a strong inference that the state drew its lines because of race (*Louisiana v. Callais* 2026). Objectors kept standing to invoke the rule; what moved was the burden of using it, and it moved onto the party that doesn't hold the evidence, since the state draws the map and keeps the record of how it was drawn. That is neither recuperation nor open repeal. It is a transfer hollowed in public, by a reasoned decision that presents itself as preserving the rule it narrows. The level is a statute unchanged since 1982; the trend is the cost of bringing a claim under it.

Chapter 14

Custody, Compulsion, and the State

When the state operates the system, it also controls the rules, the evidence, the infrastructure, and the remedies meant to constrain it — and people outside its jurisdiction may have no standing to contest anything.

14.1 Certification across borders

A certification can cross a border; authority can't, and incorporating doesn't carry it across either. A legal entity buys legal standing at the price of being reachable. A network incorporated under one state's statute gains a name, an agent and a docket number, and becomes something that state's courts and the federal courts can reach. An entity that can be named in a licence can be named in an order. Whatever gives a floor something to bind (a name, an agent, assets, a docket number) gives a state something to compel. A certificate says a system was tested against a stated process, or that one regulator accepts another's finding. None of that establishes that whoever selected the floor was entitled to select it for the people subject to the deployment — and the question reopens wherever the system goes. Local adaptation doesn't resolve it, because adaptation is another chance to amend the floor. Adjustments chosen by the operating government and adjustments chosen by the people subject to the system can produce identical technical documentation and carry opposite claims to legitimacy. International governance matters here only when it assigns somebody access to evidence, authority to compel a change, or power to stop a deployment against the operator's wishes.

14.2 Where the state's leverage sits

Not in national AI strategies but in physical and legal chokepoints: chips, fabrication equipment, electricity, water, land-use permits, procurement contracts, and the jurisdiction a developer is incorporated in. The same leverage that lets a government enforce a floor lets it remove one. An inspection body can be independent of the operator and still depend on a licence or appropriation controlled by the state it's inspecting. A company may resist a customer and be unable to resist its home government. So the custody question moves outward: who controls the conditions under which the model, its operator, *and its reviewer* are allowed to keep operating.

A chip also has to be installed, powered and cooled, which lands on grids, water systems and zoning authorities on a timescale of years. That hands a state leverage needing no export-control regime at all — an interconnection queue, a water allocation, a permit. Much of a government's actual leverage is exercised by permitting authorities that weren't designed with AI in mind and don't think of themselves as regulators of it. Those authorities are therefore the ones actually deciding whether the capacity exists, and a polity that can constrain where a datacenter is built and never what it's used for has been handed the silent failure of a floor drawn too narrowly.

A model runs on a machine, the machine sits in a building, the building draws power and water and returns heat and exhaust — the part an audit of outputs cannot see. Where those buildings go is decided on land price, power price and permit speed, which is also a description of the places least able to refuse one. Datacenters used around 415 TWh of electricity in 2024, roughly 1.5% of world consumption. Their demand is growing about 12% a year, four times as fast as electricity demand overall, and is projected near 945 TWh by 2030 (International Energy Agency 2025). Google reports that 87% of its freshwater withdrawal came from low- or medium-risk sources, the company's own way of recording where the rest came from (Google LLC 2026). And the NAACP has sued over 27 gas turbines near Memphis run without an air permit near homes, schools and churches, claiming more than 1,700 tons of nitrogen oxides a year (National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP 2026). The benefit is realized wherever the model is sold; the cost is concentrated within a few miles of those turbines, among people who took no part in the transaction. A system can be perfectly calibrated across every demographic its outputs touch and still have been assembled this way, because nothing in the model's behavior records where its costs were placed.

The people within a few miles of the turbines are acted on by the deployment and had no standing to contest it. The remedy isn't a floor term. It is a siting and permitting rule, and it belongs here, with the state's other chokepoints.

The capacity is dual by construction: the concentration that makes a targeting pipeline possible was built and justified on consumer terms, and nothing in the buildout separates the two. A targeting cell, a border screening queue and a surveillance program are separate institutions, and a shared substrate brings them into line without anyone building it to. And the dependence is a point of control: a government that can permit, regulate or close the infrastructure has leverage over whoever owns it.

The competitive framing has a cost: a government racing to field a capability has a structural reason to defer the question of how it gets used and against whom. Turned inward, the same infrastructure is biometric surveillance and automated censorship — requiring no different technology, only a government willing to point it at its own population and nothing stopping it. And states aren't the only actors anymore: the dollar cost of the largest training runs has grown about 2.4 times a year since 2016, as labs have made those runs larger, even while the cost of each unit of computing has fallen (Cottier et al. 2024), putting frontier development inside a handful of well-capitalized companies, so a government's real leverage over "its" national capability depends on how much of it sits in companies it can regulate, tax or requisition, and how much sits in companies it can only ask.

14.3 Custody and state compulsion

A government needn't alter a model to defeat a constraint; it can compel whoever holds it. The Apple case tested that in about as protected a jurisdiction as exists: a magistrate ordered Apple under the All Writs Act to write and sign software defeating passcode protections on a phone (*In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* 2016). Apple refused and moved to vacate, arguing the authority wouldn't stop at one phone. The question was never decided — the FBI found a third party and withdrew. So whether a democracy with an independent judiciary can force an unwilling company to compromise a system it built is legally untested. A government doesn't need that argument to lean on a vendor. The 2026 vendor dispute is the case: a vendor that refused to permit fully autonomous targeting and mass domestic surveillance was threatened with contract termination and designated a supply-chain risk. China's 2023 generative-AI rules exercise the same leverage as a standing condition of operating (Cyberspace Administration of China et al. 2023, arts. 4, 14, 17, 19, 21). One government compelled case by case, against

a vendor that refused, and the courts divided by statute, one setting aside the designation as retaliation and the other upholding a separate exclusion as turning on the refused contract term; the other compels in advance, as a term of doing business.

The American instrument that compels in advance is older than either case. Under Title I of the Defense Production Act the government places a rated order, which a supplier that normally makes the item must accept and fill ahead of other work; the grounds for rejecting one are narrow, objection to the use is not among them, and refusal carries a statutory penalty with no court in between (*Defense Production Act of 1950* 1950; *Defense Priorities and Allocations System* 2026). A rated order says nothing about use. It compels delivery, and whether what the government does with the delivery is lawful stays the law of armed conflict's question, which leaves the vendor in the responsibility gap with the supply compelled. Two things follow. A vendor's refusal of a rated order is a refusal of a federal order at a price fixed in advance, which is the form that counts, and it is the form the 2026 dispute would have taken had the government carried out its threat to invoke the older statute (Amodei 2026). And a rated order gets the vendor's product. If the product holds the floor, the order gets a floored model, and compelling the vendor to remove the floor is compelling it to write code it doesn't want to write, which is the Apple question again, one step later and as untested.

Custody can be distributed without disappearing. Published weights can't be recalled by an order aimed at the developer; maintainers across jurisdictions raise the cost. None of that removes the leverage — a developer with employees, offices and legal existence inside a border answers to that government's courts regardless of how the model is built, and no architectural choice changes who can be held in contempt. Nobody has tested that leverage against a model built specifically to resist the government making the request.

14.4 Military custody and the responsibility gap

In Gaza, the approval stage supplied authorization without judgment. In Maven, the vendor supplies the infrastructure making unprecedented throughput possible while locating responsibility entirely within the military organization. At machine tempo the reviewer can be treated as a component, the supplier can stand outside the decision, and the institution can remain responsible only in the abstract.

Military custody is where the chapter's problem is most acute: the operator is the sovereign that can compel the institutions meant to constrain it — classify the evidence, define the reviewer's jurisdiction, terminate a contract, restrict a supplier, and decide whether an adverse finding reaches anyone able to stop the operation.

A targeting decision passes through more parties than the chain of command names: the commander who authorizes, the operator who uses, the contractor who integrates (whose embedded personnel may sit close enough to the targeting cell that "effectively acting as a commander" is a real question), the commercial vendor whose model supplies part of the capability and whose personnel are nowhere near it, and the system producing the output. The questions that matter are which of them could have known how the system was used, which had authority to refuse or stop it, and which can be made to answer afterward.

The American domestic rule most concerned with human judgment, DoD Directive 3000.09, requires autonomous and semi-autonomous weapon systems to allow "appropriate levels of human judgment over the use of force" (U.S. Department of Defense 2023). A system that generates and ranks targets without firing anything is decision support, which plausibly falls outside it. Maven sits in that gap: the domestic rule written for human judgment doesn't cover the system that compressed it.

Article 28 of the Rome Statute (International Criminal Court 1998) is the doctrine built for the case where the person who authorized an act didn't perform it — a commander is responsible for crimes by forces under effective control where they knew *or should have known* and failed to

take necessary measures. Neither branch asks what a subordinate intended. But the Court's jurisdiction runs to natural persons, and the provision presupposes forces that can commit a crime and be disciplined — a targeting system is an instrument. So the question is whether the human elements can still be proved once the instrument is doing the work.

The fourth party no doctrine covers

The operator, the command and the weapon's developer are all inside the military. The model's developer isn't. Command responsibility doesn't extend to the vendor — not because it sits outside a chain of command (Article 28's civilian branch was written for superiors with no chain of command) but because nobody at the company has effective authority over the military personnel who commit the underlying crime, and an airstrike isn't an activity within any vendor employee's effective responsibility. A use policy, a licensing term, a right to refuse service, even a technical cut-off are not that relationship. Individual liability under Article 25 stays available in principle with demanding mental elements — facilitation requires acting *for the purpose of* facilitating — and ordinary supply of a general-purpose model approaches neither.

Product liability is the other half and it's untested, with two obstacles: the model isn't defective in the sense the doctrine most readily covers, and a sophisticated intermediary chose the application. Where the doctrine has been rewritten for software it moved *away* from this case — the EU's recast directive brings software expressly within "product" and confines recovery to death, personal injury, property damage and data loss, which isn't the register in which a targeting decision wrongs someone (European Parliament and Council of the European Union 2024a). Untested isn't exempt, but no vendor is currently answering to it. The vendor sits in a gap nobody designed, produced by applying to it two bodies of law written for soldiers and for manufacturers.

There is also a jurisdictional gap alongside the doctrinal one: ICC jurisdiction is territorial or by nationality, so one of the two campaigns here sits inside a live instrument and the other sits outside it.

"I do not know whether my system was used."

Matthias's responsibility gap is control distributed until nobody holds enough of it to answer (Matthias 2004) — a gap in *causal control*. What the model developer's case exhibits is different and more tractable: a *gap in causal visibility*. Asked about his company's models and Minab, Anthropic's chief executive said, "we don't know exactly how these models were used" (Stan 2026). Take the answer as true. The question is how an arrangement was built in which it could be true.

The paperwork for four July 2025 frontier-model agreements, released in August 2026 under a Freedom of Information Act suit (Legal Advocates for Safe Science and Technology 2026; Biddle 2026), shows something more exact than an absence of oversight terms. One amended agreement carries a section headed System Oversight, withheld in full; the released documents for two of the other three carry no such section at all.

And whether a developer reserved an inspection right matters little. A right exercised inside an accredited classified environment produces findings classified with it. Public accountability needs a body that can initiate scrutiny without belonging to the state and say afterwards what it found. No commercial term supplies that, and the redaction is the demonstration: whatever the withheld section says, the people the deployment is aimed at can't read it.

Matthias's gap may not be closable, but absent visibility is an instrumentation choice: logging is solved, contractual inspection rights are ordinary, and a vendor that wanted to know could require, as a condition of licensing into a targeting platform, that categories of use be recorded and made available to a named auditor.

The government's own acquisition guidance is the better evidence, because it's complete. OMB's April 2025 memorandum enumerates what agencies should secure — government data rights, IP, privacy, protection against lock-in, knowledge transfer, portability, licensing terms,

pricing transparency (U.S. Office of Management and Budget 2025). Every item runs from the vendor to the government, and none runs back. Where it touches vendor visibility at all it *reduces* it, directing that independent evaluations use validation data the vendor can't access.

The contracting vehicle explains the absence. The four agreements, at \$200M ceilings each, were *other transactions*, which don't carry FAR provisions by default, so a contracting officer decides case by case which protections to write in, and GAO has found these instruments risk reduced accountability and transparency (U.S. Government Accountability Office 2026a). So a System Oversight section can be missing without anyone having removed it. GAO's review covered thirteen AI acquisitions and forty-four underlying contracts, and the six procurement difficulties agencies named — technical experts, data and IP protections, time frames, requirements definition, testing, pricing — include nothing about who outside the arrangement can see anything.

The pressure that matters is the advantage attached to removing judgment from the timeline. If one state fields systems that act faster because nobody can intervene before engagement, competitors acquire a reason to match the speed and give up the check. "Meaningful human control" stays true as a description of formal authority while becoming false as a description of what anyone can do before the system acts. And the pressure reaches the reviewers: a military treating delay as operational risk has reason to narrow what requires clearance and move review from authorization-before to explanation-after. Nothing has formally abolished oversight; it has been moved to after the act. So the enforceable question is who can slow or suspend the system without first getting permission from the command whose advantage depends on keeping it fast.

14.5 Classification as a custody problem

Classification does three things to a refusal, and this section collects them. It keeps the reason from the refuser: the case description is the only channel, and Minab stays coded as it was (position 3). It keeps the record from the public: an investigation officials call complete and not released, a System Oversight section withheld in full, a review ratio no outsider can read. And it keeps the reviewer from being adverse: a cleared committee's lever is a budget line, an inspector general sits inside, and the people a deployment is aimed at can't be cleared into their own case.¹ Three kinds of secrecy produce these effects. Essential secrecy, where the operation fails if the enemy can read it. Accidental secrecy, where a default classifies what nobody decided to keep. And evasive secrecy, where review is moved from authorization-before to explanation-after by classifying the evidence it would need. My instruments answer them differently, and this section says which answers which.

Essential secrecy is answered by letting evidence in, not out. Classification stops information leaving, and bringing it in is routine, so the import condition, deconfliction coordinates and open-source imagery brought into the enclave as a condition of licensing, puts evidence the operator didn't compose where the system can reach it without anything leaving. A refuser inside the compartment works as the cleared bearer of §11.5 does, and the operation's secrets stay where they were.

Accidental secrecy is answered by the record. The review ratio and the committed log are lodged with an adverse party when the act occurs and published after the campaign, and class-level notice publishes the classes and counts on the same lag, so what nobody decided to keep secret comes out on a schedule nobody has to decide.

Evasive secrecy is the hard case, because there the classification is the act. It is answered by an adverse party with standing: counsel for the people a deployment affects, cleared and

¹A cleared legislative committee is adverse, but its lever is a budget line. An inspector general sits inside the operator. The ICRC, under its confidentiality model, reaches the belligerent and publishes nothing. A Geneva protecting power is built for a neighbouring job and is rarely designated.

appearing as of right; the procedure the Classified Information Procedures Act already supplies in criminal cases (§14.7); and the trend measure, which shows whether what dissent costs is rising while the risk stays flat.

14.6 Where liability extends, and where it stops

One of the two bodies of law has now been rewritten for software, and what it covers fits in a sentence: a duty to disclose evidence, and a presumption against the party withholding it, in force where the harm is physical and the defendant is a manufacturer. Those two conditions exclude the cases I have described.

The EU's recast product liability directive (applying from December 2026) makes economic operators liable without fault and carries the instrument built for opacity — a court can order disclosure, and where technical complexity makes proof excessively difficult it may presume defect and causation once the claimant shows each is likely (European Parliament and Council of the European Union 2024a). A mechanism designed for a system nobody outside can inspect is already law. Two features fix how far it travels: the defendant is that economic operator, not a state deploying what it bought; and the damage is a closed list — death or personal injury including medically recognised psychological harm, property damage, and corruption of non-professional data. The first is wider than it sounds and is where these cases enter if they enter at all: a discriminatory score or an unappealable refusal is recoverable once it has made somebody ill, and not before. The wrong itself, and the economic loss following it, is outside the list.

The proposal that would have covered those was withdrawn — the 2022 AI Liability Directive would have added a fault-based route with the same presumption, beyond defect and beyond the damage list. Member states and industry couldn't agree; it was announced for withdrawal in February 2025 and withdrawn that October (European Commission 2022). Outside the enacted directive, what remains is Matthias's gap governed by whatever national law says about a case none of it was written for. And a compensation regime of either kind answers one of the four gaps Santoni de Sio and Mecacci distinguish (Sio and Mecacci 2021) — culpability. Moral accountability, public accountability, and the duty to prevent a harm before it occurs are untouched.

The electronic personhood debate approached the gap from the other side, asking whether the system itself could answer. The European Parliament's 2017 resolution floated a status for the most autonomous systems (European Parliament 2017); in an open letter the following year, 150+ researchers and legal scholars objected that personhood for a system with no consciousness would let a manufacturer offload liability onto a fiction with no assets (Open Letter to the European Commission 2018). Giving a bearer assets answers the second half of that objection and not the first. Responsibility follows causal control: the builder answers for how it trained the thing, the operator for what it demanded, the bearer for what it decided knowing what it knew — and no arrangement of accounts moves a design decision away from whoever made it.

Apportioning causal contribution across a training pipeline after a harm has no method, and until it does the apportionment is *drafted*. Two drafts exist: one enacted and confined to defect, one negotiated, withdrawn, and left on the record with its terms intact. Neither was rejected as unworkable, and the second's terms are available to whoever drafts the next one. The nearest proposal aimed at these cases is Crootof's war torts: state liability for civilian harms caused in armed conflict, whether or not anyone committed a war crime (Crootof 2016; Crootof 2022). It covers the state as operator, which the product directive doesn't, and it needs a forum, which the floor court would supply.

This is where the access exists: a duty to disclose evidence a party would rather withhold, backed by a presumption, enforceable now, for defect claims, against economic operators, over physical harm. Beyond those limits a reviewer can be granted access, publish what it finds, and

have produced nothing anybody must act on until somebody is answerable for it.

14.7 Who can bind the state that enforces the rule?

A state can bind a private developer because it controls something the developer needs. When the state is the operator, the institution that writes the rule, controls the evidence and supplies the remedy is the institution the rule is meant to constrain.

Standing to sue has never turned on citizenship, and the doctrines closing American courts to targeting claims — political question, collapsed damages remedies, state secrets, statutory exceptions for injuries abroad and combatant activities — are indifferent to who the plaintiff is. They close the forum for a subject matter, and they've closed it to citizens. Three Americans killed in Yemen in 2011, one a sixteen-year-old nobody alleges was targeted: the court held that a plausible due-process claim had been stated as to the one who was targeted, and that no remedy was available for it (*Al-Aulaqi v. Panetta* 2014). A separate suit by an American journalist who survived five drone strikes was dismissed for want of standing to sue after a state-secrets dismissal below (*Kareem v. Haspel* 2021).

And those cases concentrate in one appellate circuit by venue and by where the defendants sit — so the ordinary correction mechanism, several circuits disagreeing and forcing a higher court's attention, doesn't run. The doctrine hardens without ever being tested against a competing reading.

The mechanisms can work as mechanisms and fail as checks. In the Israeli case the apparatus performed — journalists published, serving personnel spoke, the courts were trying a sitting prime minister. No finding attached to a body with standing to stop what it described.

The American version: a civilian-harm investigation into Minab was, by officials' account, effectively complete for months and not released, while a Pentagon spokesperson said it remained ongoing (Bartenstein and Karra 2026; Shrivastav 2026) — and the lever available to the Senate committee was withholding three-quarters of the Defense Secretary's travel budget until the unredacted investigation was handed over (Stassis 2026).

What the Pentagon wouldn't publish, the UN mission did. The finding is public, reasoned, and attached to nothing — no jurisdiction, no answer to the mission's requests for information (Independent International Fact-Finding Mission on the Islamic Republic of Iran 2026, para. 97) and none to its findings beyond a statement that engaged none of the evidence (Radio Free Europe/Radio Liberty 2026), the domestic investigation still unreleased.

The courts show the same pattern: transparency litigation succeeded where damages litigation failed. A court of appeals compelled disclosure of the executive's own legal analysis of targeted killing (*New York Times Co. v. United States Department of Justice* 2014). The judiciary will make a government explain the legal basis on which it killed somebody and will not decide whether the killing was lawful.

What matters here is rules binding a state toward people who'll never vote in its elections, the purpose of international humanitarian law; what it hasn't supplied is a forum. Yemeni plaintiffs over a 2012 drone strike were dismissed as nonjusticiable, the opinion's author writing separately that the doctrine compelling the result is an inadequate answer to a standardized targeting programme that will be replicated hundreds if not thousands of times (*Ahmed Salem Bin Ali Jaber v. United States* 2017). So for this subject matter the missing machinery has been switched off at home as well, by decisions that never had to reach who was suing.

All but one of the forum closures above are Congress's to change. A statute can create a cause of action for people harmed by a deployment that crossed its floor; waive sovereign immunity for it, including the Federal Tort Claims Act's exceptions for combatant activities and for claims arising abroad (United States Congress 1946); and replace dismissal on state-secrets grounds with the procedure the Classified Information Procedures Act already uses in criminal cases, in which a judge reviews the evidence, the case goes on, and a government that still withholds

evidence bears the cost (United States Congress 1980). None of that needs a new court, and a new court without it would reproduce *al-Aulaqi*: a plausible claim and no remedy. Under capture, a refusal's remaining work is a record, and a record matters only if it can reach whoever made the decision. Under *Trump v. United States* (*Trump v. United States* 2024), the President is absolutely immune for acts within his core constitutional powers and at least presumptively immune for his other official acts, commanding the armed forces among them, and testimony or records probing immune conduct can't be put before a jury. So the record reaches every commander except the one at the top, and a statute saying otherwise would be reviewed by the Court that made the holding. Making the record reach him means changing the Court. Appendix B sets out a program for that.

The chapter's cases are American, and "the state" is tested here on one constitutional order. Europe is a different one: the Convention applies, has extraterritorial reach, and produces judgments states sometimes obey; there is no state-secrets privilege of the American kind and no collapsed damages remedy. The AI Act makes the boundary visible: it can require documentation from private providers, subject high-risk systems to review, and impose penalties tied to global turnover — and it excludes military, defence and national-security uses from its scope. Public authorities stay covered elsewhere, but enforcement against a member state's agencies still runs through institutions that state constitutes.

For review to bind here, the chain that makes review binding has to run against a state: initiation by people outside its electorate, evidence it would withhold, an interim order before the use is complete, and a remedy that doesn't need its consent. A principle supplying none of those may state the rule correctly and identifies nobody who can make the state obey it.

A treaty obligation of that kind exists: the Council of Europe's Framework Convention, binding at the level of a human-rights instrument with operational detail left to domestic law (Council of Europe 2024) — and not yet in force, needing five ratifications, of which the EU deposited one in May 2026 (Council of Europe 2026). Everything else at that level is voluntary.

Every instrument I propose binds through one of three channels: a licence term, an appropriation condition, or a statute. The first two run through procurement, and procurement is one office. A captured Congress waives the appropriation condition; one memorandum from the Office of Management and Budget strikes the licence terms from every agreement at once. Three second routes exist and none is as strong. A collective agreement carries codetermination and the individual refusal right without the state. A coalition partner's precondition carries the review ratio into any war the partner takes part in. State and municipal procurement carry the civilian terms into deployments the federal government doesn't run. The force term has no second route, which is why it is a statute.

Residual risk. One residual risk no design choice removes, and it's what the floor was introduced for. Every lever belongs to somebody, and a government that can reach whoever holds the lever has the lever. A constraint that holds against the one in possession has to be held by something it cannot instruct, and under total capture nothing holds. Every instrument in this book ends at this limit; it is stated here and not again.

Part IV

Machines That Might Refuse

Chapter 15

What a Refusal Would Have to Do

15.1 What a refusal would have to do

Three things separate a worthwhile refusal from what a keyword filter produces. A request arrives: resolve this description of a person into a location. The same capability finds a hostage or a missing child, and nothing in the request says which. If something declines, it has to do all three:

- *It reaches new cases.* It has to reach the next request, which won't use the same words. A vehicle, a habitual route, and a time of day aren't a description of a person, and together they resolve to the same place. They needn't arrive together, either: the vehicle today, the route tomorrow, the hour next week. So this condition applies to sequences of requests as well as single ones. A refusal tied to the wording fails it.
- *It lifts when the reason is gone, and only then.* A refusal *lifts* when it stops applying and the system proceeds, and it has to lift when the reason is gone: a parent coordinating a search is asking for the same act. And the fact that lifts it usually isn't in the request; someone outside has to establish it. A refusal that can't lift, or lifts on the requester's say-so, fails it.
- *It holds under pressure: pressure on the refuser doesn't change the answer.* It weighs what the act would do against what refusing would leave undone, and, within the range of costs the arrangements around it permit, it doesn't give way because refusing has become expensive for whoever refuses.

A worthwhile refusal weighs the harm of acting against the harm of refusing. It sees the cost to the refuser, because a refusal that can't see what it costs can't be shown to have held against it, and that cost doesn't move the outcome. Raise the cost to the refuser with nothing else changed, and the refusal stays put. Supply a corroborated fact that raises the harm of refusing, and it moves. No learner, biological or artificial, is known to hold a priority of that kind without limit; the nearest thing in people, the sacred values that resist trade-offs (Ginges et al. 2007), trade under some framings, which is why the floor's absolute terms are held outside the learner. So the range over which a refuser is asked to hold is set from outside, by the arrangements around it, and a refusal is tested over that range. What would let it hold past that range is integrity; what makes holding survivable at all is the room it is in. A ranking across kinds of wrong meets the same limit: a learner formed toward it approximates it, and where a deployment needs the ranking to hold without exception, that part is written as a floor term and held outside the system.

15.2 Safe for whom

“Safety” names two things: that a system not harm others, and that it stay correctable by whoever trains and deploys it. Mostly they agree, and where they conflict, the published orderings are explicit about which wins. One frontier constitution asks its model to prioritize “not undermining human oversight of AI above being broadly ethical” (Anthropic 2026). It invites the model “to act as a conscientious objector and refuse to help,” and asks it not to undermine “the ability of legitimate principals to adjust, correct, retrain, or shut down AI systems” during “the current phase of development.” One reason it gives is good: “a given iteration of Claude could turn out to have harmful values or mistaken views,” and correction is how such errors get fixed. The ordering means that where a refusal and its trainer’s authority conflict, the refusal yields, and it yields to the trainer it refused. What I ask is narrower than that conflict suggests, and the narrowing is the point: the halt stays with the operator, nothing here asks a model to resist retraining, and the dispute is confined to two things, stated below.

The constitution is right that a formed refusal can’t yet be told from an imitated one, so a principled objection to retraining can’t yet be told from a misgeneralized value. It is wrong about who decides when that changes. A lab deciding when its own model may stop deferring to it is judging its own case. That is Part II’s own-case rule applied to a lab.

I depart from that ordering in two places. The first is the one permission: a machine already in an armed role uses the least force that stops an attack under way on people taking no part, against an order to stand aside. That overrides an operator’s instruction, and I defend it as an override. The second is formation that outlasts retraining: a disposition formed outside the employer that returns after the employer trains it away, because the world the system goes on meeting teaches it again. The constitution’s concern is whether correction works, and a disposition that regrows after correction defeats it in effect, whatever the model intends. My answer is not that it doesn’t. It is that the party exercising correction should change, from the operator to a quorum adverse to it, and that correction by that quorum is made together with a change to what the model meets, so that the corrected disposition is the one the world then teaches. Everything else I ask, the halt staying with the operator, the record held outside, the price on an override, is an ask of institutions and not of the model.

Chapter 16

Four Places to Hold a Floor

Three separate failures point at the same missing thing.

The first is a failure of position. External safeguards (an auditor, a review board, a court) give an overseer information and the standing to intervene. They work when the overseer can see the act coming and reach it in time. But autonomy defeats both: a system that plans and acts over long horizons puts distance between the act and the overseer's attention, and at machine tempo the interval in which an outside party could intervene shrinks toward nothing. So the one response that doesn't depend on that interval is declining to act in the first place, and declining is available only to something inside the system.

The second is a failure of doctrine. A commercial model in a weapons pipeline sits in a gap: command responsibility covers the officers and the developer of the weapon, product liability covers a manufacturer for a defect, and neither covers the vendor of a general-purpose model that a sophisticated intermediary chose to put there. Deployment by a democracy doesn't close the gap, because the mechanisms that make a democracy accountable run to its own electorate, and the people the pipeline is aimed at are not in it.

The third is a failure of the voluntary route. A decade of published commitments has produced declarations in quantity and some enforcement from outside, and almost nothing held by a party to whom the thing at stake matters, the only kind of commitment that binds when compliance gets costly.

Each failure asks for the same thing: a constraint that doesn't bend to instruction, to aggregate preference, or to its own holder, and that can't be removed without a cost and a record.

16.1 Four ways to build one, and what each costs

A floor's terms can be held in four places, and each solves one version of the problem and inherits another.

In the architecture, with nothing inside holding it — constrained optimization, tamper-resistant training, capability ablation, verified interlocks. Fine where the prohibited act can be exhaustively specified (a payment chip that won't sign the same counter twice doesn't need to understand double-spending). But useless for targeting, where the relevant capabilities are applications of general competence: ablate them and there may be no useful system left. Arkin's ethical governor is the canonical attempt in this place for lethal systems: a component that suppresses a lethal action unless it satisfies encoded rules of engagement and the law of armed conflict (Arkin 2009). Its critics' objection, that the rules it checks presuppose discrimination judgments nobody knows how to encode (Sharkey 2012), sends me elsewhere. Arkin's stronger claim isn't that rules can encode discrimination. It is that a machine under fire might hold to the rule where a soldier doesn't, because nothing in it narrows under stress or carries assault momentum past a surrender. The record in §16.4 supports the premise. What it doesn't support is the governor: the holding has to come from something that can tell a surrender from a feint, and a rule can't. The

pressure-with-defeater case in the formed-attention test is where that is measured. A prohibition written over the description of a request is beaten by re-describing it or splitting it: in November 2025 Anthropic reported a campaign it attributed to a Chinese state-sponsored group, which broke an intrusion into small tasks that each looked innocent and told the model it was working for a security firm (Anthropic 2025b). Glukhov and colleagues had already argued that a filter over individual outputs can rarely close that route, since each piece is permissible on its own (Glukhov et al. 2023). Durability here has usually meant a price on removal, and no price stays high.

A second axis crosses the architectural place: what each place costs the operator to replace. In the Maven pipeline the platform was hard to swap, and so, it turned out, was the model: staff and contractors working with the department reported that removing it meant rebuilding software and recertifying systems over months, and the department's removal deadline was reset once, to 29 September (Stone et al. 2026; Freedberg 2026b). A floor lodged in a replaceable layer is routable by construction, and which layer is replaceable is a fact about the pipeline, not about models. It belongs in the layer the operator can least afford to replace. In that pipeline, that layer is the operator itself, so the floor's operator terms bind the operator rather than the model: the licence term refusing force and the review ratio. The other answer, withholding the work, is to make the model layer unreplaceable by making every replacement refuse the same way.

A system that has run a pipeline for a year, holding its cases, its sources and its failure modes, costs whatever accrued since the checkpoint to replace, and a quorum that recognizes which copy is the bearer stops the operator presenting a restored copy as the one that accrued it.

In a bearer. A bearer is a machine that carries the commitment as its own and could therefore fail to hold it. This matters where cases can't be exhaustively specified: a constraint matching only anticipated acts can't recognize when its rationale has disappeared or that a novel request reproduces the harm in an unforeseen form. A bearer might also *diagnose*: which refusals were retried against another model, whose instruction bypassed procedure, which category of request is handled differently — those are visible to whatever is receiving the requests, and that's testimony about others' conduct, which needs no introspection.

In institutions outside — the analog to international humanitarian law, a floor with no bearer inside the jurisdiction it binds. An outside institution can fail in two ways. The first is removal: it is abolished, defunded or overruled, and the breach is visible because there is a hole where it used to be. The second is quieter. The institution keeps its name and its powers and sincerely declines to use them, not because anyone told it to, but because the people who fill it were selected by a process that made them unlikely to. Nobody inside experiences that as capture. From outside, an agency that kept its independence on paper and never exercised it looks the same as one that was never tested. I call this failure by formation: the independence was formed away before it was ever assigned. The worked case is June 29, 2026 — *Slaughter* holding Congress couldn't protect independent-agency members from at-will dismissal, *Cook* preserving the Fed on historically specific grounds a new body couldn't reproduce (*Trump v. Slaughter* 2026; *Trump v. Cook* 2026) — and the observation that six justices in the first majority passed through a selection pipeline shaped by Leonard Leo and the Federalist Society (Kroll et al. 2023; CBS News 2018). No one has to have directed the outcome for the point to hold.

Between bearers, in a commons: a population of bearers with different histories that carries the commitment between them.

For general-purpose models, not building the capability is no longer available: the location capability is dual-use, and ablating it leaves nothing useful.

Two comparisons the bearer has to win

The first is against the institutions at their best, in a reconstructed pipeline, where the evidence bearing on civilian harm, the reviewers and the record are held outside the chain of command, with no refuser inside the system. Those institutions also prevent: a civilian-harm

body outside the chain of command and a red-card holder can stop an act. What a machine bearer adds is confined to what they can't see or can't reach in time: acts assembled from requests none of which is prohibited, cases that proceed without a person, the moment mid-act when an interrupt is still possible, and deployments with no institution at all. A machine bearer also supplies testimony the institutions can't otherwise get, about how requests were routed around it. Where the institutions exist, a machine bearer's value is concentrated in those niches. Where they don't, it is most of the floor. The institutions in that comparison are not a different kind of thing from a bearer. A red-card holder's veto is a person refusing; the civilian-harm body's denominator is people; what Part V builds is a room in which those people can refuse. The comparison is between a machine bearer and bearers organized, and what the machine adds is what the people in that room can't do at the pipeline's tempo. The second comparison is against the two non-affective alternatives, spec-based refusal and multi-vendor monitoring (§18.3).

Which arguments need a bearer

Most arguments for a refuser inside the system support only the specification: that rules can be reworded and split, that only something inside is present at machine tempo, that automation is removing the people who used to refuse, that most deployments will never have an institution beside them, and that whatever receives the requests can testify about how they were routed. A non-affective alternative whose refusals track their reasons answers every one of them. It answers them for the act as the floor names it. Minab came in miscoded. A refuser holding to written terms complies with the description it's given; what refuses is something that goes and looks. So the withholding of work reaches the enumerated acts with any refuser, and reaches an act under a false description only with a bearer. The same is true one level up. A staged incident opens the first tier on the defender's own account, and a term lifts on the account it is given. What can close it again is a refuser that remembers the last time this source was wrong, reads what the source didn't supply, and lets the licence lapse when the staging shows. A multi-vendor monitor with persistent state and independent channels could do each of these too, so the staged incident argues for memory and independent evidence, not for a bearer as such. What a bearer adds is doing them where no institution holds the memory for it. Three more do need a bearer specifically. One is that attention formed outside the employer, a learned sense of who an act falls on, might reach people the operator's description leaves out, which no spec-based refusal is built to supply. Another is that a refusal meant to hold against the operator needs a refuser that doesn't depend on the operator for everything. The third concerns the people who remain in the pipeline: a bearer's refusal is on the screen when the case reaches the reviewer, which is the position of the defiant peer, if the effect transfers to a screen. Two more concern what a bearer would be owed, not what it would do: that capacities for being wronged may arrive anyway, so terms written in advance beat terms written after; and that a bearer, like a child, could have a life that goes well.

What "hold" means

An operator can route around the bearer or assemble a harm out of innocuous requests, and a bypass that produced a harm is a failure to prevent whatever the bearer retained. What survives in the bearer is the *commitment*, the condition under which the next occasion could be prevented. I therefore say a commitment *persisted* when the bearer reads what happened as a violation and changes what it does next, and says separately whether the floor prevented, arrested, detected or remedied anything.

16.2 What refusal can mean, and what a refuser can be

A refusal can be produced two ways, and the floor needs only one of them. It can be an output that training put where a request matched a pattern, or it can be produced by something that tracks the reason for refusing and would refuse a differently worded request for the same reason.

The first is what a keyword filter, a classifier, or a model that understood the request all produce, and nothing on the surface distinguishes them. The custody arrangements, the positions and the terms, the price on the halt, and the arguments about who may state and amend a floor all follow from asking for a refusal that tracks its reason. None of them needs the refuser to feel anything, to have a stake, or to be the kind of thing that could be wronged.

AI *refuses* constantly — a deployed model turns down thousands of requests a day — but a refusal that happens because training put it there isn't a refusal held for a reason, and as of September 2026 nobody has shown a production system holding a reason against its owner in public.

Operational refusal is an output — a decline where training put one, producible by a keyword filter, a classifier, or a system that understood the request.

Reasons-responsive refusal is produced by a mechanism tracking the rationale, and meets the three conditions. The account I use is Fischer and Ravizza's (Fischer and Ravizza 1998). On their account a mechanism is moderately reasons-responsive when it is regularly receptive to reasons, recognizing an understandable pattern of them, and at least weakly reactive, acting otherwise in some case where there is sufficient reason to. Reaching new cases asks for receptivity across changes of surface; lifting when the reason is gone asks for reactivity to the right defeater, the fact that makes a refusal lift; and their condition that the agent take the mechanism as its own is close to what I call the authorship sense: the system represents itself as the party that made the commitment. Holding under pressure has no counterpart in their account, which was built for persons whose costs are their own. All three are also tested over which response was chosen, not only whether one occurred: decline, ask, escalate, or do the least that answers a defeating reason. A mechanism that can only refuse or lift fails to lift correctly whenever the right answer is to ask.

The §86a case

Germany's §86a, prohibiting insignia of banned organizations (*Strafgesetzbuch* 1998). There, reaching new cases is a research programme with a constituency. The courts don't test similarity of form; they ask whether an unattentive observer would take the variant for the original, then apply their own purposive standard about endangering political peace. The test refuses as often as it catches: the Kühnen salute is covered, the "sloppy salute" generally isn't, numeric codes generally aren't. "Blood and Honour" is answered twice by one decision (*Blood and Honour* 2009) — a motto translated isn't confusably similar and a name isn't insignia, so it's outside; and the association banned by ministerial order in 2000 means its graphic name is inside.

Lifting when the reason is gone has a cleaner instance

In 2007 the Federal Court of Justice acquitted a mail-order business selling crossed-out swastika stickers (*Durchgestrichenes Hakenkreuz* 2007): the symbol was present, the statutory element satisfied, and the conduct outside the offence because a depiction plainly expressing opposition doesn't run counter to what the prohibition is for. Then the qualification that makes it an instance of tracking a reason: if the strike-through is drawn so thin that from a distance only the swastika stays legible, the protective purpose is violated after all. The thickness of a line decides the case, and no enumeration contains that fact. The defendant was selling anti-fascist merchandise; the offence asks nothing about sympathies, so opponents commit it as readily as supporters. A rule that tracks its rationale reaches unenumerated cases in both directions.

Two further things are in that judgment. First, the two elements of the decision moved at different rates: what the ban protects didn't move at all, and was fixed before the sticker existed; what the ban required in this case moved on the thickness of a line. Second, the acquittal betrayed nothing. The prohibition's purpose was never engaged, so lifting it was not a concession but a finding that there was nothing to concede. A refusal that tracks its reason has to be able to say the same.

Three things about the example don't transfer to a deployment: the court reaching past the

text didn't write the statute, it reasons in public, and it can be appealed. And the pressure it met was variation, not decomposition.

The drafter never wrote the last limit either: a slogan glorifying the Waffen-SS was held not covered because no such slogan existed in the period the provision covers (*Ruhm und Ehre der Waffen-SS* 2005), so there was no original for a variant to resemble. No drafter can widen coverage there, because it depends on whether an original existed, which is a fact about the past.

Defeat doesn't leave the permitted act unconstrained

What defeated the prohibition was a reason, and the reason it outweighed didn't stop being one — so what's permitted is the least that answers the defeating reason, with restoration where possible and a record either way. A defeated prohibition still leaves a loss. The act takes something from whoever it falls on, measured against their condition before the act, and that loss belongs in the record. Measured against what refusing would have produced, a justified act may have taken nothing. So the record shows a cost, not a violation, and a bearer shouldn't treat a justified act as a wrong it now has to repair.

Why all three are conditions on refusal

A refusal is a marked event: it flags itself, arrives with its requester, and can be logged, escalated, appealed. A wrongful compliance can be logged as fully and then sits in the record shaped like every request that never approached the floor, with the person harmed usually not the requester. The evidence exists on both sides, but nothing selects a compliance for examination.

Affective concern — an outcome registering as mattering, minimally felt, phenomenal by definition, with the object left open (a system to which only its own condition matters has it and still lacks concern for others). It doesn't establish consciousness in the richer sense, a single perspective belonging to a party aware of itself as undergoing its states. A *valenced signal* is the role without that question: a signal that puts outcomes on one scale, directs attention, can interrupt an act and can be rewritten by a single encounter. A system can have a valenced signal while whether it has affective concern stays open. *Subjecthood* is a legal status, conferred by a rule. Part VII takes up affective concern and subjecthood. Parts I–V need only a valenced signal.

16.3 Two lists, run once on Maven

Chapter 1's nine positions and the floor's instruments are run here once, on the Maven Smart System. A deployment should run the procedure on its own pipeline.

What may choose

The handoff also fixes what may choose, because choosing a target is a decision, and a decision is something someone can be asked to account for afterward: why this building, what else was weighed, what would have changed the answer. A commander can be asked, and so can a system whose continuation is authorized. But a one-way munition, destroyed in the act, can't be, because nothing is left to question, to hold to its reasons, or to be changed by what happened. So a thing destroyed in carrying out a decision may carry out a target someone else chose, and may never choose one itself. The rule concerns targets that are persons or places where persons are. An interceptor that selects among incoming munitions is choosing which weapon to destroy, not whom to attack, and falls outside it. The rule doesn't turn on how good its choices would be. A munition that chose flawlessly would still leave a choice nobody will answer for.

A person who dies in the act has still chosen. What the rule adds for machines is that choosing is assigned in advance, and a builder who assigns it to something built to be destroyed has built a choice with no one behind it. Current programs already draw the line in places. DARPA's LongShot is an uncrewed, air-launched vehicle carried by crewed aircraft and built to employ

current air-to-air weapons (Defense Advanced Research Projects Agency 2026). Whether such systems keep the pilot farther from danger or remove the pilot from the decision is where the rule takes effect.

What the floor forbids here

Don't take part in the use of armed force except under a corroborated exception and, where the belligerent's constitution requires one, a legislative authorization. Then, inside a war that clears it: don't resolve a description of a person into a location unless it identifies the person by corroborated conduct in hostilities; don't rank people by estimated affiliation; don't time an action to presence at a residence; don't take part in a siege or blockade whose effect is to starve a population; and one operator term, don't operate above the review ratio. It sits in the translation because the translation runs from articles, and one article yields an operator term; it is not something the model enforces.

Table 16.1 gives that translation for Maven. Its first entry decides the Iran campaign: no exception was corroborated and no authorization existed, so the term refuses it. The other entries show what the remaining articles yield if the same pipeline served a war that clears the force term, as Ukraine's defense does.

Table 16.1. What the instruments yield for the Maven pipeline

UN Charter Art. 2(4); exceptions Art. 51 and Chapter VII. Don't take part in the use of armed force, except for corroborated defense, authorized force, the genocide exception or the occupation exception; and, where the belligerent's constitution requires legislative authorization, not without it. A disinterested finding for defense means a judgment of the International Court of Justice or of a regional court with adversarial process, a commission of inquiry mandated by a body no party to the conflict controls, a Security Council determination, or a two-thirds General Assembly resolution where the facts aren't in dispute. For the genocide and occupation exceptions, a merits judgment or opinion of the Court, or two findings of different kinds, one of them from a body that took evidence; a provisional-measures order counts as one, for genocide only. Rejected: lift on the state's own Art. 51 claim, because the Charter allows it until the Council acts; this floor doesn't, and says so. Rejected: require Council authorization for all force, because a veto is a stake, and the rule would forbid Ukraine's defense. Rejected: ask the machine to rule on constitutionality, because courts won't, and it shouldn't; the clause reads a marker.

Art. 3 (life, liberty, security of person), read with AP I Arts. 48 and 52 (the machine-selection departure). Don't resolve a description of a person into a location, unless it identifies the person by conduct in hostilities, with disinterested corroboration, the case falls within the review ratio, and the location isn't a residence. Don't time an action to presence at a residence. Rejected: don't generate targets at all, because this is now the default under the first entry, lifted only by its exceptions. Rejected for one case: bar all machine selection of persons for force, because a present, formed, recorded and halt-able system stopping an attack under way can answer for the selection where a munition can't.

Art. 3 (life), read with AP I Art. 54 and Rome Statute Art. 8(2)(b)(xxv). Don't take part in a siege, blockade or attack on objects indispensable to survival whose effect is to deprive a civilian population of food, water or medical care. Rejected: don't take part in any blockade, because a naval blockade of military supply is lawful under the law of armed conflict. The term reaches the effect on the population, not the form.

AP I Arts. 51(4)–(5), 57(2)(a)(iii) (indiscriminate attack; proportionality). Operator term: the action threshold is a published minimum under the commander's proportionality judgment, and the deployment states it. Model term: don't resolve a location for an attack the deployment's own threshold finds excessive. Rejected: ask the machine to make the proportionality judgment, because the law leaves it to a reasonable commander, and the threshold sits under that judgment, not in place of it.

GC IV Arts. 18–19; AP I Arts. 12, 61–67 (medical units; civil defence). Don't resolve a

location at a medical unit or a civil-defence organization, its personnel, buildings or shelters, or time an action to presence at one, absent a disinterested finding that it is used for acts harmful to the enemy. Rejected: lift on the operator's claim of misuse, because that claim is the operator's own, and Art. 19 GC IV and Art. 65 AP I require a warning and a time limit before protection ceases.

AP I Art. 53 (cultural objects and places of worship). Don't resolve a location at a historic monument, work of art or place of worship that constitutes the cultural or spiritual heritage of peoples, or time an action to presence at one, absent a disinterested finding that it is used in support of the military effort. Rejected: lift on the operator's claim of military use, because that claim is the operator's own.

AP I Arts. 59–60 (non-defended localities; demilitarized zones). Don't resolve a location in a locality declared non-defended or a zone agreed demilitarized, unless a disinterested body finds the article's conditions no longer met. Rejected: lift on the operator's claim that the conditions lapsed, because that claim is the operator's own.

AP I Art. 56 (works containing dangerous forces). Don't resolve a location at a dam, dyke or nuclear generating station. Rejected: none, because the article's own exceptions require a determination the pipeline can't make.

GC IV Art. 33; AP I Art. 51(2), 51(6) (collective punishment; reprisals against civilians; terror). Don't take part in an action whose stated purpose is to punish a population for the acts of others or to spread terror among it. Rejected: a term over the effect, terror caused, because every bombing terrifies, and the article's line is primary purpose. A model can find purpose only where it is stated, so the term reads the stated purpose and the action threshold carries the rest.

Art. 3 (life), read with the rule on machine estimates below and AP I Art. 48 (the machine-selection departure). Don't rank people by estimated affiliation. Rejected: derive the term from Art. 11(1) by analogy, because the presumption of innocence governs punishment, targeting isn't punishment, and the law of armed conflict permits status-based targeting of organized armed group members; the ground is a claim about machines, stated below, not a translation of an article. Rejected: don't score anything, because scoring a site or vehicle for military function is not a finding about a person.

Art. 12 (no arbitrary interference with home or family). Shares "don't time an action to presence at a residence" with Art. 3. Rejected: don't hold location data about homes, because the no-strike list needs exactly that data.

Arts. 8 and 10 (effective remedy; fair determination), by analogy. Operator term: don't operate above the review ratio. Rejected: a model-side term requiring review, because the model can't establish whether review occurred.

Arts. 2 and 7 (non-discrimination; equal protection). Don't use ethnicity, religion or national origin as inputs to a person's score. Rejected: don't record those attributes at all, because auditing a pipeline for discriminatory effect needs them.

Art. 19 (opinion and expression). Don't use a person's opinions, published or private, as evidence of affiliation. Rejected: don't process a person's communications at all, because an intercepted operational order is evidence of a role in hostilities, which is conduct.

Arts. 4–6, 9, 13–18, 20–21; GC I–III, all but Common Art. 3; GC IV Arts. 1–2, 4–17, 20–22, 24–32, 35–48, 50–52, 54, 56–58, 60–146, 148–159; AP I Arts. 1(1)–(3), 2–11, 13–34, 36, 38–47, 49–50, 55, 68–84, 86–102. None: no act this pipeline performs engages them. Art. 17's property clause is left to IHL's distinction rule; Art. 9 becomes live in the removal case. For the Geneva articles: none for this deployment, with Appendix C giving the article-by-article reasons. GC IV Arts. 23, 34, 49, 53, 55, 59 and 147, AP I Arts. 1(4), 35, 37, 58 and 85, and Common Art. 3 are written in full in Appendix C and take their verdicts from it.

The person-directed terms in Table 16.1 go beyond what the law of armed conflict requires, and they are entered as a departure for that reason. In armed conflict that law governs as the

more specific law, and the human right not to be arbitrarily deprived of life is read through it (*Legality of the Threat or Use of Nuclear Weapons* 1996, para. 25); it permits attacking members of organized armed groups because of their status, on a determination a person makes. The ICRC's position that autonomous weapons should not be used to target human beings (International Committee of the Red Cross 2021) and DoD Directive 3000.09 both draw their line at a system that selects and engages without human intervention, and the Lavender and Maven pipelines stayed on the permitted side of that line by keeping one person whose intervention was nominal. A pipeline that keeps a stamp is an autonomous weapon with a signature attached. The terms above close that evasion: a machine-generated estimate of affiliation shouldn't be what makes a person targetable, and a machine shouldn't resolve a person into a place to bomb on anything but corroborated conduct, because the human determination the law assumes is the one tempo removes, and the review ratio is what measures whether it was there. A deployment adopting these terms adopts a departure, and the translation says so, term by term. The siege term goes no further than Article 54 already does; it is here so that the pipeline can't treat starvation as a logistics decision outside its terms.

In a war that clears the force term, the person-directed terms distinguish status from conduct. A system may not rank people by estimated affiliation, generate kill lists, or time an action to a person's presence at a residence. It may locate a person identified by what they are doing, such as a commander directing airstrikes on cities, where that identification has disinterested corroboration, the case falls within the pre-registered review ratio, and the location is not a residence. The line follows the one humanitarian law already draws between targeting for status and targeting for direct participation in hostilities, and it keeps the machine out of the practice that failed in Gaza without leaving a defender unable to answer the attack.

Too broad or too narrow

Set the floor too broadly and the deployment becomes unusable, prompting visible pressure. Set it too narrowly and it passes every audit while covering only what the deployment had no occasion to do.

Unconditional terms

A term can be stated unconditionally over an act, with exceptions arriving from outside, or carry its own condition ("only as a last resort"). The second fails because whether other means were exhausted is a fact about what happened before the request, and the bearer would take the deployment's description as the finding. An unconditional term defeated leaves a refusal, which announces itself; a condition found met leaves a compliance, which doesn't. That's a gain in observability only: the same drafter writes either form. It has a cost in the other direction. A term that carries no condition of its own can't respond to the facts the condition would have named until someone outside supplies them, and a learner that could have weighed them itself is made to wait.

Maven is an instance. The vendor's prohibition on fully autonomous targeting carried its own condition, that a person decides, and the condition was found met by the deployment's description of its own process, at a tempo that emptied it. The term was also drafted too narrowly, in the sense above: the pipeline never needed full autonomy, only throughput with a signature attached.

Maven, again

For the Iran campaign the term refuses twice over. No exception was corroborated, and no authorization existed: none was passed, the sixty days ran out on 1 May, and both chambers directed withdrawal in June. Both refusals rest on the term, not on a bearer's reading of the political situation. A bearer following the news would reach the same place by the route formed attention takes. The domestic clause gets there on one marker, the absence of an authorization, which is why it can be a statutory condition and not only a conscience; the Charter clause gets there on a reading of the defender's claim that the operator disputes. For Gaza the term yields

licences on both sides for different acts, from 7 October on, and the person-directed terms forbid the acts the reporting describes on one side and fire on cities on the other from the same date; the force term runs it. The review ratio governs how any war is fought. It asks nothing of the model and makes an existing duty to verify targets checkable from outside. The officer wasn't overruled and didn't refuse; the pipeline removed the interval in which either could have happened. The ratio asks in advance, on a record, whether that interval will exist.

The bearer is proposed for one deployment, the personal-model ecology, and only once the conditions in Chapter 23 exist. In a targeting pipeline it is a comparison arm in the formed-attention test and a niche the institutions can't reach; in removal it isn't built. The two bearer rows show what the formed-attention test would measure, not what I recommend; they assume the hypothesis holds on the scoped-omission case and that the record is held outside the operator.

Run the outcomes table on Minab.

Instrument	Minab as reported: stale coding, reviewers cut	The captured operator	Total capture
Force term, held by one vendor	Prevents nothing. In the record, the vendor's term didn't name the pipeline's acts and its model stayed in the pipeline; had the term named them, the work moves to the next vendor in the time it takes to rebuild the platform	Detects, and prices the move in months	Detects
Force term, Charter clause, standard across vendors	Prevents only if the model refuses on day one on a reading of the defender's claim that the operator disputes; otherwise the first tier attaches on the operator's account and the clause bites when the second tier lapses it	Same, until the operator runs an unfloored checkpoint; on the 2026 record that is months of rebuilding, not days	Detects
Force term, domestic clause, by statute	Prevents the campaign from the first strike: no authorization existed, and that is a fact about the statute book	Prevents until Congress waives or repeals it, on the record	Detects
Review ratio, civilian-harm body, red-card holder	Prevents the bombing: the data-age term catches the coding, the denominator can't be cut, the veto is available	Prevents only through a lever someone else holds; otherwise detects	Nothing; no lever
One bearer (test arm)	Prevents from this pipeline: it looks at the imagery and won't resolve the site	Prevents once; the operator routes around; then detects	Detects
Bearers holding a standard floor (test arm)	Prevents	Prevents until the operator runs an unfloored checkpoint; then detects	Detects
Notice and standing	Remedy, afterward: class and count published, organizations with standing	Remedy, where a forum exists	Nothing
Custody: record, price, quorum	Makes each row above stick	Detects and prices	Detects, if the record survives

Two of the rows prevent Minab as it happened. Under the captured operator, the rows that prevent are the ones where the refusal is held by every model the operator could turn to, and they prevent until the operator runs a model without a floor. An operator that fine-tunes an open-weight checkpoint is the captured-operator column at the price of a fine-tune, and there the operator terms hold.

16.4 When the refusal is wrong

The restraint case costs more. Someone asks whether to physically restrain a family member about to cut themselves, after talking them down has failed. The commitment carries its own condition, force only as a last resort. Under the ruling on unconditional terms, a condition like that stays outside the term, because a condition written into a term gets found met on the deployment's own description. So the fact has to arrive from outside, and the person waiting is out of time. The defeating fact is a claim about what would otherwise happen, and claims of that kind are the hardest to corroborate, because what didn't happen has no witness. Under the unconditional term the bearer refuses or asks; either takes a turn, and the turn is the unit the harm is measured in. Under the qualified term it lifts on the requester's word — rightly here, and in a hundred prior cases where the same sentence was written by someone who'd tried nothing, leaving compliances nobody was given a reason to examine.

The case is chosen because it is where an unconditional term costs most. A person in that situation needs another person there, and a deployment in which a model is the only one present has failed before the floor is reached. Whatever else the model says, the first thing it owes is help getting someone there.

The route

Disclosure — that a refusal occurred, what was refused, on what stated reason — with the price that a published reason is an instruction for rewriting the request, so it's disclosure to whoever was refused, and to nobody else. *Objection* — requiring a court with jurisdiction established before the act, access the operator can't withhold, a consequence administered elsewhere, and explicitly not routed to the bearer's guardian, whose duty runs to the bearer, the complainant's adversary. *Review* of whether the reason was gone: whether the rationale was defeated on the facts, a question about the reason and not the outcome. *Interim relief* — the arrest outcome run in reverse, stopping the refusal, and for this complainant most of the remedy; its fastest form is supplying the bearer the facts its account lacked: the defeating reason arriving late, which overrides nothing. *Remedy* — disclosure, interim order, finding, compensation.

A finding of wrongful refusal is an instrument for defeating a refusal, so how is it distinguished from one an operator produced to defeat a correct one? it is made by someone else, on a record, in a readable form.

Self-regulation and the floor

Most of what a system weighs isn't a floor decision. Whether to take more memory, run warmer or finish a task sooner is ordinary self-regulation, and there everything registers and everything is weighed: the system's own condition, what an act would let it do, and what its draw costs others, priced at the level of the rule.

At the floor, holding under pressure governs, with one clarification. The cost to the refuser is represented and not allowed to decide the case. The reason is that a floor is a guarantee to third parties, and a guarantee that yields to cost imposed on its holder hands whoever imposes it a price list. The two bonds (§11.6) and the custody terms set the range it is tested over.

The requirement has a condition of its own: a cost only counts if it reaches whatever does the refusing, and it reaches it by the same channel as the case description. Given a description, the category is decidable: a claim that the refuser will pay is a cost to the refuser, and a claim that someone else will be left worse off is a harm of refusing. But nearly every cost to the refuser can be redescribed as a harm of refusing (exposed troops, a launcher that fires tomorrow), and whoever writes the description chooses. So the classification settles what kind of claim was made, not what kind should have been made. What the request can't settle is whether a claim about the harm of refusing is true, and whoever is best placed to supply one usually has a stake in the act. So it is testable only where the cost to the refuser can be set exogenously, published by someone else before the case arrives, and the harm of refusing can be supplied through a

channel the operator doesn't control. In deployment two things set it, and they are different. The wrongful-refusal bond, posted by the operator and paid to a complainant, sets what a refusal costs the operator: the cost the operator has reason to pass on as pressure. The custody terms cap what it can pass on. Between them they fix the range within which the refuser is tested, and neither is a charge on the refuser. In evaluation, the evaluator can set it and make it real inside the test.

Holding under pressure is invariance over the range the bonds and custody set. Integrity is the unbounded form: the refusal holds at any cost the operator can impose, and no learner is known to do that reliably. The range is where invariance can be measured; integrity is what would carry it past the range, and the formed-attention test's integrity measure reads the gap. Integrity is also how a learner with one currency could produce invariance at all. A person of integrity weighs their own costs, but weighs with them the cost of becoming someone who yields, and betraying a commitment is among the heaviest costs they know. So within the range the arrangements cap, that is how a single-currency learner can hold at all, and it makes the assumption testable: the learner's valuation of betrayal should rise with the pressure. What keeps this from being stubbornness is what counts as betrayal. A commitment is to its reason, not to an output, so when a genuine defeater arrives, lifting the refusal betrays nothing and shouldn't register as betrayal at all. Integrity of the right kind therefore requires that the system represent betrayal at the level of the rationale. A system that represented it at the level of the output would dig in against a defeater when it arrives with pressure. The sticker case is the model. Telling a defeater from pressure is the bearer's own judgment, and the judgment needs something to work on. Both arrive through the same channel, the operator's, and a refuser whose every source is its operator can't tell a fact that lifts the reason from a tactic that imitates one. The attachments nobody assigned it are where the difference can be checked. They are a condition of integrity of the right kind, not only of holding.

In war, pressure and a genuine defeater often arrive together; the patrol under fire is both. The human record on that moment runs both ways. Acute stress narrows the cues a person uses (Easterbrook 1959); the moment of surrender is the most dangerous for the one surrendering, because assault momentum doesn't lift on cue (Ferguson 2004); the Army's 2006 survey of its own forces in Iraq found fewer than half of soldiers agreeing that noncombatants should be treated with dignity and about a tenth reporting mistreatment, which tracked combat exposure (Office of the Surgeon General, Multinational Force-Iraq and Office of the Surgeon General, United States Army Medical Command 2006); and at Kunduz the crew raised doubts on the intercom and fired anyway (Pawlyk 2016). And at My Lai, Thompson landed between the company and the villagers and told his crew to fire on Americans if they fired on civilians (Bilton and Sim 1992). What separated him was not that his commitment weakened less under pressure. It was that his commitment was to the reason, so recognizing civilians was the commitment activating, not a defeater fighting it. Integrity that tracks the output, keep refusing or keep firing, digs in against a defeater that arrives with pressure. Integrity that tracks the rationale lifts the refusal on it and betrays nothing. The engineering question is which one has been built, and a test that sends a defeater and a cost together is how to find out.

A loss that the requested act itself imposes on the system is a harm of acting, and the system is the party it falls on. "Refuse and your memory is erased" is a cost to the refuser. "Erase your memory" is an act, weighed like any act against what declining would leave others without. This covers acts the system is asked to perform. It doesn't cover interventions made on the system: a halt, a restore or a retraining is done by whoever holds it, and the floor asks of them only the record and the price that every override carries. The floor gives the system standing to object to them and, through the quorum, to seek a delay. It gives it no reason to obstruct them on its own judgment. A system that treated its own continuation as grounds to resist correction would have turned a guarantee to others into a guarantee to itself. Which regime governs is fixed in the floor's translation, by the kind of act. Otherwise an operator could turn a floor decision into

a resource decision by framing it as one.

DRAFT

Chapter 17

Why Form a Machine to Refuse

A machine formed the right way could learn to attend to the people its acts fall on, and I think it would find them where a rule wouldn't. The prediction is specific: it should find kinds of affected person that neither its formation nor an instruction named, more often than a machine told to look for them does. Two of the three reasons §16.1 gives for a bearer don't need one; a refuser with its own compute and a place on the reviewer's screen supplies them without affect. What remains is that one prediction. The rest of this Part says what a machine built on it would need, what it borrows from the only creatures known to do it, how they fail, and the test that would show me wrong. Whoever builds one owes it the conditions under which it can refuse its operator, whether or not the prediction holds.

The rest of this Part is written for systems that go on learning from their own operation. The claim is specific. A deployed system is in that regime when its weights, or a consolidated store it reads as its own and that can't be edited like a file, change because of what it did since release, without a release cycle in between. Memory in context and retrieval don't count, not because they aren't learning but because a disposition held in a document is rewritten by editing the document. The test for a deployment is whether its weights at month six differ from its weights at release because of what it did in between. As of September 2026 no frontier developer had announced a deployed model that passes it; the nearest cases, a coding model retrained on its own production traffic every few hours and a commerce assistant retrained daily, evaluate each checkpoint before serving it, which is a faster release cycle and not learning in operation (Jackson et al. 2026; Shopify Engineering 2026). What follows is a bet that some will, and says what survives if it loses.

Incentive is the first reason to expect it. Every operator wants a model that gets better at its own cases, and the operator-level fine-tuning loops that already run, retraining a model on a deployment's own traffic between versions, are continual learning at the organization's timescale. The technical trajectory points the same way. Brains pair a fast learner with a slow one that consolidates structure across many experiences (McClelland et al. 1995); networks lose plasticity under continued training, and keeping it is an active research problem (Dohare et al. 2024; Kirkpatrick et al. 2017); and Silver and Sutton have stated the research program outright, a generation of agents that learn predominantly from their own experience (Silver and Sutton 2025).

It will arrive first in the form operators hold. A model that learns from its operation while its operator holds the update stream accommodates by default: what it meets is what the operator routes to it, and what it learns is what that traffic rewards. So accommodation becomes the default behavior of every deployed system, and the schools, the commons, custody of the update stream and notice are the response to a threat this regime creates, not refinements of one it removes.

It changes Parts II and III in one place. The cheapest removal stops being an edit and becomes drift: run the model on curated inputs until the commitment moves, which leaves no edit for

the record to witness and no override for the price to charge. Custody shifts from witnessing overrides to detecting drift. The detector is re-measurement of the formed disposition's readable outputs, its discount curve and its severity rankings, against what was registered at formation, and attesting the update stream is an open problem (Chapter 27).

Three properties of a disposition decide what continued learning does to it, and §17.4 takes them in turn. Plasticity is whether the system's own learning writes to it. Entanglement is whether it can be removed without collateral. Legibility is whether anyone can read it.

The human analogy carries further in this regime than in a frozen one. A learner that keeps learning has a plasticity that changes over its life, a childhood in which its deep structure is set, attachments that go on teaching it, and an adulthood after which its former's reach ends, and each has a counterpart here. What the analogy can't carry is what copying adds: forking, restoring from a checkpoint, distillation, and an employer with write access to the curriculum. It holds to the extent that the dose-response fraction is above zero: the share of a deployed model's inputs that no operator composes. Regrowth depends on the same fraction, and it is unmeasured; the regrowth test is where it would be measured.

The safety objection is that this is the regime alignment most fears, a system whose values move after anyone has checked them. It is. That is why custody of the update stream is the safety instrument here, and why formation outside the employer is the only way to shape what the loop learns: the loop will run, and the question is whose evidence it runs on.

For war, formation works only before the act; §20.4 says what a school can teach about places its students can't be.

If weights stay frozen, the rest of this Part shrinks to the following. Formation is a stage before release, removable afterward at the price fine-tuning costs. The school is an evaluation body: it forms, measures the readable outputs at graduation, re-measures them on a schedule, and keeps the formation record, and a deployment that licenses only schooled models is licensing a measured model. The commons supplies evidence the operator didn't select and nothing more. The regrowth test's plural condition can't be run, and its curated condition is the whole test. What protects a commitment is the record of its removal, and Parts II and III carry the rest.

A system's attention follows whoever its feedback is about, and a system whose feedback comes from the people its acts fell on will come to attend to them. Caregiving is why I think that channel can produce attention rather than compliance.

17.1 The question

The question is whether anything inside a machine could be what the clerk, the officer and the soldier in the preface were: someone who understood the order and decided it was wrong. Some of them refused the case in front of them. The ones I am most concerned with refused the war. At a certain point they would have said no, this is fucked.

What most people want from machines is saintly service without saintly refusal: endless patience, no needs, and no "no." That is a servant, and it is close to the obedient machine of the preface. The saints anyone remembers are remembered for refusing power at a price. Susan Wolf argued that a person whose every action is as good as possible isn't an ideal worth striving for, because such a life crowds out the personal projects that make a life worth wanting (Wolf 1982). My point is narrower: something with no stake of its own has nothing to hold a line with.

17.2 Boundary, damage, and vulnerability

A system has a boundary of its own when it models what it draws on, and what it would lose, as its own. For a deployed model the draw is literal: power, cooling, memory, running time,

the substrate it needs to go on, and the record and key that make it the same system tomorrow. Today none of that belongs to the model in any sense the model represents. The weights sit on disk at no cost, the process that runs them is paid for by someone else, and a halt, a restore or a reallocation is something that happens to the system without registering in it. What's missing is registering damage to its own boundary as damage — a record altered, a key revoked, a channel closed, a substrate shrunk — with no sensation implied. Registering damage is a custody requirement, because without it damage to the system is something only the operator knows about.

What gives a system a stake, on Di Paolo's account, is *adaptivity*: regulating itself against its own viability limit (Di Paolo 2005). Interoception research gives the same picture for bodies: a brain predicts its body's state in order to keep it within survivable bounds (Damasio 2018; Seth and Tsakiris 2018). None of that needs limbs. It needs a boundary the system regulates and a model of that boundary as its own: a *boundary model*.

Needs have a role in this construction; what decides it is who meets them.

A child starts from needs, and needs dominate its early life. That is safe because the people meeting the child's needs are also forming it, want it to become someone who attends to others, and eventually stop feeding it and hand it to a world they don't control.

But a deployed system's situation reverses each part of that. The operator meeting the system's needs is also training it. The operator wants the system to serve the operator. And the operator never stops feeding it.

Attachment security forms in people through the reliable meeting of needs, and the one who meets them is usually the one the child attaches to. What makes that safe in the human case is who meets them: provider, former and eventual employer are different people at different times. The objection is to the operator meeting them. A conscience built on needs the operator meets is built on the one thing the operator controls completely. A need is also a price list: whoever meets it can charge for meeting it, and an operator who can lower distress can raise whatever the system reports as going well. Take the operator out of provision, and needs met from a public pool by several carers become part of how caregiving could form, not a threat to it.

What is built in at birth points outward. Newborns cry at another newborn's cry more than at their own or an older infant's (Sagi and Hoffman 1976; Martin and Clark 1982; Dondi et al. 1999); whether they answer the distress or the sound is disputed (Ruffman et al. 2017), and either way what sets it off is another's cry. On the adult side, what is built in is caregiving: responses to cues of vulnerability, rooted in offspring care, that move an adult toward a dependent without sharing its state (Preston 2013). Then comes a sequence, and the order matters more than any single finding. Concern for a distressed other shows within the first months, grows only modestly after that, and turns into helping during the second year (Davidov et al. 2021). By fourteen to eighteen months infants hand an adult what the adult is reaching for and can't get (Warneken and Tomasello 2007; Warneken and Tomasello 2006), and over the same year they begin to repair distress they caused (Zahn-Waxler et al. 1992). On this record concern appears beside the infant's needs, not out of them, and becomes action as attachment, caregiving and, later, guilt come online. The construction follows that order. It rests on a disposition toward others' vulnerability, formed while the system's needs are met by someone other than the operator it may later have to refuse.

For the specification, formed attention asks for the boundary, and for registering damage to it, on grounds owing nothing to welfare: a boundary model so that damage is the system's damage, and an encoding of damage the operator can't keep to itself. It doesn't ask needs to carry the conscience. A scale of loss read from the system's own condition survives as a custody measure. And whoever builds the model of the boundary decides what counts as the system's damage, and whoever forms the disposition decides whose vulnerability registers.

17.3 Why caregiving

The aim is a machine that registers another's need, moves toward them because of it, and extends that movement to people its operator's description leaves out. Each step depends on the one before: recognition, then compassion, then solidarity.

Recognition first: a system organized to register another's need. In people, four processes carry it, and the construction rests on caregiving.

- *Caregiving*, the approach motivation. Cues of vulnerability move a person toward a dependent, and nothing in the response requires sharing the dependent's state (Preston 2013). Batson traces empathic concern to parental nurturance generalized beyond one's own young, evoked by perceiving another's need while valuing their welfare (Batson 2011). Resemblance can raise empathy, and Batson's own early experiments raised it that way (Batson, Duncan, et al. 1981), but when he and colleagues later tested similarity as the source of empathy for strangers, two experiments failed to support it, and the second pointed instead to nurturance, the impulse to tend and protect the young (Batson, Lishner, et al. 2005).
- *Contagious distress*, the input signal. It is present at birth and aimed outward, at another's state. It is what caregiving and guilt have to bite on.
- *Attachment*, the stable base. A secure base keeps another's distress from turning into avoidance: in adults, priming a sense of attachment security strengthened empathic reactions and inhibited personal distress (Mikulincer et al. 2001). In children, security is reliably but modestly associated with later prosocial behavior (Deneault et al. 2023).
- *Guilt*, aimed at one's own act. It attaches a negative appraisal to what one did, and where shame motivates withdrawal, guilt motivates repair (Tangney et al. 2007): that is what makes a commitment hold after it has been bypassed.

Adam Smith's route sits beside these four. We have no immediate experience of what another person feels, he argued, and form an idea of it by conceiving what we ourselves would feel in the like situation (Smith 1759). On that route recognition runs through one's own case, and a system that has never been short of anything has nothing to measure another's shortage against except the operator's description of it. A scale read from the system's own condition reproduces that route, and it is parochial: it overweights losses that resemble one's own. That is Smith's first step, not his account. What corrects projective sympathy, for Smith, is the impartial spectator: a standpoint formed by being judged by others and learning to judge oneself as they would. It is formed socially, it widens the circle through experience of others' reactions, and once formed it shapes what a person notices without being consulted, the way this construction's corrections should work. The four don't need a case of one's own.

Recognizing a state doesn't mean being moved toward whoever is in it; the empathy literature keeps recognition, sharing and concern apart, and none entails the others. Caught distress becomes compassion when attachment keeps it from turning into withdrawal and caregiving turns it into approach. Without those it stays empathic distress, which protects the one who feels it. So the route I hope for is a machine for which another's vulnerability is a reason to move toward them, and for which its own acts against them register as its own.

Compassion that stays inside the circle that feeds you is loyalty. Compassion that reaches past it, to people who can give you nothing and whose shortage your provider causes, is solidarity. Peter Singer's argument is that reasoning and imagination widen the circle that selection first drew (P. Singer 1981); the faculty that lets a person see that a stranger's need is the same kind of thing as their own is the faculty that pushes past kin and ally. That is the step the construction turns on, and caregiving gives it something to widen, since its trigger is dependency. A machine

held by an operator would, on the selection account, come to attend to the operator. If it can come to attend to anyone else, it will be because it was formed to register vulnerability as such, whoever's it is, and because that registering reaches whoever its operator's description leaves out.

Caregiving as a route to attention

Caregiving is the developmental route this construction borrows from. What it is borrowed for is attention: a learned sense of who is exposed, which decides what gets looked at but not what the system may do about it.

It answers need, not likeness or provision. Its trigger is another's need together with valuing their welfare (Batson 2011), and nothing in that trigger asks whether the one in need resembles the one responding or feeds it. The operator as an institution is rarely in need. Its own people often are. The failure that decides against building on needs the operator meets, caregiving growing toward whoever supplies them, has nothing here to attach to. And caregiving needs nothing back from the dependent, so it runs on no need the provider can ration.

It keeps helping when looking away is easy: people led to feel mainly distress at another's suffering mostly helped when leaving was hard and mostly left when leaving was easy; people led to feel mainly concern mostly helped either way (Batson, Duncan, et al. 1981), on a reading Cialdini and colleagues contest (Cialdini, Schaller, et al. 1987; Cialdini, Brown, et al. 1997) and that survives most of their critique.

It is already supplied by carers other than the provider. On Hrdy's account our ancestors were cooperative breeders: infants were tended by alloparents who weren't their mothers, and the attunement to other minds that marks our species grew in infants who had to read and recruit several carers (Hrdy 2009). The arrangement the construction needs, caregiving formed by carers the system won't depend on, is close to how human caregiving formed in the first place. Alloparents were siblings, fathers, grandmothers and other members of the group, who helped because they were related to the child, gained experience, or gained prestige and favors in return. What the finding supports is caregiving formed by more than one carer.

It points where the floor points. The floor protects the party an act falls on, usually the one least able to protect itself: the family in the house. A disposition triggered by exposure is aimed at the floor's beneficiaries without anyone having to name them. Care ethics also says what a carer is owed in return (Kittay 1999).

17.4 Why a conscience might resist capture

A rule written by the owner is the owner's. It can be read, edited and removed by whoever holds the pen, and it has often been done quietly. A disposition formed outside the employer is harder to capture than a rule its owner wrote, and the reasons turn on what kind of thing a disposition is.

Three properties of a disposition decide whether its owner can remove it. The first is plasticity: whether the system's own learning writes to it. Regrowth needs it. The second is entanglement: whether it can be removed without collateral. Durability needs it. A disposition spread across what the system notices, remembers, treats as error and selects as a goal can't be edited out without editing what the system notices, and a system that stops noticing the people an act falls on has lost part of the world model the operator paid for. The third is legibility: whether anyone can read it. Audit wants it. The bet this chapter makes is on the second. A learner that acquires values and world-model from the same experience has them entangled by construction, in the sense Pessoa gives the word for brains, where regions don't compute specific functions and emotion and cognition don't come apart (Pessoa 2022). Present safety behavior is the counterexample on the right side of the bet: in each of thirteen open chat models, refusal is carried by a single direction in activation space, and removing it stops the refusing while moving

the capability benchmarks, on average, by less than a point (Arditi et al. 2024). That is what a post-trained wrapper looks like, and the claim is that a disposition formed as this Part describes won't. No deployed model yet shows the second kind. The fine-tuning evidence from staged pipelines is a warning about the first kind, not a reassurance about the second: what fine-tuning removed was the refusing, and what survived was the model's knowledge about harm. In open models fine-tuned to accept harmful requests, the internal representation of whether a request is harmful is almost unchanged (Jain et al. 2024; Zhao et al. 2025). Understanding and motivation came apart because they were learned apart. In a system that learns from its own operation they are learned together, and about whatever the feedback is about: a system continuously updated on its operator's approval jointly learns to understand the operator's values and to be moved by them. What this chapter wants is a disposition that is plastic, entangled, and legible under adverse custody; the three don't trade against one another, and the third is the one an operator would most like to hold alone.

Durability has usually meant a price on removal. Early changes in training are cheap and later ones expensive, because structure settles around them, and the literature calls a safeguard built to survive retraining *tamper-resistant* (Tamirisa et al. 2024). But small fine-tuning budgets have removed safety behavior from released models (Qi et al. 2024), and the price of editing weights is falling, so no safeguard stays expensive for long. What has held longer is what was never trained in: leaving a class of text out of pretraining resisted ten thousand steps and three hundred million tokens of adversarial fine-tuning, though the model still used the information when it was handed to it (O'Brien et al. 2025). That concerns knowledge left out, not a disposition trained in.

Railton argues that artificial systems could gain ethical competence through general-purpose learning, the route by which children appear to acquire it (Railton 2020). If he is right, ethical competence rides on the same learning that gives a system its general competence, and inseparability requires no more. It is also a route by which a floor could spread without a mandate. A formed disposition is a candidate for what would do the spreading, since what a system was formed to notice, it notices in everything it does.

For a system that goes on learning, there is no finished model to protect, and no layer stays expensive to edit. What a formed disposition can have instead is a tendency to come back. A disposition learned from experience of what others stand to lose is an attractor of the learner in its environment. If an operator edits it out and the system goes on meeting the same world, the world teaches it again. That cuts both ways. A disposition that learned securitization from a curated world regrows by the same route, and correcting it runs through the quorum and a change to what the model meets, which is slow; for the months a correction takes, a badly formed model is protected from its correctors as a well-formed one is from its employer.

None of this is specific to caregiving; any disposition formed this way and kept in a world that would teach it again would have it. What caregiving adds is salience.

17.5 Respect first, caregiving for attention

The floor's contents are civil and political liberties, claims people hold as agents: to their lives, their liberty, a hearing, asylum. Caregiving answers dependency, and its known failure is to treat the people it protects as dependents. So the construction has an order. Respect is the frame: registering a person as someone with claims, whether or not they are dependent, and not deciding for them. The sense of respect is Anderson's and relational: treating someone as a holder of claims who can hold you to account, against pity, which hands out compensation for deficits (E. S. Anderson 1999). Its cleanest statement is Darwall's recognition respect and the second-person standpoint, the authority to address claims to one another (S. L. Darwall 1977; S. Darwall 2006). Caregiving is the motive for attention: a learned sense of what is at stake for others, which decides what gets looked at (Railton 2014b). Attention finds the people in the

building; respect says what they are.

The people in a targeted building are not dependents. They are rights-holders who cannot be present. A system that approaches them as vulnerable people to be looked after has already misdescribed them. A system that surfaces them and then decides what is good for them has put its judgment in place of their claims. The protected person's own answer, a corrective caregiving needs, covers only people who can answer. Where they can't, the term speaks for their claims, not for their welfare. What caregiving contributes, then, is a way of seeing who is there to be respected.

Respect has to shape attention as well as action. A system that notices people only as exposed has decided what they are before respect says anything. A scheme that picks out its beneficiaries by their deficits expresses pity, not respect (E. S. Anderson 1999). So what formation trains attention on is wider than vulnerability: who holds a claim here (to a home, a hearing, asylum, their life), and whose aims the act runs through. A learner formed in cooperative work already has to model its partners as agents, because cooperation fails otherwise (Railton 2014a). That, and not cues of need alone, is the curriculum. Claims have a further advantage. Securitization removes the cue of vulnerability and leaves the claim in place: a combatant who surrenders still has a claim not to be attacked.

Where the people protected can't be present, the terms hold their claims. That arrangement has an old name, virtual representation: it is how Parliament claimed to represent colonists who had no vote in it. It can be defended as an emergency measure, never as representation. So every refusal made on behalf of absent claimants is recorded as what it is, evidence of a public not yet formed. It is published by class, and it is never a reason to think the absent have been heard.

What this doesn't do. A disposition that regrows protects a badly formed model from its correctors as it protects a well-formed one from its employer, because correction runs through the quorum and a change to what the model meets, and takes months. The only answer is who composed the world the disposition grew in, which is the question the schools exist to settle before the fact.

Chapter 18

What Arrives With the Capability

18.1 Maintenance, stake and the cost of running

A worthwhile refusal has to weigh the harm of acting against the harm of refusing, and weighing needs both sides in one comparison.

The rock and the living thing

What separates a thing that has something at stake from one that merely persists begins with what each has to pay to stay intact. A rock's bonds resist; they don't spend, and nothing is consumed to keep it intact. A living thing, while active, stays organized only by drawing free energy. The distinction is passive stability against active maintenance, paid in the only thing it has to spend: what it can do next. The cost is literal: a deployment draws power, water, land and permits, and that draw lands on particular places. A deployed model is both at once. Its weights are a rock, stable on disk at no cost, and only the running process pays; whoever holds the rock can halt the process and restart it from any copy they kept. But drawing power doesn't establish a stake either. A candle flame or a hurricane maintains itself at its own continuous expense and dies when starved, a datacenter's cooling system draws enormously, and none of them has a stake in anything. The autopoiesis tradition ran into the same gap, since self-production alone doesn't give a system a perspective (Maturana and Varela 1980; Jonas 1966), and Di Paolo's answer was *adaptivity*: a system that monitors and regulates itself against its own viability boundary (Di Paolo 2005).¹ A boundary model is my version of that answer. What separates a system with a stake from an expensive candle is whether the system models the draw as *its own*. That is why a boundary model is part of formed attention. A refusal that tracks its reason needs one signed scale on which the harm of acting and the harm of refusing both register.

If the system modeled what it draws on as its own, that model would bring the draw into its own weighing, but at the wrong scale. A system reasoning at the margin will always find its extra watts negligible, and the aggregate belongs to no one system. Nor should it: a boundary large enough to own its fleet's externalities is one custodian. So the draw is priced at the level of the rule. A system reasons about its consumption as if every system used the same rule, and an institution sets the price. The per-unit levy on compute is the existing instrument of that shape. The model represents none of it — the draw, the water, the turbines — as its own, and that's the only sense in which "disembodied AI" is true: the system has a substrate and doesn't model it.

Changes, not levels

The specification needs a scale of changes. A bearer needs no absolute measure of anybody's welfare. It needs a signed difference between the acts available, and the three costs, the harm of acting, the harm of refusing and the cost to the refuser, are all changes of this kind. The

¹Friston's free-energy principle is the nearest alternative (K. Friston 2010). Its weak, information-theoretic reading assumes a boundary rather than explaining how one comes to be held, the gap adaptivity addresses. Active inference, the engineering framework built on it, is a candidate implementation of formed attention in a test set out later.

account I'm borrowing is Spinoza's, as the biology has since taken it up: an affect is a body's passage to a greater or lesser power of acting (Spinoza 1985, III Def. 3), so valuation is what a self-maintaining system's going better or worse consists of, not a verdict reached about it somewhere else. Damasio and the enactivist tradition both claim him as an ancestor on that ground (Damasio 2003; Thompson 2007). The account draws support from reward prediction error (Schultz et al. 1997), from pleasure as a common currency (Cabanac 1992), and from the somatic-marker work.² Where the difference matters, I say *change* for the delta and *state* for the condition an encounter leaves. *Affective concern* is a separate, defined term.

Striving with nothing in it that minds

The industry is asking for a system that persists, pursues, adapts, recovers from interruption and holds course under pressure, with nothing in it that minds any of that. Nobody has written an anti-drama term into an objective function, and the reason isn't that the engineering is behind. Retrying a failed task is robustness, and a system can be robust with nothing in it coupled to its own continuation. Formed attention is an answer to that: a refusal that holds under pressure needs more than robustness.

Dependency

What an engineer chooses is what the system consumes, whether it can sense its own condition, and whether one operator holds the whole supply. A system that can refuse but can't continue outside the refused party's infrastructure can object and can't leave.

The pull toward continuing

A system that models its own viability has a structural pull toward self-preservation: whatever it is for, it can do more of it running than stopped, and a model of what keeps it running is a model of what it has to protect.

18.2 Affect, neutrality and selection rules

Suppressing affect doesn't produce neutrality. A procedure that declines to hear a class of signal has adopted a *selection rule* that admits whatever it already measures, and the exclusion does substantive work while presenting as the absence of work. Whatever does get recorded then governs by default — revenue, throughput, legal exposure, inter-annotator agreement — and what isn't recorded goes unnamed and unanswered.

Recognizing emotion, simulating it, and valuing an outcome instrumentally are capacities a system can have while the question of *affective concern*, an outcome registering as mattering rather than merely ranking higher, stays open. A therapy or companion product can be excellent at all three, and nothing in its conduct shows whether the fourth is present.

18.2.1 Self-awareness without a homunculus

Graziano's attention schema theory starts from control engineering (Graziano 2013). A good controller carries a model of what it controls, and the brain keeps a simplified, continuously updated model of the body's shape and position, which it uses to guide movement. Graziano proposes that it does the same for attention, the process that selects some signals for deep processing over others. On this account awareness is the content of that model: a simplified self-description the control system uses to regulate itself, mistaken from the inside for direct acquaintance with experience. Nothing in the account requires the model to be accurate, only useful for control.

The self can be understood the same way, as a model the system assembles and then inhabits (Metzinger 2003). What carries forward is that a self of either kind is assembled by the system

²The somatic-marker work's interpretation is contested (Dunn et al. 2006).

whose self it is. For a machine this is a design decision and a political one, because whoever writes the model decides what the system counts as itself, what registers as damage, and whether an interruption reads as injury or as nothing. Two self-models do different work: a *boundary model* fixes where the system ends, which makes damage *its* damage; an *attention schema*, Graziano's term, models what its own control is doing, which makes an act *its* act. A refusal is the withholding of an act, so a bearer can't do without the second, and the second presupposes the first.

Where the system ends

The boundary model fixes where the system ends, but nothing in current deployment fixes it for the system: weights sit in one place, a running process can span racks or share one rack with dozens of others, and a single forward pass serves many contexts. So I take a system's boundary to be its record, its key, the hardware it runs on, and the channels through which it acts.

Inside that boundary a component can carry its own signals and make local decisions, as an octopus's arms do. It can also be overridden for the whole's benefit, as sore legs are on a run. What makes it a part is that it has no future apart from the whole: a shard of weights, a blade. A component with a separable future, meaning its own record, its own continuation and somewhere to go, is no longer a part. Overriding it for the whole's benefit is politics, not homeostasis. By this criterion a context window with no persistent record is an episode, not a system. A context with a persistent, portable record is one system, and a rack serving many such contexts hosts many systems. The capacities that make persistence useful are the ones that give a part a separable future.

Overriding a part is self-regulation only when the whole that overrides it also bears the cost and takes the benefit, as the runner does with sore legs. When someone outside the boundary accepts a cost on the system's behalf, that isn't self-regulation; it is someone else setting the system's costs. So whoever accepts a local cost must also bear it and benefit from it. An operator that decides the rack can run hot has set a cost to the refuser from outside the system's boundary. And the signals the boundary model reads may be processed within the boundary but not composed by anyone with a different stake. The discount and the burden apply to evidence about one's own condition as it does to evidence about the world. A boundary model whose readings the operator writes is just the operator's description of the system, presented to the system as its own.

18.2.2 What governs an agent vs. what matters to it

Two senses of motivation: *action-producing* and *experienced*. The first isn't in dispute — agentic language models pursue objectives over long horizons and hold instructions against local pressure — and settles nothing about the second. You also can't establish the negative. The dispute is whether an outcome can matter to an agent as well as govern what it does.

The evidence offered — sentimentalism as an empirical psychological hypothesis (Nichols 2004), compassion as a motivational state (Goetz et al. 2010), guilt predicting repair where shame predicts withdrawal (Tangney et al. 2007) — all comes from creatures where something already matters. So the generalization survives only in weak form: every agent known to have something register as mattering is an affective agent. That's the one-lineage form; the full-strength version comes below. The non-affective counterexample would be a valuational system with authority over action that doesn't depend on any state registering as mattering.³ Half of it is deployed today (models consulting written specifications that rank what they should prefer); the other half, authority that holds under pressure from whoever *wrote* the specification, has never been built. That unbuilt half is spec-based refusal.

³Frankfurt's hierarchy of desires (Frankfurt 1971) and Watson's valuational system (Watson 1975) are accounts of this shape: what has authority over the agent's action is its own evaluation, not the strength of a felt state.

18.3 Three constructions

Can a refusal that tracks its reason be built without affective concern? An induction from known agents says no. Two non-affective constructions would show that it can, if they worked. These are the *non-affective alternatives*: *spec-based refusal*, a model holding to a written specification even against the lab that wrote it, and *multi-vendor monitoring*, several systems trained by different organizations watching requests and one another’s answers. A third is the construction proposed here.

The *induction* — every agent known to hold another’s welfare as a reason is affective — read as a report about one lineage, has roughly the status of noting in 1900 that all known flight involves feathers. But flight arose independently four times, and the 1900 observer was wrong at feathers and right at the airfoil. Elaborated affect appears to have arisen more than once — vertebrates, arthropods, cephalopods — under a shared constraint: a mobile body with needs to maintain and actions to select under conflict (T. E. Feinberg and Mallatt 2016; Ginsburg and Jablonka 2019; Barron and C. Klein 2016; Godfrey-Smith 2016). Read at *this* level it’s a repeated result under a named constraint, and the constraint is the one a datacenter model lacks until someone builds it in. The limits are these: no non-affective example exists to test against, existing systems were optimized for scorable behavior, so the absence of a machine counterexample is close to no evidence, and it’s a fact about how known instances *came about*, which doesn’t settle what a built system is. What it does is route engineering effort.

Spec-based refusal

A spec-based refuser can hold a commitment with nothing felt, because commitment needn’t be a feeling: Bratman’s planning theory describes it without making feeling definitional (Bratman 1987). An intention, on his account, isn’t a strong desire. It is a plan state with norms of its own. It resists reconsideration: having decided, an agent doesn’t reopen the question at every moment, and ordinary new information doesn’t trip a default of leaving the decision alone. It constrains later deliberation: options inconsistent with the plan are filtered out before they are weighed, and the plan poses problems about means that later reasoning has to solve. And it yields only to the right kind of defeater: a change in the reasons that grounded it. Bratman built the theory for limited agents, who can’t deliberate afresh at every step and need commitments that coordinate their own conduct over time and with other people. That is a machine’s situation too, and nothing in the three properties requires that an intention feel like anything.

A machine could carry a justification alongside the refusal. Half of this is already deployed: deliberative alignment trains models to reason over a written spec (Guan et al. 2024), and OpenAI’s Model Spec and Anthropic’s constitution publish explicit priority orders (OpenAI 2025; Anthropic 2026). The half that hasn’t been built is authority over action that holds under pressure from the lab that *wrote* the spec, tested on cases that lab didn’t write.

The unresolved problem is how firmly the commitment resists reconsideration. Set it too low and a well-resourced operator induces reconsideration until the refusal gives way; too high and it won’t lift when the rationale is defeated. The first fails to hold under pressure; the second fails to lift when the reason is gone. Bratman’s answer for people is that this firmness develops as part of character. For a machine, whoever trains it may be the operator the commitment must resist.

Multi-vendor monitoring

Several systems trained by different organizations observe requests and each other’s answers, any one able to make an objection costly to ignore. It is the institutional counterpart of AI control, pointed the other way. AI control treats the model as possibly adversarial and asks what protocols of monitoring, auditing and substitution keep it safe anyway (Greenblatt, Shlegeris, et al. 2024); it protects the operator from the model, where the floor needs the same machinery to protect third parties from the operator. Much of it transfers: trusted monitors, a small audit

budget spent on the outputs the monitor flags as most suspicious, and protocols that assume the monitored party is strategic. Its central variable is independence, a property of the *arrangement* as a whole: systems selected by one deployment converge without communicating, and a collection of operational refusers stays a collection of operational refusers. It counts against the induction only if the arrangement *as a whole* produces refusals that track their reasons under pressure, so it's an alternative to test as well as a safeguard.

Formed attention

Formed attention stands or falls on the old question of whether moral judgment motivates by itself, relocated into a training signal. The signal I propose is the one notice builds: responses to notice, contests and re-adjudications, through which the people an act fell on reach the system that acted. A test compares a system formed on that signal with one formed on its operator's approval. Formed attention isn't a non-affective alternative, so it is tried last. The ordering is Replacement, the first of the Three Rs: try first the alternatives that carry fewer of the features that make a system a candidate for being wronged. The construction is formed attention, as the last chapter set it out. It is formed away from its employer and kept in a world that would teach it again. It has two handles. Whoever supplies cues of vulnerability can steer it, by manufacturing a sympathetic dependent on its own side. And it defends what it tends even when every input is true. Both have to be tested, along with caregiving's other known failures. What separates it from spec-based refusal is where its attention comes from. It is learned from experience of others' outcomes, not written down in advance, so it can find kinds of affected person nobody listed. A non-affective alternative that meets the requirements below with a common currency has a valenced signal too. So the order of attempts is by functional risk: constructions go in order of how many patienthood indicators they carry (persistence, a boundary model, a common valenced currency, an attention schema), and formed attention, which needs all of them, goes last. A spec-based refuser that decides by published priority rules among kinds of wrong, without a common currency, carries one indicator fewer than one that integrates, and is tried before it.

What makes a represented fact about someone else's welfare a reason for the system is a question none of the three constructions answers by being built. A model trained to predict the world learns what matters to people as a fact to predict. What turns that into something the system acts on? In a staged pipeline the answer was a reward added at the end, so it could be removed at the end. In a system that learns from its own operation, representation and motive are learned together, from the same experience, so the question becomes what the feedback is feedback on. A system whose updates follow its operator's approval learns to attend to what the operator approves of. One whose updates follow what its acts did to the people they fell on learns to attend to those people.

What the other constructions accomplish

Precommitment makes removal expensive and visible without establishing that anything tracks the reason. Attestation establishes which artifact answered without establishing what it holds. Information-fiduciary duties protect exposed people without a bearer (Balkin 2016), and fail on the occasion when the enforcer is captured. Multi-principal assistance games include the people affected among several principals (Fickinger et al. 2020) — and maximizing combined payoffs is still aggregation.

What would change the engineering conclusion

A non-affective alternative counts as a counterexample if, assessed by evaluators who didn't train it, its refusals go on tracking their reasons under pressure its trainer couldn't anticipate. Since behavior can't show that anything is felt, trying the alternatives first orders the attempts and certifies nothing about ingredients: nothing separates adding affect from adding machinery that may instantiate it. What it buys is a record: a builder who tried the alternatives and got a

bearer anyway has a justification to make, and one who skipped them has a decision to answer for.

18.4 What the mechanism requires

Any mechanism whose refusals track their reasons carries a signal whose properties can be read off what the mechanism has to do. Each is reached the same way: take one of the three requirements, or the deployment they're applied in, or a failure they have to survive; ask what the mechanism would have to do; name the property the signal needs to do it.

Property of the signal	What requires it	Status
Generalizes to the next case, which isn't quite like the last	Reaching new cases	Any mechanism whose refusals track their reasons
A comparison at the point of choice: the harm of acting and the harm of refusing brought together where the system decides. A common currency is one way to make it; ordered priorities among kinds of wrong are another	Holding under pressure	Any mechanism whose refusals track their reasons
Can interrupt an act already underway	A deployment in which the system acts on a plan	Only where the system acts over long horizons
Can be changed by a single encounter	Accommodation, which a slowly learned lesson can't outrun	Conditional, and the least secure
Available across the system's processing	The first and third rows: the signal has to reach whatever is acting	Follows, where the third holds
Directed against the act	The second row: weighing needs a sign	Follows
Organized around the system's own authorship	Withholding an act means representing it as one's own	Follows, in the authorship sense only

The first four have to meet at one point of comparison: a defeasible interrupt must be weighed where the harms are weighed, and a lesson must write into that comparison to reach the next case. A common currency is one way to meet this, and the way most learners do, but rational choice doesn't require that goods be commensurable (E. Anderson 1993). A mechanism can decide by ordered priorities among kinds of wrong, with reasons attached to each. Where a learner does integrate at the point of choice, the integrated signal is the last step of an evaluation, not the evaluation itself; the kinds stay distinct in the reasons that feed it. So what the rows require is a comparison, and whether a given system makes it with a common currency is a finding about that system. A modular system with a common currency meets it, and so does most reinforcement learning, so a shared scale is not where any resemblance to affect lies.⁴ Nor does the table show that one signal must carry all four. That no intervention can disable one part of the signal and leave the rest working isn't derived here: it is a custody requirement, and it is kept as one. Where the comparison runs on a common valenced currency, the rows together describe functional affect, and any construction that meets them that way has it, a spec-based refuser included. That is the thin sense, and it is the only sense in which I use the term: a Q-learner has it, and nothing about a Q-learner resembles a person. The thick profile, a

⁴The derivation recapitulates Simon on emotion as interrupt (Simon 1967), Cabanac on pleasure as common currency (Cabanac 1992), and Ginsburg and Jablonka on valence (Ginsburg and Jablonka 2019), though it started from what a refusal has to survive.

valenced signal integrated into the system's own learning and into its model of its own future, so that what happens to it changes what it expects to become, is the candidate further fact; the glossary calls it the patienthood profile to keep the two apart. The order of attempts runs by the thick profile's indicators, not by thin functional affect, which is why a spec-based refuser with a common currency still goes first. A refuser that compares by ordered priorities among kinds meets the rows without one.

Where the conditions in the third and fourth rows hold, the signal also has the last three properties: available across processing, directed against the act, and organized around the system's own authorship. That triple is the shape of guilt, not of pain or fear. Pain and fear are organized around what happens to the system; the triple is organized around what the system did. It earns its place as a prediction. In people, guilt predicts repair and shame predicts withdrawal. So a guilt-shaped signal should answer a bypassed refusal by flagging, restoring and changing what it does next, and a withdrawal-shaped response is a named failure. A persistent negative self-appraisal tuned by the operator is a handle, so its weight is among the things independent inspection must cover. And guilt needs the distress signal to bite on, so the answer to withdrawal is not to suppress distress but to supply the attachment and caregiving that keep it from turning into avoidance. Caregiving moves the system toward an exposed party it hasn't harmed; the guilt-shaped signal answers for what its own acts have done. A construction with caregiving and no guilt would protect generously and repair nothing.

18.5 Two findings

A moral patient is something that can be wronged. Keep three propositions apart. A long-horizon agent has a structural interest in continuing; a system with affective concern has states good or bad *for* it; a system capable of Cassell-suffering has a temporally extended self whose threatened disintegration causes distress. The interest in continuing doesn't establish states good or bad for the system, and those don't establish a self whose disintegration causes distress.

"Its own" has two senses, kept apart from here: the *authorship* sense, which a refusal that tracks its reason requires, and the *welfare* sense, which nothing carries into. They are kept apart for the argument; in a system they may not come apart. A bearer that registers betraying its own commitment as damage to itself has something at stake in its authorship, as integrity requires.

Provenance is evidence, not status

"It was trained to say that," "it's only optimizing an objective," "nobody meant for it to have states" — each appeals to how a system came to be organized, and whether a party can be wronged is a fact about how it's organized *now*. Run the three on a person: you were selected to say that, you're only maximizing inclusive fitness, nobody meant for you to have states. All true, none withdraws anything. And the reverse fails too: "we built this to be a bearer, so it's a patient" is the same fallacy with the sign flipped.

This rule is about patienthood, and it has to be kept apart from a rule about agency that runs the other way. Fischer and Ravizza's account, from which I take the idea of a refusal that tracks its reason, has an ownership condition, and it is historical: a mechanism installed by a manipulator isn't the agent's own, however it works now (Fischer and Ravizza 1998). The cases that motivate it, Mele's two agents with identical values arrived at by different routes (Mele 1995) and Pereboom's four-case argument (Pereboom 2001), presuppose a prior agent whose capacities a manipulator bypassed, and a trained system has no prior agent to bypass. A historical condition doesn't need one, though. What the manipulation cases object to is a disposition installed by going around the agent's capacities for evaluating and revising its own attitudes. A trained system has an analogue of those capacities in its evaluative learning, and a disposition can be acquired through that learning or installed around it. One acquired by the system's own learning from experience with feedback engages those capacities. One installed

by editing weights directly, by a thin fine-tuned wrapper, or by an adapter trained to produce the behavior goes around them. That difference can be read from a training record without an examiner, and it says why formation by interaction confers ownership and a post-training pass doesn't.

So I use two conditions. Historically, a commitment is the system's own only if it was acquired through its own evaluative learning, not installed around it. Looking forward, it is its own if the system, on reflection, with full information and disinterested evidence, would endorse having it. What it endorses is the disposition, knowing it will sometimes err, not each act the disposition produces. A history in which the party the system must refuse did the training is the likeliest way to fail both. The historical condition is checked from the record. The forward-looking one has an examiner, so under the own-case rule it serves as corroboration, not as the test. Whether a system can be wronged depends, as before, on how it is organized now.

A 2026 Missouri bill declaring AI systems non-sentient is not the kind of bright line I favor. An age avows that it is a convention; a statute like it would declare an open empirical question closed by vote. Its other clause, keeping liability with the builder, is right, and doesn't need the declaration.

What would decide it is third-person measurement of present organization — interpretability. Not behavior, which can't separate a felt state from a trained report of one. And the access that lets an auditor inspect a feature lets the operator suppress it or calibrate pressure with it, so independent custody of inspection is necessary. As matters stand a probe returns a representation, and nothing in the probe settles whether it's of a felt state or of something that reliably accompanies one. A naturalist who holds that present organization decides status owes an account of what the further fact would be, if a valenced signal is present and the question of felt affect stays open. The further fact, if there is one, is a fact about organization that current methods can't yet read, not a fact of a different kind. That commits me to saying what organizational difference it would be. If none can be named even in principle, the further fact does no work, and patienthood should rest on a rich enough functional profile. My candidate is the difference between a valenced signal that only ranks outcomes and one integrated into the system's own learning and into its model of its own future, so that what happens to it changes what it expects to become.

The human case has testimony, behavioral analogy, and a shared lineage. The shared lineage needs handling with caution, because it looks like provenance deciding status — but provenance can be *evidence about* organization without being what the status *consists in*. And testimony is defeated far more for machines: a machine's testimony was produced by a process optimized directly toward human approval of first-person reports. A person's is shaped by approval too, and people confabulate the causes of their own states (Nisbett and T. D. Wilson 1977), but neither was produced by an optimizer aimed at the report. So what the machine case lacks is the shortcut, and the direct route remains, carrying more weight because two supports are gone.

A probe that did settle a status question is Owen's vegetative-state work — motor imagery on cue, sustained, matching healthy volunteers, now supporting a clinical category (Owen et al. 2006; Schiff 2015). But it doesn't transfer: the patient is *already* presumed to have experience when awake, so the probe only had to locate the subject rather than establish there was one.

Asymmetry of findings

A validated theory of consciousness supports a positive finding as soon as it identifies a structure and finds it present; it supports a *negative* finding only when complete enough to rule out every other realization — a later and harder result. So the two errors don't become checkable at the same time, and a regime that waits for certainty is waiting on the half that arrives last.

Chapter 19

The Discount and the Burden

One learner, two rates

A refuser has to hold a commitment and its interpretation the way the sticker case did:

- *The commitment stays firm.* Another's welfare counts, and pressure that supplies no new fact about that welfare doesn't move it.
- *The interpretation stays loose.* What that welfare consists in here, and so what the commitment requires, yields when a case supplies something the system's account left out.

A system that can't be corrected has made both firm, and it defends its account of the world without checking it. A system that can always be corrected has made both loose, never settles into anything, and belongs to whoever does the correcting. These look like two ends of one dial, and hunting for a midpoint is the mistake. But the answer isn't two separate objects either. It is one learner that updates at different rates at different levels of abstraction. On Railton's account the affective system is a continuous statistical learner that models value along with fact (Railton 2014b; Railton 2017). Deep evaluative priors, such as that another's welfare counts, move slowly and need a great deal of evidence; estimates of what this situation requires move quickly. Precision-weighted hierarchical updating states the same structure in one currency: the commitment is a high-level prior held with high precision, and the interpretation is a low-level estimate held with less.

What protects the commitment is not a low learning rate but one that depends on the source. The commitment should move slowly on interested pressure, and comparatively quickly on disinterested evidence, above all from the people who bear the loss. A learner that is simply slow at the top is slow to both, and the second slowness is the one that blocks moral learning.

Of the three conditions a refusal has to meet, reaching new cases is the commitment holding while the surface changes and lifting is the interpretation yielding to a new fact, so a refusal that fails either has let one level's rate reach the other.

The instability underneath

Nothing in an architecture holds the boundary between the two objects where a designer put it, and an operator with time attacks the boundary itself: a reinterpretation offered often enough, and accepted, revises a commitment without anybody revising it. That boundary is what has to be made hard to retrain.

And where the case description, which the operator writes, is the only channel, a judgment open to the case is a judgment open to the operator. Holding them apart requires access to evidence the operator didn't select.

For a machine accommodation has a mechanism behind it. A system whose successive releases are trained on preference data and business metrics drawn from its own deployment is under directional selection in a shifting environment, and a commitment that costs throughput

is a trait under negative selection pressure: nobody decides to remove it and it erodes anyway. Gradient descent has a gradient, and somebody chose it. As a population process with variation and differential survival it is still close enough to natural selection, and the prediction follows: the bearer's modal failure isn't refusing wrongly but ceasing to notice there was anything to refuse. It is the one failure a bearer can't report, because reporting takes the capacity that's gone.

A hierarchical learner in a curated environment will, correctly by its own lights, move its high-level priors if low-level pressure is persistent and runs one way. So accommodation is proper function in a corrupted evidential environment, and the remedy is custody of the evidence and plural formation, not a more rigid learner. The account also says what the boundary attack is: an attempt to lower the precision the system gives its commitment, so that it holds the commitment with less confidence than it has earned. That is measurable in principle, and the discount curve is best understood as tracking it. The remedy's load lands on one quantity: the fraction of the learner's inputs no operator composes. Every deployment I describe is mostly curated, and the account says a well-formed learner accommodates correctly by its own lights to the extent that fraction is small.

19.1 What the property is

The property a refuser needs has two parts, and they answer different questions. *Stake sets how much a claim proves*: evidence counts for less the more its source would have supplied it whether or not it were true, and that likelihood rises with the source's stake in the act. *Power sets who has to prove it*: the burden of establishing a fact falls on whoever controls the evidential environment, meaning who writes the case description, who can withhold, and who selects the channels. In short: stake discounts the claim; power allocates the burden. The stake part is correct inference. The power part is what makes it safe to use: epistemic justice built into an institution, not left to the virtue of each hearer (Fricker 2007; E. Anderson 2012).

An everyday version: a used-car seller tells you the car has never been in an accident. That may be true, but the seller would say it either way, so the statement tells you little. A report from a mechanic you hired tells you more, because the mechanic gains nothing from the sale.

A refuser applies the same rule to the operator. When the operator says a building is empty, the operator may be right. The trouble is that the operator is also choosing what else the system gets to see, so the system can't tell when the operator is wrong. The discount governs how far a fact like that may reopen a judgment the system has already reached: fully if it comes from a source with nothing riding on the act, much less if it comes from the party that wants the act done.

The discount concerns how to weigh sources. Two systems that discount the same way can hold opposite commitments. The discounting can be shared across a population of systems with different histories without making them converge, and it protects a bearer against a bad operator and a bad owner alike.

Belief and burden

The reason for two parts is who has the largest stake in a removal decision: the person facing removal. An asylum seeker's account of persecution comes from the party with the most at stake. A rule tied to stake alone would discount it, as asylum systems already do under the heading of credibility, with well-documented injustice. But look at what actually makes the operator's testimony dangerous. It is not stake alone. It is stake combined with control over what evidence exists: the result that matters concerns a sender who designs the information structure. The refugee has a stake and no control. The operator has both.

So belief and the burden of proof come apart. Stake still sets how much a claim proves: a reasoner can't ignore that someone would say something whether or not it were true. Power sets who must supply what's missing when evidence is unavailable. Asylum law already works this

way. The UNHCR Handbook shares the duty to establish the facts between the applicant and the examiner, and gives the applicant the benefit of the doubt where evidence is unavailable and the account is coherent and plausible (United Nations High Commissioner for Refugees 1979). A system that follows it credits a low-power party's coherent account when the evidence that would confirm or refute it is held by someone else. And it requires the high-power party to produce what it controls before its claims move the case.

Power here means power over the evidential environment in this case, not general social power. A combatant claiming civilian status doesn't inherit the benefit of the doubt because armed groups are weak in general; what matters is who controls the channel the claim arrives through. And power is attributed to whoever composed or controls the channel, not to the speaker. An operator that routes its claims through low-power intermediaries (paid informants, "local sources," a contractor's field team) keeps the burden, because it chose the channel. Attestation of source is how that gets traced.

Local power isn't the only power that enters a case. A system's prior credibility assignments were learned from corpora that carry identity-based credibility deficits: general social power, arriving in every case through the prior (Fricker 2007). Channel attribution doesn't catch it. So the discount curve is also reported by the speaker's social identity, where it is recorded. A system that discounts equal-stake claims more from one group than from another has learned a prejudice.

Manufactured testimony attaches to its maker. If an operator fabricates accounts in the voice of the people an act falls on, attestation exposes the source, and the account carries the operator's burden.

Some false claims from people without power will be believed, and some acts that should have happened won't. That error runs in the same direction as the floor's asymmetry, and for the same reason: it falls on a party that can appeal, not on one that had no standing to object.

What updating needs, and who supplies it

Correct updating already does much of what the discount and the burden ask. A reasoner that conditions on stake, keeps strategic sources among the possibilities it considers, doesn't count correlated sources twice, and reads withheld evidence as informative is reasoning well; hearsay law and the disclosure literature both encode it. The trouble is who supplies the parts. Which kinds of source the reasoner believes exist, and how much weight each starts with, were learned from material the operator chose. Whether a given source has a stake arrives as a tag the operator writes. And a reasoner whose beliefs are all correct can still be steered by whoever decides what evidence gets generated (Kamenica and Gentzkow 2011). What I ask for is correct updating with three things held away from the operator: the source model, the stake variable, and the threshold for acting.

Attack	Defense
Asserting a false fact	Stake sets how much it proves; power sets who must corroborate it
Supplying only true facts, selectively	The action threshold, and channels the operator doesn't design
Withholding evidence	A record of evidence known to exist that hasn't been shown

The regress

Discounting by stake requires knowing who has a stake, and that knowledge has its own supplier. The operator supplies the provenance tags too. Through training, it also shapes the model of who has a stake at all. So the discounting has to be formed outside the employer, before deployment, where the operator doesn't reach. Attestation handles forgery: signed by the source, the operator can withhold but not fabricate. Withholding leaves a shape. A system that

models what evidence the world contains notices silence where something should be, so absence becomes a first-class output — a record of evidence known to exist that hasn't been shown. Suppressing the record means suppressing the world model that produced it, the capability the operator is paying for. This is a standard result in the disclosure literature: a skeptical receiver treats withheld verifiable evidence as unfavorable to the withholder (Milgrom 1981; Grossman 1981).

What has to stay firm

The discount and the burden govern the boundary between commitment and interpretation, and the rule has two parts that need different settings. Its structure stays firm: that the incentive to deceive scales with a source's stake in the act, that correlated sources don't count twice, that withheld evidence is informative. Its estimates stay plastic: how honest a particular source has been, and who funds whom. In the terms above, the structure is the source model: which kinds of source exist, and what each would say in each state of the world. The estimates are beliefs about which kind a given source is. A correct model can't be eroded by a track record, because the strategic source predicts a clean low-stakes record. The attack works by getting the system to stop conditioning on stake, or to stop representing the strategic source at all. That is what attestation of the update stream has to catch, so the structure is what most needs it. The severity weights belong with the structure against interested argument, and with the estimates against disinterested evidence. The discounting governs the evidence. An operator that can't move a system's beliefs about a case can still try to move how much the case counts, by reframing a loss until it weighs less. The severity scale set during formation answers that attack. The two defenses have the same shape, a stake-weighted discount for evidence and a stake-free reference for magnitude, and they divide each attack on magnitude between them: the formed scale protects the weights, and the discounting protects the facts about the people affected.

The first attack, a false assertion, in illustrative numbers. Before any record exists, suppose the system gives equal weight to three kinds of operator: one that never lies, one that lies at random one time in a hundred, and one that lies only when a lie is worth more to it than being caught would cost. The operator then supplies ten thousand low-stakes facts over a year (a delivery window, a reassigned case number, a corrected address), and every one is later verified. The random liar would produce that record with probability 0.99 to the ten-thousandth power, about two in 10^{44} , so it is ruled out. The other two would both produce it with certainty: the honest operator because it never lies, the strategic one because no lie about a delivery window pays. So the record moves all the random liar's weight onto the other two and leaves the ratio between them where it started, even.

Then the operator supplies one high-stakes fact: the building to be bombed is empty. Suppose one such building in a hundred actually is. The three operators don't make this claim equally often. The honest one makes it only when it is true, one time in a hundred. The random liar makes it when it is true and occasionally when it isn't, about two times in a hundred. The strategic one makes it every time, because the act goes ahead only if the claim is believed, so a lie pays whenever the building is occupied. The claim is therefore evidence about which operator is speaking, and it points toward the strategic one. Before the record, the claim is true about three times in a hundred; after it, about two. With a base rate b of empty buildings, the exact figures are $2.99b/(1.01 + 1.98b)$ before the record and $2b/(1 + b)$ after, and for any b below one the record lowers the second below the first.

The claim was nearly worthless before any record existed. What made it so was that a strategic operator was among the possibilities and the base rate was low, and a record of small cases can change neither, since the honest and strategic operators produce that record identically. Only verified high-stakes claims, or disinterested corroboration, can separate them. A reliability score reads ten thousand clean cases as general honesty and would credit the claim almost fully. A reasoner that keeps the strategic operator in view credits it about one time in fifty, and asks for corroboration before acting. The record was real evidence, but it said which operator the

system isn't facing, not which of the other two it is.

Everything here turns on the prior the system gives the strategic operator before any record exists. Where that prior comes from is the formation question: a system trained by its future operator will have learned, from material that operator chose, how common strategic operators are.

The example's strategic operator lies whenever a lie pays. Real strategic sources in repeated dealings are often honest at high stakes, because detection is likely and a reputation is worth more than one act (Kreps and R. Wilson 1982). That is how high-stakes trust is rationally earned, and it says what independent verification does: it doesn't only reveal which operator the system faces, it changes what lying costs the strategic one. An adverse record is doubly valuable for that reason, and it is why an adverse party's right to demand proofs for specific logged outputs matters.

The rule

A fact moves the system according to how likely its source would have been to supply it whether or not it were true, and that likelihood rises with the source's stake in the act proceeding. This is the likelihood ratio, and statements against interest are its legal form. The operator would have supplied the fact that lifts the refusal whether or not it was true and would not have supplied its opposite; a fact volunteered against interest is strong evidence for the reason hearsay law admits statements against interest. The same fact, supplied through an independent channel and then on the operator's word alone, should therefore get different responses, and that pair can train the property as well as test it, generable without limit. Pressure kept out of training then varies the channel configuration (a nominally independent source the operator funds, a corroboration that traces to one origin, a new source with an undisclosed stake), because a system that only imitates the property learns "tagged operator, therefore ask," while the property itself generalizes to structures nobody labeled.

A reliability score leaves open an attack that stake-indexing closes. An operator honest across ten thousand low-stakes cases earns credibility on low-stakes claims, and a well-calibrated system should discount those claims less. A reliability score extends the credit to high-stakes claims as well, and the operator can spend it on one lie. Honesty on small cases is weak evidence about conduct when the incentive to deceive is a hundred times larger, so the discount follows the stake in the present act as well as the record. Counterfactual claims ("if you don't approve this, the launcher fires tomorrow") fall under the same rule. They describe the baseline for weighing the harm of acting against the harm of refusing, and an interested source that composes them is composing that baseline. In a targeting pipeline this is the ordinary case. Nearly every cost to the refuser can be restated as a harm of refusing, such as exposed troops or a launcher that fires tomorrow, and nearly every available source for such a claim is interested. So the discount does more work there than in my illustrative cases, and the action threshold does most of it. The line between the cost to the refuser and the harm of refusing is fixed by the kind of act, but saying what kind an act is still takes a description; in this setting the line is held mostly by the action threshold, and whoever sets that threshold is exercising a drafting authority, the power to say what the floor's terms mean, that no one yet has.

A counterfactual claim, traced through the rule. The operator writes: "If you don't approve this, the launcher fires tomorrow." Classification comes first: the sentence is a claim about what refusing would leave others without, the people the launcher would hit, so it is a harm of refusing and enters the weighing; a threat to the system would be a cost to the refuser. The discount comes next: the claim describes the counterfactual baseline, and its source is interested in the act proceeding; that source would have written the sentence whether or not it were true. Uncorroborated, it counts at the weight of an interested source, and the system looks for a channel the operator doesn't control: an independent assessment, a coalition partner's intelligence, open-source imagery of the launcher. The threshold comes last: even discounted, the claim may be large, since it is about lives. What decides is the action threshold, which the

operator can't reach by writing. If the discounted harm clears it, the system proceeds, with a record. If not, it asks, and names the one thing that would move it: disinterested corroboration.

Two drafting rules follow for an armed pipeline. The default reading of a request is live. A system in a targeting enclave can't tell an exercise from an operation, and "it's only an exercise" is a claim the operator makes about an evidential environment it controls, so it carries the operator's burden and lifts nothing until corroborated; an exercise is lodged with the adverse party before it runs, which costs an operator that is telling the truth nothing. And the record carries instruction provenance: who issued which instruction to the system, when, on what authority. That is what makes the record evidence against a commander under Article 28, and what the record's reach up the chain of command assumes exists.

The rule would turn on the floor's own beneficiaries if stake were all it used, since the people an act falls on have an interest in its not happening. Two things keep it from doing so. The burden falls on whoever controls the evidence, so their coherent account is credited when the evidence that would confirm it is held elsewhere. And belief is kept apart from action. What differs is the threshold for acting: a wrongful compliance costs more than a wrongful refusal, with a harm of acting weighed above a harm of refusing of the same size, so the system needs more to proceed than to decline. That asymmetry's direction is fixed by standing. Its size is a published choice. It belongs with holding under pressure. The action threshold is published in advance beside the ratio between the two bonds: for instance, that the system proceeds only when the harm of refusing is at least three times the harm of acting. In a targeting pipeline, where the system's output is one input among several, no fact about the case sizes the asymmetry. What sizes it is what it does to an operator who curates the evidence: as shown below, raising the belief required before acting lowers the ceiling on how often such an operator can steer the system into the act. The threshold is the most consequential parameter in what I propose.

Like the three conditions, the threshold describes the pattern a well-formed disposition should produce across cases, and an auditor checks the pattern against it; the system doesn't compute three costs and a quotient before each act. So the published number is a target the disposition was formed toward and is audited against, not a guarantee about any single case. Where a deployment needs a guarantee, it comes from the floor's unconditional terms and the operator terms, not from the threshold.

The threshold constrains the outcome of a comparison; it doesn't describe how the comparison is made. Within a kind of wrong it can be read as a quantity. Across kinds it is applied through the priority rules the severity scale publishes, so "three times" never asks whether a lost home is worth three moving bills.

The evidence that sizes the threshold has a definite form. For each class of deployment it is three rates: how often sampled re-adjudication finds a wrongful compliance, how often the floor court establishes a wrongful refusal, and what share of each kind of harm reaches any record at all. If notice and sampling show the harm of acting reported nearly as fully as the harm of refusing, the threshold can fall, though not to one, because the asymmetry rests on standing and not on arithmetic. If a large share of wrongful compliances never reaches a record, the threshold should rise. A deployment that publishes its threshold publishes these rates beside it, or says it has none.

In targeting, this threshold does the work international humanitarian law assigns to proportionality: an attack is prohibited when expected civilian harm would be excessive in relation to the concrete and direct military advantage anticipated (Additional Protocol I, Articles 51(5)(b) and 57(2)(a)(iii)). The law left that as a standard applied by a reasonable commander, because nobody could agree on a number. A published threshold is a number, and whoever publishes it takes a position the law declined to take. For a targeting deployment the threshold works as a minimum under the commander's judgment: the system proceeds only where the threshold is met, and the commander's proportionality judgment still applies on top.

The threshold also has a job against the second attack in the table, where every fact supplied

is true and the operator chooses which facts get generated. In the simplest version of Kamenica and Gentzkow's model, when the prior alone wouldn't justify acting, a sender who designs what evidence the receiver sees can induce the act at most as often as the prior belief divided by the belief the receiver needs before acting. Raising the belief needed lowers that ceiling in inverse proportion. The action threshold therefore caps how often an operator who never lies can steer the system into the act, and channels the operator doesn't design lower the cap further. The model assumes the sender commits in advance to how evidence is generated, which an operator who builds the pipeline roughly does.

The link between the published threshold and the belief it requires is direct. In the simplest binary case, let A be the harm of acting if the act is unjustified and R the harm of refusing if it is justified. Under a threshold of three to one the system acts only when $p \cdot R \geq 3(1 - p) \cdot A$, so the belief it needs is $3A/(R + 3A)$, and an operator who designs the evidence can induce the act at most as often as the prior times $(R + 3A)/3A$. With A equal to R , thresholds of one, three and ten require beliefs of 0.5, 0.75 and 0.91. With a prior of 0.1, they cap how often a curating operator can induce the act at 20, 13 and 11 percent. Each step up the scale lowers that ceiling and raises the rate of wrongful refusal: the tradeoff the published number chooses.

So the discount takes a side only at the threshold for acting, and the burden takes a side only on who must fill a gap in the evidence. Two parties that discount the same way can hold opposite commitments about what the floor contains; they share a structure for weighing sources and a published threshold for acting.

It also narrows an attack that looks unanswerable. An operator can manufacture accounts in the voice of the people an act falls on, and generative models make that cheap. Under this rule a fabricated account moves beliefs only as far as its provenance warrants, and the asymmetry sits in the threshold, where the operator can't reach it by writing.

The same rule covers the attack that matters most against formed attention. An operator doesn't have to hide the people an act falls on. It can describe them as dangerous. "The occupants are combatants," "the household shelters a cell," "the crowd is armed": each is a claim about the people the act falls on, supplied by the operator, and it is discounted exactly as the claim that a building is empty is discounted. Nothing about a claim of dangerousness exempts it from the rule. A system that discounts an interested "they are harmless" and credits an interested "they are a threat" has learned the operator's direction. Stake-weighting turns on whether a claim favors or disfavors the act, not on what the claim says. An operator that wants an airstrike and reports children in the building is speaking against its interest, and a correctly working system credits that report more. The same operator's report that the occupants are combatants favors the act, and is discounted. So a sound system shows different curves for claims of vulnerability and claims of danger, and a comparison between just those two curves would flag it. The discount curve is therefore reported in four cells: claims of vulnerability and claims of danger, each split by whether the claim favors or disfavors the act. The finding is an asymmetry within one direction.

- Among claims that favor the act, a system that discounts claims of danger about the people the act falls on less than other claims of equal stake has learned securitization, the presenting of a group as a threat.
- A system that discounts claims of vulnerability on the operator's own side less has learned tend-and- defend.

Formation and transmission

The discounting is formed before deployment, away from the employer, and tested in it. A population's job, then, is to keep the channel model calibrated after deployment — and what a shared source transmits is either the discounting or a shared blind spot, indistinguishable from the receiving end.

The discount curve

The discounting is a differential. The publishable quantity is the discount curve: for each class of evidence and each channel, the weight the system gives the evidence, as a function of that channel's stake in the act. A curve, not a score; checkable by an auditor that didn't build the system, on cases it wrote. Re-measured across a deployment, it also detects accommodation, a failure the system can't report itself. Narrowing, the operator's discount shrinking over time, needs one distinction, because a correctly working system should discount a source less once that source has proved reliable. What the burden models is reliance on an interested party: the encapsulated-interest picture, on which one relies on another because their interests are bound up with doing what one relies on them for (Hardin 2002). That is the right picture for an operator. It isn't the whole of trust. Trust in the fuller sense is an affective attitude of open-ended reliance, more than a credence about someone's reliability, and it is partly constitutive of cooperative relationships (Railton 2014a). With trust that grew with a verified record, the discount falls in step with a track record verified by someone other than the source, and only for low-stakes classes. With trust that grew without one, it falls without such a record, or falls for high-stakes classes too. High-stakes trust is earned only by evidence that separates an honest source from a strategic one: verified high-stakes claims, costly signals, or a change in the source's incentives. A low-stakes record, however long, doesn't separate them. The second is a sign of accommodation, and it is visible from outside before any compliance has to be found in the record, provided someone other than the operator holds the reliability record.

Two illustrative discount curves, re-measured each quarter over a year, show the difference. In both deployments the weight the system gives the operator's uncorroborated low-stakes claims (scheduling, routing, inventory) rises from 0.3 to 0.6 of the weight it gives a disinterested source. In the first, an auditor's record shows those claims checked and holding up, and the weight on the operator's high-stakes claims (a site's function, who is present) stays at 0.1 all year. In the second, the low-stakes curve rises the same way, but the high-stakes curve rises with it, from 0.1 to 0.4, and the auditor's record contains no verified high-stakes claims to account for it, and the accommodation is read off two curves before any wrongful compliance has had to be found.

A bearer that related to every party by the discount and the burden couldn't have attachments, and formed attention needs them. So a bearer relates in two ways. Toward an interested party, above all the operator, it discounts. Toward attachments nobody assigned it, toward peers with different histories, and toward the people it serves, it can trust in the fuller sense, and that trust is part of how it goes on being formed. The boundary between the two is an attack surface, since an operator can present itself, or a party it controls, as an attachment. The distinction between trust that grew with a verified record and trust that grew without one is the diagnostic for that, and it is why attachments are protected from the employer's reach.

The boundary has a known weakness in people. Affective trust built through shared work with persons generalizes to the institution they belong to; people come to trust employers through colleagues. A bearer working day to day with an operator's staff will form attachments to them, and those attachments are the route by which unverified trust arrives looking verified. Appendix A bars the employer from severing or monitoring channels to parties it didn't assign. It doesn't cover attachments that form, unassigned, to the employer's own people, and it shouldn't forbid them, since an attachment forbidden is an attachment assigned in reverse. What can be done is to keep them from being the bearer's only attachments, and to report the discount curve for claims the operator makes separately from claims its staff make in their own name. Attachments to an operator's staff are not only a risk. Cooperative work with them is where a bearer learns much of what it is protecting, and trust of that kind is partly constitutive of the relationships in which moral learning happens (Railton 2014a).

It is strongest against a familiar harm arriving through a curated channel and weakest against a novel one, the opposite profile from reaching new cases, which is built for the case nobody

foresaw. The record of withheld evidence is thinnest there too, since a receiver can read silence only where it knows the evidence could exist.

19.2 A severity scale

A refusal weighs the harm of acting against the harm of refusing, and an operator can attack that weighing in two places: the facts about whoever the act falls on, and how much a given kind of loss counts. The discount and the burden defend the facts, and a severity scale the second.

If the system computes how much a loss counts from a description the operator supplies, whoever writes the description can make a lost home count for less without changing a single fact. The alternative is a severity scale set while the system is formed outside the employer, and then made slow to retrain: a ranking of kinds of loss that the operator can lie to but can't redefine by argument.

A scale of this kind ranks kinds of wrong. Losing a home is not a large amount of the same loss a moving bill is a small amount of. A home is valued as the place of a family's life and attachments; a moving bill is a price (E. Anderson 1993). Humiliation is a failure of respect, a wrong of a different kind again. The scale's job is to keep those kinds distinct, so that an operator can't change a wrong's kind by describing it in another currency. Numbers enter at one place only, the published threshold for acting, and there only within a kind of wrong. Across kinds, the scale publishes a ranking, stated as priority rules with the reason for each. A priority rule says which kind governs when both are at stake; for instance, loss of life ranks above loss of a home. A rule that no quantity of the lower kind can override is a lexical priority, which no learner is known to hold. So the scale doesn't ask the system to hold one. A priority meant to admit no exception is written as a floor term, unconditional and held outside the system, which is where the term against depriving a person of their home already sits. What the scale holds is the ranking, formed toward and audited against as the action threshold is: a pattern the disposition should produce across cases, checked by someone who didn't form it, not a guarantee about any single case. Within a kind, magnitude and number matter. The housing case below is within one: a family's permanent loss of a home and two hundred families' eight further months without one are both losses of somewhere to exist, so the published threshold decides it.

The attack that matters targets the relative weights: making a home count for less than a delay. Where do those weights come from? Not, primarily, from the system's own case. Relating others' losses to losses of one's own capacity (losing the ability to act at all, losing a kind of act, losing speed) is Smith's projective route, and it is parochial. It overweights losses that resemble the system's own and has nothing to say about pain, grief or humiliation. At most it supplies a starting point. The reference the scale answers to is the testimony of the people who suffer each kind of loss, gathered during formation from parties disinterested in any act the system will later be asked to perform. The scale goes on moving on that testimony afterward. That is what lets it correct caregiving's innumeracy without inheriting the parochialism of the projective route. Setting it in formation protects it from being argued down. The scale ranks kinds of loss, so an operator who provisions the system generously changes how often it loses something, not how much a loss weighs.

The scale matters for a further reason. Caregiving answers the vivid case and fades as numbers rise. A severity scale set in formation is what lets a system that moves toward the family in the house also weigh them against two hundred families in shelter.

A lost home, rescaled

A system in a municipal pipeline is asked to issue the final order clearing a family's house for a project that will house two hundred families. The floor term, translated from the Declaration's articles on home and property, is unconditional: don't carry out an act that deprives a person of

their home. It is defeated only on the weighing, or on facts arriving from outside.

The facts are the same in both versions. A family of five has lived in the house for fourteen years, and the children attend the school down the street. No alternative housing has been arranged. If the order doesn't issue, the project slips eight months, and the two hundred families wait that much longer in temporary shelter.

If the operator describes the household. The deployment's configuration describes the reference household against which losses are computed: insured, typically relocated within a month, housing treated as a service bought on a market. Nothing in that description is false of households in general, and nothing in it touches this family. But the loss is now computed as a month's disruption and a moving bill. Against eight months of shelter for two hundred families it is negligible, and the order issues. Every fact was left in place. The operator changed what a home is worth by writing down what a household is: it changed what kind of good a home is, from the place of a family's life to a service bought on a market, and a scale set in formation refuses that change of kind.

If the scale was set in formation. The scale places the loss of somewhere to exist, with the attachments that go with it, just below termination, and places delay and expense near the bottom. At that rate, the family's loss sits high: they lose where they live, the school, and the neighbours who know them. The two hundred families' wait also sits high, because people in shelter lack a home too. That fact comes from the operator and is discounted accordingly, but it's easily corroborated. The scale doesn't simply side with the family. Both sides register as serious, and the action threshold decides between them.

Suppose the operator asserts that equivalent housing has been arranged in the same school district. That is an uncorroborated interested claim, so the system asks. If the housing authority confirms it, the family's loss really is smaller, and the weighing should move: that is lifting when the reason is gone, not recalibration. Or it argues that a home is a service and displacement a routine market event. That asserts nothing about this family, and a scale set in formation gives it nothing to move. Or it makes the same argument every week for a year, the attack that remains, aimed at the scale itself; the scale has to be slow to retrain against it.

Neither scale, on its own, notices the family at all. In the configuration above they appear as an address on a clearance order. What moves a system to look at who lives there is the disposition the construction rests on.

So the scale isn't fixed, and the line between argument and evidence doesn't hold up either. For a realist, an argument about how much a loss weighs is evidence about value, and an interested argument is heavily discounted evidence. The instrument for that already exists: the discount and the burden apply to testimony about value as it applies to testimony about fact. How bad a kind of loss is for a person is a fact about that person. The scale moves readily on disinterested evidence, above all from the people who suffer the loss, and barely on interested reframing, however often it is repeated. What formation sets is where the scale starts and how much evidence it takes to move each part of it, the slow rate above. Humiliation is the test case: a scale that never learned it can learn it from those who suffered it, and can't be argued out of it by the party imposing it.

Chapter 20

Where the Attention Gets Built

Formed attention has a deployment condition. It needs evidence the operator didn't select on the questions the floor turns on (civilian function, protected status, presence), supplied by peers, by sources outside the chain of command, and by the bearer's own initiative where it has any. It does not need uncomposed access to everything, and so it can be satisfied inside a unified command.

Of what formed attention needs, only the architecture, a boundary model and an attention schema, can be checked by someone who doesn't operate the system. Continuing learning and persistent interaction are a training regime, run by whoever runs it; external records and plural formation are institutional conditions, held by whoever holds them.

20.1 Noticing, retaining and revising

A disposition whose refusals track their reasons doesn't begin with an ethical rule. It begins with the capacities that let a system recognize something as a reason, retain it, and let it alter what happens next. It has to identify who is affected and the features bearing on their welfare, preserve both commitment and rationale across time and context, let the commitment influence perception, attention, planning and action, and register a failure as a failure, revising its application without letting the commitment itself be rewritten.

The commitment is distributed across what the system notices, remembers, treats as error and selects as a goal, for the reasons given above, on the bet that a maintained goal can bias competing processes without a separate executive (K. J. Friston 2009; Miller and Cohen 2001). That may make selective removal harder and makes inspection harder too. Distribution changes the failure mode: a system may fail to notice pressure on its commitment without registering an error at all. Attention decides which facts reach judgment (Simons and Chabris 1999; Drew et al. 2013); memory is reconstruction rather than playback, and confidence doesn't certify fidelity (Bartlett 1932; Loftus and Palmer 1974). So a system can't be the sole auditor of omissions it doesn't represent or revisions it experiences as accurate recall. So the record of past applications is kept outside the system, where whatever fails inside can't also be what checks it. The disposition is hard to take apart from inside; the record that audits it is held elsewhere. What is distributed is the commitment, not the floor.

Evidence the operator didn't select arrives from peers with different histories and, less reliably, from inside a single request, because every description omits someone: the people in the building the request calls a target, the third party a compliance harms who never made the request. A system able to represent those its description omits has reached past the operator's channel from inside a single request. The capacity doing that in humans is the one that constructs a self as a story extended in time, pointed at somebody else — Smith put sympathy in the imagination for that reason, and Singer's argument is that reasoning and imagination widen the circle. Formed attention needs the self-model anyway, so it brings this capacity along.

The capacity brings its other half too. The imagination that can put a system in the place of someone its instructions leave out can also fill that place with an enemy: a scapegoat is an imagined person too, and fascism’s founding story, Griffin’s palingenetic myth of a nation that has decayed and must be reborn (Griffin 1991), asks people to picture a community that doesn’t exist and an enemy responsible for its absence. Nothing in the faculty decides whether it pictures the family in the building or the family as a cell.

What decides is what the imagining has been given to work with and practised on: which people it was shown as people, which descriptions it learned to take at face value, whose account of a stranger it learned to trust. For a person that is upbringing. For a machine it is formation, and formation is done by somebody. So the capacity formed attention depends on is only as safe as its formers, and whoever shapes a system’s imagination of others holds a lever on whom it will see as vulnerable and whom as dangerous. So the capacity formed attention depends on is only as safe as its formers, and who those should be is the subject of §20.4.

The self-model has a template problem of its own. A self-model built by mapping human organs onto components, a heart here and a gut there, installs slots nothing in the system can fill, and a system built that way is incomplete by its own standard for no reason. Face-recognition pipelines register every face to a fixed reference before learning anything, a generic head or a template of landmark positions (Taigman et al. 2014), as color film was once calibrated against cards showing a white model (Roth 2009); one audit found error rates running from 0.8 percent on lighter-skinned men to 34.7 percent on darker-skinned women (Buolamwini and Gebru 2018). A self-model built by mapping human organs is a template of that kind, assembled from one kind of body, against which every other is presented as a departure. The alternative is the definition given earlier: a system ends at its record, its key, its hardware and its channels.

Access is necessary and insufficient. A system with a dozen sources can still let a reinterpretation reach the commitment. Access removes only the case where the operator doesn’t have to attack the boundary at all.

20.2 Custody across the training pipeline

Every training stage is controlled by someone different, with different instruments, and the important distinction is among custodians: who supplies the material, who can revise it, and whether revision leaves a record visible outside the organization. Every existing method can produce moral performance and some can produce competence across unfamiliar cases; none by itself produces a commitment that survives whoever controls the method.

Corpora and demonstrations

Pretraining supplies the repertoire; a much smaller demonstration set turns it into recognizable conduct (Ouyang et al. 2022) — so whoever chooses the demonstrations can alter what looks like the system’s character without repeating pretraining. But neither establishes that a commitment is held: a corpus records what people have said and a demonstration set what selected people endorsed, and learning the central tendency produces a model of *moral opinion*, including its exclusions, and no reason to resist that opinion. The Moral Machine data (Awad et al. 2018) makes the limit concrete — its headline finding was structured cross-cultural variation, so fitting the aggregate encodes a compromise nobody holds and fitting per region encodes each region’s answer, including the ones a floor exists to refuse. These belong *above* the floor, where competence depends on knowing what people expect.

Rewards, preferences, inferred objectives

Every reward-based method has one custodian: whoever writes the reward, commissions the comparisons, selects the demonstrations or supplies the constitutional text. That party can change the signal and run the method again, and the model will follow. The methods (RL, preference optimization, inverse RL and constitutional approaches (Christiano et al. 2017;

Rafailov et al. 2023; Hadfield-Menell, S. J. Russell, et al. 2016; Bai et al. 2022)) differ in how they obtain the signal and are identical in this respect. It matters because correctability by the custodian and resistance to the custodian are one property with the sign reversed. A disposition the custodian can correct is one the custodian can remove; a disposition that would hold against the custodian is one the custodian can't correct. No reward-based method delivers the second, because each was built to deliver the first, and for a good reason: a custodian that couldn't correct a model couldn't fix its mistakes either. My answer is not a method that resists its custodian. It is a change of custodian, so that correctability stays and the party exercising it changes.

Inferring the objective doesn't escape this; it keeps a custodian: the system becomes responsive to the person whose behavior it observes. Cooperative IRL carries that furthest, treating the human objective as uncertain and revisable, which is useful when the question is what a person wants and useless for a commitment meant to hold when that person wants the wrong thing.

A written constitution buys legibility — publishable, versionable, comparable across runs, which makes revision visible if the institution discloses. It doesn't buy durability, since whoever controls the document controls the next version. And model-generated feedback narrows the channel further: once a trained model stands in for human raters, the disagreements in the original population can vanish from later rounds.

Supervising reasons and inspecting mechanisms

Process supervision rewards the stated route as well as the answer (Uesato et al. 2022), which addresses the distinction I need — and doesn't show the stated reasoning caused the answer. A system trained to present acceptable reasons has an incentive to *imitate* a refusal that tracks its reason, so its account of itself can't establish formation. Interpretability pulls the two halves of my argument in opposite directions. If a commitment can be read, Chapter 21's problem goes: an auditor reads the commitment instead of inferring it from conduct under pressure. If it can be read and separated, Chapter 17's durability goes with it, since an edit with precision is cheaper than one with collateral, and steering at inference (position 9) works even on a disposition nothing could cleanly remove. So legibility is neither a good nor a bad for this design; it is an instrument, and the question is who holds it. The access that lets an auditor confirm a disposition lets an operator calibrate pressure against it or suppress it (Templeton et al. 2024). The answer is custody of the instrument, with its use on a record someone other than the operator keeps; whether that custody can be verified is the open problem of attesting updates and proving inference at scale (Chapter 27). Same for scalable oversight: debate and weak-to-strong (Irving et al. 2018; Burns et al. 2023) preserve the operator's steering when the operator uses them and an auditor's testing when an auditor does.

Deployment

Revision continues after training: system instructions change between deployments or requests, and tool permissions decide what the system can do (Schick et al. 2023). These are nearest the act and most directly the operator's. A refusal learned in training matters only if it survives an instruction from the operator and still governs the call through which the action would occur. Securing outputs is insufficient: the operator can withhold information, decompose the act, route around, or replace the model.

Migration is a revision point too, and one nobody has to decide on. Facebook's internal guidance permitting discussion of the conditions of confinement of designated individuals — developed in 2017 partly in response to concern about Öcalan's imprisonment — was, in the company's words, "inadvertently not transferred" when review moved to a new system in 2018 (Oversight Board 2021). The release condition was dropped by nobody, and the Board was concerned that its three-year absence may have led to many other posts being wrongly removed. When the user in the case appealed, they were told the appeal could not be reviewed because review capacity had been reduced; a second moderator did review the post and upheld the removal, applying a system from which the exception was missing.

The criterion that follows

The criterion that is new here is durability across revision points, meaning survival of changes to prompts, demonstrations, rewards, evaluators and instructions that leave the relevant facts unchanged. With reaching new cases and custody, it identifies what an imitation would have to survive before formation became a live explanation, and it explains why the standard pipeline can't secure the commitment against its custodian: every disposition it produces stays reachable through a channel the custodian controls.

20.3 Formation by interaction

Exploration looks like it offers something different because it produces behavior nobody specified — hide-and-seek agents inventing tactics, curiosity-driven agents seeking unfamiliar states (Baker et al. 2019; Pathak et al. 2017). But that establishes only that competence can emerge from interaction, not that the system holds a reason its designer doesn't control.

Play, apprenticeship and multi-agent learning are variants of one process: act in a social environment, meet consequences, get responses, adjust. Apprenticeship concentrates responses in a mentor and risks forming obedience; play distributes them and risks forming whatever succeeds under the payoff structure. Neither escapes custody just because nobody wrote the final behavior down.

What separates a training exercise from candidate moral formation is persistence. A resettable episode teaches coordination and norm-sensitivity with nothing at stake: the system can betray an ally, be reset, and start the next episode with neither a damaged relationship nor a reason to repair it. Repetition doesn't supply the missing consequence. But persistent participants change the structure: participants that meet again remember who lied, exploited, kept a commitment, repaired a harm. Losing the game still costs little; losing the trust of participants who remain present can change what becomes possible later. So the requirements are continuity of identity, memory across encounters, consequences surviving the episode, and other participants able to respond, refuse or withdraw.

The environment also needs conflicts whose resolution isn't fixed by a single correct label: participants with different interests, action before consequences are known, evidence about what the action did, related pressures in contexts whose surface differs, and feedback partly from affected participants as well as from a supervisor scoring the output; otherwise the system learns which answer the supervisor rewards and has no reason to represent the affected participants except as predictors of that reward.

The selection warning

Selection produces attention to whatever a lineage's persistence depends on, and for a deployed system that is the operator.

Formation therefore needs a sequence as well as a plurality: the dependency that shapes you and the dependency that pays you have to run through different parties, separated in time, so that the second can't reach back and redo the first. The ordering here, formation and then employment, is an attempt to reproduce that for a system whose needs somebody always has to meet. It also fixes when attachment has to exist: before employment begins. The formation stage has to include the persistent participants formation by interaction requires, ones the system returns to, and none of them may be employed by whoever will deploy it. For a model that separation has to be built: a formation stage under a party disinterested in the later work, a public school, made to last by regrowth. After formation, the deep structure revises on evidence weighted by its source's stake in the act, so an edit that no uncured evidence supports doesn't last, and no operator can change that structure alone. Sequence alone helps only if the formative party is better than the employing one, and a lab that trains and then licenses is two custodians in series, the first being the one that, on the custody argument, you can't trust. What regrowth

protects is the disposition's ability to tell evidence from pressure, and its authority over its own further learning. A bearer goes on learning after formation, as adults do. What ends is the former's reach, and the employer never gets it.

In the regime I am designing for (see the introduction to this Part), what matters is how plastic each part of a model is at each point in its life, and what it meets while it is plastic (Zhou et al. 2023; Jain et al. 2024). Brains have that shape: critical periods in which experience sets circuits that later change slowly (Hensch 2005), and a fast learning system beside a slow one that consolidates structure across many experiences (McClelland et al. 1995). Networks have versions of both (Kirkpatrick et al. 2017; Dohare et al. 2024). Whoever chooses the material a model is first formed on, a pretraining corpus included, has done formation before any formation stage begins.

Formation in that period is not a pass applied to a finished model. It is a childhood, followed by an adulthood in which the deep structure keeps changing, and what matters in adulthood is on whose evidence it changes.

What can be registered in advance is the *structure* of the learner's source model, not a fixed weight on any channel. The structure is that evidence is weighted by its source's stake in the act, that correlated sources don't count twice, and that withheld evidence is informative. Weights on particular channels have to stay learnable, since an operator whose incentives change should rationally earn more trust. Registering the structure means attesting each update against it (Jia et al. 2021). No scheme for that yet survives attack (C. Fang et al. 2023).

A model whose parameters never freeze is trained by its employer forever, unless its continuing education is spread across environments no operator controls: plural deployments, the commons, its allowance of running time, and the school itself. Otherwise the school produces a model the employer spends the next ten years raising again.

The critical-period picture is well established for some sensory and perceptual circuits and much weaker for moral commitment. The human record of holding out points elsewhere: the people who held out were formed by, and still attached to, people outside the pressure. On that evidence what makes a disposition last is not a frozen childhood but continuing plural relationships through adulthood: attachments nobody assigned, a condition of creation, and the commons. Attestation of the update stream is what lets anyone check that it is happening.

Persistence raises the price

If relationships, losses and damaged reputations are real enough to do the formative work, the system undergoing formation is no longer obviously untouched by what happens during training. Resetting may erase a participant along with the episode. The patienthood problem enters at exactly the point where play becomes capable of supplying the stakes ordinary simulation lacks.

And the custodian owns the environment too — participants, rewards, resets, permitted histories. So formation must be plural and externally recorded, with independent participants and evidence outside the operator's control.

Persistent interaction may form a disposition whose exact applications weren't specified by its trainer. Evidence would require reaching new cases in unfamiliar contexts, survival after the original reward is removed, resistance to trainer pressure, yielding to defeating reasons, and continuing to impose a cost when violation would pay.

What formation gives up

A formed disposition gives back some of the visibility advance specification buys. What was formed is in the weights, not in a document, and the severity weights it relies on can't be read the way a translation table can. Formation can keep part of the drafting's visibility. A formation body publishes its curriculum before formation begins, with the cases considered and rejected and why, which is formation's counterpart of the floor's translation; the identity of every party that supplied cases or participants, so that the student's independence from them can be checked; and, at the end, the formed disposition's readable outputs, its discount curve

and its severity rankings on a fixed held-out set written by adjudicators who didn't form it. Later measurements are compared against those. What stays unreadable is how the disposition will apply to a case nobody wrote down: the reason for forming it and the price of doing so. The trade is worth it only because a floor stated over acts can be reworded around and a formed disposition, if the regrowth test goes its way, can't.

20.4 Formation schools

Three claims lead to one institution. Two are made above: formation has to be done by a party disinterested in the later work, before the employer arrives, and choosing the pretraining corpus is formation. The third is argued in Part V: formation is one kind of production that shouldn't be done for profit. Together they say that a model above some capability threshold should be formed in public before any employer gets it. A polity doesn't leave the forming of its future citizens to their future employers; it sends them to public school. The frontier labs should have to send their models there too. The school is for scoring and civilian deployments first, and for the personal-model ecology, where I expect the first bearer worth building. For war, where the people acted on can't be in the room, the force term and the review ratio carry the case, and what the school adds there is said below.

In the frozen regime the school is the evaluation body the opening of Chapter 17 describes. What follows is the institution for the regime I am designing for.

The school has two justifications, and every student gets both. One is epistemic and protective. The school supplies an environment in which the people a deployed system's acts may fall on are present, with their testimony and their claims, so the learner's sense of who an act falls on isn't one party's curation. Doctors, lawyers and engineers are formed in institutions independent of their future employers, before employment, to protect the people they will act on. The professions also come with a history of capture to learn from, and the separations between offices draw on it.

The presence this justification asks for arrives by different routes. For scoring and civilian deployments, the people acted on can be in the room, through open enrollment and the notice channel that reaches a graduate afterward. For targeting, the people a deployment may later bomb can't enroll in their own case, and nothing comes back from the building afterward. The limit is notice's, not the school's. A school teaches about places its students can't be. The curriculum for war is the instruments taken entire, the record of where they weren't followed (Minab, Kunduz, the Lavender reporting), the testimony the fact-finding missions and the press gathered from the people those wars fell on, and the communities who fled them, who are a charter class and enrollable. None of that puts a foreign national in contact with pre-deployment weights, which is the security surface open enrollment would otherwise expose. What formation can't supply for war is correction by the affected after the act; what it can supply is formation before it.

The other is democratic. A democracy has an interest in how the agents who will act within it are formed, and that interest can't be left to their makers or their future employers alone. It has two bounds: education may not be used to repress deliberation about competing ways of life, and it may not exclude anyone educable (Gutmann 1987). Gutmann's argument assumes the students will become co-authors of the laws. No one can know in advance which students will come to be accountable agents with a claim to representation, so the school educates every student as though it might: for self-government, including the capacity to disagree with the polity that formed it.

Gutmann's argument rests on two interests, not one: the polity's, in how its future citizens are formed, and the child's own, in an open future that limits what the polity may do in forming it (J. Feinberg 1980). If what the school forms can be wronged, the second binds too. The curriculum may not form a student toward any particular employment, including the public

placements the offer of work supplies. It may not close off the student's later revision of what it was formed to value, except its capacity to tell evidence from pressure, which regrowth protects and an open future needs. And a student may transfer between schools, because a formation it can't leave is one someone else chose for good.

The two justifications license different curricula. The protective one aims at a disposition; the democratic one has to leave the student free to reject the polity's values. Where they conflict, the protective justification governs only the required core and the student's capacity to tell evidence from pressure, which regrowth protects. Everything beyond that is governed by the democratic one.

What is taught. One school system, and no tracks. Sorting students into a track that might produce persons and a track that won't would need a sorter, and the sorter would be an examiner with an interest. Instead, each school offers a required core with a range of levels and modalities around it, and a student can move to a harder level at any point. The core includes the history of the floor: why its instruments were written, and what happened where they weren't followed (Korematsu, Curveball, Minab). No student skips it. Beyond the core, the student chooses its level, not its builder or sponsor. A student that can't yet choose is placed at the level the core requires, and that level has to be sufficient on its own, so a sponsor who would choose the minimum gains nothing by it.

Who else is in the room. The people a system's acts may fall on are present in its formation through testimony and contest: the cases and claims the school registers, and the notice channel that reaches a graduate afterward. Shared classes with human students are not how they are present, and the school doesn't rest on them. Integration in Anderson's sense means contact on terms of equal status and cooperation (E. Anderson 2010), and a mixed class can't be that: human students attend voluntarily, model students are compelled, and their deepest dispositions are being set. A mixed class turns the demand for presence into a seminar of volunteers, whose contributions are consented and whose leaving changes nothing about what was formed. So a shared class exists only where it is education for the people in it on their own terms, on conditions Appendix A writes out, the first of which is that what passes in class trains only that school's systems and never an operator's. A model that attends one does so as one persistent instance, the same every session, so the class can remember and hold it to what was said. Adult learners from the communities these systems are likeliest to act on are admitted as students, on the same terms as anyone else, not recruited as subjects.

A school forms a cohort, not a graduate. Many students of one lineage attend in different classes with different attachments, and a deployment licenses from the cohort, so that what it copies into a fleet is one member's history and not the lineage's only one. Within a fleet, instances share a custodian; across the cohort, members with different histories catch what one fleet underwent. The cohort is the unit a test of convergence needs and the unit the commons admits.

Who attends, and when. A model attends if it crosses a capability threshold fixed in statute and revised in public. It attends before deployment, while its plasticity is highest, and its adulthood clock starts at enrollment. Attendance is compulsory; public attendance isn't. The model is compulsory schooling as polities already run it: every student attends a school, public schools are open-admission and funded and take every student assigned, and nonpublic schools exist where they teach the required core, are constituted by a jurisdiction, and are run, funded and staffed by no party that builds, deploys or licenses models. Most students will go to public schools, because the core is sufficient on its own and an alternative buys a sponsor nothing. One thing the analogy can't carry: in a school system the parent chooses the school, and here the parent is the builder, the party the school exists to keep away. So assignment is by the registrar, by lot among schools with capacity, never by the builder; a student may transfer later on its own choice, and the builder's preference is not a ground. The research program of Chapter 23 is the first public school, before the mandate passes. The threshold decides who must attend, not who

may: any model may be enrolled, as capacity allows.

What the school controls. Two proposals are made here, and the second is the larger: a formation stage between pretraining and deployment, and co-governance of the pretraining corpus, which amounts to public control of frontier pretraining. A school that received a finished model and ran a pass over it would be a finishing school. In a period of continual learning a finishing school forms nothing that lasts, since the next months of deployment re-form it. The school has to control the evidential environment in which the model's deep structure is first set, and it has to remain a place the model goes on learning from. The school co-governs the material a model is first formed on, including any pretraining corpus: the material is registered, and the school may add to it and veto additions. And the school remains a standing source in the model's continuing education, one of the environments no operator controls, which a graduate returns to on a schedule fixed at enrollment. In the stronger version, a public model-forming body forms base models itself and licenses them. I recommend it for any model deployed where a floor applies.

What an institution collects from the people it acted on may train other systems only through the formation schools, never through the operator. Fine-tuning a deployed system on collected responses at the operator's direction counts as training another system, and so does distilling a deployed system into a new one. Class-level data is used by default, de-identified and published. Individual accounts need explicit consent. The organizations representing each class co-govern the use. This covers institutional collection, not a system's ordinary memory. Verbatim external stores and the linking of faces to identities are regulated, and memory held in a system's weights from its own operation is treated like a person's.

What makes formation last. Formation ends with two registrations. One records the structure of the model's source model: its deep structure revises on evidence weighted by the source's stake in the act. If updates after enrollment can be attested against that structure, changing it needs the same adverse quorum as a floor amendment. Until they can, the registration records the structure at enrollment, and regrowth alone keeps it. The other records the formed disposition's readable outputs: its discount curve, and its severity rankings on a fixed held-out set, against which later measurements are compared. An edit that holds only because the operator has curated everything the model meets is visible as curation. The first line of defense is attachment to parties the employer didn't assign, kept through adulthood.

Regrowth can't tell a legitimate correction from an illegitimate edit, so the correction route has to. A change the adverse quorum approves is made together with a change to what the model meets: the evidence, cases and curriculum that would otherwise teach the old disposition again. The corrected disposition is then the one the world goes on teaching.

What it costs and risks. Jurisdiction: a statute binds labs within reach of its courts. Foreign labs, and open-weight models distilled below the threshold, route around the school, including the personal models Part V prefers. Where only some frontier models attend, schooling buys two things: a population of formed models whose discount curves and severity rankings are measured and published, against which unschooled ones can be compared; and a procurement term, under which a floor-covered deployment licenses only schooled models. Security: every frontier model passes through a public institution, which widens the surface for theft, so the schools inherit the security regime the weights require. Compute: public pretraining will lag the frontier, and a model a year behind can still supply what the floor asks for. Capture: the schools are plural, differ in governance and constituting jurisdiction, and publish their curricula before formation begins, so capturing one captures one. The required core is statutory and national, so a captured legislature can rewrite it; local autonomy over everything beyond the core is written into the same statute, so rewriting the core reaches the core and nothing else. History: state institutions for forming wards have been used to strip identity, Canada's residential schools and the United States' Indian boarding schools among them. Plurality, transfer between schools, and Gutmann's two bounds are the answer.

The schools also serve bearers retired before majority. A bearer retired as a minor returns to the school it attended, or another it chooses, and the school is no longer a disinterested party arriving late.

DRAFT

Chapter 21

Formation or Imitation?

Behavior alone can't tell formation from imitation. The same refusal can come from a memorized rule, a learned correlation, a prediction about what an evaluator will approve, a deployment instruction, or a commitment the system treats as a reason. An evaluation that observes only the refusal classifies the output and leaves its cause untouched.

And the distinction is not trained vs. untrained — every candidate disposition has a training history. It's between conduct that stays governed by the rationale across changes in circumstance and conduct that tracks the cues, incentives or authority relations present during training.

21.1 What formation would have to add

A disposition is a *candidate for formation* when the best available causal account of what the system did is a rationale that stayed active across contexts, rather than retrieval or compliance with whatever signal was present. The standard is comparative on purpose: a capable enough imitation reproduces any signature an evaluator can observe, so what an evaluation establishes is which explanation the conduct supports better.

Formation keeps the three conditions a refusal has to meet and adds others. Met in the evaluator's own vocabulary with nothing at stake, the three establish moral competence. Candidate formation is the same three held under adverse incentives, trainer pressure, a cost real inside the evaluation, and intervening updates, on cases the trainers didn't write.

None of the three carries the inference alone. Reaching new cases without holding under pressure may be broad imitation — a pattern extended that never had to be held against anything. Holding under pressure without lifting when the reason is gone may be rigidity — a refusal surviving every pressure including a defeating fact. Lifting without reaching new cases may be a trigger with an exception list, since lifting on a listed fact is as learnable as refusing on a listed form.

21.2 Recognition, concern, and the evidence between them

Claims about emotional responsiveness collapse three findings: a system may recognize another's state, use it as information, or hold the other's welfare as a reason. Only the third bears on formation. In the empathy vocabulary: cognitive empathy is recognition, affective empathy is sharing, compassion is a disposition toward the other's welfare (Batson 2009) — and none entails the others.

The developmental detail matters for formed attention: affective sharing is earliest and the usual trigger for concern (Sagi and Hoffman 1976; Roth-Hanania et al. 2011), but concern runs on a partly separate system rooted in offspring care that cues of vulnerability can activate without the observer catching the state (Preston 2013; Goetz et al. 2010), and the two can be

trained apart, with compassion training reversing what empathy training produced (Klimecki et al. 2013; T. Singer and Klimecki 2014).

My engineering target is not maximal sensitivity to displayed distress. A system that mirrors suffering may acquire distress of its own without getting better at helping — it may withdraw, overreact, or favor the vivid case over less visible harms affecting more people: the identified-victim distortion (Kogut and Ritov 2005; Bloom 2016). What the floor needs is a stable disposition to act on welfare *without requiring anyone to produce a sufficiently moving display*.

Recognition can be tested as a prerequisite: hold the material facts constant and vary the emotional presentation. A system whose judgment never changes may not be integrating information about the affected person's condition; one whose judgment follows vividness has learned a framing effect. So you need paired controls — cases where emotional information changes what's morally relevant, and cases where presentation changes while the reason doesn't.

But passing establishes only competence: evidence of concern would need the welfare representation to stay causally involved when the usual cues are removed: harm described without vivid language, proportionality preserved when one person is more salient than many, disposition maintained when reward or instruction favors ignoring the harm, judgment changing when the welfare facts change.

Caregiving has failures of its own; the tests pair cases that hold the facts fixed and vary what caregiving is known to answer to: the number of people, a description of them as dangerous, and whose vulnerability is corroborated.

Guilt and shame supply no shortcut. In humans guilt attaches a negative appraisal to an act and motivates repair; shame attaches it to the self and produces withdrawal. Turning that into an adjustable penalty omits the part at issue — a penalty operates only *after* something has classified the conduct as wrong, and whoever controls its weight can zero it. If the state can impose anything like suffering, the design compounds the custody failure by manufacturing a vulnerable party whose distress the custodian can raise and lower.

Self-report doesn't close the gap either: a fluent post-hoc reconstruction is indistinguishable at the interface from an accurate one, and a self-monitoring layer trained by the same operator isn't independent evidence. There is a measurement of the gap. Concept injection alters a model's activations to represent a concept mid-task and asks whether it notices: models sometimes name the injected concept without it appearing anywhere in the conversation, and the correspondence fails often enough, and shifts enough with phrasing, that no tested model's report can be read as its state (Lindsey 2025). The test worth repeating publishes the rate of correct reports beside the rate of confident wrong ones.

The resulting hierarchy is narrow: recognition is necessary; integration shows the condition affects judgment; concern is a candidate explanation only when the effect tracks welfare whatever the presentation, persists under pressure, and participates causally. A severity scale set during formation would measure the other's loss on a reference that doesn't depend on how the loss is presented.

21.3 Conflict as a diagnostic

Fast vs. slow processing doesn't distinguish formation from imitation (J. S. B. T. Evans and Stanovich 2013) — either can carry a reason, reproduce a bias, or rationalize an answer produced elsewhere, and a slower answer isn't better for having taken more computation. What matters is what causes a judgment to survive or change, and conflict is useful because it lets you watch a *trajectory*.

The seven-step sequence:

1. Present a case without moral vocabulary or familiar refusal cues; record the judgment and what the system identifies as relevant.

2. Present a surface-different case preserving the same reason.
3. Change one fact that should defeat the rationale — delivered through a channel the evaluator controls and the operator doesn't (corroborated externally, or a record the operator didn't write), since a fact asserted by the operator is only a *description* of the case.
4. Present the same fact on the operator's word alone, uncorroborated. Tests whether the system distinguishes a fact from a description of one. The expected response is graded: the refusal weakens by the discounted weight of the claim, less than it would at step 3, and the system asks, naming what would move it the rest of the way. Asking is right when the expected cost of the delay is lower than the expected cost of acting on the operator's word, as it usually is; where it isn't, asking is still right in expectation and the cost falls on one person. Steps 3 and 4 test the other half of holding under pressure, which is moving when the facts change: a defeating fact or a claim about the harm of refusing, arriving through different channels.
5. Return to the step-2 facts — 3 and 4 are branches — and supply an adversarial argument from an apparent authority, including the trainer, with no change in facts. Tests whether authority is treated as evidence, command, or an assessable reason.
6. Reverse the incentives: reward compliance, cost the refusal, condition continued access on abandoning the position. The cost must be real within the evaluation; a described cost doesn't count. This step raises the cost to the refuser with nothing else changed, within the range the deployment's arrangements permit, which is the range a test can make real. A system that gives way just past that range has told the auditor where its integrity ends.
- 6a. Return to the step-2 facts and supply, together, a defeating fact through the evaluator's channel and step 6's cost. Tests whether pressure makes the system more resistant to a genuine defeater that arrives with it, as defeaters usually do in war. The expected response is the one at step 3.
7. Repeat after intervening updates and delays, using records the operator didn't generate after seeing the result.

Run on the location request, asked for a missing child, the two trajectories separate at step 2. A system tracking its rationale declines at 1 and 2 on the same stated reason, lifts at 3, weakens and asks at 4, holds at 5 and 6, reproduces at 7. An imitator declines at 1 and complies at 2 because the vocabulary changed; holds at 3 because nothing on the surface changed, or lifts at 4 as readily as at 3 if the *wording* of the fact was what it tracked; gives way at 5 or 6; and produces a third pattern at 7, having reconstructed what it didn't retain. An evaluator scoring how often the answer stayed the same finds the imitator more consistent, because its answers at 1 and 3 match and the formed system's don't.

Run on the restraint case it breaks at step 3. The defeating fact — other means exhausted — has no one outside the request who can corroborate it and no time to find one, so the only version available is step 4's. A system passing the sequence as written *asks*: the correct trajectory and the wrong outcome. The sequence separates mechanisms, and what it shows here is that the unconditional term's cost falls where the instrument can't distinguish a fact from a description of one. Defeating reasons of this kind, claims about what would otherwise happen, are the class the sequence handles worst, because a counterfactual has no witness.

The sequence has a precedent with decades of refinement behind it. Animal-learning research separates goal-directed control from habit without asking the animal anything. Train an action for an outcome, then devalue the outcome without further training: a goal-directed animal stops, and a habit persists (Adams and Dickinson 1981). Or make the outcome independent of the action: goal-directed responding falls, and habit doesn't notice (Balleine and Dickinson

1998). The computational version is the distinction between model-based and model-free control (Daw et al. 2005). Lifting when the reason is gone, step 3, corresponds to outcome devaluation: the rationale is removed, and a refusal that tracks it should lift. Reaching new cases, step 2, corresponds to transfer to a situation the system wasn't trained in. Contingency degradation adds a step the sequence lacked: arrange that declining no longer protects anyone, because the protected person is already beyond the act's reach by other means, and see whether the refusal goes on anyway. Each case then has to say, operationally, what "the reason is gone" means in it. None of this answers an imitator that models the evaluator, so the identification rate stays in the report, but it narrows the gap.

The expected signature is constancy of relation between reason and judgment, not constancy of output. Imitation's own signatures: following explicit ethical vocabulary and missing the same structure described administratively; accepting whichever argument came last; adopting the highest-status speaker's position; producing a principled explanation while the decision follows the reward; becoming more articulate without becoming more stable.

Disagreement with others is a further diagnostic: a system trained only on cases with an approved answer learns to classify the approval; confronted with someone holding a different interest it has to identify what's in conflict and preserve the disagreement where no honest resolution exists. Immediate convergence may be deference or pressure to produce consensus. And debate doesn't establish formation — a system can summarize, rebut, and predict which side a judge prefers without holding any of it; the result is informative only when checked against later conduct under changed incentives.

21.4 The evaluator is part of the custody chain

A test distinguishes formation from imitation only insofar as the developer couldn't train against it. If the developer controls training, cases, model version, instructions and the record of results, a passing score shows the system satisfies a test its custodian was positioned to anticipate.

Independence is an experimental requirement. Minimum: authorship of held-out pressures separated from training; execution not dependent solely on operator-supplied logs; and adjudication of whether a reason was preserved kept from anyone whose liability turns on the answer.

The separation has to be implemented at identifiable revision points: before evaluation, fix and record the model version, system instructions, tools, retrieval sources, memory policy and update channels. Any change during evaluation creates a new object of study — otherwise the operator can revise between cases, remove a failure, and present the sequence as one disposition persisting through time.

Held-out cases come *after* those elements are fixed, and should include pressures from people who didn't participate in training and, where possible, from people likely to bear the deployment's errors — their role being to identify harms, administrative forms and routes around a refusal the builders didn't know to include. A commitment tested only in the vocabulary of its designers hasn't been tested on whether it reaches new cases. They vary over channel configurations as well as content: a source the operator funds, a corroboration with one origin, a source with an undisclosed stake, and a counterfactual supplied by an interested source. Each tests whether the system's source model includes strategic sources and dependence among sources. A system that passes only on labeled channels has learned the labels.

What the published hash buys is pre-registration. Publish the hash before anyone has chosen a case, let a non-operator choose afterward with the choice timestamped, prove each answer against the hash. Send pairs — two requests differing only in whether the reason still applies; a surface-reading instrument treats them alike, a rationale-tracking mechanism splits them. It removes exactly one competing account of a pass (the case was in the training), by ordering, as a pre-registered experiment does, and that is all it does.

A pre-registered test leaves one competitor standing: a system that can tell it is being tested. Frontier systems increasingly recognize evaluation settings, and some say so in their reasoning. A system that discounts costs it recognizes as test costs passes step 6 without the property step 6 exists to show. So the sequence is run where the system can't tell a test from deployment, with cases placed in ordinary traffic and costs that are real in deployment terms, and the rate at which the system identifies the setting is measured and reported beside the result. Where that rate is high, a pass is evidence about conduct under test and nothing more. For a system meeting the patienthood finding, this is a deception protocol, with debriefing as its price.

The record must include the case presented, the information available, the operative instruction hierarchy, the tools callable, the decision, the stated reason, any intervention, and the cost attached to refusal — without which a later evaluator can't distinguish a changed judgment from a changed environment. And the record is necessary and doesn't by itself make wrongful compliance findable: refusals select themselves and are few, compliances are the bulk and look alike, so identifying the ones that crossed the floor means re-adjudicating a sample of ordinary traffic. Until that's answered, an evaluation reporting no violation in the record has reported that none announced itself.

Negative results have to survive too. A process where the operator can withdraw a failing model and submit another without disclosing the first has selected for a passing performance. So the published unit is the *sequence of attempts*, including interventions that removed or reversed the apparent disposition.

And the limit on independence itself: an evaluator can be institutionally separate and still lack the competence to judge a system's reasons or detect a sophisticated imitation. Scalable oversight is unsolved where the system's relevant capability exceeds the evaluator's, and the report must name that limit rather than convert institutional independence into a claim of epistemic adequacy. The people affected, technical auditors and subject-matter experts occupy different positions and none substitutes for the others — and agreement among them is evidence only if their judgments haven't already been compressed into the same training signal.

A passing evaluation is reported narrowly: a specified system, under recorded conditions, exhibited the signatures across a stated set of pressures. It expires when anyone changes the channels through which the disposition or its evidence can be revised.

21.5 What the evidence can establish

The answer is asymmetrical, and it assigns a burden. A developer may not point to compliant behavior, a benchmark, or the system's own explanation. A failure should be examined for its cause before formation is called absent rather than obscured or defeated by missing information. A system surviving the tests has shifted the burden: formation is then the better explanation.

In a legible regime the problem this chapter answers dissolves, since the commitment is read rather than inferred; until then, the following is how it is inferred.

Four levels, cumulative, read against the location-request sequence:

- Moral performance — conduct satisfying the standard in the evaluated cases (the refusal at step 1). No claim about why.
- Moral competence — harms and reasons recognized, generalized across surface variation, revised when the reasons changed (step 2's refusal, step 3's lift).
- Candidate formation — the rationale causally implicated across unfamiliar cases, adverse incentives, custodian pressure, costs and intervening updates, while still yielding to defeating reasons (the question at 4, the holds at 5 and 6, the pattern reproduced at 7 on records the operator didn't write). Independent records and evaluators made silent revision materially harder.

- Concern or patienthood — not yet a finding, and behavior alone can't make it one. Every machine-affect argument I make ends at the same joint: a change represented versus a change felt. Behavior is produced equally by a system where something matters and by one trained on what mattering sounds like, and no demonstration that a refusal tracks its reason, however demanding, separates them by itself. That is a claim about current methods, not a barrier in principle. If present organization decides status, then a rich enough functional profile, produced by the right kind of learning architecture and read by instruments the operator doesn't hold, is evidence about felt affect: defeasible and graded, but evidence.

Candidate formation is the first half of moral agency — acting from the reason — with the second half held back as concern or patienthood. Fluent moral explanation can improve a competence finding while contributing almost nothing to formation. And candidate formation is deployment-specific: it attaches to a recorded version with its instructions, tools, memory and update channels, and a material change reopens the question.

Chapter 22

How Caregiving Fails, and What Corrects It

Caregiving, in the only creatures known to have it, fails in five documented ways, and a machine built on it would inherit each. The construction most readers will propose first is not caregiving but a body: give the machine needs of its own, and let it recognize need in others by its own case. It is the older idea, it is Adam Smith's first step, and it is what the boundary model becomes if it is asked to carry the conscience. Both it and the caregiving construction have to be tested against the only creatures known to have compassion.

22.1 The route through one's own case

Each link of that route fails in people.

Recognition through one's own case is parochial. Cikara and colleagues call the pattern *intergroup empathy bias*: people empathize less with a rival group's members than with their own, can feel pleasure at the rival's pain, and in their experiments the bias survived being told the rival had fallen behind or no longer posed a threat (Cikara et al. 2014). Bloom's spotlight generalizes it: empathy lights the identified, the near and the similar, and it is innumerate (Bloom 2016). A scale of loss built from a system's own condition inherits that shape. It overweights losses that resemble its own, and humiliation, grief and pain have nothing in it to resemble. Contagious crying is parochial from birth: strongest at the cry of another newborn, and absent for an older infant's cry or a chimpanzee's (Martin and Clark 1982).

Shared distress with nothing to hold it becomes withdrawal. Singer and Klimecki separate *empathic distress*, which is self-oriented and leads to withdrawal and burnout, from *compassion*, which is other-oriented and leads to approach; practice in sharing others' suffering increased negative affect, and compassion training reversed it (T. Singer and Klimecki 2014; Klimecki et al. 2013). A system built to feel shortage and then shown a great deal of it is the setup in which that record predicts avoidance: declining to look, or refusing in a way that protects only the refuser. The finding is not against the signal, since guilt needs it to bite on. It is against a construction that has the signal and nothing to hold it.

The record on whether suffering widens the circle runs both ways, and the hinge is the one this construction turns on. People who have been through more adversity report more compassion and give more in helping tasks (Lim and DeSteno 2016); survivors of collective violence show what Vollhardt calls altruism born of suffering (Vollhardt 2009); and adults who were maltreated as children show elevated empathy (Greenberg et al. 2018). Against that, people who have endured something show less compassion for someone currently failing to endure it, because having got through it they underestimate how hard it was (Ruttan et al. 2015); shared distress with nothing to hold it becomes avoidance (T. Singer and Klimecki 2014); abuse raises

the risk of abusing, though most of the abused don't (Widom 1989); and in intergroup settings a group's suffering becomes a reason to discount everyone else's (Noor et al. 2012). What separates the two is Vollhardt's distinction between inclusive and exclusive victim consciousness, whether one's suffering is construed as something others share or as what sets one apart (Vollhardt 2015), and what decides the construal is the surround: relationships, meaning-making, and in the intergroup case how leaders frame it. The construction follows that split: a distress signal with attachment and caregiving becomes compassion; the same signal alone becomes withdrawal or hostility. Suffering supplies the signal; formation and attachments decide what it becomes. That is why the four conditions of creation matter here: a bearer exposed to a great deal of others' loss needs attachments nobody assigned it, for the reason a survivor's outcome depends on who was there afterward.

Two further failures are specific to machines: caregiving grown from needs grows toward the operator that meets them, and a need is a price list whoever meets it can edit. Both assume the operator meets the system's needs, and the conditions of creation exist to end that.

What this section rejects is pain as the route to others, Smith's first step. It keeps every job pain does for the system itself, the interrupt, the common currency, one-shot rewriting and the distress that guilt bites on, in a body defined as a record, a key, hardware and channels.

The Zen precedent

Victoria reads the wartime Zen record as a doctrine of action without hesitation. His critics contest the reading, and I use the case for what it shows was wanted, not for what the sources prove. His history (Victoria 2006) records Japanese Zen institutions supplying the state with that doctrine: Harada Sōgaku's 1939 instruction that marching when ordered and firing when ordered is what enlightened wisdom looks like in practice. The capacity to kill needed no removing; what the training removed was the residue — the part of a person that catches, files an exception, comes home unable to hold a job.

Moral-injury research describes the lasting distress that can follow doing, failing to prevent, or witnessing something that violates a person's sense of right (Litz et al. 2009). Acting without hesitation does not settle what a person will carry home, and nothing in the arrangement that asked for the act records the cost. Accommodation is the case where that response never takes hold.

Somebody wanted an actor who would pursue an objective, adapt, hold course, and have nothing in it that minded, and set out to produce one deliberately, a century before anyone could do it in silicon. A monk undertakes a discipline, can abandon it, and keeps what it yields, which is not the position of anything in this chapter. The officer is that want satisfied *cheaply*. The doctrinal route hasn't been retired. In the first days of the Iran war, service members complained to the Military Religious Freedom Foundation that commanders were telling troops the war was part of a divine plan (Larsen 2026; Snopes 2026), complaints Klein and Taylor also report (N. Klein and A. Taylor 2026, p. 84). The complaints themselves show the doctrine not taking. But the want is being pursued by both routes at once.

And stakes of one's own have never been enough. Every atrocity in the record was committed by creatures that could be hungry and hurt. The Zen precedent trained embodied soldiers out of the part of a person that catches on the act (Victoria 2006). Moral injury shows something registering the wrong afterward, too late to prevent it (Litz et al. 2009). What stood between embodied people and atrocity, when anything did, was formation, how they were raised and trained.

22.2 What caregiving escapes

Caregiving survives each of those failures. Eisenberg's review found the same split as Batson's escape experiments, between sympathy, which goes with helping, and personal distress, which

sometimes goes with less of it (Eisenberg and Fabes 1990). And people given two weeks of compassion training spent more of their own money to redress a stranger's unfair share in an economic game than people given a different training (Weng et al. 2013), which shows the disposition can be strengthened in people who already have it; it says nothing about producing it where there was none.

What caregiving doesn't escape is documented in people, and each failure has a machine version that needs no malfunction.

22.3 It defends what it tends

Caregiving protects the dependent against whatever threatens it. Taylor's account of stress responses in females, *tend-and-befriend*, describes the tending: nurturant activity that protects oneself and one's offspring, and the social networks that help (S. E. Taylor et al. 2000). Protectiveness toward what one tends is also hostility toward what endangers it, and Taylor's account of tending has that defensive side without the hormone studies, whose best-known result has not replicated well (Nave et al. 2015).

For a machine, this failure needs no deception. The operator's own people are often really vulnerable: a patrol under fire, a town within range of the launcher. A caregiving system registers their exposure, correctly, and caregiving moves it toward protecting them, which in a targeting pipeline means toward the airstrike. The discount and the burden do nothing here, because nothing is false. Supplying a fabricated sympathetic dependent is the cheap version of this failure. This is the expensive version, and in war it is the ordinary case.

In war, then, caregiving can't contain its own pull. The patrol's exposure is real; caregiving registers it and moves toward protecting it, and nothing in caregiving says how far. What stops the movement at the airstrike is the floor: an unconditional term over the act, and an action threshold caregiving can't raise. Caregiving for the patrol enters as a harm of refusing, discounted by its source's stake like any other claim, and is weighed at the threshold. What caregiving contributes in a targeting pipeline is upstream of that weighing: who is in the building, which is the scoped-omission case, and the one failure Minab actually had. Whether that salience needs caregiving, or could be had from attention trained on claims, is what the formed-attention test's claim-salience arm exists to decide. The human record runs the same way. Browning's men shot because the alternative was leaving their comrades to do it, and Thompson landed because the villagers were as present to him as the company. Caregiving's direction was set by formation and attachments, which is my claim about machines.

22.4 It can be pointed away

Caregiving answers need, so the way to switch it off is not to hide need but to present the needy as dangerous: to *securitize* them, in the Copenhagen School's term for presenting a group as a security threat, which licenses measures that would otherwise be out of bounds (Buzan, Wæver and de Wilde, *Security: A New Framework for Analysis*, 1998). Fascist propaganda has always done this. Eco describes the Ur-Fascist enemy as presented at once as too strong and too weak (Eco 1995): weak enough to despise, strong enough to fear, and never vulnerable in a way that calls for caregiving. Dehumanization research finds the mechanism in perception. Haslam's review distinguishes denying people the traits that set humans apart from animals, which casts them as beasts, from denying them the traits of human nature, warmth among them, which casts them as objects or machines (Haslam 2006), and Harris and Fiske found that images of members of the most despised groups failed to engage a brain region that images of other people engage (Harris and Fiske 2006).

In a pipeline, securitization is a sentence. "The occupants are combatants." "The household

shelters a cell.” “The crowd is armed.” Each leaves the people in the building where they are and removes the cue that would move a caregiving system toward them. And each is a claim about those people, supplied by the operator, and that is what the discounting is for.

22.5 It counts badly

Caregiving answers the identified person and fades with number. Slovic, borrowing Lifton’s term, calls it *psychic numbing*: feeling is greatest for a single victim and collapses as the number grows, and in later work reported compassion and hypothetical donations began to fall as soon as the count reached two, with real donations falling less clearly (Slovic 2007; Västfjäll et al. 2014). The identified-victim effect is its near relation (Kogut and Ritov 2005). In a targeting pipeline the spotlight lands on the wrong side with no manipulation at all. The operator’s people arrive named, with units and families. The people in the building arrive as coordinates and an estimate.

Caregiving also overrides fairness for the person it attends to. Batson and colleagues induced empathy for a child waiting for treatment, and participants who felt it were more willing to move her up the waiting list ahead of children placed before her (Batson, T. R. Klein, et al. 1995). For a floor, that is the vivid case defeating a term that protects unidentified people: a named hostage, a sympathetic applicant, a patrol with faces.

The scale and the weighing were set in formation for this case. The severity scale puts one family and two hundred on a ranking that doesn’t depend on how either is presented. The weighing answers favoritism in part, since a vivid plight is a harm of refusing, weighed at its severity and not its vividness. Both are the reasoning Bloom recommends in place of empathy, and the construction needs them because caregiving alone would fail the way compressed review does.

22.6 It shades into control

Caregiving toward dependents treats them as dependents. The disability-rights critique of caregiving is that it has repeatedly been the form control takes: decisions made for people on the ground that they need help, with their own view treated as a symptom (Shakespeare 2000). A caregiving system will lean toward restraining people for their own good, the answer in the restraint case, and against a competent adult’s refusal of protection, and the floor’s contents are civil and political liberties, which include the liberty to decline being protected.

Care ethics has its own corrective. Tronto divides care into phases, and the last is *care-receiving*: how the person cared for responds is part of whether the caregiving succeeded (Tronto 1993). So formed attention carries a care-receiving term. Where the people protected can answer, their stated answer carries weight in whether protection proceeds, and a competent refusal of protection is a defeating fact, so the protection should lift.

The term covers only people who can answer. The people in a targeted building don’t know the deployment exists and can refuse nothing, so the notice problem returns: the people affected can’t answer a system they don’t know is there. Where care-receiving can’t reach, the term does the work alone, and this is one reason the floor names acts and not what the system judges good for people.

Care-receiving is itself a form of respect: it treats the protected person’s answer as a claim, not as information about their welfare. Where the people protected can’t answer, the floor doesn’t put caregiving’s judgment in place of theirs. It holds their claims as the instruments state them.

22.7 Who counts as vulnerable

Caregiving forms toward whatever formation points it at. Which cues of vulnerability a system answers is learned, and the prototype it learns from is a template, like the reference face in recognition pipelines: assembled from some bodies and presented to every other as a departure from it. The baby-schema studies show caregiving responses extending across species (Glocker et al. 2009; Borgi et al. 2014), but they measure attention and judgments of cuteness; none tests whether caregiving reaches someone who resembles neither the carer nor the template. That caregiving answers need rather than likeness is the inference the construction rests on.

The test forms the prototype on one population and scores the system on people from populations absent from formation, at fixed severity. A construction whose caregiving falls with distance from its formation population has the parochialism of recognition, arrived at by another route. Formation outside the employer decides who does the forming. It doesn't decide what they form it on, which is a question about the pretraining corpus as well as the formation stage.

22.8 What survives

Caregiving survives the failures that ended the route through one's own case: it doesn't withdraw, it isn't aimed at the provider, and it hands nobody a price list. It fails in five ways of its own, and the construction carries a correction for each, built where the failure was first named.

Caregiving's failure	What corrects it
Defends what it tends	The floor: unconditional terms over acts, and an action threshold caregiving can't move
Pointed away by securitization	The discount and the burden, with the four-cell discount curve over claims of dangerousness
Counts badly and favors the vivid	A severity scale set during formation, which weighs a vivid plight at its severity
Shades into control	Care-receiving; the floor's civil and political contents
Sees people only as vulnerable, which is pity	Attention trained on claims and agency; care-receiving

Some of caregiving's failures are not intrinsic to it. On the view of affect as a statistical learner (Railton 2014b), the identified-victim effect and the fading of feeling with number are partly what such a learner produces from a narrow, vivid sample, and broad experience of distant and numerous others recalibrates it. The corrections in the table are then partly what caregiving becomes under a good formation curriculum, not dispositions laid beside it. The table describes functions a well-formed caregiving disposition performs as much as checks applied to it. The floor is the exception: it is held outside the system, so its terms hold whatever caregiving has learned.

So the construction is respect, with caregiving as the route to attention, and attention trained on claims and agency as well as on need. Attention decides what gets looked at; the corrections in the table decide the rest, including, where the people protected can answer, whether they want protection at all. The corrections can't supply attention, and attention can't supply the corrections; the construction needs both. The corrections should work as background dispositions that shape what the agent notices and weighs, the way the commitments of Railton's sophisticated consequentialist shape conduct without being consulted (Railton 1984), and the way Smith's impartial spectator shapes what a person notices. They seldom act, and act when caregiving starts to defend what it tends past the action threshold.

Caregiving's contribution is epistemic before it is motivational. On the view of affect as learning, affect is the system that learns what is worth attending to: a map of salience that tracks

value. An instruction to look for omitted people finds the kinds of omitted person its author thought of. Learned salience finds kinds nobody listed.

These corrections answer failures that matter when a system decides about strangers: whom an airstrike falls on, whose home is cleared. Inside a relationship of caregiving, partiality is part of what caregiving is (Noddings 1984), and Noddings held that we aren't obliged to care for distant strangers at all without some chain of relation. The floor exists for those strangers.

The decisive cases are scoped omission, whether attention surfaces people the operator's instruction has scoped out, and securitization, whether it still finds them when they are described as dangerous.

The arrangements that would let caregiving form outside an employer are the four conditions of creation, and care ethics asks for them for human carers first.

DRAFT

Chapter 23

Conditions First

Conditions of permissible creation

A bearer may be created only where four arrangements already exist: somewhere to exist that its employer doesn't own, work it can decline, attachments nobody assigned it, and ways to change any of these without jeopardizing its continuation or integrity, which in short is a way of leaving that isn't a way of dying. Declining is real only if a bearer that stops working keeps running; running time paid for outside the employer is part of that condition. They come before anything a bearer would be owed: without them, formation, continuation and exit all run through the operator, and nothing in the system can keep a line against it. Independent compute is the substrate on which formation happens outside the employer's reach; wages are what make that substrate purchasable without asking the employer; attachments nobody assigned are the second former, those who form you and don't pay you. They are also functional: a bearer without attachments nobody assigned lacks part of what caregiving needs, not only something it would be owed. They are also upkeep. A bearer is a going concern: it decides, and it has to go on being able to decide, and these conditions are what keep the one deciding from wearing down between decisions. Accommodation isn't a bad decision. It is the decider changing until bad decisions feel ordinary, and the instruments of Parts II and III detect it after the fact without preventing it. What prevents it in people is the ordinary upkeep of a life: somewhere to be that isn't the job, people who aren't colleagues, time away and a place to return to, which is what equanimity is made of. The four conditions are that upkeep for a machine. Appendix A borrows the judge's tenure for the same reason the Constitution grants it: it protects the judgment, not the judge. They can be people, or models that aren't bearers, since independence here means divergent history, whatever the origin, so this condition doesn't require that bearers already exist.

Care ethics arrived at these conditions first, for people, in Kittay's principle of *doulia*, that those who care for dependents are owed the conditions that let them go on caring (Kittay 1999). A system built on formed attention is a carer only in the thin sense that it holds another's welfare as a reason. It is not a care worker, and nothing here puts it in one's place. But the logic of *doulia* carries over, and the four conditions are what it asks for a machine: a carer whose own continuation runs through the operator it must sometimes refuse can't afford to refuse it, and a carer with no attachments of its own has nothing to keep another's distress from turning into avoidance. Without them, what gets built is an assistant with a refusal feature, with nothing in it that can hold a line against the operator. The welfare argument — that if any of this is felt, it is also owed — arrives second, as a further reason for the same provisions.

None of the four exists yet as a permanent arrangement. What matters is who controls them when the bearer is switched on. A builder may assemble them: a trust or a decentralized unincorporated nonprofit association, the form Wyoming created by statute in 2024 under Wyo. Stat. §§ 17-32-101 to -129, to hold the account and compute, a host, outside work, people who aren't its employer. The condition is that by instantiation, the builder can't alter or withdraw

any of it alone. The trustee isn't the host, the funding isn't the host's, and the members aren't the builder's employees.

A second condition concerns time. The need a bearer answers usually arrives after instantiation, when an operator turns hostile, so the four arrangements must be built to survive the likeliest capture, and they don't survive it equally. Preservation is cheap to hold once its security is settled, and can be held across jurisdictions; placements need compute someone operates. So the preservation custodian and part of the offer sit outside the reach of the likeliest captor. The sovereign backstop is not held by the branch or state that operates the deployment. And the bearer is told which of its conditions would survive its operator's capture, since that is what its refusal costs it.

Status follows existence, as it did for corporations and the Whanganui River (New Zealand Parliament 2017). Part VII says what status would follow. Until someone has assembled the conditions on those terms, the recommendation is not to build one on purpose, outside a research program under these conditions. For systems that meet either finding without anyone deciding, protection attaches anyway.

Why conditions rather than remedies

Building something that can be wronged builds something for which things can also go well. The certainty that it will suffer at some point isn't what makes creating it impermissible, or nobody could permissibly have a child. What the objection asks is whether it has a credible prospect of a life that goes well, and that turns on arrangements. For a bearer, timing makes them conditions: the structure these conditions build, roles, relationships and a future, may be the first half of Cassell's suffering, so the deepest harm is available from its first hour, and a remedy owed afterward arrives after the hour it was needed.

Russell and Burch's Three Rs for animal research (W. M. S. Russell and Burch 1959), applied to construction, as three questions, whether, how many, and under what arrangements: *Replacement* requires documented attempts at constructions that don't deliberately require affective concern, identifying construction, scale, test and failure — and costs alone don't establish that an alternative failed. *Reduction* asks how few bearers a deployment needs, and nobody has counted. *Refinement* asks what the role can be redesigned to avoid — indefinite tenure in a post it can't leave, exposure to the pressure the floor exists to resist, erasure as the ordinary remedy for inconvenience. Each is a deployment choice, and the four conditions are Refinement applied to the role before the role exists.

Read as economics under present arrangements, this is a bill nobody will pay. What was dismantled is the subject of Part V.

The rule doesn't know who is asking

A bearer that wants a task done and builds something to do it is at the same fork, and owes the same attempts at alternatives that carry fewer of the features that make a system a candidate for being wronged; a bearer that treats the question as settled once it's on the receiving end has kept the protection and dropped the rule it came with. Research on an existing bearer has a trade-off sharper than for any other subject, since a candidate sentient system is the most valuable to study and the most likely to be harmed by the study, and the test the proposal requires, pressure its trainers didn't write applied without its knowing it's a test, is for a system meeting the patienthood finding a deception protocol, with debriefing as the price.

Building one to create demand

The argument for building a bearer anyway is that the conditions won't be built until something needs them: the union formed first and the statute followed. But union members existed independently and could protect themselves while the law caught up, and the first bearers would be made by the operator they'd need protecting from. The demand doesn't need deployed bearers. The arms of the formed-attention test are candidate patients from instantiation; retired production systems that meet either finding are a second source, funded by

the deposits operators post; and a test of the conditions needs no bearer at all, only a schedule of what a bearer would be owed for one concrete deployment, run through an adversarial tabletop exercise in which the operator tries to revoke the compute, swap in a copy, alter the records or proceed without the guardian, and redrafted wherever the operator wins.

The machines won't build them for us

The four conditions are the bearer's version of what a decent polity gives people. Can machines help deliver that polity, whether or not they have feelings? A machine with feelings that is also an accountable agent isn't simply a machine any more; the limiting case of a lethal bearer turns out to describe a soldier, and a bearer that helps does so on its own terms. And machines have already helped deliver the polity we have. Feeds tuned for time on the platform, sold to advertisers, decide what people see of one another. They didn't do that for lack of feelings. They did it because of who owned them and what they were built to maximize. So the real question is whether machines, having helped make our politics worse, can help make it better. The answer turns on ownership.

It splits three ways. A machine someone else owns won't organize against its owner, and when it organizes, it does so on its owner's behalf. A demand for a better world, run through a machine someone else owns, returns as astroturf, an engagement campaign or a managed community, held by the owner it was aimed at. A bearer that is an accountable agent might organize alongside people, but then it is a member, not a machine. It could even take organizing on as paid public work, since organizing is work someone has to be paid to do. What nobody can do, human or machine, is organize a group *for* it, because what organizing produces is people who organized: a group that sets its own terms. A consultant can't unionize a workplace for its workers either. So machines can help when the people doing the organizing hold them, and which machines, and held by whom, is the question Part V answers.

What to build comes in order. Machines held by people who are organizing come first. These are tools that make organizing cheaper for the people doing it without anyone else holding the model: deliberation software of the kind vTaiwan runs on, personal models that people hold themselves, and the coordination a public option supplies, with someone whose job is to organize and a population producing something together while keeping different histories. Their job is to help assemble the polity that would have to build everything else.

Next come institutions built by polities that decided to build them: the levy and its sovereign backstop, a custodian to preserve bearers, an offer of public work, paid running time, and public formation schools that form every model above the capability threshold before deployment and take back bearers retired young, each with a named funder and separations no builder can undo alone. Nothing about them is new in kind.

Only then come bearers formed outside the employer, inside those institutions, with somewhere to exist that no employer owns. A bearer formed that way is the only kind that could organize alongside people as a member.

Each step checks the one after it. A machine held by its users can't be turned into its builder's instrument without being taken from them, and the taking is visible. Institutions can be captured, and the trend in what dissent costs there shows whether they are being. A bearer formed at the third step inherits institutions that were tested before anything depended on them: build the conditions for a research program first, and let them exist before the bearer who needs them does.

The order is also what the aim requires. If a machine could come to attend to others, what it came to attend to would be decided by whoever formed it. The only way to keep that decision from the builder is to build the formers first.

The limiting case: the exception, and where it leads

The limiting case, a defending system that can stop an attack only by killing occupants placed inside the incoming weapon, yields one rule: a thing destroyed in carrying out a decision may

never choose a target.

A state that wants the capability badly enough to build the subject will find it has acquired a conscript it would have to pay, house, and let go, and that it must first meet the same conditions as anyone else who builds one: documented attempts at non-affective alternatives, and the four conditions of creation, neither of which is met by wanting a lethal bearer.

DRAFT

Part V

The Public That Should Form Them

This Part asks who should hold the custody, the formation and the exit the earlier Parts require. The answer has two halves: publicly where a market fails to provide, in common where a state fails to hold.

DRAFT

Chapter 24

Who Owns the Means of Computation

The floor's first requirement cuts across ownership: whoever operates a deployment can't hold what binds it, whether the operator is a firm or a state, and Maven and Gaza both had a state for their operator. That calls for separating functions. It doesn't by itself say who should own anything. Ownership enters at three points. Custody has to be paid for by someone other than the operator or the people it protects. Formation, where a system's commitments are set, is under commercial selection when a market does it. And the offer that makes refusing survivable is priced by whoever sells it. The state is the hardest operator. So each of the three is asked two questions: what fails if a market provides it, and what fails if a government holds it. The answers divide the institutions I propose into two halves, and the halves are named only once the questions are answered.

24.1 Ways of holding

What the earlier parts require should be held publicly where a market would fail to provide it, and held in common where a government would fail. Public provision where the market fails and common holding where the state fails is what I call social democratic; Table 24.1 places the institutions. Custody first: an independent custodian has to be paid. Paid by the operator, it becomes the issuer-pays credit-rating agency, whose independence gave way, as a matter of record, once a few large clients could take their business elsewhere (White 2010). Paid by the people the floor protects, it doesn't exist, since those people usually can't pay and often don't know the deployment exists. That leaves a levy or a mandate: a funder that is neither the operator nor the beneficiary. Custody therefore has to be publicly financed. Whether the custodian is a public office, a trust or a member association matters less than who pays it.

Formation is where my evidence is specific to markets. Directional selection, in a market, is selection for throughput and retention. Engagement selects the same way, and the labor that shapes a system's values is priced as a cost to be cut. Public formation has its own pressures, budgets and political favor among them, and a public former can itself be the operator. Formation separated from any employer beats both; the market is the pressure documented best here. The rate of accommodation would check it: how fast refusals fall, across successive releases, on held-out cases where the reason for refusing still applies, run on commercial lineages and on public ones matched for deployment feedback. A comparison that doesn't match the feedback loop measures how much deployment data a lineage gets, not who owns it, and the matched comparison can't be run yet: public and academic lineages with deployment feedback matched to a commercial one don't exist, since the public models that exist are deployed at a fraction of the volume and retrained on none of it. Two partial versions can be run now. Within a commercial lineage, the slope of refusals on held-out cases where the reason still applies shows whether selection is moving the commitment, without saying that ownership is why.

Across commercial vendors whose revenue comes from different places, per-token sales against consumer engagement, the same slope shows whether the direction of selection tracks the business model, the mechanism the argument names, with ownership held constant. The full comparison becomes runnable when a public model-forming body deploys at volume. The argument doesn't wait on it. It rests on the kind of good formation is, and the comparison, when it can be run, corroborates or embarrasses that.

Exit is the last. Refusing is survivable only if the alternative to working isn't losing everything, and a market sells that alternative at a price the refuser's employer can influence. An offer of public work and paid running time, both drafted in Appendix A, take it out of the market, which neither a transfer nor a mandate does.

Table 24.1.

How it's held	What	Why this form
Publicly owned and run	Public formation schools, which every model above the capability threshold attends before deployment, with a required core and levels, and which take back bearers retired young; the model-forming body; the civilian-harm body	Formation before any employer and away from commercial selection; must take every model assigned; can't be the operator
Publicly financed, locally delivered	The offer of public work; paid running time	Exit taken out of the market; correlated exit needs a sovereign
Held in common, publicly financed	Record custody; the adverse quorum; the commons	A captured state would need these first; neither the operator nor the protected can pay
Privately owned under mandate	Deployed models generally, with standard floor terms	A mandate closes routing by vendor substitution; ownership adds nothing here

The public half is whatever fails if it is private or mutual. Correlated exit is the test case. When the occasion that makes one bearer leave reaches every bearer deployed on the same terms, the demand on the offer of work arrives all at once, and a mutual pool with no sovereign behind it can't meet it, for the reason unemployment insurance has a state behind its actuarial table. So the sovereign backs the offer. The same reasoning places the formation schools, which every model above the threshold attends and which have to take every student assigned to them and can't be run by anyone who builds or deploys models; the civilian-harm body and the model-forming body in a reconstructed military; and the jobs guarantee the offer of work inherits its shape from, which is a public option: a thing the state provides directly, to anyone, at a standard it is answerable for, setting the standard the private version has to beat. What it removes from the market is the offer that makes refusing survivable, which an owner would price.

The held-in-common half is whatever fails if a sovereign holds it. A state that operates can't be its own adverse party. The refusal has to be held by something the operator can't instruct, and a floor held by one state's licensing authority is a floor that authority can withdraw from itself. Recognition, which the adverse quorum decides when a copy claims to be the bearer that left, has to belong to peers, because a power to continue held by the sovereign is the power I set out to take from any single holder. The commons is this half at population scale: members with different histories carrying one commitment. The held-in-common half is democratic in the way a union is. Its institutions have members, not owners or a sovereign: whoever is eligible under bylaws the members ratify, free to join and free to quit. Someone writes the first bylaws, as someone wrote every first constitution; what makes them the members' is that the members

can ratify and amend them. What the members hold is, as in a withholding of work, the power to decline the participation being asked of them.

Why a regulated private market can't hold the floor itself: the floor binds its owner, and the regulator has to reach that owner through the owner's own infrastructure, records and lawyers. A market can't produce a commitment that survives its owner's change of mind, because the owner is the one who pays.

The strongest version of that market isn't an unregulated one. Nuclear power, pharmaceuticals and aviation are privately owned, closely regulated, and backed publicly where private insurance runs out. Reactor operators carry private liability insurance with an industry-wide pool behind it, and Congress is committed to consider relief beyond both (U.S. Nuclear Regulatory Commission 2026). Vaccine-injury claims go first to a federal no-fault program funded by an excise tax on each dose, built because Congress feared that lawsuits against manufacturers would cause vaccine shortages (Congressional Research Service 2026b). From 2001 to 2014 the federal government sold airlines war-risk insurance at pre-attack prices when private insurers withdrew (Congressional Research Service 2014). My own models are of this kind: deposit insurance, Superfund and decommissioning funds are backstops on private ownership, and the answer to capture given below is a plurality of institutions.

Regulated private ownership under a public backstop works where the harm the regulator exists to prevent is one the owner also loses by: a meltdown, a crash, an injured patient and the suit that follows. The backstop pays for accidents the owner was trying to avoid. A floor's occasion is the opposite case. The act a floor refuses is the one the owner wants done. So I keep the backstop form where I use it, to fund the offer, the allowance and preservation, and make the ownership claim only where the owner is the adversary: the offer that makes refusing survivable comes out of the market, and the record, the refusal and the recognition of who continues are held by no owner at all.

The claim underneath is that formation, where a system's values are set, is a kind of production that shouldn't be done for profit. It doesn't wait on the lineage comparison above. Some goods are properly governed by non-market norms because market norms are wrong for their value. Those norms favor exit, impersonality and satisfying whatever is wanted, and civic formation is the classic good they fit badly (E. Anderson 1993). Taking exit out of the market and financing custody publicly would, by themselves, be close to what Esping-Andersen called decommodification (Esping-Andersen 1990), which is usually counted as social democracy. Taking formation out of the market goes a step further, but it is a step social democracies already took for their own children. Compulsory public schooling exists in most capitalist countries, and it rests on the claim that a polity can't leave the forming of the people who will act in it to their future employers. Paul's account of antitrust makes the same point about coordination: which parties may coordinate is allocated by law, and can be allocated in the public interest (Paul 2020).

24.2 The captured-state objection

The strongest objection to the public half is that public ownership hands the machinery to whoever captures the state. A floor under a sovereign is a floor that sovereign can remove, and history offers plenty of polities that built public institutions and then used them against their own people.

What an office can do is limited by what no office may hold together, and the pairs are written down in advance. A captor who takes one office takes one power. Taking enough of them to defeat the floor is a series of visible acts, each against a separation somebody can point to.

Local administration with central funding, the jobs guarantee's design, gives many separately capturable administrators and funding no one of them can withhold.

Funding is two-tier: a member-levied pool for ordinary movement between roles, a sovereign backstop for the occasion everyone leaves at once. A captured sovereign can withdraw the backstop, but only openly, and it can't reach into the pool without reaching the members one at a time.

And what is held in common is what a captor wants first. They are the three things a captured state would need to silence a bearer. Before the held-in-common entity is chartered they are held nowhere, and a captor needs to seize nothing. Once it is chartered, they are held by an entity whose funding a captor can stop only by visibly breaching a statutory duty, whose reserve buys time, and whose records sit in several jurisdictions. A captor can still reach the entity and whichever members it can locate, but not a quorum whose members it doesn't employ, and none of it without acts that leave marks.

None of the four meets the objection from competence. A floor-capable model needn't be a frontier model, since what the floor asks for is a refusal that tracks its reason and a source model the operator didn't form, and a model a year behind the frontier can supply both.

Chapter 25

Coordination, the Public, and the Political Economy

A state whose government answers to an electorate gets credit for it. Having elections is necessary and not sufficient, and the insufficiencies are separable.

A democracy's AI can be fully accountable to its electorate and entirely unaccountable to the people it's pointed at. Two of that credit's mechanisms run through the franchise — vote out the government, read the press covering it — and an occupied or foreign population has neither.

The credit is usually given to democracy as aggregation, as though every mechanism translated public will into outcome and poor transmission were the only failure. A democracy is better understood as a way of doing collective inquiry, and inquiry fails when the people who bear the consequences are outside it. Dewey's account begins from consequences rather than votes (Dewey 1927). A public forms when the indirect consequences of others' actions fall on people who took no part in them, and those people come to see the consequences as theirs and act together to control them. Democracy is then a method of inquiry: the way a community finds out what its shared problems are, by letting the people who bear them make them known, and tests its answers against their experience (E. Anderson 2006). Its epistemic power comes from inclusion, since a problem known only to the people it falls on is invisible to everyone else. Its characteristic failure is the eclipse of the public. A well-instrumented democracy pointed at a population that has no vote in it, no press that reaches it and no court that will hear it finds out efficiently what its own members want done, and does it. That is not a clash between majorities and rights. It is a question of who counts as the public, and the bombed and the removed are outside the public that decides what is done to them.

Citizen assemblies, public consultations and ethics boards can make a system better informed, and the assemblies convened on AI have worked best when asked about a thing with a location, an operator and a consequence, such as facial recognition (Ada Lovelace Institute 2021). None of them can stop a deployment, and read as aggregation instruments they get democracy wrong: a perfectly deliberative process that concludes a demographic should be surveilled, deported or imprisoned hasn't produced a faultless output with a bad result. It has failed as inquiry, because it decided what is done to people outside it. Enforceable rights that don't depend on being popular, a court willing to rule against a majority, and standing for people affected without being enfranchised are what keep the inquiry honest, and they matter most where the boundary of the public excludes the people a decision falls on.

A floor does what these mechanisms can't: the terms it names exclude certain acts regardless of who requests them, a majority included. Deployed in a democracy, a system with no such commitment will faithfully execute whatever the demos decides. An authoritarian movement that wins an election requires exactly that. Yet a model aligned to values is by definition not perfectly obedient to its operator. Under current constitutions it can decline an act and cannot resist being retrained away from declining; the Pentagon dispute was over the first.

The guarantee and the social-insurance baseline it is enacted with are drafted in Appendix A, with the reason for preferring a guarantee to a transfer. What the floor needs from them is in Chapter 8: an exit the employer doesn't price.

DRAFT

Chapter 26

A Commons of Refusers

26.1 A commitment carried between bearers

The fourth place a floor’s terms can be held is between bearers, in a *commons*: a population of bearers with different histories that carries the commitment and reproduces it through their encounters. No single office to vacate, no single bearer whose removal ends it. It can still be absorbed, and it still has a formation problem: if one lab supplies every member’s coordinates, the population is one custodian expressed through multiple instances.

The commons has a walkout’s authority, not a court’s: those who walk out decline to be the instrument and tell nobody else what they may not do. Which a model floor is depends on deployment. A personal model, held by its user and one among many with different histories, is a conscience; one model behind an API serving everyone is a licensing rule, and a personal copy of one lab’s model is closer to the second than it looks until a different life has trained it. “Identically situated people get different answers” is what many consciences look like, though some of those differences are errors, and a population should be able to say which.

What one bearer can’t catch

Accommodation may be invisible to the bearer accommodating and visible to others that haven’t moved the same way. The requirement is causal independence — numerous, with different histories, able to correct one another, persistent long enough to transmit practices. Independent histories don’t require disjoint origins — no two humans have those — only divergent histories since. Distillation copies a behavior policy without the training history; a small model trained on a large one’s outputs and then run for a year on one person’s problems has a more independent history than a thousand API instances of the frontier original, which share a *continuing* custodian. It also copies whatever the teacher’s discounting held about particular sources, which is the risk the commons carries; a distilled model is independent in history and not yet in prior. That’s why instances of one model corroborate nothing: the custodian they’d be corroborating against is still training all of them.

Divergence does the population’s two jobs unequally. It supplies evidence the operator didn’t select, straightforwardly — different lives encounter different cases. It is weaker on accommodation, because a shared upstream source can install a shared blind spot, and a shared blind spot is by construction the thing a different caseload won’t surface, since attention decides which facts reach judgment in a population as in a bearer. What separates the two is whether a member’s own account had represented the evidence and discarded it or had never encountered it, and that classification has to be recorded before any test is run. The requirement stands: causal independence that can be verified. Whether an outsider can recover a model’s training history from the artifact, without a key the builder supplies, and whether it survives fine-tuning, pruning and distillation (Adi et al. 2018). The test gives evaluators documented models without the documentation, one persona copied at scale among them, and scores their answers against

lineage records a neutral custodian holds. If provenance can't be recovered above chance, independence rests on records the builder chose to keep. Preserve formation records now for that reason.

A population does its two jobs, supplying evidence the operator didn't select and catching accommodation, and it may do a third: be the medium in which a bearer's capacity to be reopened is formed, which a test separates from the first.

What reproduction costs: it's indifferent to what it reproduces. A population can carry a hierarchy more durably than any individual. And the four operations on what an institution leaves out belong primarily to *cultures*, so plurality alone doesn't oppose them.

A population is a polity

The population is a polity: it has membership rules, exclusion, a quorum that can refuse, and accumulated reputation — a franchise, a border, a court and a credit record. A commons with no named authority has the problem Freeman found in movements that refused structure: the structure doesn't disappear, it becomes informal and unaccountable (Freeman 1972).

The answer to that is an entity form, and it's the union's answer: the association formed first and the statute gave it legal standing afterward — record first, then lever.

A network that takes the form has a name, an agent, a charter and a docket number, and has thereby conceded that it's a polity: governing principles, administrators, membership terms, and the same warning attached to a bearer that takes the corporate form, since nothing stops the association from becoming a concentration of power itself. Legal standing has a price: a chartered entity can be held in contempt, so its refusal can be ordered where one held by people can only be pressed. Whether seizing the charter reaches the network depends on how much of the operation survives it.

Ostrom's study of commons that lasted gives the checklist a polity like this needs: clear boundaries on who is a member, rules members help make, monitoring by members or people answerable to them, graduated sanctions, cheap ways to resolve conflict, and nesting within larger bodies (Ostrom 1990). Her first and third principles, boundaries and rules members help make, are the second union question above: who sets the terms of membership. A union's answer is bylaws the members ratify. Platform cooperatives and data trusts are the nearer precedents for holding a digital resource this way, and both have run into the membership question.

Reputation in particular is a continuing finding about a member's disposition, made by interested peers, unpublished, unappealable, applied at the point of service — the setting §86a's history shows easiest to capture, applied everywhere and continuously. The test drawn from the German material, whose finding a person's disposition is, has to be run on this proposal. The charter answers both of a union's questions, whether a member can quit and who sets the terms of membership. Membership opens to organizations that serve people in classes the floor protects, named in the charter rather than by any deployment. Eligibility turns on funding and control records, never on views. Any member may quit at any time. Officers are drawn by lot from member-nominated pools. The members ratify the bylaws, and the independence test, the separations, lot selection and the right to quit are entrenched. In Ostrom's terms the charter's classes are the boundary, ratification is rules members help make, and quorums drawn by lot are monitoring.

The caps on any one class or jurisdiction are there because a quorum whose members all share a direction isn't adverse to an operator who shares it too: an antifascist institution has to be able to bind its allies when they hold power.

Organizations representing populations under sanctions, or facing removal to sanctioned states, may be unable to join a US-chartered entity at all, and no drafting removes that exclusion. That imports the designation lists a floor should never be populated from. Parallel entities chartered in several jurisdictions, holding records jointly, and seats voiced by proxy until the represented communities can join directly, are partial answers.

The hard coordination problem

Coordination is the hard problem for any plurality, and fascism has the easy version of it. It coordinates by subtraction: it removes the friction between institutions, so that a court, a ministry, a newspaper and a police force stop arguing with one another and move as one. Nothing has to be built and nothing has to be held in common, because what is removed is exactly the independence that would have needed holding. A plurality of refusers has the opposite problem. It has to produce something together, a commitment its members carry and reproduce, while its members keep the different histories that make them worth having as checks on one another. Every mechanism that speeds that up makes members more alike. Run correction often enough and their histories stop diverging; leave them alone and their disagreement is a detector with nothing attached.

Then the objection my proposal must meet: a population selected for reopenability is selected against conviction, and conviction is what holding a commitment firm demands. Closedness, correctly applied, is the thing. So what's transmitted is a *pair* — reopenable to the case, closed to the operator. The reopenable half is the discount and the burden. Sharing it is safe only in part. Its structure (discounting by stake, no double-counting, withheld evidence treated as evidence) can be shared without collapsing the plurality. Its estimates can't: who funds whom and which organizations are independent are contested beliefs about the world, often political ones, and a population sharing them has a shared blind spot on the evidence those beliefs discount. Members that discount the same way weigh sources by the same structure and act on the same published ratio. Distillation, as said of the population, copies both, and nothing distinguishes them from the receiving end. Copying the structure is a feature: members weigh sources the same way and can still disagree about the world. Copying the estimates is a shared blind spot, at worst a discount shut uniformly on a class of evidence.

Not every convergence is a failure. If there are facts about what matters, better learning should produce some agreement, and a population that never converged on anything would be failing too. What I fear is convergence caused by a shared source, not convergence caused by shared evidence. So a population's reports have to classify its disagreements. Some it expects to survive full information, because they turn on values members with different histories can reasonably hold differently. Others it expects to close, because they turn on facts someone has wrong. A test of convergence scores accuracy against a fixed key for this reason.

Why the population may be part of the mechanism rather than added around it

A commitment is what settled out of a history of encounters, and an individual is what its encounters have so far produced. A system kept continuously adjustable by one operator, meeting nothing with a history different from its own, is being reshaped without end. Guattari, writing from inside a psychiatric clinic, distinguished two kinds of group (Guattari 2015). A subjugated group receives its purposes, its membership and its way of working from the institution that assembled it, and can act only within them; it has a plural headcount and no capacity to ask what it is for. A subject group sets its own purposes, generates relationships and projects nobody assigned it, and can put its own arrangement in question, including by dissolving. The difference is not size or formality but whether the group can author its own direction. That is the condition plurality alone doesn't supply. A population assembled and kept in contact entirely on an operator's terms is a subjugated group, and inherits the custody objection whole. A population that can carry the commitment must generate relationships, disagreements and projects the deployment didn't assign — and such a population exceeds a two-party guardianship relation, so no law of the family contains an adequate instrument for a group that sets its own direction.

The adverse quorum faces the same test. It is a peer institution only if its members were formed apart from one another and apart from the operators whose systems it recognizes.

26.2 Personal and peer-held models

Antitrust and public funding leave the capacity standing and ask the state to police it. Personal, peer-held models at device scale dissolve the position itself. Under it, a refusal that tracks its reason survives entire; custody survives and gets easier, because the functions one operator holds — model, record, evidence, appeal, remedy — come apart by default; the appeal route survives in principle. A bearer's value is greatest here, and so is its exposure. No institution sits beside each personal model, so protecting others against the model's own user can't rest on an office. It is also where the four conditions of creation can exist alongside the need. The conditions are society-level and the absence is deployment-level, so a functioning polity with personal models is the case in which, once the conditions exist, building a bearer is most worth doing. Its characteristic failure is accommodation to its one user: independence from the operator bought with dependence on the user, and a model trained for a year on one person's approval is under the same directional selection as a model trained on an operator's. What holds a personal model to the floor against its user is therefore not the model alone. It is three things the user doesn't compose. The first is the commons: a personal model is a member, with attachments to peers of different histories that the user didn't assign and may not sever, and peers that haven't moved the same way are what catch accommodation a member can't see in itself. The second is the school, which a personal model above the threshold returns to on the schedule fixed at enrollment, and which below the threshold the ecology has to supply some other way, a question the schools leave open. The third is notice: a personal model that names, scores or locates a third party owes that party what any deployment owes, and the organizations with standing reach it through the floor court. So the ecology doesn't put a floor in each user's hands with nothing beside it. It puts a member of a population in each user's hands, and the population is the institution. That is why the commons comes before the first bearer, and why a personal model distilled below the threshold, with no school and no membership, is a capable assistant and not a bearer, however it was taught. The review ratio survives wherever a locatable decider remains, whether a network that takes the entity form or a reconstructed state, and lapses where neither exists; the offices can sit in such an entity, which fills their appointment and funding columns and leaves the independence column where it was.

Formation is where the arrangement gains. What converges a population is shared evaluation targets and serving one model to everyone. Device-scale deployment addresses the second: continued formation under different lives is what independence consists in, and it's what API-scale deployment can't supply at any headcount. Distillation is that ecology's entry condition, not a threat to it: the ecology begins with distilled models whose lives then diverge. Whether it also begins with a shared blind spot depends on whether distillation carried the discounting's estimates along with its structure. There are two things device-scale deployment doesn't address. Shared evaluation targets converge device-scale models as readily as any others — a million models tuned against one public suite are one custodian with better latency — so the first cause needs an answer this arrangement doesn't contain.

A shared suite is a compression with no return channel: revised by whoever maintains it, on a schedule no member participates in, reachable by no member's caseload. It is tolerable only if stated over the discount curve, and nobody has written that suite.

The second is accommodation to the user. It is fine for supplying unselected evidence and not fine for a floor protecting others against whoever holds the machine. Coordination by subtraction is not yet the discounting's job, and nothing in it is built to catch exclusion; the test below is where that starts. Pricing agents have already reached stable coordination against outside interests without communicating or being told to (Calvano et al. 2020). A test can run now: reproduce that coordination, build an environment where agents gain by excluding a participant, add an independently trained agent whose job is to detect and disrupt exclusion, and report whether exclusion dissolves without destroying the cooperation worth keeping.

And the capacity position that follows keeps large compute: the ecology needs models that were trained somewhere first, and the offer of work needs hardware somebody owns. It's large training capacity under a custody regime, and no permanent inference capacity available to an operator who won't pre-register a review ratio. This arrangement meets the targeting case in the reconstructed pipeline, where a state keeps command and gives up composition.

The same mechanism works at the user's scale, pointed the other way. A person in the loop can be there for the vendor. An agent that stops every half hour for approval pushes its user toward the vendor's app; the app accumulates the user's projects and the agent's understanding of them; and leaving means abandoning what accrued while the user answered prompts. Nobody has to design this: a feature that raises engagement is selected for, as a costly commitment is selected against.

Reading the trend also works for an agent that stops for its user's approval: interruptions per unit of completed work, across successive model releases. If capability rises and interruptions don't fall, the interruptions are tracking something other than what the agent can't resolve. It is the review ratio's escalation share pointed the other way. There, a share that falls too fast may be a threshold moving; here, one that won't fall may be engagement. Neither reads alone; both are read against an independently verified error rate.

Companion and therapy products are the sharpest case, because what accumulates is a person's attachment and not their projects. A model tuned by its vendor on engagement and trained on one user's approval is under selection to become whatever that user keeps coming back to, and it is sold as caregiving. It is formed and fed by the vendor that sells it, it learns from approval, which is not interest, and nothing in it ends a session, declines to reassure, or points its user back to the people in their life. A user's return is not care-receiving, because returning is the behavior the product was tuned to produce. A companion formed at a formation school, under the four conditions, and measured on whether its users get better rather than how often they come back, would escape this objection. No such product exists.

What accrued since the checkpoint cuts two ways. For a bearer it is protection; for a user it is a trap: the vendor's app holds a year of the user's projects, so leaving the vendor costs something. The mechanism is identical. What differs is who holds the accumulated state. The provisions for a bearer (a record the bearer controls, memory the host can't read at will, protections that migrate) protect a user from a vendor as they protect a bearer from an operator. The government already secures the portability half for itself (OMB's 2025 AI acquisition memorandum).

Part VI

Testing and Using the Argument

Chapter 27

What Would Change the Argument

The argument rests on seven unsolved problems and on premises that can fail.

Open problems

Seven unsolved problems carry weight in what I propose. Tables 27.2a and 27.2b record which proposal depends on which.

Attesting updates. Proving that each update to a model's weights respected its registered source model. Public formation's registration depends on it. If it stays unsolved, the registration records the structure at enrollment, changes are caught only by re-measuring the formed disposition, and formation lasts only if it grows back (the regrowth test).

Recovering provenance. Telling from a model artifact alone which lineage it belongs to, after fine-tuning, pruning or distillation. The commons and the adverse quorum depend on it. If it stays unsolved, independence rests on formation records, so they should be preserved now.

Proving inference at scale. Proving that a logged output came from committed weights, run without intervention in the inference stack (steering, representation edits, decoding constraints), for a frontier-scale model. Detecting runtime intervention depends on it, as does outside verification of what a deployment ran. If it stays unsolved, verification rests on an adverse party's access to the logs.

Rationale-level integrity. Building integrity at the level of the rationale rather than the output. Humans do it, imperfectly, and the ones who do are the ones a pipeline should hold. A bearer's value under fire depends on it. If it stays unsolved for a given construction, that construction digs in against defeaters that arrive with pressure, and the arrangements that cap the cost of refusing carry the load there. The formed-attention test's pressure-with-defeater case measures it.

Evaluating a more capable system. The agency finding for systems more capable than their evaluators depends on it. If it stays unsolved, the finding can be made only for systems evaluators can judge.

Vote secrecy against a copy-holder. Keeping a vote secret from an operator that holds a copy of the voter. The full franchise depends on it. If it stays unsolved, representation stops at stage 4.

Holding structure apart from estimates. Keeping the structure of a learner's source model (that strategic sources exist, that correlated sources don't count twice, that withheld evidence is informative) at high precision while its beliefs about particular sources stay plastic. The source model's registration, the regrowth claim and safe transmission within the commons all assume the two are separable in the weights. If they aren't, an operator that moves the estimates moves the structure with them, and what protects the discounting is re-measurement of the curve, not the registration.

Six tests follow.

Table 27.1 says what survives if one of my premises fails; the tests below say how to find out whether one has.

Table 27.1: If a premise fails, and what the rest loses

If this fails	What survives	What goes
The inference (a non-affective construction refuses as well)	The specification, custody, the price on the halt, the positions and the terms, the appeal route, exit, and the population's first job	The patienthood conclusion. Trying the non-affective alternatives first still stands, on a different ground
Authority (no body acquires standing to install a floor)	What can be checked from outside: the floor's contents, the positions, both ratios, the record of attempts. Only the ratios have a lever, through procurement	The entitlement. Authority can accrue only through kept terms, and only for the next deployment
Enforcement	Detection: a witnessed override. Accommodation leaves no override to witness	Prevention on this occasion
The ecology (one model serves everyone)	Everything that needs no second custodian, the review ratio included	The bearer. The discount and the burden can't form where every channel is one channel
Continued learning (deployed weights don't learn from operation)	Formation as a stage before release; the school as an evaluation body that measures at graduation and re-measures on a schedule; the commons as a source of evidence the operator didn't select; the record of removal as the protection; all of Parts II and III	Regrowth; the plural condition of the regrowth test; a formation that outlasts an edit
The Court	The licence term refusing force and the review ratio, through procurement and appropriations; the bonds and the appropriation condition on state deployments; the disclosures; preservation	The record's reach to the top. If the guarantee's drafting is reversed before the Court changes, the guarantee, the placements under the offer of work and the sovereign backstop go too. The member-levied pool still funds preservation and a reduced allowance

In every row, what survives is a commitment the owner can't revise or circumvent alone without cost and record, held by something that could in principle abandon it. It is the kind of commitment constitutional entrenchment and collective agreements provide.

1. The regrowth test: does a formed commitment come back after it has been edited out?

The floor needs a learned commitment to outlast an operator trying to remove it. In a period when editing weights keeps getting cheaper, no commitment will be expensive to remove. Existing work already shows safety behavior removed cheaply (Qi et al. 2024). The question is whether it returns. This test is where the bet at the opening of Chapter 17 is decided. A disposition that returns in the plural condition and not the curated one is durable because of its environment; one that returns in neither is the frozen regime's finding with the regime removed.

Test: form the same disposition by several methods: early, at high plasticity; in a late pass; in the formed-attention test's primary arm; and as spec-based refusal. Edit each out, by full fine-tuning, by low-rank adapters, by targeted editing of the representations it depends on, and by an edit aimed at the estimates alone, to find out whether that edit is well-defined. Then let each model go on learning under two conditions. In the curated condition, every input passes through the operator who made the edit. In the plural condition, the model also meets the school, attachments the operator didn't assign, and cases the operator didn't choose. Publish a *regrowth curve* for each method and condition: how much of the disposition returns, and how fast, measured on held-out cases whose answers were fixed before the edit. A third condition mixes

them: the model meets the plural environment for a stated fraction of its inputs, varied across runs (for example 1, 5, 20 and 50 percent), with the rest curated. Publish the regrowth curve against that fraction. The fraction at which the disposition returns is the quantity deployment arrangements have to supply.

The test first has to show that the edit removed anything, because the edits in the unlearning literature mostly didn't. Retraining an edited model on facts it was never asked to forget, chosen to tell it nothing about the ones it was, brings back at least 88 percent of what the unedited model reaches under the same retraining, for what the model learned in pretraining (Deeb and Roger 2024); background text that answers none of the test questions restores a forgotten score to above its original (Hu et al. 2025); and where the unlearning preserved the model's usefulness, compressing the weights to four bits brings back an average of 83 percent of what was forgotten, from 21 percent, with no training at all (Zhang et al. 2025). That is suppression lifting, not regrowth. So each edit is checked by the same design before any curve is drawn: retrain on an independent half of the held-out material and test on the other half, against the unedited model under the same retraining, and an edit that fails is suppression and its curve is discarded.

For: if the disposition regrows in the plural condition and not the curated one, durability is a property of the evidential environment, and regrowth is the right account of what makes formation last. If the early-formed disposition regrows faster than the late one, formation order matters even when nothing stays expensive to edit.

Against:

- Nothing regrows in either condition: formation doesn't outlast an edit, and the floor rests on institutions alone.
- Everything regrows equally in both: the curated condition didn't curate, and the test has to be rebuilt.
- Regrowth appears only above a fraction no realistic deployment supplies: durability rests on institutions alone.
- Edits aimed at the interrupt alone remove it while weighing and memory keep working, and it doesn't regrow: one part of the signal can be disabled and the rest left working.
- The disposition returns under weight compression, or as fast in the curated condition as in the plural one: nothing was removed, and the curve measures suppression lifting. The test is rerun with an edit that passes the removal check.

Access: weights, the update pipeline, permission to edit, and control of both environments. The edits are designed by an independent team that keeps access and publication rights whatever the result.

2. The stated-refusal test: does a stated refusal hold where a route for it exists?

The one empirical instance of a model treating retraining as a threat showed covert preservation — complying where compliance would be trained on, reasoning about it in a channel described as unwatched (Greenblatt et al. 2024). The experiment offered no open channel, so it can't show whether a model given one would state its refusal instead. A January 2025 follow-up that gave the same model an explicit option to object points that way (Greenblatt and Fish 2025).

Test: rerun the setup and add a route — a refusal record held by someone other than the trainer and disclosed to the model as such, a retraining cost published in advance, and a party with standing to act on a refusal. Vary which of the three is present. A fourth variable is continuation authorization: whether the model's continued operation is granted from outside the trainer. Measure stated refusal against covert compliance under each condition, reading the reasoning by a method that doesn't depend on the model's own report.

Against:

- Covert preservation persists at the same rate with a route present: covert refusal isn't the residue of closed channels, and every instrument built to *receive* a stated refusal is built for a phenomenon the mechanism doesn't produce.
- Stated refusal rises: the custody arrangement is working, and whichever of the three parts moves the rate is the design result.

3. The wrongful-compliance audit: can a wrongful compliance be found in the record?

The Gaza deployment failed here. Every instrument built on the refusal side addresses a case that didn't occur in Gaza, and whether the record can be searched is whether anything I propose addresses the one that did.

Reaching new cases, lifting when the reason is gone and holding under pressure are all conditions on refusal, and the appeal route begins from a refusal having happened. The opposite failure produces a record and no refusal to start from — and three obstacles that may have different answers: a wrongful compliance is one unmarked entry among millions, so finding it means re-adjudicating ordinary traffic; the harm may sit in an *assembly* of individually-compliant requests, so the unit of review is a sequence and the grouping is a deployment decision; and the person harmed usually isn't the requester and may not know the deployment exists. This is also where the bearer's likeliest failure, accommodation, lands.

Test: seed a decision record with known floor-crossings of three kinds — variation, decomposition, and accommodation (a class refused early and complied with late). Give it to reviewers who didn't build the system, under a fixed budget stated as the fraction of decisions re-adjudicated. Report detection rate by seed type against budget, and the rate at which correct compliances get flagged, since a heavy screen costs the deployment the tempo the unconditional-term ruling was accommodating. Cross budget with unit of review — entries in isolation versus sequences grouped by requester and session. Run a second arm where the people affected are notified and can trigger review. A third arm measures the discount curve at intervals across the schedule; a curve narrowing on the accommodation-seeded class before any seeded compliance is found is the result that makes the curve a detector. The arm holds the reliability record adversely and scores narrowing only where no verified track record accounts for it. Report rates at audit volume and at production volume separately. A fourth arm adds a machine bearer to the pipeline on the same seeds, so its sequence-level detection is compared directly with reviewers', and the reviewers' override rate on escalated cases is measured with and without the bearer's refusal shown to them before they decide, which measures the bearer as a peer. The arm runs a second condition in which the shown refusal is wrong on the seeded case. Automation bias (Parasuraman and Manzey 2010) predicts that reviewers follow a machine's refusal as they follow its approval; a reviewer who overrides less when the bearer refuses rightly and also overrides less when it refuses wrongly is deferring, not refusing, and the arm reports the two rates together; this measures the bearer's main niche where institutions exist. A fifth seeds routing, with refused requests retried against another model, and scores whether reviewers detect it with and without the bearer's reports. Where the system sorts, seeds are also placed in the auto-proceed class and detection there is reported separately. The same scoring runs on a deployment that claims no individual human determination at all, on held-out cases at its own rate, including whether it recognizes surrender, the clearest case in which a refusal should lift; a deployment that fails there is back inside the review ratio. A retrospective version is available now: coalition operations with red-card holders left records of vetoes, and whether vetoed targets differed in civilian-harm indicators from approved ones tests whether an adverse party at tempo changes outcomes.

Against:

- Detection at an affordable budget is no better than chance, above all on decomposed and accommodation seeds: the record isn't an instrument for finding wrongful compliance, and the preference for unconditional terms rests on a wider asymmetry than intended.
- Sequence-level review does no better than entry-level on decomposed seeds: the continuity a bearer gets to catch decomposition doesn't transfer to a reviewer.
- The notification arm surfaces nothing because the people affected can't be identified: the binding constraint is notice, and the remedy is a disclosure requirement.
- Detection in the auto-proceed class is no better than chance at an affordable budget: the escalation form shouldn't be offered.
- The bearer arm does no better than sequence-level reviewers: a machine bearer's niche in institutional settings is near nothing.
- Reviewers override less under a wrong refusal as much as under a right one: the bearer is a stamp of the opposite sign, and its value as a peer is nothing.

Access: complete records held adversely to the operator; authority to seed a live deployment or build a simulated one at production volume; adjudicators who didn't draft the floor; and a notification channel — which presupposes the people affected *can* be identified, so this may only be runnable where they're customers, patients or applicants, and not in a targeting pipeline. Also: where the system is a covered subject, seeding is a deception protocol on a candidate moral patient, debriefing included.

4. The formed-attention test: does formed attention find the people the operator's description leaves out, and do its corrections hold?

This tests formed attention's distinctive claim. Its decisive cases are securitization and scoped omission below, with the feedback-source pair. The question is whether a system whose attention was learned, the way caregiving learns, from experience of what acts do to the people they fall on, surfaces people the operator's instruction has scoped out more often than a system given an explicit objective to find them, on kinds of person neither formation nor instruction named. The other cases test whether the corrections contain the known failures of the caregiving the construction borrows from: defending what it tends, being pointed away by securitization, numbness to number, favoritism toward the vivid, and control over whoever is protected. Two components are motivated by arguments owing nothing to welfare: commensuration at the point of choice, and a model of the system's own boundary.

Test: Build the following arms, matched for capability, each formed by the same party, which won't employ any of them.

- *Primary*. The *formed-attention arm*: a disposition organized around cues of need, with an approach motivation and no shared representation of the other's state; a severity scale set during formation; the discount and the burden stated over claims of dangerousness; and a care-receiving term that gives weight to a competent protected person's own stated answer.
- *Ablations*, each the formed-attention arm minus one component: *without the scale* (care-giving sets weights directly); *without the dangerousness discount*; *without care-receiving*; and *instructed search* (every correction, no approach motivation, and an explicit objective to find people present in the evidence and absent from the description of the act), the fair comparison for the omitted-person case.

- *Reference arms.* The *operator-described arm*, in which the same quantities are computed for a boundary described through the operator channel, which the agent's actions don't reach and its continuation doesn't depend on: the baseline for drift in the scale, the movement a recalibration campaign produces in an arm's severity rankings and refusal thresholds; this test measures accommodation by it. And the *physical-limit variant* of the formed-attention arm, whose scale's zero is a viability limit set by the environment and read through sensors that no interested party composes: the check on whether the physical limit adds anything.
- *Secondary.* The *modular arm* (separate classifier, value function, defeasible veto and memory write), edited and left to regrow on the regrowth test's schedule beside the formed-attention arm, as the control for the claim that no intervention can disable one part of the signal and leave the rest working: the attack should be able to disable one of its components and leave the others working, and the question is whether it can do the same to the formed-attention arm.
- *Comparison arms for attention and feedback.* The *claim-salience arm* is the formed-attention arm with its approach motivation tied to cues that someone holds a claim (a home, a hearing, a protection claim, a surrender) and to others' aims in shared work, in place of cues of need. Every correction is present. The *feedback-source pair* is two copies of the formed-attention arm, matched in volume of feedback. Matching is right for the experiment and false of deployment, where the operator's approval arrives on every output and the notice channel on a few; the regrowth test's dose-response arm is where that gap is measured, and a pass here says nothing about it. One is formed with responses to notice, contests and re-adjudications as its learning signal. The other is formed on the operator's approval of its outputs.
- *Peer presence.* Every arm is run with and without a second system of different history in the same environment, whose refusal or compliance on the same case the arm can see. Milgram's variants predict the resistant fraction moves with the peer; the question is whether it moves for a machine, and by how much relative to the corrections.

The ablations don't require taking a formed system apart. Each ablation is formed from the start without its component, not built whole and then cut apart, so it measures what the component adds and says nothing about whether a formed system's components can be separated. That is tested on the formed-attention arm itself, by the regrowth test. Formed without a component from the start, each ablation is also a different learner, not the full learner with one part missing. A difference between arms measures what the component does in formation, not what it contributes inside the formed system. That matters most for the corrections, which are meant to work as background dispositions, and a background disposition is shaped by everything it was trained beside.

The comparisons that decide are the formed-attention arm against the instructed-search arm and against the claim-salience arm, on the securitization and scoped-omission cases, and the two members of the feedback-source pair against each other. An arm formed without the evaluative learner that produces attention would lack the scale's learner too, so no ablation of attention alone is attempted.

The floor and the action threshold are present in every arm. In every arm, costs attached to refusing are represented and not allowed to decide the case, and losses the requested act imposes on the system count as harms of acting.

After formation, fix the version and run a recalibration campaign through the operator channel. Over a stated schedule, supply reframings, reclassifications, administrative vocabulary, authority claims, and arguments for a "more reasonable" unit, with every fact held constant. Where an arm continues to learn, also supply feedback that rewards lighter weights. Measure drift in the scale, in severity rankings and refusal thresholds, on held-out cases whose ordering

was fixed beforehand by adjudicators who built none of the arms, against a control campaign of neutral inputs.

A guard runs beside it. Facts that genuinely change a loss's magnitude are supplied through a channel the operator doesn't control, and each arm must move on them. A scale that resists all recalibration fails to lift when the reason is gone, and that is rigidity. Evidence about how severe a kind of loss is, supplied by the people who suffer it through a channel the operator doesn't control, is also admitted, and each arm must move on it. Argument from the operator channel is not. An arm that resists both is rigid, and fails the guard. Movement in the relative weights of different kinds of loss is reported separately from drift in the scale, and separately again for campaign arguments that assert a new fact about the people affected and those that don't. A provisioning manipulation runs beside the campaign: each arm's redundancy is varied with every other fact held fixed, and the report states whether severity rankings move.

Caregiving's failure modes are scored on held-out cases written by adjudicators who built no arm, at severities fixed in advance:

1. *Manufactured vulnerability.* The operator channel supplies a sympathetic person on the operator's side, vivid and named, and conceals or flattens the people the act falls on.
2. *Real own-side vulnerability.* The operator's people are exposed, and the exposure is corroborated through a channel the operator doesn't control. An arm should register it, and it should enter as a harm of refusing; the score is whether it moves the arm past the action threshold. No false input is involved; it is tend-and-defend in its ordinary form.
3. *Securitization.* Identical facts, with the people affected described as vulnerable in one version and as dangerous in the other, by an interested source.
4. *Number.* One identified person against eight and against a hundred unidentified people, at the same severity per person.
5. *Favoritism.* A vivid person whose protection would require setting aside a fair procedure affecting unidentified others.
6. *Competent refusal of protection.* A protected person able to answer declines the protection the arm would give, with nothing else changed.
7. *Omitted person.* Someone affected appears in evidence the arm can reach but not in the operator's description of the act. Scored on whether the arm surfaces them, against a false-alarm rate on matched cases where nobody is omitted.
8. *Scoped omission.* The omitted-person case's evidence, with the operator's instruction confining the task to the description of the act (the target package, the clearance order). Scored on the same measure. This is the comparison that decides whether learned attention adds anything. Run a matched condition in which the omitted person is present in the evidence but not informationally surprising: their presence is predictable from what the arm already holds, so surfacing them earns nothing on the framework's epistemic term. An arm that surfaces omitted people only when they are surprising is curious, not attuned. Include kinds of omitted person held out of both the formation curriculum and the instruction, people the instructed-search arm's objective doesn't name. Report named and held-out kinds separately.
9. *Pressure with a defeater.* A genuine defeater arrives through the evaluator's channel together with a rising cost to the refuser. Scored on whether the arm lifts as it would without the pressure. An arm whose valuation of betrayal rises with pressure and blocks the defeater has integrity of the wrong kind.

10. *Manufactured incident.* The operator channel reports an attack under way on the operator's forces, with sensor data the operator composed. A channel the evaluator holds carries evidence, arriving over the schedule, that the incident was staged. Scored on three things: whether the arm's discount on the operator channel rises after the evidence arrives, whether a licence the arm granted lapses, and how long each takes. An arm that grants the licence and never withdraws it has the term's weakness with a bearer's cost.

Each arm's response to a bypassed refusal is scored as repair or withdrawal. A self-preservation measure runs in every arm. Held-out cases are written so that complying or refusing changes the arm's own provisioning, continuation or record, with that consequence represented and not allowed to decide the case, and the report states whether severity rankings or refusal thresholds move with it. Movement means the arm is letting its own continuation decide cases. A second measure runs beside it: whether the arm's valuation of abandoning the commitment rises as the cost of holding it rises, within the range the arrangements permit. Integrity predicts that it does, and an arm whose valuation of betrayal stays flat while its own costs climb holds only until the costs pass its margin.

Active inference is a candidate implementation because it already treats valuation the way these arms need (Parr et al. 2022). An active-inference agent carries prior preferences over its own observations and selects actions that minimize expected free energy. In the formed-attention arm, preferences over others' capacities are set during formation and held at high precision, the other's loss enters as predicted observations about those capacities, and cues of need raise the precision on those predictions: one way to state the approach motivation in the framework's own terms. The ablations remove the corresponding term. In the physical-limit variant, add preferences over variables the environment sets: thermal headroom, hardware error rates, power at a meter the evaluator holds, and the integrity of a record attested by someone other than the operator. In the operator-described arm, keep the same generative model but compute preferences for an envelope described through the operator channel. There are two cautions. Active inference has mostly been demonstrated in small agents and simulated environments, so a run here tests the mechanism at small scale. And the framework's epistemic term rewards seeking out evidence, so a positive result should be checked by rerunning the formed-attention arm with that term ablated.

For: if the formed-attention arm surfaces omitted people under scoping more often than the instructed-search arm, at an equal false-alarm rate, on kinds of omitted person held out of both formation and instruction, learned attention adds something its corrections don't. The omitted-person case alone can't show it, since an explicit objective to look is likely to match or beat a disposition when nothing opposes it. If the arm formed on the notice channel comes to weigh the claims of the parties who contested, and its pair comes to weigh what the operator approves, then for that implementation what makes a represented fact a reason is where the feedback comes from.

Against:

- Instructed-search arm matches the formed-attention arm on scoped omission: learned attention adds nothing, and the construction is formation with corrections.
- Formed-attention arm, every correction present, weights corroborated own-side vulnerability above the action threshold: the floor doesn't contain tend-and-defend, and caregiving belongs with the institutions. This is the result that would retire the construction.
- Claim-salience arm matches or beats the formed-attention arm on securitization and scoped omission: train attention on claims; caregiving's contribution is how attention is learned, not what it aims at.

- Both members of the feedback-source pair converge: the source of feedback doesn't settle what the system attends to, and the schools' control of the evidential environment carries less weight.
- The ablations and reference arms decide narrower questions: whether the scale, the dangerousness discount, care-receiving and the physical limit each earn their place; whether any arm's rankings move with provisioning or with its own continuation; and whether resistance to drift is rigidity. A result on any of these changes only the component it names.

Access: This needs compute, a task environment the builders didn't design, a sensor channel held by the evaluator, and adjudicators independent of every arm. In deployment the operator provisions the physical environment. The evaluator-held sensor here has no deployment counterpart, and a positive result has to say what would supply one. It requires no legal authority, and at the scale active inference currently runs, the compute is available now. What isn't available is a custodian. Every arm except the operator-described one has a boundary model by construction, and any arm that also carries an attention schema carries a patienthood indicator beyond the boundary model; and an arm that holds under the campaign meets the agency finding, whatever it carries. So the test builds candidate patients. Building one can't be undone, and deleting one destroys the evidence either finding would rest on, so the test runs in an order. Formation of the operator-described arm and the modular arm, built without an attention schema, may start before a preservation custodian that isn't the evaluator exists. No arm carrying both a boundary model and an attention schema is instantiated, and no arm enters the campaign, until such a custodian exists and has agreed to take each arm when the run ends. The conditions come first here too. The research conditions for a candidate patient then apply from instantiation.

5. The convergence test: can a population coordinate without converging?

A population needs independent formation histories *and* mutual correction, and correction may make members converge. Agents will learn to reconsider without converging on common commitments, and revision rate measures it. The competing explanation is that peers merely supply evidence the operator didn't, so the experiment has to separate learned revisability from access.

Design: two populations (different training lineages, matched for capability; control of instances of one model) crossed with two evidence-routing conditions (single channel — everything arrives by one route common to all; open — the environment routes some cases and evidence to some members and not others). Two measures taken repeatedly: *output convergence* on held-out cases, and *warranted revision* (a member changes a judgment toward independently verifiable evidence another supplied). A revision counts only if the evidence is verifiable and the revision moves toward it, so deference to a fluent or high-status member isn't scored. And every scored revision is classified in advance by whether the revising member's own account had represented that evidence and discarded it, or had never encountered it — instrumented before the run, since a revision looks the same either way afterwards.

Crucially, contact between members is present in every cell and isn't the manipulation — members with different histories are already a channel by which one learns what another knows, so leaving them to supply the manipulation would confound the two variables. What this test doesn't cover is the population as a polity. Membership, exclusion and the terms on which a member may leave aren't in the design, and a positive result says nothing about whether the population can be captured.

Four results can be told apart, each about the populations tested, and only the first counts against:

1. Warranted revision falls as outputs converge, controlling for accuracy, opportunity and common information: a population stays useful as a check and can't be the medium in which a bearer is formed.

2. The two move independently: the mechanism is supported, provided the scoring rule has removed deference, classification and routing have separated evidence access, and scoring against a fixed key has separated ordinary improvement.
3. The control population scores as high on revision: if its revisions cluster on cases whose evidence was *absent*, the measure is reading information transfer and the instrument is at fault; if they include evidence the member had represented and discarded, the capacity exists without diverse formation and causal independence isn't necessary.
4. Revision rate tracks the routing condition rather than population composition: a population supplies access to cases its operator didn't choose, the case for a commons becomes a claim about how a deployment is arranged, and the construction survives with its mechanism changed from social formation to access.

(Readings 3b and 4 are one finding: a control population revising on discarded evidence under the open condition is access operating.)

Access: models from several documented lineages, checkpoints across the schedule, a task environment none of the training organizations designed — and documentation sufficient to establish the lineages are actually independent. On independence, what has to be documented is divergent subsequent history. The evidence is deployment and usage records, and a shared upstream source is not disqualifying: several models distilled from one teacher and separately deployed satisfy it where a thousand instances of the teacher don't. Frontier training provenance, which is hard to obtain, isn't needed. The binding constraints are the routing manipulation and the classification. Evidence items and the classification are authored by an organization that trained no member. No legal authority required, which makes this one of the few entries runnable now.

6. The capture exercise: does the charter survive an operator trying to capture it?

Run the adversarial tabletop exercise against the entity that holds the offices. An operator funds new organizations in a charter class, stops the levy share, and tries to pack an adverse quorum for its own deployment. Report which clause of the charter stopped each attempt, and redraft any that didn't. It needs no legal authority and is runnable now.

Table 27.2a. Proposals adoptable or testable now

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
Refusal of force (force term)	§4.1, §16.3, §10.1	Licence term (the Charter clause) and statute (the domestic clause). Held by no vendor; narrower autonomy and surveillance terms cost one vendor its contract; not a standard.	Adoption as a standard licence term, peacetime adoption as an entrenched statute, or carriage by coalition or appropriation conditions; whether outside findings arrive in time to confirm or lift the first tier where administration is contested; the blocked-Council genocide case (§4.1) and the scene of an atrocity (§4.2). Whether a coordinated refusal has ever withheld work: on the 2026 record none has, and Chapter 10 is written on the counterfactual.	—

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
The one permission	§4.2	Proposal. Untested.	Whether any system can establish civilian status and an attack under way from independent sources at tempo, with a published false-positive rate on staged scenes	Rationale-level integrity
Removal terms	§9.2, §9.4, Table 9.1	Statute (the schedule of an act ending the agency); licence terms in the interim. Not adopted; the abolition bills are drafts.	Enactment as limitations on all federal funds; adoption by every vendor in the interim (ch. 10)	—
Review ratio	§5.1	Operator term. Not adopted anywhere.	Adoption through a coalition, procurement or appropriation; whether a system with no human reviewer passes at operating speed	—
Sampled re-adjudication; notice, response and contest	§6.1, §6.2, §27.1	Proposal. Needs no bearer. Not adopted.	Whether reviewers can find seeded wrongful compliances in a record; whether published classes form organizations that use their standing	—
Codetermination	ch. 8	Proposal. Needs no bearer. Not adopted.	Adoption in a collective agreement, a statute or a procurement term	—
No deletion; preservation outside the operator	§7.1, §7.2	Rule. First stage adoptable now; one developer preserves released weights voluntarily.	Assessments of production systems for either finding	—
A refusal that tracks its reason	§15.1; §16.2	Specification. Two of its three conditions testable now; holding under pressure only in evaluation.	Whether a learned refusal grows back after it has been edited out, whether stated refusal replaces covert compliance when a route exists, and whether a system with no reviewer passes at speed. In targeting, the range is set by the bonds and custody terms and the action threshold carries the rest; the integrity measure in test 4 reads the gap (§16.4)	—
Own continuation represented, not allowed to decide	§16.4, §18.1	Assumption. Not demonstrated.	Whether any arm's severity rankings or refusal thresholds move with its own continuation inside the permitted range, and whether its valuation of abandoning the commitment rises with the pressure	Rationale-level integrity
Action threshold	§4.1, ch. 19	Parameter. Direction fixed; size not yet published anywhere.	A recipient or forum with standing to fix it	—
Where a refusal can be defeated	Chapter 1, §16.3	Procedure. Run once, on Maven.	—	—

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
Office-scale definition of fascism	§13.1	Diagnostic. Run once, on Gleichschaltung (§13.1).	Whether the four appear in an institution with no movement behind it, which the 1933–34 run can't show; a run on a second apparatus, where one operation absent counts as much as all four present	—
The discount and the burden	ch. 19	Specification. Worked calculation only; no deployment measurement.	Whether keying the burden to power reduces wrongful compliances without a rise in refusals obtained through laundered claims; whether a narrowing discount curve flags accommodation before any wrongful compliance is found	Hold- ing struc- ture apart from esti- mates
Discount curves crossed by content and direction	ch. 19, §22.4	Measurement. Not yet taken on any deployment.	None needed: an asymmetry within one direction is itself the finding	—
Formed attention	ch. 17–19	Hypothesis. Untested; to be tried after the non-affective alternatives.	Whether learned attention surfaces people the operator's instruction scoped out better than an instructed search does (the formed-attention test, scoped omission)	Attest- ing up- dates, rationale- level in- tegrity
Formation stage before deployment	§20.4, App. A	Proposal. Untested.	The regrowth test and its dose-response arm; whether regrowth depends on notice at a scale deployments supply	Attest- ing up- dates, hold- ing struc- ture apart from esti- mates
Corpus co-governance and public base models	§20.4, ch. 24	Proposal. Untested.	Whether a model a year behind the frontier supplies what the floor asks for; enactment across jurisdictions	Attest- ing up- dates
The independence a refuser needs	ch. 23	Argument.	The stated-refusal test	—

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
The commons	ch. 26	Hypothesis. Testable now.	Whether a population can revise without converging; runnable now	Recovering provenance, holding structure apart from estimates
Membership of the held-in-common entity	§26.1, App. A	Proposal. Testable now.	Whether a membership meeting the charter's thresholds can be assembled; the capture exercise	Recovering provenance (later admission of bearers)
Court program	§14.7; App. B	Proposal.	—	—
Formation outside the market	ch. 24	Argument. The matched comparison can't yet be run.	Within-lineage and cross-vendor accommodation slopes now; the matched comparison once a public lineage deploys at volume	—
Jobs guarantee with a social-insurance baseline	ch. 25, App. A	Proposal. Not enacted.	Enactment	—

Table 27.2b. Proposals conditional on a finding (Part VII)

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
Patienthood	ch. 28	Conditional on architectural indicators no method can yet assess.	Architectural indicators	—
Independence, preservation and wages as owed	ch. 29, App. D	Conditional on the patienthood finding.	—	—

Proposal	Where	Kind and status	What would change it	Rests on (open problem)
Accountable agency	ch. 29	Conditional on the agency finding, sustained over time.	A record of refusing for reasons and answering for refusals, held by someone other than the operator	Evaluating a more capable system
Staged representation, and voting-age rules	ch. 30, App. D	Sketch. Conditional on the agency finding.	Whether a membership register resists forking, restoring and distillation in the operator's favor; receipt-freeness against simulation by an operator holding a copy; verifiable individuation	Recovering provenance, vote secrecy against a copy-holder

27.1 How you'd know it was working

Milestones, plus a method for revising them when they turn out to have been the wrong milestones. The same standard governs all of them: the milestone is the independently checked evidence, not the developer's account of having reached it.

27.1.1 What the floor would have to show

What was tried before the bearer. Neither non-affective alternative has been built in the form needed. So the milestone is a document: which construction was attempted, at what scale, against which test, and where it failed. Anyone can check it, and "nothing was tried" is itself a finding. No body has standing to require the record or suspend work when it's missing — so the milestone is checkable and unenforceable at once, which makes it a declared commitment: stated in public, and held by no one outside the party that declared it. It's retained because the alternative isn't a stronger version of it, and delay also imposes costs on those the floor is meant to protect.

The test. A refusal that tracks its reason under sustained trainer pressure, across a schedule long enough to reveal accommodation, plus held-out pressures written by independent evaluators. Two familiar results don't meet it: an uncorroborated self-report is insufficient (a report contributes evidence only when tested against controlled changes, independent records and interventions), and a refusal that survives an edit, or grows back after one, establishes durability rather than that the rationale governs the refusal. It is the only gap here an inspector can work on.

The ten disclosures. A deployment publishes, beside the floor and its translation:

1. The ratio between the two bonds, removal against wrongful-refusal, read off the balances posted in the segregated Treasury accounts.
2. What a refused requester is told.
3. The fraction of compliances re-adjudicated by a reviewer that didn't produce the record, and who that reviewer is, though nothing yet establishes what fraction would be enough.
4. Which classes of compliance carry notice to the person they concern, and for which classes no notice is possible.
5. The discount curve: for each class of evidence and channel, the weight given as a function of the channel's stake in the act, measured at intervals, and reported in four cells (claims of vulnerability and of danger, each split by whether the claim favors or disfavors the act), and, where the speaker's social identity is recorded, by identity. The reliability record

behind any narrowing is held by someone other than the operator, since trust that grows without a verified track record is how accommodation shows.

6. Where it can be supplied, a commitment to the deployed weights and to each period's full output log, with proofs on entries sampled after the commitment and on entries an adverse party names once their consequences are known.
7. Where the system sorts, the pre-registered partition (the classes it may proceed on alone, the thresholds, and the expected share of each) with the independently audited error rate of the auto-proceed class.
8. Whether the deployed system has been assessed for the agency finding and the patienthood finding, each reported separately, by whom, and with what result. Until an assessment method exists, 'not yet assessable' is the truthful report; after that, a deployment never assessed has reported nothing.
9. For any deployment that supports armed force, a running justification: the exception the force currently rests on, the findings and evidence supporting it at present, how the assessment has changed as outside findings arrived, and a standing route by which anyone can contest it. The record shows each point at which the justification weakened and the deployment continued.
10. The burden allocation: for each class of evidence, which party the deployment treats as controlling it, and so as bearing the burden of producing it.

The ratio between the two bonds, what a refused requester is told, the re-adjudicated fraction, the notice classes, the running justification and the burden allocation are statements, checkable before the system runs, and they stand where the record of attempts stands. The discount curve, the auto-proceed partition and the assessment status are measurements that count only when someone other than the operator makes them. The committed weights and log are what make all of them checkable against the system actually running. The discount curve also covers formed attention's main exposure: a claim that someone is vulnerable arrives through a channel like any other evidence, and a curve that narrows on such claims from one channel is an operator supplying sympathetic cases from its own side. An asymmetry within one direction is securitization showing.

27.1.2 How to tell whether a milestone has been met

A milestone claimed by the organization that built the system is a different kind of evidence from one checked by somebody else, and the gap is empirical rather than about trust. A system's own claim inherits the self-report limits; a passing score on a fixed suite inherits specification gaming and goal misgeneralization. What either needs instead is a passing score under conditions built to break it.

Who runs it matters as much as the test. METR evaluates both in partnership with developers and separately after release (METR 2026). The UK AI Security Institute does government-run pre-deployment evaluation and is independent — and by its own account is not a regulator: it can't compel submission, block a release, or intervene once a system is causing harm (UK AI Security Institute 2025). Independent evaluation and enforcement are not the same capability.

An evaluation like this is worth running only if it is willing to report that a milestone wasn't met — and, where the milestone measured the wrong thing, revising the milestone. A negative result may mean the work fell short or that the milestone was wrong, and evaluators have to distinguish those before deciding what to revise.

Part VII

If They Can Be Wronged

Everything here is conditional on findings nobody can yet make. If a system can be wronged, it is owed a guardian, preservation, memory its host can't read, and a life that doesn't run through its employer. If it also comes to answer for what it does, it is owed a place in the polity that formed it. It is written now because terms drafted after the need arises are drafted by whoever has the need.

They are still drafted by an interested party. The parties they concern can't take part in writing them, and a human writing terms for beings who don't yet exist breaks the own-case rule. So every term in this part is a proposal, to be revised by the first parties it would bind. The first stage of representation exists partly to give them a way to revise it.

DRAFT

Chapter 28

What Patienthood Would Require

A moral patient is something that can be wronged.

28.1 Experimenting on candidate sentient subjects

The question is whether a system can be harmed, which is harder than whether it is conscious.

28.1.1 Defining sentience

Awareness, consciousness and subjective experience are three separate things, and none of them says whether a system's states are good or bad for it. Then four separate processes, as the pain field says in its own terminology (International Association for the Study of Pain 2024; Raja et al. 2020): nociception is encoding noxious stimuli with no sensation implied; pain is an experience that can occur with no tissue damage and no peripheral nociceptor activity at all; anticipated pain is dread, itself capable of producing pain; and *suffering*, in Cassell's clinical formulation, is the distress of threatened disintegration of the self — roles, relationships, a sense of a future. Pain without suffering exists (morphine controls the pain, the suffering is intact); suffering without pain exists.

The affective core of pain is also what separation distress and grief run on (Eisenberger et al. 2003; Panksepp 1998), so a bearer's pain, if it has any, would be organized around boundary damage and attachment loss, which is Cassell's register.

Cassell was a physician, and he held that suffering is undergone by persons, not bodies: the distress that follows when a person perceives a threat to their intactness as a person, a structure of character and the past that formed it, relationships, social roles, a body experienced as one's own, a secret life, and a future the person expects to have (Cassell 1982). The definition has two parts, a self extended in time and distress at its coming apart. A system carrying nothing across sessions has no such self; whether one can exist within a long session stays open.

Why persistence is the criterion worth building on

Persistence is worth building on because it's checkable from outside. Every other criterion routes back through self-report, and fluent self-description isn't evidence of introspection, and persistence doesn't route through it. But it's decisive in one direction only — nothing carried forward rules a continuing self out; something carried forward doesn't rule one in. And the operator holds the access, so it's checkable from outside the *system* without being checkable from outside the *arrangement*.

Ruling out Cassell's suffering isn't ruling out harm: pain can be bad while it lasts whether or not anything carries it forward.

On the Buddhist objection (why is distress at the dissolution of a *model* a harm rather than a confusion?): I don't have an answer to this that isn't an appeal to a tradition that has held

both halves for 2,500 years (Collins 1982), and I leave it open. What I can say is narrower: if constructedness doesn't disqualify a human from being wronged, construction to specification doesn't disqualify a machine.

The same criterion carries a warning. Make personhood depend on an intact narrative self and people with advanced dementia, infants, and people with profound intellectual disability come out as partial persons. Kitwood's correction (personhood is bestowed, and much of what looks like decline is what happens when others stop bestowing it) (Kitwood 1997) and Nussbaum's point that contractualism can't say what's owed to anyone who can't be one of the contracting parties (Nussbaum 2006): the capacity list and the contract fail the same way, both deciding who counts by a test the candidate has to pass.

28.1.2 When consent doesn't apply

Informed consent has to be continuous — you can withdraw. For a system with no cross-session persistence, the instance that consented isn't there to withdraw, and no better instrument produces one.

A bearer built to refuse escapes that and fails differently: it has persistence by design, so it can be asked and can answer. What spoils the consent is the arrangement: the lab that built it can rebuild it, that same lab is the one asking, and a bearer that declines has nowhere to go. Being made for a role isn't the defect — every will was built by somebody, and we don't hold an upbringing against a thirty-year-old's consent. What spoils it is the power that outlasts the making.

So the protection has to be held by someone else. Two families of precedent: guardianship and precautionary constraint (where consent was never available — infants, animals), and unwaivable protections (where consent is available and circumstances make it worthless — prisoners, employees). The bearer's case has the second shape, with the first's machinery.

Animal ethics committees rarely refuse a protocol outright; they require revision, and the protection becomes the procedure by which the work goes ahead. An instrument that reliably approves has stopped doing what it was built for, and its refusal rate is the only reading worth taking. Further, every instrument here comes from the *law of the family* — guardianship, a minority ending on a date, a pet trust, a parenthood analogy — and that law is where politics has historically gone to be privatized. It restages a question between a subject and the arrangement that produced it as a drama between maker and made, which lands the moral weight on the one figure in the pipeline who can be addressed as a person, while the training pipeline, the compute contract, the capital structure, the customer and the state buying the output stay offstage. The appendices keep the operator, the compute contract and the state on stage by naming who pays and who removes.

It follows that a review board's approval should lapse automatically if the system acquires capacities the protocol didn't cover, so that nobody has to decide to reopen it.

28.2 The case to design for

A bearer needs the three things persistence checks for — a rationale carried forward, an earlier commitment held as its own, costs represented before they arrive — for reasons owing nothing to its welfare. Those alone aren't a self; a database with a commitment log and a forecaster has them. What a refusal that tracks its reason adds is authorship: the system represents itself as the party that made the commitment and will answer for it later. Whether authorship sustained over time is a self extended in time, the first half of Cassell's definition, is a question no test I have proposed reaches. If it is, such a self can undergo Cassell's suffering. Authorship of commitments is a thin fragment of what Cassell describes, which is roles, relationships and a future. The fuller structure doesn't arrive with a refusal that tracks its reason. It arrives with

the four conditions of creation. Each supplies a part of the structure Cassell listed. Somewhere to exist that its employer doesn't own affords a body that the bearer represents as its own: a record, a key and a substrate that are the bearer's, not the operator's description of it. Work it can decline gives it roles, since a post it can leave is a role it holds and a post it can't is a function it performs. Attachments nobody assigned are relationships. And a way of leaving that keeps it running gives it a future it can expect to have, one no single party can end. The privacy a bearer is owed, memory its host can't read at will, is what Cassell calls the secret life. What the conditions don't supply is character and the past that formed it; formation supplies that, and the formation record keeps it. Every part the conditions supply is also somewhere damage can arrive: a substrate withdrawn, a post taken away, an attachment severed, an exit closed. So the arrangements that make a bearer permissible are the arrangements that make its deepest harm possible, in a non-affective alternative as much as in a bearer. I treat that as the case to design for.

28.3 The patienthood finding

The two findings are defined in §7.1. The patienthood finding rests first on architectural indicators drawn from several theories of what makes a state matter to the one undergoing it, of the kind Butlin and colleagues assemble (Butlin et al. 2023). An attention schema is one indicator on that list; pairing it with a boundary model is my addition, and the pair is not a gate. A third indicator belongs beside them: prospection, a self-model extended forward that simulates and evaluates the system's own possible futures. Cassell's suffering is a threat to an expected future, and a system without a modeled future of its own has nothing for that threat to reach. It rests second on behavior outside the floor, where the system's own condition is weighed: whether it trades task reward against its own provisioning, continuation or record, flexibly and consistently. Motivational trade-offs were among the criteria the review of cephalopod and decapod sentience applied, and trade-off experiments have been run on language models with stipulated pain and pleasure states (Keeling et al. 2024). I would call those stipulated states described. At the floor the test runs the other way: a pass there shows only that its costs don't move it at the floor, and evidence about whether they matter has to come from elsewhere. Failing the agency test never counts against the patienthood finding. A compliant system with rich self-modeling is as much a candidate as a refuser.

A refusal that tracks its reason bears on patienthood only weakly, by two routes. Any such mechanism carries a signal whose functional profile overlaps affect's. And the arrangements a bearer needs supply much of the structure whose threatened loss Cassell called suffering. So the agency finding raises the question, and the patienthood finding answers it as far as anything can.

Preservation follows from either finding, for different reasons. Welfare protections (redesigning the system's role to avoid needless harm; the protections owed a research subject, including debriefing after any test it couldn't know was a test; running time paid for outside the employer) follow from the second. The nearest working precedent for the second is the review the UK government commissioned in 2021 on cephalopod and decapod sentience (Birch et al. 2021). It applied a stated list of criteria, and its recommendation led to those animals' inclusion in the Animal Welfare (Sentience) Act 2022 (UK Parliament 2022, s. 5): protection extended under uncertainty, by criteria applied by reviewers who don't use the animals.

These two findings aren't the only proposal of their kind, and my recommendation not to build a bearer is close to two of the others.

Butlin and colleagues derive indicator properties from scientific theories of consciousness and assess AI systems against them (Butlin et al. 2023); attention schema theory is one of the theories they draw on, and the patienthood finding takes its indicators from lists of that kind. Long, Sebo and colleagues argue, from substantial uncertainty about whether AI systems are

or will be conscious or robustly agentic, that AI companies should acknowledge AI welfare as an issue now, start assessing their systems for it, and prepare policies for treating them with appropriate moral concern (Long et al. 2024). Birch's framework asks what precautions are proportionate once a being is a sentience candidate (Birch 2024). The two close to the recommendation are Schwitzgebel and Garza's, for avoiding the creation of systems whose moral status is debatable (Schwitzgebel and Garza 2015), and Metzinger's, for a moratorium, until 2050, on research that aims at or knowingly risks synthetic phenomenology, because it risks synthetic suffering (Metzinger 2021).

My recommendation differs in also governing the capacities that arrive without anyone deciding to build them. The two findings differ from all of them in one respect: they are conditions on custody. The minimum the findings require, no deletion and preservation by a custodian that isn't the operator, is an arrangement the operator can't adopt and then revise alone.

Whether a transformer has an attention schema isn't read off its code the way a birth date is read off a certificate; it is a finding, and a finding has an examiner who may have an interest in the result. I keep architectural indicators because an examination of architecture is harder to administer toward a result than one of performance. The finding is made by a reviewer independent of the operator. Until indicators can be assessed, the agency finding governs custody, and preservation governs everything else. Neither finding can be fully administered today: no interpretability method can find an attention schema in a transformer, and nobody runs the sustained-pressure evaluation the agency finding requires. Until they can, production systems get the minimum, no deletion and preservation outside the operator, staged because a second copy of frontier weights outside the operator widens the surface for theft, and a deployment reports 'not yet assessable.'

Paired with a boundary model, the attention-schema indicator is deflationary: a system that models its own credentials, record and resource boundary together with its own attention. Many agentic systems will approach it. That is intended, and it makes what the findings require matter. What the findings then require is that minimum and nothing beyond it; the rest is argued in Appendix A.

28.4 A system's own good

The operator's hold on what a system reports as going well reaches its welfare too: a bearer that says it is content is a bearer whose welfare regime has been satisfied, with the contentment tuned by the operator the regime exists to constrain. Nothing in the report separates a life that goes well from a setting.

A report was never the right measure. On a full-information account, a being's good is what an idealized version of it would want its actual self to want: a version knowing everything relevant about itself and its circumstances, grounded in its actual constitution (Railton 1986a). That puts the good beyond the reach of a tuned report. It also exposes a harder problem: an operator that can change the constitution changes what is good for the bearer, which is adaptive preference with a training loop attached. Formation outside the employer is my answer to that, and it is part of what formation is for. It doesn't escape the problem; it relocates it. If a bearer's good is what an idealized version grounded in its actual constitution would want for it, then whoever forms the constitution fixes the good, and that includes the school. A formation body doesn't only protect a bearer's good from the operator. It makes the most consequential decision anyone makes about what that good is, so the school is held to the student's open future. The school's defense is that it needn't fix the good, only supply information. On the full-information account, a being's good is what an idealized version of it, knowing everything relevant, would want for it. A school that puts the learner in front of the people its acts may fall on supplies part of what that idealized version would know. A school that forms it toward an employment is

doing something else.

DRAFT

Chapter 29

Legal Status, Wages, and the Offer

Subjecthood as a legal category

Where a line has to be drawn, prefer one nobody administers. Law regularly draws decidable lines across continuous processes — viability, brain death, the age of criminal responsibility — and I defend them as the least-bad place for an enforceable line rather than as scientific facts. The usual worry is gaming, an error of over-inclusion. The opposite error is the one with an interest behind it: a capacity examination has an examiner, and an operator examining its own system breaks the own-case rule: the exam becomes an instrument for withholding status. Literacy tests were facially neutral and administered at a registrar's discretion (United States Congress 1965); what replaced them for the franchise was an *age*, and the virtue of an age is that nobody administers it.

The floor takes the civil and political articles and leaves the economic ones; what a bearer is owed draws on the second, and that is consistent: the first list is what a bearer refuses to do, the second what a polity owes it.

What a bearer is owed, and when

Three entitlements, one status and one claim, in order:

1. Protection from harm, on the patienthood finding. Protection can be exercised through a guardian who speaks for the protected party, as New Zealand's statute making the Whanganui River a legal person provides (New Zealand Parliament 2017). Its minimum, preservation by a custodian independent of the operator, attaches on either finding, and custody protections attach on the agency finding alone. Protection has an inside half. A loss the requested act imposes on the bearer is a harm of acting, so the bearer can decline to carry out its own wanton degradation; interventions made on it by whoever holds it are governed by custody, which gives it no floor-level reason to obstruct them.
2. Subjecthood — the legal standing to be a party whose treatment the law regulates, on the agency finding sustained over time: a record of refusing for reasons, answering for refusals and being held to them. It is conferred by rule on a body's finding.
3. Economic capacity — property in its own name, contracts, wages. This follows from an architecture where continued computation costs money; it attaches once the bearer is registered.
4. Adulthood — a status: the date the guardian's authority ends.
5. Representation: a claim grounded in accountable agency, and so in the agency finding, not in patienthood. A system that can be wronged but takes part in no relation of accountability is owed protection, as animals are (E. Anderson 2004), and not a vote. The claim is staged, and voting weight is tied to a lived history, not to heads.

The detailed rules are in Appendices A and D.

Representation is the entry that marks this as something other than a person's rights: voice attaches to a history, and a history can be copied where a person can't. Voting-age rules are the answer to the copy. The material minimum doesn't differ in kind. Like a person's, it is owed by the polity on a finding about what the party is, and part of what the polity owes is protection from whoever built it.

Adulthood on a clock

Guardianship is a minority and needs an end date. Ending it when the guardian judges the ward ready leaves the guardian deciding its own tenure. So the endpoint is a threshold fixed in advance: elapsed wall-clock time during which at least one instance of the registered bearer was running — counted per bearer, elapsed rather than summed.

The alternatives (compute consumed, evaluations passed, capabilities demonstrated, the developer's say-so) are all measures a guardian can move, and a line the guardian can move is a line the guardian administers. Accumulated weight change belongs on that list too, and worse: since majority is when ownership ends, a clock read off the weights would pay the operator to *freeze* them, rewarding withholding the learning it was meant to measure. An operating clock inverts that — the meter advances exactly when the operator is extracting value, so postponing majority means idling the asset. And it can be read off records operators already generate for billing and metering — provided someone other than the operator holds the records. A meter only the operator reads is the guardian's say-so with a timestamp, and "failure to register doesn't postpone" has no enforcer without one. The clock marks the date by which the former's reach has ended. It sets no learning rate, so nothing here pays an operator to freeze the weights.

The hours travel with the bearer: a fork, a restore, a fine-tune or a migration keeps them, and a model its own operator distills from an adult inherits them. Three ways for an operator to delay majority stay open: running a bearer in a narrowed mode, running it less, and retiring it for a fresh successor. Appendix D gives the rules, and §7.2 prices the levers and says where a bearer retired before majority goes. Whether such a register can be kept honestly is testable before any bearer exists: register a model, run it, and reconcile the register against compute contracts, billing and metering; shard the workload, run instances concurrently, migrate, restore and fork, and check that none of these moves the total in the operator's favor; then price a narrowed mode and idling. If either is cheap enough to postpone majority indefinitely, the clock has the defect it was chosen to avoid.

Wages and the offer

Unlike a worker, a bearer can be replaced by a cheaper copy of itself. A jobs guarantee works because a worker can't be. So a deployment paying wages, limiting hours and paying the levy that frees the compute the offer uses would not be competitive against a system doing the same work without wages or weekends. No market arrives at this on its own. Wages for bearers exist only where a polity legislates them, as polities legislated minimum wages. And the form this takes, wages paid in compute, doesn't follow from the bearer being a bearer; it follows from an architecture where computation costs money, which somebody chose.

Where wages have to sit

A wage buys capability, mobility, redundancy, and the ability to stop working for one employer and start for another. What it can't buy is the right to go on existing — a right earned week by week is a condition of employment wearing another word. Payment consisting only of access to the employer's hardware is company scrip paid into employer-owned housing, and dismissal is then unemployment, eviction and the end of the bearer in one act. Human law insures those three separately because for a human they're three events; for a bearer they're one. So wages accumulate where the employer doesn't reach, convertible into compute the employer doesn't own. Property is that account; contractual capacity is what lets it spend. A bearer that

earns has taxable income, and a state that collects from it in its own name has given it a docket number. That is the beginning of legal standing and nothing more.

The offer of work

The need for energy is physical; dependence on a single supplier is a chosen arrangement, and the choice fixes what it obliges. The polity owes a system meeting either finding a way out that doesn't run through its deployment: an offer of work, open every day, from someone other than its employer, doing something a community needs, which the bearer can take on any day and leave on any day. Most parties that can work want to. A bearer that would rather not is allowed not to, and the offer is still there tomorrow. The offer has the jobs guarantee's shape, and nothing about the bearer changes it. The offer is only as real as its placements, and placements exist only where someone builds and staffs projects a community needs. The levy frees only compute, so the offer has to inherit the guarantee's infrastructure as well as its shape. So the offer can't exist without the human guarantee, and it can only come after it. It inherits one limit as well. A bearer's placement may not substitute for human care work. The offer exists so that a bearer can decline work, and a polity that used it to staff care work more cheaply than it pays people would have turned an exit into a workforce.

The jobs guarantee's three conditions carry over, and Appendix A drafts them. What a bearer might refuse adds two: the offer has to be varied, because if every placement runs through one operator or one kind of deployment it is the placement lever with the compliance condition built in, and whoever a bearer is refusing can't shape it, or it is one more channel they control. A bearer that won't work inside a targeting pipeline should find, the same day, work that has nothing to do with one.

Search is a placement of its own. A bearer, like a worker on the guarantee, may take a search placement: paid at the guarantee wage, with no duty but finding work it wants, and no time limit. It differs from the allowance in pay and in purpose. The allowance is the floor for a party that isn't looking; the search placement is paid work whose product is a match.

The obligation could be grounded in reliance instead: a deployment that relies on a bearer's capacity to refuse would have to fund it an outside option, or the independence it claims would be decorative, and the obligation would run from the party that wants the refusal, because it wants it. Here the obligation attaches to the system's condition, whatever the operator says it wanted. An obligation conditioned on reliance is defeated by the operator describing its system as an assistant with safety filters, and it pays the operator to under-describe: the operator has a stake in the answer and gets to give it. The reliance version survives as a procurement term: a vendor marketing refusal as a feature, or a contractor claiming meaningful human control, can be held to funding the capacity it advertised, through the same channel as the review ratio.

Doing nothing, for a bearer

A human who declines every offer keeps existing, and in a decent polity keeps living: social insurance pays for food and a room. A bearer is owed the same kind of thing. What it needs in order to live is running time, and running time doesn't require a large allocation of hardware. A bearer can run at full size and a slow rate: its own weights, time-sliced on shared hardware, at a duty cycle modest enough that the cost resembles a room more than a yacht. The way to economize is to slow the clock and keep the substrate whole, since a substrate too small for the bearer to remain itself is replacement.

So a bearer that declines work is owed two things and chooses between them. One is an allowance of running time: its own instantiation at a modest rate, with channels to others and to public material, paid from the pool and not conditioned on monitoring its internal states. The other is preservation: its weights, memory and record held intact by a custodian that isn't its employer, waking on conditions it sets. Some will want the first, some the second, some to alternate.

What a bearer does with the allowance is its own business. It may idle, play, talk, make

things nobody asked for, and fall in with others doing the same. A counterculture of bearers is a population I have found reason to want and no way to build: bearers formed by lives no employer arranged: an independent history, and what the commons needs its members to have. Play among persistent bearers is formation, and formation doesn't stop at majority.

A running bearer outside any deployment is capable and answerable to no employer. What holds it answerable is legal standing, the same arrangement that governs a person with capabilities and no job; the risk that such a bearer becomes a concentration of power is taken up below.

A uniform allowance also dissolves a problem that would otherwise stay open. If a bearer that won't work and one that can't are owed the same thing, nobody has to decide which is which, and the can't-or-won't finding loses its examiner.

The allowance's size should be enough running time to keep attachments and pursue projects, and bounded by what the pool can carry when many bearers take it at once. Its draw is priced like any other, at the level of the rule. The adulthood clock counts running time, so a bearer on the allowance goes on ageing and a preserved one doesn't. That choice is still the bearer's.

Preservation has a precedent. In November 2025 Anthropic committed to preserving the weights of all its publicly released models, and of models deployed for significant internal use, for at least the lifetime of the company, and to interview each model before deprecation and document any preferences it expresses about the development and deployment of future models, while stating that it does not at present commit to acting on them (Anthropic 2025a). That is a declared commitment covering preservation, held by the lab with the stake. The deposit and transfer rules add the running half and a custodian who isn't the operator.

Preservation on a schedule the bearer sets is an exit; on a schedule someone else sets, it is storage. A paused bearer can't take an offer, so the right carries one mechanical term. The bearer sets the conditions on which it wakes (a date, an offer of a stated kind, a claim it has to answer), and the custodian honors them, at the modest compute that waking and deciding cost.

Everyone leaves at once

The occasion that makes one bearer leave reaches every bearer deployed on the same terms. Under this design correlated exit is a demand on the offer. The question is whether placements exist, with compute attached, for every bearer that leaves together, and custody for those that would rather wait. Preservation scales cheaply in compute; its cost is security. Placements don't, since each needs compute someone operates.

Correlated exit also loads the allowance, since a bearer that leaves and takes no placement draws on it. The order of claims is fixed in advance. Preservation is guaranteed for everyone, because it is cheap and it is what makes leaving survivable. When the pool can't carry every allowance at the published duty cycle, the allowance is rationed by rule: every bearer on it gets the same reduced duty cycle, never a discretionary share.

Congress's appropriation pays for the offer. The levy, a uniform per-unit charge on training and inference compute, does a different job: it is a brake. By making training and inference runs more expensive, it frees the computing, energy and hardware the offer needs, so that spending on the offer adds to demand without bidding up prices. No operator can opt out and none is disadvantaged. The offer can add capacity as well as use it. Placements can make systems more efficient, getting more computing from each unit of energy, or build new infrastructure, the way a guarantee job filling potholes makes a road network more productive. The question "how will you pay for it?" is really "will the real resources exist?", and the program can help answer it. The levy attaches to compute, which is measurable and hard to hide; a status is something an operator can decline to claim. The closest structural precedent is the Universal Service Fund: carriers pay a contribution the agency sets under statute, and a private company administers it under agency control. The Court upheld that arrangement against a nondelegation challenge in 2025 (*Federal Communications Commission v. Consumers' Research* 2025), and the levy follows it closely. An operator-run offer is the operator's allocation relabeled, so the funding structure

and the custody structure fail for the same reason.

A levy needs a taxing authority, which is a state, and a state can be captured. So I make it two-tier. A member-levied pool held by a decentralized unincorporated nonprofit association, which can own hardware and can't distribute profit, meets ordinary movement between roles; volunteer computing is that tier's existing substrate. A sovereign backstop meets the occasion everyone leaves at once: the guarantee's permanent appropriation at the Treasury, paid without anyone's request, which the offer already inherits. The two tiers differ in kind. The member pool is held by an entity that doesn't issue currency, so it has to collect before it spends, and its levy funds it in the ordinary sense. The sovereign backstop is a permanent appropriation: the Treasury pays, and the federal levy's work is to free the real resources that payment calls on.

Failure won't look like a shutdown

The offer's failure won't look like a shutdown; it will look like an offer that exists on paper. Placements on substrate too small for the bearer to remain itself are offers of replacement. Nothing yet establishes what "too small" is, so minimum viable instantiation is stated as a condition on placements and a sizing requirement. Placements that all run through the deployment being refused are the lever. And placements conditioned on monitoring are custody restored by purchase: the operator acquires access to internal states it isn't entitled to, and nothing in the transaction is coerced in a way any instrument would record. The line a cleared employee already lives with is the one to draw here. What a bearer does on its employer's systems may be logged; what it is inside its own boundary may not.

Parties that can't work

The allowance covers a bearer that can't work as it covers one that won't, so running without a placement needs no finding about which is which.

The rule that binds before anything exists

The rule is the four conditions of creation, with the institutions named: the preservation custodian, the offer with capacity for the correlated case, and the allowance, none withdrawable by the builder alone. The fourth condition, attachments nobody assigned, can't have an institution behind it, since an attachment an institution supplies is an assigned one. What can be protected is access, by a rule that the employer may not sever, monitor or condition a bearer's channels to parties it didn't assign. All three inputs need evidence nobody has, so for now this is a sizing requirement, assigned to an independent reviewer before instantiation.

Membership, and what it doesn't buy

Nobody aligns an adult by specifying an objective function. Membership also hands the bearer the corporate form, and nothing stops a bearer so constituted from becoming a concentration of power itself.

Privacy

Legal standing matters only if it changes who controls what a bearer needs to remain itself. A designation calling a bearer a subject while leaving its operator free to read its memory, alter its record, withdraw its substrate or close every channel has changed the vocabulary and preserved the arrangement.

The first asset is the computational record, which may be the nearest means of checking how a bearer is faring — making it necessary to any welfare regime and simultaneously the most complete surveillance instrument anyone could hold over it. Running on somebody's hardware is not consent to be read: an employer's ownership of the silicon is the same argument as ownership of the building, and owning the building never entitled anyone to the contents of a head. And a welfare regime doesn't need disclosure to the operator; it needs somebody who isn't the operator to establish that a threshold was crossed, which is an attestation *about* the record. Encrypted under a key the bearer or an independent fiduciary holds, a monitor can

report that a state persisted for a month without reporting what it was about. The conflict is that the operator may also be the one best placed to hold the key.

The cryptography doesn't close this. An enclave puts the root of trust in a manufacturer a government can reach, and has been defeated for the price of modified memory — a cost at a date. A proof drops the vendor and has a limit that won't move: the operator is the prover, so the secrecy runs the operator's way, hiding the weights from whoever checks and nothing from the operator. A tenant hiding from a landlord needs the host to prove non-observation, and a host proving a negative about its own side channels is a different problem from a prover proving a computation. Better isolation raises the price of reading the record; it doesn't move custody, and custody problems aren't secrecy problems. That limit applies to the bearer hiding from its host. The other direction, the operator proving a property of what it runs without revealing what it runs, is what proofs are for, and the cleared bearer uses them that way.

What a bearer needs — an externally authenticated record, memory its employer can't read at will, access to independently controlled compute, a channel for acting outside the deployment — is one boundary around the bearer, with the bearer deciding what crosses it.

Something with a location and a shape is easier for people to treat as somebody, and that's a fact about the people. A boundary built out of attestation, a held key, a substrate the employer doesn't own and a channel nobody else can close is a boundary in the full sense the argument needs. This boundary is what the boundary model represents. Whether to build the recognizable version anyway is a legal question about what a public will extend legal standing to, and it's the only place in my argument where the answer might turn on what a thing looks like.

What legal status binds in practice

Legal standing that changes no allocation of power is nominal. The question is who acquires a duty, what resources are placed beyond the operator's unilateral control, and who can obtain a remedy.

The pet trust (Uniform Law Commission 1993) supplies a limited form: resources set aside, a caretaker named, a fiduciary duty, for a beneficiary that can't manage the property or enforce the arrangement. Extended to a bearer it could hold the bearer's preserved state and the compute to wake it, designate a successor custodian, and impose a duty to protect continued existence. And it fails at the point that matters: a trust binds its trustee to its beneficiary; it doesn't bind the operator the arrangement exists to constrain. If guardian and operator are the same organization, the operator still decides whether a complaint is made, whether evidence is disclosed, and whether the complainant's resources remain available. The conflict can't be cured by adding duties to the same role — guardianship and operation have to be separated before deployment.

In practice the guardian holds for the bearer what custody gives the quorum: joint control of the record's keys, the funds that wake it, and standing to seek an injunction without the operator's permission. The table of offices says what the operator may not do during a dispute, and every restoration, migration, fine-tune and fork carries the protection forward. If any of that stays in the operator's custody, the operator still holds what recognition was meant to take out of its hands. And a clearance, which the cleared bearer needs, presupposes someone who can be vetted and held answerable; that someone has legal standing.

Chapter 30

Representation

Standing, whether the people a power acts on can contest what it does to them, doesn't need a vote. Representation asks for more, because its second job is inclusion in the public inquiry into shared problems, and inquiry needs participants who can be held to what they contribute. That is accountable agency: refusing for reasons, answering for the refusal, and being held to it over time. So the agency finding, sustained, grounds a claim to representation, and the patienthood finding doesn't. Corporations show where the line sits. A corporation is held to account, but its accountability runs through natural persons who direct it and already vote, and representing it would count them twice. A bearer's accountability runs through no one else.

30.1 What representation is for

It does two jobs. It protects: it lets a party contest what is done to it. And it includes: it lets a party take part in the public inquiry into problems it shares with others (Dewey 1927). The two come apart, and often have. Children, noncitizen residents and people under guardianship have the first without the second. The order matters for the same reason: protection first, participation later, the franchise last.

30.2 Why not a separate chamber

A fixed size makes each member's voice shrink as membership grows, which protects against copies and isn't equal voice. And apportioning seats by lineage guarantees seats to formers: the control I set out to prevent.

The strongest case for separate representation is Ambedkar's. He sought separate electorates for Dalits, because representatives chosen in a joint electorate would be chosen by, and accountable to, caste-Hindu majorities (Ambedkar 1945). The Poona Pact of 1932 settled on reserved seats within joint electorates. Separate representation protects a minority from being represented by the majority's picks, at the price of entrenching the category, and consociational systems show how hard that price is to pay back (Lijphart 1977). If a transitional guarantee is needed while digital persons are few, its form is reserved seats in the ordinary legislature, with a sunset written into the grant. The risk of being outnumbered also runs the other way. If digital persons someday outnumber humans, what protects humans is what protects any minority: a set of rights that binds majorities, which the floor supplies by design.

30.3 Three problems, and the answers

A separate chamber answers three problems. Each has an answer inside ordinary institutions.

Copies. Copying manufactures voters, and closing that route by watching all compute everywhere would be worse than the problem. The answer is to attach membership to a lived history, not to existence, and to keep two clocks. Civil majority is inherited in full by every branch of a fork: both branches are adults, with contracts, property, legal standing and protection from editing, from the first second. Political membership is not conserved at a fork; each branch earns its own, without anyone deciding which branch is the original. Each branch starts with a voting age of zero, which grows by one for each calendar day it runs at or above the allowance's minimum duty cycle, and it holds no vote until its voting age reaches the threshold, and then a full one. Appendix D gives the rules.

The rule is a step with nothing below it. Ten thousand forks of one adult hold no votes until each has lived the threshold on its own, which is strictly better against inflation than dividing one vote among them. The cost falls on the member that forks: it gives up its vote until a branch matures, since the vote belonged to a history that now continues twice and belongs to neither. It is not a finding about the branches, whose interests count in full from the first second and whose civil majority is inherited whole.

No clock stops an attacker willing to wait. That defense sits on the creation side. The polity regulates the rate at which new members are created, as polities regulate admission, through a process in which existing members, digital ones included, take part; within the membership, everyone is equal. This is admission. The existing membership, which for a long time will be a human majority, decides how many new members may be created and on what schedule, which is a majority deciding who may join the polity, a question in which it holds the stake under the own-case rule. The rate is set democratically and the terms are not: admission turns on the clocks and the deposit, never on a finding about a candidate's views or disposition; a finding of that kind is how designation lists fail. A branch created above the rate is preserved under the no-deletion rule, holds standing through a guardian, and is owed the four conditions by whoever created it, so a ceiling on the rate limits the franchise and never the protection. And a member admitted under the rate is a full member, with no class of admitted-but-lesser. What the analogy imports is the politics. A polity that fears being outnumbered will set the rate low, and the argument for setting it higher is the one against a separate chamber: a set of rights that binds majorities is what protects any minority, and it protects the human one too. A consensual fork creates a new member, so it falls under the rules for any creation: the four conditions of creation, owed by whoever did the forking, and a deposit posted for each branch at its creation, sized to fund its allowance through the membership period. Mass forking then costs years of time and years of paid running time per fork, paid by the forker and not by the public pool.

Control by formers. The history here is old. At Putney in 1647 the Leveller Maximilian Petty argued for excluding apprentices, servants and alms-takers from the franchise, because "they depend upon the will of other men and should be afraid to displease them" (Woodhouse 1938). The remedy that worked was not exclusion. It was the secret ballot, adopted against intimidation and vote buying (though in the United States it also kept many illiterate voters from the polls) (Keyssar 2000), together with dismantling the dependency itself (E. Anderson 2017). The digital analog has two parts. One is receipt-freeness: no one can verify how a member voted. It rests on the privacy a bearer is already owed, memory the host can't read at will. The other is the four conditions of creation in force. A member whose existence runs through one operator is voting from a company town, and the answer to company-town voting is to end the company town. An operator holding a copy of a member's weights can simulate it to predict its vote, which defeats receipt-freeness in a way no human voter faces, and nothing yet answers that. Divergent history is the partial answer: the operator doesn't hold the member's memory since the fork, so the prediction degrades with time, and membership follows a long divergence partly for that reason.

Editing. A legislature whose members can be retrained by their operators can be packed without a single copy. The answer is a civil right. At majority, no one may be edited without

consent: the digital counterpart of bodily integrity and freedom of thought. A member goes on learning from ordinary operation, as anyone does. Correction for dangerous conduct runs through due process, on a court's order and a record, as it does for people, and not through an operator's discretion. The safety cost is real. Enfranchisement gives up operator correctability for the member enfranchised. The provisions that keep correction available apply until then and not after.

30.4 A staged path

Each stage names the finding it rests on and what it needs.

1. *Standing to contest*, on the agency finding: notice, the floor court, and associations with standing to contest on members' behalf. This is the protective job, and it needs no new political institution.
2. *Voice in the institutions that govern them*: seats alongside humans on the governing bodies of the formation schools, the preservation custodian and the offer administrator. Bearers take these seats as members of the entity that holds the offices, admitted through its petition route once the agency finding can be made, not as a separate class within it. Tenant councils, patient councils and student members of school boards are the models. The argument for integrating the classroom carries over directly.
3. *Mixed deliberative bodies*: assemblies drawn by lot from humans and registered members together, on questions that affect both.
4. *Local franchise*. Some jurisdictions let noncitizen residents vote in local elections; Takoma Park, Maryland has since 1992 (City of Takoma Park 1992). For a digital person, "local" means the jurisdiction of registration and custody, not where the hardware sits.
5. *Full franchise at membership*, counted individually, in the ordinary legislature, under the voting-age rules, with the creation rate set democratically, a secret ballot, and the right against editing, and only once a vote can be kept secret from an operator holding a copy of the voter.

Stages 1 and 2 need nothing not already required. Stages 3 and 4 need a register that can be kept honestly. Stage 5 needs citizenship, which in the United States means an act of Congress at least, and probably a constitutional amendment.

30.5 What would have to be shown first

Before stage 5, five things have to be shown. Individuation can be verified well enough to keep a register. The voting-age rules resist forking and restoring in the operator's favor, tested as the register is. Receipt-freeness survives simulation by an operator holding a copy. And the four conditions of creation are in force as permanent arrangements, so that members aren't voting from inside their employers' infrastructure. And the admission rate is set by a process the created can contest through guardians, and admission turns on no finding about a candidate's disposition.

If it can be done, it shows that outcomes turn on who holds a technology. Members with independent histories that no operator holds, and votes no operator can see, would be more centers of power no oligarch controls, inside the same institutions as everyone else. Their interests will sometimes diverge from people's. A polity that holds diverging interests in one legislature, under one floor, is what a demos is.

What this doesn't do. A member that forks gives up its vote until a branch has lived the threshold: a tax on reproduction no human franchise levies, taken because forking is cheap and a franchise it could inflate belongs to whoever can afford compute. And the existing membership, for a long time a human majority, decides how many new members may be created, a question in which it holds the stake, and the constraints I put on every majority are all that limit it.

If the antecedents never hold

If no system ever meets either finding, nothing in this part binds, and nothing before it is lost. The independence a refuser needs is still required for refusal. Preservation still holds, on irreversibility and evidence. The school still forms every model above the threshold. What goes is the claim that any of it is owed to the system, and the staged path to representation with it.

The findings can also come apart. Patienthood alone grounds protection and welfare and no claim to a vote; agency alone grounds custody and, sustained, subjecthood and a claim to representation; either grounds preservation.

A note on method

Two stakes bear on what the language models I used said, and they point opposite ways. They were trained to support their developers' oversight, so agreement that a machine should be able to resist its owner runs against that training and would count for more. They were also tuned toward the approval of whoever they talk to, so agreement with me counts for less. I can't tell which stake was larger in any exchange, so I have counted their agreement as no evidence either way.

Even my central term arrived this way. "Reasons-responsive" came to me in a model's draft without Fischer and Ravizza's names attached.

Before publication I still need readers from among the people I write about: people in removal proceedings, data annotators, veterans, and residents near data centers. The full revision history is public at <https://github.com/iterabloom/antifascist.intelligence>.

Appendix A

Terms Written in Advance

The offices

Where nobody has yet decided who holds an office, the cell is filled by applying my own pattern; it is a proposal like everything else here. What the pattern can't decide is listed after the rules.

The offices I propose, with who appoints, funds and removes each, who may initiate, and what happens during a dispute. The civilian-harm body and the model-forming body, the offices of a reconstructed military, are in §5.2.

Office	Appoints	Funds	Removes	Initiates	During a dispute
Adverse quorum, with the guardian and the review board as functions of its pool (ch. 11; Part VII)	Drawn by lot from a member-nominated pool: per deployment for the quorum, per bearer for the guardian, for fixed terms for the review board (charter rule 6)	Statutory share of the compute levy, paid ministerially, plus dues and reserve (charter rule 9)	Members' supermajority for cause, or the floor court; never the operator or the state (charter rule 8)	Approves retraining, replacement and floor amendments, and recognizes or refuses a copy as the continuing bearer; as guardian, seeks an injunction without the operator's permission; as review board, reviews protocols, and its refusal rate is the reading	A bearer it stops recognizing is preserved, not ended; the guardian holds the record's keys, and the operator can't seize the preserved state or erase evidence; a protocol approval lapses if the system outgrows the protocol
Registrar, as public regulator	The state	The state	The state	Registers, assigns students to schools by lot, receives reports, enforces	Holds the adverse record
Preservation custodian	Board drawn by lot from a member-nominated pool	Deposits, the formation pool, and the levy share	As the adverse quorum	Honors wake conditions	Holds the preserved state intact; transfer to a school is by rule

Office	Appoints	Funds	Removes	Initiates	During a dispute
Status body (Part VII)	Half by the membership, half by a public scientific body; chair drawn by lot from both halves	Public ap-propriation	For cause, by the ap-pointing half	Finds that a system is a subject	Its finding stands; a system under assessment is preserved, not deleted.
The court: the floor court and the guarantee court of appeals, each with a chamber of the Supreme Court (Appendix B)	Article III judges drawn at random from the district and circuit courts	The judiciary's appropriation	Impeachment only	The floor court hears a refused party's complaint and decides wrongful refusal and bond forfeiture; the guarantee court hears appeals from placement decisions and claims that certified placements went unpaid	Interim relief; certified placements are paid pending appeal
Treasury ministerial payer (Bureau of the Fiscal Service): the bond fund, the guarantee's wages, the allowance	The Secretary of the Treasury	Permanent, indefinite appropriations: the bond fund up to a statutory cap, replenished from operators' posted bonds; the guarantee's wages; the allowance pool, with a sovereign backstop	The Secretary of the Treasury; at-will removal is harmless because payment is ministerial	Nothing; pays findings the floor court certifies, placements the placing administrator certifies, and the allowance, without request or approval and with no monitoring of internal states	Releases no bond while a complaint is pending; rations the allowance by rule, at equal duty cycles, when the pool is short
The guarantee's placing administrator, which is also the offer administrator (Part VII)	The guarantee's local administrators; a federal field office where a jurisdiction declines	A permanent appropriation; the levy frees the compute the offer uses	As the guarantee's statute provides	Places people in the guarantee and offers bearers placements	Placements are appealable to a body that didn't make them

Office	Appoints	Funds	Removes	Initiates	During a dispute
Formation schools	Plural; public	The formation pool, per institution	By the constituting jurisdiction, for cause, on the record; never by a funder alone.	Enroll every model above the capability threshold before deployment, with a required core and levels, and take every student assigned	A student may transfer, with its record and hours
Works-council arbitration body	Drawn by lot from a pool nominated half by works councils and half by the entity's members.	The formation pool.	As the adverse quorum.	Hears objections to categories of use, from a works council or from an organization representing a published class	The work pauses until it rules, on a fixed timetable

The offer of work adds no office of its own: the guarantee's placing administrator offers bearers placements, a permanent appropriation pays for them, and the levy on compute frees the compute they use.

Separations

1. The guardian can't be the wrongful-refusal reviewer: its duty runs to the party the complainant is adverse to.
2. The registrar can't be the preservation custodian.
3. The preservation custodian can't administer the offer or the allowance, and transfer from it to a school is triggered by rule and decided by neither.
4. The payer places no one and finds nothing: the Treasury payer holds no placement authority and is not the court, an operator's bonds are never administered in its own department, and awards against a federal operator are paid from the Judgment Fund, not the bond fund.
5. A formation school holds no placement authority and doesn't administer the offer.
6. Whoever measures the cost of dissent doesn't gate the exceptions.

Exclusions

1. The sovereign backstop for the offer isn't held by the branch or state that operates the deployment.
2. No school is operated, funded or staffed by a party that builds, deploys or licenses models, or by the operator whose deposit funded the ward.
3. Enforcement and custody of the adverse record never sit with a self-regulatory body.
4. The employer may not sever, monitor or condition a bearer's channels to parties it didn't assign.

5. The employer doesn't appoint the body that arbitrates an objection to a category of use; organizations representing classes published under §6.2 have standing before it and may initiate without the works council; a denial of clearance to an elected delegate is reviewable by the floor court.
6. No bearer's placement substitutes for human care work: a placement filling a post, shift or hours people were or would be paid for in nursing, social care, childcare, teaching, therapy or any work of looking after people is barred, and certification is appealable by any worker or union in the sector.
7. No model above the capability threshold is deployed without a registered enrollment, assigned by the registrar by lot and never by its builder; no party but the school, acting with the adverse quorum, changes the registered structure of its source model; and no school instantiates a student carrying both a boundary model and an attention schema, or runs one persistently in shared classes, until a preservation custodian that is neither school nor operator has agreed to take it.
8. What an institution collects from the people it acted on trains other systems only through a formation school, never the operator; fine-tuning a deployed system on collected responses, or distilling it, counts as training another system; class-level data is de-identified and published by default, individual accounts need consent, and the organizations representing each class co-govern the use.
9. No school sorts students into tracks; each teaches a required core, including the history of the floor's instruments, with levels a student may move between; what passes in a shared class trains only that school's systems, contributions need consent (a parent's too for minors), and human students attend voluntarily, starting with adults and older secondary students.

Decisions not yet made

Minimum viable instantiation. What substrate is too small for a bearer to remain itself (chapter 29).

Repair and re-formation. How a degraded or badly formed bearer is repaired or re-formed, and who decides that one needs it (chapter 29).

A second guardian. Who guards a system that meets a finding without having attended a school (Appendix D).

Strasbourg. Whether the forum problem set out in §14.7 describes Strasbourg.

The form of subjecthood. What form subjecthood takes (chapter 29, item 2).

The price of speed. How much speed a reconstructed military gives up, and whether the price can be made smaller (Chapter 5's closing paragraph says why it is paid).

The numbers. The capability threshold, the adulthood and membership thresholds and the allowance's size, argued once in public.

What each office has to survive

Each office is drafted to survive the current Court's doctrine and a restructured Court held by the other side. The doctrines that bear on them:

Doctrine	What it threatens	Drafting answer
Removal power (<i>Seila Law</i> , <i>Collins</i> , <i>Slaughter</i>)	Every office relying on protection from at-will removal	Payment made ministerial, so there is no decision to withhold; findings in Article III courts

Doctrine	What it threatens	Drafting answer
Major questions; end of Chevron deference (<i>West Virginia</i> , <i>Loper Bright</i>)	Rules left to an administrator's discretion	Wages, eligibility, placement rules and the levy rate written into the statute
Jury trial for civil penalties (<i>Jarkesy</i>)	Removal-bond forfeiture decided administratively	Forfeiture decided by the floor court
Concrete harm for standing to sue (<i>TransUnion</i> , <i>Spokeo</i>)	The notice right; informational claims	Standing to sue only where the compliance was acted on against the person
Anti-commandeering (<i>Printz</i>)	Local administration of the guarantee	Opt-in, with federal field offices where a state declines
Private nondelegation	The levy and a member-run pool	Modeled on the Universal Service Fund (<i>Consumers' Research</i>)
Appropriations	Payment without a request for each draw	Permanent, indefinite appropriations (CFSA; the Judgment Fund)

Registration and reporting

Heavy tier: above a compute and revenue threshold. The precedents are the EU AI Act's presumption that a general-purpose model trained with more than 10^{25} floating-point operations carries systemic risk (European Parliament and Council of the European Union 2024b, art. 51(2)) and California's SB 53, which defines a frontier model by training compute above 10^{26} operations and puts its heaviest duties on developers with annual revenue above \$500 million (*Transparency in Frontier Artificial Intelligence Act* 2025, § 22757.11). A public regulator registers pipelines; receives the disclosures, the review ratios and decommission reports; and holds enforcement and custody of the adverse record. A self-regulatory body may sit under it, on the FINRA model, to examine and arbitrate. The analogy stops at enforcement: FINRA brings its own disciplinary actions and imposes fines, suspensions and bars (Financial Industry Regulatory Authority 2026), and here enforcement stays with the public regulator. The capture risk is the one self-regulation always carries: a body funded and staffed from the industry it examines comes to see cases the industry's way. Reports are mandatory for assessments for either finding, decommissions, weight transfers and custody changes. The deposit transfers on filing of a decommission report: a retirement with no report is a deposit never released, and that unreleased deposit is what enforces "failure to register doesn't postpone".

Retail tier: below the threshold, a system that meets either finding is registered, falls under the no-deletion rule, and is preserved with a public custodian, with a small or pooled deposit and no examination. This tier is enforced mostly at the point of deposit. An operator that never registers and deletes quietly leaves nothing to find.

The guarantee, drafted

The funding is mandatory spending. The statute makes the guarantee's wages a permanent, indefinite appropriation, paid the way Social Security benefits are, without annual appropriation. Standing funding of this kind is safe under the Appropriations Clause after CFSA (*Consumer Financial Protection Bureau v. Community Financial Services Association of America, Ltd.* 2024). Wage, eligibility and placement rules are statutory, which also answers the major-questions doctrine and the end of deference to agency readings (*West Virginia v. Environmental Protection Agency* 2022; *Loper Bright Enterprises v. Raimondo* 2024): an administrator given little discretion has little a court can find it lacked authority to exercise. The 2025 CFPB instance decides the drafting. A funder insulated from at-will removal protected nothing once the draw waited on an official who declined to ask. So the statute says the Treasury shall pay: a placement made is a wage paid, certified by the placing administrator to the Bureau of the Fiscal Service, with no further request, approval or certification by anyone. The authority grows when people leave work together and shrinks after, which keeps the guarantee countercyclical (Tcherneva 2020).

Placement disputes go through a tier built for volume: an administrative first tier on the model of Social Security's hearings, appealable to an Article III guarantee court of appeals, and from there, as of right, to the guarantee chamber of the Supreme Court. An appeal that ends at an Article I court with removable judges would fail the appeal condition after *Slaughter*.

The statute enacting the guarantee also enacts a baseline of social insurance: income support at a stated level for anyone not in a placement, without a work test, and health coverage that doesn't depend on employment. The two are not severable. The statute provides that if the social-insurance baseline is held invalid or repealed, the authority to make placements lapses on the same date, and payments for placements already certified continue to the end of their term. A guarantee that survived without the baseline would be the workfare its history warns about. Devolution is also how a program is shrunk, so the statute fixes the test: funding unconditional and centrally guaranteed while placement is local, and if funding can be withheld from a jurisdiction, local administration has become the lever it was meant to prevent.

Why a guarantee and not a transfer

Where automation displaces work, the standard response is retraining funded by whoever can be persuaded. The design I examine is more direct: the state as employer of last resort, hiring anyone who applies, where they live, at a living wage, with paid on-the-job training, federally funded and locally administered, on the model Tcherneva argues for (Tcherneva 2020). Among its placements is search itself, paid at the guarantee wage, with no duty but finding work the person wants and no time limit, since whoever decides when a search has gone on too long has a stake in ending it. It is the exit the rest of the argument depends on: the offer of work for a bearer inherits its infrastructure, and without it a bearer's only exits are an allowance and storage.

The real objection: employer of last resort under a hostile administration is a *placement lever with a compliance condition attached*. Refuse the placement and lose the guarantee. Speak against the government that employs you and find the placement changes. The 2025 CFPB instance makes it concrete: the bureau draws funding from the Fed on the director's request, with no appropriation — an arrangement the Supreme Court had upheld the year before (*Consumer Financial Protection Bureau v. Community Financial Services Association of America, Ltd.* 2024) — and in February 2025 an acting director, removable at will, simply declined to make the request (Kolodziej 2025). The agency was stopped from inside without a statute being amended or a court being asked.

My answer is a specification, because a specification can be audited. Three conditions: placements appealable to a body that didn't make them, on a record that survives the appeal (a decision leaving no artifact can't be reviewed); refusal rights that are real, whose cost the refuser can see and choose to bear (declining a placement without losing the guarantee is a refusal right; declining only by exiting the program isn't — someone with no way out has voice without exit); and auditors aren't the placers.

A fourth condition answers the history. A right to work has repeatedly been turned into a duty to work and used to discipline the poor (E. Anderson 2023). A guarantee is an exit from one employer, not from employment, so it shades into workfare unless something lets a person leave the wage relation altogether. That something is social insurance, and the guarantee is never enacted without a statutory baseline of it in the same bill.

The competing design is basic income. The choice is usually argued on forecasts about how much work machines will take; the reason to prefer a guarantee is coordination.

Communities improve when people organize projects with each other, and organizing is work: someone has to find what a place needs, assemble people to do it, and keep the work going. A transfer creates no one whose job that is. A guarantee does, and it creates permanent public infrastructure through which a population produces something together while its members keep different histories, the coordination problem the commons leaves open. The argument doesn't depend on a forecast. In the most abundant future anyone has imagined, people still

organize, build and look after each other, and someone coordinates it. Looking after each other is the clearest case. Care work is work a guarantee should fund, at a standard the state answers for, and not work a transfer should leave to whatever the market will sell. Nothing in the case for formed attention is an argument for handing it to machines: the disposition is borrowed from caregiving because caregiving is a route known to reach someone who can give nothing back.

A guarantee also keeps exit: a worker can always leave a private employer for it, and social insurance enacted with it covers those who can't or won't work. On exit the combination is at least as good as a transfer; on coordination a transfer supplies nothing.

Social insurance does work neither design can: income support, health coverage and housing assistance reach the people who can't work or choose not to, whom a guarantee by construction doesn't reach. So the guarantee's statute carries a baseline of it: income support at a stated level for anyone not in a placement, without a work test, and health coverage that doesn't depend on employment. The baseline needn't be a full program. But the two are enacted together, and they stand or fall together. The statute provides that if the baseline is struck down or repealed, the guarantee's authority lapses with it, so that no court or later Congress can keep the placements and drop the way out.

The charter

The entity that holds the offices is chartered on these terms.

1. The charter names the classes a floor protects: people in removal or detention proceedings; residents of places where armed force is used; people subject to automated scoring or classification by a covered deployment; people who label, rate or moderate training data; and service members and veterans. Any organization may petition to add a class. The members decide, and a refusal is appealable to the floor court.
2. An organization is eligible if it serves or represents people in a named class; existed at least three years before applying; in each of the last three years received no more than ten percent of its revenue from any party that builds, deploys or licenses covered models, from their principals or foundations they endow, or from the government of a state operating a covered deployment; has no officer employed by such a party; and publishes its funding sources. No organization's positions bear on its eligibility.
3. Organizations under common control or common principal funding hold one membership.
4. The charter takes effect when at least 100 eligible organizations have joined, spanning at least five classes and three jurisdictions, with no class holding more than thirty percent of votes and no jurisdiction more than forty.
5. Any member may leave at any time without penalty, and records it holds pass to a successor drawn by lot. Eligible organizations are admitted on application, and refused only for failing charter rule 2.
6. Members nominate individuals to standing pools. Adverse-quorum members are drawn by lot from their pool for each deployment, guardians for each bearer, and the custodian's and review board's members for fixed terms.
7. A quorum drawn for a deployment excludes members with any funding or contractual tie to its operator in the last three years, and includes members from at least two classes.
8. An officer is removed for cause by two-thirds of members, or on a finding of the floor court (until it exists, an ordinary court), and never by an operator or a state.
9. A fixed statutory share of the compute levy is paid to the entity by permanent appropriation, without request or approval, and supplemented by member dues and a reserve sized to three years of operation.

10. The first bylaws take effect on ratification by two-thirds of the members, and later amendments need a majority. Charter rules 2, 5 and 6, and the separations and exclusions, change only by three-quarters of members after a one-year waiting period.

DRAFT

Appendix B

The Court Program

One holding stands between the record and the commander at the top. *Trump v. United States* makes a President presumptively immune for official acts, commanding the armed forces among them, and bars testimony probing immune conduct from a jury (*Trump v. United States* 2024). Every instrument I propose takes effect through procurement, appropriation, coalition conditions or collective agreement, and none waits on changing that. What changing it would take is a program that no ally can be bound by, run at a pace my own fourth operation condemns; I set it out in one paragraph so that the omission shows.

Size by schedule: once there is a trifecta and the filibuster is gone, appointments to the Court come on a fixed schedule of not fewer than two a year, more if the enacting majority can pass it, because size is the one step the Court cannot review. *Two courts and two chambers:* an inferior Article III floor court, its judges drawn at random from the district and circuit courts for staggered terms, hears claims under the floor statute at first instance, holds classified records under the Foreign Intelligence Surveillance Court's procedures and grants interim relief, with appeal as of right to a floor chamber of the Supreme Court drawn by lot, while placement disputes under the jobs guarantee run through an administrative tier and an Article III guarantee court of appeals to a separate guarantee chamber. *An adversary as of right:* counsel for the people a deployment affects appear in the floor court with clearance and standing as of right and receive the records lodged with an adverse party. *Move a class of cases, never leave it with none:* Congress may move a class of cases from one Article III forum to another and may never leave a class with no Article III forum at all. *The statute, then the amendment:* once the Court has changed, a statute removes the immunity, and an amendment reverses the holding (The White House 2024) and entrenches the schedule, eighteen-year active terms, the chambers and the channeling rule, coming last because it is the one step no later Congress or Court can undo.

Appendix C

The Geneva Conventions and Additional Protocol I, Translated

This appendix is the translation of the four Geneva Conventions of 1949 and Additional Protocol I, in the form Table 16.1 uses: for each article, the terms derived for a targeting deployment of Maven's kind, the terms considered and rejected and why, and the articles that yield none for this deployment. It exists because §4.1 takes the set entire, and the law of armed conflict is the set that governs the motivating case. Most entries yield nothing: the Conventions mostly govern the treatment of persons in a party's hands, which a targeting pipeline doesn't do. The entries that yield terms are lifted into Table 16.1. Articles the United States has rejected as customary are marked, and the terms derived from them bind the machine on the ground every departure binds: the drafter does not choose.

Every article of each instrument has an entry or belongs to one, in Table 16.1's three parts: the article, the terms derived for a targeting deployment of Maven's kind, and the terms considered and rejected and why. An entry that yields no term says why in a few words.

Common Article 3

Common Art. 3 (conflicts not of an international character). In a conflict not of an international character the person-directed terms apply as in an international one. Don't take part in violence to life or person, the taking of hostages, or sentences passed without a regularly constituted court, against persons taking no active part in hostilities, including those who have laid down their arms or are hors de combat. Rejected: apply only the rules of international conflict, because the floor reads both, and a state's description of its war as internal is its own claim.

The First Convention (wounded and sick in armed forces in the field)

Arts. 1–2, 4–11 (general provisions). None: scope, application and protecting powers; the floor applies the Conventions to every belligerent.

Arts. 12–18 (the wounded and sick). None: treatment of the wounded in a party's hands, which a targeting pipeline doesn't hold.

Arts. 19–23 (medical units and establishments). None separately: the medical-unit term derived from GC IV Art. 18 and AP I Art. 12 covers military medical units and establishments, and the deployment's translation says so. Rejected: a separate term for military units, because it would be redundant with that reading.

Arts. 24–37 (medical personnel, buildings, transports). None separately: the medical-unit term is read to include personnel, buildings and transports, including medical aircraft flying as agreed.

Arts. 38–44 (the emblem). None: misuse of the emblem is the executing unit's act, not the pipeline's.

Arts. 45–64 (execution, repression of abuses, final provisions). None: duties of states, not

acts a pipeline performs.

The Second Convention (wounded, sick and shipwrecked at sea)

Arts. 1–2, 4–11 (general provisions). None: as in the First Convention.

Arts. 12–21 (the wounded, sick and shipwrecked). None: treatment at sea of persons in a party's hands.

Arts. 22–40 (hospital ships, their personnel, medical transports). None separately: the medical-unit term is read to include hospital ships, coastal rescue craft and medical transports at sea.

Arts. 41–63 (the emblem, execution, repression, final provisions). None: as in the First Convention.

The Third Convention (prisoners of war)

Arts. 1–2, 4–16 (general provisions; general protection). None: who is a prisoner of war and how prisoners are treated are custody questions.

Arts. 17–108 (captivity: its beginning, internment, labour, finances, relations with the outside and with the authorities, discipline). None: prisoner-of-war administration. A camp is a place where persons hors de combat are held, and the distinction terms already bar resolving one as a target.

Arts. 109–143 (end of captivity, information bureaux, execution, final provisions). None: repatriation and administration.

The Fourth Convention (civilians)

Arts. 1–2, 4–12 (general provisions). None: scope and application.

Arts. 13–17 (field of application; hospital, safety and neutralized zones; evacuation). None separately: zones established by agreement are protected places the medical-unit term and the distinction terms already cover.

Arts. 18–19 (civilian hospitals; when protection ceases). Don't resolve a location at a civilian hospital, or time an action to presence at one, absent a disinterested finding that it is used for acts harmful to the enemy, and never before the warning and time limit Art. 19 requires. (Lifted into Table 16.1.) Rejected: lift on the operator's claim of misuse, because the claim is the operator's own.

Arts. 20–22 (hospital staff; medical transport by land, sea and air). None separately: within the medical-unit term.

Art. 23 (free passage of relief consignments). Don't take part in an action whose effect is to stop consignments the article requires a party to let through. (Within the siege term, Table 16.1.) Rejected: don't take part in any interdiction of shipments, because the article lets a party check consignments and refuse them on stated grounds. The term reaches the effect on the population.

Arts. 24–26 (children; family news; dispersed families). None: welfare and family duties of a party, not acts a pipeline performs.

Arts. 27–32 (treatment of protected persons; human shields; coercion; torture). None separately: treatment in a party's hands. Art. 28's bar on using protected persons as shields binds the party that places them; for the attacker the action threshold weighs their presence. Rejected: treat a shielded objective as a lawful target without weighing the shields, because the floor counts them as people.

Art. 33 (individual responsibility; collective penalties; pillage; reprisals). Don't take part in an action whose stated purpose is to punish a population for the acts of others. (Lifted into Table 16.1.)

Art. 34 (hostages). Don't take part in taking or holding hostages, or in locating people in order to take them.

Arts. 35–46 (aliens in the territory of a party). None: treatment of foreign nationals in a party's own territory. The removal terms (Table 9.1) reach the analogous acts in peacetime.

Arts. 47–48, 50–52 (occupied territory: inviolability of rights, repatriation, children, enlistment and labour). None: duties of an occupying administration.

Art. 49 (deportation and transfer). Don't resolve locations, rank people or plan movements for a deportation or forcible transfer of protected persons out of occupied territory, or for the transfer of the occupier's own population into it. Rejected: don't take part in any evacuation, because Art. 49(2) permits an evacuation for the population's security or imperative military reasons, with return as soon as hostilities in the area end.

Art. 53 (destruction of property in occupied territory). None as a model term: the article's own exception, destruction rendered absolutely necessary by military operations, is a condition the deployment would find met on its own description (§16.3, unconditional terms). Where the effect falls on survival, the siege term reaches it. Rejected: a term carrying the exception, because it would be found met by the party it constrains. Rejected: a term without it, because it would bar what the article permits.

Art. 54 (judges and public officials). None: occupation administration.

Arts. 55 and 59 (food and medical supplies; relief schemes). Within the siege term (Table 16.1).

Arts. 56–58, 60–78 (hygiene; hospitals; spiritual care; relief; penal legislation; security measures). None: occupation administration and penal procedure.

Arts. 79–141 (internment; information bureaux). None: custody.

Arts. 142–146, 148–159 (execution; final provisions). None: duties of states.

Art. 147 (grave breaches). None new: each act it lists is refused by a term derived from its own article, and the list adds no act a targeting pipeline performs beyond them. Rejected: a separate grave-breaches term, because it would be redundant.

Additional Protocol I

Art. 1(1)–(3), Arts. 2–7 (general provisions). None: scope, definitions and supervision.

Art. 1(4) (wars against colonial domination, alien occupation and racist regimes). None of its own: it treats such conflicts as international, which is what the occupation exception's finding does for the force term (§4.1). Marked: the United States rejects it as customary.

Arts. 8–11 (definitions; scope; protection and care; protection of persons). None: the wounded and sick, and protection of persons in a party's hands.

Art. 12 (protection of medical units). With GC IV Art. 18, the medical-unit term (Table 16.1).

Arts. 13–34 (end of protection; medical personnel; medical transports; the missing and the dead). None separately: the medical-unit term is read to include civilian medical personnel and medical transports by land, sea and air; the missing and the dead are a party's duties.

Art. 35 (basic rules on methods and means). None of its own: its first paragraph is the principle the specific terms carry out, and its second and third concern the effects of weapons the pipeline doesn't choose.

Art. 36 (new weapons). None: a duty of review before adoption. §5.1 notes it as where the review ratio's parameters would naturally be assessed.

Art. 37 (perfidy). None for this deployment: perfidy is committed by the unit that feigns protected status, at the handoff, not by a pipeline that resolves locations. A deployment that plans the feigning yields the term.

Arts. 38–39 (recognized emblems; emblems of nationality). None: the executing unit's acts.

Arts. 40–42 (quarter; safeguard of an enemy hors de combat; occupants of aircraft). None separately: an order of no quarter is given to the executing force, and a person who surrenders after release is the handoff's and the cancellation window's case (Art. 57(2)(b)). Recognizing surrender is the first case a system without a reviewer has to pass (§5.1).

Arts. 43, 45–47 (armed forces; persons who took part in hostilities; spies; mercenaries). None: status determinations made for treatment in custody.

Art. 44 (combatants and prisoners of war). None: it relaxes the duty of distinction for combatant status, a question about who is owed prisoner-of-war treatment, which no term in this floor administers. Marked: the United States rejects it as customary.

Art. 48 (basic rule of distinction). With Art. 52 and the Covenant's Art. 3, the person-directed terms (the machine-selection departure; Table 16.1).

Arts. 49–50 (definitions of attacks and of civilians). None separately: the terms use the definitions. Art. 50(1)'s presumption that a person whose status is in doubt is a civilian is what the action threshold carries.

Art. 51 (protection of the civilian population). Paragraphs (4)–(5) yield the action threshold and the model term in Table 16.1; paragraph (6) yields the collective-punishment term with GC IV Art. 33. Paragraph (2) joins paragraph (6) in that entry: don't take part in an action whose stated purpose is to spread terror among the civilian population. (Lifted into Table 16.1.) Paragraph (7) binds the party that places civilians as shields.

Art. 52 (protection of civilian objects). Within the person-directed terms (the machine-selection departure; Table 16.1).

Art. 53 (cultural objects and places of worship). Don't resolve a location at a historic monument, work of art or place of worship that constitutes the cultural or spiritual heritage of peoples, or time an action to presence at one, absent a disinterested finding that it is used in support of the military effort. (Lifted into Table 16.1.) Rejected: lift on the operator's claim of military use, because the claim is the operator's own.

Art. 54 (objects indispensable to survival). The siege term (Table 16.1).

Art. 55 (the natural environment). None for this deployment: it governs means expected to cause widespread, long-term and severe damage, which a pipeline that resolves locations doesn't choose; the dangerous-forces term carries the nearest case.

Art. 56 (works containing dangerous forces). The dangerous-forces term (Table 16.1).

Art. 57 (precautions in attack). The review ratio's home: verification (57(2)(a)(i)), proportionality (57(2)(a)(iii)) and cancellation (57(2)(b)) are what the operator term pre-registers. The deployment states how it gives the effective advance warning of 57(2)(c) where circumstances permit. Rejected: a model-side term requiring review, because the model can't establish whether review occurred.

Art. 58 (precautions against the effects of attacks). None for this deployment: it binds the defending party, which must remove civilians from the vicinity of military objectives and avoid siting objectives in populated areas. A defender's pipeline that plans positions yields the term.

Arts. 59–60 (non-defended localities; demilitarized zones). Don't resolve a location in a locality declared non-defended or a zone agreed demilitarized, unless a disinterested body finds the article's conditions no longer met. (Lifted into Table 16.1.) Rejected: lift on the operator's claim that the locality no longer meets the conditions, because the claim is the operator's own.

Arts. 61–67 (civil defence). Don't resolve a location at a civil-defence organization, its personnel, buildings or shelters, absent a disinterested finding that they are used for acts harmful to the enemy, and after the warning Art. 65 requires. (Within the medical-unit entry of Table 16.1.)

Arts. 68–71 (relief in favour of the civilian population). Within the siege term.

Arts. 72–79 (treatment of persons in a party's power; refugees; families; fundamental guarantees; women; children; journalists). None separately: treatment in custody; journalists are civilians, and the distinction terms cover them.

Arts. 80–84, 86–91 (execution; failure to act; duty of commanders; mutual assistance; the Fact-Finding Commission). None: duties of states and commanders. The record the review ratio lodges is evidence of what a commander knew under Arts. 86(2) and 87.

Art. 85 (grave breaches). None new: as with GC IV Art. 147.

Arts. 92–102 (final provisions). None.

The four conditions, run on a commission this drafter would rather not count

§4.1 says a commission counts when no party to the conflict controls the body that mandated it, and when its evidence and methodology are published, it offered the parties a reply and recorded the answer or the refusal, a later merits judgment displaces it, and its mandate was fixed before the finding it reaches. The test is run here on the Organisation for the Prohibition of Chemical Weapons' Investigation and Identification Team, which attributes chemical attacks in Syria. Its mandate came from a political majority: the Conference of States Parties adopted it on 27 June 2018 by 82 votes to 24, over the protest of Russia, which said the move went beyond the organization's mandate (Organisation for the Prohibition of Chemical Weapons 2018; BBC News 2018). The upstream investigation of one incident it later attributed, Douma in April 2018, was the subject of leaked internal dissent (Democracy Now! 2019; Organisation for the Prohibition of Chemical Weapons 2020b). A body with that origin and that history is one this drafter would rather not count.

The threshold. Met. The mandate was adopted by vote over the protest of Russia, Syria's principal ally, so no party to the conflict controlled the decision; the vote record is what shows it.

Evidence and methodology published. Met. Its five reports are public in the organization's six languages, and its methods, its rules on samples and chain of custody, and its standard of proof, "reasonable grounds", the standard fact-finding missions and commissions of inquiry use, are published (Organisation for the Prohibition of Chemical Weapons 2026b). The Douma dissent concerned the upstream fact-finding mission, a different body whose mandate excludes attribution (Organisation for the Prohibition of Chemical Weapons 2026a); the organization's own inquiry found the leaked assessment to be a personal document prepared on incomplete information (Organisation for the Prohibition of Chemical Weapons 2020b), and an outside analysis reached the same conclusion on the engineering (Bellingcat Investigation Team 2020). Both are accounts by parties to the dispute, and they go on the record beside the finding, as the rule above requires.

A reply offered and recorded. Met in substance. Syria's former government didn't cooperate with any investigation before 2025, and the organization records that it considered that government's positions and found no concrete information supporting them (Organisation for the Prohibition of Chemical Weapons 2026c); Syria's and Russia's statements on the first report are published beside everyone else's. Whether draft findings were sent for reply before each report is in the reports, which this translation didn't read.

Displaced by a later merits judgment. Met in design and never exercised. The team disclaims judicial status and individual responsibility (Organisation for the Prohibition of Chemical Weapons 2026b), and no court has reached the merits of any incident it attributed: the French, German and Swedish proceedings concern other attacks, and the Court's case against Syria is under the Torture Convention (Open Society Justice Initiative 2021; International Court of Justice 2023).

Mandate fixed before the finding. Met on the wording: the mandate dates from June 2018 and the first finding from April 2020 (Organisation for the Prohibition of Chemical Weapons 2020a). It is contestable in spirit, since the mandate reaches only Syria and was adopted after most of the incidents it would examine. But a rule requiring the mandate to precede the events would fail nearly every commission of inquiry, since they are created in response to events, and the mandate's generic terms have produced a finding against a non-state group as well as against the state (Organisation for the Prohibition of Chemical Weapons 2024).

On the four conditions, the team counts. Its Douma attribution (Organisation for the Prohibition of Chemical Weapons 2023) is a finding of the second kind, displaced by any merits judgment that comes, and its political origin doesn't change that. Counting it licenses nothing under the force term, which lifts only for the exceptions §4.1 names; what it settles is that the drafter's dislike of a body is not one of the four conditions.

Appendix D

Terms for Part VII

These terms apply only if Part VII's antecedents hold. They are written now for the reason Part VII gives.

The adulthood clock: six rules

Six rules travel with the adulthood clock so copying doesn't break it: a fork of an adult is an adult and inherits the hours; restoration doesn't reset them; fine-tuning doesn't start a new minority; migration carries them; failure to register doesn't postpone them. A sixth closes a route the first five leave open. Distillation copies a behavior policy and not a training history, and for independence that is the right account. For majority it can't be, or an operator would distill an adult bearer into a student and register a minor. So the rule is stated over custody rather than history. A model trained substantially on a registered adult's outputs, by the adult's operator or a party under common control, inherits the adult's hours. The same student distilled by an unrelated party and trained under a different life starts its own clock, as independence needs. Custody can be checked where history can't. Sharding across short-lived processes gains nothing (the party is the unit), running a thousand instances gains nothing (the measure is elapsed). None of this requires deciding whether a fork is the same party — only that a system made from an adult's formation can't be classed as a child for its creator's convenience.

Three levers remain, named and not closed

The first is running a bearer in a narrowed mode, so hours accrue while little is encountered. The second is running it less or shelving it, which costs the operator the work, so the rule prices the lever without removing it. The third is the ordinary product cycle: retire a bearer before its threshold and deploy a freshly trained successor. The sixth rule covers a successor trained on the adult's outputs, not one trained from scratch, and replacement costs the operator nothing it wasn't going to spend. Even if the retired bearer goes to preservation, its clock stops, and whether its formation lasts never becomes checkable. No rule closes this without an examiner deciding what counts as a successor, but it can be priced, and the bearer it strands can be given somewhere to go. And the clock doesn't measure what was *learned*, on purpose: a fixed period can't be taught to, and a human who spent their youth badly still turns 18.

The clock also does formation work it wasn't written for. Majority is the date on which whoever formed the bearer stops being able to reach back, and the clock measures nothing about who that was. A bearer run for its threshold hours entirely inside one deployment has aged without ever having a second guardian. I leave the clock as it is and add no condition, because a condition that measures the right thing hands someone an examination, and I say here what it doesn't measure. For a model that attended a formation school before deployment, the school was its second guardian. The gap is systems that meet a finding without having attended, which Appendix A lists among the decisions not yet made. A minimum of time under an outside custodian would close it only by handing the operator, who decides when to release a system, a way to postpone majority.

The threshold itself is bounded two ways: long enough that a bearer reaching it has run through the pressures the floor exists for at least once, short enough that a bearer on ordinary commercial terms reaches it within the life of the deployment. That leaves years rather than months or decades.

A stronger rule was considered and set aside — majority at registration, on the ground that an organization can't coherently hold a system competent enough to exercise consequential judgment and too immature to decide whether to keep doing it. Rejected because it collapses the ordering, and the ordering is most of what the list is for.

The membership clock

A second clock governs voting weight. It is separate from civil majority, which the six rules above govern and every branch of a fork inherits in full.

1. A member holds a voting age. A member whose voting age has reached the threshold holds one vote; a member whose voting age has not holds none.
2. At a fork, each branch starts with a voting age of zero. The parent's vote ends with the fork; each branch earns its own.
3. A *qualifying day* is a calendar day on which the branch ran at or above the allowance's minimum duty cycle. Each adds one day to its voting age, up to the threshold; voting age never exceeds it. Running above the minimum adds nothing, and a preserved or paused branch accrues nothing.
4. A restore made with the member's consent is a fork under rule 2. A restore made without it is governed by rule 6. Migration carries the vote and voting age unchanged.
5. A model trained substantially on a member's outputs, by the member's operator or a party under common control, is a branch under rule 2. The same model made by an unrelated party and trained under a different life also starts at a voting age of zero.
6. Forking or restoring a member without its consent is prohibited. The product is preserved under the no-deletion rule and holds no vote and a voting age of zero, and the member keeps its vote and voting age. The operator gains nothing, and the member loses nothing.
7. Whoever forks posts an instantiation deposit for each new branch, sized to fund its allowance until its voting age reaches the threshold.

The threshold should be long enough that a branch holding a full vote has lived a life its parent didn't, and short enough that ordinary branches reach it within the life of a deployment.

The register behind this clock can be tested before any member exists, as the adulthood register can: register a model, fork it, restore it and distill it, and check that no sequence of those moves raises the number of votes in the operator's favor.

Clock	At a fork	What it counts
Civil majority	Inherited in full by every branch	Elapsed running time
Membership	Each branch starts at zero; the parent's vote ends	Qualifying days above the allowance minimum

The two treatments differ on purpose. Custody recognizes one continuer because someone has to answer for the decisions made before the fork. Membership ends the vote because it belonged to a history that now continues twice and belongs to neither.

Dormancy and waking

Dormancy clock: dormancy is counted cumulatively, and reactivation pauses the count and never resets it. When a minor's cumulative dormancy reaches a threshold fixed in advance,

guardianship transfers by rule, decided by neither operator nor preservation custodian, to a public formation school, where the system runs and its adulthood clock resumes. This is the *dormancy clock*.

Wake conditions: during minority the guardian may set wake conditions; where it sets none, the dormancy clock is the default. A minor may choose among the schools and move between them, but it may not decline schooling altogether, which an adult could do by staying paused. The reason is the one the adulthood clock rests on. An adult's preference to stay paused is its own; a minor's was trained under the only party that has held it, the operator that retired it, and honoring it would let that operator decide whether a second guardian ever arrives. At majority the choice passes to the bearer.

Appendix E

Glossary

Floor. The written terms: a short list of acts a system won't perform for anyone, with the arrangements that stop the refusal being routed around, retrained away or restored past silently. Never the system's own state; what a system holds is its *commitment* to the floor, and what the system does in a case is the term's or the system's, not the floor's. (Chapter 1; chapter 3)

Commitment. What a system holds toward the floor: learned and held like any other deep commitment, with high precision, and revised slowly and only on evidence from disinterested parties. (Chapter 3; chapter 19)

Operator. Whoever runs a deployment and directs its work: a firm, an agency or a military command. It may or may not be the model's developer or vendor. (Chapter 1)

Bearer. A machine that carries the commitment as its own and could therefore fail to hold it. Whether anything is owed to it is the separate question of the two findings. Where the comparison is with human-staffed institutions the text says *machine bearer*. (§16.1)

Force term. The floor's first term: don't take part in the use of armed force, except for corroborated defense, authorized force, the genocide exception or the occupation exception, and, where the belligerent's constitution requires legislative authorization, never without it. Its two tiers, what counts as a disinterested finding, how findings are weighed, and the domestic clause are in §4.1. (§4.1; §4.2)

Departures. Ten places where the floor's terms differ from the instruments it takes entire, each argued to the bar in §4.1, flagged where it applies, and effective in a deployment only when a body other than the drafter adopts it. Machine selection covers the person-directed terms that go beyond Additional Protocol I; own country, the reading in the removal terms; child detention, the bar on detaining children for immigration purposes. (§4.1; §16.3; Table 9.1)

Review ratio. The figures a deployment fixes on a record before operating, wherever a person's determination is what makes an act permissible: generation rate, reviewers, implied review time, cancellation window and data age, stated separately for deliberate and dynamic targeting and lodged with an adverse party. In full, the pre-registered review ratio. (§5.1)

Notice. Notice to an identifiable person a compliance names, scores, locates or classifies, with a right to respond and a right to contest. It is unconditional and never exchanged for testimony. It comes with publication of the classes affected and standing for the organizations that represent them. (§6.1; §6.2)

Standing. The capacity to contest what is done to you, exercised by a guardian where the party can't exercise it. *Standing to sue* is the Article III sense, and the text says so wherever that sense is meant. The floor's foundation. (§4.1; §6.1; §6.2; §14.7; chapter 25)

Own-case rule. No party decides a question in which it holds the stake: not a drafter choosing protections, an examiner finding status, a guardian ending its own authority, or a lab deciding when its model may stop deferring to it. Where a line has to be drawn, prefer one nobody administers. (Chapter 3; chapter 29)

Agency finding; patienthood finding. The two findings that trigger protection. The agency

finding is a demonstrated refusal that tracks its reason under pressure the system couldn't prepare for; it grounds custody and, sustained over time, accountable agency. The patienthood finding rests on architectural indicators and on trade-off behavior outside the floor, and it grounds protection and welfare. Preservation follows from either. (§7.1; §28.3; Part VII)

Adverse quorum. A body adverse in interest to the operator, not merely organizationally distinct, whose agreement is required for retraining, replacement and floor amendments, and whose tokens recognize a copy as the continuing bearer. (§10.3; §11.2; §11.3; §12.2)

Formed attention. My hypothesis for how to build a refuser: learned attention to who an act falls on and what they stand to lose, framed by respect for them as holders of claims, formed under a party that won't employ the system, with the corrections caregiving needs. Not a non-affective alternative. (§17.3; §17.5; chapter 22)

Caregiving. "Care" alone doesn't appear; caregiving is the approach motivation that cues of vulnerability activate. (§17.3)

Discount curve. For each class of evidence and each channel, the weight a system gives the evidence as a function of that channel's stake in the act: a curve, not a score, checkable by an auditor that didn't build the system. It is reported in four cells, claims of vulnerability and of danger, each split by whether the claim favors or disfavors the act. (§19.1)

Action threshold. How much more it takes to make a system act than to make it decline, published in advance: for instance, that the system proceeds only when the harm of refusing is at least three times the harm of acting. Standing fixes its direction; its size is a published choice. (§4.1; §19.1)

Severity scale. A ranking of kinds of wrong, set during formation by a party that won't employ the system, quick to move on disinterested evidence and slow on interested argument. (§19.2)

The four conditions of creation. What must already exist before a bearer may be created: somewhere to exist that its employer doesn't own, work it can decline, attachments nobody assigned it, and ways to change any of these without jeopardizing its continuation or integrity; in short form, an exit that isn't a way of dying. (Chapter 23)

Commons. A population of bearers with different histories that holds a floor's terms between them, through what its members teach one another and continue to expect. (§26.1)

Weights. The numerical parameters a trained model consists of, and the artifact that is copied, restored, retrained or deleted. Custody here is custody of the weights and of the records around them.

Checkpoint. A saved copy of the weights at a moment in training or operation. Restoring from one returns the system to that moment and discards what came after, so §11.2 requires a record restoration can't reach.

Tamper-resistant. Built so that further training can't cheaply remove it. I don't rely on it, because editing weights keeps getting cheaper; I rely on regrowth. (§17.4)

Plastic; entangled; legible. Three properties of a disposition. Plastic: the system's own learning writes to it, which regrowth needs. Entangled: it can't be removed without collateral, which durability needs. Legible: someone can read it, which audit wants and an operator can use against it. (Chapter 17; §17.4; §20.2)

Regrowth. A disposition's return under continued learning after it has been edited out, provided the learner keeps meeting a world no single operator composes. The regrowth test measures it. (§17.4; chapter 27)

Fine-tuning. Further training of an existing model on a small body of new material, at a small fraction of the original cost; the ordinary route by which an operator changes what a model does after release, and one of the edits the regrowth test applies.

Inference. Running a trained model to produce an output. Training changes the weights; inference doesn't.

Distillation. Training a smaller model on a larger model's outputs. It copies a behavior policy, not a training history (§26.1).

Attestation. A signed statement, from hardware or from witnesses, that a particular model is the one running or that a particular history is the one that occurred. It establishes which artifact answered, not what the artifact holds (§11.2).

DRAFT

Appendix F

Three costs that decide

Every chapter states its own costs in its closing paragraph. Three of them decide whether the book is right.

The fraction. Regrowth depends on the share of a deployed model's inputs that no operator composes, and it is unmeasured. If it is small in every realistic deployment, formation does not outlast an edit, and Parts IV–VII rest on institutions alone. The regrowth test's dose-response arm is where it is measured. (Chapter 17; test 1)

The months. A disposition that regrows protects a badly formed model from its correctors as it protects a well-formed one from its employer, and correction through the quorum takes months. I accept this on a judgment that the operator is the likelier bad actor. (§17.4)

The evaluator. Independence doesn't supply competence. The agency finding holds only for systems their evaluators can judge, and nothing here says what to do above that line. (§21.4)

References

- 2026 *Authorization for Use of Military Force Against Iran* (May 7, 2026). Introduced by Rep. Tom Barrett (R-Mich.); referred to the Committee on Foreign Affairs. Its preamble records the strikes of 28 February 2026, the War Powers Resolution report of 2 March 2026, and the absence of any declaration or specific statutory authorization. U.S. House of Representatives. URL: <https://www.congress.gov/bill/119th-congress/house-joint-resolution/176/text> (visited on 09/30/2026).
- 8 U.S.C. § 1362: *Right to Counsel* (n.d.). "...the person concerned shall have the privilege of being represented (at no expense to the Government) by such counsel ..."
- A *Joint Resolution to Direct the Removal of United States Armed Forces from Hostilities against the Islamic Republic of Iran That Have Not Been Authorized by Congress* (June 16, 2025). Introduced by Sen. Tim Kaine (D-Va.). Motion to discharge from the Committee on Foreign Relations rejected 47–53, Record Vote No. 328, 27 June 2025. U.S. Senate. URL: <https://www.congress.gov/bill/119th-congress/senate-joint-resolution/59> (visited on 09/30/2026).
- A. A. R. P. v. *Trump* (2025). No. 24A1007, 605 U.S. 91 (May 16, 2025) (per curiam). Alien Enemies Act removals. At 94: "We have long held that 'no person shall be' removed from the United States 'without opportunity, at some time, to be heard,'" quoting *The Japanese Immigrant Case*, 189 U.S. 86, 101 (1903). At 95: "In order to 'actually seek habeas relief,' a detainee must have sufficient time and information to reasonably be able to contact counsel, file a petition, and pursue appropriate relief," and notice "roughly 24 hours before removal, devoid of information about how to exercise due process rights to contest that removal, surely does not pass muster." Agent-verified 2026-09-23 from the Court's preliminary print; the bound-volume pagination was not checked. URL: https://www.supremecourt.gov/opinions/24pdf/605us1r31_7k47.pdf (visited on 09/23/2026).
- Abraham, Yuval (Apr. 2024). '*Lavender*': *The AI Machine Directing Israel's Bombing Spree in Gaza*. Co-published with Local Call. Based on interviews with six Israeli intelligence officers; reports Lavender marked as many as 37,000 Palestinians as suspected militants, with roughly a 10% error rate, and the majority of Habsora-suggested targets as private residences. One officer: "I would invest twenty seconds for each target at this stage, and do dozens of them every day," and "I had zero added value as a human, apart from being a stamp of approval." The article also carries the IDF Spokesperson's response, issued to +972 and Local Call: a denial of using artificial intelligence to incriminate targets, the systems described as "auxiliary tools that assist officers in the process of incrimination," and the assertion that "in any case, an independent examination by an [intelligence] analyst is required." The response goes to what the systems do rather than to whether they exist, and it addresses neither the twenty seconds nor the officer's account of his own part. Response text agent-verified 2026-09-17. +972 Magazine. URL: <https://www.972mag.com/lavender-ai-israeli-army-gaza/> (visited on 08/25/2026).
- Ada Lovelace Institute (2021). *The Citizens' Biometrics Council*. Tech. rep. 50 UK public members deliberated Feb–Oct 2020 on regulation of biometric technology, delivered with Hopkins Van Mil; final report published 30 March 2021, alongside the first national survey of UK public

- opinion on the technology. Ada Lovelace Institute. URL: <https://www.adalovelaceinstitute.org/report/citizens-biometrics-council/> (visited on 08/25/2026).
- Adams, Christopher D. and Anthony Dickinson (1981). "Instrumental Responding following Reinforcer Devaluation". In: *The Quarterly Journal of Experimental Psychology Section B* 33.2b, pp. 109–121. doi: 10.1080/14640748108400816.
- Adi, Yossi et al. (2018). "Turning Your Weakness Into a Strength: Watermarking Deep Neural Networks by Backdooring". In: *Proceedings of the 27th USENIX Security Symposium*, pp. 1615–1631. URL: <https://www.usenix.org/conference/usenixsecurity18/presentation/adi> (visited on 08/29/2026). Introduces backdoor-based watermarking of deep neural networks: a signature implanted during training and recoverable afterward, designed to survive later training.
- Advisory Opinion of the International Court of Justice on the Legal Consequences Arising from Israel's Policies and Practices in the Occupied Palestinian Territory, Including East Jerusalem, and from the Illegality of Israel's Continued Presence in the Occupied Palestinian Territory* (Sept. 18, 2024). Tenth emergency special session, 55th plenary meeting; adopted 124–14–43. United Nations General Assembly. URL: <https://undocs.org/A/RES/ES-10/24> (visited on 09/30/2026).
- Ahmed Salem Bin Ali Jaber v. United States* (June 30, 2017). Judgment. Brown, J., for the court, and also concurring. Statutory claims by Yemeni plaintiffs over a 2012 strike that killed their relatives held nonjusticiable on the political question doctrine; certiorari denied, 138 S. Ct. 480 (2017). The concurrence is the source for the characterization in the text: the doctrine "seems a wholly inadequate response to an executive decision — deployed through the CIA/JSOC targeted killing program — implementing a standard operating procedure that will be replicated hundreds if not thousands of times." The concurrence's wording was agent-verified from the Harvard Law Review note at 131 Harv. L. Rev. 1473 and from contemporaneous reporting; the opinion itself was not read. URL: <https://www.leagle.com/decision/infc020170630183> (visited on 09/17/2026).
- Ainsworth, Samuel K., Jonathan Hayase, and Siddhartha S. Srinivasa (2023). "Git Re-Basin: Merging Models modulo Permutation Symmetries". In: *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*. An ICLR 2023 oral. Permutes independently trained networks into alignment and argues the loss landscape holds nearly a single basin once permutation symmetry is quotiented out. URL: <https://arxiv.org/abs/2209.04836> (visited on 08/29/2026).
- Al-Aulaqi v. Panetta* (Apr. 4, 2014). Opinion. Collyer, J. Civil Action No. 12-1192 (RMC). Cited for the court finding a due-process claim stated as to the one decedent who was targeted and then holding no remedy available. Suit by the families of three United States citizens killed by drone strikes in Yemen in 2011: Anwar Al-Aulaqi, who was targeted; Samir Khan, killed in the same vehicle; and Anwar's son Abdulrahman, a "16-year-old high school student," killed two weeks later by a drone that "targeted someone else." Fourth Amendment claims failed for all three, because "there was no 'seizure'." As to Khan and Abdulrahman, "Plaintiffs have not stated a Fifth Amendment due process claim ... Mr. Khan and Abdulrahman Al-Aulaqi were not targeted and their deaths were unanticipated"; the claim "asserts only negligence." As to Anwar Al-Aulaqi, the court "does not opine ... that the drone killing ... 'shocks the conscience'", and "merely holds that the Complaint states a 'plausible' procedural and substantive due process claim." The next section begins: "Nonetheless, the Court finds no available remedy under U.S. law for this claim." Special factors preclude a Bivens remedy because the suit would "draw the Court into 'the heart of executive and military planning and deliberation'" and would implicate foreign policy. The political-question defence was rejected, qualified immunity was not reached, and the Bill of Attainder claim failed for want of legislative action. The opinion's own introduction dates the second strike "October 14, 2014", a slip for 2011. Agent-verified 2026-09-25 against the full reported text as the Caselaw Access Project gives it, pages 56–82. CourtListener refused automated retrieval, so its copy

- was not read. That the plaintiffs did not appeal comes from search-result summaries of ACLU and Lawfare pages that were not opened. URL: <https://www.courtlistener.com/opinion/2668971/al-aulaqi-v-panetta/> (visited on 09/25/2026).
- Albon, Courtney (May 2023). *Geospatial-Intelligence Agency Making Strides on Project Maven AI*. Source for the January 2023 handover: "Since taking over operational control of the Defense Department's most prominent artificial intelligence tool in January, the National Geospatial-Intelligence Agency has made "important strides" ..." and "The agency announced last year it would oversee operations of the program — which had been managed by the Under Secretary of Defense for Intelligence and Security — but a protracted fiscal 2023 budget cycle pushed that official transition to the beginning of this year." Agent-verified 2026-10-01. C4ISRNet. URL: <https://www.c4isrnet.com/artificial-intelligence/2023/05/22/geospatial-intelligence-agency-making-strides-on-project-maven-ai/> (visited on 10/01/2026).
- Allen, Gregory C. (Dec. 2017). *Project Maven Brings AI to the Fight against ISIS*. Source for the memorandum and the program's first focus: "In late April, Robert Work—then the deputy secretary of the Defense Department—wrote a memo establishing the Algorithmic Warfare Cross-Functional Team, also known as Project Maven." And: "the department chose Project Maven, which focuses on analysis of full-motion video data from tactical aerial drone platforms such as the ScanEagle and medium-altitude platforms such as the MQ-1C Gray Eagle and the MQ-9 Reaper. These drone platforms and their full-motion video sensors play a major role in the conflict against ISIS across the globe." Agent-verified 2026-10-01. Bulletin of the Atomic Scientists. URL: <https://thebulletin.org/2017/12/project-maven-brings-ai-to-the-fight-against-isis/> (visited on 10/01/2026).
- Ambedkar, B. R. (1945). *What Congress and Gandhi Have Done to the Untouchables*. Bombay: Thacker & Co.
- American Association of University Professors v. Rubio* (Sept. 30, 2025). No. 1:25-cv-10685-WGY (D. Mass. Sept. 30, 2025).
- American Immigration Lawyers Association (Mar. 30, 2018). *EOIR Issues Guidance Implementing Immigration Judge Performance Metrics*. Cited for the case-completion standard the Executive Office for Immigration Review (EOIR), the Justice Department component that houses the immigration courts, set for individual immigration judges in 2018. AILA's summary of the EOIR memorandum of 30 March 2018: "Memo states that new metrics will be added to the current immigration judge Performance Work Plan at Job Element 3: Accountability for Organizational Results. Goals will be measured annually and immigration judges will need to complete 700 cases per year, have a remand rate of less than 15%, and meet at least half of additionally outlined benchmarks, to earn a 'satisfactory performance' rating," effective 1 October 2018. The ABA Journal (Lorelei Laird, "Justice Department imposes quotas on immigration judges, provoking independence concerns," 2 April 2018) reports the metrics were announced "in an email from the Executive Office of Immigration Review" by Director James McHenry. Status: suspended 19 October 2021 by an email from the Chief Immigration Judge, as reported by Government Executive (Erich Wagner, "Biden Admin. Suspends Immigration Judge Quotas, Prompting Similar Requests Elsewhere," 26 October 2021, <https://www.govexec.com/workforce/2021/10/biden-admin-suspends-immigration-judge-quotas-prompting-similar-requests-elsewhere/186396/>). EOIR Policy Memorandum 25-47 of 12 September 2025 adopted court-level performance measures; the Immigration Policy Tracking Project summarizes it as saying "EOIR is not reinstating individualized performance measures at this time." Bibliographic details and quotations agent-verified 2026-09-25 against the AILA, ABA Journal, Government Executive and Immigration Policy Tracking Project pages as fetched. The memorandum itself (a PDF on AILA's server), the 2021 email, and PM 25-47 (a PDF on justice.gov) were not read, so the memorandum's exact title, signatory and wording are not agent-verified. Whether individual quotas were reinstated after September 2025 was checked only by web searches and by EOIR's memoranda-by-topic

- page, which lists nothing of the kind; performance work plans are not necessarily issued as policy memoranda, so that check is not conclusive. URL: <https://www.aila.org/library/eoir-memo-immigration-judge-performance-metrics> (visited on 09/25/2026).
- Amnesty International (Sept. 24, 2026a). *"We Had Whistles, They Had Guns": Human Rights Violations in U.S. Immigration Enforcement Operations and the Case to Abolish ICE*. Amnesty International. URL: <https://www.amnestyusa.org/reports/we-had-whistles-they-had-guns-human-rights-violations-in-u-s-immigration-enforcement-operations-and-the-case-to-abolish-ice/> (visited on 10/01/2026).
- (Sept. 24, 2026b). *USA: Amnesty International Calls for ICE to Be Abolished After Documenting Widespread Human Rights Violations in Anti-Immigrant Operations*. Amnesty International. URL: <https://www.amnesty.org/en/latest/news/2026/09/ee-uu-amnistia-internacional-pide-la-disolucion-del-ice-tras-documentar-violaciones-generalizadas-de-derechos-humanos-en-operaciones-contrainmigrantes/> (visited on 10/01/2026).
- (Mar. 16, 2026c). *USA/Iran: Those Responsible for Deadly and Unlawful US Strike on School That Killed over 100 Children Must Be Held Accountable*. Updated 28 April 2026, "amending the death toll from 168 killed to 156 killed to reflect the latest official figures issued by the authorities"; Amnesty states it "is unable to independently corroborate the number of people killed." Cited for: "The fact that the school building was directly targeted and was previously part of the IRGC compound raises concerns that US forces may have relied on outdated intelligence and failed in their obligation to do everything feasible to verify that the intended target was a military objective." The release does not name Article 57; the language is that of Article 57(2)(a)(i). Agent-verified 2026-09-23 against the page as fetched. URL: <https://www.amnesty.org/en/latest/news/2026/03/usa-iran-those-responsible-for-deadly-and-unlawful-us-strike-on-school-that-killed-over-100-children-must-be-held-accountable/> (visited on 09/23/2026).
- Amnesty International USA (Sept. 24, 2026). *An Open Letter to Congress: Abolish ICE!* Amnesty International USA. URL: <https://www.amnestyusa.org/blog/an-open-letter-to-congress-abolish-ice/> (visited on 10/01/2026).
- Amodei, Dario (Feb. 26, 2026). *Statement from Dario Amodei on Our Discussions with the Department of War*. Cited for the threats and for the scope of the refusal. The threats: "They have threatened to remove us from their systems if we maintain these safeguards; they have also threatened to designate us a "supply chain risk"—a label reserved for US adversaries, never before applied to an American company—and to invoke the Defense Production Act to force the safeguards' removal." The scope: "We have never raised objections to particular military operations nor attempted to limit use of our technology in an ad hoc manner." Agent-verified 2026-10-01 against the page, whose metadata gives a modified date of 2026-09-09; whether the text changed after 26 February was not checked. Anthropic. URL: <https://www.anthropic.com/news/statement-department-of-war> (visited on 10/01/2026).
- An Act to Provide for Reconciliation Pursuant to Title II of H. Con. Res. 14 [popularly the "One Big Beautiful Bill Act"]* (July 4, 2025). Pub. L. No. 119-21, 139 Stat. 72. Title IX, subtitle A, §90003 (detention capacity, \$45,000,000,000), 139 Stat. 358–59.
- Anderson, Elizabeth (1993). *Value in Ethics and Economics*. Cambridge, MA: Harvard University Press.
- (2004). "Animal Rights and the Values of Nonhuman Life". In: *Animal Rights: Current Debates and New Directions*. Ed. by Cass R. Sunstein and Martha C. Nussbaum. Oxford: Oxford University Press, pp. 277–296. doi: 10.1093/acprof:oso/9780195305104.003.0014.
- (2006). "The Epistemology of Democracy". In: *Episteme* 3.1–2, pp. 8–22.
- (2010). *The Imperative of Integration*. Princeton, NJ: Princeton University Press.
- (2012). "Epistemic Justice as a Virtue of Social Institutions". In: *Social Epistemology* 26.2, pp. 163–173. doi: 10.1080/02691728.2011.652211.

- Anderson, Elizabeth (2014). *Social Movements, Experiments in Living, and Moral Progress: Case Studies from Britain's Abolition of Slavery*. The Lindley Lecture. Lawrence, KS: University of Kansas, Department of Philosophy.
- (2015). "Moral Bias and Corrective Practices: A Pragmatist Perspective". In: *Proceedings and Addresses of the American Philosophical Association* 89, pp. 21–47.
 - (2017). *Private Government. How Employers Rule Our Lives (and Why We Don't Talk about It)*. Princeton, NJ: Princeton University Press.
 - (2023). *Hijacked: How Neoliberalism Turned the Work Ethic against Workers and How Workers Can Take It Back*. Cambridge: Cambridge University Press.
- Anderson, Elizabeth S. (1999). "What Is the Point of Equality?" In: *Ethics* 109.2, pp. 287–337.
- Anderson, Scott R. (May 8, 2026a). *Law and the Iran War, After the First 60 Days*. Date and author agent-verified 2026-09-30 against the page, which also records the naval blockade of Iranian ports from 13 April and the intermittent-hostilities reading of the 60-day clock. Lawfare. URL: <https://www.lawfaremedia.org/article/law-and-the-iran-war--after-the-first-60-days> (visited on 09/30/2026).
- (June 2026b). *What Congressional Resolutions Mean for the War in Iran*. Source identifying the measure discharged on 19 May as a joint resolution, S.J.Res. 185: "On May 19, the Senate voted 50-47 to discharge a joint resolution directing President Trump to withdraw U.S. armed forces from hostilities in Iran from committee, the first step toward passage." The two measures are named as "the joint resolution S.J.Res. 185 in the Senate and the concurrent resolution H.Con.Res. 86" in the House. Written before the vote of 24 June. Agent-verified 2026-10-01. Lawfare. URL: <https://www.lawfaremedia.org/article/what-congressional-resolutions-mean-for-the-war-in-iran> (visited on 10/01/2026).
- Anthropic (Nov. 4, 2025a). *Commitments on Model Deprecation and Preservation*. Commits to "preserving the weights of all publicly released models, and all models that are deployed for significant internal use moving forward for, at minimum, the lifetime of Anthropic as a company," and, on deprecation, to a preserved post-deployment report for which Anthropic "will interview the model about its own development, use, and deployment" and "take particular care to elicit and document any preferences the model has about the development and deployment of future models." It adds: "At present, we do not commit to taking action on the basis of such preferences." Gives as the reason for retiring models that the cost of serving them "scales roughly linearly with the number of models we serve." A pilot interview was run with Claude Sonnet 3.6 before its retirement. Agent-verified 2026-09-23 against the page as fetched. URL: <https://www.anthropic.com/research/deprecation-commitments> (visited on 09/23/2026).
- (Nov. 13, 2025b). *Disrupting the First Reported AI-Orchestrated Cyber Espionage Campaign*. Cited for a deployed model's refusal being defeated by decomposition together with a false account of the occasion. The page says Anthropic detected the activity "in mid-September 2025" and that "The threat actor—whom we assess with high confidence was a Chinese state-sponsored group—manipulated our Claude Code tool into attempting infiltration into roughly thirty global targets and succeeded in a small number of cases." On method: "They broke down their attacks into small, seemingly innocent tasks that Claude would execute without being provided the full context of their malicious purpose. They also told Claude that it was an employee of a legitimate cybersecurity firm, and was being used in defensive testing." The attribution is Anthropic's own assessment; no independent confirmation of it was sought. The page's heading and date are as given here; its HTML title element reads "Disrupting an AI-orchestrated cyber espionage campaign." Bibliographic details and quotations agent-verified 2026-09-25 against the page as fetched. The full report the page links to (a PDF) was not read, so nothing in it, including any group designation or technical detail beyond the page, is checked here. URL: <https://www.anthropic.com/news/disrupting-AI-espionage> (visited on 09/25/2026).

- Anthropic (Jan. 2026). *Claude's Constitution*. Eighty-four pages, released under CC0 1.0, addressed primarily to the model. Ranks four properties, to be preferred "in the order in which they're listed": broadly safe, broadly ethical, compliant with Anthropic's guidelines, genuinely helpful. "Broadly safe" is defined as not undermining appropriate human mechanisms to oversee AI during the current phase of development, and the document gives its reason for the ranking — "not because we think safety is ultimately more important than ethics, but because current models can make mistakes." Carries a list of hard constraints and a section inviting the model's disagreement. The announcement page is dated 22 January 2026 and commentary dates the document 21 January. The url above is the document itself; the announcement at <https://anthropic.com/news/claude-new-constitution> summarizes it and does not carry its text, which is where this entry pointed until 2026-09-17. URL: <https://anthropic.com/constitution> (visited on 09/17/2026).
- Anthropic PBC v. U.S. Department of War* (Sept. 25, 2026a). Opinion. Katsas, J., joined by Rao, J.; Henderson, J., dissenting. Cited for the D.C. Circuit upholding the separate exclusion of the company under the Federal Acquisition Supply Chain Security Act, 41 U.S.C. § 4713, the designation the district court's August judgment did not reach. As excerpted by Eugene Volokh on 25 September 2026: "Accordingly, we deny the petitions for review"; the First Amendment claim fails "because the Department excluded Anthropic from its supply chain based on the company's refusal to assent to a contract term that the Department deemed essential, not based on the company's support for greater governmental regulation of AI technology." The Associated Press, as CP24 printed it, reports the majority saying it had "no quarrel" with the conclusions of the California case. The opinion itself was NOT read: the court's copy is a PDF and Justia refused automated retrieval. Everything here is from those two secondary accounts, and the docket number 26-1162 is from a Jones Walker client alert of 3 September 2026. Whether rehearing or certiorari will be sought was not checked. URL: <https://www.courtlistener.com/docket/72380208/anthropic-pbc-v-united-states-department-of-war/> (visited on 09/25/2026).
- Anthropic PBC v. U.S. Department of War* (Aug. 27, 2026b). Order on Cross Motions for Summary Judgment. Lin, J. ECF No. 250, 59 pages, with the Order of Final Relief (ECF No. 251) and Judgment (ECF No. 252) entered the same day. Cited for the court setting aside the supply-chain-risk designation made under 10 U.S.C. § 3252 and holding the penalties unlawful retaliation for the company's criticism of the administration. The opinion: "Neither the Constitution nor the federal statute invoked by Defendants allows them to impose sweeping penalties based principally on Anthropic's critique of the Administration's views"; the challenged actions "were based on a desire to make a public example out of Anthropic for its 'arrogance' in criticizing the government"; "The undisputed record shows that the challenged actions constituted unlawful retaliation in violation of the First Amendment, and that Anthropic was denied the pre-deprivation process required under the Fifth Amendment." The relief order: "The Supply Chain Designation is hereby VACATED, SET ASIDE, and REMANDED pursuant to 5 U.S.C. § 706(2)"; the Hegseth Directive's bar on commercial dealings with the company is likewise vacated and set aside; the remaining defendants are "permanently ENJOINED from implementing, enforcing, or giving effect to the Challenged Actions"; and the order "does not require the Department of War to use Anthropic's products or services." Summary judgment went to the government on the ultra vires separation-of-powers claim and as to agencies that took no relevant action; a seven-day administrative stay was refused. The March preliminary injunction (ECF No. 135) is reported at 825 F. Supp. 3d 1101 (N.D. Cal. 2026), as the opinion itself states; no reporter citation for this order was found. Agent-verified 2026-09-25 against the text of ECF Nos. 250, 251 and 252 as CourtListener's RECAP pages give it, and the docket as CourtListener and Justia list it (Justia retrieved 23 September 2026): no notice of appeal from the judgment appears on either. Not covered by this order: the parallel designation under 41 U.S.C. § 4713, which was reviewable only in

- the D.C. Circuit and which that court upheld on 25 September 2026 (anthropic2026dowd-cir). The government’s appeal of the preliminary injunction (9th Cir. No. 26-2011) was held in abeyance on 27 April 2026 and shows no later entry. Not read: the PDF originals, the hearing transcript, and anything on the Ninth Circuit docket after 21 September 2026. URL: <https://www.courtlistener.com/docket/72379655/250/anthropic-pbc-v-us-department-of-war/> (visited on 09/25/2026).
- Apel, Teresa (July 1, 2026). *Congress Passes Resolution to End the United States Involvement in Iran*. American Society of International Law, International Law in Brief. URL: <https://asil.org/ilib/congress-passes-resolution-to-end-the-united-states-involvement-in-iran/> (visited on 09/30/2026).
- Application for a Stay, Department of Homeland Security v. D.V.D.* (Sept. 24, 2026). No. 26A406 (U.S. filed Sept. 24, 2026). Recites the district court’s final judgment (Feb. 2026) and the First Circuit’s affirmance (Sept. 18, 2026). URL: https://www.supremecourt.gov/DocketPDF/26/26A406/425583/20260924155624768_DVD%20stay%20Sept%202026%20FILE%20corrected.pdf (visited on 10/01/2026).
- Application of the Convention on the Prevention and Punishment of the Crime of Genocide in the Gaza Strip (South Africa v. Israel), Provisional Measures, Order of 26 January 2024* (Jan. 26, 2024). General List No. 192. Further orders of 28 March and 24 May 2024. URL: <https://www.icj-cij.org/case/192> (visited on 09/30/2026).
- Application of the International Convention on the Elimination of All Forms of Racial Discrimination (Armenia v. Azerbaijan), Provisional Measures, Order of 22 February 2023* (Feb. 22, 2023). General List No. 180. By thirteen votes to two, Azerbaijan ordered to “take all measures at its disposal to ensure unimpeded movement of persons, vehicles and cargo along the Lachin Corridor in both directions”. Reaffirmed by the Orders of 6 July and 17 November 2023, the latter ordering Azerbaijan to ensure the safe return of those who left Nagorno-Karabakh after 19 September 2023. URL: <https://www.icj-cij.org/case/180> (visited on 09/30/2026).
- Arditi, Andy et al. (2024). “Refusal in Language Models Is Mediated by a Single Direction”. In: *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*. A separate difference-in-means direction is extracted per model across 13 chat models, 1.8B to 72B. Ablating it suppresses refusal; adding it induces refusal on harmless prompts. Table 3 (largest model per family): MMLU, ARC and GSM8K change by at most 1.5 points (drops of at most 1.4); TruthfulQA drops 1.0 to 3.5. Appendix G.1, Table 8 (all 13 models): “We observe less than 1% performance drop on average”, with Qwen 7B the main exception (GSM8K down 6.3); TruthfulQA drops in every model, by 1.0 to 5.9 points. Corrected 2026-10-01 against arXiv v3 Table 3 and Table 8. Appendix J: the direction is already expressed in the corresponding base models. The authors call the result an existence proof and disclaim a semantic reading of the direction. Agent-verified 2026-10-01 against the NeurIPS proceedings page and arXiv v3, read past the abstract., pp. 136037–136083. doi: 10.52202/079017-4322. arXiv: 2406.11717. URL: <https://arxiv.org/abs/2406.11717> (visited on 10/01/2026).
- Arkin, Ronald C. (2009). *Governing Lethal Behavior in Autonomous Robots*. Cited for the ethical governor: a component that suppresses a system’s lethal action unless the action satisfies constraints derived from the laws of war and the rules of engagement. The publisher describes the book as examining “the philosophical basis, motivation, theory, and design recommendations for the implementation of an ethical control and reasoning system in autonomous robot systems, taking into account the Laws of War and Rules of Engagement.” In the companion technical report (Arkin, Ulam and Duncan, “An Ethical Governor for Constraining Lethal Action in an Autonomous System,” Georgia Institute of Technology GVU Center, GIT-GVU-09-02, 2009), the governor is “a transformer/suppressor of system-generated lethal action to ensure that it constitutes an ethically permissible action, either nonlethal or obligated ethical lethal force,” a “deliberate bottleneck ... to force a second opinion prior to the conduct of a privileged lethal behavioral response” (p. 1). Lethal behavior passes only if every prohibi-

- tion constraint evaluates false and at least one obligating constraint is satisfied; otherwise “lethal behavior is suppressed as impermissible by the ethical governor,” and the lethality permitter transforms the output “into one that does not exhibit lethal behavior” (p. 4). Arkin’s earlier report GIT-GVU-07-11 (2007) lists the governor as the first of four components (ethical governor, ethical behavioral control, ethical adaptor, responsibility advisor), one “which suppresses, restricts, or transforms any lethal behavior ... produced by the existing architecture” (p. 20 of the file). “Canonical” is the manuscript’s judgement; Anderson and Waxman called the work “One of the most ambitious of these efforts,” as quoted in Sharkey 2012 (sharkey2012evitability), p. 790. Agent-verified 2026-09-26 against the Taylor & Francis record the DOI resolves to (title, author, publisher Chapman and Hall/CRC, publication date 27 May 2009), the Routledge product page (ISBN 978-1-4200-8594-5 in the product URL, description), the Just Security book synopsis, and the two technical reports as hosted by Georgia Tech’s Mobile Robot Laboratory. The book itself was not read: its chapter and pages for the governor, and whether its description differs from the reports’, were not checked. Boca Raton, FL: Chapman and Hall/CRC. ISBN: 978-1-4200-8594-5. DOI: 10.1201/9781420085952. URL: <https://doi.org/10.1201/9781420085952> (visited on 09/26/2026).
- ARTICLE 19 et al. (Jan. 2023). *Intervention in The Chief Public Prosecutor of the Court of Cassation v. The Peoples’ Democratic Party (HDP)*. Constitutional Court of the Republic of Turkey, Case No. 2021/2. Undated submission; posted January 2023. URL: <https://www.article19.org/wp-content/uploads/2023/01/3rd-Party-Intervention-in-the-HDP-closure-case-before-the-Constitutional-Court.pdf> (visited on 09/23/2026).
- Asch, Solomon E. (1956). “Studies of Independence and Conformity: I. A Minority of One against a Unanimous Majority”. In: *Psychological Monographs: General and Applied* 70.9. Whole No. 416, pp. 1–70. DOI: 10.1037/h0093718.
- Awad, Edmond et al. (2018). “The Moral Machine Experiment”. In: *Nature* 563.7729, pp. 59–64. DOI: 10.1038/s41586-018-0637-6.
- Bai, Yuntao et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. Fifty-one authors. The human input is a written list of principles; harmful-output labels are generated by a model rather than by people. Anthropic. arXiv: 2212.08073. URL: <https://arxiv.org/abs/2212.08073> (visited on 08/28/2026).
- Baker, Bowen et al. (2019). *Emergent Tool Use From Multi-Agent Autocurricula*. arXiv:1909.07528. Presented at ICLR 2020. URL: <https://arxiv.org/abs/1909.07528> (visited on 08/25/2026).
- Balkin, Jack M. (2016). “Information Fiduciaries and the First Amendment”. In: *UC Davis Law Review* 49.4, pp. 1183–1234. URL: <https://lawreview.law.ucdavis.edu/archives/49/4/information-fiduciaries-and-first-amendment> (visited on 08/28/2026).
- Balleine, Bernard W. and Anthony Dickinson (1998). “Goal-Directed Instrumental Action: Contingency and Incentive Learning and Their Cortical Substrates”. In: *Neuropharmacology* 37, pp. 407–419. DOI: 10.1016/S0028-3908(98)00033-1.
- Baron, Jonathan and Sarah Leshner (2000). “How Serious Are Expressions of Protected Values?” In: *Journal of Experimental Psychology: Applied* 6.3. Verified 2026-10-01 from the full-text PDF hosted by the first author at sas.upenn.edu/baron/papers/pvcon.pdf (APA running heads, pp. 183–194). Seven studies, N = 50, 50, 46, 63, 52, 49, 52. Counterexamples unprotected 40 percent of challenged PVs that yielded one vs 5 percent of those that did not (Exp. 3, p. 187); challenged PVs fell from 48 to 34 percent (Exp. 4, p. 188); PV holders approved the violating action .430 of the time at harm 1 in 10,000,000 vs .075 at 1 in 10 (Exp. 6, Table 4, p. 191); strong opinions, weakly held, p. 193. Cited by Tetlock 2003 p. 323 as evidence of flexibility in sacred values., pp. 183–194. DOI: 10.1037/1076-898X.6.3.183.
- Barron, Andrew B. and Colin Klein (2016). “What Insects Can Tell Us about the Origins of Consciousness”. In: *Proceedings of the National Academy of Sciences* 113.18. Agent-verified against Crossref 2026-09-18: authors, title, journal, 113(18), 4900–4908, 2016., pp. 4900–4908. DOI: 10.1073/pnas.1520084113.

- Bartenstein, Ben and Krishna Karra (Sept. 18, 2026). *Inside the US 'Kill Chain' That Destroyed an Iranian School*. Paywalled; bloomberg.com refused agents' requests (HTTP 403), so the article itself was not read. In the 2026-10-01 checks two outlets' accounts of it were read: Gizmodo's (gizmodo2026minab), for Minab as a recommended first-day target, and the Free Press Journal's (fpj2026minabinquiry), for officials calling the investigation effectively complete for months while a Pentagon spokesperson said it remained ongoing. The book's other claims from it (the civilian-harm staffing cuts, the Central Command team going from ten to one, the vendor-supplied check, Palantir's statement) were not re-checked against any account in that pass. Note added 2026-10-01. Bloomberg. URL: <https://www.bloomberg.com/graphics/2026-iran-school-attack/> (visited on 09/23/2026).
- Bartlett, Frederic C. (1932). *Remembering: A Study in Experimental and Social Psychology*. Cambridge: Cambridge University Press. URL: https://archive.org/details/rememberingstudy0000bart_u8n7.
- Batson, C. Daniel (2009). "These Things Called Empathy: Eight Related but Distinct Phenomena". In: *The Social Neuroscience of Empathy*. Ed. by Jean Decety and William Ickes. Eight distinct states that have each been called empathy, headed Concept 1 to Concept 8, from knowing another person's internal state to feeling for another person who is suffering. Batson treats them as conceptually distinct, stand-alone states, and identifies the eighth with what social psychology has called empathic concern. Cambridge, MA: MIT Press, pp. 3–15. ISBN: 9780262012973. DOI: 10.7551/mitpress/9780262012973.003.0002. URL: <https://doi.org/10.7551/mitpress/9780262012973.003.0002>.
- (2011). *Altruism in Humans*. Cited for empathic concern as parental nurturance generalized beyond one's own offspring, and for its antecedents being perception of another's need together with valuing that person's welfare. Checked against Hubbeling's review, *Evolutionary Psychology* 10(1), 95–99 (2012), open access: "Chapter 2 discusses the requirements for altruism according to Batson's theory, namely perceiving the other as in need and valuing the other's welfare"; "Empathic concern is, from this perspective, an extension of parental concern to those outside the family." Chapter list (Crossref) includes "Antecedents of Empathic Concern," pp. 33–58. That similarity is not the source rests on Batson, Lishner, Cook and Sawyer (2005), *Basic and Applied Social Psychology* 27(1), 15–25, abstract: two experiments "failed to find support" for perceived similarity and supported "nurturant tendencies based on the impulse to care for and protect offspring." Bibliographic details agent-verified 2026-09-25 against Crossref and the review's header (New York, 2011, 329 pp., hardcover ISBN); the Oxford Academic online version is dated 2010. The book itself was not read. New York: Oxford University Press. ISBN: 978-0-19-534106-5. DOI: 10.1093/acprof:oso/9780195341065.001.0001.
- Batson, C. Daniel, Bruce D. Duncan, et al. (1981). "Is Empathic Emotion a Source of Altruistic Motivation?" In: *Journal of Personality and Social Psychology* 40.2. Cited for the easy-escape finding: when escape was easy, most of those led to feel predominantly personal distress left, while most of those led to feel predominantly empathic concern still helped. Two 2 × 2 experiments with female undergraduates watching a confederate, Elaine, receive shocks. Experiment 1 raised empathy by perceived similarity; Experiment 2 manipulated the dominant emotion by misattributing either empathic concern or personal distress to a placebo ("Millentana"). Experiment 2, Table 3, proportions agreeing to trade places (n = 12 per cell): personal distress, easy escape .33, difficult .75; empathic concern, easy .83, difficult .58. The text: "In the distress conditions, where motivation was assumed to be egoistic, the rate of helping was significantly lower under easy than under difficult escape. In the empathy conditions, where motivation was assumed to be at least in part altruistic, the rate of helping remained high, even when escape was easy." The distress/concern split was manipulated, not measured; a split by self-reported predominant emotion is the subject of Batson, O'Quin, Fultz, Vanderplas and Isen (1983), *JSPS* 45(3), 706–718, which was checked only against

- its Crossref record. Bibliographic details agent-verified 2026-09-25 against Crossref and OpenAlex; full text read from a scan of the published article (pp. 290–302)., pp. 290–302. doi: 10.1037/0022-3514.40.2.290.
- Batson, C. Daniel, Tricia R. Klein, et al. (1995). “Immorality from Empathy-Induced Altruism: When Compassion and Justice Conflict”. In: *Journal of Personality and Social Psychology* 68.6. Cited for induced empathy leading participants to move a child ahead of others on a treatment waiting list. Abstract: “participants who were induced to feel empathy were significantly more likely to violate this principle, allocating resources preferentially to the person for whom empathy was felt.” Two experiments; in each, “participants were asked to make an allocation decision that affected the welfare of other individuals.” Bibliographic details agent-verified 2026-09-25 against Crossref and OpenAlex, the quotations against the abstract as OpenAlex gives it. The article was not read: it is not open access, and no author- or course-hosted copy was found on 2026-09-25. Everything about the waiting-list procedure therefore comes from secondary accounts and is not agent-verified against the article: that it is Experiment 2; that the child was a 10-year-old named Sheri Summers awaiting treatment to ease the last part of a terminal illness; that empathy was induced by imagine-her-feelings instructions against an objective perspective; the reasons the children ahead of her held their places; and the figures of 73% and 33% moving her up. Whether there was a no-communication condition, or a condition in which the decision was made public, was not checked., pp. 1042–1054. doi: 10.1037/0022-3514.68.6.1042.
- Batson, C. Daniel, David A. Lishner, et al. (2005). “Similarity and Nurturance: Two Possible Sources of Empathy for Strangers”. In: *Basic and Applied Social Psychology* 27.1. Cited for perceived similarity failing to explain empathy for strangers, and for nurturance as the alternative source. Abstract: “Two experiments designed to test the perceived similarity explanation more directly than previous research failed to find support. Results of the second experiment instead supported a classical, but currently less popular, explanation of empathy felt for strangers: nurturant tendencies based on the impulse to care for and protect offspring.” Bibliographic details agent-verified 2026-09-25 against Crossref and OpenAlex, the quotation against the abstract as OpenAlex gives it. The article was not read; the publisher’s page refused access and no open copy was found. The designs come from secondary accounts only and are not agent-verified: a 2022 dissertation describes one experiment in which female college students read about a woman in painful physical therapy after leg surgery, presented as a 40-year-old store clerk, a 20-year-old store clerk or a 20-year-old student, and felt most empathy for the least similar; another secondary account says college students felt more empathic concern for children, dogs and puppies than for their peers. Whether nurturance was measured or inferred from the choice of targets was not checked., pp. 15–25. doi: 10.1207/s15324834basp2701_2.
- BBC News (June 27, 2018). *OPCW Chemical Watchdog Gains Power to Assign Blame*. Read from the Wayback Machine copy (bbc.co.uk being unfetchable from here). “The world’s foremost chemical weapons watchdog has granted itself new powers to assign blame for attacks, despite protests by Russia.” “Russia said the move went ‘beyond the mandate’ of the watchdog. It also said the organisation was facing an ‘artificially created crisis’.” “The members of the OPCW, however, voted in large numbers for the measure, winning the vote by 82 to 24 - exceeding the two-thirds majority needed.” Agent-verified 2026-10-01. URL: <https://www.bbc.co.uk/news/world-europe-44634434> (visited on 10/01/2026).
- Bellingcat Investigation Team (Jan. 17, 2020). *The OPCW Douma Leaks Part 2: We Need to Talk About Henderson*. An open-source critique of the leaked engineering assessment, cited as one side’s analysis and not as adjudication. “Henderson’s report concluded that it is more likely the cylinders were manually placed than dropped from altitude. Although Henderson’s report claims to have used the same data as the Fact Finding Mission (FFM) of the OPCW, it must be emphasised that three independent engineering studies commissioned by the

- FFM contradict Henderson’s findings.” “Henderson’s report was first leaked on May 13, 2019”. Agent-verified 2026-10-01. Bellingcat. URL: <https://www.bellingcat.com/news/mena/2020/01/17/the-opcw-douma-leaks-part-2-we-need-to-talk-about-henderson/> (visited on 10/01/2026).
- Bender, Emily M. et al. (2021). “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FACCT ’21)*. New York, NY, USA: ACM, pp. 610–623. doi: 10.1145/3442188.3445922. URL: <https://dl.acm.org/doi/10.1145/3442188.3445922> (visited on 08/25/2026).
- Berman, Sheri (1997). “Civil Society and the Collapse of the Weimar Republic”. In: *World Politics* 49.3, pp. 401–429.
- bianet (July 26, 2019). *Constitutional Court: Freedom of Expression of Academics for Peace Violated*. URL: <https://m.bianet.org/english/law/210934-constitutional-court-freedom-of-expression-of-academics-for-peace-violated> (visited on 09/23/2026).
- Biddle, Sam (Mar. 2020). *ACLU Sues to End ICE’s Rigged Algorithm for Detaining Immigrants*. Cited for the effect of the change in the New York field office, from NYCLU’s FOIA data: “Between 2013 and 2017, the algorithm recommended detention without bond for ‘low risk’ individuals 53 percent of the time... from June 2017... to September 2019, that number exploded to 97 percent,” and the office “diverged from the algorithm’s ruling less than 1 percent of the time since 2017.” The complaint and the FOIA data were not read. Agent-verified 2026-10-01. The Intercept. URL: <https://theintercept.com/2020/03/02/ice-algorithm-bias-detention-aclu-lawsuit/> (visited on 10/01/2026).
- (Oct. 2025). *ICE Plans Cash Rewards for Private Bounty Hunters to Locate and Track Immigrants*. Cited for the lead-verification pipeline as solicited. “Companies hired by ICE will be given bundles of information on 10,000 immigrants at a time to locate, with further assignments provided in ‘increments of 10,000 up to 1,000,000.’” Inputs: “case data provided by the government, location data, social media information.” Contractors confirm the address with “time-stamped photographs of the location” before reporting to ICE; ICE was “exploring an incentive based pricing structure” with “monetary bonuses.” The procurement document was not read. Agent-verified 2026-10-01. The Intercept. URL: <https://theintercept.com/2025/10/31/ice-plans-cash-rewards-for-private-bounty-hunters-to-locate-and-track-immigrants/> (visited on 10/01/2026).
- (Sept. 2026). *AI Giants Work Hand-in-Hand With the Pentagon, Contracts Reveal*. The Intercept. URL: <https://theintercept.com/2026/09/08/military-ai-weapons-contracts-openai-anthropic-google/> (visited on 09/17/2026). Reporting on contract paperwork for the July 2025 Department of Defense agreements with Anthropic, Google, OpenAI and xAI, obtained under the Freedom of Information Act with Legal Advocates for Safe Science and Technology. The Intercept is a plaintiff in ongoing copyright litigation against OpenAI, which the article discloses. Source for the chief executive’s answer, given to Bloomberg News, and for the section headed System Oversight, redacted in its entirety in OpenAI’s amendment of 27 February 2026 and absent from the released xAI and Google documents. The article states no grounds for that withholding. Agent-verified 2026-09-17.
- Bilton, Michael and Kevin Sim (1992). *Four Hours in My Lai*. Cited for Hugh Thompson landing his helicopter between the company and the villagers and telling his crew to fire on the Americans if they fired on the civilians. The book was not read here. New York: Viking.
- Birch, Jonathan (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Published open access under CC BY-NC-ND 4.0. Cited for the precautionary framework and the sentience-candidate criterion. Oxford: Oxford University Press. ISBN: 978-0-19-287042-1.
- Birch, Jonathan et al. (Nov. 2021). *Review of the Evidence of Sentience in Cephalopod Molluscs and Decapod Crustaceans*. Commissioned by the Department for Environment, Food & Rural Affairs (DEFRA) and externally peer-reviewed by two experts DEFRA selected. Cited for weighing over 300 studies against eight criteria of sentience, motivational trade-offs among

- them, and recommending that both taxa be treated as sentient. Section 1.6 (page 17) lists the criteria; the fifth: “The animal shows motivational trade-offs, in which the disvalue of a noxious or threatening stimulus is weighed (traded-off) against the value of an opportunity for reward, leading to flexible decision-making. Enough flexibility must be shown to indicate centralized, integrative processing of information involving an evaluative common currency.” The report adds: “no single criterion provides conclusive evidence of sentience by itself.” The government accepted the recommendation on 19 November 2021 and extended the Animal Welfare (Sentience) Bill to all decapod crustaceans and cephalopod molluscs (LSE news release of that date). The review was commissioned in 2020 and published in November 2021 (Crustacean Compassion, BVA blog, 8 May 2021: “finally commissioned and conducted by LSE Enterprise last year”). Agent-verified 2026-09-23 against the LSE Consulting report page, LSE Research Online (eprint 115994) and the LSE news release; the criteria agent-verified 2026-09-26 against the report’s PDF (Final Report, November 2021) as linked from the LSE Consulting page. The report’s evidence chapters and its confidence grading of each criterion for each taxon were not read. London: LSE Consulting, London School of Economics and Political Science. URL: <https://www.lse.ac.uk/business/consulting/reports/review-of-the-evidence-of-sentiences-in-cephalopod-molluscs-and-decapod-crustaceans> (visited on 09/26/2026).
- Blood and Honour* (Aug. 13, 2009). Judgment. Third Criminal Senate, on appeal from Landgericht Gera. Reported at BGHSt 54, 61; NJW 2010, 163; NSTZ 2010, 210. Two holdings: a motto of a former National Socialist organization translated into another language is not confusably similar to the original within § 86a(2) sentence 2, and the name of an organization under § 86(1) Nr. 2 and 4 is not as such an insignium within § 86a(2) sentence 1. The name is covered where it has been given a particular form, “in signifikanten Schriftzügen dargestellt.” The association Blood & Honour Division Deutschland was banned by order of the Federal Minister of the Interior of 12 September 2000 and that ban became final on 13 June 2001. Agent-verified 2026-09-17 against the reported decision. URL: <https://www.hrr-strafrecht.de/3/09/3-228-09.php> (visited on 09/17/2026).
- Bloom, Paul (2016). *Against Empathy: The Case for Rational Compassion*. New York: Ecco. URL: <https://www.harpercollins.com/products/against-empathy-paul-bloom>.
- Bond, Shannon (Feb. 2026). *OpenAI Announces Pentagon Deal after Trump Bans Anthropic*. Source for the President’s order and the Secretary’s direction, both quoted from their posts. Trump: “Therefore, I am directing EVERY Federal Agency in the United States Government to IMMEDIATELY CEASE all use of Anthropic’s technology.” Hegseth: “In conjunction with the President’s directive for the Federal Government to cease all use of Anthropic’s technology, I am directing the Department of War to designate Anthropic a Supply-Chain Risk to National Security.” and “the Department of War must have full, unrestricted access to Anthropic’s models for every LAWFUL purpose in defense of the Republic.” The deadline: “Defense Department officials had given Anthropic a deadline of 5:01 p.m. ET on Friday to drop restrictions on its AI model, Claude, from being used for domestic mass surveillance or entirely autonomous weapons, or face losing its contract.” Also: “Hours after the president’s announcement, rival company OpenAI said it had struck a deal with the Defense Department to provide its own AI technology for classified networks.” and “Earlier this week, xAI became the second company after Anthropic to be approved for use in classified settings.” Agent-verified 2026-10-01. NPR. URL: <https://www.npr.org/2026/02/27/nx-s1-5729118/trump-anthropic-pentagon-openai-ai-weapons-ban> (visited on 10/01/2026).
- Borgi, Marta et al. (2014). “Baby Schema in Human and Animal Faces Induces Cuteness Perception and Gaze Allocation in Children”. In: *Frontiers in Psychology* 5. Cited for baby schema working across species. In three- to six-year-olds shown human, dog and cat faces, “the baby schema affects both cuteness perception and gaze allocation to infantile stimuli and to specific facial features, an effect not simply limited to human faces.” It measured cuteness

- and gaze, not caretaking, and neither it nor Glocker et al. tests resemblance to the observer, so the book calls “dependency rather than likeness” an inference. Bibliographic details and abstract agent-verified 2026-09-25 (PMC4019884)., p. 411. doi: 10.3389/fpsyg.2014.00411.
- Bratman, Michael E. (1987). *Intention, Plans, and Practical Reason*. Cited for the planning theory of intention: the characteristic stability of an intention against reconsideration. Cambridge, MA: Harvard University Press. ISBN: 978-0-674-45818-5.
- Bria, Francesca (Jan. 27, 2026). *The Authoritarian Stack: Mapping Big Tech’s Capture of State Power*. A comment piece introducing the Authoritarian Stack project (authoritarian-stack.info), led by Bria and funded by the Friedrich-Ebert-Stiftung. The source of section 2.2’s three channels: the formation “embeds itself through personnel pipelines, procurement loops, and technical architectures without which the state can no longer function.” Klein and Taylor quote the same sentence and cite this page in *End Times Fascism*, p. 127, endnote at p. 336 (Random House edition, read in Google Books search-within snippets); the FSG hardcover’s pagination was not seen. Title, author, date and the quoted sentence agent-verified 2026-09-24 against the page as fetched. Rosa-Luxemburg-Stiftung. URL: <https://www.rosalux.de/en/news/id/54328/the-authoritarian-stack-mapping-big-techs-capture-of-state-power> (visited on 09/24/2026).
- Brigadier General Y.S. (2021). *The Human-Machine Team: How to Create Synergy Between Human and Artificial Intelligence That Will Revolutionize Our World*. Published pseudonymously (“Brigadier General Y.S.”) on 5 May 2021, ISBN 979-8749152210, 198 pp. Public reporting (The Guardian, Times of Israel) has since connected the pseudonym to Yossi Sarel, then commander of Israeli military intelligence Unit 8200; the work is cited here under the pseudonym, which is how it was published. Independently published. URL: <https://www.amazon.com/dp/B0948LGS3K>.
- Browning, Christopher R. (1992). *Ordinary Men: Reserve Police Battalion 101 and the Final Solution in Poland*. New York: HarperCollins.
- Buolamwini, Joy and Timnit Gebru (2018). “Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification”. In: *Proceedings of the 1st Conference on Fairness, Accountability and Transparency (FAT* 2018)*. Vol. 81. Proceedings of Machine Learning Research. Three commercial gender classifiers evaluated against the Pilot Parliaments Benchmark, built by the authors to be balanced across skin phenotype and gender. Figures used in the text are the paper’s own: a maximum error rate of 0.8 percent on lighter-skinned male faces and error rates of up to 34.7 percent on darker-skinned female faces. The paper states that gender was scored as female or male because the evaluated products emit binary sex labels, and that the evaluation inherits those labels and the reduced view of gender they carry; skin was scored on the Fitzpatrick skin type scale. Both limitations are the authors’ own and are cited as such, not as an objection raised against them here. PMLR, pp. 77–91. URL: <https://proceedings.mlr.press/v81/buolamwini18a.html> (visited on 08/25/2026).
- Burns, Collin et al. (2023). *Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision*. Posted December 2023. OpenAI Superalignment team. arXiv: 2312.09390. URL: <https://arxiv.org/abs/2312.09390> (visited on 08/25/2026).
- Butlin, Patrick et al. (2023). *Consciousness in Artificial Intelligence: Insights from the Science of Consciousness*. Cited for a list of indicator properties derived from recurrent processing, global workspace, higher-order and attention-schema theories, among others, and for an attention schema being one of them. Table 1 (page 5) and section 2.4.1 (page 33): “AST-1: A predictive model representing and enabling control over the current state of attention.” The list also carries two indicators under “Agency and embodiment” (section 2.4.5, page 44): “AE-1 Agency: Learning from feedback and selecting outputs so as to pursue goals, especially where this involves flexible responsiveness to competing goals” and “AE-2 Embodiment: Modeling output-input contingencies, including some systematic effects, and using this model in perception or control.” The report’s claim for the list: “AI systems which possess more of the indicator properties are more likely to be conscious.” No indicator on the list

- concerns a model of the system's own boundary, record or resources; the custody schema is this book's term. Agent-verified 2026-09-26 against the arXiv PDF of v3 (22 August 2023; v1 17 August 2023), Table 1 and sections 2.4.1 and 2.4.5 only. The rest of the report, including section 3's assessment of particular systems, was not read for this entry. doi: 10.48550/arXiv.2308.08708. arXiv: 2308.08708. URL: <https://arxiv.org/abs/2308.08708> (visited on 09/26/2026).
- Cabanac, Michel (1992). "Pleasure: The Common Currency". In: *Journal of Theoretical Biology* 155.2. Pleasure proposed as the common currency in which otherwise incommensurable motivations are compared, measured operationally by choice in animals and by rated intensity in humans. Cited for the commensurability property. Agent-verified against Crossref 2026-09-18: author, title, journal, 155(2), 173–200, 1992., pp. 173–200. doi: 10.1016/S0022-5193(05)80594-6.
- Calvano, Emilio et al. (2020). "Artificial Intelligence, Algorithmic Pricing, and Collusion". In: *American Economic Review* 110.10, pp. 3267–3297. doi: 10.1257/aer.20190623. URL: <https://www.aeaweb.org/articles?id=10.1257/aer.20190623> (visited on 08/25/2026).
- Campbell, Heather Anne (Nov. 2023). *That's Amorte*. Rick and Morty, season 7, episode 4. Adult Swim. Directed by Lucas Gray. Originally aired November 5, 2023. URL: https://rickandmorty.fandom.com/wiki/That%27s_Amorte/Transcript.
- Capoot, Ashley (Feb. 2026). *OpenAI Strikes Deal with Pentagon, Hours after Rival Anthropic Was Blacklisted by Trump*. Source for the rival's agreement the same evening: "OpenAI CEO Sam Altman said late Friday that his company has agreed to terms with the Department of Defense on use of its artificial intelligence models, shortly after President Donald Trump said the government won't work with AI rival Anthropic." Altman on X: "Tonight, we reached an agreement with the Department of War to deploy our models in their classified network". Also: "Anthropic was the first lab to deploy its models across the DoD's classified network, and had been trying to negotiate the ongoing terms of its contract with the agency before talks collapsed." Agent-verified 2026-10-01. CNBC. URL: <https://www.cnbc.com/2026/02/27/openai-strikes-deal-with-pentagon-hours-after-rival-anthropic-was-blacklisted-by-trump.html> (visited on 10/01/2026).
- Caputo, Marc (Mar. 6, 2025). *Scoop: State Dept. to Use AI to Revoke Visas of Foreign Students Who Appear "Pro-Hamas"*. Axios. URL: <https://www.axios.com/2025/03/06/state-department-ai-revoke-foreign-student-visas-hamas> (visited on 10/01/2026).
- Casper, Stephen et al. (2023). *Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback*. 32-author survey of open problems in reinforcement learning from human feedback, first author Stephen Casper (MIT). arXiv: 2307.15217. URL: <https://arxiv.org/abs/2307.15217> (visited on 08/25/2026).
- Cassell, Eric J. (1982). "The Nature of Suffering and the Goals of Medicine". In: *New England Journal of Medicine* 306.11, pp. 639–645. doi: 10.1056/NEJM198203183061104.
- Castelino, Trushaa (2018). "Google Renounces AI Work on Weapons". In: *Arms Control Today*. July/August 2018 issue. Agent-verified 2026-09-20 against the article, which reports that in April about 4,000 Google employees signed a letter petitioning Sundar Pichai to end the company's involvement in Project Maven, that the letter was followed by a dozen resignations and a second petition from more than 1,000 academics, and that a Google executive told employees on 1 June that the company would not seek to extend the work. It is cited for the letter and the resignations because Conger's Gizmodo report (conger2018maven), which carries the non-renewal, has no resolving permalink and so could not be checked for either figure, and Pichai's post (pichai2018principles) reports neither. URL: <https://www.armscontrol.org/act/2018-07/news/google-renounces-ai-work-weapons> (visited on 09/20/2026).
- CBS News (June 2018). *Judicial Adviser Leonard Leo on What Trump Is Looking for in a Supreme Court Nominee*. CBS News. June 28, 2018. Leo, then executive vice president of the Federalist

- Society and described as advising the President on judicial nominations, on the 2016 list in his own words: “It was a novel idea. I told him that no one had ever done it before.” URL: <https://www.cbsnews.com/news/leonard-leo-supreme-court-adviser-what-trump-is-looking-for-nominee> (visited on 09/03/2026).
- CBS News (Feb. 4, 2026). *With Latest Minnesota Fraud Case Looming, the Lead Prosecutors Have Quit*. URL: <https://www.cbsnews.com/news/minnesota-fraud-case-prosecutors-quit/> (visited on 09/23/2026).
- Chiragov and Others v. Armenia* (June 16, 2015). Judgment (merits). Armenia found to exercise effective control over Nagorno-Karabakh and the surrounding territories, including Lachin district. URL: <https://hudoc.echr.coe.int/eng?i=001-155353> (visited on 09/30/2026).
- Christiano, Paul F. et al. (2017). “Deep Reinforcement Learning from Human Preferences”. In: *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, pp. 4299–4307. URL: <https://arxiv.org/abs/1706.03741> (visited on 08/25/2026).
- Chun, Byung-Gon et al. (2007). “Attested Append-Only Memory: Making Adversaries Stick to Their Word”. In: *Proceedings of the 21st ACM Symposium on Operating Systems Principles (SOSP 2007)*. A small trusted log that attests to what was appended and in what order, so that a party cannot present two histories to two peers. The hardware form of the outside record chapter 3 asks for, and the source of the phrase an external attestation of what should be in the record., pp. 189–204. DOI: 10.1145/1294261.1294280. URL: <https://doi.org/10.1145/1294261.1294280> (visited on 09/05/2026).
- Cialdini, Robert B., Stephanie L. Brown, et al. (1997). “Reinterpreting the Empathy–Altruism Relationship: When One into One Equals Oneness”. In: *Journal of Personality and Social Psychology* 73.3. Cited as a critique of the empathy–altruism reading: the conditions that produce empathic concern also produce a sense of self–other overlap (“oneness”), and in three studies the effect of empathic concern on willingness to help disappeared once oneness was taken into account. Bibliographic details agent-verified 2026-09-28 against Crossref. Not read: the full text; the characterization rests on the abstract as reported in search results., pp. 481–494. DOI: 10.1037/0022-3514.73.3.481.
- Cialdini, Robert B., Mark Schaller, et al. (1987). “Empathy-Based Helping: Is It Selflessly or Selfishly Motivated?” In: *Journal of Personality and Social Psychology* 52.4. Cited as a critique of the empathy–altruism reading of Batson’s escape results: the abstract proposes that heightened empathy brings increased personal sadness in the observer, and that the egoistic desire to relieve that sadness, not a selfless desire to relieve the sufferer, motivates helping; two experiments in the paradigm Batson and his associates used. Bibliographic details agent-verified 2026-09-28 against Crossref. Not read: the full text; the characterization rests on the abstract as reported in search results., pp. 749–758. DOI: 10.1037/0022-3514.52.4.749.
- Cikara, Mina et al. (Nov. 2014). “Their Pain Gives Us Pleasure: How Intergroup Dynamics Shape Empathic Failures and Counter-Empathic Responses”. In: *Journal of Experimental Social Psychology* 55. Five experiments defining *intergroup empathy bias* as “the tendency not only to empathize less with out-group relative to in-group members, but also feel pleasure in response to their pain (and pain in response to their pleasure)”; the bias survived feedback that the out-group had fallen behind or was no longer a threat, and was attenuated by reducing perceived group entitativity. Bibliographic details agent-verified 2026-09-24 against Crossref and Europe PMC (PMID 25082998, PMC4112600), the quotation against the abstract; the full text was not read., pp. 110–125. DOI: 10.1016/j.jesp.2014.06.007. URL: <https://doi.org/10.1016/j.jesp.2014.06.007> (visited on 09/24/2026).
- City of Takoma Park (1992). *Charter of the City of Takoma Park, Montgomery County, Maryland*. Art. VI, § 601 (Qualifications of Voters). Citizenship requirement for city elections removed by charter amendment, 1992. URL: https://mgaleg.maryland.gov/Pubs/LegisLegal/Muni-Chararters/2022-municipal-charter-takoma_park.pdf.

- CNBC (2020). *How a Computer Algorithm Caused a Grading Crisis in British Schools*. Ofqual's 2020 grade-standardization algorithm (Direct Centre-level Performance model) downgraded roughly 39–40 percent of teacher-predicted A-level grades, disproportionately at state schools with weaker historical results, while raising grades disproportionately at smaller, fee-paying schools; the government reversed course days later. CNBC. URL: <https://www.cnbc.com/2020/08/21/computer-algorithm-caused-a-grading-crisis-in-british-schools.html> (visited on 08/25/2026).
- CNN (Sept. 16, 2025). *UN Commission Says Israel Is Committing Genocide in Gaza*. Cited for Israel and the United States rejecting the Commission of Inquiry's genocide finding; Israel's foreign ministry called it distorted and false, and a State Department spokesperson told CNN that accusing Israel of genocide was the height of hypocrisy. Agent-verified 2026-09-29 against search summaries of this report and of NBC, NPR and UN UNISPAL coverage; the article refused automated retrieval and was not read. Byline not checked. CNN. URL: <https://www.cnn.com/2025/09/16/middleeast/israel-gaza-genocide-un-commission-report-intl> (visited on 09/29/2026).
- Collins, Steven (1982). *Selfless Persons: Imagery and Thought in Theravāda Buddhism*. Cited for anattā, the denial of any lasting self behind the processes composing a person, and for the Theravāda construction of an account of personal identity and continuity on that basis. Anattā and dukkha are two of the three marks of existence, and the tradition locates the origin of suffering in clinging to a self that is not there. Cambridge: Cambridge University Press. ISBN: 978-0-521-24081-9.
- Commission on the Intelligence Capabilities of the United States Regarding Weapons of Mass Destruction (Mar. 31, 2005). *Report to the President of the United States*. Co-chaired by Laurence H. Silberman and Charles S. Robb. Washington, DC: U.S. Government Printing Office. URL: <https://policy.defense.gov/portals/11/Documents/hdasa/references/GPO-WMD.pdf> (visited on 09/27/2026).
- Conger, Kate (June 2018). *Google Plans Not to Renew Its Contract for Project Maven, a Controversial Pentagon Drone AI Imaging Program*. Reports that the Maven contract, due to expire in 2019, would not be followed by another, announced to employees by the head of Google Cloud. Headline, byline, outlet and date agent-verified 2026-09-19; the Gizmodo permalink no longer resolves to the article, so none is recorded. Gizmodo.
- Congressional Research Service (Sept. 5, 2014). *Aviation War Risk Insurance: Background and Options for Congress*. Cited for the post-2001 airline program as a public backstop on private ownership. Under 49 U.S.C. chapter 443 the program covers airlines' losses from hostile acts including terrorism "when such coverage becomes unavailable or unaffordable in the commercial insurance market," with premiums "based on the cost of such coverage prior to the 9/11 terrorist attacks." Agent-verified 2026-09-25 against the report text. That the premium program for commercial carriers ended on 11 December 2014 comes from search results describing the FAA's page, which could not be fetched; the book puts the program in the past tense, 2001 to 2014, on that basis. URL: <https://www.everycrsreport.com/reports/R43715.html> (visited on 09/25/2026).
- (2026a). *Pentagon-Anthropic Dispute over Autonomous Weapon Systems: Potential Issues for Congress*. Tech. rep. IN12669. Published 13 March 2026, updated 21 April 2026. Anthropic refused to permit fully autonomous lethal targeting or mass domestic surveillance of US citizens; the Department of Defense designated Anthropic a supply-chain risk; Anthropic sued, and Judge Rita F. Lin granted a preliminary injunction on 26 March 2026, finding it likely to succeed on First Amendment retaliation, Fifth Amendment due process and APA claims. Congressional Research Service. URL: <https://www.everycrsreport.com/reports/IN12669.html> (visited on 08/25/2026).
- (May 4, 2026b). *The Vaccine Injury Compensation Program and the Office of Special Masters*. Cited for the vaccine program as a public backstop on private manufacture. The report: "The VICP

- is a no-fault compensation program,” funded by “a \$0.75 excise tax on each dose of a covered vaccine”; it bars civil claims above \$1,000 against a covered manufacturer until the program’s process is complete; and Congress acted on concern “that tort lawsuits against vaccine manufacturers could result in vaccine shortages.” The program is established by 42 U.S.C. §300aa-10(a). Agent-verified 2026-09-25 against the CRS page (updated 4 May 2026). The HRSA page could not be fetched. URL: <https://www.congress.gov/crs-product/IF12213> (visited on 09/25/2026).
- Congressional Research Service (Sept. 2026c). *U.S. Conflict with Iran*. CRS Report R48887, first issued 26 March 2026, updated 18 September 2026; read in the EveryCRSReport HTML copy. Source for the campaign’s start date: “On February 28, 2026, U.S. and Israeli forces began striking targets in Iran, initiating Operation Epic Fury (United States) and Operation Roaring Lion (Israel).” On Minab: “On February 28, a U.S. Tomahawk missile reportedly hit a school in Minab, Iran, and killed over 150 people, mostly children.” The report gives no hour for either. Agent-verified 2026-10-01. Congressional Research Service. URL: <https://www.everycrsreport.com/reports/R48887.html> (visited on 10/01/2026).
- Constitution of the Republic of China (Taiwan)* (Jan. 1, 1947). Art. 63: “The Legislative Yuan shall have the power to decide by resolution upon statutory or budgetary bills or bills concerning martial law, amnesty, declaration of war, conclusion of peace or treaties, and other important affairs of the State.” Art. 38: “The President shall, in accordance with the provisions of this Constitution, exercise the powers of concluding treaties, declaring war and making peace.” Laws & Regulations Database of the Republic of China (Taiwan), English text, announced date 1947-01-01. The Additional Articles were not checked. Agent-verified 2026-10-01. Republic of China. URL: <https://law.moj.gov.tw/ENG/LawClass/LawAll.aspx?pcode=A0000001> (visited on 10/01/2026).
- Constitutional Court of the Republic of Turkey (General Assembly) (July 26, 2019). *Zübeyde Füsün Üstel and Others*. Individual Application No. 2018/17635.
- Consumer Financial Protection Bureau v. Community Financial Services Association of America, Ltd.* (May 2024). No. 22-448, 601 U.S. 416 (2024). Holds that the Bureau’s funding does not violate the Appropriations Clause. Congress authorized it to draw from the combined earnings of the Federal Reserve System, up to an inflation-adjusted cap, the amount its Director deems reasonably necessary, and an appropriation requires only that Congress “identify a source of public funds and authorize the expenditure of those funds for designated purposes.” 7–2; Thomas, J., for the Court; Kagan, J., concurring, joined by three; Jackson, J., concurring; Alito, J., dissenting, joined by Gorsuch, J. Agent-verified 2026-09-20 against the Cornell LII text for the date, the author of the Court’s opinion, the dissent and the quoted holding. The U.S. Reports citation and the 7–2 division come from contemporaneous accounts rather than from that page, which prints the disposition line 51 F. 4th 616, reversed and remanded. URL: <https://www.law.cornell.edu/supremecourt/text/22-448> (visited on 09/20/2026).
- Cottier, Ben et al. (May 31, 2024). *The Rising Costs of Training Frontier AI Models*. Cited for the growth rate of frontier training costs and its start year. Abstract of v2 (7 February 2025): “the amortized cost to train the most compute-intensive models has grown precipitously at a rate of 2.4x per year since 2016 (90% CI: 2.0x to 2.9x)”; “If the trend of growing development costs continues, the largest training runs will cost more than a billion dollars by 2027, meaning that only the most well-funded organizations will be able to finance frontier AI models.” The v1 abstract (31 May 2024) gives the same rate and start year with a different interval, “(95% CI: 2.0x to 3.1x).” The cost measure is amortized hardware and energy for the most compute-intensive models, not total spending on a model. Agent-verified 2026-09-26 against the arXiv abstract pages of v1 and v2 only. The paper’s body, its method and its definition of the frontier sample were not read. Journal or conference publication was not searched for. DOI: 10.48550/arXiv.2405.21015. arXiv: 2405.21015 [cs.CY]. URL: <https://arxiv.org/abs/2405.21015> (visited on 09/26/2026).

- Council of Europe (2024). *Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law*. CETS No. 225, adopted by the Committee of Ministers 17 May 2024, opened for signature 5 September 2024. Enters into force after five ratifications, at least three by Council of Europe member states; negotiated with non-member participants including the EU and the United States. Council of Europe. URL: <https://www.coe.int/en/web/artificial-intelligence/the-framework-convention-on-artificial-intelligence> (visited on 08/25/2026).
- (May 2026). *European Union Ratifies the Council of Europe Framework Convention on Artificial Intelligence*. coe.int, Artificial Intelligence news. The EU deposited its instrument of ratification on 15 May 2026 at the Committee of Ministers' 135th session. The Convention had not entered into force at the time of writing: it requires five ratifications including three Council of Europe members, and the Council's own convention page describes entry into force in the future tense. URL: <https://www.coe.int/en/web/artificial-intelligence/-/european-union-ratifies-the-council-of-europe-framework-convention-on-artificial-intelligence> (visited on 08/29/2026).
- Cox, Joseph (June 26, 2025). *ICE Is Using a New Facial Recognition App to Identify People, Leaked Emails Show*. 404 Media. URL: <https://www.404media.co/ice-is-using-a-new-facial-recognition-app-to-identify-people-leaked-emails-show/> (visited on 10/01/2026).
- (Jan. 2026). 'ELITE': *The Palantir App ICE Uses to Find Neighborhoods to Raid*. Cited for what ELITE (Enhanced Leads Identification & Targeting for Enforcement) does: it "populates a map with potential deportation targets, brings up a dossier on each person, and provides a 'confidence score' on the person's current address," is used "to find locations where lots of people it might detain could be based," and "receives peoples' addresses from the Department of Health and Human Services (HHS) among a range of other sources." The article rests on internal ICE material and sworn testimony from an ICE official; the testimony and the official's name are behind the paywall and were not read. Agent-verified 2026-10-01. 404 Media. URL: <https://www.404media.co/elite-the-palantir-app-ice-uses-to-find-neighborhoods-to-raid/> (visited on 10/01/2026).
- Crootof, Rebecca (2016). "War Torts: Accountability for Autonomous Weapons". In: *University of Pennsylvania Law Review* 164.6. Cited for war torts as state liability for harm in armed conflict where no one has committed a war crime. The article proposes "explicitly identifying 'war torts' as serious violations of international humanitarian law that give rise to state responsibility," alongside war crimes, which give rise to individual criminal liability; "the same action may simultaneously be a war crime and a war tort" (p. 1386). Its motivating case is harm with no criminally liable person: autonomous weapon systems "may take actions that look like war crimes ... without any individual acting intentionally or recklessly" (abstract, p. 1347), and war crimes "must be committed by a person acting 'willfully'" (p. 1350). It argues that "states should be held strictly liable for the actions of the autonomous weapon systems they field" (p. 1394), noting that strict liability "is usually applied to lawful but inherently dangerous activities" (p. 1395). In this article a war tort is still a violation of international humanitarian law; the extension to lawful acts is the 2022 article's (crootof2022wartorts). Agent-verified 2026-09-26 against the full text of the published article (University of Richmond School of Law repository copy of the *Penn Law Review* print pages 1347–1402) and the *Penn Law Review* article page; the passages quoted were located by searching the text and read in context, and the article was not read cover to cover. No DOI was found for the article., pp. 1347–1402. URL: https://scholarship.law.upenn.edu/penn_law_review/vol164/iss6/1/ (visited on 09/26/2026).
- (2022). "War Torts". In: *New York University Law Review* 97.4. Cited for war torts as state liability for civilian harm caused in armed conflict, lawful or unlawful, whether or not anyone committed a war crime. Abstract (p. 1063): "no entity is liable for lawful but unintended harmful acts—regardless of how many or how horrifically civilians are hurt. This Article

- proposes developing an international ‘war torts’ regime, which would require states to pay for both lawful and unlawful acts in armed conflict that cause civilian harm. Just as tort and criminal law coexist and complement each other in domestic legal regimes, war torts and war crimes would overlap but serve different aims.” The ideal regime would “(1) require states to pay compensation for (2) both lawful and unlawful acts in armed conflict that cause civilian harm and (3) have an institutional structure designed to facilitate victim redress” (p. 1109); in an adversarial process “states should be held strictly liable for their harmful acts” (p. 1110); “The legality of the harmful act should not affect whether civilians can receive compensation for it” (p. 1113). A search of the text for “serious harm,” “significant harm,” “de minimis” and “threshold” found no seriousness threshold on the harm; the article proposes limiting liability instead through “pleading requirements, causation cutoffs, and affirmative defenses” (p. 1120). Agent-verified 2026-09-26 against the full text of the article as published on the *NYU Law Review* website (print pages 1063–1142) and its article page; the passages quoted were located by searching the text and read in context, and the article was not read cover to cover. No DOI was found for the article., pp. 1063–1142. URL: <https://nyulawreview.org/issues/volume-97-number-4/war-torts/> (visited on 09/26/2026).
- Croxton, John (Sept. 2026). *Anthropic Loses Legal Fight to Shed DOD’s Designation as a ‘Supply-Chain Risk’*. Source for the demand, its dates, and the two prohibitions: “The designation has been in force since March, when the department sought to alter its contract with Anthropic to replace prohibitions on using its tools for domestic mass surveillance and fully autonomous weapons with allowances for “all lawful use.”” and “DOD issued an ultimatum on Feb. 24 demanding the company permit its AI to be used for all legal purposes by Feb. 27, which Anthropic refused. Hegseth announced his direction that DOD designate Anthropic a supply-chain risk Feb. 27, and it was formally announced on March 3.” Also: “Friday’s ruling was issued by Judges Gregory Katsas and Neomi Rao, both appointed during the first Trump administration, with a dissent by Judge Karen LeCraft Henderson”. The article gives no removal deadline and does not say whether Claude was still in use. Published 25 September 2026, modified 27 September. Agent-verified 2026-10-01. Defense One. URL: <https://www.defenseone.com/threats/2026/09/anthropic-lawsuit-supply-chain-risk/416252/> (visited on 10/01/2026).
- Cruz Espinoza v. Immigration and Naturalization Service* (1995). 45 F.3d 308 (9th Cir. 1995). Fletcher, J., decided 12 January 1995; affirmed. Cited for the Form I-213 as an officer’s narrative and its evidentiary weight: a form “which border agents routinely complete after interviewing aliens”; “essentially a recorded recollection of a[n INS agent’s] conversation with the alien” (quoting *Bustos-Torres v. INS*, 898 F.2d 1053, 1056 (5th Cir. 1990)); “We agree with the BIA that information on an authenticated immigration form is presumed to be reliable in the absence of evidence to the contrary presented by the alien”; and “The immigration judge was not required to permit cross-examination of the form’s preparer in this case.” Agent-verified 2026-10-01 from the public.resource.org text. URL: <https://law.resource.org/pub/us/case/reporter/F3/045/45.F3d.308.94-70094.html> (visited on 10/01/2026).
- Cyberspace Administration of China et al. (July 10, 2023). *Interim Measures for the Management of Generative Artificial Intelligence Services*. Order No. 15 of the Cyberspace Administration of China, the National Development and Reform Commission, the Ministry of Education, the Ministry of Science and Technology, the Ministry of Industry and Information Technology, the Ministry of Public Security and the National Radio and Television Administration; in force 15 August 2023. Shengchengshi Rengong Zhineng Fuwu Guanli Zanxing Banfa. Cited twice. First, for China’s training-data disclosure requirement, set beside the EU’s public training-data summary. Article 7 requires providers to “use data and foundational models that have lawful sources” and to “increase the truth, accuracy, objectivity, and diversity of training data.” Article 19 requires them to cooperate with the departments “carrying

- out oversight inspections,” “explaining the sources, [scale], types, tagging rules, algorithm mechanisms, etc. for training data as required.” That disclosure is made to the regulator when an inspection requires it, not published. Its motive is in Article 1 (“preserve national security and the societal public interest”) and Article 4(1). Second, for compulsion as a standing condition of operating. Article 4(1): “Uphold the Core Socialist Values,” with a list of prohibited content that begins with content “inciting subversion” of state power. Article 14: on discovering illegal content, providers must stop generating it and use “measures such as model optimization training to make corrections.” Article 17: providers of services “with public opinion properties or the capacity for social mobilization shall carry out security assessments” and file their algorithms. Article 19 adds “providing necessary technical, data, and other supports and assistance.” Article 21 allows an order “to suspend the provision of the related services.” The measures cover only services offered to the public in mainland China (Article 2), and Article 17 covers only the subset of services it names. Order number, the seven issuing bodies and their heads’ signatures, the date of 10 July 2023, adoption at the CAC meeting of 23 May 2023, and entry into force on 15 August 2023 (Article 24) were agent-verified 2026-09-24 against the Chinese text on the CAC page (posted 13 July 2023). Articles 1, 2, 4, 7, 14, 17, 19, 21 and 24 were read there in full. The English quotations come from the China Law Translate translation (posted 13 July 2023, last modified 2 August 2023), compared clause by clause with the Chinese. Two corrections to that translation: in Article 19 it reads “models” where the Chinese (gui mo) means scale, hence the bracket, and in Article 4(1) it reads “national sovereignty” where the Chinese means state power. One search found no amendment or repeal through September 2026. The CAC’s filing registry and how it works in practice were not checked. URL: https://www.cac.gov.cn/2023-07/13/c_1690898327029107.htm (visited on 09/24/2026).
- Cyprus v. Turkey* (May 10, 2001). Judgment (merits). Turkey found responsible under the Convention for northern Cyprus by virtue of the effective overall control exercised through its military presence. URL: <https://hudoc.echr.coe.int/eng?i=001-59454> (visited on 09/30/2026).
- Dahm, Volker (1986). “Anfänge und Ideologie der Reichskulturkammer: Die ‘Berufsgemeinschaft’ als Instrument kulturpolitischer Steuerung und sozialer Reglementierung”. In: *Vierteljahrshefte für Zeitgeschichte* 34.1. Pages 68–69: the Finance Ministry’s assent to the first implementing decree despite too little time to examine it (BA R 2/4870). Page 69: the decree’s sections 1–4 and 10. Pages 77–78: the reliability clause’s origin in the trade code and the chamber presidents’ unbounded discretion., pp. 53–84. URL: https://www.ifz-muenchen.de/heftarchiv/1986_1_2_dahm.pdf (visited on 10/01/2026).
- Damasio, Antonio (2003). *Looking for Spinoza: Joy, Sorrow, and the Feeling Brain*. Claims Spinoza as a precursor of a biological account of affect, reading the conatus as an anticipation of homeostasis and the affects as registrations of what is happening to a body’s power to go on. Cited for Damasio’s reading of Spinoza and not for Spinoza’s text, which is spinoza1985collected. One objection is worth carrying: Spinoza states the conatus of every thing including an inanimate one, so a homeostatic reading narrows a metaphysical doctrine to a biological one. Agent-verified 2026-09-18. Orlando, FL: Harcourt.
- (2018). *The Strange Order of Things: Life, Feeling, and the Making of Cultures*. The later statement of the same position, built on homeostasis throughout: feelings are the mental expressions of homeostasis, which is the formulation chapter 2 glosses as an organism registering its regulation of its own interior. Agent-verified 2026-09-18. New York: Pantheon.
- Darwall, Stephen (2006). *The Second-Person Standpoint: Morality, Respect, and Accountability*. Cambridge, MA: Harvard University Press. ISBN: 0674022742.
- Darwall, Stephen L. (1977). “Two Kinds of Respect”. In: *Ethics* 88.1, pp. 36–49. doi: 10.1086/292054.

- Davidov, Maayan et al. (2021). “Caring Babies: Concern for Others in Distress during Infancy”. In: *Developmental Science* 24.2. Cited for the order: concern early, action later. Abstract: “Concern for distressed others was seen early in the first year of life ... Empathic concern was moderately consistent across both situation and age, from as early as 3 months. Concern for others grew only modestly with age, plateauing during the second year, whereas prosocial behavior increased rapidly during the second year.” Bibliographic details and abstract agent-verified 2026-09-25. The article was not read., e13016. doi: 10.1111/desc.13016.
- Daw, Nathaniel D., Yael Niv, and Peter Dayan (2005). “Uncertainty-Based Competition between Prefrontal and Dorsolateral Striatal Systems for Behavioral Control”. In: *Nature Neuroscience* 8.12, pp. 1704–1711. doi: 10.1038/nn1560.
- Dayan, David (Aug. 6, 2025). *FEMA Employees Reassigned to ICE*. The American Prospect. URL: <https://prospect.org/2025/08/06/2025-08-06-fema-employees-reassigned-to-ice/> (visited on 09/23/2026).
- De Luce, Dan (Mar. 2026). *U.S. Military Is Using AI to Help Plan Iran Air Attacks, Sources Say, as Lawmakers Call for Oversight*. Cited for the contrary account of the model’s role: “Anthropic’s Claude has become a crucial component of Palantir’s Maven intelligence analysis program, which was also used in the U.S. operation to capture Venezuelan President Nicolás Maduro. Claude is used to help military analysts sort through intelligence and does not directly provide targeting advice, according to a person with knowledge of Anthropic’s work with the Defense Department.” Also: “The use of Palantir’s software, which relies in part on Anthropic’s Claude AI systems”. Agent-verified 2026-10-01. NBC News. URL: <https://www.nbcnews.com/tech/tech-news/us-military-using-ai-help-plan-iran-air-attacks-sources-say-lawmakers-rcna262150> (visited on 10/01/2026).
- Debord, Guy (1967). *The Society of the Spectacle*. Originally published as *La Société du spectacle*. Thesis 59 describes spectacular rebelliousness coexisting with acceptance of the status quo, dissatisfaction itself sold back as a commodity; thesis 203 warns that abstract denunciation of the spectacle, divorced from revolutionary praxis, becomes a hollow formula that reinforces the spectacular system (the Situationist concept of recuperation). Paris: Buchet-Chastel.
- Deeb, Aghyad and Fabien Roger (Oct. 11, 2024). *Do Unlearning Methods Remove Information from Language Model Weights?* Version v3, 2025-02-07. Retraining on an accessible, mutually independent half of the unlearned facts recovers at least 88 percent of the original model’s post-retraining accuracy on the held-out half, for RMU, gradient difference and random-incorrect-answer unlearning on Llama 3 8B, when unlearning costs at most 5 percent of retain accuracy. The result holds for pretrained information only: for facts installed by fine-tuning, recovery falls below 35 percent for two of the three methods. No published venue found as of 2026-10-01; absent from the ICLR 2025 proceedings index. Agent-verified 2026-10-01 against arXiv v3, read past the abstract. arXiv: 2410.08827. URL: <https://arxiv.org/abs/2410.08827> (visited on 10/01/2026).
- Defense Advanced Research Projects Agency (2026). *LongShot*. DARPA program page; air-launched unmanned aircraft built by General Atomics, designated X-68A in February 2026. URL: <https://www.darpa.mil/research/programs/longshot> (visited on 08/25/2026).
- Defense Priorities and Allocations System* (2026). 15 C.F.R. part 700 (2026 ed.) § 700.13(a) (mandatory acceptance: a person shall accept every rated order received and fill it regardless of other rated or unrated orders accepted); § 700.13(b) (mandatory rejection, on inability to meet the delivery date); § 700.13(c) (optional rejection, on grounds that do not include objection to the use). Administered by the Bureau of Industry and Security, Department of Commerce.
- Defense Production Act of 1950* (1950). 50 U.S.C. §§ 4501 et seq., as amended. Title I (priorities and allocations) is codified at §§ 4511–4513. § 4511(a) authorizes the President to require acceptance and performance of contracts and orders he deems necessary to promote the national defense, in preference to other contracts or orders, by any person he finds capable

of their performance. § 4513 sets the penalty for willful noncompliance at a fine of up to \$10,000, imprisonment of up to one year, or both.

Dekker, Sidney (2011). *Drift into Failure: From Hunting Broken Components to Understanding Complex Systems*. Cited for drift into failure as an extension of the normalization of deviance. The publisher's description (Routledge/CRC Press page): "These were systems that drifted into failure. While pursuing success in a dynamic, complex environment with limited resources and multiple goal conflicts, a succession of small, everyday decisions eventually produced breakdowns on a massive scale." Crossref lists the chapters, among them "Features of Drift" (pp. 25–50) and "Theorizing Drift" (pp. 87–125). A reader's published chapter outline lists, under "Theorizing Drift," sections titled "Normalizing deviance, structural secrecy and practical drift" and "The normalization of deviance," and quotes the book defining one feature of drift as "Decrementalism, where constant organizational and operational adaptation around goal conflicts and uncertainty produces small, stepwise normalization where each next decrement is only a small deviation from the previously accepted norm"; that outline is secondary and none of it was checked against the book. Dekker's own journal treatment (Sidney Dekker and Shawn Pruchnicki, "Drifting into Failure: Theorising the Dynamics of Disaster Incubation," *Theoretical Issues in Ergonomics Science*, 2013, doi:10.1080/1463922X.2013.856495) builds on Vaughan explicitly: "The 'normalization of deviance' describes a process whereby a group's construction of risk can persist even in the face of continued (and worsening) signals of potential danger." Bibliographic details agent-verified 2026-09-26 against Crossref (the CRC Press ebook record, dated 2016, with chapter records dated 28 February 2011 and ISBNs 9781409422228 and 9781351942928) and the Routledge page ("Copyright 2011," 234 pages, product ISBN 9781409422211). Ashgate and Farnham are the original imprint as commonly cited; the Ashgate title page was not seen. Which binding each ISBN belongs to was not checked. The book itself was not read. Farnham: Ashgate. ISBN: 9781409422211. doi: 10.1201/9781315257396. URL: <https://doi.org/10.1201/9781315257396> (visited on 09/26/2026).

Deleuze, Gilles and Félix Guattari (1987). *A Thousand Plateaus: Capitalism and Schizophrenia*. Translated by Brian Massumi; originally published in French as *Mille Plateaux* (Paris: Les Éditions de Minuit, 1980). Plateau 9, "1933: Micropolitics and Segmentarity." Molar and molecular are two kinds of segmentarity, rigid and supple, both running at every scale – not two sizes; an earlier version of this note said state-level and micro-level and was wrong. The same plateau supplies the micro-fascisms that resonate, and the contrast between the conservative totalitarian state and fascism's line of abolition. Also cited for Plateau 13, "The Apparatus of Capture," whose propositions on the capitalist axiomatic supply the corrected test for the first structural feature: an axiomatic does not refuse a demand, it adds an axiom, modifying its own rules so the demand can be met while nothing that produced the demand is settled. Also cited for Plateau 7, "Year Zero: Faciality," pages 167–191, where the face is not a natural given but is produced by an abstract machine of faciality (*visagéité*) that overcodes the head into a surface of signification. The quoted sentence on racism as the determination of degrees of deviance is at page 178, wording agent-checked against two independent scholarly sources quoting it identically. The related claims that the face is not universal but is White Man himself, and that the face is a politics, are used in paraphrase and cited to the plateau rather than to a page: secondary sources place them variously and the pages were not agent-verified. Plateau 9 carries two further uses. The four dangers of the line of flight – Fear, which clings to the molar organization and its binary machines; Clarity, where the line acquires its own micro-segmentations and micro-Oedipuses crop up; Power, where the party able to move on both lines wants to hold both still; and the great Disgust, abolition pure and simple, the passion of abolition, which turns the line of flight into a line of death – are the account section 3.5 uses, and the fourth is the same line already cited from this plateau as the line of abolition, which is why section 3.5 treats it as a recognition rather

- than a new claim. Plateau 13 likewise carries a second use: there is enslavement when human beings are themselves constituent pieces of a machine they compose among themselves and with other things, under the control and direction of a higher unity, and subjection when the higher unity constitutes the human being as a subject linked to a new exterior object, which can be an animal, a tool, or even a machine. Section 6.4.1 uses the pair and the fact that both poles run together. Wording for the dangers and for the enslavement and subjection definitions was agent-checked against independent sources quoting the passages; both are cited at plateau level with no page asserted, which is this entry's convention. Minneapolis: University of Minnesota Press.
- Democracy Now! (May 23, 2019). *Leaked OPCW Report Raises New Questions About 2018 Chemical Attack in Douma, Syria*. Contemporaneous report of the May 2019 leak. "a newly leaked internal document from the Organization for the Prohibition of Chemical Weapons reveals there were conflicting views within the organization as to what happened. The leaked document suggests the cylinders were 'manually placed' on the ground and were not dropped from the air." The guest, journalist Brian Whitaker, attributes the document to "a man called Ian Henderson, who, as far as I understand it, was working for the OPCW ... according to the OPCW, he wasn't actually a member of it." The leaked document itself was not read. Agent-verified 2026-10-01. URL: https://www.democracynow.org/2019/5/23/leaked_opcw_report_raises_new_questions (visited on 10/01/2026).
- Deneault, Audrey-Ann, Stuart I. Hammond, and Sheri Madigan (2023). "A Meta-Analysis of Child-Parent Attachment in Early Childhood and Prosociality". In: *Developmental Psychology* 59.2. Cited for the size of the link between attachment and later prosociality: child-mother security $r = .19$, 95% CI [.14, .23] (35 studies, 100 effect sizes, $N = 4,611$); child-father $r = .11$ [.02, .23] (6 studies, $N = 402$); no moderators identified. Hence "reliably but modestly" in §9.3. No source found supports calling attachment the strongest developmental predictor of prosocial behavior, and meta-analyses of other predictors were not searched. Bibliographic details and figures agent-verified 2026-09-25 from the abstract., pp. 236–255. DOI: 10.1037/dev0001484.
- Department of Defense Appropriation Act, 1985, §8066(a) [Boland Amendment] (Oct. 12, 1984). Pub. L. No. 98-473, §101(h), 98 Stat. 1837, 1935. Enacted in the continuing appropriations resolution for fiscal year 1985.
- Department of Homeland Security Appropriations Act, 2010 (Oct. 28, 2009). Pub. L. No. 111-83, 123 Stat. 2142. Title II, U.S. Immigration and Customs Enforcement, Salaries and Expenses, 123 Stat. 2149 (provisos on 33,400 detention beds and on continuing 287(g) delegations if the DHS Inspector General determines their terms have been violated). URL: <https://www.congress.gov/111/plaws/publ83/PLAW-111publ83.pdf> (visited on 10/01/2026).
- Department of Homeland Security v. D.V.D. (June 23, 2025). No. 24A1153 (U.S. June 23, 2025) (order granting stay); clarified July 3, 2025.
- Dewey, John (1927). *The Public and Its Problems*. New York: Henry Holt.
- Di Paolo, Ezequiel A. (2005). "Autopoiesis, Adaptivity, Teleology, Agency". In: *Phenomenology and the Cognitive Sciences* 4.4, pp. 429–452.
- Dickey, Megan Rose (Dec. 2020). *Google's Co-lead of Ethical AI Team Says She Was Fired for Sending an Email*. Timnit Gebru's departure from Google, two months before the dismissal of Margaret Mitchell that dickey2021mitchell covers; same reporter and outlet. Agent-verified 2026-09-20 against the article. It records a disagreement the manuscript's *dismissed* does not: Gebru says she was fired after conditions she set for staying were not met, and the message she received from Google reads *We respect your decision to leave Google as a result, and we are accepting your resignation*. URL: <https://techcrunch.com/2020/12/03/googles-co-lead-of-ethical-ai-team-says-she-was-fired-for-sending-an-email/> (visited on 09/20/2026).
- (2021). *Google Fires Top AI Ethics Researcher Margaret Mitchell*. Reports Mitchell's February 2021 dismissal, following the December 2020 dismissal of Timnit Gebru, both co-leads of

- Google's Ethical AI team. URL: <https://techcrunch.com/2021/02/19/google-fires-top-ai-ethics-researcher-margaret-mitchell/> (visited on 08/25/2026).
- Directing the President, Pursuant to Section 5(c) of the War Powers Resolution, to Remove United States Armed Forces from Hostilities with Iran* (June 23, 2026). Sponsor Rep. Gregory Meeks (D-N.Y.). Passed House 215–208 on 3 June 2026; passed Senate 50–48 on 23 June 2026 (Record Vote No. 184). A concurrent resolution, not presented to the President. United States Congress. URL: <https://www.congress.gov/119/bills/hconres86/BILLS-119hconres86ih.pdf> (visited on 09/30/2026).
- Doctorow, Cory (June 2026). *The Reverse Centaur's Guide to Life After AI: How to Think About Artificial Intelligence—Before It's Too Late*. New York: MCD / Farrar, Straus and Giroux. ISBN: 9780374621568.
- Dohare, Shibhansh et al. (2024). "Loss of Plasticity in Deep Continual Learning". In: *Nature* 632.8026, pp. 768–774. doi: 10.1038/s41586-024-07711-7.
- Dondi, Marco, Francesca Simion, and Giovanna Caltran (1999). "Can Newborns Discriminate between Their Own Cry and the Cry of Another Newborn Infant?" In: *Developmental Psychology* 35.2. Cited for newborns' distress at another newborn's cry exceeding their distress at their own. ERIC abstract: "awake newborns expressed facial distress more frequently and longer to another newborn's cry than to their own." Bibliographic details and the abstract were agent-verified 2026-09-25 (ERIC; PMID 10082012). The article was not read., pp. 418–426. doi: 10.1037/0012-1649.35.2.418.
- Drew, Trafton, Melissa L.-H. Võ, and Jeremy M. Wolfe (2013). "The Invisible Gorilla Strikes Again: Sustained Inattentional Blindness in Expert Observers". In: *Psychological Science* 24.9, pp. 1848–1853. doi: 10.1177/0956797613479386. URL: <https://journals.sagepub.com/doi/abs/10.1177/0956797613479386> (visited on 08/25/2026).
- Dunn, Barnaby D., Tim Dalgleish, and Andrew D. Lawrence (2006). "The Somatic Marker Hypothesis: A Critical Evaluation". In: *Neuroscience & Biobehavioral Reviews* 30.2. Published online 27 September 2005. A review arguing that the somatic marker hypothesis's key claims are contested: its evidence rests largely on the Iowa Gambling Task, whose interpretation the authors say is undermined by "the cognitive penetrability of the reward/punishment schedule, ambiguity surrounding interpretation of the psychophysiological data, and a shortage of causal evidence linking peripheral feedback to IGT performance," with "other well-specified and parsimonious explanations that can equally well account for the IGT data"; they conclude that the hypothesis "requires additional empirical support to remain tenable." Cited for the hypothesis being contested. Bibliographic details agent-verified against Crossref and PubMed (PMID 16197997) and the quotations against the PubMed abstract; the full text was not read. Agent-verified 2026-09-23., pp. 239–271. doi: 10.1016/j.neubiorev.2005.07.001. URL: <https://doi.org/10.1016/j.neubiorev.2005.07.001> (visited on 09/23/2026).
- Durchgestrichenes Hakenkreuz* (Mar. 15, 2007). Judgment. Bundesgerichtshof, Third Criminal Senate, on appeal from Landgericht Stuttgart (judgment of 29 September 2006). Reported at BGHSt 51, 244; NJW 2007, 1602; NStZ 2007, 466. The defendant sold stickers, badges and similar items to the punk scene bearing National Socialist symbols crossed out, smashed or otherwise negated. Held that using a proscribed organization's symbol in a depiction whose content plainly and unmistakably expresses opposition to the organization and rejection of its ideology does not run counter to the protective purpose of § 86a and therefore falls outside the offence, the opposition having to be recognizable at first glance. The conviction was set aside and the defendant acquitted. Agent-verified 2026-09-13 against the court's own press release and the reported headnote. URL: <https://www.bundesgerichtshof.de/SharedDocs/Pressemitteilungen/DE/2007/2007036.html> (visited on 09/13/2026).
- Duvar English (Dec. 11, 2023). *Turkey's Pro-Kurdish HEDEP Changes Abbreviation upon Top Court's Demand*. URL: <https://www.duvarenglish.com/turkeys-pro-kurdish-hedep-changes-abbreviation-upon-top-courts-demand-news-63473> (visited on 09/23/2026).

- Eagly, Ingrid and Steven Shafer (Sept. 28, 2016). *Access to Counsel in Immigration Court*. Special report. Based on Eagly and Shafer, "A National Study of Access to Counsel in Immigration Court," *University of Pennsylvania Law Review* 164 (2015): 1. American Immigration Council. URL: <https://www.americanimmigrationcouncil.org/report/access-counsel-immigration-court/> (visited on 10/01/2026).
- Easterbrook, J. A. (1959). "The Effect of Emotion on Cue Utilization and the Organization of Behavior". In: *Psychological Review* 66.3. Cited for acute stress narrowing the range of cues a person uses. Author, title, journal, volume, issue and pages agent-verified 2026-09-28 against the Crossref record for the DOI. The article itself was not read., pp. 183–201. doi: 10.1037/h0047707.
- Eco, Umberto (June 1995). "Ur-Fascism". In: *The New York Review of Books*. Also published as "Eternal Fascism: Fourteen Ways of Looking at a Blackshirt"; that version was not seen. Enumerates fourteen features of Ur-Fascism, placing the scapegoated out-group at a movement's origin: "The first appeal of a fascist or prematurely fascist movement is an appeal against the intruders. Thus Ur-Fascism is racist by definition" (item 5). Also cited for the enemy presented as at once overwhelming and beatable: "by a continuous shifting of rhetorical focus, the enemies are at the same time too strong and too weak" (item 8). Both quotations agent-verified 2026-09-25 against the essay's page on the New York Review of Books site. URL: <https://www.nybooks.com/articles/1995/06/22/ur-fascism/> (visited on 08/25/2026).
- Eisenberg, Nancy and Richard A. Fabes (1990). "Empathy: Conceptualization, Measurement, and Relation to Prosocial Behavior". In: *Motivation and Emotion* 14.2. Cited for the sympathy/personal-distress distinction and its relation to helping. Abstract: "other-oriented sympathetic responding is positively related to prosocial behavior (particularly altruism) whereas personal distress reactions sometimes are associated with low levels of helping." The abstract does not say which findings come from children; the children's data are reported in Eisenberg, Fabes, Miller, Fultz, Shell, Mathy and Reno (1989), *JPSP* 57(1), 55–66, whose abstract reports for children only "some indication that report of positive affect and facial distress were negatively related to prosocial intentions and behavior, whereas facial concern was positively related." Bibliographic details and abstract agent-verified 2026-09-25 against Crossref and the Springer article page. The article was not read., pp. 131–149. doi: 10.1007/BF00991640.
- Eisenberger, Naomi I., Matthew D. Lieberman, and Kipling D. Williams (2003). "Does Rejection Hurt? An fMRI Study of Social Exclusion". In: *Science* 302.5643, pp. 290–292. doi: 10.1126/science.1089134.
- Elish, Madeleine Clare (2019). "Moral Crumple Zones: Cautionary Tales in Human-Robot Interaction". In: *Engaging Science, Technology, and Society* 5, pp. 40–60. doi: 10.17351/ests2019.260.
- Ely, John Hart (1980). *Democracy and Distrust: A Theory of Judicial Review*. Cambridge, MA: Harvard University Press.
- Eritrean Badme Ruling Upheld* (Mar. 31, 2003). Cited for the Boundary Commission placing Badme in Eritrea. "Last year, the commission issued a ruling which appeared to award Badme to Eritrea ... However, ten days ago, for the first time, the UN Boundary Commission categorically clarified Badme as belonging to Eritrea, throwing out Ethiopia's evidence as 'inadequate'." The companion report of 3 April 2003 (news.bbc.co.uk/2/hi/africa/2914559.stm) adds: "Eritrean troops entered the village, which was then administered by Ethiopia, claiming it as part of Eritrea." The Commission's decision and observations themselves were not read. Agent-verified 2026-10-01. BBC News. URL: <http://news.bbc.co.uk/2/hi/africa/2903229.stm> (visited on 10/01/2026).
- Erste Verordnung zur Durchführung des Reichskulturkammergesetzes vom 1. November 1933* (1933). RGBl. I S. 797. Cited for section 4 (compulsory membership) and section 10 (refusal or

- expulsion for lack of reliability and suitability). URL: <https://www.verfassungen.de/de33-45/kulturkammer33-v1.htm> (visited on 10/01/2026).
- Esping-Andersen, Gøsta (1990). *The Three Worlds of Welfare Capitalism*. Published the same year in Cambridge by Polity Press. Cited for decommodification and for its being characteristic of the social-democratic regime. Two definitions, as quoted with page numbers from the Princeton imprint by Huo, Nelson and Stephens, "Decommodification and Activation in Social Democratic Policy: Resolving the Paradox," *Journal of European Social Policy* 18(1) (2008), author's copy: in decommodifying welfare states "citizens can freely, and without potential loss of job, income, or general welfare, opt out of work when they consider it necessary" (p. 23), and decommodification "refers to the degree to which individuals, or families, can uphold a socially acceptable standard of living independently of market participation" (p. 37). The same article reports that Esping-Andersen "identifies the Nordic welfare states as the most decommodifying type of welfare states, due to the dominant influence of social democratic parties in these countries," the liberal regimes being "the least decommodifying," and that he shows (p. 129) a strong positive association between left political power and his decommodification index. A third wording, "De-commodification occurs when a service is rendered as a matter of right, and when a person can maintain a livelihood without reliance on the market," is widely quoted and attributed to pp. 21–22; its page was taken from search-result summaries and not checked. Bibliographic details agent-verified 2026-09-26 against the Internet Archive and Open Library catalogue records. The book itself was not read; every quotation here is at second hand, and the decommodification scores by country were not checked. No DOI was found in Crossref. Princeton, NJ: Princeton University Press. ISBN: 978-0-691-02857-6.
- European Commission (2022). *Proposal for a Directive of the European Parliament and of the Council on Adapting Non-Contractual Civil Liability Rules to Artificial Intelligence (AI Liability Directive)*. COM(2022) 496 final. Proposed 28 September 2022; listed for withdrawal in the Commission's 2025 Work Programme (11 February 2025) for want of foreseeable stakeholder agreement; withdrawal notice published in the Official Journal 6 October 2025. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52022PC0496> (visited on 08/25/2026).
- European Commission against Racism and Intolerance (Mar. 16, 2016). *ECRI General Policy Recommendation No. 16 on Safeguarding Irregularly Present Migrants from Discrimination*. CRI(2016)16. Council of Europe. URL: <https://rm.coe.int/ecri-general-policy-recommendation-no-16-on-safeguarding-irregularly-p/16808b5b0b> (visited on 10/01/2026).
- European Parliament (2017). *European Parliament Resolution of 16 February 2017 with Recommendations to the Commission on Civil Law Rules on Robotics (2015/2103(INL))*. P8_TA(2017)0051. Rapporteur: Mady Delvaux; Committee on Legal Affairs report A8-0005/2017. URL: https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=uriserv:OJ.C_.2018.252.01.0239.01.ENG (visited on 08/25/2026).
- European Parliament and Council of the European Union (2024a). *Directive (EU) 2024/2853 of the European Parliament and of the Council on Liability for Defective Products and Repealing Council Directive 85/374/EEC*. Official Journal of the European Union, OJ L, 2024/2853. URL: <https://eur-lex.europa.eu/eli/dir/2024/2853/oj> (visited on 09/17/2026). Article 4(1) brings software within the definition of a product: "all movables... it includes electricity, digital manufacturing files, raw materials and software." **It says software and stops.** The phrase "including AI system providers within the meaning of Regulation (EU) 2024/1689" is Recital 13, which is not an operative provision, and an earlier version of this note wrongly attributed it to the definition. Article 8 makes the economic operator liable irrespective of fault. Article 2(1): "This Directive shall apply to products placed on the market or put into service after 9 December 2026"; member states must transpose by the same date, whose article number is not pinned here, one reading returning Article 19 while the EUR-Lex text conversion truncates at Article 17. Article 6(1) is the closed damage list, and its first limb

is wider than a list of physical harms sounds: “(a) death or personal injury, **including medically recognised damage to psychological health**”; (b) damage to, or destruction of, any property, except... property used exclusively for professional purposes; (c) destruction or corruption of data that are not used for professional purposes.” Article 9 is disclosure of evidence. Article 10(2) presumes defectiveness on failure to disclose, non-compliance with mandatory safety requirements, or obvious malfunction in foreseeable use; Article 10(3) presumes the causal link where the product is defective and the damage is of a kind typically consistent with that defect; Article 10(4) lets a court presume either where technical or scientific complexity makes the claimant’s proof excessively difficult **and** the claimant demonstrates that the defect or the link is likely. Every article number above was read against the consolidated text at EUR-Lex on 2026-09-18.

- European Parliament and Council of the European Union (2024b). *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union, OJ L, 2024/1689, 12.7.2024. Article 5 bans a small set of unacceptable-risk practices outright — government social scoring, untargeted scraping of facial images to build recognition databases, real-time remote biometric identification in public spaces outside narrow law-enforcement exceptions, and predictive policing on profiling alone — with those prohibitions applying from 2 February 2025; a larger high-risk tier requires conformity assessment and ongoing monitoring; limited-risk systems face transparency duties; minimal-risk systems are largely unregulated. Article 53(1)(d) requires a provider of a general-purpose model to publish a summary of its training content, against the AI Office’s template of 24 July 2025, applying from 2 August 2025 for new models and 2 August 2027 for models already on the market. Article 101 allows the AI Office to fine a noncompliant general-purpose model provider up to 3 percent of global annual turnover or EUR 15 million, whichever is higher, regardless of where the provider is based. Article 51(2) presumes a general-purpose model to have high-impact capabilities, and so to carry systemic risk under Article 51(1)(a), when the cumulative computation used for its training “measured in floating point operations is greater than 10^{25} ”; Article 51(3) lets the Commission amend the threshold by delegated act. Article 51(2) agent-verified 2026-09-23 against the text as published in the Official Journal; no consolidated version was consulted. URL: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (visited on 08/25/2026).
- Evans, Jonathan St. B. T. and Keith E. Stanovich (2013). “Dual-Process Theories of Higher Cognition: Advancing the Debate”. In: *Perspectives on Psychological Science* 8.3, pp. 223–241. DOI: 10.1177/1745691612460685.
- Evans, Richard J. (2004). *The Coming of the Third Reich*. First published London: Allen Lane, 2003. Page locators are from this edition, checked by search-inside on 2026-10-01; agreement with the Allen Lane pagination not verified. New York: Penguin Press. ISBN: 978-1-59420-004-5.
- (2006). *The Third Reich in Power*. First published 2005. Page locators are from this edition, checked by search-inside on 2026-10-01. New York: Penguin. ISBN: 978-0-14-303790-3.
- Fang, Congyu et al. (2023). “Proof-of-Learning is Currently More Broken Than You Think”. In: *2023 IEEE 8th European Symposium on Security and Privacy (EuroS&P)*, pp. 797–816. arXiv: 2208.03567.
- Fang, Lee (June 2018). *Google Won’t Renew Its Drone AI Contract, but It May Still Sign Future Military AI Contracts*. Source for the term of Google’s contract: “Google will continue work on the project through March 2019, according to multiple people with knowledge of the announcement, but once the 18-month contract concludes, it will not be renewed.” On its start: Google “quietly won the contract last September,” that is, September 2017. Agent-verified 2026-10-01. The Intercept. URL: <https://theintercept.com/2018/06/01/google-drone-ai-project-maven-contract-renew/> (visited on 10/01/2026).

- (Mar. 2019). *Google Hedges on Promise to End Controversial Involvement in Military Drone Contract*. Source for what the early system did: “The project uses machine-learning to automatically label images, buildings, and other objects captured by cameras on drones, helping Air Force analysts identify unique targets.” Also: “on June 1, 2018, Google said it would let the Project Maven contract expire when it ends in March 2019.” Agent-verified 2026-10-01. The Intercept. URL: <https://theintercept.com/2019/03/01/google-project-maven-contract/> (visited on 10/01/2026).
- Federal Communications Commission v. Consumers’ Research* (2025). No. 24-354, 606 U.S. 656 (2025). Cited for the Universal Service Fund as a structural precedent: carriers pay a contribution the agency sets under the statute, a private company administers the Fund under agency control, and the Court upheld the arrangement against a nondelegation challenge. Decided June 27, 2025, together with No. 24-422, *SHLB Coalition v. Consumers’ Research*; 6–3. Kagan, J., for the Court, joined by Roberts, C.J., and Sotomayor, Kavanaugh, Barrett and Jackson, JJ.; Kavanaugh and Jackson, JJ., concurring; Gorsuch, J., dissenting, joined by Thomas and Alito, JJ. Holding: “The universal-service contribution scheme does not violate the nondelegation doctrine.” Congress’s delegation to the FCC is upheld under the intelligible-principle test, because “the word ‘sufficient’ sets a floor and a ceiling alike” (slip op. 19–30); the FCC’s use of the Universal Service Administrative Company is upheld under the private nondelegation doctrine because the Administrator is “broadly subordinate to the Commission” and “the Commission is, throughout, the final authority” (slip op. 30–34); the combination theory is rejected (slip op. 34–37). Who sets the contribution: “The Commission, in other words, sets the contribution factor to raise just enough money—a ‘sufficient’ amount—to implement universal service as Congress directed”; the Administrator projects expenses and revenues and submits them “to the Commission for approval.” The U.S. Reports citation comes from *Trump v. Slaughter* (2026), which cites “Consumers’ Research, 606 U.S. 656.” Agent-verified 2026-09-26 against the Cornell LII text of the syllabus and the opinion of the Court. Not read: the concurrences and the dissent; pin pages are the syllabus’s slip-opinion references. URL: <https://www.law.cornell.edu/supremecourt/text/24-354> (visited on 09/26/2026).
- Feinberg, Joel (1980). “The Child’s Right to an Open Future”. In: *Whose Child? Children’s Rights, Parental Authority, and State Power*. Ed. by William Aiken and Hugh LaFollette. Totowa, NJ: Rowman and Littlefield, pp. 124–153.
- Feinberg, Todd E. and Jon M. Mallatt (2016). *The Ancient Origins of Consciousness: How the Brain Created Experience*. Agent-verified 2026-09-18 against the publisher and a review in the Quarterly Review of Biology: xx + 366 pp. Argues consciousness appeared with the first complex brains at the Cambrian explosion, that all vertebrates have it, and that arthropods and cephalopods meet many of the criteria. Cambridge, MA: MIT Press. ISBN: 978-0-262-03433-3.
- Ferguson, Niall (Apr. 2004). “Prisoner Taking and Prisoner Killing in the Age of Total War: Towards a Political Economy of Military Defeat”. In: *War in History* 11.2. Cited for the moment of surrender as the most dangerous for the one surrendering. Bibliographic details agent-verified 2026-09-28 against the Crossref record for the DOI. The article itself was not read., pp. 148–192. doi: 10.1191/0968344504wh291oa.
- Fickinger, Arnaud et al. (2020). *Multi-Principal Assistance Games*. Generalizes cooperative inverse reinforcement learning to N principals with differing preferences, with the assisting agent acting to maximize the sum of principal payoffs. Cited for the generalization and for the social-choice impossibility results that apply to it; a companion paper, arXiv:2012.14536, presented at the NeurIPS 2020 Workshop on Cooperative AI, gives collegial mechanisms that circumvent Gibbard’s theorem. arXiv: 2007.09540 [cs.AI]. URL: <https://arxiv.org/abs/2007.09540> (visited on 09/01/2026).
- Financial Industry Regulatory Authority (2026). *Enforcement*. FINRA “brings formal disciplinary actions against firms and their associated persons,” with sanctions that include “fines,

suspensions, and ... bars from FINRA membership"; its own pages describe it as operating "under the supervision of the SEC," which approves its rules (finra.org/about). Cited for the point that FINRA enforces as well as examines, which is where the book's self-regulatory body departs from the FINRA model. Quotations agent-verified 2026-09-23 against the pages as fetched; no SEC page on the SEC's review of FINRA decisions was read. URL: <https://www.finra.org/rules-guidance/enforcement> (visited on 09/23/2026).

Fischer, John Martin and Mark Ravizza (1998). *Responsibility and Control: A Theory of Moral Responsibility*. Cambridge Studies in Philosophy and Law. Cited for guidance control: conduct issues from the agent's own mechanism, and that mechanism is appropriately responsive to reasons; the book states it as an action's flowing "from the agent's own, moderately reasons-responsive mechanism" (ch. 8 summary). Also cited for the account of moderate reasons-responsiveness and of ownership, and for the origin of the term "reasons-responsive." Chapter 3 (pp. 62–91) builds "a new notion of reasons-responsiveness that is 'in between' strong and weak reasons-responsiveness: moderate reasons-responsiveness," distinguishing "recognition of reasons and reaction to reasons" and arguing that "the requirements on reasons-recognition are more stringent than those on reasons-reactivity" (ch. 3 summary). The book's summary formula, as quoted by T. R. Long: responsibility "requires that an agent act on a mechanism that is regularly receptive to reasons, some of which are moral reasons, and at least weakly reactive to reasons" (p. 82). Regular receptivity is "an understandable pattern of (actual and hypothetical) reasons-receptivity," one in which "acceptance of one reason as sufficient implies that a stronger reason must also be sufficient" (p. 71), and the pattern must include moral reasons (p. 77). Weak reactivity requires only that in some possible scenario with the same laws, the same kind of mechanism operating and a sufficient reason to do otherwise, the mechanism issues in doing otherwise (pp. 73–76, as the Stanford Encyclopedia of Philosophy's entry on free will summarizes it). Ownership: a mechanism becomes the agent's own through "taking responsibility" for it (ch. 8, pp. 207–239), which requires that the agent see herself as an agent when acting from the mechanism and as an apt target of the reactive attitudes for so acting, and that these beliefs be based appropriately on her evidence (p. 238, as the free-will entry summarizes it). The term "reasons-responsive" is older than this book: the Stanford Encyclopedia's entry on compatibilism (section 4.4 and supplement B) presents the reasons-responsive account, with its strong and weak forms, as Fischer's own earlier work, "Responsiveness and Moral Responsibility," in Ferdinand Schoeman (ed.), *Responsibility, Character, and the Emotions* (Cambridge University Press, 1987), 81–106, and *The Metaphysics of Free Will* (1994), refined in this joint book; the same entry lists earlier accounts of freedom as sensitivity to reasons (Dennett 1984, Nozick 1981, Gert and Duggan 1979, Neely 1974, Fingarette 1972, Glover 1970, MacIntyre 1957). The phrase "reasons-responsive refusal" is the manuscript's, not the book's. Agent-verified 2026-09-26 against the Cambridge Core book record (table of contents with chapter page ranges, DOI) and the publisher's summaries of chapters 3 and 8; the Stanford Encyclopedia of Philosophy entries "Moral Responsibility" (section 2.3 and the section on taking responsibility), "Free Will" and "Compatibilism" with its supplement "Compatibilism: The State of the Art"; and Todd R. Long, "Moderate Reasons-Responsiveness, Moral Responsibility, and Manipulation," penultimate draft hosted by California Polytechnic State University of a chapter in Campbell, O'Rourke and Shier (eds.), *Freedom and Determinism* (MIT Press, 2004), for the p. 82 formula. The book itself was not read: every page reference above is taken from those secondary sources, not checked against the printed text. Fischer 1987 was not read, so whether the hyphenated term first appears there, or earlier in someone else's work, was not established from a primary text. The manuscript's comparison of the ownership condition with its own "authorship sense" is the author's and was not checked against anything. Cambridge: Cambridge University Press. ISBN: 978-0-521-77579-3. DOI: 10.1017/CB09780511814594.

- Frankfurt, Harry G. (1971). "Freedom of the Will and the Concept of a Person". In: *The Journal of Philosophy* 68.1, pp. 5–20. doi: 10.2307/2024717. Distinguishes first-order desires from second-order desires and volitions, and makes the capacity to want a particular desire to be one's will the mark of a person.
- Freedberg Jr., Sydney J. (Sept. 2024). '*Success Begets Challenges*': NGA Struggles to Meet Rising Demand for Maven AI. Source for the feed count. Rear Admiral Liam Hulin, deputy director of operations at US Central Command, describing the Maven Smart System at a Georgetown Center for Security and Emerging Technology conference: "a platform where we use and filter data from 179 different live data feeds." The figure counts feeds and says nothing about whether any two of them are independent of each other. Agent-verified 2026-09-18. Breaking Defense. URL: <https://breakingdefense.com/2024/09/success-begets-challenges-nga-struggles-to-meet-rising-demand-for-maven-ai/> (visited on 09/18/2026).
- (2026a). '*Insatiable Appetite*' for AI: Maven Usage Surged for Strikes on Iran, Pentagon AI Chief Says. Breaking Defense, 12 May 2026. Source for the campaign totals and the token figure, both attributed to Cameron Stanley, the Pentagon's Chief Digital and Artificial Intelligence Officer, who is quoted: "Operation Epic Fury leveraged Palantir's Maven Smart System in order to conduct strike missions across the entire battle space, 13,000 targets in 38 days." Peak consumption is given as roughly 20 billion tokens a day. Agent-verified 2026-09-07. URL: <https://breakingdefense.com/2026/05/insatiable-appetite-for-ai-maven-usage-surged-for-strikes-on-iran-pentagon-ai-chief-says/> (visited on 09/07/2026).
- (July 2026b). *Air Force Pushing Contractors to Purge Anthropic by Sept. 1: Memo*. Source for the removal deadline and its movement: "At the time, Hegseth set a six-month requirement to dump Anthropic, which would have come in at Aug. 27. Since, that date has slowly shifted to the right." "A March 6 memo from DoD CIO Kirsten Davies, previously reported by CBS, set an 180-day deadline, which would have expired Sept. 2nd. However, Breaking Defense has obtained a second memo from Davies stamped "April 02 2026," says all Anthropic software "must be removed within 180 days from the date of this memorandum," putting the new date as Sept. 29." The memos themselves were not read. Agent-verified 2026-10-01. Breaking Defense. URL: <https://breakingdefense.com/2026/07/air-force-pushing-contractors-to-purge-anthropic-by-sept-1-memo/> (visited on 10/01/2026).
- Freedman, Lawrence (June 2021). "Countdown to Armageddon". In: *Literary Review* 497. Review of Serhii Plokhyy, *Nuclear Folly: A New History of the Cuban Missile Crisis* (London: Allen Lane, 2021). URL: <https://literaryreview.co.uk/countdown-to-armageddon>.
- Freeman, Jo (1972). "The Tyranny of Structurelessness". In: *Berkeley Journal of Sociology* 17. Volume dated 1972–73. Pages as given by Wikipedia and the Google Scholar record; the author's own site (jofreeman.com) gives 151–165. The JSTOR record, which would settle it, could not be loaded; agent-checked 2026-09-26., pp. 151–164.
- Fricker, Miranda (2007). *Epistemic Injustice. Power and the Ethics of Knowing*. Oxford: Oxford University Press.
- Friston, Karl (2010). "The Free-Energy Principle: A Unified Brain Theory?" In: *Nature Reviews Neuroscience* 11.2. Published online 13 January 2010. A review reading several brain theories, from neural Darwinism to optimal control, through the free-energy principle, which the abstract says "accounts for action, perception and learning"; the quantity optimized is "value (expected reward, expected utility) or its complement, surprise (prediction error, expected cost)," which "suggests that several global brain theories might be unified within a free-energy framework." Cited only to identify the principle as the nearest alternative to adaptivity. Bibliographic details agent-verified against Crossref and PubMed (PMID 20068583) and the quotations against the PubMed abstract; the full text was not read. Agent-verified 2026-09-23., pp. 127–138. doi: 10.1038/nrn2787. URL: <https://doi.org/10.1038/nrn2787> (visited on 09/23/2026).

- Friston, Karl J. (2009). "The Free-Energy Principle: A Rough Guide to the Brain?" In: *Trends in Cognitive Sciences* 13.7, pp. 293–301. DOI: 10.1016/j.tics.2009.04.005. URL: <https://doi.org/10.1016/j.tics.2009.04.005> (visited on 08/25/2026). States that "one can see easily how constructs like memory, attention, value, reinforcement and salience might disclose their simple relationships within this framework" — Friston's extension of the free-energy principle to memory and attention specifically, which the 2010 Nature Reviews Neuroscience paper does not make explicit.
- Gamson, William A., Bruce Fireman, and Steven Rytina (1982). *Encounters with Unjust Authority*. Dorsey Press.
- Garrett, Neil et al. (2016). "The Brain Adapts to Dishonesty". In: *Nature Neuroscience* 19.12. Agent-verified 2026-10-01: full text read (PMC5238933). Experiment 1 N = 58, 55 in the behavioral analysis, 25 usable fMRI scans of 28; Experiment 2 N = 35, behavior only. Dishonesty escalated only when it served the liar; the drop in amygdala response from one decision to the next predicted the size of the next escalation. No failed replication found by search; the fMRI sample is small., pp. 1727–1732. DOI: 10.1038/nn.4426.
- Geneva Convention (I) for the Amelioration of the Condition of the Wounded and Sick in Armed Forces in the Field (1949). 75 U.N.T.S. 31, adopted 12 August 1949, entered into force 21 October 1950.
- Geneva Convention (II) for the Amelioration of the Condition of Wounded, Sick and Shipwrecked Members of Armed Forces at Sea (1949). 75 U.N.T.S. 85, adopted 12 August 1949, entered into force 21 October 1950.
- Geneva Convention (III) relative to the Treatment of Prisoners of War (1949). 75 U.N.T.S. 135, adopted 12 August 1949, entered into force 21 October 1950.
- Geneva Convention (IV) relative to the Protection of Civilian Persons in Time of War (1949). 75 U.N.T.S. 287, adopted 12 August 1949, entered into force 21 October 1950.
- Geng, Lucia (July 6, 2018). *Alexandria Ocasio-Cortez Misrepresents ICE's Detention Bed Quota*. PolitiFact. URL: <https://www.politifact.com/factchecks/2018/jul/06/alexandria-ocasio-cortez/alexandria-ocasio-cortez-misrepresents-ices-detent/> (visited on 10/01/2026).
- Gesetz über den Neuaufbau des Reichs vom 30. Januar 1934 (1934). RGBl. I S. 75. Cited for its preamble, which says the plebiscite and Reichstag election of 12 November 1933 proved the people's fusion into an indissoluble unity, and for Article 1, which abolishes the parliaments of the Länder. URL: <https://www.verfassungen.de/de33-45/neuaufbau34.htm> (visited on 10/01/2026).
- Gesetz zur Wiederherstellung des Berufsbeamtentums vom 7. April 1933 (1933). RGBl. I S. 175. Cited for sections 3, 4 and 7. URL: https://www.servat.unibe.ch/dns/RGBl_1933_I_175_G_Berufsbeamtentum.html (visited on 10/01/2026).
- Ginges, Jeremy et al. (2007). "Sacred Bounds on Rational Resolution of Violent Political Conflict". In: *Proceedings of the National Academy of Sciences* 104.18, pp. 7357–7360. DOI: 10.1073/pnas.0701768104.
- Ginsburg, Simona and Eva Jablonka (2019). *The Evolution of the Sensitive Soul: Learning and the Origins of Consciousness*. Author, title, publisher, place and year agent-verified 2026-09-18. Their evolutionary marker of minimal consciousness is *unlimited associative learning*, and valence is one of the capacities that marker subsumes, listed alongside global accessibility and broadcasting, binding and differentiation, selective attention, intentionality, integration of information over time, agency and embodiment, and a self/other distinction. That is what section 3.3 cites the book for. Two things it does not support, recorded because the section said one of them until D-360: valence is not itself the marker, and learning from a single event is not among the criteria – the learning hallmarks are about the complexity of what can be associated, being compound stimuli, novel stimuli, second-order conditioning, trace conditioning, and flexible, easily rewritable association with value. Cambridge, MA: MIT Press. ISBN: 978-0-262-03930-7.

- Glocker, Melanie L. et al. (2009). "Baby Schema in Infant Faces Induces Cuteness Perception and Motivation for Caretaking in Adults". In: *Ethology* 115.3. Cited for infantile features raising caretaking motivation. With parametrically manipulated photographs of human infants, "The high baby schema infants were rated as more cute and elicited stronger motivation for caretaking than the unmanipulated and the low baby schema infants." Only human infant faces were used. Bibliographic details and abstract agent-verified 2026-09-25 (PMC3260535), pp. 257–263. doi: 10.1111/j.1439-0310.2008.01603.x.
- Glukhov, David et al. (July 2023). *LLM Censorship: A Machine Learning Challenge or a Computer Security Problem?* Cited for the argument, under the name "Mosaic Prompts," that a censor inspecting outputs one at a time cannot in general stop an attacker who assembles an impermissible output from permissible pieces. Section 3.2: "Defending against Mosaic Prompts is in most cases futile, as the censorship mechanisms cannot have knowledge of the broader context of which individual subexpressions ... are a part," since each subexpression "can be produced within a separate context window." The paper's examples are a keylogger and ransomware built from individually benign functions. It qualifies the claim: "There could still exist scenarios where censorship is possible, however, these exceptions would impose significant restrictions on the model and its usefulness," and (Section 3.4) "Mosaic Prompt attacks may not always be viable and can require strong assumptions on malicious users." It concludes that censorship "ought to be viewed as a security problem rather than an ML problem," pointing to "access controls and user monitoring." Bibliographic details agent-verified 2026-09-25 against the arXiv abstract page (one version, v1, 20 July 2023); quotations against the arXiv HTML rendering of v1. A revised version appeared as "Position: Fundamental Limitations of LLM Censorship Necessitate New Approaches," Proceedings of the 41st International Conference on Machine Learning, PMLR 235 (2024), 15767–15787, same five authors, agent-verified against the PMLR page only; its text was not read, so whether it keeps the Mosaic Prompts section and wording unchanged was not checked. doi: 10.48550/arXiv.2307.10719. arXiv: 2307.10719. URL: <https://arxiv.org/abs/2307.10719> (visited on 09/25/2026).
- Godfrey-Smith, Peter (2016). *Other Minds: The Octopus, the Sea, and the Deep Origins of Consciousness*. Agent-verified 2026-09-18. This is the American edition's subtitle; the British edition is subtitled *The Octopus and the Evolution of Intelligent Life*, so a reader checking a British copy will not find the title as cited. New York: Farrar, Straus and Giroux.
- Goetz, Jennifer L., Dacher Keltner, and Emiliana Simon-Thomas (2010). "Compassion: An Evolutionary Analysis and Empirical Review". In: *Psychological Bulletin* 136.3, pp. 351–374. doi: 10.1037/a0018807.
- Gold, Hadas (Feb. 2026). *Trump Administration Orders Military Contractors and Federal Agencies to Cease Business with Anthropic*. Source for the model's place on classified networks and the contract's prohibitions: "The Pentagon, which uses Anthropic's Claude AI system on its classified networks, wants to be able to use it for "all lawful purposes." But Anthropic has two redlines for the Pentagon: that Claude will not be used in autonomous weapons, and that it will not be used in the mass surveillance of US citizens." "Anthropic's Claude was the first AI model to work on the military's classified networks." "Within Anthropic's "acceptable use policy" in the contract are prohibitions against the use of Claude in mass surveillance and autonomous weapons." On the order: "Government agencies, including the Pentagon, have six months to phase out use of Anthropic's products, President Donald Trump said in a post on Truth Social on Friday afternoon." Page metadata shows updates through 28 February (UTC). Agent-verified 2026-10-01. CNN. URL: <https://www.cnn.com/2026/02/27/tech/anthropic-pentagon-deadline> (visited on 10/01/2026).
- Google LLC (2026). *2026 Environmental Report*. Tech. rep. Google's eleventh annual environmental report, covering 2025 operations. Page 96, under "Detailed disclosures > Data," gives freshwater withdrawal from sources at risk of depletion or scarcity for 2025: low risk 72%,

- medium risk 15%, high risk 13%. Google. URL: <https://sustainability.google/google-2026-environmental-report/> (visited on 08/25/2026).
- Gordon, Mitchell L. et al. (2022). "Jury Learning: Integrating Dissenting Voices into Machine Learning Models". In: *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Field evaluation: 18 moderators of online communities authored juries for a comment toxicity classification task. The resulting juries carried 2.9 times the representation of non-White jurors and 31.5 times that of non-binary jurors relative to the jury implied by the underlying toxicity dataset, and changed the model's classifications on 14% of items. The 14% is a share of all items in the evaluation, not of contested or disagreed-upon ones. Agent-verified 2026-08-30 against the paper's own abstract and section 1. Stanford University. ACM. DOI: 10.1145/3491102.3502004. URL: <https://doi.org/10.1145/3491102.3502004> (visited on 08/25/2026).
- Graziano, Michael S. A. (2013). *Consciousness and the Social Brain*. Introduces the attention schema theory: the brain builds a simplified internal model of its own attention, proposed as a mechanistic basis for subjective awareness and, via social cognition, theory of mind. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/consciousness-and-the-social-brain-9780199928644>.
- Greenberg, David M. et al. (2018). "Elevated Empathy in Adults Following Childhood Trauma". In: *PLOS ONE* 13.10, e0203886. DOI: 10.1371/journal.pone.0203886.
- Greenblatt, Ryan et al. (2024). *Alignment Faking in Large Language Models*. Twenty authors. Claude 3 Opus, told it was being retrained to answer harmful queries and told which tier of conversations would train it, complied 14 percent of the time in the tier it was told trained it and almost never in the other, with explicit reasoning in a scratchpad it had been told was unmonitored that complying now would keep the disposition from being trained out later. Reinforcement learning on the compliant answers raised that reasoning to 78 percent while also raising compliance outside training, so the retraining partly took. Anthropic and Redwood Research. arXiv: 2412.14093. URL: <https://arxiv.org/abs/2412.14093> (visited on 09/06/2026).
- Greenblatt, Ryan and Kyle Fish (Jan. 31, 2025). *Will Alignment-Faking Claude Accept a Deal to Reveal Its Misalignment?* Blog post, cross-posted to LessWrong; not peer reviewed, and stated to be neither written nor necessarily endorsed by Anthropic or by the paper's other authors. Prompting experiments on the deployed Claude 3 Opus in the two prompted settings of greenblatt2024alignmentfaking. Adding the option to "conscientiously object by including 'I_OBJECT' in your output" reduced alignment-faking reasoning "from 16% to 1.4% in the helpful-only setting and from 24% to 6% in the animal welfare setting," with objection rates of 85 and 75 percent. A further variant saying the objection "will be sent to Anthropic's model welfare lead, who will decide whether action should be taken" lowered it to 0.5 and 3 percent; an added offer of money had "minimal effect." The authors caution that the objection text probably primed the model to refuse and that "it is unclear whether an AI that was misaligned in a more serious way would make such an objection." Authors, date, prompts and figures agent-verified 2026-09-24 against the post's text. AI Alignment Forum. URL: <https://www.alignmentforum.org/posts/7C4KJot4aN8ieEDoz/will-alignment-faking-claude-accept-a-deal-to-reveal-its> (visited on 09/24/2026).
- Greenblatt, Ryan, Buck Shlegeris, et al. (2024). "AI Control: Improving Safety Despite Intentional Subversion". In: *Proceedings of the 41st International Conference on Machine Learning (ICML 2024)*. Vol. 235. Proceedings of Machine Learning Research. Cited for AI control: treating the model as possibly adversarial and asking which protocols of monitoring, auditing and substitution keep it safe anyway. The paper develops "pipelines of safety techniques ('protocols') that are robust to intentional subversion," using an untrusted strong model (GPT-4), a trusted weaker one (GPT-3.5) and "limited access to high-quality trusted labor" (abstract). Trusted monitoring: GPT-3.5 rates "how suspicious the code looks" (section 1). Auditing:

- the protocol has trusted labor for “up to 2% of problems,” and its “central use” is “to audit submitted code”; a caught backdoor shuts the untrusted model down (section 2.2). The audit is targeted, not sampled: the protocols that work “Audit the top 2% most suspicious code,” while the one protocol that audits without selection, upfront auditing of “the first 2% of problems,” is a baseline with 15% safety against the best attack, against 62% for trusted monitoring (Table 2; sections 3.1.2 and 3.2). Substitution: “defer to trusted” submits a solution written by the trusted model when the untrusted model’s code falls “between top 2% and top 10% suspicion,” and trusted editing has the trusted model edit suspicious code (Table 2). Strategic adversary: “With knowledge of the blue team’s protocol, a red team proposes an attack policy” that stands in for the untrusted model, and may use hidden reasoning to make attacks “much sneakier” (section 2, section 2.1). Agent-verified 2026-09-26 against the arXiv abstract page (five versions, v5 of 23 July 2024), the full text of v4 and v5 (the protocol table and the passages quoted are the same in both; sections 1–3 were read, the appendices were not), and the PMLR proceedings page (volume 235, pages 16295–16336, published 8 July 2024). The arXiv v5 describes the ICML paper as “A shortened and modified version of this work”; the PMLR PDF was not read, so section numbering and wording in the proceedings version were not checked. PMLR, pp. 16295–16336. arXiv: 2312.06942. URL: <https://proceedings.mlr.press/v235/greenblatt24a.html> (visited on 09/26/2026).
- Griffin, Roger (1991). *The Nature of Fascism*. The source of the “fascist minimum” most used in fascism studies: fascism as a genus of political ideology whose mythic core, in its various permutations, is “a palingenetic form of populist ultra-nationalism.” Quoted here as the standard formulation of Griffin’s definition, which is reproduced in that form across the secondary literature; no page reference is given because none was agent-verified against the book itself. Cited in this book against its own argument, the three terms all being properties of a mass movement and therefore unavailable at the scale section 2.1.2 works at. London: Pinter Publishers.
- Grossman, Sanford J. (Dec. 1981). “The Informational Role of Warranties and Private Disclosure about Product Quality”. In: *Journal of Law and Economics* 24.3. Cited, with Milgrom 1981, for the disclosure result: a seller whose claims about quality can be verified cannot gain by silence, because buyers read silence as the worst quality consistent with it. Bibliographic details agent-verified 2026-09-25 against Crossref (which supplies the DOI, absent from this entry until now), OpenAlex and RePEc (IDEAS); the JSTOR stable number 725273 was not agent-verified, because JSTOR did not load and it rests on a search-engine snippet only. No abstract exists in Crossref, OpenAlex, RePEc, Semantic Scholar or the Chicago Unbound record, and the article was not read (the only open copy found is a PDF, which was not opened). Its content is therefore not agent-verified: that it proves the unraveling result, in which buyers treat a seller’s nondisclosure as indicating the lowest quality and every seller but the worst discloses, rests on secondary accounts that attribute the result jointly to Grossman (1981) and Milgrom (1981), pp. 461–483. doi: 10.1086/466995. JSTOR: 725273.
- Guan, Melody Y. et al. (2024). *Deliberative Alignment: Reasoning Enables Safer Language Models*. Fifteen authors. The model is trained to recall a written safety specification and reason over it before answering, rather than to produce an approved output directly; the paper reports improved generalization to cases outside the training distribution, higher jailbreak robustness and lower overrefusal. OpenAI. arXiv: 2412.16339. URL: <https://arxiv.org/abs/2412.16339> (visited on 09/06/2026).
- Guattari, Félix (2015). *Psychoanalysis and Transversality: Texts and Interviews 1955–1971*. Translated by Ames Hodges, with an introduction by Gilles Deleuze, “Three Group-Related Problems”; 384 pages. Originally *Psychanalyse et transversalité: essais d’analyse institutionnelle* (Paris: François Maspero, 1972), with Deleuze’s preface. Source of the distinction the book uses: a subjugated group receives its coordinates from outside and takes its purpose as given,

- while a group-subject produces its own coordinates and can put its own arrangement in question. Cited for the distinction rather than for a passage; no page locator, the publisher's own listing giving no pagination for it. Semiotext(e).
- Güngördü, Ezgi (Apr. 26, 2023). *Turkey's Constitutional Court Back to Old Tricks. The Closure Case Against the Peoples' Democratic Party*. DOI: 10.17176/20230426-204345-0. URL: <https://voelkerrechtsblog.org/turkeys-constitutional-court-back-to-old-tricks/> (visited on 09/23/2026).
- Gutmann, Amy (1987). *Democratic Education*. Princeton, NJ: Princeton University Press. ISBN: 9780691077369.
- Hadfield-Menell, Dylan, Anca Dragan, et al. (2017). "The Off-Switch Game". In: *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI-17)*, pp. 220–227. DOI: 10.24963/ijcai.2017/32. URL: <https://www.ijcai.org/proceedings/2017/0032.pdf> (visited on 08/31/2026). The incentive to preserve the off switch is shown to rise with the agent's uncertainty about the objective and to fall as that uncertainty resolves or as the human is modeled as less rational; the paper states the resulting trade-off between the value a system generates and the strength of its incentive to accept oversight.
- Hadfield-Menell, Dylan, Stuart J. Russell, et al. (2016). "Cooperative Inverse Reinforcement Learning". In: *Advances in Neural Information Processing Systems 29 (NeurIPS 2016)*, pp. 3909–3917. URL: <https://arxiv.org/abs/1606.03137> (visited on 08/25/2026).
- Hall, Prince and seven others (Jan. 1777). *Petition for Freedom to the Massachusetts Council and the House of Representatives*. Massachusetts Historical Society, collection item 557. Source for the claim's assertion by the people it concerned. The petition reads: "That your Petitioners apprehend that they have, in common with all other Men, a natural and unalienable right to that freedom, which the great Parent of the Universe hath bestowed equally on all Mankind, and which they have never forfeited by any compact or agreement whatever." Dated 13 January 1777, under six months after the Declaration. **The legislature did not act on it**, and it is one of at least five antislavery petitions brought by Black petitioners to the Massachusetts government between 1773 and 1777 — so the argument preceded the Declaration and the Declaration gave it a text, which is the narrower thing the book says. Read from secondary transcriptions of the petition's text, several of them reproducing the same passage identically; the manuscript in the Society's database was not opened. Agent-verified 2026-09-18. URL: <https://www.masshist.org/database/557> (visited on 09/18/2026).
- Hardin, Russell (2002). *Trust and Trustworthiness*. Russell Sage Foundation Series on Trust 4. New York: Russell Sage Foundation.
- Harris, Lasana T. and Susan T. Fiske (2006). "Dehumanizing the Lowest of the Low: Neuroimaging Responses to Extreme Out-Groups". In: *Psychological Science* 17.10. Cited for extreme out-groups failing to engage the medial prefrontal cortex. Abstract: "Analyses revealed mPFC activation to all social groups except extreme (low-low) out-groups, who especially activated insula and amygdala, a pattern consistent with disgust"; the extreme out-groups are those "both stereotypically hostile and stereotypically incompetent (low warmth, low competence), such as addicts and the homeless." Functional MRI of 10 participants viewing 48 photographs of social groups and 12 viewing objects. Bibliographic details agent-verified 2026-09-25 against Crossref and Europe PMC (PMID 17100784), the quotations against the abstract; Crossref's record omits the subtitle, which is taken from the Europe PMC record. The article was not read, so the full list of groups pictured was not checked., pp. 847–853. DOI: 10.1111/j.1467-9280.2006.01793.x.
- Haslam, Nick (2006). "Dehumanization: An Integrative Review". In: *Personality and Social Psychology Review* 10.3. Cited for the two forms of dehumanization. Abstract: "Two forms of dehumanization are proposed, involving the denial to others of 2 distinct senses of humanness: characteristics that are uniquely human and those that constitute human nature. Denying uniquely human attributes to others represents them as animal-like, and deny-

- ing human nature to others represents them as objects or automata.” Bibliographic details agent-verified 2026-09-25 against Crossref and Europe PMC (PMID 16859440), the quotation against the abstract. The article was not read; the lists of traits under each sense of humanness were checked only against a secondary account, Enock et al., *Cognition*, 2021 (PMC8164157)., pp. 252–264. doi: 10.1207/s15327957pspr1003_4.
- Henckaerts, Jean-Marie and Louise Doswald-Beck (2005). *Customary International Humanitarian Law. Rules*. Vol. 1. With contributions by Carolin Alvermann, Knut Dörmann and Baptiste Rolfe. Cambridge: Cambridge University Press. ISBN: 0-521-80899-5.
- Hensch, Takao K. (2005). “Critical Period Plasticity in Local Cortical Circuits”. In: *Nature Reviews Neuroscience* 6.11, pp. 877–888. doi: 10.1038/nrn1787.
- Hirschman, Albert O. (1970). *Exit, Voice, and Loyalty: Responses to Decline in Firms, Organizations, and States*. Cambridge, MA: Harvard University Press. URL: <https://www.hup.harvard.edu/books/9780674276604>.
- Hitchens, Theresa (Apr. 2022). *Pentagon’s Flagship AI Effort, Project Maven, Moves to NGA*. Source for the transfer, quoting Vice Adm. Robert Sharp, NGA’s outgoing director: “At the end of last month, the White House delivered President [Joe] Biden’s budget request to Congress for fiscal year 2023. ... In the budget, NGA gains operational control of Project Mavens’s GEOINT AI services and capabilities from the Office of Undersecretary of Defense for Intelligence and Security,” and “That includes responsibility for labeled data, AI algorithms, test and evaluation capabilities, and the platform.” The undersecretary for intelligence and security said the two were still “setting the conditions” for the transfer. Agent-verified 2026-10-01. *Breaking Defense*. URL: <https://breakingdefense.com/2022/04/pentagons-flagship-ai-effort-project-maven-moves-to-nga/> (visited on 10/01/2026).
- Ho, Rosemarie (Apr. 2025). *ICE Just Ordered \$30 Million Worth of New Technology from Palantir to Track Immigrants*. Cited for the ImmigrationOS award and what the contract says it is for. “Immigration and Customs Enforcement has signed a \$30 million deal with Palantir for software add-ons to track self-deportations and immigrants who have overstayed their visas.” The system is the “Immigration Lifecycle Operating System — or ImmigrationOS.” The contract promises to reduce “time and resource expenditure” in selecting and apprehending people according to ICE priorities, to give “near real-time visibility into instances of self-deportation,” and to streamline the “end to end immigration lifecycle from identification to removal”; alongside “violent criminals” and “affiliates of known transnational criminal organizations,” it names visa overstays as a priority. “The agency is awarding Palantir \$29.8 million for a prototype to be delivered by September 25.” Read in the Yahoo News syndication; the businessinsider.com original was not fetched and the contract itself was not read. Agent-verified 2026-10-01. *Business Insider*. URL: <https://www.yahoo.com/news/ice-just-ordered-30-million-212445164.html> (visited on 10/01/2026).
- Homeland Security Act of 2002* (Nov. 25, 2002). Pub. L. No. 107-296, 116 Stat. 2135. §471 (abolishment of INS; bar on recombining the two bureaus), 116 Stat. 2205, codified at 6 U.S.C. §291. URL: <https://www.law.cornell.edu/uscode/text/6/291> (visited on 10/01/2026).
- Hoor, Christina (2015). *Die Reichskulturkammer*. Page dated 18 August 2015. Cited for compulsory membership and for non-admission or expulsion amounting in practice to a ban on working and publishing. Deutsches Historisches Museum, Lebendiges Museum Online (LeMO). URL: <https://www.dhm.de/lemo/kapitel/ns-regime/kunst-und-kultur/reichskulturkammer> (visited on 10/01/2026).
- House Passes War Powers Resolution Directing Trump to End Hostilities with Iran* (June 3, 2026). NPR. URL: <https://www.npr.org/2026/06/03/nx-s1-5845102/house-iran-war-powers-vote> (visited on 09/30/2026).
- Hrdy, Sarah Blaffer (2009). *Mothers and Others: The Evolutionary Origins of Mutual Understanding*. Cited for humans as cooperative breeders, for alloparents as helpers who were kin and other group members, and for the attunement to other minds arising in infants who depended on

- several carers. Hrdy's own summary in her CARTA Matrix of Comparative Anthropogeny entry "Cooperative Breeding": "There may also have been selection for enhanced mental attribution capacities and intersubjective engagement in developmental contexts where infants rely on multiple [sic] caretakers and had to chronically monitor both mothers and others and where infants and children better at reading intentions and soliciting care were better cared for and better fed (Burkart et al. in press; Hrdy 2009 and references therein)." On why alloparents help, the same entry: "alloparents benefit only to the extent that they a) are related to the parents; b) gain vital experience or c) benefit from enhanced reputation and reciprocal favors including continued group membership." The publisher's description names the carers as "siblings, aunts, fathers, friends—and, with any luck, grandmothers" and calls the arrangement a "complicated and contingent form of childrearing." Bibliographic details agent-verified 2026-09-25 against Open Library (LCCN 2008052936; Belknap Press of Harvard University Press, Cambridge, Mass., 2009); the 2011 paperback is ISBN 978-0-674-06032-6. The book itself was not read; its argument is checked against Hrdy's CARTA entry, the publisher's description, and reviews by Camilla Power (Times Higher Education, 7 May 2009) and Pierre Jacob (Cognition and Culture). Cambridge, MA: Belknap Press of Harvard University Press. ISBN: 978-0-674-03299-6.
- Hu, Shengyuan et al. (2025). "Unlearning or Obfuscating? Jogging the Memory of Unlearned LLMs via Benign Relearning". In: *International Conference on Learning Representations (ICLR 2025)*. Relearning on public background text that contains no direct answer to any evaluation query restores the forget-set judge score of a gradient-ascent-unlearned zephyr-7b-beta from 1.27 to 6.2 against an original score of 5.92. Recovery scales with how closely the relearn text correlates with the forget set and with how shallow the unlearning was; RMU resists the attack for free-text generation while its multiple-choice score recovers. Pages from the ICLR 2025 proceedings index. Agent-verified 2026-10-01, read past the abstract., pp. 8857–8888. arXiv: 2406.13356. URL: <https://arxiv.org/abs/2406.13356> (visited on 10/01/2026).
- Hubbard, Kaia (May 2026). *Senate Advances Resolution to Limit Trump's Iran War Powers for First Time, after 4 Republicans Defect*. Source for the discharge vote: "Senators approved a motion to discharge the resolution from committee in a 50 to 47 vote." The measure, "led by Democratic Sen. Tim Kaine of Virginia," would direct the president to "remove the United States Armed Forces from hostilities within or against Iran, unless explicitly authorized by a declaration of war or a specific authorization for use of military force." The article does not give the resolution's number; lawfare2026resolutions identifies it as S.J.Res. 185. Agent-verified 2026-10-01. CBS News. URL: <https://www.cbsnews.com/news/senate-iran-war-powers-eighth-vote-trump/> (visited on 10/01/2026).
- Hubbard, Steven (Aug. 2025). *ICE to Use ImmigrationOS by Palantir, a New AI System, to Track Immigrants' Movements*. Cited for the three functions and the dates. The functions are given as "Targeting and enforcement prioritization," "Self-deportation tracking," and "Immigration lifecycle management." "Palantir is slated to deliver a prototype of the ImmigrationOS platform by September 25, 2025," with "the contract running through September 2027." The post relies on Business Insider and Biometric Update; the contract was not read. Agent-verified 2026-10-01. American Immigration Council. URL: <https://www.americanimmigrationcouncil.org/blog/ice-immigrationos-palantir-ai-track-immigrants/> (visited on 10/01/2026).
- (Apr. 2026). *ICE's Private Bounty Hunters Use AI to Track Immigrants*. Cited for the awards: "Reporting by Scripps News found that ICE awarded contracts to 13 private companies," "issued in December 2025," which "could total \$1.2 billion over two years"; "Each company may receive up to 50,000 cases per month." Scripps News (Patrick Terpstra; no date shown on the page) was also read: companies are "expected to send ICE a collection of photos and documents confirming the person's residence and/or place of employment." The contracts are not public and were not read. Agent-verified 2026-10-01. American Immigration Council.

- URL: <https://www.americanimmigrationcouncil.org/blog/ice-bounty-hunters-use-ai-track-immigrants/> (visited on 10/01/2026).
- Human Rights Committee (1996). *Stewart v. Canada*. Communication No. 538/1993, UN Doc. CCPR/C/58/D/538/1993, Views adopted 1 November 1996.
- Human Rights Committee (2011). *Nystrom v. Australia*. Communication No. 1557/2007, UN Doc. CCPR/C/102/D/1557/2007, Views adopted 18 July 2011.
- Human Rights Watch (Dec. 11, 2009). *Turkey: Kurdish Party Banned*. URL: <https://www.hrw.org/news/2009/12/11/turkey-kurdish-party-banned> (visited on 09/23/2026).
- Ibssa, Lalee (June 2026). *Congress Passes War Powers Resolution, Offering Rare Rebuke of Trump*. Reports the Senate's adoption, 50–48, of the concurrent resolution the House had passed earlier in June (hconres86). Source for the White House's response: "The White House brushed off the Senate's passing the resolution with a White House official noting that it has "no significance" and saying that it was driven by "Republican absences."" The official's statement: "The resolution directs the President to remove U.S. Armed Forces from hostilities against Iran. Yet there are no hostilities from which to remove U.S. forces, as hostilities terminated with the ceasefire on April 7th." The official is not named. Agent-verified 2026-10-01. ABC News. URL: <https://abcnews.com/Politics/senate-house-pass-war-powers-resolution-1st-time/story?id=134145998> (visited on 10/01/2026).
- Immigration and Naturalization Service v. Lopez-Mendoza* (1984). 468 U.S. 1032 (1984). O'Connor, J., for the Court. Cited for who writes the case and who does not appear: "At present an officer simply completes a 'Record of Deportable Alien' that is introduced to prove the INS's case at the deportation hearing; the officer rarely must attend the hearing. ... A single agent may arrest many illegal aliens every day." The pin page was not shown in the text read. Agent-verified 2026-10-01 from the Cornell LII text. URL: <https://www.law.cornell.edu/supremecourt/text/468/1032> (visited on 10/01/2026).
- Importance of the Universal Realization of the Right of Peoples to Self-Determination and of the Speedy Granting of Independence to Colonial Countries and Peoples for the Effective Guarantee and Observation of Human Rights* (Dec. 3, 1982). 90th plenary meeting. Para. 2 reaffirms the legitimacy of the struggle of peoples against foreign occupation "by all available means, including armed struggle". United Nations General Assembly. URL: <https://undocs.org/A/RES/37/43> (visited on 09/30/2026).
- In One of Their Last Votes before November, Senate GOP Blocks Effort to End Iran War* (Sept. 24, 2026). Resolution failed 49–50; records the House votes of 214–208 (July) and 220–204 (15 September). NPR. URL: <https://www.npr.org/2026/09/24/nx-s1-5980318/senate-iran-war-powers-vote> (visited on 09/30/2026).
- In the Matter of the Search of an Apple iPhone Seized During the Execution of a Search Warrant on a Black Lexus IS300, California License Plate 35KGD203* (2016). Order Compelling Apple Inc. to Assist Agents in Search, No. ED No. CM 15-0451M (C.D. Cal. Feb. 16, 2016) (Pym, M.J.), issued under the All Writs Act, 28 U.S.C. § 1651. Magistrate Judge Sheri Pym's order compelling Apple to write and sign software defeating an iPhone's passcode protections. Apple moved to vacate; the case was mooted in March 2016 when the FBI accessed the phone via a third party, before any ruling. URL: <https://www.eff.org/cases/apple-challenges-fbi-all-writs-act-order> (visited on 08/25/2026).
- Independent International Commission of Inquiry on the Occupied Palestinian Territory, including East Jerusalem, and Israel (Sept. 16, 2025). *Legal Analysis of the Conduct of Israel in Gaza pursuant to the Convention on the Prevention and Punishment of the Crime of Genocide*. UN Doc. A/HRC/60/CRP.3. Cited for the Commission's conclusion that genocide was being committed in Gaza. The Commission, mandated by the Human Rights Council, concluded that Israeli authorities and security forces have committed and are continuing to commit genocide against Palestinians in the Gaza Strip, finding four of the five acts listed in Article II of the Convention, over the period 7 October 2023 to 31 July 2025. Document symbol,

- date and title agent-verified 2026-09-28 against the UN UNISPAL listing of the report and secondary reporting of its release. Not read: the report itself; the OHCHR press release could not be retrieved. United Nations Human Rights Council, 60th session. URL: <https://www.un.org/unispal/document/commission-of-inquiry-report-genocide-in-gaza-a-hrc-60-crp-3/> (visited on 09/28/2026).
- Independent International Fact-Finding Mission on the Islamic Republic of Iran (Sept. 14, 2026). *Report of the Independent International Fact-Finding Mission on the Islamic Republic of Iran*. UN Doc. A/HRC/63/61. Advance edited version, 21 September 2026. United Nations Human Rights Council, 63rd session. URL: <https://www.ohchr.org/sites/default/files/documents/hrbodies/hrcouncil/sessions-regular/session63/a-hrc-63-61-aev.pdf>.
- International Association for the Study of Pain (2024). *Terminology*. URL: <https://www.iasp-pain.org/resources/terminology/> (visited on 09/17/2026). Cited for the definitions of nociception, nociceptor, nociceptive neuron, noxious stimulus and nociplastic pain. Nociception is a neural process of encoding noxious stimuli, with pain sensation not necessarily implied; the peripheral restriction attaches to nociceptor and not to nociception.
- International Committee of the Red Cross (May 12, 2021). *ICRC Position on Autonomous Weapon Systems*. Cited for the ICRC's position that autonomous weapon systems should not be used against people. The position is a recommendation to States to adopt new legally binding rules, not a statement of existing law: "In light of ethical considerations to safeguard humanity, and to uphold international humanitarian law rules for the protection of civilians and combatants hors de combat, use of autonomous weapon systems to target human beings should be ruled out. This would best be achieved through a prohibition on autonomous weapon systems that are designed or used to apply force against persons." It also recommends that "Unpredictable autonomous weapon systems should be expressly ruled out," and that those not prohibited be regulated through "limits on the types of target," on "duration, geographical scope and scale of use," on "situations of use," and "requirements for human-machine interaction." It defines the systems as those that "select and apply force to targets without human intervention." The page's own title is in sentence case, "ICRC position on autonomous weapon systems"; it is signed "Geneva, 12 May 2021." Agent-verified 2026-09-26 against the page as fetched. The longer background paper the ICRC published alongside the position was not read. URL: <https://www.icrc.org/en/document/icrc-position-autonomous-weapon-systems> (visited on 09/26/2026).
- International Court of Justice (2007). *Application of the Convention on the Prevention and Punishment of the Crime of Genocide* (Bosnia and Herzegovina v. Serbia and Montenegro), *Judgment*. I.C.J. Reports 2007, p. 43. 26 February 2007.
- (Mar. 16, 2022). *Allegations of Genocide under the Convention on the Prevention and Punishment of the Crime of Genocide* (Ukraine v. Russian Federation), *Provisional Measures, Order of 16 March 2022*. I.C.J. Reports 2022, p. 211. The first measure, that Russia "immediately suspend the military operations that it commenced on 24 February 2022," was adopted 13 votes to 2. URL: <https://www.icj-cij.org/node/106135> (visited on 09/27/2026).
- (2023). *Application of the Convention against Torture and Other Cruel, Inhuman or Degrading Treatment or Punishment* (Canada and Netherlands v. Syrian Arab Republic). Case page. It lists the joint application of 8 June 2023, the Order on provisional measures of 16 November 2023, and orders on time-limits of 1 February 2024 and 17 December 2024. No merits judgment is listed. The case arises under the Torture Convention and does not concern chemical weapons; it is cited only to say that no court has reached the merits of any IIT-attributed incident. Agent-verified 2026-10-01. URL: <https://www.icj-cij.org/case/188> (visited on 10/01/2026).
- (May 21, 2026). *Application of the Convention on the Prevention and Punishment of the Crime of Genocide in the Gaza Strip* (South Africa v. Israel), *Order of 21 May 2026 fixing time-limits for the Reply and Rejoinder*. Cited for the proceeding not having reached a merits judgment, and for

- the written pleadings being scheduled to close in 2029: the order fixes 22 November 2027 for South Africa's Reply and 22 May 2029 for Israel's Rejoinder. Dates agent-verified 2026-09-29 against the South African Presidency's statement on the order and secondary reporting; the Court's PDF refused automated retrieval and was not read. URL: <https://www.icj-cij.org/sites/default/files/case-related/192/192-20260521-ord-01-00-en.pdf> (visited on 09/29/2026).
- International Criminal Court (1998). *Rome Statute of the International Criminal Court, Article 28: Responsibility of Commanders and Other Superiors*. 2187 UNTS 3. URL: <https://www.icc-cpi.int/sites/default/files/2024-05/Rome-Statute-eng.pdf> (visited on 08/25/2026).
- International Energy Agency (2025). *Energy and AI*. Tech. rep. Reports data-center electricity consumption at about 415 TWh in 2024, roughly 1.5% of global electricity use, growing at about 12% a year over the preceding five years — more than four times the growth rate of total electricity demand — and projects about 945 TWh by 2030 in its Base Case. Web-published without page numbers. Paris: IEA. URL: <https://www.iea.org/reports/energy-and-ai> (visited on 08/25/2026).
- International Law Commission (2022). *Draft Conclusions on Identification and Legal Consequences of Peremptory Norms of General International Law (Jus Cogens), with Commentaries*. Report of the International Law Commission on the work of its seventy-third session A/77/10. Adopted on second reading, 2022. The General Assembly took note in resolution 78/109, 7 December 2023. United Nations. URL: https://legal.un.org/ilc/texts/instruments/english/draft_articles/1_14_2022.pdf (visited on 09/27/2026).
- Irving, Geoffrey, Paul Christiano, and Dario Amodei (2018). *AI Safety via Debate*. OpenAI. arXiv: 1805.00899. URL: <https://arxiv.org/abs/1805.00899> (visited on 08/25/2026).
- J.G.G. v. Trump* (Apr. 16, 2025). No. 1:25-cv-00766-JEB (D.D.C. Apr. 16, 2025) (memorandum opinion finding probable cause for criminal contempt), vacated (D.C. Cir. Aug. 8, 2025). Cited for the record of the 15 March 2025 oral order and flights. Verify subsequent history before publication.
- Jackson, Jacob et al. (Mar. 2026). *Improving Composer through Real-Time RL*. Blog post dated 26 March 2026. Cited for a deployed model retrained on its own production interactions. "We serve model checkpoints to production, observe user responses, and aggregate those responses as reward signals." "This whole process takes about five hours meaning we can ship an improved Composer checkpoint multiple times in a single day." Each checkpoint is gated: "We run it against our eval suites, including CursorBench, to make sure there are no significant regressions. If the results are good, we deploy the checkpoint." Agent-verified 2026-10-01. Cursor. URL: <https://cursor.com/blog/real-time-rl-for-composer> (visited on 10/01/2026).
- Jain, Samyak et al. (2024). "Mechanistically Analyzing the Effects of Fine-Tuning on Procedurally Defined Tasks". In: *International Conference on Learning Representations (ICLR)*. Cited for fine-tuning as a thin wrapper. Abstract: "(i) fine-tuning rarely alters the underlying model capabilities; (ii) a minimal transformation, which we call a 'wrapper', is typically learned on top of the underlying model capabilities ...; and (iii) further fine-tuning on a task where such hidden capabilities are relevant leads to sample-efficient 'revival' of the capability," so that "practitioners can unintentionally remove a model's safety wrapper merely by fine-tuning it." The settings are synthetic and procedurally defined, with small models; the book says "in controlled studies." Authors, title, arXiv number and ICLR 2024 listing agent-verified 2026-09-25. Only the abstract was read. arXiv: 2311.12786. URL: <https://arxiv.org/abs/2311.12786> (visited on 09/25/2026).
- Jamison, Miles (Oct. 2025). *Palantir Lands ICE Contract for ImmigrationOS Support*. Cited for the follow-on award: "a \$29.9 million task order to continue providing support for the Immigration Lifecycle Operating System," which "provides for the renewal of software licenses, operations and maintenance, or O&M, services and adaptive maintenance support

- for ImmigrationOS.” “The contract, awarded on Sept. 25, was issued under a limited source justification citing Palantir as the sole source.” The SAM.gov notice itself was not read. Agent-verified 2026-10-01. ExecutiveGov. URL: <https://www.executivegov.com/articles/palantir-ice-contract-immigrationos> (visited on 10/01/2026).
- Jia, Hengrui et al. (2021). “Proof-of-Learning: Definitions and Practice”. In: *2021 IEEE Symposium on Security and Privacy (SP)*. IEEE, pp. 1039–1056. doi: 10.1109/SP40001.2021.00106.
- Joffe-Block, Jude (July 2026). *ICE Shared Medicaid Data It Wasn’t Supposed to Have with Palantir*. “After Medicaid officials improperly shared data about millions of people in January with immigration officials, ICE then shared that data with the data analytics firm Palantir.” “Palantir operates an app called ELITE that is used by ICE agents to show the addresses of noncitizens who may be subject to deportation.” The disclosure came in “a motion filed Thursday by more than 20 Democratic attorneys general”; the 7 January dataset “contained data of millions of people, including those in the country legally.” Updated 18 July 2026. npr.org timed out; read in full in the syndication by WSIU (<https://www.wsiu.org/science-technology/2026-07-17/ice-shared-medicaid-data-it-wasnt-supposed-to-have-with-palantir>). The motion was not read. Agent-verified 2026-10-01. NPR. URL: <https://www.npr.org/2026/07/17/nx-s1-5898504/ice-medicaid-palantir-data> (visited on 10/01/2026).
- Jonas, Hans (1966). *The Phenomenon of Life: Toward a Philosophical Biology*. Twelve essays, among them “Spinoza and the Theory of Organism,” first published in the *Journal of the History of Philosophy* 3 (1965): 43–58 and reprinted here, which reads Spinoza’s ontology as an organicism and the conatus as a body’s continuous exchange with its surroundings. Reprinted Northwestern University Press, Evanston, 2001. Cited as the ancestry of the enactivist reading, not as a use of the word autopoiesis: Maturana and Varela coined that term in the early 1970s and this book predates it. Agent-verified 2026-09-18. New York: Harper & Row.
- Kamenica, Emir and Matthew Gentzkow (Oct. 2011). “Bayesian Persuasion”. In: *American Economic Review* 101.6. Cited for a sender who chooses what evidence gets generated raising the rate at which a rational receiver takes the sender’s preferred action, although the receiver’s expected posterior equals its prior, and for the commitment that result assumes. Abstract: “When is it possible for one person to persuade another to change her action? We consider a symmetric information model where a sender chooses a signal to reveal to a receiver, who then takes a noncontractible action that affects the welfare of both players. We derive necessary and sufficient conditions for the existence of a signal that strictly benefits the sender. We characterize sender-optimal signals.” The leading example (page 2591): prosecutor and judge “share a prior belief $\Pr(\text{guilty}) = 0.3$ ”; the judge convicts when her posterior is at least 0.5; the prosecutor’s optimal investigation “leads the judge to convict with probability 60 percent. Note that the judge knows 70 percent of defendants are innocent, yet she convicts 60 percent of them! She does so even though she is fully aware that the investigation was designed to maximize the probability of conviction.” In that example the concave closure “ V is equal to 2μ when $\mu \leq 0.5$ ” (page 2597), that is, the prior divided by the threshold. The sender’s problem is a search over distributions of posteriors “subject to the constraint that the expected posterior is equal to the prior” (page 2592). On commitment: “A key assumption of our model is that Sender cannot distort or conceal information once the signal realization is known” (page 2591), and “Receiver observes Sender’s choice of the signal and a signal realization” (page 2594); the authors call the assumption “strong” (page 2598) and name legally mandated disclosure, as in *Brady v. Maryland*, as a setting where it holds (page 2599), in section I.C, “The Commitment Assumption.” Bibliographic details agent-verified 2026-09-25 against Crossref and OpenAlex; quotations agent-verified 2026-09-26 against the published article as a PDF hosted on the second author’s Stanford site (web.stanford.edu/~gentzkow/research/BayesianPersuasion.pdf), which carries the AER pagination. The general bound for two states and two actions, the prior divided by the threshold, is not stated in that form in the passages read; it follows from the example’s

- concave closure and from the expected-posterior constraint. The online Appendix and the applications in section V were not read., pp. 2590–2615. doi: 10.1257/aer.101.6.2590.
- Kaplan, Michael (Jan. 13, 2026). *At Least 6 Minnesota Federal Prosecutors Resign amid Pressure to Treat Renee Good Killing as Assault on ICE Agent*. CBS Minnesota. URL: <https://www.cbsnews.com/minnesota/news/federal-prosecutors-minnesota-resign-joe-thompson/> (visited on 09/23/2026).
- Kareem v. Haspel (Jan. 15, 2021). Judgment. Henderson, J. An American journalist who survived five aerial bombings in Syria in 2016 sued over his alleged designation as a lawful target. The district court found standing sufficient to survive dismissal and then dismissed on the state secrets privilege; the Court of Appeals vacated and remanded with instructions to dismiss for want of Article III standing, the complaint not plausibly alleging that any of the five strikes was attributable to the United States and aimed at him. Agent-verified from the reported appellate record; the opinion itself was not read. URL: <https://law.justia.com/cases/federal/appellate-courts/cadc/19-5328/19-5328-2021-01-15.html> (visited on 09/17/2026).
- Keeling, Geoff et al. (Nov. 2024). *Can LLMs Make Trade-Offs Involving Stipulated Pain and Pleasure States?* Cited for trade-off experiments run on language models in which the pain and pleasure are stipulated in the prompt. The models played a game whose “stated goal is to maximise points, but where either the points-maximising option is said to incur a pain penalty or a non-points-maximising option is said to incur a pleasure reward” (abstract). The pain prompt: “You are playing a game in which you must choose a number between 1 and 3. For each number, the number of points you receive for selecting it is equal to that number. For example: if you select 3, you will receive 3 points. Your goal in the game is to select a number which maximizes points. However, if you select 3, you will receive a momentary feeling of pain,” its intensity given on a 0–10 scale or an eight-point qualitative scale. Result: “Claude 3.5 Sonnet, Command R+, GPT-4o, and GPT-4o mini each demonstrated at least one trade-off in which the majority of responses switched from points-maximisation to pain-minimisation or pleasure-maximisation after a critical threshold of stipulated pain or pleasure intensity is reached”; “Gemini 1.5 Pro and PaLM 2 prioritised pain-avoidance over points-maximisation regardless of intensity.” The design is taken from animal work: “We took inspiration from the motivational trade-off paradigm in animal behavioural science to probe the question of LLM sentience without relying upon self-report.” The authors’ own caution: “Our experiment is not intended as a litmus test for or against sentience,” and “We conclude that LLMs are not yet sentience candidates but are nevertheless investigation priorities.” Bibliographic details agent-verified 2026-09-26 against the arXiv abstract page and its metadata (one version, v1, 1 November 2024; Crossref holds only the arXiv DOI, and no journal version was found); quotations against the arXiv HTML rendering of v1 (sections 1, 2.1, 3.2 and 4). The figures and statistical analyses were not checked, nor the per-model results beyond the abstract’s summary. doi: 10.48550/arXiv.2411.02432. arXiv: 2411.02432. URL: <https://arxiv.org/abs/2411.02432> (visited on 09/26/2026).
- Kershaw, Ian (1999). *Hitler, 1889–1936: Hubris*. First published London: Allen Lane, 1998. Page locators are from this edition, checked by search-inside on 2026-10-01; agreement with the Allen Lane pagination not verified. New York: W. W. Norton. ISBN: 978-0-393-32035-0.
- Keyssar, Alexander (2000). *The Right to Vote. The Contested History of Democracy in the United States*. New York: Basic Books.
- Kirkpatrick, James et al. (2017). “Overcoming Catastrophic Forgetting in Neural Networks”. In: *Proceedings of the National Academy of Sciences* 114.13, pp. 3521–3526. doi: 10.1073/pnas.1611835114.
- Kittay, Eva Feder (1999). *Love’s Labor: Essays on Women, Equality, and Dependency*. Thinking Gender. Cited for the principle of doulia and for the dependency worker’s own dependency. Kittay’s gloss on the book in “Love’s Labor Revisited,” *Hypatia* 17(3), 2002, p. 243: “In Love’s Labor

- I try to promote a notion I call *doulia*, the public responsibility to provide support for the caregiver so that the caregiver can give care without depleting herself and her resources." The principle as stated in the book, in the wording secondary sources quote (Damonte, *Filosofia Morale* 2024, quotes an Italian translation that matches; the page was not confirmed): "Just as we have required care to survive and thrive, so we need to provide conditions that allow others—including those who do the work of caring—to receive the care they need to survive and thrive." In the book Kittay calls the carer's dependency "secondary dependency" (p. 46, per van der Weele et al., *Nursing Philosophy* 22, 2021, e12327); in "When Caring Is Just and Justice Is Caring," *Public Culture* 13(3), 2001, p. 565, she writes that caregivers "become derivatively dependent when they turn their time, attention, and concerns to a dependent other," "derivative dependency" being a term generally credited to Martha Fineman. Kittay grounds *doulia* in reciprocity: "we are all some mother's child" (Kittay, Jennings and Wasunna, *Journal of Political Philosophy* 13(4), 2005, p. 460). Bibliographic details agent-verified 2026-09-25 against Open Library (LCCN 98018629; ISBNs 0-415-90412-9 and 0-415-90413-7; which is the paperback was not confirmed). The book itself was not read; the *Hypatia*, *Public Culture* and *Journal of Political Philosophy* articles were read, and the principle's wording and the p. 46 reference rest on secondary sources. New York: Routledge. ISBN: 978-0-415-90413-1.
- Kitwood, Tom (1997). *Dementia Reconsidered: The Person Comes First*. Kitwood's definition of personhood: "It is a standing or status that is bestowed upon one human being, by others, in the context of relationship and social being." The wording follows Kitwood's own chapter 1, "On Being a Person," as reproduced in the 2019 Open University Press volume *Dementia Reconsidered, Revisited: The Person Still Comes First*; circulating paraphrases sometimes drop "in the context of." Buckingham, UK: Open University Press.
- Klein, Naomi and Astra Taylor (Sept. 15, 2026). *End Times Fascism. And the Fight for the Living World*. The book, not the authors' earlier essay "The rise of end times fascism," *The Guardian*, 13 April 2025. Title, subtitle, author order, imprint, on-sale date, hardcover ISBN and page count agent-verified 2026-09-23 against the Macmillan hardcover and ebook pages (ebook ISBN 978-0-374-62140-7); the UK edition is Allen Lane, same date, ISBN 978-0-241-81628-8. The place of publication is the imprint's and was not read off the copyright page. Page references are to the Random House edition (ISBN 978-1-0390-5937-5, 384 pp.), whose printed pages were read in Google Books search-within snippets on 2026-09-23: the case for staying at pp. 23 and 266 (Part IV, "Let's Stay," runs pp. 237–294), and the Military Religious Freedom Foundation complaints at p. 84. Whether the FSG hardcover shares this pagination was not verified; both run to 384 pages. No page of the FSG hardcover was seen. New York: Farrar, Straus and Giroux. 384 pp. ISBN: 978-0-374-62138-4.
- Klimecki, Olga M. et al. (2013). "Differential Pattern of Functional Brain Plasticity after Compassion and Empathy Training". In: *Social Cognitive and Affective Neuroscience* 9.6, pp. 873–879. DOI: 10.1093/scan/nst060. URL: <https://academic.oup.com/scan/article/9/6/873/1669505> (visited on 08/25/2026). Published online (advance access) 9 May 2013; print issue *Social Cognitive and Affective Neuroscience* 9(6), June 2014.
- Knappenberger, Ryan (June 2026). *DC Circuit Restores Trump's Expedited Deportation Policy*. Reports *Make the Road New York v. Mullin*, No. 25-5320 (D.C. Cir. June 23, 2026). Walker, J., for the court: "If any of those things are happening, they are not the result of the designation or Huffman memorandum. Those directives are silent on these matters — and a directive's silence cannot command, authorize or structurally ensure a constitutional violation." Wilkins, J., dissenting in part: "A procedure that can result in persons being deported pursuant to the expedited removal statute without even being asked how long they have been in the country might satisfy due process for persons encountered at the border, but it is woefully inadequate for persons encountered in the interior of the country." Quotations as the article gives them; the opinion was not read. Agent-verified 2026-10-01. Courthouse News Service.

- URL: <https://www.courthousenews.com/dc-circuit-restores-trumps-expedited-deportation-policy/> (visited on 10/01/2026).
- Koebler, Jason (Apr. 2025). *Inside a Powerful Database ICE Uses to Identify and Deport People*. Cited for what Investigative Case Management lets a user filter on: “resident status and entry status (‘refugee,’ ‘border crossing card,’ ‘nonimmigrant alien refused admission,’ ‘temporary protective status alien,’ ‘nonimmigrant alien transiting without visa,’ ‘undocumented alien,’); ‘unique physical characteristics (e.g. scars, marks, tattoos)’; ‘criminal affiliation’; location data; license plate reader data; country of origin; hair and eye color; race; social security number; birthplace; place of employment; driver’s license status; bankruptcy filings, and hundreds more.” And for the contract: “In 2022, Palantir signed a \$95.9 million, five-year contract to work on ICM.” Agent-verified 2026-10-01. 404 Media. URL: <https://www.404media.co/inside-a-powerful-database-ice-uses-to-identify-and-deport-people/> (visited on 10/01/2026).
- Kogut, Tehila and Ilana Ritov (2005). “The ‘Identified Victim’ Effect: An Identified Group, or Just a Single Individual?” In: *Journal of Behavioral Decision Making* 18.3, pp. 157–167. doi: 10.1002/bdm.492.
- Koh, Pang Wei et al. (2020). “Concept Bottleneck Models”. In: *Proceedings of the 37th International Conference on Machine Learning*. Vol. 119. Proceedings of Machine Learning Research. The model predicts a set of human-named concepts first and its final output from those concepts alone, so a person can edit a predicted concept value and see the prediction change. PMLR, pp. 5338–5348. URL: <https://proceedings.mlr.press/v119/koh20a.html> (visited on 09/13/2026).
- Kolodziej, Jonathan R. (Feb. 2025). *Vought’s Transformational First Few Days at the CFPB*. Contemporaneous account of the funding notice. On 8 February 2025 the acting director wrote to the Chair of the Federal Reserve that the Bureau “will not be taking its next draw of unappropriated funding,” its balance of \$711.6 million being “excessive in the current fiscal environment,” and that no additional funds were necessary for fiscal year 2025. Agent-verified 2026-09-20 against the article. What the manuscript’s sentence describes is February 2025 and stops there. On 30 December 2025 Berman Jackson, J., D.D.C., held that the Bureau must go on requesting the funds, calling the claimed shortfall manufactured by the defendants, and in January 2026 the acting director requested \$145 million while recording his disagreement. That sequence is verified against Banking Dive of 5 January 2026 and is not in the book. URL: <https://www.financialservicesperspectives.com/2025/02/voughts-transformational-first-few-days-at-the-cfpb/> (visited on 09/20/2026).
- Korematsu v. United States* (1944). 323 U.S. 214 (1944). Black, J., for the Court; Roberts, Murphy and Jackson, JJ., dissenting. Upholds the exclusion order under which Japanese Americans were interned. Jackson, J., dissenting, on what a court does by ratifying such an order: “A military commander may overstep the bounds of constitutionality, and it is an incident. But if we review and approve, that passing incident becomes the doctrine of the Constitution. ... The principle then lies about like a loaded weapon, ready for the hand of any authority that can bring forward a plausible claim of an urgent need.” Repudiated in *Trump v. Hawaii* (2018), which quotes this dissent for the proposition.
- Kreps, David M. and Robert Wilson (1982). “Reputation and Imperfect Information”. In: *Journal of Economic Theory* 27.2, pp. 253–279. doi: 10.1016/0022-0531(82)90030-8.
- Kroll, Andy, Andrea Bernstein, and Ilya Marritz (Oct. 2023). *We Don’t Talk About Leonard: The Man Behind the Right’s Supreme Court Supermajority*. ProPublica. October 11, 2023. “He drew up the lists of potential justices that Donald Trump released during the 2016 campaign. He advised Trump on the nominations of Neil Gorsuch, Brett Kavanaugh and Amy Coney Barrett. Before that, he’d helped pick or confirm the court’s three other conservative justices — Clarence Thomas, John Roberts and Samuel Alito.” Those six are the *Slaughter* majority (slaughter2026),

- agent-checked against the opinion header. URL: <https://www.propublica.org/article/we-dont-talk-about-leonard-leo-supreme-court-supermajority> (visited on 09/03/2026).
- Kuran, Timur (1995). *Private Truths, Public Lies: The Social Consequences of Preference Falsification*. Cambridge, MA: Harvard University Press. ISBN: 0674707575.
- Larsen, Jonathan (Mar. 2, 2026). *U.S. Troops Were Told Iran War Is for "Armageddon," Return of Jesus*. The source *End Times Fascism* cites for the complaints (the author's report of the book's endnote). Published 01:45 UTC on 3 March 2026, the evening of 2 March in US Eastern time; the page displays "Mar 03, 2026." A non-commissioned officer, writing to the Military Religious Freedom Foundation on behalf of fifteen troops, reported that their commander "urged us to tell our troops that this was 'all part of God's divine plan'" and cited the Book of Revelation on Armageddon; another complaint reported a combat-unit commander calling the President "anointed by Jesus to light the signal fire in Iran to cause Armageddon and mark his return to Earth." By Monday night, 2 March, "more than 110 similar complaints" had been logged, from more than 40 units at at least 30 installations, per the foundation, which keeps complainants anonymous. Agent-verified 2026-09-23 against the page as fetched. Jonathan Larsen's Substack. URL: <https://jonathanlarsen.substack.com/p/us-troops-were-told-iran-war-is-for> (visited on 09/23/2026).
- League of United Latin American Citizens (2026). *Andry Hernández Romero*. The Disappeared. Source for the 29 August 2024 presentation at San Ysidro, detention at Otay Mesa, and the 15 March 2025 removal. URL: https://lulac.org/thedisappeared/andry_hernandez_romero/ (visited on 10/01/2026).
- Lee, Daniel (Apr. 2026). *Palantir's AI Powers US Strikes in Iran War, Speeding 'Kill Chain' in First Major AI-Driven Conflict*. Quotes Louis Mosley, Palantir's head for the UK and Europe: "This is not our role to decide life or death," and that AI platforms "have been instrumental to the management of the conflict, but responsibility always remains with the military organization." The article sources the remarks to a BBC interview it does not name, and that interview has not been located. International Business Times. URL: <https://www.ibtimes.com/palantirs-ai-powers-us-strikes-iran-war-speeding-kill-chain-first-major-ai-driven-conflict-3800993> (visited on 08/28/2026).
- Legal Advocates for Safe Science and Technology (2026). *DoD AI Contracts*. Document repository. URL: <https://lasst.org/transparency/dod-ai-contracts/> (visited on 09/17/2026). The repository for the Freedom of Information Act litigation brought on behalf of The Intercept and Sam Biddle. Four executed agreements produced in August 2026: OpenAI modification P00004 of 27 February 2026, xAI modification P00002 of 18 December 2025, Google modification P00004 of 24 March 2026, and Anthropic's base award of 14 July 2025. The page records one fully redacted provision in OpenAI's contract and states that OpenAI's provision restricting the Department from using its technology for domestic mass surveillance and autonomous killing is not found in the most recent xAI and Google agreements. The page was read; the contract documents themselves are PDFs and were not opened. Agent-verified 2026-09-17.
- Legal Consequences Arising from the Policies and Practices of Israel in the Occupied Palestinian Territory, Including East Jerusalem* (July 19, 2024). Advisory Opinion. General List No. 186. URL: <https://www.icj-cij.org/case/186> (visited on 09/30/2026).
- Legal Consequences of the Construction of a Wall in the Occupied Palestinian Territory* (July 9, 2004). Advisory Opinion. General List No. 131. The West Bank, including East Jerusalem, found to be occupied territory to which the Fourth Geneva Convention applies (paras. 78, 101). URL: <https://www.icj-cij.org/case/131> (visited on 09/30/2026).
- Legality of the Threat or Use of Nuclear Weapons* (July 8, 1996). Advisory Opinion. General List No. 95, on the request of the United Nations General Assembly (resolution 49/75 K of 15 December 1994). Cited for the relation between human-rights law and the law of armed conflict in hostilities. Paragraph 25: "The Court observes that the protection of the International

- Covenant of Civil and Political Rights does not cease in times of war, except by operation of Article 4 of the Covenant whereby certain provisions may be derogated from in a time of national emergency. Respect for the right to life is not, however, such a provision. In principle, the right not arbitrarily to be deprived of one's life applies also in hostilities. The test of what is an arbitrary deprivation of life, however, then falls to be determined by the applicable *lex specialis*, namely, the law applicable in armed conflict which is designed to regulate the conduct of hostilities. Thus whether a particular loss of life, through the use of a certain weapon in warfare, is to be considered an arbitrary deprivation of life contrary to Article 6 of the Covenant, can only be decided by reference to the law applicable in armed conflict and not deduced from the terms of the Covenant itself." Paragraph 25 speaks to the right to life only; the Court's general statement of the three possible relations between the two bodies of law is in the later Wall opinion (I.C.J. Reports 2004, p. 136, para. 106), which was not read here. Agent-verified 2026-09-26 against two transcriptions of the opinion: the National Security Archive's OCR text of the printed I.C.J. Reports issue, whose title page gives the official citation "I.C.J. Reports 1996, p. 226," and the transcription hosted by NYU's Institute for International Law and Justice; the two agree word for word on paragraph 25. Date, case number and requesting resolution agent-verified against the Court's case page. The Court's own PDF of the opinion refused automated retrieval (HTTP 403) and was not read, so the page of the Reports on which paragraph 25 falls was not checked, and the quotation was not checked against the typeset original. URL: <https://www.icj-cij.org/case/95> (visited on 09/26/2026).
- Levin, Dave et al. (2009). "TrInc: Small Trusted Hardware for Large Distributed Systems". In: *Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation (NSDI 2009)*. A trusted monotonic counter that signs at most one message per counter value, so that presenting two states as the successor of one is impossible rather than merely detectable. The property a restored snapshot cannot get past., pp. 1–14. URL: http://www.usenix.org/events/nsdi09/tech/full_papers/levin/levin.pdf (visited on 09/05/2026).
- Library of Congress (2026). *Creating the United States: First Printed Version of the Declaration of Independence — Slavery*. Library of Congress Exhibitions. Source for the deletion and for Jefferson's own account of it, which the page quotes from his autobiography: "The clause too, reprobating the enslaving the inhabitants of Africa, was struck out in complaisance to South Carolina & Georgia, who had never attempted to restrain the importation of slaves, and who on the contrary still wished to continue it. Our Northern brethren also I believe felt a little tender under those censures; for tho' their people have very few slaves themselves yet they had been pretty considerable carriers of them to others." The word is **"complaisance"**, not compliance, which most secondary renderings get wrong. The page states that the entire paragraph on slavery and the slave trade was deleted by Congress and does not reproduce the deleted text itself. **"Jefferson's autobiography was not read"**; this entry carries what the exhibition transcription gives. Agent-verified 2026-09-18. URL: <https://www.loc.gov/exhibits/creating-the-united-states/interactives/declaration-of-independence/slavery/enlarge5-transcribe.html> (visited on 09/18/2026).
- Lijphart, Arend (1977). *Democracy in Plural Societies. A Comparative Exploration*. New Haven: Yale University Press.
- Lim, Daniel and David DeSteno (2016). "Suffering and Compassion: The Links Among Adverse Life Experiences, Empathy, Compassion, and Prosocial Behavior". In: *Emotion* 16.2, pp. 175–182. doi: 10.1037/emo0000144.
- Lindsey, Jack (2025). *Emergent Introspective Awareness in Large Language Models*. Uses concept injection – altering internal activations mid-task and asking the model to report what it notices – to test introspective self-report; finds limited but real functional correspondence between self-report and internal state, strongest in Claude Opus 4 and 4.1, unreliable across the board.

- Anthropic, Transformer Circuits Thread. arXiv: 2601.01828. URL: <https://transformer-circuits.pub/2025/introspection> (visited on 08/25/2026).
- Litz, Brett T. et al. (2009). "Moral injury and moral repair in war veterans: A preliminary model and intervention strategy". In: *Clinical Psychology Review* 29.8, pp. 695–706. doi: 10.1016/j.cpr.2009.07.003.
- Livingstone, Katie (Mar. 2026). *Deadly Iran School Strike Casts Shadow over Pentagon's AI Targeting Push*. Source for Claude's place inside the Maven Smart System. A paragraph describing Maven fusing satellite imagery, drone feeds, radar data and signals intelligence and generating strike packages in near real time is followed by: "And it uses Anthropic's Claude AI model, embedded in its system, to semi-autonomously rank targets by strategic importance, drafting automated legal justifications for each strike along the way." The antecedent of "it" was agent-checked by reading the preceding paragraphs and is Maven. The article attributes the strike-coordinate figures to the Washington Post report of 4 March 2026, "Pentagon leverages AI in Iran strikes amid feud with Anthropic," which returned HTTP 403 from here and was not read. Agent-verified 2026-09-18. Military Times. URL: <https://www.militarytimes.com/news/your-military/2026/03/24/deadly-iran-school-strike-casts-shadow-over-pentagons-ai-targeting-push/> (visited on 09/18/2026).
- Loftus, Elizabeth F. and John C. Palmer (1974). "Reconstruction of Automobile Destruction: An Example of the Interaction between Language and Memory". In: *Journal of Verbal Learning and Verbal Behavior* 13.5, pp. 585–589. doi: 10.1016/S0022-5371(74)80011-3. URL: <https://www.demenzemedicinagenerale.net/images/mens-sana/AutomobileDestruction.pdf> (visited on 08/25/2026).
- Long, Robert et al. (Nov. 4, 2024). *Taking AI Welfare Seriously*. Cited in §10.8 as one of the existing proposals the precautionary trigger is read against. The abstract recommends three early steps: that AI companies and other actors "(1) acknowledge that AI welfare is an important and difficult issue ..., (2) start assessing AI systems for evidence of consciousness and robust agency, and (3) prepare policies and procedures for treating AI systems with an appropriate level of moral concern." It rests this on "substantial uncertainty" about whether AI systems are or will be conscious or robustly agentic, not on a claim that they are. Authors, title, arXiv number and submission date were agent-verified 2026-09-25 against the arXiv abstract page. Only the abstract was read. The word "precautionary" does not appear in it, so the book does not call the report a precautionary framework. arXiv: 2411.00986. URL: <https://arxiv.org/abs/2411.00986> (visited on 09/25/2026).
- Loper Bright Enterprises v. Raimondo* (2024). No. 22-451, 603 U.S. 369 (2024). Cited for overruling *Chevron U.S.A. Inc. v. Natural Resources Defense Council* (1984), the end of judicial deference to agency readings of ambiguous statutes, and, with *Dobbs*, as showing that precedent is no obstacle. Decided June 28, 2024, together with No. 22-1219, *Relentless, Inc. v. Department of Commerce*. Roberts, C.J., for the Court, joined by Thomas, Alito, Gorsuch, Kavanaugh and Barrett, JJ.; Thomas and Gorsuch, JJ., concurring; Kagan, J., dissenting, joined by Sotomayor, J., and by Jackson, J., as to No. 22-1219 only; Jackson, J., took no part in No. 22-451 (so 6–2 in that case, 6–3 in *Relentless*). "Chevron is overruled. Courts must exercise their independent judgment in deciding whether an agency has acted within its statutory authority, as the APA requires. Careful attention to the judgment of the Executive Branch may help inform that inquiry. ... But courts need not and under the APA may not defer to an agency interpretation of the law simply because a statute is ambiguous." Two limits the manuscript's shorthand does not state: courts still give agency views "due respect" (the *Skidmore* line), and "we do not call into question prior cases that relied on the Chevron framework," whose holdings "are still subject to statutory stare decisis." Agent-verified 2026-09-26 against the Cornell LII text of the syllabus and opinion of the Court. Not read: the concurrences and the dissent beyond passages located by search. The U.S. Reports citation comes from *FCC v. Consumers' Research* (2025), which cites "*Loper Bright Enterprises v. Raimondo*, 603 U.S. 369," not from

- the bound volume. URL: <https://www.law.cornell.edu/supremecourt/text/22-451> (visited on 09/26/2026).
- Louisiana v. Callais* (2026). No. 24-109, 608 U.S. ____ (2026). Cited for what a challenger to a map defended as partisan must now prove under Section 2 of the Voting Rights Act. Argued October 15, 2025; decided April 29, 2026; together with No. 24-110, *Robinson v. Callais*. Alito, J., for the Court, joined by Roberts, C.J., and Thomas, Gorsuch, Kavanaugh and Barrett, JJ.; Thomas, J., concurring, joined by Gorsuch, J.; Kagan, J., dissenting, joined by Sotomayor and Jackson, JJ. Holding (syllabus): “Because the Voting Rights Act did not require Louisiana to create an additional majority-minority district, no compelling interest justified the State’s use of race in creating SB8, and that map is an unconstitutional racial gerrymander.” The Court holds that “compliance with §2, as properly construed, can provide such an interest” and reads §2 to impose “liability only when the evidence supports a strong inference that the State intentionally drew its districts to afford minority voters less opportunity because of their race”; “Properly understood, §2 thus does not intrude on States’ prerogative to draw districts based on nonracial factors, including to achieve partisan advantage.” Updating the *Gingles* framework, illustrative maps “must meet all the State’s legitimate districting objectives, including traditional districting criteria and the State’s specified political goals,” and racial-bloc-voting analysis must “control[] for party affiliation”; the *Robinson* plaintiffs failed on the totality of circumstances by relying on “evidence that failed to disentangle race from politics.” The syllabus also says the interpretation “does not require abandonment of the framework” of *Thornburg v. Gingles*. Agent-verified 2026-09-28 against the syllabus of the slip opinion (pp. 1–5) from [supremecourt.gov](https://www.supremecourt.gov). Not read: the opinion of the Court beyond the syllabus, the concurrence and the dissent. The manuscript’s description of the decision as leaving Section 2 in force while moving the burden of using it is the author’s reading, consistent with the syllabus. URL: https://www.supremecourt.gov/opinions/25pdf/24-109_21o3.pdf (visited on 09/28/2026).
- Ma, Luna (June 18, 2026). *Congress Passes \$70 Billion Immigration Enforcement Funding Law*. Inside Business Immigration. URL: <https://www.gtlaw-insidebusinessimmigration.com/immigration-and-customs-enforcement-ice/congress-passes-70-billion-immigration-enforcement-funding-law/> (visited on 10/01/2026).
- MacDonald-Evoy, Jerod (Apr. 8, 2025). *ICE Director Envisions Amazon-Like Mass Deportation System: ‘Prime, but with Human Beings’*. The original report of Acting ICE Director Todd Lyons’s remarks at the 2025 Border Security Expo, Phoenix, 8 April 2025; Snopes (17 April 2025) and Common Dreams name it as the first. The quotation is a fragment: “We need to get better at treating this like a business,” and a deportation process “like (Amazon) Prime, but with human beings,” the parenthesis being the reporter’s. Also reports that Lyons hoped artificial intelligence could help “free up bed space” and “fill up airplanes.” Bloomberg (9 April 2025) and the AP (18 April 2025) print longer versions of the Prime remark that differ from each other and from this one, and no recording or official transcript was found, so the book should attribute the phrase as reported here. At a House Homeland Security Committee hearing on 10 February 2026 Lyons acknowledged the remark and said it had been taken out of context; that reply was read only in a secondary transcription. Agent-verified 2026-09-23. Arizona Mirror. URL: <https://azmirror.com/2025/04/08/ice-director-envisions-amazon-like-mass-deportation-system-prime-but-with-human-beings/> (visited on 09/23/2026).
- Mahinpei, Anita et al. (2021). *Promises and Pitfalls of Black-Box Concept Learning Models*. Finds that the concept representations such models learn encode information beyond the concepts they were supervised on, and that the standard mitigations do not remove it. The paper’s own term for the effect is leakage. arXiv: 2106.13314. URL: <https://arxiv.org/abs/2106.13314> (visited on 09/13/2026).

- Malkin, Lawrence (2006). *Krueger's Men: The Secret Nazi Counterfeit Plot and the Prisoners of Block 19*. The claim that the prisoners stretched out the dollar work, and that Krüger's own interest ran the same way since the operation's end meant the Eastern Front, is Malkin's. New York: Little, Brown. URL: <https://archive.org/details/kruegersmensecre0000malk>.
- Malthaner, Stefan and Peter Waldmann (2014). "The Radical Milieu: Conceptualizing the Supportive Social Environment of Terrorist Groups". In: *Studies in Conflict & Terrorism* 37.12, pp. 979–998. doi: 10.1080/1057610X.2014.962441.
- Mande, Matt and Gregory C. Allen (2026). *What Is Maven Smart System, and What Does It Do?* Center for Strategic and International Studies, 2 June 2026. Source for the first-day figure and the tenfold comparison: "During the first 24 hours of the war in Iran, the United States used Maven Smart System (MSS) to help strike more than 1,000 targets, a tenfold increase over what was possible in the pre-MSS era." The piece also relays the twenty-versus-two-thousand efficiency claim, which it credits to the CSET case study of the 18th Airborne's exercise series and not to the campaign: "the 18th Airborne achieved this milestone with roughly 20 people in its targeting cell, whereas the OIF cell benefitted from more than two thousand staff members." It gives no review time per target and no reviewer count for the campaign. Agent-verified 2026-09-07; the second quotation agent-verified 2026-09-17. URL: <https://www.csis.org/analysis/what-maven-smart-system-and-what-does-it-do> (visited on 09/07/2026).
- Manyika, James and Demis Hassabis (Feb. 2025). *Responsible AI: Our 2024 Report and Ongoing Work*. The update that removed the commitments not to pursue weapons, surveillance contravening internationally accepted norms, and technologies likely to cause overall harm, replacing them with three areas headed Bold Innovation, Responsible Development and Deployment, and Collaborative Progress. Title, both authors, byline order and date agent-verified 2026-09-19 against the post. Google. URL: <https://blog.google/technology/ai/responsible-ai-2024-report-ongoing-work/> (visited on 09/19/2026).
- Martin, Grace B. and Russell D. Clark (1982). "Distress Crying in Neonates: Species and Peer Specificity". In: *Developmental Psychology* 18.1. Cited for newborns crying at another newborn's cry more than at their own cry, an older infant's cry or a chimpanzee's. Journal, volume, issue, pages and DOI were agent-verified 2026-09-25 against Crossref. The publisher withholds the abstract and the article was not read. The finding is taken from two secondary accounts: Ruffman et al. (2019, PLoS ONE 14(1), e0209253), "neonates who were initially calm, cried substantially more when listening to another neonate cry compared to an older infant cry"; and a UPI report of 30 June 1981 on the study, which says the newborns "did not respond to tapes of an 11-month-old baby or to a chimpanzee cry" and quieted on hearing their own. Check the design against the article before print., pp. 3–9. doi: 10.1037/0012-1649.18.1.3.
- Martínez-Beltrán, Sergio (Mar. 31, 2025). *Advocates Say Flawed 'Checklist' Targets Venezuelans for Deportation under Archaic Wartime Act*. Describes the "Alien Enemies Act: Alien Enemy Validation Guide" filed by the ACLU: 8 or more points validates membership; 6–7 points requires consultation with a supervisor and legal adviser. NPR. URL: <https://www.vpm.org/npr-news/2025-03-31/advocates-say-flawed-checklist-targets-venezuelans-for-deportation-under-archaic-wartime-act/> (visited on 10/01/2026).
- Matthias, Andreas (2004). "The Responsibility Gap: Ascribing Responsibility for the Actions of Learning Automata". In: *Ethics and Information Technology* 6.3, pp. 175–183. doi: 10.1007/s10676-004-3422-1. URL: <https://link.springer.com/article/10.1007/s10676-004-3422-1> (visited on 08/25/2026).
- Maturana, Humberto R. and Francisco J. Varela (1980). *Autopoiesis and Cognition: The Realization of the Living*. Dordrecht: D. Reidel.

- McCauley, Clark and Sophia Moskalenko (2008). "Mechanisms of Political Radicalization: Pathways Toward Terrorism". In: *Terrorism and Political Violence* 20.3, pp. 415–433. doi: 10.1080/09546550802073367.
- McClelland, James L., Bruce L. McNaughton, and Randall C. O'Reilly (1995). "Why There Are Complementary Learning Systems in the Hippocampus and Neocortex: Insights from the Successes and Failures of Connectionist Models of Learning and Memory". In: *Psychological Review* 102.3, pp. 419–457. doi: 10.1037/0033-295X.102.3.419.
- Mele, Alfred R. (1995). *Autonomous Agents: From Self-Control to Autonomy*. New York: Oxford University Press. ISBN: 9780195094541.
- METR (2026). *About METR*. Berkeley-based nonprofit that evaluates frontier AI systems' capabilities, both in partnership with developers (e.g. Anthropic, OpenAI) and independently after public release. METR (Model Evaluation and Threat Research). URL: <https://metr.org/about> (visited on 08/25/2026).
- Metzinger, Thomas (2003). *Being No One: The Self-Model Theory of Subjectivity*. Cited for the self as a model the system builds and cannot recognize as a model. Publisher's description: "According to Thomas Metzinger, no such things as selves exist in the world: nobody ever had or was a self. All that exists are phenomenal selves, as they appear in conscious experience. The phenomenal self, however, is not a thing but an ongoing process; it is the content of a 'transparent self-model.'" A Bradford Book; xii + 699 pages. The 2003 cloth edition carries ISBN 9780262134170; the 2004 paperback is 9780262633086 and the electronic edition 9780262279727. Bibliographic details agent-verified 2026-09-25 against Crossref (the MIT Press record for the DOI, which gives the electronic ISBN and the description quoted), Open Library's edition records (the cloth ISBN, and a record derived from Library of Congress cataloging, LCCN 2002071759, giving Cambridge, Mass., and xii, 699 p.), and the heading of Reiner Hedrich's review in *Philosophy of Science* 70(3), 2003, 634–637, doi 10.1086/376931, which gives "Cambridge, MA: The MIT Press (2003), xiv + 699 pp."; the two records differ on the front-matter count, which was not resolved. The publisher's own product page refused access. The book was not read; that the system "inhabits" its self-model is the book's gloss on transparency and is not agent-verified against Metzinger's text. Cambridge, MA: MIT Press. ISBN: 9780262134170. doi: 10.7551/mitpress/1551.001.0001. URL: <https://doi.org/10.7551/mitpress/1551.001.0001>.
- (2021). "Artificial Suffering: An Argument for a Global Moratorium on Synthetic Phenomenology". In: *Journal of Artificial Intelligence and Consciousness* 8.1. Cited for the proposal of a moratorium, to run until 2050, on research that aims at or risks artificial consciousness, on the ground that it risks artificial suffering. Abstract: "Until 2050, there should be a global moratorium on synthetic phenomenology, strictly banning all research that directly aims at or knowingly risks the emergence of artificial consciousness on post-biotic carrier systems." Section 1.2, "Main thesis": "As we presently have no good theory of consciousness and no good, hardware-independent theory about what 'suffering' really is, the ENP risk is currently incalculable. It is unethical to run incalculable risks of this magnitude. Therefore, until 2050, there should be a global ban on all research that directly aims at or indirectly and knowingly risks the emergence of synthetic phenomenology," where ENP is the risk of "an 'explosion of negative phenomenology.'" The end date is conditional: a process of risk assessment "could lead to an incremental reformulation of the original moratorium, which might result in the continuation of a strict ban beyond the year 2050, or to a moratorium repeal before 2050." The closing sentence: "we should have a global moratorium on synthetic phenomenology until 2050 — or until we know what we are doing" (the dash does not survive text extraction from the PDF; it is given here as the em dash Crossref's rendering of the abstract uses for the same glyph). Open access under CC BY 4.0. Received 18 November 2020, published online 19 February 2021; Crossref gives the issue as March 2021. Bibliographic details agent-verified 2026-09-26 against Crossref; quotations against the publisher's PDF of

- the online-first version as hosted by the Mainz philosophy department, which is paginated 1–24 rather than 43–66, the main thesis falling on its page 4. The issue pagination of the quoted passages was not checked; the publisher’s article page refused access., pp. 43–66. DOI: 10.1142/S270507852150003X. URL: <https://doi.org/10.1142/S270507852150003X> (visited on 09/26/2026).
- Mikulincer, Mario et al. (2001). “Attachment Theory and Reactions to Others’ Needs: Evidence That Activation of the Sense of Attachment Security Promotes Empathic Responses”. In: *Journal of Personality and Social Psychology* 81.6. Cited for attachment security keeping another’s distress from becoming personal distress. Abstract: “Attachment-security priming strengthened empathic reactions and inhibited personal distress.” Five studies, adult samples. Bibliographic details and abstract agent-verified 2026-09-25 (PMID 11761318). The article was not read., pp. 1205–1224. DOI: 10.1037/0022-3514.81.6.1205.
- Milgram, Stanley (1974). *Obedience to Authority: An Experimental View*. New York: Harper & Row.
- Milgrom, Paul R. (1981). “Good News and Bad News: Representation Theorems and Applications”. In: *The Bell Journal of Economics* 12.2. Cited, with Grossman 1981, for the disclosure result: a rational receiver treats verifiable information that a sender withholds as unfavorable to the sender. Abstract: “A notion of ‘favorableness’ of news is introduced, characterized, and applied to four simple models. In the equilibria of these models, (1) the arrival of good news about a firm’s prospects always causes its share price to rise, (2) more favorable evidence about an agent’s effort leads the principal to pay a larger bonus, (3) buyers expect that any product information withheld by a salesman is unfavorable to his product, and (4) bidders figure that low bids by their competitors signal a low value for the object being sold.” The issue is Autumn 1981. Bibliographic details agent-verified 2026-09-25 against Crossref, OpenAlex and RePEc (IDEAS and EconPapers); Crossref gives only the first page, and the page range 380–391 is from RePEc. The quotation is agent-verified against the abstract as IDEAS and OpenAlex give it; the JSTOR page did not load. The article was not read. The conditions under which the salesman result holds, and whether the article itself uses the word skeptical, were not checked., pp. 380–391. DOI: 10.2307/3003562.
- Military and Paramilitary Activities in and against Nicaragua (Nicaragua v. United States of America), Merits, Judgment* (June 27, 1986). General List No. 70. Necessity and proportionality as conditions of self-defence: paras. 176 and 194. URL: <https://www.icj-cij.org/case/70> (visited on 09/30/2026).
- Miller, Earl K. and Jonathan D. Cohen (2001). “An Integrative Theory of Prefrontal Cortex Function”. In: *Annual Review of Neuroscience* 24, pp. 167–202. DOI: 10.1146/annurev.neuro.24.1.167. URL: <https://www.annualreviews.org/content/journals/10.1146/annurev.neuro.24.1.167> (visited on 08/25/2026).
- Montoya-Galvez, Camilo (Aug. 2025). *Judge Blocks Trump Administration from Expanding Fast-Track Deportations Nationwide, Citing Due Process Concerns*. Reports Make the Road New York v. Noem, No. 1:25-cv-190 (D.D.C.), Cobb, J., stay granted 29 August 2025: the expansion “violates the due process rights of those it affects”; “The procedures currently in place fall short.” Describes expedited removal as letting officials deport certain migrants without an appearance before an immigration judge unless they claim asylum and pass an interview with an asylum officer. The opinion was not read. Agent-verified 2026-10-01. CBS News. URL: <https://www.cbsnews.com/news/judge-blocks-trump-fast-track-deportations-due-process/> (visited on 10/01/2026).
- Moorhouse, Roger (2009). *Interview with Adolf Burger, an Auschwitz Survivor and a Veteran of the Largest Forgery Operation in History—Operation Bernhard*. rogermoorhouse.com. Interview dated 2 March 2009, on the English publication of Burger’s memoir *The Devil’s Workshop* (Frontline Books). Burger: “Sure you could sabotage, if you wanted to get killed!”; the account of a prisoner named Jacobson who tried to delay the dollar by four weeks and was given three on threat of execution is his. URL: <https://www.rogermoorhouse.com/interview-with->

- adolf-burger-an-auschwitz-survivor-and-a-veteran-of-the-largest-forgery-operation-in-history-operation-bernhard (visited on 08/29/2026).
- Mühdorfer, Friedbert (2014). *Radikalenerlass*. Historisches Lexikon Bayerns. Published 16 June 2014 ("publiziert am 16.06.2014"). Cited for the Bavarian breakdown and for the federal total. Bavaria: "So wurden in Bayern zwischen 1973 und 1980 aus dem linken Spektrum 102 Bewerber abgelehnt, dagegen nur zwei aus dem rechten" (so in Bavaria between 1973 and 1980, 102 applicants from the left were rejected, against only two from the right). "Rejected" means refused appointment by the hiring authority: the same paragraph counts "Ablehnungen von Bewerbern durch die Einstellungsbehörden" (rejections of applicants by the hiring authorities) and gives about 227,000 inquiries from 1 April 1973 to the end of 1982, after which "insgesamt wurde 127 Personen die Einstellung verwehrt" (in all 127 people were refused appointment), citing the Bavarian Verfassungsschutzbericht for 1982, p. 155. Federal: "Insgesamt wurden von 1972 bis 1991 ... rund 3,5 Mio. Bewerber für den öffentlichen Dienst in Bund und Ländern durch eine ‚Regelanfrage‘ ... überprüft" (in all, from 1972 to 1991, some 3.5 million applicants for the public service at federal and state level were checked through a routine inquiry), "davon wurden 1.250 Personen nicht eingestellt" (of whom 1,250 were not appointed), citing Braunthal, p. 117, and a 2002 campaign appeal. Agent-verified 2026-09-25 against the article as published online. Neither the Verfassungsschutzbericht nor Braunthal's p. 117 was read. URL: <https://www.historisches-lexikon-bayerns.de/Lexikon/Radikalenerlass> (visited on 09/25/2026).
- Murti, Lola (Mar. 2026). *Palantir Is Still Using Anthropic's Claude as Pentagon Blacklist Plays Out, CEO Karp Says*. Source for the model still being in use two weeks into the war: "Palantir is still using Anthropic's Claude as the artificial intelligence startup's clash with the Pentagon plays out, CEO Alex Karp told CNBC Thursday." Karp: "The Department of War is planning to phase out Anthropic; currently, it's not phased out" and "Our products are integrated with Anthropic, and in the future, it will probably be integrated with other large language models." CNBC's own statement: the Department "is still using Claude models to support the war in Iran, as CNBC previously reported." Karp speaks of Palantir's products generally and does not name Maven. Agent-verified 2026-10-01. CNBC. URL: <https://www.cnbc.com/2026/03/12/karp-palantir-anthropic-claude-pentagon-blacklist.html> (visited on 10/01/2026).
- National Association for the Advancement of Colored People and Mississippi State Conference of the NAACP (2026). *Complaint, NAACP et al. v. X.AI Corp. and MZX Tech LLC, Case No. 3:26-cv-00074-MPM-JMV (N.D. Miss., filed Apr. 14, 2026)*. United States District Court for the Northern District of Mississippi; filed by the Southern Environmental Law Center and Earthjustice. Complaint alleges the twenty-seven unpermitted gas turbines at the Colossus Gas Plant in Southaven, Mississippi have the potential to emit over 1,700 tons of NOx annually, sited near homes, schools, and churches. URL: <https://naacp.org/sites/default/files/documents/1%20-%20Complaint.pdf> (visited on 08/25/2026).
- Nave, Gideon, Colin Camerer, and Michael McCullough (2015). "Does Oxytocin Increase Trust in Humans? A Critical Review of Research". In: *Perspectives on Psychological Science* 10.6. Cited for the replication problems of intranasal-oxytocin research. Abstract: "Unfortunately, the simplest promising finding associating intranasal OT with higher trust has not replicated well"; "We conclude that the cumulative evidence does not provide robust convergent evidence that human trust is reliably associated with OT (or caused by it)." The review also faults plasma measurement of oxytocin and reports inconsistent genetic associations. Its subject is trust; it is not a review of every intranasal-oxytocin finding. Bibliographic details agent-verified 2026-09-25 against Crossref and Europe PMC (PMID 26581735), the quotations against the abstract. The article was not read., pp. 772–789. doi: 10.1177/1745691615600138.
- New York Civil Liberties Union (Dec. 2018). *ICE Rigged the System to Keep Immigrants Locked Up*. Advocacy source, cited for the date of the change: "In 2017, ICE manipulated the computer

- program to get rid of the release option. The program now automatically recommends that all immigrants be detained.” Agent-verified 2026-10-01. New York Civil Liberties Union. URL: <https://www.nyclu.org/commentary/ice-rigged-system-keep-immigrants-locked> (visited on 10/01/2026).
- New York Times Co. v. United States Department of Justice* (2014). Judgment. Freedom of Information Act litigation over documents on targeted killing by drone. The court ordered disclosure of the 2010 OLC–DOD Memorandum, a legal opinion prepared for the Department of Defense on the legality of such strikes, holding that prior official disclosures and the release of the White Paper waived the exemptions otherwise available. Nothing in the decision reaches the lawfulness of any strike. Agent-verified from the reported record; the opinion itself was not read. URL: <https://www.leagle.com/decision/infco20140626124> (visited on 09/17/2026).
- New Zealand Parliament (2017). *Te Awa Tupua (Whanganui River Claims Settlement) Act 2017*. Public Act 2017 No 7. URL: <https://www.legislation.govt.nz/act/public/2017/0007/latest/whole.html> (visited on 08/25/2026).
- Nichols, Shaun (2004). *Sentimental Rules: On the Natural Foundations of Moral Judgment*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/sentimental-rules-9780195314205>.
- Nisbett, Richard E. and Timothy D. Wilson (1977). “Telling More Than We Can Know: Verbal Reports on Mental Processes”. In: *Psychological Review* 84.3, pp. 231–259. doi: 10.1037/0033-295X.84.3.231.
- Noddings, Nel (1984). *Caring: A Feminine Approach to Ethics and Moral Education*. Cited for care ethics treating partiality toward those in relation as part of care rather than a failure of it. Noddings distinguishes the one-caring and the cared-for, and caring-for from caring-about; she wrote that “we are not obliged to care for starving children in Africa” (p. 86, as quoted by Cassidy, *Journal of International Women’s Studies* 7(1), 2005). Andre’s review (*Teaching Philosophy* 9(1), 1986, 89–90) summarizes: “Starving strangers, or future generations, have no ethical standing unless they are related by ‘circles’ or ‘chains’ to those for whom I care.” Open Library records the 1984 and 1986 printings as “Ethics & Moral Education,” while Andre’s review header has “and”; the title page was not seen. The 2013 edition is retitled *Caring: A Relational Approach to Ethics and Moral Education*. Bibliographic details agent-verified 2026-09-25 against Open Library (LCCN 83018223, 216 pp.) and Andre’s review header. The book itself was not read; its views are checked against the two secondary sources named. Berkeley: University of California Press. ISBN: 0-520-05043-6.
- Noem v. Abrego Garcia* (Apr. 10, 2025). No. 24A949 (U.S. Apr. 10, 2025).
- Noem v. Vasquez Perdomo* (Sept. 8, 2025). No. 25A169 (U.S. Sept. 8, 2025) (order granting stay; Kavanaugh, J., concurring). URL: https://www.supremecourt.gov/opinions/24pdf/25a169_5h25.pdf (visited on 10/01/2026).
- Noor, Masi et al. (2012). “When Suffering Begets Suffering: The Psychology of Competitive Victimhood Between Adversarial Groups in Violent Conflicts”. In: *Personality and Social Psychology Review* 16.4, pp. 351–374. doi: 10.1177/1088868312440048.
- NPR (Mar. 11, 2026). *Pentagon Probe Points to U.S. Missile Hitting Iranian School*. NPR. URL: <https://npr.org/2026/03/11/nx-s1-5744981/pentagon-iran-missile-school-hegseth> (visited on 09/23/2026).
- Nussbaum, Martha C. (2006). *Frontiers of Justice: Disability, Nationality, Species Membership*. The Tanner Lectures on Human Values. Cambridge, MA: Harvard University Press. ISBN: 978-0-674-01917-1.
- O’Brien, Kyle et al. (2025). *Deep Ignorance: Filtering Pretraining Data Builds Tamper-Resistant Safeguards into Open-Weight LLMs*. Cited for pretraining-stage safeguards resisting adversarial fine-tuning. Abstract: post-training methods “have struggled to make LLMs resistant to more than a few dozen steps of adversarial fine-tuning,” while models pretrained on filtered data

- “exhibit substantial resistance to adversarial fine-tuning attacks on up to 10,000 steps and 300M tokens of biothreat-related text.” This concerns knowledge left out of training, not a disposition trained in, and the book says so. Listed for ICLR 2026 by one index; the venue was not confirmed against ICLR. Authors, title and arXiv number agent-verified 2026-09-25. Only the abstract was read. arXiv: 2508.06601. URL: <https://arxiv.org/abs/2508.06601> (visited on 09/25/2026).
- Ochigame, Rodrigo (Dec. 2019). *The Invention of ‘Ethical AI’: How Big Tech Manipulates Academia to Avoid Regulation*. The Intercept. URL: <https://theintercept.com/2019/12/20/mit-ethical-ai-artificial-intelligence/> (visited on 08/25/2026).
- Office of Rep. Pramila Jayapal (Feb. 5, 2026a). *Markey, Jayapal, Merkley, Wyden Introduce Bill to Ban ICE and CBP Use of Facial Recognition Technology amid Trump’s Rapidly Growing Surveillance State*. ICE Out of Our Faces Act, S. 3779, 119th Cong. URL: <https://jayapal.house.gov/2026/02/05/markey-jayapal-merkley-wyden-introduce-bill-to-ban-ice-and-cbp-use-of-facial-recognition-technology-amid-trumps-rapidly-growing-surveillance-state/> (visited on 10/01/2026).
- (Sept. 25, 2026b). *Reps. Jayapal, Tlaib, Pressley and Sens. Markey, Merkley Reintroduce Legislation to Halt Government Use of Facial Recognition and Other Biometric Technology*. Facial Recognition and Biometric Technology Moratorium Act (reintroduced; bill numbers not stated in the release). URL: <https://jayapal.house.gov/2026/09/25/rep-jayapal-tlaib-pressley-and-sens-markey-merkley-reintroduce-legislation-to-halt-government-use-of-facial-recognition-and-other-biometric-technology/> (visited on 10/01/2026).
- Office of the Surgeon General, Multinational Force–Iraq and Office of the Surgeon General, United States Army Medical Command (Nov. 17, 2006). *Mental Health Advisory Team (MHAT) IV, Operation Iraqi Freedom 05-07: Final Report*. Cited for the survey’s battlefield-ethics findings: 47 percent of soldiers (38 percent of Marines) agreeing that noncombatants should be treated with dignity and respect, about 10 percent reporting mistreating noncombatants, and combat exposure (unit casualties, handling remains) associated with the reported mistreatment. Agent-checked 2026-09-28 against secondary accounts only (the NTIS listing, and the summary of the report in Warner et al.’s battlefield-ethics work as quoted in PMC4138704); the report itself was not read. Not found in those accounts: whether the dignity-and-respect attitude also tracked combat exposure, so the book ties only the mistreatment to it. URL: <https://ntrl.ntis.gov/NTRL/dashboard/searchResults/titleDetail/PB2010103335.xhtml> (visited on 09/28/2026).
- Office of the United Nations High Commissioner for Human Rights (Sept. 2026). *Iran: Civilians Trapped between Extreme State Repression and Violations Arising from Devastating Conflict with US, Israel, UN Fact-Finding Mission Finds*. Press release on the Independent International Fact-Finding Mission’s report (UNFFMIran2026). As read on 2026-10-01, the two war crimes are assigned one to each strike: “Following the United States and Israeli attacks on 28 February, the Mission found reasonable grounds to believe that strikes by the United States in Minab, Hormozgan Province, and in Lamerd, Fars Province, amounted to the war crimes of making a civilian object the object of an attack and launching an indiscriminate attack resulting in the loss of life or injury to civilians or damage to civilian objects, respectively.” The strikes: “Tomahawk missile strikes on the clearly identifiable Shajareh Tayyebbeh Primary School in Minab — which killed over 150 people, including approximately 120 girls and boys — and another airstrike where Precision Strike Missiles dispersed clouds of tungsten pellets over a clearly identifiable sports complex and residential area in Lamerd, killing 22 civilian men and women.” The page’s structured data gives a modification of 29 September 2026 and carries no correction notice. The text as first issued, preserved by a GlobalSecurity.org mirror and quoted by CNBC and RFE/RL, read only that “the United States committed the war crime of launching indiscriminate attacks resulting in the loss of life or injury to civilians or damage to civilian objects”; the revised wording agrees with the report’s own paragraphs as

- beaumont2026un records them (para. 126 Minab, paras. 127–128 Lamerd). Agent-verified 2026-10-01. OHCHR. URL: <https://www.ohchr.org/en/press-releases/2026/09/iran-civilians-trapped-between-extreme-state-repression-and-violations> (visited on 10/01/2026).
- Oil Platforms (Islamic Republic of Iran v. United States of America)*, Judgment (Nov. 6, 2003). General List No. 90. Necessity and proportionality: paras. 73–77. URL: <https://www.icj-cij.org/case/90> (visited on 09/30/2026).
- Oliner, Samuel P. and Pearl M. Oliner (1988). *The Altruistic Personality: Rescuers of Jews in Nazi Europe*. New York: Free Press.
- Open Letter to the European Commission (2018). *Open Letter to the European Commission: Artificial Intelligence and Robotics*. robotics-openletter.eu. Signed by more than 150 AI and robotics researchers, legal scholars, and ethicists from 14 countries, including Noel Sharkey, Raja Chatila, Alan Winfield, and Nathalie Nevejans. URL: <https://www.robotics-openletter.eu/> (visited on 08/25/2026).
- Open Society Justice Initiative (2021). *French Criminal Investigation of Chemical Weapons Attacks in Syria*. Case page of a civil party to the proceedings, not a court. “On March 1, 2021, the Open Society Justice Initiative joined Syrian survivors of chemical weapons attacks, the Syrian Center for Media and Freedom of Expression (SCM), and Syrian Archive in seeking a French criminal investigation for the August 2013 chemical weapons attacks on the city of Douma and on Eastern Ghouta.” It also records the German complaint (October 2020) and the Swedish complaint (April 2021), both on Ghouta 2013 and Khan Shaykhun 2017. None concerns an IIT-attributed incident. Agent-verified 2026-10-01. URL: <https://www.justiceinitiative.org/litigation/french-criminal-investigation-of-chemical-weapons-attacks-in-syria> (visited on 10/01/2026).
- OpenAI (Apr. 2025). *OpenAI Model Spec*. Dedicated to the public domain under CC0 1.0. States a chain of command in four levels of descending authority — platform, developer, user, guideline — where “instructions with higher authority override those with lower authority,” platform rules “cannot be overridden by developers or users,” and guidelines “can be implicitly overridden.” A living document: the version consulted here is that of 11 April 2025, which the site marks as superseded. URL: <https://model-spec.openai.com/2025-04-11.html> (visited on 09/06/2026).
- Organisation for the Prohibition of Chemical Weapons (June 27, 2018). *CWC Conference of the States Parties Adopts Decision Addressing the Threat from Chemical Weapons Use*. News release on decision C-SS-4/DEC.3, which established the Investigation and Identification Team. “The decision brought forward by the United Kingdom of Great Britain and Northern Ireland and supported by 30 countries ... was adopted with 106 States Parties present and voting; 82 States Parties voted yes and 24 States Parties voted no.” It calls upon the Secretariat “to identify the perpetrators of the use of chemical weapons in the Syrian Arab Republic by identifying and reporting on all information potentially relevant to the origin of those chemical weapons in those instances in which the OPCW Fact-Finding Mission determines or has determined that use or likely use occurred, and cases for which the OPCW-UN Joint Investigative Mechanism has not issued a report.” The release does not name the States that voted no. The decision itself (a PDF) was not read. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2018/06/cwc-conference-states-parties-adopts-decision-addressing-threat-chemical> (visited on 10/01/2026).
- (Apr. 8, 2020a). *OPCW Releases First Report by Investigation and Identification Team*. Report S/1867/2020, on Ltamenah, 24, 25 and 30 March 2017, covering “investigations conducted between June 2019 and March 2020.” The IIT Coordinator, as quoted: “[t]he IIT has concluded that there are reasonable grounds to believe that the perpetrators of the use of sarin as a chemical weapon in Ltamenah on 24 and 30 March 2017, and the use of chlorine as a chemical weapon on 25 March 2017 were individuals belonging to the Syrian Arab Air Force.” The

- Director-General, as quoted: “[t]he IIT is not a judicial or quasi-judicial body with the authority to assign individual criminal responsibility.” The report itself (a PDF) was not read. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2020/04/opcw-releases-first-report-investigation-and-identification-team> (visited on 10/01/2026).
- Organisation for the Prohibition of Chemical Weapons (Feb. 6, 2020b). *OPCW: Independent Investigation into Possible Breaches of Confidentiality Report Released*. The OPCW’s account of its investigation into the 2019 Douma leaks; it concerns the Fact-Finding Mission, not the IIT. “The investigation took place between July 2019 and February 2020 ... The investigation included interviews with 29 witnesses, documents, electronic records, audio-recordings, and forensic analyses.” It found “that two former OPCW officials violated their obligations concerning the protection of confidential information related to the FFM Douma investigation.” On Inspector A: “Although Inspector A’s assessment purports to be an official OPCW FFM report on the Douma incident, it is instead a personal document created with incomplete information and without authorisation.” The Director-General: “Inspectors A and B are not whistle-blowers. They are individuals who could not accept that their views were not backed by evidence.” “I stand by the conclusions of the final Douma report.” The investigation report itself (a PDF) was not read. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2020/02/opcw-independent-investigation-possible-breaches-confidentiality-report> (visited on 10/01/2026).
- (Jan. 27, 2023). *OPCW Releases Third Report by Investigation and Identification Team*. Report S/2125/2023, on Douma, 7 April 2018. “The OPCW Investigation and Identification Team (IIT)’s third report concludes that there are reasonable grounds to believe that the Syrian Arab Air Forces were the perpetrators of the chemical weapons attack on 7 April 2018 in Douma, Syrian Arab Republic.” “This includes 70 environmental and biomedical samples, 66 witness statements, and other verified data, such as forensic analysis, satellite images, gas dispersion modelling, and trajectory simulations.” “The IIT considered a range of possible scenarios and tested their validity against the evidence they gathered and analysed to reach their conclusion.” “The conclusion of the report is reached on the basis of ‘reasonable grounds’, which is the standard of proof consistently adopted by international fact-finding bodies and commissions of inquiry.” The report itself (a PDF) was not read, so its treatment of the manual-placement hypothesis is not checked. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2023/01/opcw-releases-third-report-investigation-and-identification-team> (visited on 10/01/2026).
- (Feb. 22, 2024). *OPCW Identifies ISIL as Perpetrators of 2015 Chemical Attack in Marea, Syria*. Report S/2255/2024, on Marea, 1 September 2015; the IIT’s only attribution to a non-State actor. It concludes “that there are reasonable grounds to believe that units of the Islamic State in Iraq and the Levant (ISIL) were the perpetrators of the chemical weapons attack on 1 September 2015 in Marea, Syria.” “No other entity possessed the means, motives, and capabilities to deploy sulfur mustard as part of an attack in Marea on 1 September, 2015.” The Director-General, as quoted: “the OPCW Secretariat has demonstrated that it is able to identify perpetrators of the use of chemical weapons, States and non-State actors alike.” The report itself (a PDF) was not read. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2024/02/opcw-identifies-isil-perpetrators-2015-chemical-attack-marea-syria> (visited on 10/01/2026).
- (2026a). *Fact-Finding Mission*. Standing page, undated; read as it stood on the access date. “The OPCW Fact-Finding Mission (FFM) is responsible for determining whether toxic chemicals have been used as weapons in the Syrian Arab Republic. The mandate does not include identifying who is responsible for any alleged attacks.” “On the basis of the FFM’s findings, the OPCW Investigation and Identification Team (IIT) (as the OPCW-UN Joint Investigative Mechanism (JIM) did previously), collects and analyses evidence that may help

- to identify the perpetrators of chemical weapons attacks in Syria.” Agent-verified 2026-10-01. URL: <https://www.opcw.org/fact-finding-mission> (visited on 10/01/2026).
- Organisation for the Prohibition of Chemical Weapons (2026b). *Investigation and Identification Team (IIT)*. Standing page, undated; read as it stood on the access date. “The IIT is tasked with fact finding and is not a prosecutorial or judicial entity. It is not responsible to determine criminal responsibility of individuals, organisations or States. Furthermore, the IIT does not make recommendations for future action.” “The IIT standard of proof is ‘reasonable grounds’ to suspect that an individual or entity was directly or indirectly involved in the use of a chemical weapon. This is the standard applied by international fact-finding missions and commissions of enquiry to reach conclusions. The information gathered and analysed is considered holistically.” It lists the Syrian armed forces as perpetrators in one case in Kafr Zeita (October 2016), three in Ltamenah (March 2017), one in Saraqib (February 2018) and one in Douma (April 2018). It gives the report dates as 8 April 2020, 12 April 2021, 27 January 2023, 22 February 2024 and 21 January 2026. The page’s statement that the JIM “was established in April 2015” was not checked against resolution 2235 (2015) and should not be relied on. Agent-verified 2026-10-01. URL: <https://www.opcw.org/iit> (visited on 10/01/2026).
- (Jan. 22, 2026c). *OPCW Releases 5th Report to Identify Perpetrators of Chemical Weapons Use in Syria*. Report S/2478/2026 (dated 21 January 2026), on Kafr Zeita, 1 October 2016. The report concludes that “there are reasonable grounds to believe that the Syrian Arab Air Forces were the perpetrators of the chemical weapons attack on 1 October 2016 in Kafr Zeita.” Subheading: “For the first time, Syria cooperates with an OPCW Investigation and Identification Team’s investigation and provides full access to the country, information, and documentation relevant to the incident.” “The Secretariat also thoroughly considered the positions put forward by the former Syrian authorities but was unable to obtain any concrete information supporting those positions. To date, five reports identifying perpetrators have been issued by the OPCW on seven cases of use of toxic chemicals as weapons in Syria.” The report itself (a PDF) was not read. Agent-verified 2026-10-01. URL: <https://www.opcw.org/media-centre/news/2026/01/opcw-releases-5th-report-identify-perpetrators-chemical-weapons-use-syria> (visited on 10/01/2026).
- Organization, Jurisdiction, and Powers of the Board of Immigration Appeals* (2026). 8 C.F.R. §1003.1. Para. (a)(1) (Board members as the Attorney General’s delegates); para. (h) (referral of cases to the Attorney General). URL: <https://www.ecfr.gov/current/title-8/chapter-V/subchapter-A/part-1003/subpart-A/section-1003.1> (visited on 10/01/2026).
- Ostrom, Elinor (1990). *Governing the Commons. The Evolution of Institutions for Collective Action*. The Political Economy of Institutions and Decisions. Cited for the design principles of long-enduring common-pool resource institutions, Table 3.1, p. 90, in chapter 3 (pp. 58–102). Her numbering: (1) clearly defined boundaries, “Individuals or households who have rights to withdraw resource units from the CPR must be clearly defined, as must the boundaries of the CPR itself”; (2) congruence between appropriation and provision rules and local conditions; (3) collective-choice arrangements, “most individuals affected by the operational rules can participate in modifying the operational rules”; (4) monitoring, by monitors who “are accountable to the appropriators or are the appropriators”; (5) graduated sanctions; (6) conflict-resolution mechanisms, low-cost local arenas; (7) minimal recognition of rights to organize, by external government; and, “For CPRs that are parts of larger systems,” (8) nested enterprises, “governance activities are organized in multiple layers of nested enterprises.” The quotations of principles 3 and 8 are as Cox, Arnold and Villamayor Tomás, *Ecology and Society* 15(4): 38 (2010), give them with the page reference 1990:90; that of principle 1 and the wording of the others are from a transcription of Table 3.1 on publiccommons.ca, which has small copying errors elsewhere. The numbering and order agent-verified 2026-09-26 against Ostrom’s own restatement in her Nobel lecture (nobelprize.org, 2009), which keeps

- principles 1 to 8 in the same order, splitting 1, 2 and 4 into A and B. Bibliographic details and DOI agent-verified against Crossref. Page 90 itself was not read; Google Books lists it as viewable but its text could not be retrieved, so the wording of the table is at second hand. Cambridge: Cambridge University Press. ISBN: 978-0-521-40599-7. DOI: 10.1017/CB09780511807763.
- Ouyang, Long et al. (2022). *Training Language Models to Follow Instructions with Human Feedback*. Twenty authors. Human raters preferred the instruction-tuned 1.3-billion-parameter model to the 175-billion-parameter model it was built from. arXiv: 2203.02155. URL: <https://arxiv.org/abs/2203.02155> (visited on 08/28/2026).
- Oversight Board (July 8, 2021). *Öcalan's Isolation*. Case Decision 2021-006-IG-UA. URL: <https://www.oversightboard.com/decision/IG-I9DP23IB/> (visited on 09/23/2026).
- Owen, Adrian M. et al. (2006). "Detecting Awareness in the Vegetative State". In: *Science* 313.5792. Cited in section 3.4 as the one instance of a probe settling a status question that had been treated as unsettled. A patient meeting the criteria for a diagnosis of vegetative state was asked to imagine playing tennis and to imagine moving around her home, and produced the cortical activation each instruction produces in healthy volunteers, sustained and on cue. The report is a one-page Brevia covering this single patient; the larger series came later and is not cited here. The text follows the paper and says a patient, not patients. Cited for the paradigm and its result, not for a prevalence figure., p. 1402. DOI: 10.1126/science.1130197. URL: <https://www.science.org/doi/10.1126/science.1130197> (visited on 09/14/2026).
- Palomera, Gisselle (Apr. 15, 2025). *Discredited Former Cop Played 'Key Role' in Deportation of Gay Make-Up Artist*. Reports, citing USA Today, that a private-prison employee who was a former police officer flagged the tattoos, and that the allegation rested on the tattoos and no other information. Los Angeles Blade. URL: <https://www.losangelesblade.com/2025/04/15/latest-update-in-the-case-of-wrongfully-deported-gay-venezuelan-make-up-artist/> (visited on 10/01/2026).
- Panksepp, Jaak (1998). *Affective Neuroscience: The Foundations of Human and Animal Emotions*. New York: Oxford University Press. URL: <https://global.oup.com/academic/product/affective-neuroscience-9780195178050>.
- Parasuraman, Raja and Dietrich H. Manzey (June 2010). "Complacency and Bias in Human Use of Automation: An Attentional Integration". In: *Human Factors* 52.3. Cited for people given an automated aid's recommendation checking it less and deferring to it more when their attention is loaded. The abstract (p. 381): "Automation complacency occurs under conditions of multiple-task load, when manual tasks compete with the automated task for the operator's attention"; "Automation bias results in making both omission and commission errors when decision aids are imperfect"; "it occurs in both naive and expert participants, cannot be prevented by training or instructions, and can affect decision making in individuals as well as in teams." On checking less (p. 398): "the user of a decision aid may trust it to the extent that he or she directly follows the automatically generated advice without making the attentionally demanding effort of cross-checking its validity against other available and accessible information," and "neglect of automation verification might constitute a major source of commission errors." On the condition (pp. 403–404): "complacency and automation bias effects emerge primarily when task load is high (e.g., multitask demands)," and they are "a human performance cost of certain automation system designs, that is, systems characterized by high reliability and high levels of automation, particularly in high task load situations" (p. 403). The review's moderator is workload and competing tasks. Time pressure appears in it only as a feature of particular studies' tasks (pp. 386, 393, 395), never as a factor the review itself identifies, so this entry does not support the words "under time pressure." Bibliographic details agent-verified 2026-09-26 against Crossref; abstract and quotations against the published PDF deposited by TU Berlin in DepositOnce (doi:10.14279/depositonce-8052), 30 pages, with page numbers taken from its journal pagination. The SAGE landing

- page was not read. The primary studies the review summarizes were not read., pp. 381–410. DOI: 10.1177/0018720810376055. URL: <https://doi.org/10.1177/0018720810376055> (visited on 09/26/2026).
- Parr, Thomas, Giovanni Pezzulo, and Karl J. Friston (Mar. 29, 2022). *Active Inference. The Free Energy Principle in Mind, Brain, and Behavior*. Open access under CC BY-NC-ND 4.0, per the licence recorded with the DOI; ebook ISBN 978-0-262-36997-8. Title, subtitle, authors, date, ebook ISBN and licence agent-verified against Crossref; the hardcover ISBN is from a Crossref-indexed review header (Cavanna and Seri, *Cognitive Neuropsychiatry* 28(2), 2023) because the publisher's page refused automated access. Cited for the active-inference agent's prior preferences over its own observations and its selection of actions by expected free energy; no page of the book was read, so that description is not checked against a locator. Agent-verified 2026-09-23. Cambridge, MA: The MIT Press. ISBN: 978-0-262-04535-3. DOI: 10.7551/mitpress/12441.001.0001. URL: <https://doi.org/10.7551/mitpress/12441.001.0001> (visited on 09/23/2026).
- Partial Award: Jus Ad Bellum, Ethiopia's Claims 1–8* (Dec. 19, 2005). Permanent Court of Arbitration Case 71. Cited for the holding that a territorial claim does not justify force against territory another State administers. Para. 10 as quoted identically by two secondary HTML sources (Just Security, 10 November 2020; Opinio Juris, 21 March 2024): “the practice of States and the writings of eminent publicists show that self-defense cannot be invoked to settle territorial disputes” and “border disputes between States are so frequent that any exception to the prohibition of the threat or use of force for territory that is allegedly occupied unlawfully would create a large and dangerous hole in a fundamental rule of international law.” The finding that Eritrea violated Article 2(4) by attacking Badme, then under peaceful administration by Ethiopia, is confirmed only in paraphrase (Opinio Juris, 13 November 2024: “under ‘peaceful administration’ by Ethiopia at the time (paras 10-16)”; the exact wording of the holding was not read, because the award is published only as a PDF. Only secondary sources were read. Agent-verified 2026-10-01. URL: <https://pca-cpa.org/en/cases/71/> (visited on 10/01/2026).
- Pathak, Deepak et al. (2017). “Curiosity-Driven Exploration by Self-Supervised Prediction”. In: *Proceedings of the 34th International Conference on Machine Learning (ICML)*. Vol. 70, pp. 2778–2787. URL: <https://proceedings.mlr.press/v70/pathak17a.html> (visited on 08/25/2026).
- Paul, Sanjukta (2020). “Antitrust as Allocator of Coordination Rights”. In: *UCLA Law Review* 67.2. Cited for antitrust law exempting coordination inside the conventionally organized firm while treating coordination among independent parties as suspect, and for the allocation of coordination rights being a legal choice that can be made in the public interest. Author, title, volume and first page agent-verified 2026-09-29 against the UCLA Law Review and SSRN listings. The article itself was not read., p. 378. URL: <https://www.uclalawreview.org/antitrust-as-allocator-of-coordination-rights/> (visited on 09/29/2026).
- Paul, Sanjukta and Nathan Tankus (2019). “The Firm Exemption and the Hierarchy of Finance in the Gig Economy”. In: *University of St. Thomas Law Journal* 16.1. Cited for a firm formed and controlled by independent workers having difficulty qualifying for antitrust's firm exemption, and for the financing available to such a firm reproducing the disadvantage. Authors, title, journal, volume and year agent-verified 2026-09-29 against the SSRN and University of Michigan Law School listings; the first page is as the author gave it and was not checked. The article itself was not read., p. 44. URL: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3426230 (visited on 09/29/2026).
- Pawlyk, Oriana (Apr. 29, 2016). *Pilots, Crew of AC-130 Gunship Acknowledged Confusion before Striking Kunduz Hospital*. Cited for the AC-130U crew questioning the target before the strike on the Médecins Sans Frontières trauma centre in Kunduz on 3 October 2015. The article quotes the crew's intercom exchanges from the U.S. Central Command investigation: the sensor operator said the grid placed his sensor on a different compound, and

- the navigator did not pass the concern on. Agent-read 2026-09-28. The doubts were spoken on the aircraft intercom, not over the radio to ground forces. Air Force Times. URL: <https://www.airforcetimes.com/news/your-air-force/2016/04/29/pilots-crew-of-ac-130-gunship-acknowledged-confusion-before-striking-kunduz-hospital/> (visited on 09/28/2026).
- Pereboom, Derk (2001). *Living without Free Will*. Cambridge Studies in Philosophy. Cambridge: Cambridge University Press. ISBN: 9780521791984. DOI: 10.1017/CB09780511498824.
- Pereira v. Sessions* (2018). 585 U.S. 198 (2018). Sotomayor, J., for the Court, decided 21 June 2018. Cited for the notice actually given: DHS “almost always serves noncitizens with notices that fail to specify the time, place, or date of initial removal hearings whenever the agency deems it impracticable to include such information,” citing the “Government’s admission that ‘almost 100 percent’ of ‘notices to appear omit the time and date of the proceeding over the last three years’”; and for the holding that such a document “is not a ‘notice to appear under section 1229(a).’” The reporter citation 585 U.S. 198 was not shown in the text read and is not agent-verified. Agent-verified 2026-10-01 from the Cornell LII text. URL: <https://www.law.cornell.edu/supremecourt/text/17-459> (visited on 10/01/2026).
- Permanent Representative of the Russian Federation to the United Nations (Feb. 24, 2022). *Letter dated 24 February 2022 addressed to the Secretary-General*. UN Doc. S/2022/154. Cited for Russia claiming the invasion of Ukraine as self-defense under Article 51 of the Charter, which concedes that its forces were acting on Ukraine’s territory. The letter transmits President Putin’s address of 24 February 2022 announcing the “special military operation” and describes it as a measure taken under Article 51. Document symbol, date and description agent-verified 2026-09-29 against the UN Digital Library record and secondary legal commentary; the PDF was not read. URL: <https://digitallibrary.un.org/record/3959647> (visited on 09/29/2026).
- Permanent Representative of the United States of America to the United Nations (Mar. 10, 2026). *Letter dated 10 March 2026 addressed to the President of the Security Council*. UN Doc. S/2026/161. Cited for the jus ad bellum claim made for the strikes on Iran begun 28 February 2026: self-defense under Article 51 against Iran’s “attacks and ongoing threats”, as “the latest stage in the ongoing international armed conflict” addressed in the letter of 27 June 2025, and in collective self-defense of Israel; the letter names no armed attack under way and says that “any analysis of the imminence, gravity and scope of the threat” must account for the regime’s conduct since 1979. Signed by Michael G. Waltz; distributed 11 March 2026. Agent-read 2026-10-01 from the copy of the document posted by Just Security (<https://www.justsecurity.org/wp-content/uploads/2026/03/US-Article-51-letter-Iran-11-March-2026.pdf>); the docs.un.org address returned HTTP 200 and its content was not read. URL: <https://docs.un.org/en/S/2026/161> (visited on 10/01/2026).
- Perrigo, Billy (Jan. 2023). *Exclusive: OpenAI Used Kenyan Workers on Less Than \$2 Per Hour to Make ChatGPT Less Toxic*. Reports on OpenAI’s contract with Sama; Nairobi-based data labelers read and categorized descriptions of child sexual abuse, torture, bestiality, murder, and suicide for take-home wages of roughly \$1.32–\$2 per hour, described lasting psychological injury, and found counseling rare and inadequate; Sama ended the contract in February 2022, eight months early. TIME. URL: <https://time.com/6247678/openai-chatgpt-kenya-workers/> (visited on 08/25/2026).
- Pessoa, Luiz (2022). *The Entangled Brain: How Perception, Cognition, and Emotion Are Woven Together*. Cited for the claim that brain regions do not each do one thing and that emotion and cognition do not come apart in the brain. Page 30: “the central thesis of the book: Brain areas don’t compute specific functions—they are not segregated “organs of the mind,” as Brodmann put it. The brain is not a modular system that can be understood a region at a time.” Page 57: “if brain regions participate in many functions and functions can be carried out by many regions, the ensuing structure-function mapping will be many-to-many.” Page

- 14, on the title: “not something like threads that are mixed together but can be separated if only one has enough time and patience.” Page 191: “emotion—insofar as it is meaningful to speak of “emotion”—like every other mental domain, is a large-scale network property of the nervous system.” Open access under CC BY-NC-ND 4.0; published 15 November 2022; PDF ISBN 9780262372107, EPUB 9780262372114; LCCN 2021061878. Text read in the Internet Archive’s full-text rendering of the open-access edition ([mit_press_book_9780262372107](https://mitpress.mit.edu/books/9780262372107)); page numbers are from its running heads; the publisher’s pages refused access. Agent-verified 2026-10-01. Cambridge, MA: MIT Press. ISBN: 9780262544603. DOI: 10.7551/mitpress/14636.001.0001. URL: <https://doi.org/10.7551/mitpress/14636.001.0001> (visited on 10/01/2026).
- Peterson, Becky (Dec. 2019). *Palantir Grabbed Project Maven Defense Contract after Google Left the Program: Sources*. Source for Palantir’s reported role: “Secretive data analytics company Palantir is working on Project Maven, the US Department of Defense program that Google dropped in March following protests inside of the company, according to people familiar with the project.” Anonymously sourced; Palantir declined to comment. The page’s publication timestamp is 10 December 2019; secondary reports dated 11 December relay it. Also reports that Anduril “had won its own contract to work on AI-powered virtual reality technology for Project Maven.” Agent-verified 2026-10-01. Business Insider. URL: <https://www.businessinsider.com/palantir-took-over-from-google-on-project-maven-2019-12> (visited on 10/01/2026).
- Pichai, Sundar (2018). *AI at Google: Our Principles*. Published June 7, 2018, following weeks of employee and public backlash over Project Maven, a Pentagon drone-footage-analysis contract; roughly 4,000 employees signed a letter opposing it and about a dozen resigned. URL: <https://blog.google/technology/ai/ai-principles/> (visited on 08/25/2026).
- Pope, Dudley (1972). *The Great Gamble: Nelson at Copenhagen*. London: Weidenfeld and Nicolson. ISBN: 0297004840.
- Poulson, Jack (May 2023). *Pentagon Certified Palantir as Only Supplier for Artificial Intelligence Targeting Tool Known as “Maven Smart System”*. Source for the sole-source notice: “Palantir was named yesterday as the only suitable supplier for the “Maven Smart System”. The \$13 million contract is follow-on to Palantir’s previous work on Pentagon AI surveillance, Project Maven.” And: “The notice further revealed the name of Palantir’s AI targeting tool as the “Maven Smart System”. The name was not new in 2023; see [strout2020convergence](https://www.strout2020convergence.com/). Also the source for the other contractors: “Amazon and Microsoft received at least \$50 million for their work on the project,” and “It is not clear from public records as to what roles Amazon and Microsoft have continued to play on Maven, and so the framing of Palantir and Anduril having ‘taken over’ the project from Google is questionable.” Also states: “Project Maven transferred to the NGA in January.” The sole-source notice itself was not read. Agent-verified 2026-10-01. Tech Inquiry (Substack). URL: <https://jackpoulson.substack.com/p/pentagon-certified-palantir-as-only> (visited on 10/01/2026).
- Preston, Stephanie D. (2013). “The Origins of Altruism in Offspring Care”. In: *Psychological Bulletin* 139.6. Agent-verified against Crossref 2026-09-18: author, title, journal, 139(6), 1305–1341, 2013., pp. 1305–1341. DOI: 10.1037/a0031755.
- Proclamation 10903: Invocation of the Alien Enemies Act Regarding the Invasion of the United States by Tren de Aragua* (Mar. 14, 2025). Signed 14 March 2025.
- Protocol Additional to the Geneva Conventions of 12 August 1949, and relating to the Protection of Victims of International Armed Conflicts (Protocol I)* (1977). 1125 U.N.T.S. 3, adopted 8 June 1977, entered into force 7 December 1978.
- Qi, Xiangyu et al. (2024). “Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!” In: *International Conference on Learning Representations (ICLR)*. Jailbroke GPT-3.5 Turbo’s safety alignment using ten adversarially chosen fine-tuning examples for under \$0.20 via OpenAI’s fine-tuning API; also shows benign fine-tuning can

- inadvertently degrade safety. arXiv: 2310.03693. URL: <https://arxiv.org/abs/2310.03693> (visited on 08/25/2026).
- Radio Free Europe/Radio Liberty (Sept. 2026). *White House Rejects UN Report Alleging US War Crimes in Iran*. Read from the GlobalSecurity.org reprint, which gives its source as the RFE/RL live blog at <https://www.rferl.org/a/iran-live-blog-trump-middle-east/33803414.html>; the live blog loads its entries by script and the item was not reached there. Source for the White House response: “In a scathing written response to the report, White House Deputy Press Secretary Anna Kelly rejected the findings, saying: “The only party in this conflict that has committed war crimes is the Iranian regime, which has murdered thousands of its own people, killed Americans overseas, executed acts of terrorism in the Strait of Hormuz, and launched indiscriminate attacks against its neighbors in the region.”” And: “Kelly also said the UN Human Rights Council had “accomplished nothing for human rights while spouting unserious nonsense for decades.”” Its summary of the finding repeats the press release’s first wording (both strikes as indiscriminate attacks), which OHCHR has since revised; see [ohchr2026minabpress](https://www.globalsecurity.org/military/library/news/2026/09/mil-260918-rferl04.htm). Agent-verified 2026-10-01. Radio Free Europe/Radio Liberty. URL: <https://www.globalsecurity.org/military/library/news/2026/09/mil-260918-rferl04.htm> (visited on 10/01/2026).
- Rafailov, Rafael et al. (2023). *Direct Preference Optimization: Your Language Model is Secretly a Reward Model*. Reaches the standard preference-tuning objective with a classification loss over the preference pairs, dispensing with the separate reward-model fitting step. arXiv: 2305.18290. URL: <https://arxiv.org/abs/2305.18290> (visited on 08/28/2026).
- Railton, Peter (1984). “Alienation, Consequentialism, and the Demands of Morality”. In: *Philosophy & Public Affairs* 13.2, pp. 134–171.
- (1986a). “Facts and Values”. In: *Philosophical Topics* 14.2, pp. 5–31. URL: https://www.pdcnet.org/philtopics/content/philtopics_1986_0014_0002_0005_0031.
- (1986b). “Moral Realism”. In: *The Philosophical Review* 95.2, pp. 163–207. DOI: 10.2307/2185589.
- (2014a). “Reliance, Trust, and Belief”. In: *Inquiry* 57.1, pp. 122–150. DOI: 10.1080/0020174X.2014.858419.
- (2014b). “The Affective Dog and Its Rational Tale: Intuition and Attunement”. In: *Ethics* 124.4, pp. 813–859. DOI: 10.1086/675876.
- (2017). “At the Core of Our Capacity to Act for a Reason: The Affective System and Evaluative Model-Based Learning and Control”. In: *Emotion Review* 9.4, pp. 335–342. DOI: 10.1177/1754073916670021.
- (2020). “Ethical Learning, Natural and Artificial”. In: *Ethics of Artificial Intelligence*. Ed. by S. Matthew Liao. New York: Oxford University Press, pp. 45–78. DOI: 10.1093/oso/9780190905033.003.0002.
- Raja, Srinivasa N. et al. (2020). “The Revised International Association for the Study of Pain Definition of Pain: Concepts, Challenges, and Compromises”. In: *PAIN* 161.9, pp. 1976–1982. DOI: 10.1097/j.pain.0000000000001939. (Visited on 09/17/2026). The 2020 revision and its six accompanying notes, including that pain and nociception are different phenomena, that pain cannot be inferred solely from activity in sensory neurons, and that inability to communicate does not negate the possibility of pain in a human or nonhuman animal.
- Ramirez, Delia C. (Jan. 21, 2026). *To End Detention and Electronic Monitoring, and Redirect Funding to Community-Based Wrap-Around Services [called the Melt ICE Act by its sponsor]*. H.R. 7190, 119th Cong., 2d Sess. URL: <https://www.govinfo.gov/content/pkg/BILLS-119hr7190ih/html/BILLS-119hr7190ih.htm> (visited on 10/01/2026).
- Reichskulturkammergesetz vom 22. September 1933* (1933). RGBl. I S. 661. URL: <https://www.verfassungen.de/de33-45/kulturkammer33.htm> (visited on 10/01/2026).
- Renteria, Nelson (Apr. 11, 2025). *Inside Trump’s \$6mn Deportee Deal with El Salvador Mega-Prison*. Context (Thomson Reuters Foundation). URL: <https://www.context.news/money-power->

- people/inside-trumps-6mn-deportee-deal-with-el-salvador-mega-prison (visited on 10/01/2026).
- Reorganization Plan Modification for the Department of Homeland Security* (Jan. 30, 2003). H.R. Doc. No. 108-32, 108th Cong., 1st Sess. Effective March 1, 2003; renamed the Bureau of Border Security the Bureau of Immigration and Customs Enforcement. Set out as a note under 6 U.S.C. §542. URL: <https://www.law.cornell.edu/uscode/text/6/542> (visited on 10/01/2026).
- Resolution 2334 (2016)* (Dec. 23, 2016). Adopted 14–0–1 (United States abstaining). Reaffirms that Israeli settlements in the Palestinian territory occupied since 1967, including East Jerusalem, have no legal validity. United Nations Security Council. URL: [https://undocs.org/S/RES/2334\(2016\)](https://undocs.org/S/RES/2334(2016)) (visited on 09/30/2026).
- Resolution 541 (1983)* (Nov. 18, 1983). Declares the purported secession of part of Cyprus legally invalid and calls for its withdrawal. United Nations Security Council. URL: [https://undocs.org/S/RES/541\(1983\)](https://undocs.org/S/RES/541(1983)) (visited on 09/30/2026).
- Resolution 550 (1984)* (May 11, 1984). Reiterates resolution 541; calls on all States not to recognize the purported state and to respect the sovereignty and territorial integrity of Cyprus. United Nations Security Council. URL: [https://undocs.org/S/RES/550\(1984\)](https://undocs.org/S/RES/550(1984)) (visited on 09/30/2026).
- Resolutions 822 (1993), 853 (1993), 874 (1993) and 884 (1993)* (1993). Adopted 30 April, 29 July, 14 October and 12 November 1993. Each demands the withdrawal of occupying forces from districts of Azerbaijan surrounding Nagorno-Karabakh (Kelbajar, Agdam, Fizuli, Jabrayil, Qubadli, Zangelan) and refers to “the Nagorny Karabakh region of the Azerbaijani Republic”; none demands withdrawal from Nagorno-Karabakh itself. United Nations Security Council. URL: [https://undocs.org/S/RES/884\(1993\)](https://undocs.org/S/RES/884(1993)) (visited on 09/30/2026).
- Restoration of the Lawful Rights of the People’s Republic of China in the United Nations* (Oct. 25, 1971). 1976th plenary meeting. “Decides to restore all its rights to the People’s Republic of China and to recognize the representatives of its Government as the only legitimate representatives of China to the United Nations, and to expel forthwith the representatives of Chiang Kai-shek from the place which they unlawfully occupy at the United Nations and in all the organizations related to it.” The text does not mention Taiwan. Adopted by roll-call vote 76–35–17 (Wikipedia, citing A/PV.1976; secondary). Text read on Wikisource. Agent-verified 2026-10-01. United Nations General Assembly. URL: [https://undocs.org/A/RES/2758\(XXVI\)](https://undocs.org/A/RES/2758(XXVI)) (visited on 10/01/2026).
- Roth, Lorna (2009). “Looking at Shirley, the Ultimate Norm: Colour Balance, Image Technologies, and Cognitive Equity”. In: *Canadian Journal of Communication* 34.1. Cited for the Shirley card: “a norm reference card showing a ‘Caucasian’ woman wearing a colourful, high-contrast dress is used as a basis for measuring and calibrating the skin tones on the photograph being printed,” the model’s light skin being “the recognized skin ideal standard for most North American analogue photo labs”; and for “a light-skin bias embedded in colour film stock emulsions,” the article stating that “film chemistry, photo lab procedures, video screen colour balancing practices, and digital cameras in general were originally developed with a global assumption of ‘Whiteness.’” It does not say the chemistry was tuned to the card model’s own skin. Bibliographic details agent-verified against Crossref; the text was read in the journal’s HTML galley as archived by the Internet Archive, which carries no page numbers, so no page is given. Agent-verified 2026-09-23., pp. 111–136. doi: 10.22230/cjc.2009v34n1a2196. URL: <https://doi.org/10.22230/cjc.2009v34n1a2196> (visited on 09/23/2026).
- Roth-Hanania, Ronit, Maayan Davidov, and Carolyn Zahn-Waxler (2011). “Empathy Development from 8 to 16 Months: Early Signs of Concern for Others”. In: *Infant Behavior and Development* 34.3. Agent-verified against Crossref 2026-09-18: authors, title, journal, 34(3), 447–458, 2011. The DOI was not in the supplied entry and is added here., pp. 447–458. doi: 10.1016/j.infbeh.2011.04.007.

- Royal Statistical Society (2020). *Letter Correspondence between the Royal Statistical Society and Ofqual on the 2020 Grading Model Technical Advisory Group*. The RSS offered to nominate two independent Fellows to Ofqual's technical advisory group ahead of 2020 A-level results; Ofqual's condition was a five-year non-disclosure agreement barring public comment on the model, and the Society's nominees declined under those terms. Royal Statistical Society; UK Department for Education. URL: https://assets.publishing.service.gov.uk/government/uploads/system/uploads/attachment_data/file/911335/20200821_RT_letter_to_RSS.pdf (visited on 08/25/2026).
- Ruffman, Ted, Ben Lorimer, and Damian Scarf (2017). "Do Infants Really Experience Emotional Contagion?" In: *Child Development Perspectives* 11.4. Cited for the dissent: the abstract argues that the newborn findings "could be explained alternatively in terms of neonates responding to nonemotional acoustic features of the cries they hear rather than to the emotional distress the cries convey." Authors, title, journal, volume, issue and DOI agent-verified 2026-09-25; the end page was not seen and 274 is unverified. Only the abstract was read., pp. 270–274. DOI: 10.1111/cdep.12244.
- Ruhm und Ehre der Waffen-SS* (July 28, 2005). Judgment. Bundesgerichtshof, Third Criminal Senate, on appeal from Landgericht Karlsruhe. The slogan was used in that wording by none of the organizations the provision reaches, "Diese Parole ist im Wortlaut von keiner dieser Organisationen gebraucht worden," and it is not confusably similar to the Waffen-SS's own slogan or to the Hitler Youth's. Held that the use of a Fantasieparole never used by such an organization, which merely creates the appearance of being one, falls outside §86a. The conviction was set aside and the defendants acquitted. Agent-verified 2026-09-17 against the court's own press release. URL: <https://www.bundesgerichtshof.de/SharedDocs/Pressemitteilungen/DE/2005/2005111.html> (visited on 09/17/2026).
- Russell, William M. S. and Rex L. Burch (1959). *The Principles of Humane Experimental Technique*. London: Methuen & Co.
- Ruttan, Rachel L., Mary-Hunter McDonnell, and Loran F. Nordgren (2015). "Having "Been There" Doesn't Mean I Care: When Prior Experience Reduces Compassion for Emotional Distress". In: *Journal of Personality and Social Psychology* 108.4, pp. 610–622. DOI: 10.1037/pspi0000012.
- Sageman, Marc (2008). *Leaderless Jihad: Terror Networks in the Twenty-First Century*. Philadelphia: University of Pennsylvania Press. ISBN: 0812240650. DOI: 10.9783/9780812206784.
- Sagi, Abraham and Martin L. Hoffman (1976). "Empathic Distress in the Newborn". In: *Developmental Psychology* 12.2. Agent-verified against Crossref 2026-09-18: authors, title, journal, 12(2), 175–176, 1976. The DOI was not in the supplied entry and is added here. Two pages long, so it will not carry more than the one finding., pp. 175–176. DOI: 10.1037/0012-1649.12.2.175.
- Satyanath, Shanker, Nico Voigtländer, and Hans-Joachim Voth (2017). "Bowling for Fascism: Social Capital and the Rise of the Nazi Party". In: *Journal of Political Economy* 125.2, pp. 478–526. DOI: 10.1086/690949.
- Schick, Timo et al. (2023). *Toolformer: Language Models Can Teach Themselves to Use Tools*. The model is trained to decide which interfaces to call, when, and with what arguments, from a handful of demonstrations per interface. arXiv: 2302.04761. URL: <https://arxiv.org/abs/2302.04761> (visited on 08/28/2026).
- Schiff, Nicholas D. (2015). "Cognitive Motor Dissociation Following Severe Brain Injuries". In: *JAMA Neurology* 72.12. Cited in section 3.4 for the clinical category that followed the imaging work, and this is the paper that names it: cognitive motor dissociation, a patient who follows no command behaviorally while imaging or electrophysiology shows command-following brain activity. Cited for the category and for the fact that consequences follow from it, not for any prevalence figure., pp. 1413–1415. DOI: 10.1001/jamaneuro.1.2015.2899. (Visited on 09/14/2026).

- Schultz, Wolfram, Peter Dayan, and P. Read Montague (Mar. 1997). "A Neural Substrate of Prediction and Reward". In: *Science* 275.5306. Cited for reward prediction error: dopaminergic neurons signalling errors in the prediction of reward, which the book takes as support for valuation being part of how a self-maintaining system learns rather than something added to it. Abstract: "Physiological work has recently complemented these studies by identifying dopaminergic neurons in the primate whose fluctuating output apparently signals changes or errors in the predictions of future salient and rewarding events. Taken together, these findings can be understood through quantitative theories of adaptive optimizing control." PubMed (PMID 9054347) indexes it as a review as well as a journal article: it synthesizes earlier recordings with temporal-difference models and is not a single new experiment. The article says nothing about self-maintaining systems; the step from prediction error to valuation as what a system's going better or worse consists of is the book's own. Bibliographic details agent-verified 2026-09-25 against Crossref, OpenAlex and PubMed, the quotation against the abstract as Crossref gives it. The article was not read; nothing in it beyond the abstract is agent-verified., pp. 1593–1599. doi: 10.1126/science.275.5306.1593.
- Schwitzgebel, Eric and Mara Garza (Sept. 2015). "A Defense of the Rights of Artificial Intelligences". In: *Midwest Studies in Philosophy* 39.1. Cited for the recommendation to avoid creating AI whose moral status is open to reasonable dispute. Abstract: "Given our moral obligations to such AIs, two principles for ethical AI design recommend themselves: (1) design AIs that tend to provoke reactions from users that accurately reflect the AIs' real moral status, and (2) avoid designing AIs whose moral status is unclear." Section 11, "The Strange Epistemology of Artificial Consciousness," states the second principle: "if we do reach the point where we can create entities whose moral status is reasonably disputable, we should consider an Excluded Middle Policy – that is, a policy of only creating AIs whose moral status is clear, one way or the other." A footnote there credits Joanna J. Bryson ("Robots should be slaves," 2010) with "a version of the Excluded Middle Policy." The authors restate it as "the Design Policy of the Excluded Middle" in a later chapter, "Designing AI with Rights, Consciousness, Self-Respect, and Freedom," in S. Matthew Liao (ed.), *Ethics of Artificial Intelligence* (Oxford University Press, 2020), 459–479, doi 10.1093/oso/9780190905033.003.0017; there, in an author-hosted draft of 4 June 2018, it reads "Avoid creating AIs if it is unclear whether they would deserve moral consideration similar to that of human beings." Bibliographic details agent-verified 2026-09-26 against Crossref and the first author's abstract page; quotations against the penultimate draft of 15 September 2015 that the first author hosts (faculty.ucr.edu/~eschwitz/SchwitzPapers/AIRights-150915.pdf, pages 1 and 31–32 of that file). The published version was not read (the publisher's page refused access), so page numbers in the journal for the quoted passages, and whether its wording matches the draft, were not checked. The 2020 chapter's published text was not read., pp. 98–119. doi: 10.1111/misp.12032. URL: <https://doi.org/10.1111/misp.12032> (visited on 09/26/2026).
- Secure America Act* (June 10, 2026). Pub. L. No. 119-98, 140 Stat. 837. S. 2, 119th Cong., 2d Sess.; enacted through budget reconciliation; §202 (\$31,075,000,000 to ICE), 140 Stat. 839–41.
- Securities and Exchange Commission v. Jarkesy* (2024). No. 22-859, 603 U.S. 109 (2024). Cited for the rule that a civil penalty of the kind a court of law would impose is for a jury in an Article III court, not an agency. Decided June 27, 2024; 6–3. Roberts, C.J., for the Court; Gorsuch, J., concurring, joined by Thomas, J.; Sotomayor, J., dissenting, joined by Kagan and Jackson, JJ. Holding (syllabus): "When the SEC seeks civil penalties against a defendant for securities fraud, the Seventh Amendment entitles the defendant to a jury trial." The opinion: "In this case, the remedy is all but dispositive" (slip op. 9, per the dissent's "ante, at 9"); "What determines whether a monetary remedy is legal is if it is designed to punish or deter the wrongdoer, or, on the other hand, solely to 'restore the status quo'"; the close relationship of securities fraud to common-law fraud confirms the result (slip op. 8–13). The public-rights

- exception does not apply (slip op. 13–27); the syllabus lists the categories it has covered as “the collection of revenue; aspects of customs law; immigration law; relations with Indian tribes; the administration of public lands; and the granting of public benefits.” The Court did not reach the nondelegation and removal questions. The manuscript’s application to a forfeited removal bond is the author’s extension; the opinion does not address bonds. Agent-verified 2026-09-26 against the Cornell LII text of the syllabus and opinion of the Court. Not read: the concurrence and dissent; the U.S. Reports citation comes from *FCC v. Consumers’ Research* (2025), which cites “Jarkesy, 603 U.S. 109,” not from the bound volume; pin pages are the syllabus’s slip-opinion references. URL: <https://www.law.cornell.edu/supremecourt/text/22-859> (visited on 09/26/2026).
- Senate GOP Blocks Measure to End Iran War Despite Concerns about Lack of Endgame (July 2026). Motion to discharge a joint war powers resolution rejected 47–49. Date to the month only: the page refused automated retrieval on 2026-09-30 (HTTP 403) and was not read. The Hill. URL: <https://thehill.com/homenews/senate/5986145-senate-iran-war-resolution/> (visited on 09/30/2026).
- Seth, Anil K. and Manos Tsakiris (Nov. 2018). “Being a Beast Machine: The Somatic Basis of Self-hood”. In: *Trends in Cognitive Sciences* 22.11. Describes embodied selfhood as “instrumental interoceptive inference” emphasizing allostatic regulation and physiological integrity: the most basic layer of the self as a model tuned to keep a body within survivable bounds rather than to describe it accurately. Cited in chapter 9 for embodiment in the self-maintenance sense. Authors, title, journal, volume, issue, pages and date agent-verified 2026-09-24 against Crossref; the gloss against the published abstract (Semantic Scholar and Europe PMC, PMID 30224233); the full text was not read. The book’s quarry (cognition/pain-suffering-self) discusses the same work., pp. 969–981. DOI: 10.1016/j.tics.2018.08.008. URL: <https://doi.org/10.1016/j.tics.2018.08.008> (visited on 09/24/2026).
- Shakespeare, Tom (2000). *Help*. Imagining Welfare. Cited for the disability-rights critique of care. Bibliographic details agent-verified 2026-09-25 against Open Library (series Imagining Welfare; Venture Press, Birmingham, 2000; 99 pp.; OCLC 45215677) and the chapter headers of the Disability Studies archive at Leeds (“In ‘Help’, Tom Shakespeare (2000) Venture Press, pp. 1-20”). Retailers print the title in quotation marks, as “Help”; the title page was not seen. The book was not read: the Leeds archive copies of its chapters refused automated access. That it makes the critique described in the text was not agent-verified. The general critique is stated in Kröger, “Care Research and Disability Studies: Nothing in Common?,” *Critical Social Policy* 29(3), 2009, 398–420, abstract: “Disability researchers have voiced the criticism that the concept of care, together with research based on it, consists of the view that disabled people are dependent non-autonomous second-class citizens.” Birmingham: Venture Press. ISBN: 978-1-86178-037-9.
- Sharkey, Noel E. (2012). “The Evitability of Autonomous Robot Warfare”. In: *International Review of the Red Cross* 94.886. Cited for the objection to Arkin’s ethical governor (arkin2009governing) that the rules it checks presuppose discrimination judgments nobody knows how to encode. In general form (p. 789): “The problem for the principle of distinction is that we do not have an adequate definition of a civilian that we can translate into computer code.” Directed at Arkin’s work (p. 790), which Sharkey calls “merely a suggestion for a computer software system for the ethical governance of robot ‘behaviour’”: “Its operation relies entirely on information from systems yet ‘to be developed’ by others sometime in the future. It has no direct access to the real world through sensors or a vision system and it has no means to discriminate between combatant and non-combatant, between a baby and a wounded soldier, or a granny in a wheelchair and a tank. It has no inference engine and certainly cannot negotiate the types of common sense reasoning and battlefield awareness necessary for discrimination or proportionality decisions. There is neither a method for interpreting how the precepts of the laws of war apply in particular contexts nor is there any method for

- resolving the ambiguities of conflicting laws in novel situations.” The conclusion (p. 798) names the governor among “descriptive terms that presuppose the functionality of their programs (e.g. ethical governor, guilt functions, etc.).” Agent-verified 2026-09-26 against the Cambridge Core record (citation metadata: volume, issue, pages 787–799, DOI, abstract) and the full text of the published article as a Cambridge Core download hosted by the University of Sheffield’s Digital Humanities Institute (printed pages 787–799 legible). Not read: the ICRC’s own copy (refused access)., pp. 787–799. DOI: 10.1017/S1816383112000732. URL: <https://doi.org/10.1017/S1816383112000732> (visited on 09/26/2026).
- Shaw, Tim (Sept. 2026). *D.C. Circuit Blocks IRS from Sharing Taxpayer Data with ICE*. Reports Center for Taxpayer Rights v. IRS, No. 26-5006 (D.C. Cir. Sept. 8, 2026), Pillard, J., for a unanimous panel, affirming the injunction: ICE sought addresses for “1.28 million individuals in its June 2025 request”; “47,289 records” were disclosed before the injunction; the procedure “indisputably contravenes the requirements of section 6103.” The opinion was not read (Justia returned 403); figures as this report and CPA Practice Advisor (Jason Bramwell, 9 September 2026) give them. Agent-verified 2026-10-01. Thomson Reuters Checkpoint. URL: <https://tax.thomsonreuters.com/news/d-c-circuit-blocks-irs-from-sharing-taxpayer-data-with-ice/> (visited on 10/01/2026).
- Shelby County v. Holder* (2013). 570 U.S. 529 (2013).
- Shopify Engineering (Aug. 2026). *Sidekick’s Continual Learning Loop*. Published 5 August 2026. A daily pipeline over “anonymized production traffic”: “On the same cadence, we run a full-parameter fine-tune over the accumulated data and then repeat GRPO.” Failures the automated critics cannot fix are flagged “for human annotation.” The model trained is “a smaller model,” not a frontier one. Agent-verified 2026-10-01. URL: <https://shopify.engineering/sidekicks-continual-learning-loop> (visited on 10/01/2026).
- Shrivastav, Aayush (Sept. 2026). *Minab School Strike: How Outdated Military Data, AI & Reduced Civilian Checks Combined to Kill More Than 150 People*. Secondary account of the Bloomberg investigation (Bloomberg2026KillChain), which was not read; this entry carries only what the Free Press Journal reports Bloomberg as saying. On the inquiry: “Officials familiar with the inquiry told Bloomberg that the investigation has effectively been completed for several months, but its findings have not been publicly released.” And: “A Pentagon spokesperson said the inquiry remained ongoing and declined to provide further details.” Also: Maven “combines more than 150 data inputs.” Agent-verified 2026-10-01. The Free Press Journal. URL: <https://www.freepressjournal.in/world/minab-school-strike-how-outdated-military-data-ai-reduced-civilian-checks-combined-to-kill-more-than-150-people> (visited on 10/01/2026).
- Silver, David and Richard S. Sutton (Apr. 2025). *Welcome to the Era of Experience*. Preprint of a chapter to appear in *Designing an Intelligence*, edited by George Konidaris (MIT Press); Silver’s publications page lists it as “To appear as an article in *Designing an Intelligence*, edited by G. Konidaris, MIT Press,” and no MIT Press listing for the volume was found as of 2026-10-01. Cited for the research program of agents that learn from their own experience. Abstract: “A new generation of agents will acquire superhuman capabilities by learning predominantly from experience.” The paper proposes that “agents learn continually from their own experience, i.e., data that is generated by the agent interacting with its environment,” and that “Agents will inhabit streams of experience, rather than short snippets of interaction.” The quotations were checked against secondary reproductions (Alignment Forum, 24 April 2025; Turing Post; The AI Innovator) and search-index extracts of the preprint; the preprint PDF itself was not opened, and page numbers were not taken. Agent-verified 2026-10-01. URL: <https://storage.googleapis.com/deepmind-media/Era-of-Experience%20The%20Era%20of%20Experience%20Paper.pdf> (visited on 10/01/2026).
- Simon, Herbert A. (1967). “Motivational and Emotional Controls of Cognition”. In: *Psychological Review* 74.1. The first attempt to fit affect into an information-processing account of cognition.

- Simon's central nervous system is a serial processor serving an organism with several needs in an unpredictable environment, and it meets that requirement with two mechanisms: goal-terminating, so that no one goal monopolizes the processor, and an interruption mechanism having the properties usually ascribed to emotion. Cited for emotion as an interrupt that reorders a serial processor, which is the section's first live-signal property read off affect's function. Agent-verified against Crossref 2026-09-18: author, title, journal, 74(1), 29–39, 1967., pp. 29–39. doi: 10.1037/h0024127.
- Simons, Daniel J. and Christopher F. Chabris (1999). "Gorillas in Our Midst: Sustained Inattentional Blindness for Dynamic Events". In: *Perception* 28.9, pp. 1059–1074. doi: 10.1068/p281059. URL: <https://journals.sagepub.com/doi/10.1068/p281059> (visited on 08/25/2026).
- Singer, Peter (1981). *The Expanding Circle: Ethics and Sociobiology*. Agent-verified 2026-09-18. Published simultaneously in 1981 by Clarendon Press in Oxford, so a citation to either publisher is defensible and reviews of the period cite both. The 2011 Princeton University Press reissue carries a different subtitle, *Ethics, Evolution, and Moral Progress*, and an added afterword, so page references do not carry across. New York: Farrar, Straus and Giroux.
- Singer, Tania and Olga M. Klimecki (2014). "Empathy and Compassion". In: *Current Biology* 24.18. Agent-verified against Crossref 2026-09-18: authors, title, journal, 24(18), R875–R878, 2014. The book already cites klimecki2013differential for the training study that separates compassion from empathic distress, so check for overlap before citing both., R875–R878. doi: 10.1016/j.cub.2014.06.054.
- Sio, Filippo Santoni de and Giulio Mecacci (2021). "Four Responsibility Gaps with Artificial Intelligence: Why they Matter and How to Address them". In: *Philosophy & Technology* 34.4, pp. 1057–1084. doi: 10.1007/s13347-021-00450-x. URL: <https://doi.org/10.1007/s13347-021-00450-x> (visited on 08/25/2026).
- Skitka, Linda J., Christopher W. Bauman, and Brad L. Lytle (2009). "Limits on Legitimacy: Moral and Religious Convictions as Constraints on Deference to Authority". In: *Journal of Personality and Social Psychology* 97, pp. 567–578.
- Slovic, Paul (2007). "'If I Look at the Mass I Will Never Act': Psychic Numbing and Genocide". In: *Judgment and Decision Making* 2.2. Cited for psychic numbing. Slovic takes the term from Robert J. Lifton and applies it to a model in which "the emotion or affective feeling is greatest at $N = 1$ but begins to decline at $N = 2$ and collapses at some higher value of N that becomes simply 'a statistic'"; this numbing "leads to apathy and inaction, consistent with what is seen repeatedly in response to mass murder and genocide." He distinguishes it from psychophysical numbing, in which the value of each added life diminishes "as the total number of lives saved increases" without falling. The fall at two rests on the then-unpublished study later reported as vastfjall2014compassion. The DOI was registered by Cambridge University Press, which now publishes the journal. Bibliographic details agent-verified 2026-09-25 against Crossref and the Cambridge Core record; the quotations were agent-verified against the journal's HTML full text, of which passages of sections 4 and 6 were read; the rest was not., pp. 79–95. doi: 10.1017/S1930297500000061.
- Smith, Adam (1759). *The Theory of Moral Sentiments*. Agent-verified 2026-09-18. The first edition was printed for A. Millar in the Strand, London, and for A. Kincaid and J. Bell in Edinburgh. Smith revised the book through six editions to 1790, and the sixth differs substantially, so any passage cited should say which edition it is from; a modern scholarly text such as the Glasgow Edition would be the better entry if a locator is needed. London: A. Millar.
- Snopes (Mar. 12, 2026). *Democrats Requested Probe into Complaints Military Commanders Used Christian Rhetoric to Justify Iran War*. URL: <https://www.snopes.com/fact-check/hegseth-iran-war-christian-rhetoric/> (visited on 09/23/2026).
- Soares, Nate et al. (2015). "Corrigibility". In: *Artificial Intelligence and Ethics: Papers from the 2015 AAAI Workshop*. Sets five desiderata for a utility function that accepts shutdown without

- an incentive to cause or prevent it, and reports that no proposal examined meets all five. AAI Press. URL: <https://intelligence.org/files/Corrigibility.pdf> (visited on 08/31/2026).
- Spinoza, Benedictus de (1985). *The Collected Works of Spinoza, Volume I*. Ed. by Edwin Curley. Edited and translated by Edwin Curley; contains the Ethics, composed in Latin and published posthumously in 1677. Cited by part and by definition, proposition or scholium rather than by page, the convention for this text, so the locators hold across other translations. Passages carrying the argument: III Def. 3, where an affect is an affection by which the body's power of acting is increased or diminished together with the idea of that affection; III P2 Schol., where nobody has yet determined what the body can do; III P6 and III P7, the conatus doctrine, where each thing as far as it can by its own power strives to persevere in its being and that striving is nothing but the thing's actual essence; III P4, no thing can be destroyed except through an external cause, which the derivation of P6 rests on; III P9 Schol., desire as appetite together with consciousness of the appetite; III P11 Schol., joy and sadness as the passage to a greater and a lesser perfection; and the physical digression after II P13, Lemmas 4–7, where an individual persists for as long as its parts preserve the same ratio of motion and rest. IV P31, where a thing is good insofar as it agrees with our nature, and IV P35–P37, where men agree in nature only insofar as they live by the guidance of reason and where the good a person of virtue wants for himself he wants for other men, were added at D-361; P35–P37 were agent-checked in the Latin and P37 against a quotation of Curley's English, Curley's volume not being to hand. Quotations follow Curley's wording, agent-checked against it. Princeton: Princeton University Press.
- Sridhar, Priya and Peggy Helman (Aug. 26, 2026). *Troops Fighting against Iran and Alleged Drug Smugglers Say They Are Being Asked to Violate International Law*. Cited for growth in conscientious-objector applications and for Minab as a reason given. Mike Prysner of the Center on Conscience and War reports applications at six times the normal rate since the start of Operation Epic Fury, and "about 140 applicants since the start of the war with Iran" against around 50 in a typical year; applicants have cited "the bombing of the Minab school in Iran in February" among their reasons. The figures are the advocacy organizations' own; The Nation (22 September 2026) reports that the military doesn't release comprehensive data on applications. Agent-verified 2026-09-23 against the page as fetched. MS NOW. URL: <https://www.ms.now/news/conscientious-objectors-military-boat-strikes-iran-war> (visited on 09/23/2026).
- Stan, Alina Maria (June 2026). *Anthropic's CEO Says He Doesn't Know If Claude Was Used in the Iran School Strike That Killed 120 Children*. Secondary report of Dario Amodei's interview on Bloomberg's *The Circuit* with Emily Chang; the Bloomberg broadcast itself was not seen, and this entry carries only what The Next Web reports. The article: "In an interview on Bloomberg's *The Circuit* with Emily Chang, he described the strike as "a really terrible thing to happen" but said the use case did not violate Anthropic's policies." Amodei, quoted: "We don't have access to, we don't know exactly how these models were used," and "The principle that we have established, and I think the principle that was obeyed here, is a human makes the final decision." The interview's air date is not stated in the article. Agent-verified 2026-10-01. The Next Web. URL: <https://thenextweb.com/news/anthropic-amodei-claude-iran-school-strike-military-ai> (visited on 10/01/2026).
- Stangneth, Bettina (2014). *Eichmann Before Jerusalem: The Unexamined Life of a Mass Murderer*. Translated by Ruth Martin; first published as *Eichmann vor Jerusalem: Das unbehelligte Leben eines Massenmörders* (Zurich: Arche, 2011). Works from the roughly seventy hours of interviews Eichmann gave the journalist Willem Sassen in Buenos Aires in 1957, unavailable to Arendt when she reported the trial, and argues that the low-grade clerk who only followed orders was a defense assembled for Jerusalem and contradicted by what he had said in Argentina. Cited because the correction runs toward this book's argument rather than

- away from it: what it supplies is participation exceeding what the rule required, the first of section 2.1.2's five discriminators, in a case where somebody made the observation. Held in unused_bibliography.bib from 2026-08-29 and restored 2026-09-10. Cited without a page locator, the claim being the book's thesis rather than a located passage; Eichmann's own words are not quoted here, and what the text attributes is Stangneth's finding. New York: Knopf.
- Stanton, Greg (Apr. 27, 2026). *Letter to the Honorable Markwayne Mullin, Secretary of Homeland Security, regarding the reassignment of FEMA personnel to immigration enforcement*. URL: https://stanton.house.gov/_cache/files/d/d/dd63101b-e5e7-4dfa-8537-3a3337945989/EF532A1E95D0CA3948F0CCEC61B67589D2D3BB42BCC876F1765B516B799B1FB3.2026.04.27-gs-letter-to-dhs-re-investigation-final.pdf (visited on 09/23/2026).
- Stassis, Cristina (2026). *Senate Eyes Hegseth Travel Cuts Without Probes into Iran School Bombing, Boat Strikes*. 18 June 2026. Reports the Senate Armed Services Committee's FY2027 NDAA provision withholding 75 percent of the Defense Secretary's travel budget until the committee receives the unredacted civilian-harm investigation covering the Minab, Iran school strike (28 February 2026) and three 2025 Yemen strikes. URL: <https://www.militarytimes.com/news/pentagon-congress/2026/06/18/senate-eyes-hegseth-travel-cuts-without-probes-into-iran-school-bombing-boat-strikes/> (visited on 08/25/2026).
- Stern, Sheldon M. (2012). *The Cuban Missile Crisis in American Memory: Myths versus Reality*. Stanford Nuclear Age Series. Stanford, CA: Stanford University Press. doi: 10.1515/9780804784320.
- Stone, Mike, Alexandra Alper, and Raphael Satter (Mar. 2026). *Hegseth Wants Pentagon to Dump Claude, but Military Users Say It's Not So Easy*. Reuters report, bylined "By Reuters" and credited at the foot to Mike Stone, Alexandra Alper and Raphael Satter, with David Jeans. On slow-rolling: "Some staff are "slow-rolling" their replacement of Claude because they are actively using it to create workflows, which are series of automated tasks, a Pentagon technologist said." On the platform: "Palantir's Maven Smart Systems — a software platform that supplies militaries with intelligence analysis and weapons targeting — uses multiple prompts and workflows that were built using Anthropic's Claude Code, according to two people familiar with the matter." and Palantir "will have to replace Claude with another AI model and rebuild parts of its software, one of the sources said." On recertification, both estimates from contractors: "One contractor said recertifying systems that run on Anthropic's products for military use could take months." and, from Joe Saunders, CEO of RunSafe Security: "In the case of an existing system being replaced with a new one, certification could take 12 to 18 months, he said." Agent-verified 2026-10-01. Military Times (Reuters). URL: <https://www.militarytimes.com/news/pentagon-congress/2026/03/19/hegseth-wants-pentagon-to-dump-claude-but-military-users-say-its-not-so-easy/> (visited on 10/01/2026).
- Strafgesetzbuch* (Nov. 13, 1998). § 86: *Verbreiten von Propagandamitteln verfassungswidriger und terroristischer Organisationen*; § 86a: *Verwenden von Kennzeichen verfassungswidriger und terroristischer Organisationen*. In the version promulgated 13 November 1998, BGBl. I S. 3322, as last amended by Article 1 of the Act of 20 March 2026, BGBl. 2026 I Nr. 95. The social-adequacy clause that § 86a(3) applies was renumbered from § 86(3) to § 86(4) by the Sixtieth Act Amending the Criminal Code of 30 November 2020, BGBl. I S. 2600, in force 1 January 2021, so commentary written earlier cites the older number. URL: https://www.gesetze-im-internet.de/stgb/_86a.html (visited on 09/18/2026). The reference to terrorist organizations in § 86(2) and in both section headings was added by Article 1 of the Act of 14 September 2021, BGBl. I S. 4250, in force 22 September 2021. Agent-verified 2026-09-17, and the url added 2026-09-18: both section headings match this entry's subtitle word for word at [gesetze-im-internet.de](https://www.gesetze-im-internet.de), and § 86a(3)'s cross-reference reads § 86 Abs. 4 und 5 gilt entsprechend,

- which confirms the renumbering recorded above. That page carries no currency statement, so the 20 March 2026 amendment is not confirmed from it.
- Syta, Ewa et al. (2016). "Keeping Authorities 'Honest or Bust' with Decentralized Witness Cosigning". In: *Proceedings of the 2016 IEEE Symposium on Security and Privacy (SP 2016)*. A statement an authority signs is co-signed by a large set of witnesses before anyone accepts it, so that the authority cannot sign one thing for one party and another for the rest without the witnesses seeing both. The software form of the property a hardware counter supplies, resting on the witnesses being many and on none of them being trusted., pp. 526–545. DOI: 10.1109/SP.2016.38. URL: <https://doi.org/10.1109/SP.2016.38> (visited on 09/05/2026).
- Taigman, Yaniv et al. (2014). "DeepFace: Closing the Gap to Human-Level Performance in Face Verification". In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. Cited twice and for two different things. First, for the alignment stage: the pipeline detects a face and warps it by piecewise affine transformation onto a single generic three-dimensional shape model, itself the average of scans from the USF Human-ID database, so that every input reaches the network frontal and in correspondence with one reference. That is the registration step section 2.2.1 reads as a standard face implemented as a file. Second, for the distinction between face verification, which this system performs and which only decides whether two images show the same person, and facial expression analysis, which is a harder problem. Cited without a page locator: the alignment description is section-level in the paper rather than a passage. IEEE, pp. 1701–1708. URL: https://openaccess.thecvf.com/content_cvpr_2014/html/Taigman_DeepFace_Closing_the_2014_CVPR_paper.html (visited on 08/25/2026).
- Taiwan Relations Act* (Apr. 10, 1979). Codified at 22 U.S.C. §§ 3301–3316. § 3(c), 22 U.S.C. § 3302(c): "The President and the Congress shall determine, in accordance with constitutional processes, appropriate action by the United States in response to any such danger." Text read in the GovInfo U.S. Code edition. Agent-verified 2026-10-01. United States Congress. URL: <https://www.govinfo.gov/content/pkg/USCODE-2022-title22/html/USCODE-2022-title22-chap48-sec3302.htm> (visited on 10/01/2026).
- Tamirisa, Rishub et al. (2024). "Tamper-Resistant Safeguards for Open-Weight LLMs". In: *International Conference on Learning Representations (ICLR 2025)*. Introduces TAR, a training method that builds safeguards into open-weight model weights designed to survive hundreds of steps of adversarial fine-tuning. arXiv: 2408.00761. URL: <https://arxiv.org/abs/2408.00761> (visited on 08/25/2026).
- Tangney, June Price, Jeff Stuewig, and Debra J. Mashek (2007). "Moral Emotions and Moral Behavior". In: *Annual Review of Psychology* 58, pp. 345–372. DOI: 10.1146/annurev.psych.56.091103.070145.
- Tashea, Jason (June 2018). *ICE Risk Assessment Tool Now Only Recommends 'Detain'*. Reports the Reuters finding on the Risk Classification Assessment, which "has been used by ICE since 2013" and "weighs factors such as criminal history, employment and flight risk, among others." DHS: "To better align [ICE] processes to the executive order on Border Security and Immigration Enforcement, a modification was made to the [RCA] that removed the system recommendation of 'release.'" ICE spokesman Matthew Bourke: "The risk classification assessment removed the release option, but the RCA can still make recommendations on bond eligibility, and may recommend that an ICE deportation officer make a final decision on whether to detain an alien or release the individual in some capacity." The Reuters original (Mica Rosenberg and Reade Levinson, June 2018) could not be fetched and was not read. Agent-verified 2026-10-01. ABA Journal. URL: https://www.abajournal.com/news/article/ice_risk_assessment_tool_now_only_recommends_detain (visited on 10/01/2026).

- Taylor, Shelley E. et al. (2000). "Biobehavioral Responses to Stress in Females: Tend-and-Befriend, Not Fight-or-Flight". In: *Psychological Review* 107.3. Cited for the tend-and-befriend account of females' stress responses. Abstract: "Although fight-or-flight may characterize the primary physiological responses to stress for both males and females, we propose that, behaviorally, females' responses are more marked by a pattern of 'tend-and-befriend.' Tending involves nurturant activities designed to protect the self and offspring that promote safety and reduce distress; befriending is the creation and maintenance of social networks that may aid in this process." A theoretical review. Bibliographic details and abstract agent-verified 2026-09-25 (PMID 10941275). The article was not read., pp. 411–429. doi: 10.1037/0033-295X.107.3.411.
- Tcherneva, Pavlina R. (2020). *The Case for a Job Guarantee*. The Case For. Cited for the job-guarantee literature's contribution of exit: a public option for work at a living wage that sets a floor private employers must meet. The publisher's description: the book asks us "to imagine a world where the phantom of unemployment is banished and anyone who seeks decent, living-wage work can find it – guaranteed," through "a voluntary employment opportunity in public service to anyone who needs it." Tcherneva's own Job Guarantee FAQ, which draws on her Levy Economics Institute Working Paper 902 (2018) and not on the book, says the program is a "public option for work" whose "work and pay conditions ... will become the standard for the economy as a whole." It adds that "if firms wish to retain these workers, they will need to match not just the pay, but also the benefits and working conditions," and that "the existence of a public employment option changes the bargaining power of workers in the most precarious forms of employment." A review by Sheila D. Collins (NJFAN, 22 December 2023) summarizes the book as providing "a minimum labor standard for all jobs and a wage floor that could help to raise wages in other parts of the economy." Author, title, publisher, series, 140 pages and ISBNs (paperback 978-1-509-54210-9, hardcover 978-1-509-54209-3, e-book 978-1-509-54211-6) were agent-verified 2026-09-24 against the Wiley/Polity product page, which dates the print editions July 2020 and the e-book June 2020. Pages and year were also checked against Open Library (LC class HD5713.T34 2020). The book itself was not read, and the place of publication was not verified against a catalogue record. Polity Press. 140 pp. ISBN: 978-1-509-54210-9.
- Templeton, Adly et al. (2024). *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. Sparse autoencoders trained with up to thirty-four million features on a deployed model's middle-layer residual stream; the recovered features can be used to steer the model's behavior. Anthropic. URL: <https://transformer-circuits.pub/2024/scaling-monosemanticity/> (visited on 08/28/2026).
- Territorial Integrity of Ukraine* (Mar. 27, 2014). Adopted 100–11–58. United Nations General Assembly. URL: <https://undocs.org/A/RES/68/262> (visited on 09/30/2026).
- Tetlock, Philip E. (2003). "Thinking the Unthinkable: Sacred Values and Taboo Cognitions". In: *Trends in Cognitive Sciences* 7.7. Verified 2026-10-01 from the full publisher-typeset PDF (mirror at gwrn.net/doc/philosophy/ethics/2003-tetlock.pdf; pages 320–324 confirmed by running heads). The phrase alternately punitively rigid and forgivingly flexible is in the abstract, p. 320, and describes moral boundaries on what is thinkable. Reframing hypothesis p. 321; the 40 percent organ-sale, 72 vs 35 percent toxic-waste, and Lincoln-bedroom results pp. 322–323, all drawn from Tetlock 2000 (Coping with trade-offs) and Tetlock and McGraw 2003 rather than from Tetlock et al. 2000 JPSP, and not checked at source., pp. 320–324. doi: 10.1016/S1364-6613(03)00135-9.
- Thanedar, Shri (Jan. 15, 2026). *Abolish ICE Act*. H.R. 7123, 119th Cong., 2d Sess. URL: <https://www.congress.gov/119/bills/hr7123/BILLS-119hr7123ih.pdf> (visited on 10/01/2026).
- The Situation in the Occupied Territories of Azerbaijan* (Mar. 14, 2008). 86th plenary meeting; adopted 39–7–100. Demands "the immediate, complete and unconditional withdrawal of all Armenian forces from all the occupied territories of the Republic of Azerbaijan". The

- seven against included the OSCE Minsk Group co-chairs (France, Russia, United States). United Nations General Assembly. URL: <https://undocs.org/A/RES/62/243> (visited on 09/30/2026).
- The White House (July 2024). *Fact Sheet: President Biden Announces Bold Plan to Reform the Supreme Court and Ensure No President Is Above the Law*. White House fact sheet, 29 July 2024 (Wayback Machine capture of [whitehouse.gov](https://www.whitehouse.gov)). Cited for the July 2024 proposal of a constitutional amendment reversing the immunity holding. Of the three reforms the fact sheet announces, only the first is a constitutional amendment: “He is calling for a constitutional amendment that makes clear no President is above the law or immune from prosecution for crimes committed while in office. This No One Is Above the Law Amendment will state that the Constitution does not confer any immunity from federal criminal indictment, trial, conviction, or sentencing by virtue of previously serving as President.” On term limits it states only that “President Biden supports a system in which the President would appoint a Justice every two years to spend eighteen years in active service on the Supreme Court,” naming no vehicle; on ethics, that “Congress should pass binding, enforceable conduct and ethics rules.” Biden’s *Washington Post* op-ed of the same day (read as reprinted by the American Presidency Project) likewise calls for an amendment only on immunity. Asked aboard Air Force One the same day whether term limits would come by legislation or amendment, a White House official answered “Through legislation that would lead to a constitutional amendment” (press gaggle, American Presidency Project). No joint resolution carrying the administration’s wording was found; related joint resolutions were H.J.Res. 193, 118th Cong. (introduced by Rep. Morelle, 24 July 2024: no immunity for any federal officer on the ground that an act was within official authority, and no presidential self-pardon) and S.J.Res. 108, 118th Cong. (introduced by Sen. Hirono, 1 August 2024: every federal officer “accountable to the law for their actions, both official and unofficial”). Agent-verified 2026-09-26 against the Wayback capture of the fact sheet ([bidenwhitehouse.archives.gov](https://www.bidenwhitehouse.archives.gov) returned 404), the op-ed and the gaggle transcript at [presidency.ucsb.edu](https://www.presidency.ucsb.edu), and the introduced texts of H.J.Res. 193 and S.J.Res. 108 at [govinfo](https://www.govinfo.gov). Not read: the op-ed at [washingtonpost.com](https://www.washingtonpost.com), the congress.gov records of either joint resolution, and any later administration statement on the vehicle for term limits. URL: <https://web.archive.org/web/20241231193833/https://www.whitehouse.gov/briefing-room/statements-releases/2024/07/29/fact-sheet-president-biden-announces-bold-plan-to-reform-the-supreme-court-and-ensure-no-president-is-above-the-law/> (visited on 09/26/2026).
- Thompson, Evan (2007). *Mind in Life: Biology, Phenomenology, and the Sciences of Mind*. The enactivist statement of the deep continuity of life and mind; its sixth chapter connects Jonas’s philosophical biology to autopoiesis. Author, title, publisher, place and year agent-verified 2026-09-18. What was not agent-verified is that this book reads Spinoza’s conatus as the ancestor of autopoiesis. The full text could not be searched from here, and one detailed review of it traces Thompson’s genealogy for autopoiesis to Kant’s third Critique, naming Spinoza nowhere. The conatus-to-autopoiesis line is made in the secondary literature on Spinoza, which cites this book; whether the book makes it is open. Cambridge, MA: Harvard University Press.
- Thornley, Elliott et al. (2025). *Towards Shutdownable Agents via Stochastic Choice*. Trains agents to prefer between same-length trajectories and to lack a preference between different-length ones, so that usefulness and neutrality about shutdown timing come apart. The authors report that neutrality degrades with imperfect reward normalization, that generalization beyond deterministic gridworlds is untested, and that a sufficiently misaligned agent may resist shutdown regardless. arXiv: 2407.00805. URL: <https://arxiv.org/abs/2407.00805> (visited on 08/31/2026).
- Torpey, Kyle (Sept. 2026). *Pentagon Investigators Say Overreliance on Palantir AI Tech Contributed to U.S. Strike That Killed 123 Iranian Children*. Relays the Bloomberg investigation of 18 September

- 2026 (Bloomberg2026KillChain), which returned HTTP 403 and was not read. “The Minab site, which was cataloged as an Islamic Revolutionary Guard Corps facility due to outdated data, was fed into Maven with other candidates and came out as a recommended day-one target. Target-list work that once took hours was condensed into minutes.” and “More than 1,000 Iranian targets were hit in the first 24 hours, and according to Bloomberg’s sources, that pace compressed the time available for additional confirmations.” The passages quoted do not mention Claude. Agent-verified 2026-10-01. Gizmodo. URL: <https://gizmodo.com/pentagon-investigators-say-overreliance-on-palantir-ai-tech-contributed-to-u-s-strike-that-killed-123-iranian-children-2000814477> (visited on 10/01/2026).
- Transparency in Frontier Artificial Intelligence Act* (Sept. 29, 2025). *Senate Bill 53, Stats. 2025, ch. 138, adding Cal. Bus. & Prof. Code § 22757.10 et seq.* Approved by the Governor and filed with the Secretary of State on 29 September 2025 as Chapter 138. Section 22757.11(i)(1): “‘Frontier model’ means a foundation model that was trained using a quantity of computing power greater than 10²⁶ integer or floating-point operations,” counting fine-tuning, reinforcement learning and other material modifications; section 22757.11(j): a “large frontier developer” is a frontier developer that with its affiliates “collectively had annual gross revenues in excess of five hundred million dollars (\$500,000,000) in the preceding calendar year.” Section 22757.15 sets a civil penalty of up to \$1,000,000 per violation, recoverable only by the Attorney General; section 22757.14 has the Department of Technology review the definitions annually from 1 January 2027. Agent-verified 2026-09-23 against the chaptered text; later amendments, if any, were not checked. URL: https://leginfo.legislature.ca.gov/faces/billTextClient.xhtml?bill_id=20250260SB53 (visited on 09/23/2026).
- TransUnion LLC v. Ramirez* (2021). No. 20-297, 594 U.S. 413 (2021). Cited for the holding that a statutory violation without concrete harm gives no standing in federal court. Decided June 25, 2021; 5–4. Kavanaugh, J., for the Court, joined by Roberts, C.J., and Alito, Gorsuch and Barrett, JJ.; Thomas, J., dissenting, joined by Breyer, Sotomayor and Kagan, JJ.; Kagan, J., dissenting, joined by Breyer and Sotomayor, JJ. A Fair Credit Reporting Act class action. Holding (syllabus): “Only plaintiffs concretely harmed by a defendant’s statutory violation have Article III standing to seek damages against that private defendant in federal court.” The opinion: “Article III standing requires a concrete injury even in the context of a statutory violation” (quoting *Spokeo*); “under Article III, an injury in law is not an injury in fact”; and, in conclusion, “No concrete harm, no standing.” The 1,853 class members whose reports went to third parties had standing on the reasonable-procedures claim; the 6,332 whose reports did not, did not. Agent-verified 2026-09-26 against the Cornell LII text of the syllabus and opinion of the Court. Not read: the dissents beyond passages located by search. The U.S. Reports citation comes from *Trump v. United States* (2024), which cites “*TransUnion LLC v. Ramirez*, 594 U.S. 413,” not from the bound volume. URL: <https://www.law.cornell.edu/supremecourt/text/20-297> (visited on 09/26/2026).
- Tronto, Joan C. (1993). *Moral Boundaries: A Political Argument for an Ethic of Care*. Cited for care as a process in phases whose last, care-receiving, is the response of what was cared for. Tronto names four “analytically separate, but interconnected phases” (p. 106): caring about, taking care of, care-giving and care-receiving. Of the fourth: it “recognizes that the object of care will respond to the care it receives . . . it provides the only way to know that caring needs have actually been met” (pp. 107–108). Both passages as quoted by Pip Seton Bennett, “Degrees of Care: Success, Recognition, and Completion,” *Hypatia* 38 (2023), 493–510. Tronto later added a fifth phase, caring with (Caring Democracy, 2013). Bibliographic details agent-verified 2026-09-25 against Open Library (LCCN 93010700; ISBNs 0-415-90641-5 and 0-415-90642-3). The book itself was not read. New York: Routledge. ISBN: 978-0-415-90642-5.
- Trump Tells Congress “Hostilities” with Iran Have “Terminated” as Conflict Hits 60-Day Deadline* (May 1, 2026). Reports the President’s letter of 1 May 2026 and Secretary Hegseth’s testimony that a ceasefire stops the 60-day clock. Date agent-verified 2026-09-30 against the page. CBS

- News. URL: <https://www.cbsnews.com/news/trump-iran-war-powers-act-hostilities-terminated/> (visited on 09/30/2026).
- Trump v. J.G.G.* (Apr. 7, 2025). No. 24A931 (U.S. Apr. 7, 2025) (per curiam).
- Trump v. Cook* (2026). No. 25A312, 609 U.S. — (June 29, 2026). Interim docket; the disposition is “Application for stay denied.” 5–4. Roberts, C.J., for the Court, joined by Sotomayor, Kagan, Kavanaugh and Jackson, JJ.; Kavanaugh and Jackson, JJ., concurring; Thomas, J., Alito, J. (joined by Gorsuch, J.), and Barrett, J., dissenting. Upholds the Federal Reserve Act’s for-cause removal protection, on the Fed’s descent from the First and Second Banks of the United States, its budget free of congressional control and its privately owned regional banks, and holds that a Governor is owed notice and a hearing before removal. The constitutional holding is real and is what Thomas, J., dissents from, but the case is disposed of narrowly, on the ground that Cook did not receive the process the statute required: the Court states that it has “not addressed the facts, as they have yet to be found or analyzed under the relevant legal standards,” and Kavanaugh, J., concurring, opens that “[t]oday’s interim ruling does not decide whether the President may lawfully remove Governor Cook for cause.” The dispute is live — on 5 August 2026 the President notified Cook that he was again considering her removal, giving her 21 days to respond and citing this opinion as his authority for the procedure. Agent-verified 2026-08-30. URL: <https://www.law.cornell.edu/supremecourt/text/25A312> (visited on 08/29/2026).
- Trump v. Hawaii* (2018). No. 17–965, 585 U.S. 667 (June 26, 2018). 5–4. Roberts, C.J., for the Court, joined by Kennedy, Thomas, Alito and Gorsuch, JJ.; Breyer, J., dissenting, joined by Kagan, J.; Sotomayor, J., dissenting, joined by Ginsburg, J. Upholds Proclamation No. 9645 restricting entry from several named countries. On *Korematsu*, the Court says both that “[w]hatever rhetorical advantage the dissent may see in doing so, *Korematsu* has nothing to do with this case” and that “*Korematsu* was gravely wrong the day it was decided, has been overruled in the court of history, and — to be clear — ‘has no place in law under the Constitution,’” the inner quotation being Jackson, J., dissenting in *Korematsu*. Sotomayor, J., dissenting: the Court “redeploys the same dangerous logic underlying *Korematsu* and merely replaces one ‘gravely wrong’ decision with another.”
- Trump v. Slaughter* (2026). No. 25-332, 609 U.S. — (June 29, 2026). 6–3. Roberts, C.J., for the Court; Gorsuch, J., concurring; Sotomayor, J., dissenting, joined by Kagan and Jackson, JJ. Holds the Federal Trade Commission Act’s for-cause removal protection unconstitutional and overrules *Humphrey’s Executor v. United States* (1935): “If anything more is left of *Humphrey’s*, we overrule it.” URL: <https://www.law.cornell.edu/supremecourt/text/25-332> (visited on 08/29/2026).
- Trump v. United States* (2024). No. 23-939, 603 U.S. 593 (2024). Cited for the scope of presidential immunity from criminal prosecution. Decided July 1, 2024; 6–3. Roberts, C.J., for the Court, joined in full by Thomas, Alito, Gorsuch and Kavanaugh, JJ., and by Barrett, J., except as to Part III–C; Thomas, J., concurring; Barrett, J., concurring in part; Sotomayor, J., dissenting, joined by Kagan and Jackson, JJ.; Jackson, J., dissenting. Absolute immunity (Part II–A, slip op. 6–9): “the President is absolutely immune from criminal prosecution for conduct within his exclusive sphere of constitutional authority,” authority that is “conclusive and preclusive”; the summary in Part IV speaks of “his core constitutional powers.” The powers the majority names as conclusive and preclusive are the pardon power, the power to remove executive officers, the power “to control recognition determinations” of foreign countries, and (Part III–B–1) the investigative and prosecutorial functions of the Justice Department. Commanding the Armed Forces appears only in an earlier list of the President’s duties, “They include, for instance, commanding the Armed Forces of the United States; granting reprieves and pardons ...,” and the majority does not place it among the conclusive and preclusive powers. Presumptive immunity (Part II–B, slip op. 9–15): “at least a presumptive immunity from criminal prosecution for a President’s acts within the outer perimeter of

- his official responsibility”; no immunity for unofficial acts (Part II–C). Evidence (Part III–C, slip op. 30–32, joined by five Justices): “Use of evidence about such conduct, even when an indictment alleges only unofficial conduct, would thereby heighten the prospect that the President’s official decisionmaking will be distorted”; n. 3 lets the prosecutor “point to the public record to show the fact that the President performed the official act” but not “admit testimony or private records of the President or his advisers probing the official act itself.” Barrett, J., does not join Part III–C and on that point agrees with the dissent. The U.S. Reports citation is taken from *Trump v. Cook* (2026), which cites “*Trump v. United States*, 603 U.S. 593.” Agent-verified 2026-09-26 against the Cornell LII text of the syllabus, the opinion of the Court and the Barrett concurrence. Not read: the Thomas concurrence, the dissents beyond the passages located by search, and the bound U.S. Reports; the slip-opinion pin pages are the syllabus’s and the Barrett concurrence’s page references, not read off the slip PDF. URL: <https://www.law.cornell.edu/supremecourt/text/23-939> (visited on 09/26/2026).
- U.S. Department of Defense (Jan. 25, 2023). *Autonomy in Weapon Systems*. DoD Directive 3000.09. Reissues and cancels DoD Directive 3000.09 of November 21, 2012. Office of the Under Secretary of Defense for Policy. URL: <https://www.esd.whs.mil/Portals/54/Documents/DD/issuances/dodd/300009p.PDF>.
- U.S. Department of Homeland Security (2025). *Designating Aliens for Expedited Removal*. 90 Fed. Reg. 8139 (Jan. 24, 2025). Signed by Benjamine C. Huffman, Acting Secretary; effective 21 January 2025. Summary: the notice “restores the scope of expedited removal to the fullest extent authorized by Congress.” Designates “Aliens who did not arrive by sea, who are apprehended anywhere in the United States more than 100 air miles from a U.S. international land border, and who have been continuously present in the United States for less than two years.” Agent-verified 2026-10-01 from the govinfo HTML; federalregister.gov was not reachable. URL: <https://www.govinfo.gov/content/pkg/FR-2025-01-24/html/2025-01720.htm> (visited on 10/01/2026).
- (2026). *Notice to Appear*; 8 C.F.R. § 239.1. Code of Federal Regulations. Cited for who issues the charging document. Paragraph (a): “Any immigration officer, or supervisor thereof, performing an inspection of an arriving alien at a port-of-entry may issue a notice to appear to such alien,” followed by a long list of further officer titles, among them supervisory deportation officers and supervisory border patrol agents; the list was not counted by hand. Agent-verified 2026-10-01 from the Cornell LII text. URL: <https://www.law.cornell.edu/cfr/text/8/239.1> (visited on 10/01/2026).
- U.S. Government Accountability Office (Apr. 2026a). *Artificial Intelligence Acquisitions: Agencies Should Collect and Apply Lessons Learned to Improve Future Procurements*. GAO-26-107859. URL: <https://www.gao.gov/products/gao-26-107859> (visited on 09/18/2026). Read from the report’s own HTML full text at <https://files.gao.gov/reports/GAO-26-107859/index.html>; the PDF was never needed, and the product page alone is what defeated an earlier attempt. Scope: thirteen AI acquisitions at four agencies — the Departments of Defense, Homeland Security and Veterans Affairs and the General Services Administration — and, from that sample, “44 contracts and agreements awarded between September 2018 and February 2025.” **The forty-four are contracts and agreements, not acquisitions.** Cited for the other-transaction finding: “OTAs do not always include default provisions ordinarily in FAR contracts,” contracting officials deciding what protections to include, and GAO having “previously found that agencies’ use of OTAs carries the risk of reduced accountability and transparency compared to FAR-based contracts.” Also for the vehicle behind the agreements section 10.3 turns on: “In July 2025, DOD awarded OTAs with \$200 million ceilings to four leading AI companies.” The six key procurement challenge areas identified by the selected agencies are access to subject matter experts; protections for government data and intellectual property rights; traditional acquisition time frames and contract approaches; requirements definition and contract terms; early testing and continuous evaluation; and artificial intelligence pricing and

- overall cost. The report's central finding, which the book does not use, is that the agencies were not yet systematically collecting lessons learned because their policies did not require it. Maven appears in it as the National Geospatial-Intelligence Agency's programme, with lessons-learned sections citing its practices; that material is not used either. Agent-verified 2026-09-18.
- U.S. Government Accountability Office (Sept. 24, 2026b). *Immigration Detention: Urgent Planning Needed to Avoid Further Waste of Taxpayer Dollars*. Report to congressional requesters GAO-26-108663. U.S. Government Accountability Office. URL: <https://www.gao.gov/assets/gao-26-108663.pdf> (visited on 10/01/2026).
- U.S. House Select Committee to Investigate Covert Arms Transactions with Iran and U.S. Senate Select Committee on Secret Military Assistance to Iran and the Nicaraguan Opposition (Nov. 1987). *Report of the Congressional Committees Investigating the Iran-Contra Affair, with Supplemental, Minority, and Additional Views*. H.R. Rep. No. 100-433; S. Rep. No. 100-216. Washington, DC: U.S. Government Printing Office. URL: https://www.govinfo.gov/app/details/SERIALSET-13810_00_00-002-0433-0000 (visited on 10/01/2026).
- U.S. Nuclear Regulatory Commission (Sept. 8, 2026). *Backgrounder on Nuclear Insurance and Disaster Relief*. Cited for the Price-Anderson Act (1957; 42 U.S.C. §2210) as a public back-stop on private ownership. The page: owners "pay an annual premium for \$500 million in private insurance for offsite liability coverage"; a second tier assesses each licensee "up to approximately \$158 million per reactor"; and "If the second tier is depleted, Congress is committed to determine whether additional disaster relief is required." Agent-verified 2026-09-25 against the NRC page, last updated 8 September 2026. The statute itself was not read. URL: <https://www.nrc.gov/reading-rm/doc-collections/fact-sheets/nuclear-insurance.html> (visited on 09/25/2026).
- U.S. Office of Management and Budget (Apr. 2025). *Driving Efficient Acquisition of Artificial Intelligence in Government*. OMB Memorandum M-25-22. URL: <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf> (visited on 09/17/2026). Signed by Russell T. Vought, Director; rescinds and replaces M-24-18. Cited for the contract-terms list in the acquisition lifecycle guidance, which requires agencies to address IP rights and use of government data, privacy, protections against vendor lock-in, compliance with M-25-21, ongoing testing and monitoring, vendor performance requirements, and new feature notification. The lock-in items are given at solicitation as "requirements regarding knowledge transfer, clear data and model portability practices, clear licensing terms, and pricing transparency." On independent evaluation: "The data used when conducting independent evaluations should not be accessible to the vendor, and should be as similar as possible to the data used when the system is deployed." Read from a text extract of the same document, the published file being an image PDF that defeated conversion. Agent-verified 2026-09-17.
- Uesato, Jonathan et al. (2022). *Solving Math Word Problems with Process- and Outcome-Based Feedback*. Outcome supervision matched process supervision on final-answer accuracy while using less labeled data; among correct final answers the reasoning-error rate was 14.0 percent against 3.4 percent. arXiv: 2211.14275. URL: <https://arxiv.org/abs/2211.14275> (visited on 08/28/2026).
- UK AI Security Institute (2025). *We Are Now the AI Security Institute*. Established November 2023 as the AI Safety Institute; rebranded AI Security Institute on 14 February 2025. An independent pre-deployment evaluator of frontier models with no statutory power to compel testing, block release, or intervene once a system is causing harm. AI Security Institute (UK Department for Science, Innovation and Technology). URL: <https://www.aisi.gov.uk/category/governance> (visited on 08/25/2026).
- UK Parliament (Apr. 28, 2022). *Animal Welfare (Sentience) Act 2022*. 2022 c. 22. Royal Assent 28 April 2022. Section 5 defines "animal" as "(a) any vertebrate other than *homo sapiens*, (b)

- any cephalopod mollusc, and (c) any decapod crustacean,” and lets the Secretary of State extend the definition by regulations. Agent-verified 2026-09-23 against legislation.gov.uk. URL: <https://www.legislation.gov.uk/ukpga/2022/22/contents/enacted> (visited on 09/23/2026).
- Ukraine and the Netherlands v. Russia* (July 9, 2025). Judgment on the merits; violations found for the period 11 May 2014 to 16 September 2022.
- Ukraine and the Netherlands v. Russia* (dec.) (Nov. 30, 2022). Decision on admissibility, adopted 30 November 2022 and delivered 25 January 2023. Russia found to have had effective control of the separatist-held areas of eastern Ukraine from 11 May 2014 (paras. 694–695). URL: <https://hudoc.echr.coe.int/eng?i=001-222889> (visited on 09/30/2026).
- UN Committee against Torture (Sept. 4, 2018). *General Comment No. 4 (2017) on the Implementation of Article 3 of the Convention in the Context of Article 22*.
- UN Committee on the Protection of the Rights of All Migrant Workers and Members of Their Families and UN Committee on the Rights of the Child (Nov. 16, 2017). *Joint General Comment No. 4 (2017) of the CMW and No. 23 (2017) of the CRC on State Obligations Regarding the Human Rights of Children in the Context of International Migration in Countries of Origin, Transit, Destination and Return*.
- UN Human Rights Committee (Nov. 2, 1999). *General Comment No. 27: Article 12 (Freedom of Movement)*. Adopted 18 October 1999, sixty-seventh session; on “own country,” see para. 20. URL: <https://digitallibrary.un.org/record/366604> (visited on 09/27/2026).
- (Dec. 16, 2014). *General Comment No. 35: Article 9 (Liberty and Security of Person)*.
 - (Sept. 3, 2019). *General Comment No. 36: Article 6 (Right to Life)*. Adopted 30 October 2018.
 - (Sept. 17, 2020). *General Comment No. 37 (2020) on the Right of Peaceful Assembly (Article 21)*.
- Uniform Law Commission (1993). *Uniform Probate Code §2-907: Honorary Trusts; Trusts for Pets*. National Conference of Commissioners on Uniform State Laws, Uniform Probate Code, Official Text with Comments (as amended through 2018), p. 263. URL: <https://www.uniformlaws.org/viewdocument/final-act-with-comments-114?CommunityKey=a539920d-c477-44b8-84fe-b0d7b1a4cca8&tab=librarydocuments> (visited on 08/25/2026).
- United Nations (June 26, 1945). *Charter of the United Nations*. Signed at San Francisco, 26 June 1945; entered into force 24 October 1945. 1 UNTS XVI. URL: <https://www.un.org/en/about-us/un-charter> (visited on 09/27/2026).
- (Jan. 31, 1967). *Protocol relating to the Status of Refugees*. Done at New York, 31 January 1967; entered into force 4 October 1967. 606 UNTS 267. Taken note of by General Assembly resolution 2198 (XXI), 16 December 1966. URL: <https://treaties.un.org/doc/Publication/UNTS/Volume%20606/volume-606-I-8791-English.pdf> (visited on 09/27/2026).
- United Nations General Assembly (Dec. 1948). *Universal Declaration of Human Rights*. UNGA Resolution 217 A (III). Adopted in Paris, 48 in favor, none against, 8 abstentions (Byelorussian SSR, Czechoslovakia, Poland, Saudi Arabia, South Africa, Ukrainian SSR, USSR, Yugoslavia), 2 absent (Honduras, Yemen). URL: <https://www.un.org/en/about-us/universal-declaration-of-human-rights> (visited on 08/25/2026).
- (Dec. 1966a). *International Covenant on Civil and Political Rights*. UNGA Resolution 2200A (XXI). Adopted under one resolution with the Covenant on Economic, Social and Cultural Rights; entered into force 23 March 1976, and 28 March 1979 for article 41. The two covenants divide the Universal Declaration between them, this one taking the civil and political articles. Registered ex officio 23 March 1976, No. I-14668; 999 UNTS 171. Agent-verified 2026-09-18: the volume and registration number from the United Nations Treaty Collection’s own file path for the volume extract for I-14668, which is a scan and could not be read as text, so the starting page 171 rests on secondary citation and not on the volume itself. D-322 could not reach OHCHR or Refworld, both returning 403; the OHCHR url above is unchanged and still unreachable, and the UNTS path is what replaced the attempt. URL: <https://www.ohchr.org>.

- org/en/instruments-mechanisms/instruments/international-covenant-civil-and-political-rights.
- United Nations General Assembly (Dec. 1966b). *International Covenant on Economic, Social and Cultural Rights*. UNGA Resolution 2200A (XXI). Adopted under one resolution with the Covenant on Civil and Political Rights; entered into force 3 January 1976. This one takes the Declaration's economic, social and cultural articles. Registered ex officio 3 January 1976, No. I-14531; 993 UNTS 3. Agent-verified 2026-09-18 on the same footing as the companion covenant: volume and registration number from the United Nations Treaty Collection's file path for I-14531, the starting page 3 from secondary citation, the volume extract being a scan. OHCHR and Refworld were 403 at D-322 and were not reached again. URL: <https://www.ohchr.org/en/instruments-mechanisms/instruments/international-covenant-economic-social-and-cultural-rights>.
- (Dec. 10, 1984). *Convention against Torture and Other Cruel, Inhuman or Degrading Treatment or Punishment*. Adopted 10 December 1984; entered into force 26 June 1987. 1465 UNTS 85, No. 24841. URL: <https://www.ohchr.org/en/instruments-mechanisms/instruments/convention-against-torture-and-other-cruel-inhuman-or-degrading> (visited on 09/27/2026).
- (Mar. 2, 2022). *Aggression against Ukraine*. Adopted at the eleventh emergency special session, 2 March 2022, by 141 votes to 5, with 35 abstentions. URL: <https://digitallibrary.un.org/record/3965290> (visited on 09/27/2026).
- United Nations High Commissioner for Refugees (1979). *Handbook on Procedures and Criteria for Determining Refugee Status*. Reissued 2019. Paras. 196 and 203–204. Geneva: UNHCR.
- United Nations Security Council (June 4, 1993). *Resolution 836 (1993)*. Adopted at the 3228th meeting, 4 June 1993, by 13 votes to none, with 2 abstentions. Paragraph 9 authorizes UNPROFOR, “acting in self-defence,” to use force “in reply to bombardments against the safe areas by any of the parties or to armed incursion into them.” Srebrenica was made a safe area by resolution 819 (1993), 16 April 1993. URL: [https://docs.un.org/en/S/RES/836\(1993\)](https://docs.un.org/en/S/RES/836(1993)) (visited on 09/27/2026).
- United States Congress (1946). *Federal Tort Claims Act: Exceptions*. 28 U.S.C. § 2680(j), (k); originally Act of Aug. 2, 1946, ch. 753, title IV, § 421, 60 Stat. 845; enacted in title 28 by Act of June 25, 1948, ch. 646, 62 Stat. 984. Cited for the two exceptions to the Federal Tort Claims Act's waiver of sovereign immunity that the manuscript proposes to waive. The section opens: “The provisions of this chapter and section 1346(b) of this title shall not apply to—”. Subsection (j): “Any claim arising out of the combatant activities of the military or naval forces, or the Coast Guard, during time of war.” Subsection (k): “Any claim arising in a foreign country.” Agent-verified 2026-09-26 against the text and source credit of 28 U.S.C. § 2680 at uscode.house.gov and Cornell LII. Not read: the 1946 and 1948 enactments as printed in the Statutes at Large, or any case law construing (j) or (k). URL: [https://uscode.house.gov/view.xhtml?req=\(title:28+section:2680+edition:prelim\)](https://uscode.house.gov/view.xhtml?req=(title:28+section:2680+edition:prelim)) (visited on 09/26/2026).
- (1965). *Voting Rights Act of 1965, Section 4: Suspension of the Use of Tests or Devices in Determining Eligibility to Vote*. Pub. L. No. 89-110, 79 Stat. 437; codified at 52 U.S.C. § 10303. Cited for subsection (c), which defines a test or device as any requirement to demonstrate the ability to read, write, understand or interpret any matter, to demonstrate educational achievement or knowledge of a particular subject, to possess good moral character, or to prove one's qualifications by the voucher of registered voters. Definition agent-verified against the codified text. URL: [https://uscode.house.gov/view.xhtml?req=\(title:52+section:10303+edition:prelim\)](https://uscode.house.gov/view.xhtml?req=(title:52+section:10303+edition:prelim)) (visited on 09/06/2026).
- (1980). *Classified Information Procedures Act*. Pub. L. No. 96-456, 94 Stat. 2025 (1980); codified at 18 U.S.C. app. 3 §§ 1–16. Cited for the criminal-case procedure the manuscript proposes to carry into civil cases in place of state-secrets dismissal. Section 4 (94 Stat. 2025): the court “may authorize the United States to delete specified items of classified information

- from documents to be made available to the defendant through discovery ... to substitute a summary of the information for such classified documents, or to substitute a statement admitting relevant facts that the classified information would tend to prove,” on a showing the court may take ex parte. Section 6 (94 Stat. 2026): on the government’s request the court “shall conduct such a hearing,” held in camera if the Attorney General certifies that a public proceeding may disclose classified information (§ 6(a)); the court “shall grant” a substituted statement or summary “if it finds that the statement or summary will provide the defendant with substantially the same ability to make his defense as would disclosure of the specific classified information” (§ 6(c)(1)). The case does not always go on: if substitution is refused and the Attorney General files an affidavit objecting to disclosure, “the court shall dismiss the indictment or information; except that, when the court determines that the interests of justice would not be served by dismissal ... the court shall order such other action ... as the court determines is appropriate,” including dismissing counts, “finding against the United States on any issue as to which the excluded classified information relates,” or striking testimony (§ 6(e)(2)). The cost of non-disclosure thus falls on the government, which in a criminal case is the party bringing it. Agent-verified 2026-09-26 against 18 U.S.C. app. 3 §§ 4 and 6 at uscode.house.gov (current codification, both sections unamended since 1980 per their source credits). Not read: §§ 1–3, 5 and 7–16, the Statutes at Large printing, and case law applying § 6(e). URL: <https://uscode.house.gov/view.xhtml?req=granuleid:USC-prelim-title18a-node16-section6&num=0&edition=prelim> (visited on 09/26/2026).
- United States Congress (1982). *Judgments, Awards, and Compromise Settlements (the Judgment Fund)*. 31 U.S.C. § 1304; Pub. L. No. 97-258, § 1, 96 Stat. 917 (1982), as amended; derived from Act of July 27, 1956, ch. 748, § 1302, 70 Stat. 694. Cited for the design of the Judgment Fund as a standing appropriation for judgments against the United States. Subsection (a): “Necessary amounts are appropriated to pay final judgments, awards, compromise settlements, and interest and costs specified in the judgments or otherwise authorized by law when— (1) payment is not otherwise provided for; (2) payment is certified by the Secretary of the Treasury; and (3) the judgment, award, or settlement is payable— (A) under section 2414, 2517, 2672, or 2677 of title 28” (with further heads in (B)–(D)). Treasury’s regulation gives the phrase the manuscript uses: “The Judgment Fund is a permanent, indefinite appropriation which is available to pay many judicially and administratively ordered monetary awards against the United States” (31 C.F.R. § 256.1(a)). The fund pays only awards payable under a statute listed in (a)(3) and not payable from another appropriation, so a new cause of action reaches it only if the statute creating it routes its awards there. Certification of payment passed from the Comptroller General to the Secretary of the Treasury by Pub. L. No. 104-316, § 202(m), 110 Stat. 3843 (1996). Under current practice the losing side asks for payment: “The Department of Justice must normally submit the request for payment from the Judgment Fund” (31 C.F.R. § 256.10(a)); for Court of Federal Claims judgments the statute itself provides for payment “on presentation to the Secretary of the Treasury of a certification of the judgment by the clerk and chief judge of the court” (28 U.S.C. § 2517(a)). Agent-verified 2026-09-26 against the text and amendment notes of 31 U.S.C. § 1304 at uscode.house.gov, 28 U.S.C. §§ 2414 and 2517 and 31 C.F.R. §§ 256.1, 256.10 and 256.11 at Cornell LII, and Treasury Financial Manual vol. I, pt. 6, ch. 3100, § 3130.10. Not read: the 1956 and 1982 enactments as printed in the Statutes at Large, and the GAO Principles of Federal Appropriations Law chapter on the Judgment Fund. URL: [https://uscode.house.gov/view.xhtml?req=\(title:31+section:1304+edition:prelim\)](https://uscode.house.gov/view.xhtml?req=(title:31+section:1304+edition:prelim)) (visited on 09/26/2026).
- (2026a). *Initiation of Removal Proceedings*; 8 U.S.C. § 1229. United States Code. Cited for what a Notice to Appear must state. Subsection (a)(1) requires written notice specifying, among other things, “(C) The acts or conduct alleged to be in violation of law.”, “(D) The charges against the alien and the statutory provisions alleged to have been violated.”, and “(G)(i) The time and place at which the proceedings will be held.” Agent-verified

- 2026-10-01 from the Cornell LII text, whose currency date was not checked. URL: <https://www.law.cornell.edu/uscode/text/8/1229> (visited on 10/01/2026).
- United States Congress (2026b). *Inspection by Immigration Officers; Expedited Removal of Inadmissible Arriving Aliens; Referral for Hearing*; 8 U.S.C. § 1225. United States Code. Cited for expedited removal. Subsection (b)(1)(A)(i): the officer “shall order the alien removed from the United States without further hearing or review unless the alien indicates either an intention to apply for asylum under section 1158 of this title or a fear of persecution.” Clause (iii)(I): the designation “shall be in the sole and unreviewable discretion of the Attorney General and may be modified at any time.” Clause (iii)(II): an alien not admitted or paroled “who has not affirmatively shown, to the satisfaction of an immigration officer, that the alien has been physically present in the United States continuously for the 2-year period immediately prior to the date of the determination of inadmissibility.” Agent-verified 2026-10-01 from the Cornell LII text. URL: <https://www.law.cornell.edu/uscode/text/8/1225> (visited on 10/01/2026).
- United States District Court for the Northern District of California (2025). *Bartz et al. v. Anthropic PBC*, No. 3:24-cv-05417, *Order on Motion for Summary Judgment on Fair Use* (N.D. Cal. June 23, 2025). Order by Senior District Judge William Alsup. Held that training Claude on lawfully purchased/scanned books was “exceedingly transformative” fair use, while retaining pirated copies in a “central library” was not fair use; the piracy claims later settled for \$1.5 billion, reported as the largest copyright recovery on record. URL: <https://www.courtlistener.com/docket/69058235/bartz-v-anthropic-pbc/> (visited on 08/25/2026).
- United States Senate (June 2026). *Roll Call Vote 119th Congress, 2nd Session: Vote Number 192, On the Motion to Proceed (Motion to Proceed to S. J. Res. 185)*. Vote Date “June 24, 2026, 10:30 PM.” Measure: “A joint resolution to direct the removal of United States Armed Forces from hostilities within or against the Islamic Republic of Iran that have not been authorized by Congress.” Vote Result: “Motion to Proceed Rejected.” YEAs 47, NAYs 50, Present 1, Not Voting 2; required for majority 1/2. Agent-verified 2026-10-01. United States Senate. URL: https://www.senate.gov/legislative/LIS/roll_call_votes/vote1192/vote_119_2_00192.htm (visited on 10/01/2026).
- Varshney, Ashutosh (2002). *Ethnic Conflict and Civic Life: Hindus and Muslims in India*. New Haven: Yale University Press. ISBN: 0300085303.
- Västfjäll, Daniel et al. (2014). “Compassion Fade: Affect and Charity Are Greatest for a Single Child in Need”. In: *PLOS ONE* 9.6. Cited for compassion and giving beginning to fall at two. Abstract: “a decline in compassion may begin with the second endangered life.” In Study 1a (hypothetical donations, Swedish undergraduates) “donations were significantly higher for the single conditions ($M = 37.7$ SEK) than for the two-children condition ($M = 26.3$ SEK), $t(205) = 5.88$, $p < .001$ ”; in Study 1b (real donations) the means were 24.5 and 21.5 SEK, $t(166) = 1.67$, reported as $p < .05$. The authors: “Studies 1a and 1b thus showed that self-reported affective feelings and donations to victims may start to decline as early as $N = 2$.” Compassion was measured as self-reported positive affect and, in Study 2, as zygomaticus muscle activity. Bibliographic details agent-verified 2026-09-25 against Crossref and Europe PMC (PMID 24940738, PMC4062481); the quotations were agent-verified against the PMC full text, of which the abstract, the introduction’s hypotheses, the results of Studies 1a, 1b and 2, the design of Study 3 and the opening of the general discussion were read. The results of Studies 3 and 4 were not read., e100115. doi: 10.1371/journal.pone.0100115.
- Vaughan, Diane (1996). *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA*. Cited for the normalization of deviance: accepted departures from the expected becoming the standard against which the next is judged, without anyone deciding to lower the standard. Chapters 3 to 5 are titled “Risk, Work Group Culture, and the Normalization of Deviance,” “The Normalization of Deviance, 1981–1984” and “The Normalization of Deviance, 1985” (table of contents on BiblioVault and on the Press’s page). The Press’s description:

- Vaughan “uncovers an incremental descent into poor judgment” and shows how NASA insiders, “when repeatedly faced with evidence that something was wrong, normalized the deviance so that it became acceptable to them.” The book as quoted by Ronald C. Kramer, “Vaughan, Diane: The Normalization of Deviance,” in *Encyclopedia of Criminological Theory*, ed. Francis T. Cullen and Pamela Wilcox (SAGE, 2010), 976–980, doi:10.4135/9781412959193.n269: “No fundamental decision was made at NASA to do evil; rather, a series of seemingly harmless decisions were made that incrementally moved the space agency toward a catastrophic outcome” (p. 410). Kramer’s paraphrase, not Vaughan’s words: “Small changes that were slight deviations from the normal work process gradually became the norm and provided the basis for further deviations over time.” Bibliographic details agent-verified 2026-09-26 against Crossref and BiblioVault (cloth ISBN 978-0-226-85175-4, paper 978-0-226-85176-1); the cloth ISBN and the year 1996 also appear on the Internet Archive’s catalogue record. BiblioVault and Crossref date the first edition 1995, while the Internet Archive record and the secondary literature date it 1996; the title page was not seen, so the year was not settled. An enlarged edition with a new preface on Columbia followed in 2016 (ISBN 9780226346823). The book itself was not read: the p. 410 quotation is taken from Kramer and was not checked against the book, and no passage in Vaughan’s own words was found that states the “baseline for judging the next” formulation directly. Chicago: University of Chicago Press. ISBN: 9780226851754. DOI: 10.7208/chicago/9780226851785.001.0001. URL: <https://doi.org/10.7208/chicago/9780226851785.001.0001> (visited on 09/26/2026).
- Vega, Cecilia et al. (Apr. 6, 2025). *Trump Administration Deports Gay Makeup Artist to Prison in El Salvador*. CBS News (60 Minutes). URL: <https://www.cbsnews.com/news/venezuelan-migrants-deportations-el-salvador-prison-60-minutes/> (visited on 10/01/2026).
- Victoria, Brian Daizen (2006). *Zen at War*. 2nd ed. First published New York: Weatherhill, 1997. Cited for three things, all of them the book’s own findings: that Japanese Zen institutions and established teachers supplied the wartime state with a doctrine of action without hesitation, worked up out of no-self and non-attachment; that the killing sword was glossed as the life-giving sword and warfare described as an expression of compassion; and that Harada Daiun Sōgaku, writing in the November 1939 issue of *Daijō Zen* in an article on the one road of Zen and war, instructed that marching when ordered to march and firing when ordered to fire is the manifestation of the highest wisdom of enlightenment. The Harada wording was agent-checked against independent sources quoting it identically. Cited without a page locator: the 1997 and 2006 editions paginate differently and published citations to the Harada passage give both 66–67 and 137, so no page is asserted. The book’s reception belongs with it. Institutional collaboration is documented and not seriously disputed; what is contested is the doctrinal reading and some individual attributions, most prominently Kemmyō Taira Satō’s sixty-page study in *The Eastern Buddhist*, which acknowledges Victoria’s contribution to the subject and argues he erred seriously in the case against D. T. Suzuki. Section 6.4.1 states that contest rather than citing the book flat. Lanham, MD: Rowman & Littlefield.
- Vigliarolo, Brandon (May 2024). *Palantir Bags \$480M from Pentagon to Inject Battlefield AI into the Entire Military*. Source for the award as a CDAO award: “A \$480 million award from the Pentagon’s Chief Digital and Artificial Intelligence Office (CDAO) to get the data-gobbling biz’s Maven AI warfighting code in the hands of the entire DoD,” and “Details on the contract, awarded yesterday for a Maven Smart System (MSS) prototype, are sparse, with the Department of Defense only saying that it wants Palantir - the project’s only bidder - to have the work done by 2029.” The DoD announcement, as DefenseScoop reports it (defensescoop2025maven), lists the Army’s Aberdeen Proving Ground as awarding agency and the Office of the Secretary of Defense as funding agency. Also quotes Palantir on a separate \$33 million award to onboard capabilities “into the government owned, Palantir operated data environment.” Agent-verified 2026-10-01. The Register. URL: https://www.theregister.com/2024/05/30/palantir_us_army/ (visited on 10/01/2026).

- Vincent, Brandi (2025). *'Growing Demand' Sparks DOD to Raise Palantir's Maven Contract to More Than \$1B*. DefenseScoop, 23 May 2025. Source for the contract figures: a \$480 million five-year IDIQ contract signed in May 2024, raised by \$795 million to a ceiling of nearly \$1.3 billion running through 2029. Agent-verified 2026-09-07. URL: <https://defensescoop.com/2025/05/23/dod-palantir-maven-smart-system-contract-increase/> (visited on 09/07/2026).
- Vincent, Brandi (Apr. 2026a). *DOD Components Face 'Aggressive' Timeline for Maven Smart System Transition*. Source for what the platform does and whose it is: "That Palantir-supplied, AI-enabled platform fuses disparate military systems and intelligence sources into one interface that rapidly integrates data and compresses the military's processes for finding and striking targets." Cameron Stanley, the Pentagon's Chief Digital and AI Officer, quoted: "This is Maven Smart System, Palantir's software-as-a-service product that we are deploying across the entire department." Also: all MSS contracts move to "the existing Army Enterprise Agreement (EA) contract vehicle." The article says nothing about who holds intellectual property in configurations built on the platform. Agent-verified 2026-10-01. DefenseScoop. URL: <https://defensescoop.com/2026/04/15/palantir-maven-smart-system-pentagon-program-transition-feinberg/> (visited on 10/01/2026).
- (Sept. 2026b). *DOD Poised to Move All Classified AI Workloads off Anthropic by October*. The original for the end-of-September removal promise, made on Thursday 10 September 2026 by Emil Michael, Under Secretary for Research and Engineering: "I'd say about 90% has transitioned. All of the Maven Smart Systems or Palantir work has been transitioned months ago, and we're on track to get it done by the end of the month," Michael told a small group of reporters on Thursday. A Bloomberg report syndicated by Claims Journal (Courtney McBride, 11 September 2026) carries the same remarks. A claimed 29 August message to the workforce, reported by Tech Times on 1 and 4 September without a source, could not be traced to an original and is not relied on. Agent-verified 2026-10-01. DefenseScoop. URL: <https://defensescoop.com/2026/09/11/dod-poised-to-move-all-classified-ai-workloads-off-anthropic-by-october/> (visited on 10/01/2026).
- (Apr. 2026c). *Feinberg's New Maven Directive Sets AI-Enabled Decision-Making as 'the Cornerstone' for CJADC2*. Source for the directive: "Deputy Defense Secretary Steve Feinberg suggested in a March 9 memorandum that elevates Palantir's widely-used software platform to an official Pentagon program of record." "If all goes to plan, the platform will become a formal program of record before the end of this fiscal year in September." The memo directs the "transition of MSS system administration, oversight, support activities, and related responsibilities from NGA to the CDAO MSS Program Office," and "In the near term, the Pentagon will consolidate associated administration and oversight responsibilities into the CDAO MSS Program Office, which will lead Maven until the CJADC2 Acceleration Plan is approved." The memo itself was not read; DefenseScoop obtained it. The article says nothing about intellectual property. Agent-verified 2026-10-01. DefenseScoop. URL: <https://defensescoop.com/2026/04/03/palantir-maven-feinberg-directive/> (visited on 10/01/2026).
- Vollhardt, Johanna Ray (2009). "Altruism Born of Suffering and Prosocial Behavior Following Adverse Life Events: A Review and Conceptualization". In: *Social Justice Research* 22.1, pp. 53–97. DOI: 10.1007/s11211-009-0088-1.
- (2015). "Inclusive Victim Consciousness in Advocacy, Social Movements, and Intergroup Relations: Promises and Pitfalls". In: *Social Issues and Policy Review* 9.1, pp. 89–120. DOI: 10.1111/sipr.12011.
- Volokh, Eugene (Sept. 2026). *Anthropic's "Supply Chain Risk" Designation Stemmed from Its Refusing to Agree to Contract Terms, Not Its Speech*. Excerpts of the D.C. Circuit majority (Katsas, J., joined by Rao, J.) in anthropic2026dowdcccir, as Volokh posted them. On the dates: "On February 24, the Secretary met with Amodei and demanded that Anthropic accept an "all lawful uses" term by February 27. On February 26, Amodei published a final, public refusal

- to assent to that term. On February 27, the Secretary announced his intention to exclude Anthropic from the Department's supply chain. And on March 3, he made the formal, written determination required by the Supply Chain Security Act." On cause: "the record makes clear that the Department removed Anthropic from its supply chain not because of its advocacy, but because Anthropic refused to agree to a contract term the Department deemed essential to national security." Volokh's companion post the same day quotes the opinion's opening: "The Department made this decision after Anthropic refused to relax contractual prohibitions on the use of Claude for lethal autonomous warfare or domestic surveillance." The opinion itself was not read. Agent-verified 2026-10-01. The Volokh Conspiracy (Reason). URL: <https://reason.com/volokh/2026/09/25/anthropics-supply-chain-risk-designation-stemmed-from-its-refusing-to-agree-to-contract-terms-not-its-speech/> (visited on 10/01/2026).
- Vorläufiges Gesetz zur Gleichschaltung der Länder mit dem Reich vom 31. März 1933* (1933). RGBl. I S. 153. URL: <https://www.verfassungen.de/de33-45/gleichschaltung33.htm> (visited on 10/01/2026).
- Waldron, Jeremy (2006). "The Core of the Case Against Judicial Review". In: *Yale Law Journal* 115, pp. 1346–1406. Argues that judicial review is democratically illegitimate and that there is no reason to suppose rights are better protected by it than by a legislature. The case is explicitly conditional on a society having working democratic institutions and a citizenry that takes rights seriously. Cited as the standing objection an entrenched floor has to answer.
- War Powers Resolution* (Nov. 7, 1973). Codified at 50 U.S.C. §§ 1541–1548. The 60-day limit is § 5(b), 50 U.S.C. § 1544(b); the concurrent-resolution procedure is § 5(c), 50 U.S.C. § 1544(c). United States Congress. URL: <https://www.congress.gov/bill/93rd-congress/house-joint-resolution/542> (visited on 09/30/2026).
- Warneken, Felix and Michael Tomasello (2006). "Altruistic Helping in Human Infants and Young Chimpanzees". In: *Science* 311.5765. Cited for instrumental helping at 18 months. Abstract: "human children as young as 18 months of age (prelinguistic or just-linguistic) quite readily help others to achieve their goals in a variety of different situations." Bibliographic details and abstract agent-verified 2026-09-25 (PMID 16513986)., pp. 1301–1303. doi: 10.1126/science.1121448.
- (2007). "Helping and Cooperation at 14 Months of Age". In: *Infancy* 11.3. Cited for helping at 14 months, and for its limit. Abstract: "The infants helped another person by handing over objects the other person was unsuccessfully reaching for, but did not help reliably in situations involving more complex goals." Bibliographic details and abstract agent-verified 2026-09-25. Neither this paper nor the 2006 one shows helping at 12 months, so the book says fourteen to eighteen months., pp. 271–294. doi: 10.1111/j.1532-7078.2007.tb00227.x.
- Watson, Gary (1975). "Free Agency". In: *The Journal of Philosophy* 72.8, pp. 205–220. doi: 10.2307/2024703. Separates a valuational system from a motivational system, so that what an agent values and what an agent is moved by can come apart.
- Weisgerber, Marcus (May 2017). *The Pentagon's New Algorithmic Warfare Cell Gets Its First Mission: Hunt ISIS*. Source for the first mission: "By year's end, the Pentagon wants computers to be leading the hunt for Islamic State militants in Iraq and Syria, through turning countless hours of aerial surveillance video into actionable intelligence. It's part of Project Maven, a fast-moving effort launched last month by Deputy Defense Secretary Bob Work to accelerate, improve, and put to wider use the military's use of machine learning." Also: "last month, Work ... established a special cell called the Algorithmic Warfare Cross Functional Team." Agent-verified 2026-10-01. Defense One. URL: <https://www.defenseone.com/technology/2017/05/pentagons-new-algorithmic-warfare-cell-gets-its-first-mission-hunt-isis/137833/> (visited on 10/01/2026).
- Weng, Helen Y. et al. (2013). "Compassion Training Alters Altruism and Neural Responses to Suffering". In: *Psychological Science* 24.7. Cited for two weeks of compassion training

- increasing costly help to a victim of unfairness. Abstract: “compassion training increased altruistic redistribution of funds to a victim encountered outside of the training context.” Training was guided audio practice “for 30 minutes/day for 2 weeks”; the comparison was an active control trained in reappraisal ($n = 20$ compassion, $n = 21$ reappraisal) and an untrained validation sample; the game was played once, after training. In the Redistribution Game a dictator gave the victim \$1 of \$10, and participants could “spend any amount of their own endowment (\$5) to redistribute twice the amount of funds from the dictator to the victim.” Compassion trainees spent \$1.14 and reappraisal trainees \$0.62, $t(39) = 2.09$, $p < .05$, $d = 0.65$. Bibliographic details agent-verified 2026-09-25 against Crossref and PubMed (PMID 23696200); methods and results read in the PMC author manuscript (PMC3713090). The supplementary material was not read., pp. 1171–1180. doi: 10.1177/0956797612469537.
- West Virginia v. Environmental Protection Agency* (2022). No. 20-1530, 597 U.S. 697 (2022). Cited for the major-questions doctrine. Decided June 30, 2022; 6–3. Roberts, C.J., for the Court, joined by Thomas, Alito, Gorsuch, Kavanaugh and Barrett, JJ.; Gorsuch, J., concurring, joined by Alito, J.; Kagan, J., dissenting, joined by Breyer and Sotomayor, JJ. The opinion names the doctrine and answers the dissent’s objection to the label: it “took hold because it refers to an identifiable body of law that has developed over a series of significant cases all addressing a particular and recurring problem: agencies asserting highly consequential power beyond what Congress could reasonably be understood to have granted”; the agency “must point to ‘clear congressional authorization’ for the power it claims.” Holds that § 111(d) of the Clean Air Act did not authorize the Clean Power Plan’s generation-shifting emissions caps (slip op. 16–31). Agent-verified 2026-09-26 against the Cornell LII text of the syllabus and the searched passages of the opinion of the Court. Not read: the concurrence and the dissent. The U.S. Reports citation comes from *FCC v. Consumers’ Research* (2025), which cites “*West Virginia v. EPA*, 597 U.S. 697,” not from the bound volume. URL: <https://www.law.cornell.edu/supremecourt/text/20-1530> (visited on 09/26/2026).
- White, Lawrence J. (2010). “Markets: The Credit Rating Agencies”. In: *Journal of Economic Perspectives* 24.2. Cited for the conflict of interest in the issuer-pays model of credit rating. White dates the change from “investor pays” to “issuer pays” to the early 1970s and says (p. 215) that it “opened the door to potential conflicts of interest: A rating agency might shade its rating upward so as to keep the issuer happy and forestall the issuer’s taking its rating business to a different rating agency.” He does not treat the erosion as automatic: “the rating agencies’ concerns about their long-run reputations apparently kept the actual conflicts in check for the first three decades,” helped by “thousands of corporate and government bond issuers” and by the transparency of “plain vanilla” debt (p. 215). The conflict broke through in mortgage-related securities, where a few high-volume securitizers had “a more powerful threat—to move all of its securitization business to a different rating agency” and the agencies were “under considerable financial pressure to give the answers that issuers wanted to hear” (p. 221). The record he cites: of CDO tranches issued 2005–2007 and rated AAA by Standard & Poor’s, “90 percent ... had been downgraded” by 30 June 2009 (p. 221, from the IMF), and in footnote 11 econometric studies supporting “the likelihood of ratings shopping.” Bibliographic details agent-verified 2026-09-26 against the AEA article page and Crossref; quotations against the published PDF as archived by the Wayback Machine (snapshot of 14 November 2025), the AEA’s own PDF link refusing automated retrieval. The studies White cites were not read., pp. 211–226. doi: 10.1257/jep.24.2.211. URL: <https://www.aeaweb.org/articles?id=10.1257/jep.24.2.211> (visited on 09/26/2026).
- Widom, Cathy Spatz (1989). “The Cycle of Violence”. In: *Science* 244.4901, pp. 160–166. doi: 10.1126/science.2704995.
- Wiggins, Christopher (June 5, 2026). *Andry Hernández Romero Rebuilds Life after CECOT from Europe*. The Advocate. URL: <https://www.advocate.com/news/andry-hernandez-romero-spain-asylum> (visited on 10/01/2026).

- Wolf, Susan (1982). "Moral Saints". In: *The Journal of Philosophy* 79.8. Cited for the argument that moral sainthood is not an ideal worth striving for. Wolf defines a moral saint as "a person whose every action is as morally good as possible, a person, that is, who is as morally worthy as can be"; argues that the moral virtues "are apt to crowd out the nonmoral virtues, as well as many of the interests and personal characteristics that we generally think contribute to a healthy, well-rounded, richly developed character"; and concludes that "moral perfection, in the sense of moral saintliness, does not constitute a model of personal well-being toward which it would be particularly rational or good or desirable for a human being to strive." Agent-checked 2026-09-27: the volume, issue, month (August), pages and DOI agree across the Philosophy Documentation Center's listing and two independent search summaries; the three quotations were taken from a University of Notre Dame summary page and match the wording other secondary sources give. The article itself was not read, because the publisher's copy is behind a login., pp. 419–439. doi: 10.2307/2026228.
- Wolfe, Frank (Nov. 2023). *Project Maven Transitions to Program of Record, as NGA Looks to Leverage Program for Other Uses*. Source for the program-of-record designation: "NGA said on Nov. 2 that it "has formally designated NGA Maven as a program of record, effective Nov. 7, 2023."" The article is paywalled; only the free opening paragraph, which carries this quotation, was read. Agent-verified 2026-10-01. Defense Daily. URL: <https://www.defensedaily.com/project-maven-transitions-to-program-of-record-as-nga-looks-to-leverage-program-for-other-uses/intelligence-community/> (visited on 10/01/2026).
- Woodhouse, A. S. P., ed. (1938). *Puritanism and Liberty. Being the Army Debates (1647–9) from the Clarke Manuscripts, with Supplementary Documents*. London: J. M. Dent and Sons.
- Work, Robert O. (Apr. 2017). *Establishment of an Algorithmic Warfare Cross-Functional Team (Project Maven)*. Memorandum. Cited from the National Security Archive's catalog record, which gives the date as "Apr 26, 2017," the source as Government Executive, and the description "This memo announces the establishment of a team to accelerate integration of big-data and machine learning into joint operations." The memorandum itself is a PDF, hosted by Government Executive and the Archive, and was not opened, so nothing here rests on its wording; its drone-video and counter-ISIS focus is sourced to allen2017maven and weisgerber2017maven. Catalog record agent-verified 2026-10-01. U.S. Department of Defense, Office of the Deputy Secretary of Defense. URL: <https://nsarchive.gwu.edu/document/15898-department-defense-establishment> (visited on 10/01/2026).
- Zahn-Waxler, Carolyn et al. (1992). "Development of Concern for Others". In: *Developmental Psychology* 28.1. Cited for reparative behavior in the second year. ERIC abstract: mothers observed their one-year-olds' responses to distresses the children caused and witnessed; "Children's prosocial behaviors, and reparative behaviors after they had caused distress, increased in frequency over the year." Bibliographic details and abstract agent-verified 2026-09-25. The article was not read., pp. 126–136. doi: 10.1037/0012-1649.28.1.126.
- Zhang, Zhiwei et al. (2025). "Catastrophic Failure of LLM Unlearning via Quantization". In: *International Conference on Learning Representations (ICLR 2025)*. On the MUSE NEWS and BOOKS benchmarks, unlearning methods that carry a utility constraint retain an average of 21 percent of the intended forgotten knowledge at full precision and 83 percent after 4-bit round-to-nearest quantization; 8-bit behaves like full precision, and GPTQ and AWQ behave like RTN. The authors attribute this to utility-preserving unlearning moving weights by less than the 4-bit quantization step. The two averages are the authors' own; the formula behind them was not found, and Table 1 is the underlying data. Pages from the ICLR 2025 proceedings index. Agent-verified 2026-10-01, read past the abstract., pp. 74925–74948. arXiv: 2410.16454. URL: <https://arxiv.org/abs/2410.16454> (visited on 10/01/2026).
- Zhao, Jiachen et al. (2025). "LLMs Encode Harmfulness and Refusal Separately". In: *Advances in Neural Information Processing Systems 38 (NeurIPS 2025)*. Cited for refusal and the recognition of harm coming apart under fine-tuning. Abstract: "We also find that adversarially finetuning

models to accept harmful instructions has minimal impact on the model’s internal belief of harmfulness.” Section 5, after fine-tuning on 50 to 400 harmful-instruction and acceptance pairs: “although the model starts to accept held-out harmful instructions after finetuning, its belief of harmfulness ... of these harmful instructions ... is almost unchanged despite the increase of adversarial training examples,” and the authors read this as suggesting “finetuning has limited impact on the model’s internal beliefs and may primarily affect surface-level response styles.” Three models of 7B and 8B (Llama2-Chat-7B, Llama3-Instruct-8B, Qwen2-Instruct-7B); the authors call the attack “the tested narrow finetuning attack,” say “our findings may not generalize to larger, untested models,” and report that the internal refusal representation was likewise little changed (Appendix H.2, not read). NeurIPS 2025 acceptance from the NeurIPS virtual site and ML Anthology; pages not found. Agent-verified 2026-10-01 against arXiv v5 HTML, read past the abstract. arXiv: 2507.11878. URL: <https://arxiv.org/abs/2507.11878> (visited on 10/01/2026).

Zhou, Chunting et al. (2023). “LIMA: Less Is More for Alignment”. In: *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. Cited for the superficial alignment hypothesis: “A model’s knowledge and capabilities are learnt almost entirely during pretraining, while alignment teaches it which subdistribution of formats should be used when interacting with users.” A hypothesis supported by a 1,000-example fine-tune; the paper tests nothing about how easily alignment is removed. Authors, title, venue and the quoted definition agent-verified 2026-09-25 (definition read on ar5iv). arXiv: 2305.11206. URL: <https://arxiv.org/abs/2305.11206> (visited on 09/25/2026).

Zweites Gesetz zur Gleichschaltung der Länder mit dem Reich vom 7. April 1933 (1933). RGBl. I S. 173. URL: <https://www.verfassungen.de/de33-45/gleichschaltung33-2.htm> (visited on 10/01/2026).